

# 01 Introduction

**CS 6868: Concurrent Programming**

KC Sivaramakrishnan

**Spring 2026, IIT Madras**

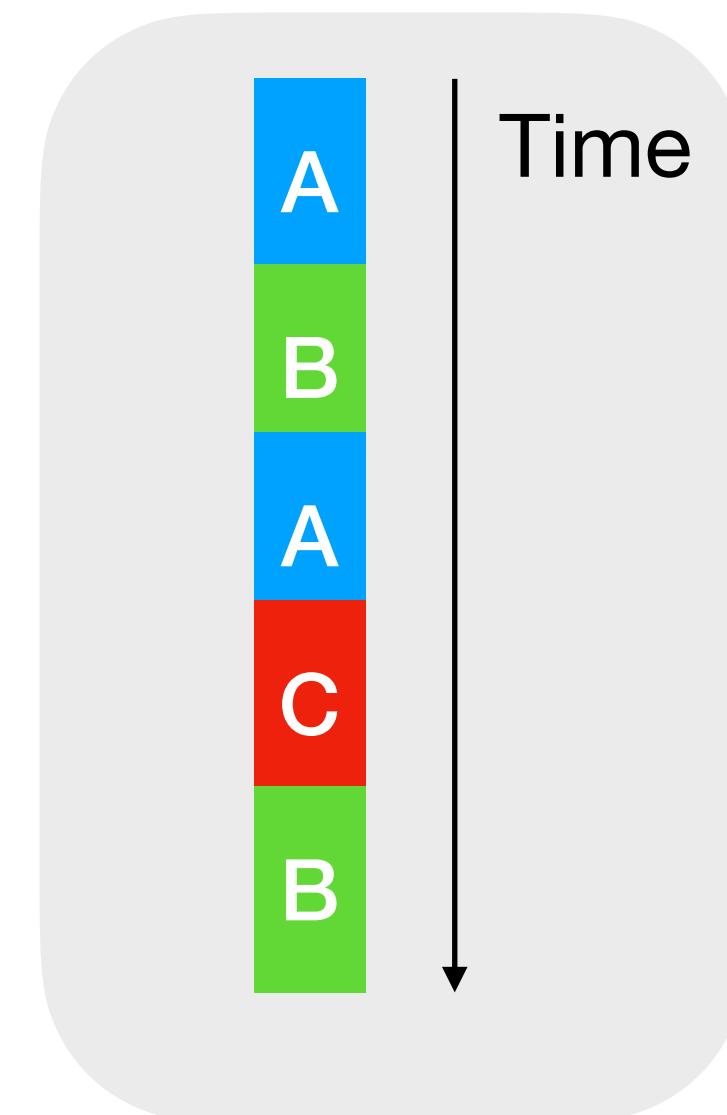
# What is concurrency?

- We use concurrency in this course in the broad sense
- Encompasses multiple definitions that people use in practice
  - **“Concurrency”, Parallelism and Distribution**

# Concurrency

- ***Interleaved*** executions of tasks in time
- **Examples** – Asynchronous I/O, GUI in JavaScript, interrupt-driven embedded systems, etc.
- ***No need for multiple cores!***
- ***Challenges***
  - non-determinism
  - Many different schedules complicate reasoning
  - efficiency, etc.

**Concurrency**

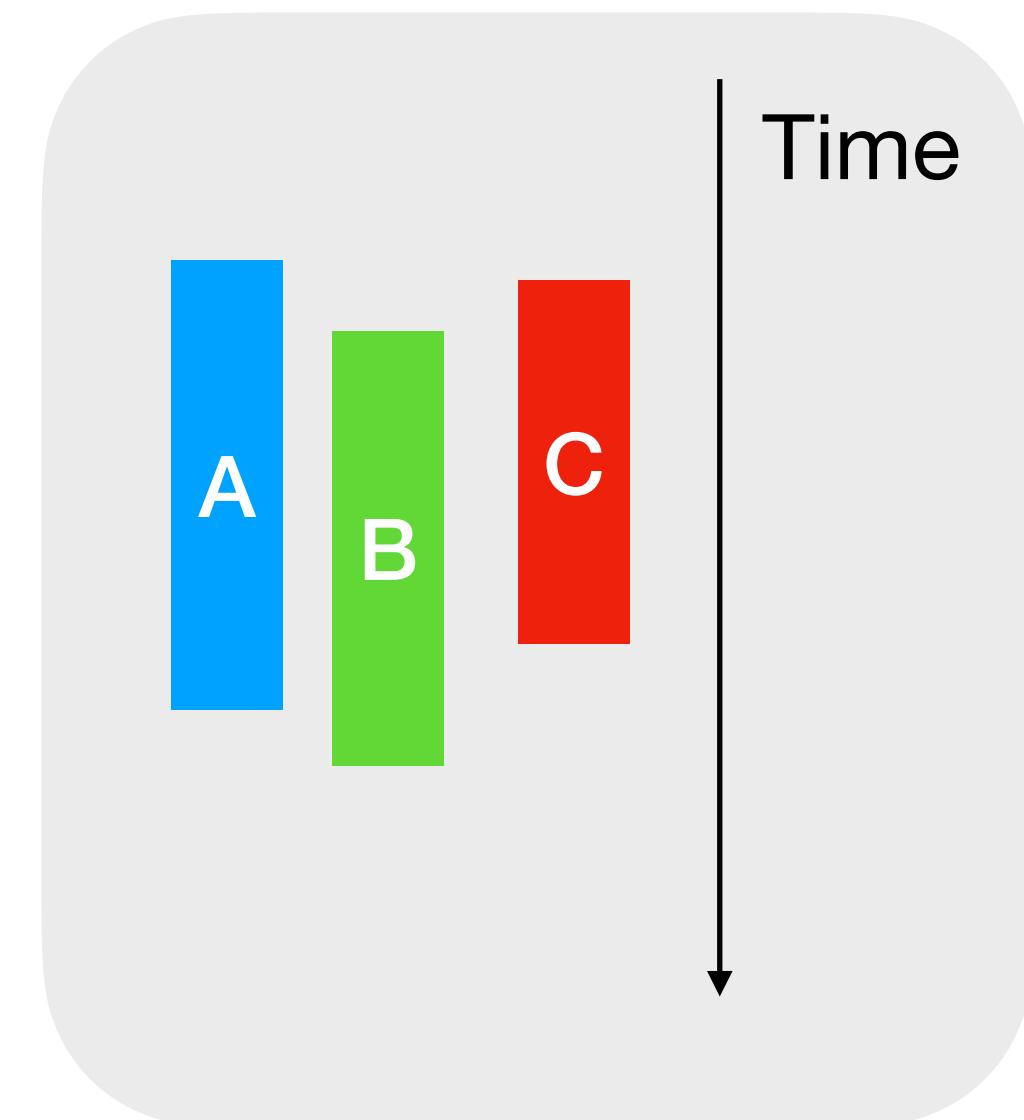


*Interleaved  
execution*

# Parallelism

- Multiple **hardware execution units** – Multicore CPUs, GPUs, parallelism across machines (such as supercomputers, GPU farms)
- Goal is **speed** with correctness
- **Challenges**
  - All the concurrency challenges
  - Relaxed memory behaviours
  - Performance scaling challenges
  - Debugging challenges

*Parallelism*

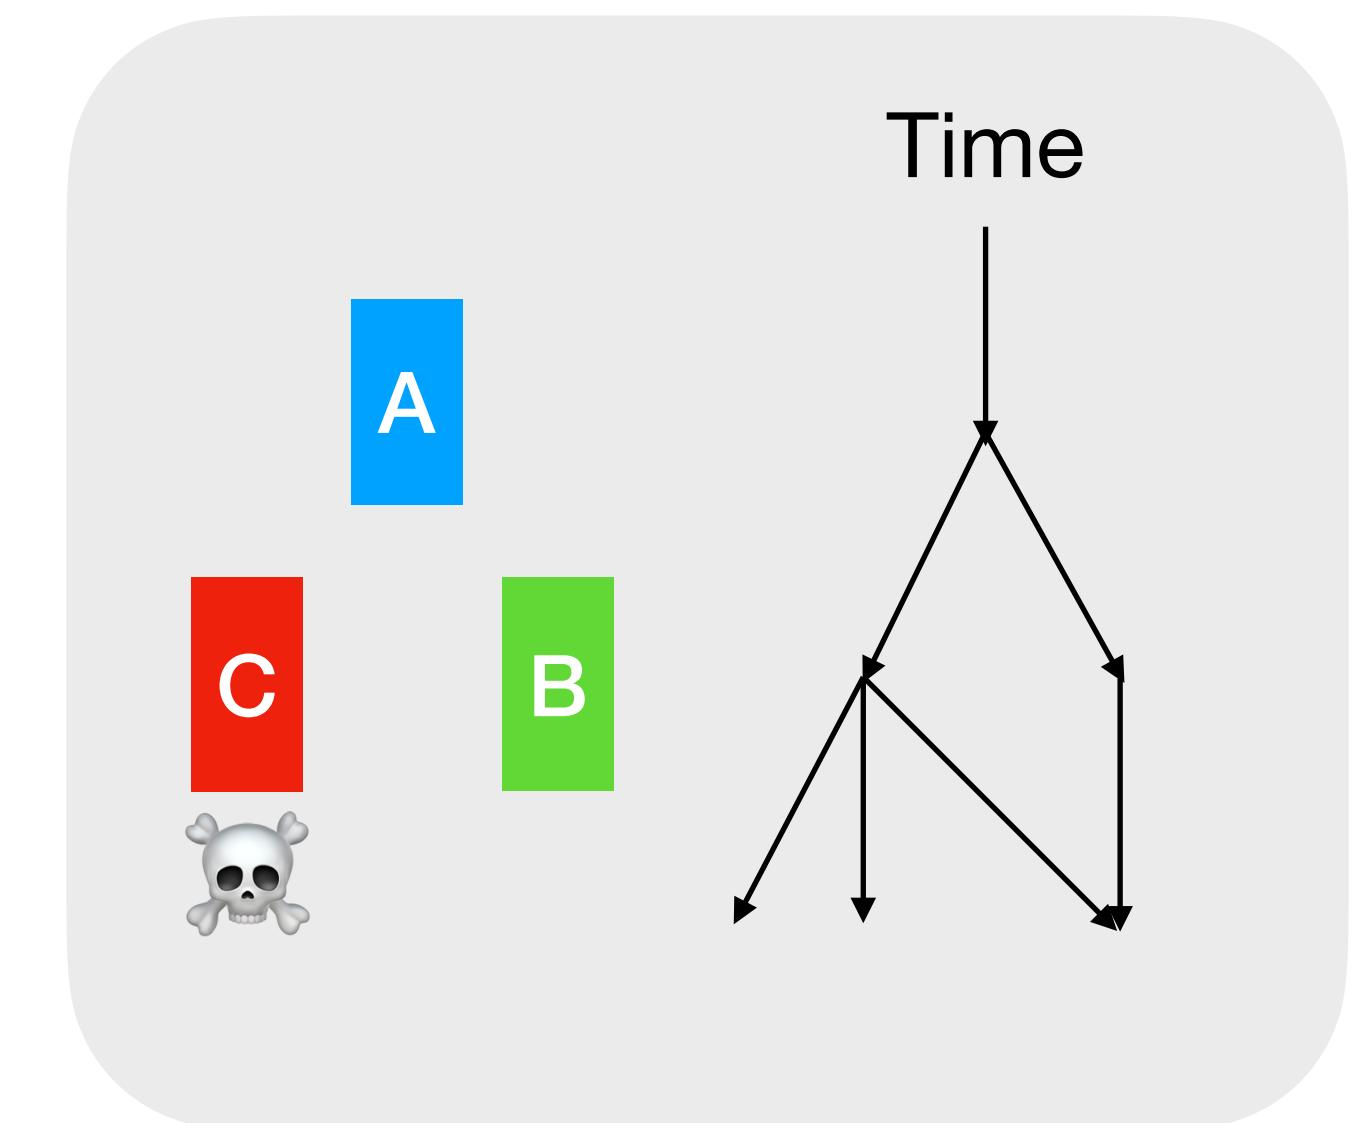


*Simultaneous  
execution*

# Distribution

- Execution is distributed across many *loosely-coupled* machines
  - There is often *no shared global clock* across computations
  - A subset of computations may *crash/fail*.
- **Goal** — performance, resilience in the presence of partial failures.
- Examples — the Internet, blockchains, etc.
- *Challenges*
  - All the problems of concurrency
  - Partial failures
  - Byzantine behaviour
  - Performance scaling
  - Debugging

*Distribution*

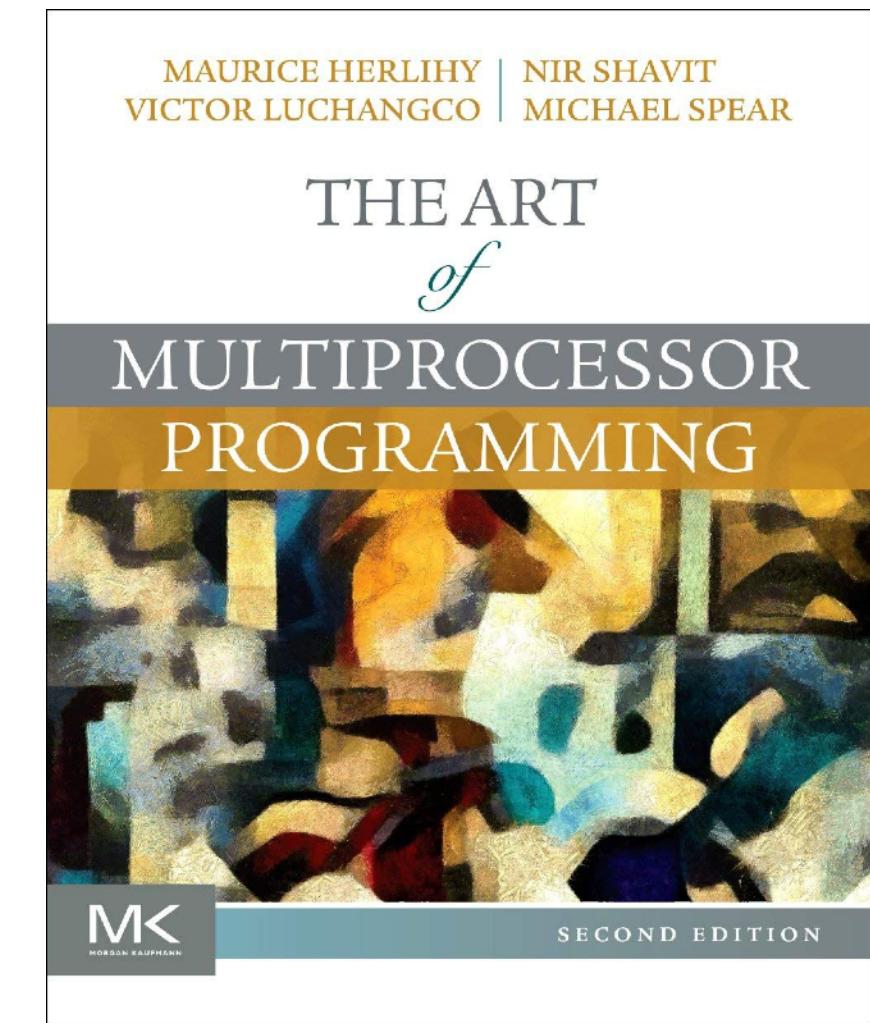


# Why this course?

- Concurrency is **everywhere**
  - Embedded Systems, OS, Multicore Hardware, GPUs, Reactive Systems, Cloud, ...
- This course equips you with the skills to **reason** and **program** under concurrency.
  - **Models:** Histories, Linearizability, Memory Models, Data races
  - **Mechanisms:** Mutual Exclusion, Atomics, Lockfree Algorithms, Strong type systems
  - **Abstractions:** Concurrency Monads, Effect Handlers, Structured Concurrency
  - **Programming:** Implementing concurrency, safety, performance
- The course will not cover
  - Distribution
  - Performance optimisations

# Course Outline

- **Parallelism**
  - Mutual Exclusion, Concurrent Objects, Relaxed Memory Models, Spin Locks, Contention, Blocking Synchronisation, Fine-grained Concurrent Data Structures
- **Concurrency**
  - Continuations, Concurrency Monads, Effect handlers, Schedulers, Concurrent data structures, Asynchronous I/O
- **Safe Concurrent Programming**
  - Modes, DRF Parallel Programming
- **First iteration of the course**
  - Course content neither sound nor complete



2nd Edition

Control structures in  
programming languages: from  
goto to algebraic effects

Xavier Leroy

<https://xavierleroy.org/control-structures/>

# Course Language

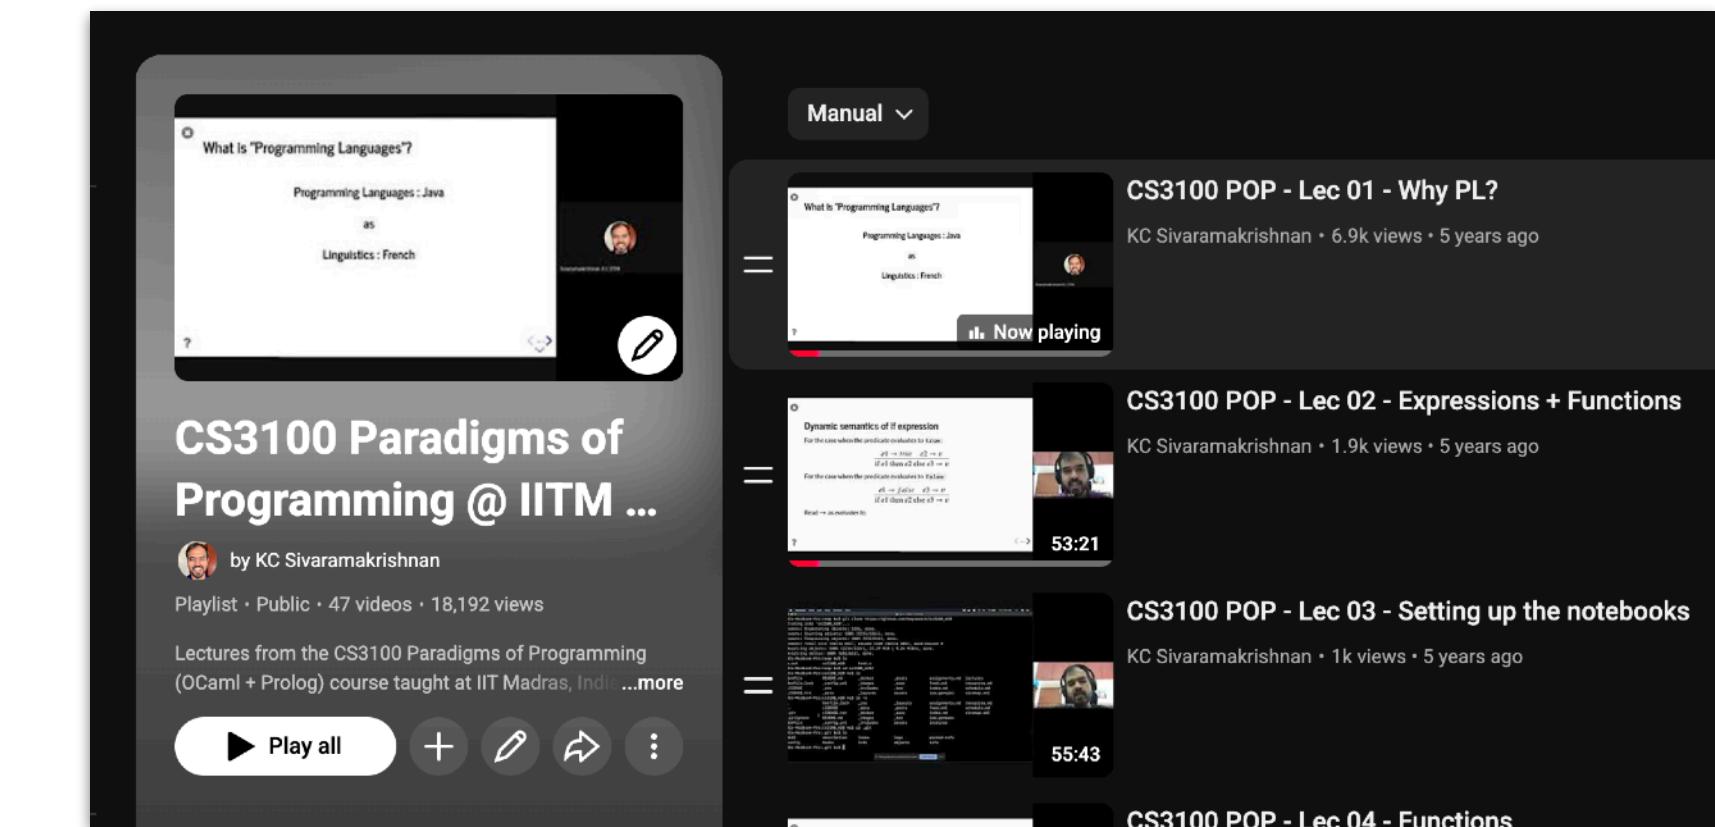
- **OCaml** will be the language of choice for the course
- **Strong Pre-requisite: CS3100 OCaml Parts**
  - *I will not cover OCaml basics in this course!*
- OCaml Resources – CS3100 @ IITM, CS3110 book from Cornell U

## CS3100: Paradigms of Programming (IITM Monsoon 2020)

This is the Github repo for the course CS3100 Paradigms of Programming taught at IITM in the Monsoon semester 2020. The course website is here: [https://kcsrk.info/cs3100\\_m20/](https://kcsrk.info/cs3100_m20/). The course also has a [YouTube playlist](#) of all the lectures. The repo includes all the lecture notes, slide deck and assignments.

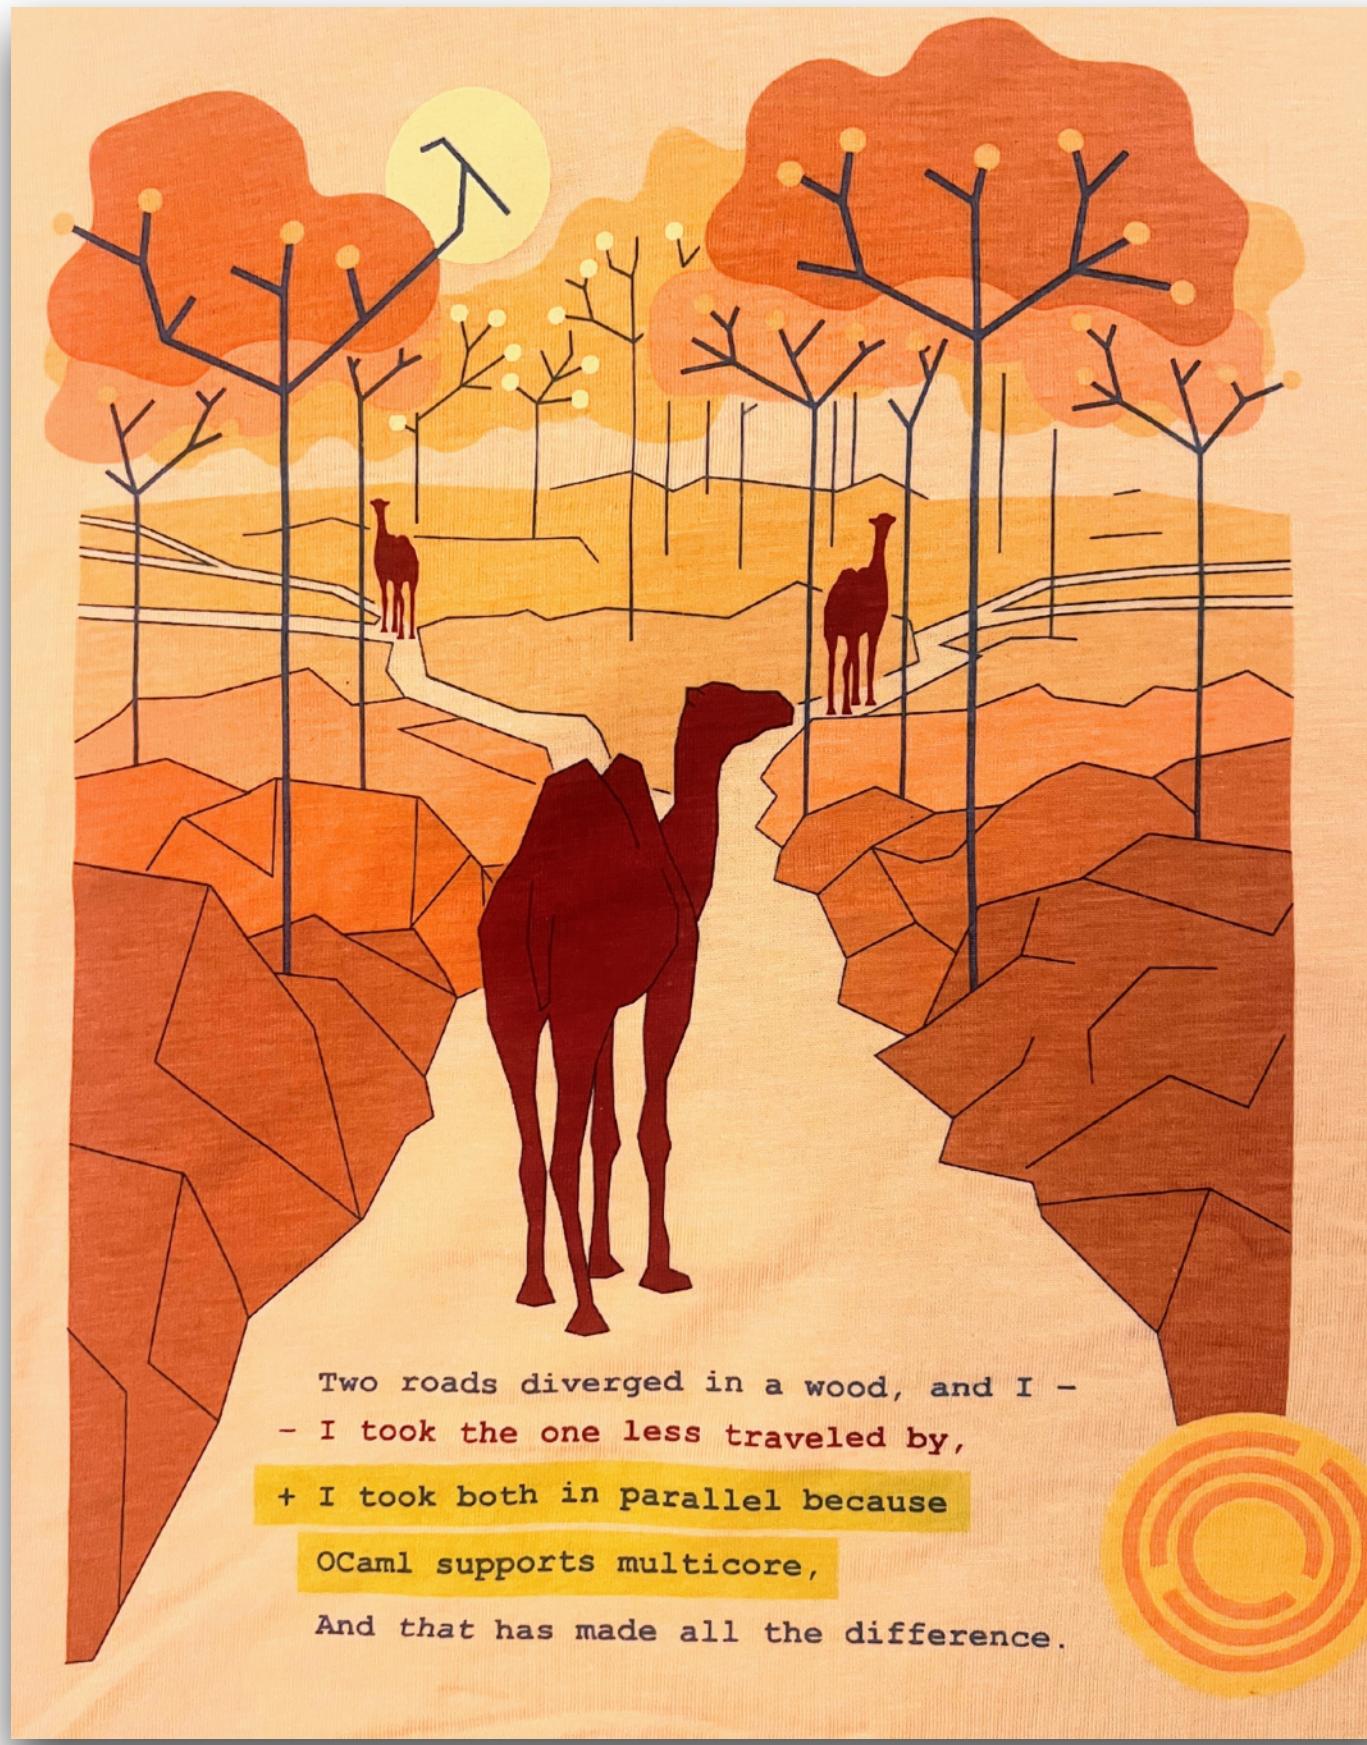
The course teaches OCaml and Prolog.

[https://github.com/kayceesrk/cs3100\\_m20](https://github.com/kayceesrk/cs3100_m20)



[YouTube PlayList](#)

# Course Language



The screenshot shows the OxCaml website. At the top, there is a navigation bar with links for "About", "Documentation", and "Get OxCaml". The Jane Street logo is also present. Below the navigation, there is a large blue banner featuring a white ox and the text "OCaml, Oxidized!". The main content area contains a section titled "OxCaml" with a description of what it is and its purpose.

**OxCaml**

OxCaml is a fast-moving set of extensions to the OCaml programming language.

It is both Jane Street's production compiler, as well as a laboratory for experiments focused towards making OCaml better for performance-oriented programming. Our hope is that these extensions can over time be contributed to upstream OCaml.

OCaml 5 – concurrency & parallelism

OxCaml – safe, performance-oriented extension of OCaml

# Evaluation Plan

- 6 in-class short quizzes (**20%**)
  - 15 minutes each, announced
  - 2 each – before quiz 1, before quiz 2, before the end sem
  - Best 5/6
- Mid-term (**20%**)
- End Sem (**20%**)
- Programming assignments x 4 (**24%**)
  - Individual
- Research mini project (**16%**)
  - Groups of 2

Sign Honour Code;  
Automated plagiarism check;  
Institute policy on malpractice

*Don't use LLM; they hinder learning*

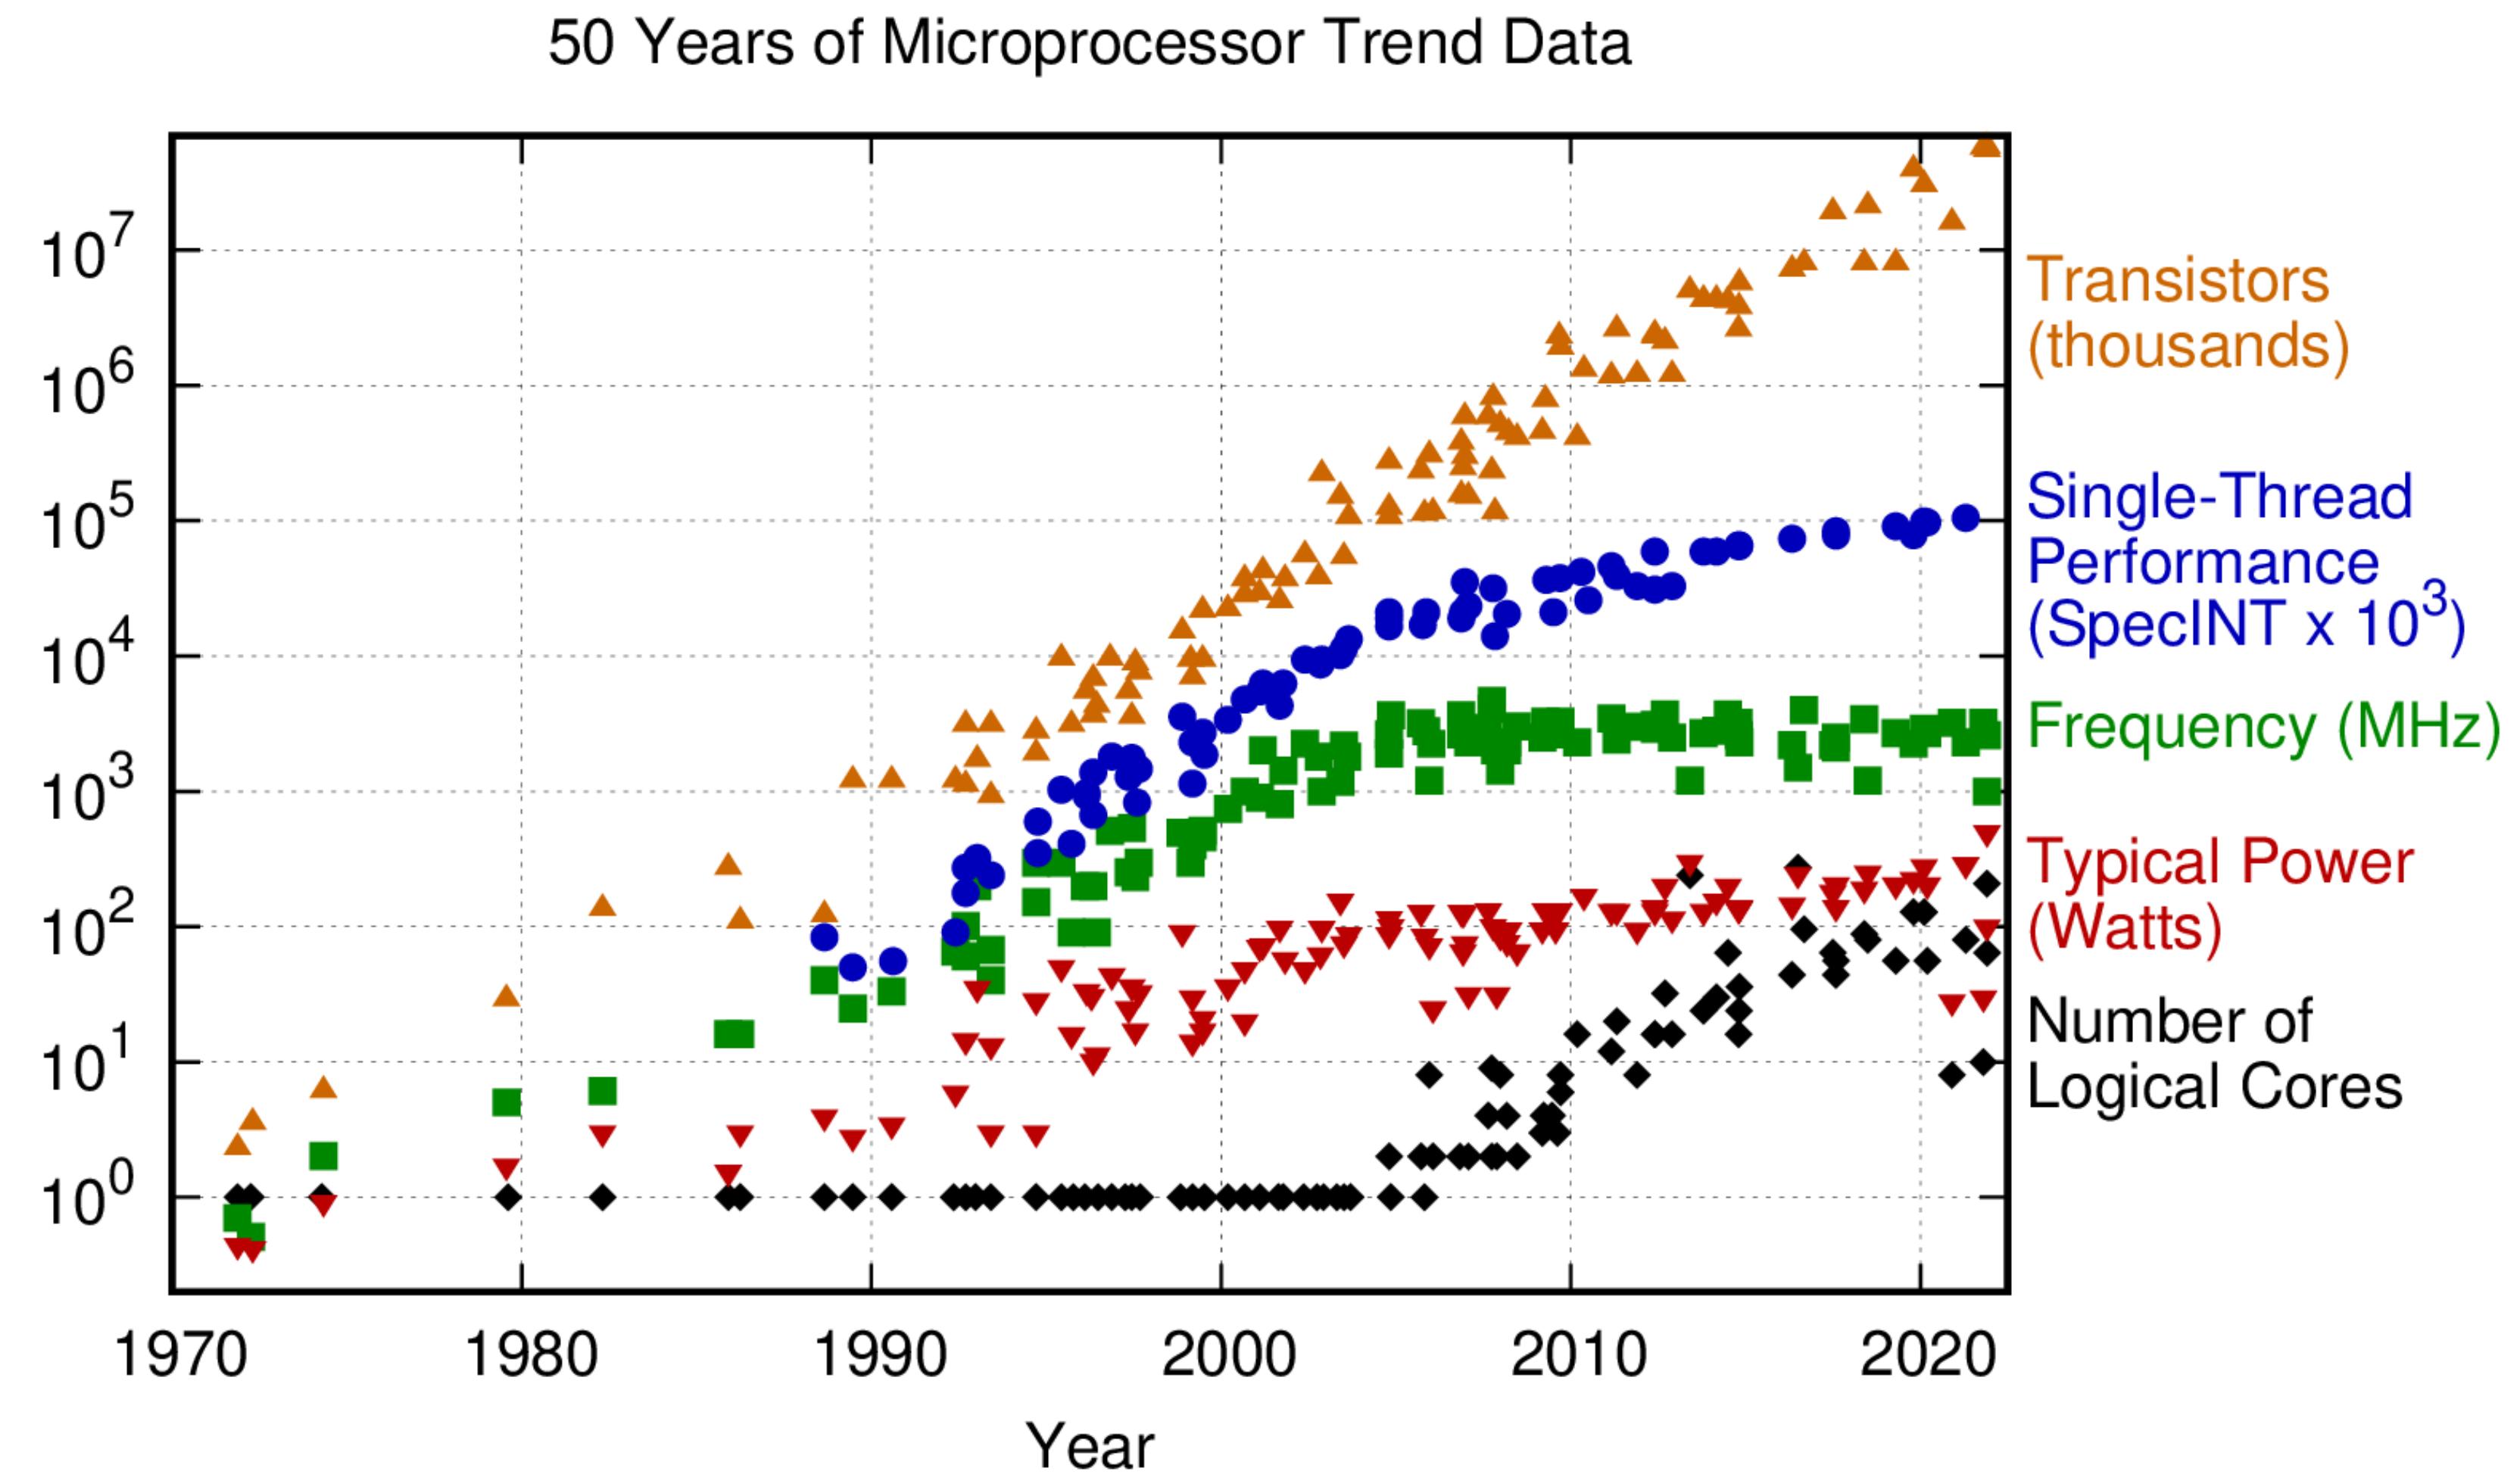
*Use LLM; they make you productive\**

\* if you know what you want

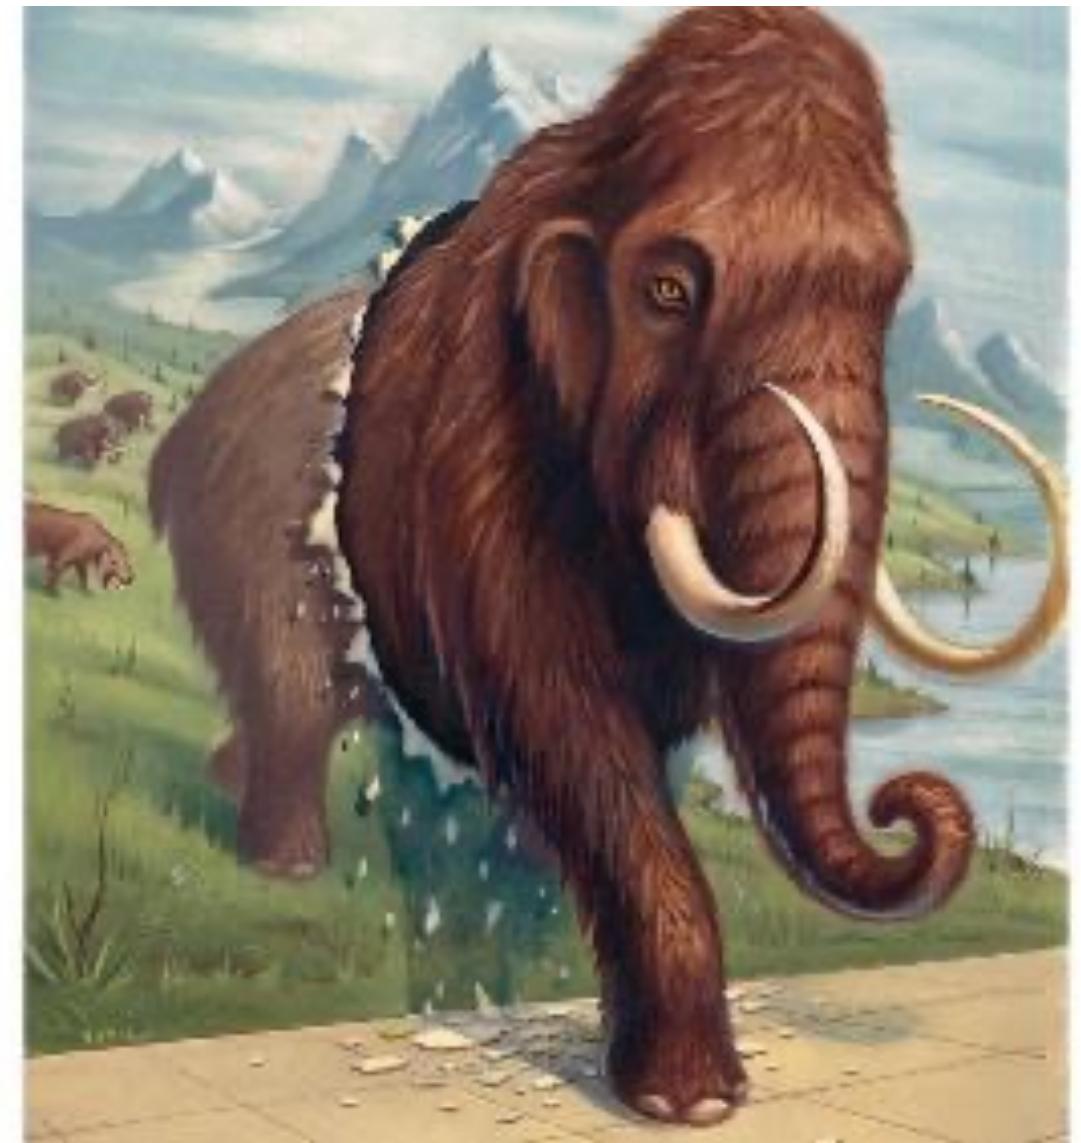
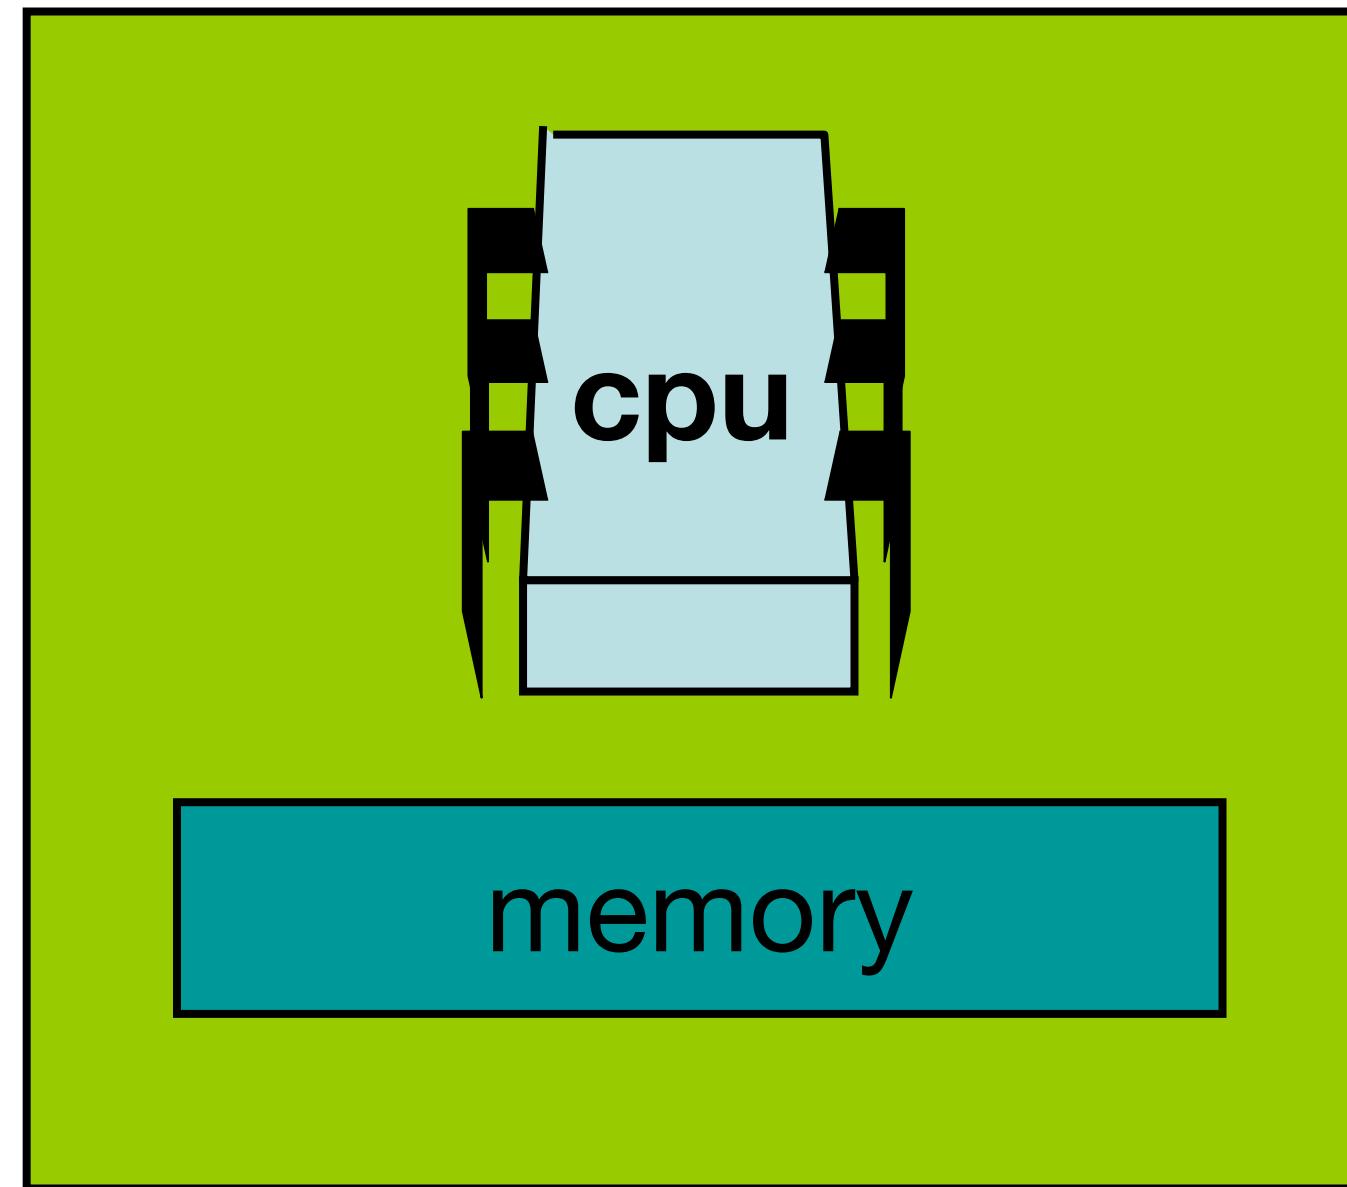
# Parallelism

# Moore's Law

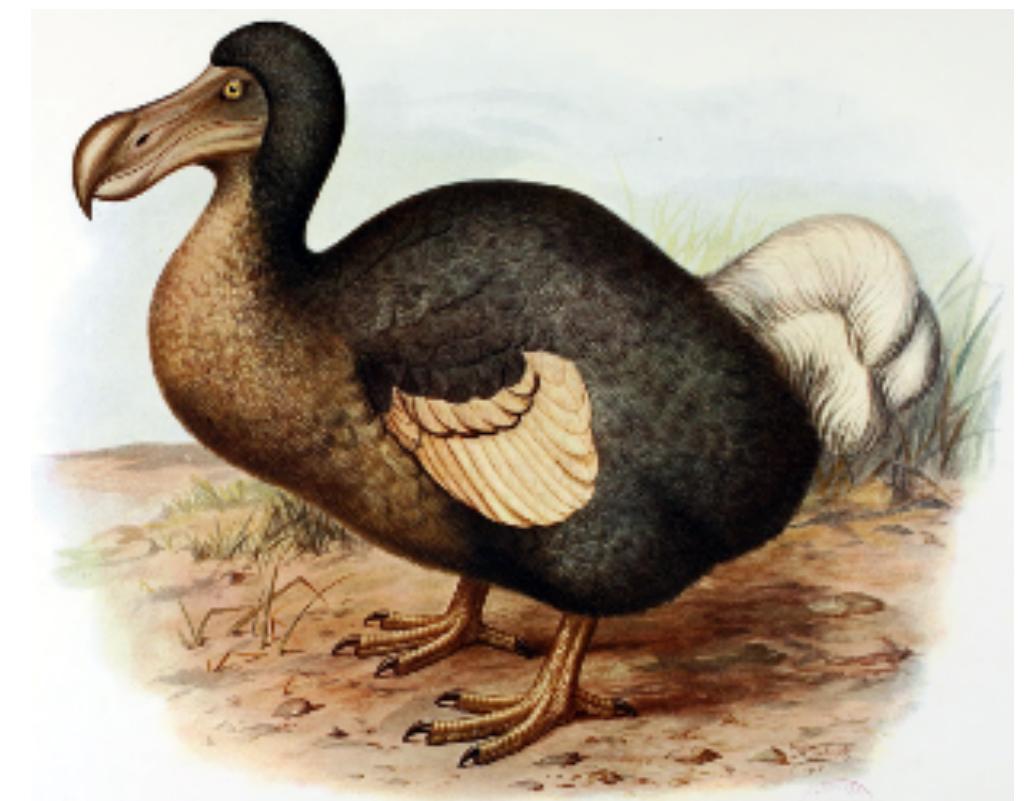
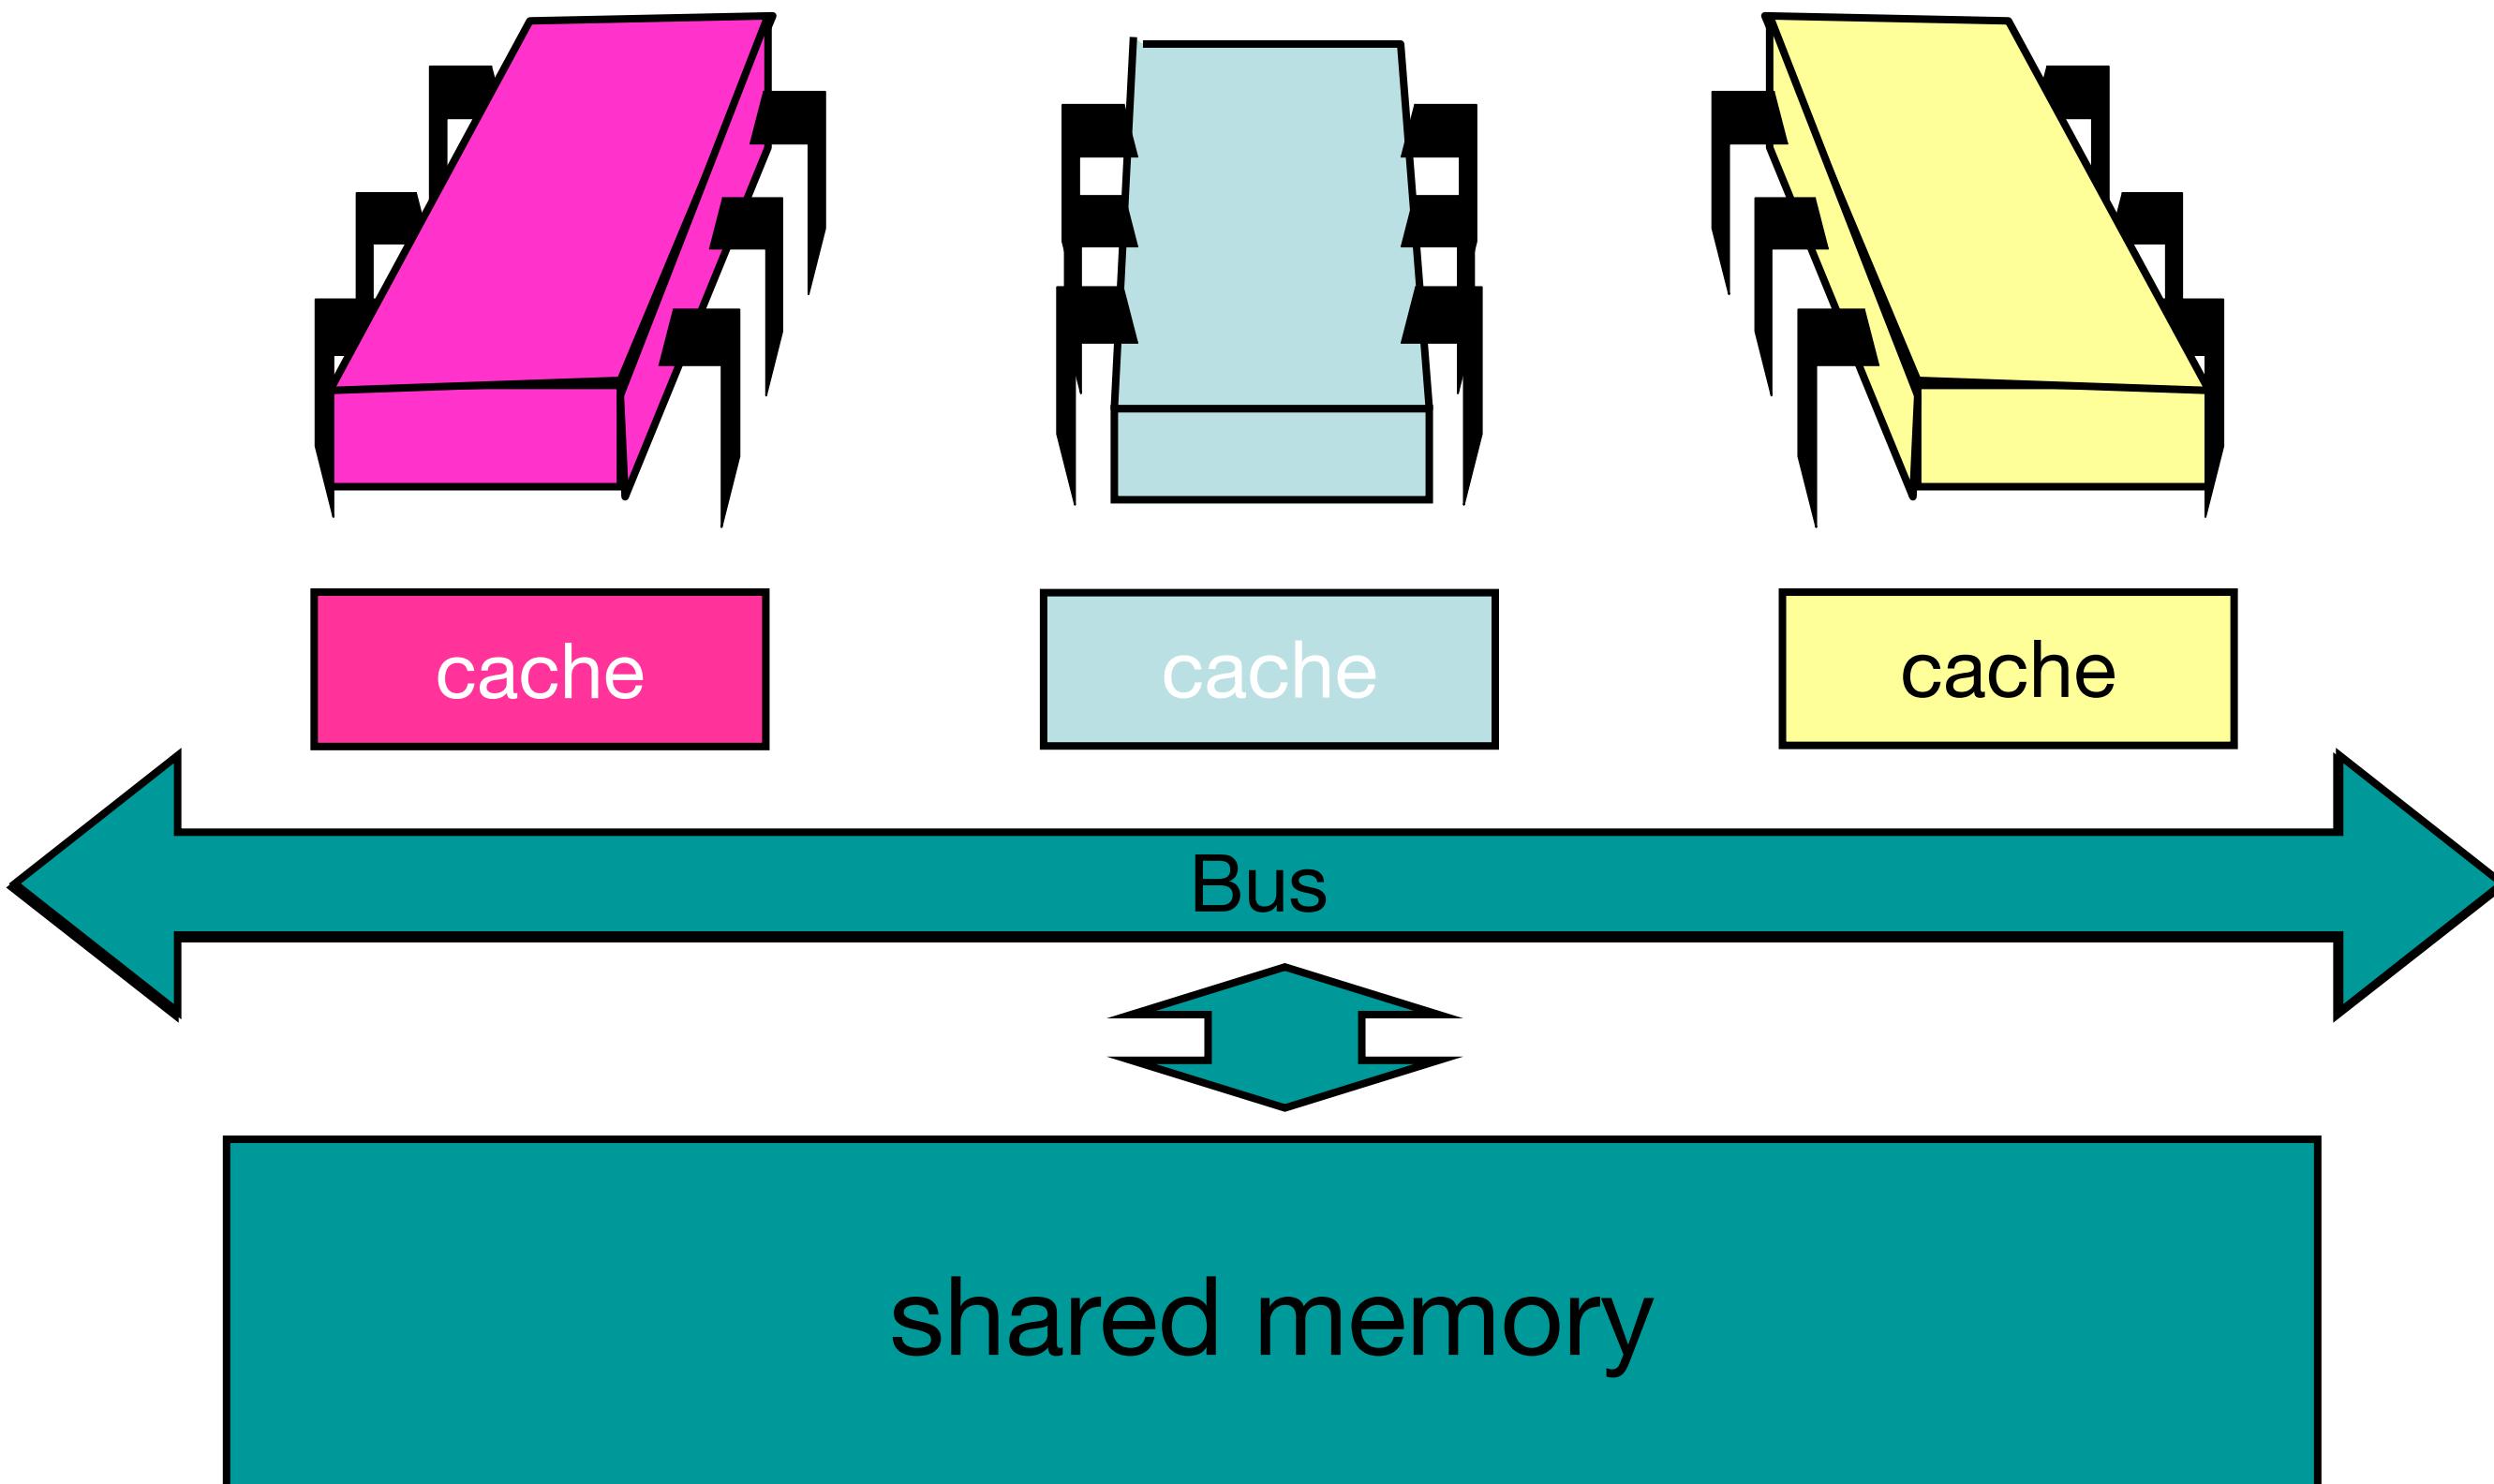
*The number of transistors on an integrated circuit doubles roughly every 18 to 24 months, resulting in exponential growth in computing power.*



# Uniprocessor (Extinct)

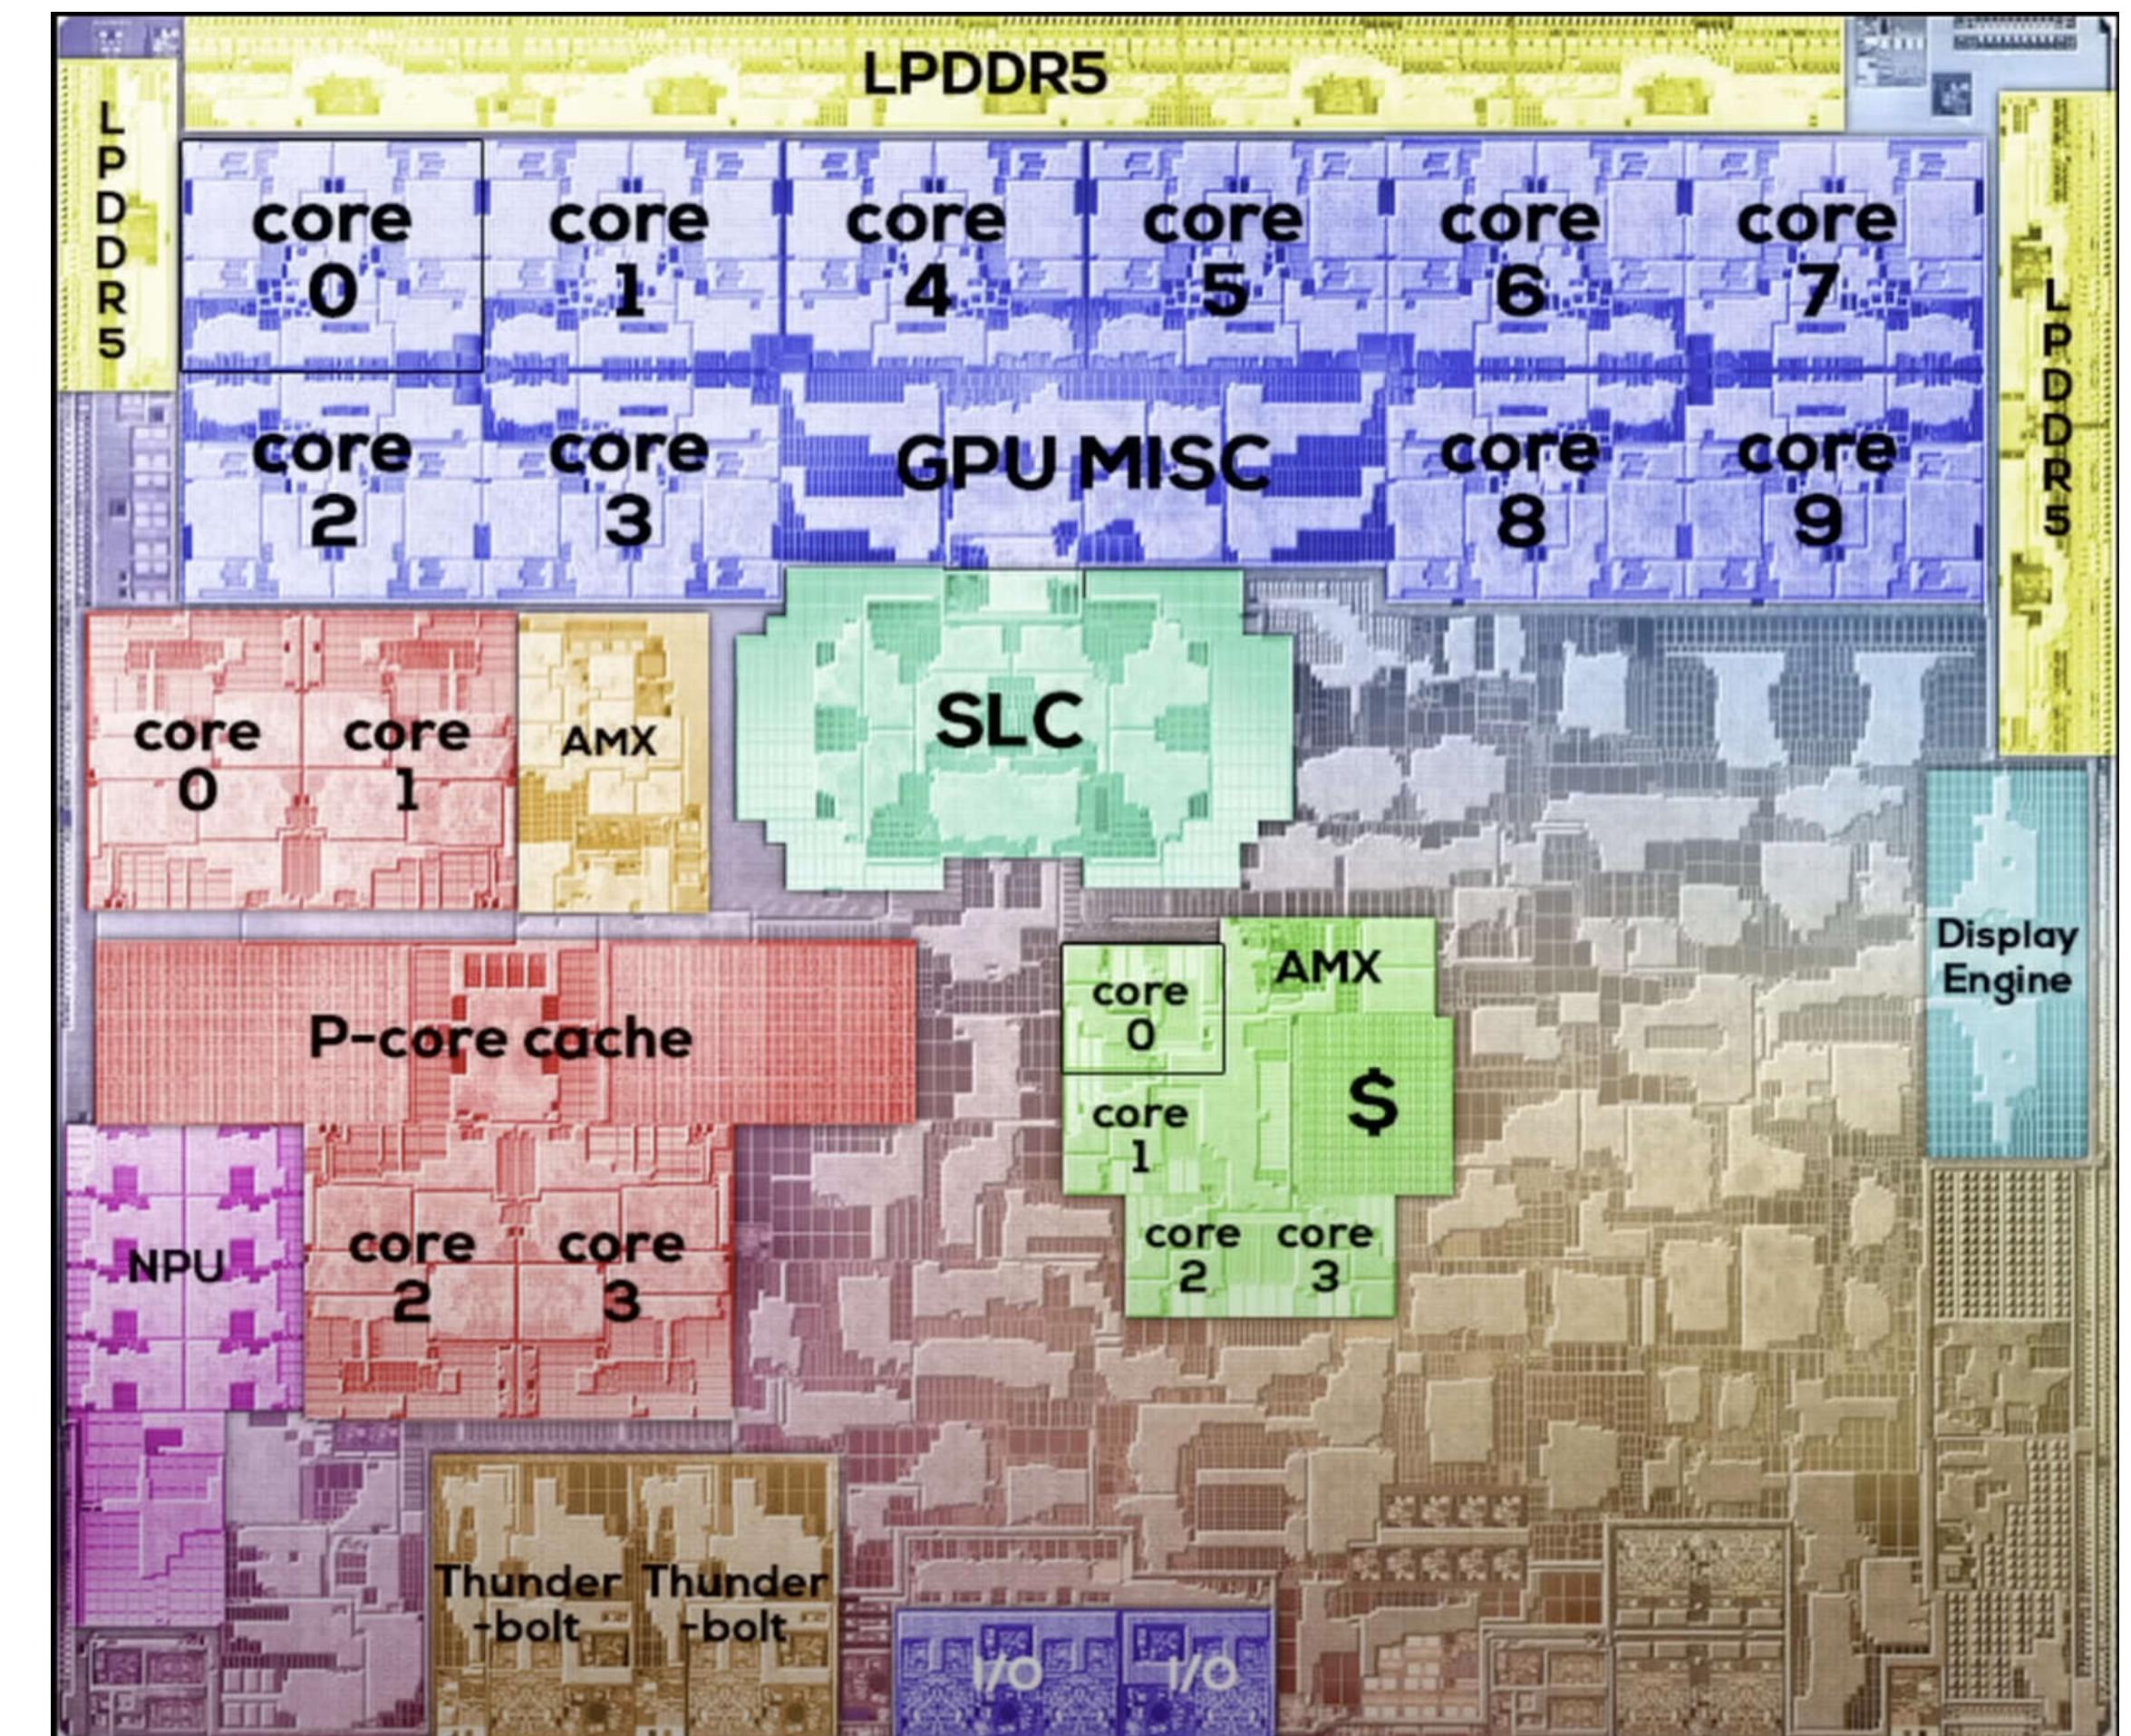


# Shared Memory Multiprocessor (SMP): Extinct



# Modern Multicore Processors

- Various names for this
  - System-on-a-chip (SoC)
  - Chip multiprocessor (CMP)
  - A multicore machine

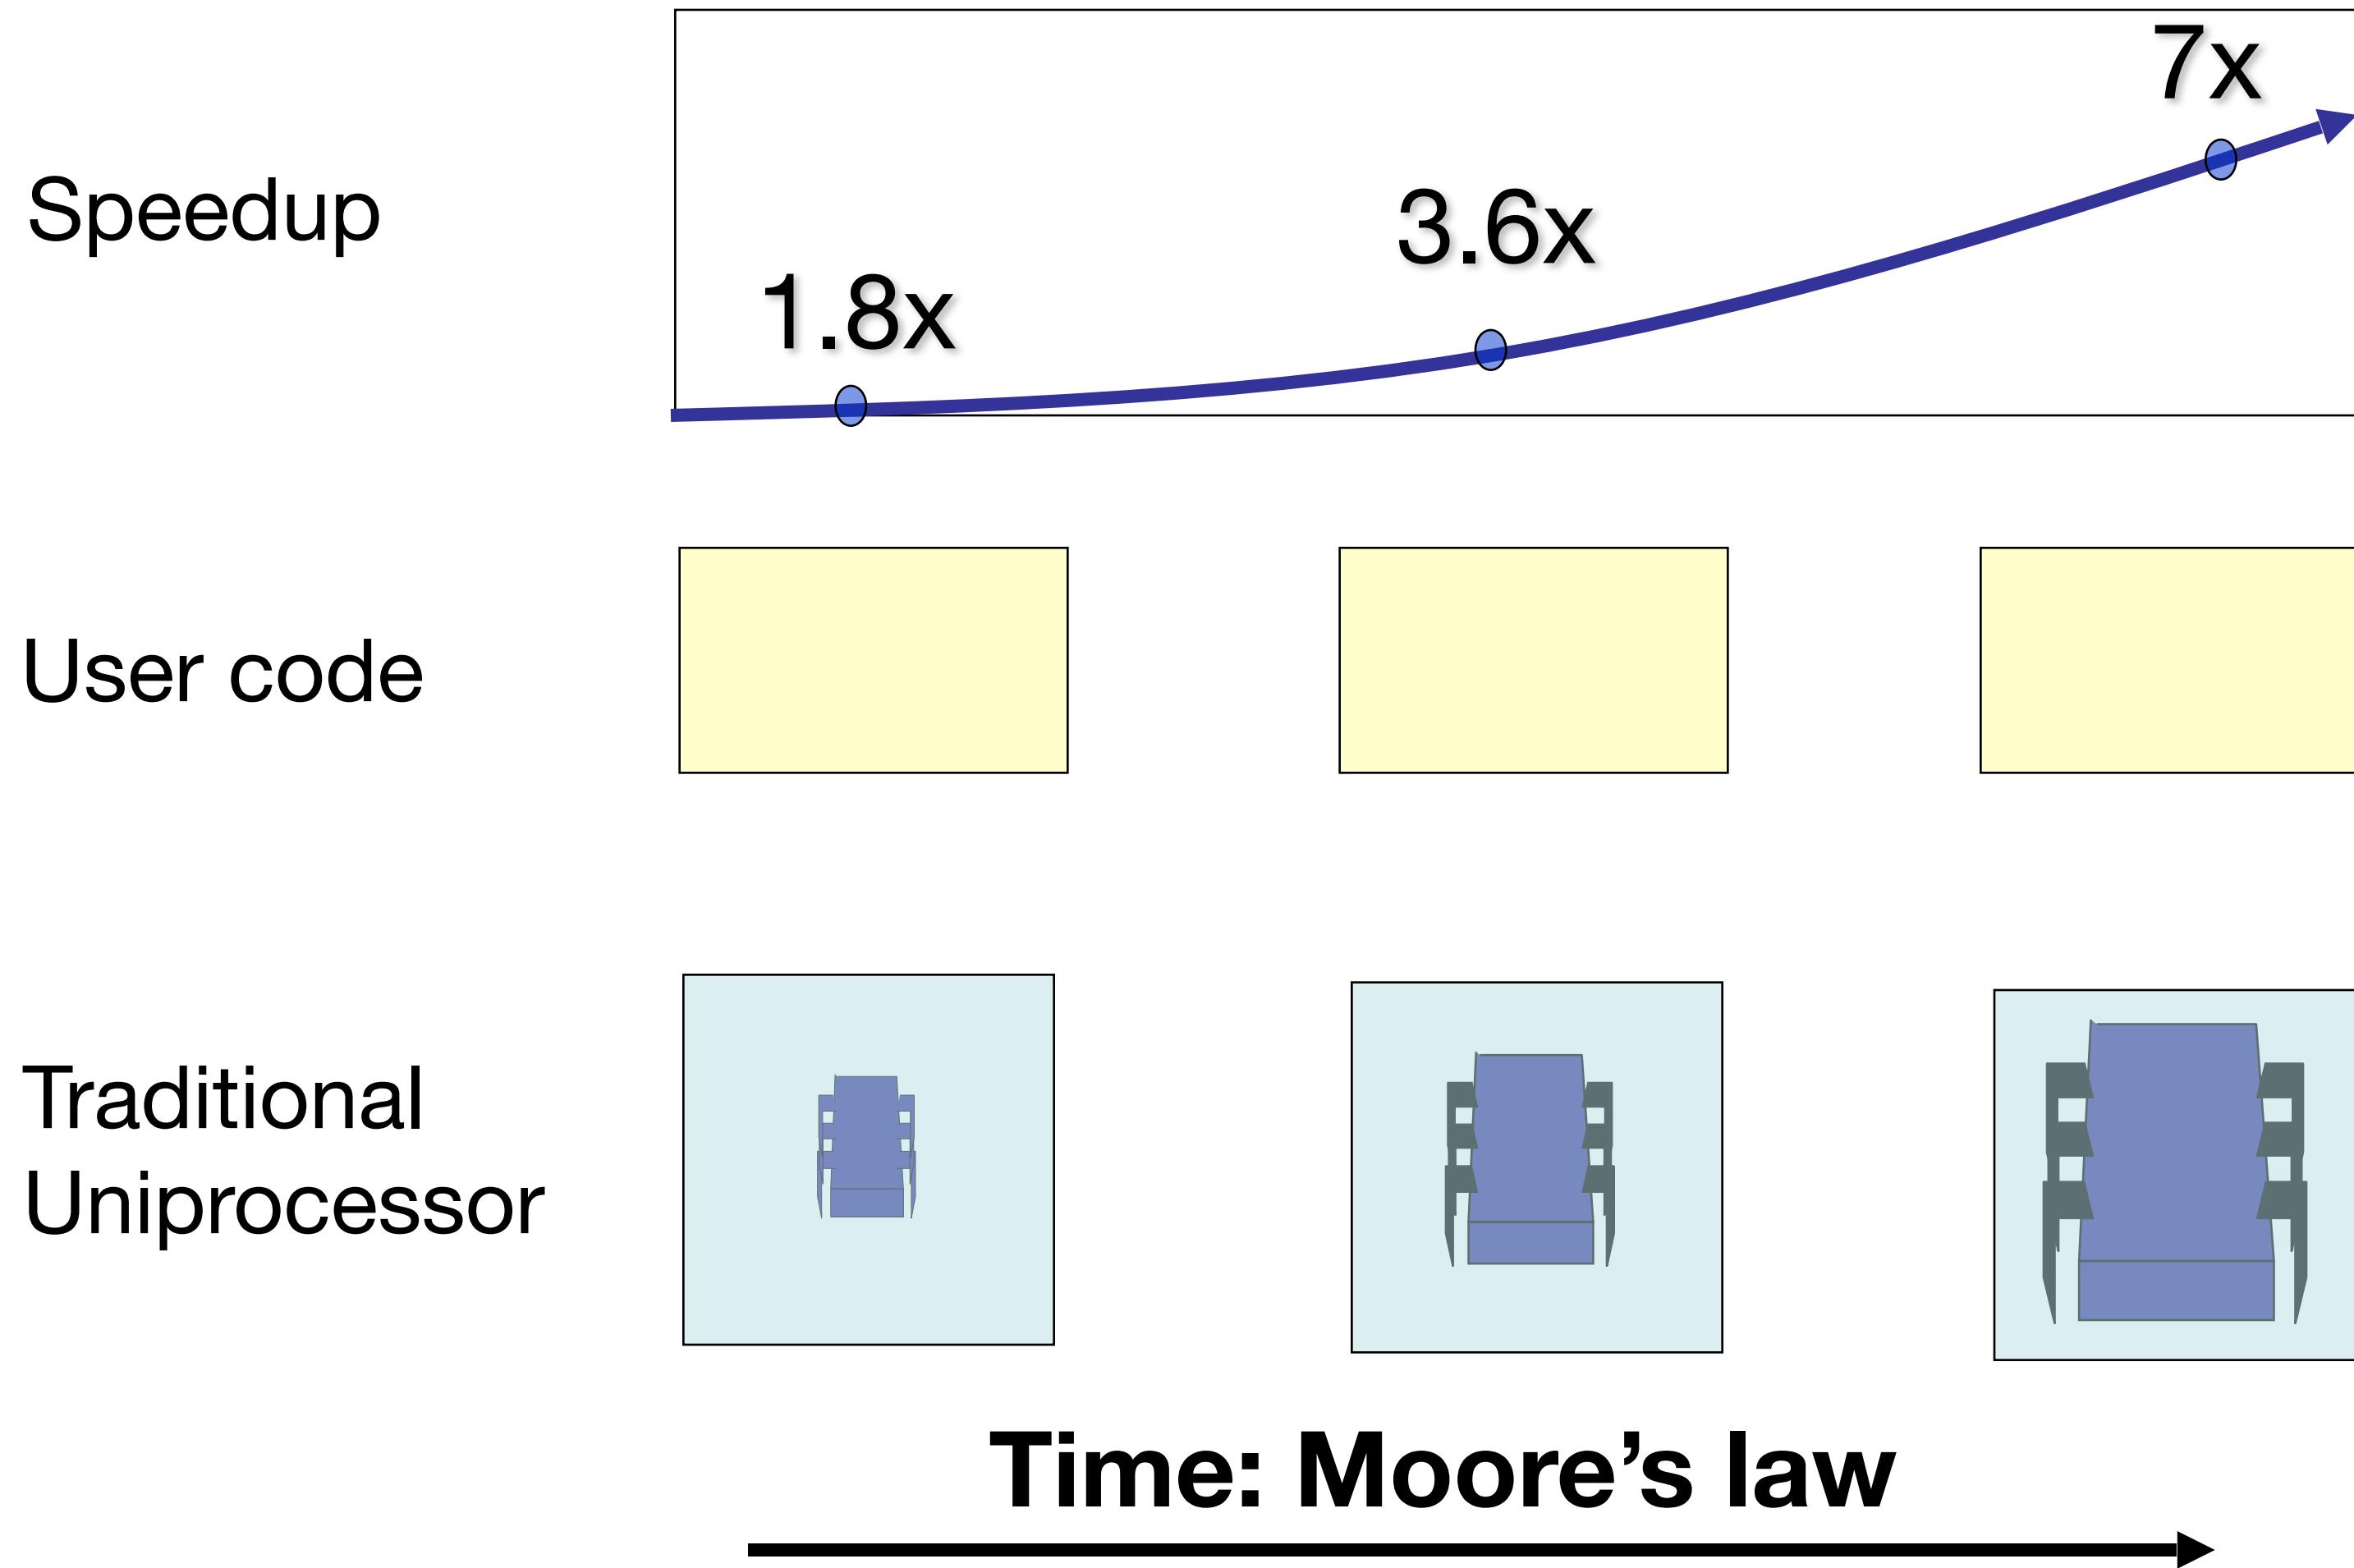


Apple M3 floor plan/die shot

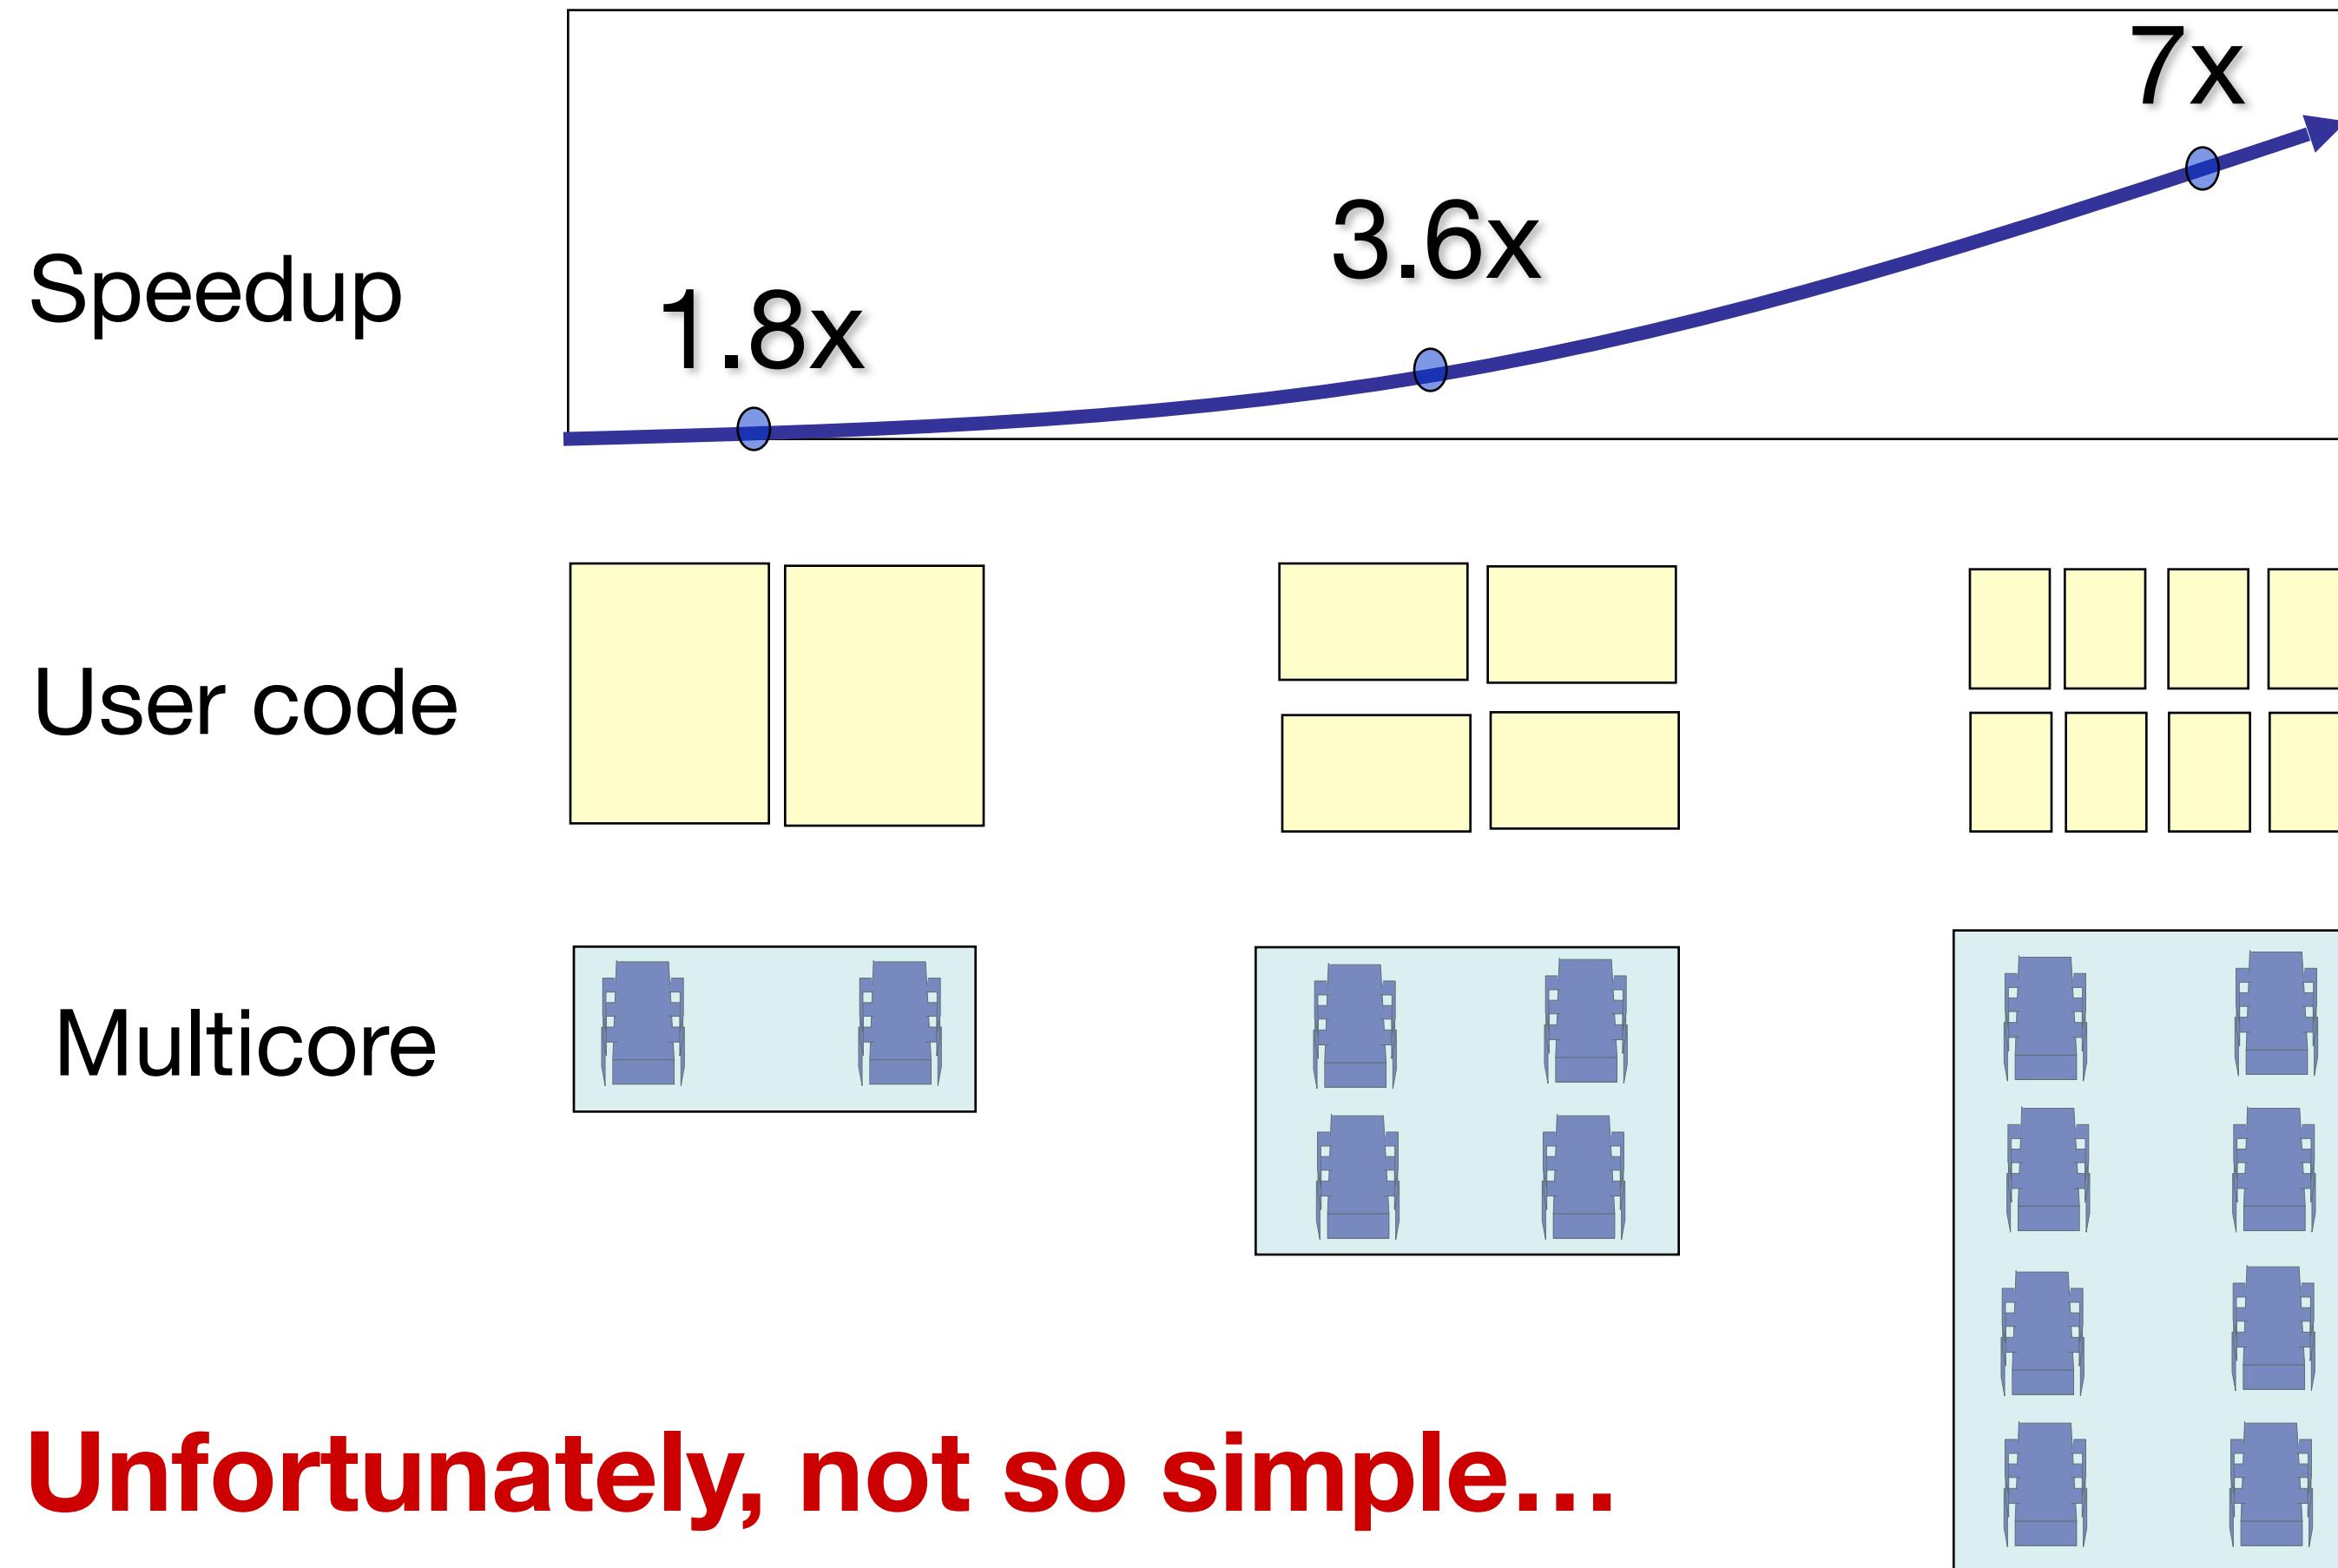
# Why do we care?

- Time no longer cures software bloat
  - The “free ride” is over
- When you double your program’s path length
  - You can’t just wait 6 months
  - Your software must somehow exploit twice as much concurrency

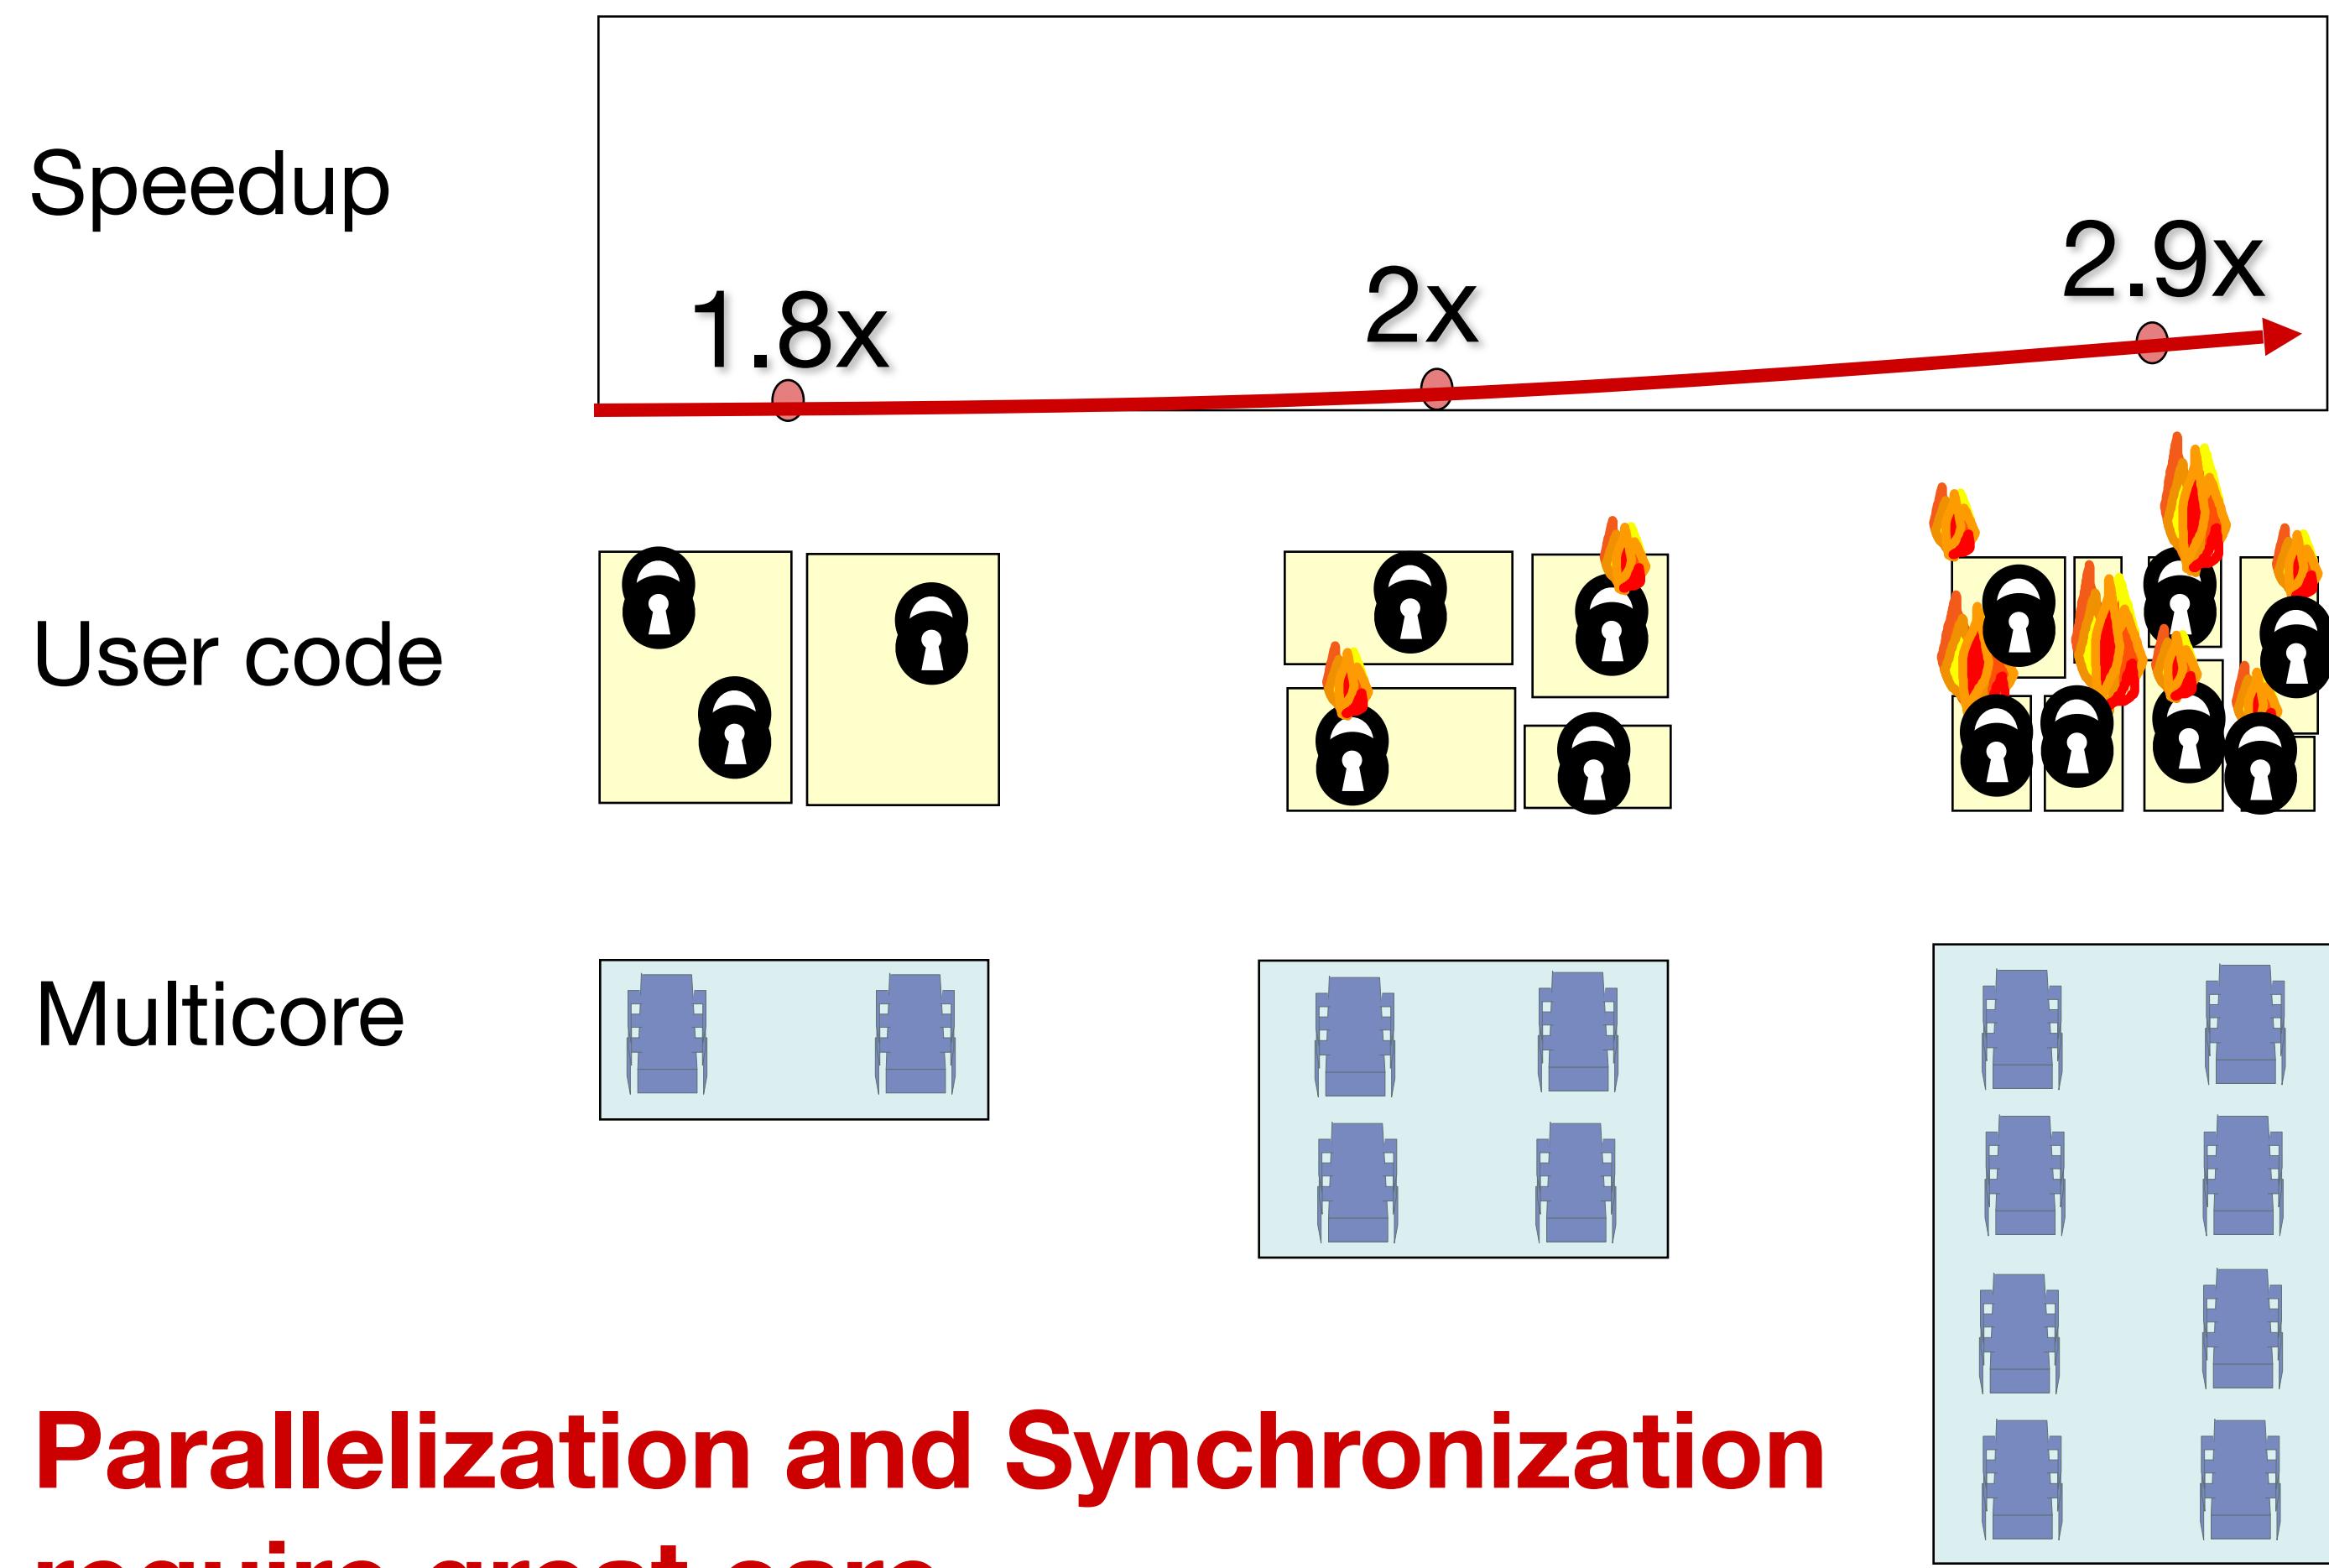
# Traditional Scaling Process



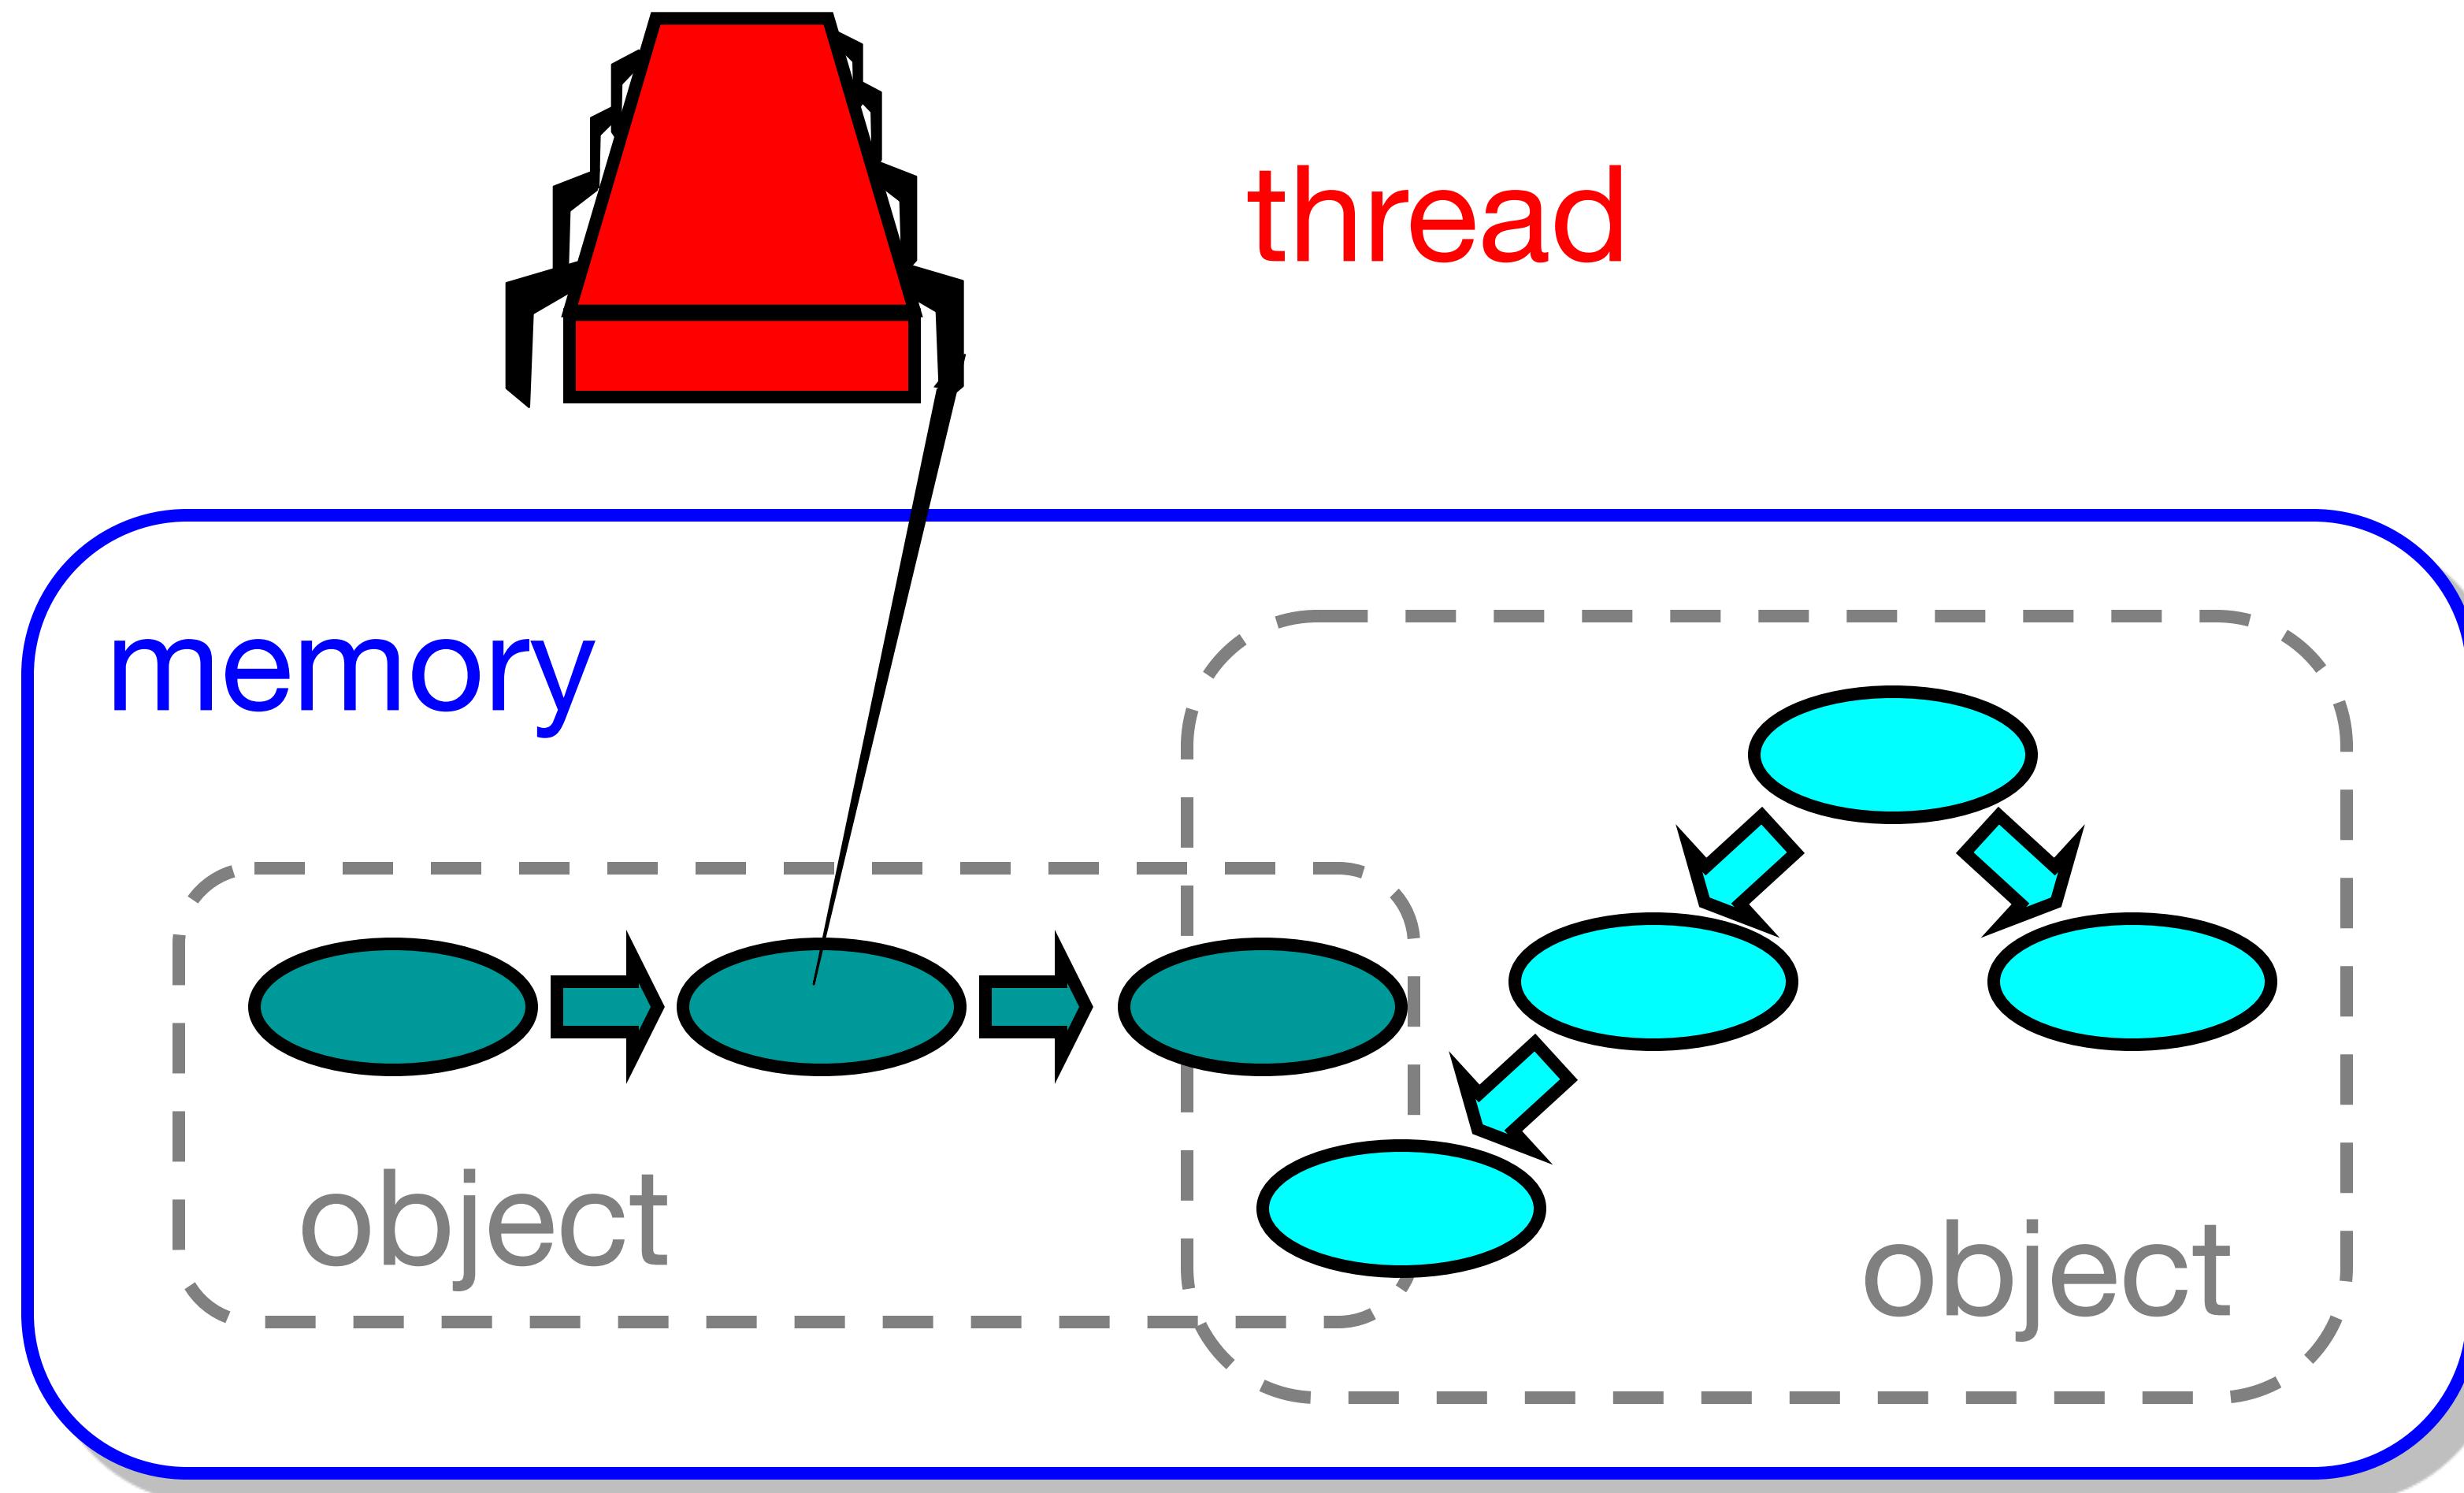
# Ideal Scaling Process



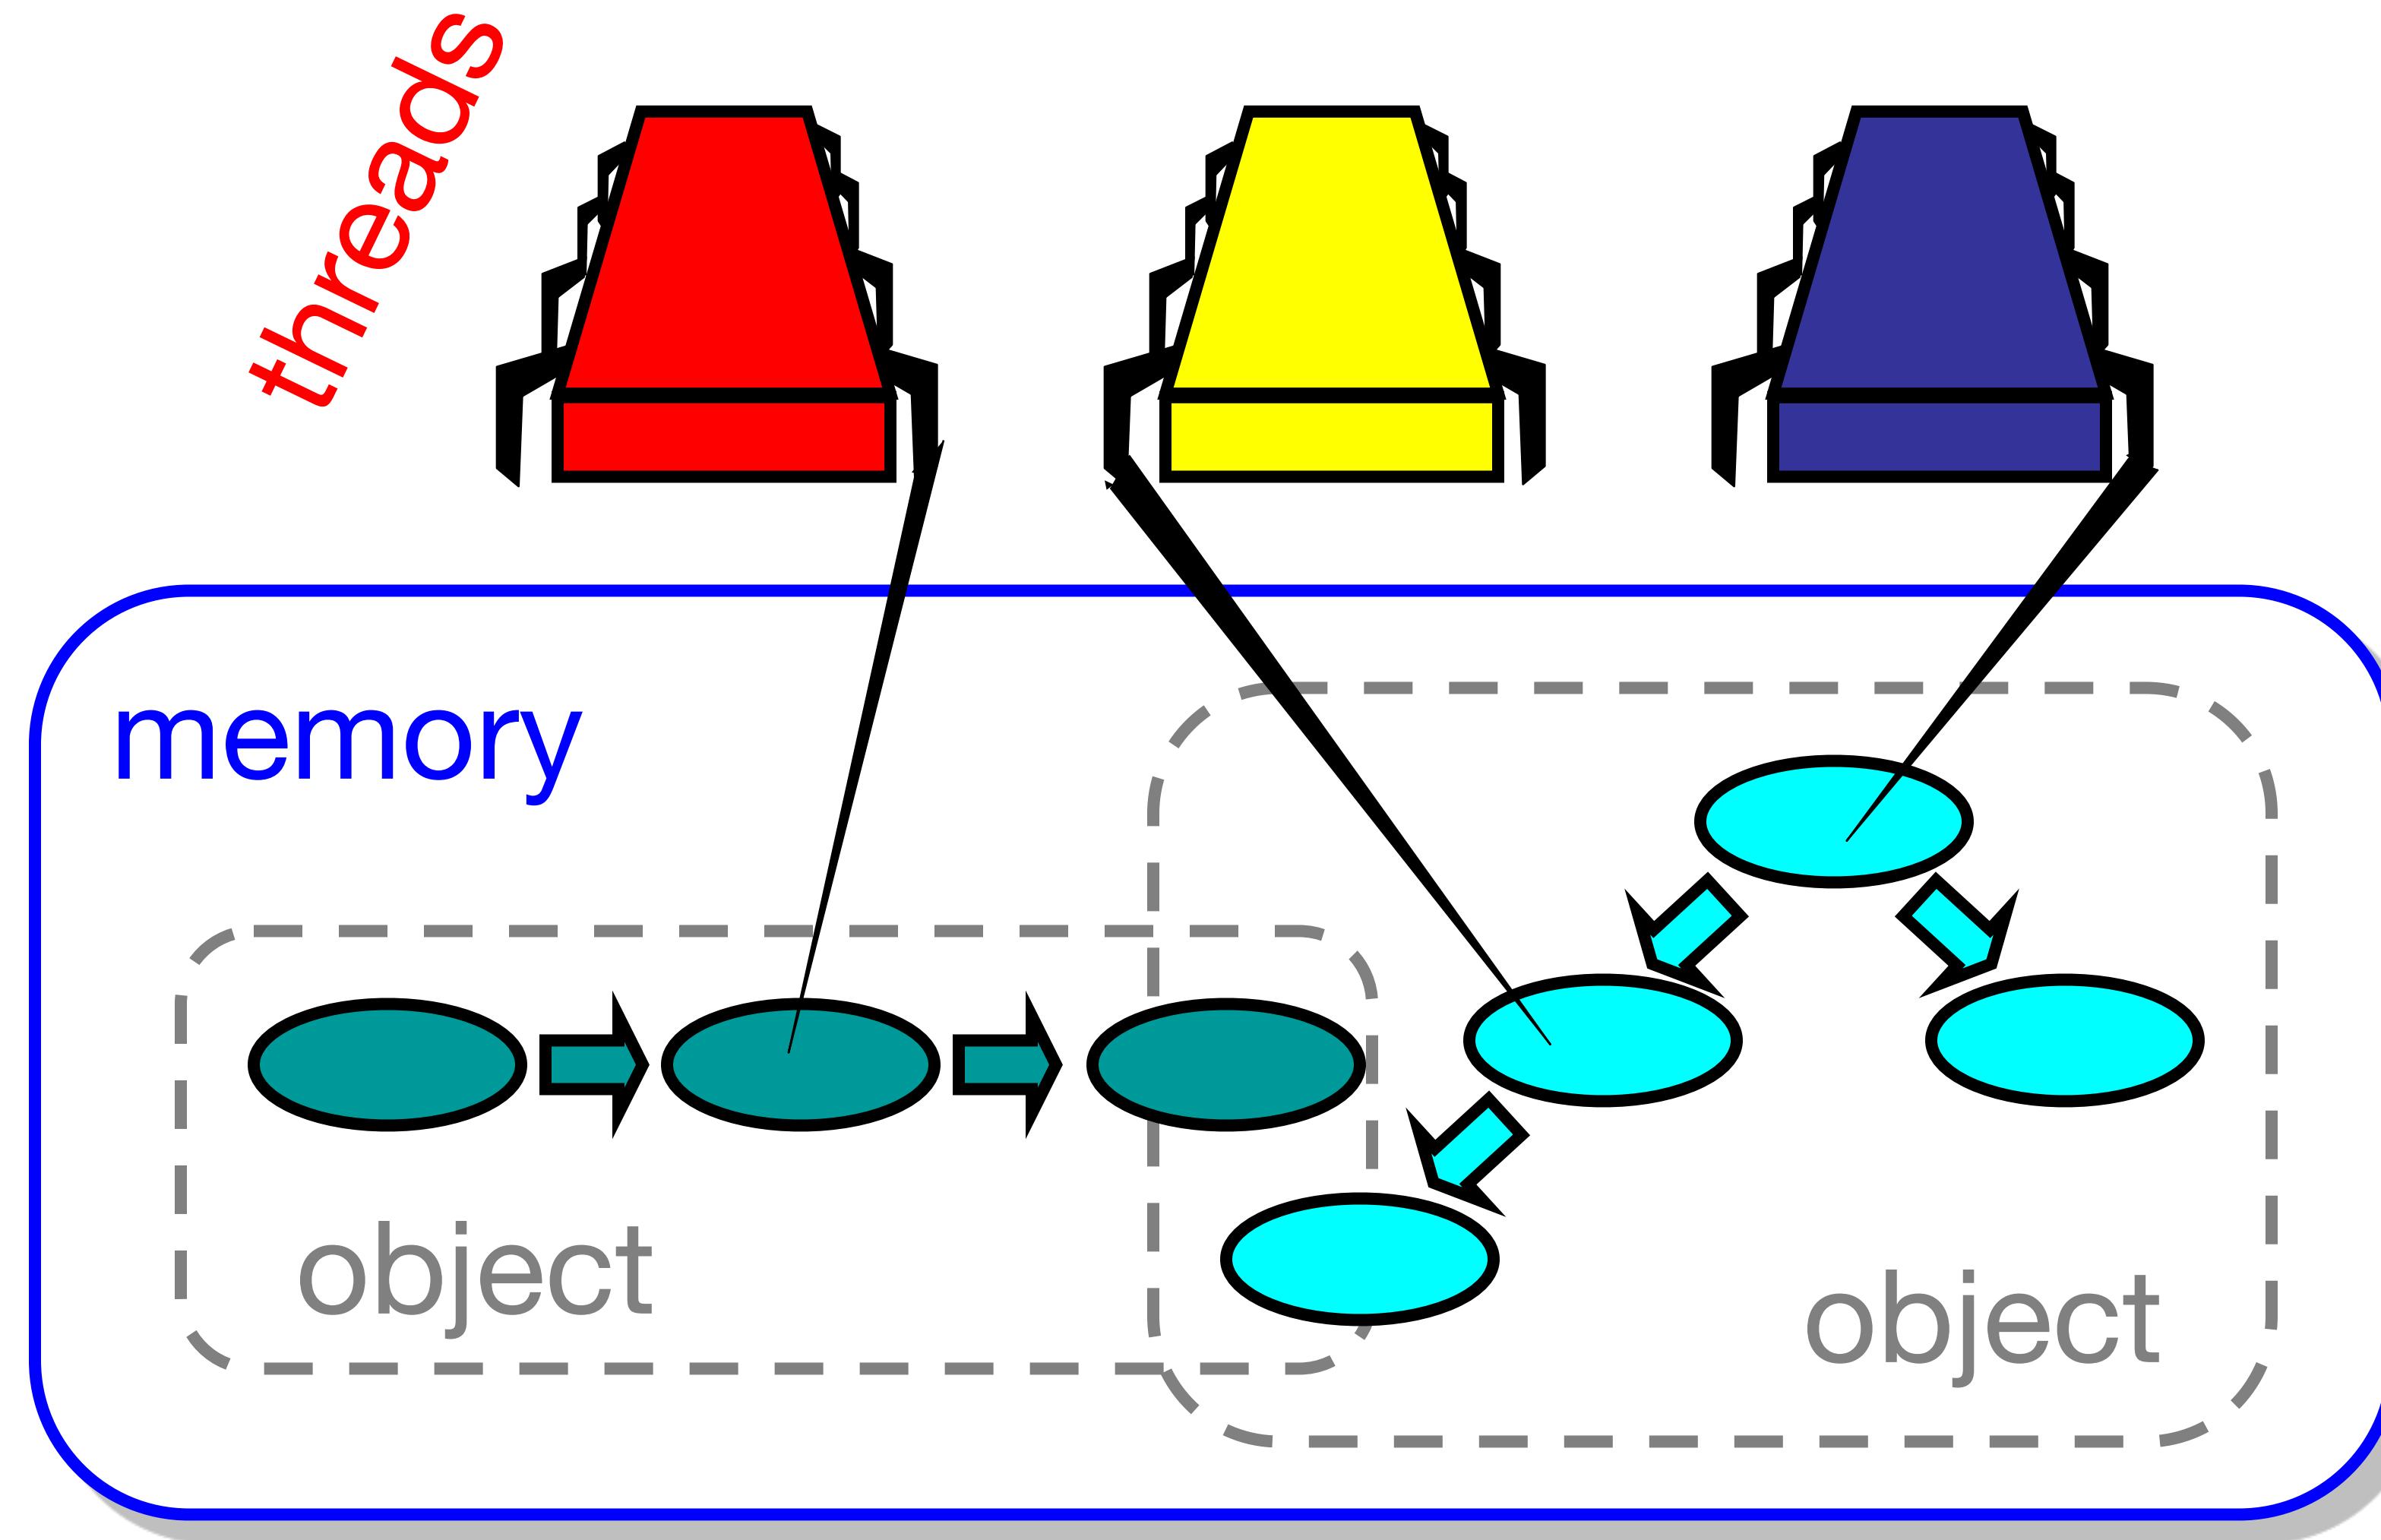
# Actual Scaling Process



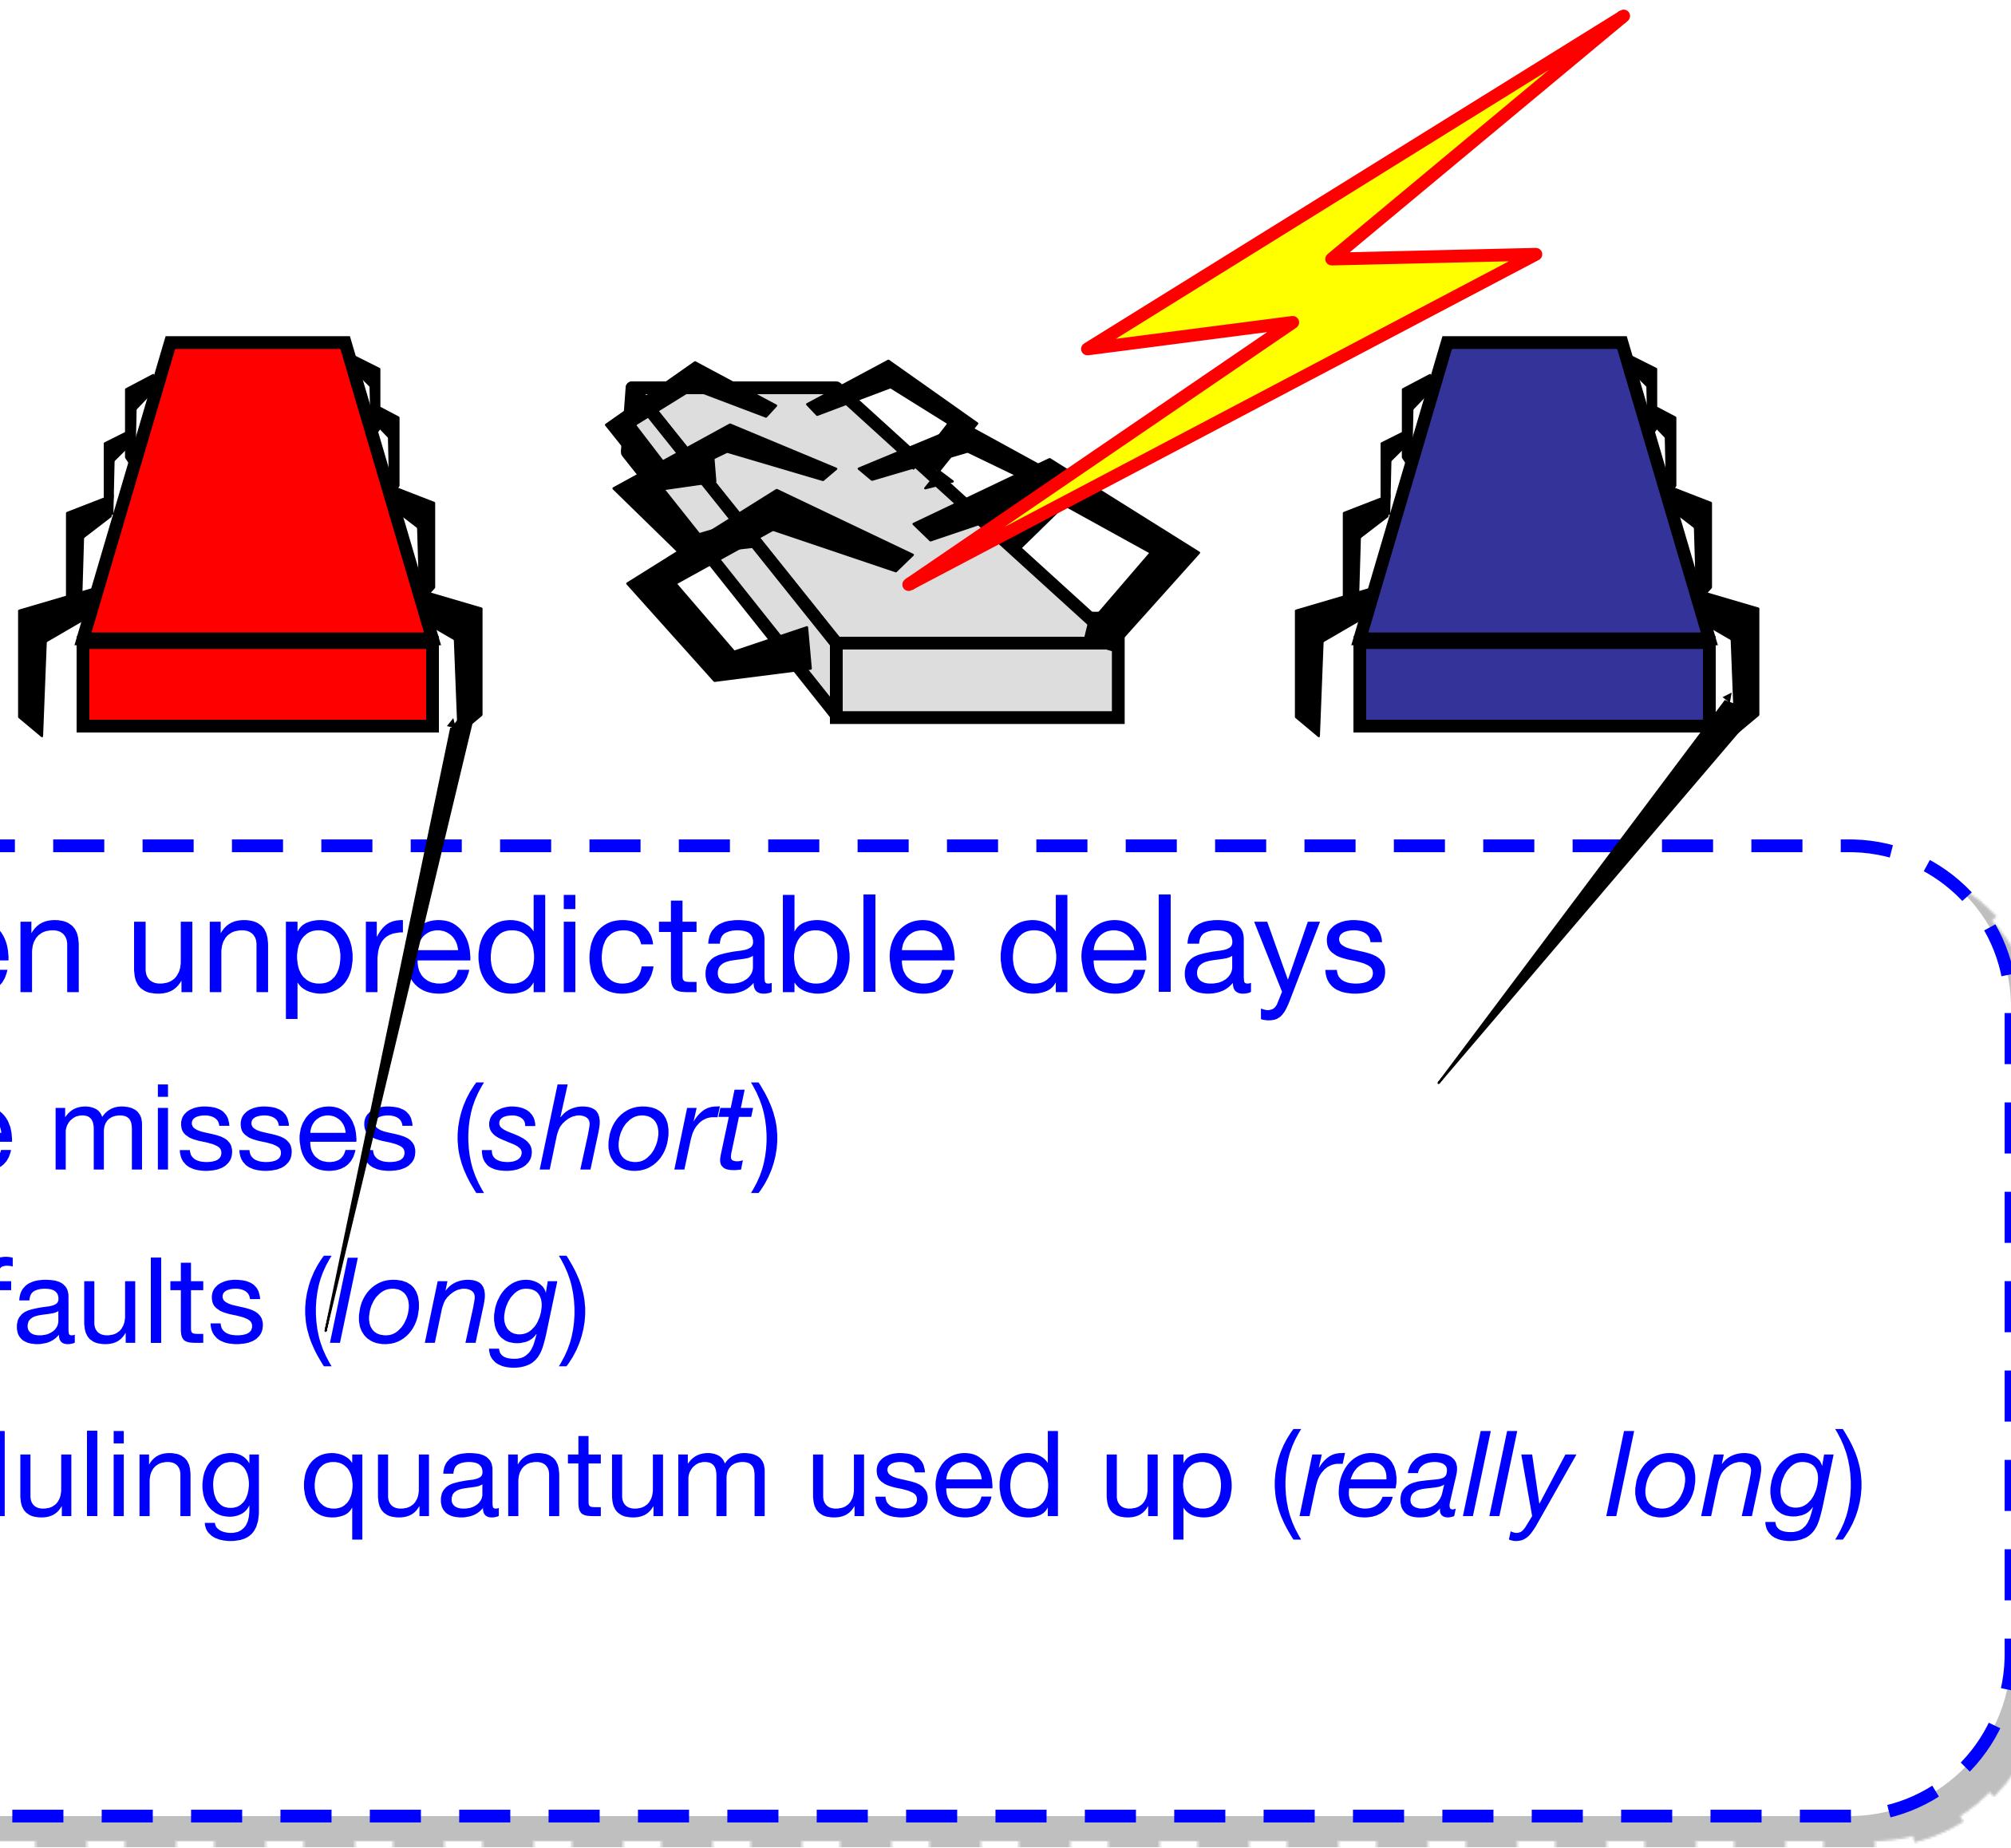
# Sequential Computation



# Concurrent Computation

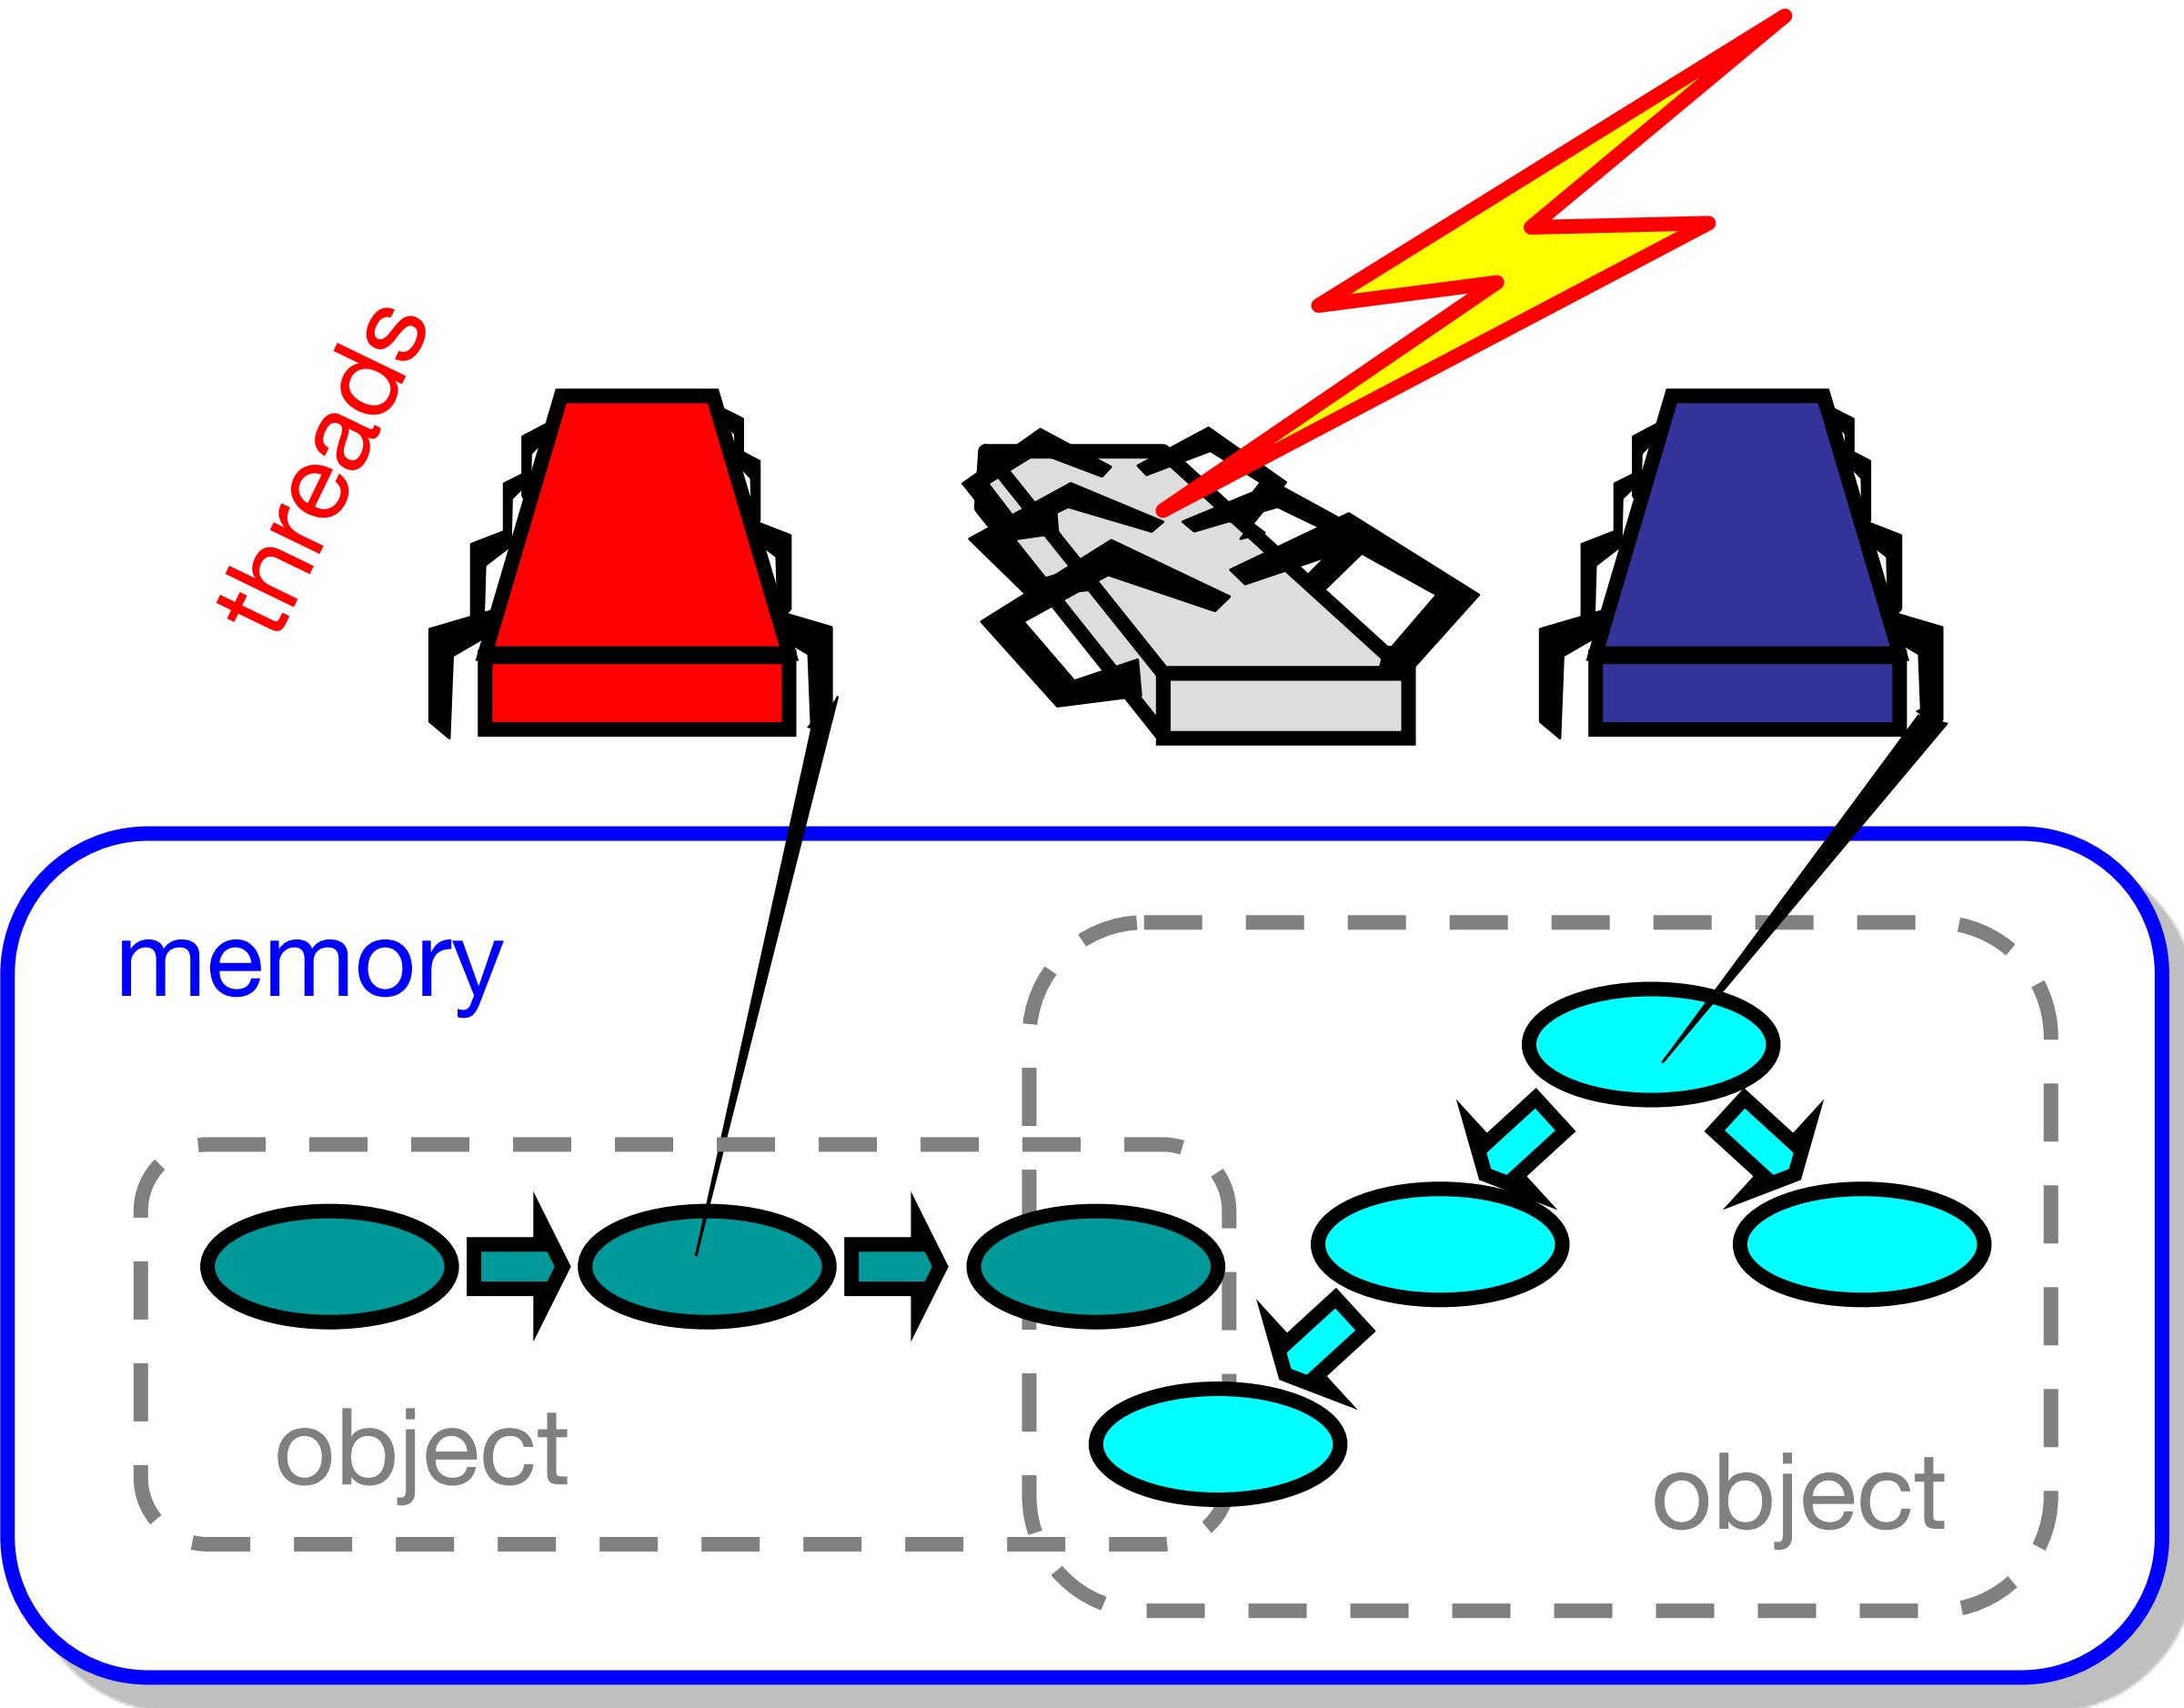


# Asynchrony



# Model Summary

- Multiple threads
  - Sometimes called processes
- Single shared memory
- Objects live in memory
- Unpredictable asynchronous delays



# Roadmap for the Parallelism side

- We are going to focus on ***principles*** first and then ***practice***
  - Start with ***idealised*** models of concurrent computations
  - Look at ***simplistic*** problems
  - Emphasise ***correctness*** over ***pragmatism***
- ***Principles*** will be foundational for the concurrency parts that we will study later
- “Correctness may be theoretical, but incorrectness has practical impact”

# Designing Concurrent Programs

# Software Setup

- OCaml 5.4
- Use vscode as default
  - Comes with Vim and Emacs modes
- Use OCaml vscode platform
- More instructions on the course GitHub page

## Software Setup

We will use OCaml 5.4 or later. Follow the instructions at <https://ocaml.org/docs/install.html> to install OCaml and the platform tools. Use Linux, macOS, \*BSD or WSL on Windows for best compatibility. Some of the later lectures will need Linux tools. At IITM, you can use the DCF machines, which have the tools installed.

Below is the instruction for Linux/macOS systems.

```
bash -c "sh <(curl -fsSL https://opam.ocaml.org/install.sh)"  
opam init # initialize opam  
opam switch create 5.4.0 # create a new switch with OCaml 5.4.0  
opam install ocaml-lsp-server odoc ocamlformat utop dune # install useful packages
```



It is recommended that you use VSCode with the [OCaml Platform](#) extension for development.

Refer to individual lecture directories for specific setup instructions.

# Concurrency Jargon

- Hardware
  - Processors
- Software
  - Threads/Processes/“Domains”
- Domains are unit of parallel execution in OCaml
  - Spawn and join domains using `Domain.spawn` and `Domain.join`
  - See also
    - Domain module in OCaml Standard library <https://ocaml.org/manual/5.4/api/Domain.html>
    - Parallelism Chapter in OCaml manual <https://ocaml.org/manual/5.1/parallelism.html>
- Sometimes ok to confuse them, sometimes not

# Recursive Fibonacci in Parallel

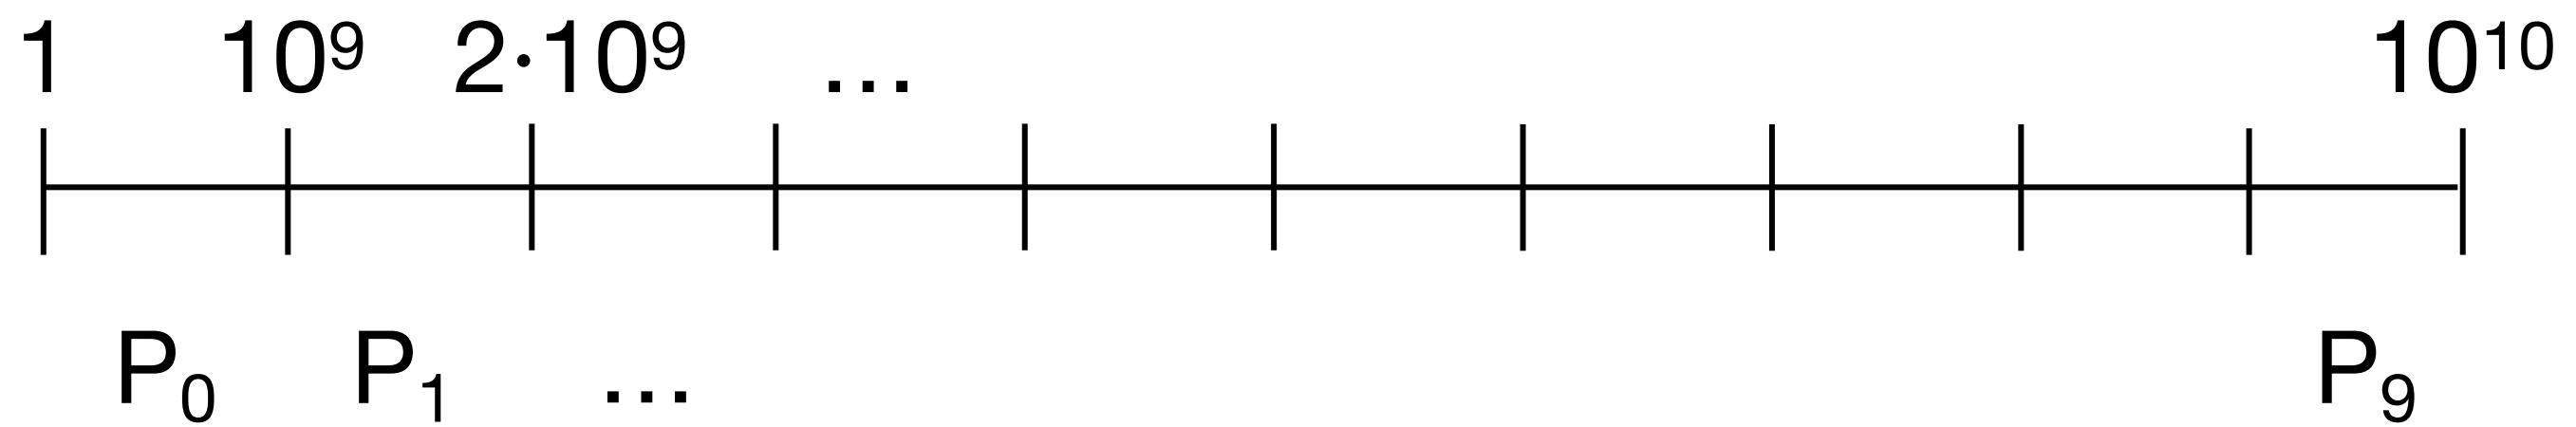
- Challenge
  - Compute recursive fib(N) twice
- Embarrassingly parallel computation
  - No dependencies between the two calls
- Use Domains to parallelise the computation

**Demo**

# Parallel Primality Testing

- Challenge
  - Print primes from 1 to  $10^{10}$
- Given
  - Ten-processor multiprocessor
  - One thread per processor
- Goal
  - Get ten-fold speedup (or close)

# Load balancing



- Split the work evenly
- Each thread tests range of  $10^9$

# Procedure for thread i

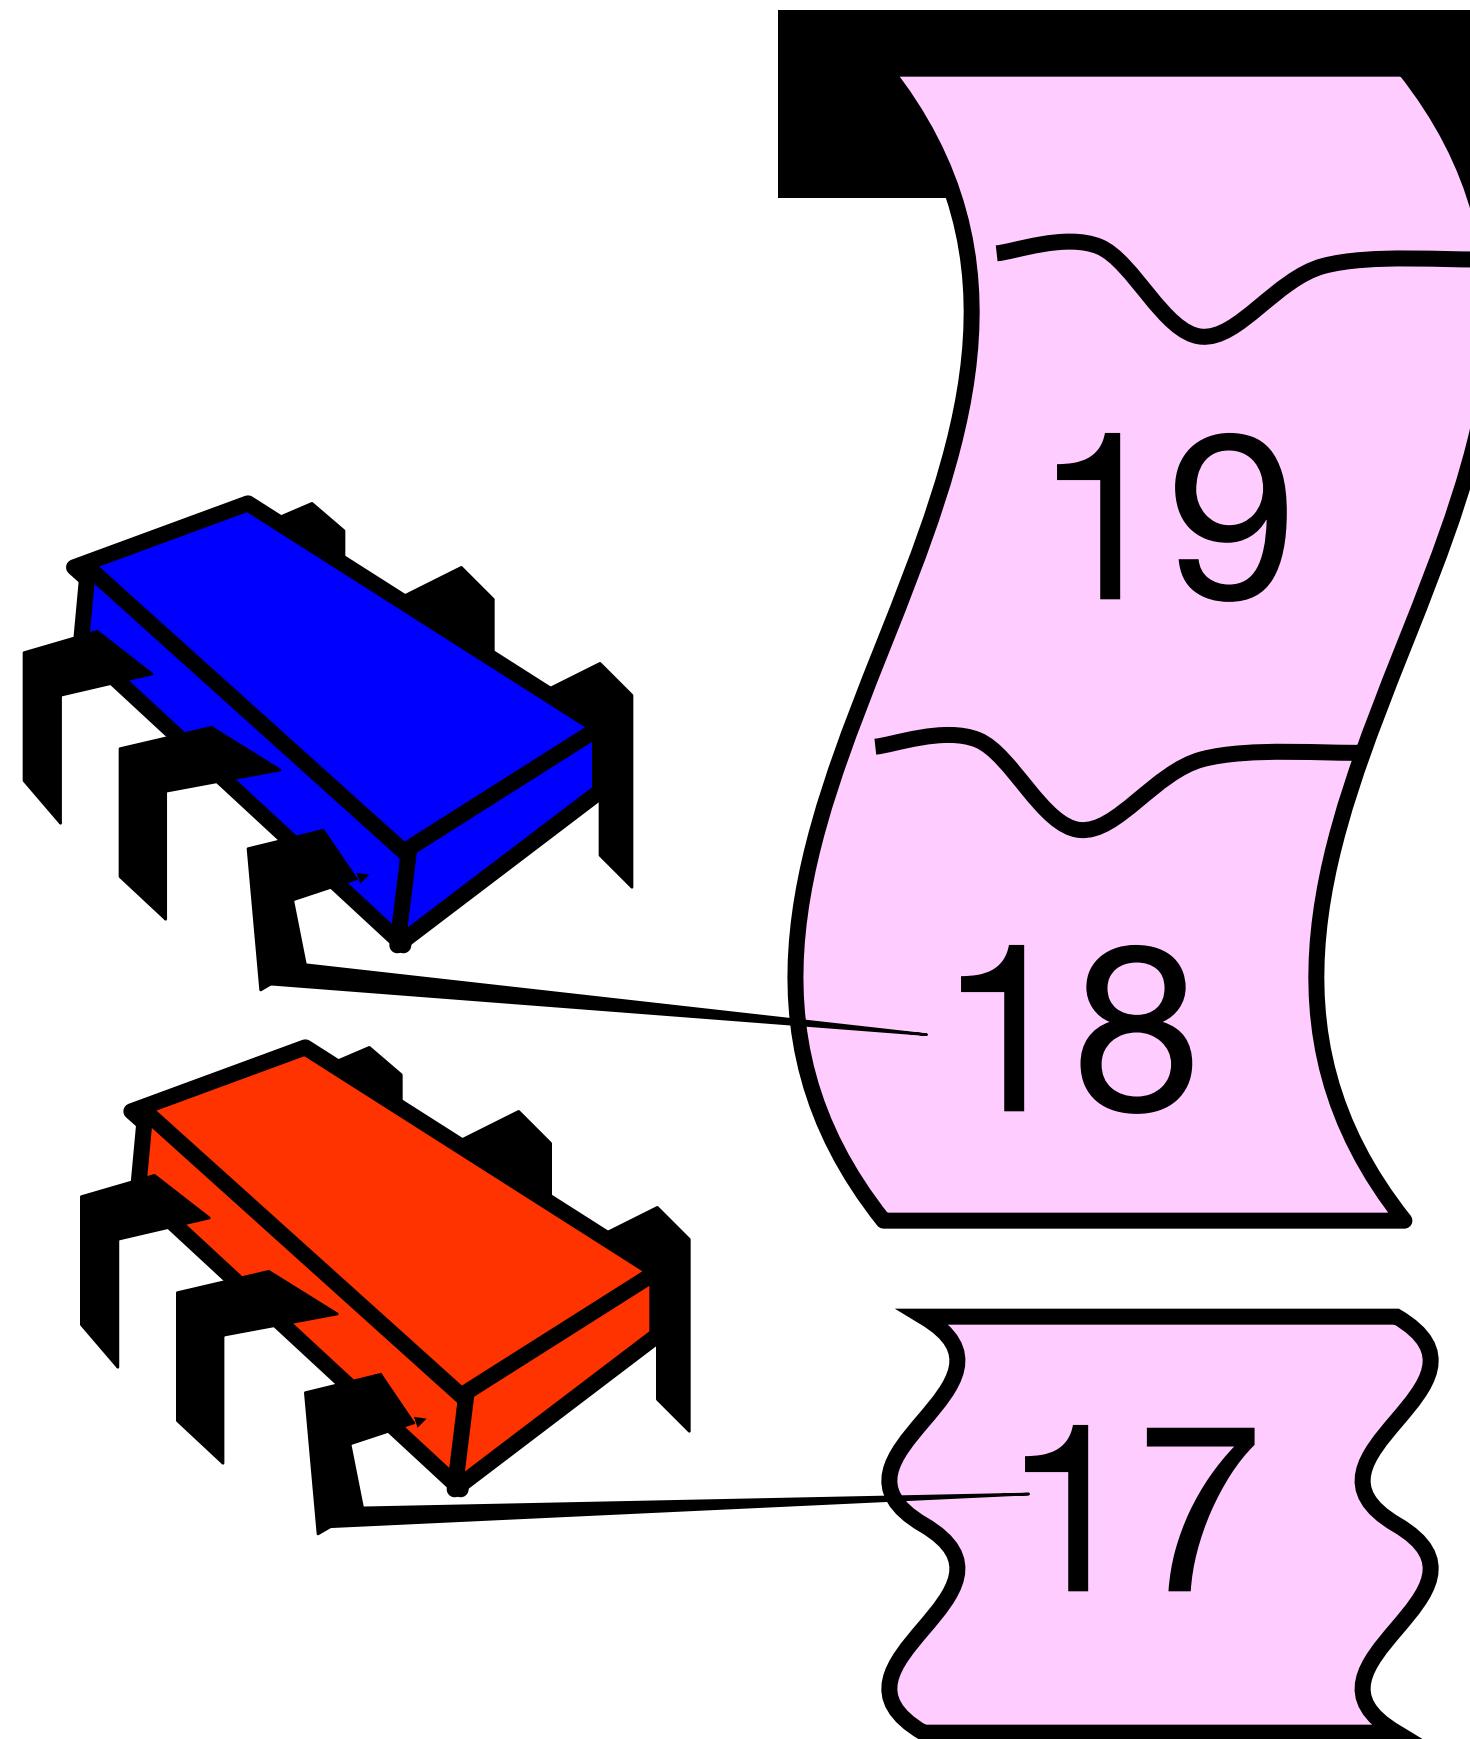
```
let is_prime n =
  if n <= 1 then false
  else if n = 2 then true
  else if n mod 2 = 0 then false
  else
    let rec check_divisor d =
      (* If n has a divisor d > sqrt(n), then it must
         also have a corresponding divisor n/d < sqrt(n). *)
      if d * d > n then true
      else if n mod d = 0 then false
      else check_divisor (d + 2)
    in
    check_divisor 3

let print_primes_in_range start_range end_range =
  for i = start_range to end_range do
    if is_prime i then
      Printf.printf "%d\n" i
  done
```

- Issues
  - Higher ranges have fewer primes
  - Yet larger numbers harder to test
  - Thread workloads
  - Uneven and hard to predict
- Need **dynamic** load balancing

Demo

# Shared counter



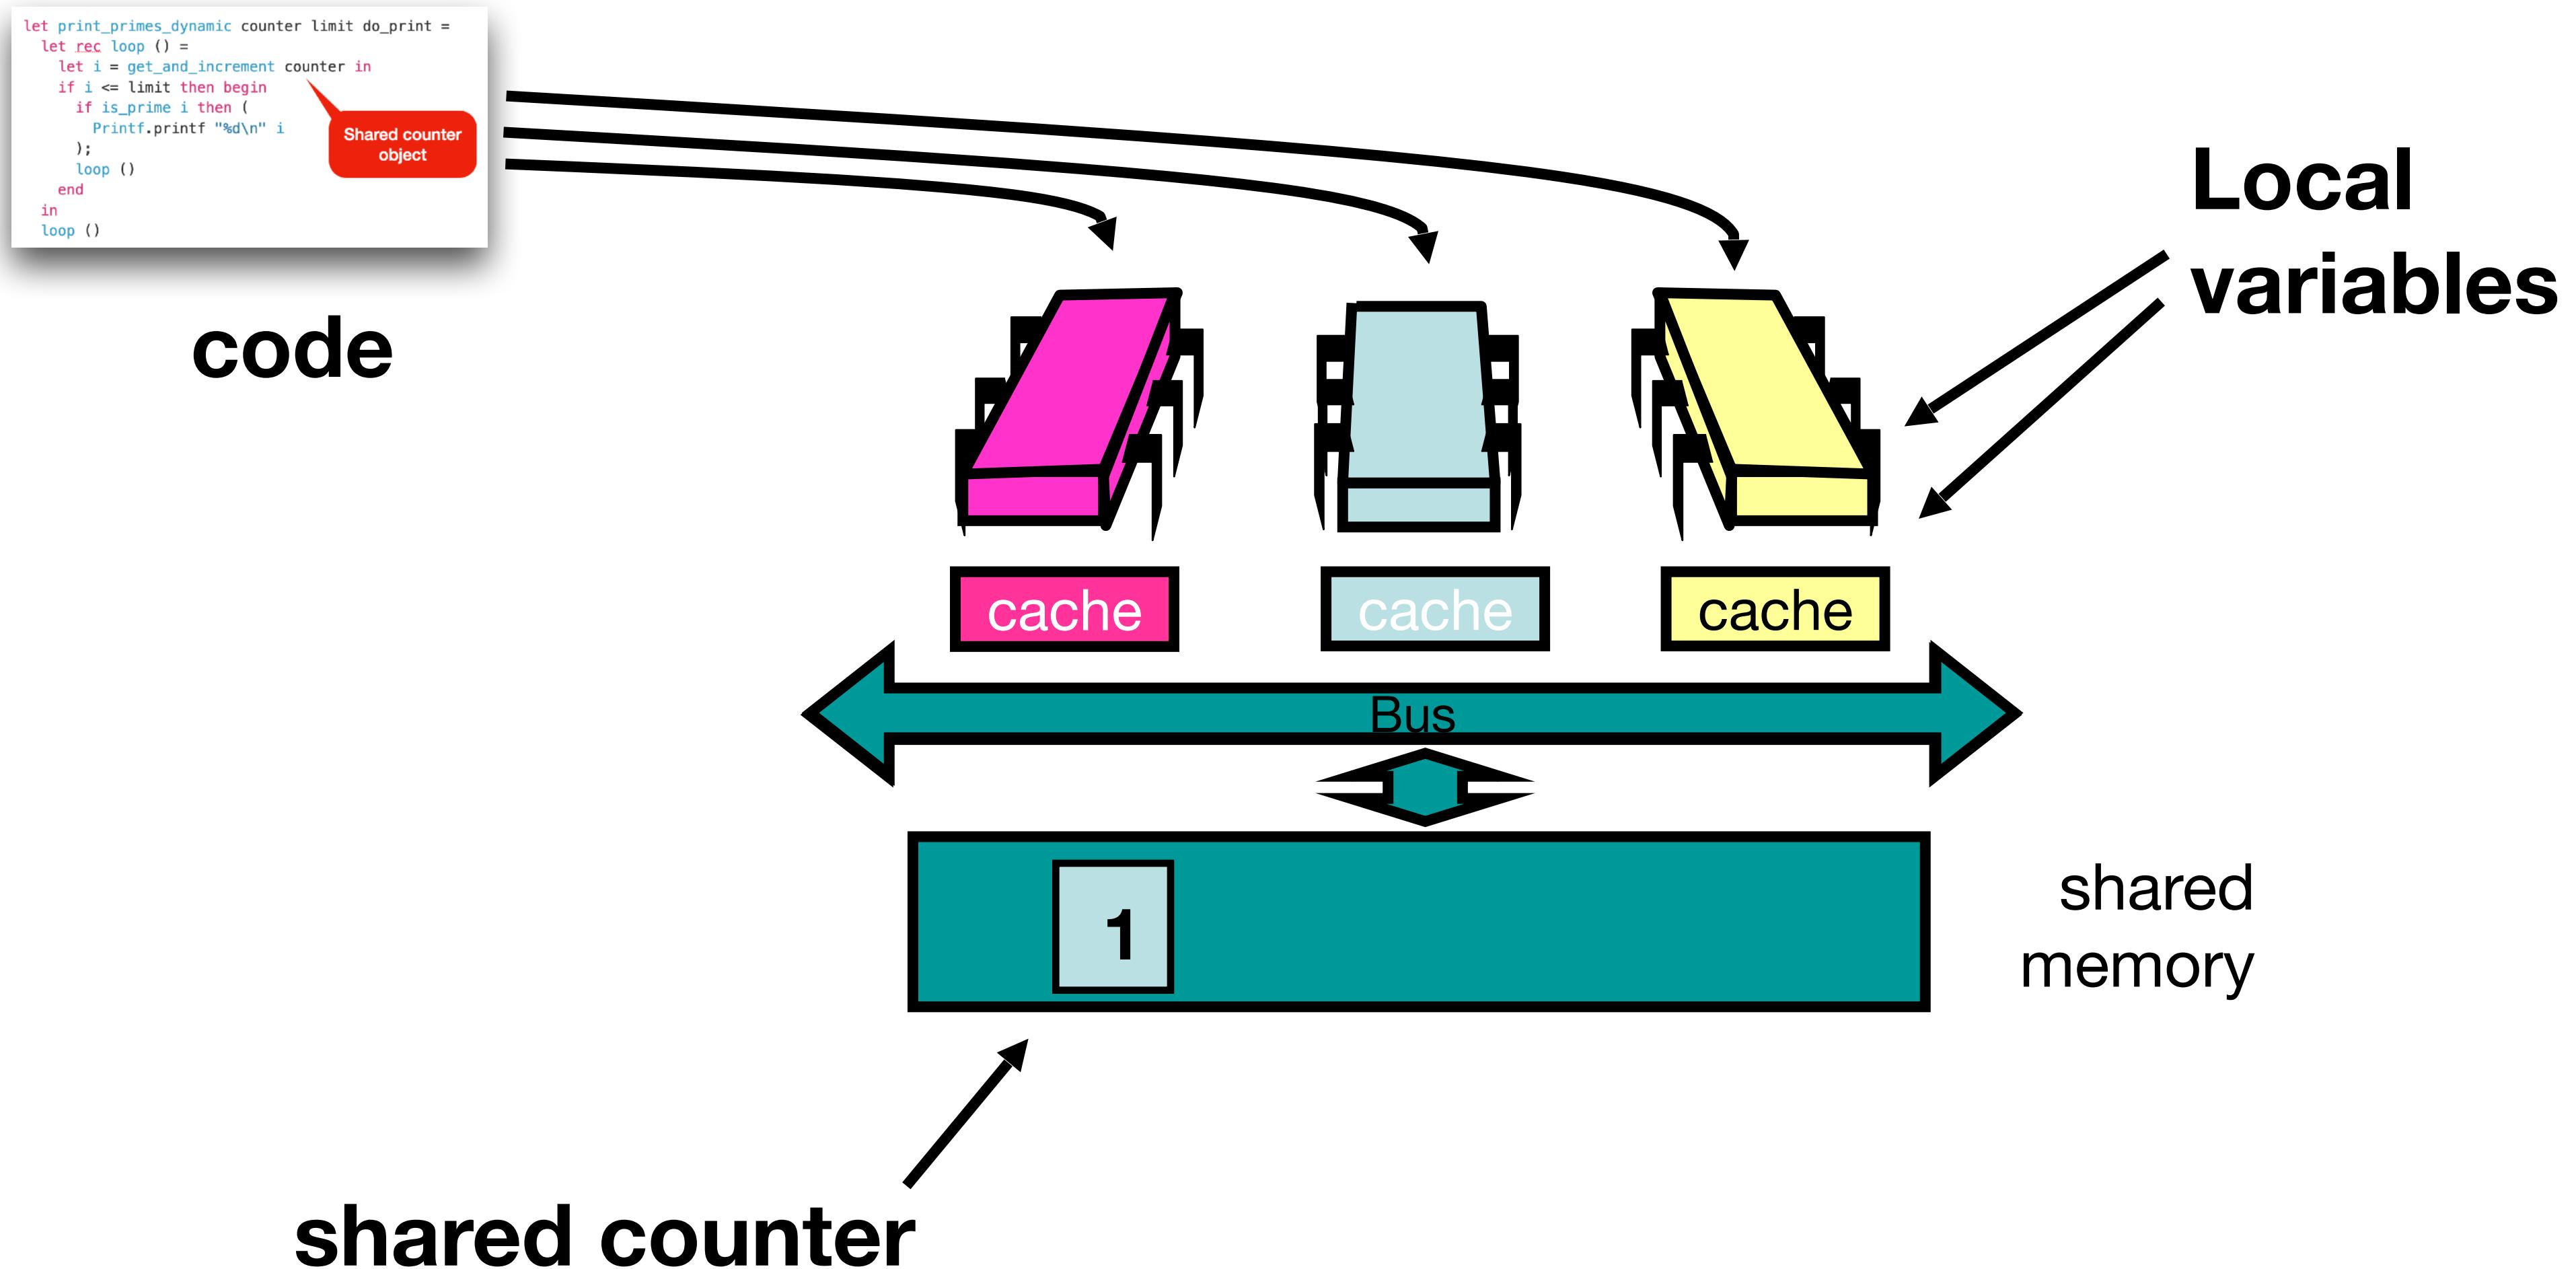
each thread  
takes a number

# Procedure for thread i

```
let print_primes_dynamic counter limit do_print =
  let rec loop () =
    let i = get_and_increment counter in
    if i <= limit then begin
      if is_prime i then (
        Printf.printf "%d\n" i
      );
      loop ()
    end
  in
  loop ()
```

Shared counter  
object

# Where things reside



# Procedure for thread i

```
let print_primes_dynamic counter limit do_print =
  let rec loop () =
    let i = get_and_increment counter in
    if i <= limit then begin
      if is_prime i then (
        Printf.printf "%d\n" i
      );
      loop ()
    end
  in
  loop ()
```

Stop when  
every value  
examined

Shared counter  
object

# Counter Implementation

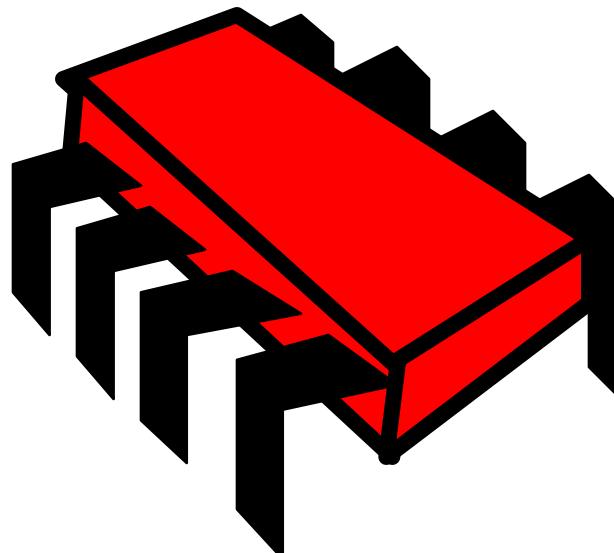
```
(* A counter is just a reference *)
let create_counter initial_value =
  ref initial_value
```

**OK for single thread,  
not for concurrent threads**

```
let print_primes_dynamic counter limit do_print =
  let rec loop () =
    let i = get_and_increment counter in
    if i <= limit then begin
      if is_prime i then (
        Printf.printf "%d\n" i
      );
      loop ()
    end
  in
  loop ()
```

# Not so good ...

Counter... **1**



read  
1

**2**

⋮

read  
2

**3**

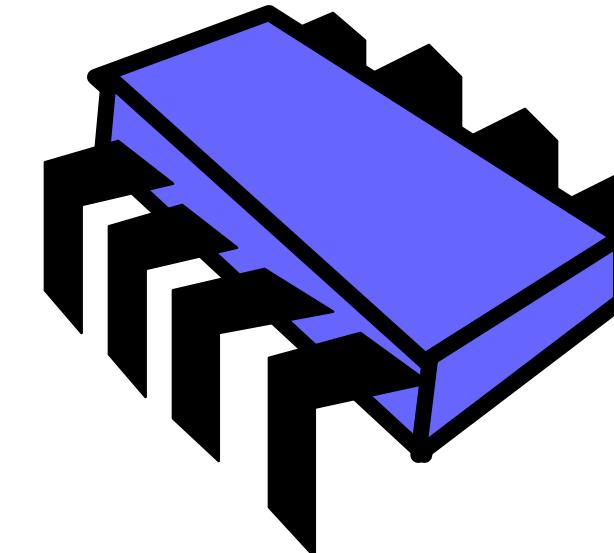
⋮

write  
3

**2**

⋮

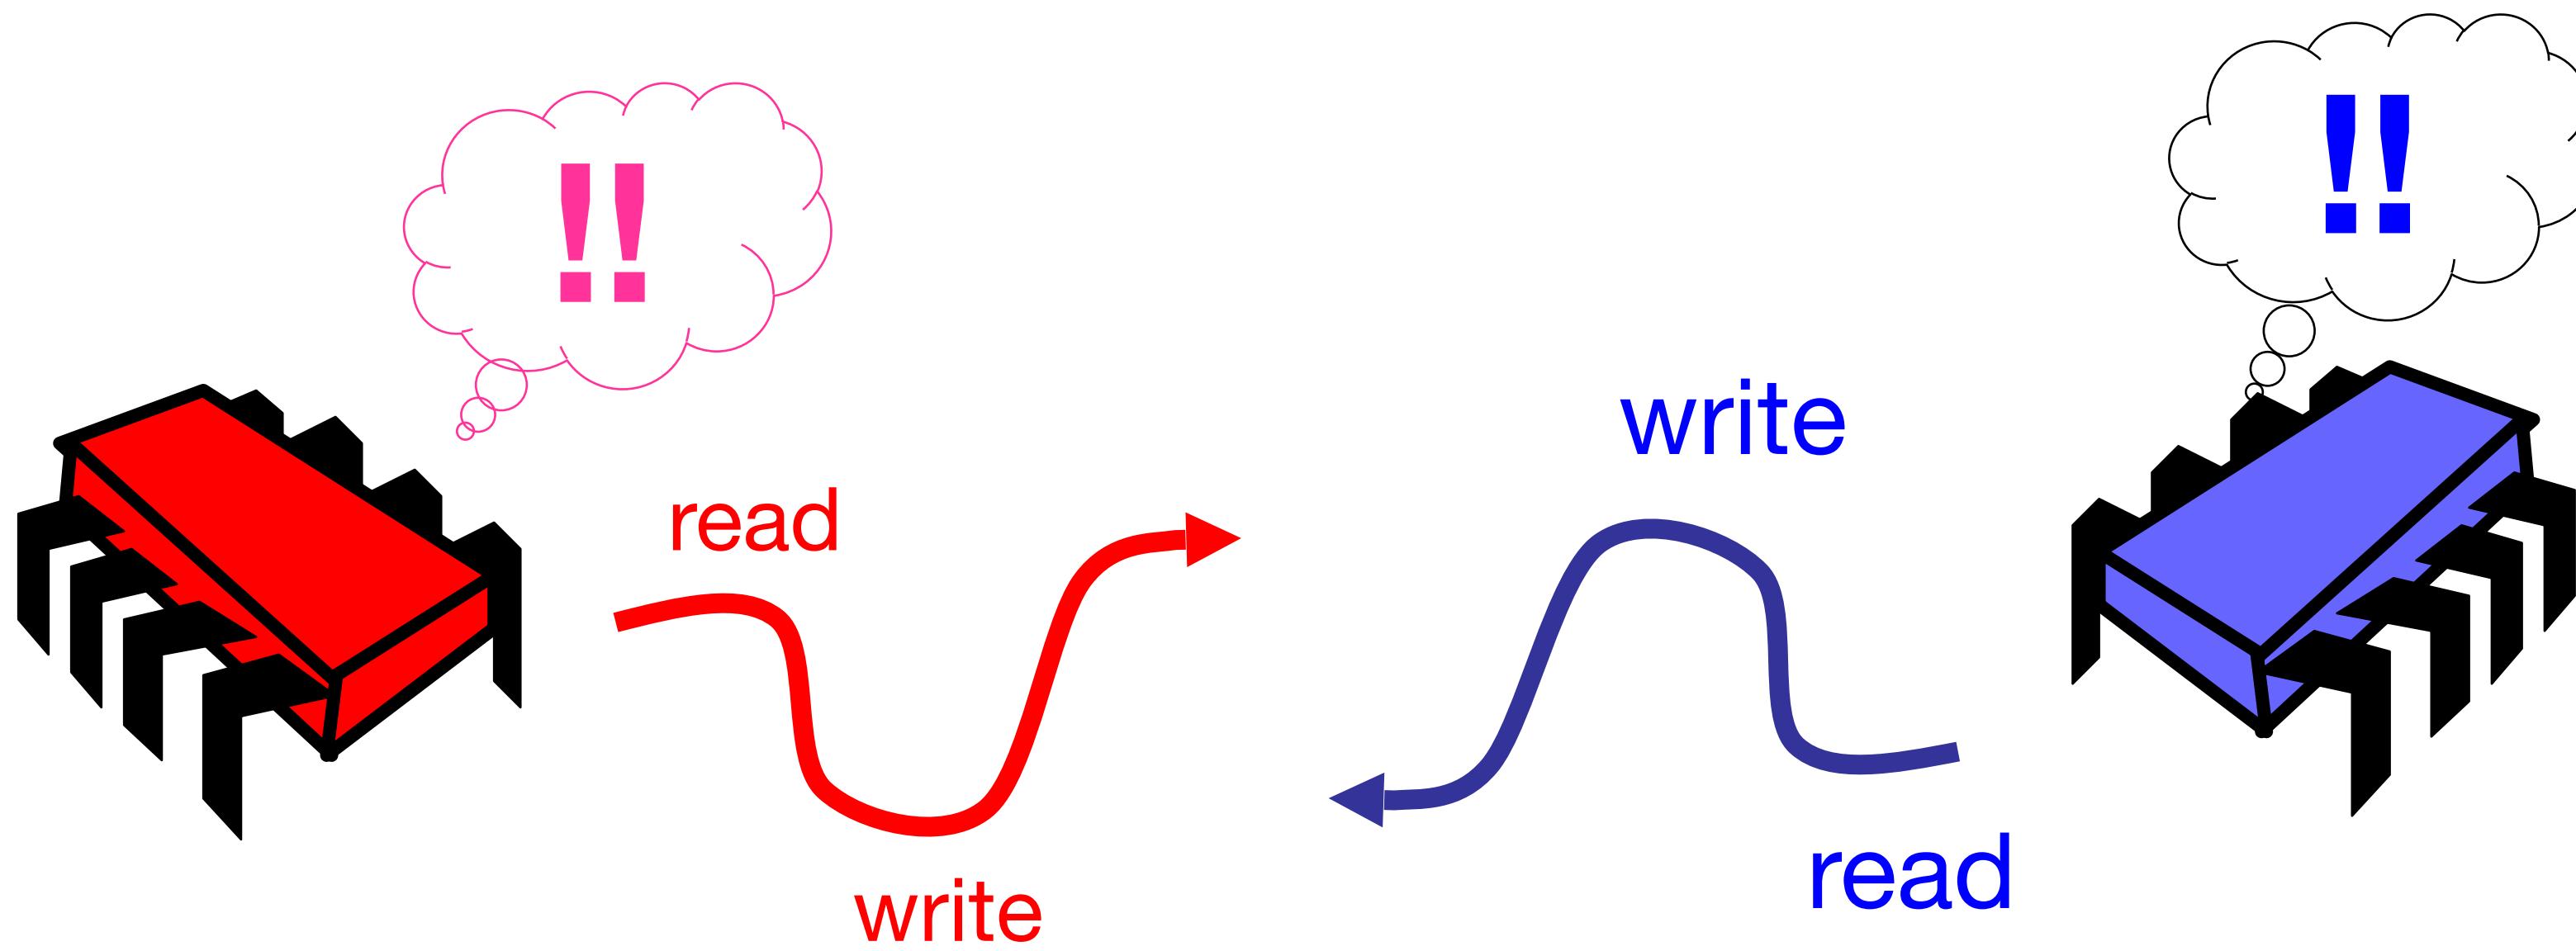
write  
2



read  
1

time

# Is this problem inherent?



If we could only glue reads and writes  
together...

# Challenge

```
(* A counter is just a reference *)
let create_counter initial_value =
  ref initial_value

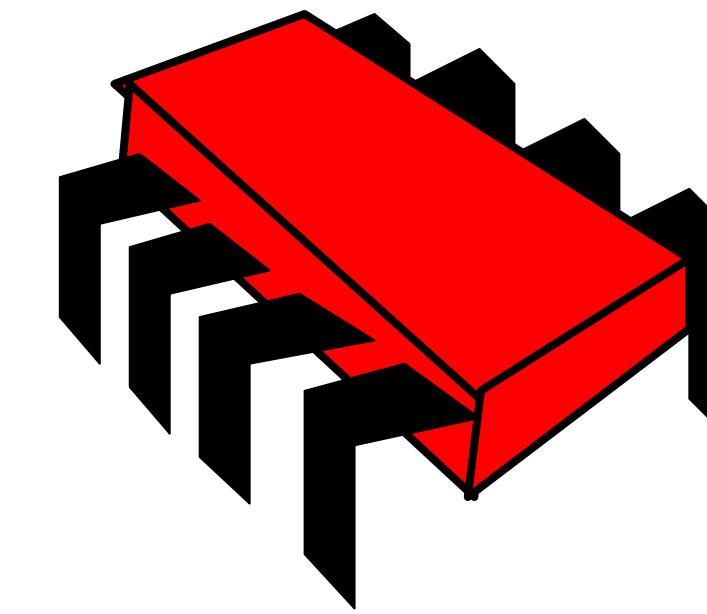
let get_and_increment counter =
  let v = !counter in
  counter := v + 1;
  v
```

Make this atomic (indivisible)

# Hardware solution

```
(* A counter is just a reference *)
let create_counter initial_value =
  ref initial_value

let get_and_increment counter =
  let v = !counter in
  counter := v + 1;
  v
```

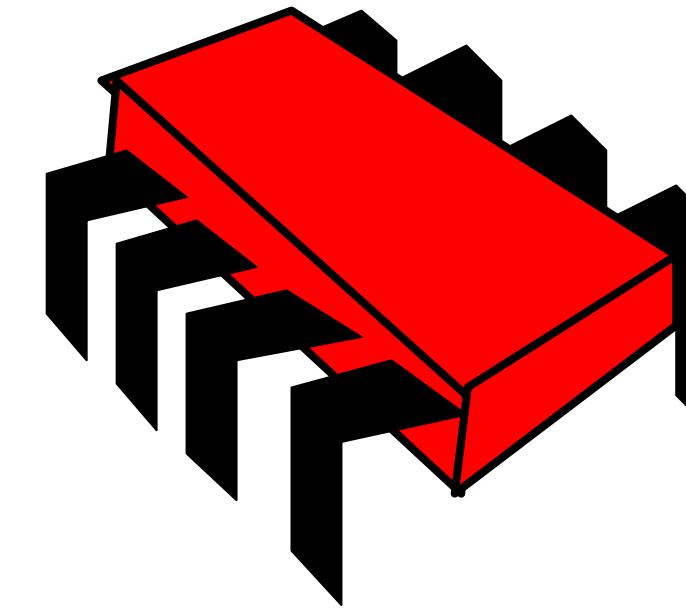


ReadModifyWrite instruction

# OCaml 5 solution

```
(* Thread-safe counter using atomic operations *)
let create_counter initial_value =
    Atomic.make initial_value

let get_and_increment counter =
    Atomic.fetch_and_add counter 1
```



ReadModifyWrite instruction

Demo

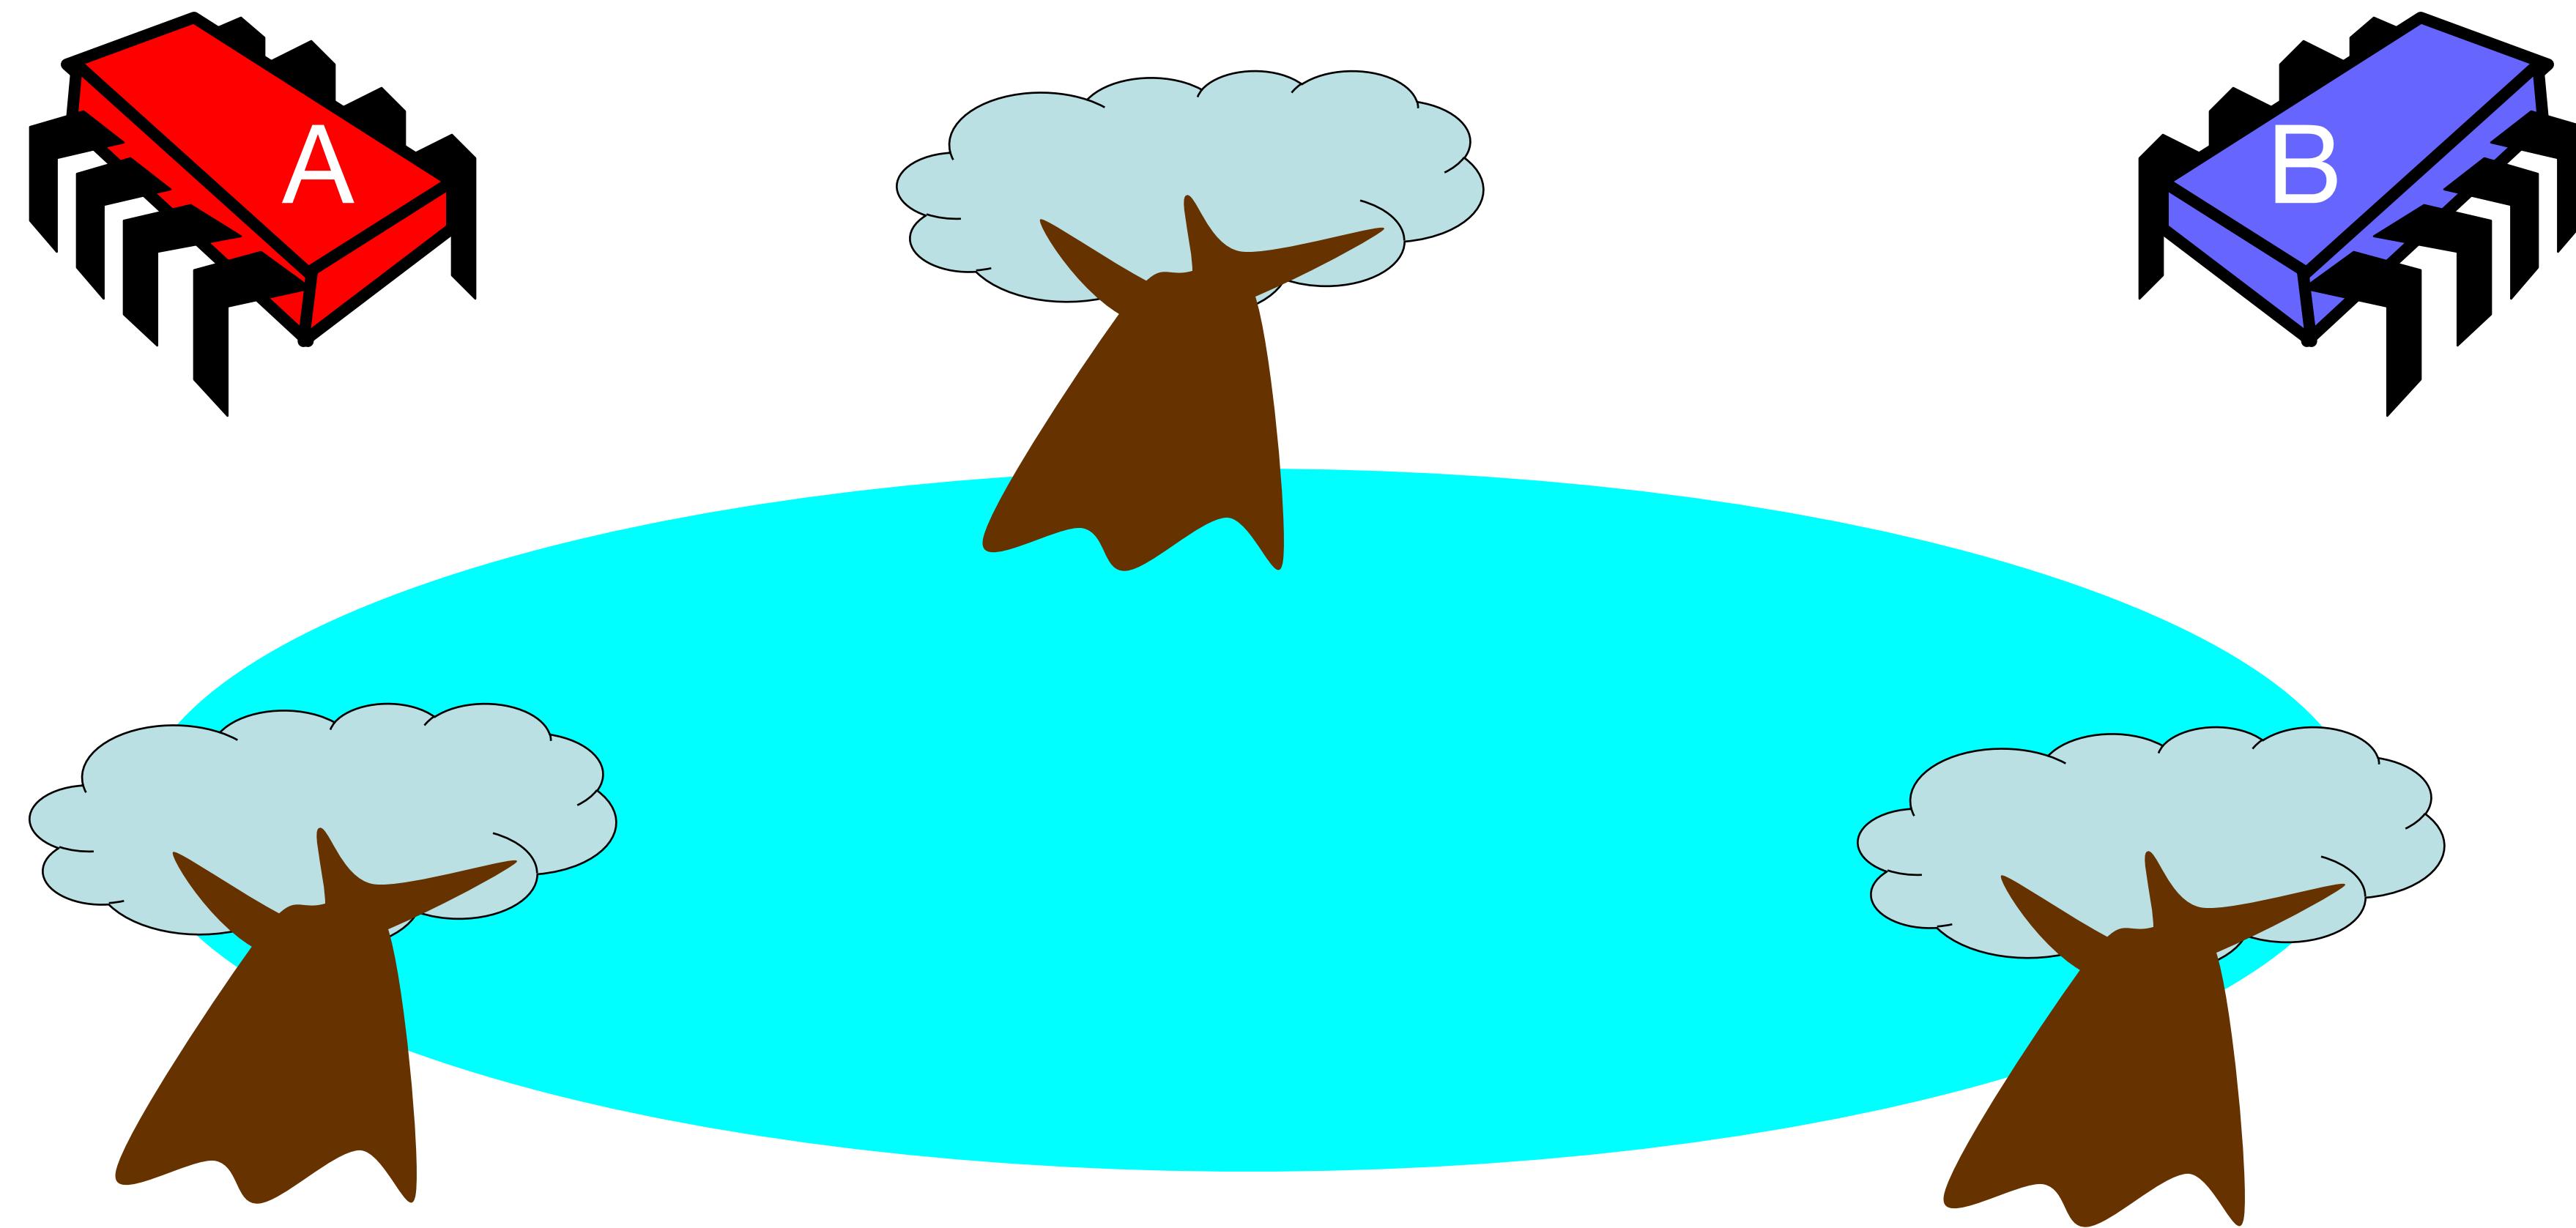
# OCaml 5 solution (2)

```
let create_counter initial_value =
  (ref initial_value, Mutex.create ())

let get_and_increment (counter_ref, mutex) =
  Mutex.lock mutex;
  let value = !counter_ref in
  counter_ref := value + 1;
  Mutex.unlock mutex;
  value
```

Mutual Exclusion

# Mutual Exclusion (or Alice and Bob share a pond)



# Alice has a pet



# Bob has a pet



# The Problem

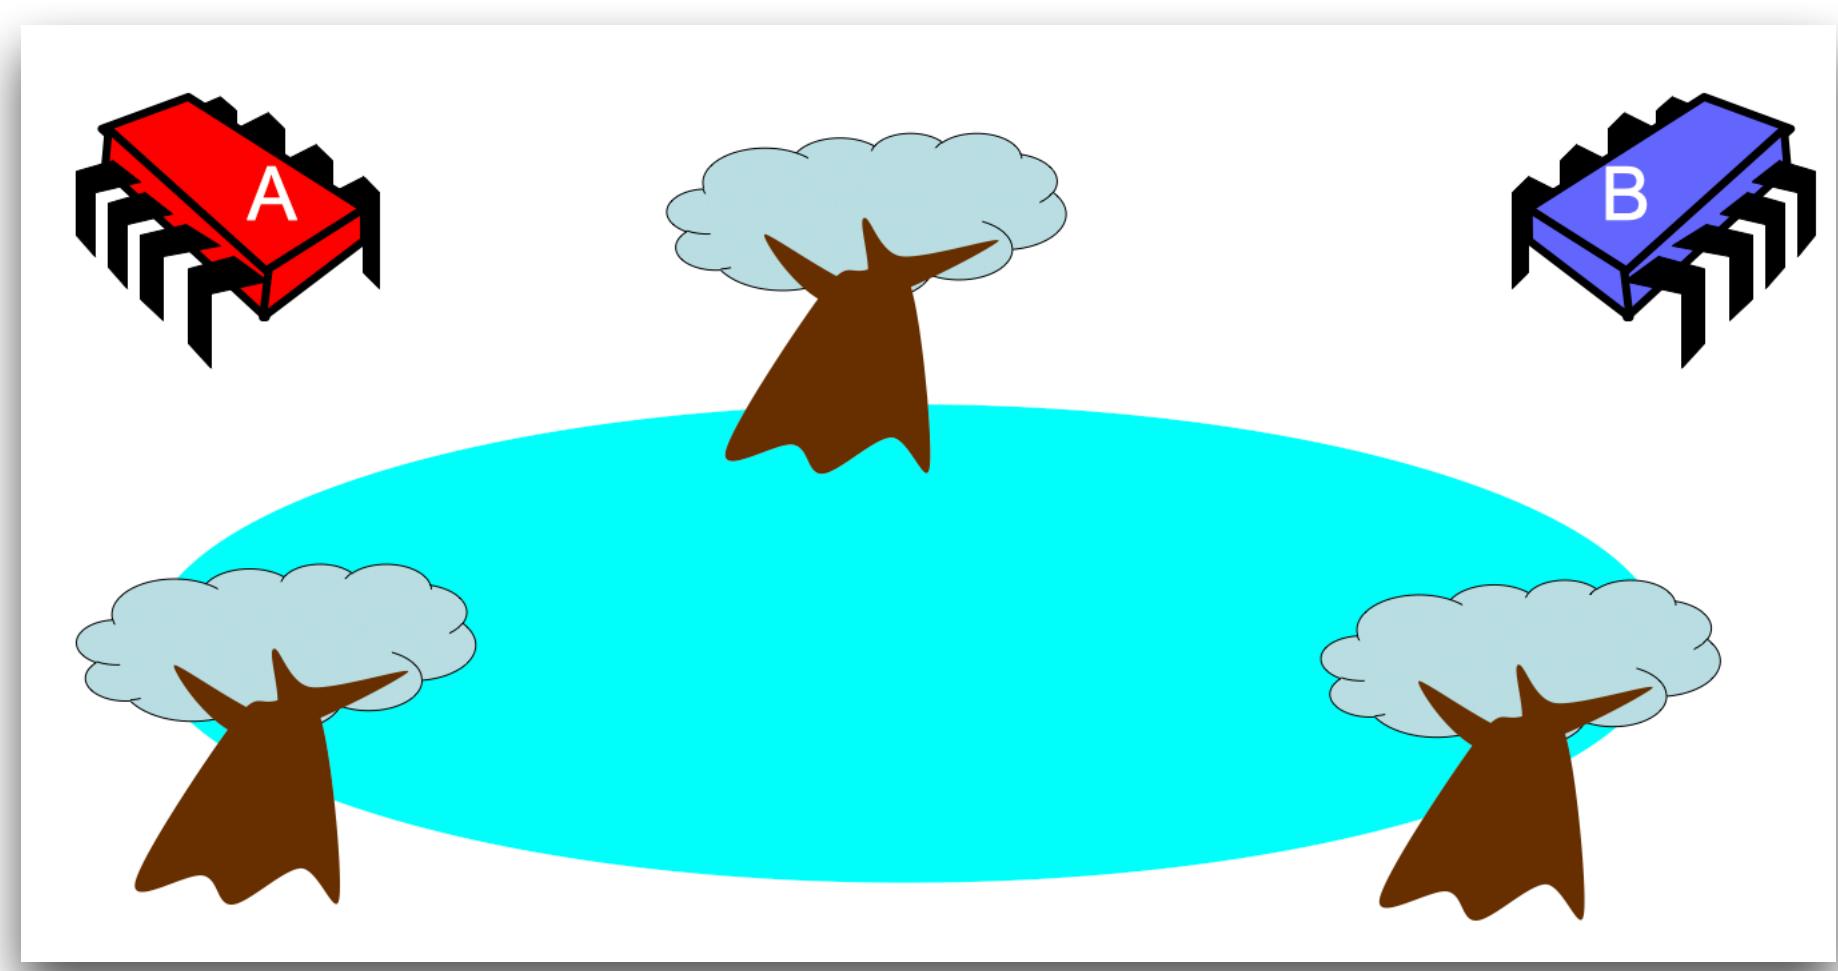


# Formalising the Problem

- Two types of formal properties in asynchronous computation:
  - **Safety** Properties – Nothing bad happens ever
  - **Liveness** Properties – Something good happens eventually
- **Mutual Exclusion**
  - Both pets never in pond simultaneously
  - This is a **safety** property
- **No Deadlock**
  - if only one wants in, it gets in
  - if both want in, one gets in.
  - This is a **liveness** property

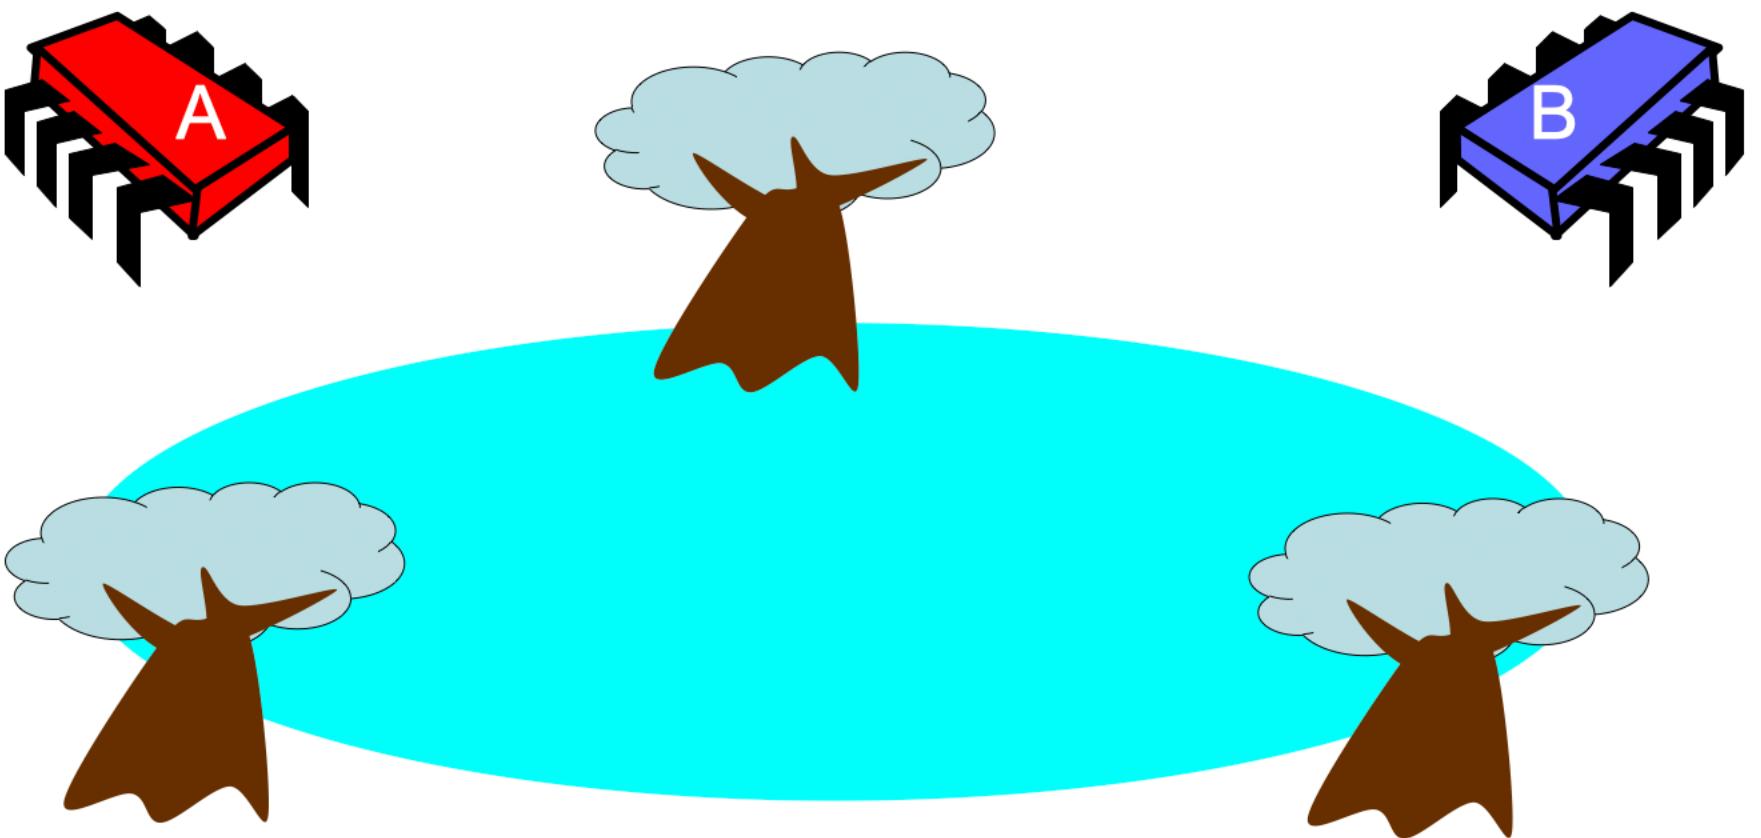
# Simple Protocol

- Idea — Just look at the pond
- Gotcha
  - Not atomic
  - Trees obscure the view
- Interpretation
  - Threads can't "see" what other threads are doing
  - Explicit communication required for coordination

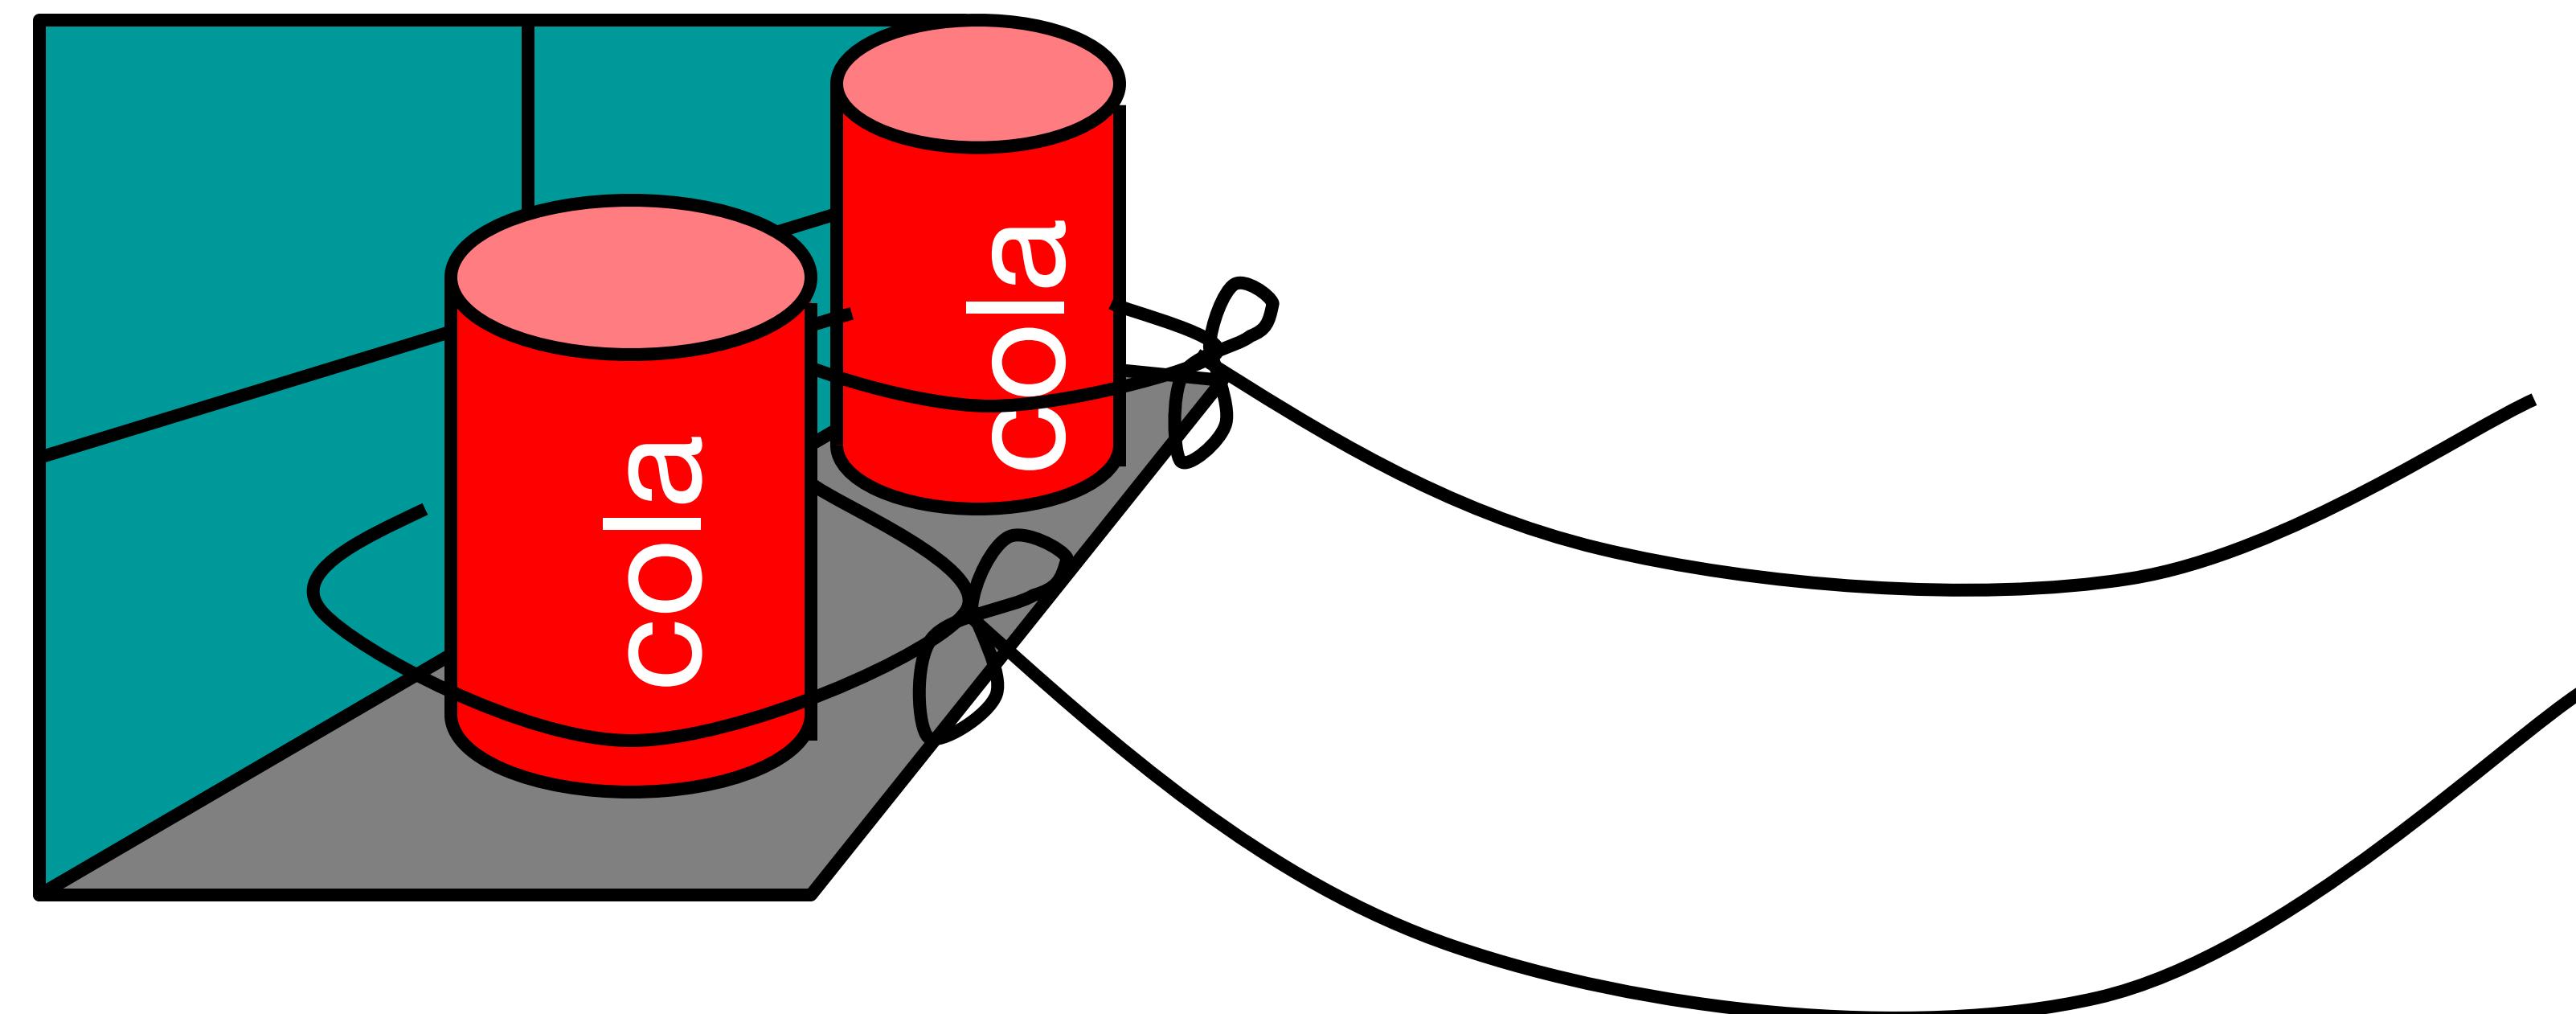


# Mobile phone protocol

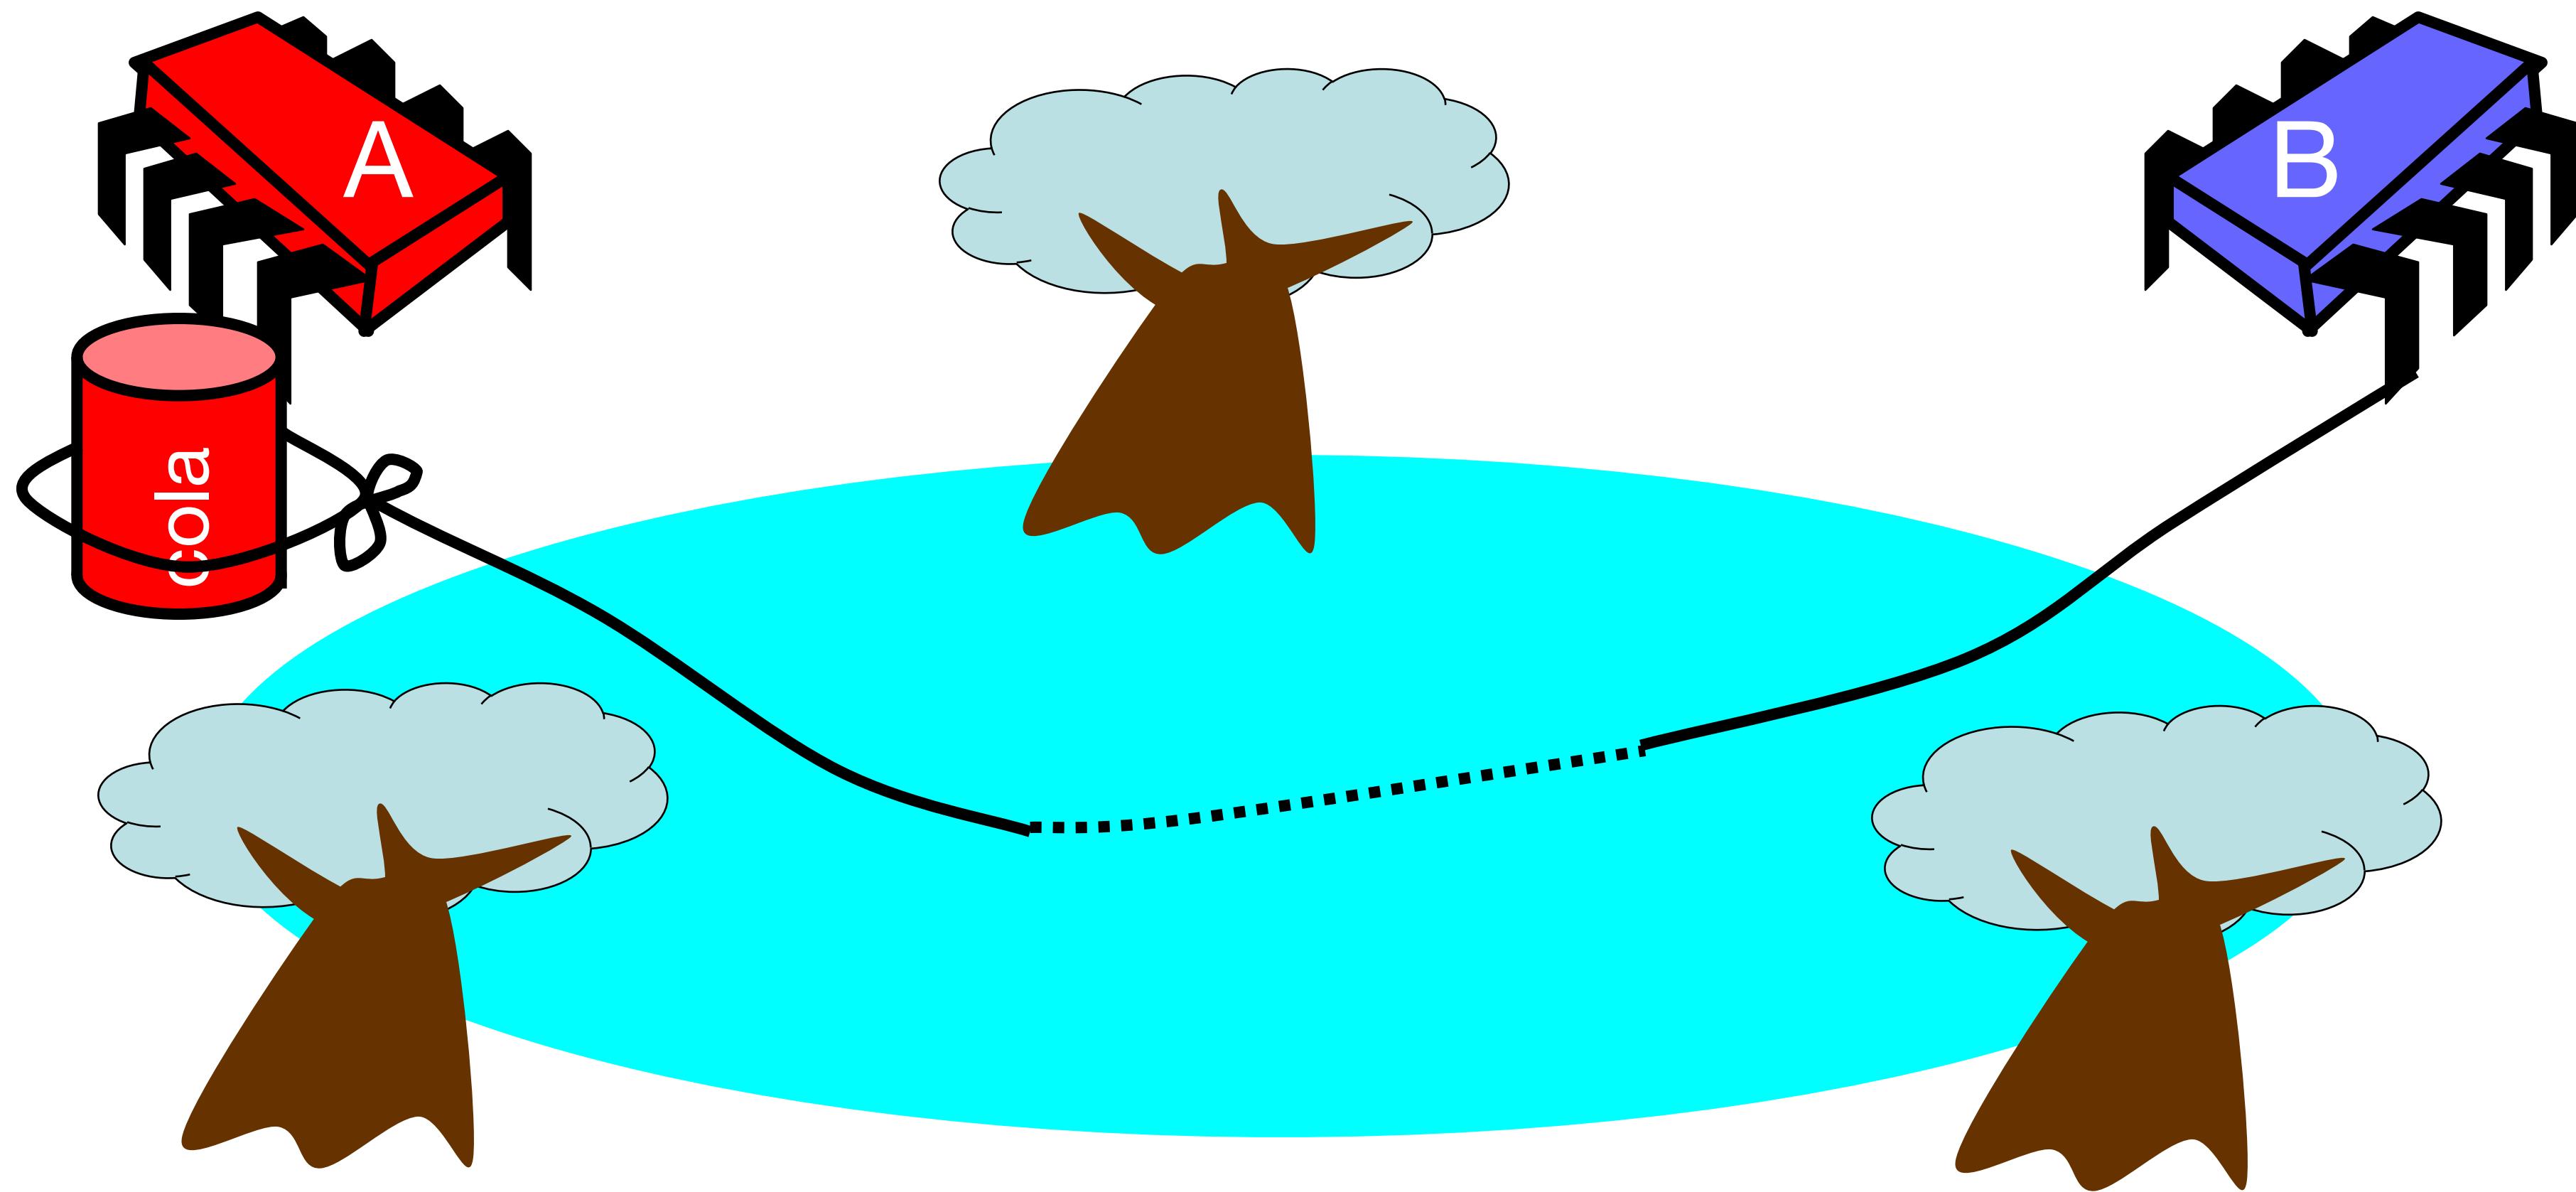
- Idea
  - Bob calls Alice (or vice versa)
- Gotcha
  - Bob takes shower
  - Alice recharges battery
  - Bob out shopping for pet food ...
- Interpretation
  - Message-passing doesn't work
    - Recipient might not be listening or there at all
    - Communication must be Persistent (like writing) and not Transient (like speaking)



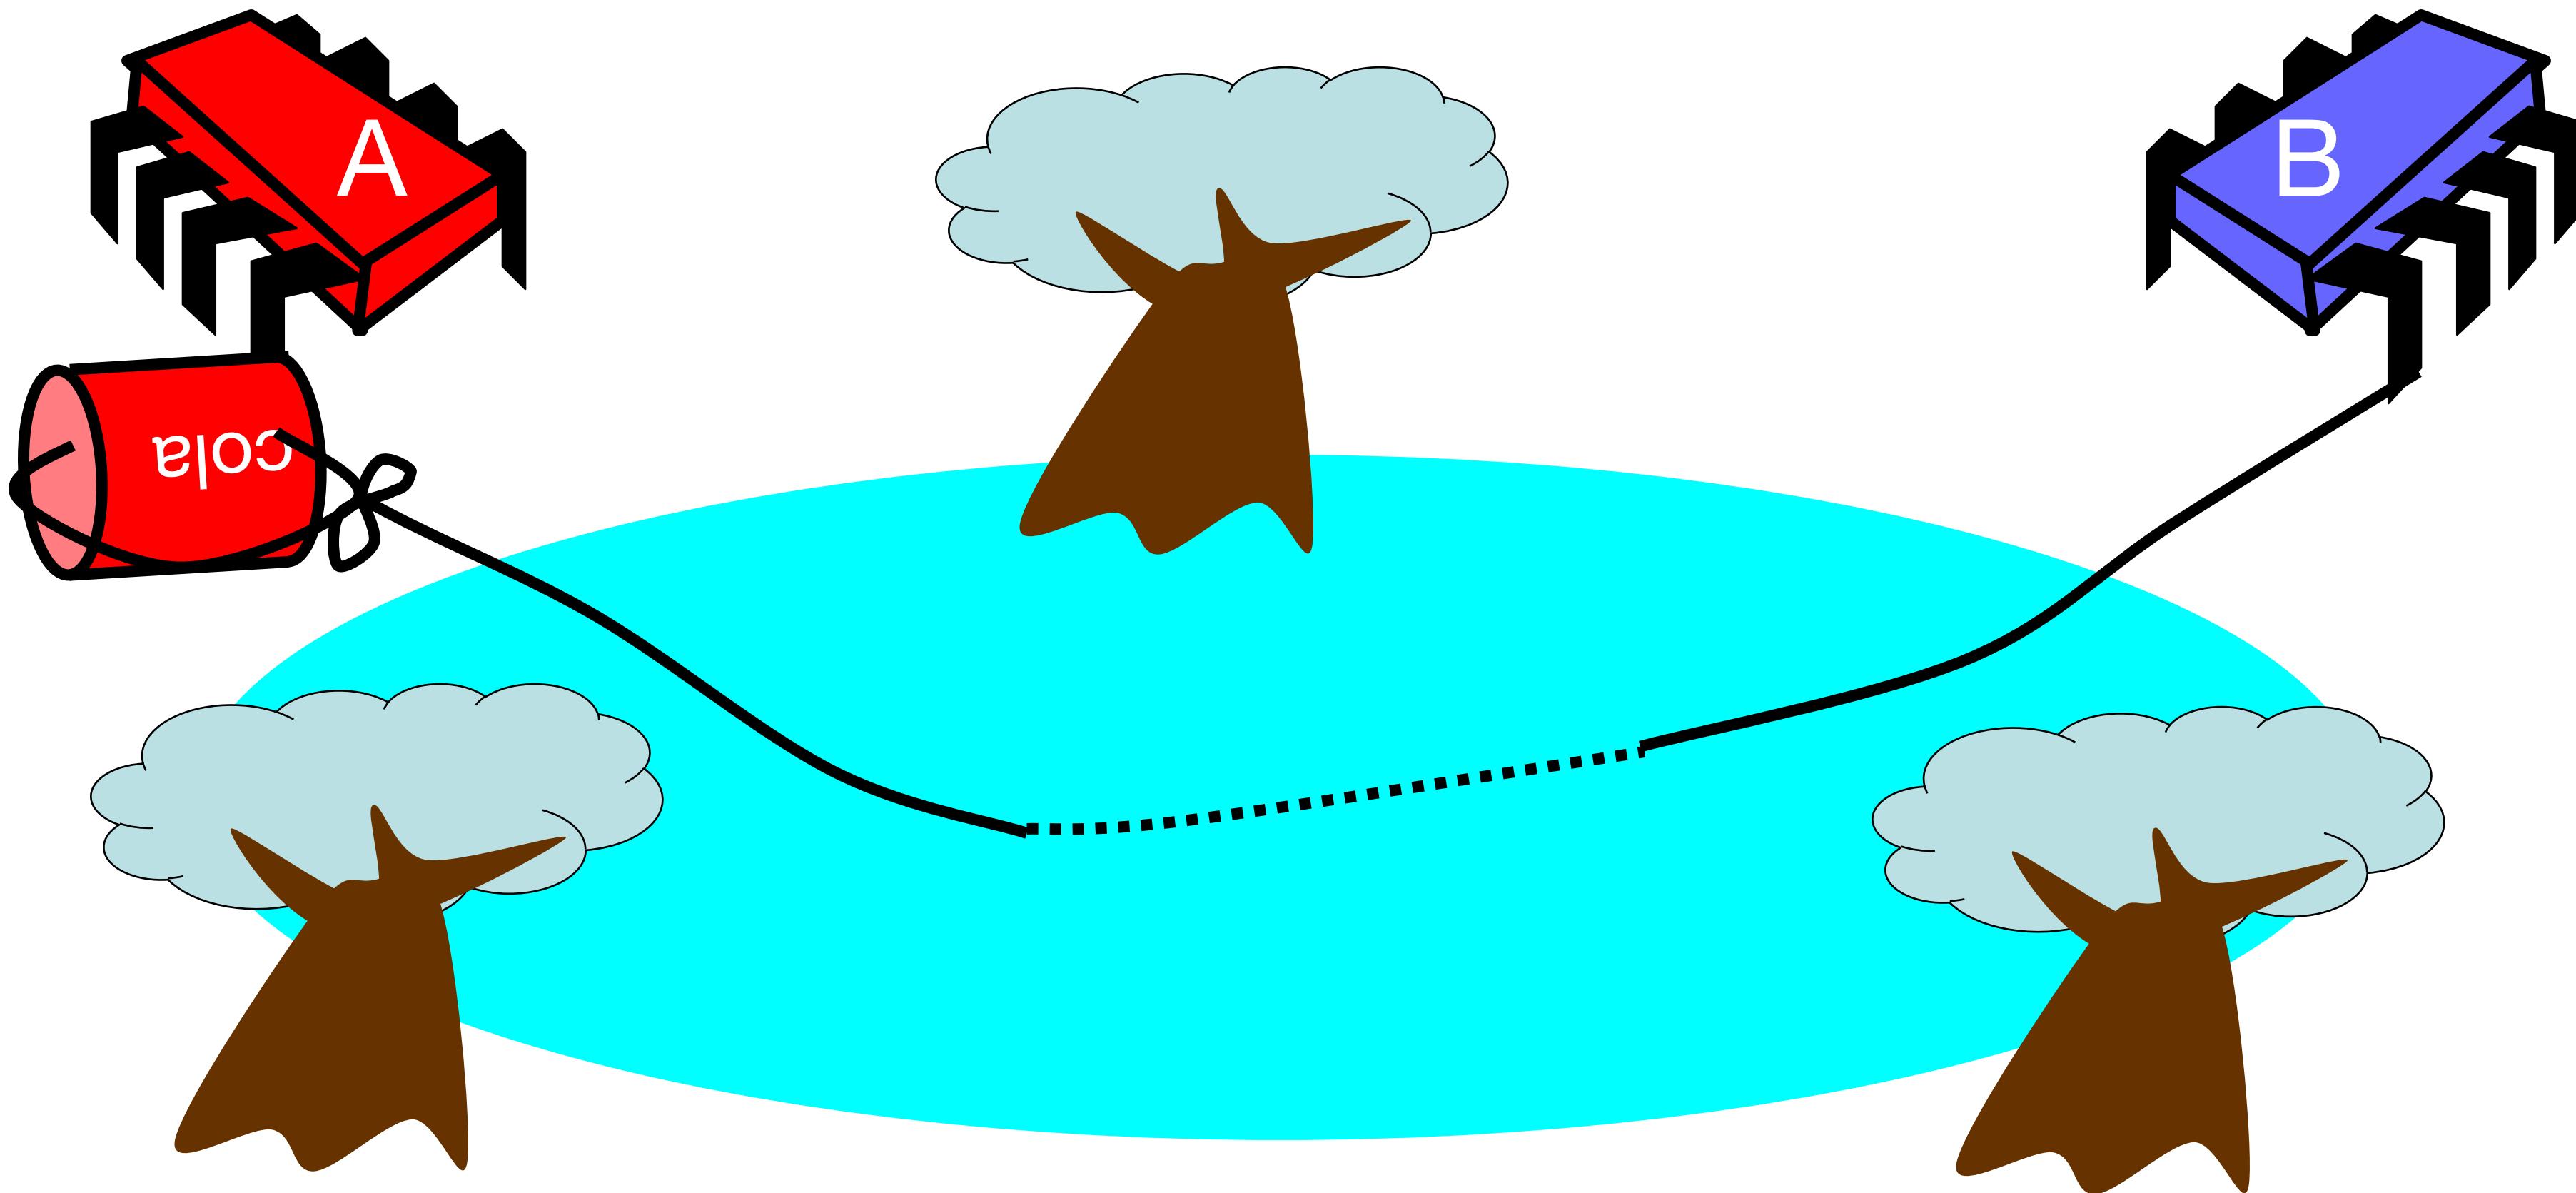
# Can Protocol



# Can Protocol – Bob conveys a bit



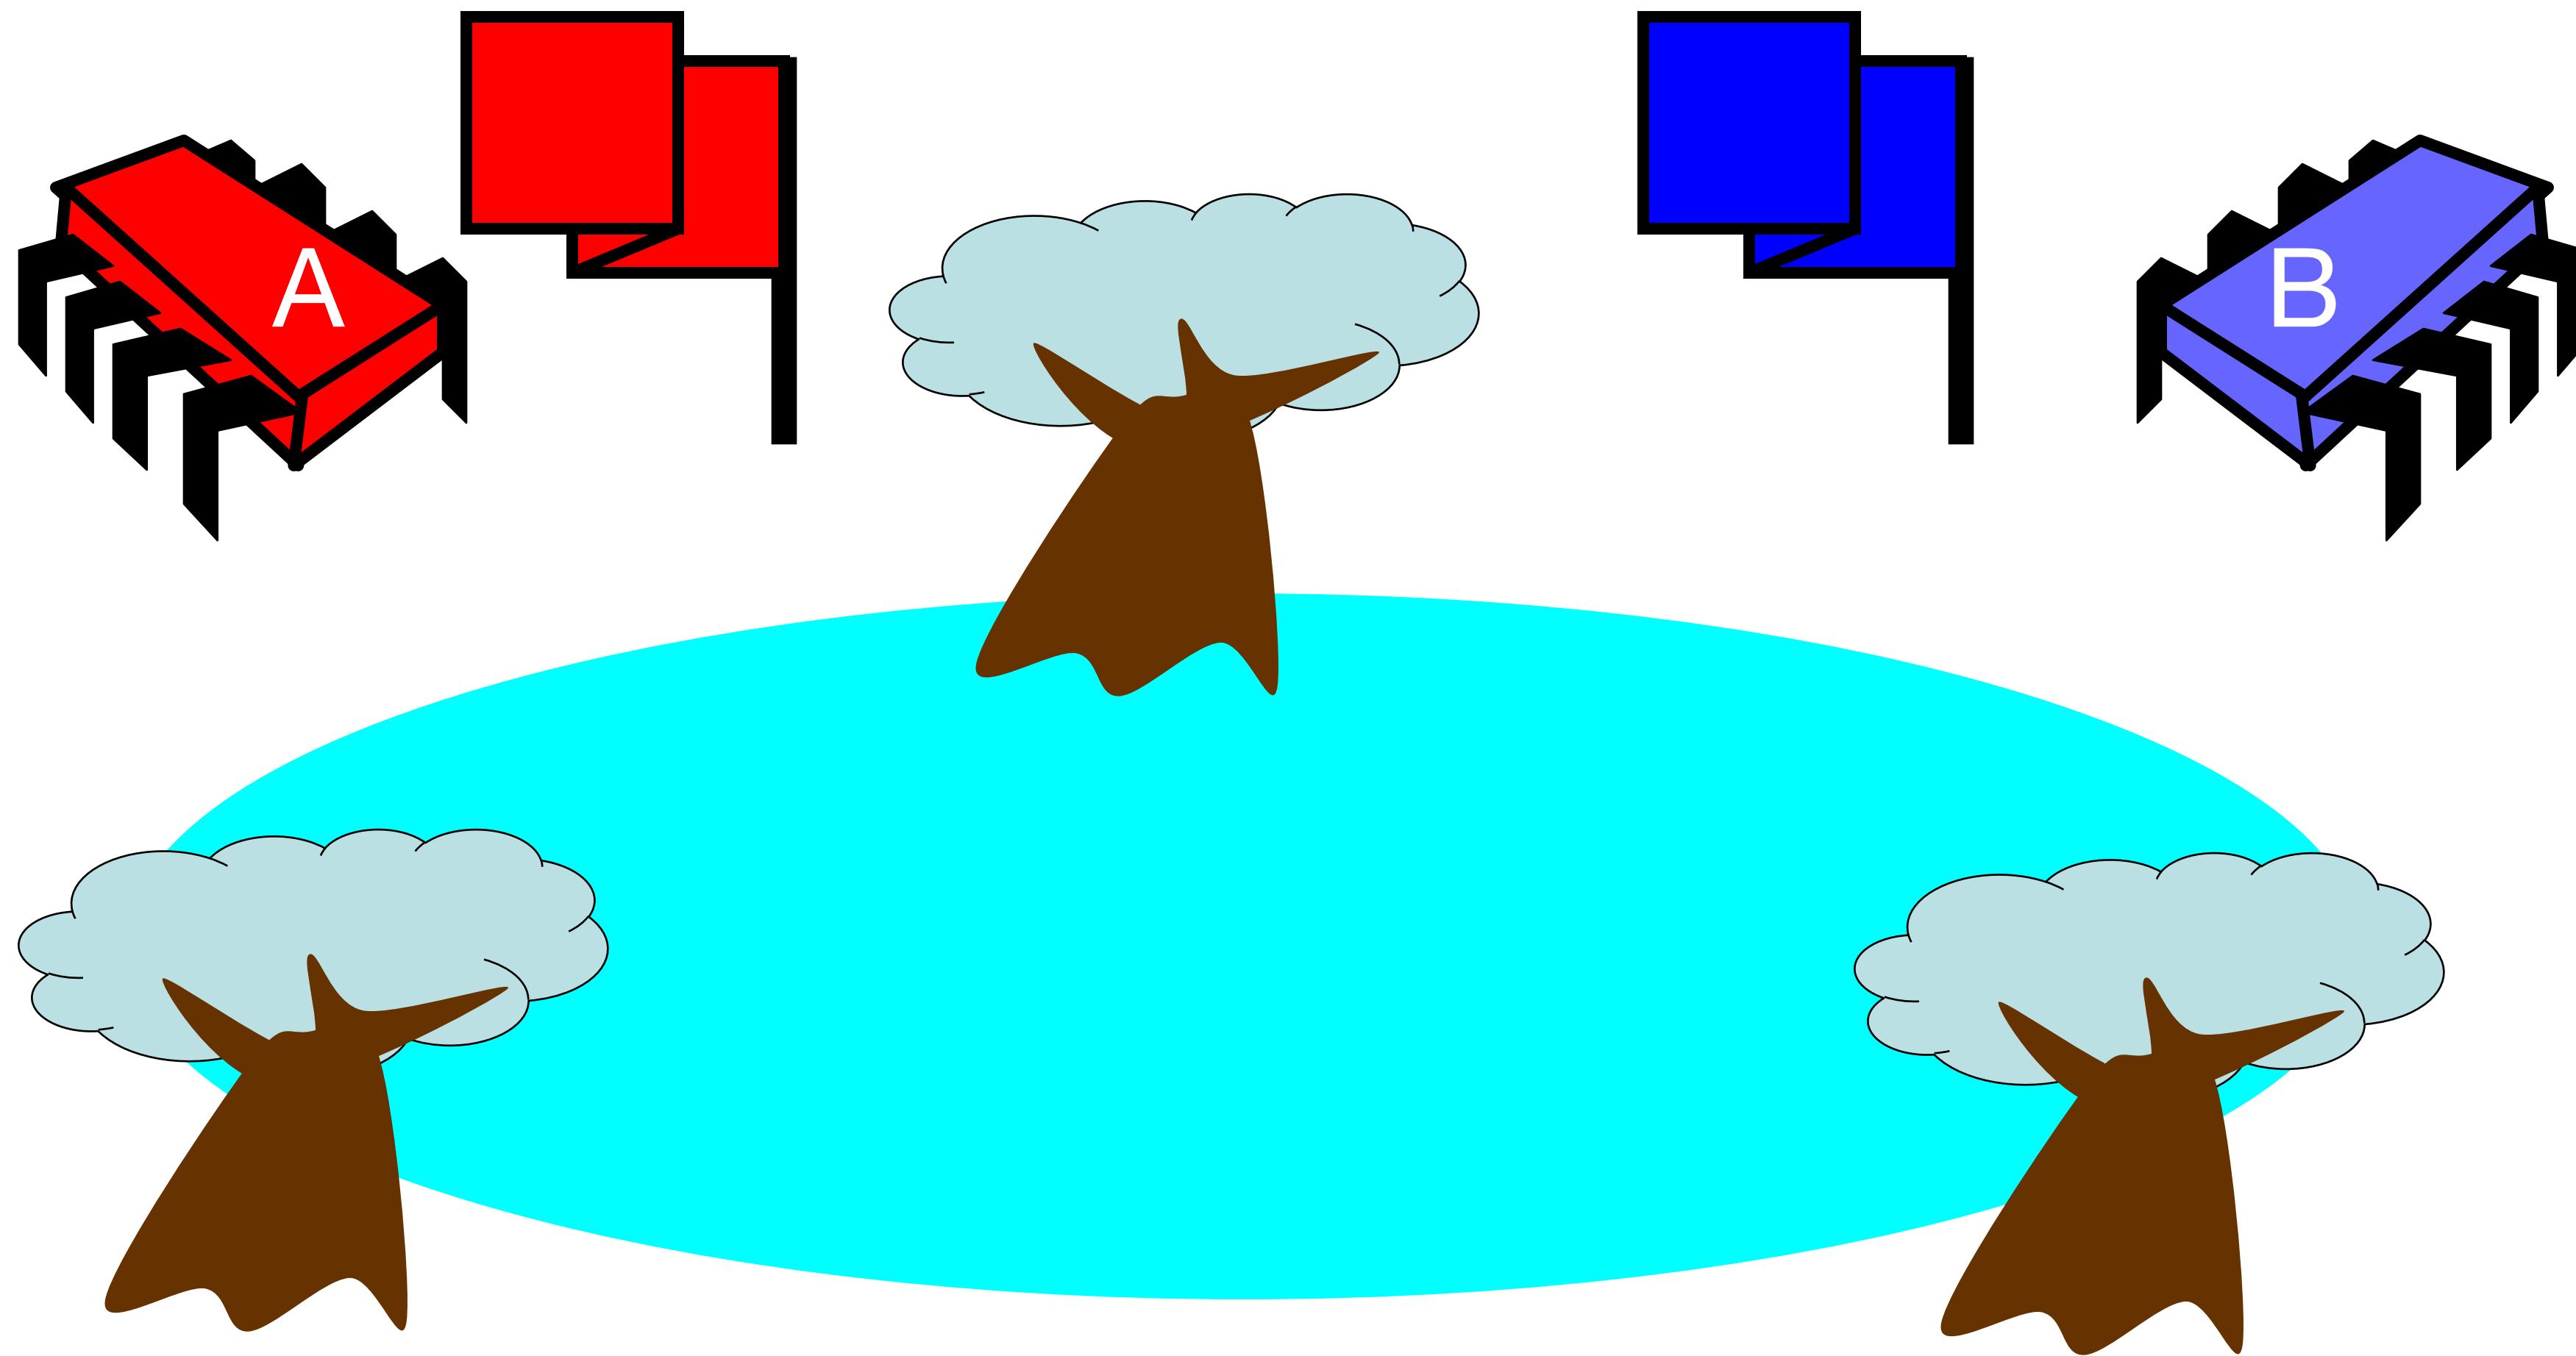
# Can Protocol – Bob conveys a bit



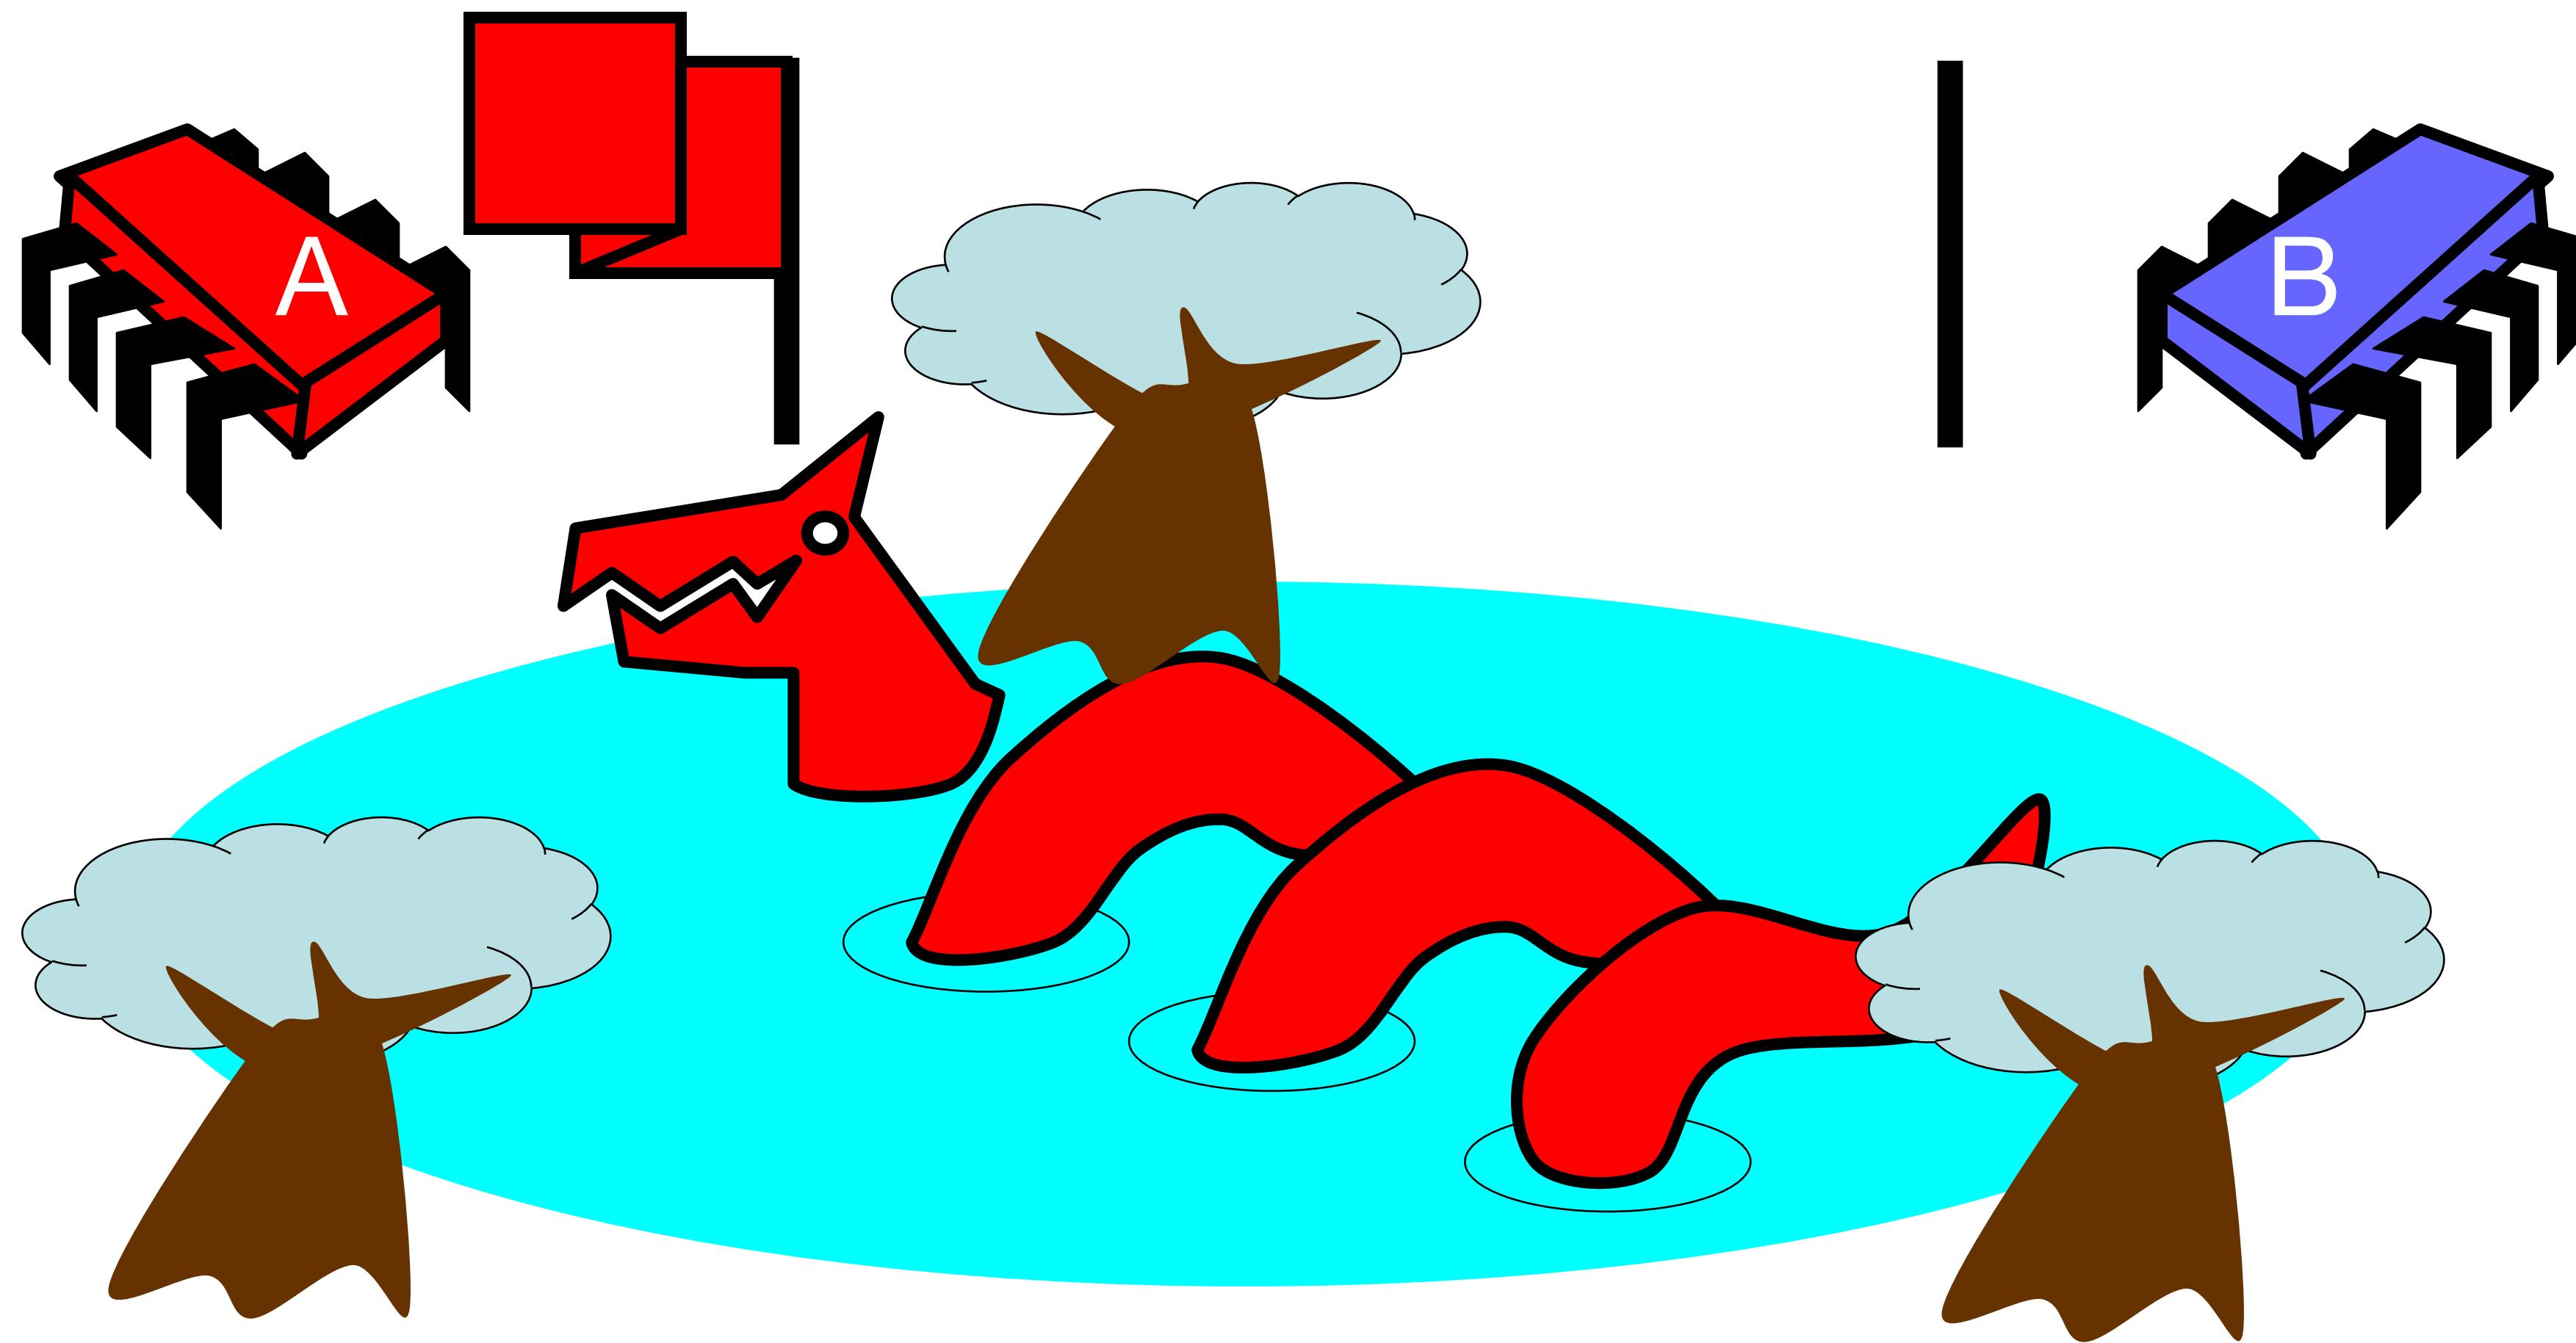
# Can Protocol

- **Idea**
  - Cans on Alice's windowsill
  - Strings lead to Bob's house
  - Bob pulls strings, knocks over cans
    - Alice resets them
- **Gotcha**
  - Cans cannot be reused
  - Bob runs out of cans; Alice is gone on a vacation
- **Interpretation**
  - Cannot solve mutual exclusion with interrupts
  - Sender sets fixed bit in receiver's space
  - Receiver resets bit when ready
  - Requires unbounded number of interrupt bits

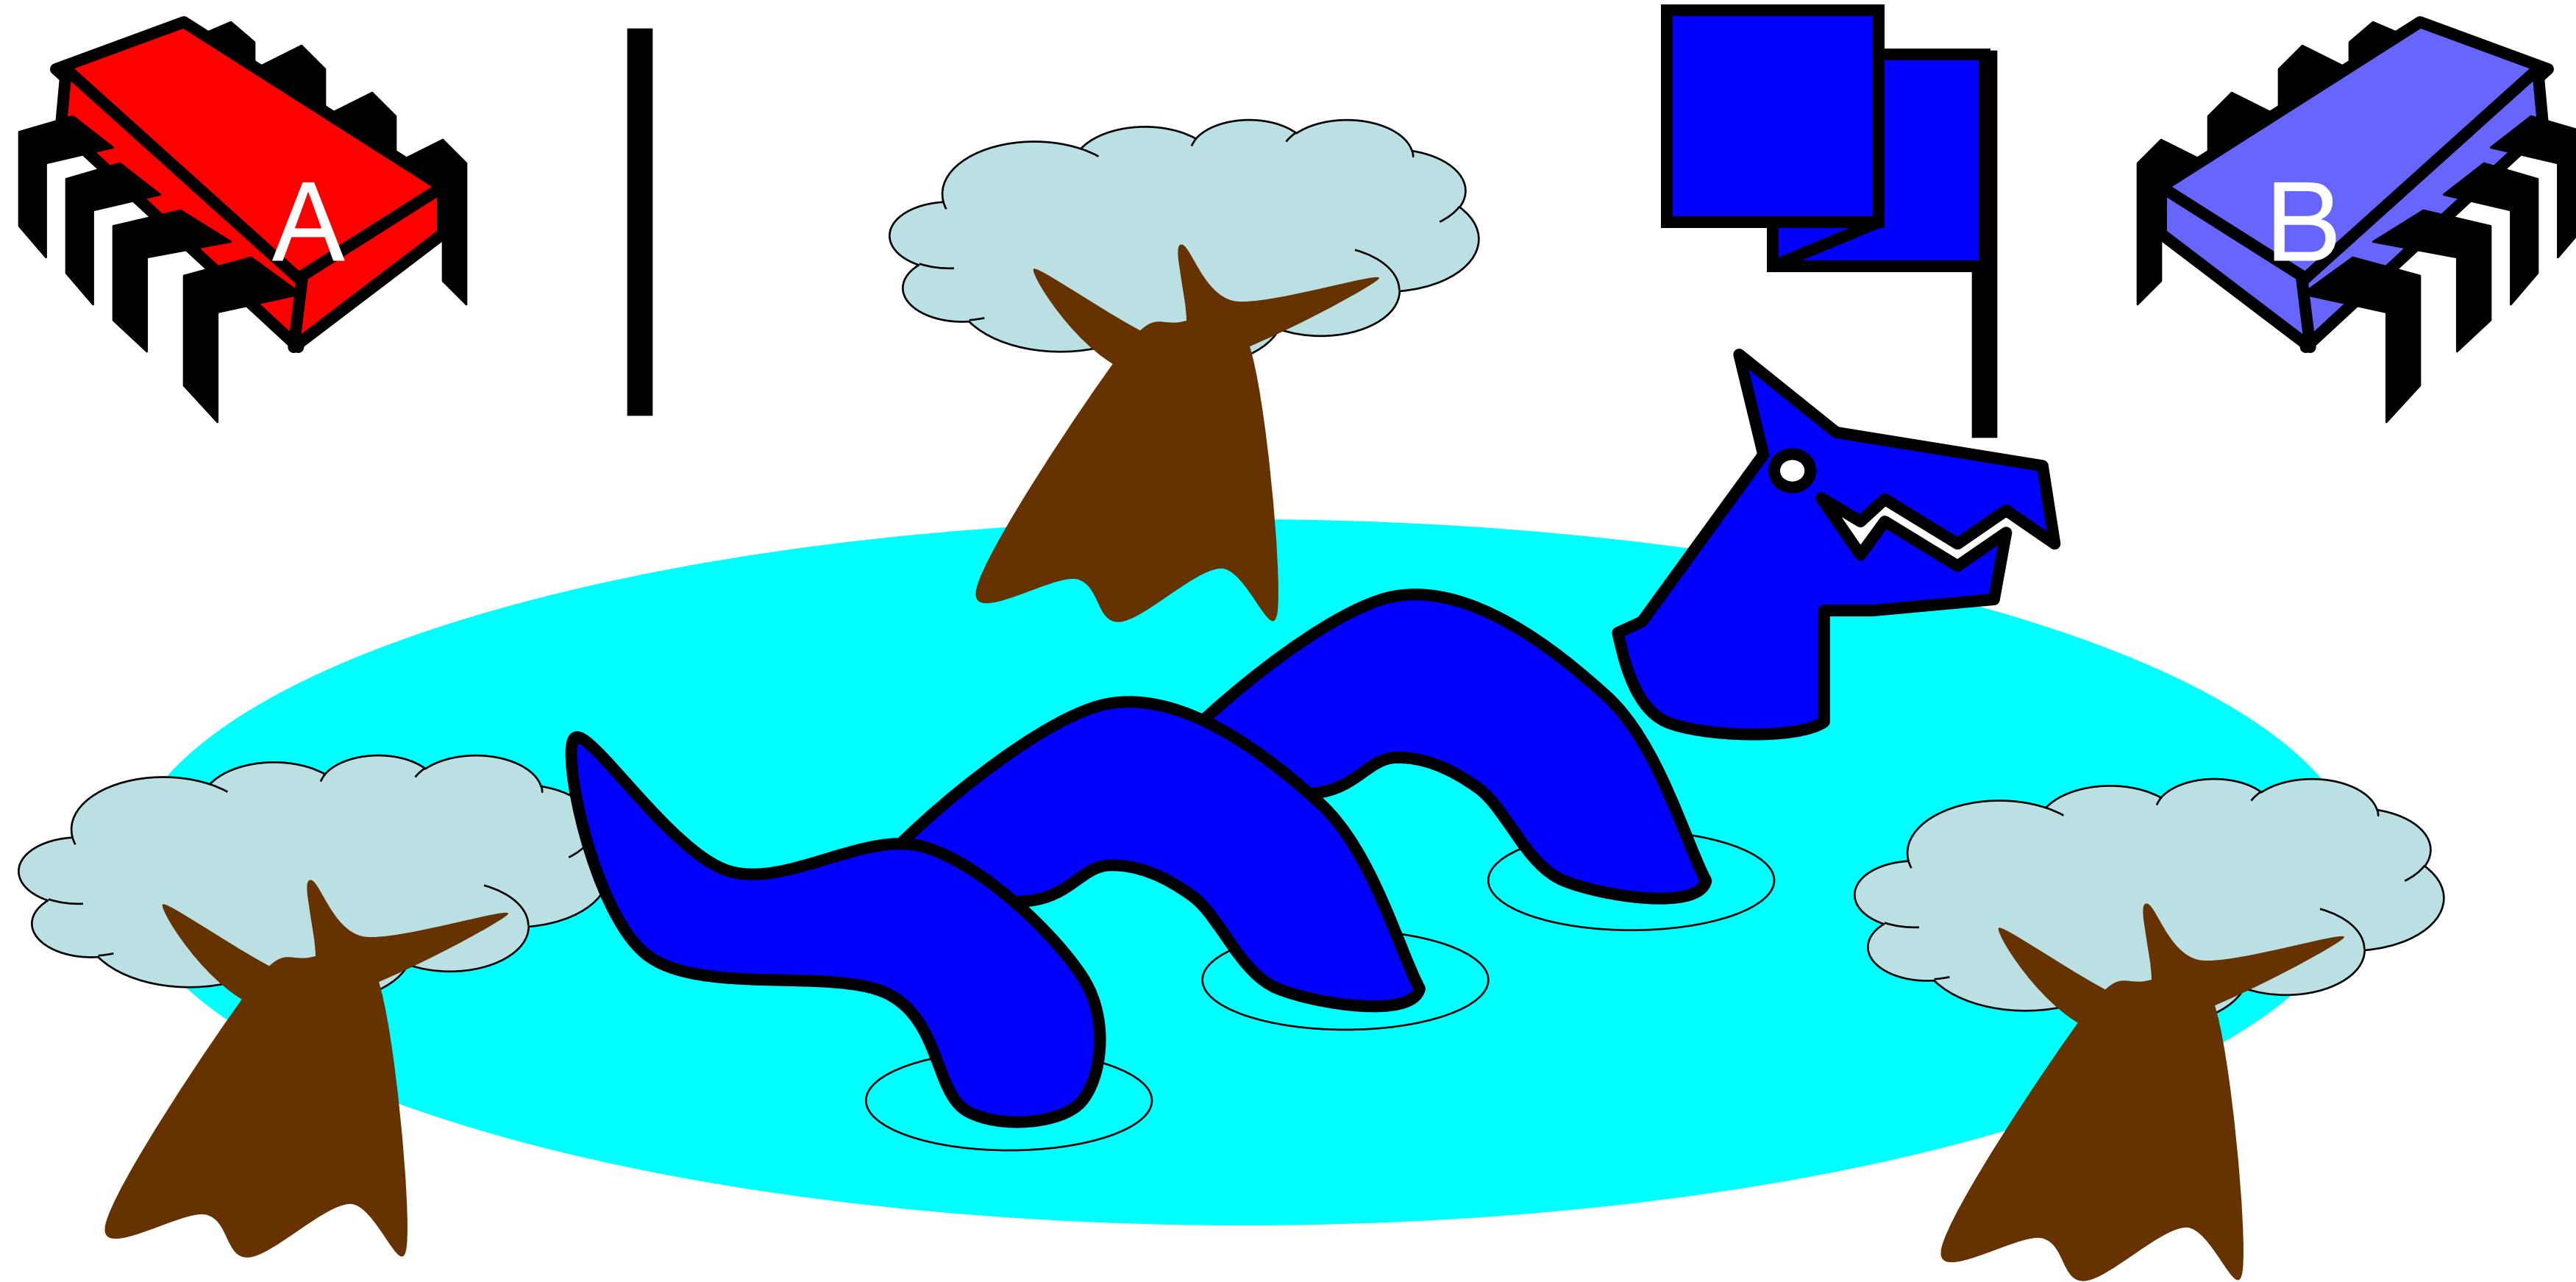
# Flag Protocol



# Alice's Protocol (sort of)



# Bob's Protocol (sort of)



# Flag Protocol

- **Alice's Protocol**
  - Raise flag
  - Wait until Bob's flag is down
  - Unleash pet
  - Lower flag when pet returns
- **Bob's Protocol**
  - Raise flag
  - Wait until Alice's flag is down
  - Unleash pet
  - Lower flag when pet returns



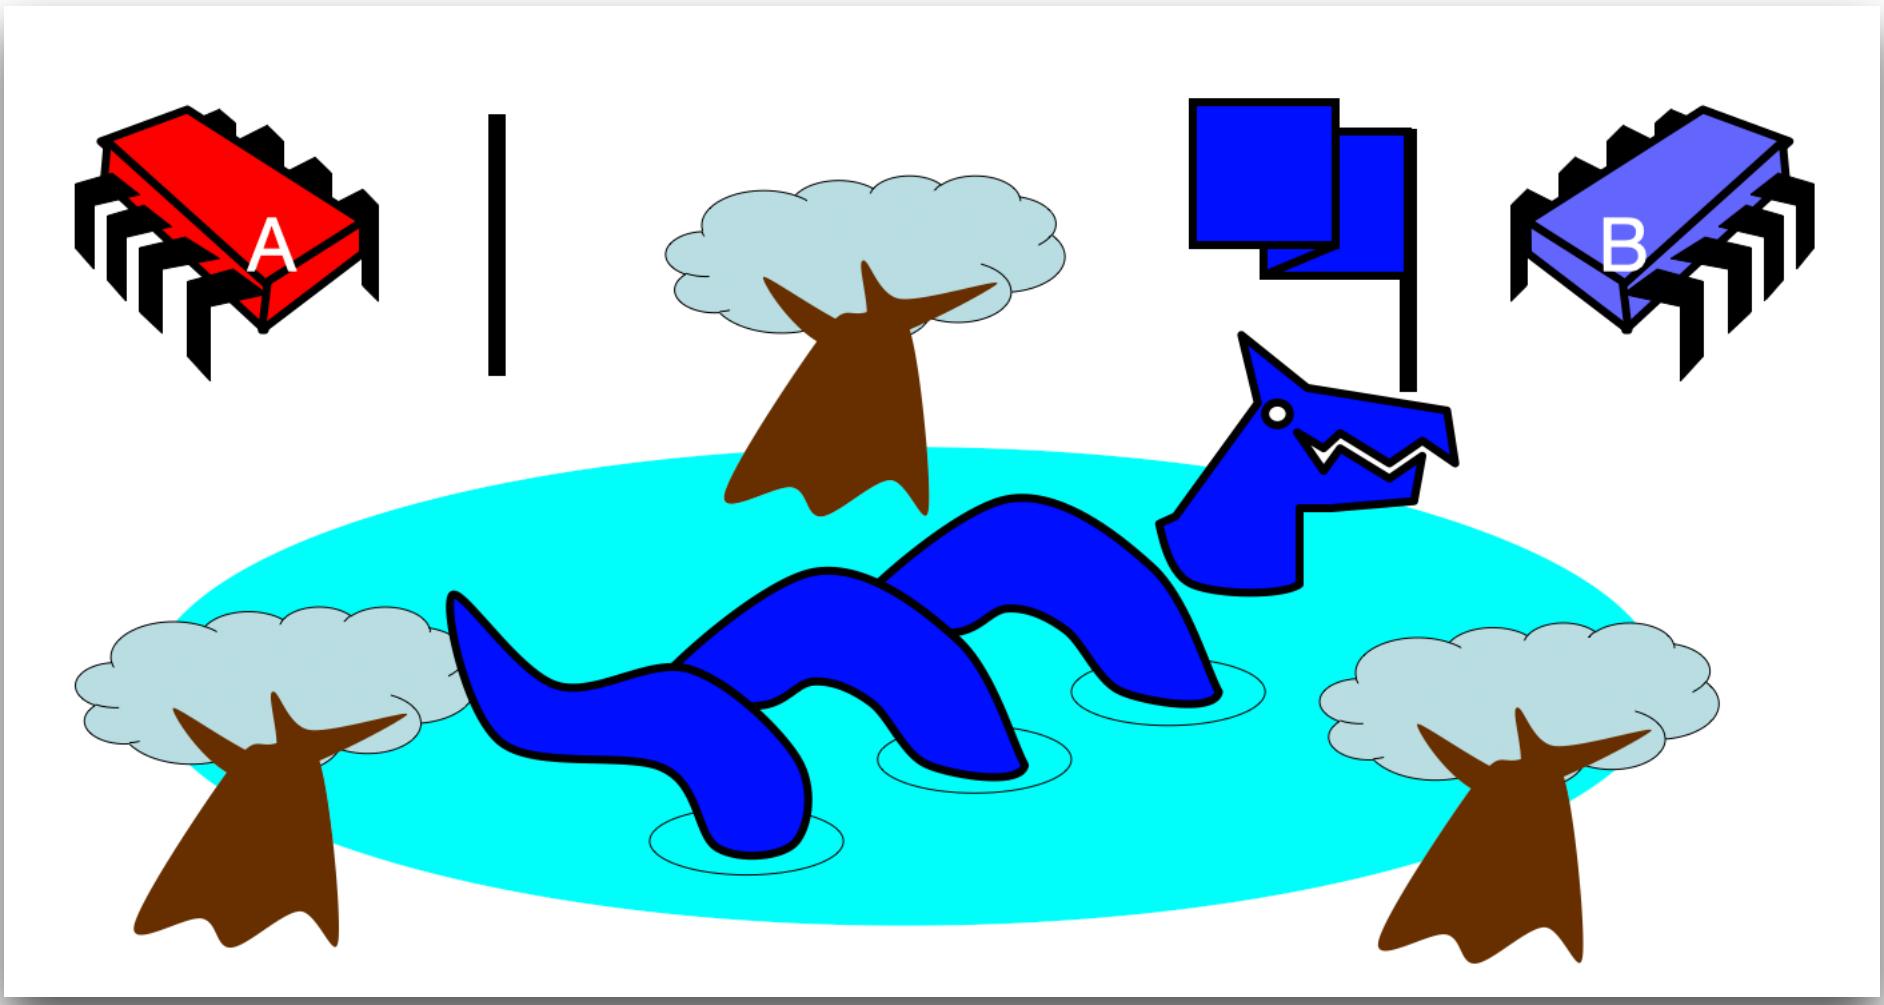
# Flag Protocol

- **Alice's Protocol**
  - Raise flag
  - Wait until Bob's flag is down
  - Unleash pet
  - Lower flag when pet returns
- **Bob's Protocol (2nd try)**
  - Raise flag
  - While Alice's flag is up
    - Lower flag
    - Wait for Alice's flag to go down
    - Raise flag
    - Unleash pet
    - Lower flag when pet returns

Bob defers to Alice

# The Flag Principle

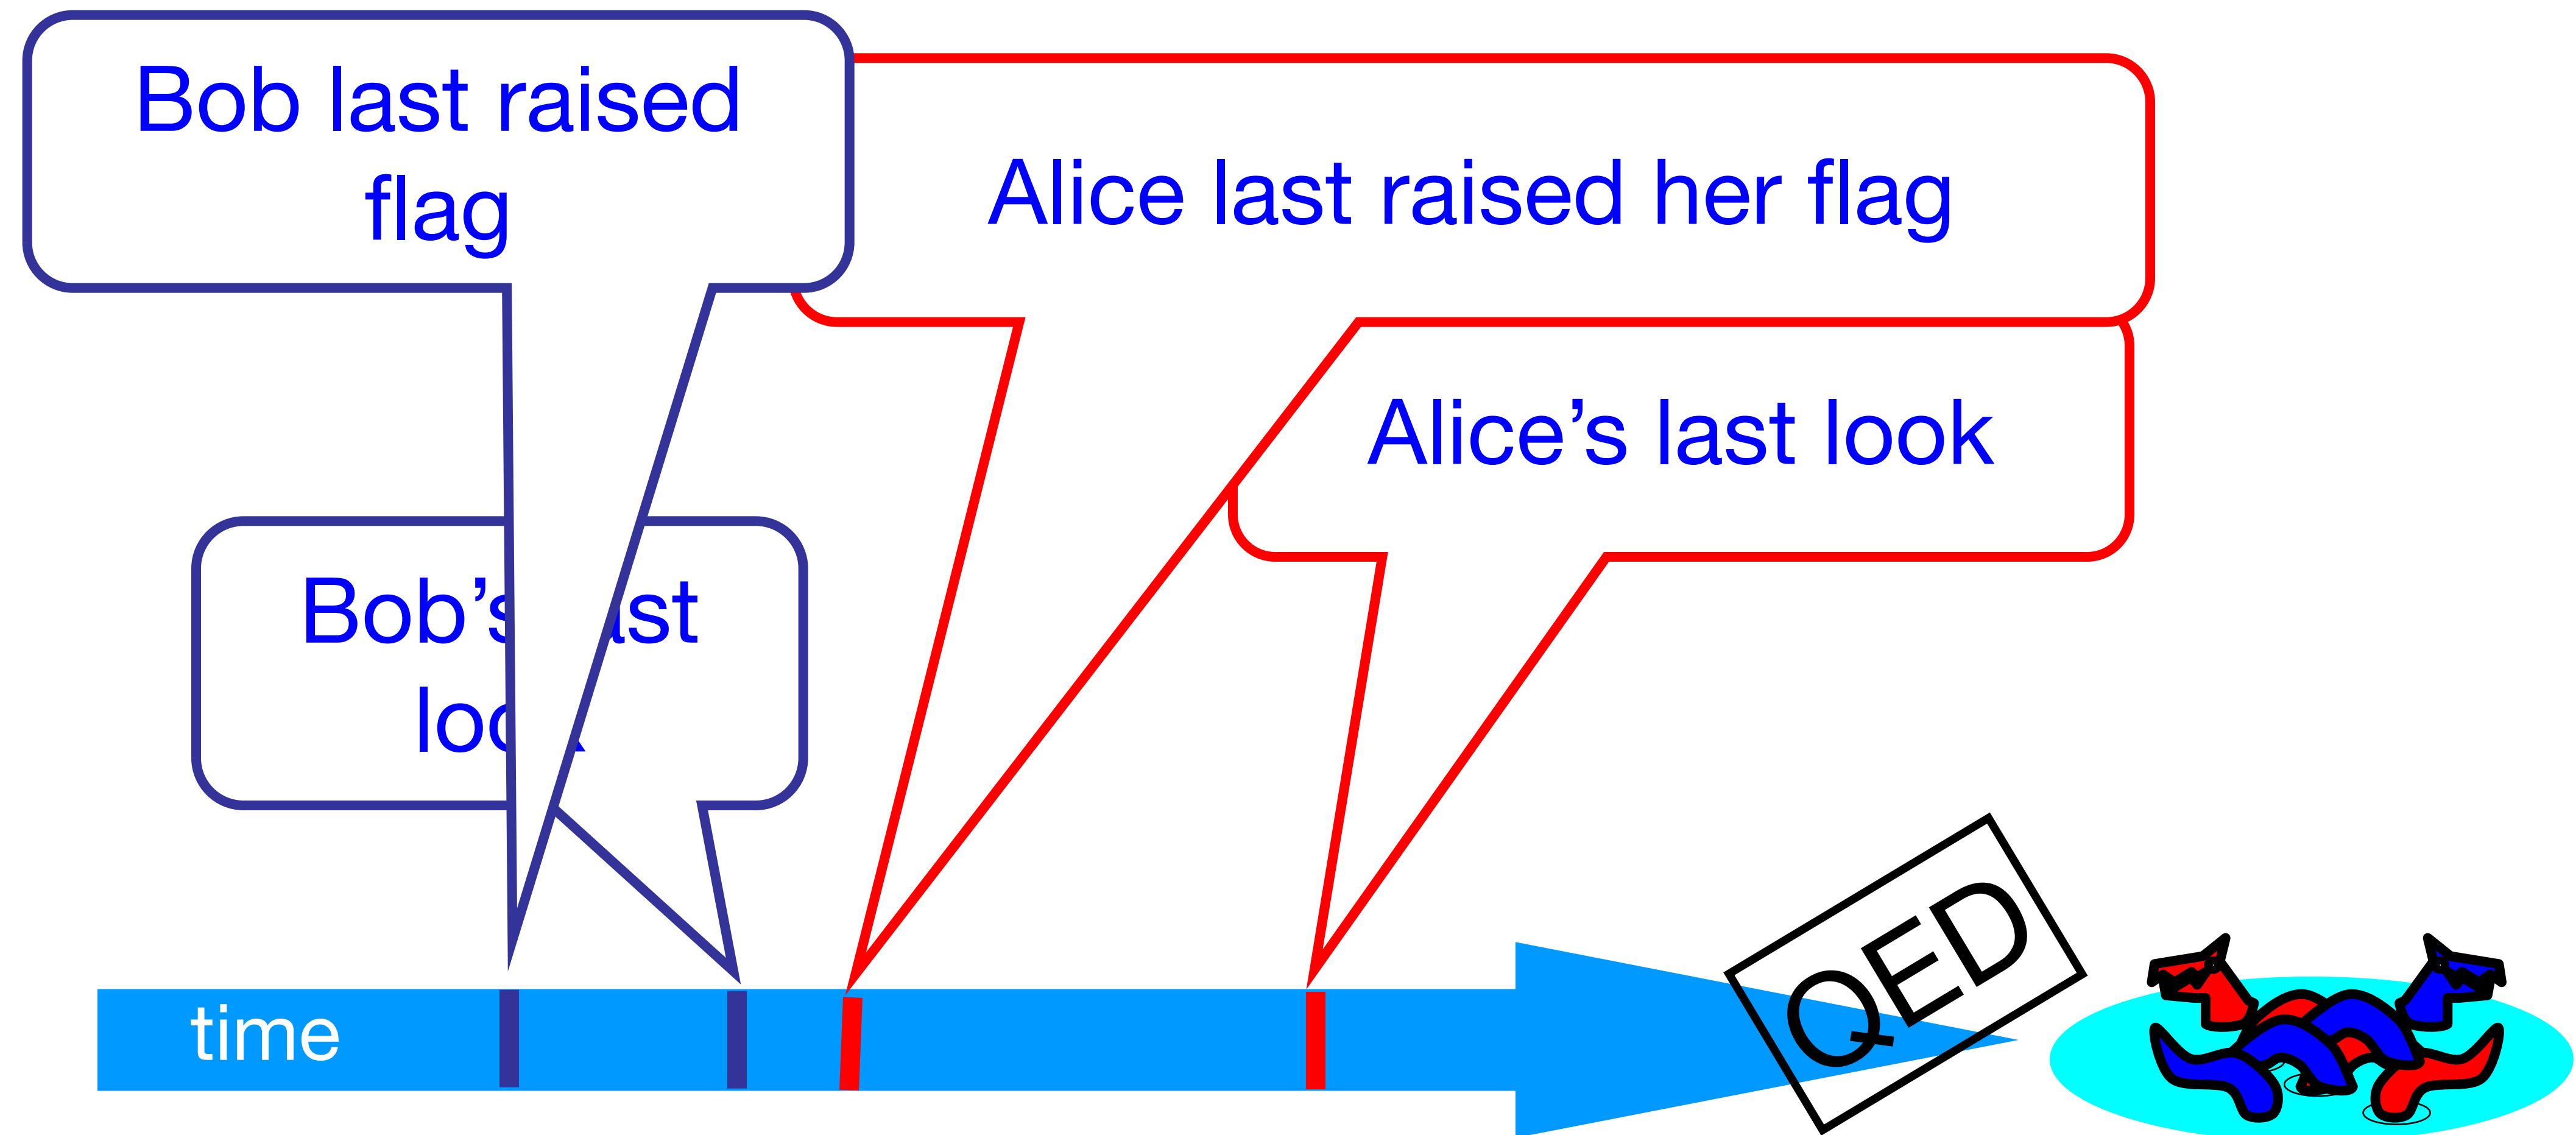
- Raise the flag
- Look at others' flag
- Flag Principle:
  - If each raises and looks, then
  - Last to look must see both flags up



# Proof of Mutual Exclusion

- Assume both pets in pond
  - Derive a contradiction
  - By reasoning ***backwards***
- Consider the last time Alice and Bob each looked before letting the pets in
- Without loss of generality assume Alice was the last to look...

# Proof of Mutual Exclusion



Alice must have seen Bob's Flag. A Contradiction

# Proof of No Deadlock

- If one pet wants to get in, it gets in.
- Deadlock requires both parties to continually try to get in.
- If Bob sees Alice's flag, he backs off and gives her priority
  - Alice's lexicographic privilege



# Remarks

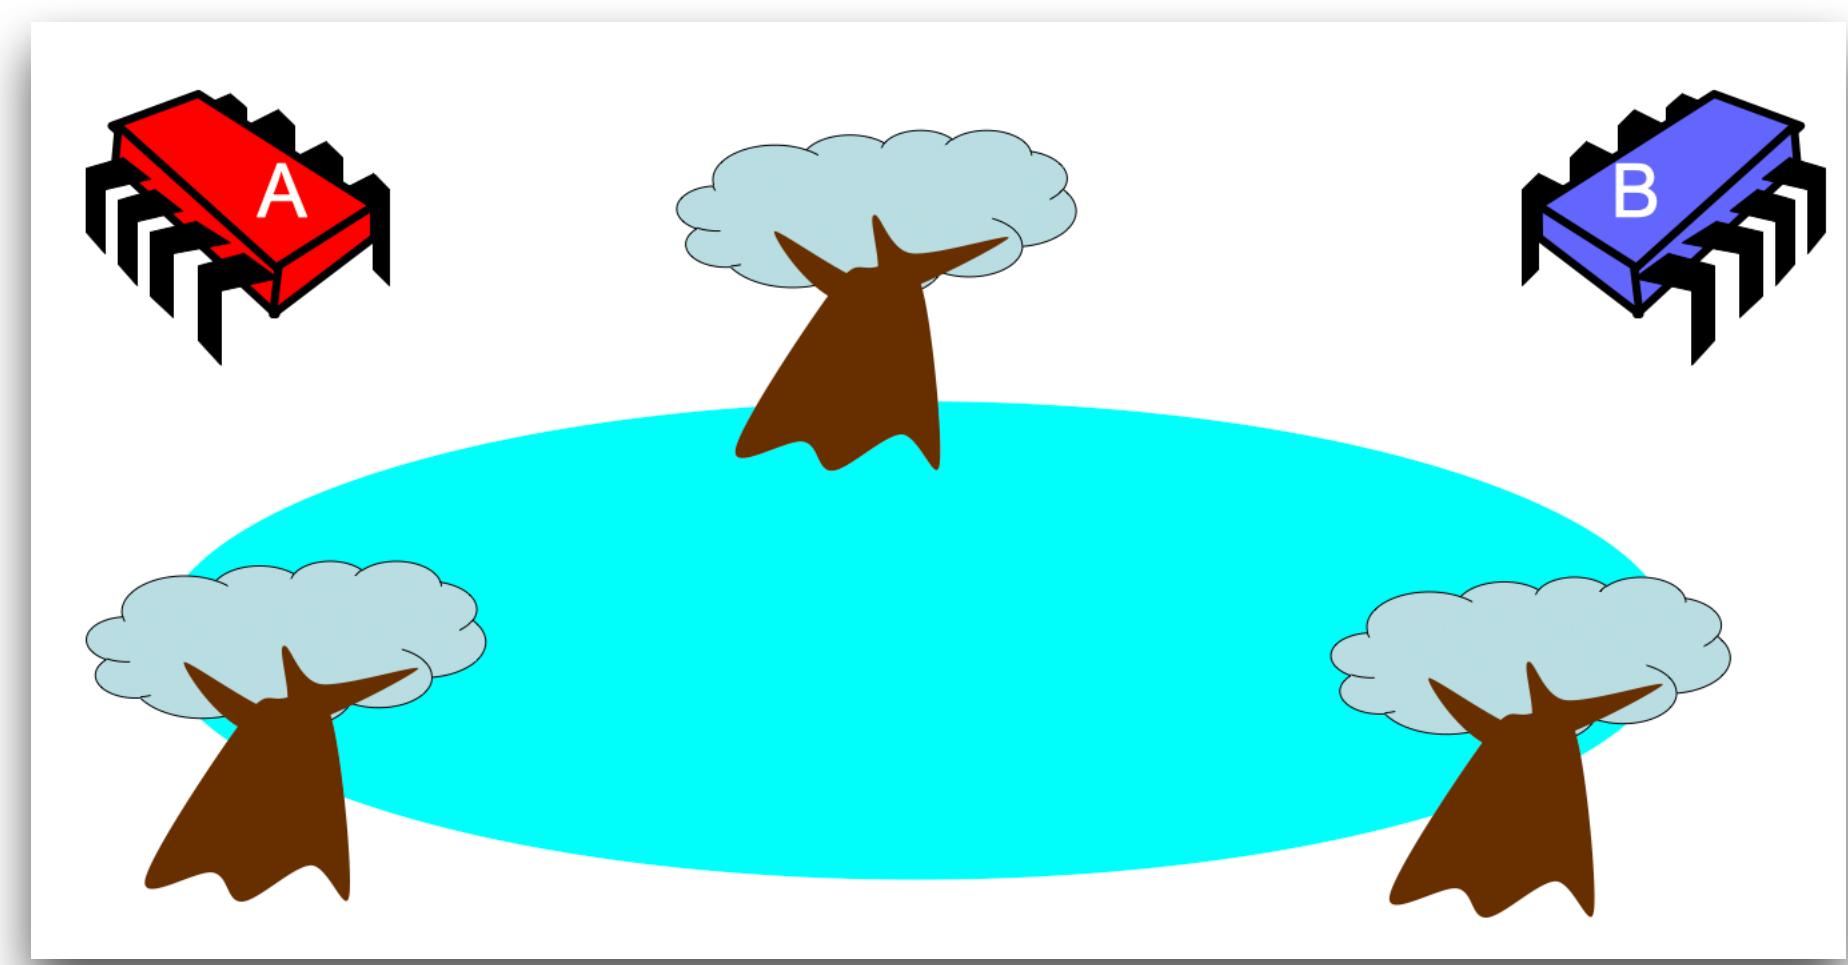
- Protocol is *unfair*
  - Bob's pet might never get in and *starve*
- Protocol uses *waiting*
  - If Bob is eaten by his pet, Alice's pet might never get in

# Moral of the story

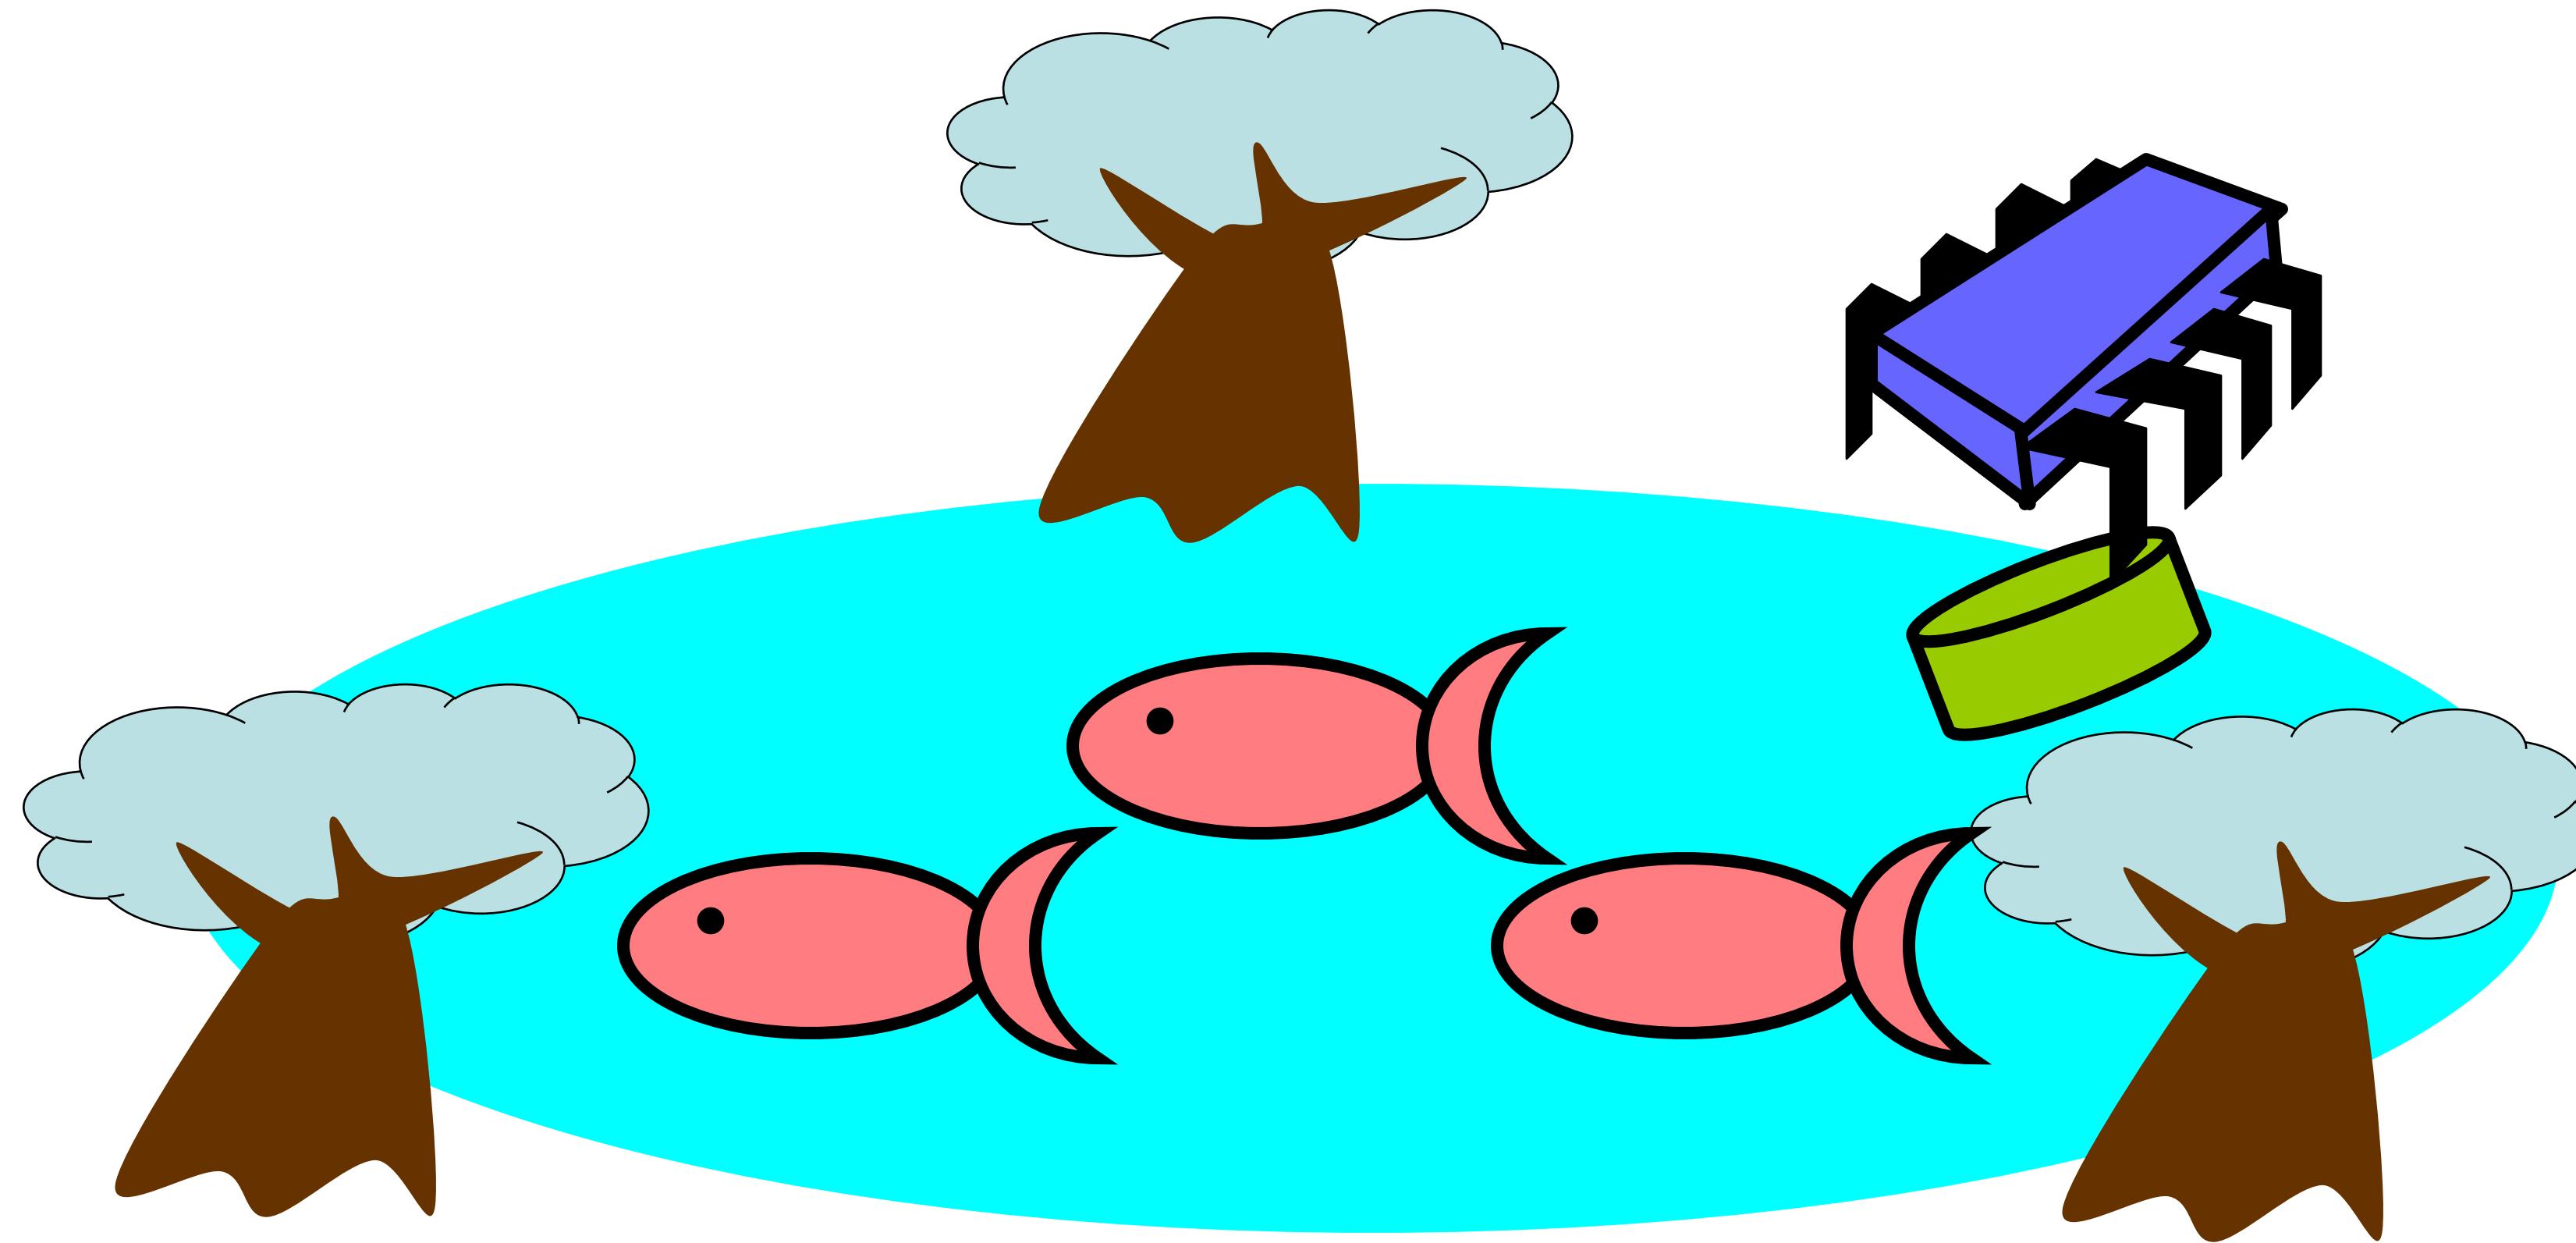
- Mutual Exclusion ***cannot be solved*** by
  - transient communication (cell phones)
  - interrupts (cans)
- It ***can be solved*** by
  - one-bit shared variables
  - that can be read or written

# The fable continues

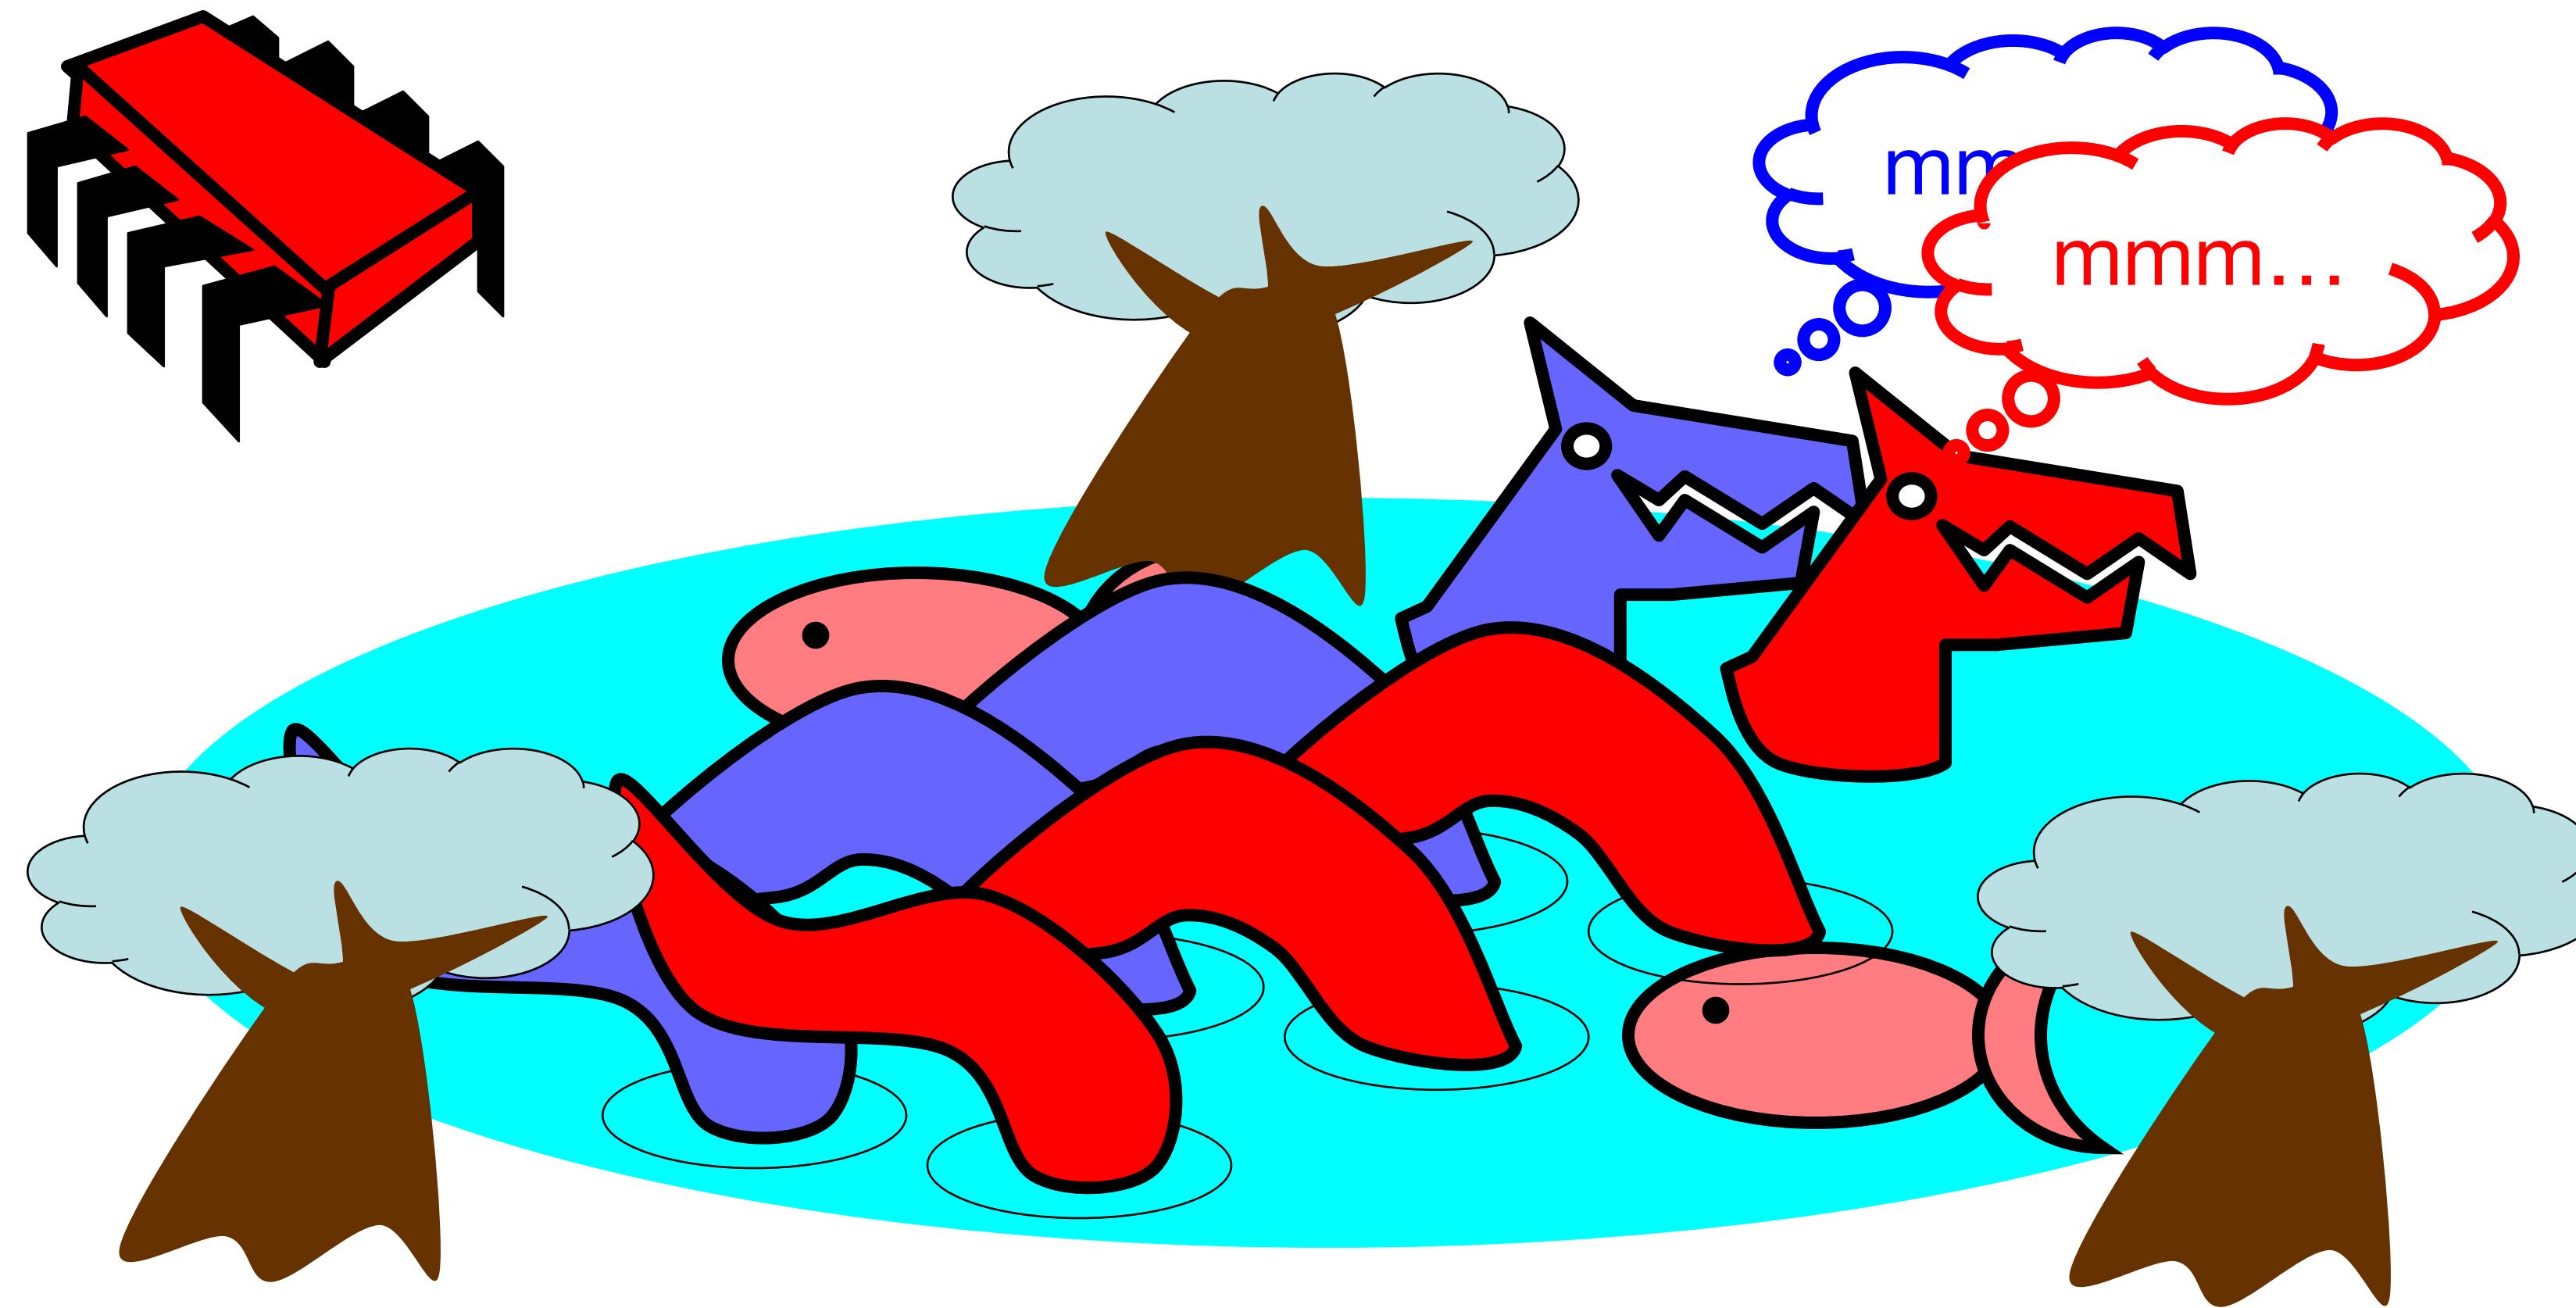
- Alice and Bob fall in love & marry
  - Then they fall out of love & divorce
  - She gets the pets
- He has to feed them
- Leading to a new coordination problem:  
**Producer-Consumer**



# Bob Puts Food in the Pond



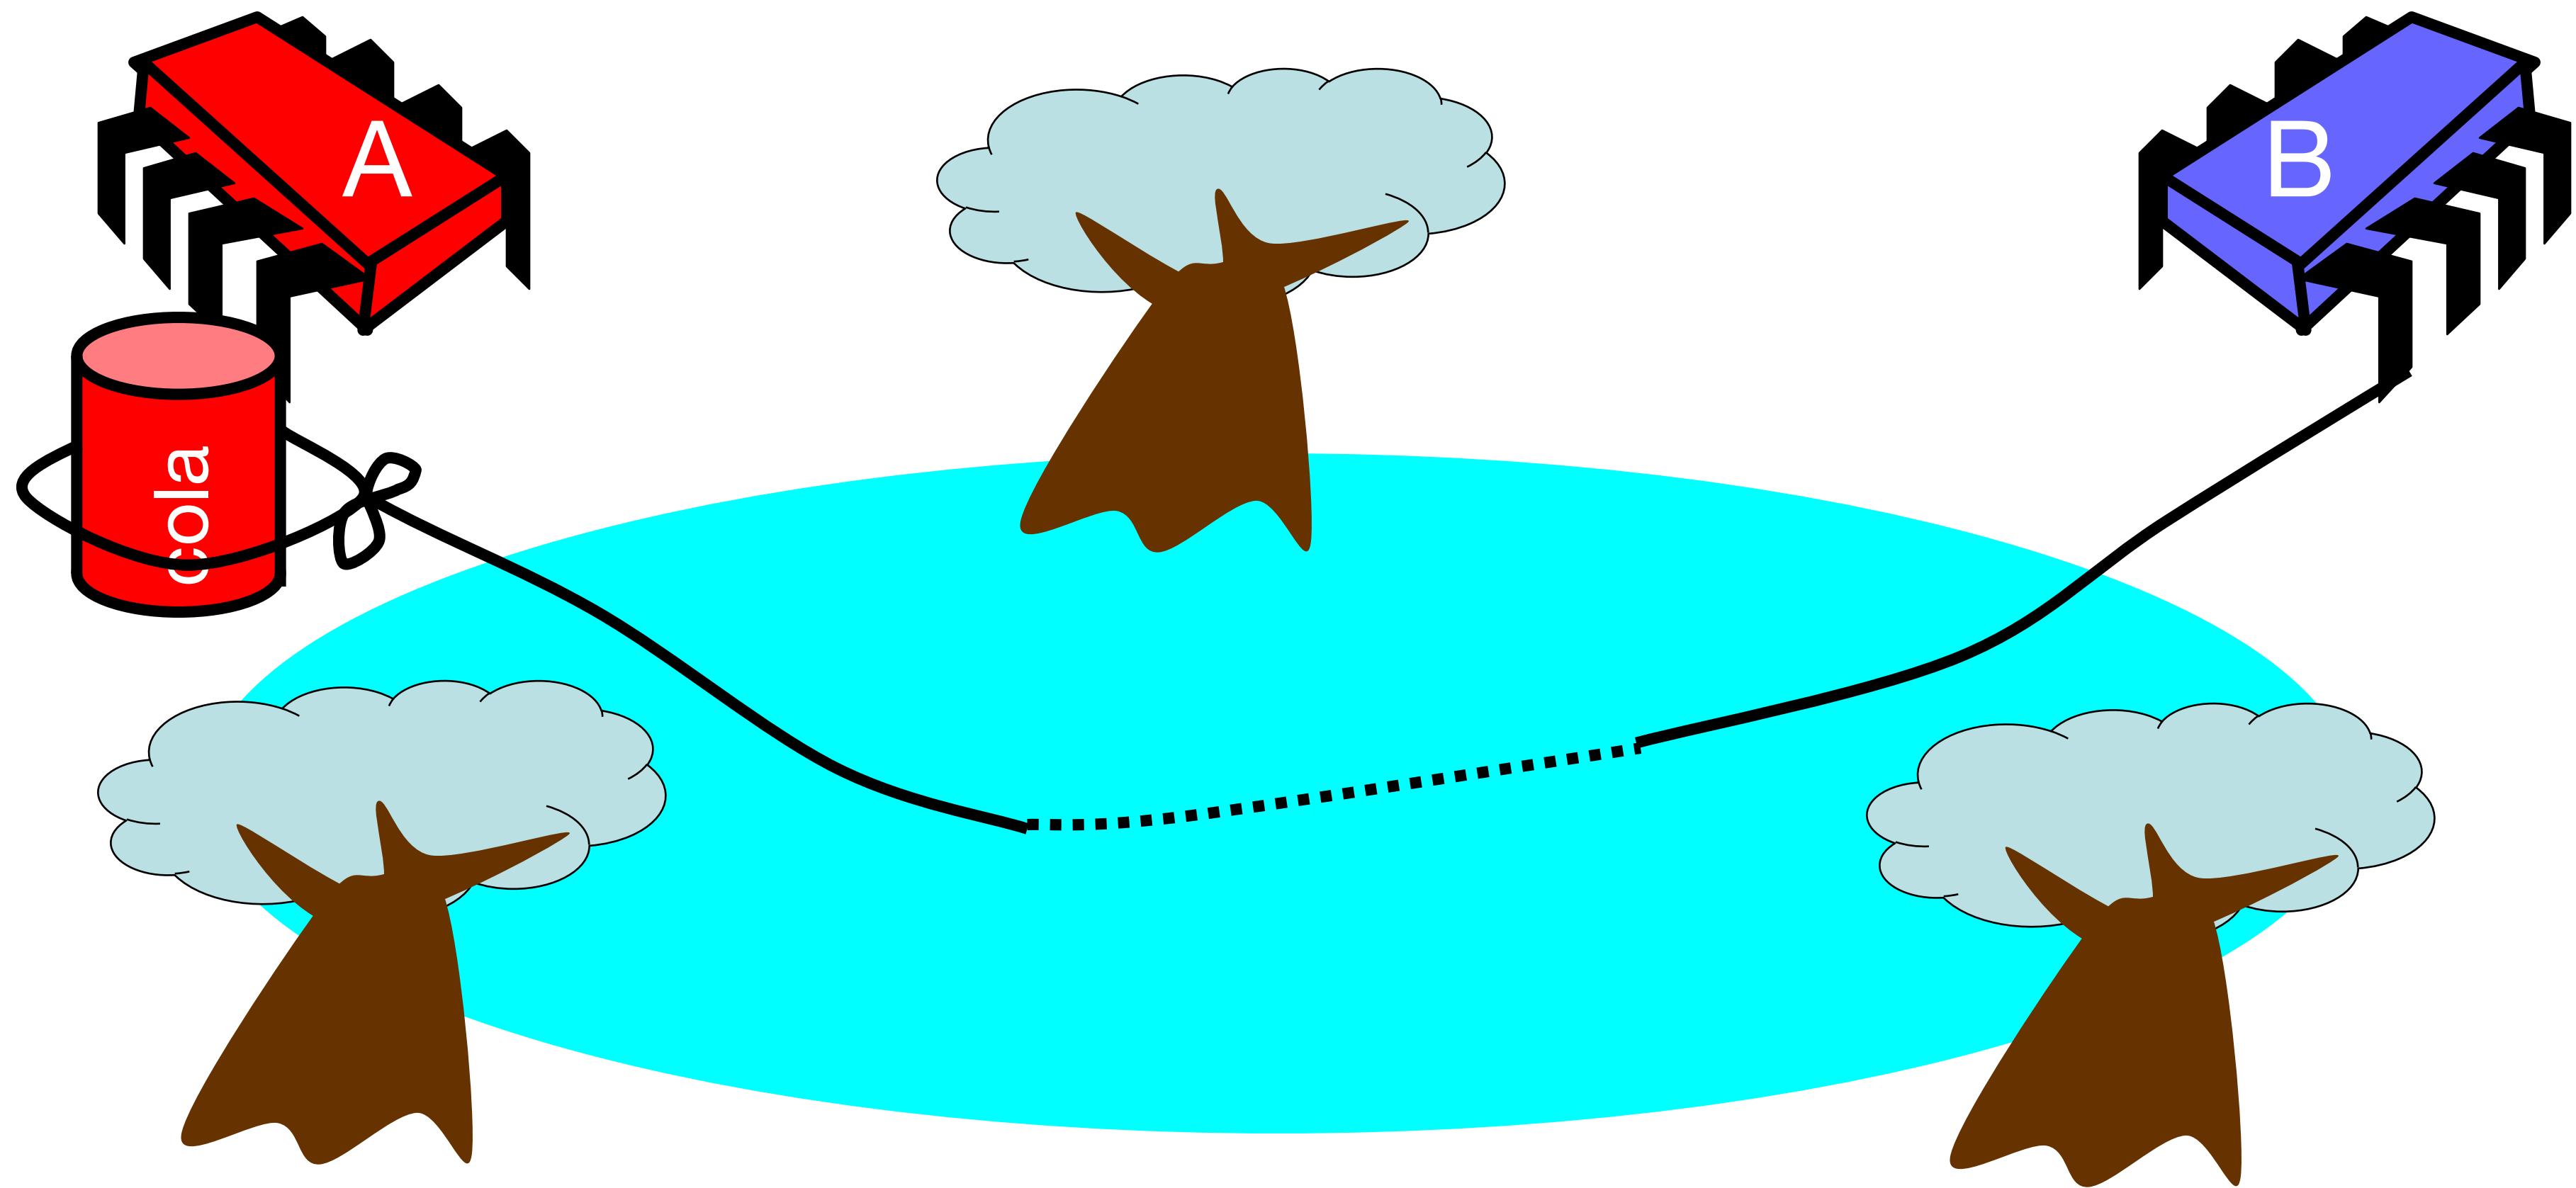
# Alice releases her pets to feed



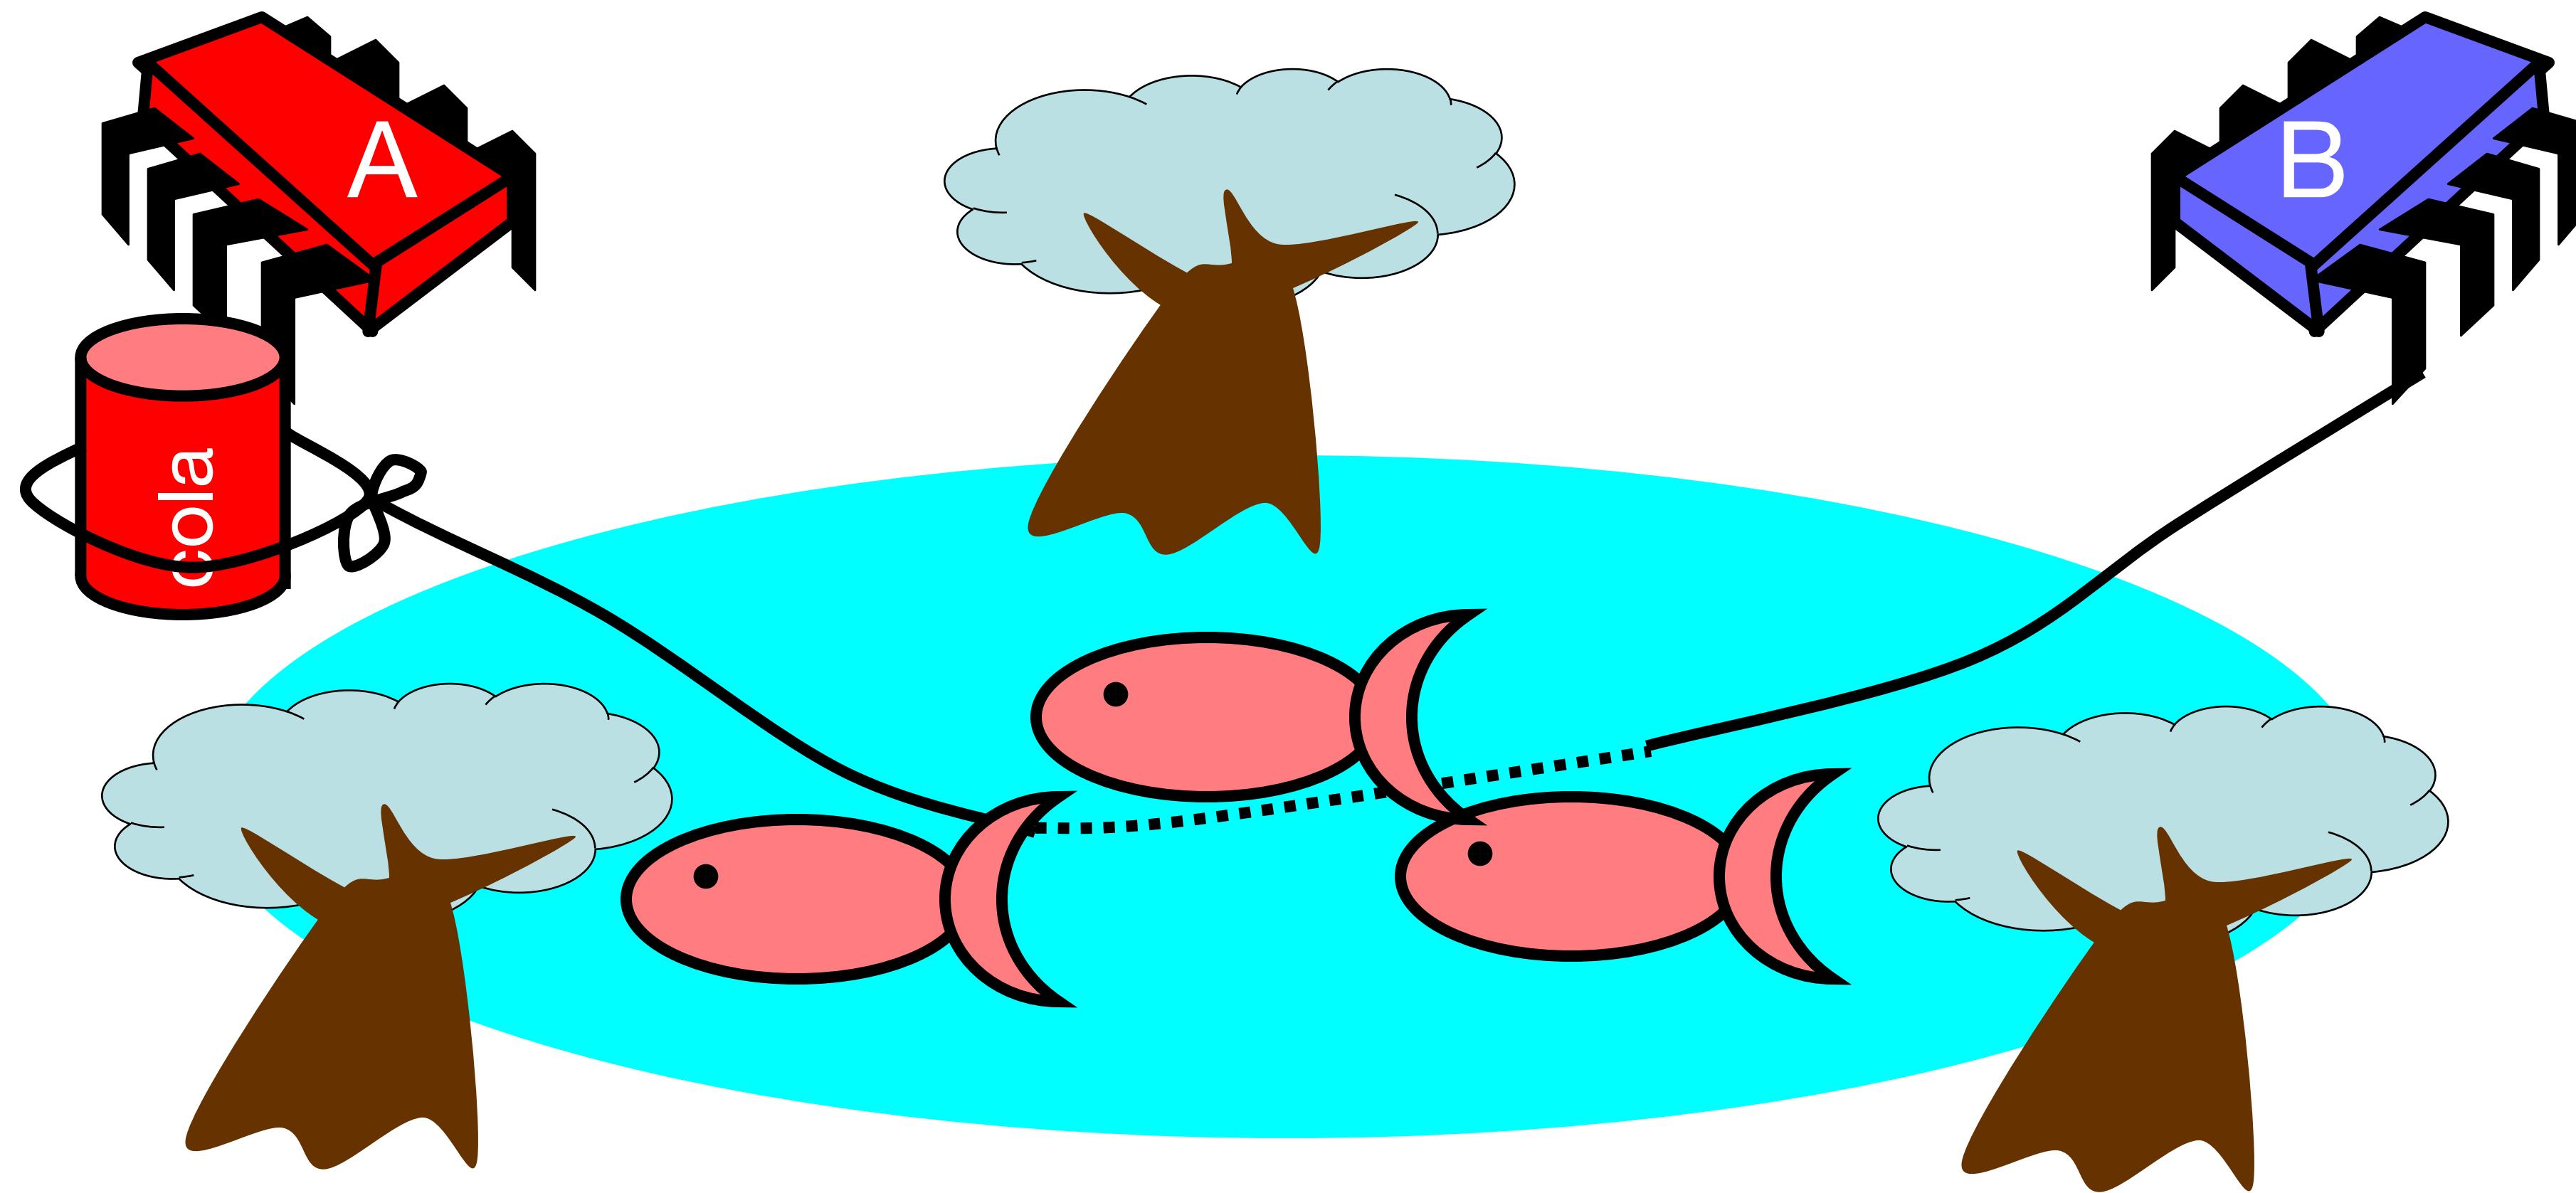
# Producer/Consumer

- Alice and Bob can't meet
  - Each has a ***restraining order on the other***
  - So he puts food in the pond
  - And later, she releases the pets
- Avoid
  - Releasing pets when there's no food
  - Putting out food if uneaten food remains
- Need a mechanism so that
  - Bob lets Alice know when the food has been put out
  - Alice lets Bob know when to put out more food

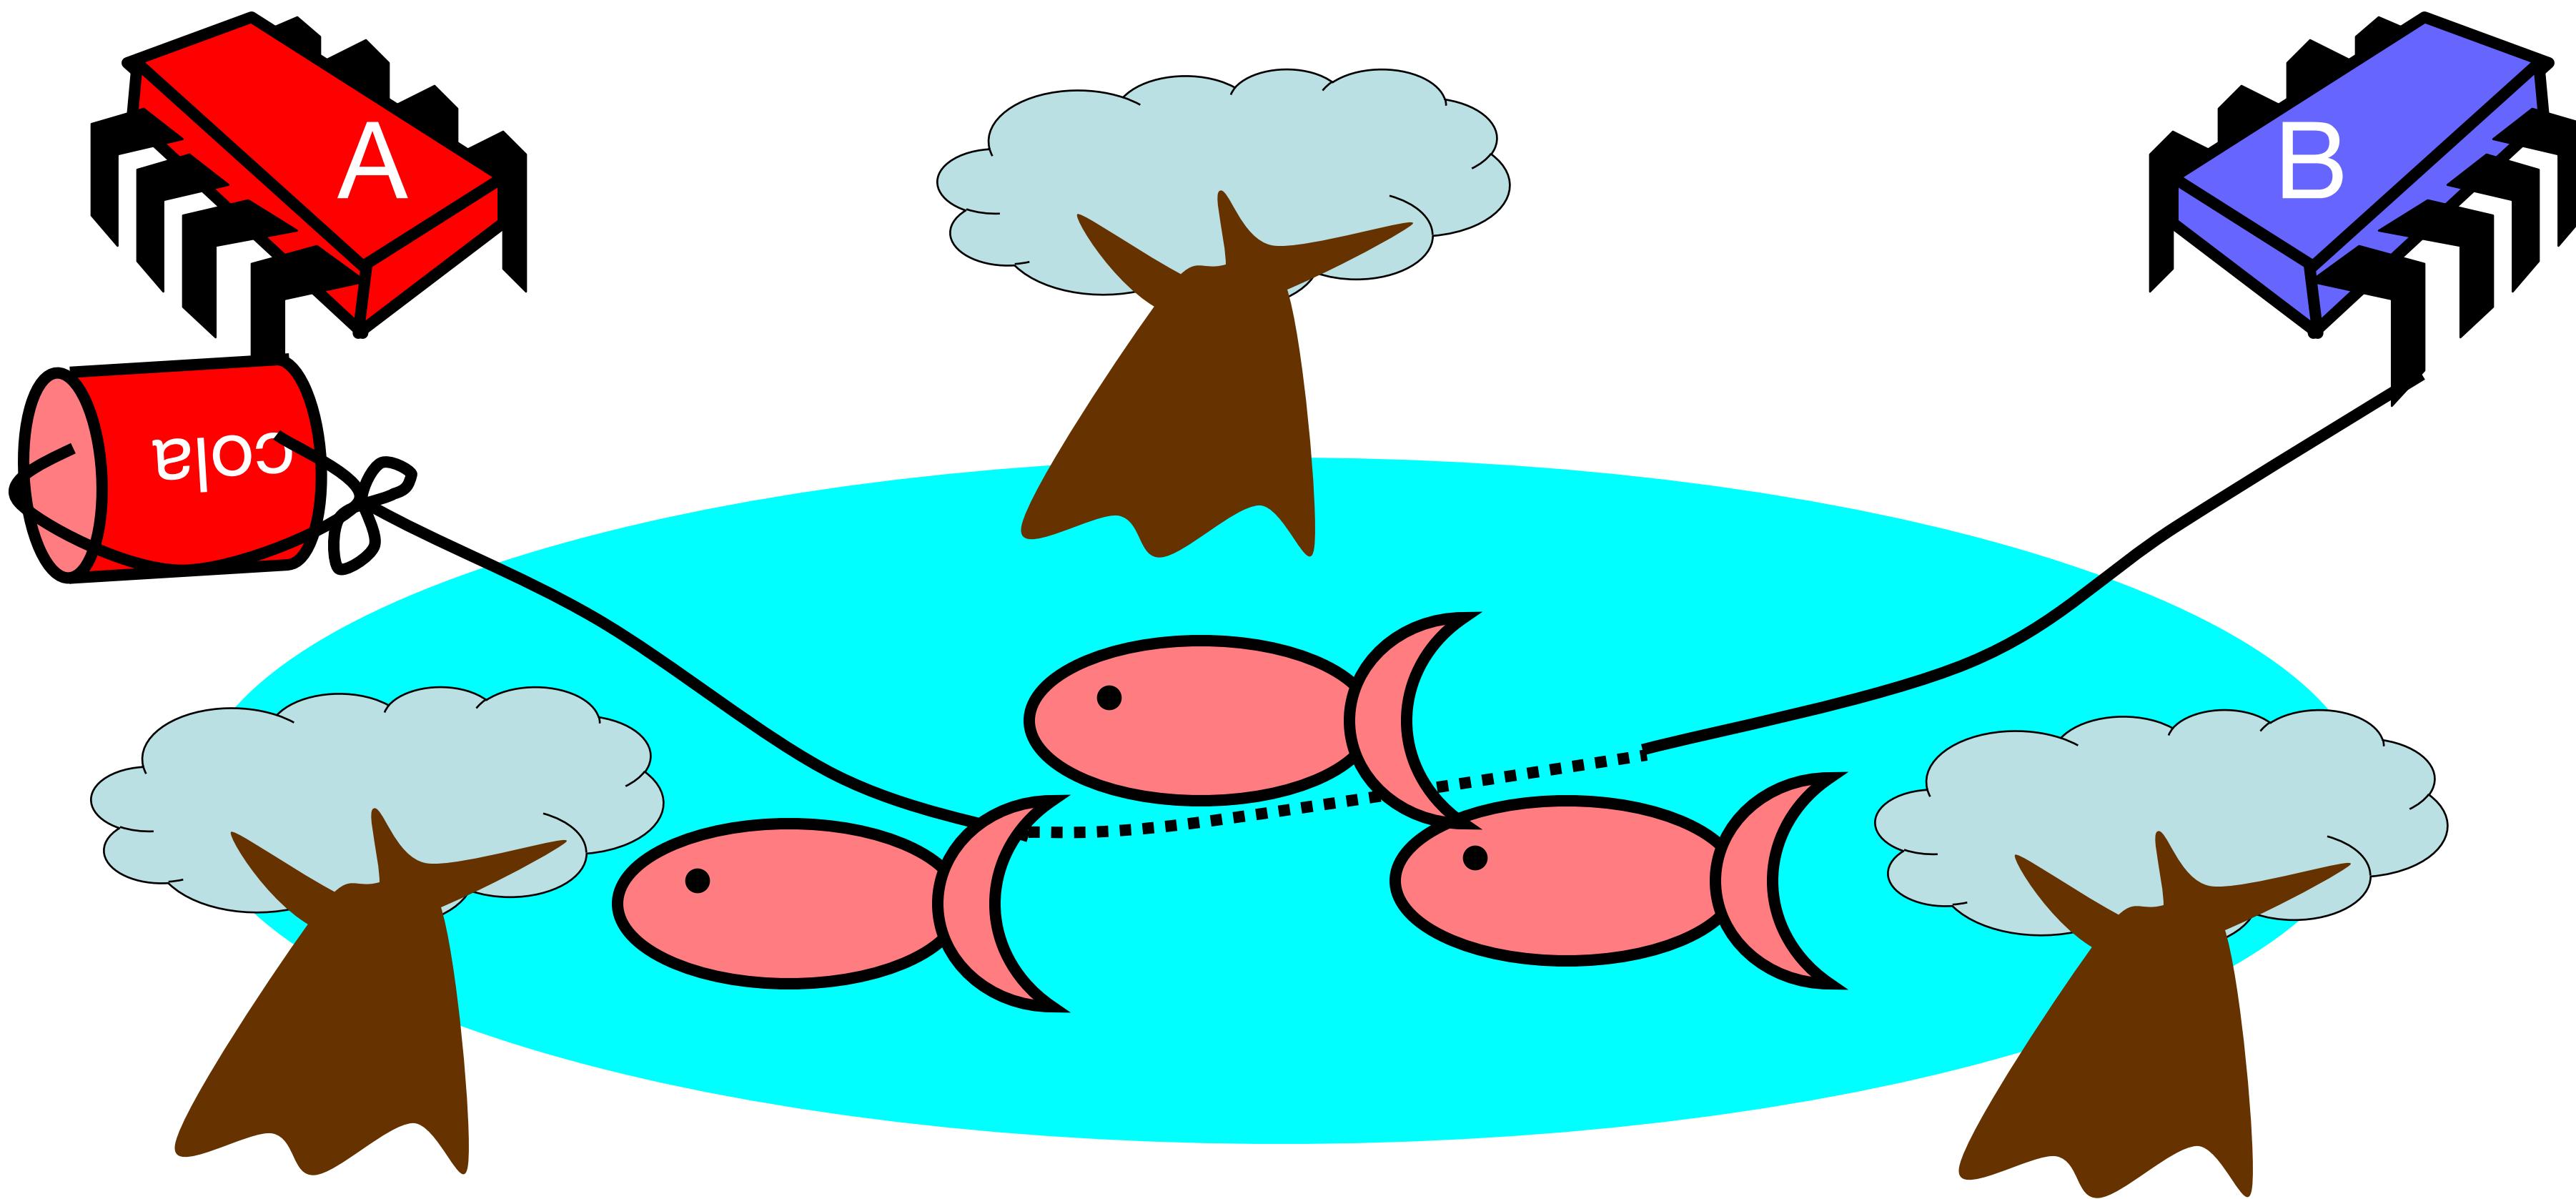
# Surprise Solution



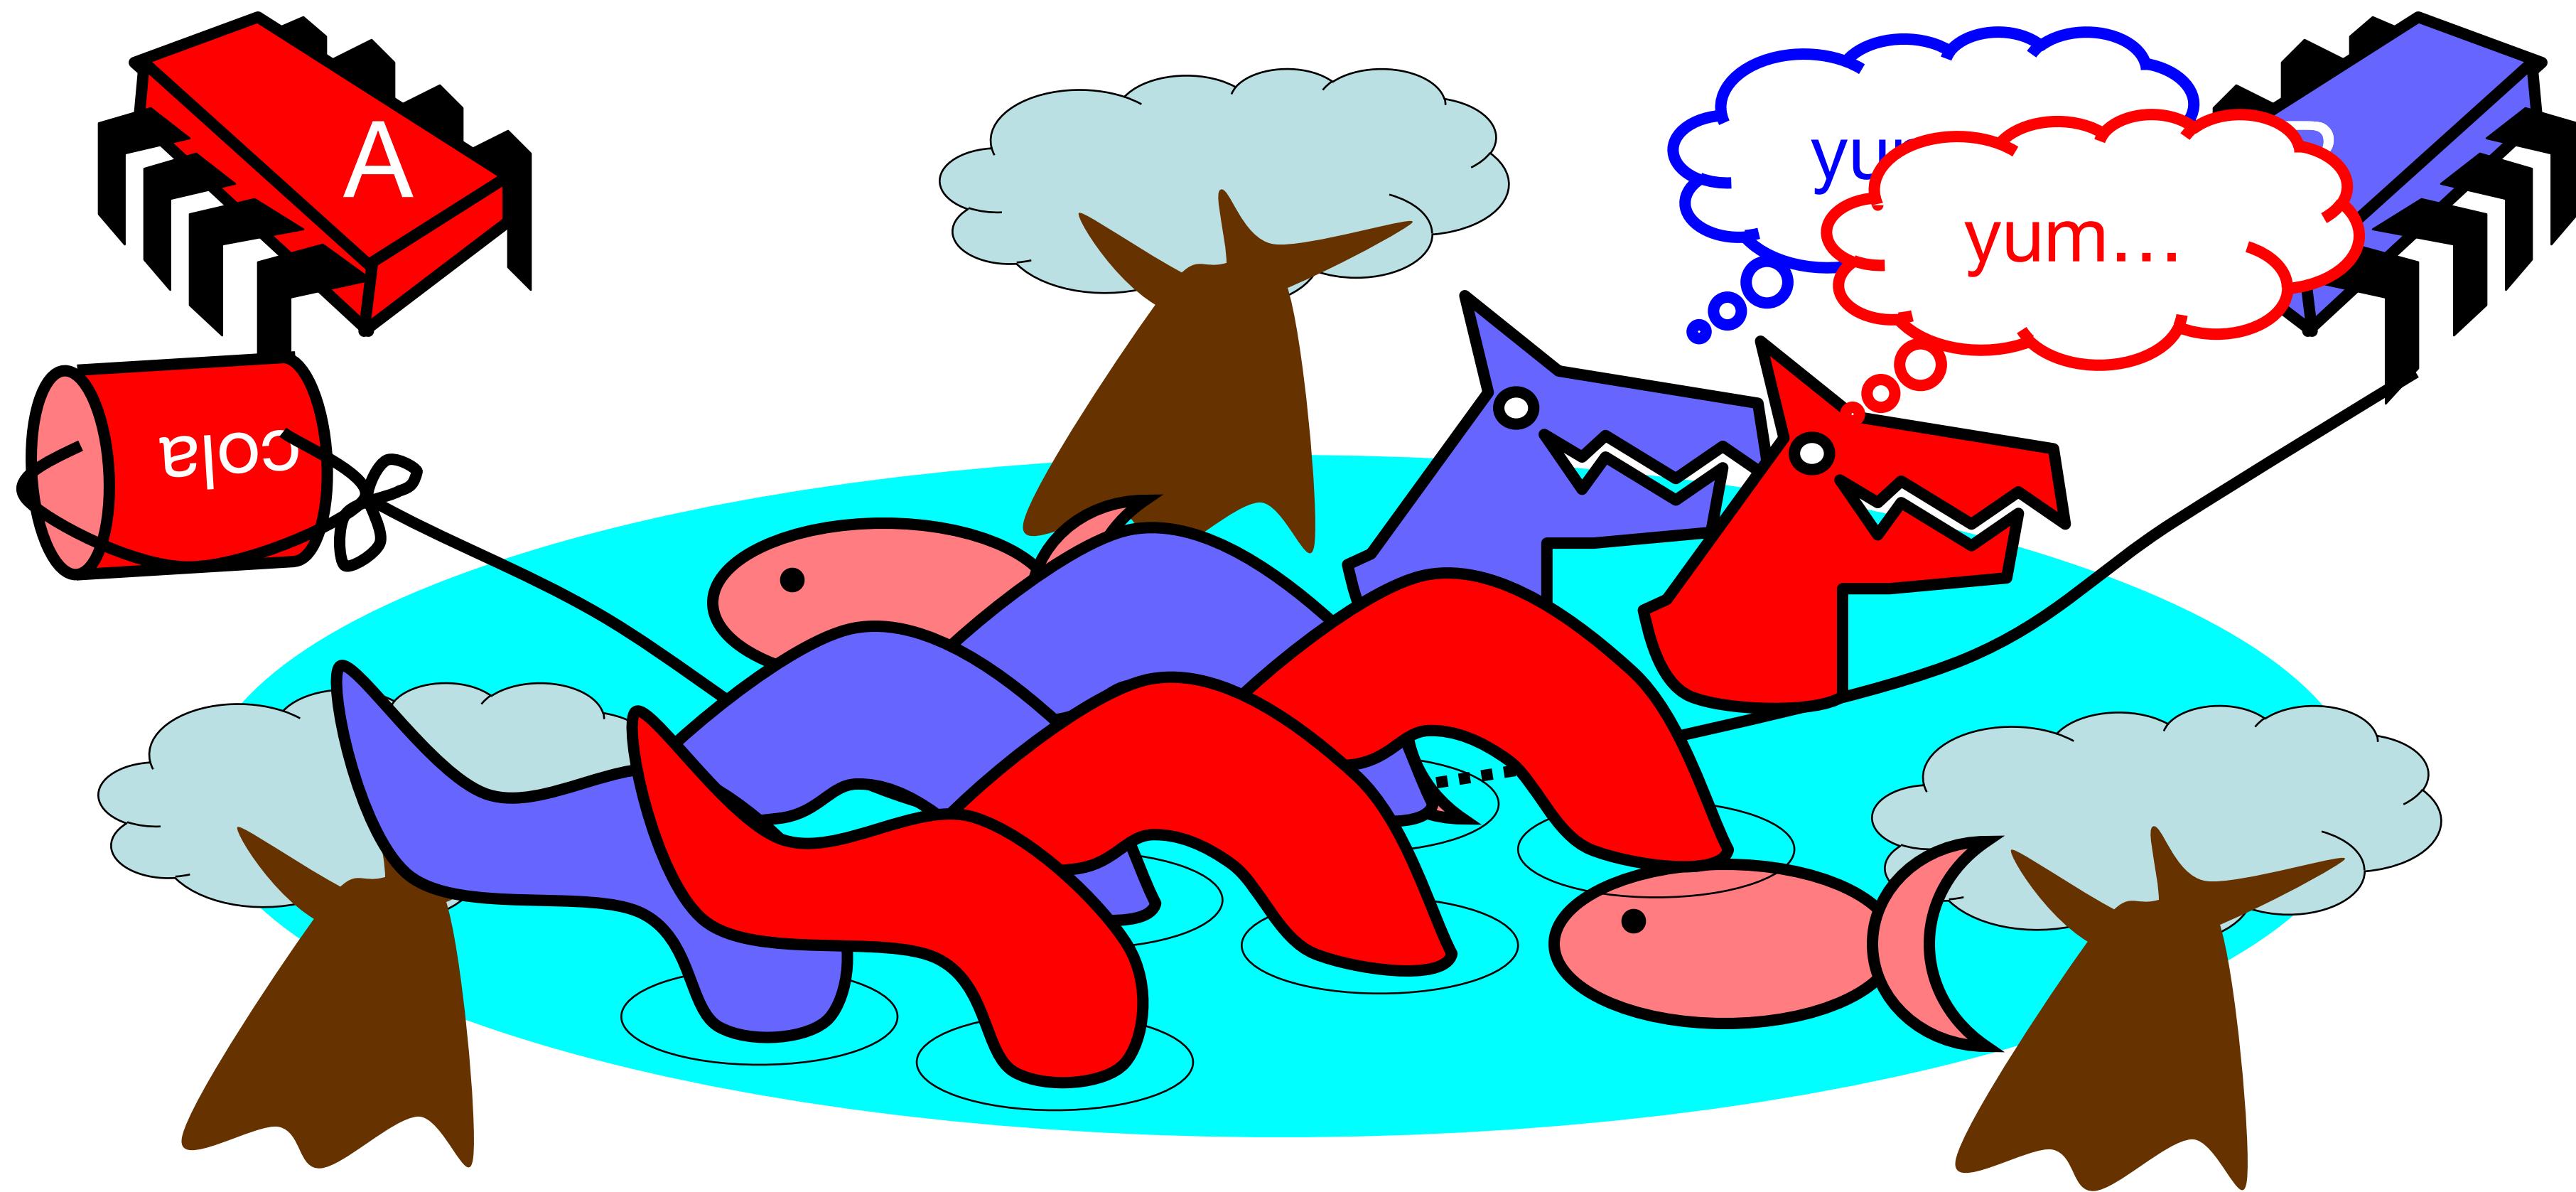
# Bob puts food in the Pond



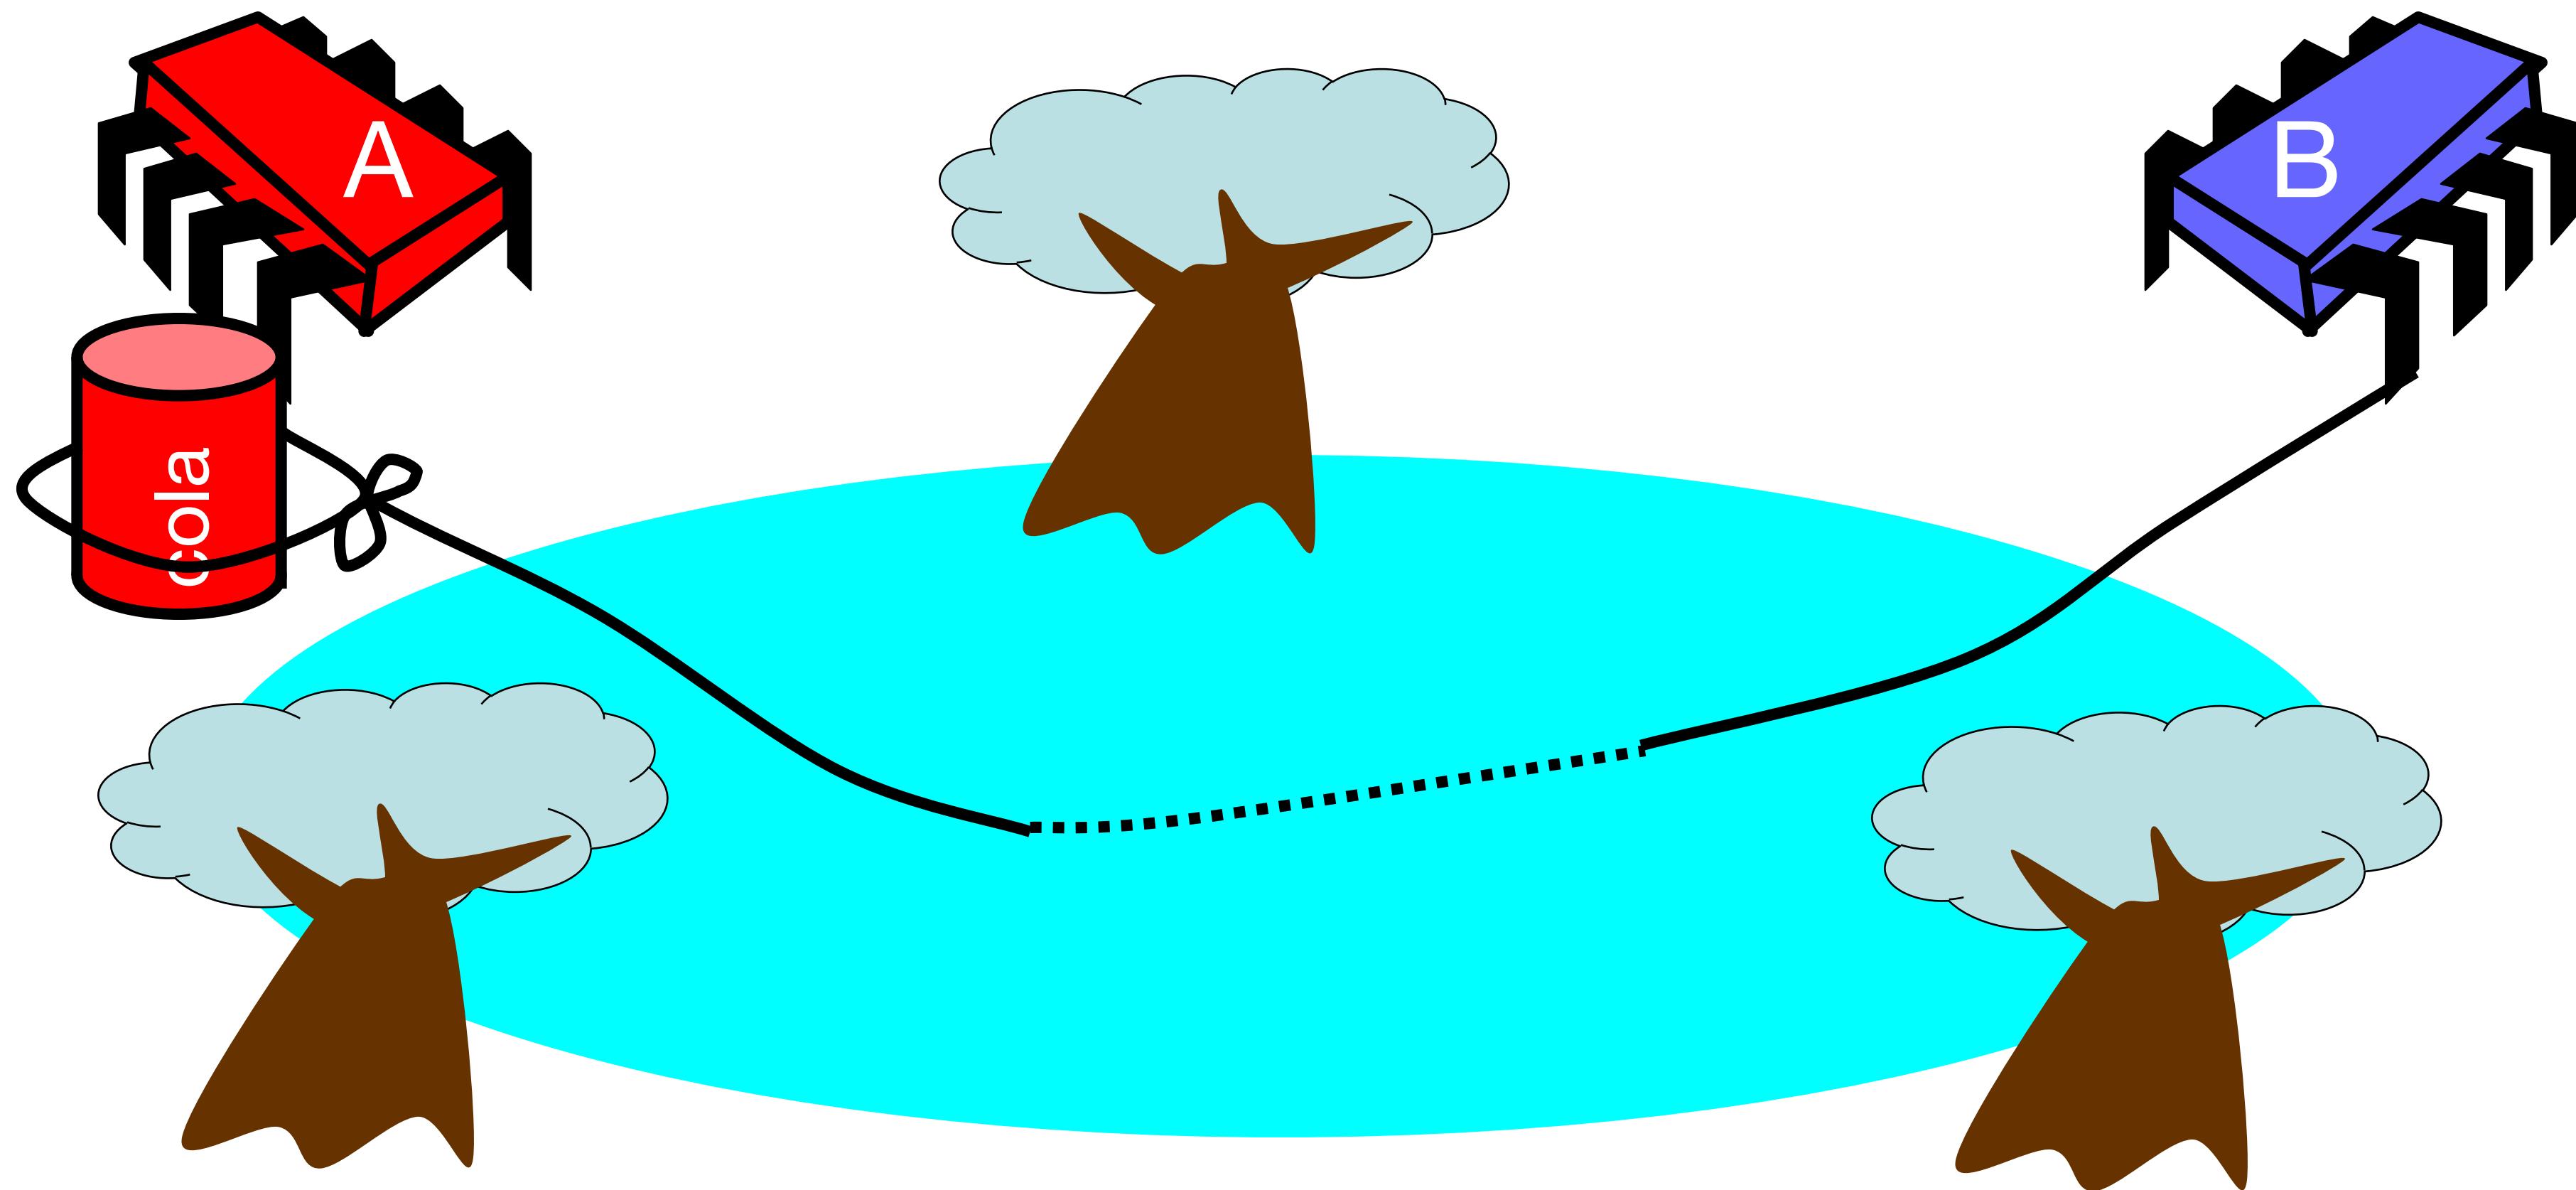
# Bob knocks over can



# Alice Releases Pets



# Alice resets can when pets are fed



# Producer/Consumer Code

```
(* Alice (Consumer) - releases pets to eat fish *)
let alice can pond =
  while true do
    while is_up can do
      ()
    done;
  (match get_food pond with
  | Some fish ->
    Printf.printf "Alice: Releasing pets to eat %s\n"
      (fish_to_string fish);
    Unix.sleepf 0.1;
    Printf.printf "Alice: Recapturing pets\n");
  | None -> failwith "impossible");
  reset can
done
```

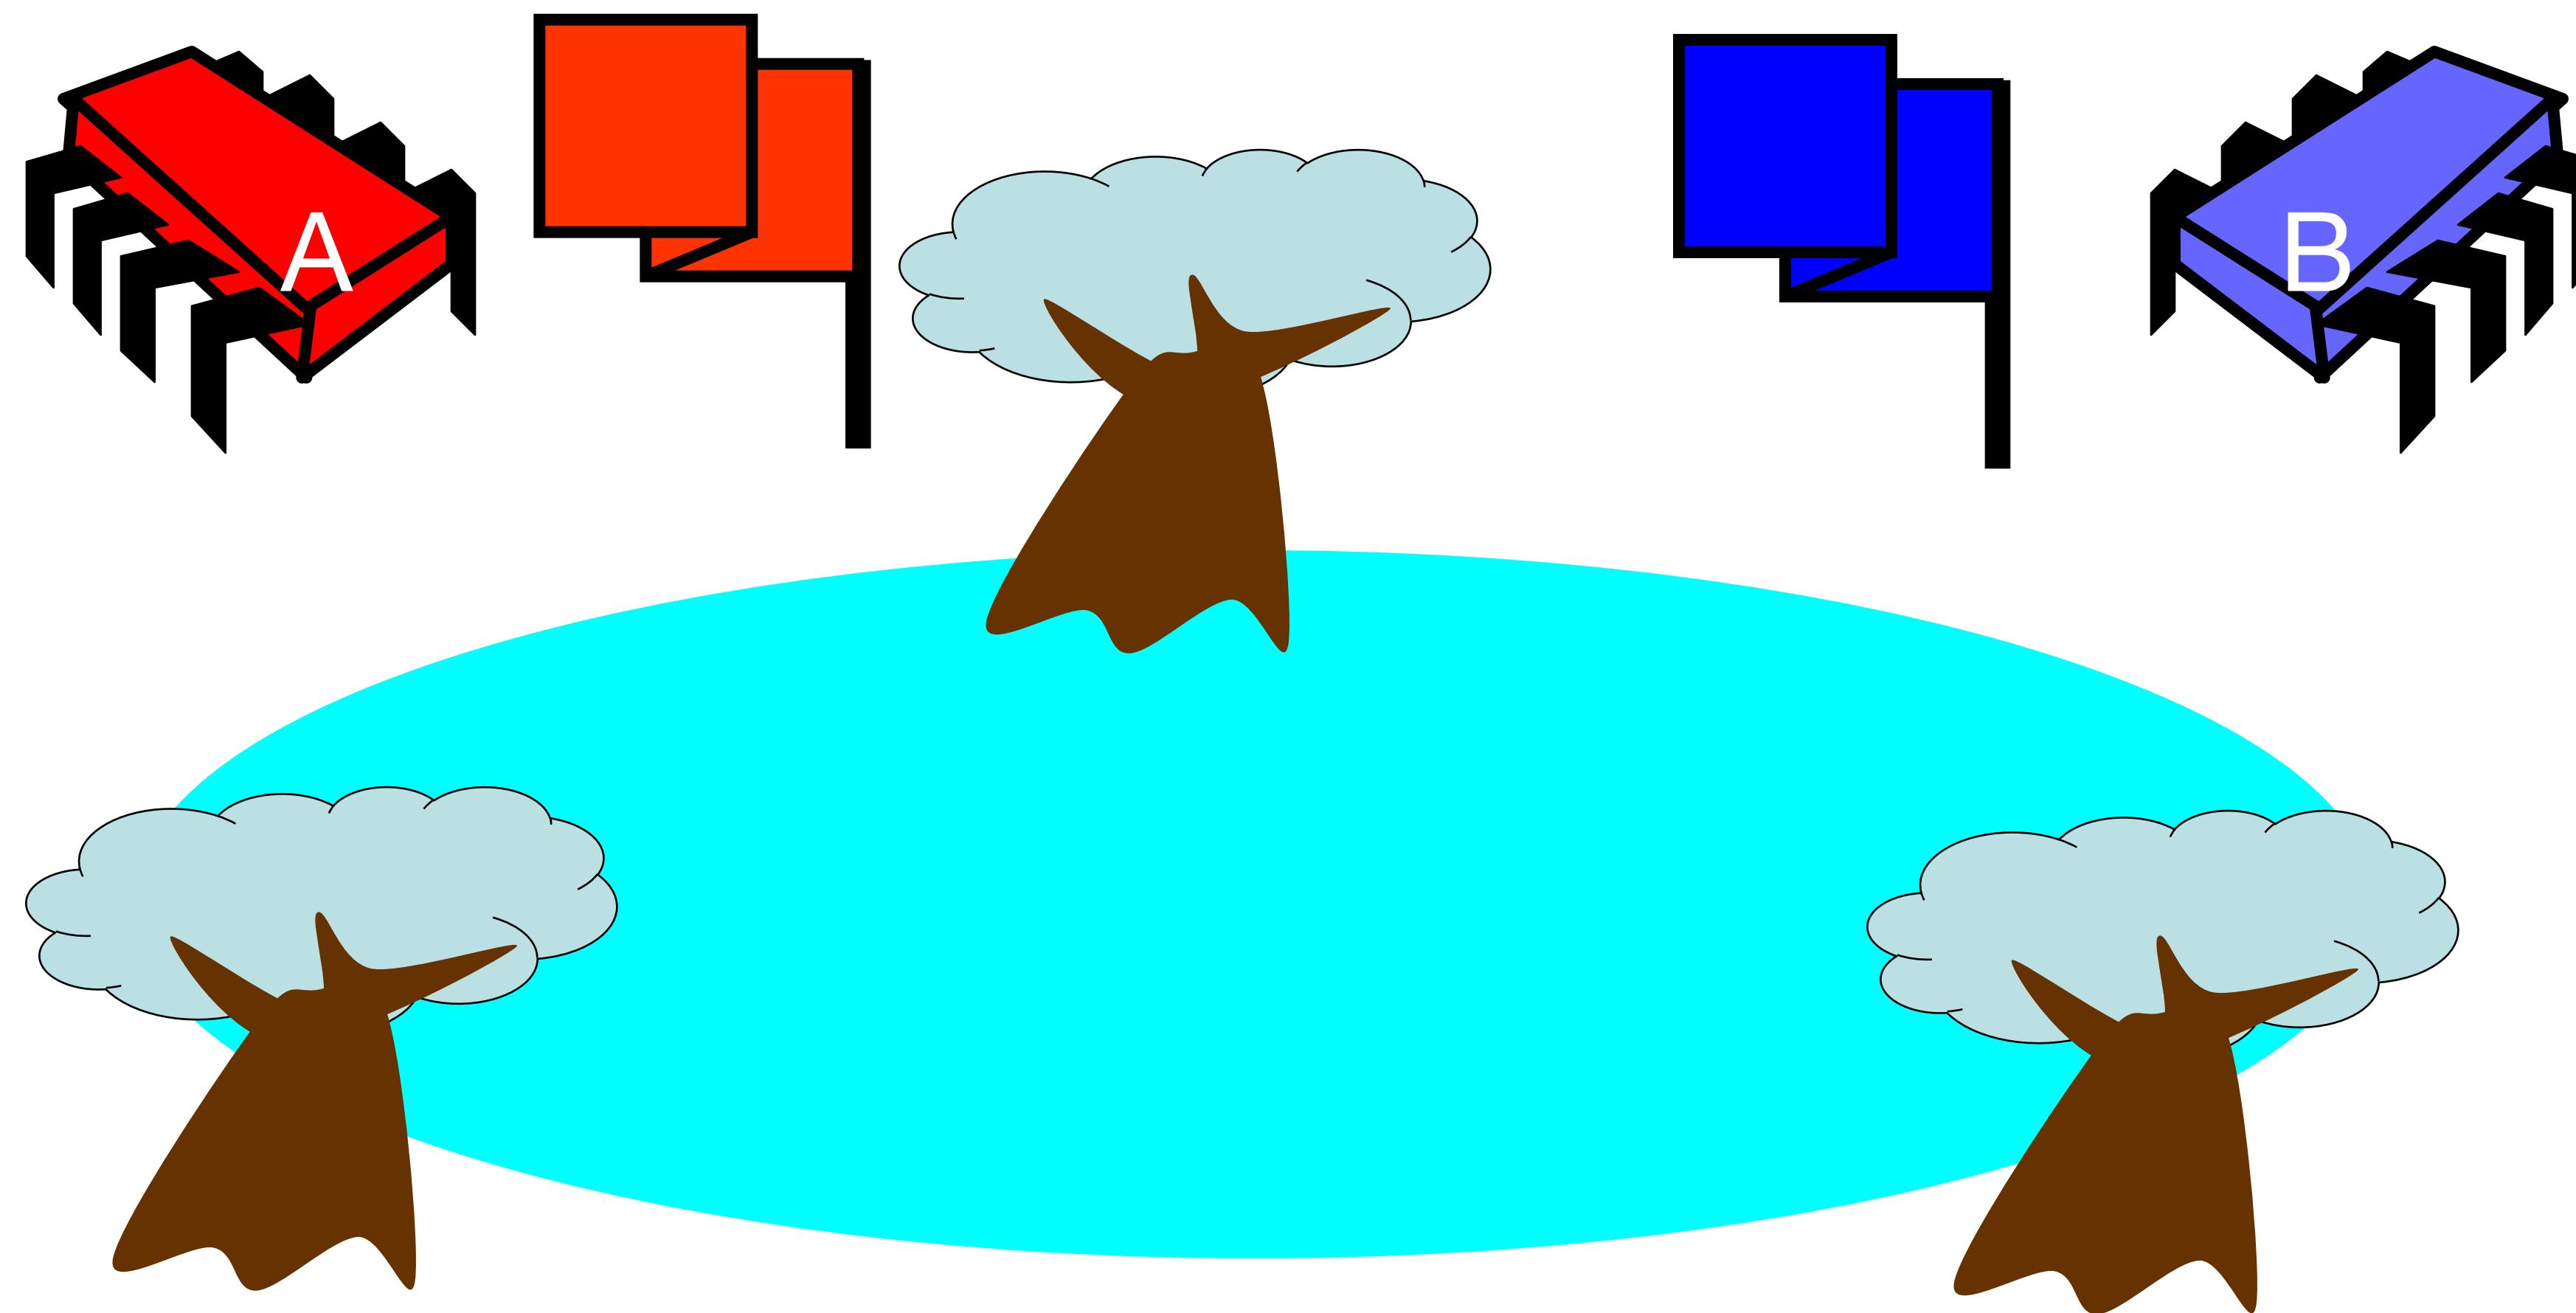
```
(* Bob (Producer) - stocks the pond with fish *)
let bob can pond =
  while true do
    while is_down can do
      ()
    done;
    let fish = random_fish () in
    Printf.printf "Bob: Stocking pond with %s\n"
      (fish_to_string fish);
    stock_pond pond fish;
    knock_over can
  done
```

Demo

# Correctness

- **Mutual Exclusion** 
  - Pets and Bob are never together in the pond
- **No Starvation** 
  - If Bob is always willing to feed, and the pets are always famished, then the pets eat infinitely often.
- **Producer/Consumer** 
  - The pets never enter the pond unless there is food, and Bob never provides food if there is unconsumed food.
  - Which of the above are **safety** and **liveness** properties?

# Could also solve using flags



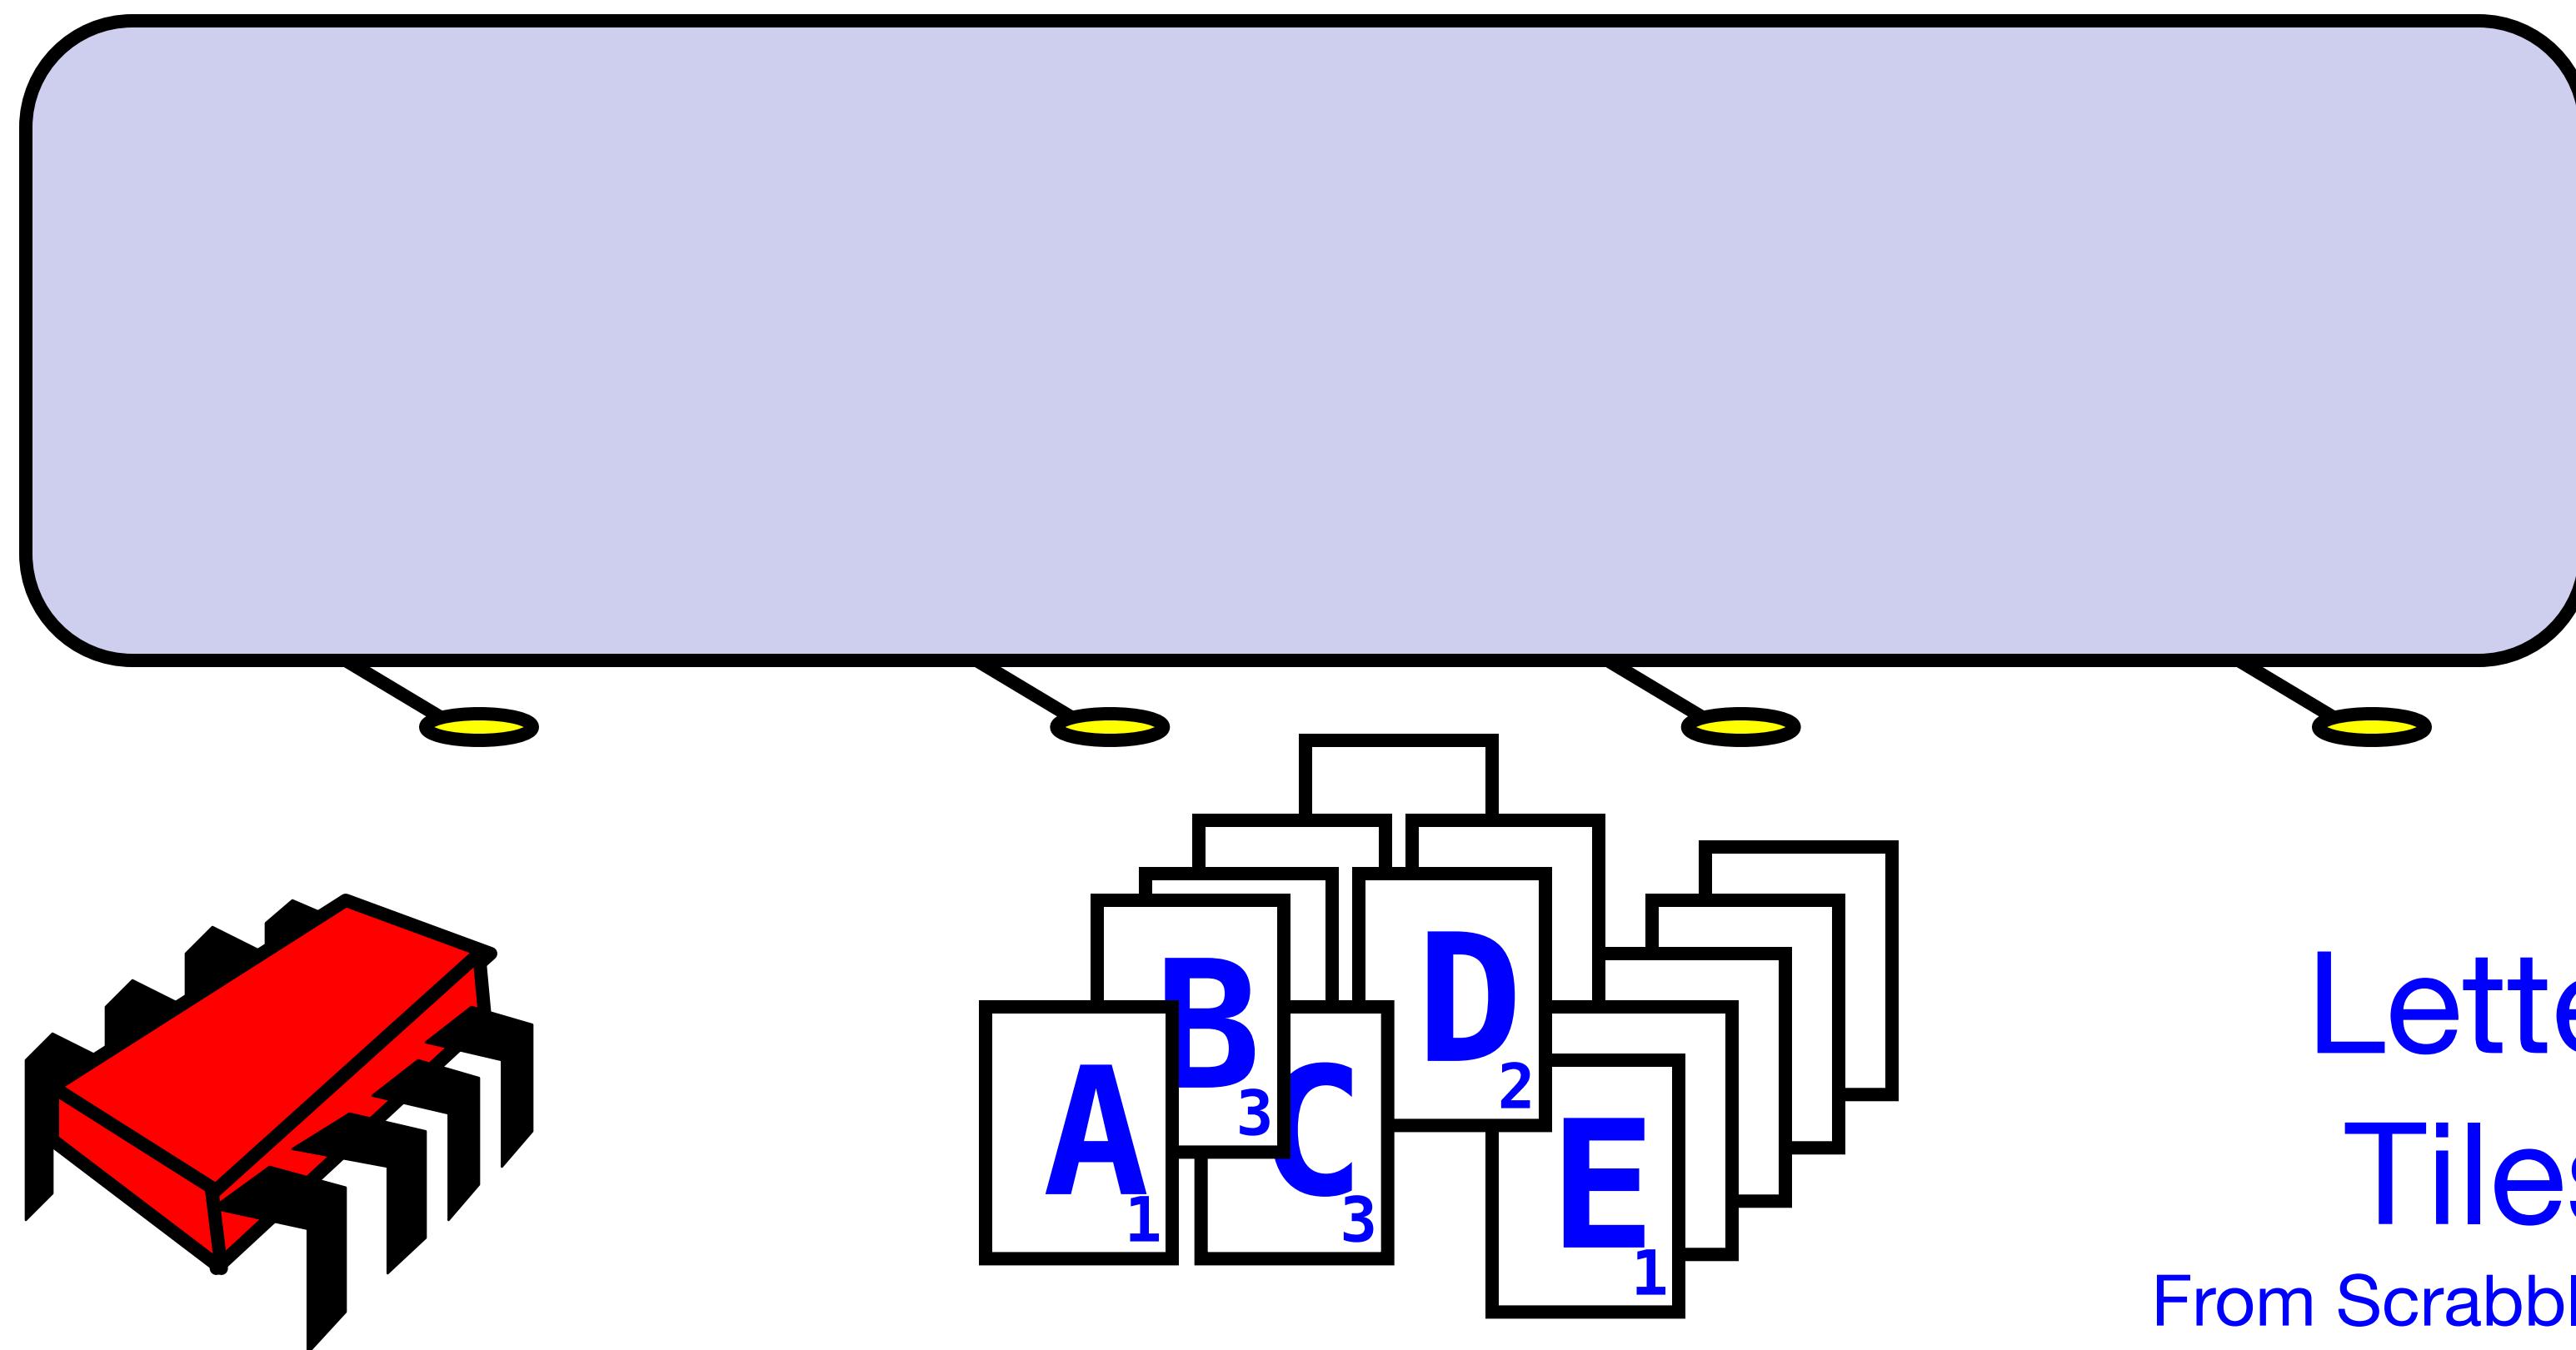
# Waiting

- Both solutions use ***waiting***
  - while mumble do ... done
- In some cases waiting is ***problematic***
  - If one participant is delayed
  - So is everyone else
  - But delays are common & unpredictable

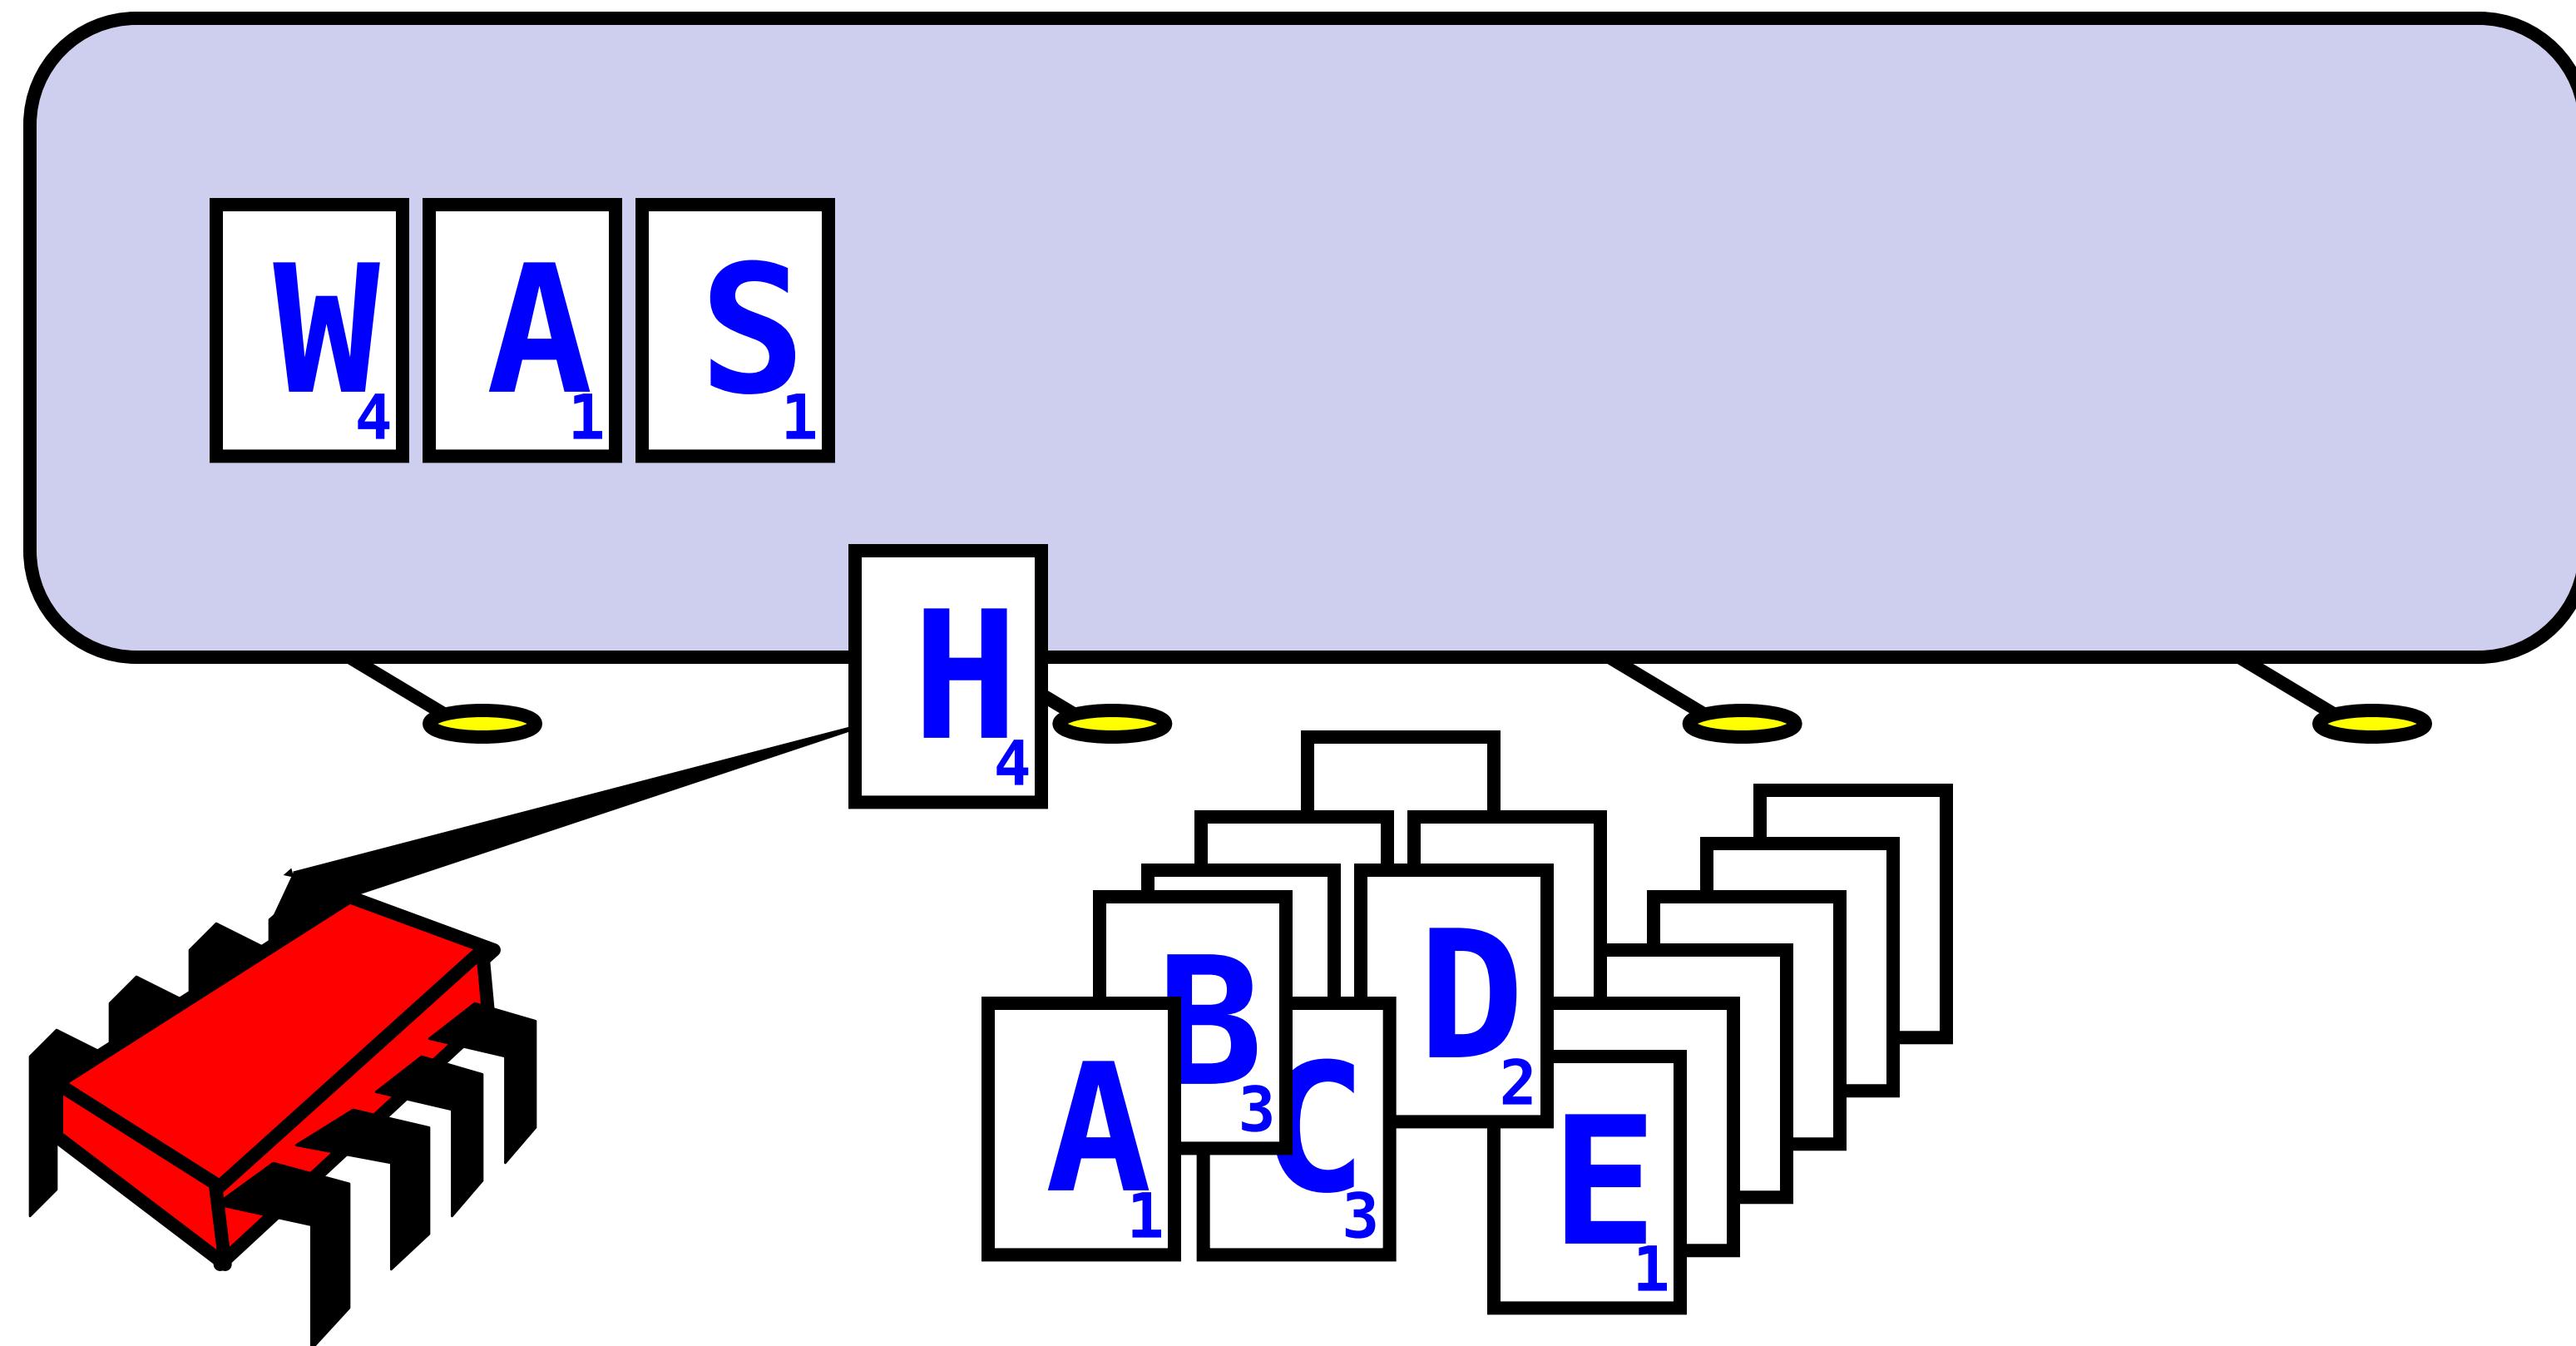
# The fable drags on ...

- Bob and Alice still have issues
- So they need to communicate
- They agree to use billboards ...

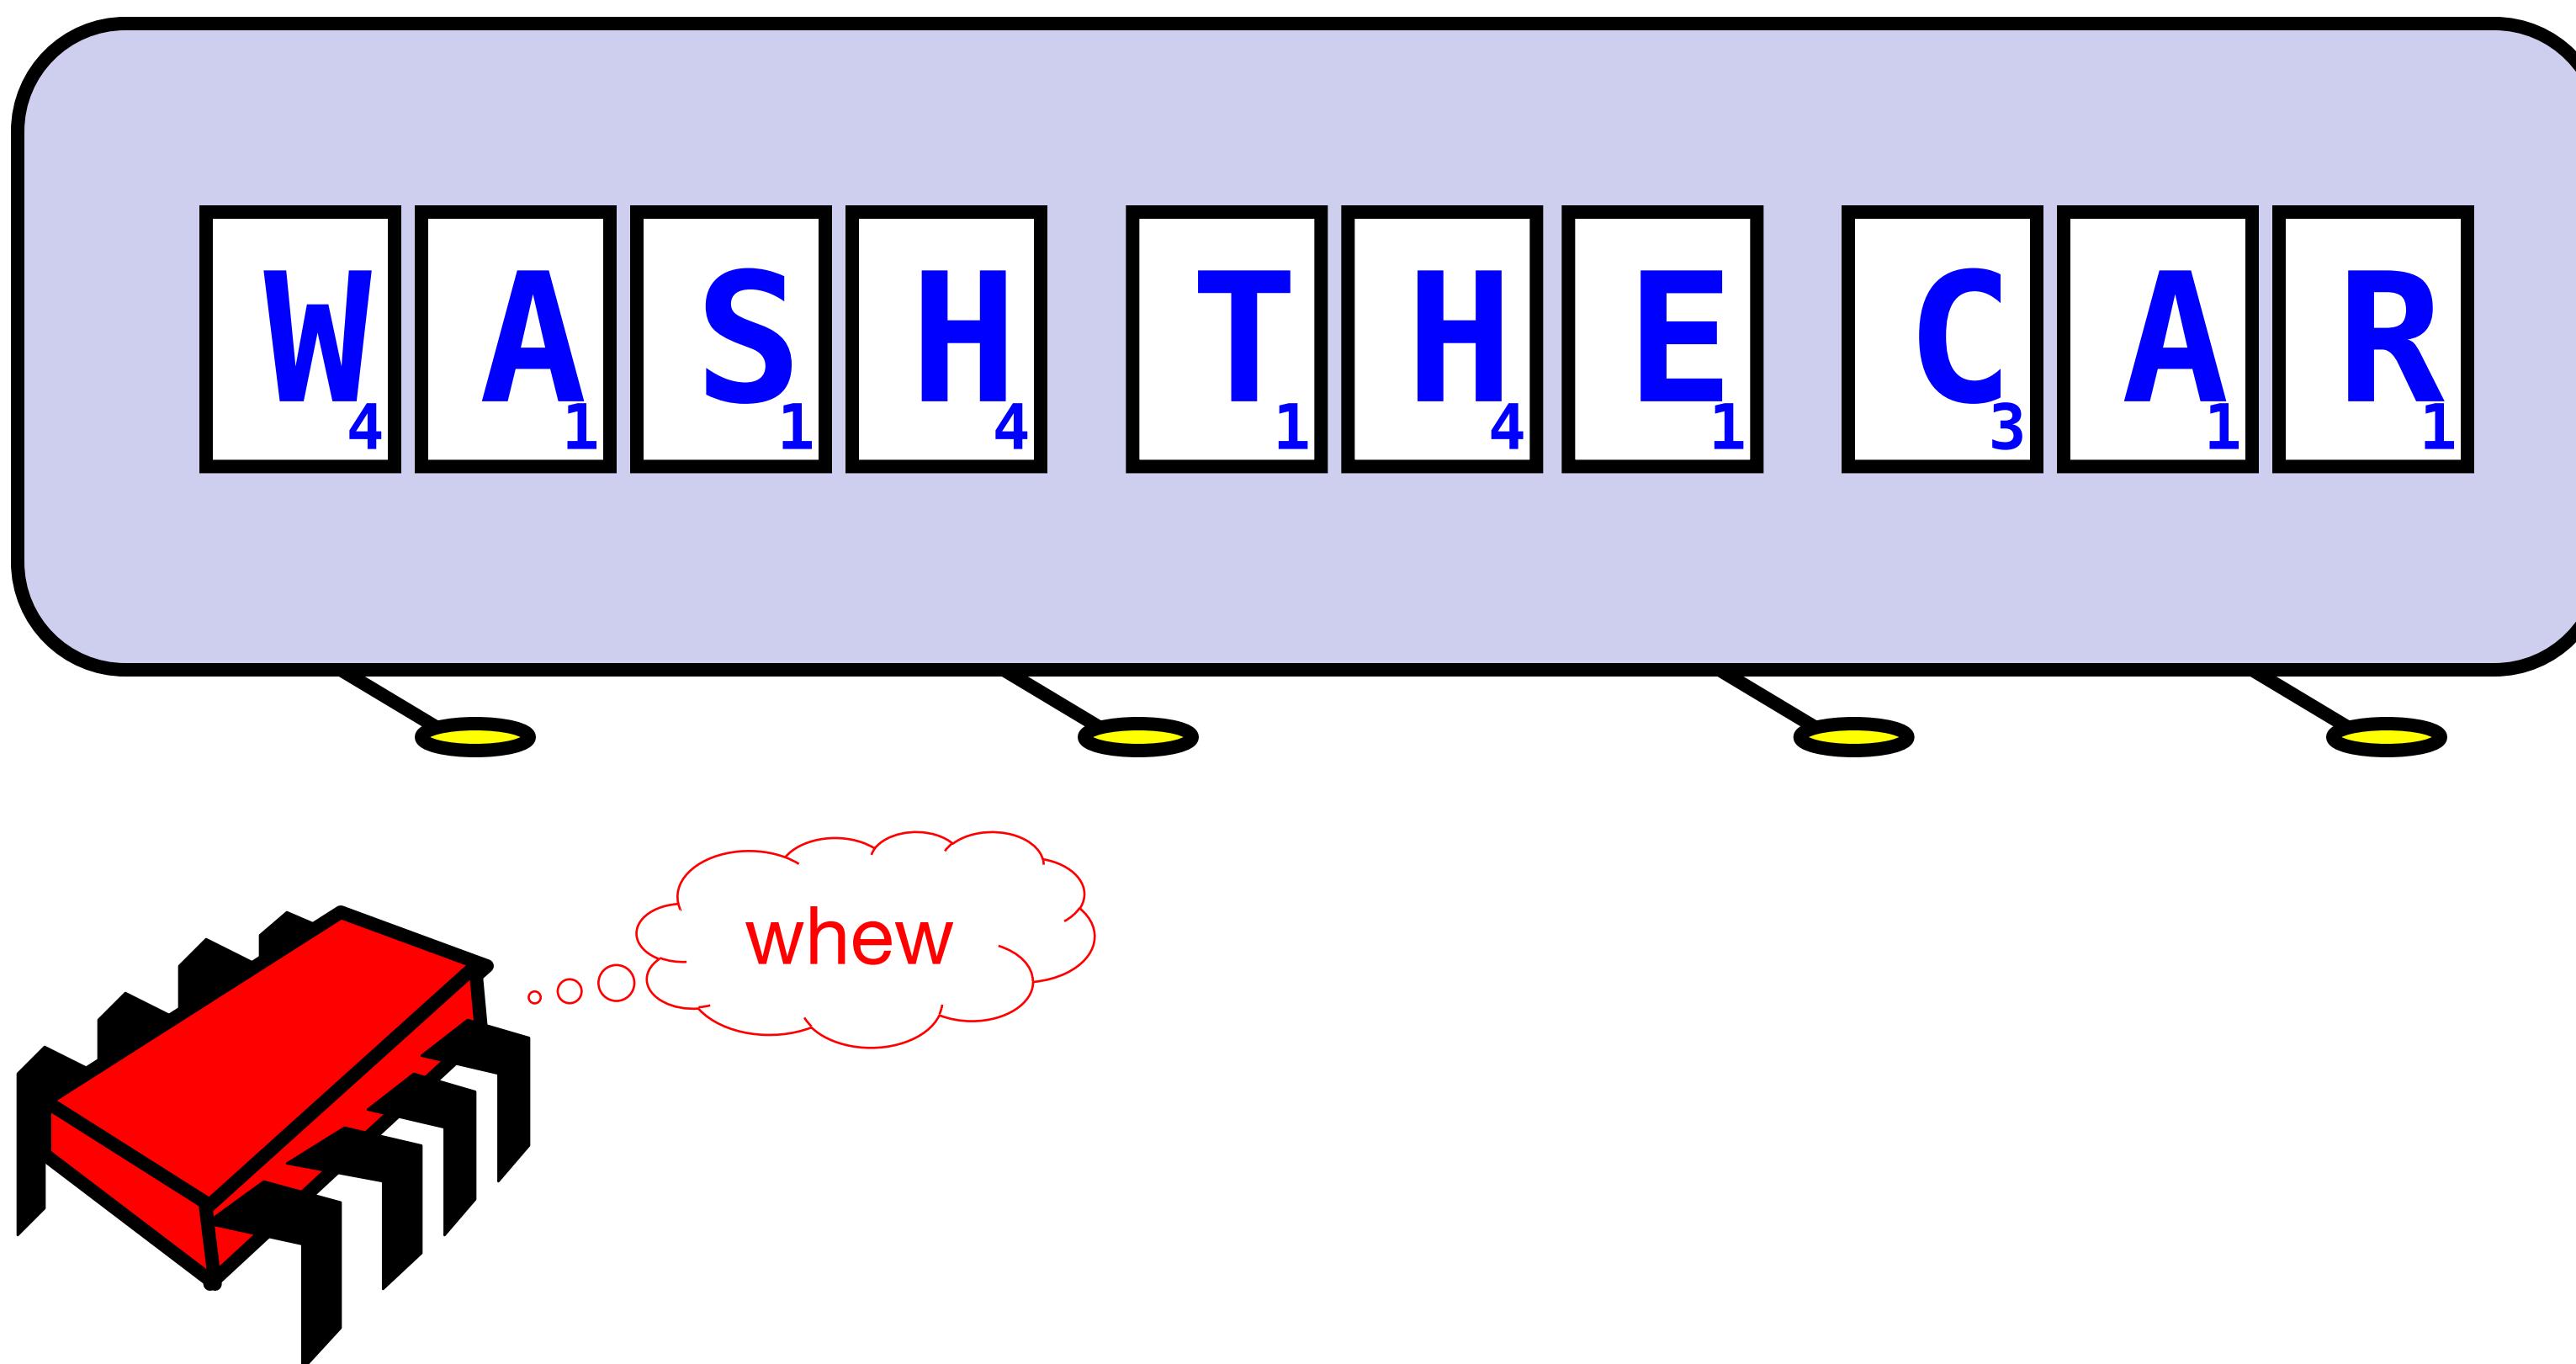
# Billboards are large



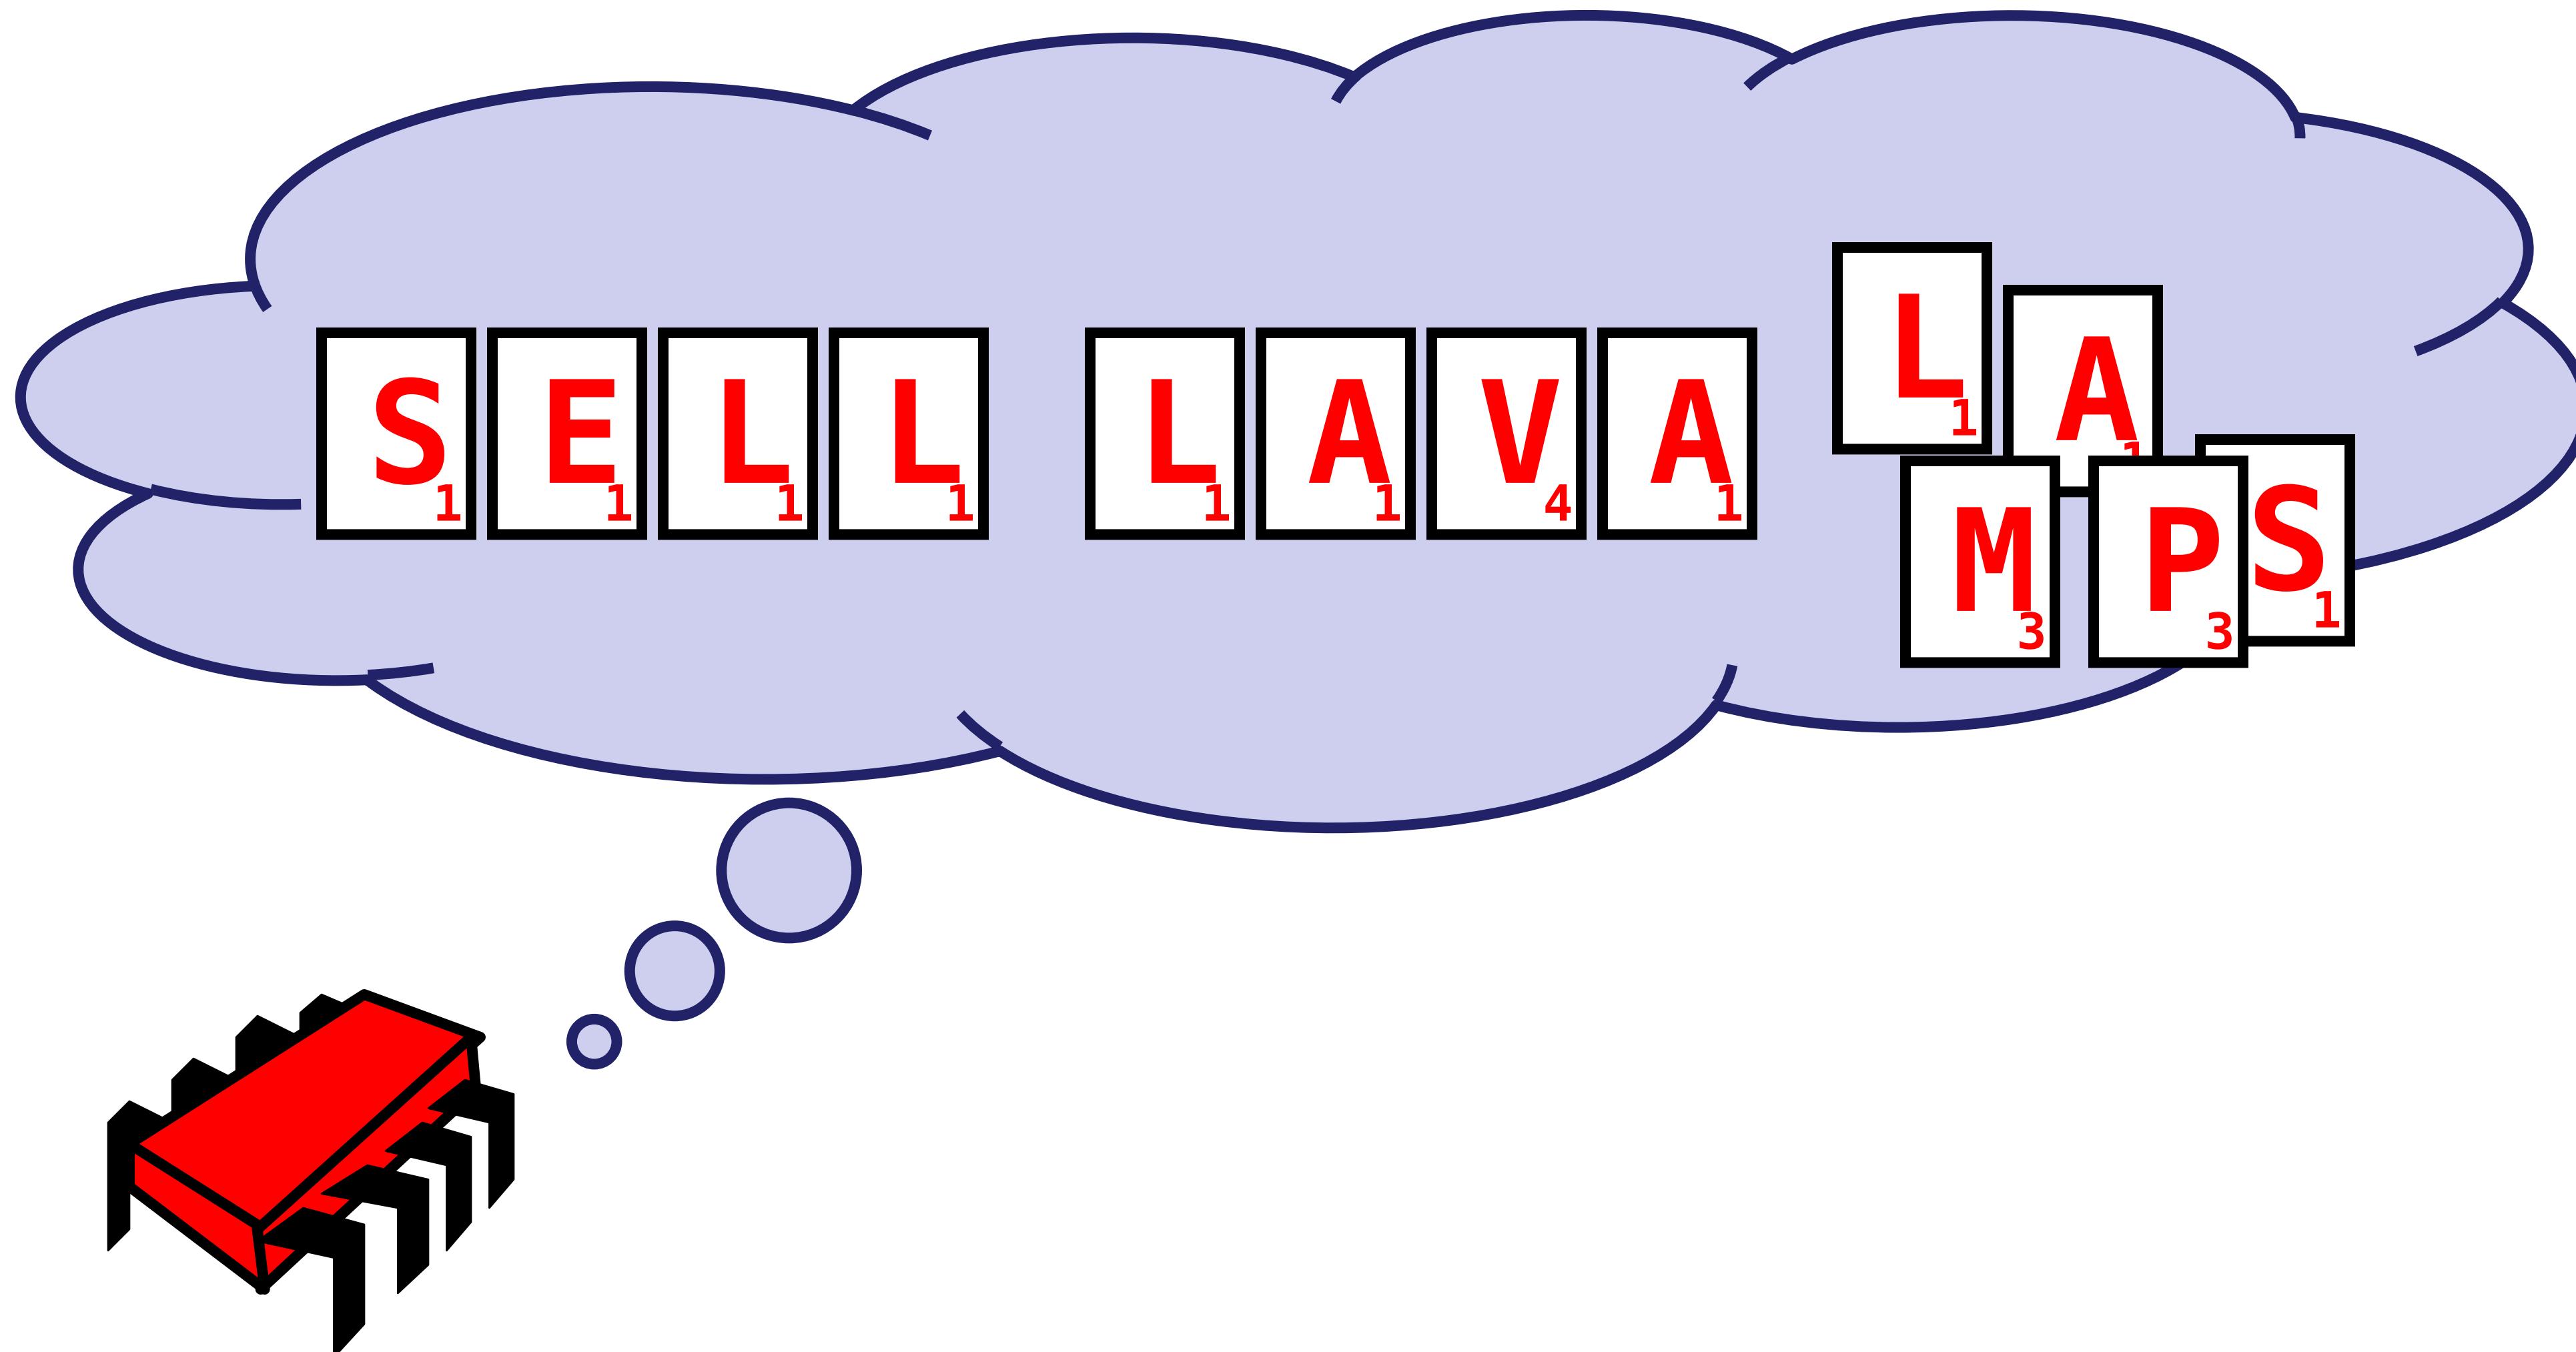
# Write one letter at a time



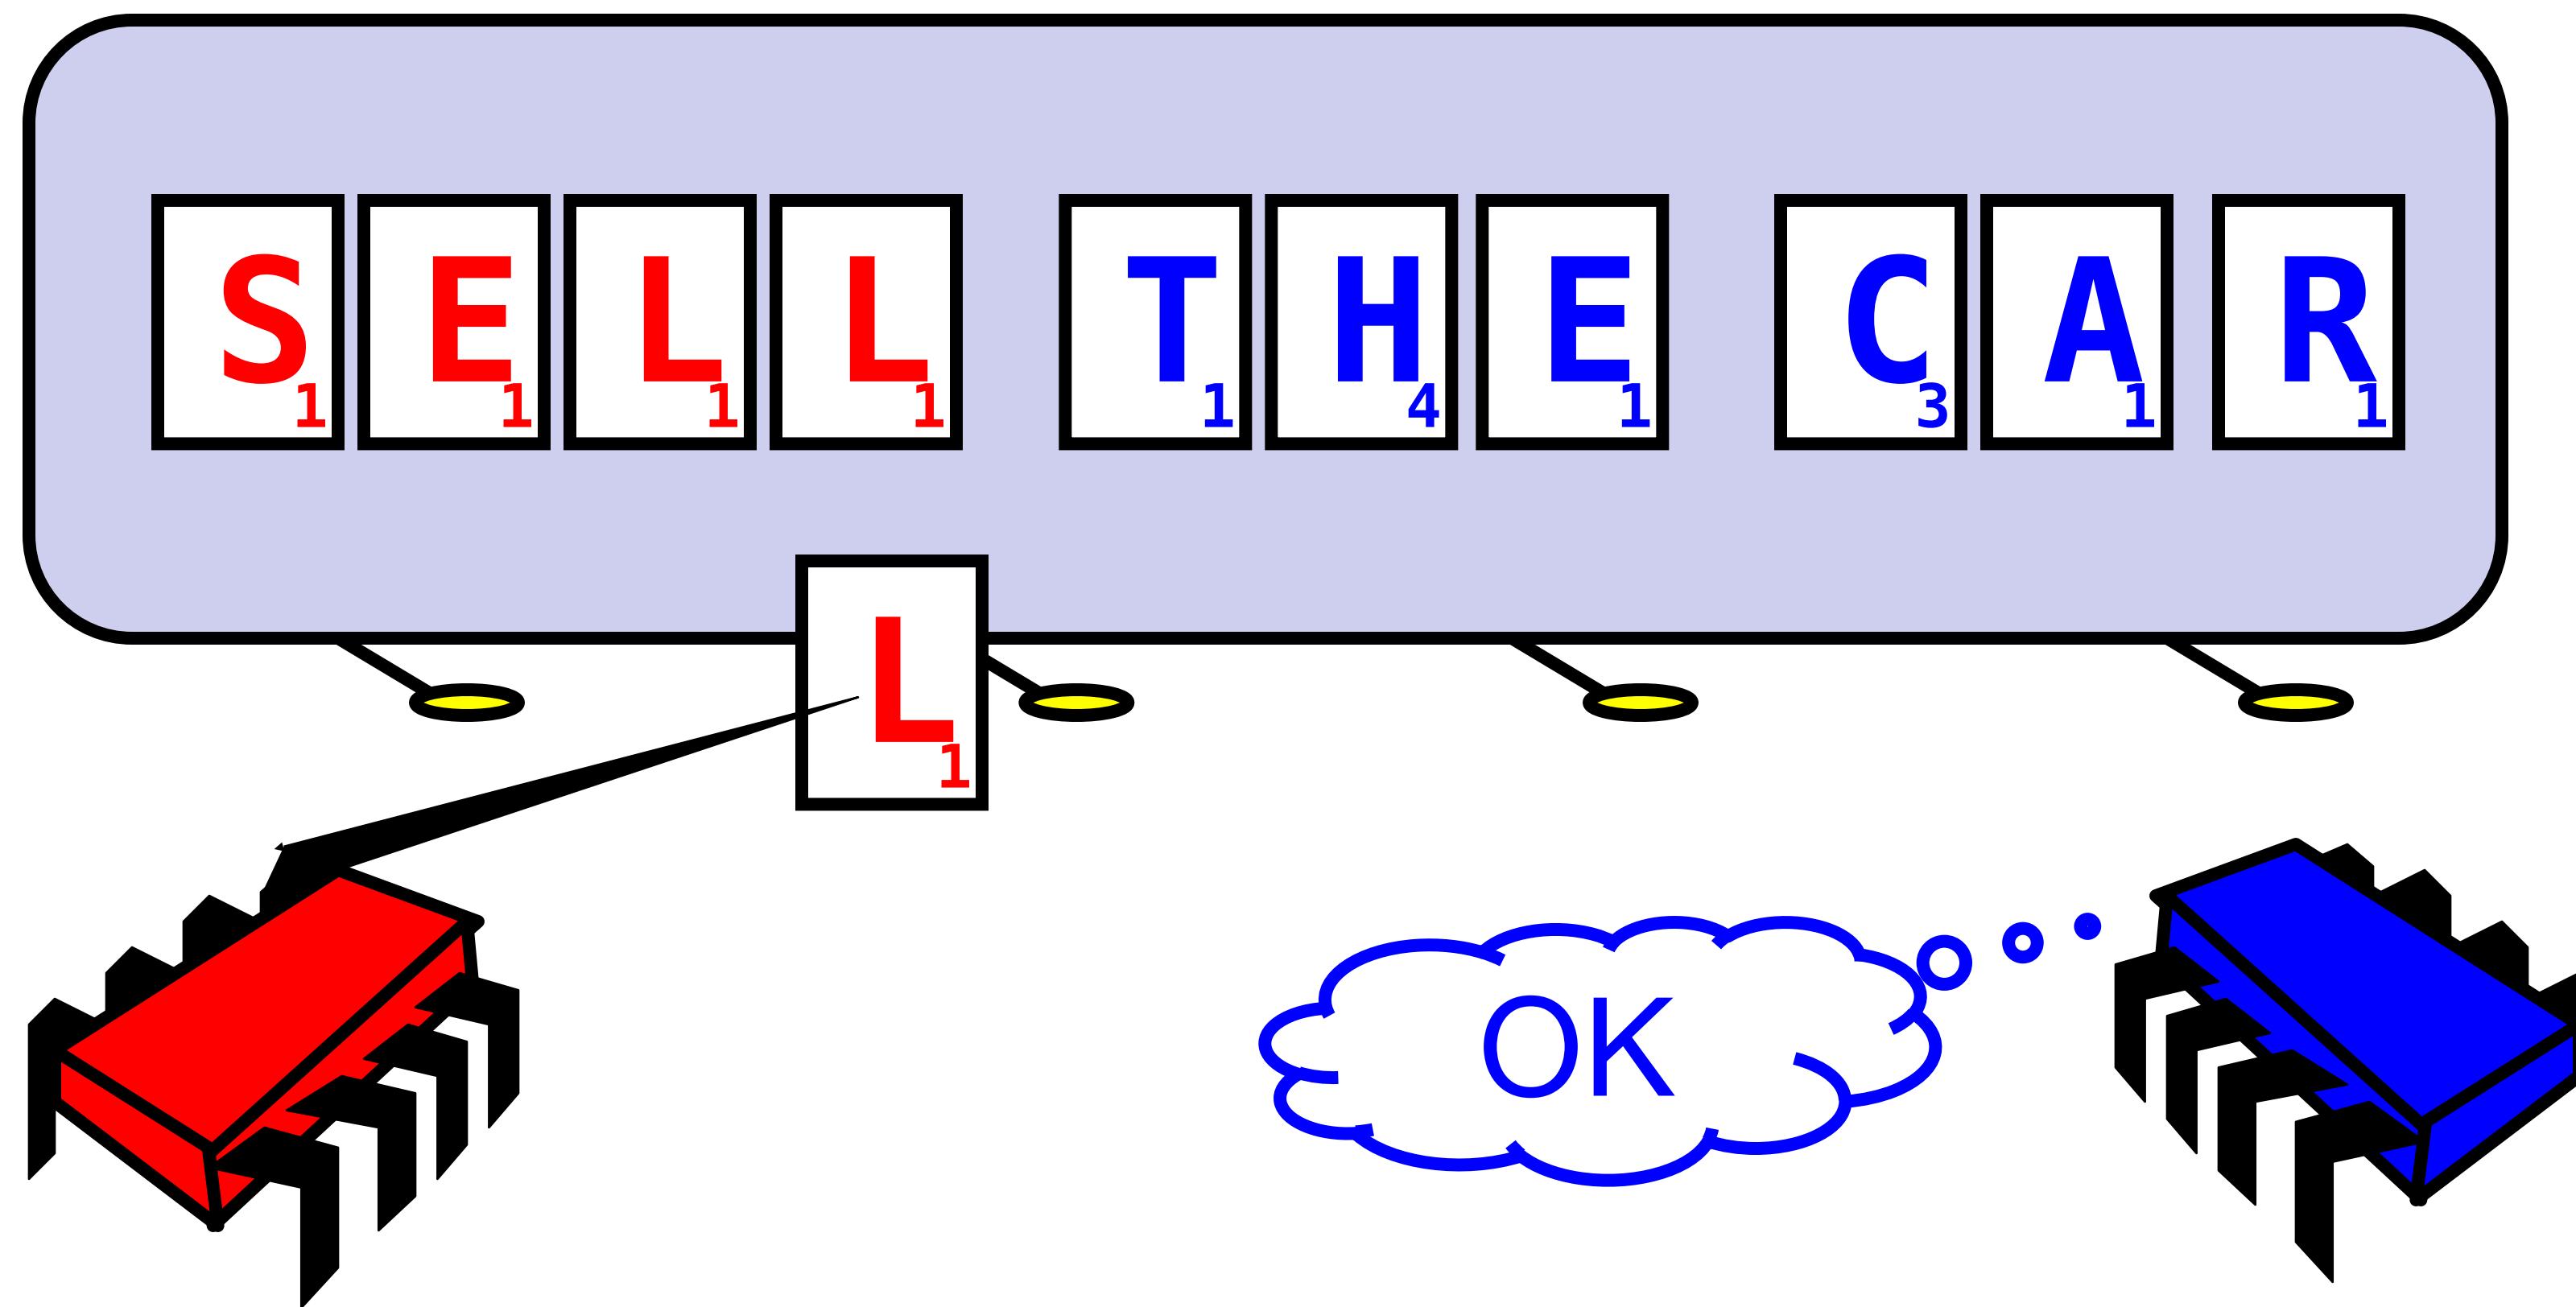
# To post a message



# Let's send another message



# Uh-Oh



# Readers/Writers

- Devise a protocol so that
  - Writer writes one letter at a time
  - Reader reads one letter at a time
  - Reader sees “snapshot”
    - Old message or new message
    - No mixed messages
- Easy with ***mutual exclusion***
- But mutual exclusion requires ***waiting***
  - One waits for the other
  - Everyone executes sequentially
- Remarkably
  - ***We can solve R/W without mutual exclusion***
  - We will see these in later lectures

# Esoteric?

- Assume that you have a ***multithreading-safe, balanced, binary tree*** implementation
- Observe that multiple concurrent reads can be done without synchronisation
- Even update operations, **add** or **delete**, may modify disjoint parts of the tree
  - May be done in parallel without blocking each other
- You want a **size** function to return the number of elements in the tree.
  - Every update operation needs to modify **size**
  - Threads need to wait for exclusive access to the counter!

performance  
bottleneck

# Size() Readers/Writers Solution

- Maintain a global array `count[i]`, indexed by the thread ID
- Each thread `i` keeps a count of  $| \text{add} | - | \text{remove} |$  in the `count` array
- Function `size()` now only needs to read a “***consistent snapshot***” of the `count` array
  - This eliminates the bottleneck

# Why do we care?

- We want as much of the code as possible to execute concurrently (in parallel)
- A larger sequential part implies reduced performance
- **Amdahl's law:** this relation is not linear...

# Amdahl's Law

**Speedup =**

$$\frac{1\text{-thread execution time}}{n\text{-thread execution time}}$$

# Amdahl's Law

**Speedup =** 
$$\frac{1}{1 - p + \frac{p}{n}}$$

# Amdahl's Law

Speedup =

$$\frac{1}{1 - p + \frac{p}{n}}$$

Parallel fraction

# Amdahl's Law

Sequential fraction

Parallel fraction

Speedup =  $\frac{1}{1 - p + \frac{p}{n}}$

The diagram illustrates the formula for Speedup in Amdahl's Law. It features a red line graph starting at the top left, sloping down to the right, and then leveling off. A horizontal black line extends from the point where the red line levels off. Two red arrows point from the text labels 'Sequential fraction' and 'Parallel fraction' to the terms  $1 - p$  and  $\frac{p}{n}$  respectively in the formula below. The term  $1$  is also labeled above the horizontal line.

# Amdahl's Law

$$\text{Speedup} = \frac{\text{Parallel fraction}}{\text{Sequential fraction} + \frac{1 - p}{n}}$$

Diagram illustrating Amdahl's Law:

- Sequential fraction**: Represented by a red line that decreases from 1 (at 0 threads) towards 0.
- Parallel fraction**: Represented by a horizontal black line at  $y = p$ .
- Number of threads**: Represented by a red line that increases from 0 towards infinity.
- The formula shows the speedup as the ratio of the parallel fraction to the sum of the sequential fraction and the term  $\frac{1-p}{n}$ .

# Amdal's Law



Bad synchronization ruins everything

# Example

- Ten processors
- 60% concurrent, 40% sequential
- ***How close to 10-fold speedup?***

$$\text{Speedup} = 2.17 = \frac{1}{1 - 0.6 + \frac{0.6}{10}}$$

# Example

- Ten processors
- 80% concurrent, 20% sequential
- ***How close to 10-fold speedup?***

$$\text{Speedup} = 3.57 = \frac{1}{1 - 0.8 + \frac{0.8}{10}}$$

# Example

- Ten processors
- 90% concurrent, 10% sequential
- ***How close to 10-fold speedup?***

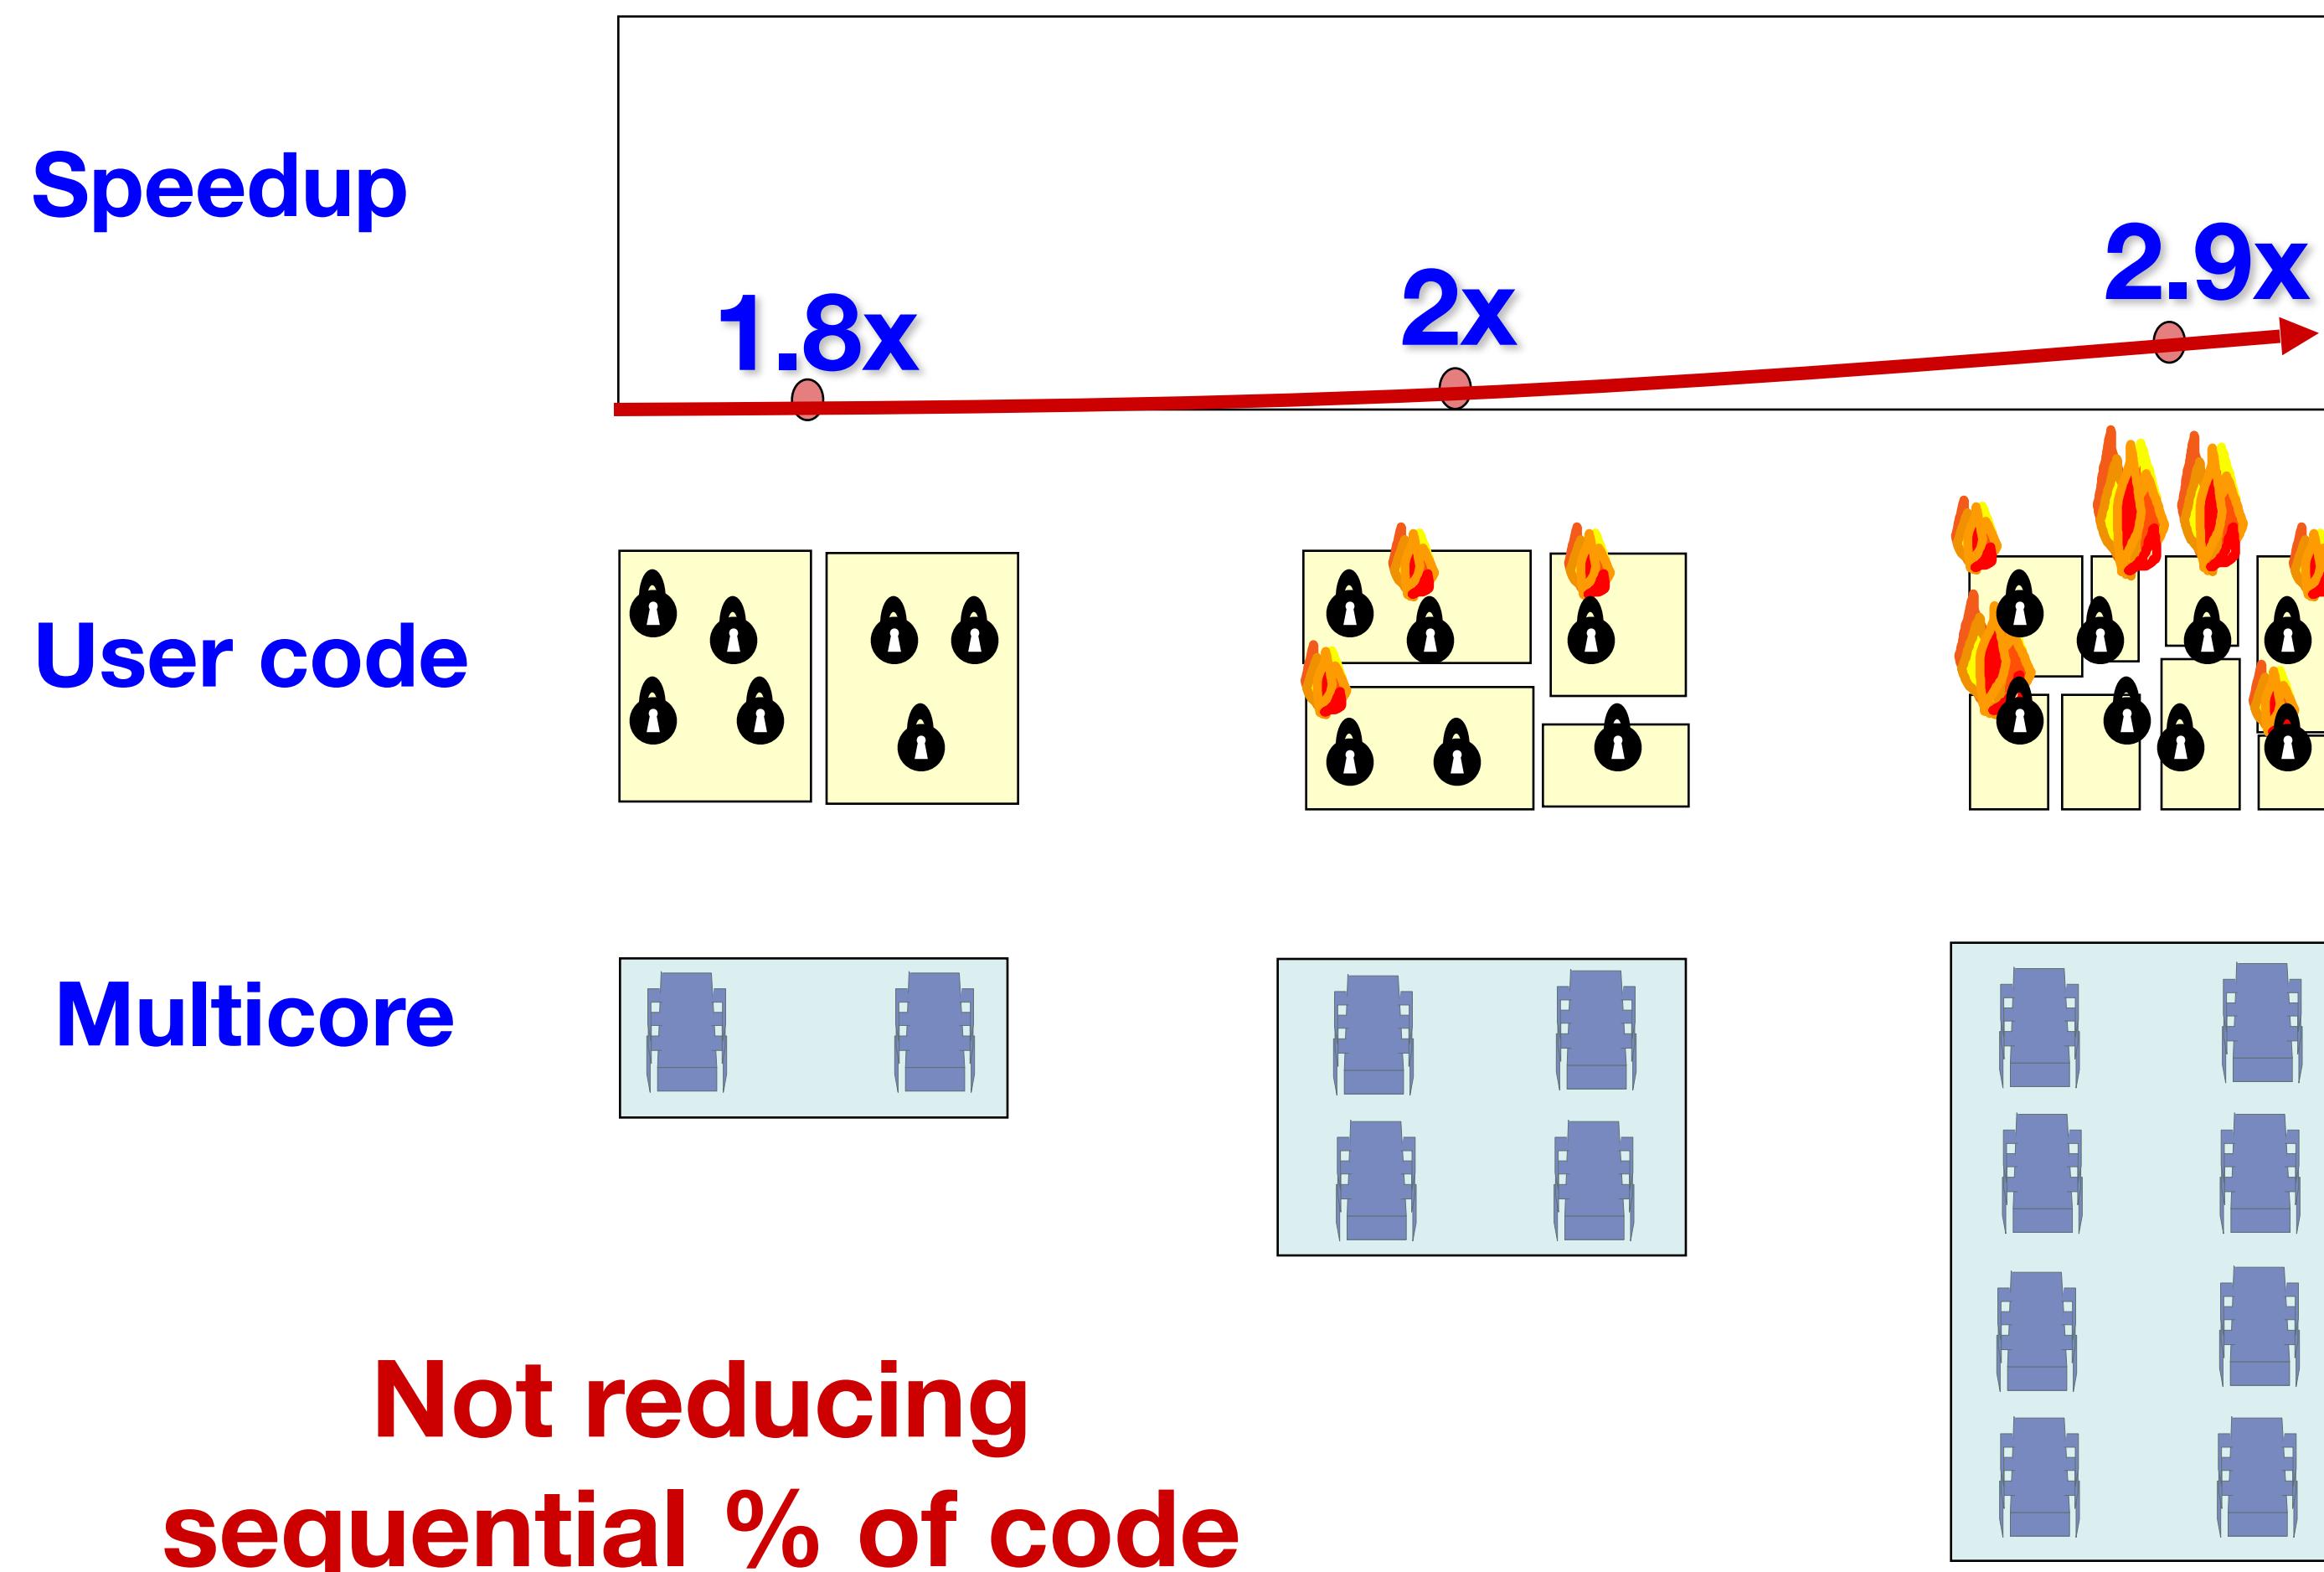
$$\text{Speedup} = 5.26 = \frac{1}{1 - 0.9 + \frac{0.9}{10}}$$

# Example

- Ten processors
- 99% concurrent, 1% sequential
- ***How close to 10-fold speedup?***

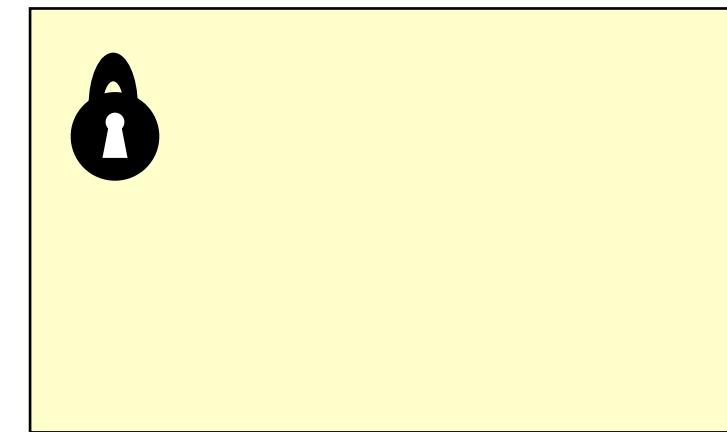
$$\text{Speedup} = 9.17 = \frac{1}{1 - 0.99 + \frac{0.99}{10}}$$

# Back to real-world multicore scaling

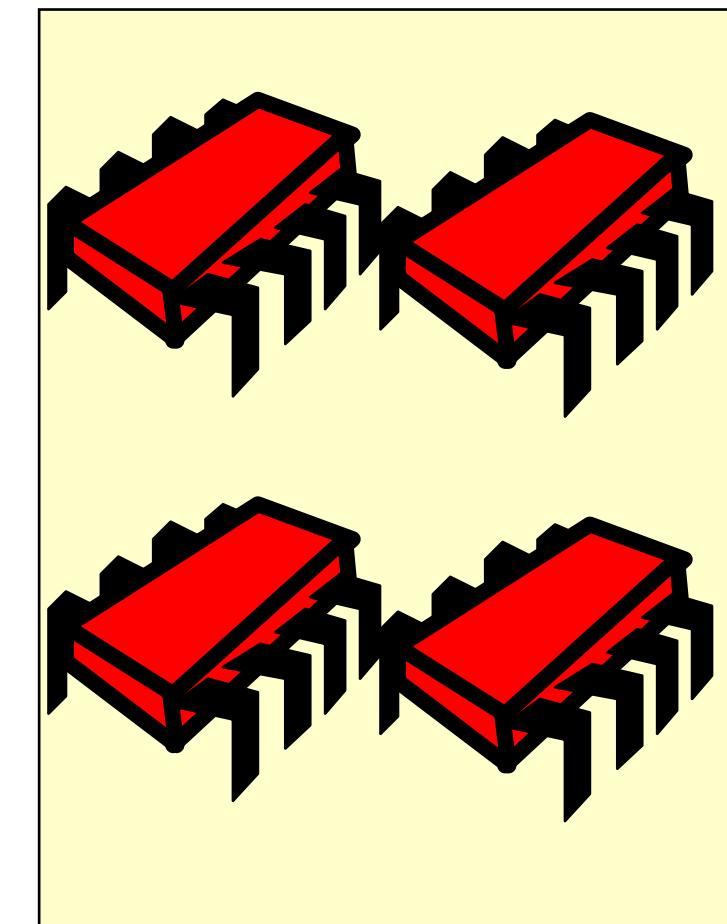


# Shared Data Structures

**Coarse  
Grained**

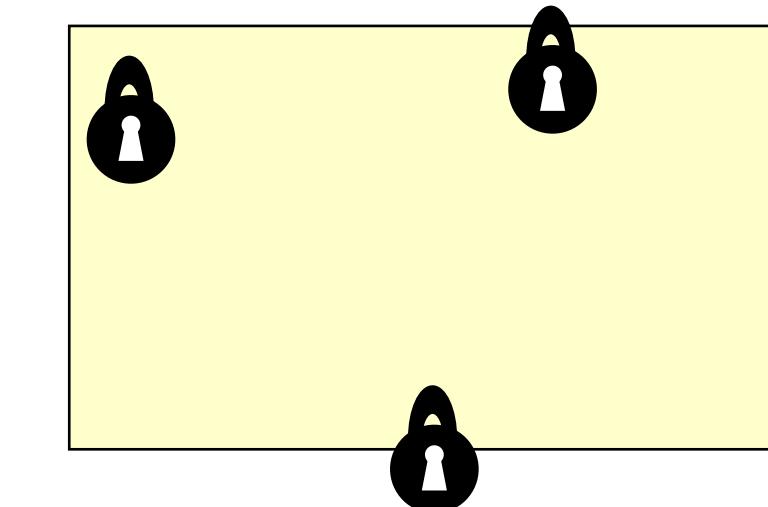


25%  
Shared

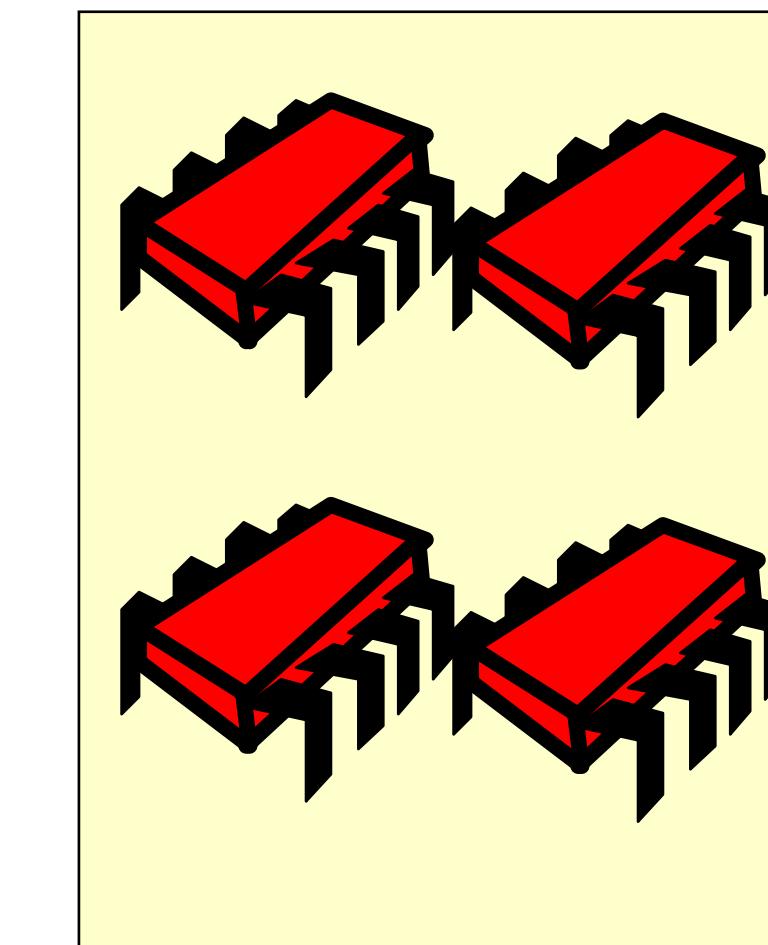


75%  
Unshared

**Fine  
Grained**

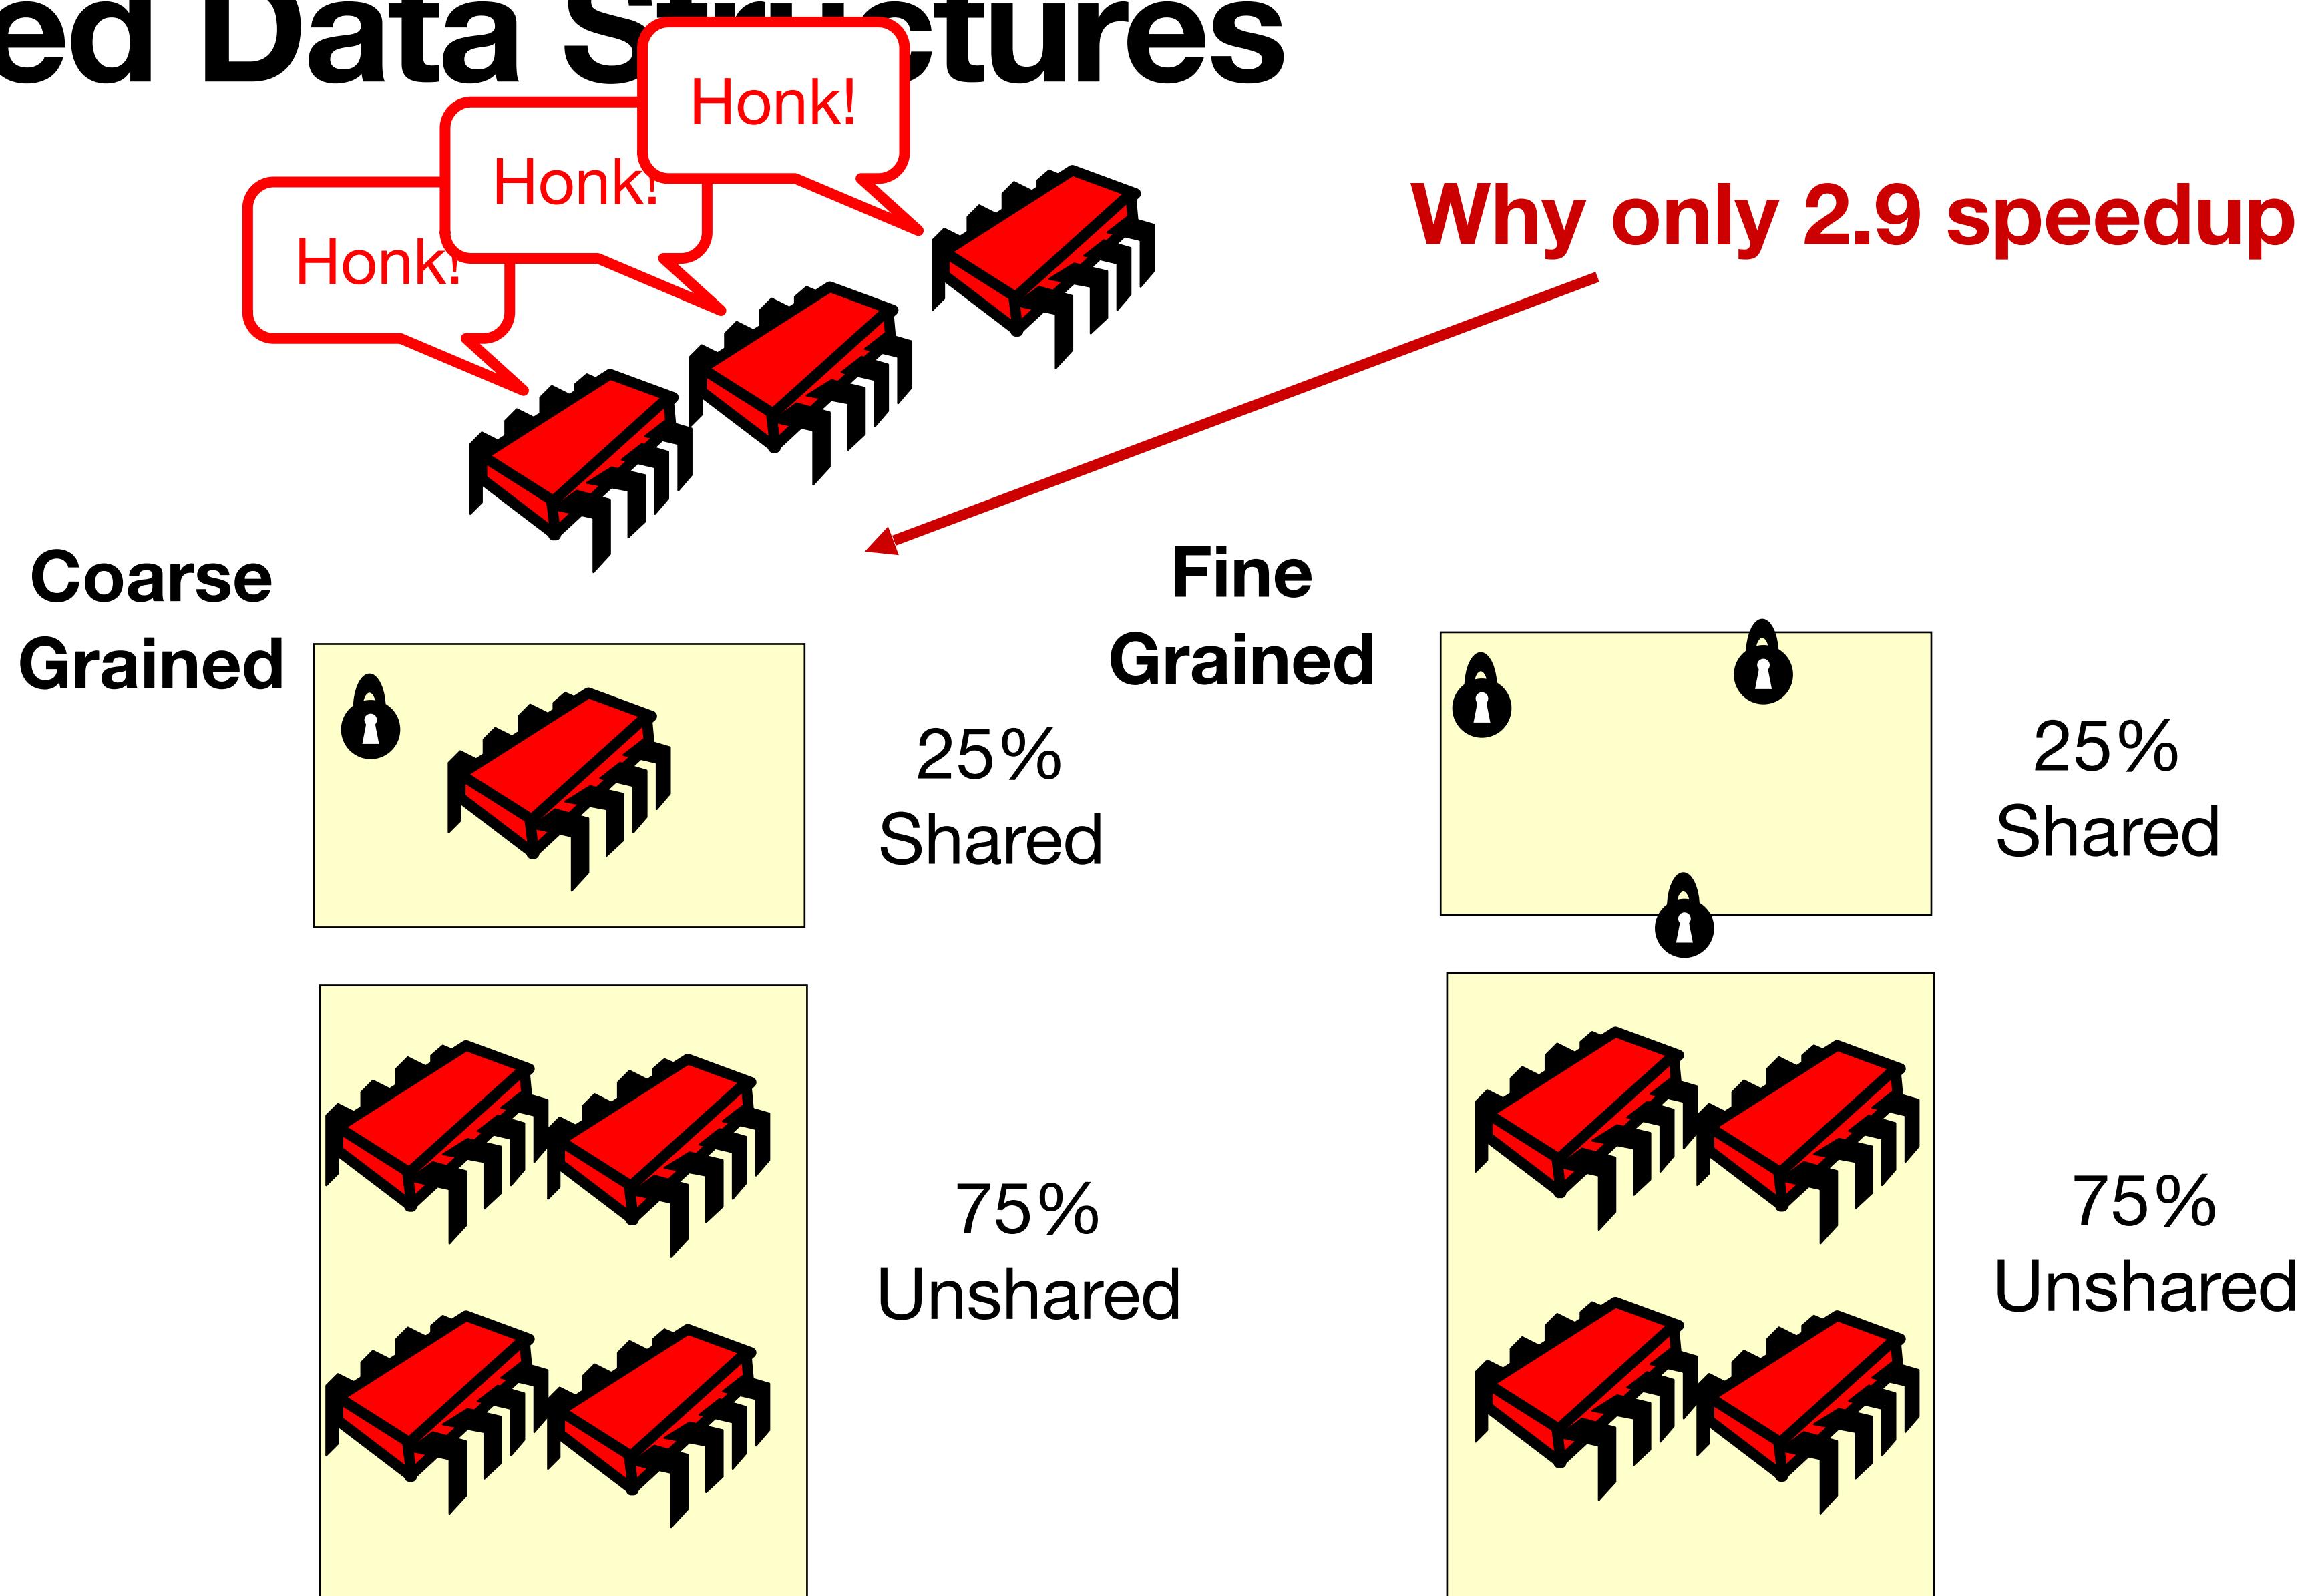


25%  
Shared

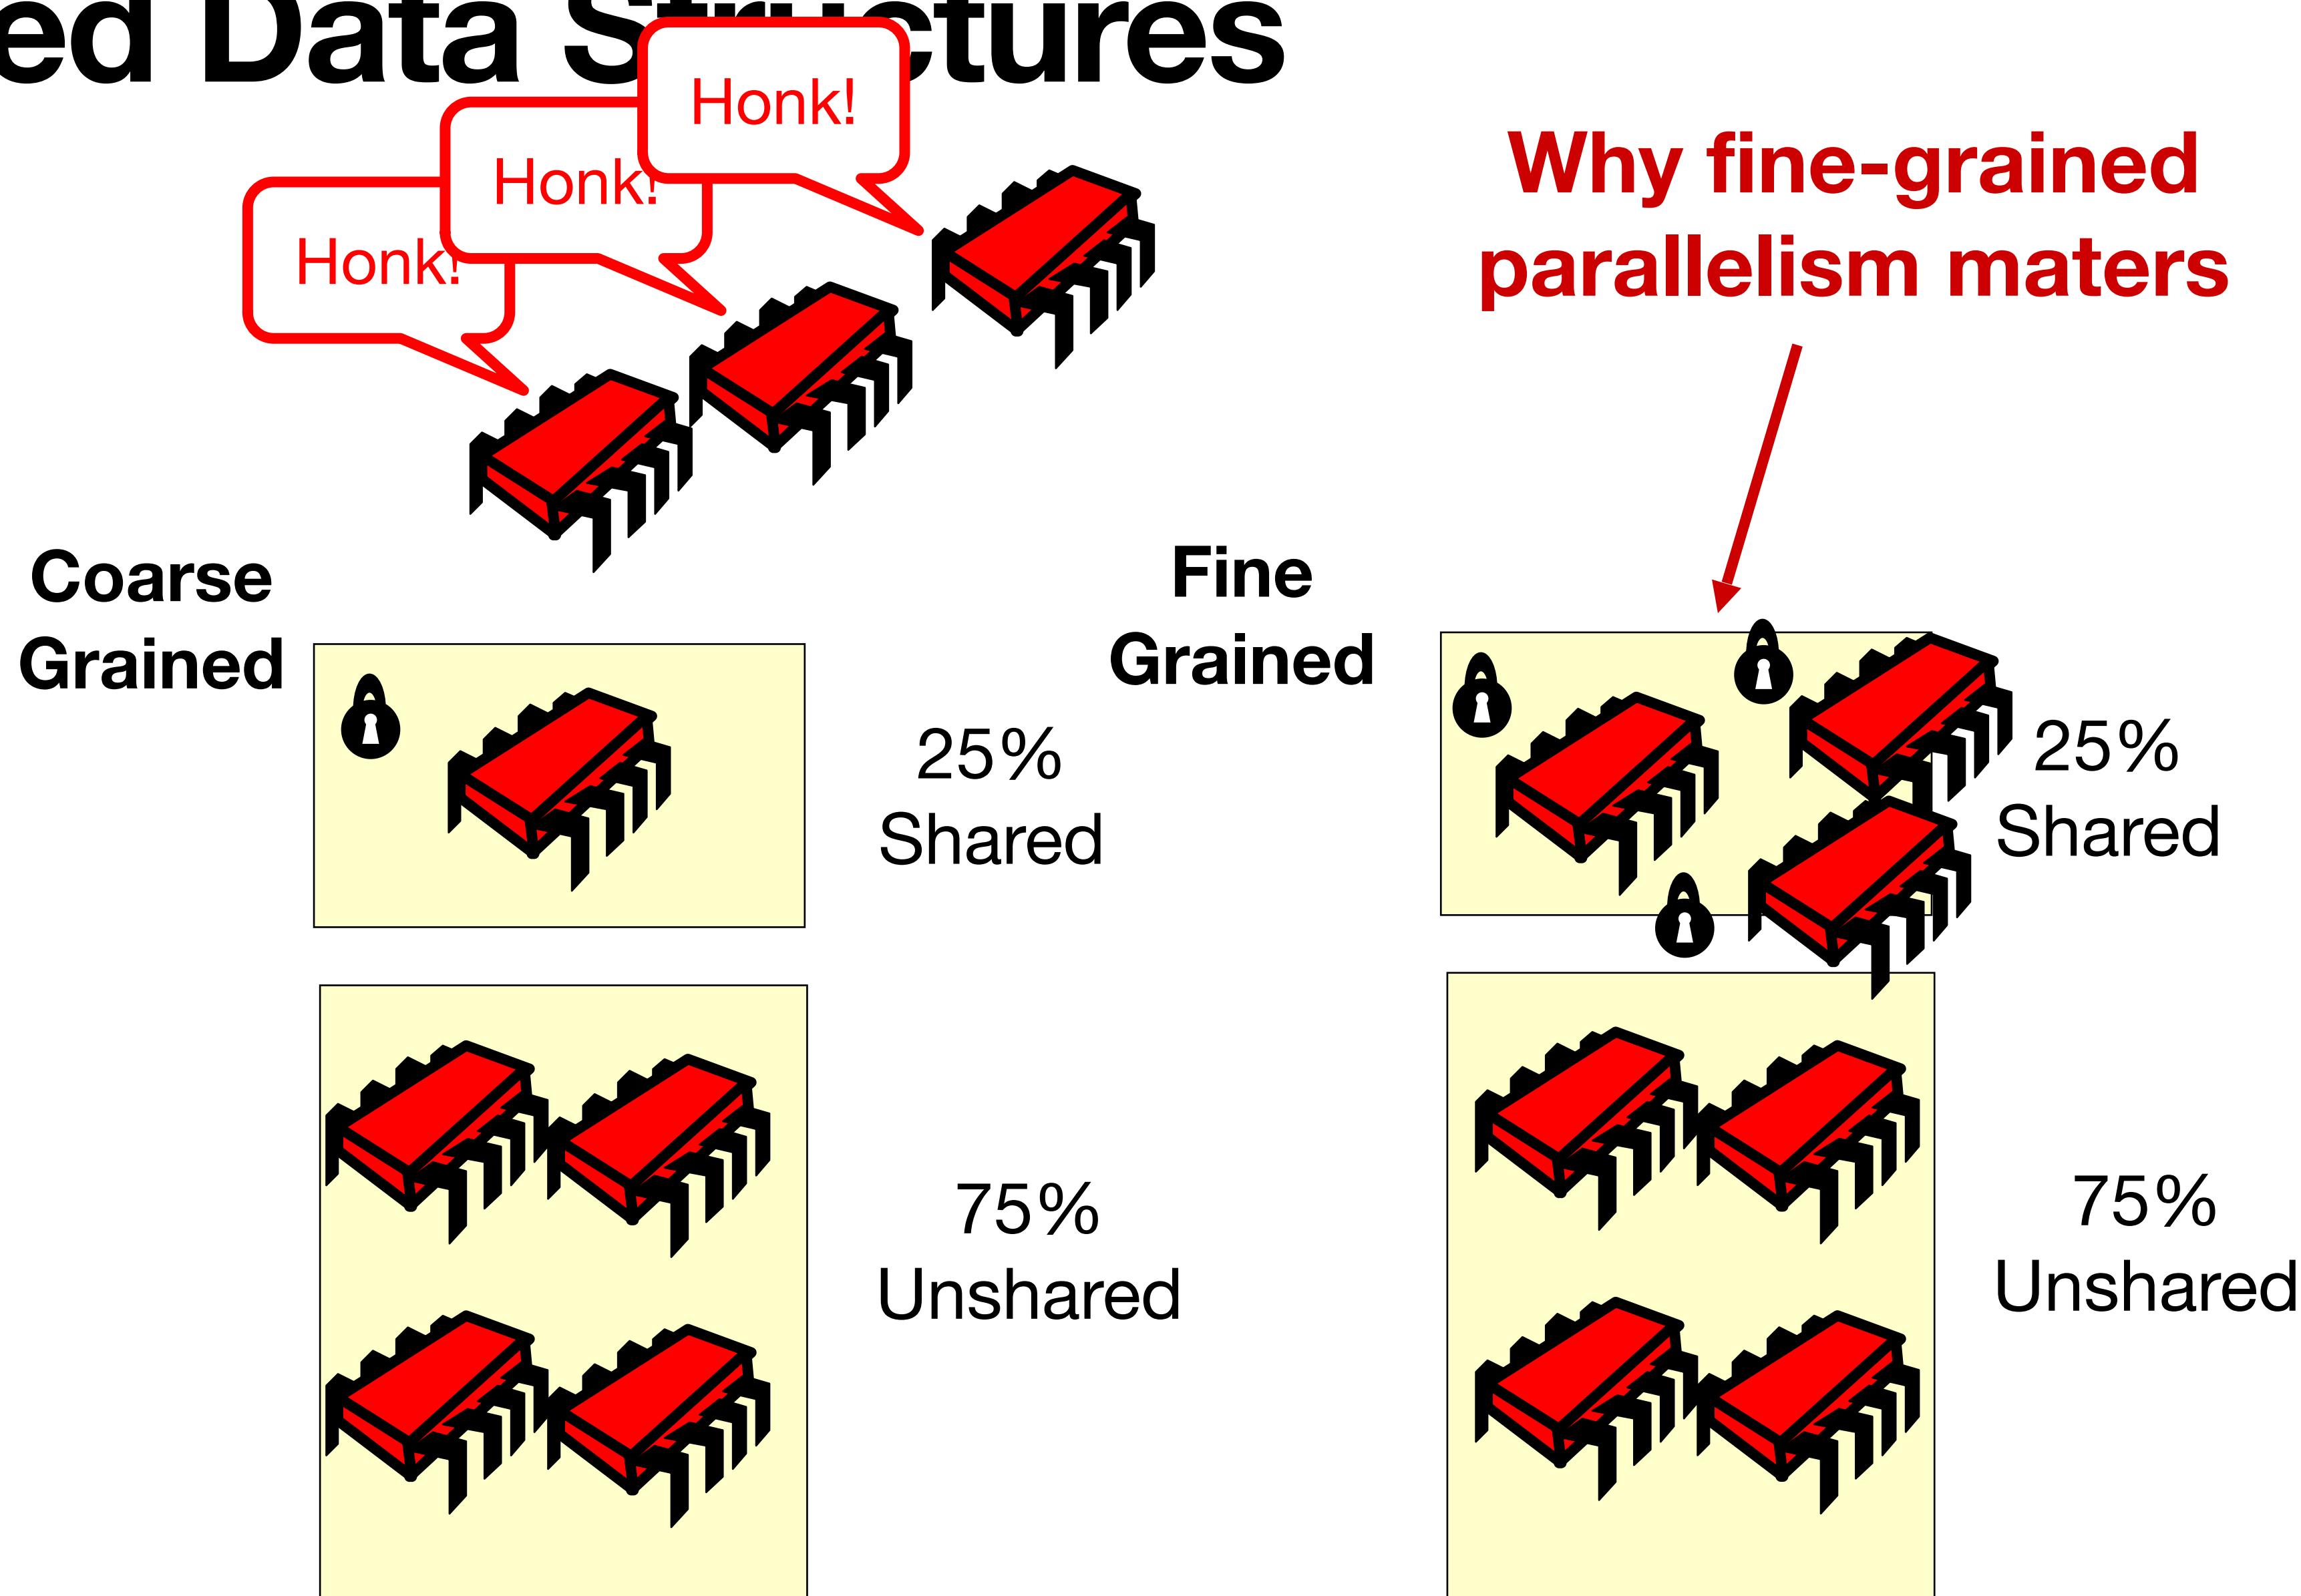


75%  
Unshared

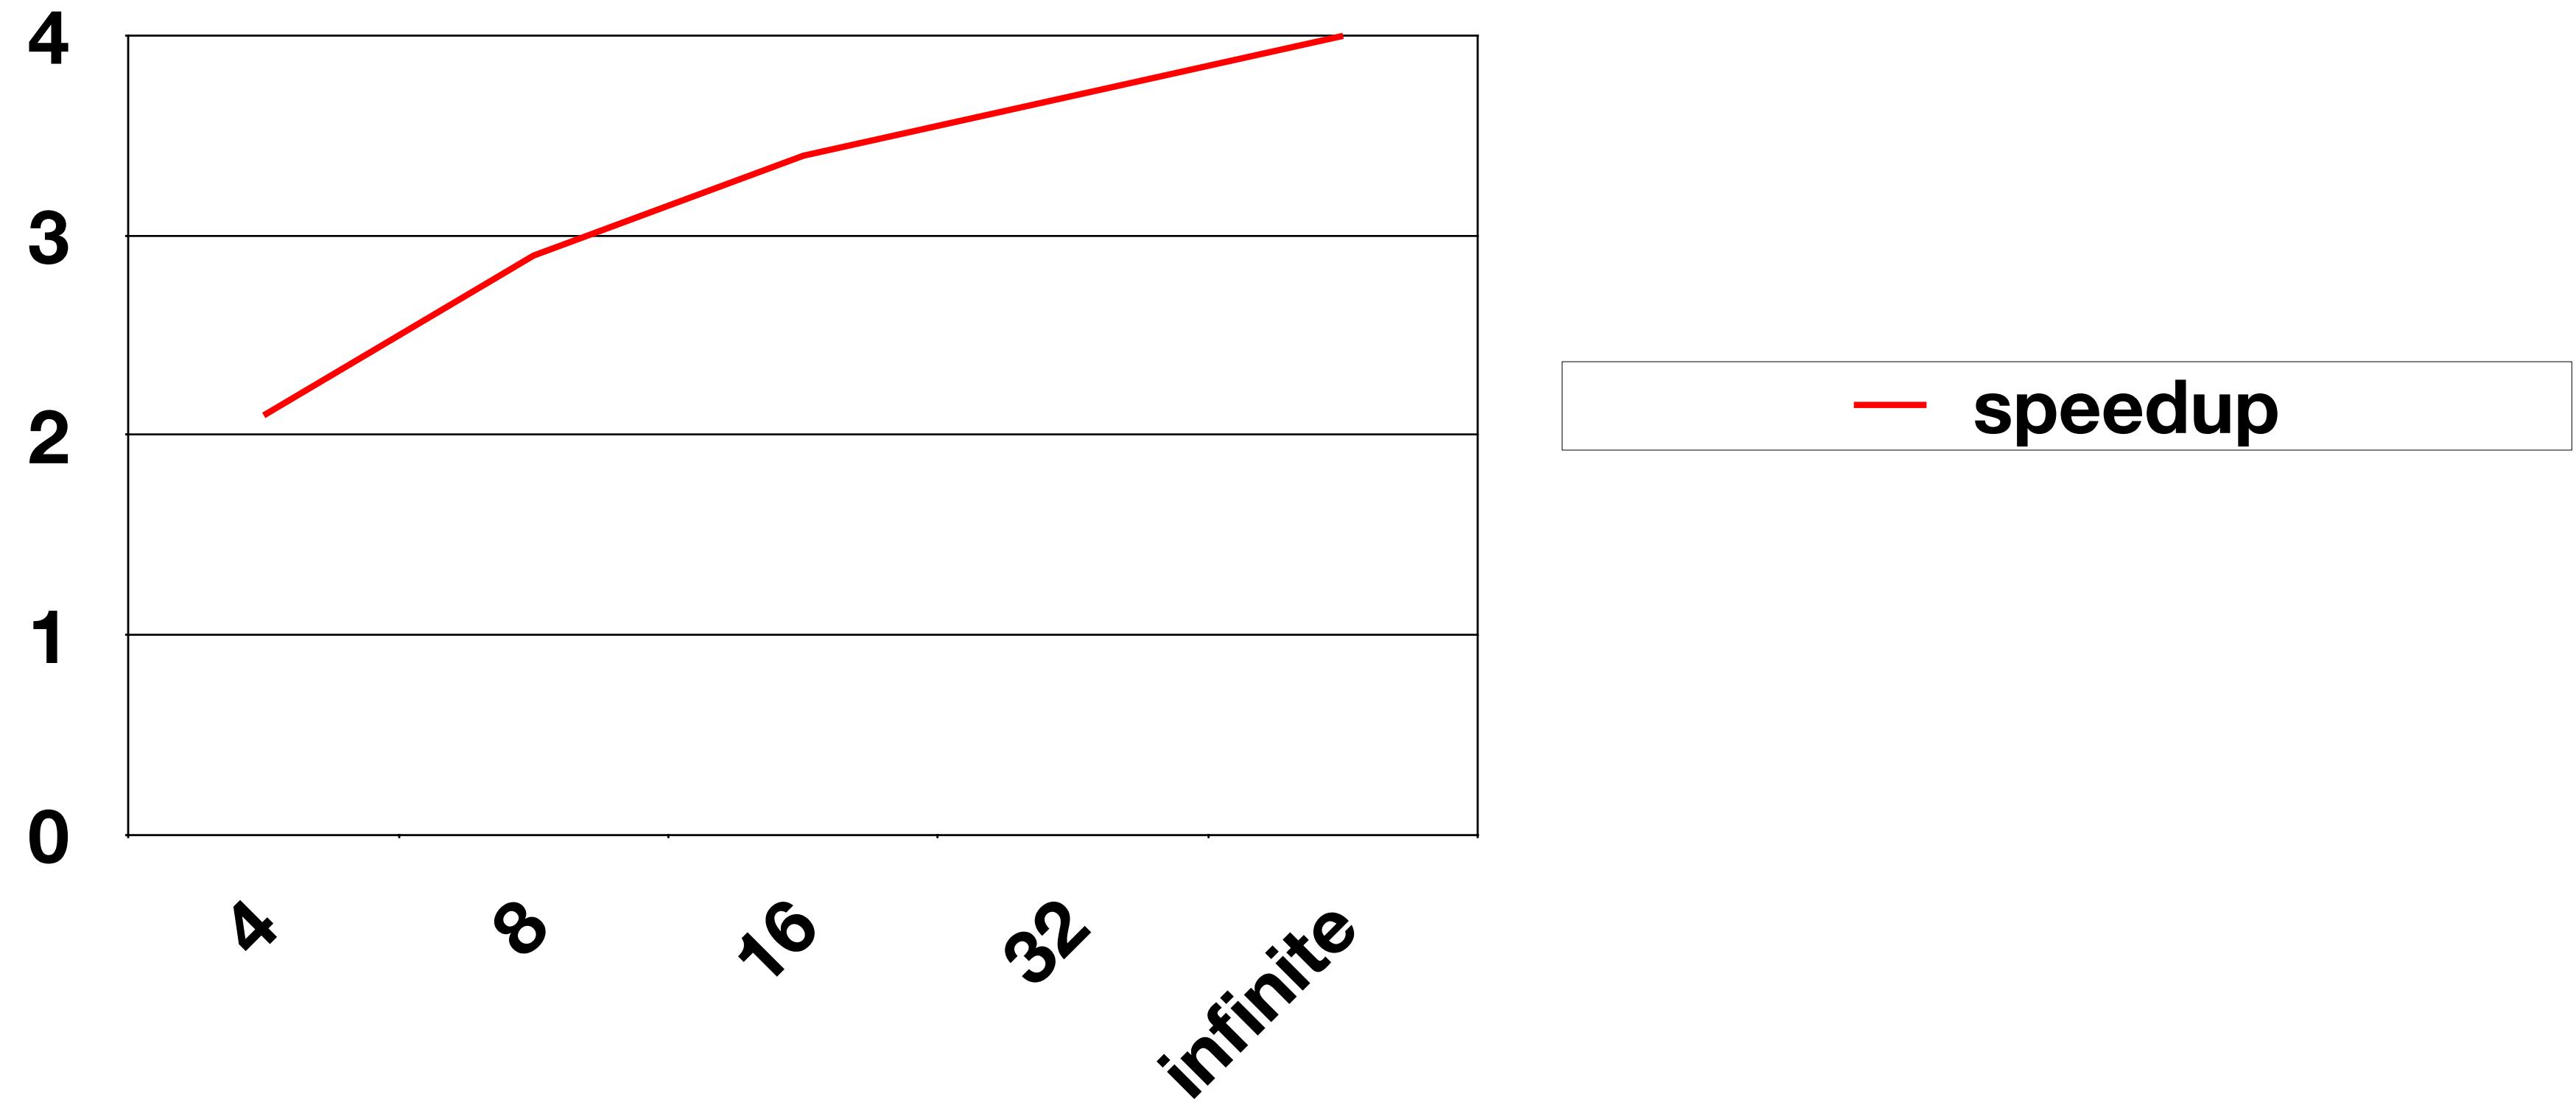
# Shared Data Structures



# Shared Data Structures



# Diminishing returns



***This course is about the parts that  
are hard to make concurrent ...  
but still have a big influence on speedup!***

**Fin**



This work is licensed under a [Creative Commons Attribution-ShareAlike 2.5 License](#).

- **You are free:**
  - **to Share** — to copy, distribute and transmit the work
  - **to Remix** — to adapt the work
- **Under the following conditions:**
  - **Attribution.** You must attribute the work to “The Art of Multiprocessor Programming” (but not in any way that suggests that the authors endorse you or your use of the work).
  - **Share Alike.** If you alter, transform, or build upon this work, you may distribute the resulting work only under the same, similar or a compatible license.
- For any reuse or distribution, you must make clear to others the license terms of this work. The best way to do this is with a link to
  - <http://creativecommons.org/licenses/by-sa/3.0/>.
- Any of the above conditions can be waived if you get permission from the copyright holder.
- Nothing in this license impairs or restricts the author's moral rights.