

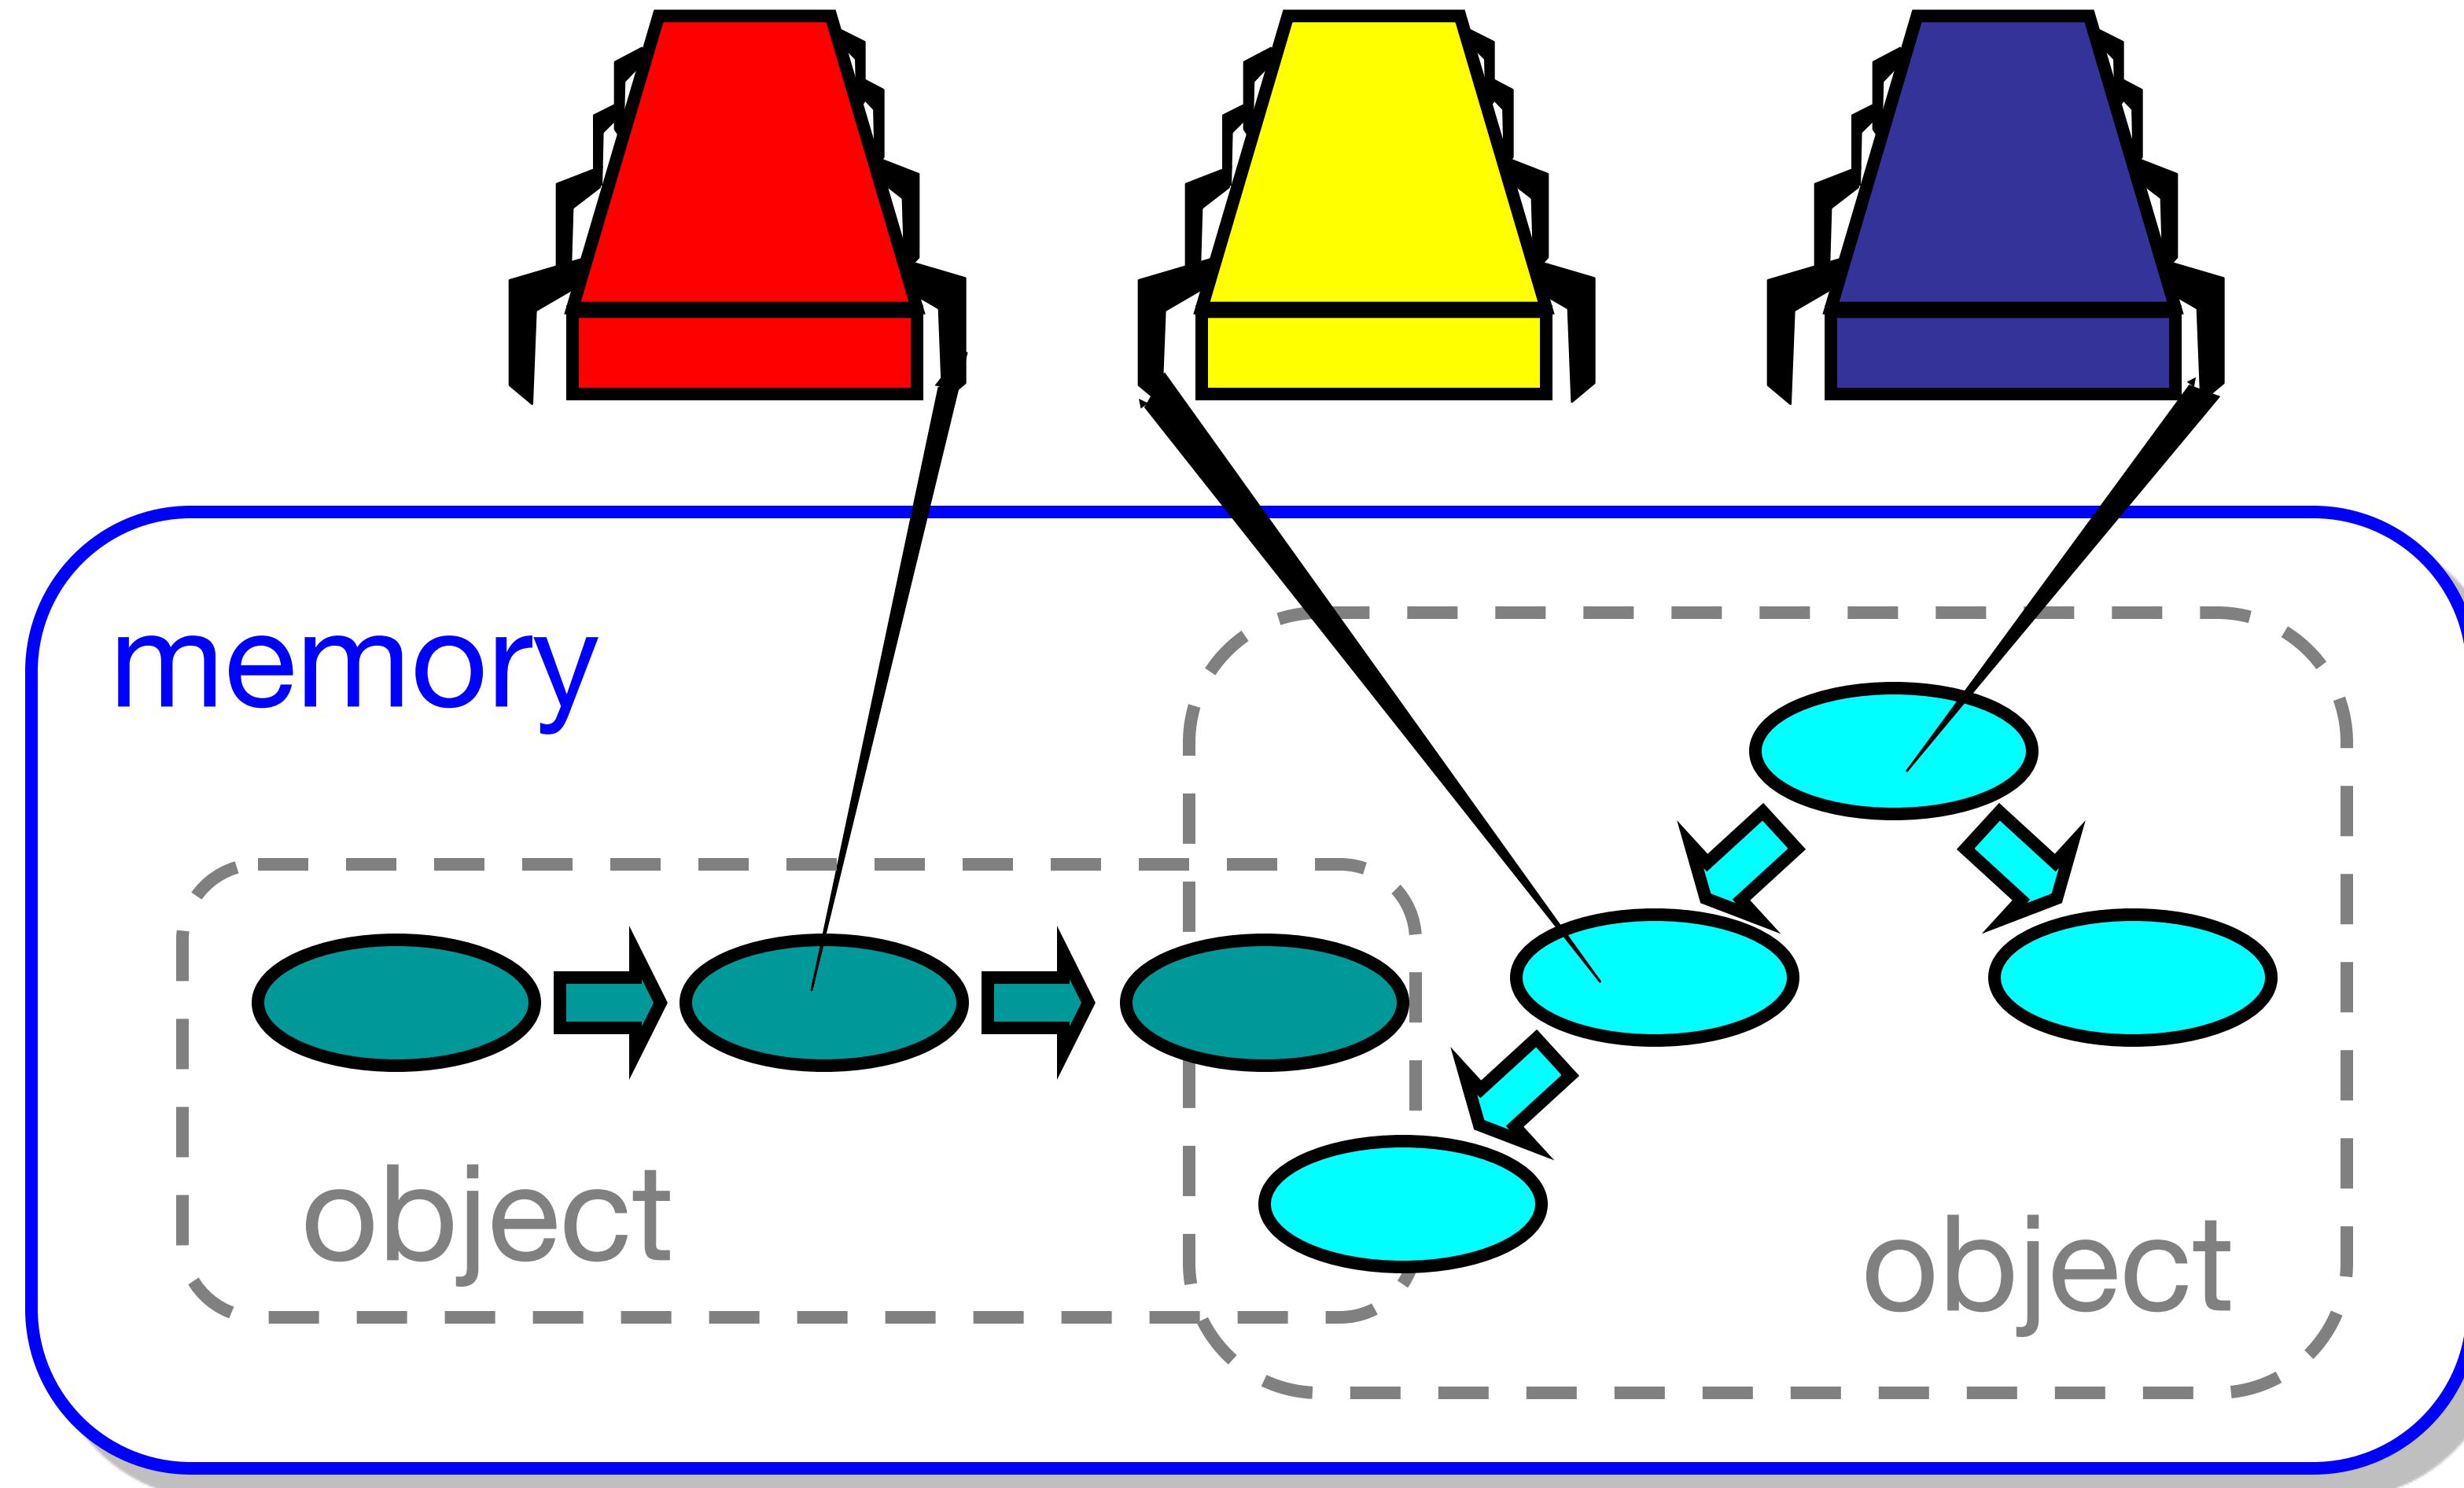
# **03 Concurrent Objects**

**CS 6868: Concurrent Programming**

**KC Sivaramakrishnan**

**Spring 2026, IIT Madras**

# Concurrent Computation



# Objectivism

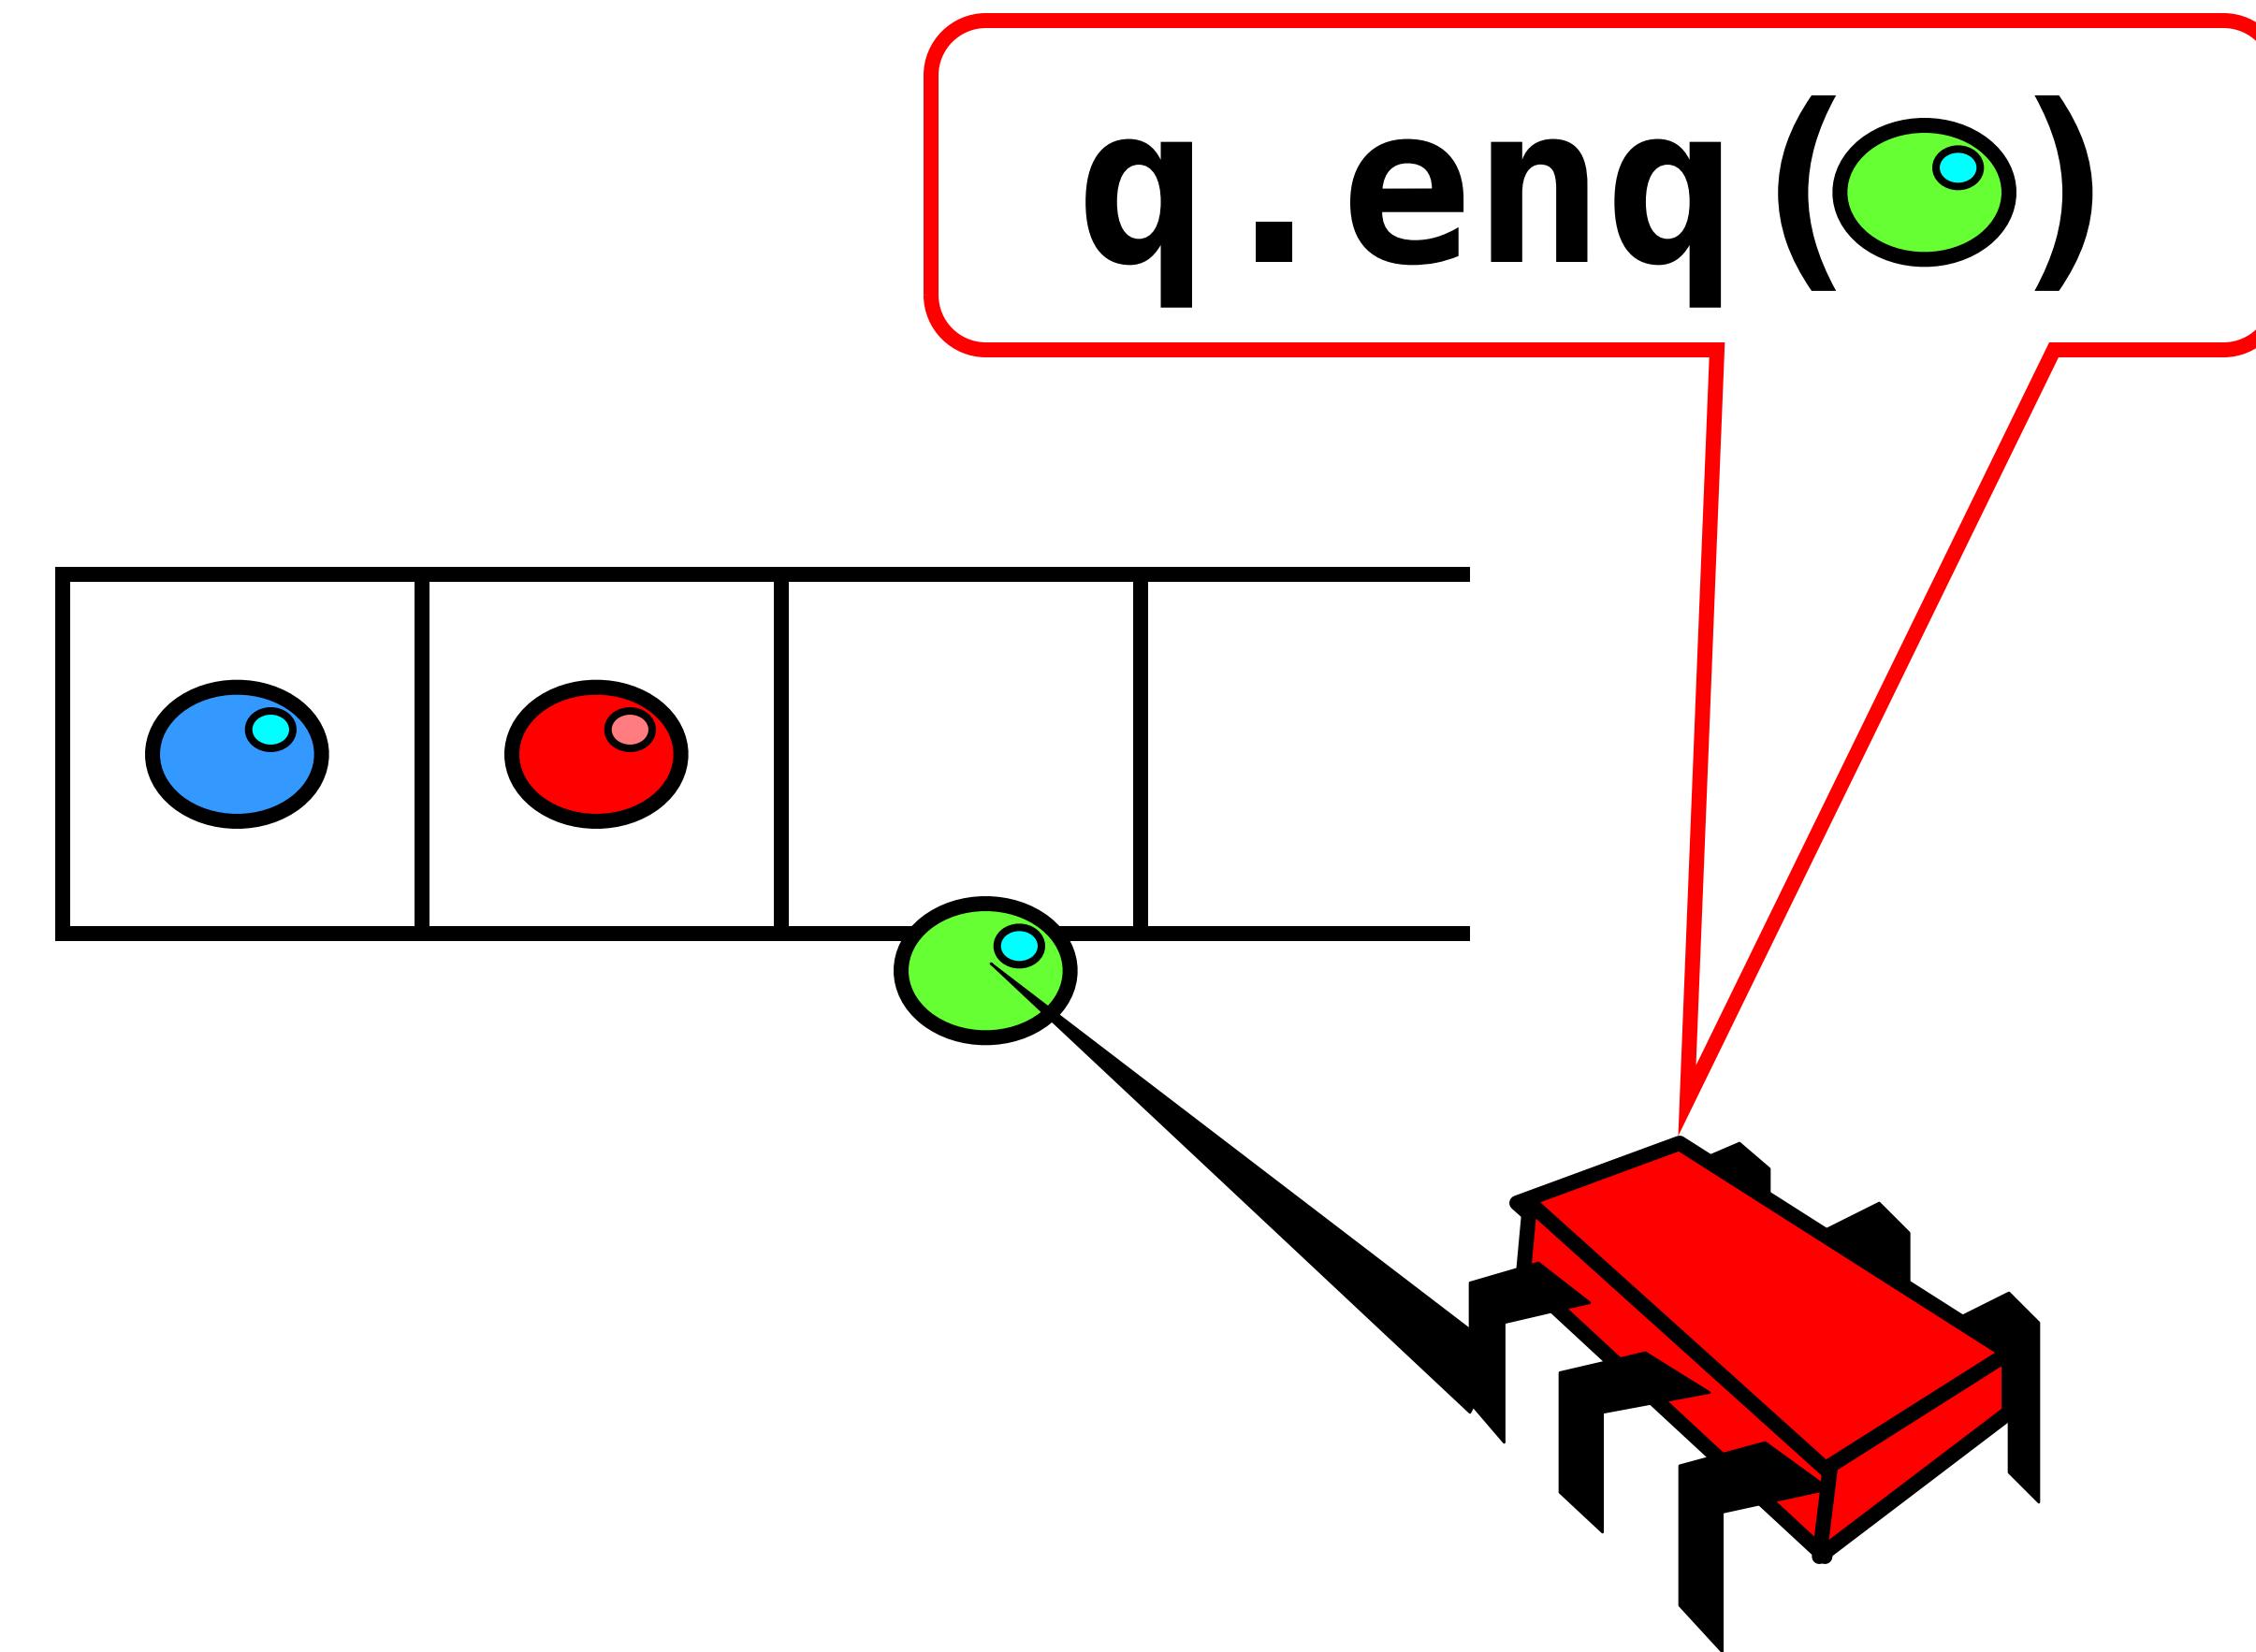
- What is a concurrent object?
  - How do we **describe** one?
  - How do we **implement** one?
  - How do we **tell if we're right**?

# Objectivism

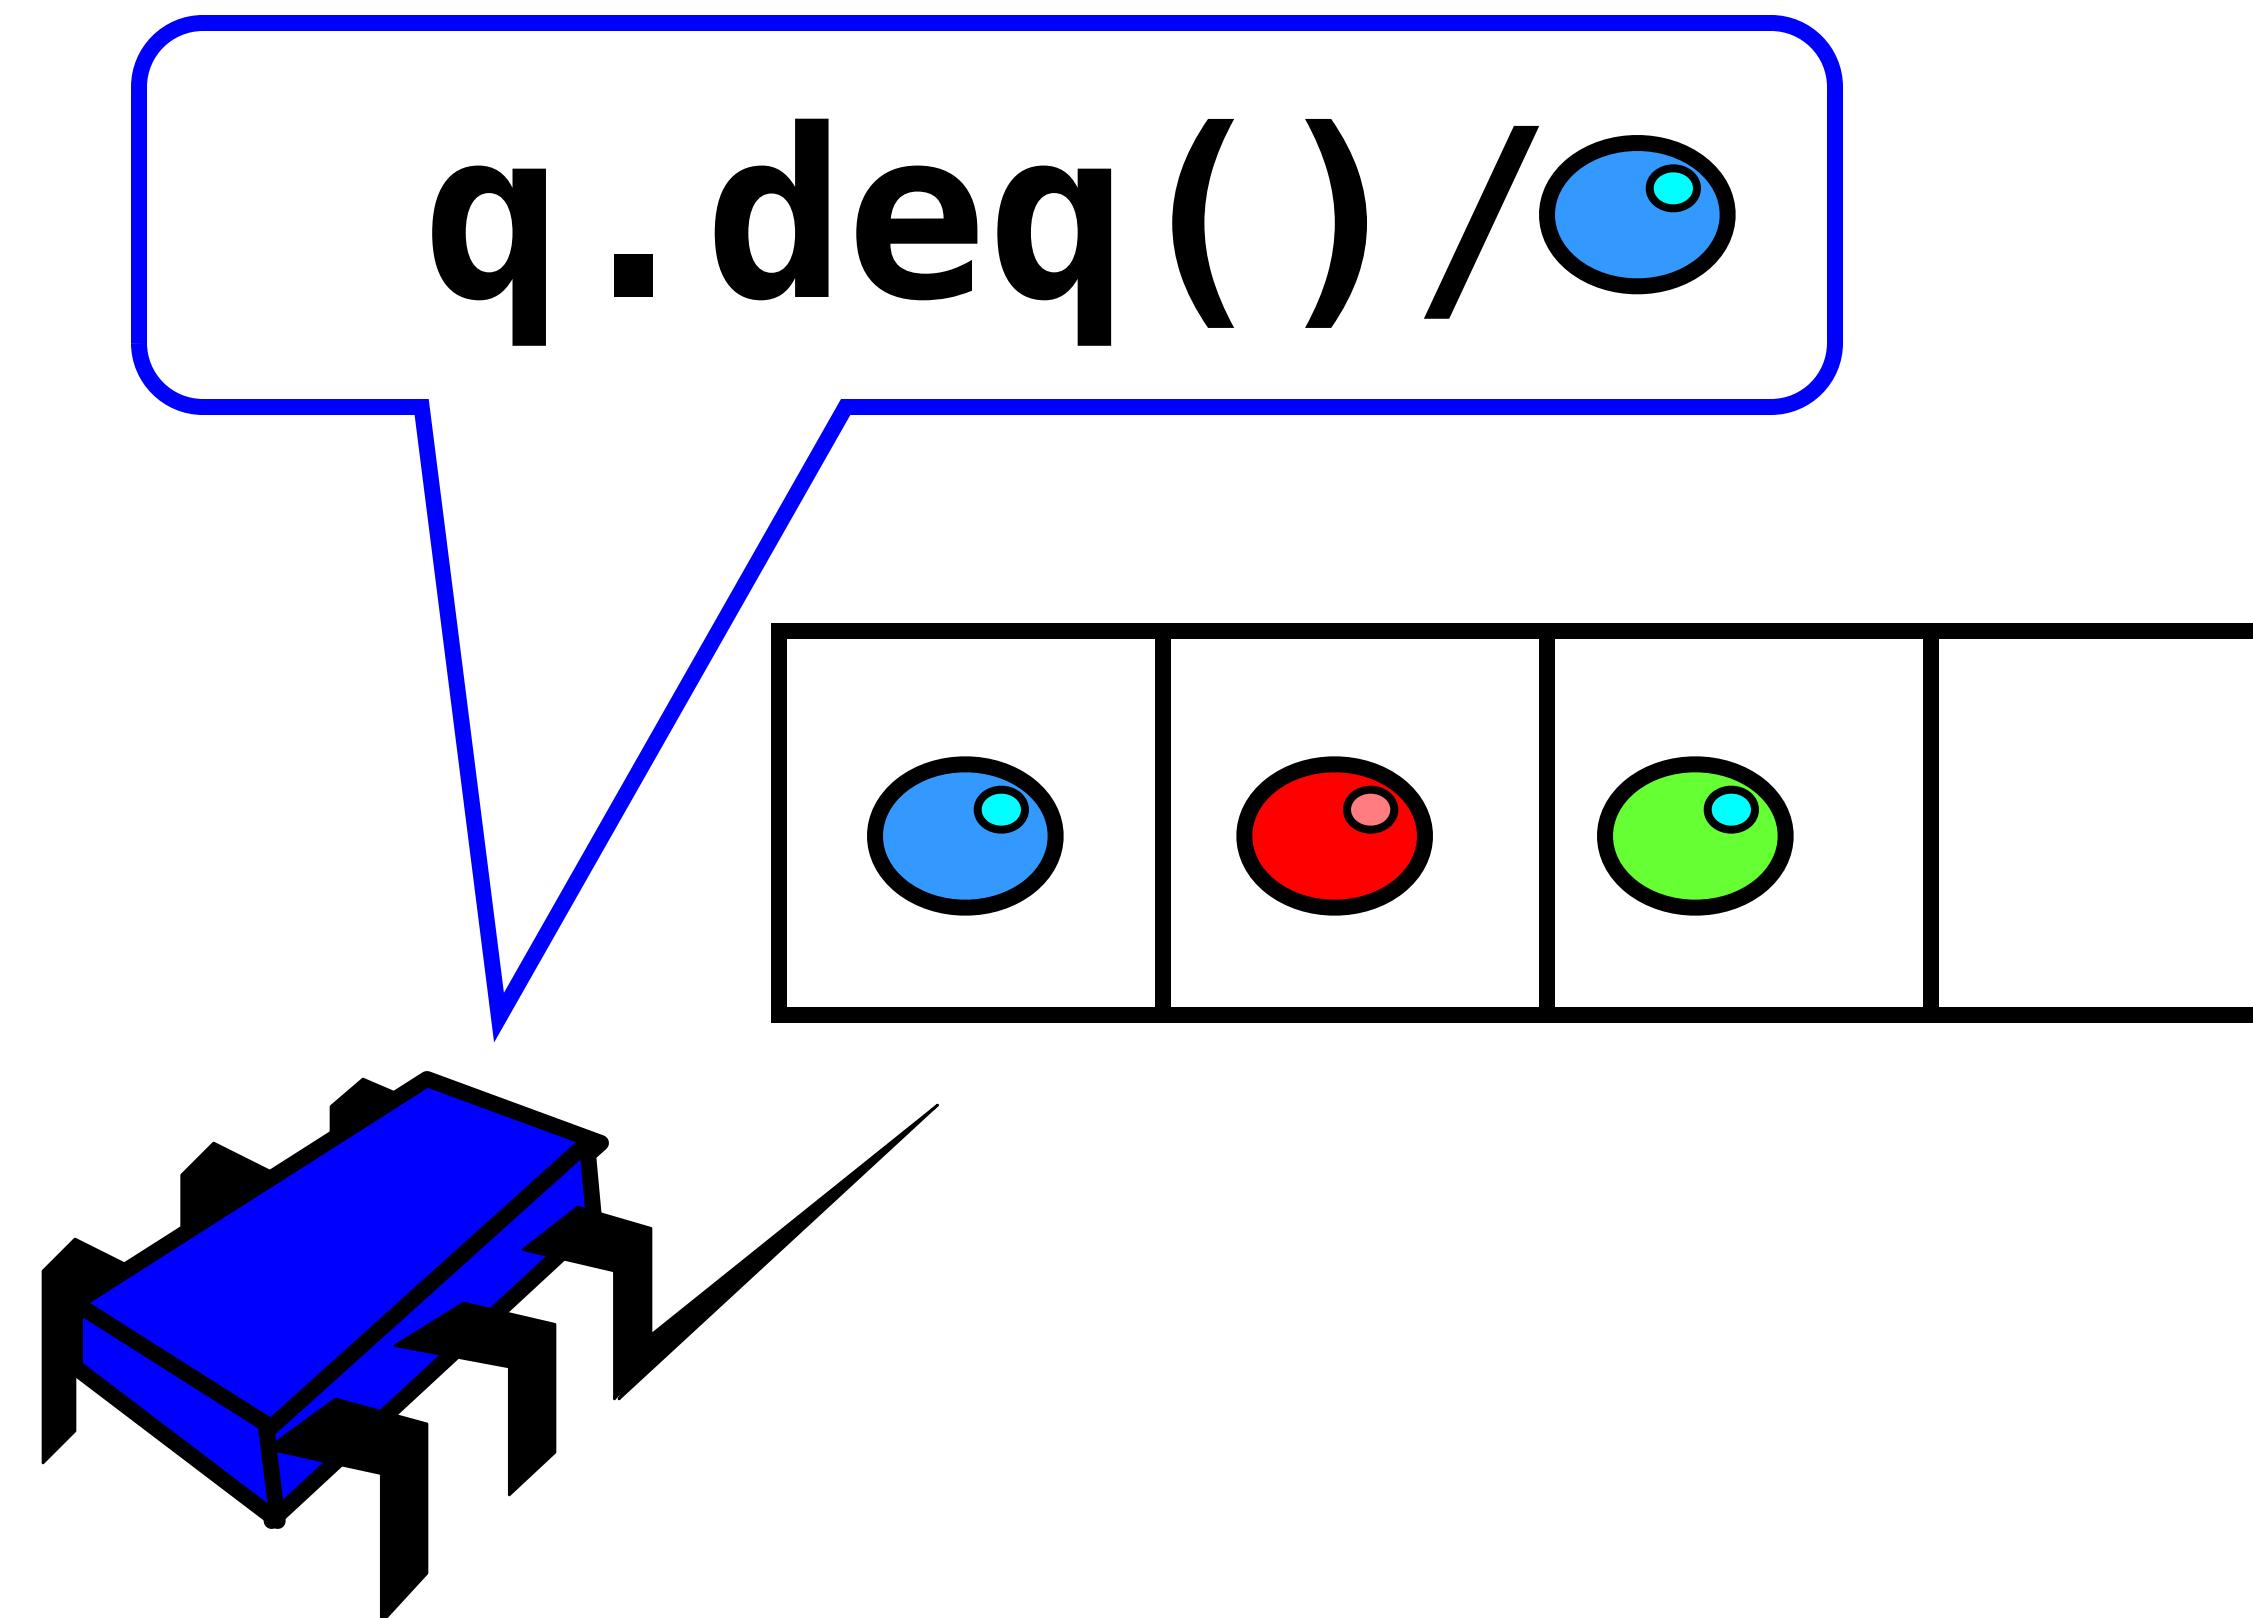
- What is a concurrent object?
  - How do we **describe** one?
  - ~~How do we **implement** one?~~
  - How do we **tell if we're right?**

# Concurrent Queues

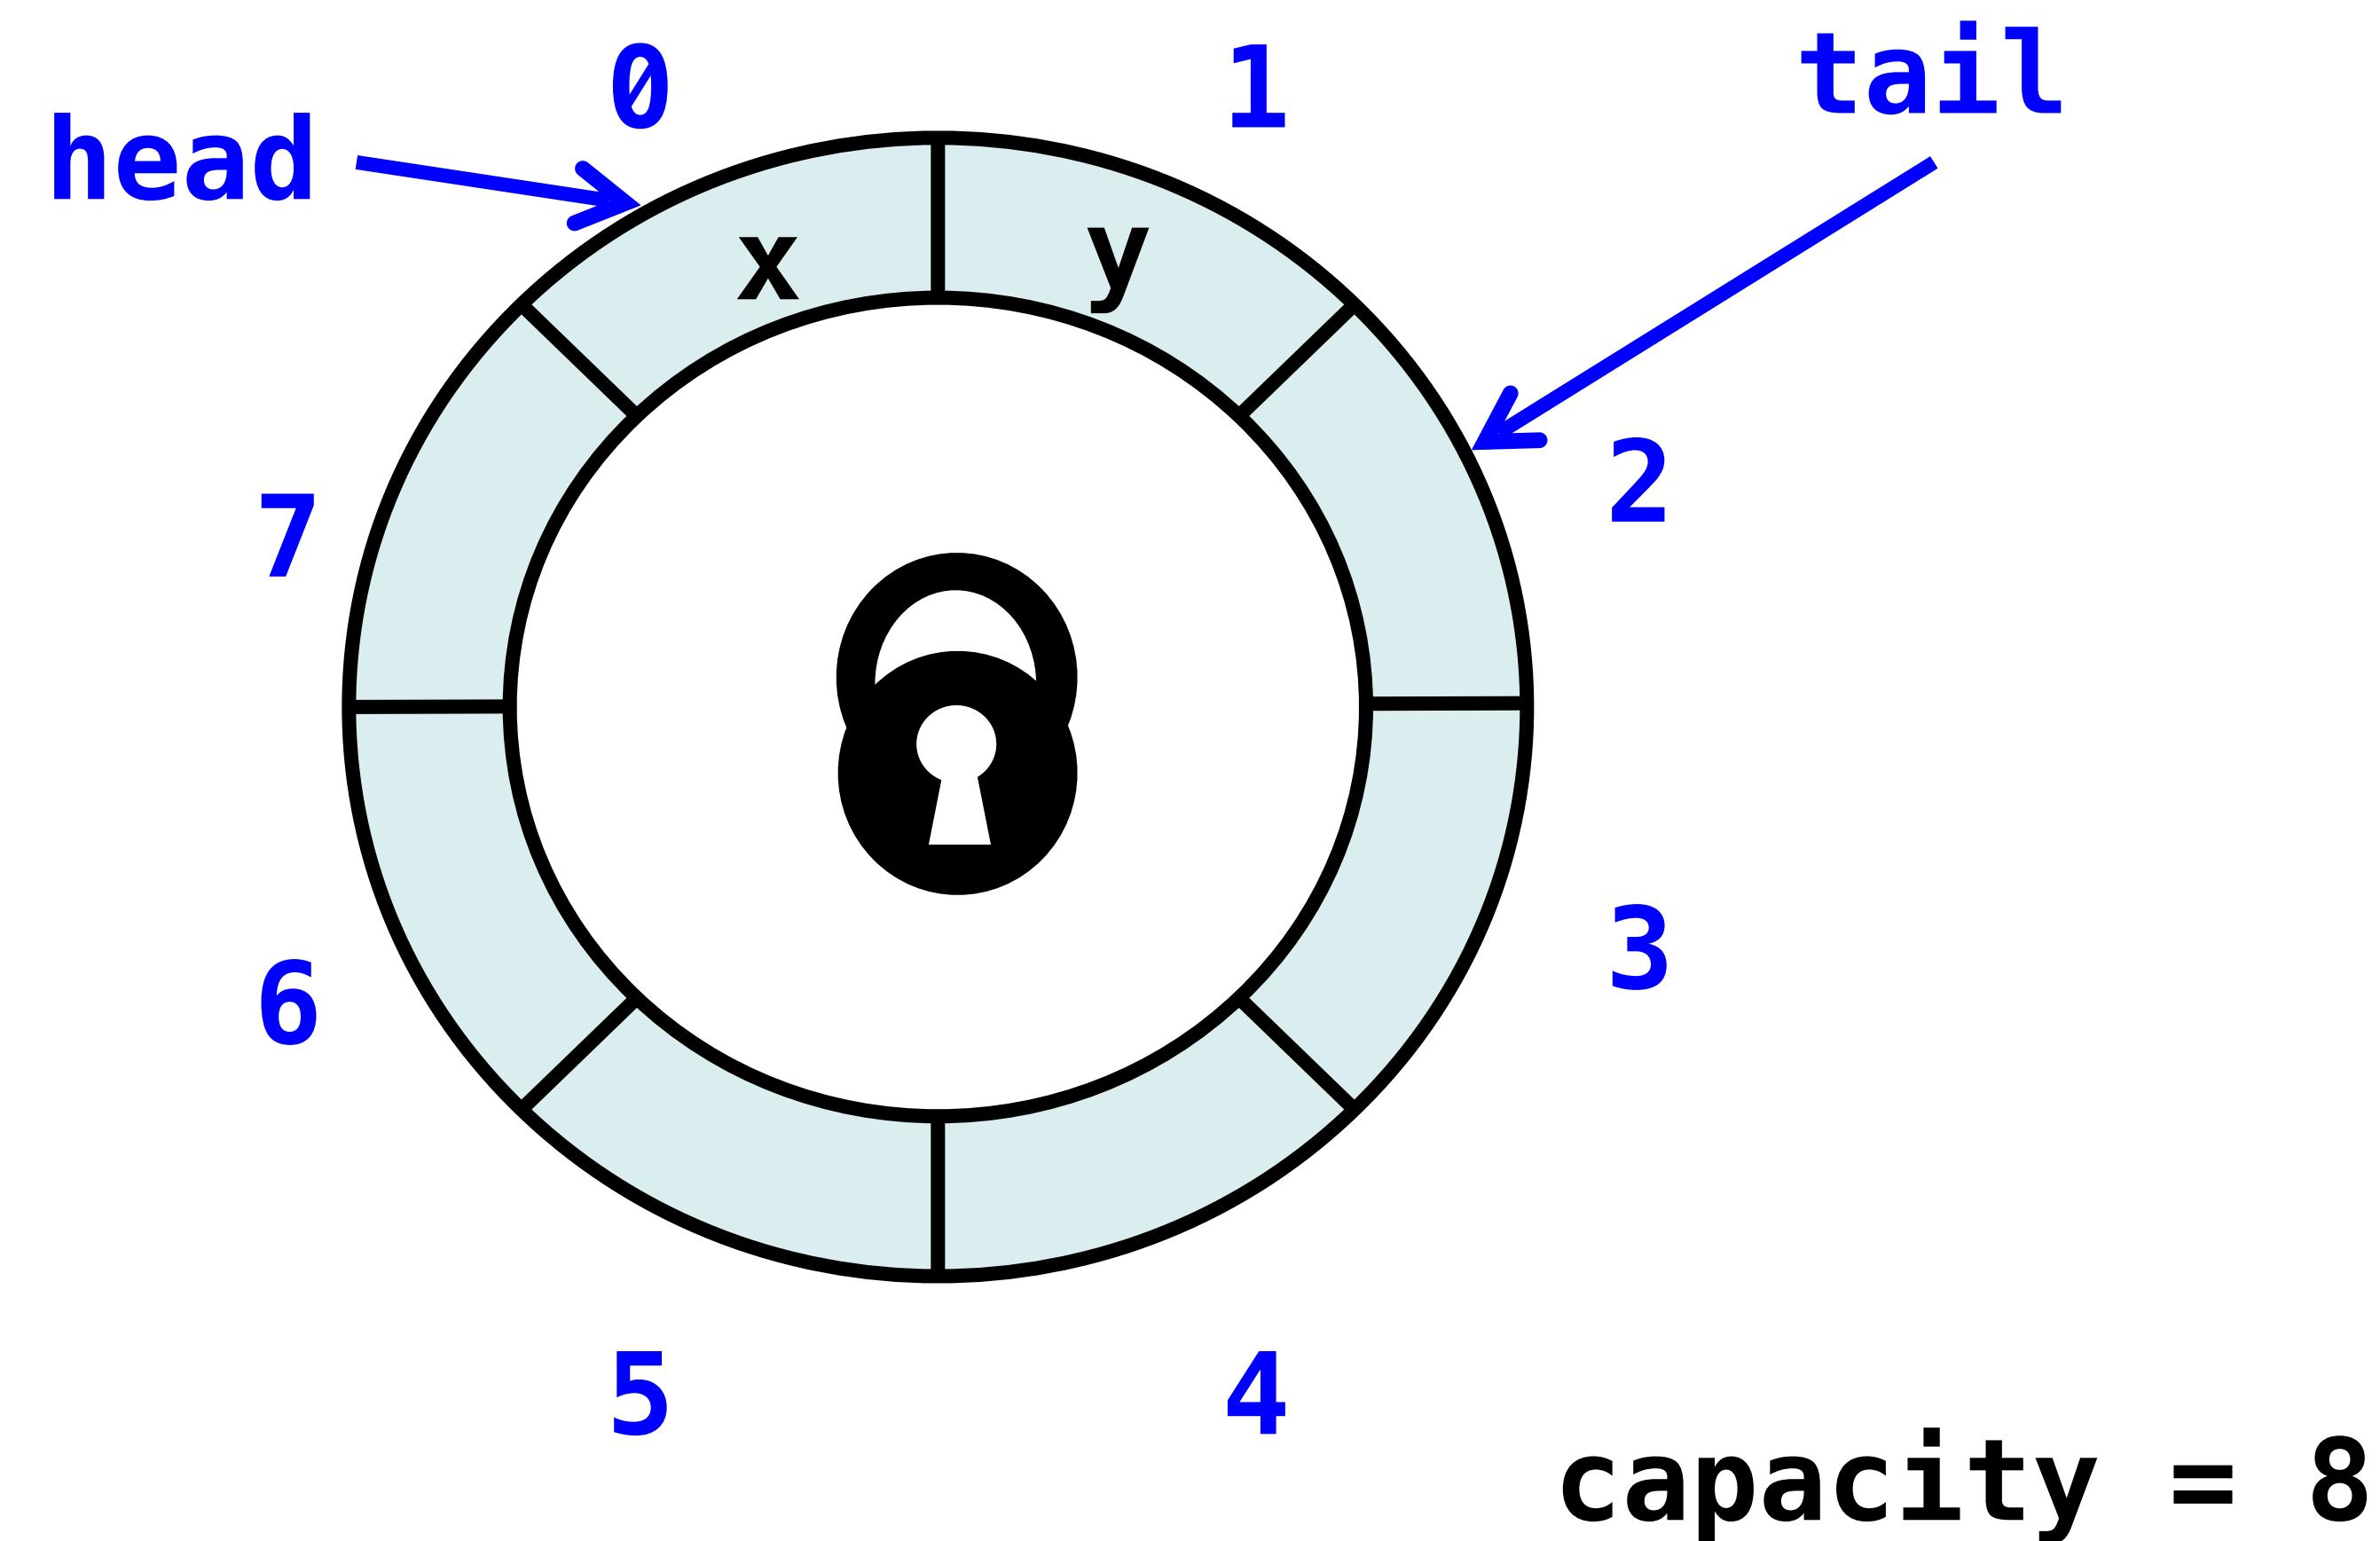
# FIFO Queue – Enqueue method



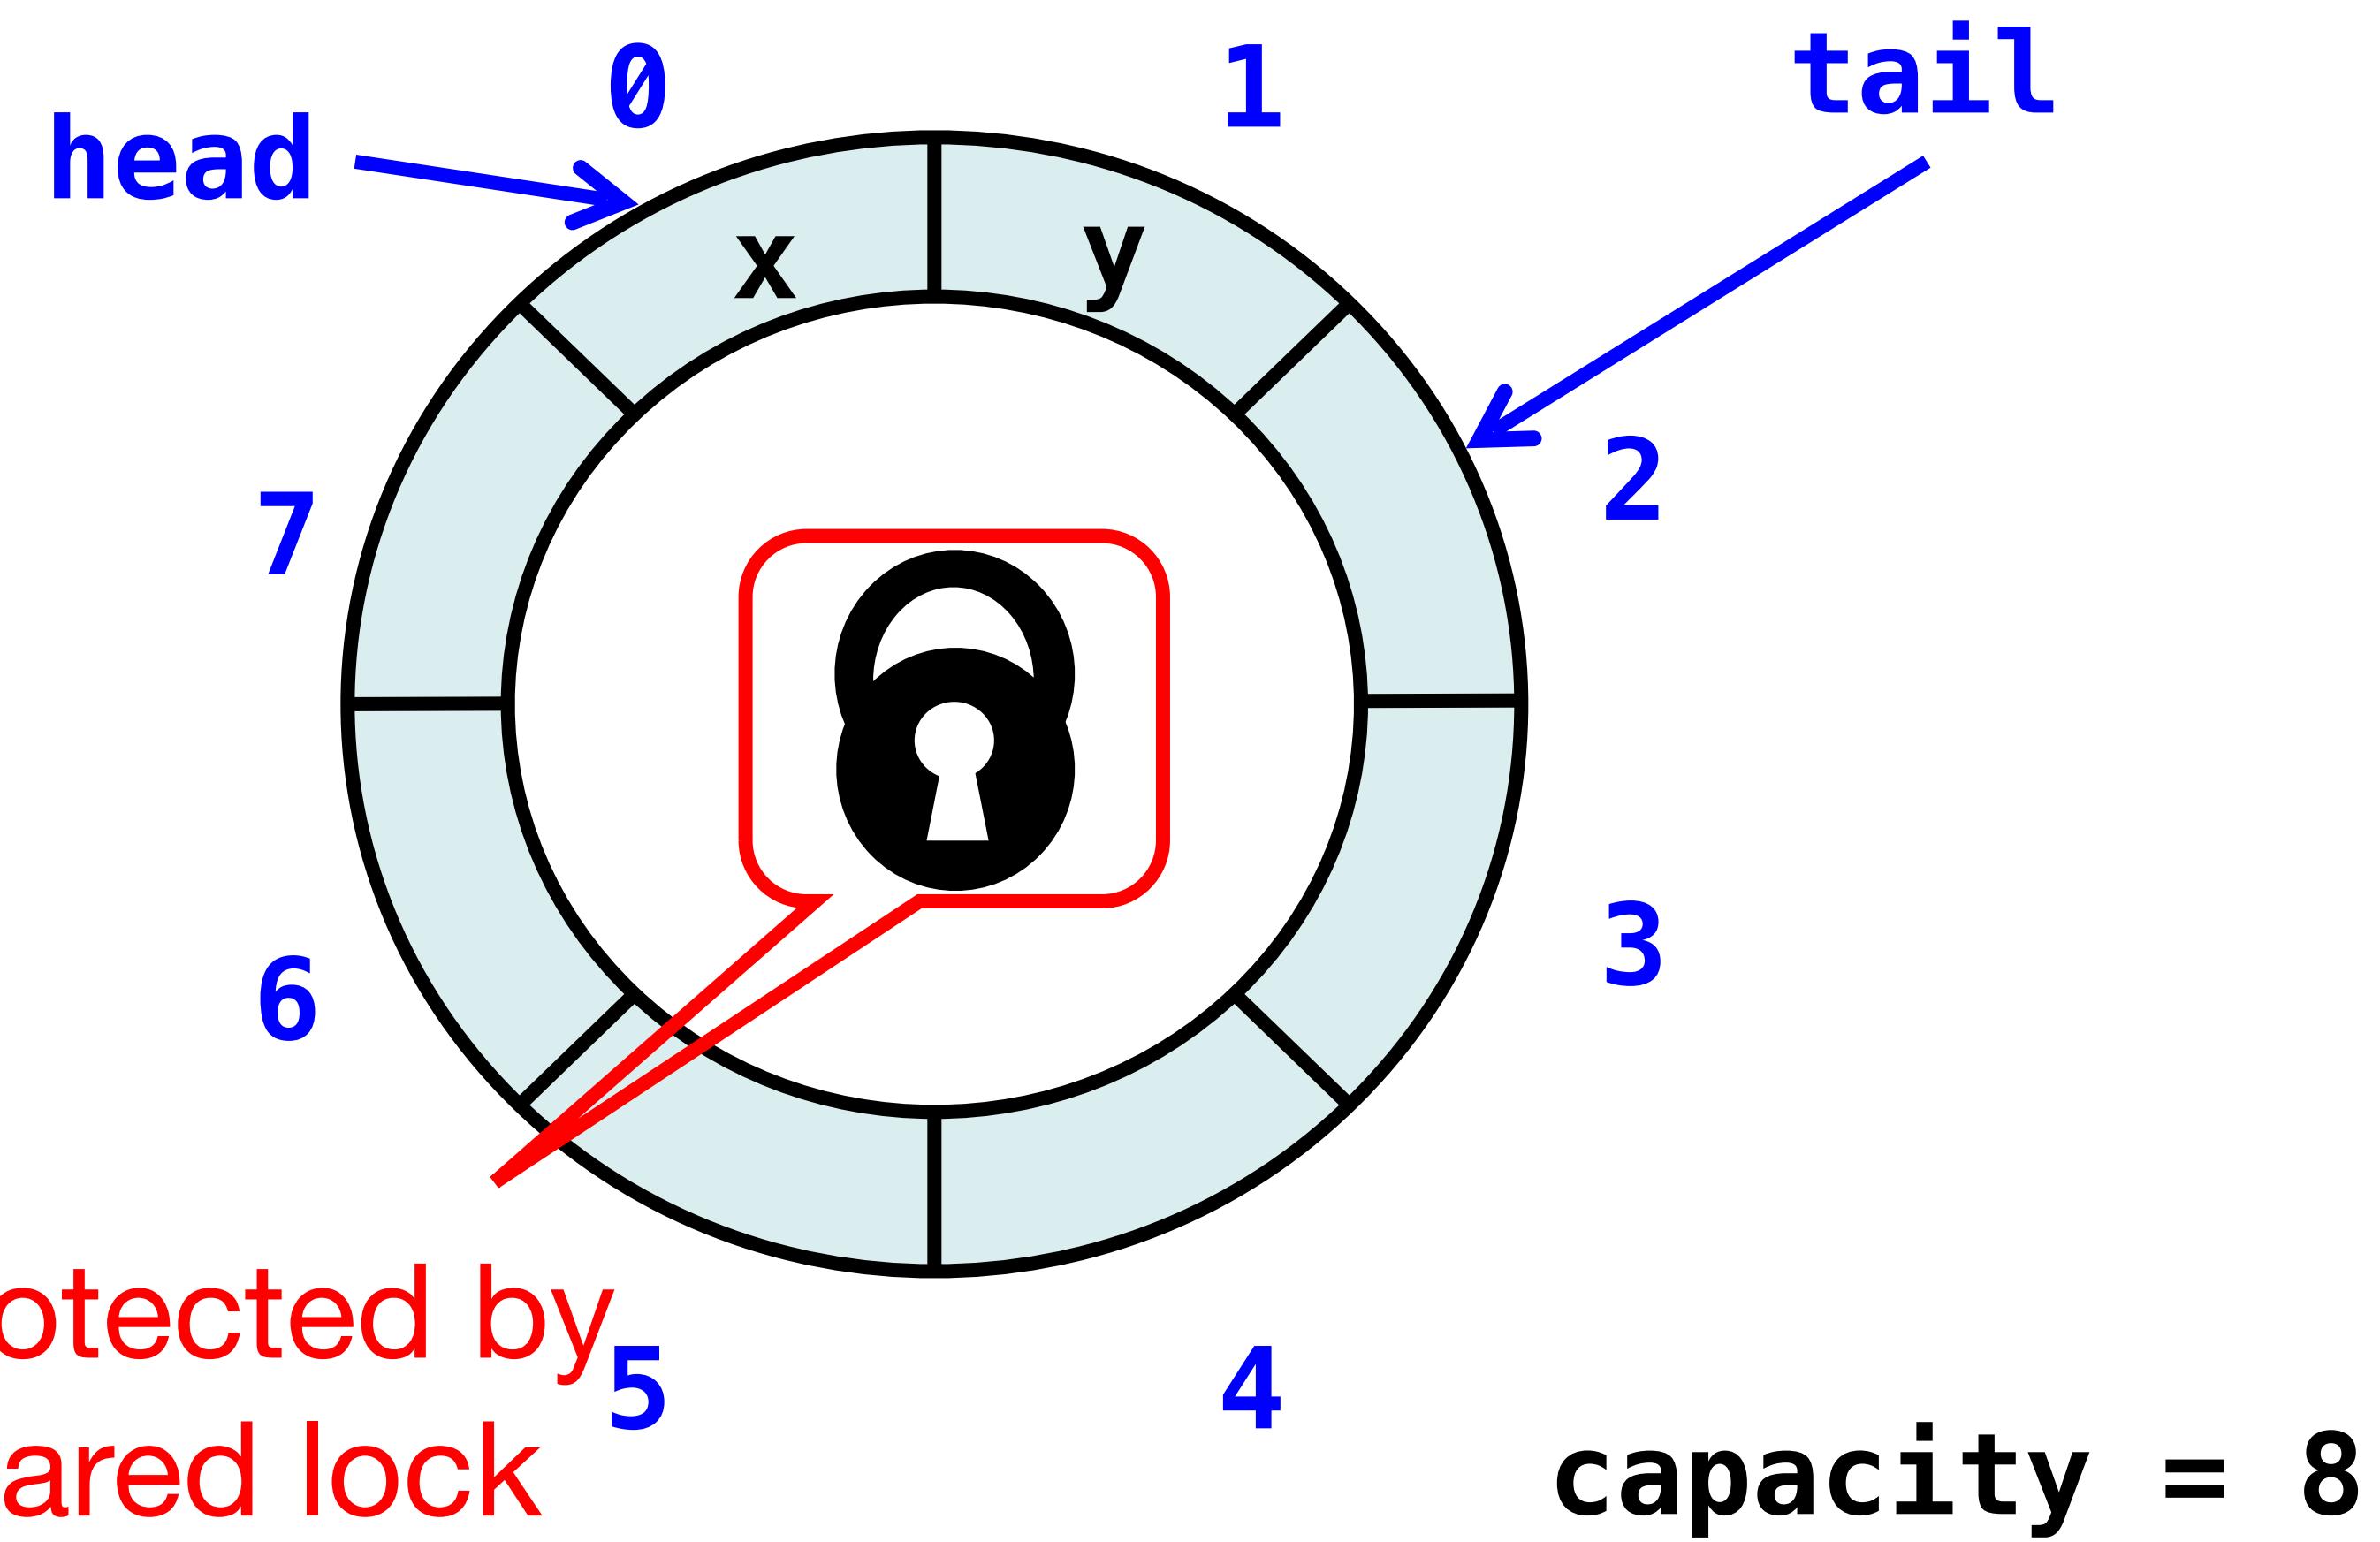
# FIFO Queue – Dequeue method



# Lock-based Queue



# Lock-based Queue



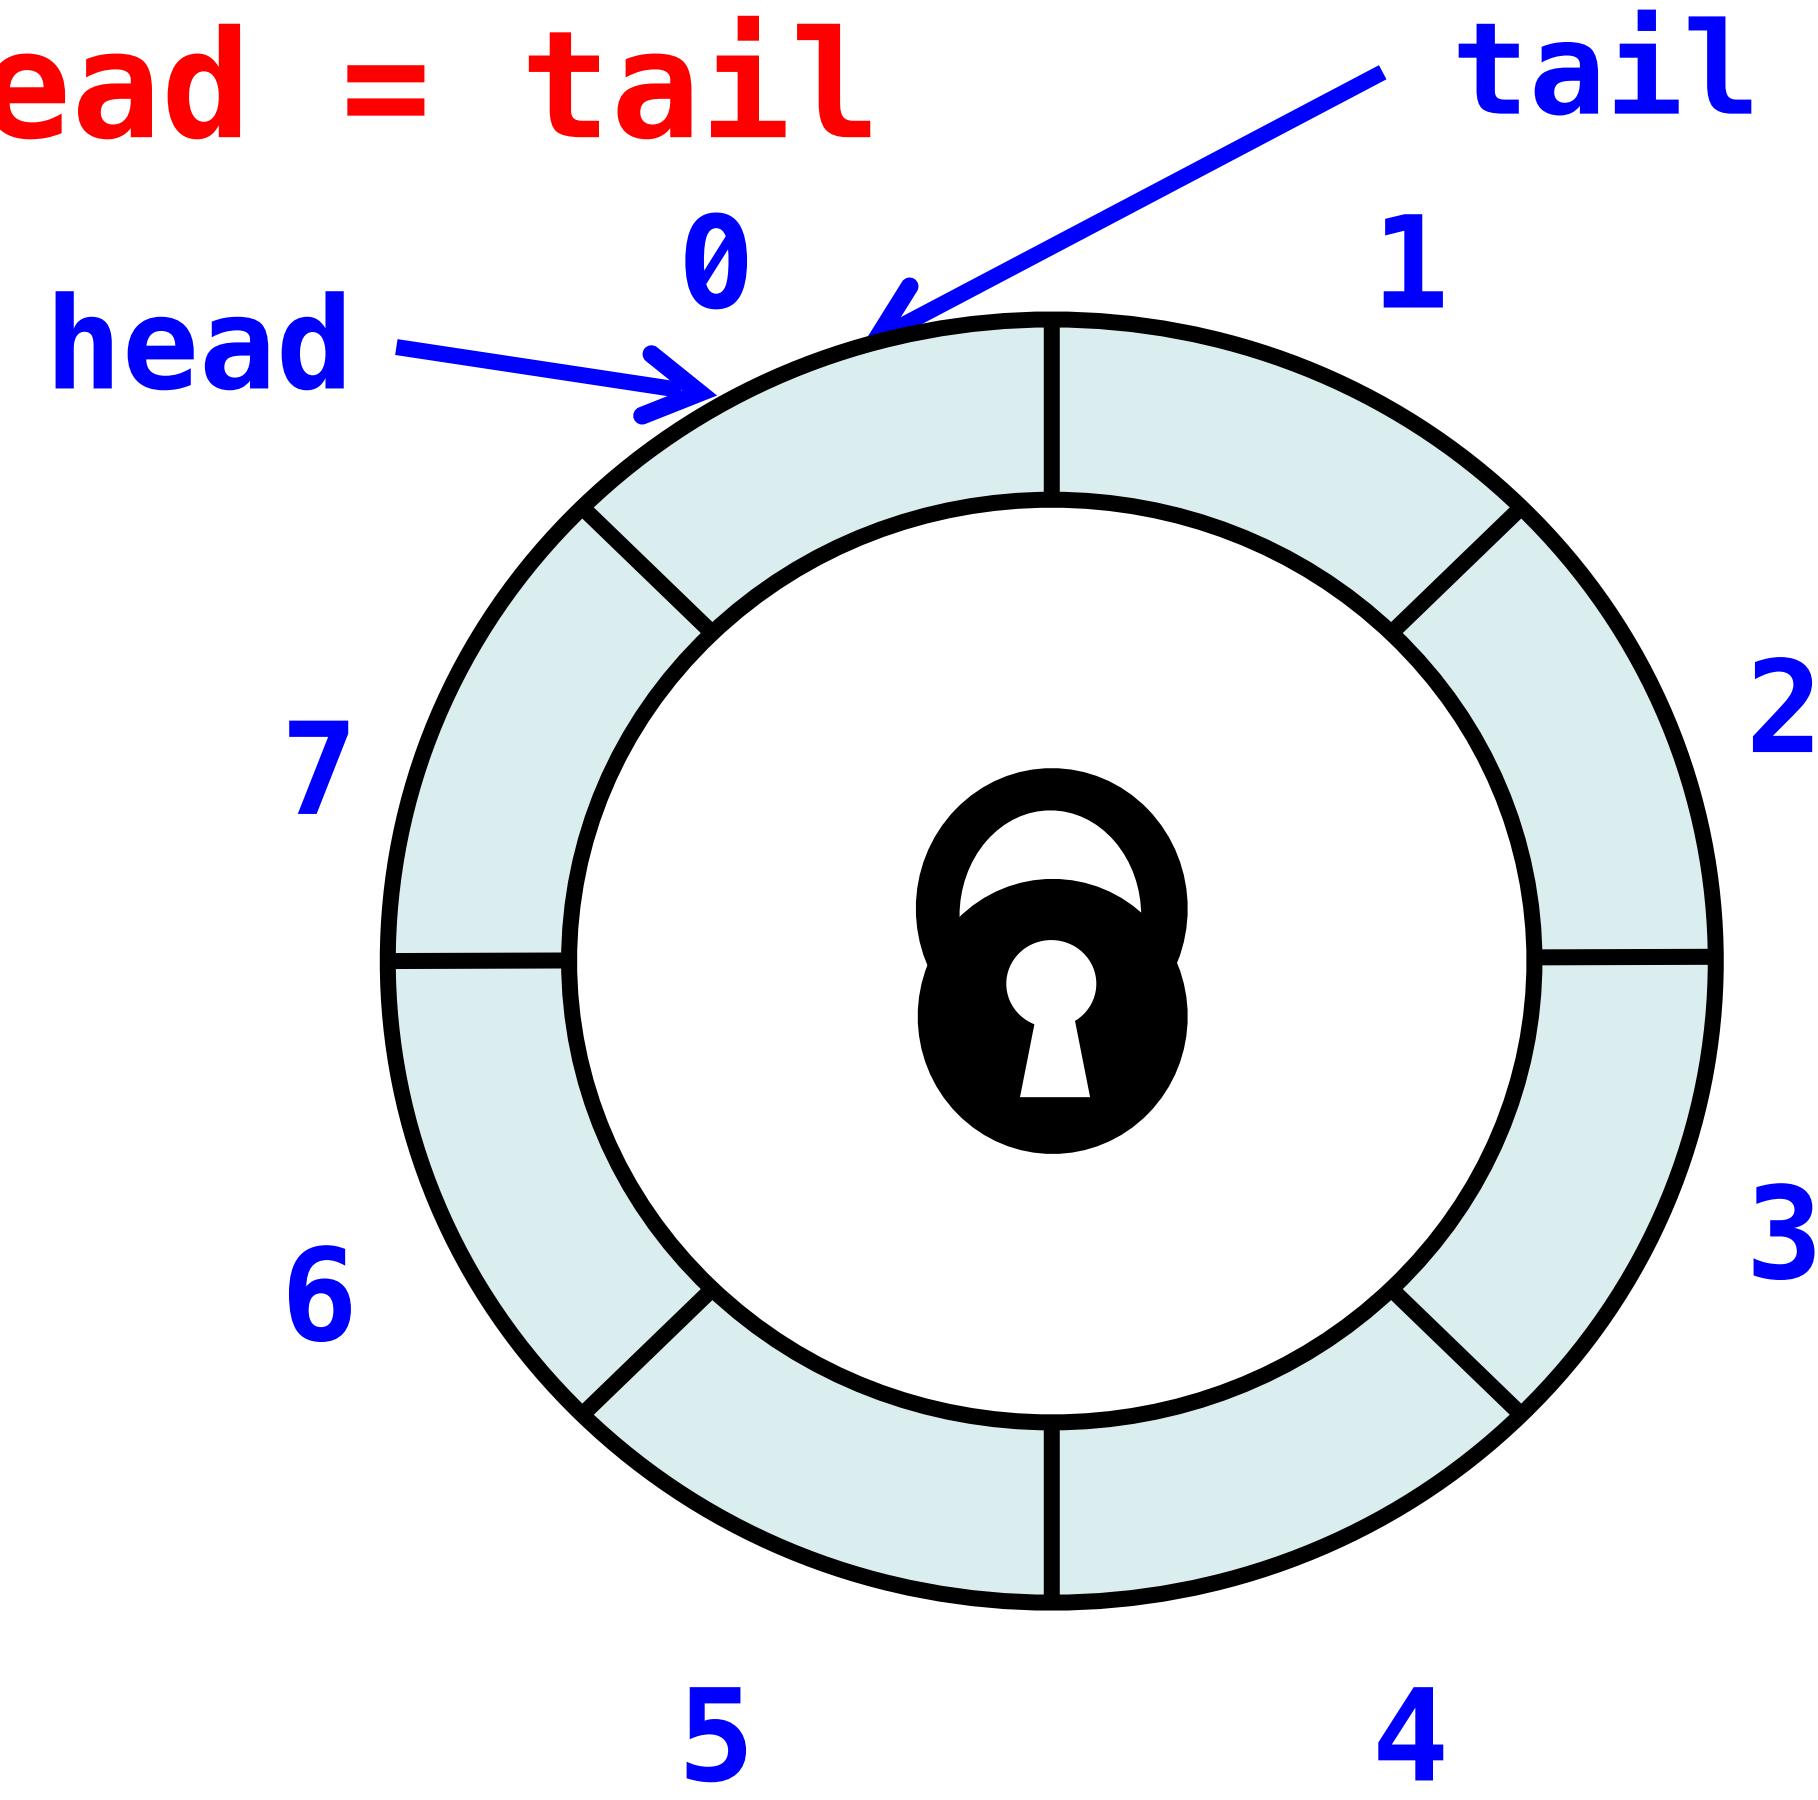
# Lock-based Queue

```
exception Full  
exception Empty
```

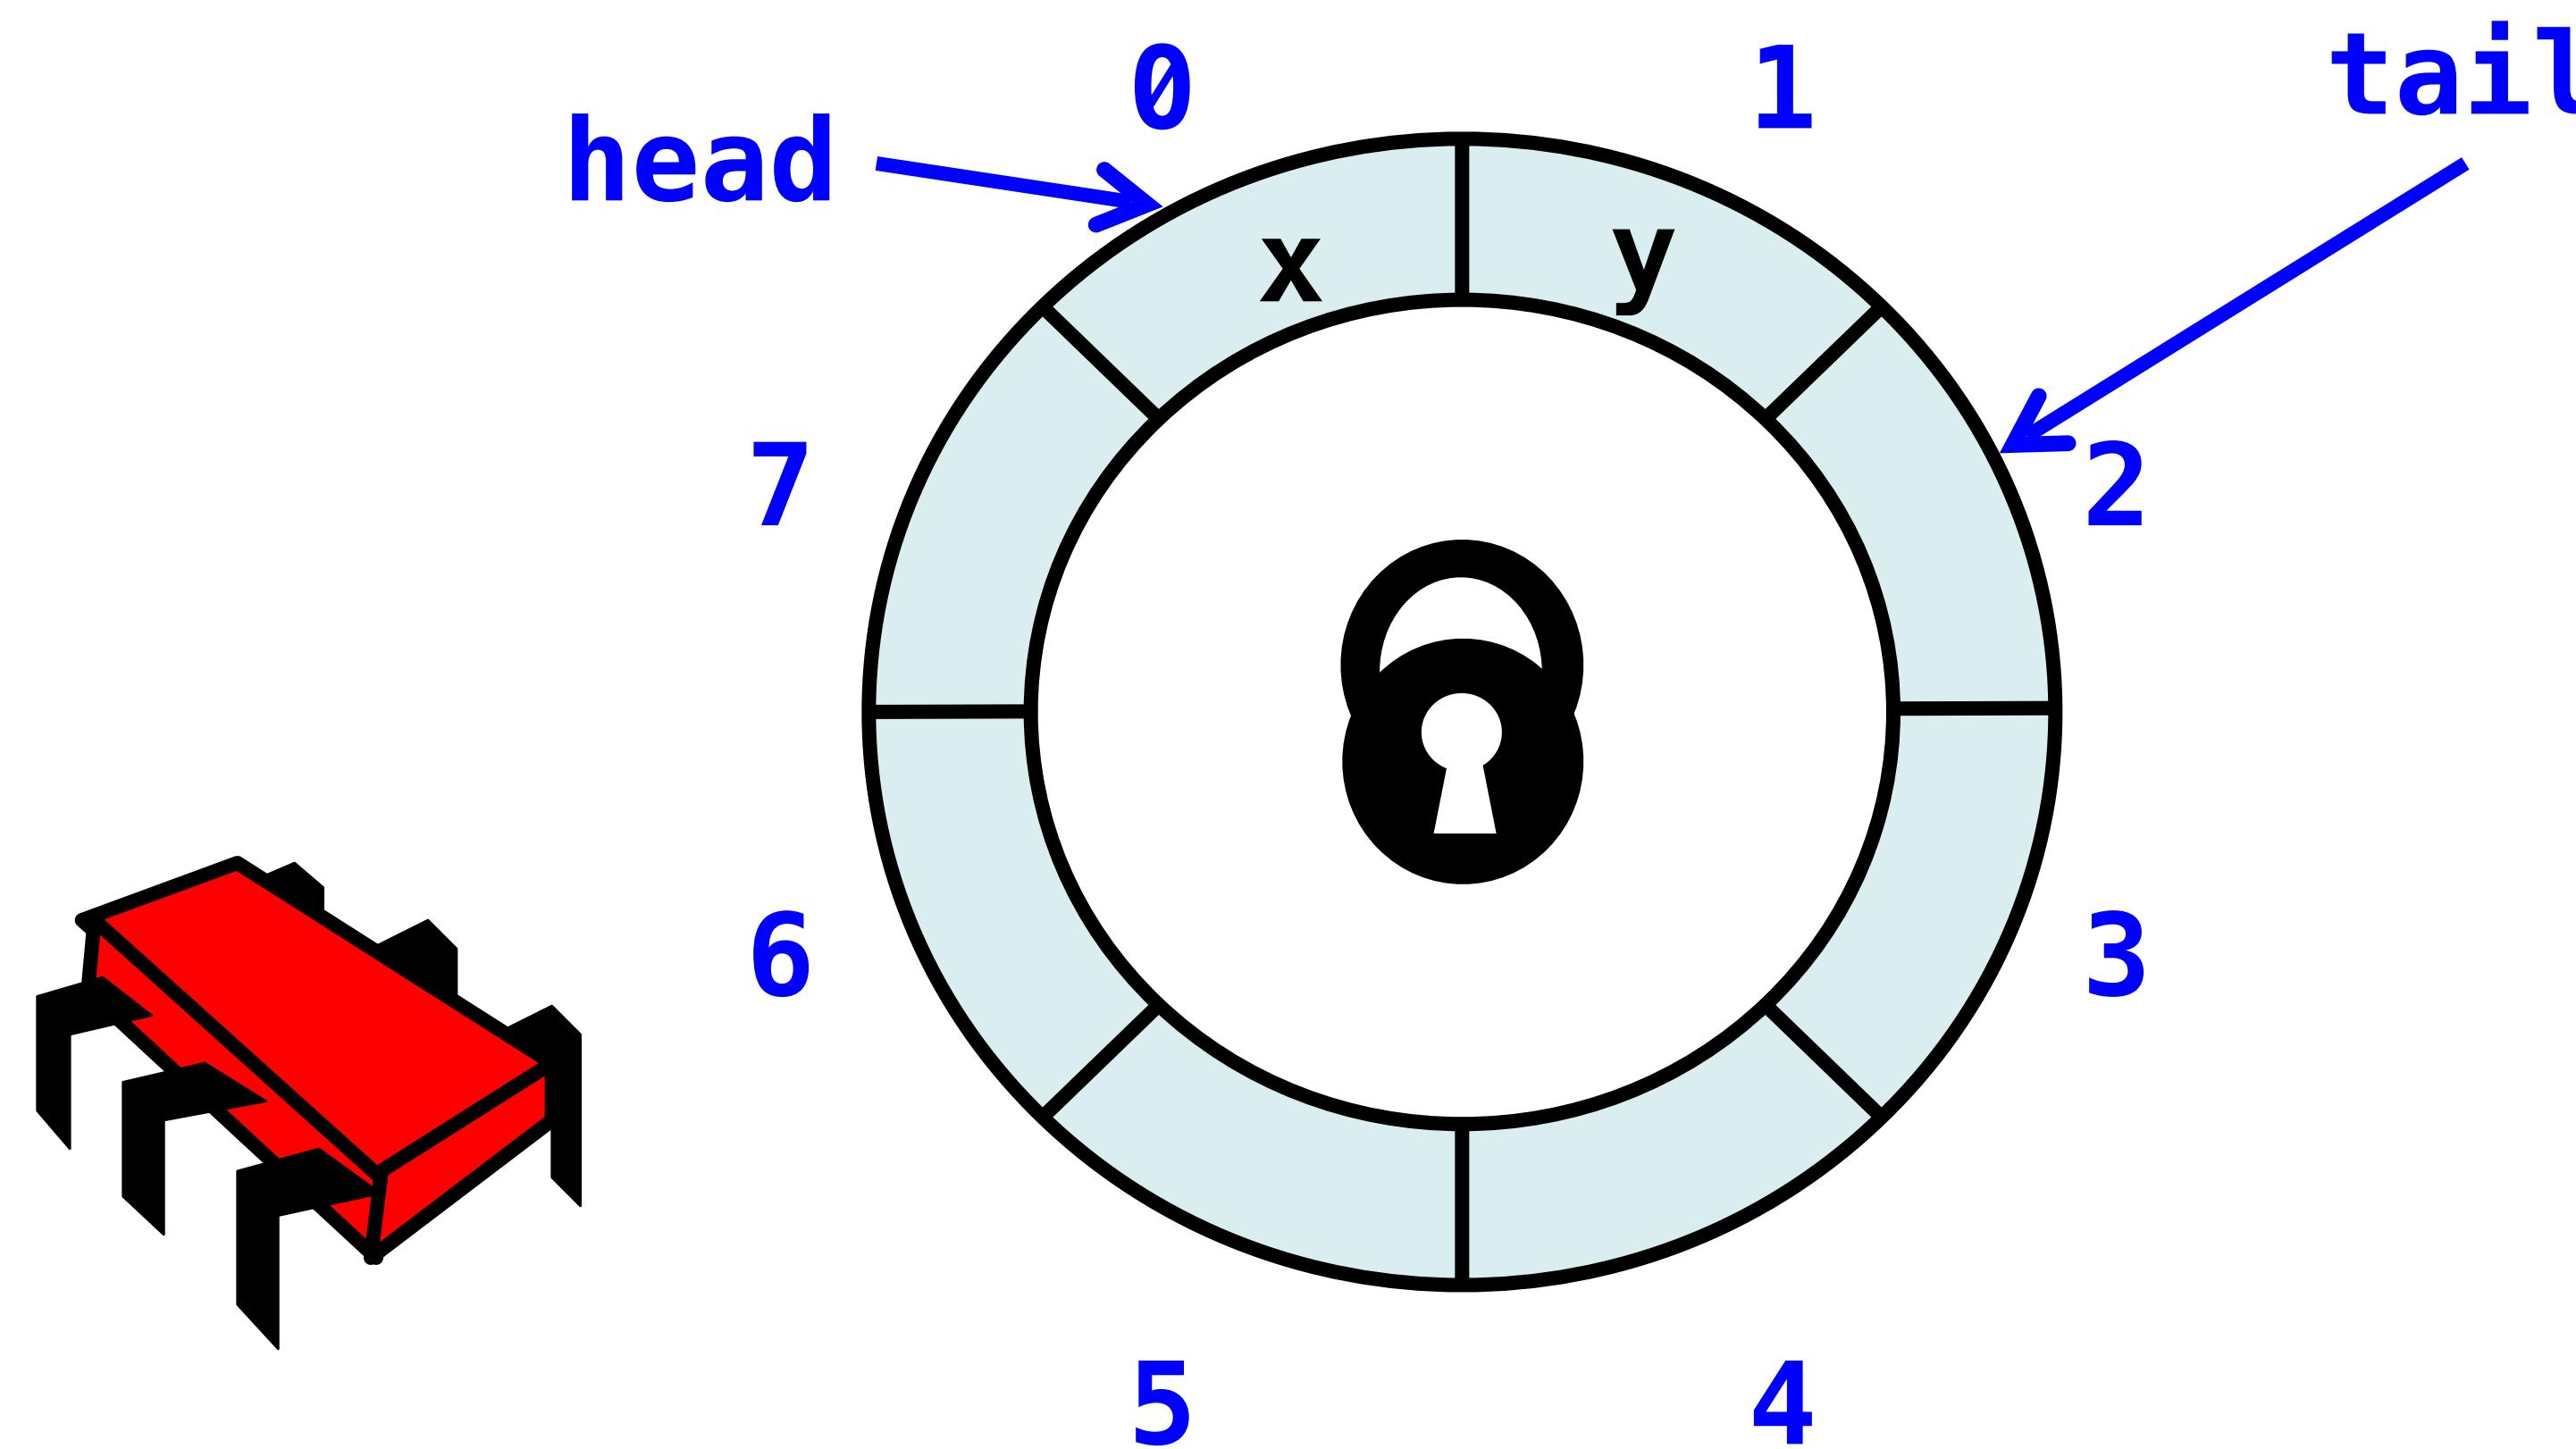
```
type 'a t = {  
    items : 'a option array;  
    capacity : int;  
    mutable head : int;  
    mutable tail : int;  
    lock : Mutex.t;  
}
```

```
let create capacity =  
{  
    items = Array.make capacity None;  
    capacity;  
    head = 0;  
    tail = 0;  
    lock = Mutex.create ();  
}
```

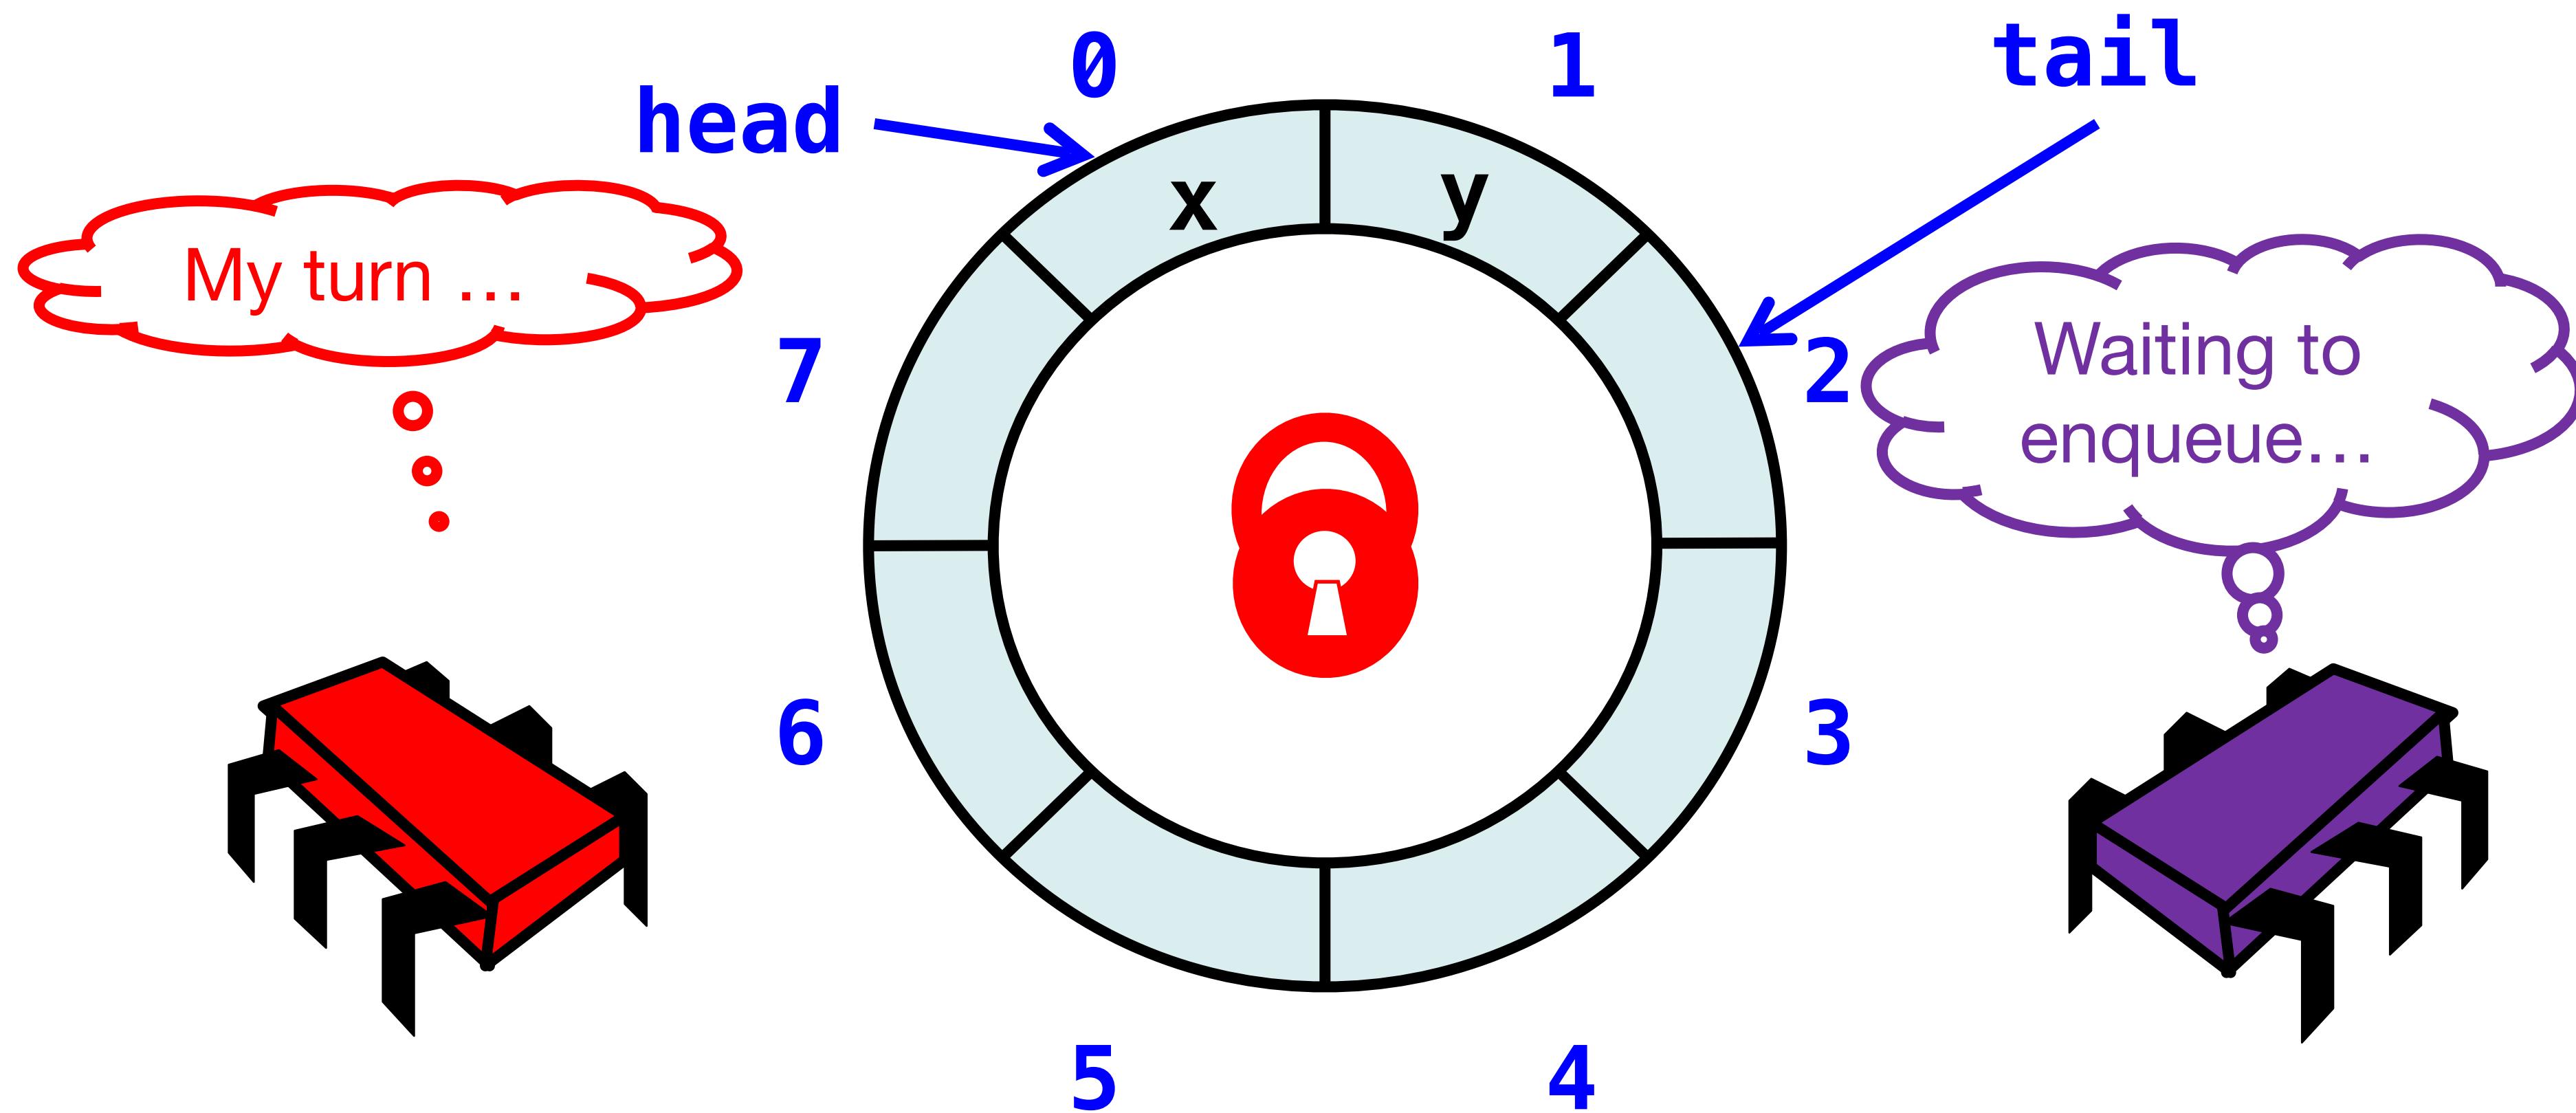
Initially: **head = tail**



# Lock-based Queue – deq( ) operation



# Acquire Lock

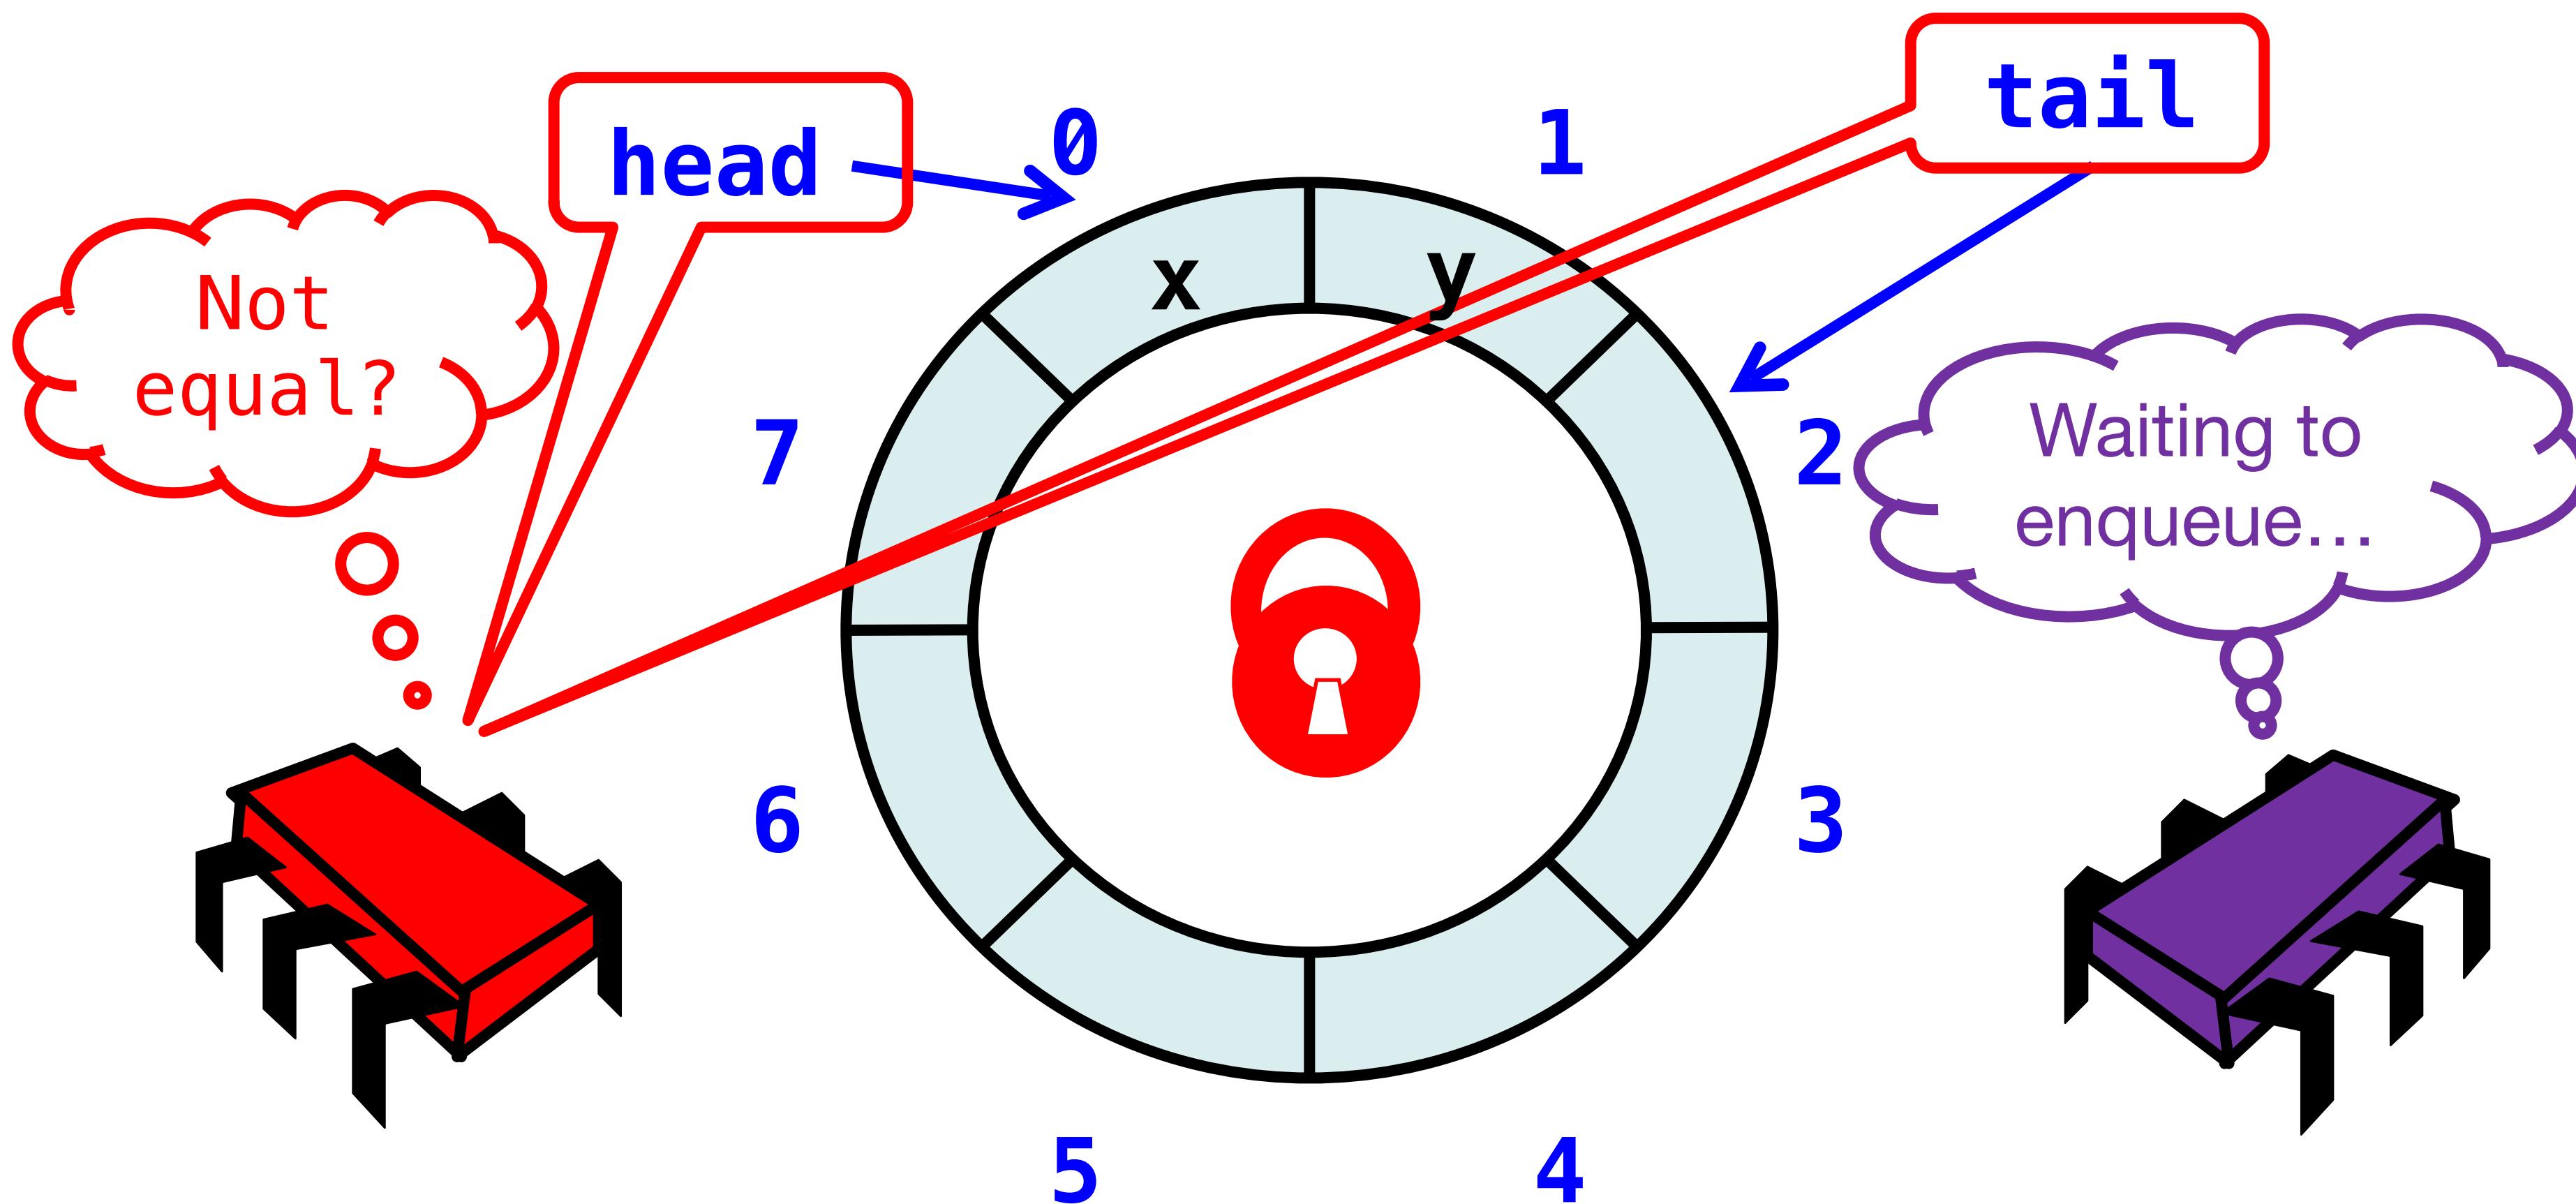


# Acquire Lock

```
let deq q =
  Mutex.lock q.lock;
  try
    if q.tail = q.head then
      raise Empty;
    match q.items.(q.head mod q.capacity) with
    | None -> assert false
    | Some x ->
        q.head <- q.head + 1;
        Mutex.unlock q.lock;
        x
  with e ->
    Mutex.unlock q.lock;
    raise e
```

Acquire lock at  
method start

# Check if non empty



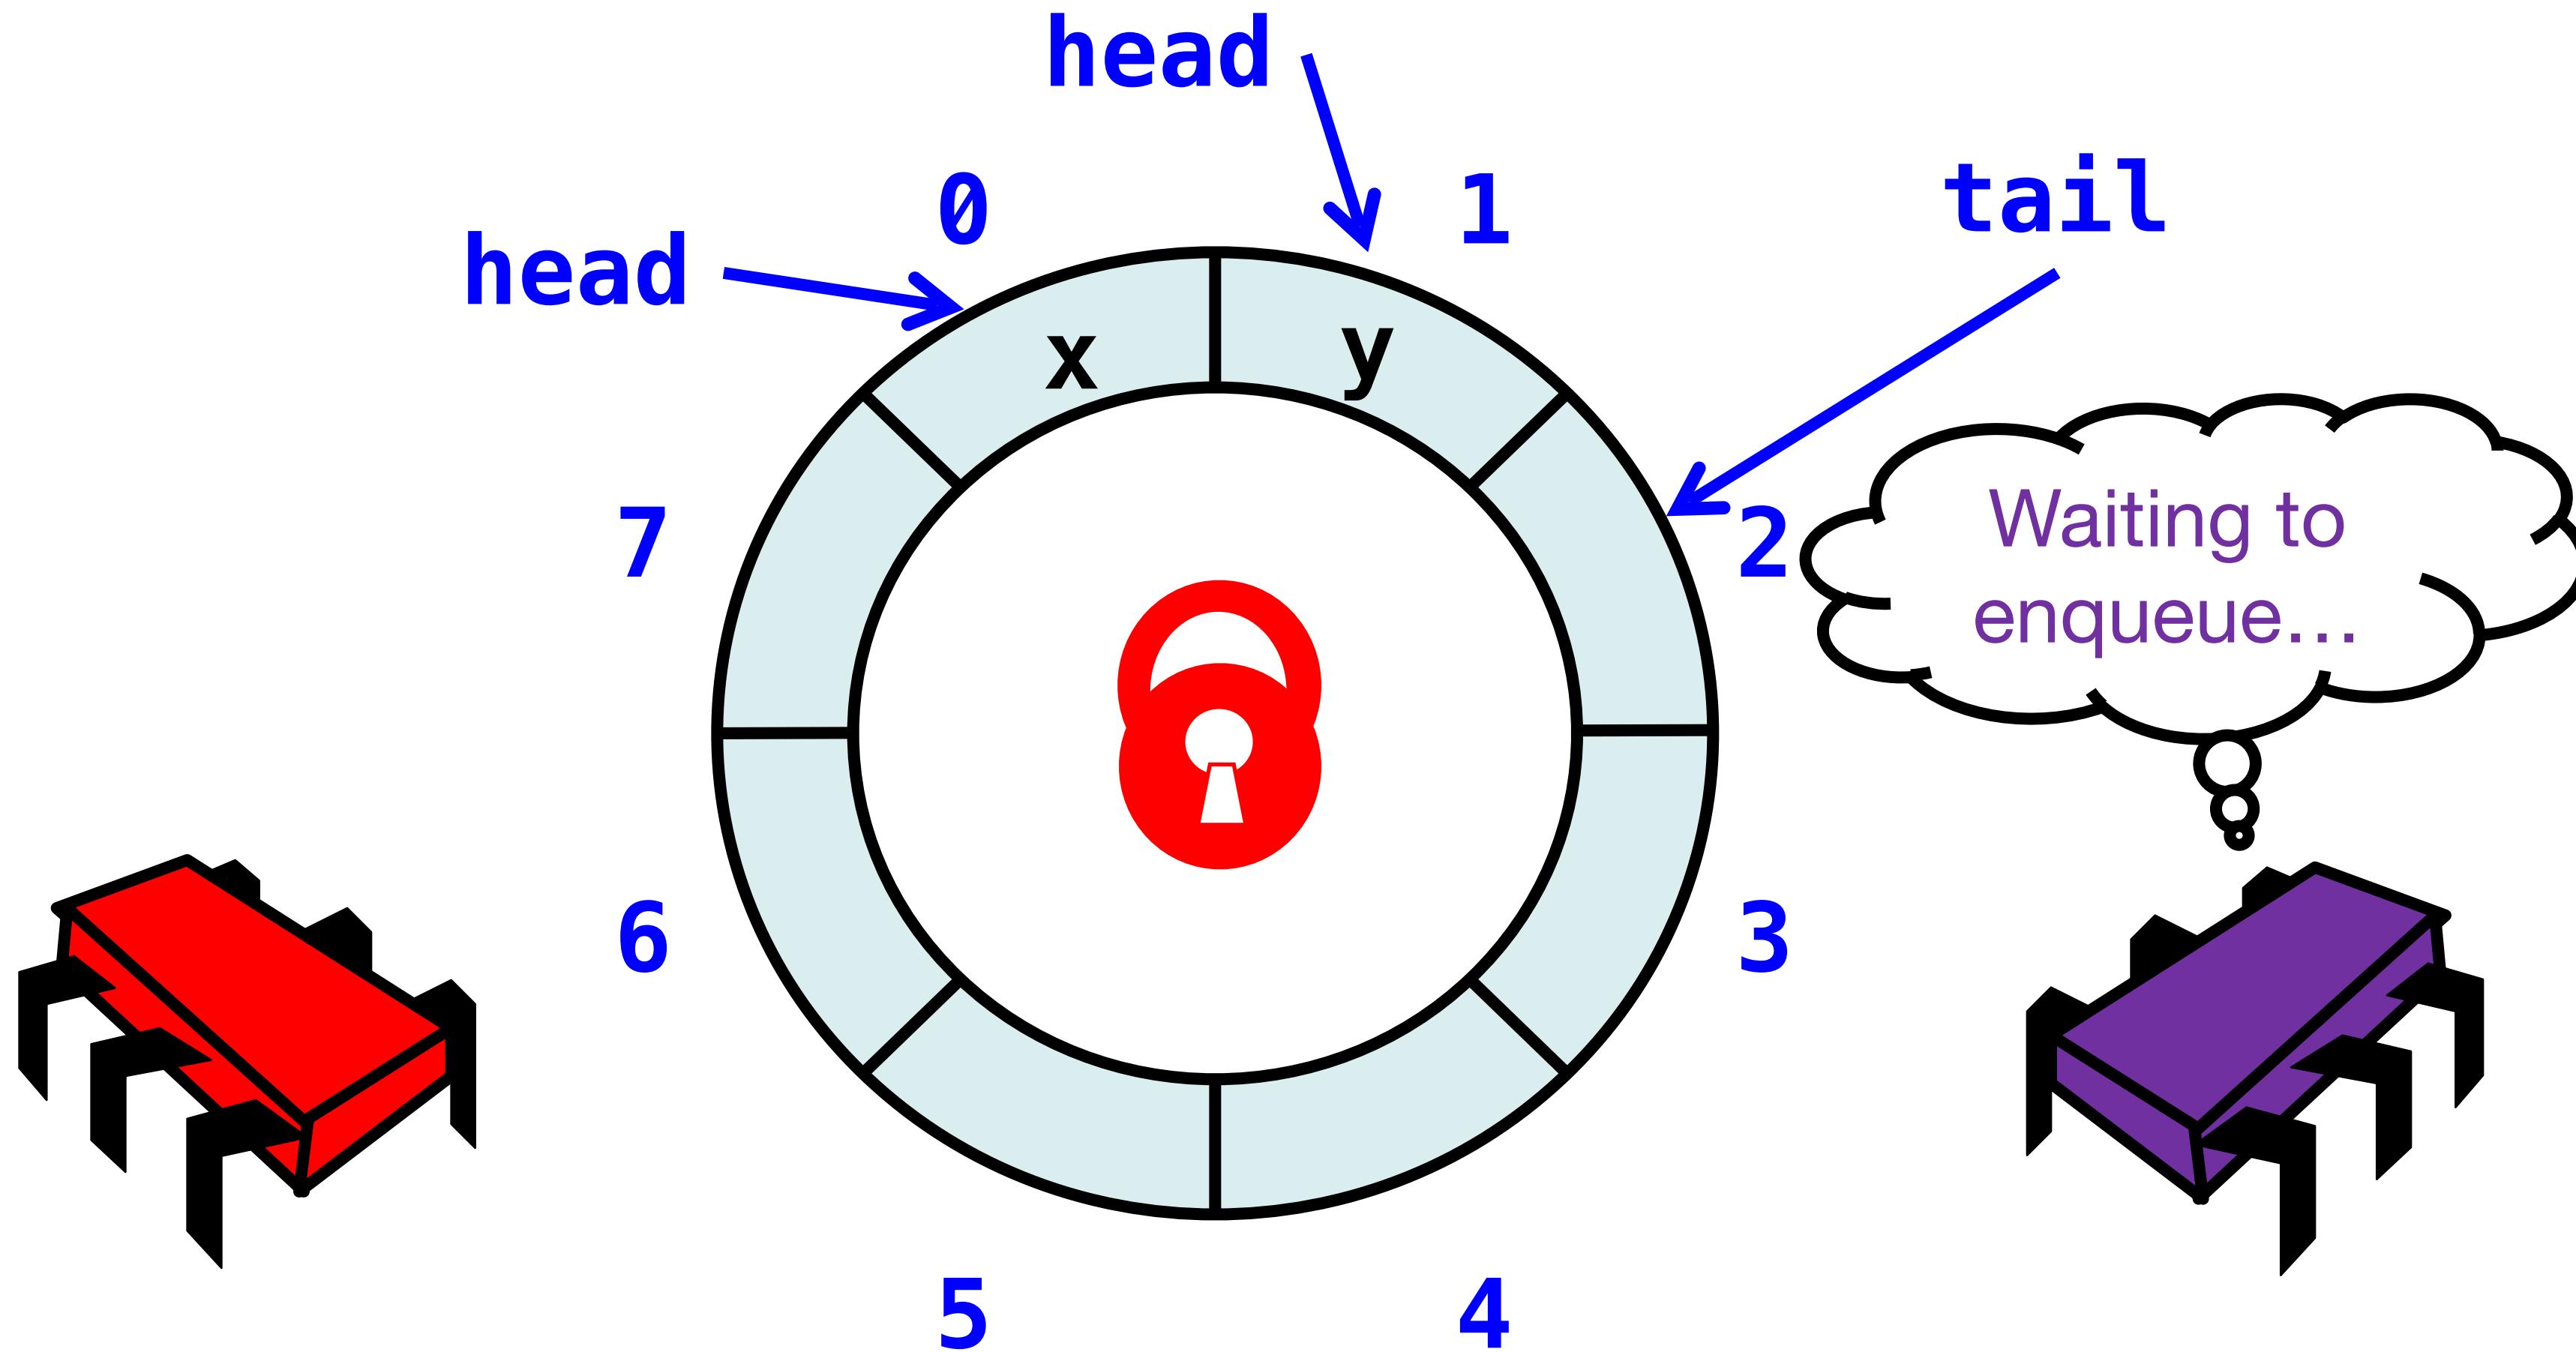
# Check if non empty

```
let deq q =
  Mutex.lock q.lock;
  try
    if q.tail = q.head then
      raise Empty;
    match q.items.(q.head mod q.capacity) with
    | None -> assert false
    | Some x ->
        q.head <- q.head + 1;
        Mutex.unlock q.lock;
        x
  with e ->
    Mutex.unlock q.lock;
    raise e
```

If queue empty  
throw exception

In case of  
exceptions,  
lock released  
here

# Modify the queue

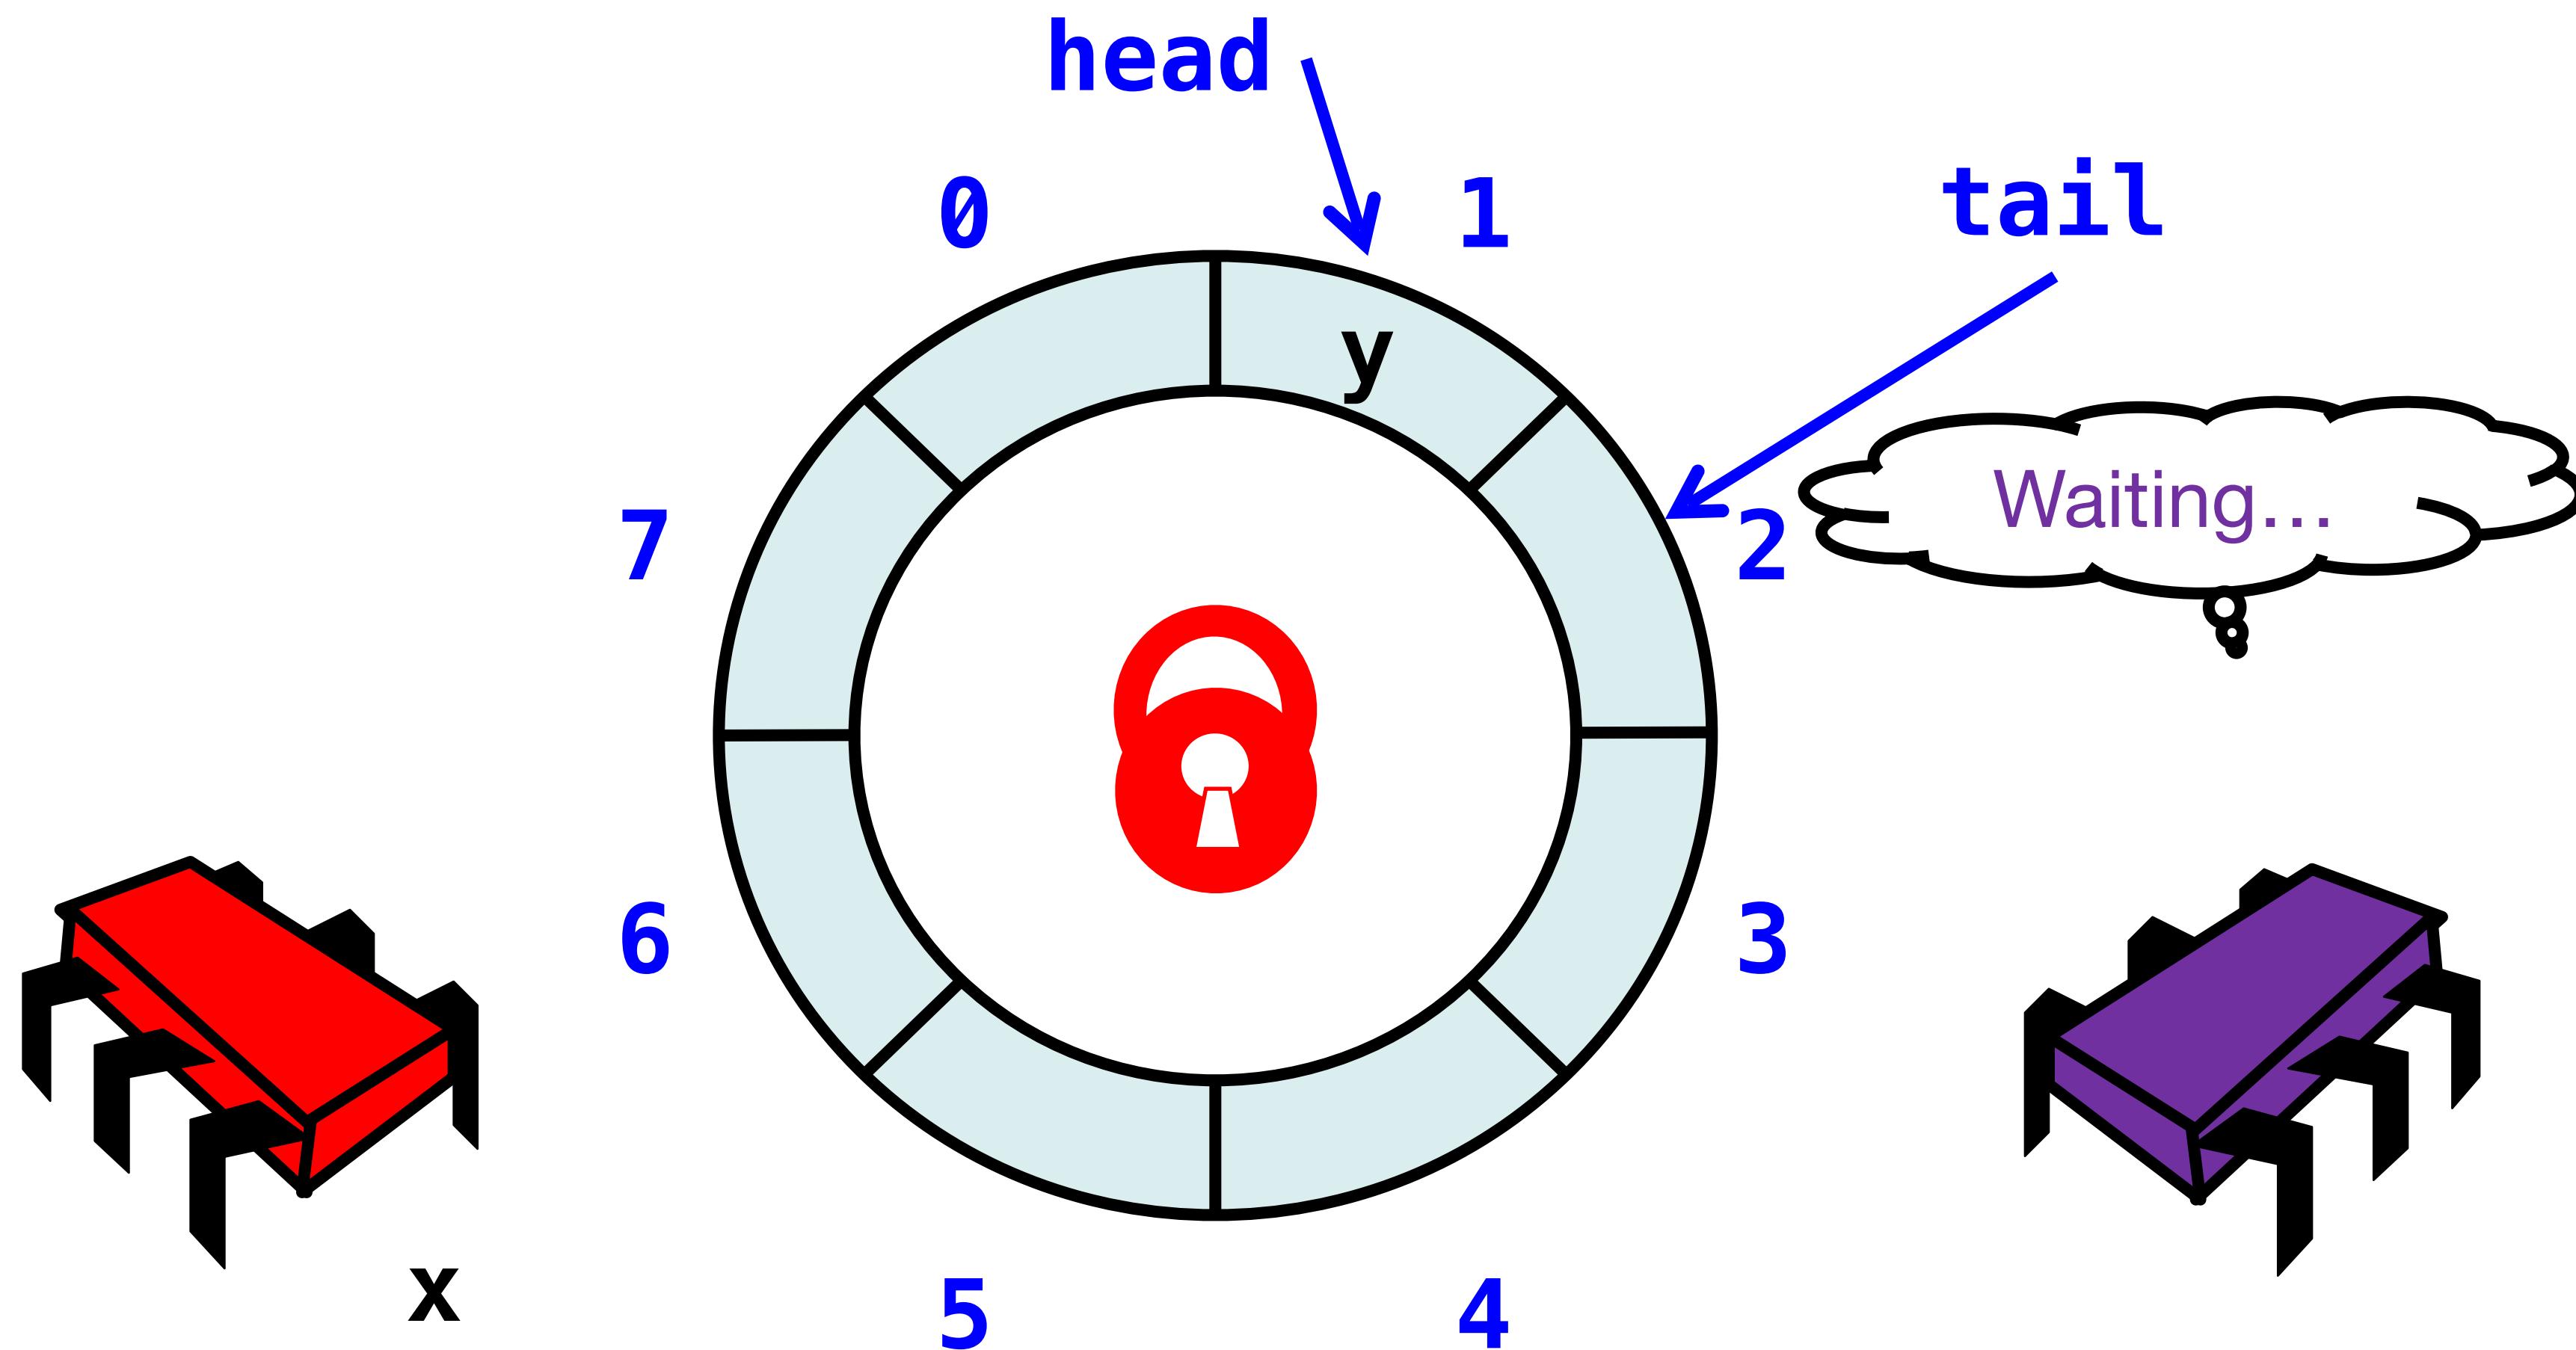


# Modify the queue

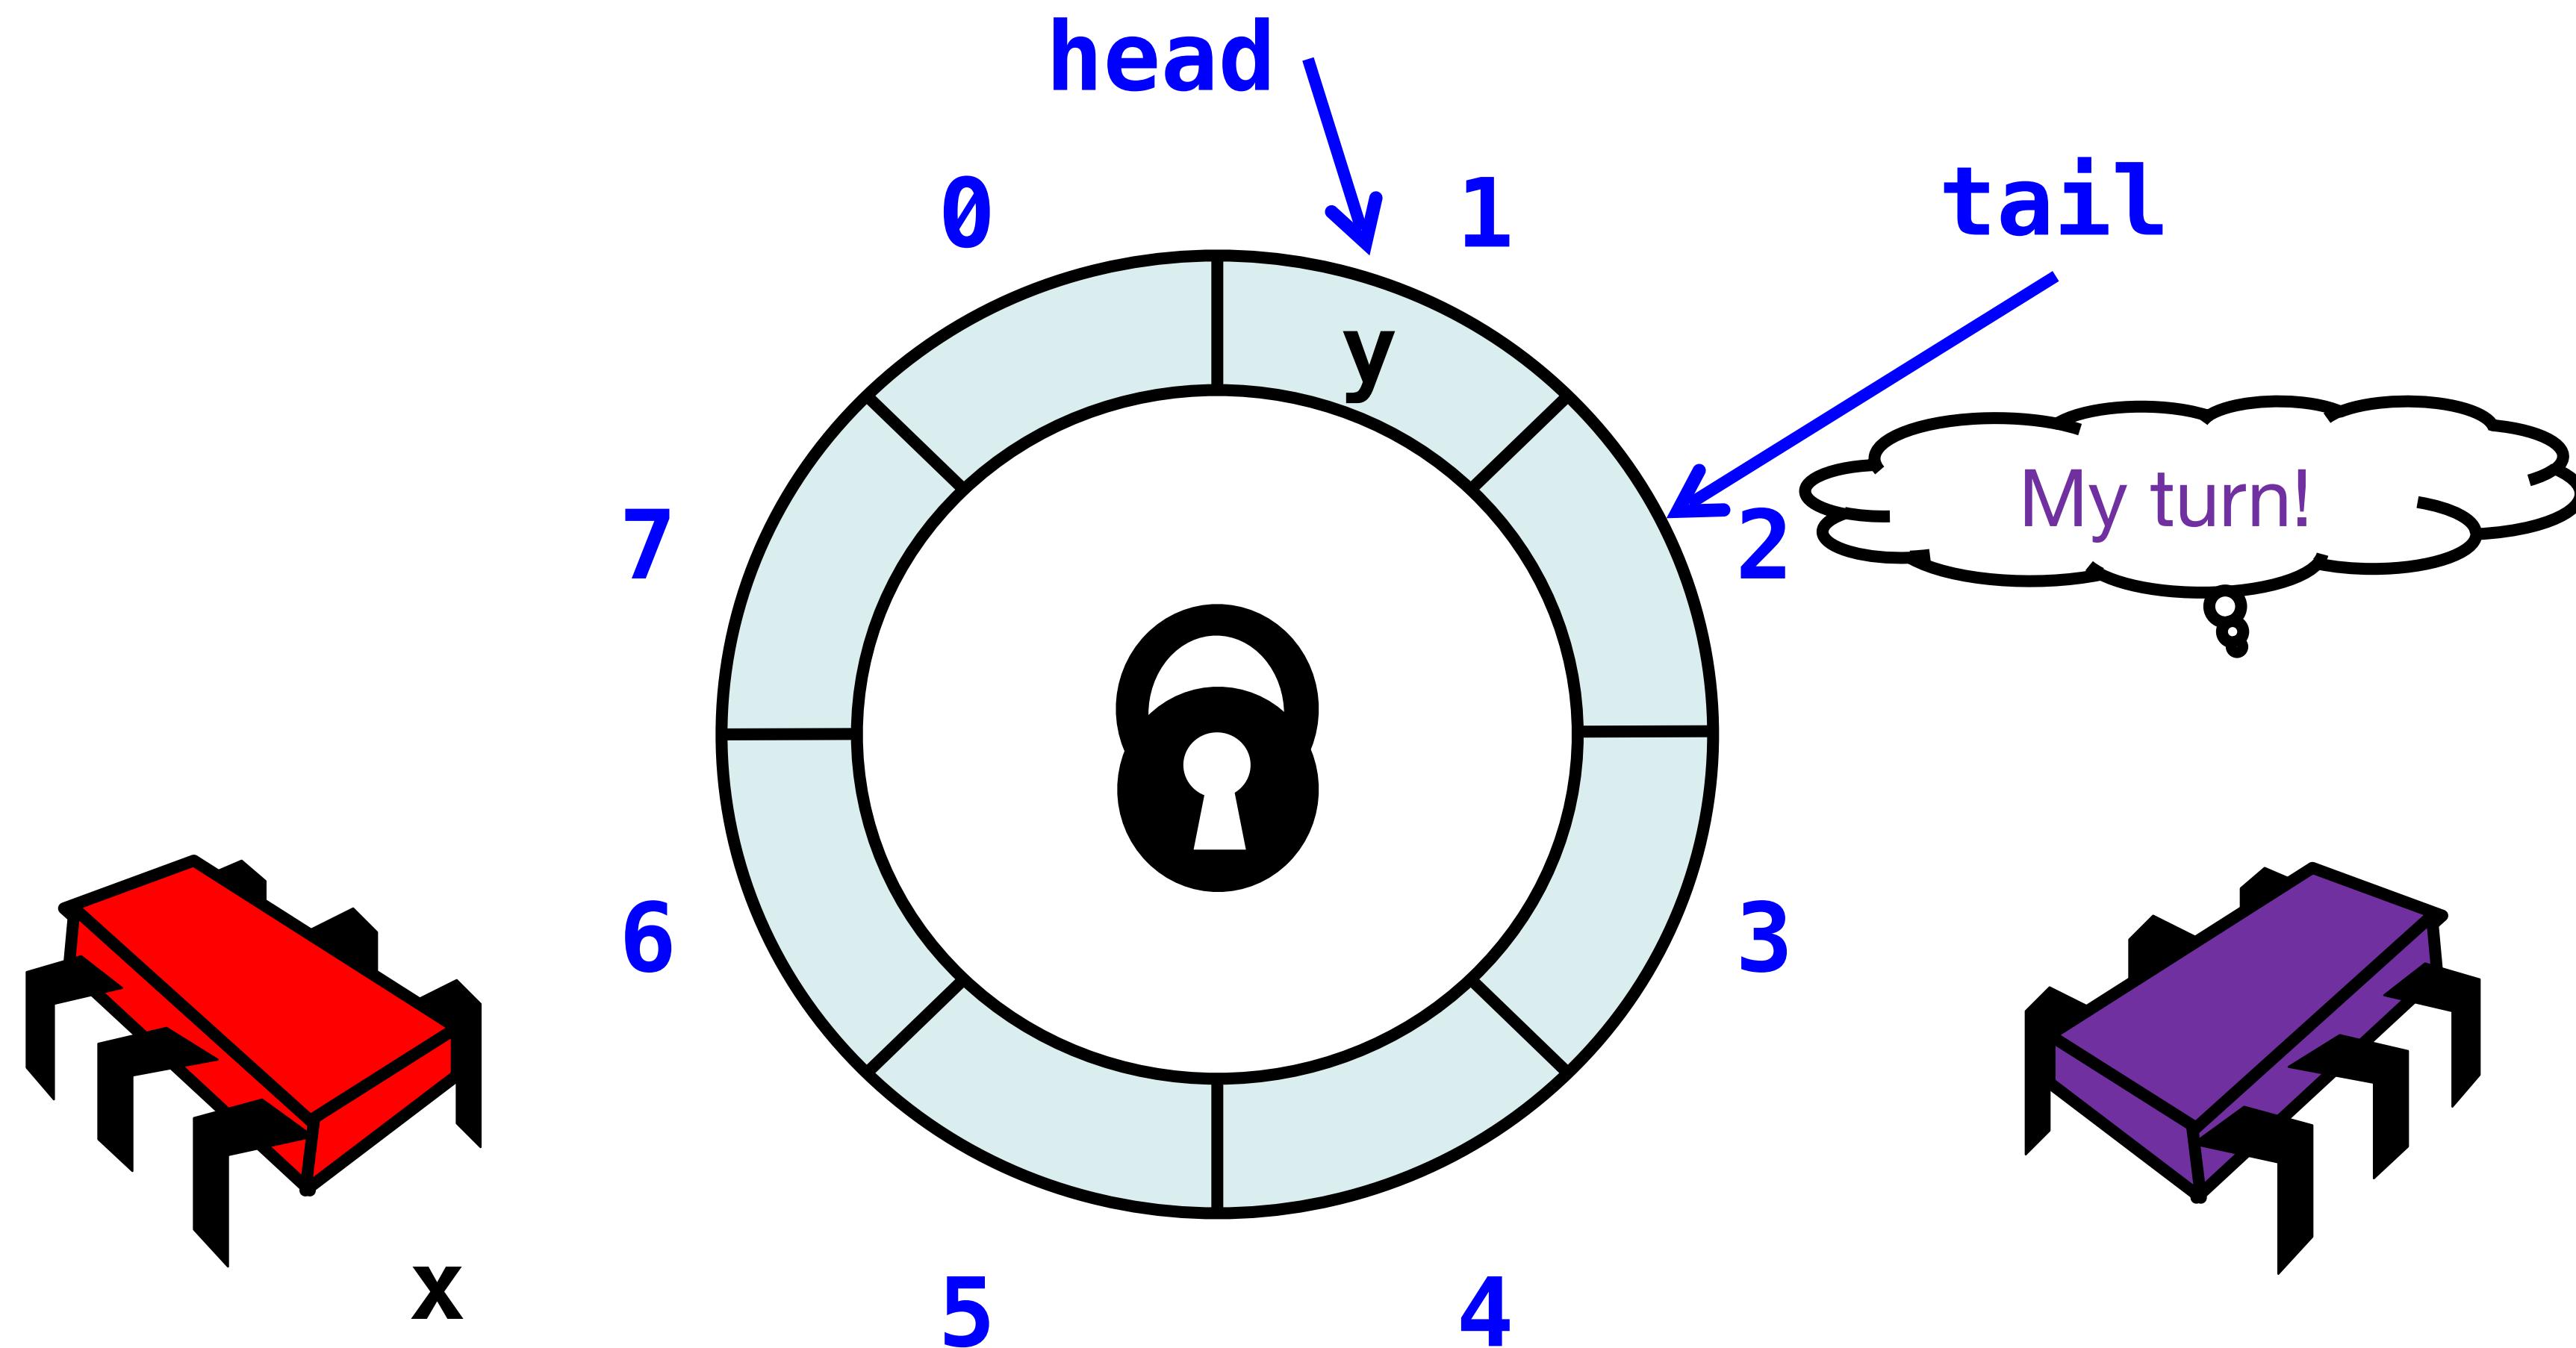
```
let deq q =
  Mutex.lock q.lock;
  try
    if q.tail = q.head then
      raise Empty;
    match q.items.(q.head mod q.capacity) with
    | None -> assert false
    | Some x ->
        q.head <- q.head + 1;
        Mutex.unlock q.lock;
        x
      with e ->
        Mutex.unlock q.lock;
        raise e
```

Queue not empty?  
Remove item “x” and update head

# Release the lock and return item



# Release the lock and return item



# Release the lock and return item

```
let deq q =
  Mutex.lock q.lock;
  try
    if q.tail = q.head then
      raise Empty;
    match q.items.(q.head mod q.capacity) with
    | None -> assert false
    | Some x ->
        q.head <- q.head + 1;
        Mutex.unlock q.lock;
        x
  with e ->
    Mutex.unlock q.lock;
    raise e
```

Unlock and return item “x”

# Implementation – deq()

```
let deq q =
  Mutex.lock q.lock;
  try
    if q.tail = q.head then
      raise Empty;
    match q.items.(q.head mod q.capacity) with
    | None -> assert false
    | Some x ->
        q.head <- q.head + 1;
        Mutex.unlock q.lock;
        x
  with e ->
    Mutex.unlock q.lock;
    raise e
```

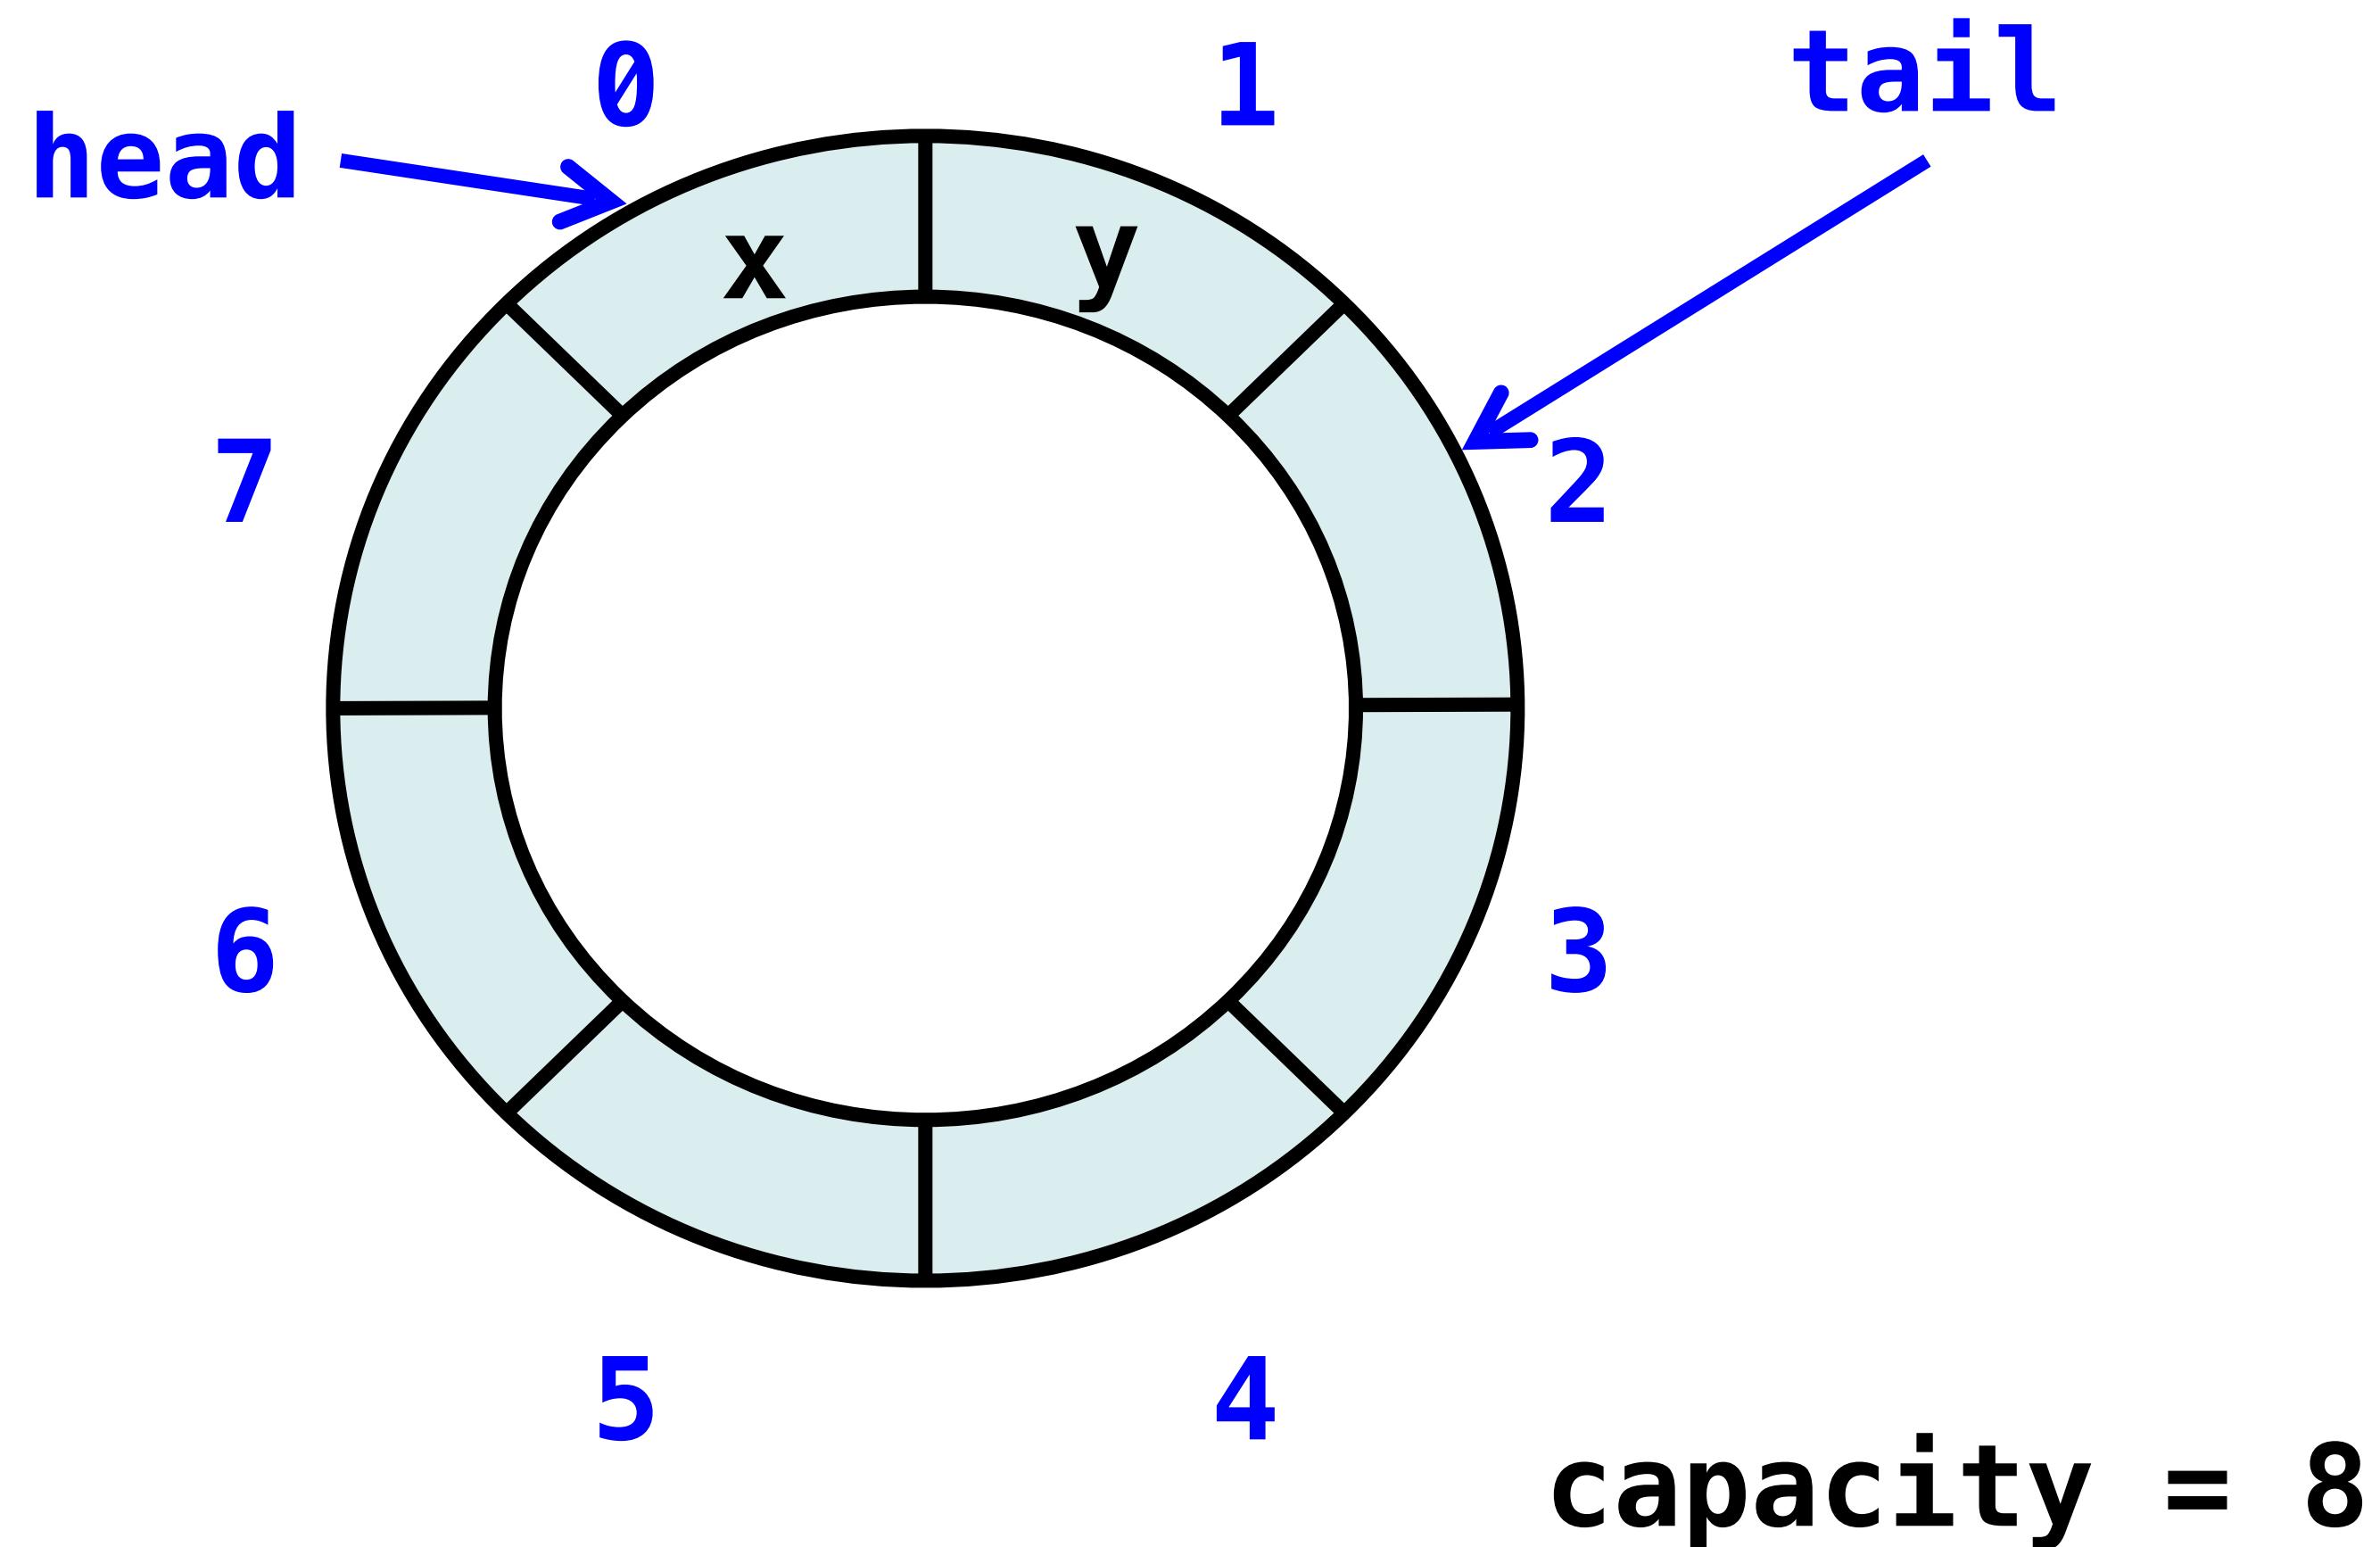
Should be correct because  
modifications are mutually exclusive...

Demo

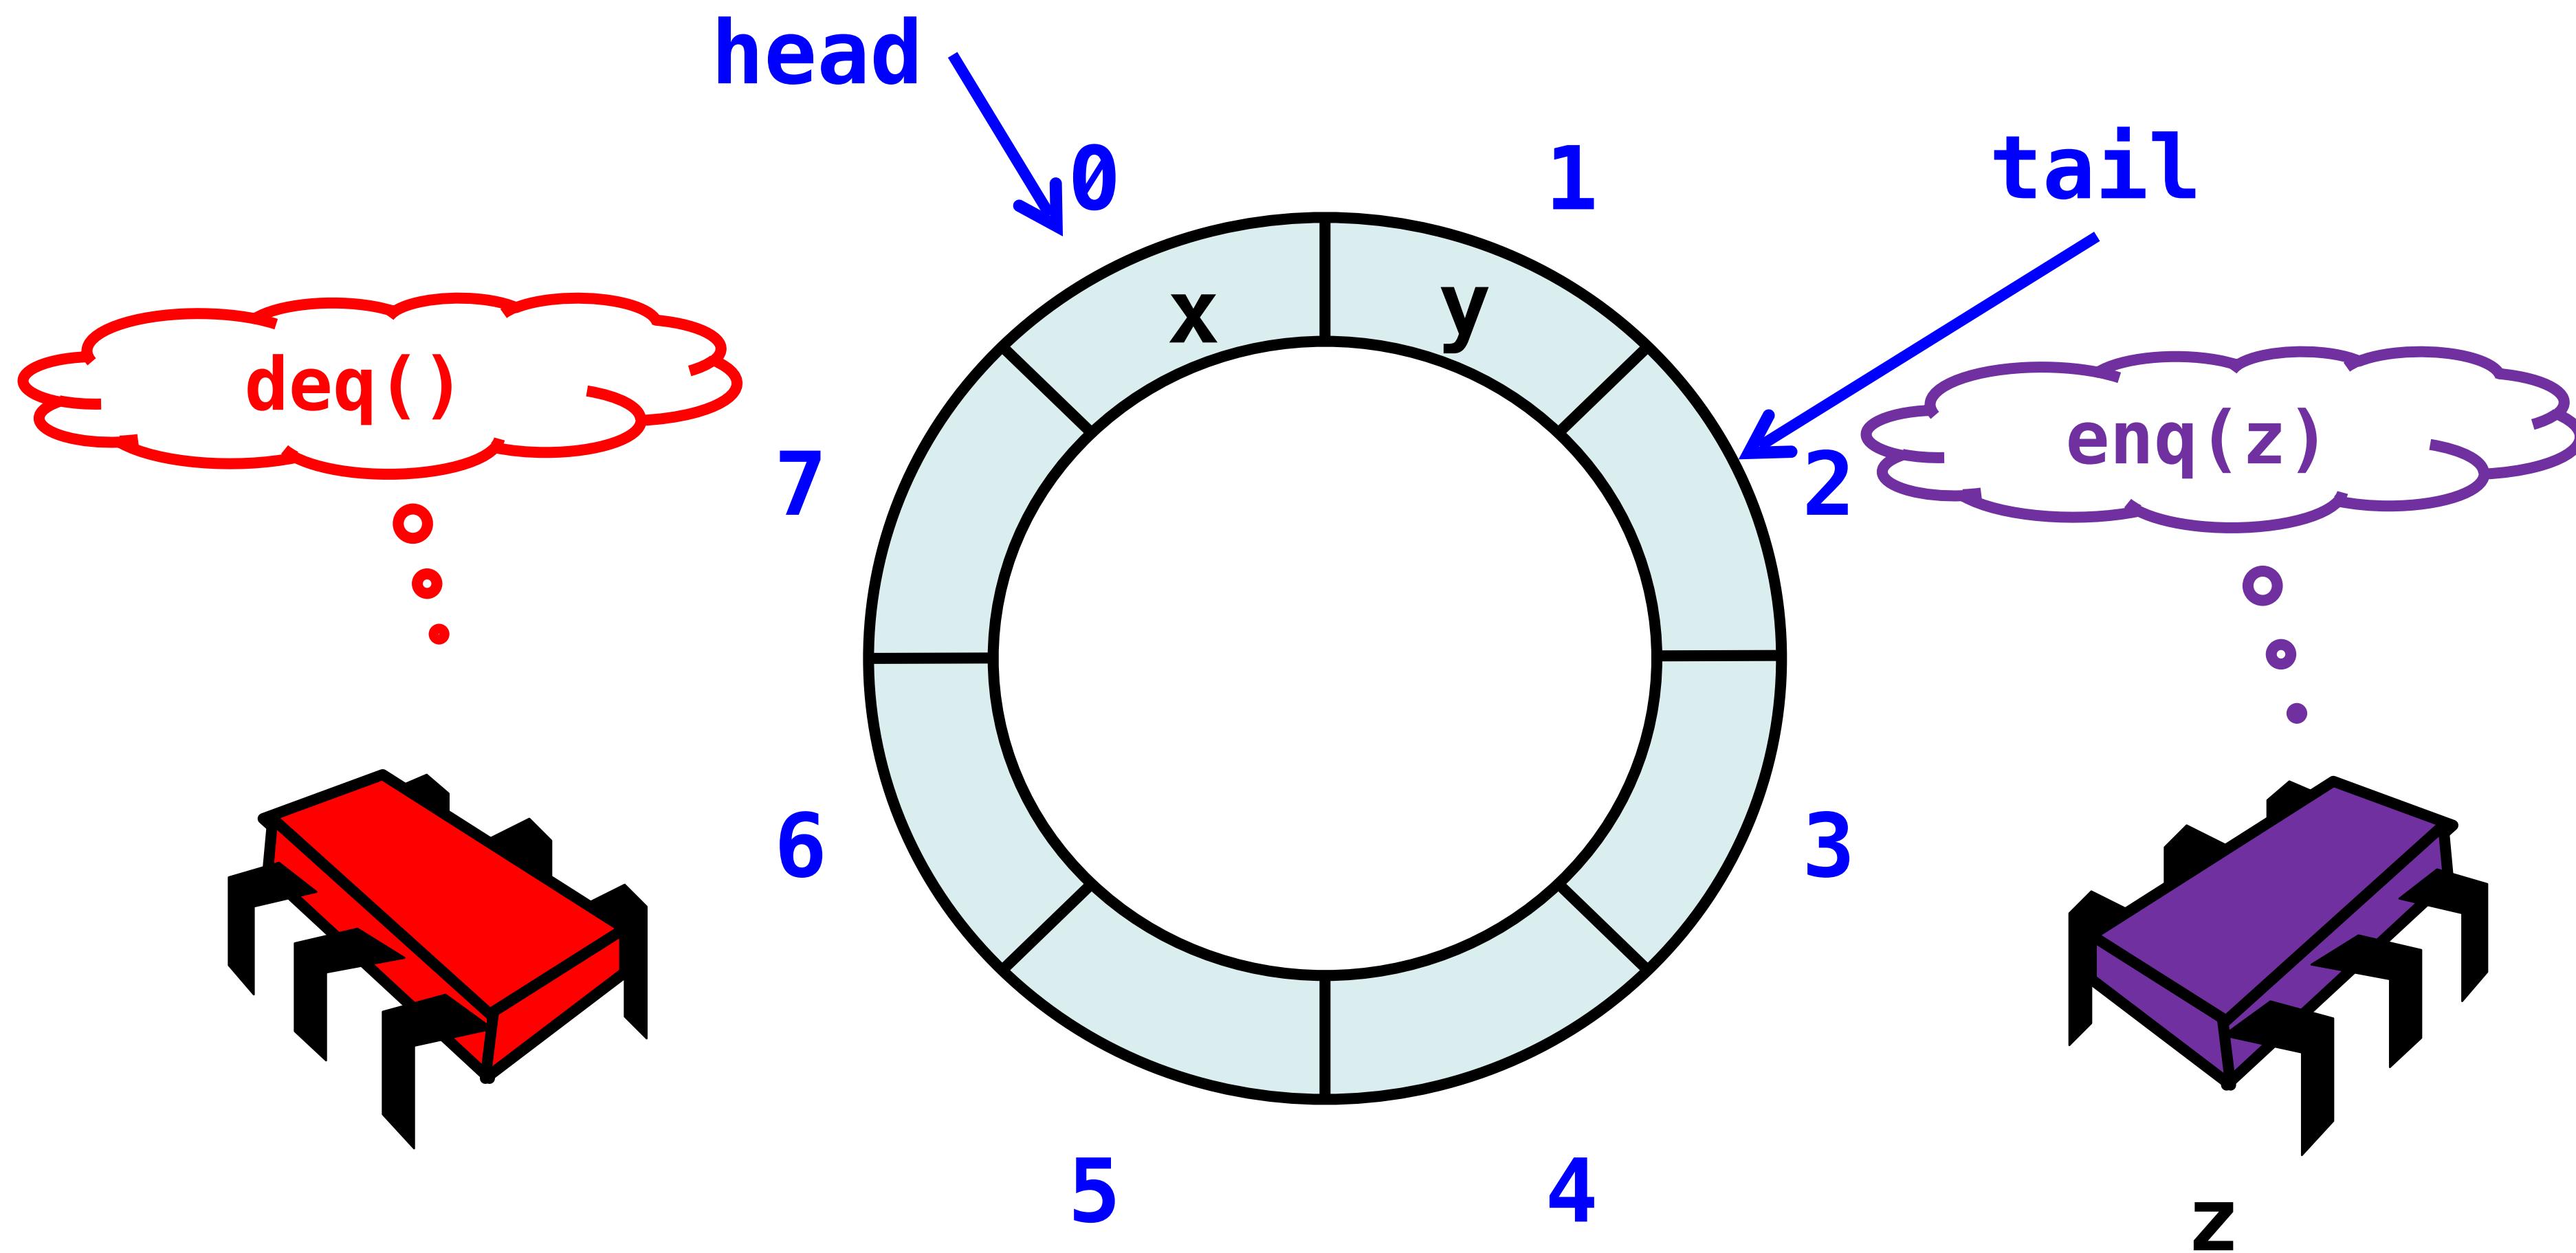
# Consider the following implementation

- The same thing without mutual exclusion
- For simplicity, only **two** threads
  - One thread **enq only**
  - The other **deq only**

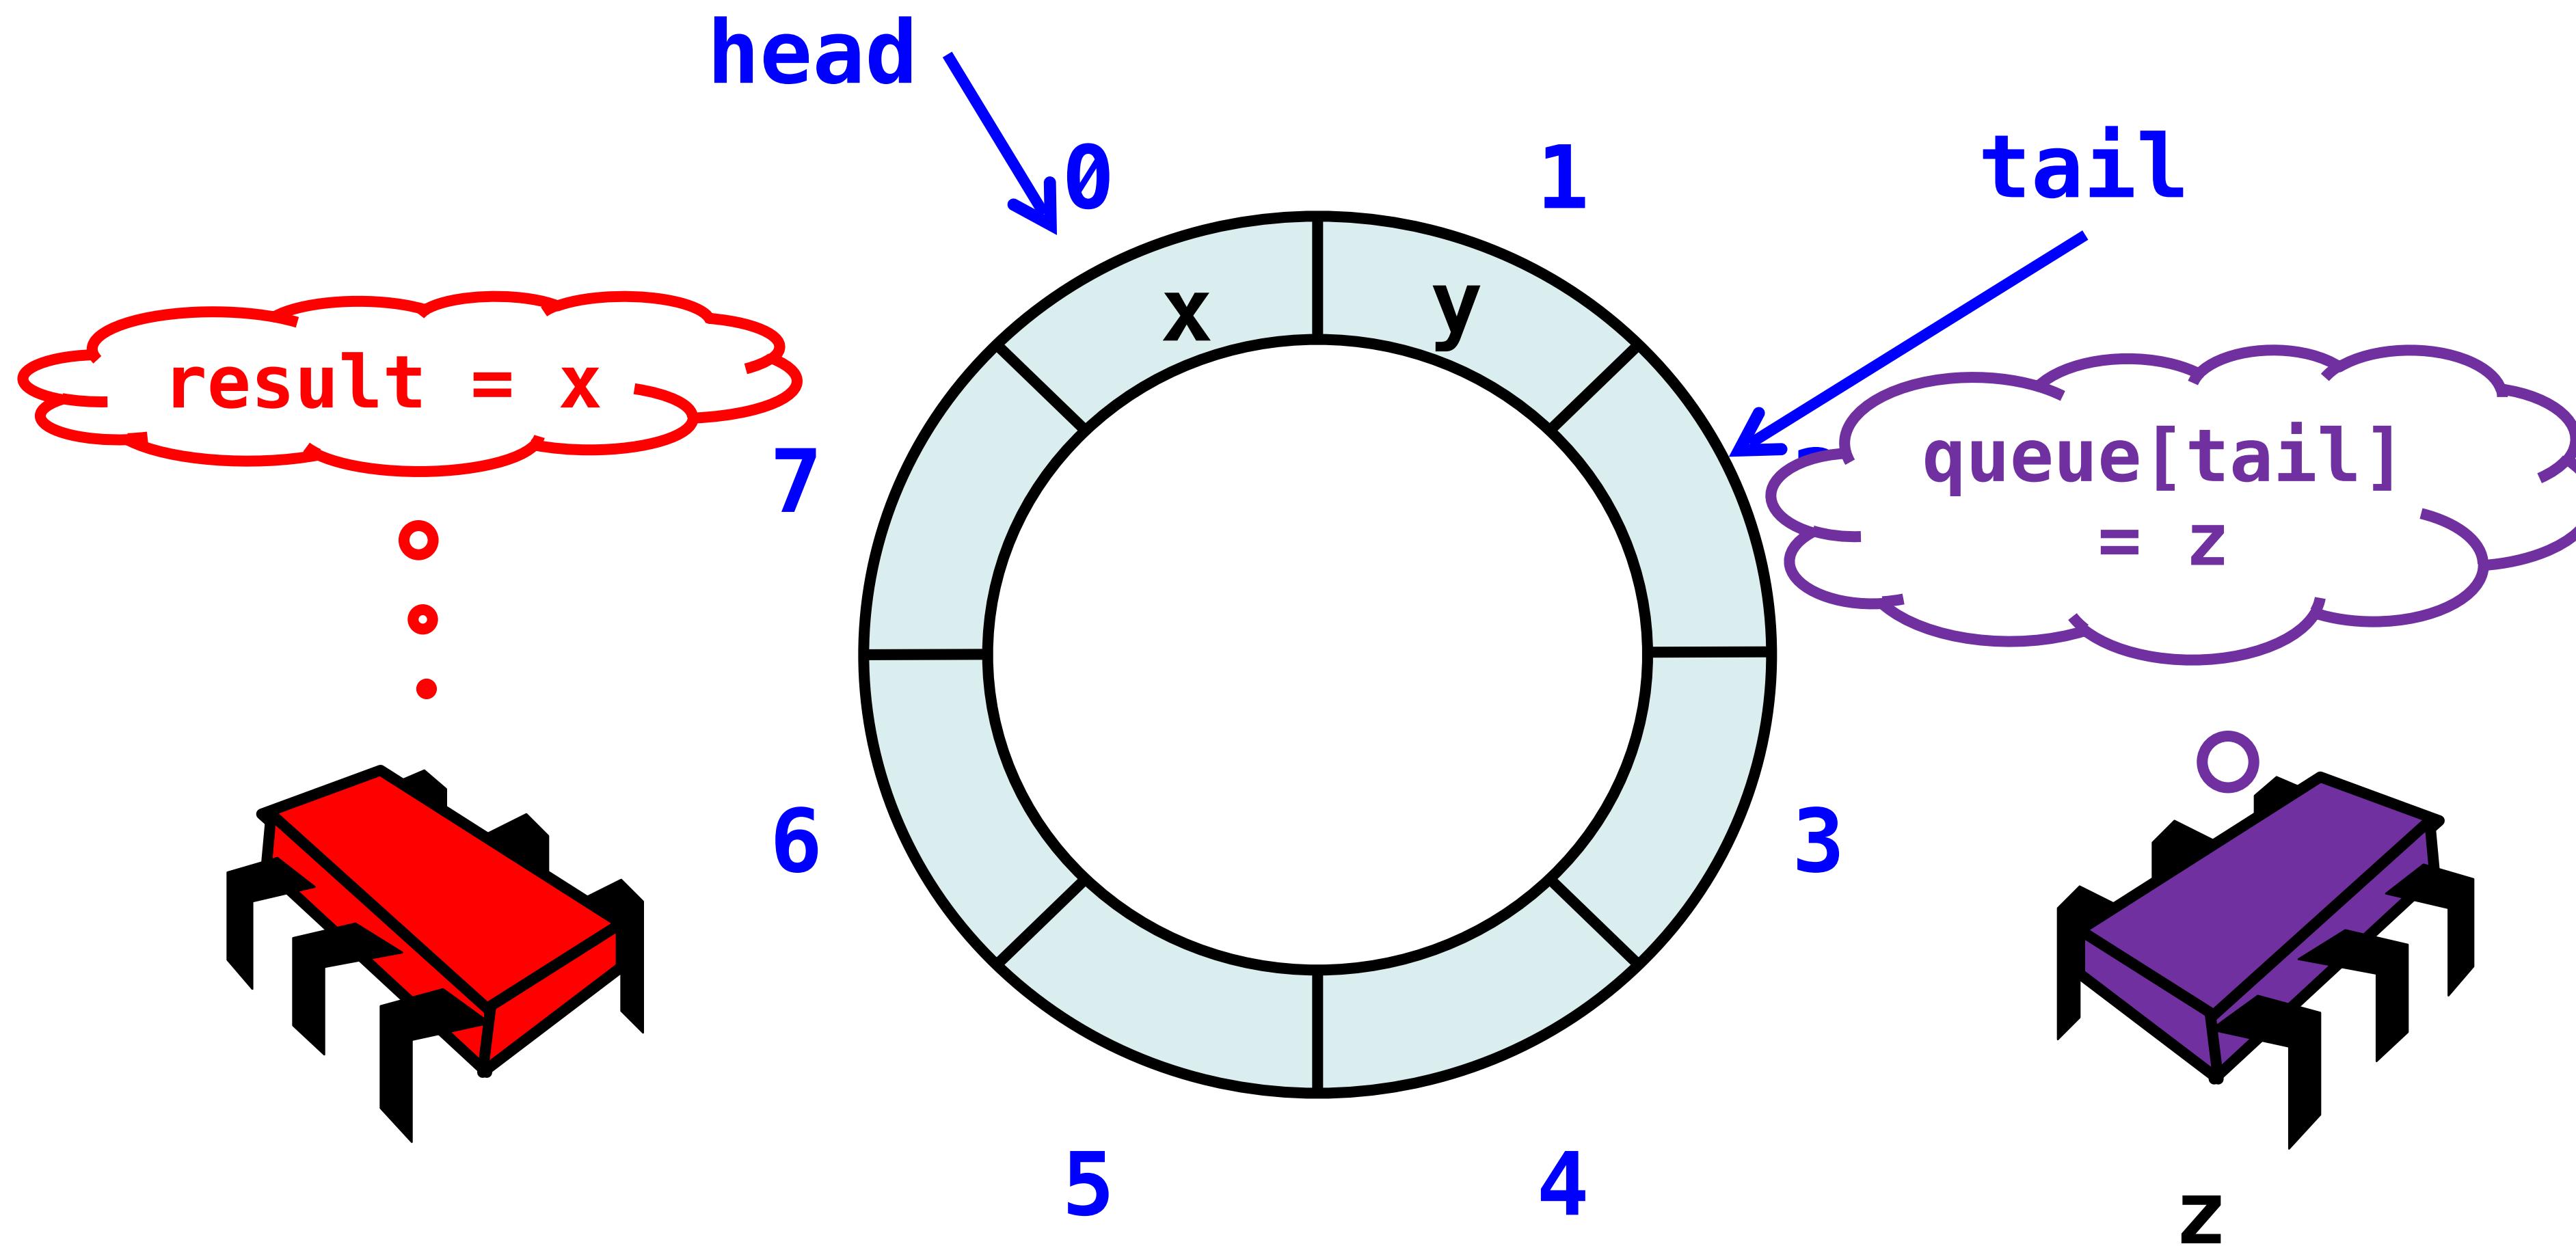
# Wait-free 2-thread queue



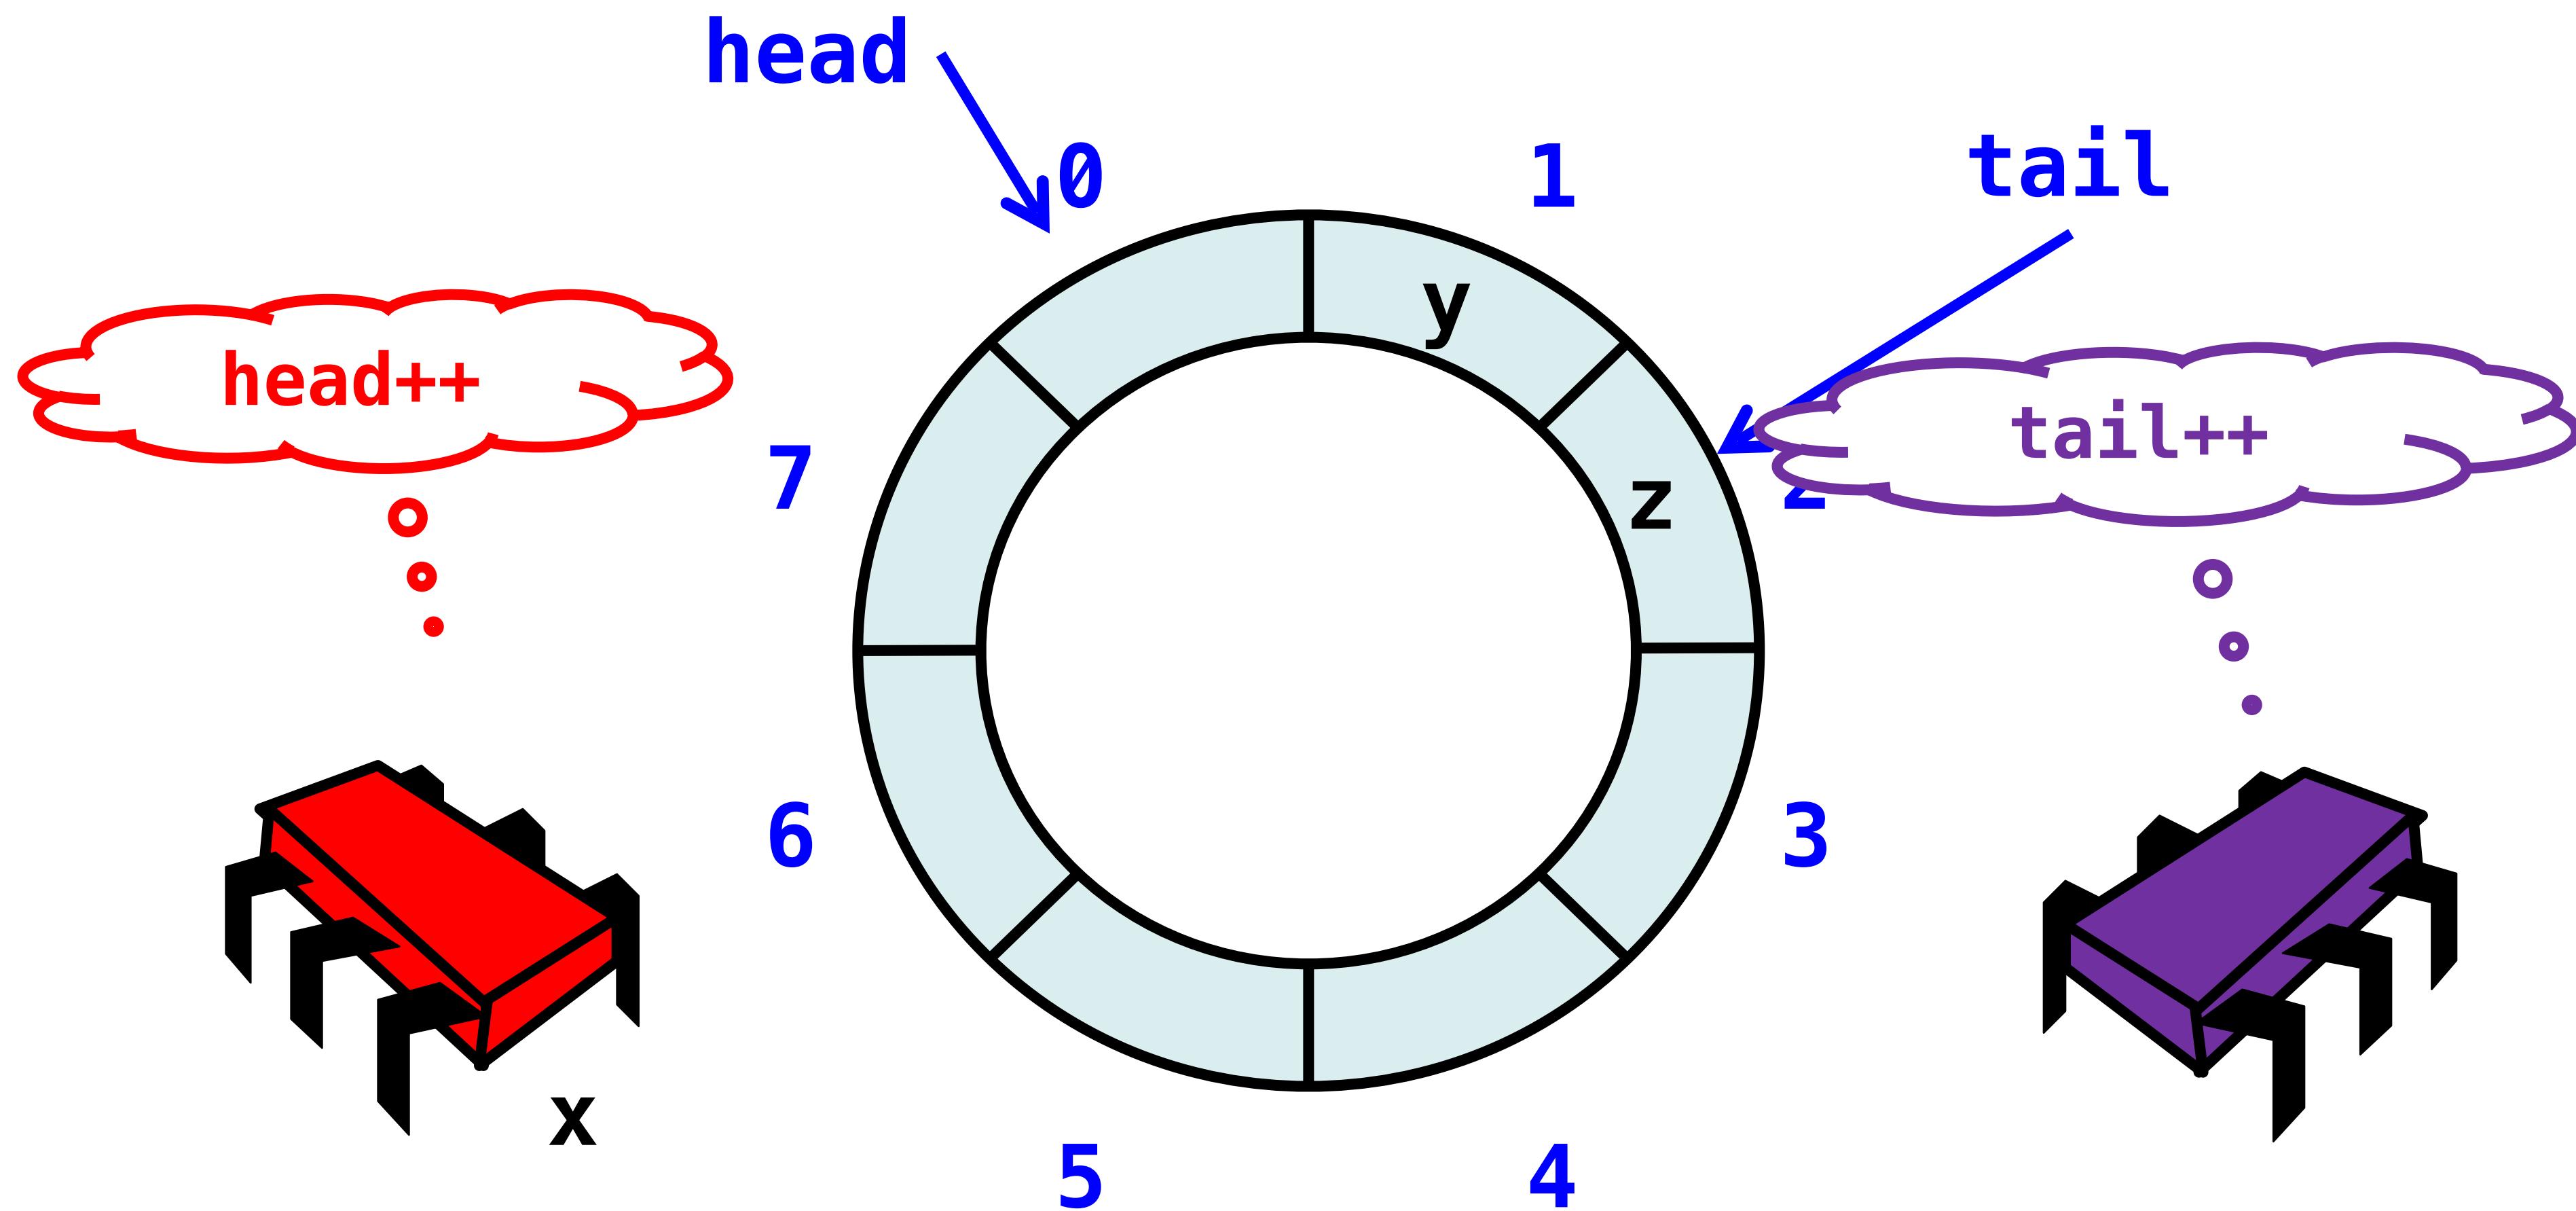
# Wait-free 2-thread queue



# Wait-free 2-thread queue



# Wait-free 2-thread queue



# Wait-free 2-thread queue

**No locks needed!**

```
(** Enqueue – should be called by only ONE thread *)    (** Dequeue – should be called by only ONE thread *)
let enq q x =
  (* Check if queue is full *)
  if q.tail - q.head = q.capacity then
    raise Full;
  (* Write to the array *)
  q.items.(q.tail mod q.capacity) <- Some x;
  (* Advance tail *)
  q.tail <- q.tail + 1

let deq q =
  (* Check if queue is empty *)
  if q.tail = q.head then
    raise Empty;
  (* Read from the array *)
  match q.items.(q.head mod q.capacity) with
  | None | Some _ as head :: rest when q.head = q.tail - 1 && q容量 = 1 then
    q.items <- rest;
    q.head <- q.head + 1;
    Some head
  | Some head :: rest when q容量 > 1 then
    q.items <- rest;
    q.head <- q.head + 1;
    Some head
  | None | Some _ as head :: rest when q容量 > 1 then
    q.items <- rest;
    q.head <- q.head + 1;
    None
```

How do we define “correct” when  
modifications are not mutually  
exclusive?

Demo

# **Concurrency Specification**

# What *is* a concurrent queue?

- Need a way to **specify** a concurrent queue object
- Need a way to **prove** that an algorithm implements the object's specification
- Lets talk about object specifications ...

# Correctness and Progress

- In a concurrent setting, we need to specify both the **safety** and the **liveness** properties of an object
- Need a way to define
  - when an implementation is **correct**
  - the conditions under which it guarantees **progress**

Lets begin with correctness

# Sequential Objects

- Each object has a ***state***
  - Usually given by a set of ***fields***
  - Queue example: sequence of items
- Each object has a set of ***methods***
  - Only way to manipulate state
  - Queue example: **enq** and **deq** methods

# Sequential Specifications

- If (*precondition*)
  - the object is in such-and-such a state
  - before you call the method,
- Then (*postcondition*)
  - the method will return a particular value
  - or throw a particular exception.
- and (*postcondition, cont*)
  - the object will be in some other state
  - when the method returns,

# Pre and Post Conditions for Dequeue

- Precondition:
  - Queue is ***non-empty***
- Postcondition:
  - Returns first item in queue
- Postcondition:
  - Removes first item in queue
- Precondition:
  - Queue is ***empty***
- Postcondition:
  - Raises Empty exception
- Postcondition:
  - Queue state is unchanged

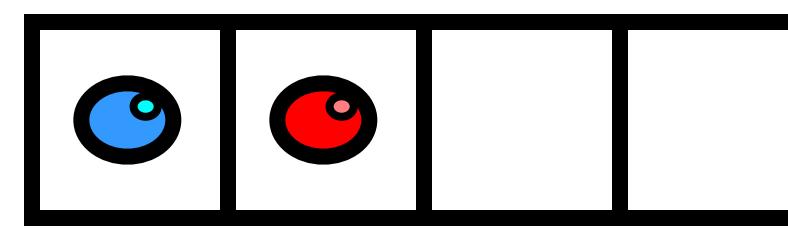
# Why Sequential Specifications Totally Rock

- Interactions among *methods* captured by side-effects on object state
  - State meaningful between method calls
- *Documentation* size is linear in the number of methods
  - Each method described in isolation
- Can add *new methods*
  - Without changing descriptions of old methods

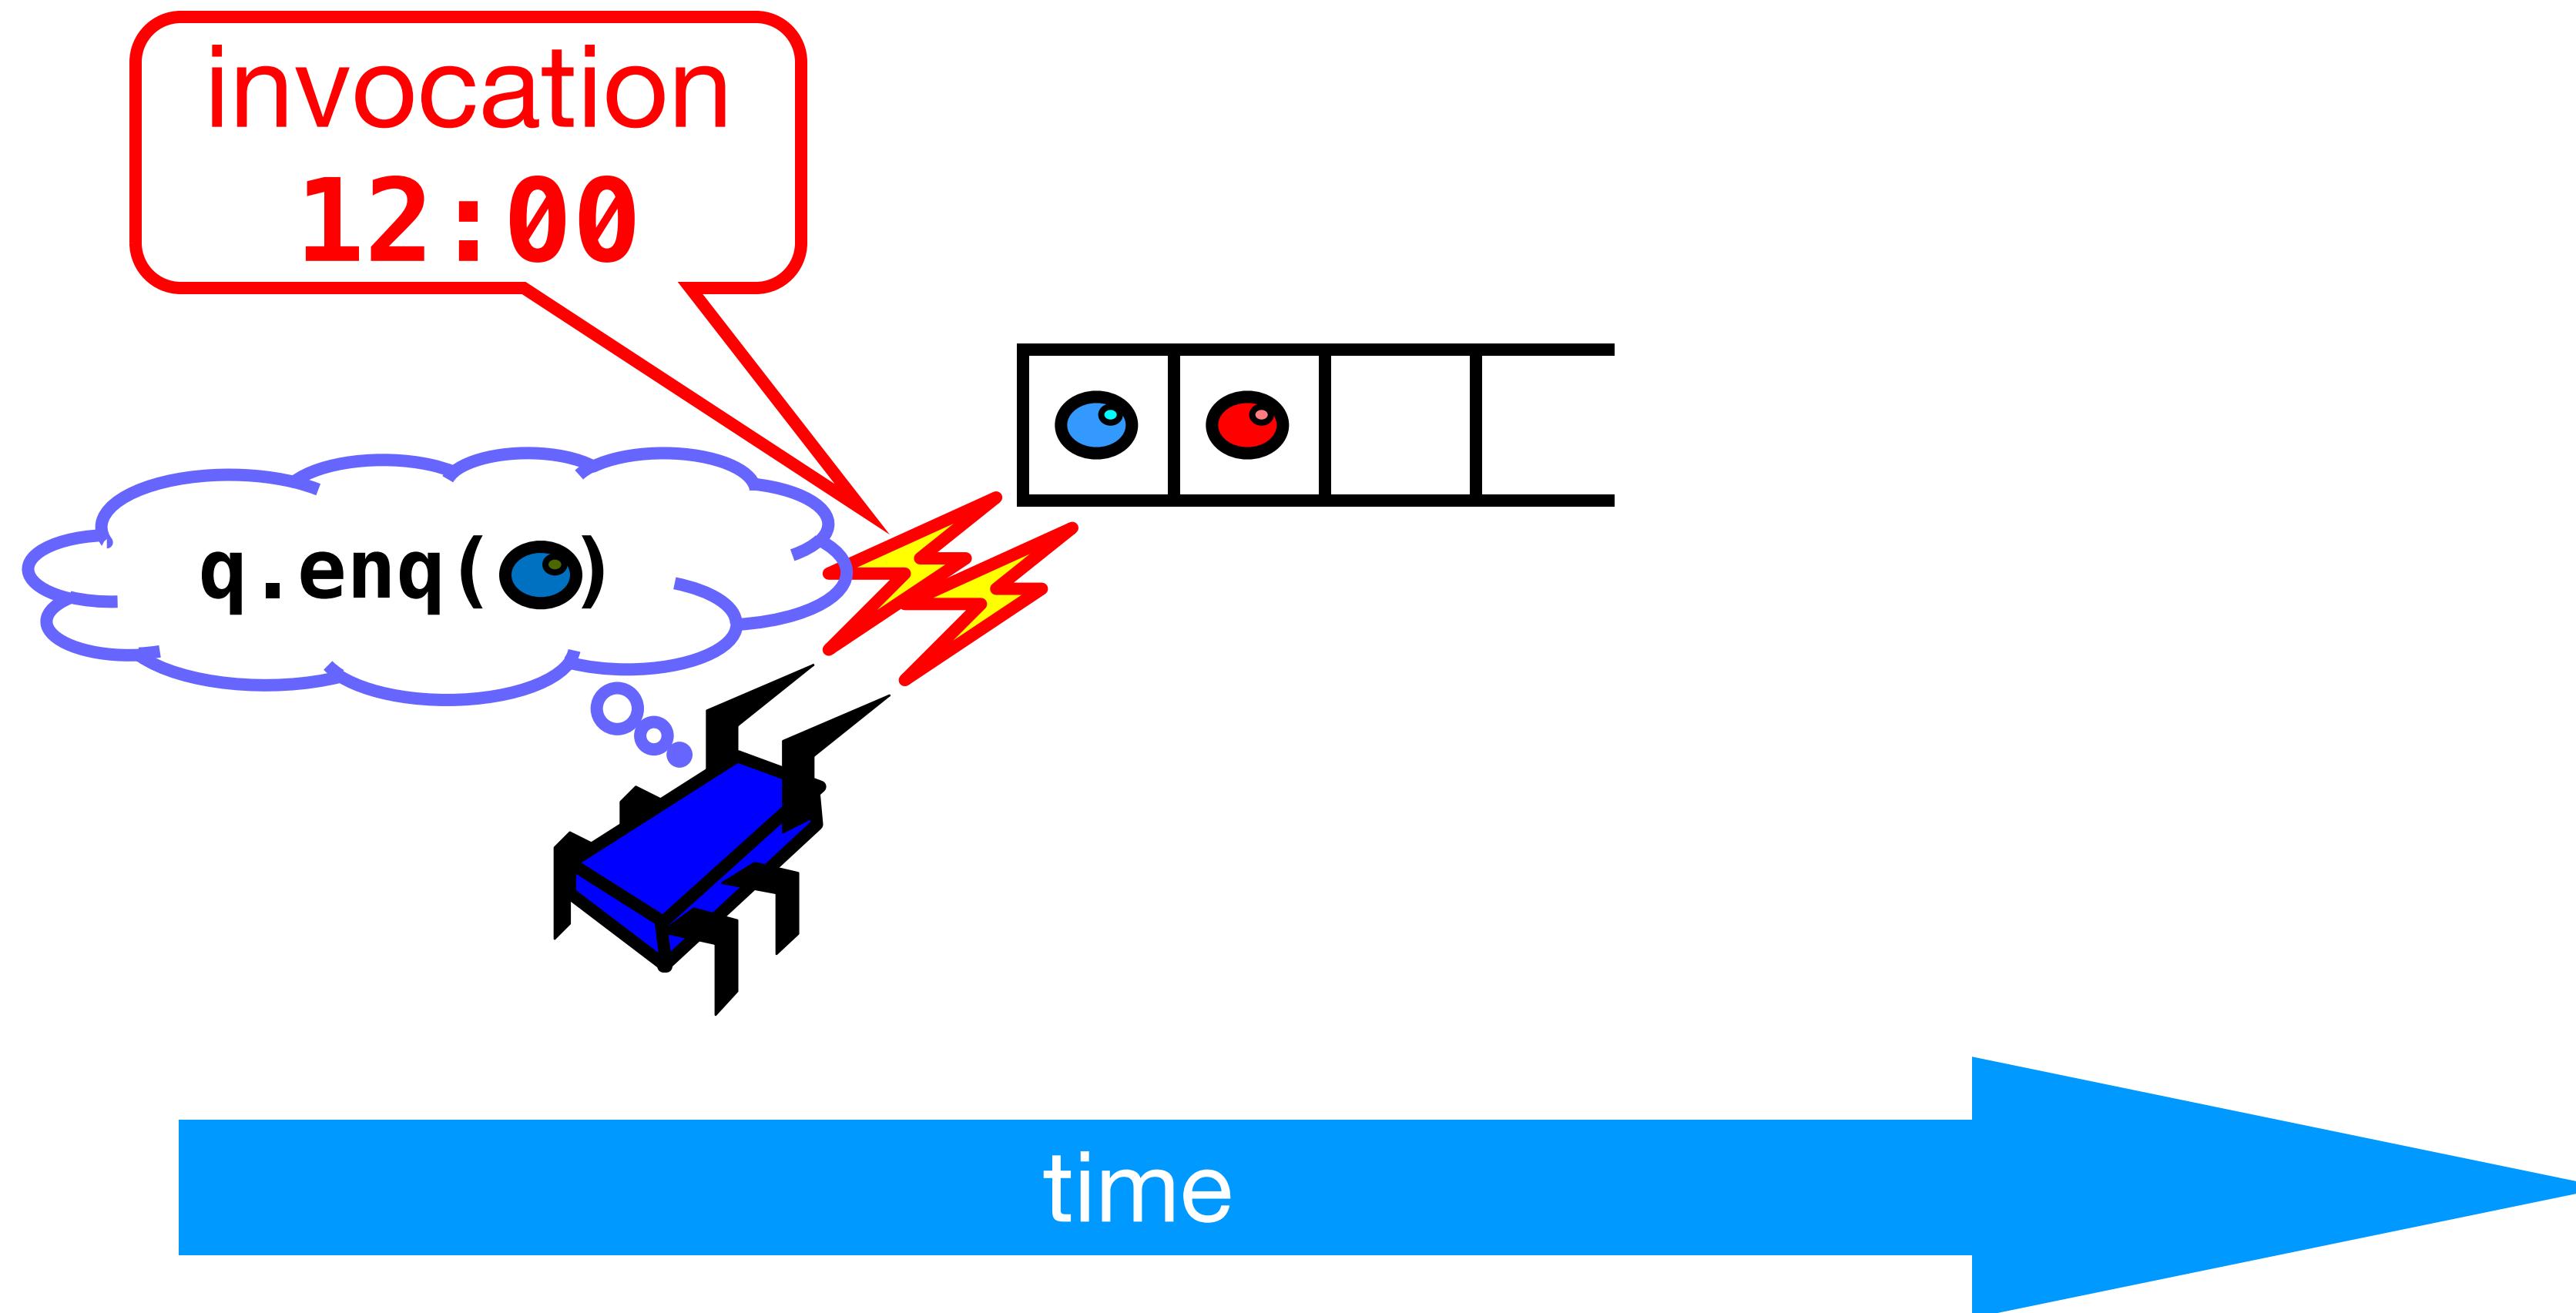
# What about concurrent Specifications?

- Methods?
- Documentation?
- Adding new methods?

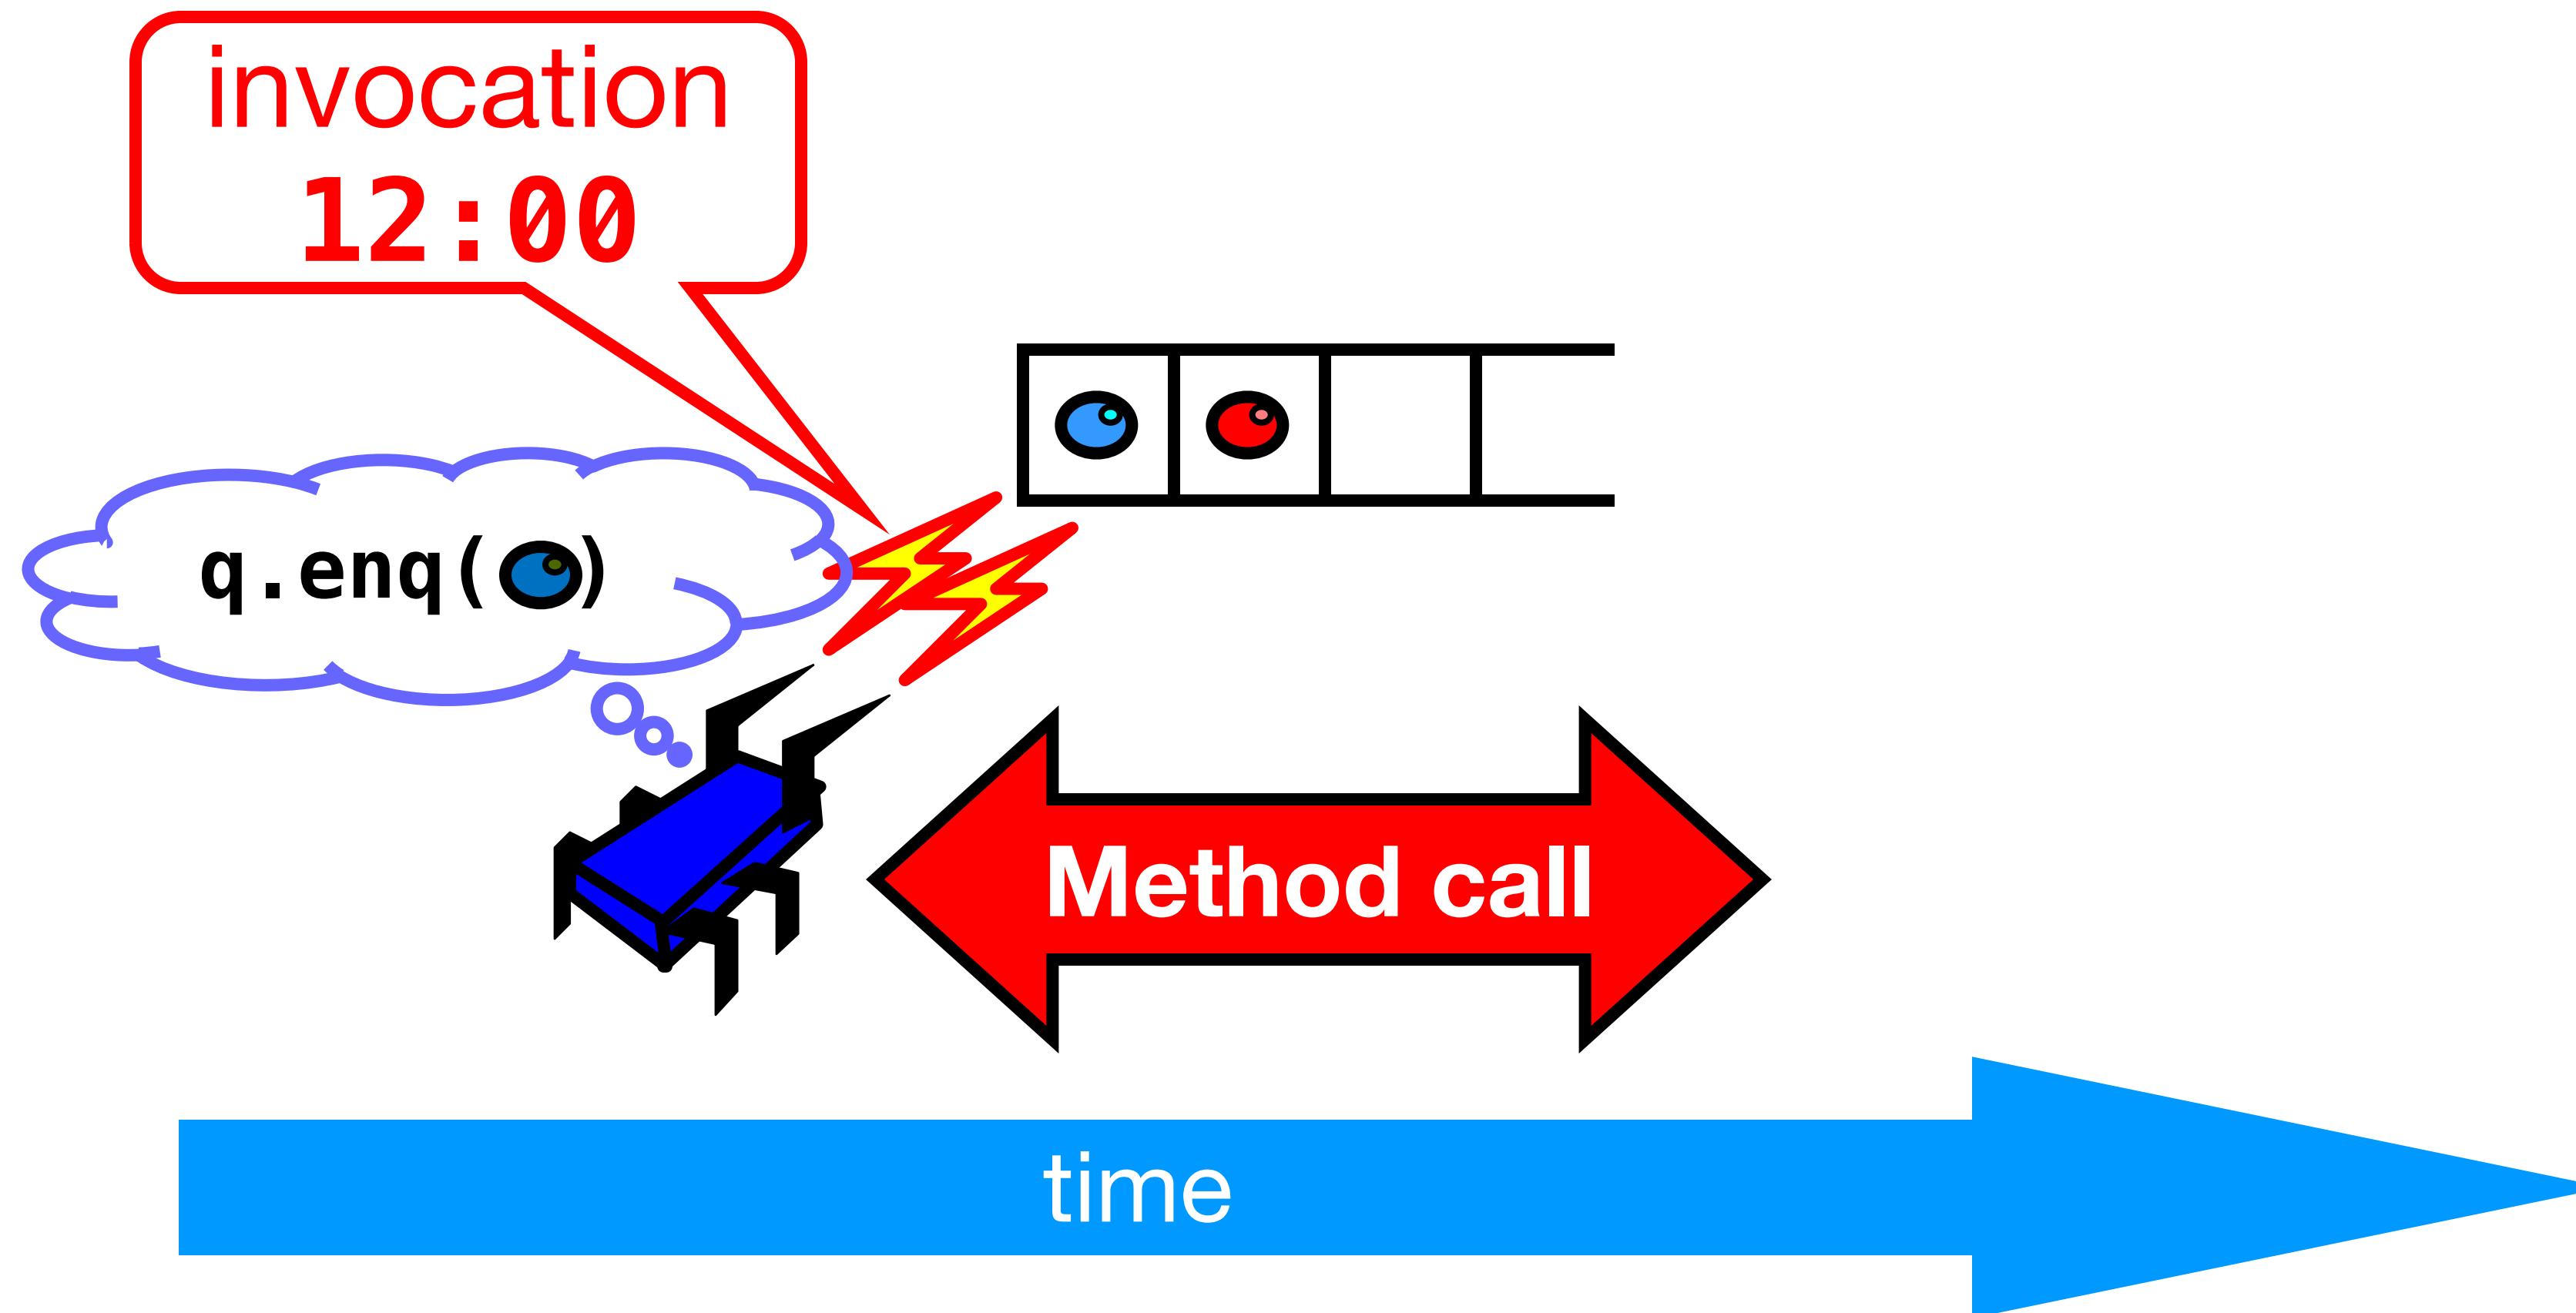
# Methods take time



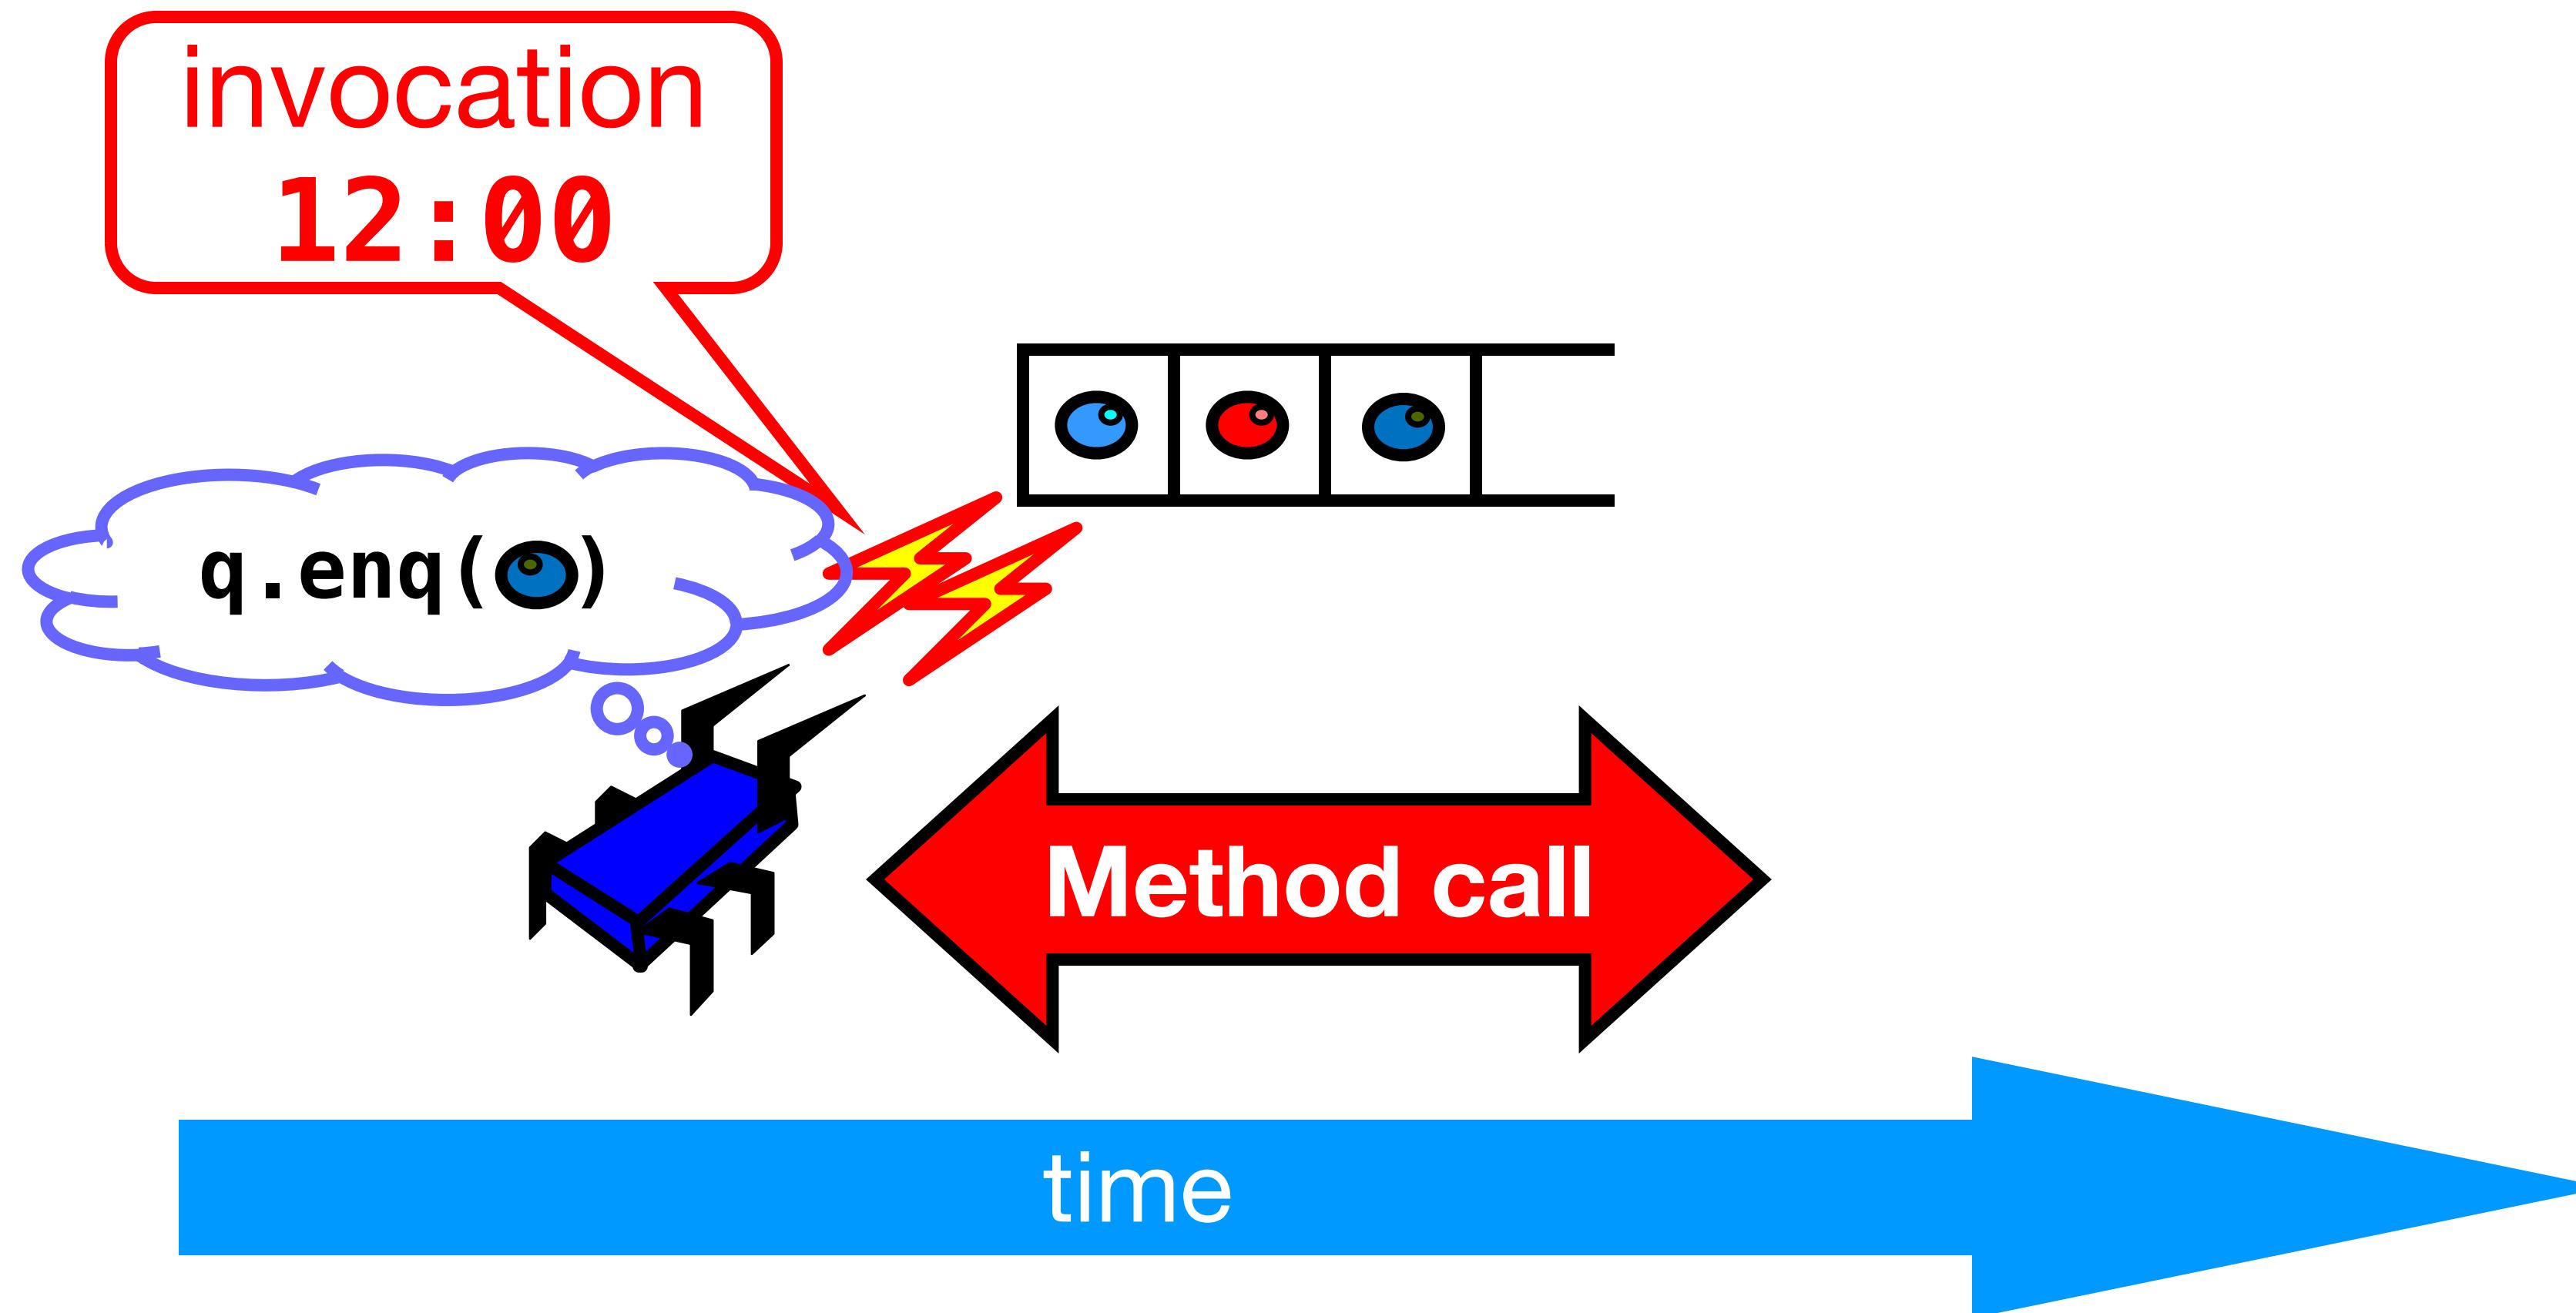
# Methods take time



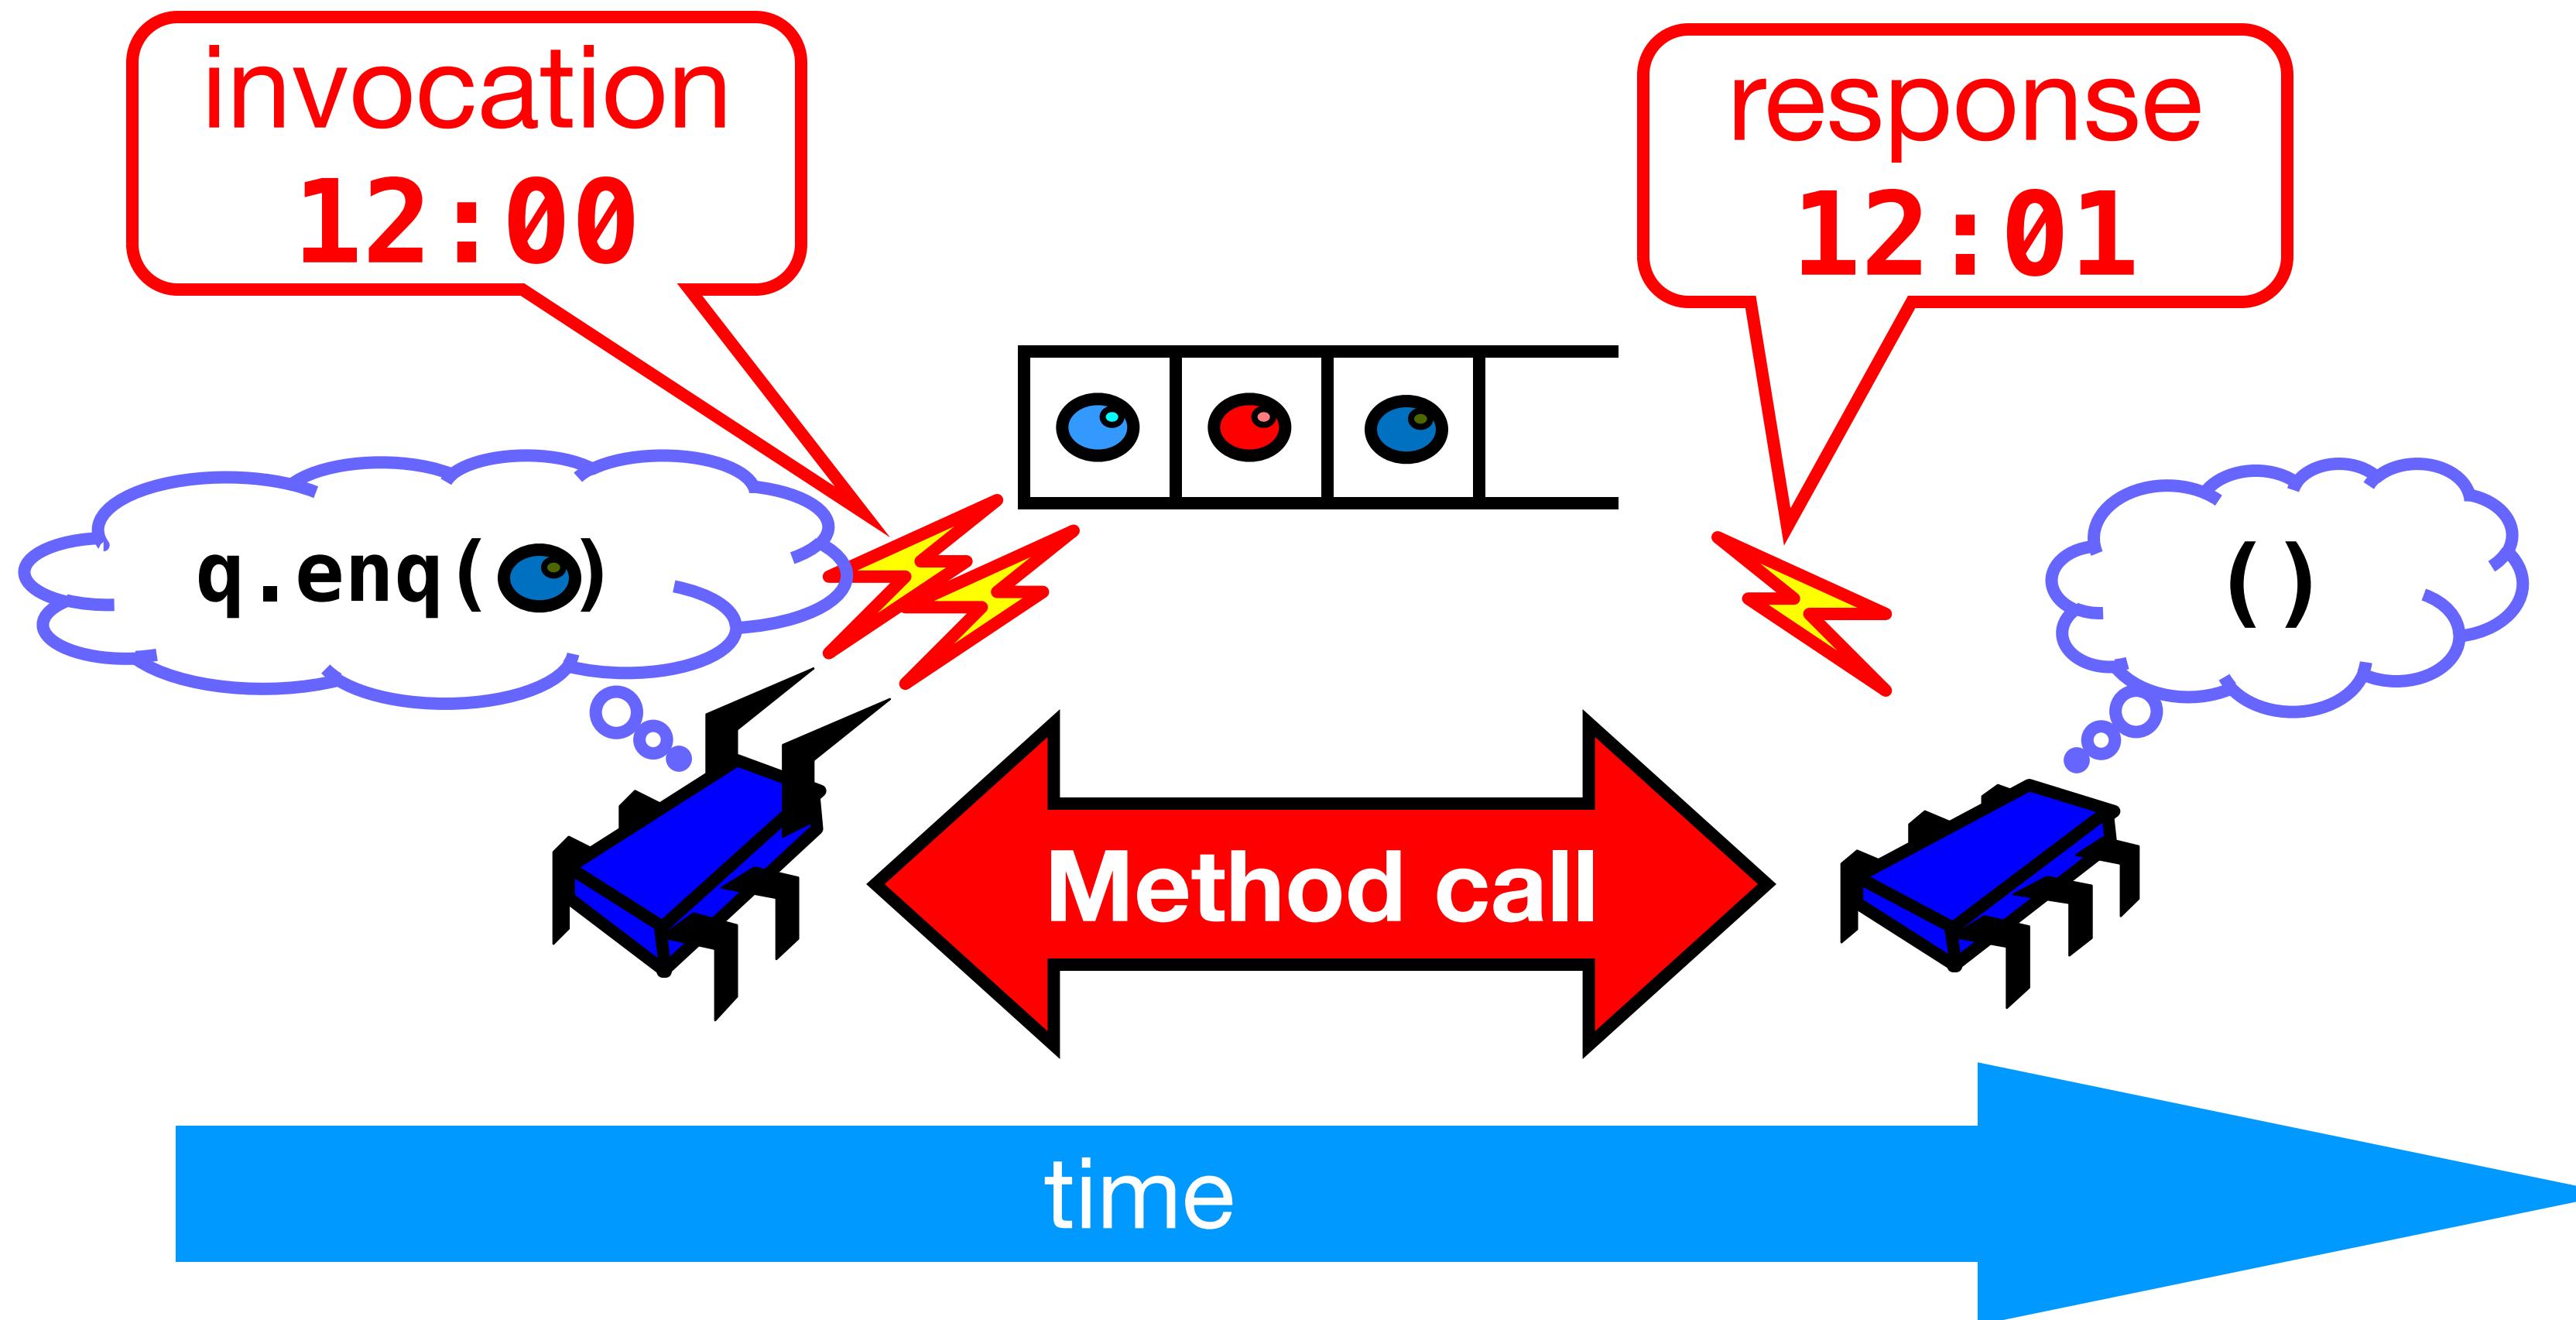
# Methods take time



# Methods take time



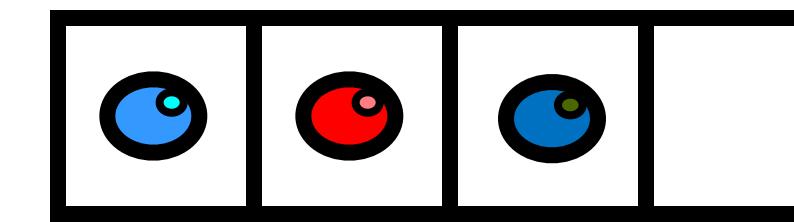
# Methods take time



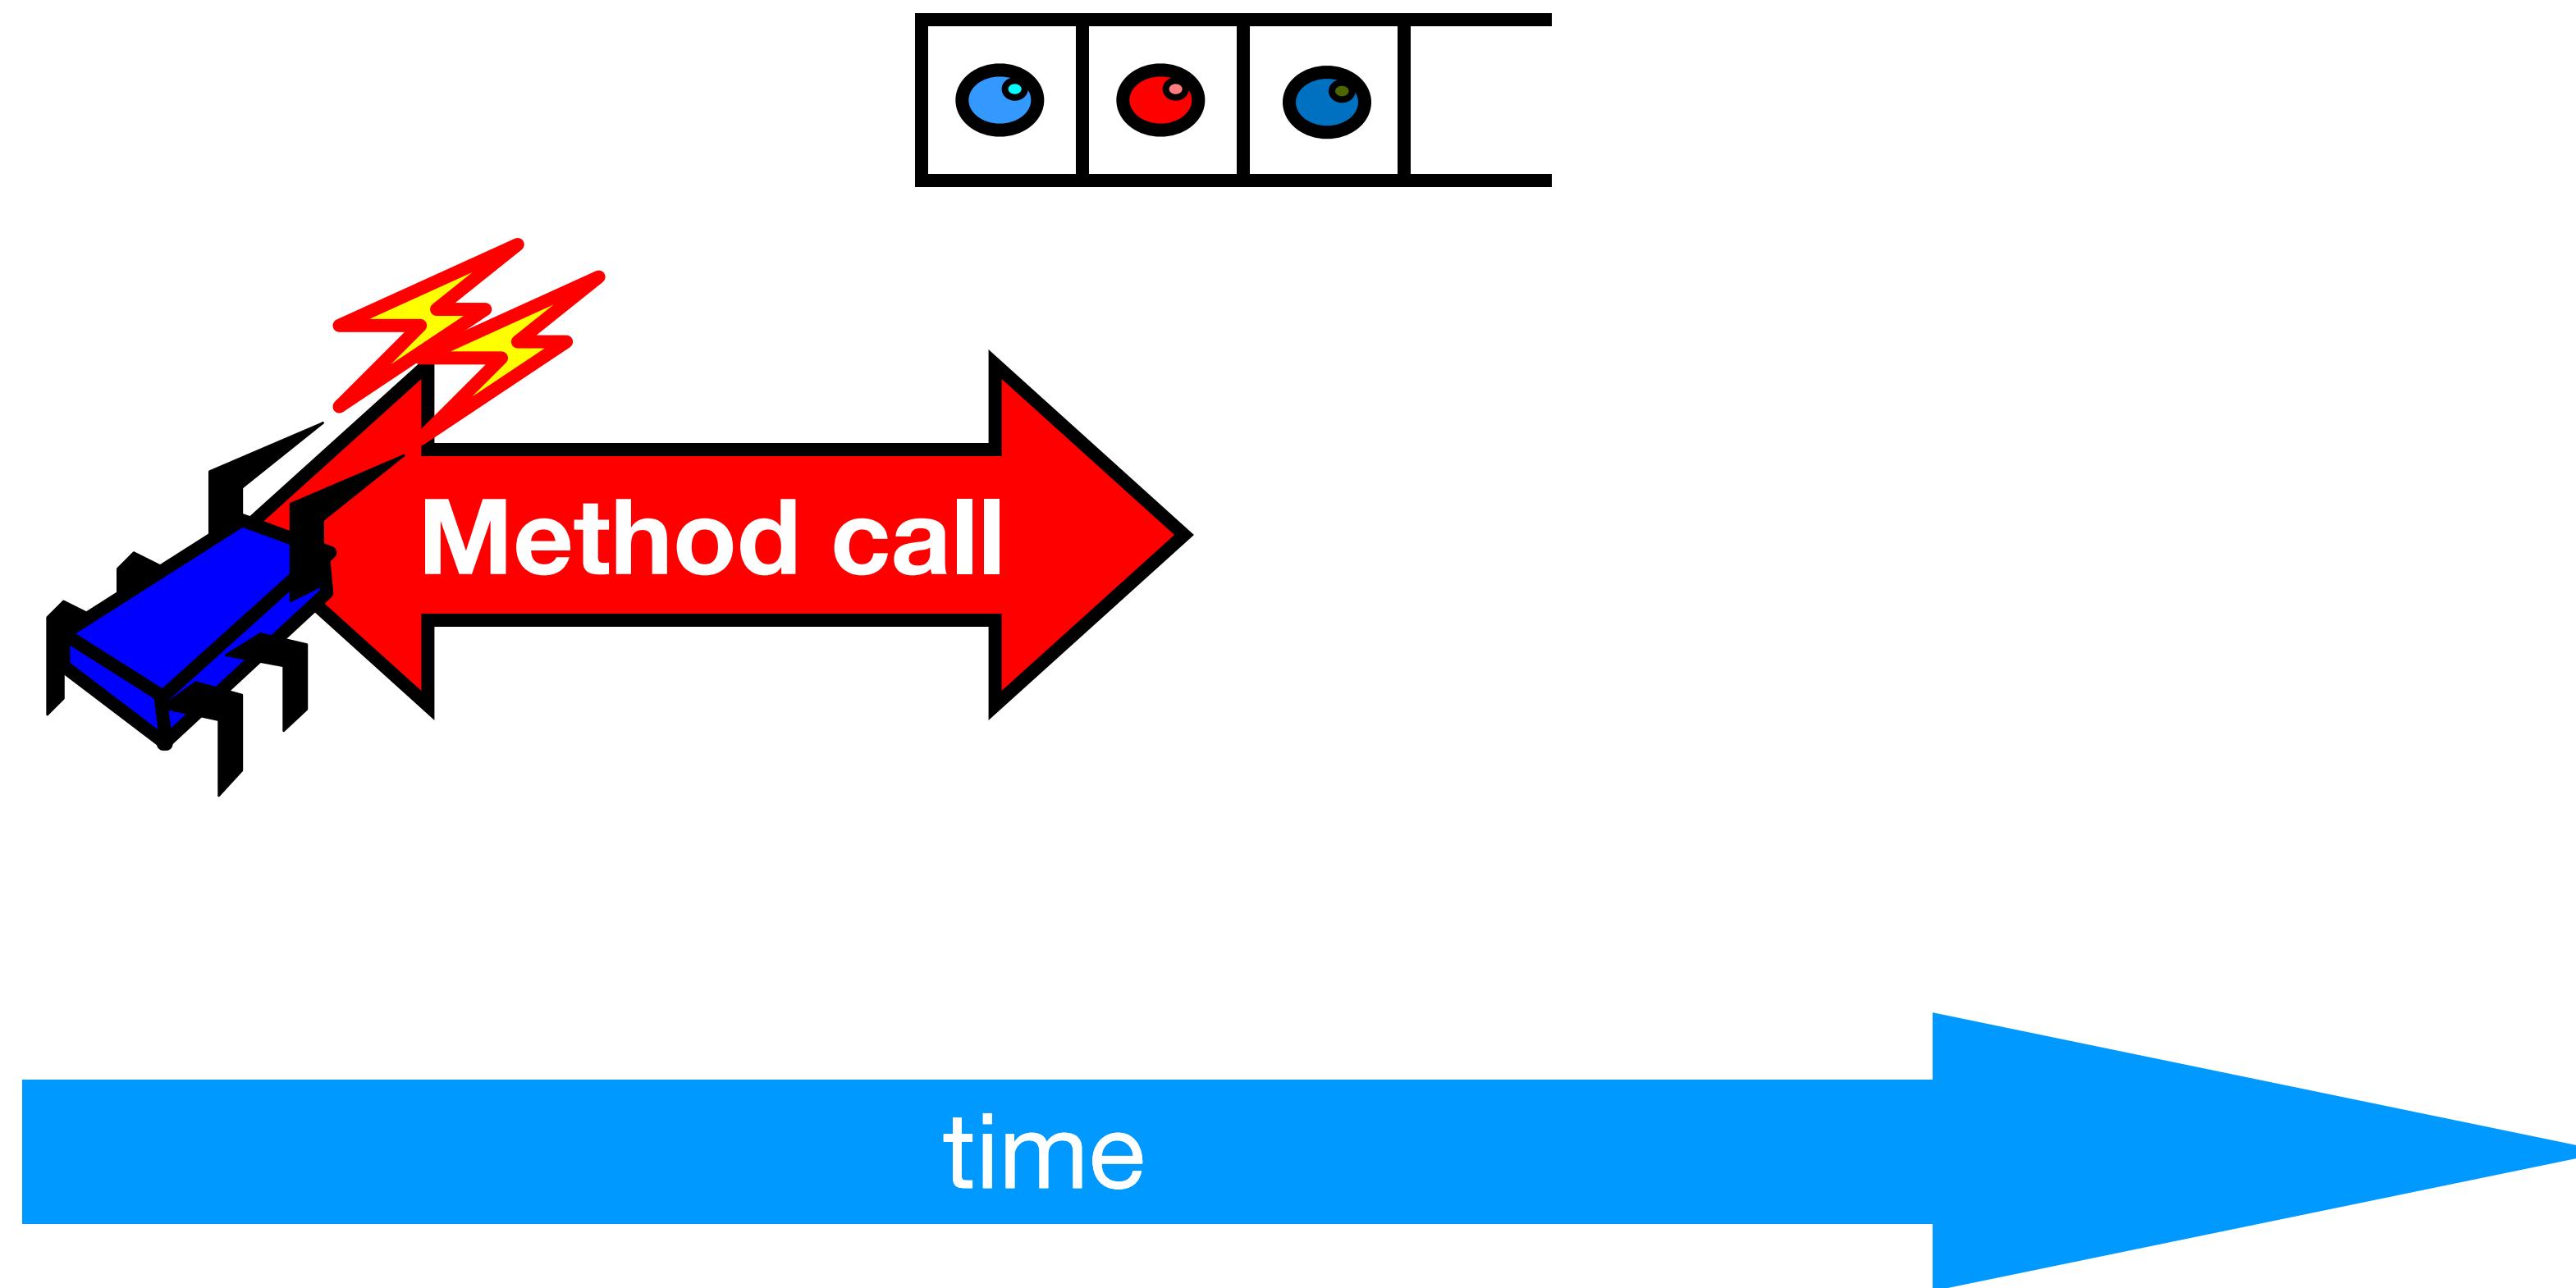
# Sequential vs Concurrent

- Sequential
  - Methods take time? Who knew?
- Concurrent
  - Method call is not an ***event***
  - Method call is an ***interval***

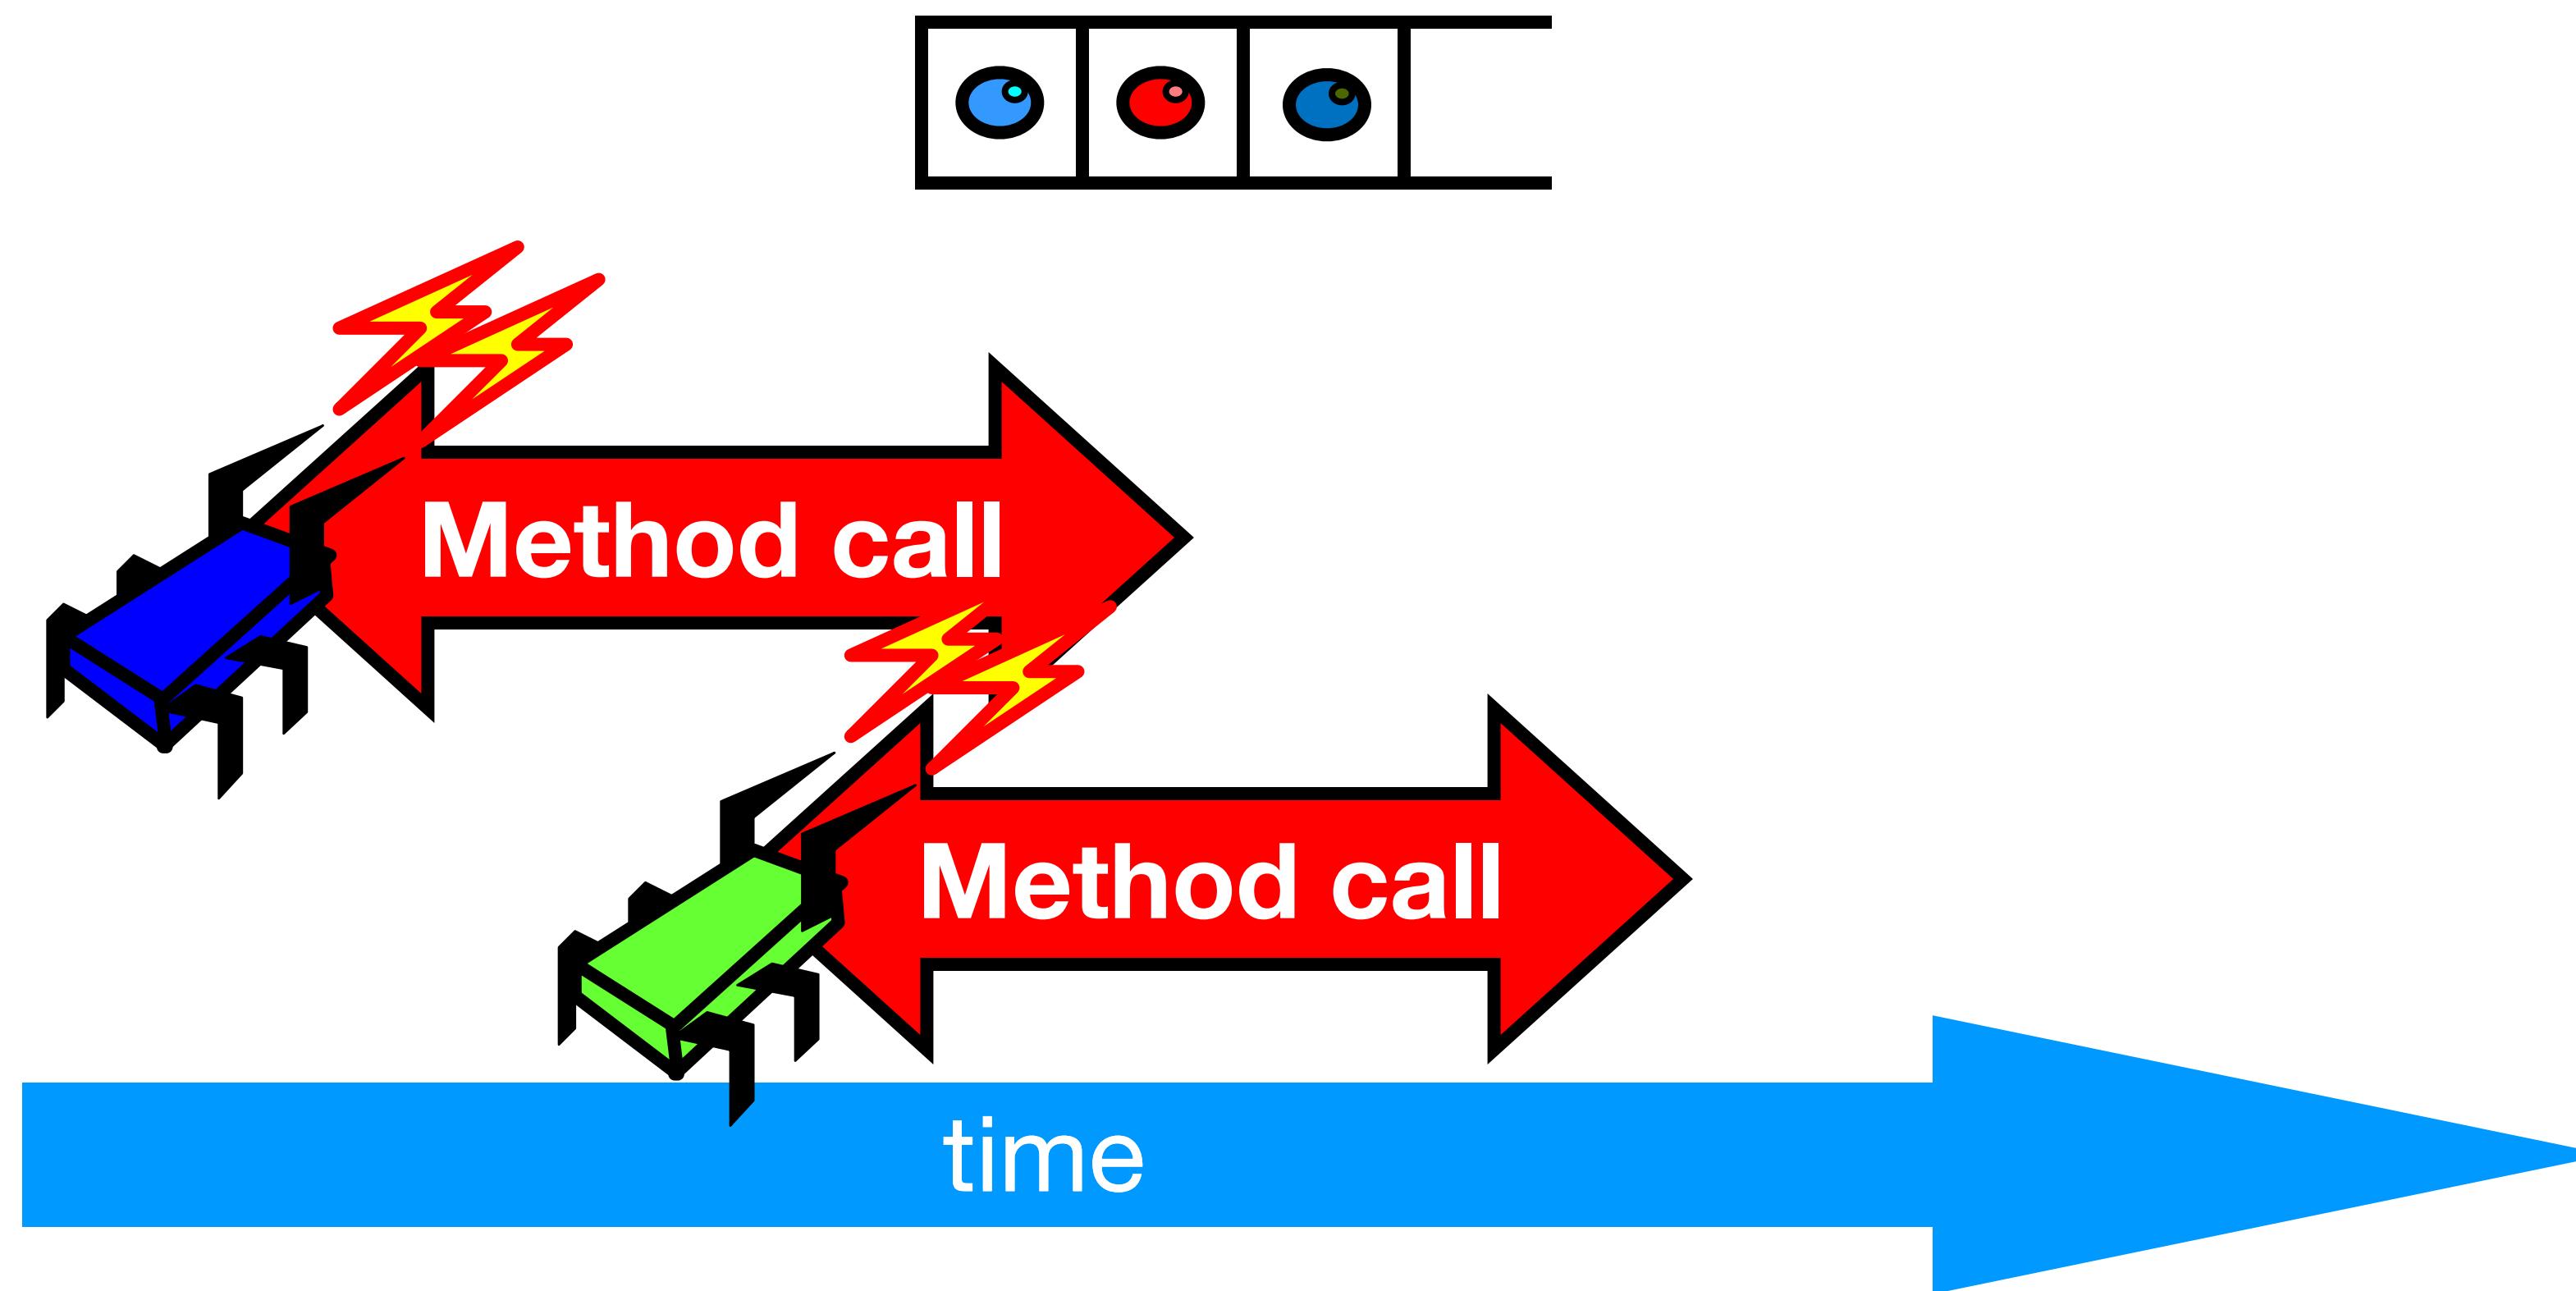
# Concurrent Methods Take Overlapping Time



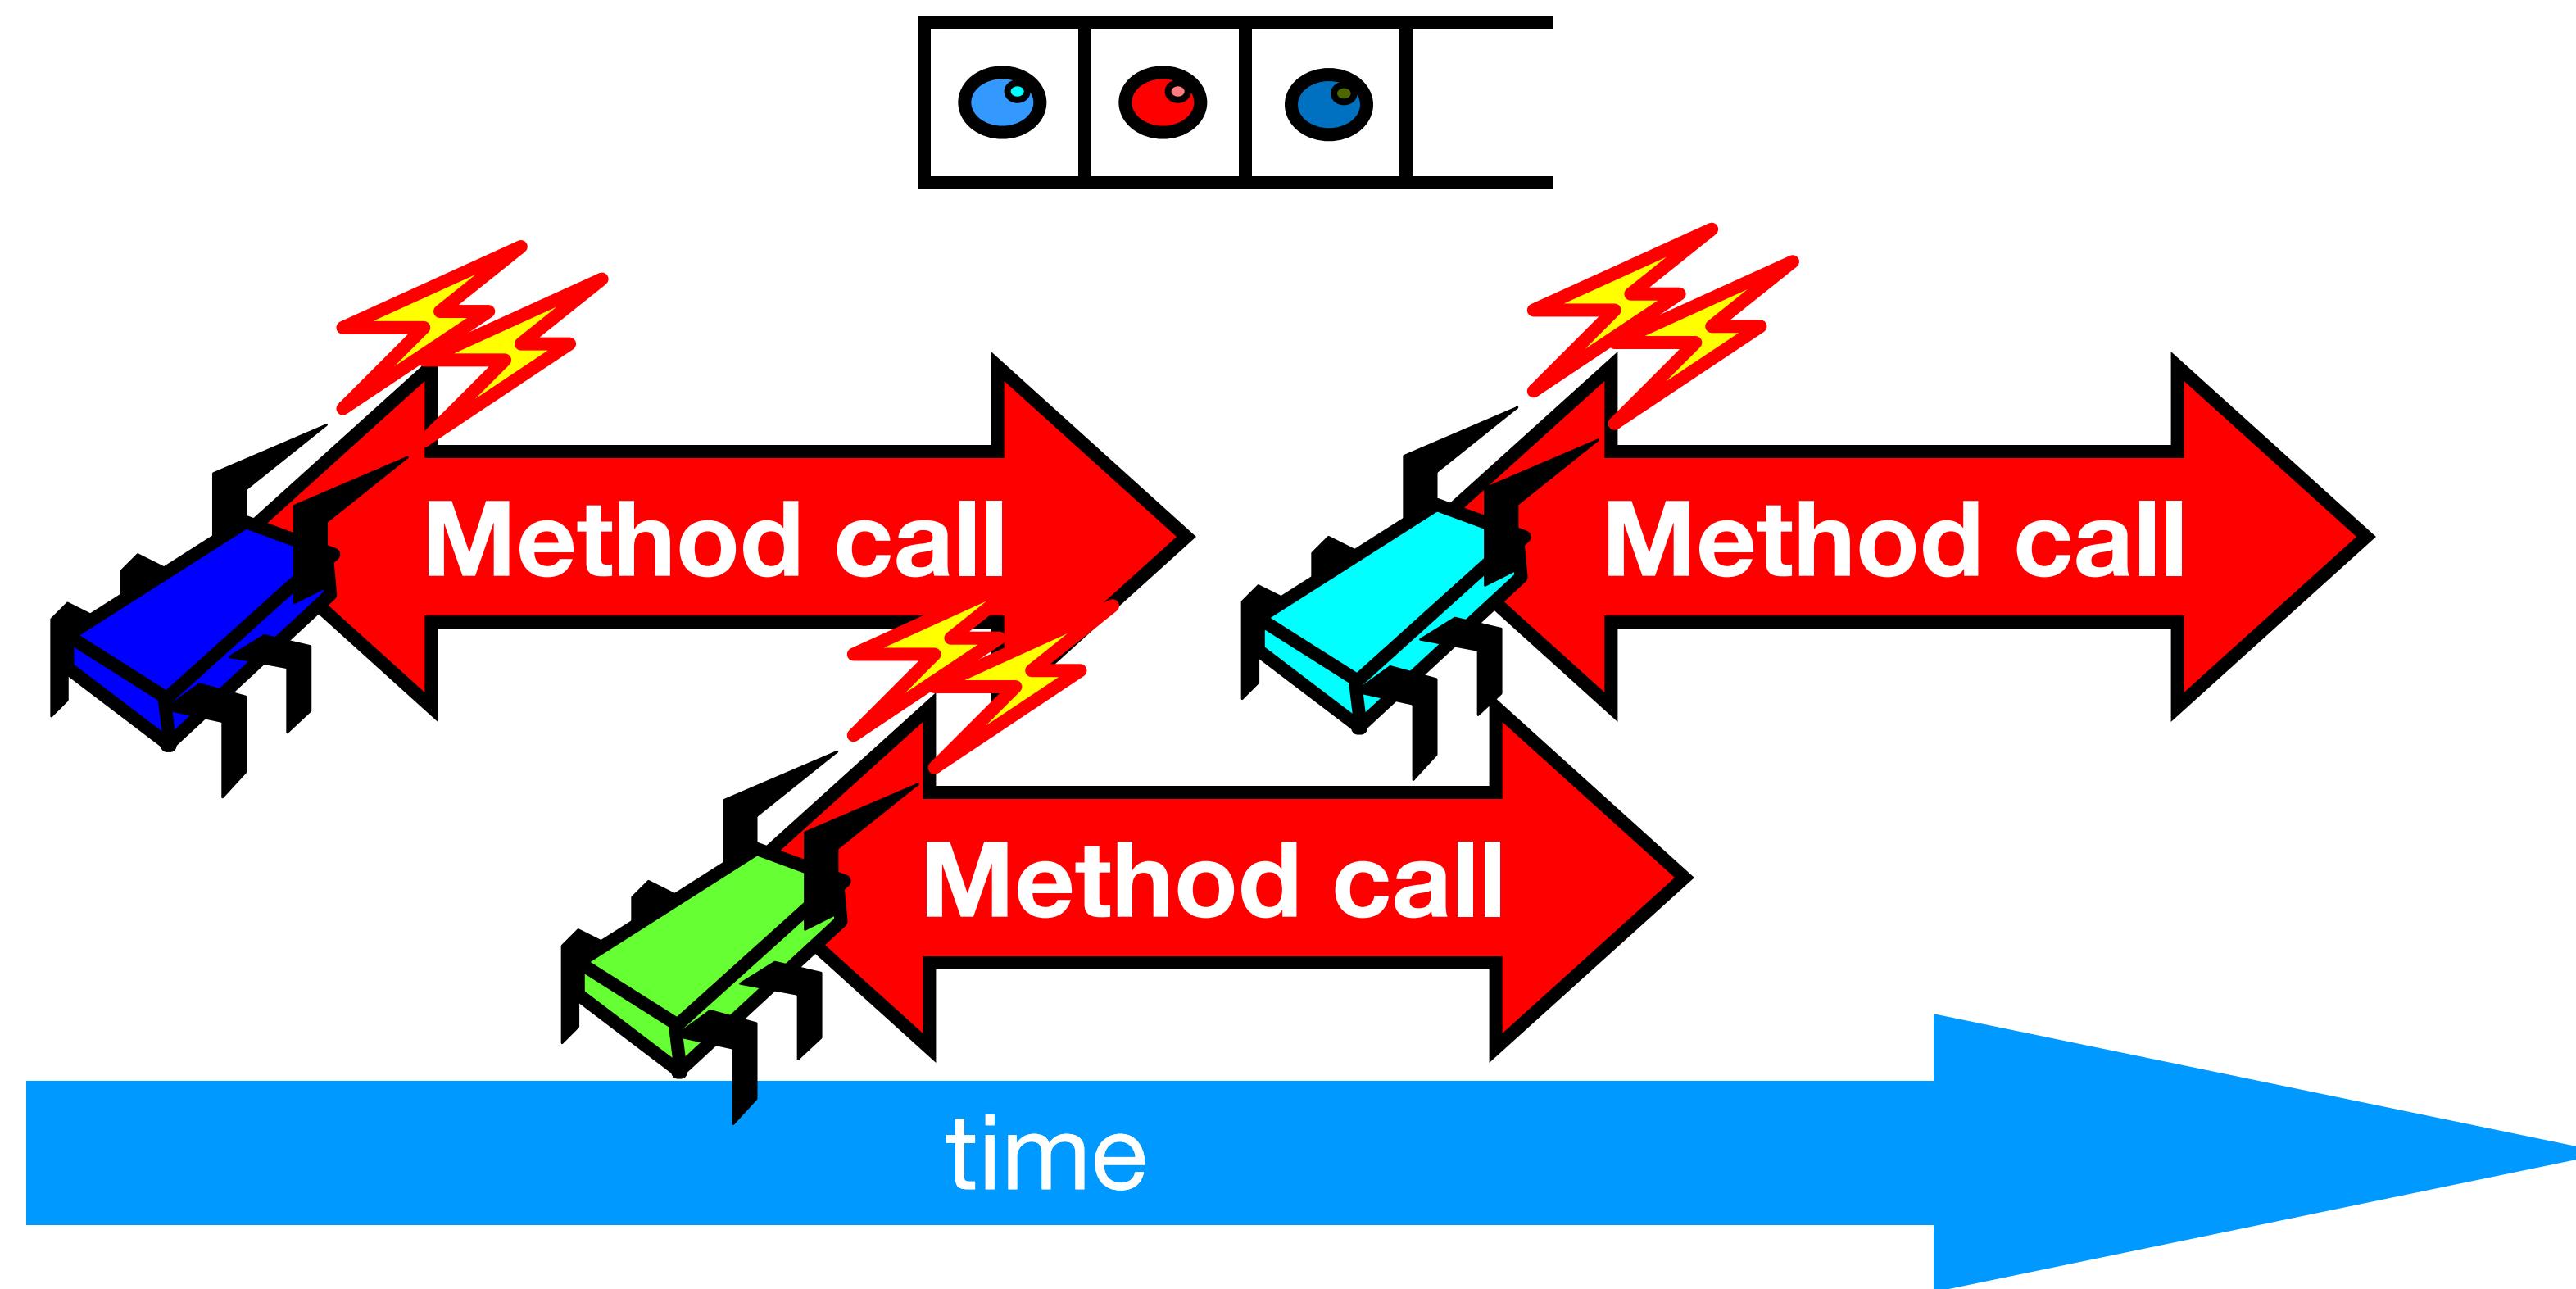
# Concurrent Methods Take Overlapping Time



# Concurrent Methods Take Overlapping Time



# Concurrent Methods Take Overlapping Time



# Sequential vs Concurrent

- Sequential
  - Object needs a meaningful state only ***between*** method calls
- Concurrent
  - Because method calls overlap, the object might ***never*** be between method calls

# Sequential vs Concurrent

- Sequential:
  - Each method described in isolation
- Concurrent
  - Must characterize ***all*** possible interactions with concurrent calls
    - What if two **enq()** calls overlap?
    - Two **deq()** calls? **enq()** and **deq()**? ...

# Sequential vs Concurrent

- Sequential:
  - Can add new methods without affecting older methods
- Concurrent:
  - Everything can potentially interact with everything else

# Sequential vs Concurrent

- Sequential:
  - Can add new methods without affecting older methods
- Concurrent:
  - Everything can potentially interact with everything else



# The Big Question

- What does it *mean* for a ***concurrent*** object to be correct?
  - What *is* a concurrent FIFO queue?
  - FIFO means ***strict temporal order***
  - Concurrent means ***ambiguous temporal order***

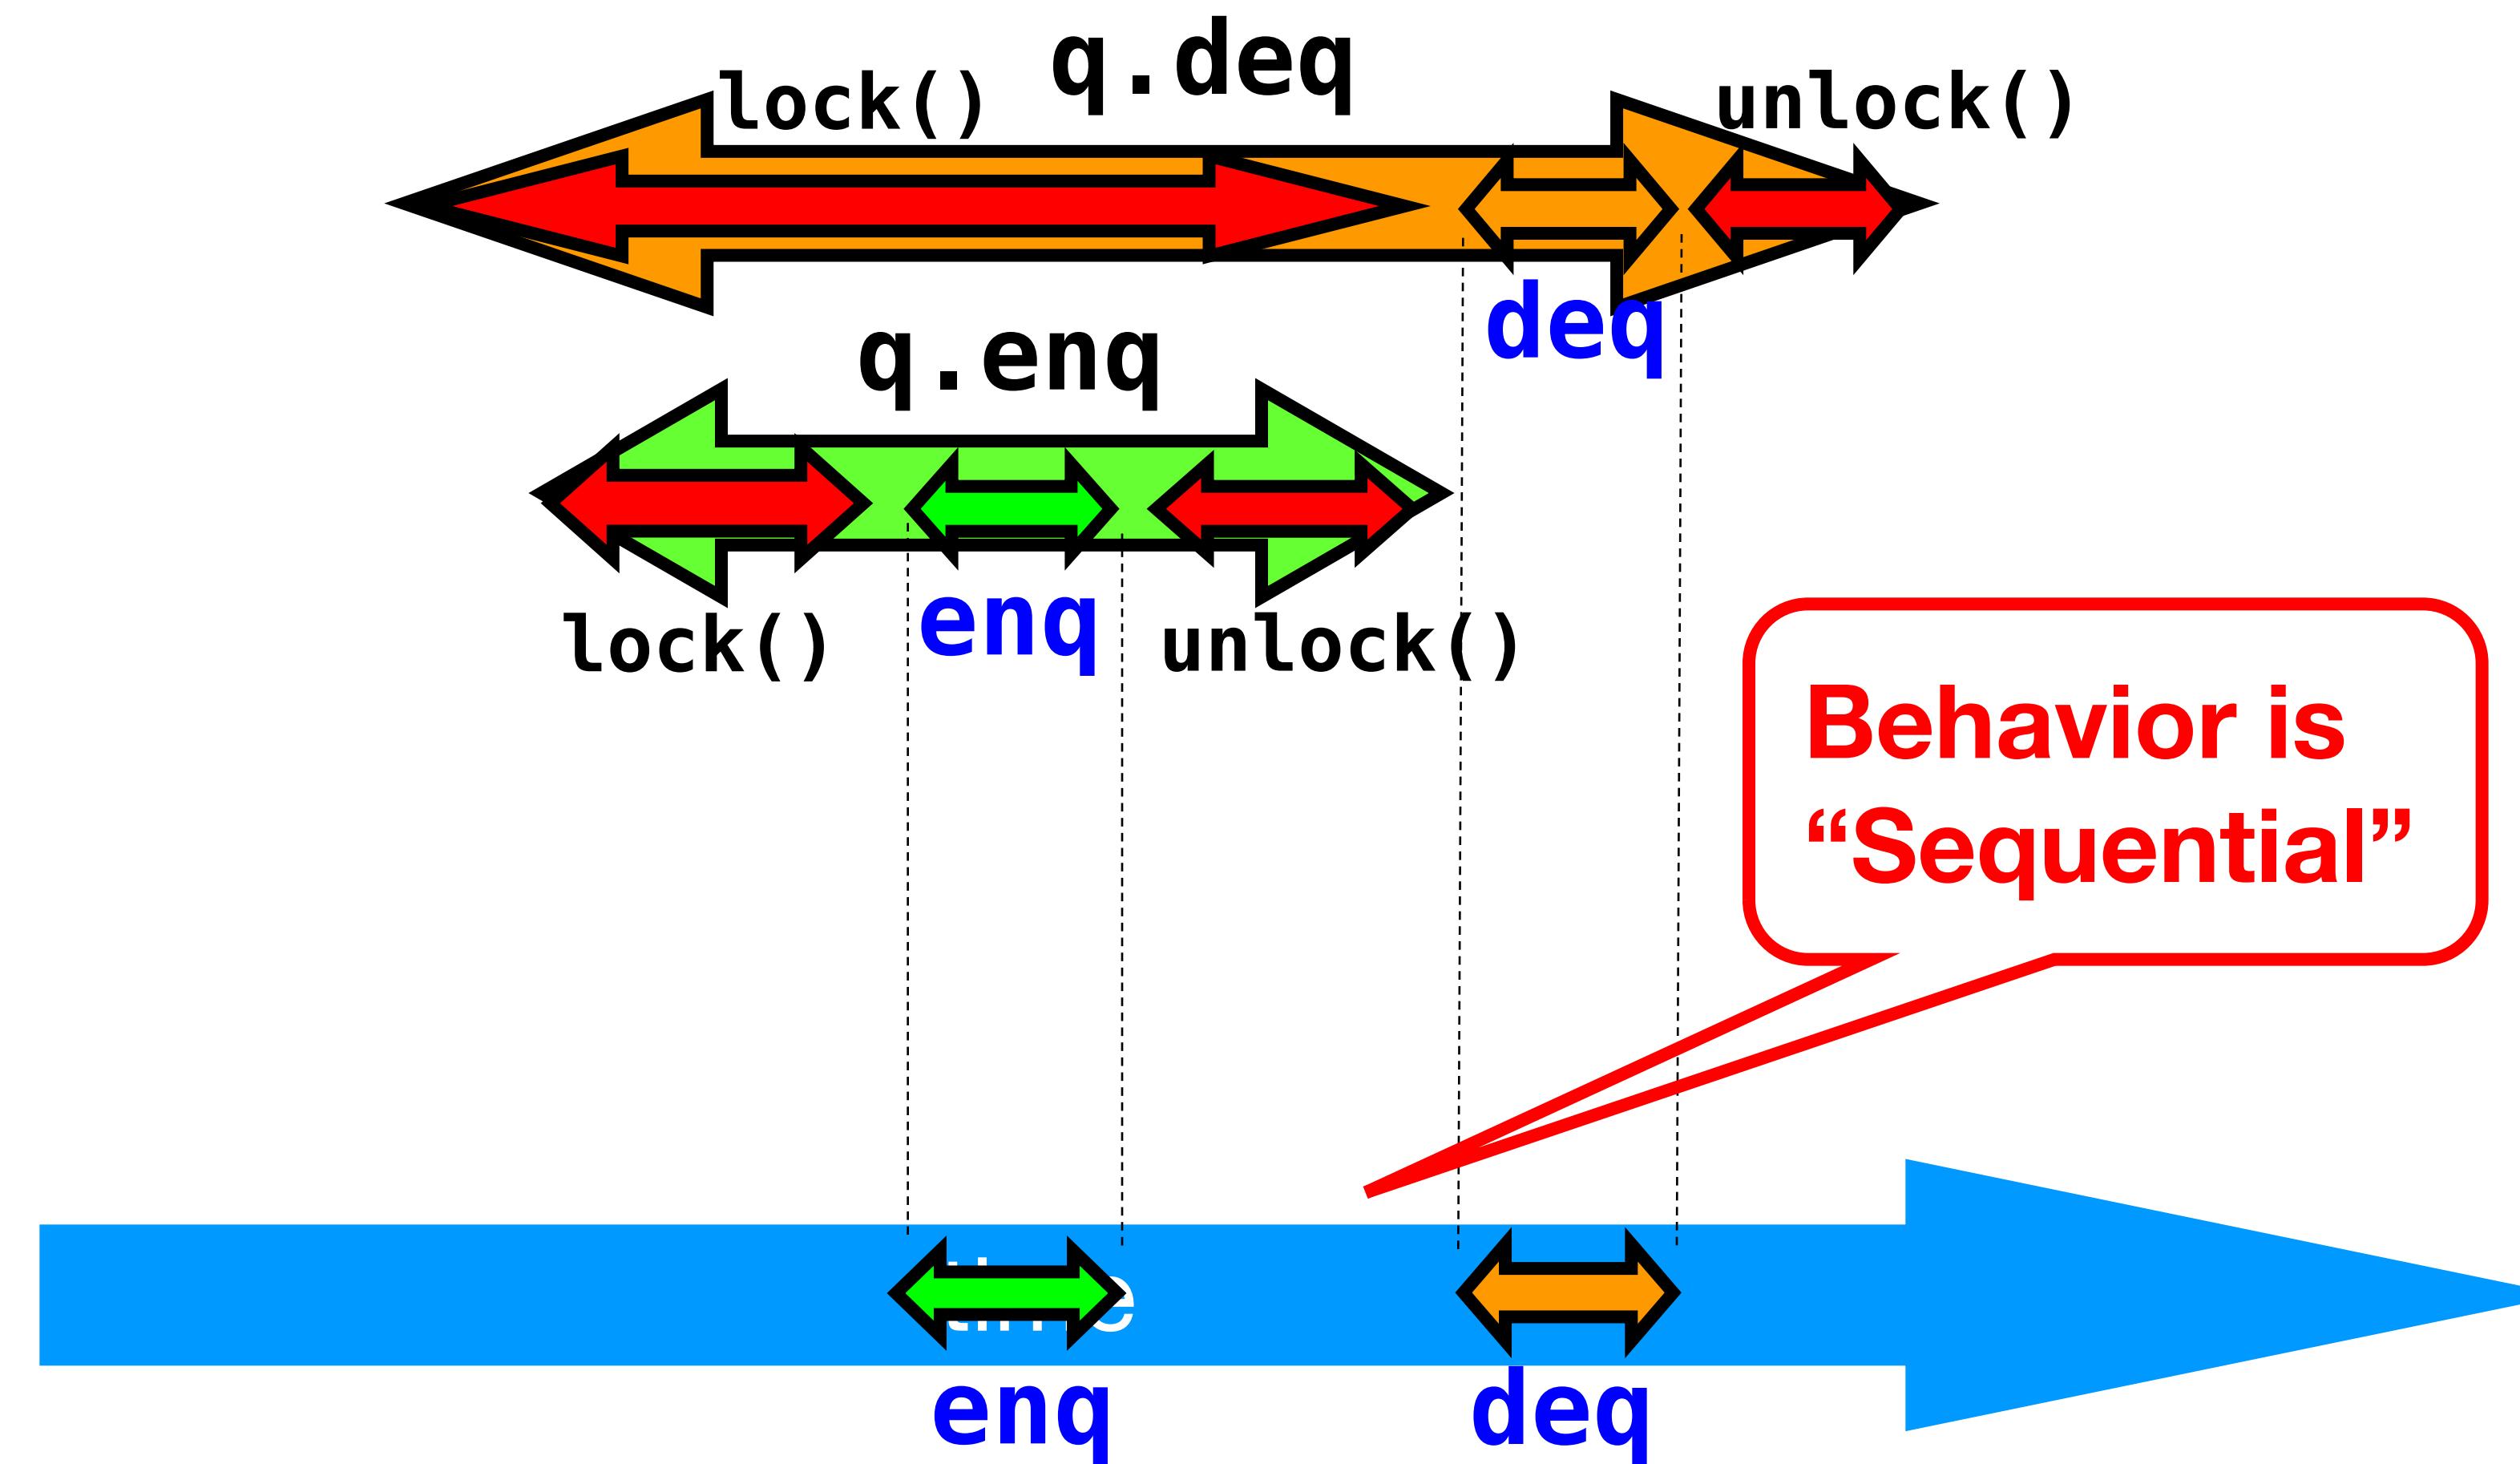
# Intuitively

```
let deq q =
  Mutex.lock q.lock;
  try
    if q.tail = q.head then
      raise Empty;
    match q.items.(q.head mod q.capacity) with
    | None -> assert false
    | Some x ->
        q.head <- q.head + 1;
        Mutex.unlock q.lock;
        x
  with e ->
    Mutex.unlock q.lock;
    raise e
```

*All queue modifications are mutually exclusive*

# Intuitively

Lets capture the idea of describing  
the concurrent via the sequential



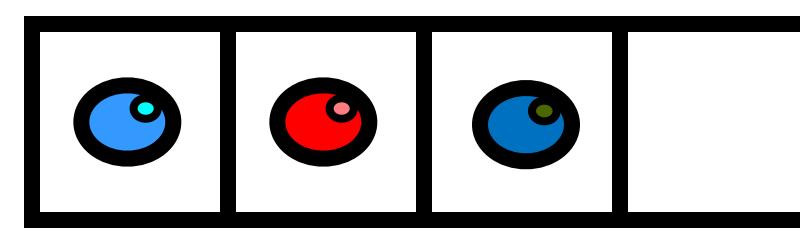
# Linearizability

- Each method should
  - “take effect”
  - Instantaneously
  - Between ***invocation*** and ***response*** events
- Object is correct if this “sequentialised” behaviour is correct
- Any such concurrent object is
  - **Linearizable™**

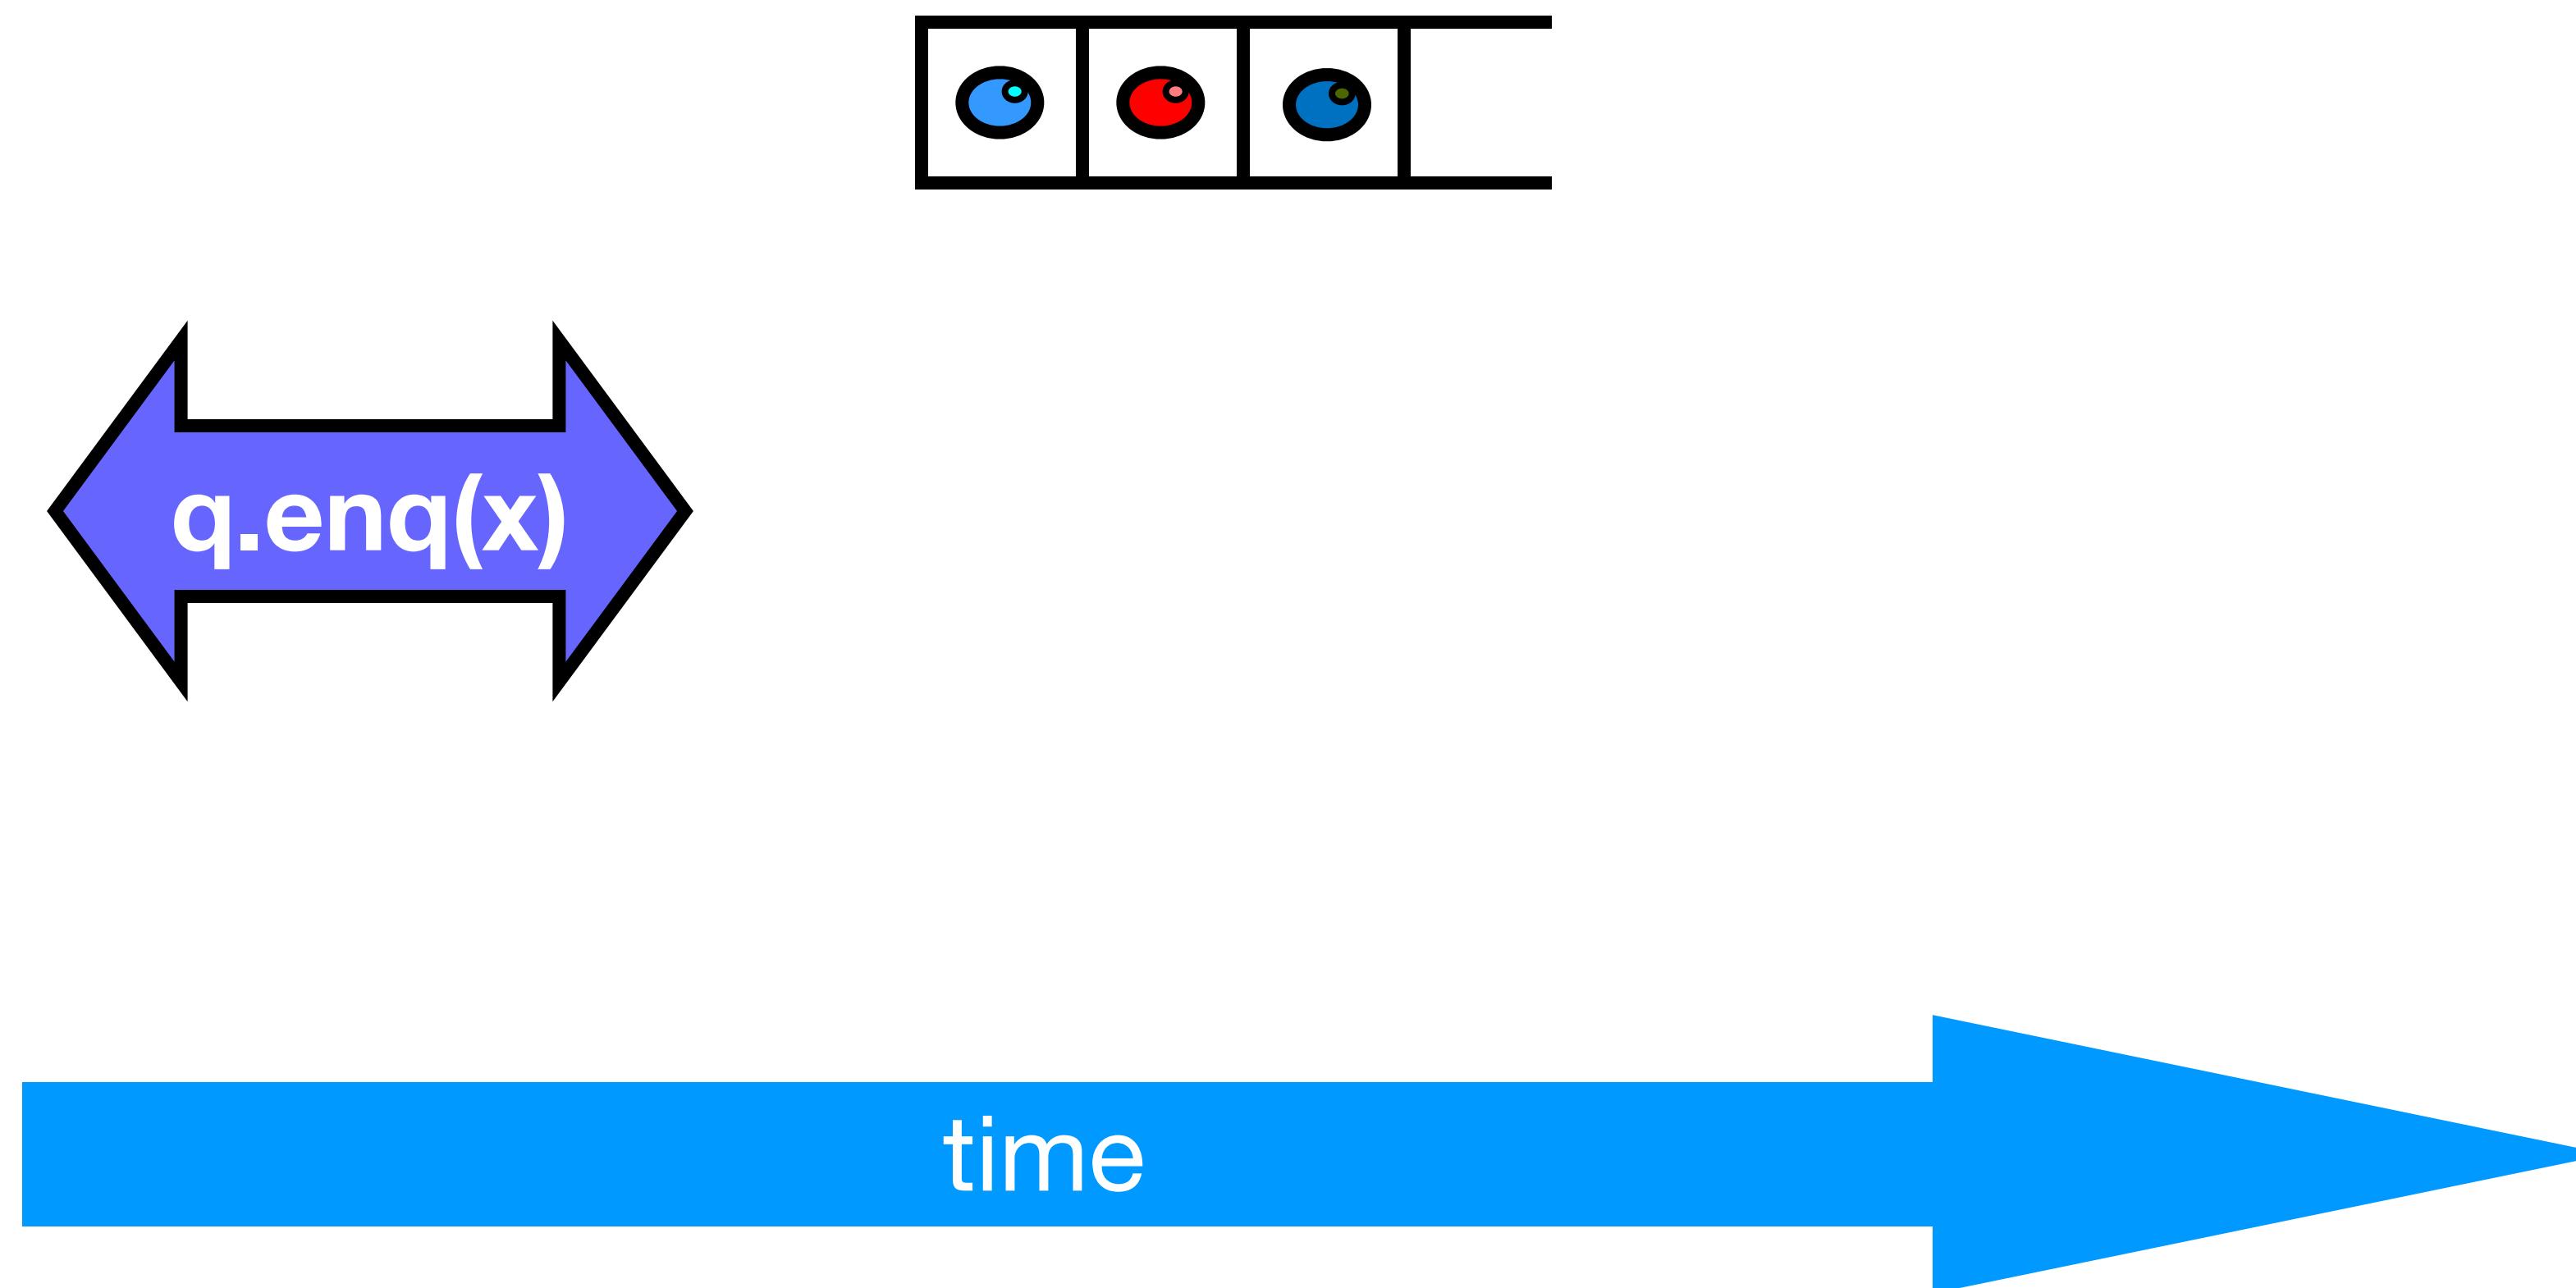
# Is it really about the object?

- Each method should
  - “take effect”
  - Instantaneously
  - Between *invocation* and *response* events
- Sounds like a property of *an execution...*
- A linearizable *object*
  - One of whose all possible executions are linearizable

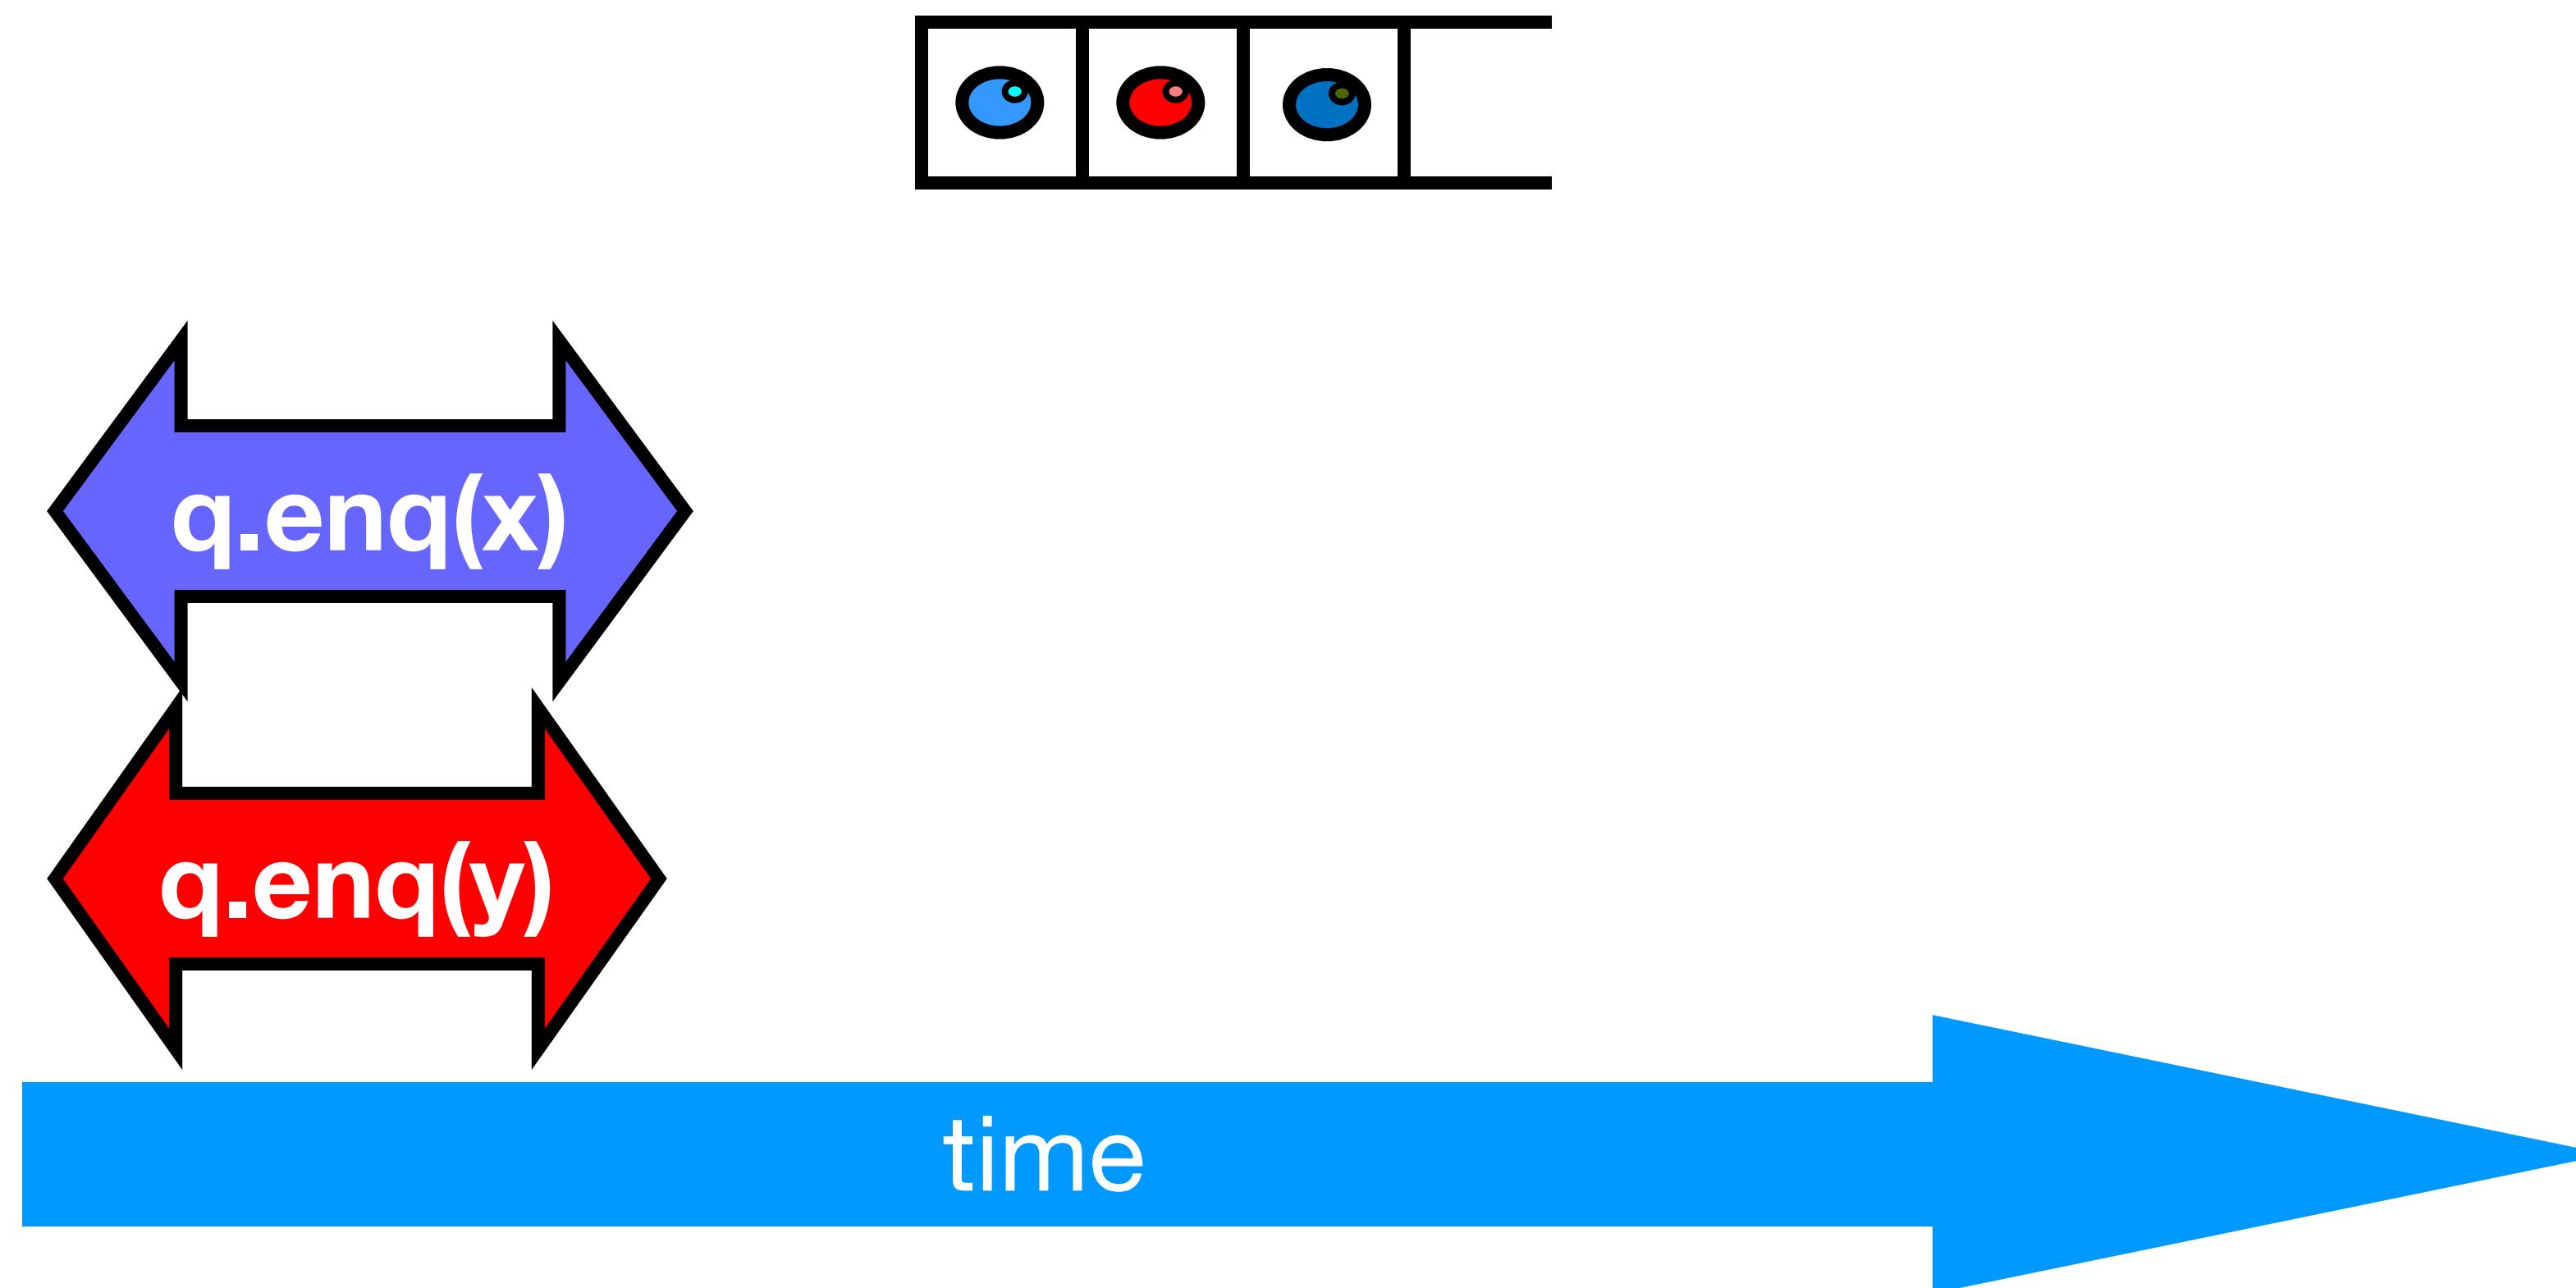
# Example



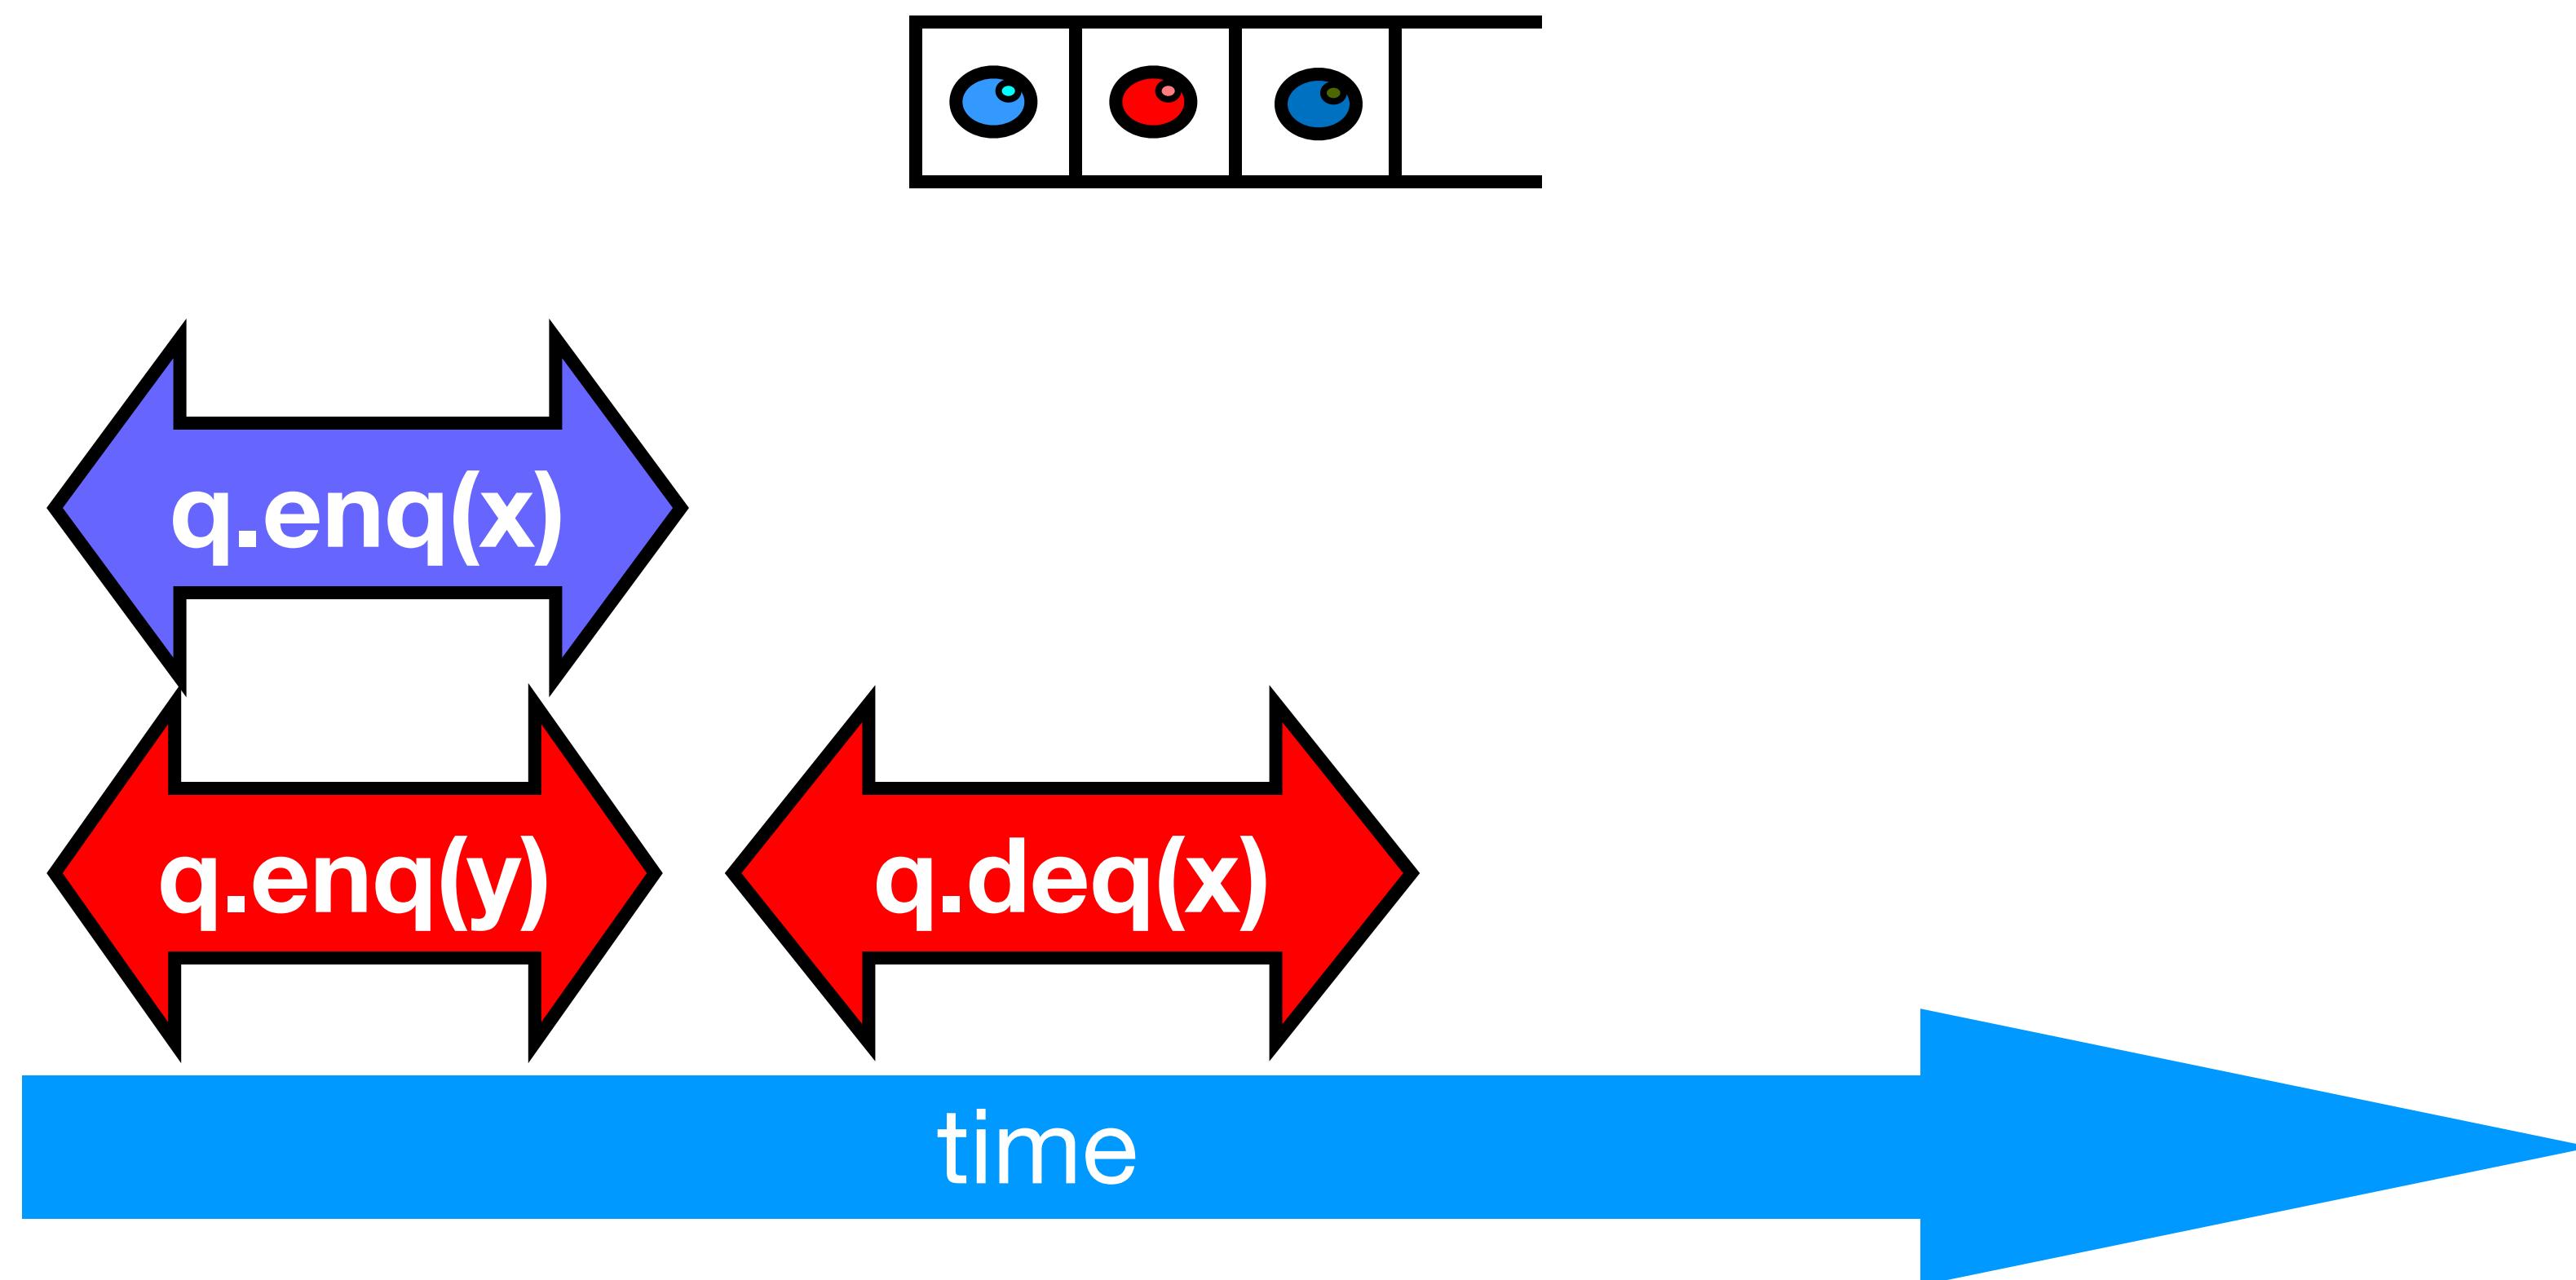
# Example



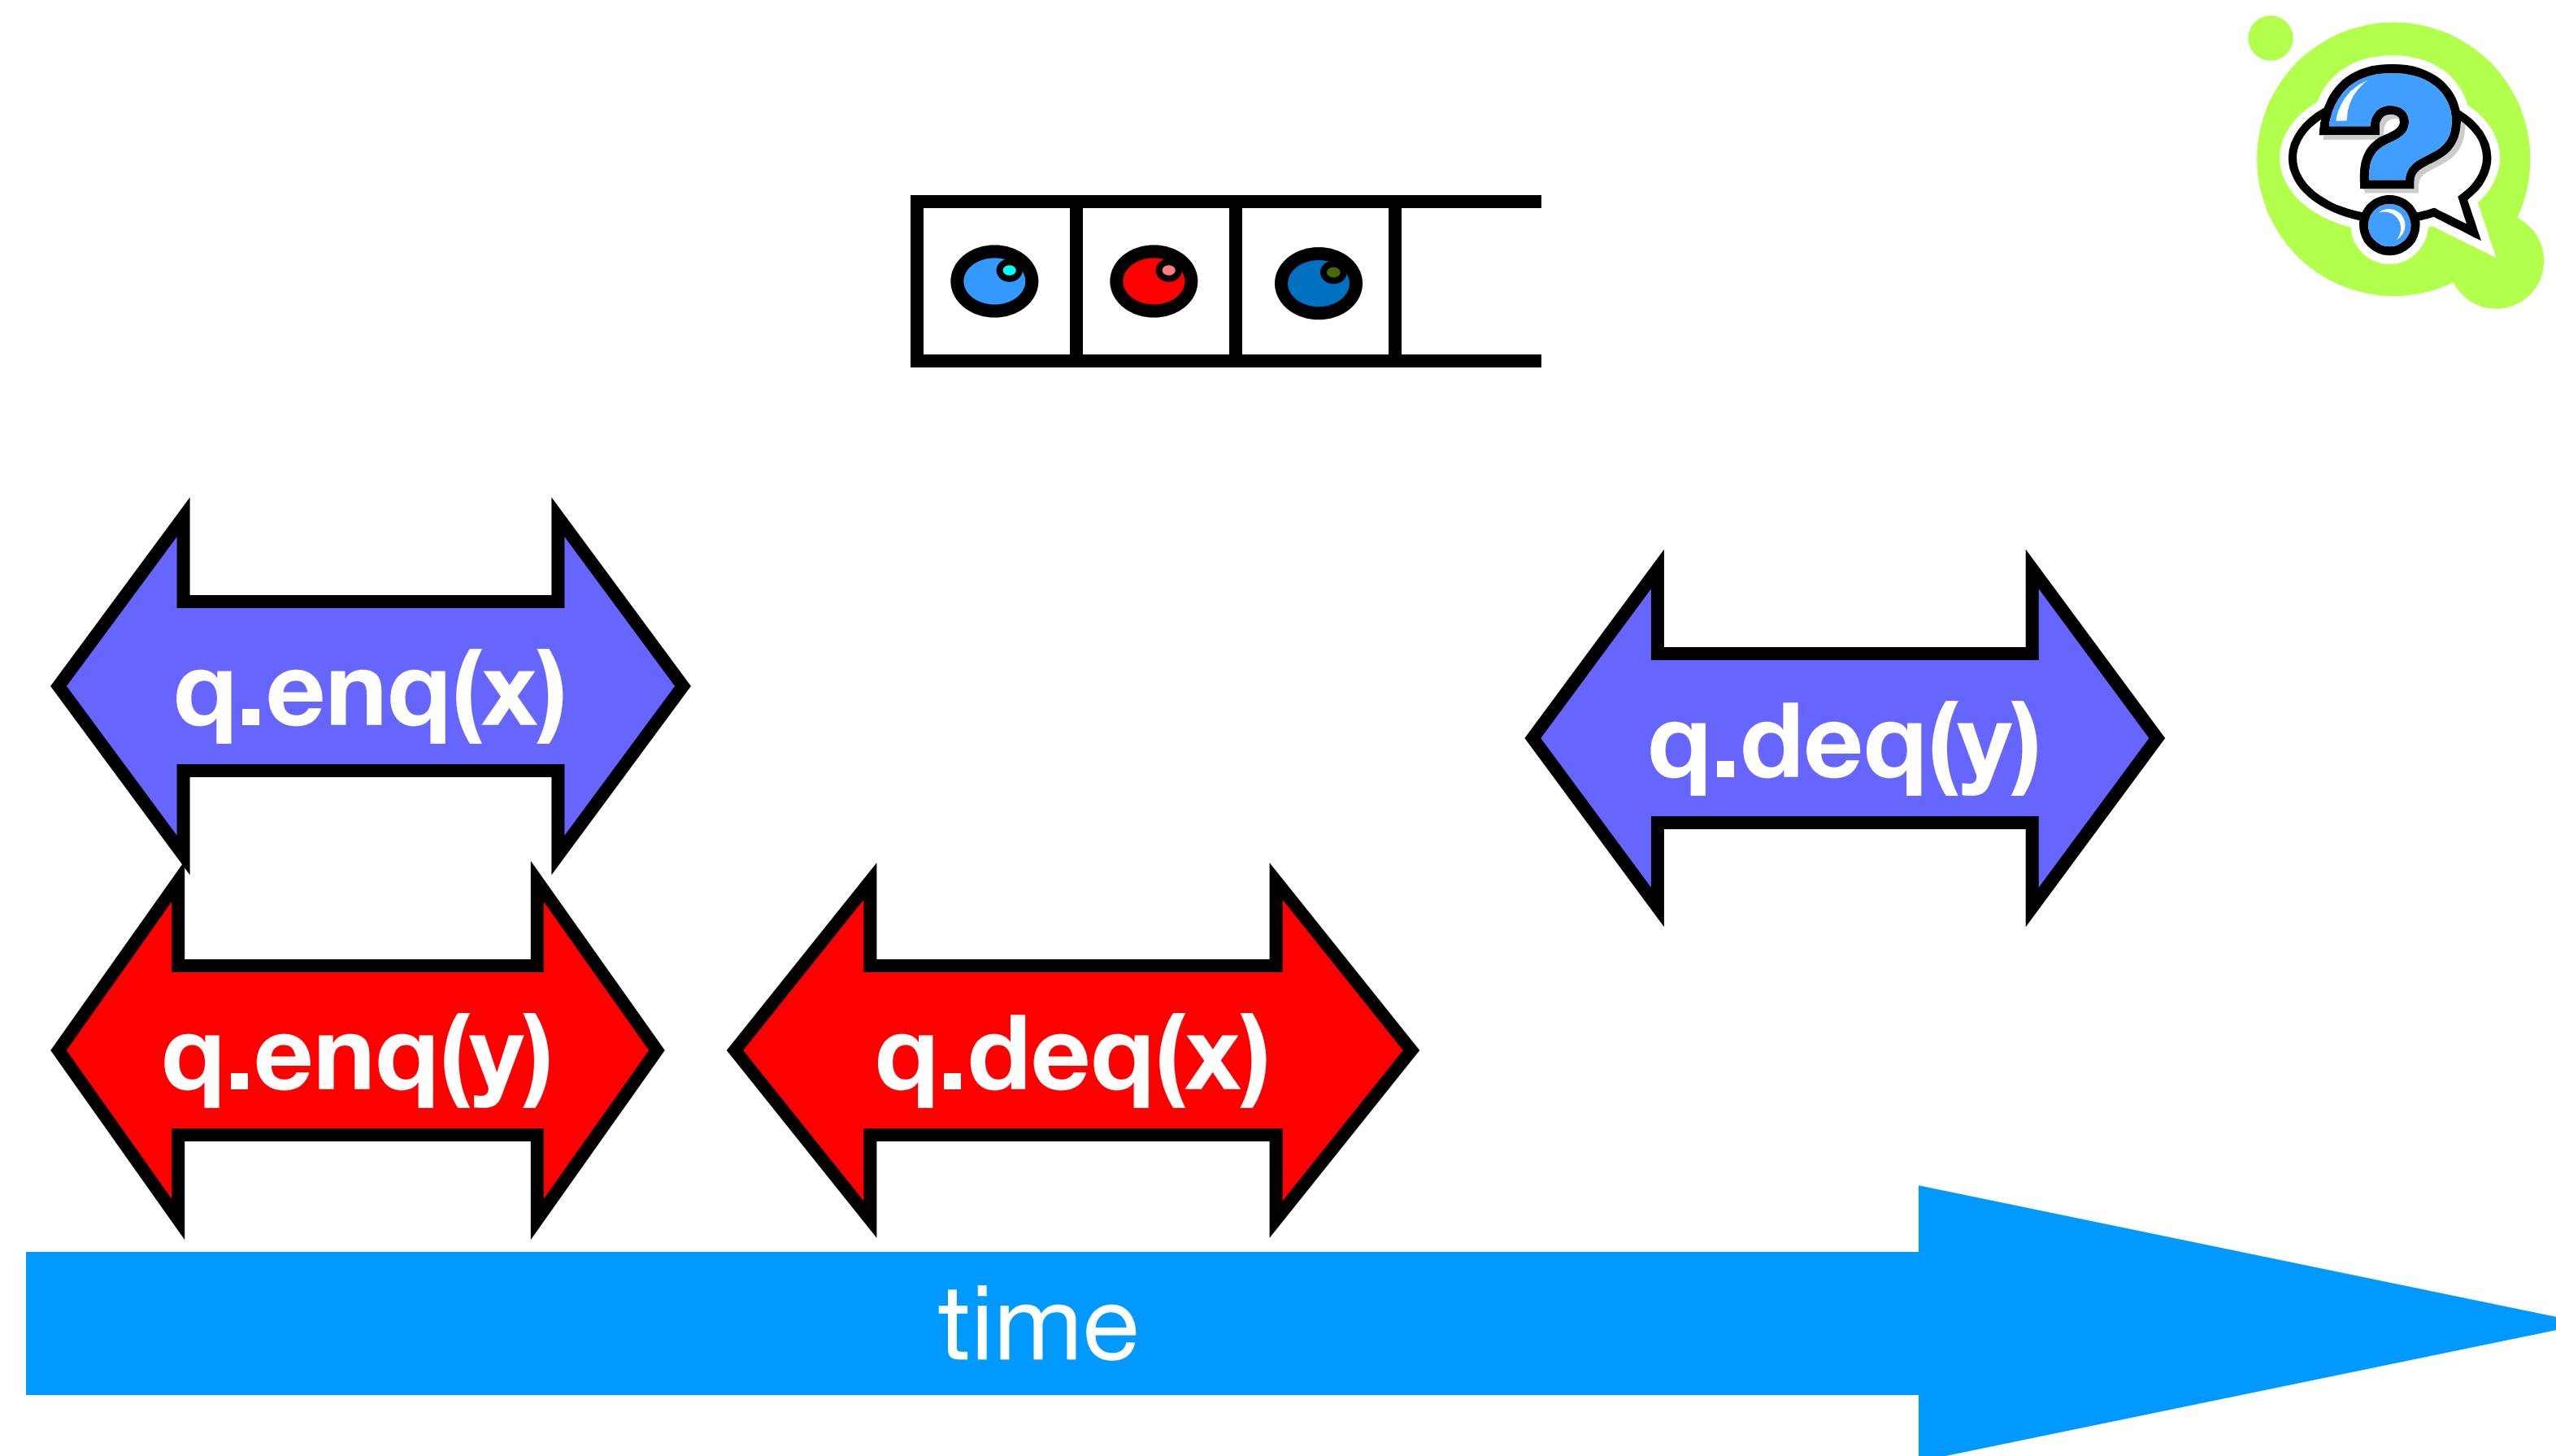
# Example



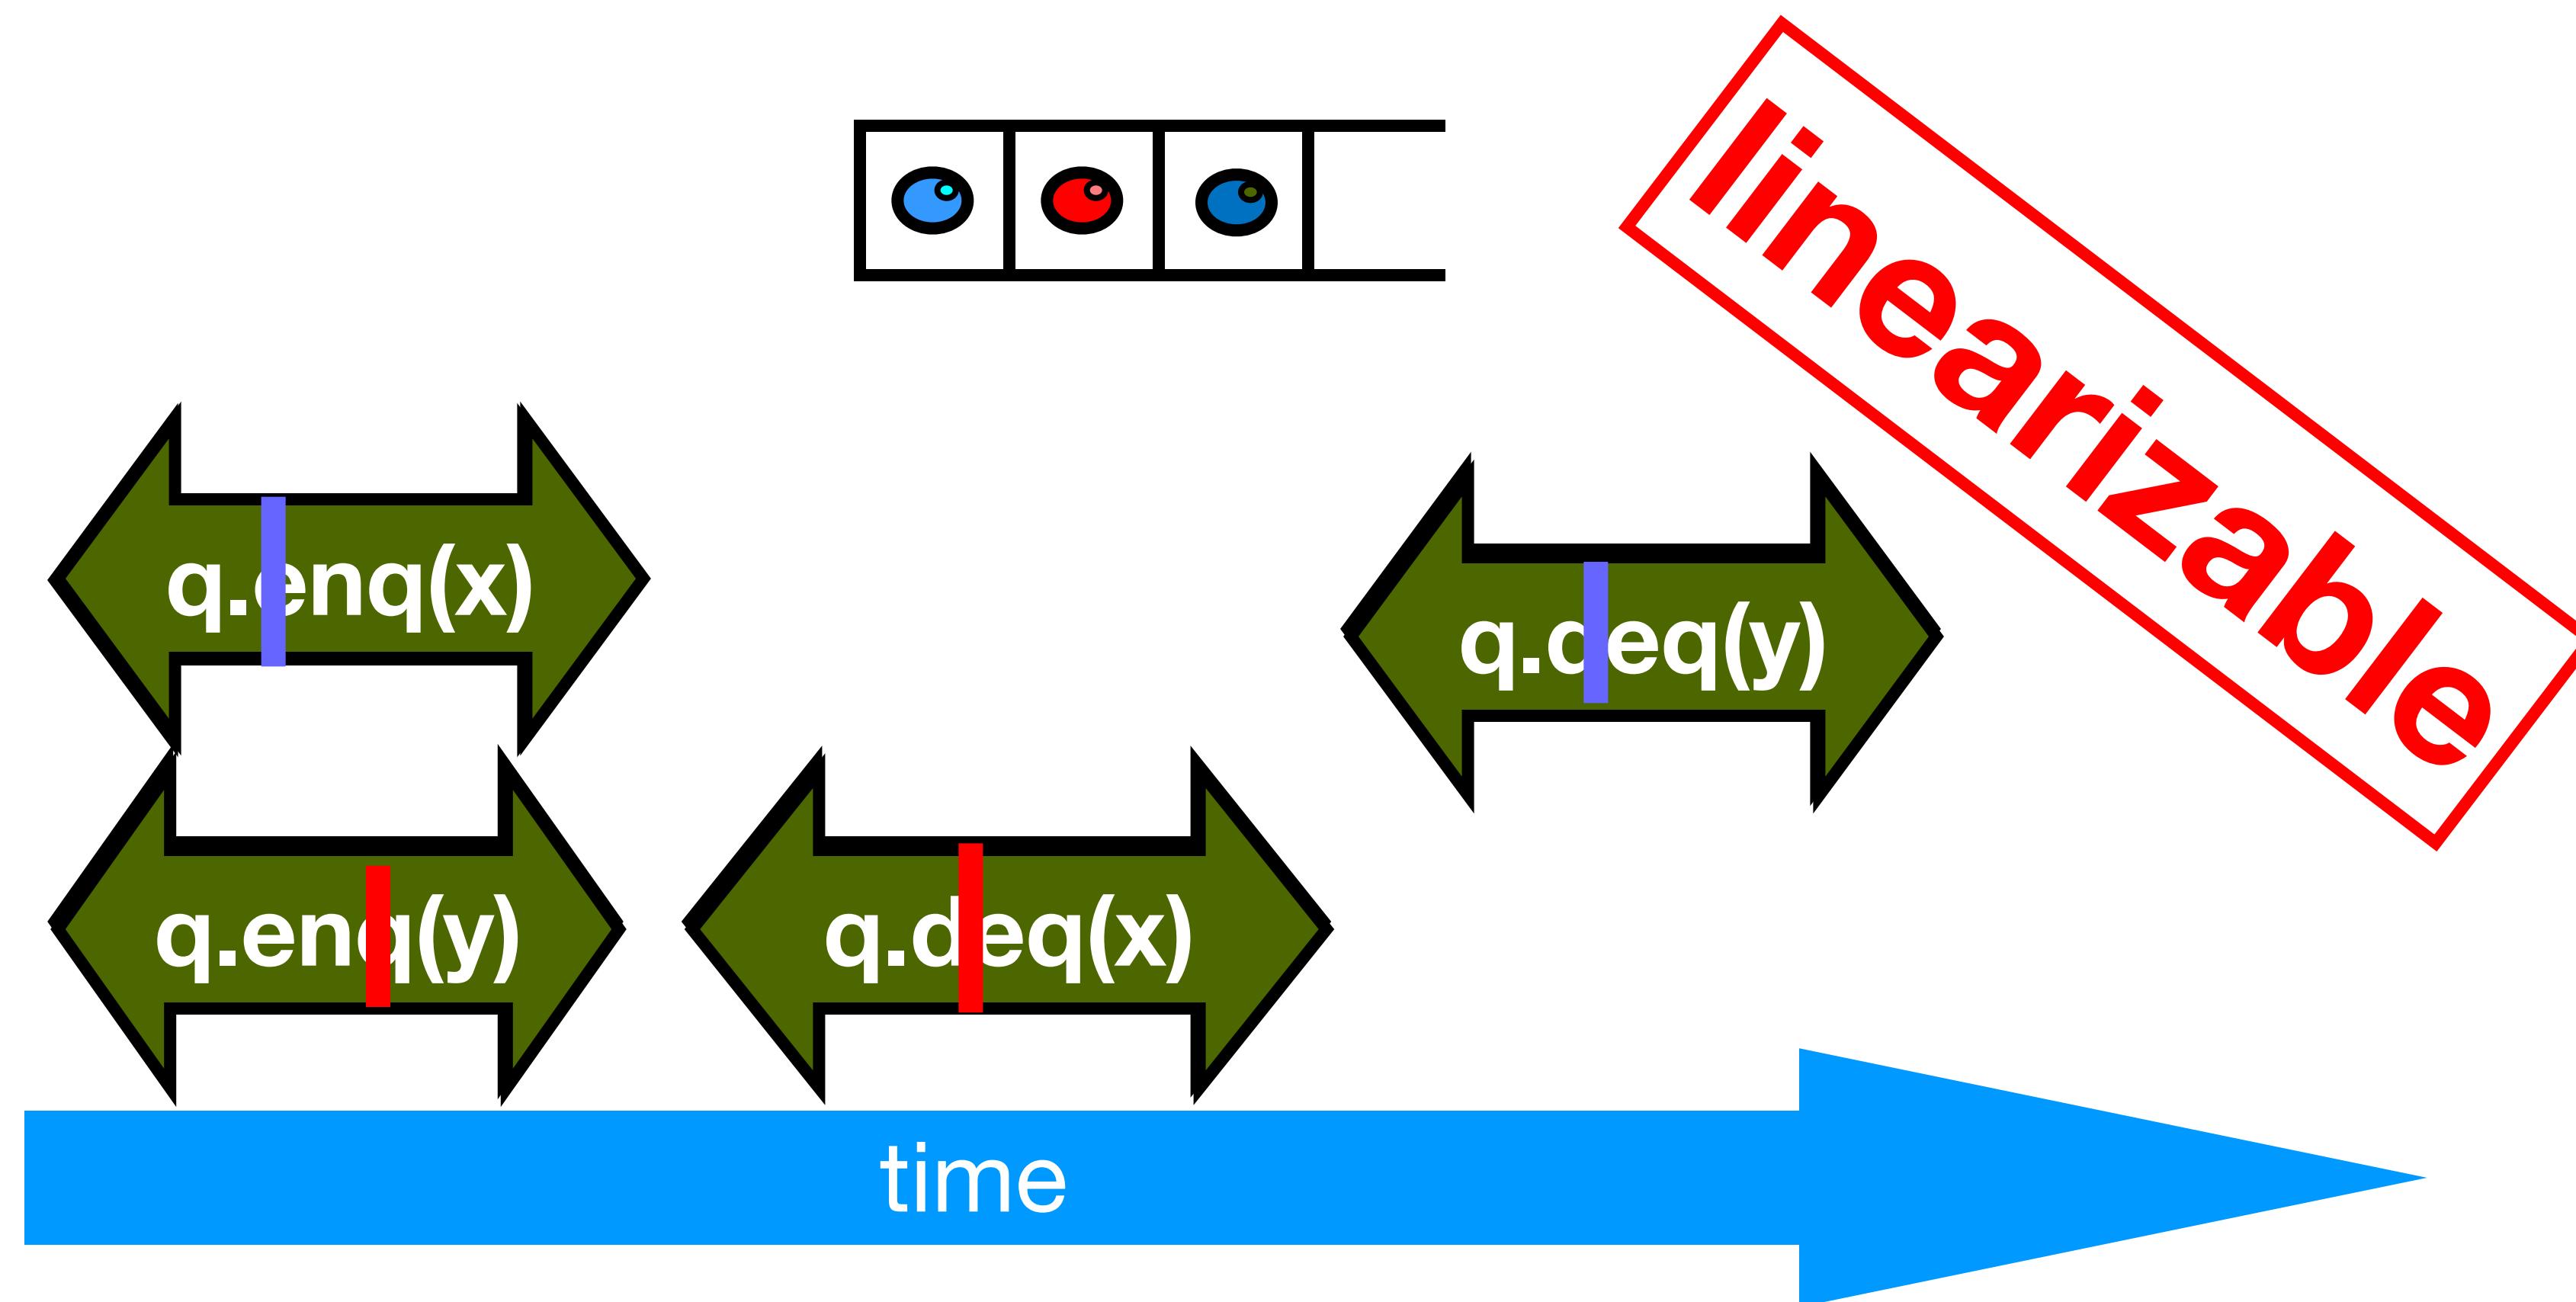
# Example



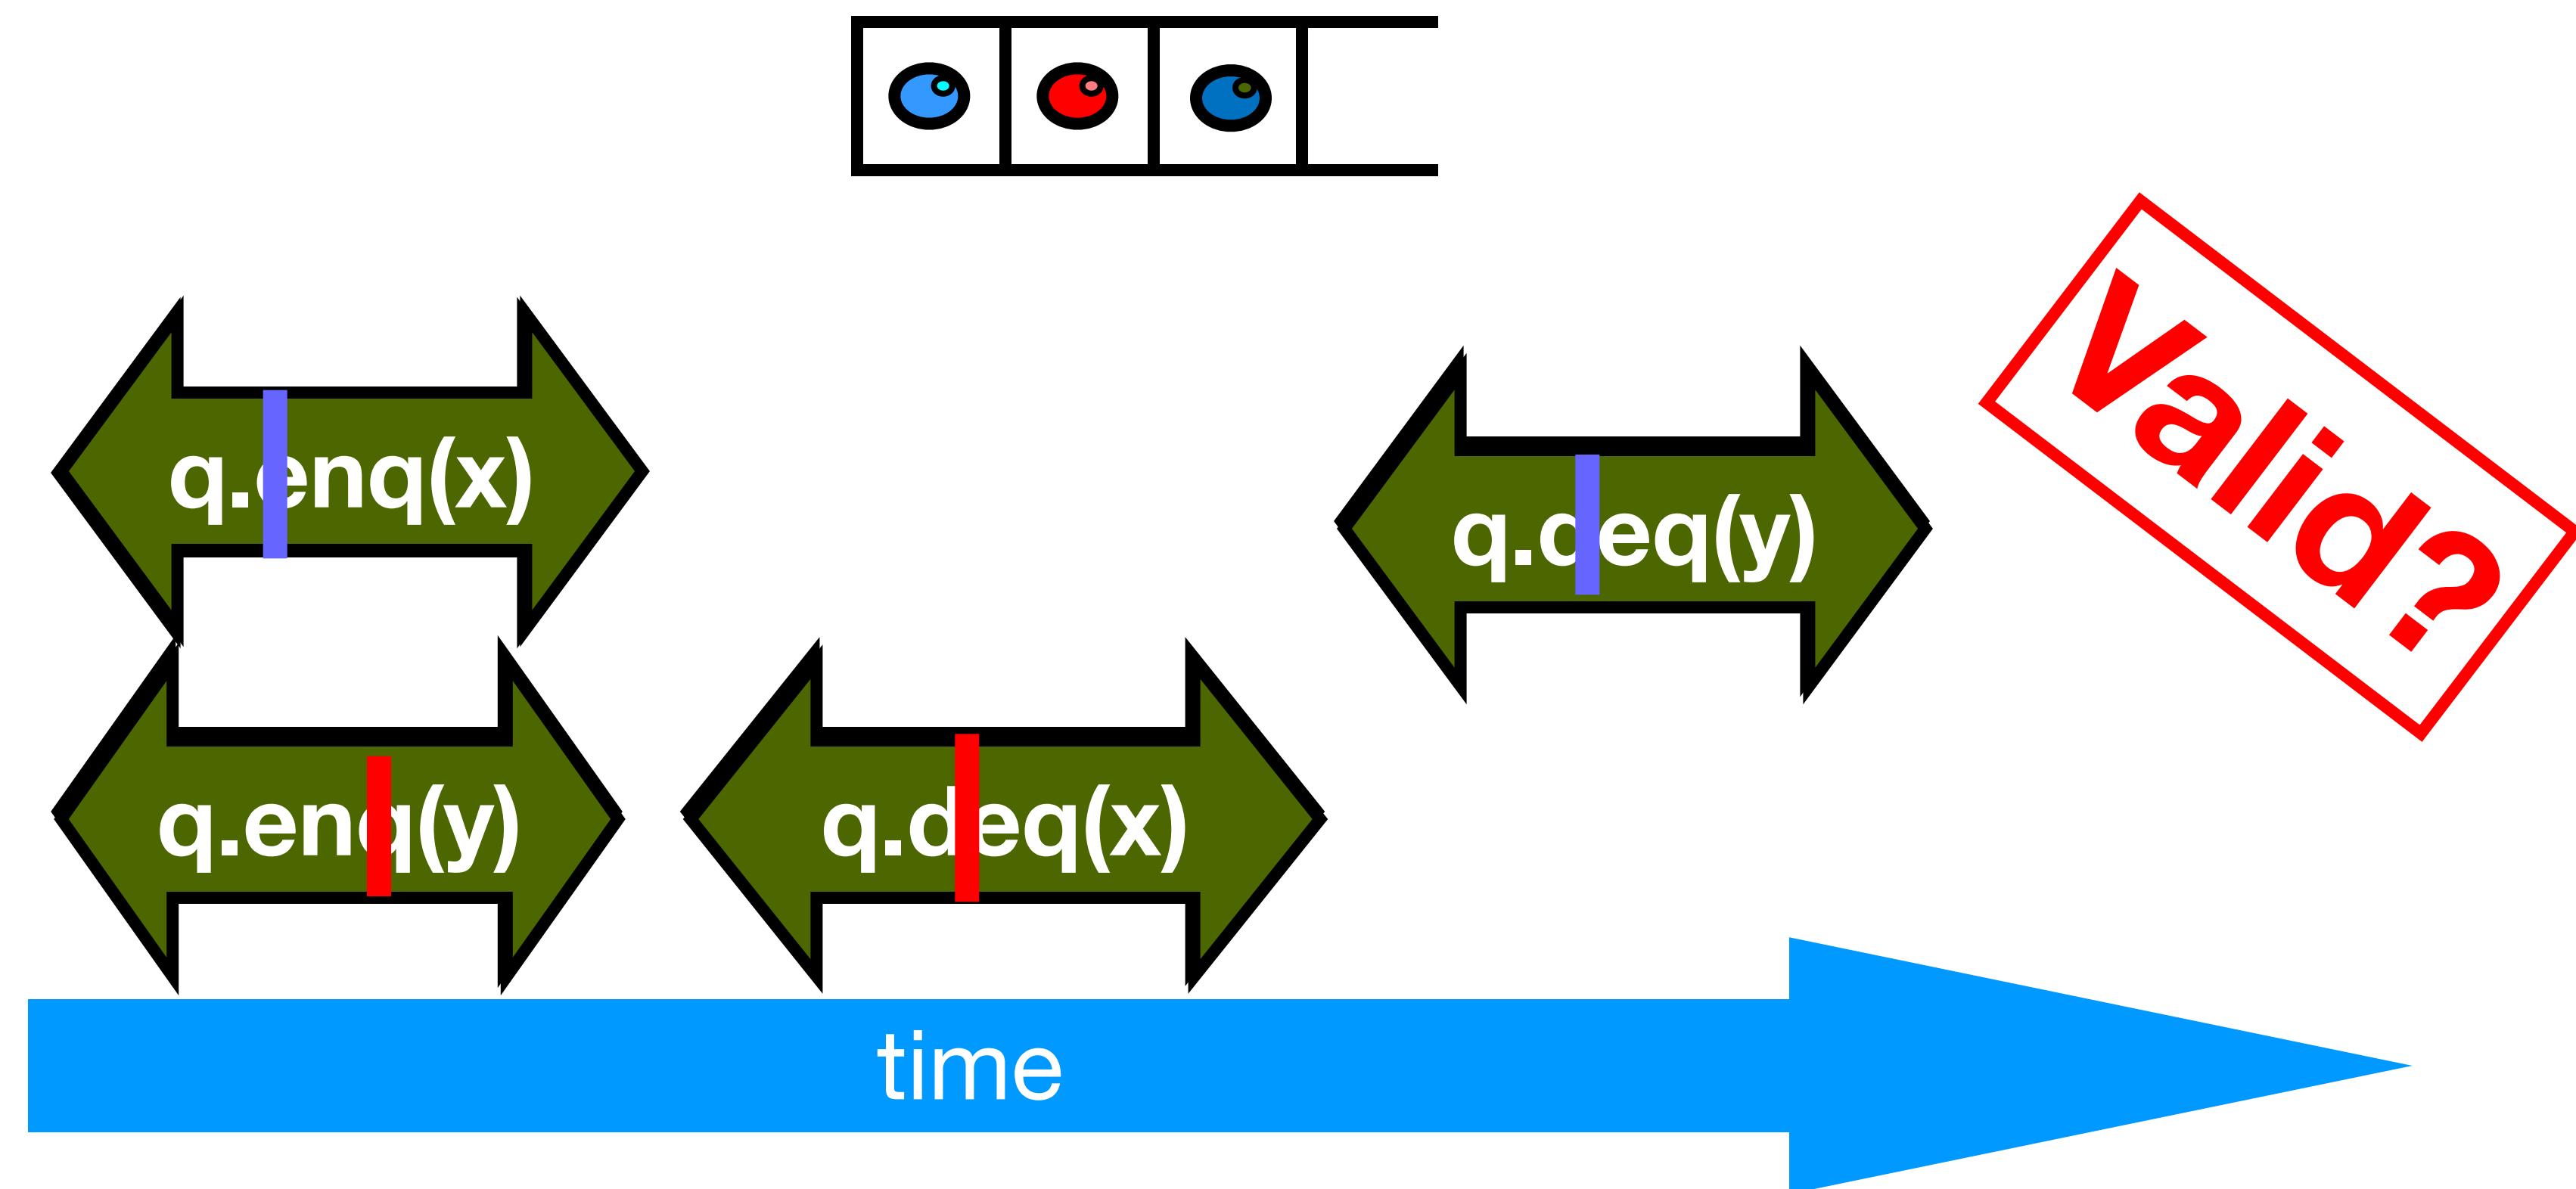
# Example



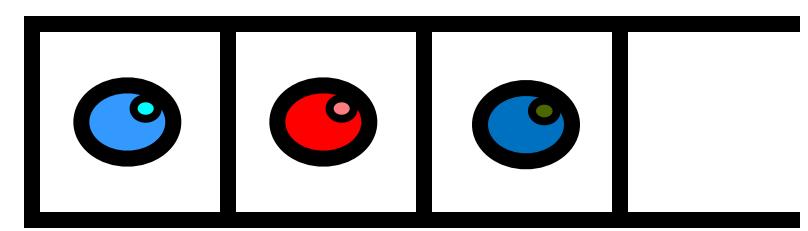
# Example



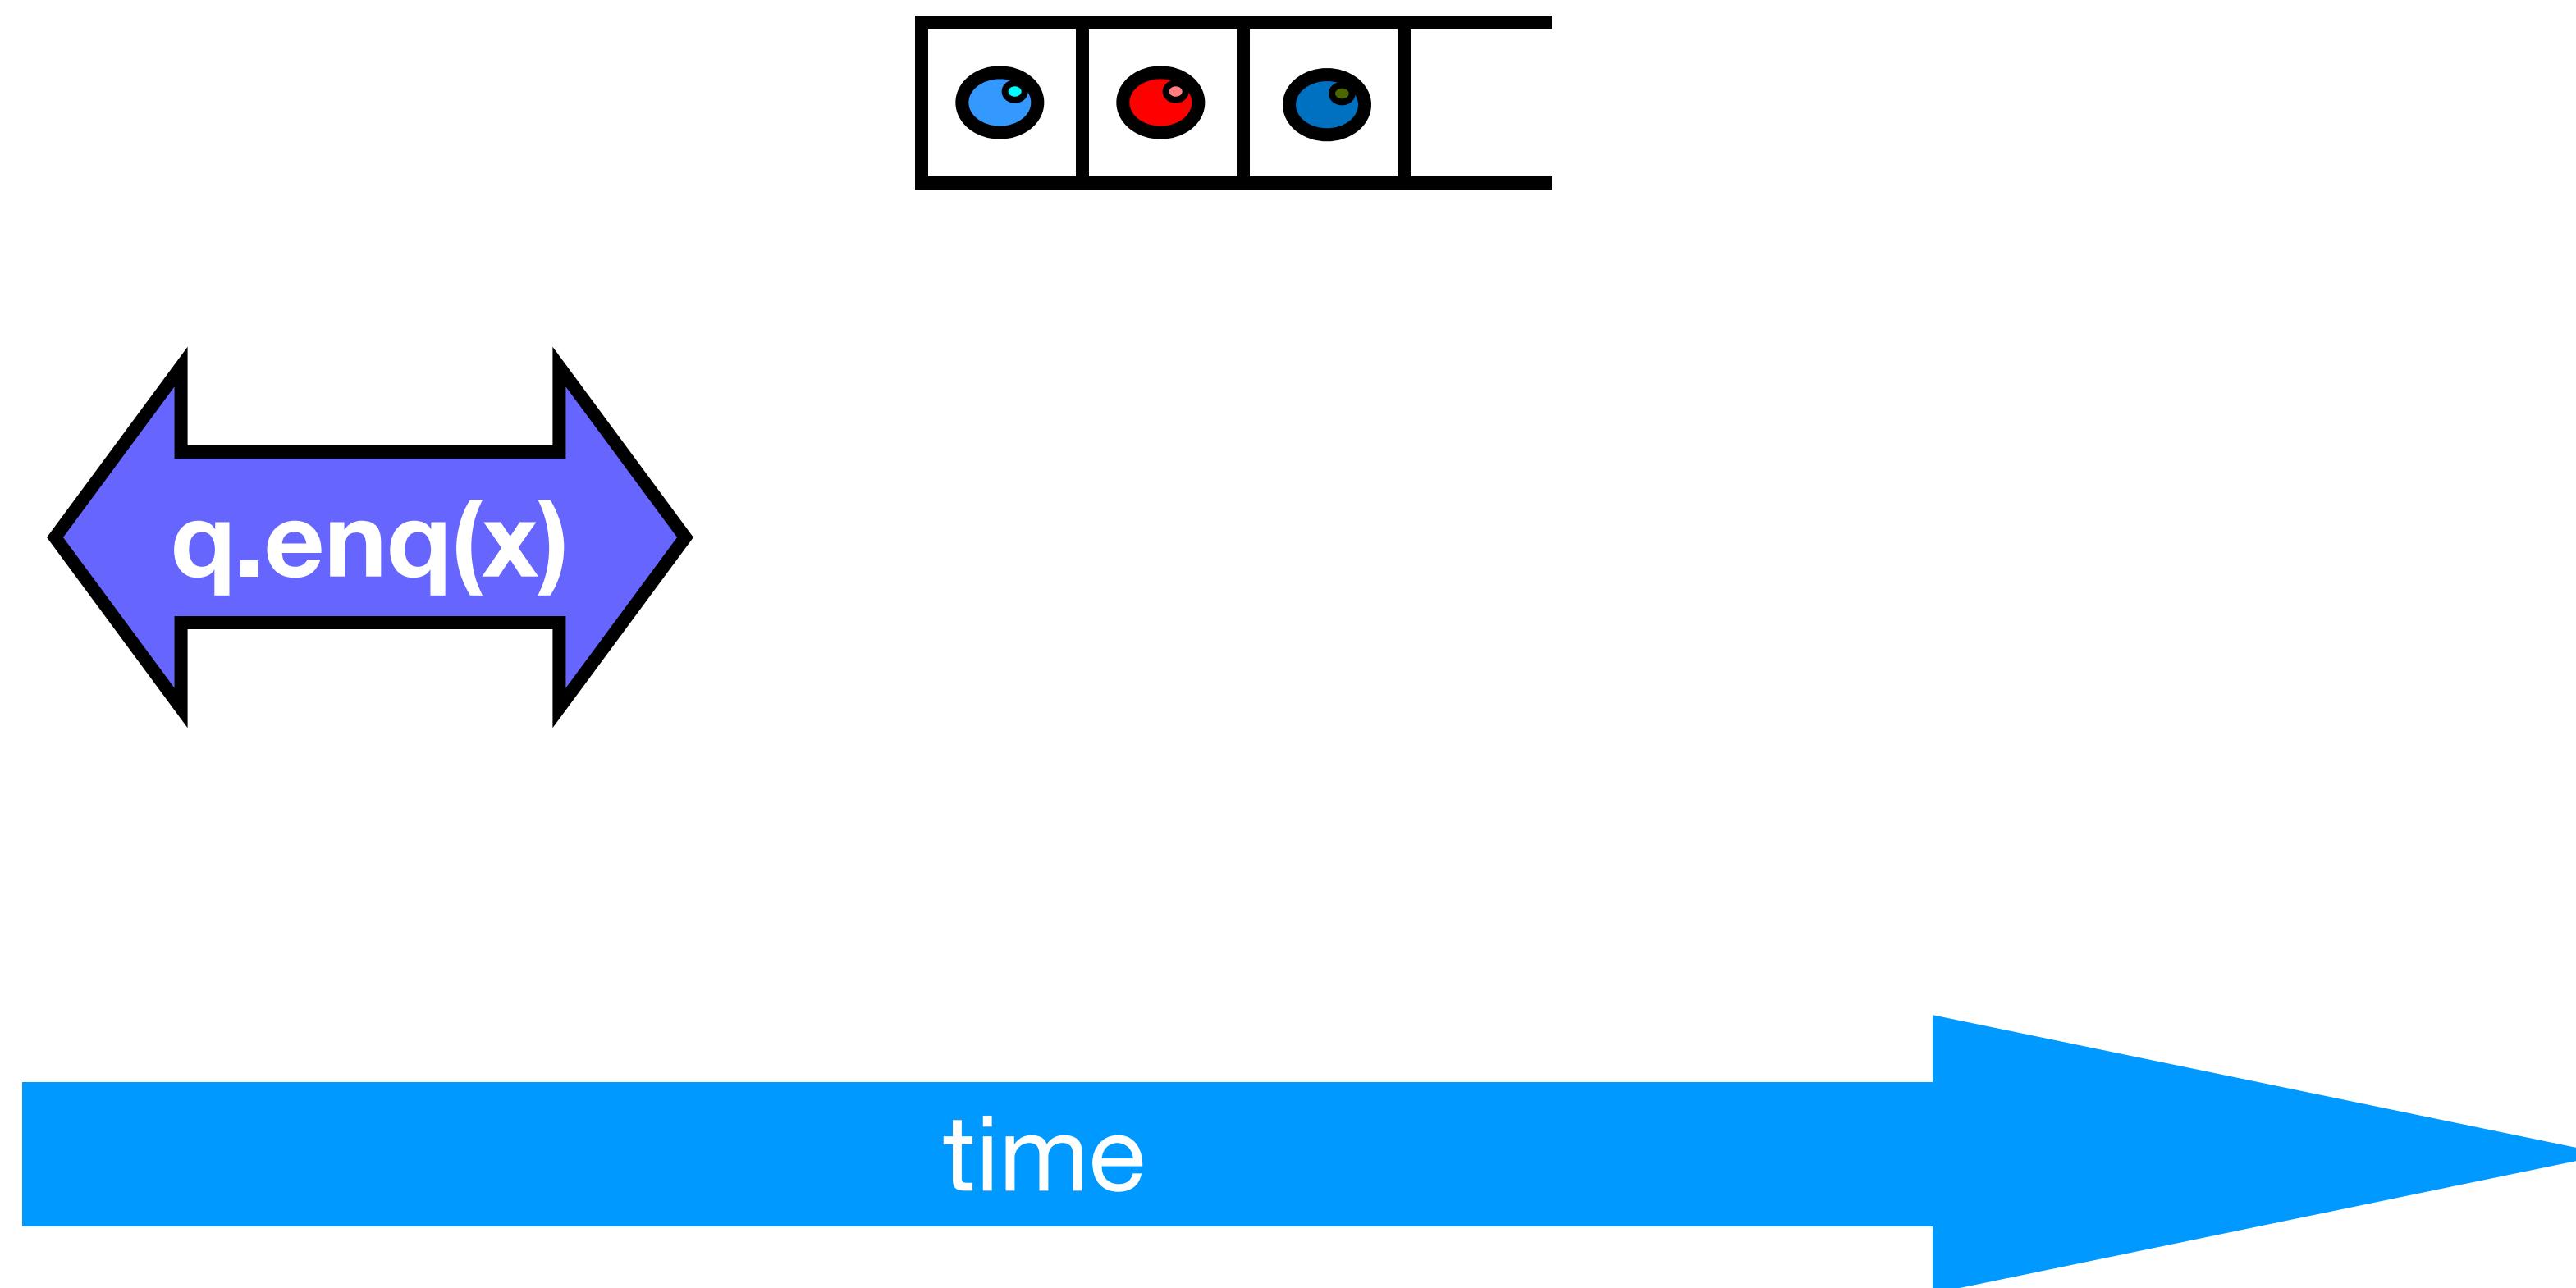
# Example



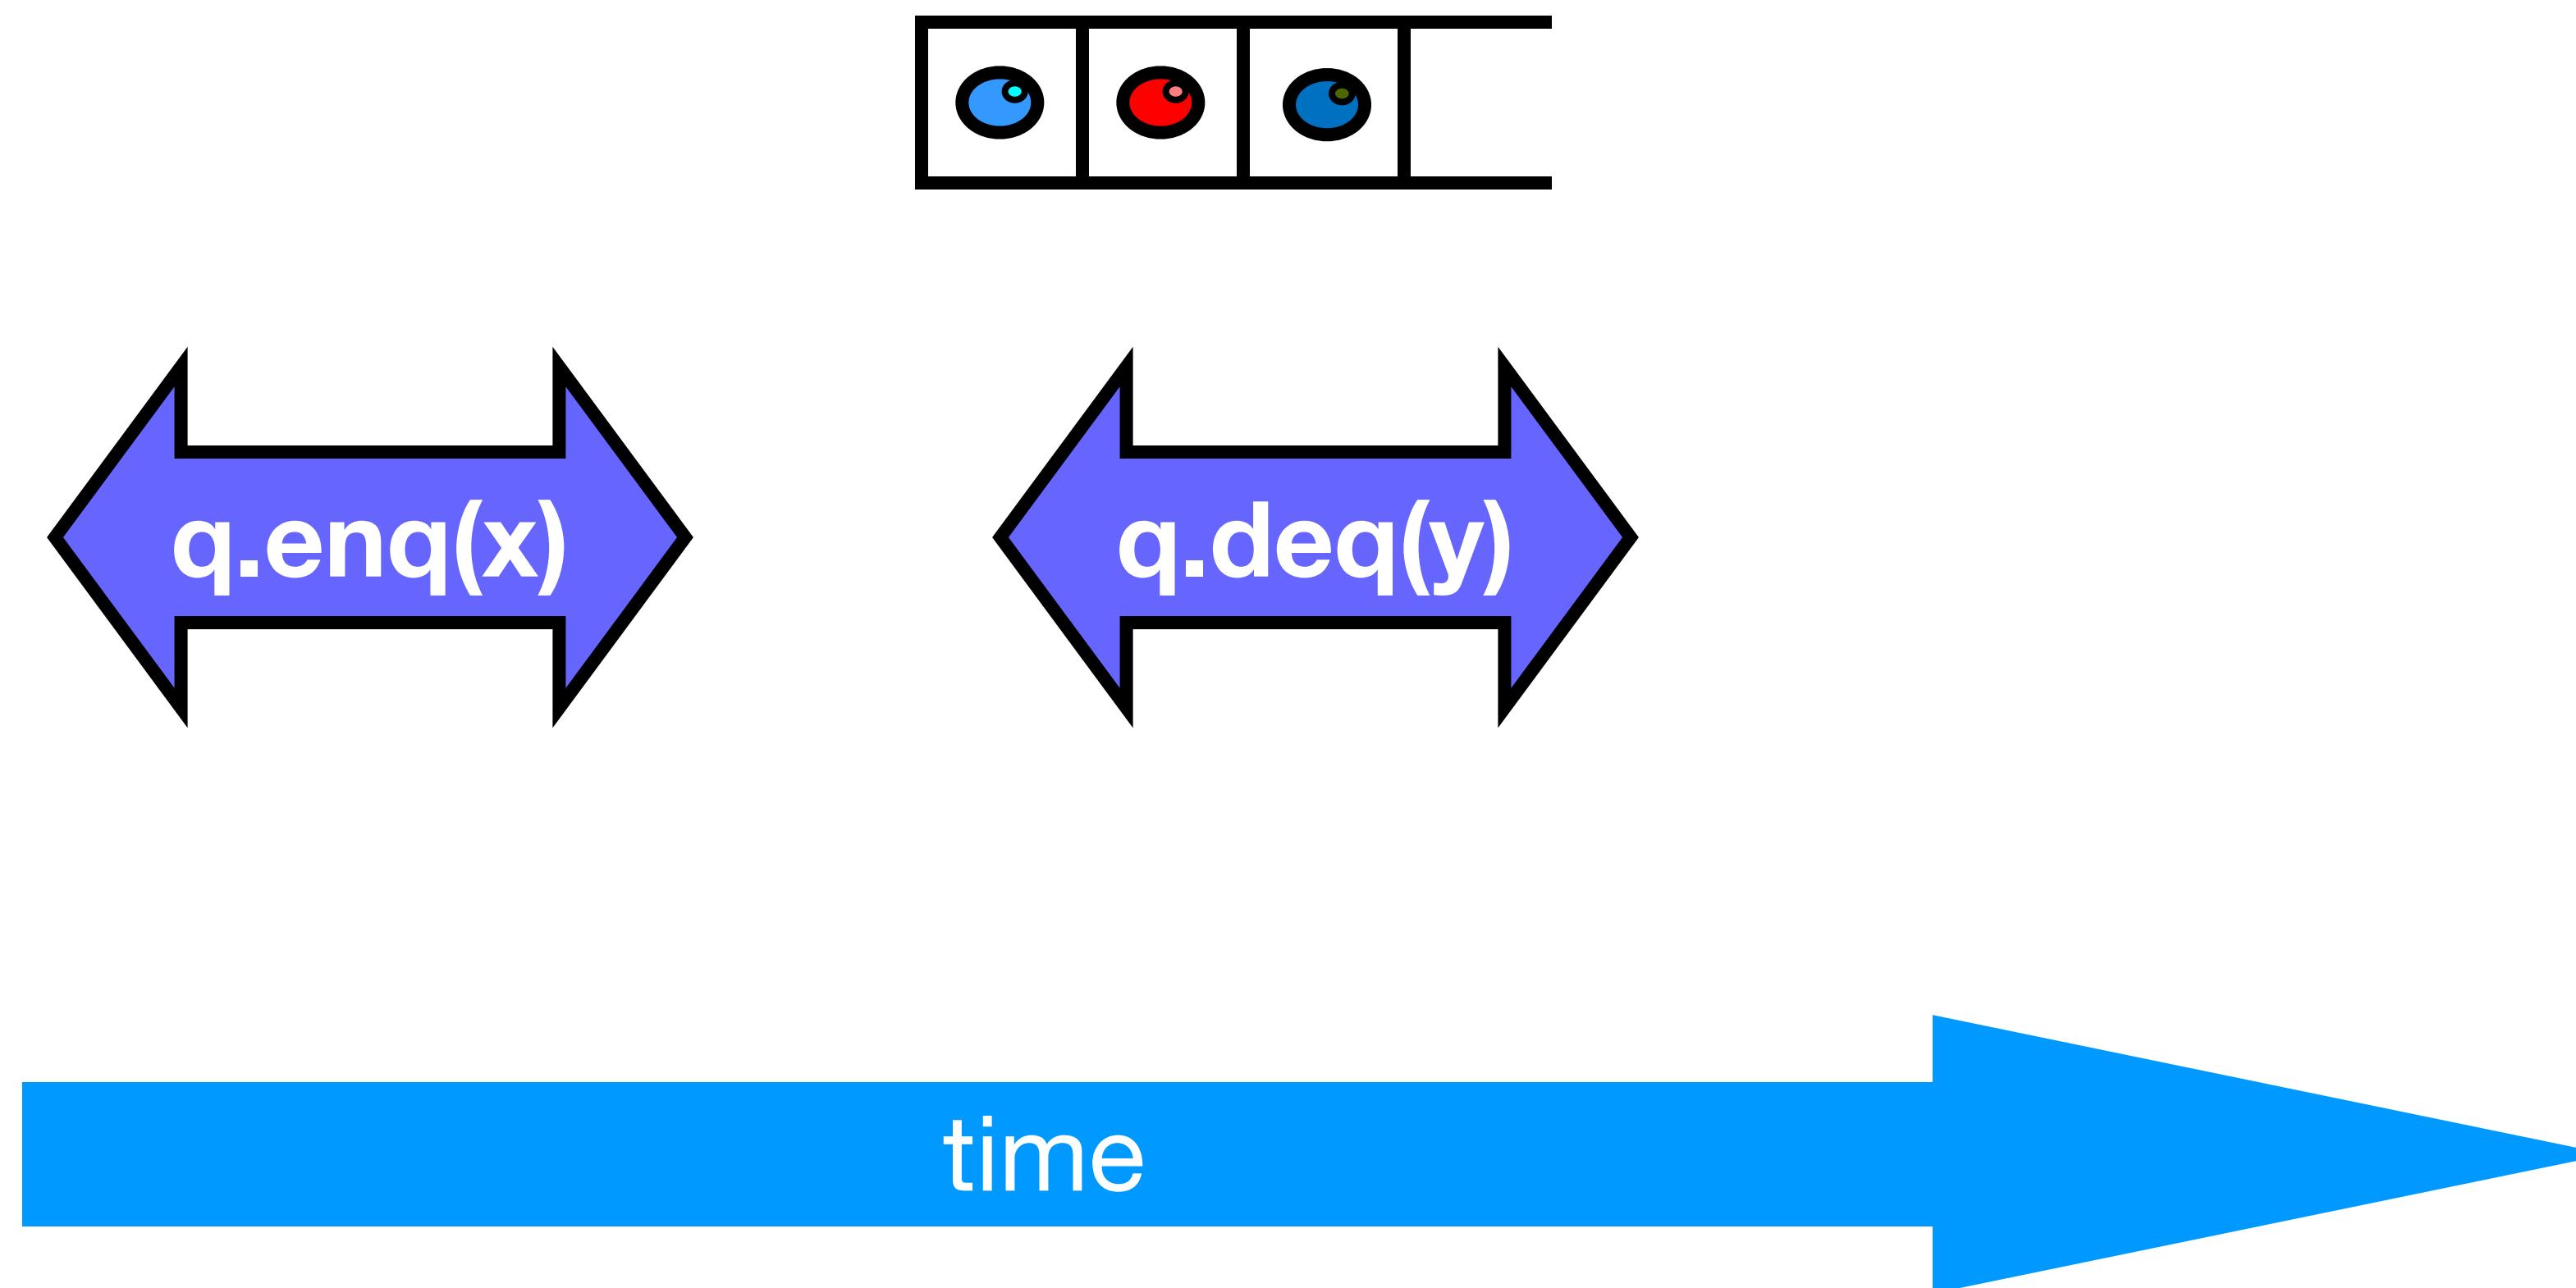
# Example



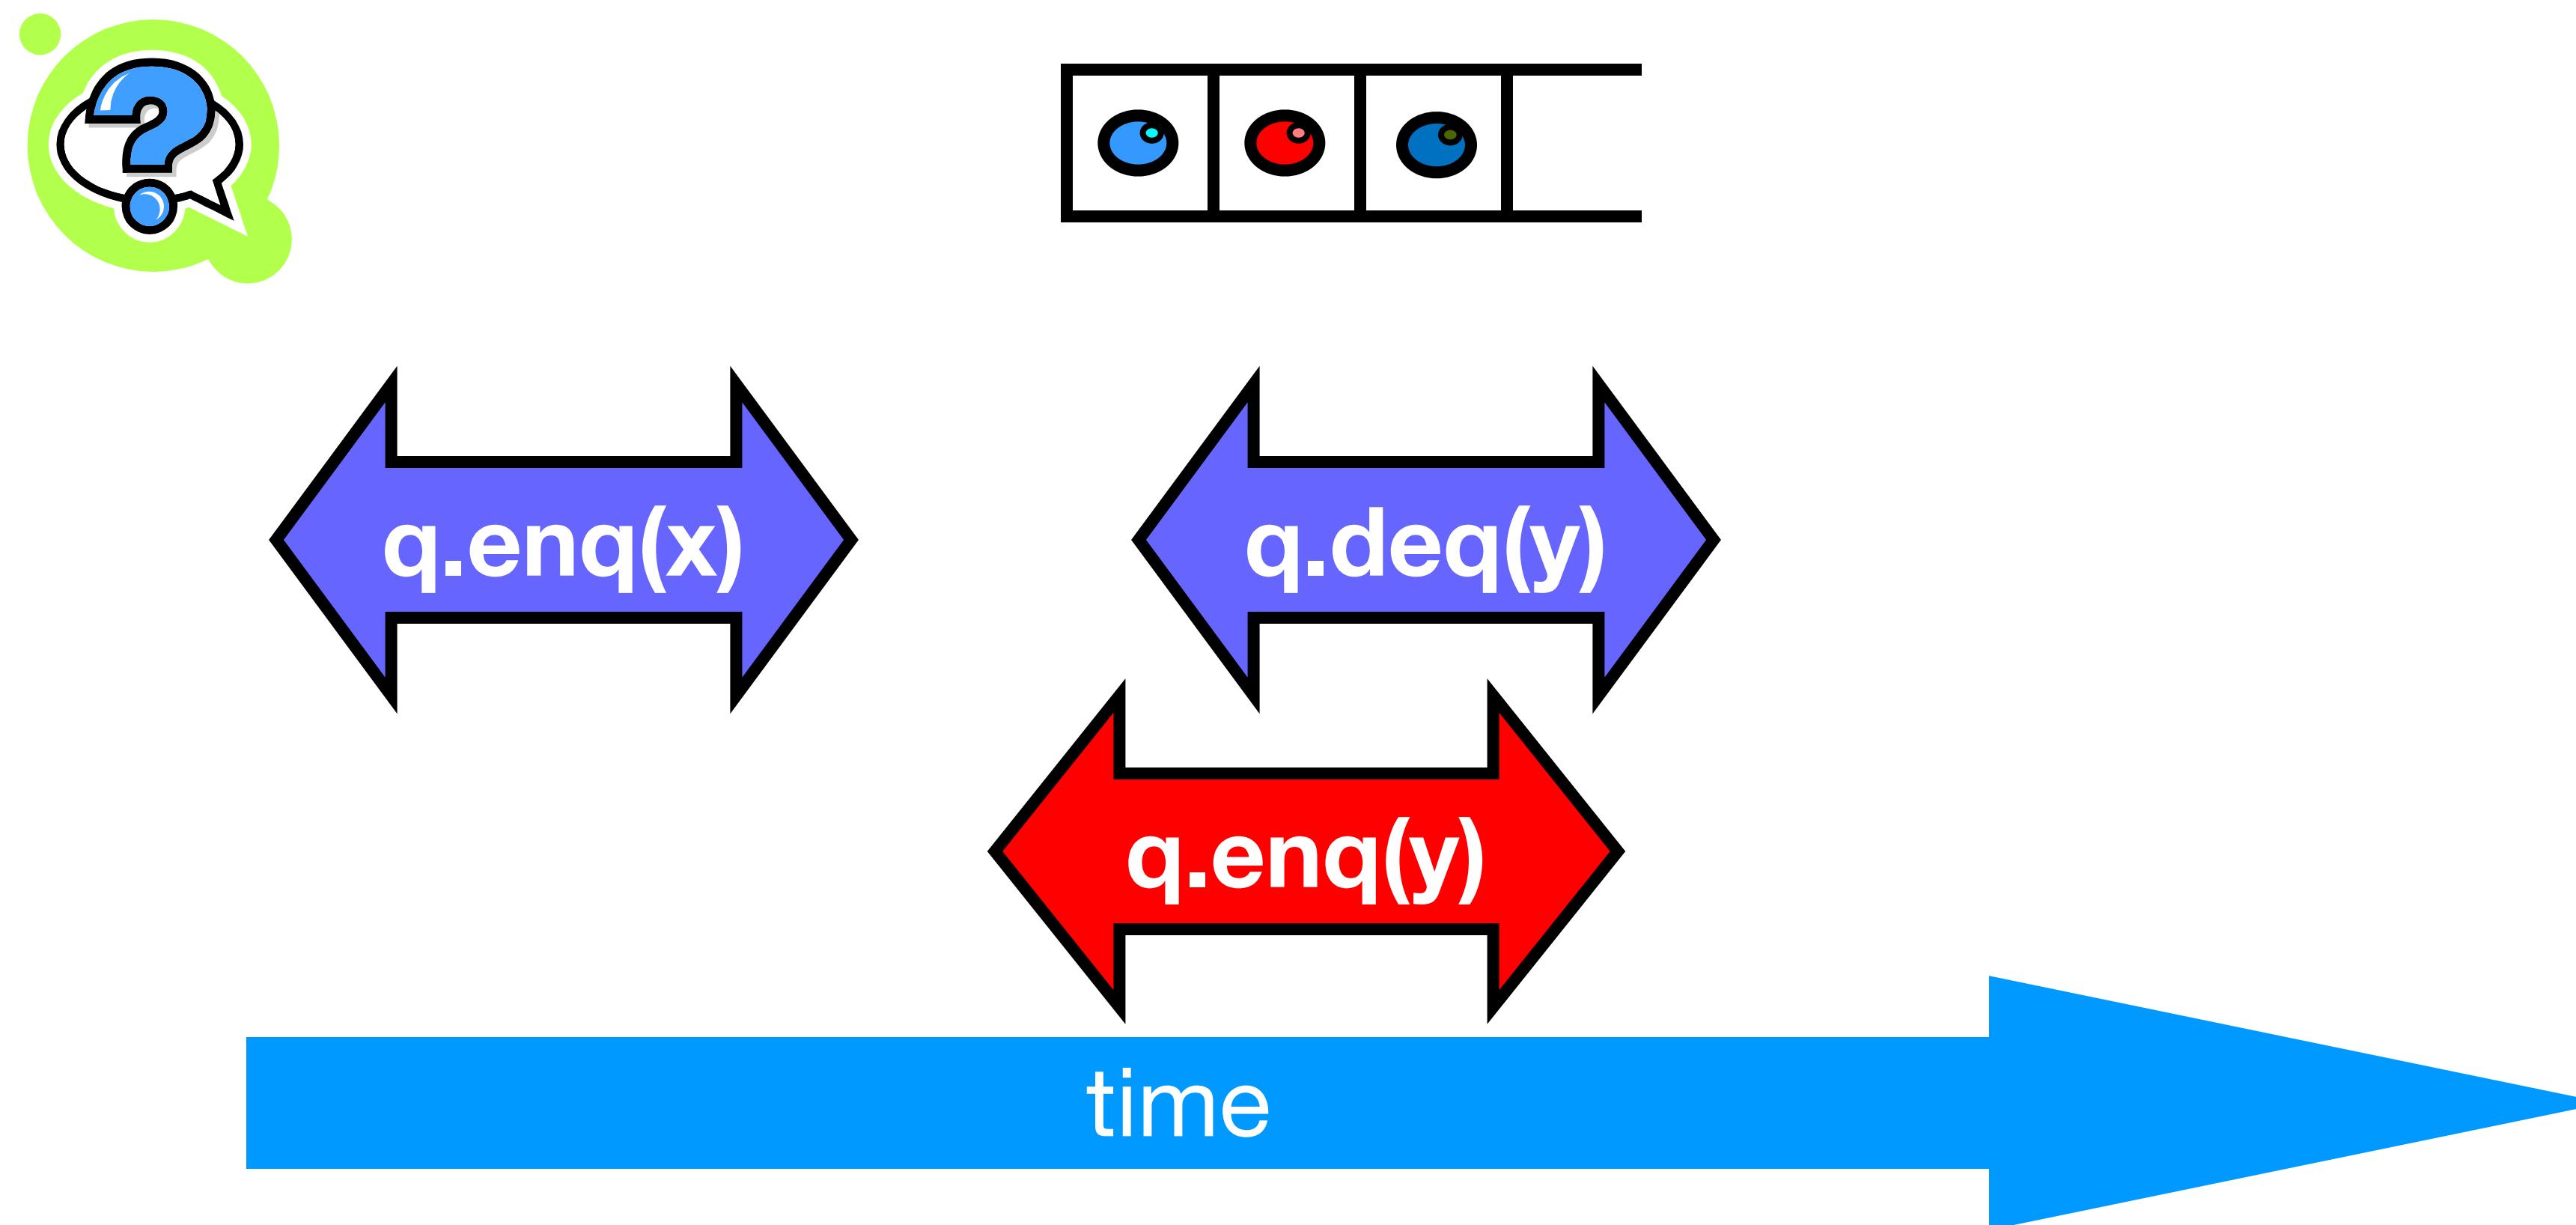
# Example



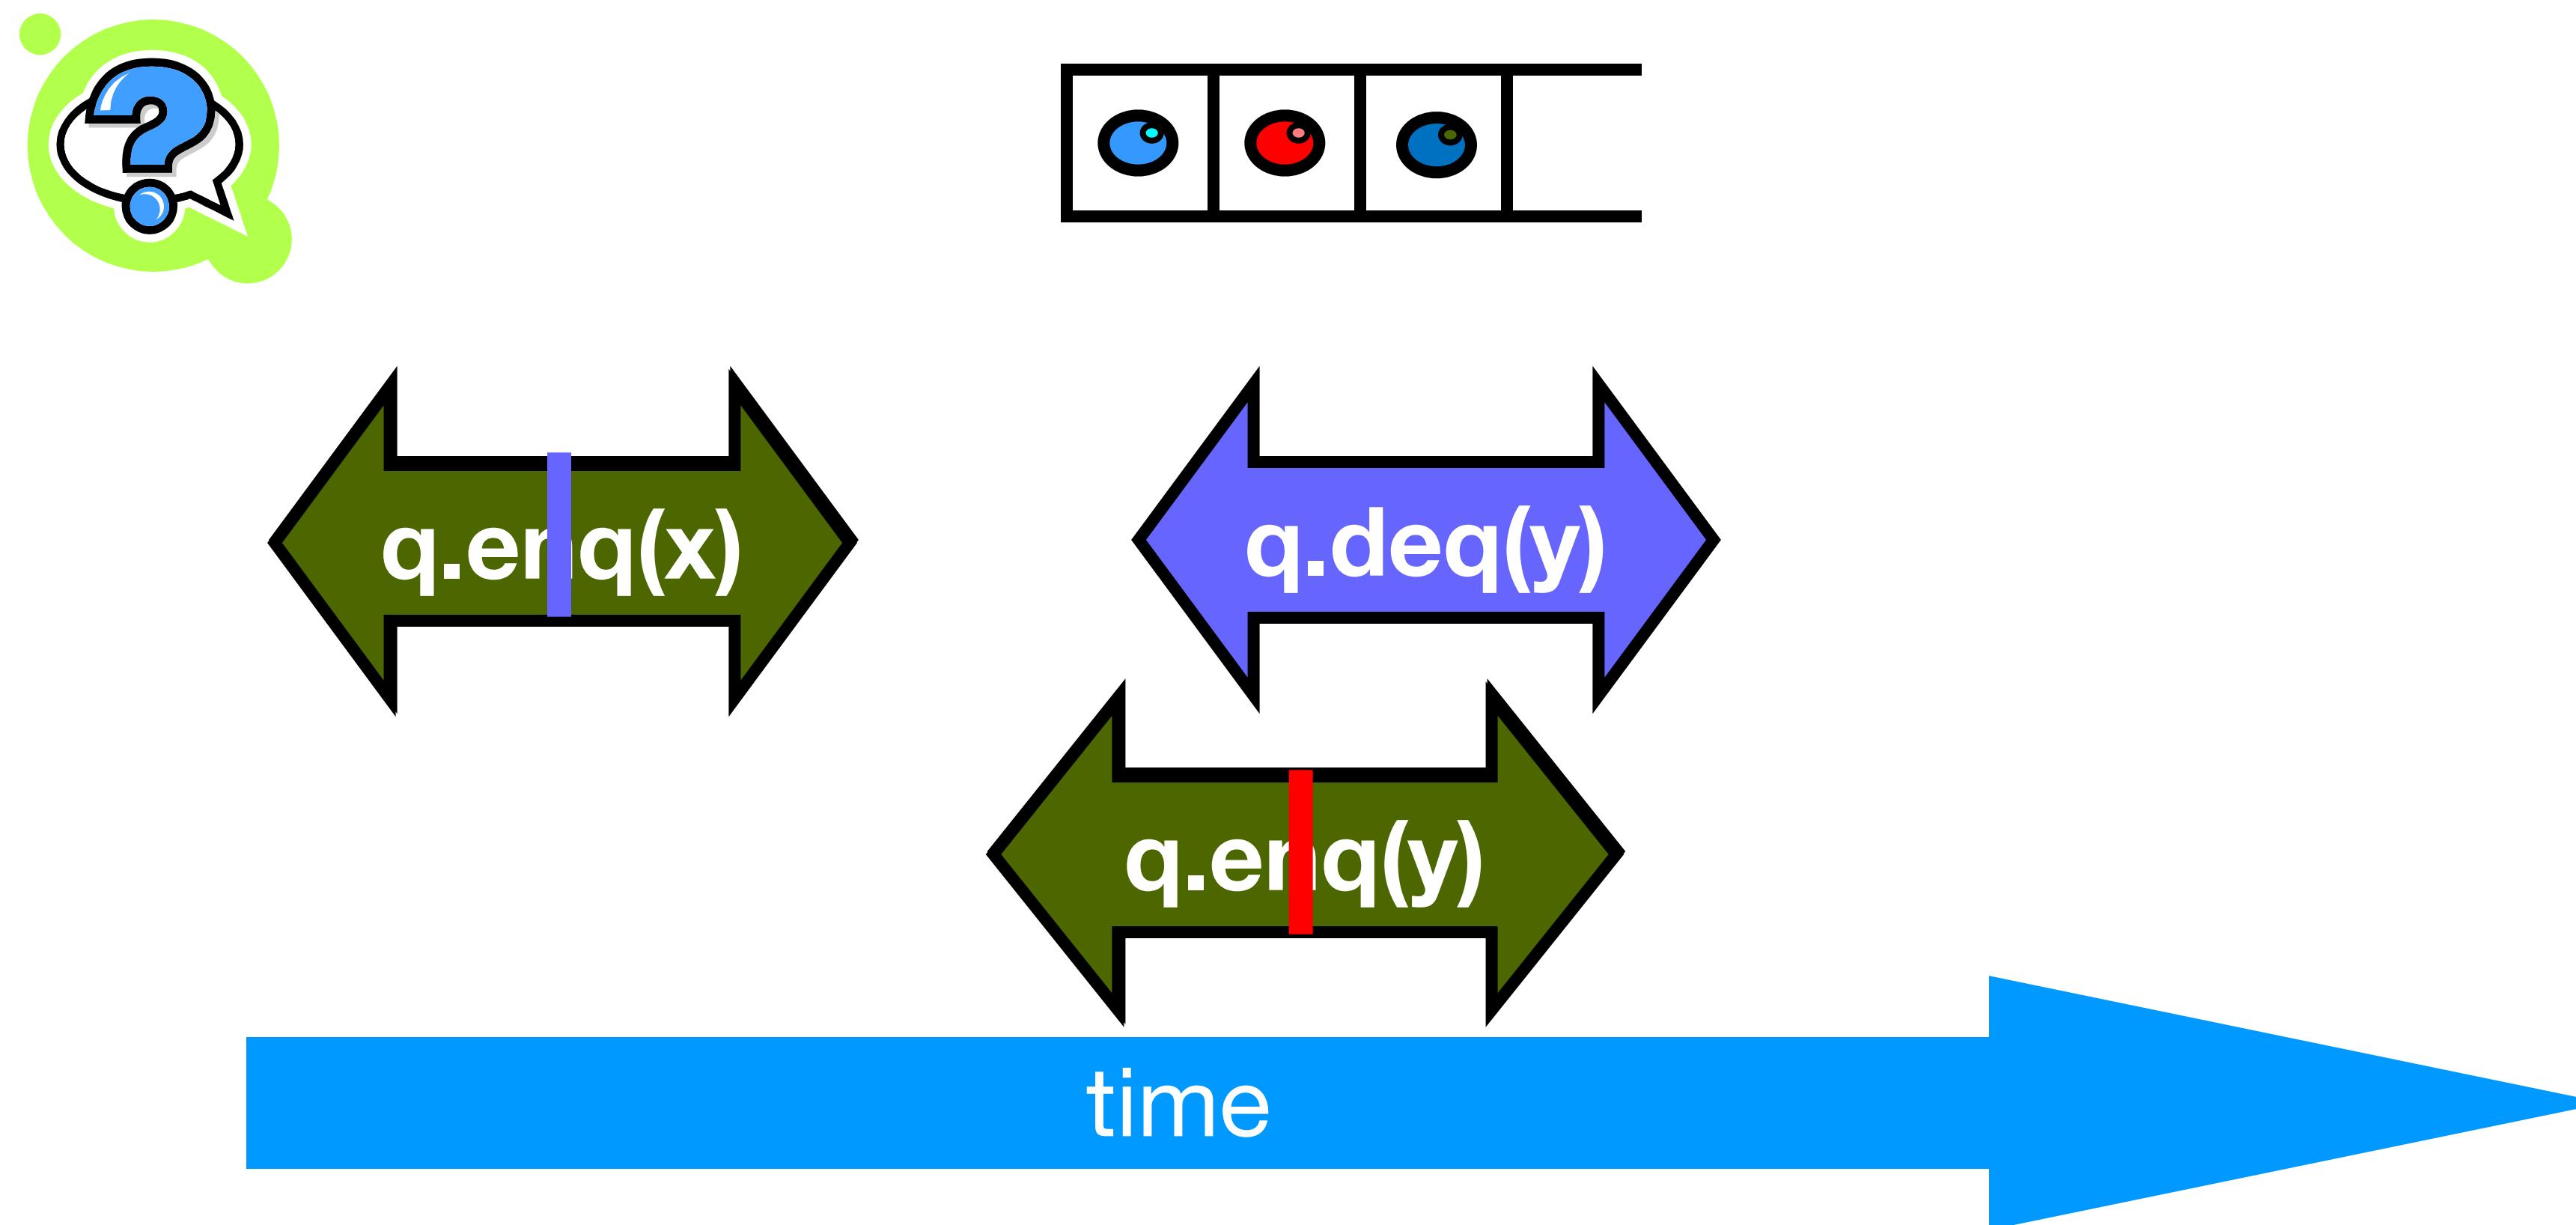
# Example



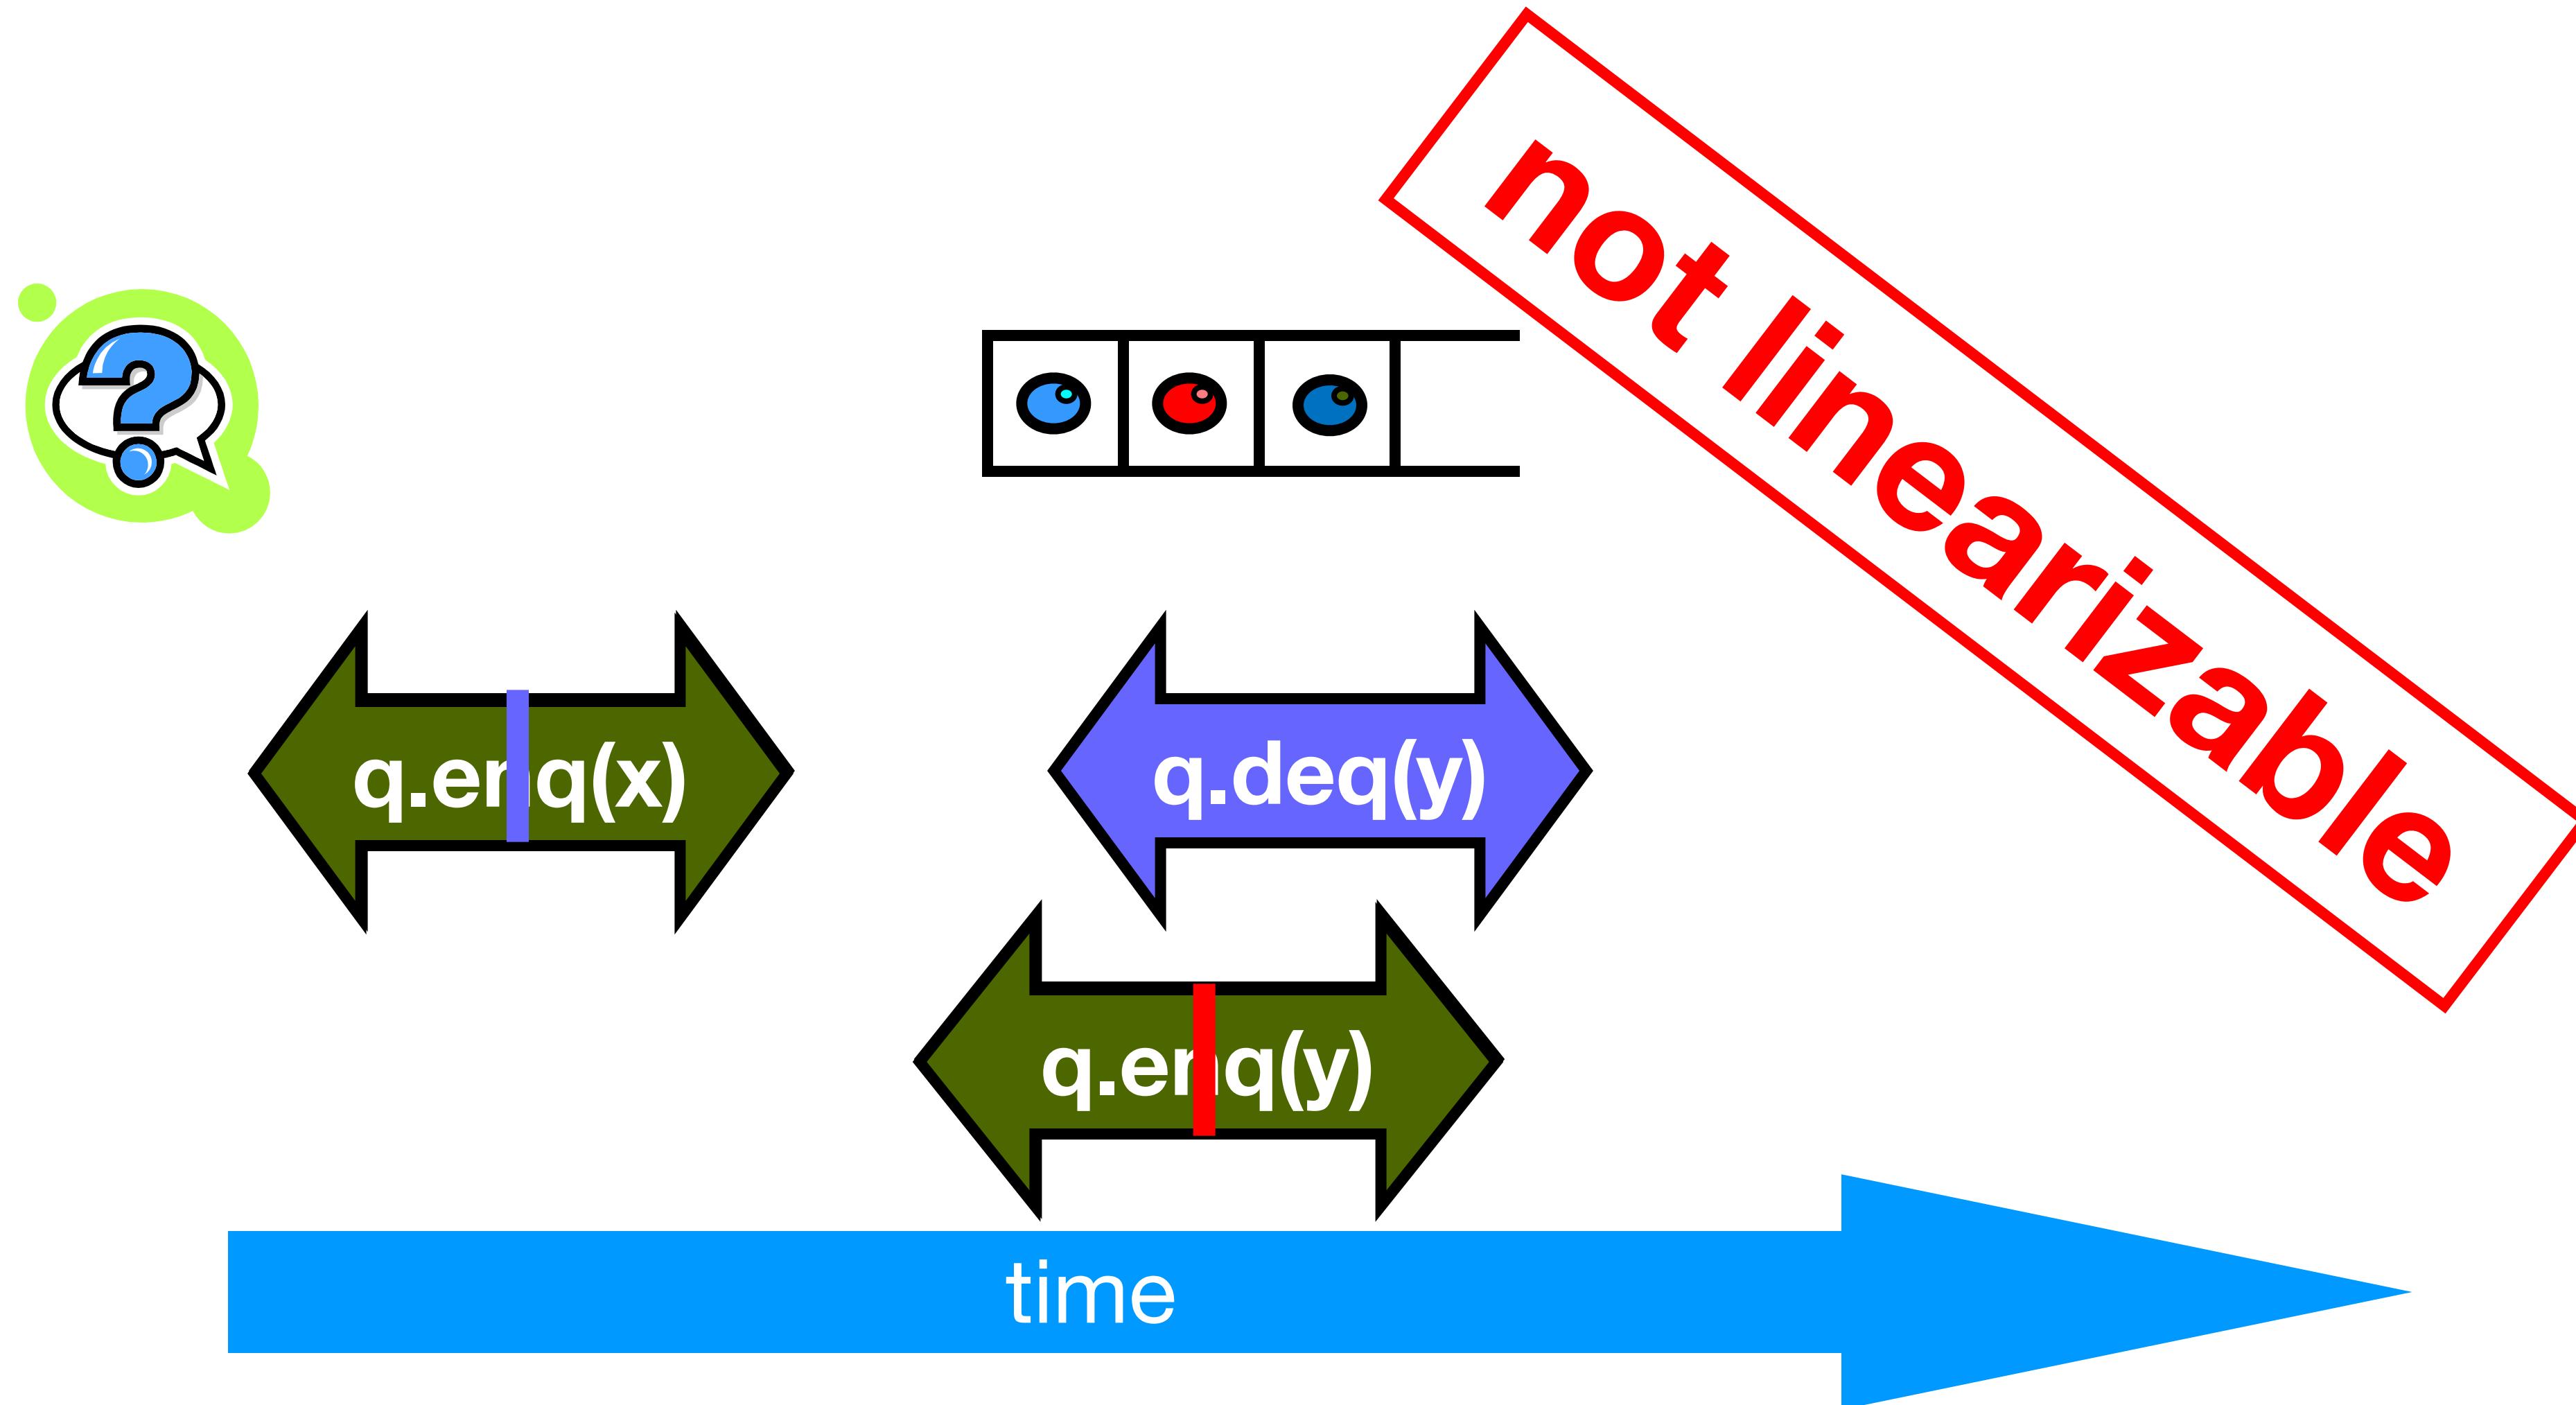
# Example



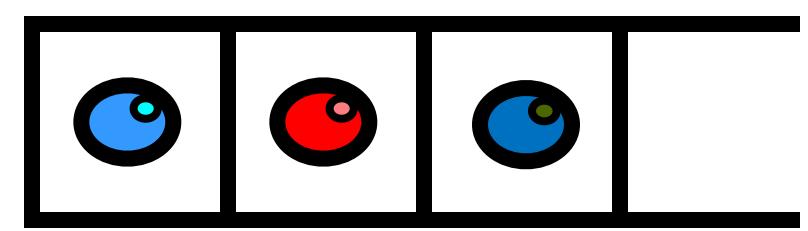
# Example



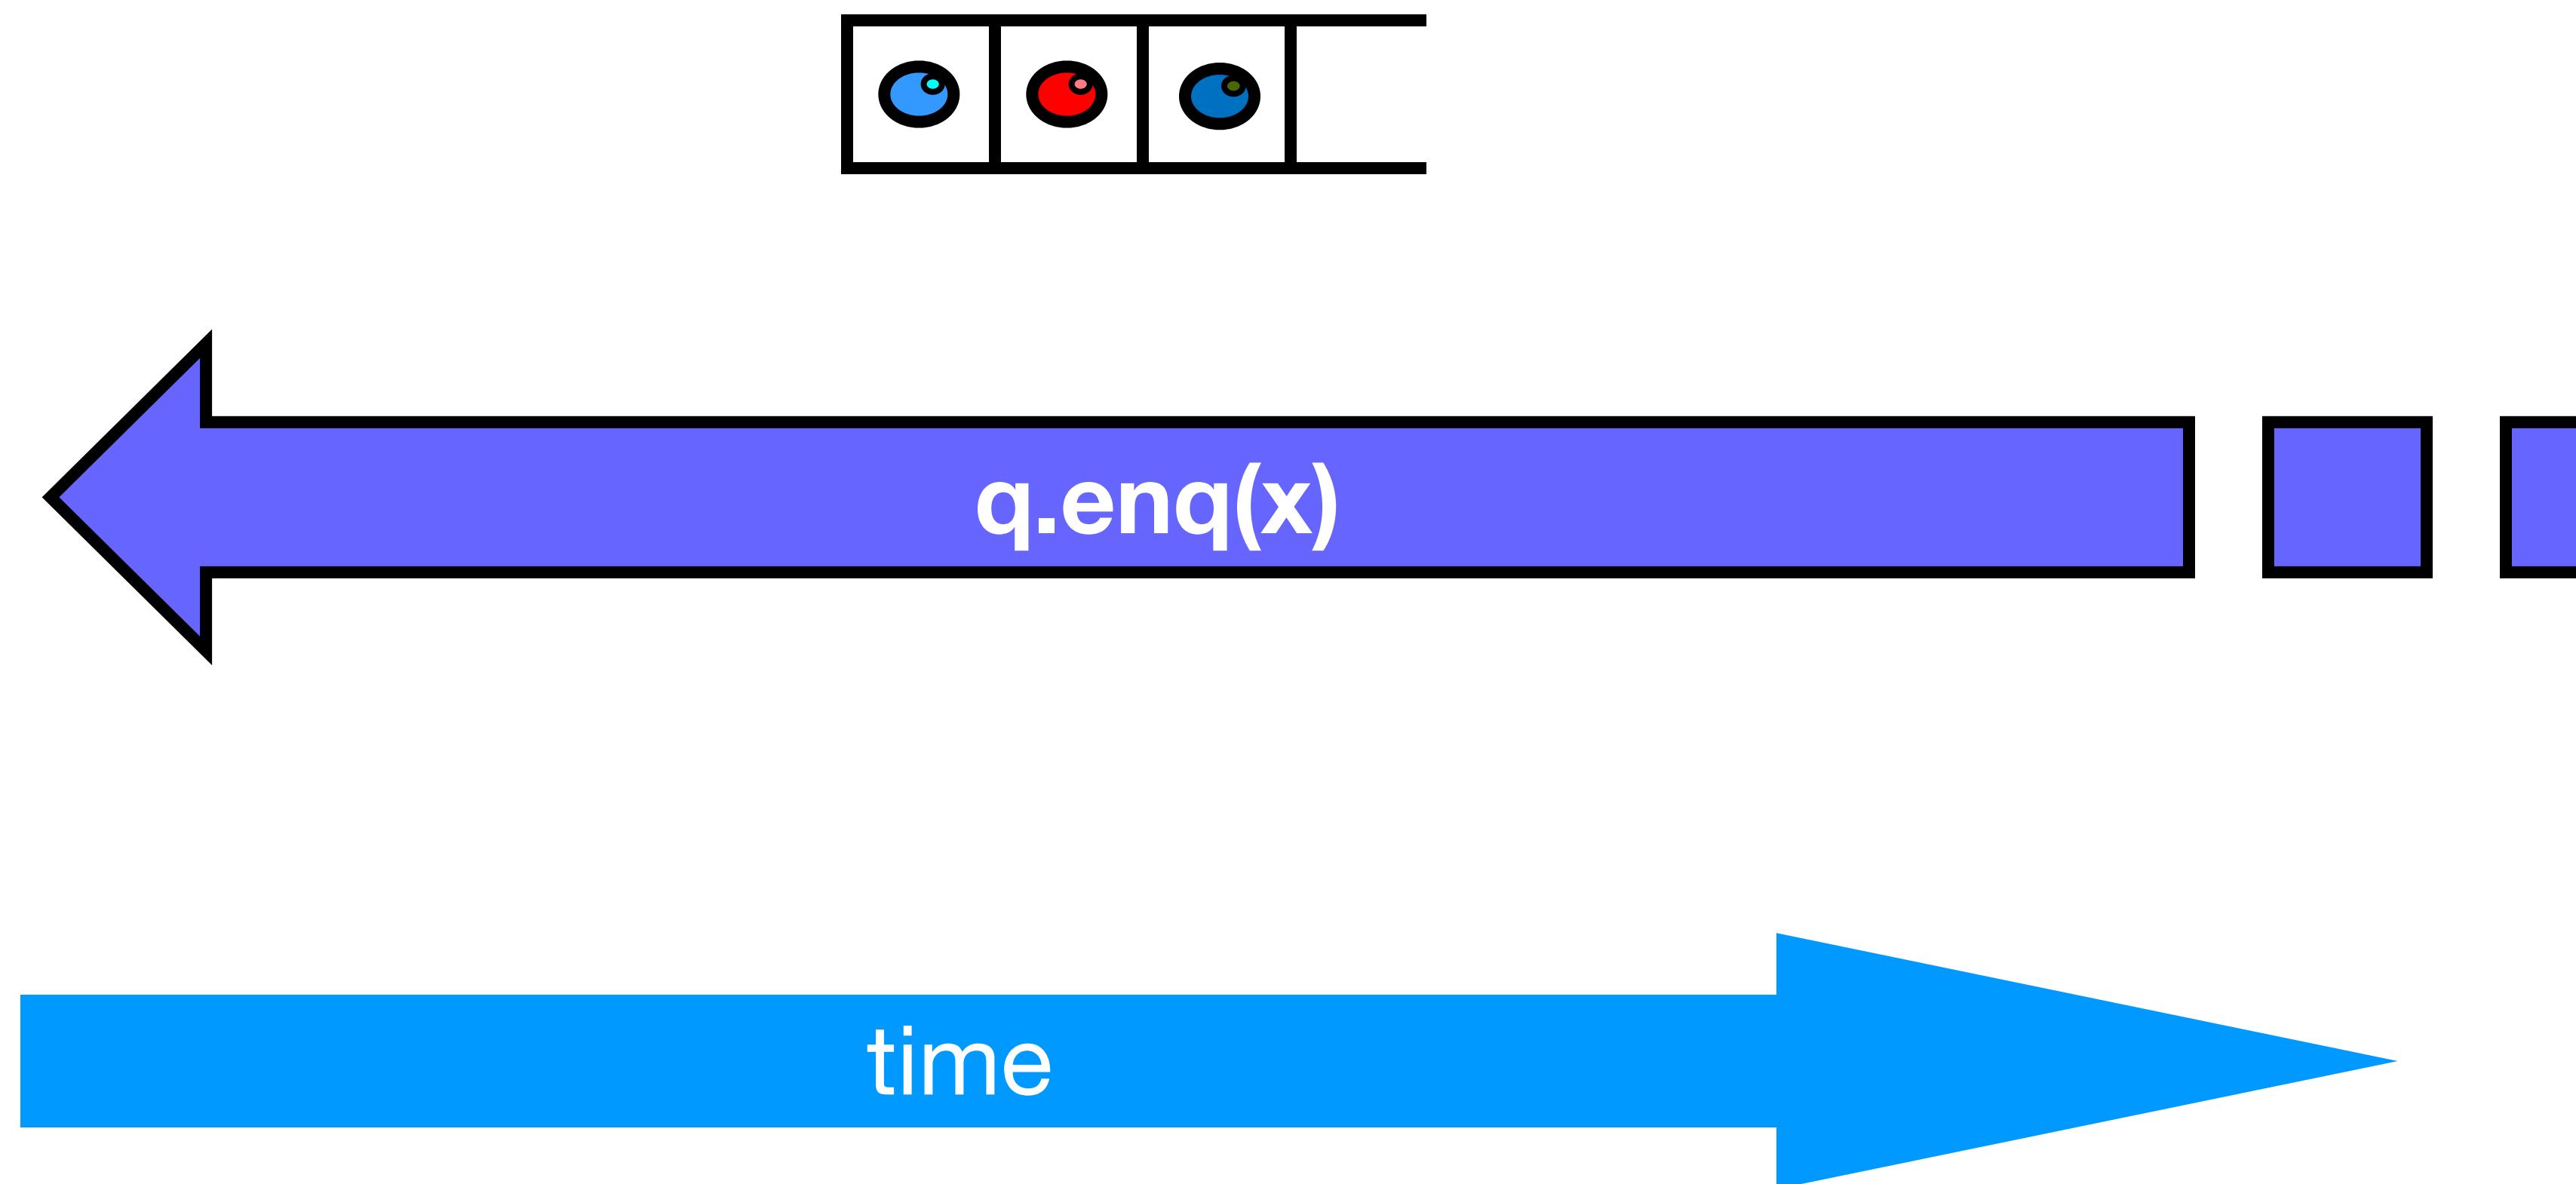
# Example



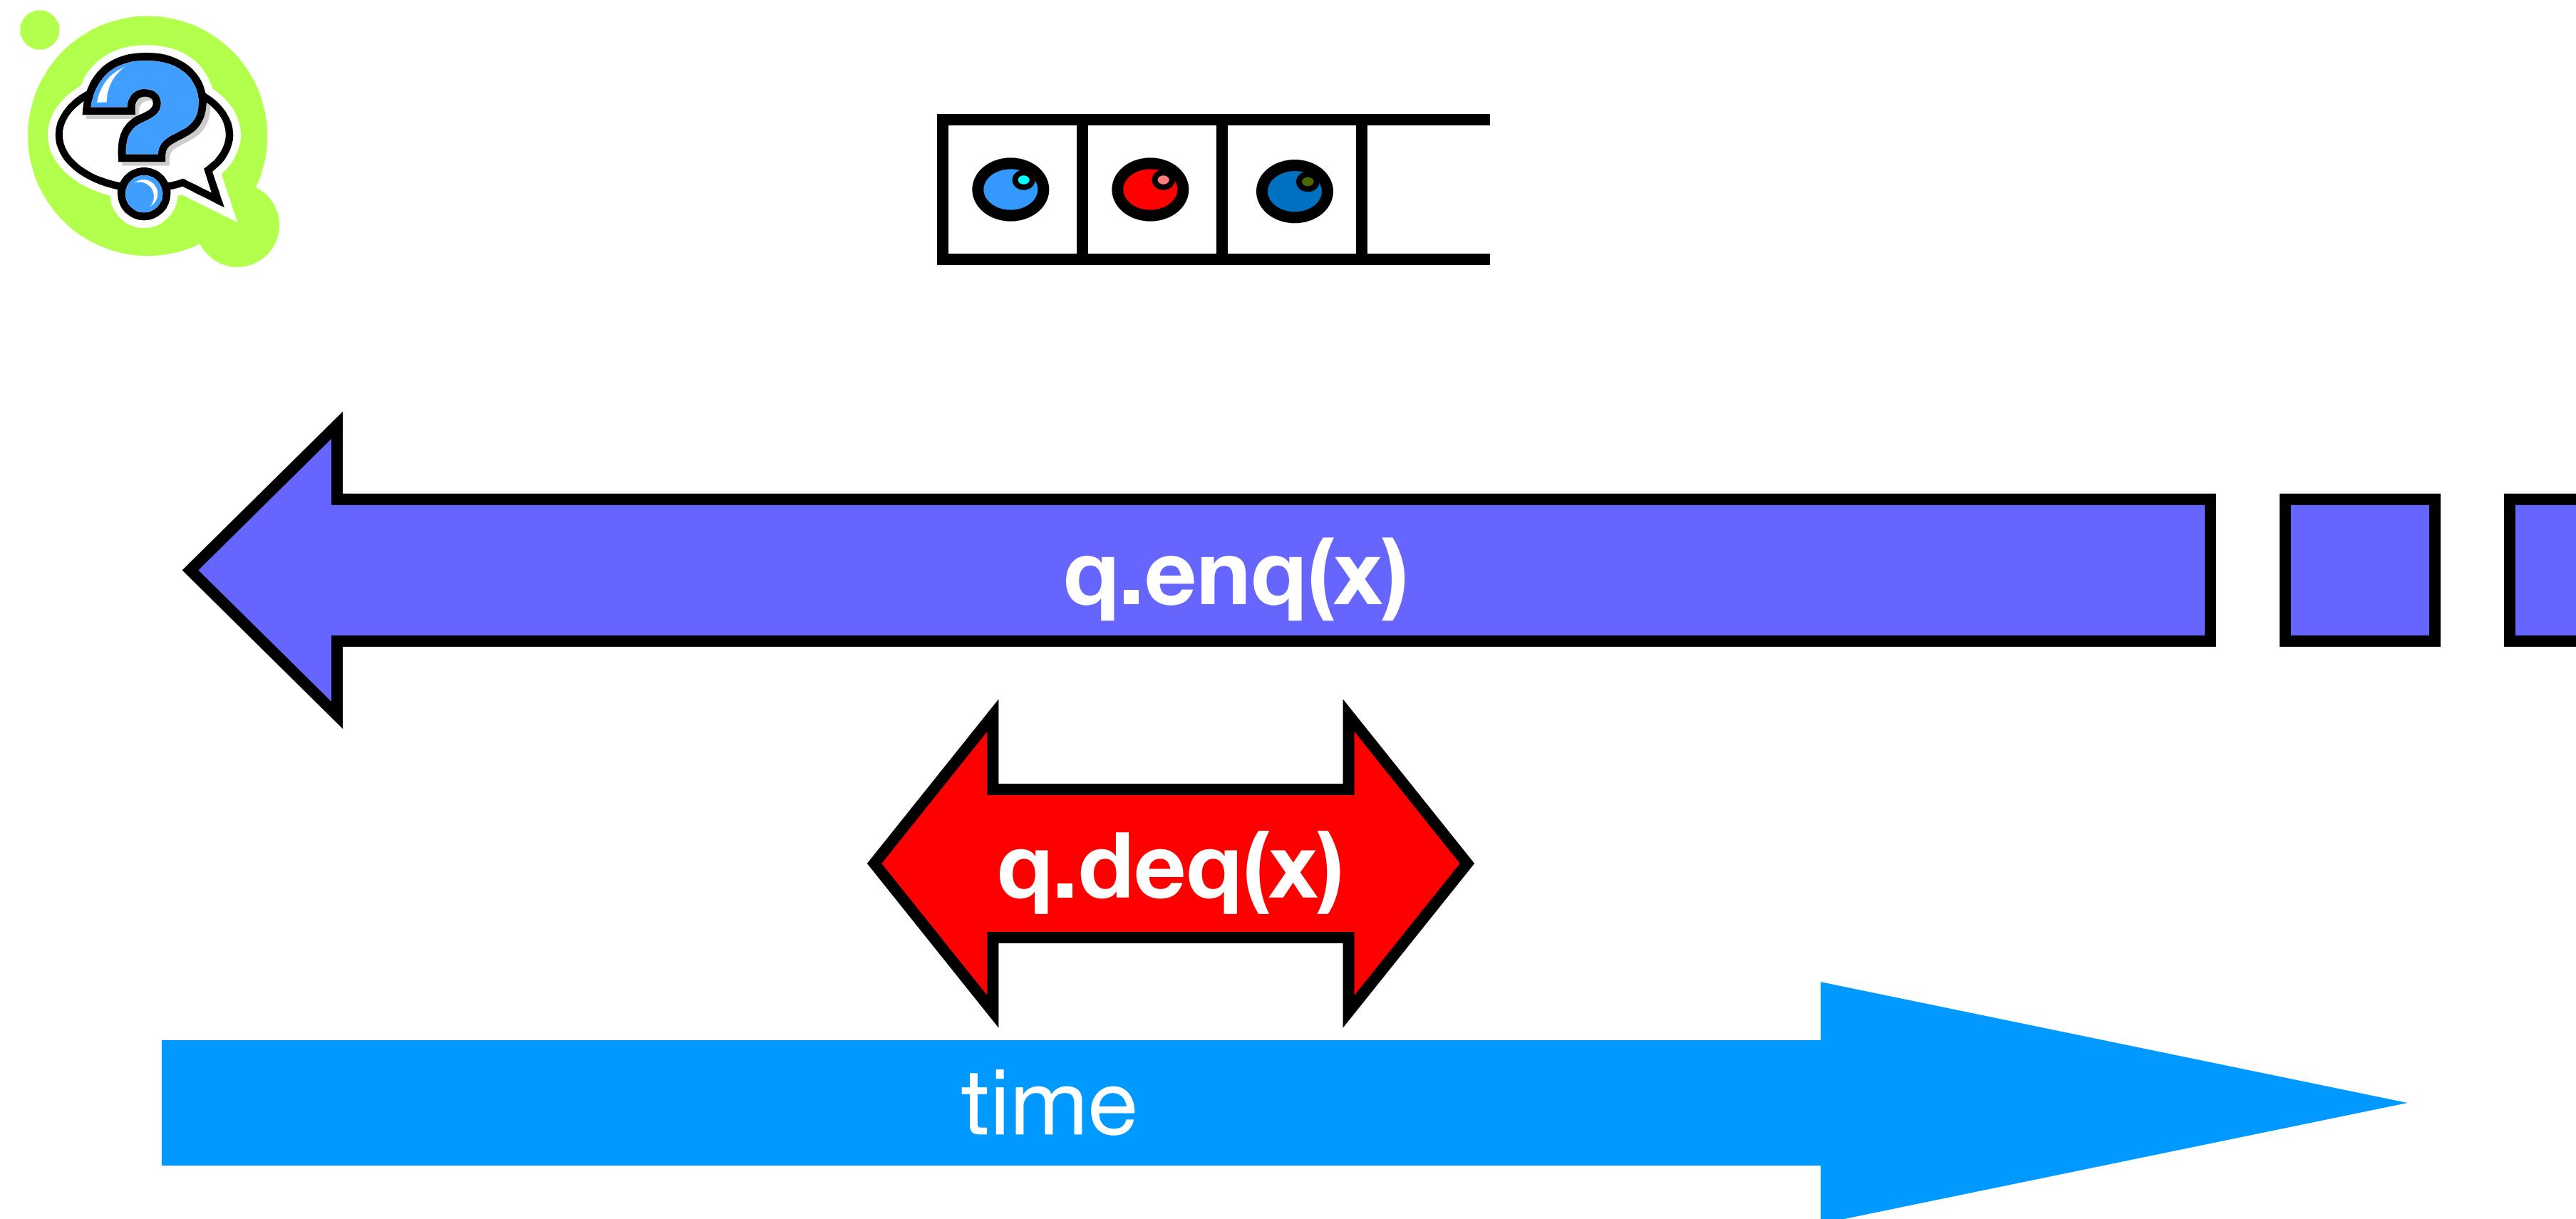
# Example



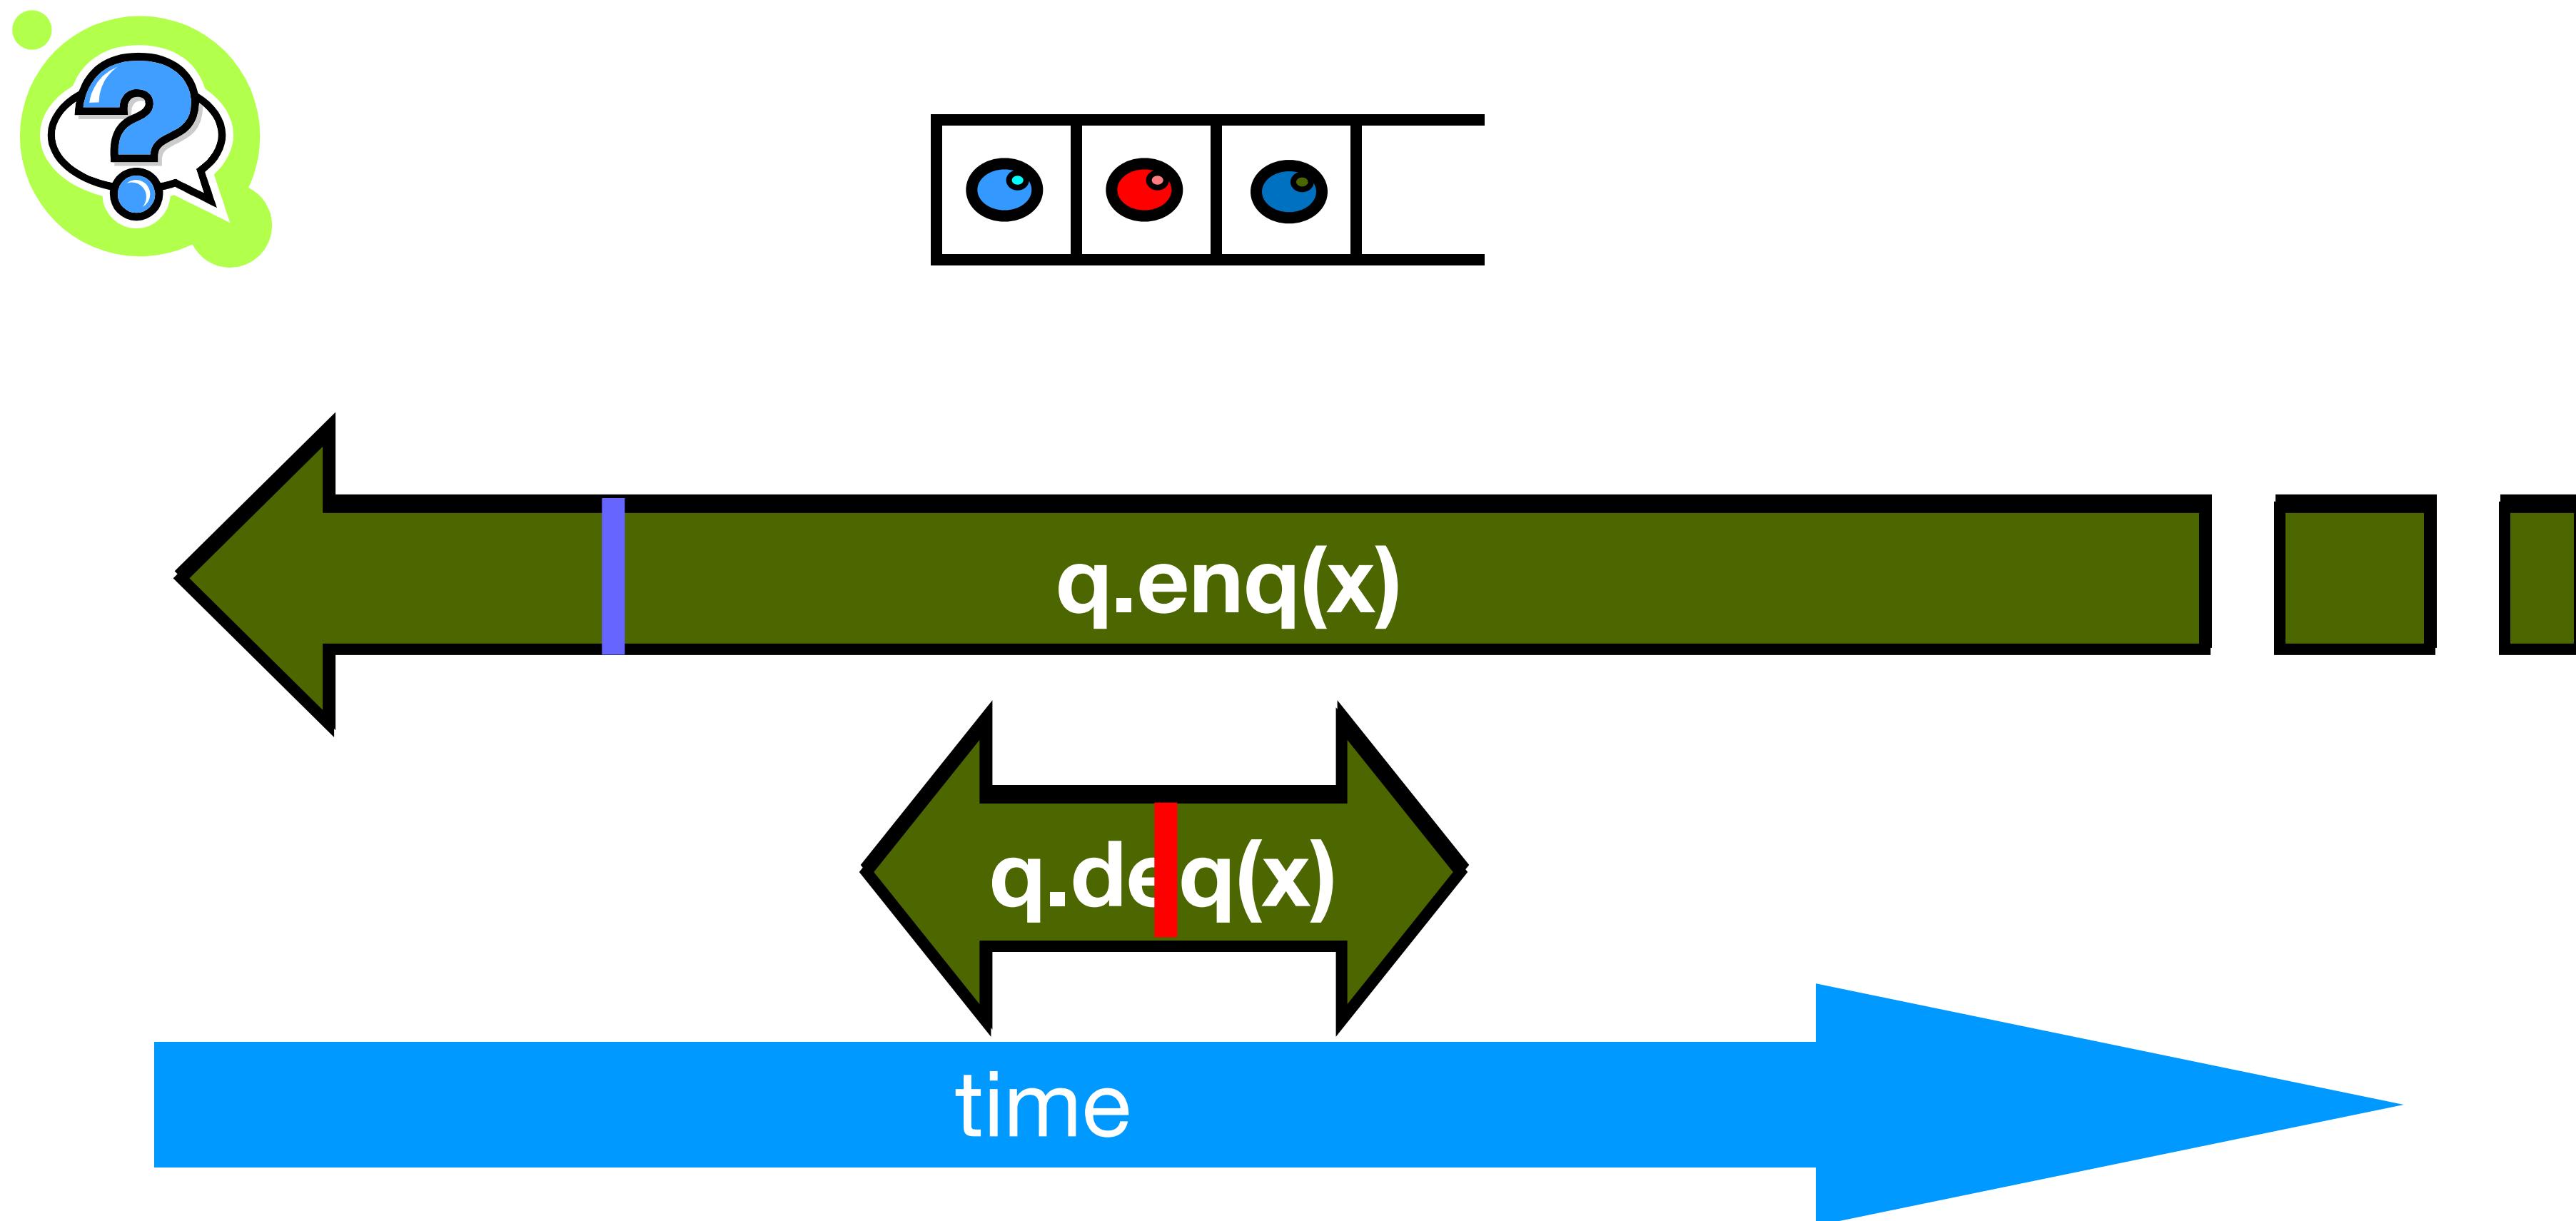
# Example



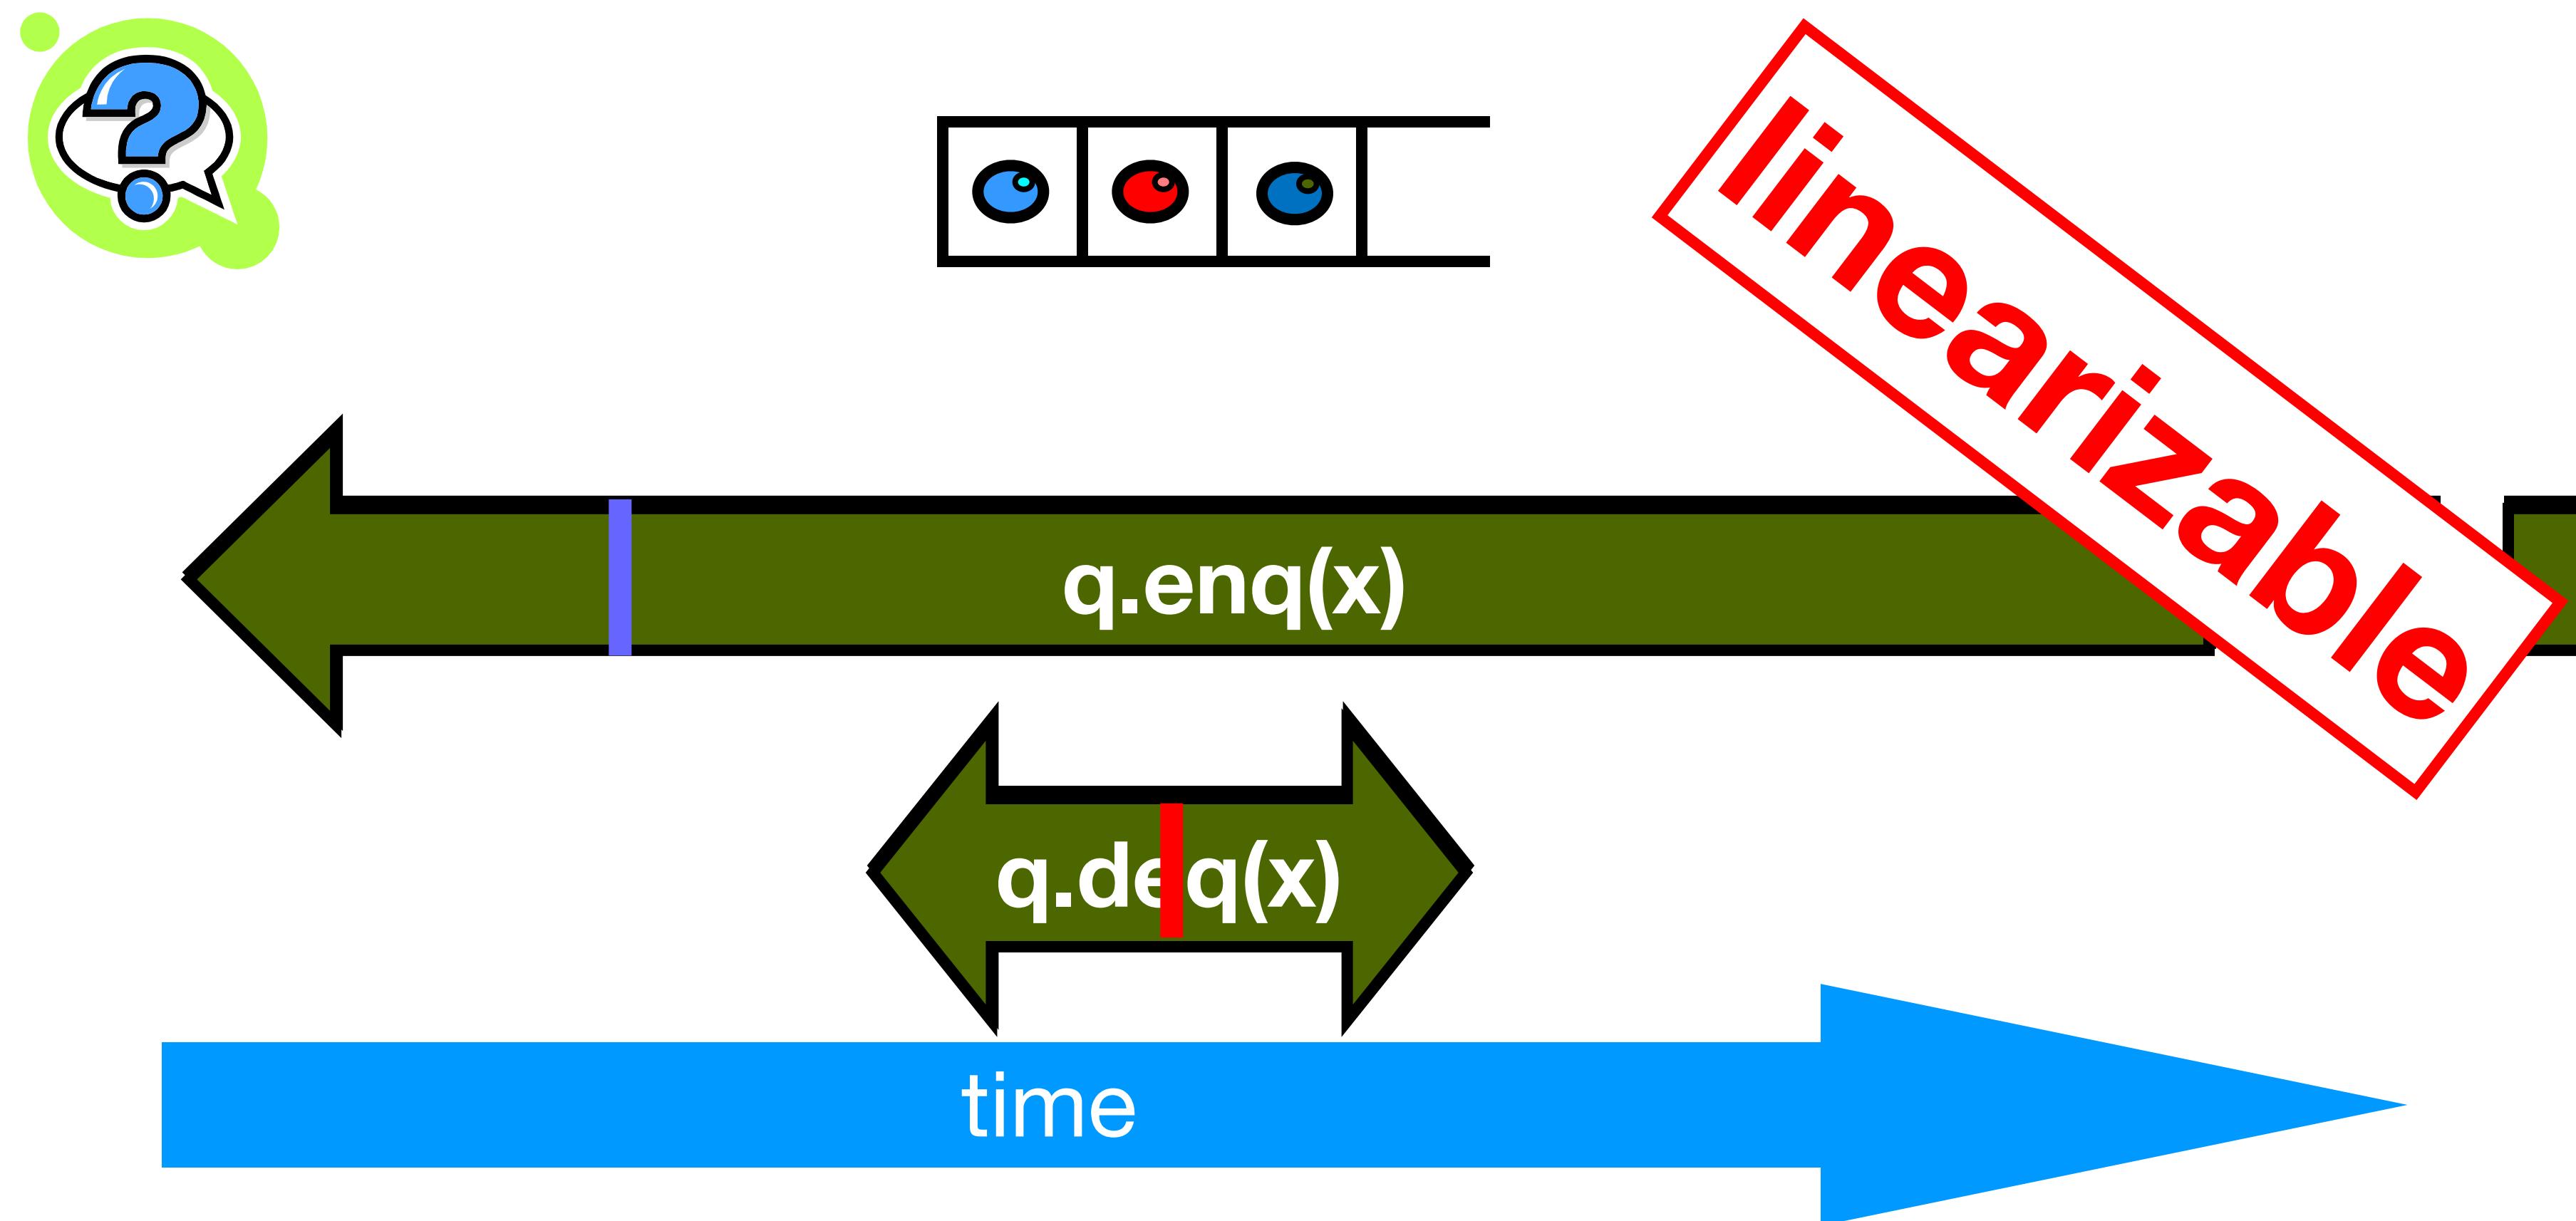
# Example



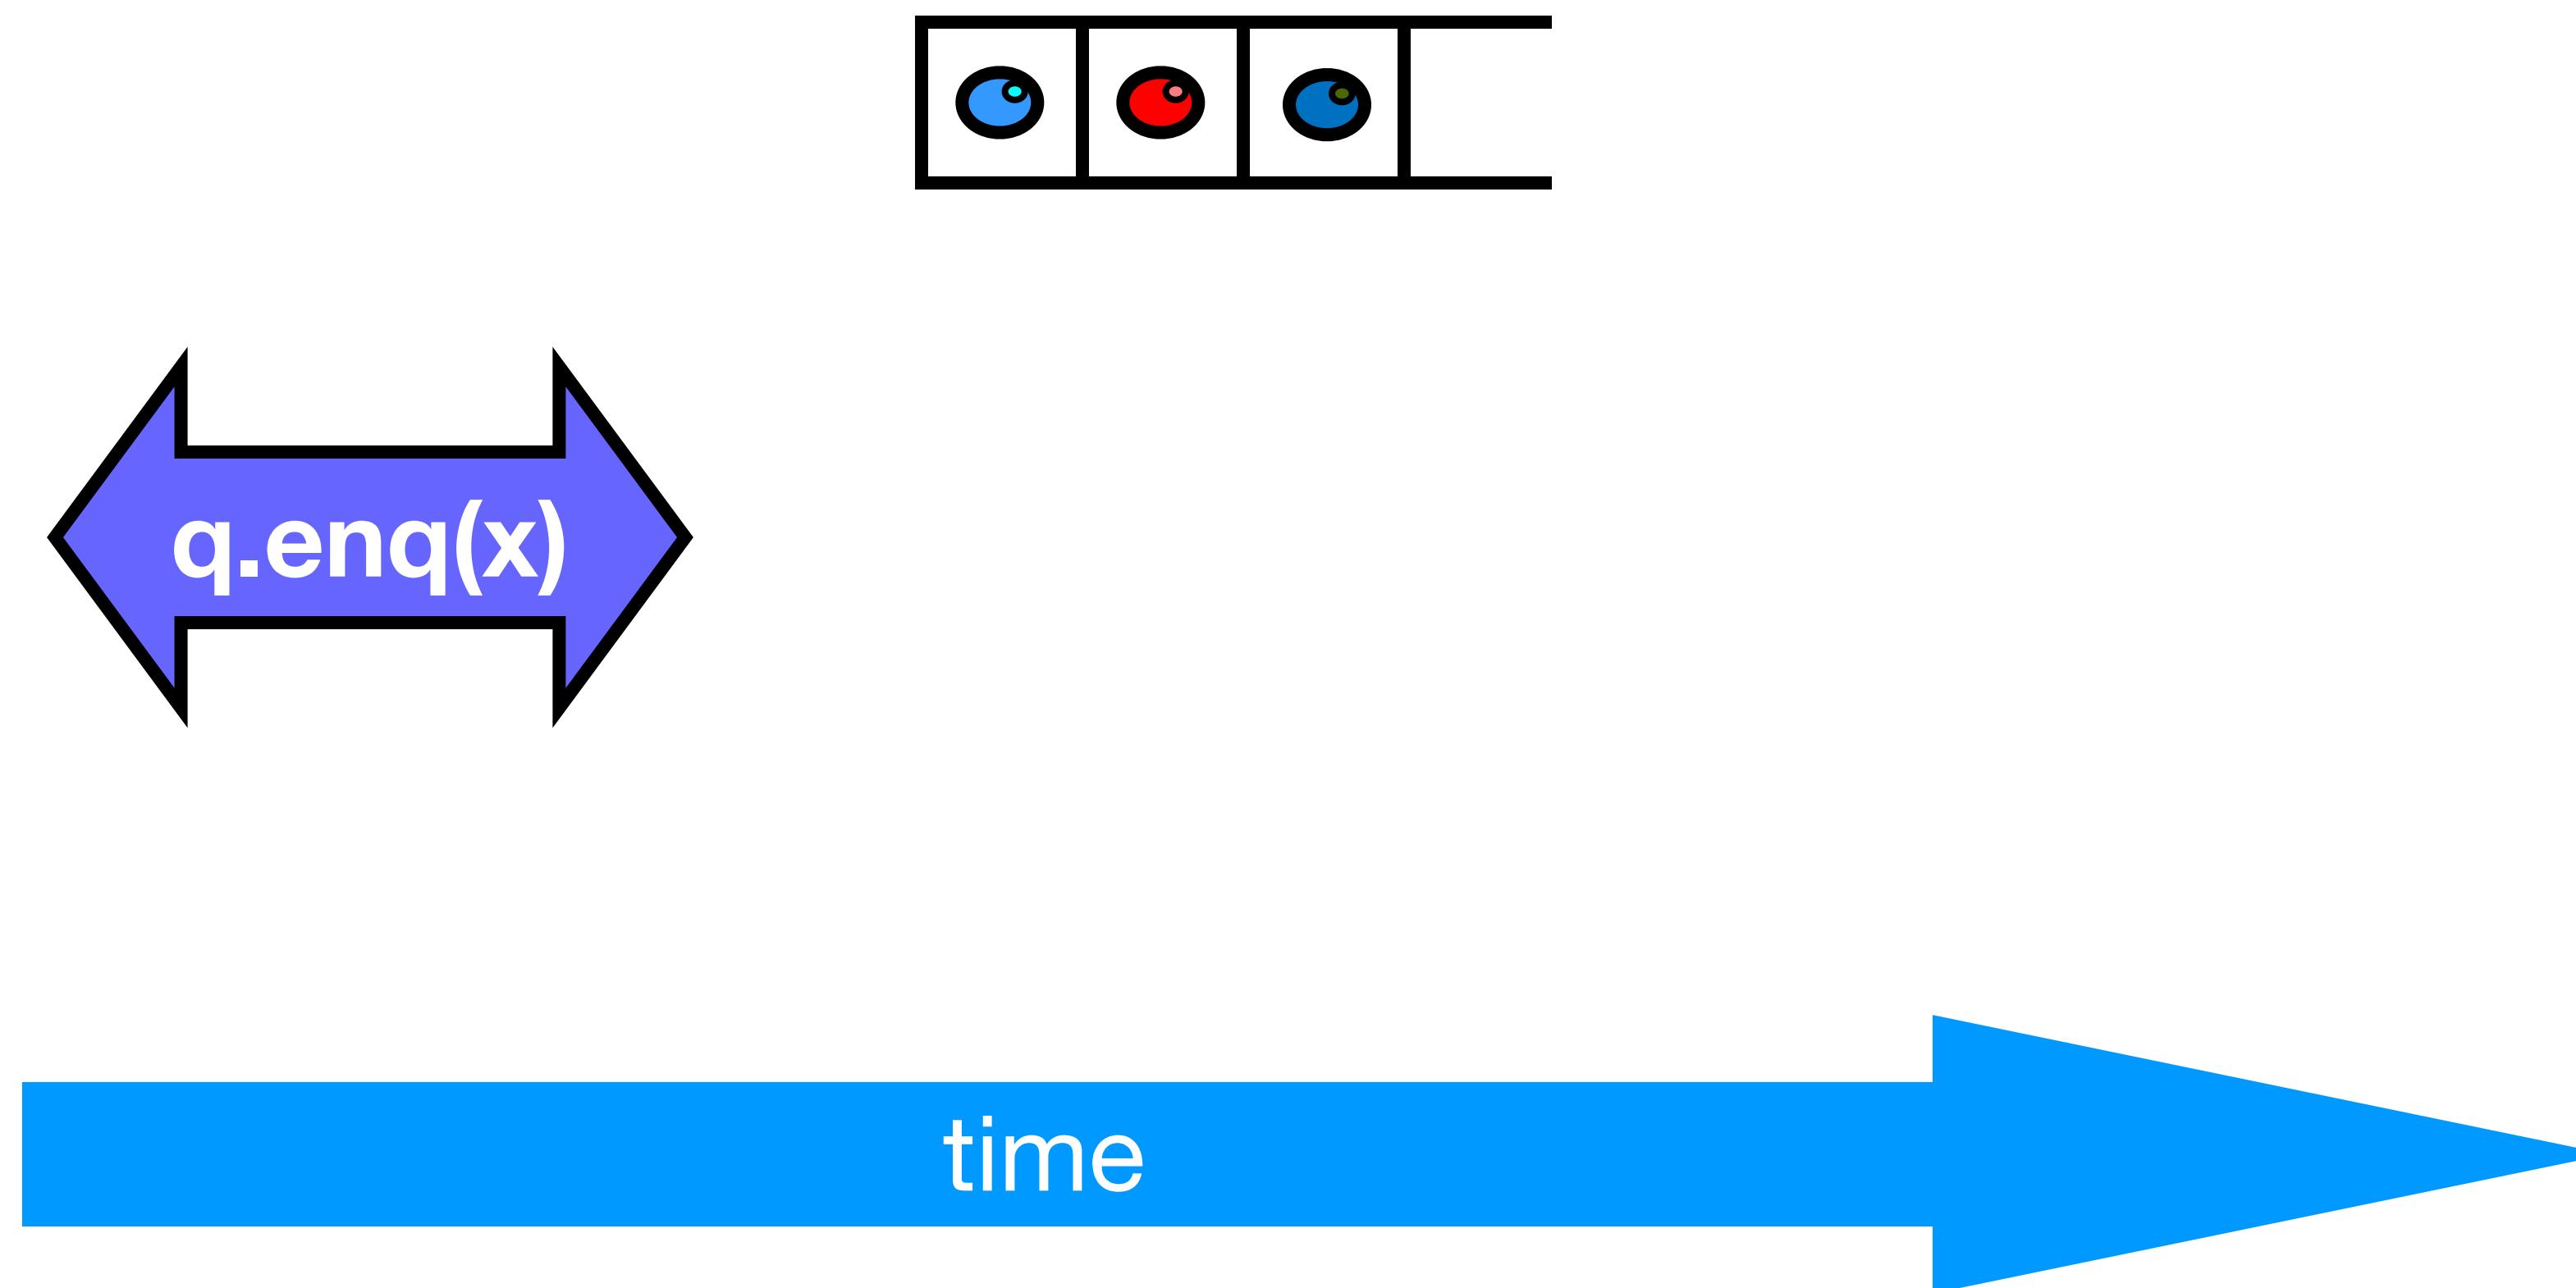
# Example



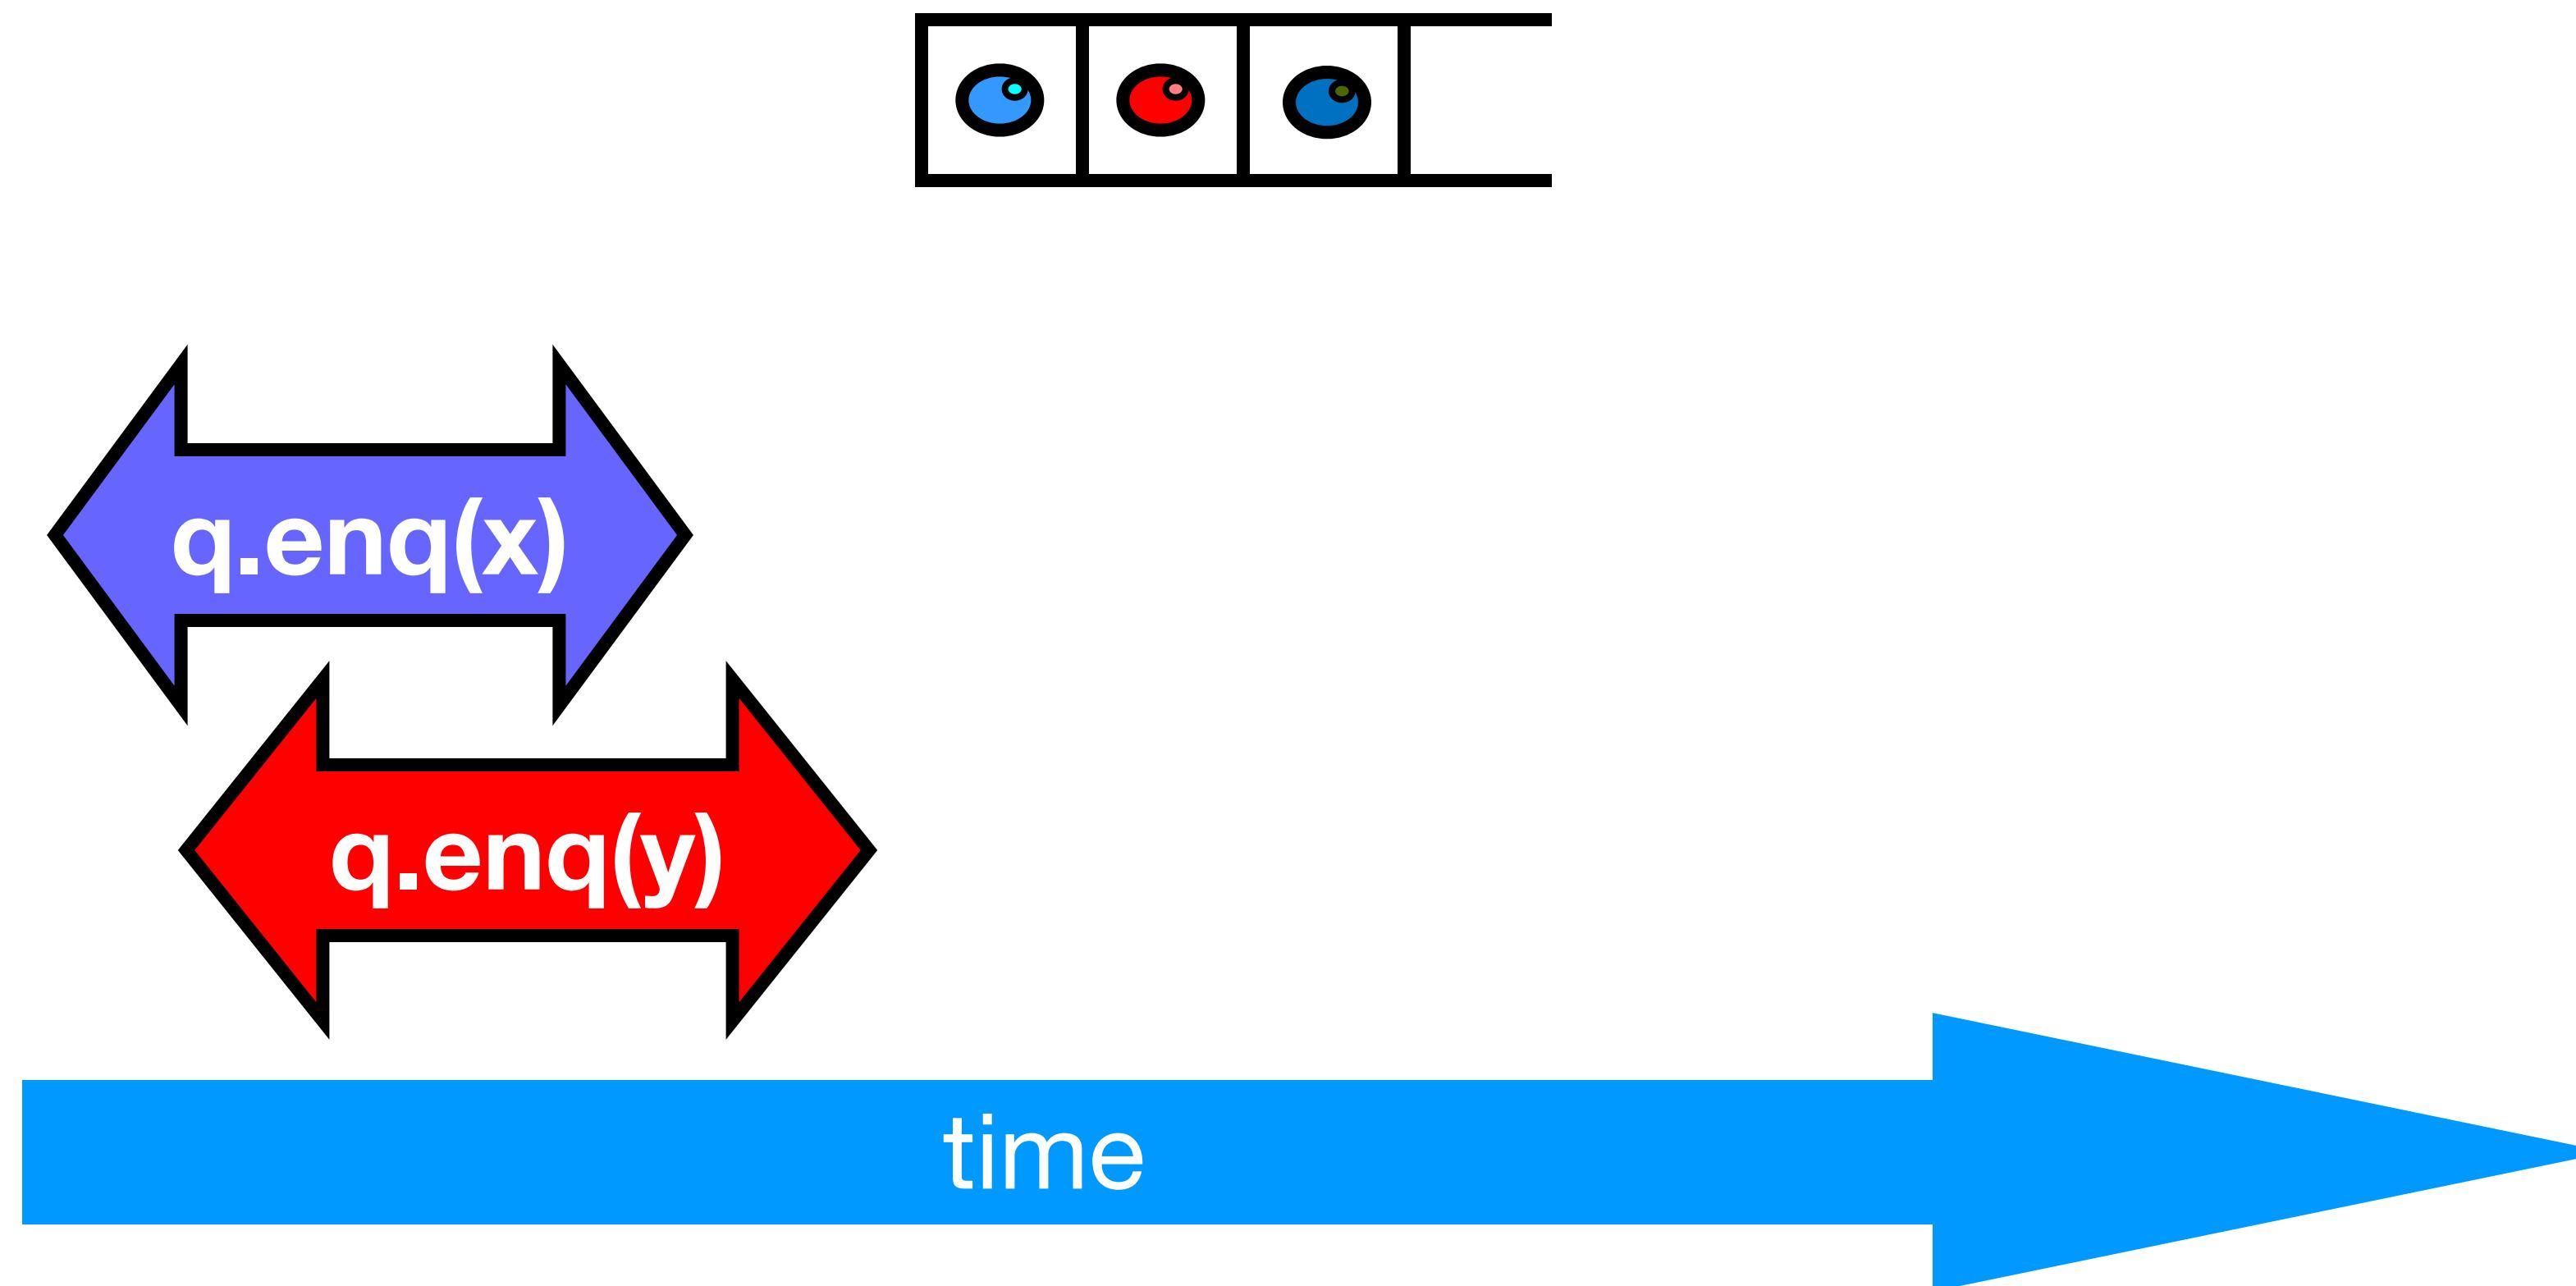
# Example



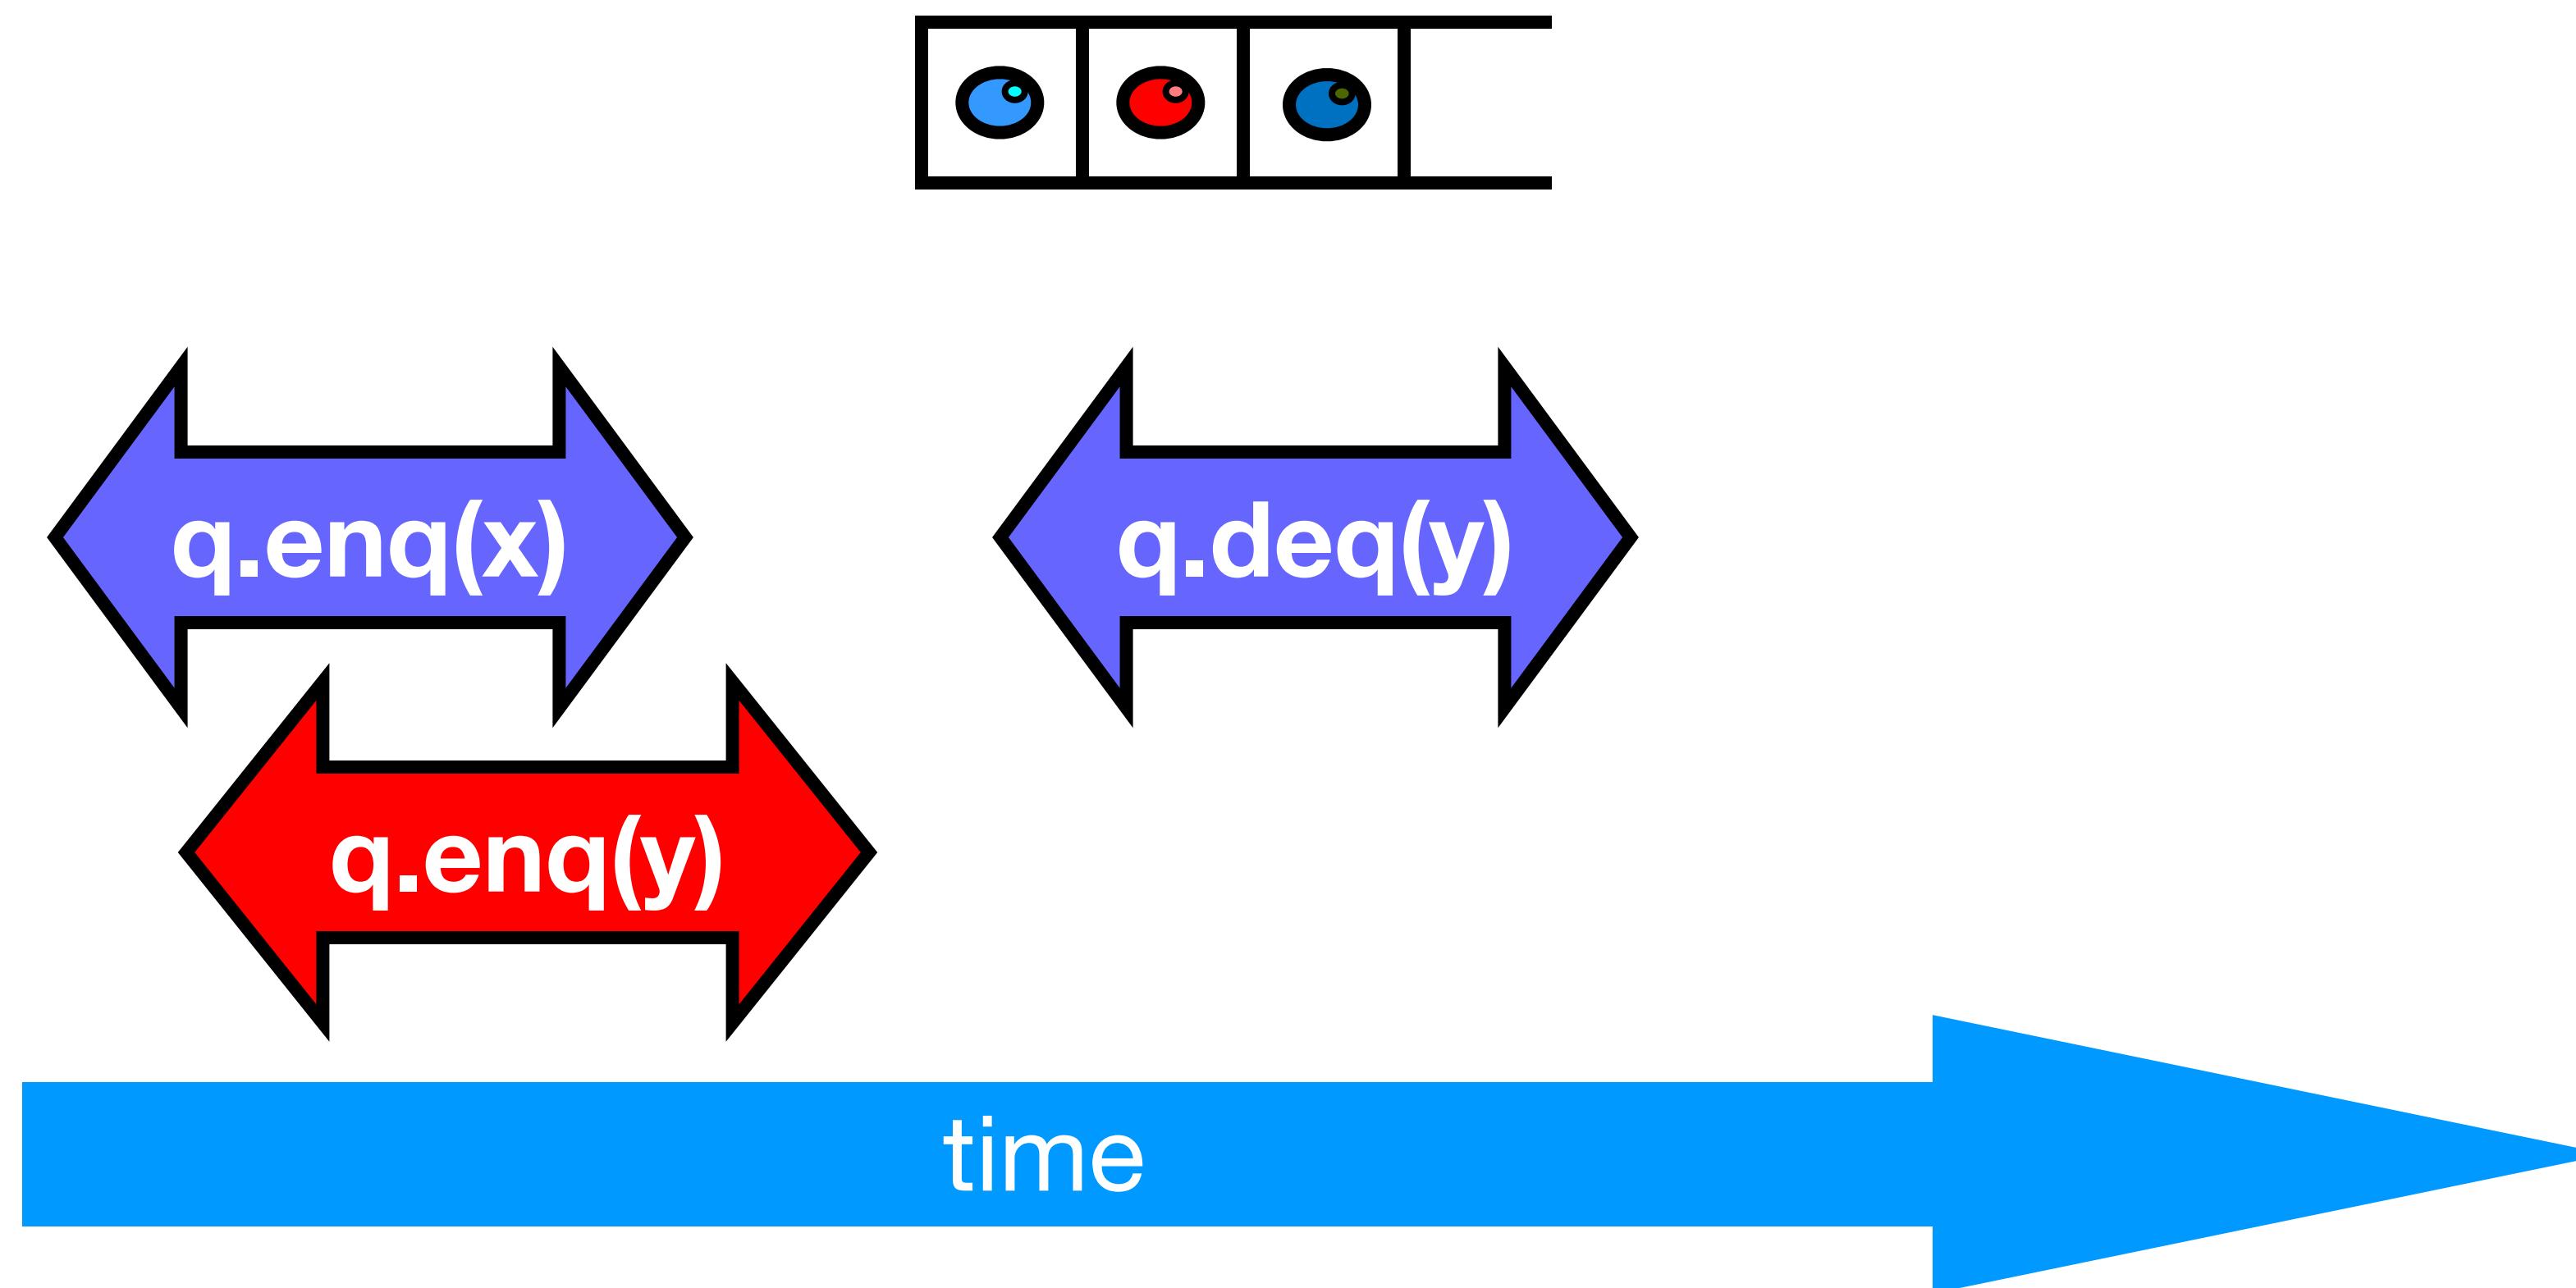
# Example



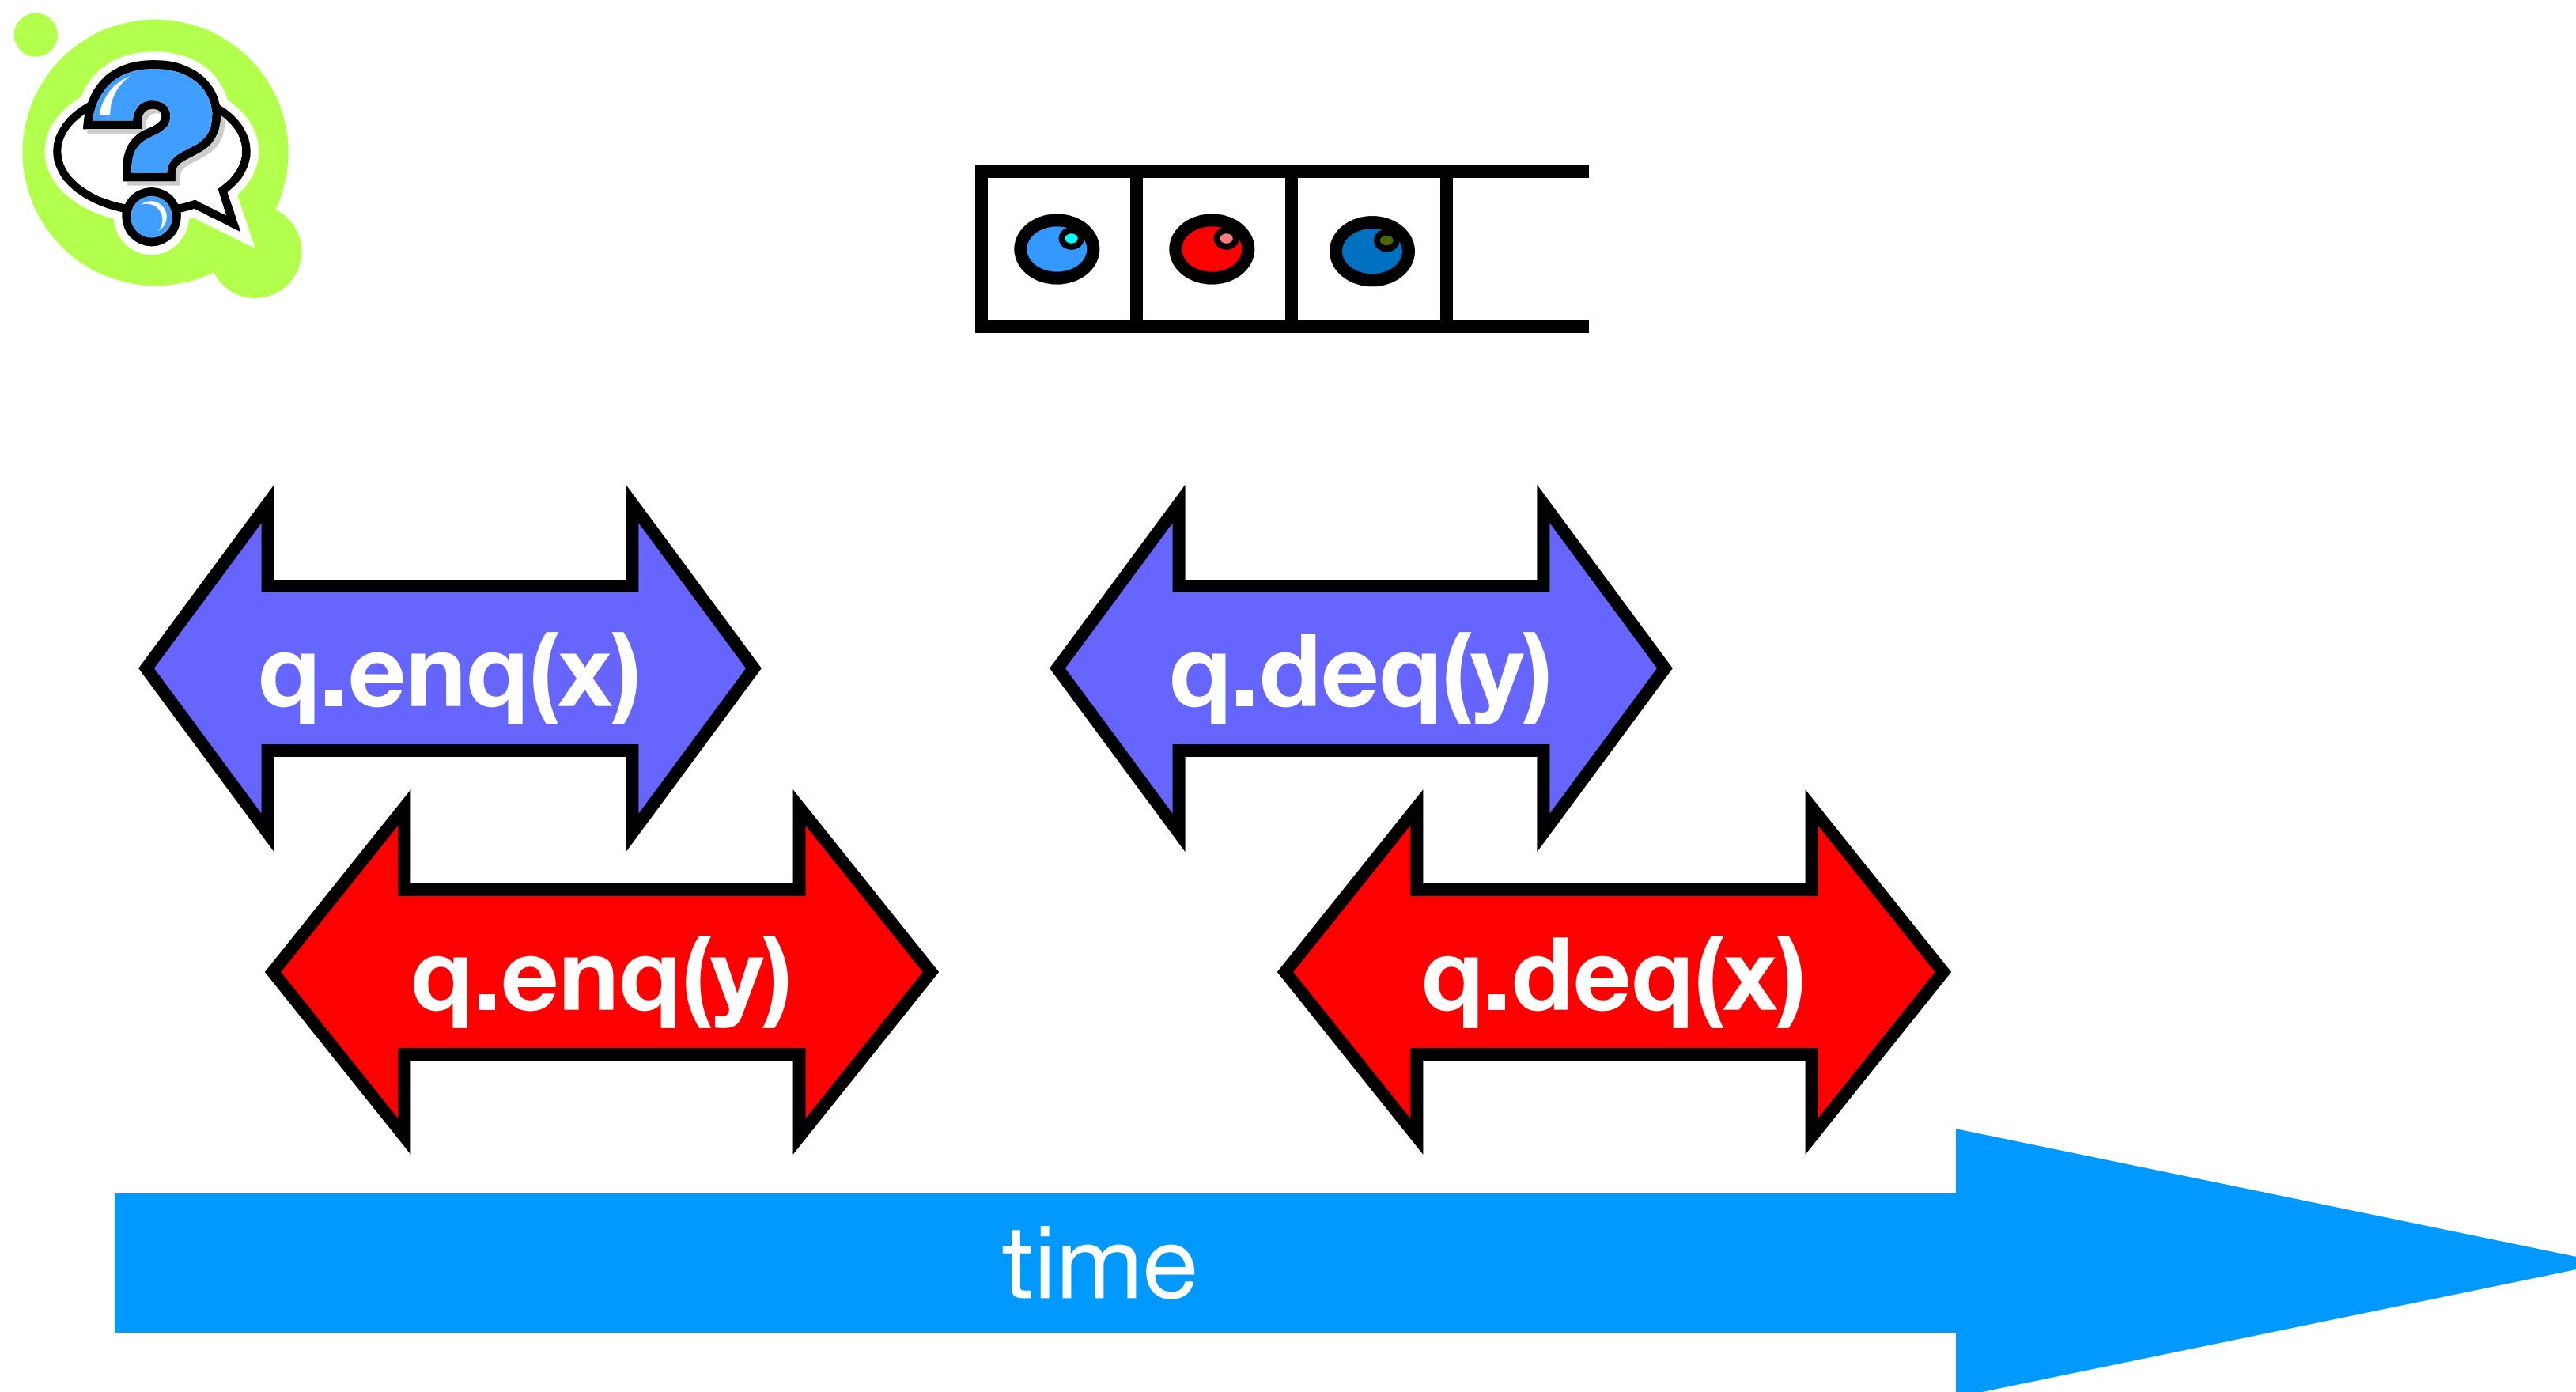
# Example



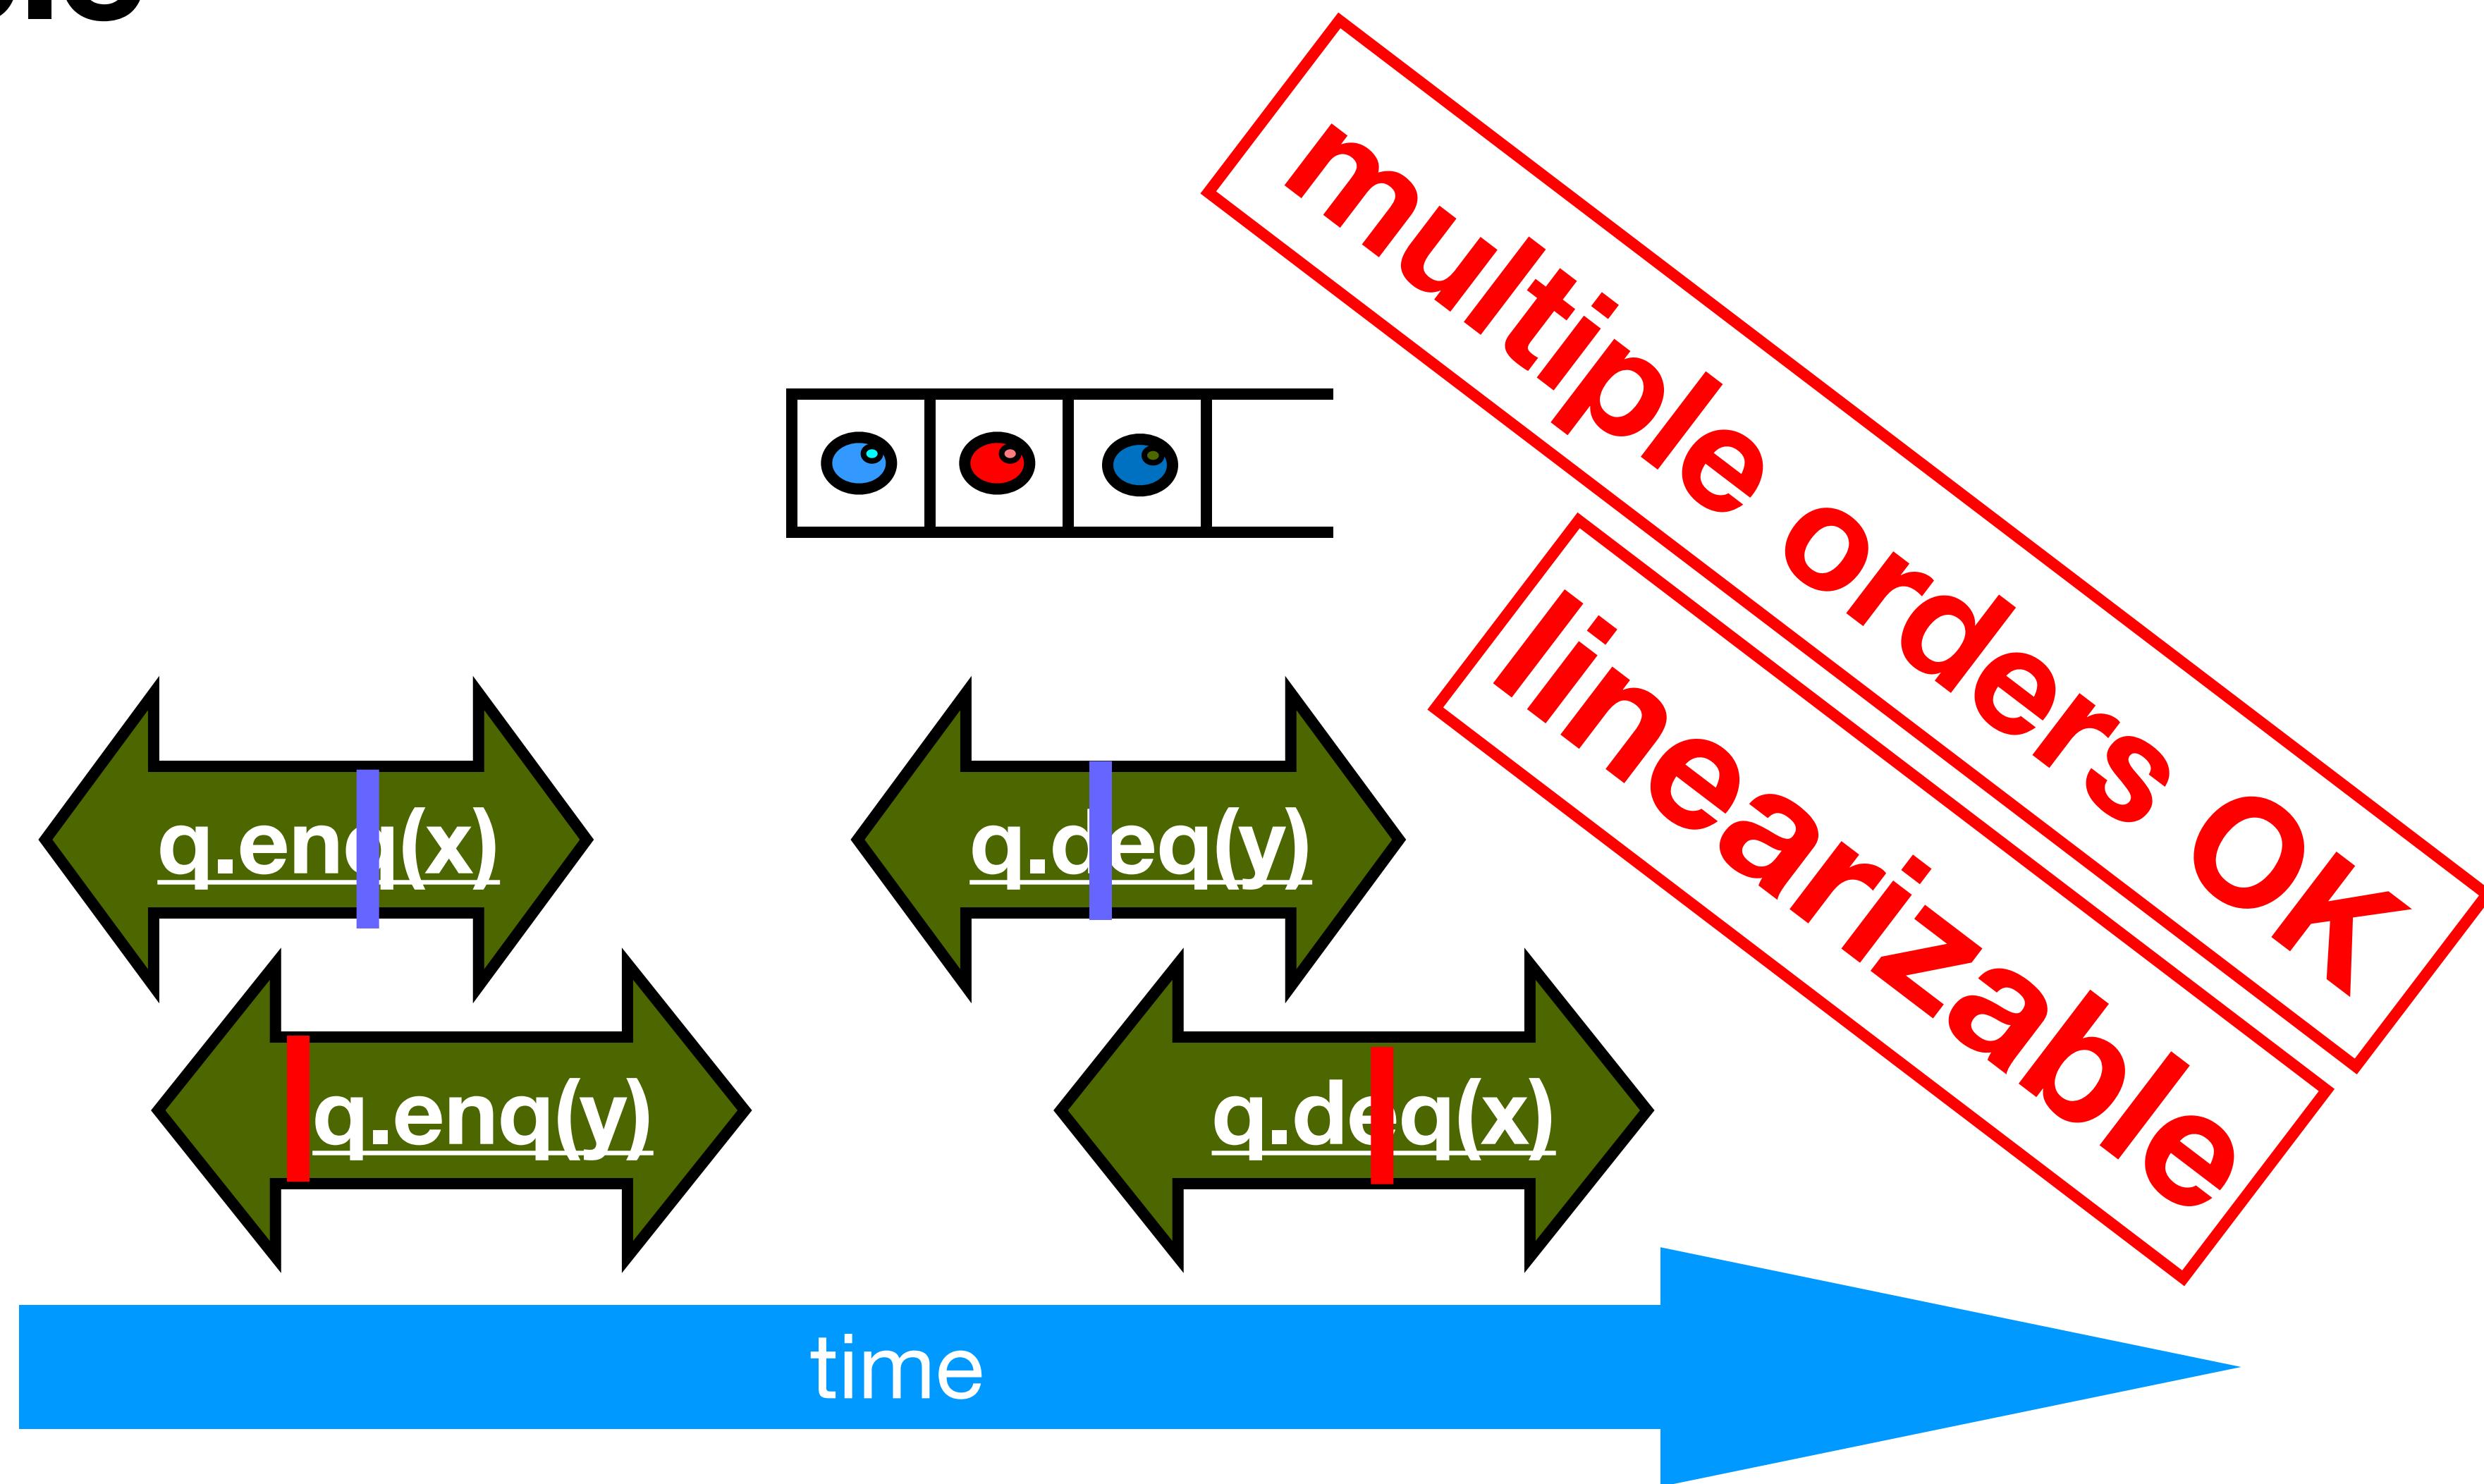
# Example



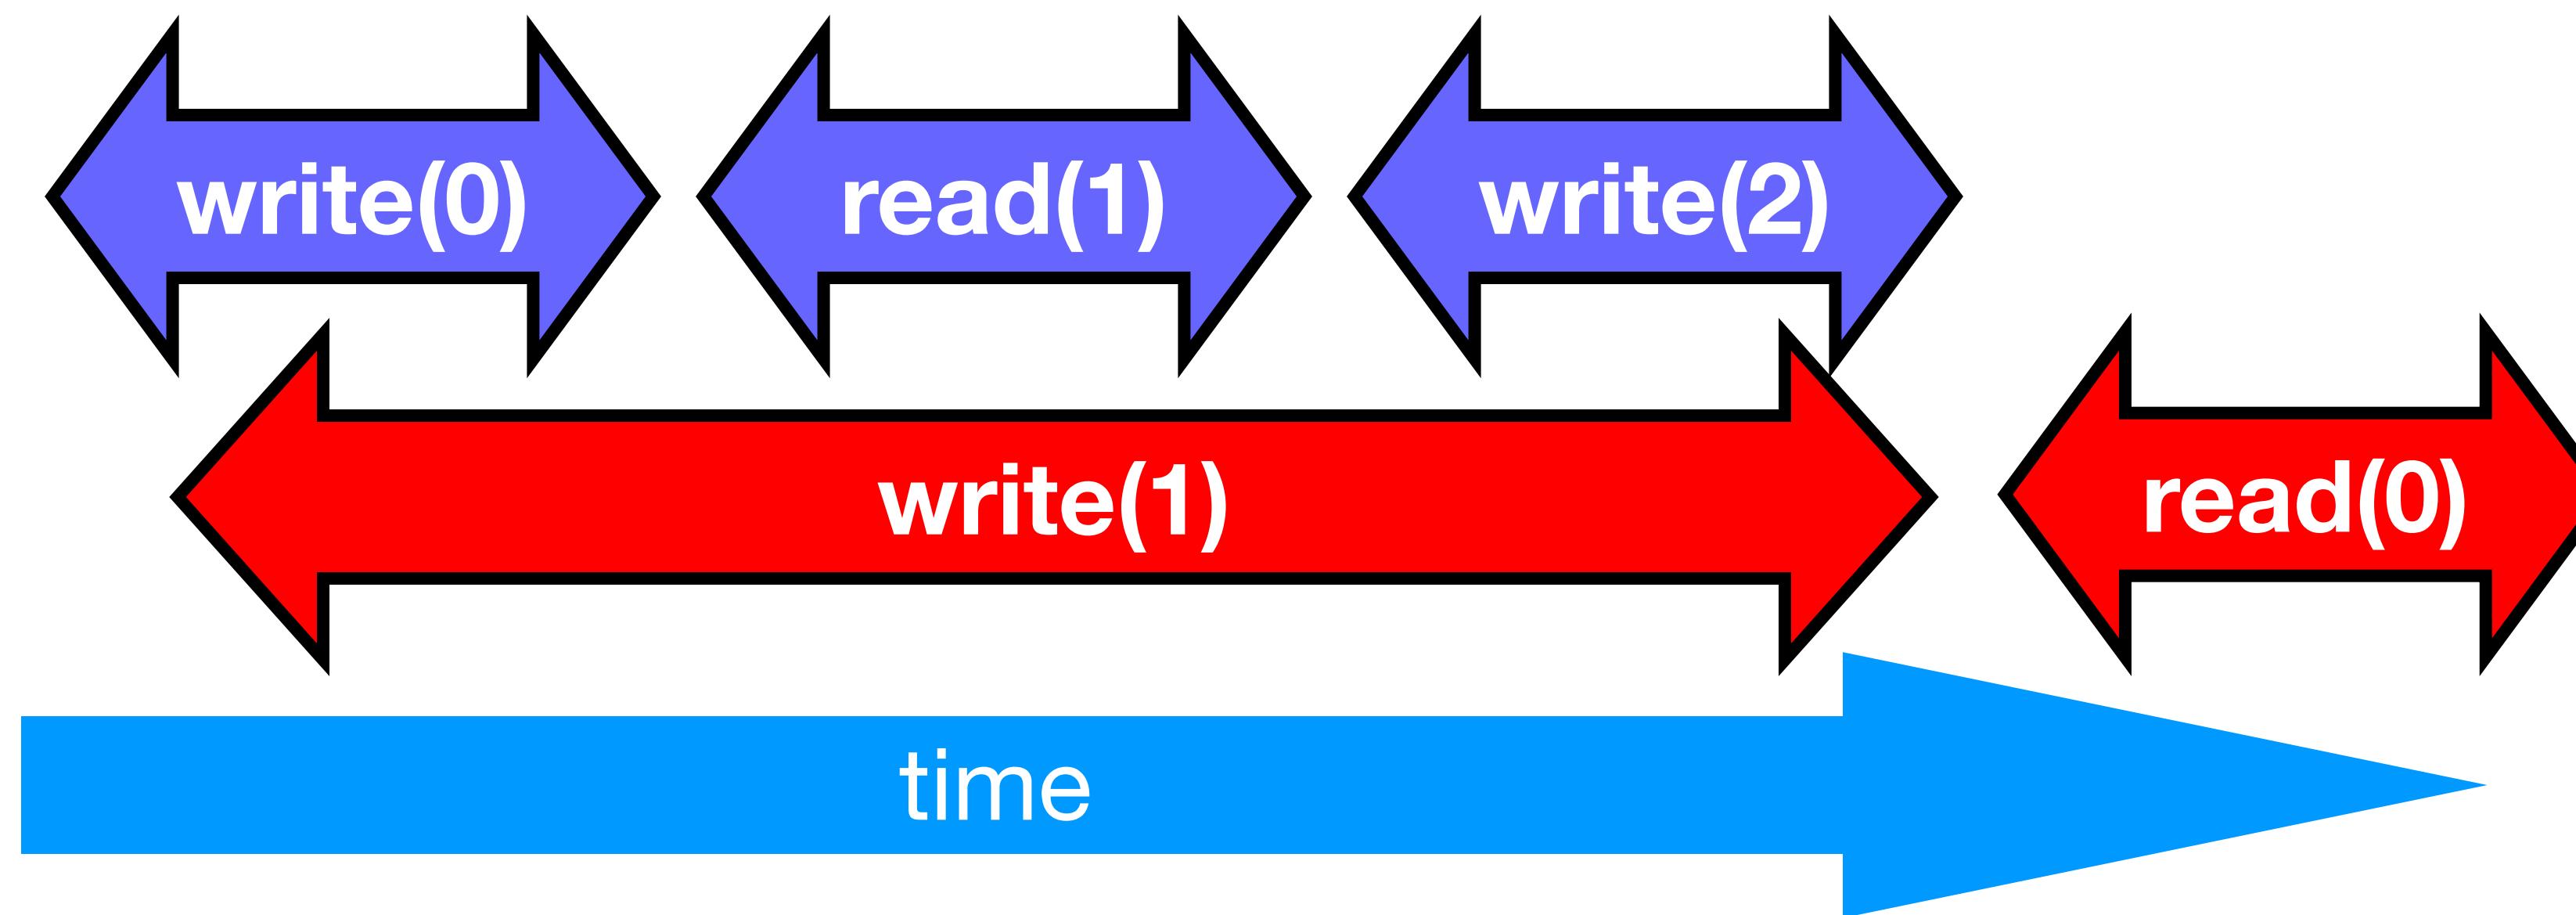
# Example



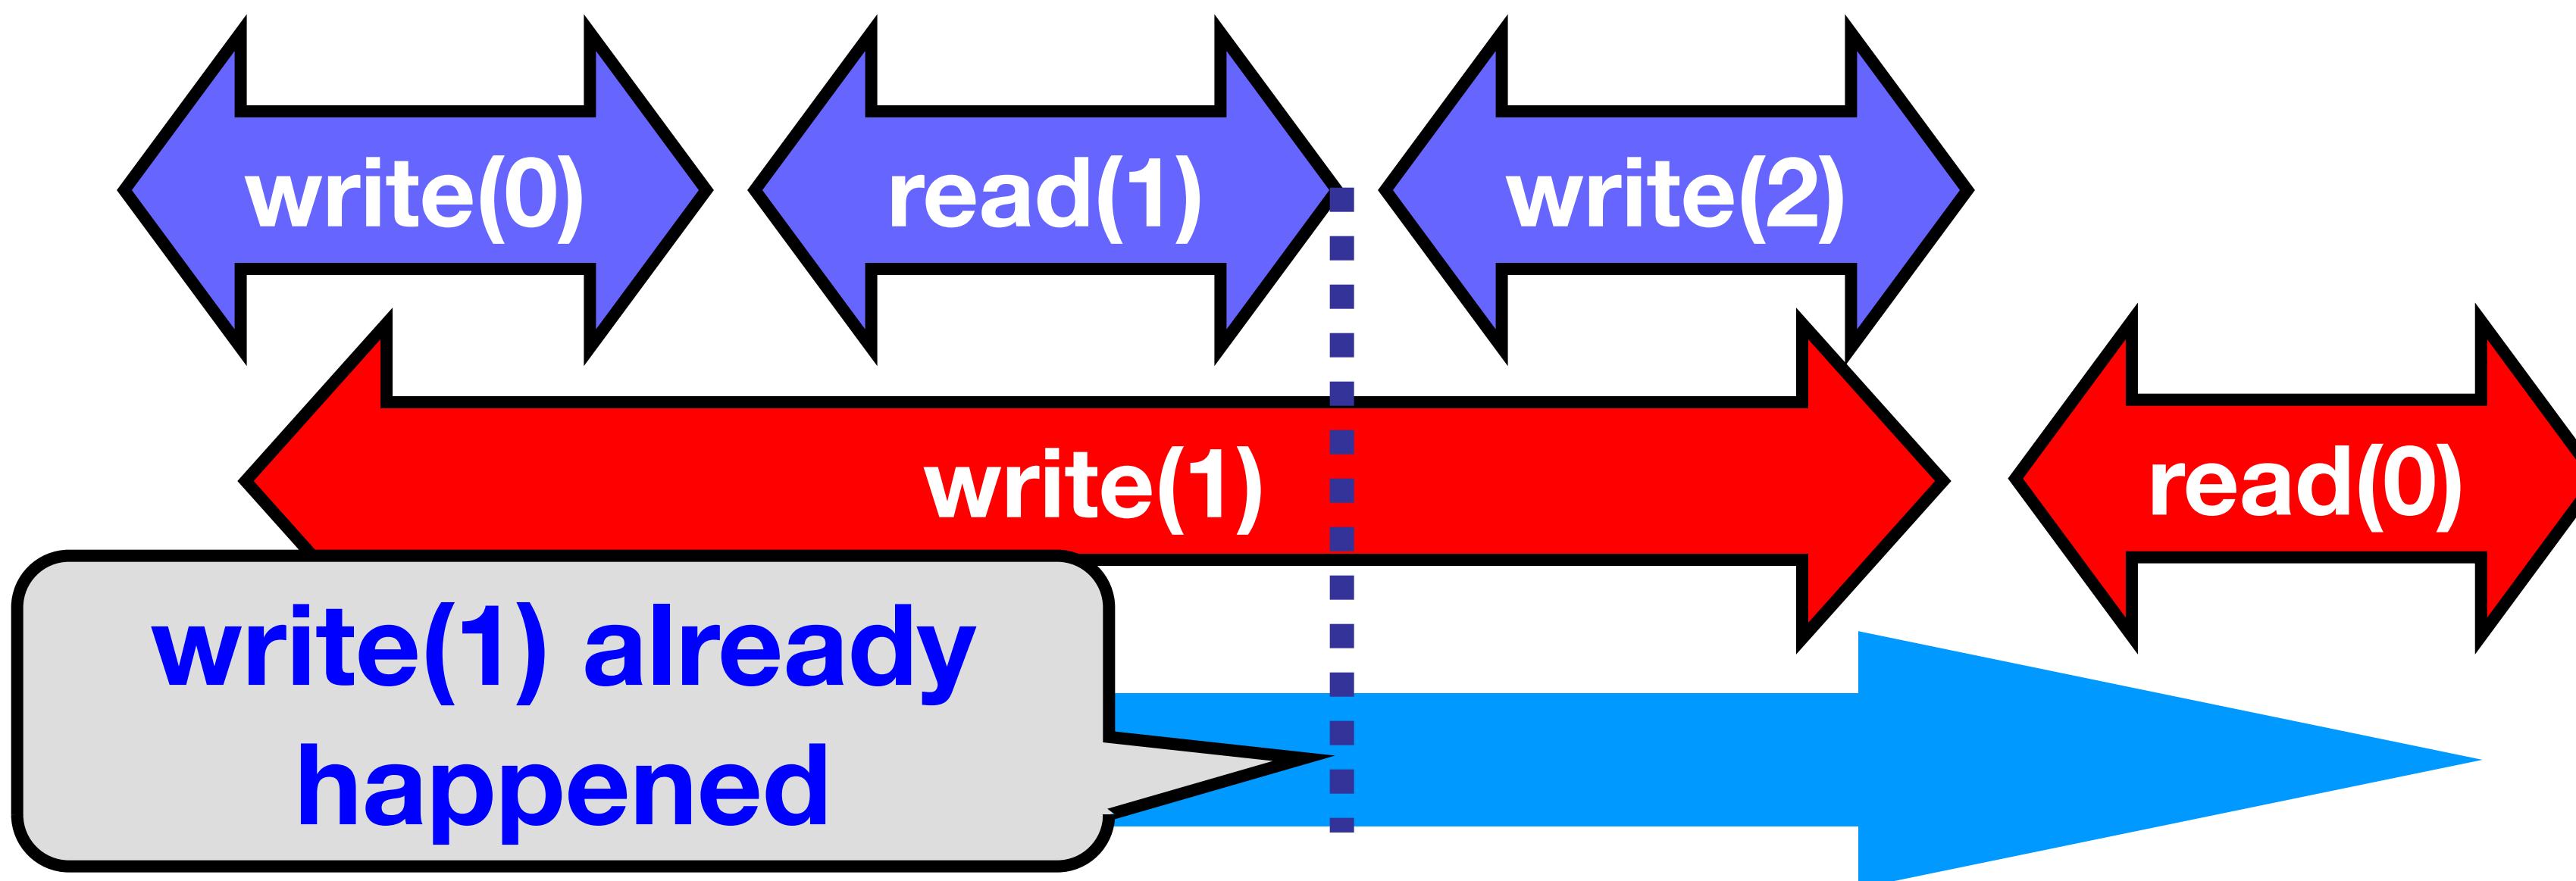
# Example



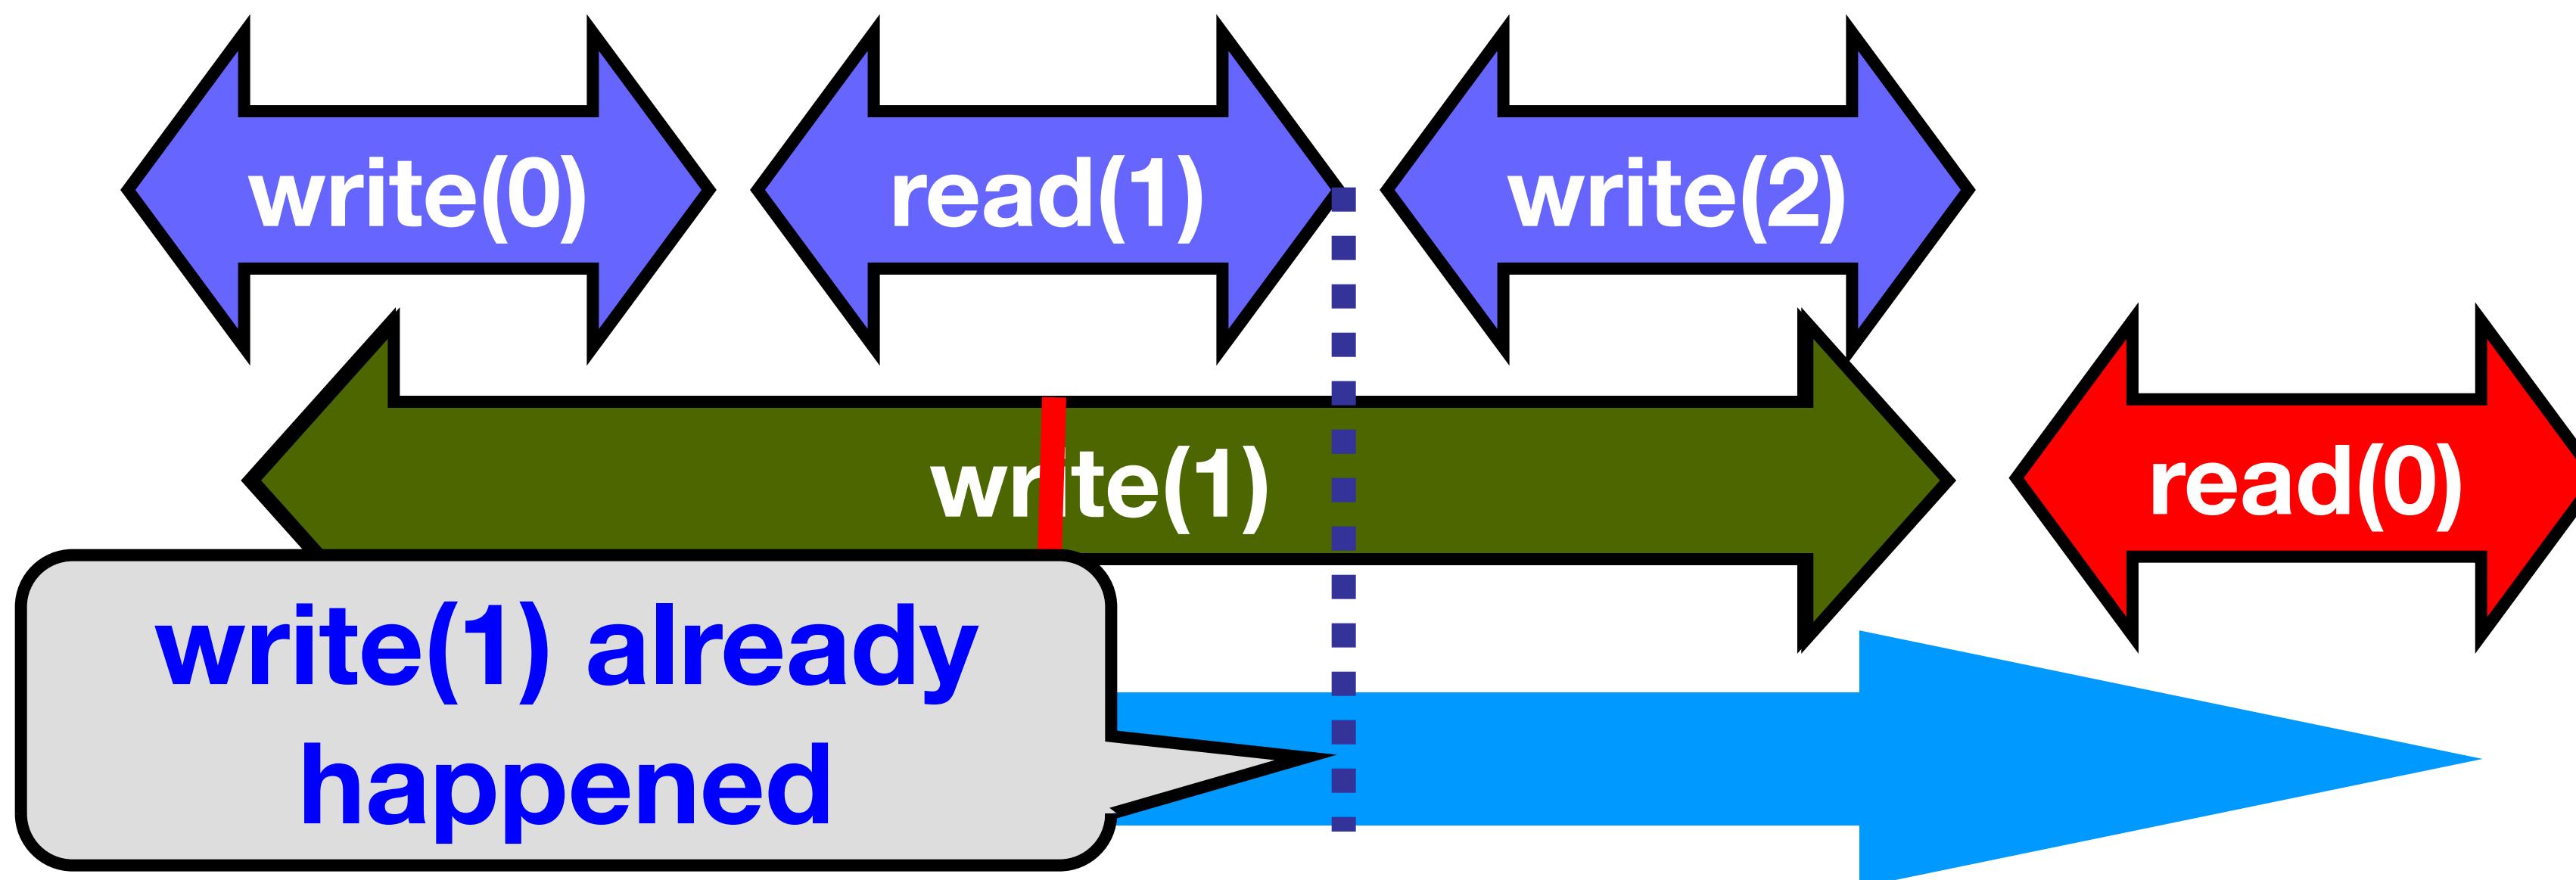
# Read/Write Register Example



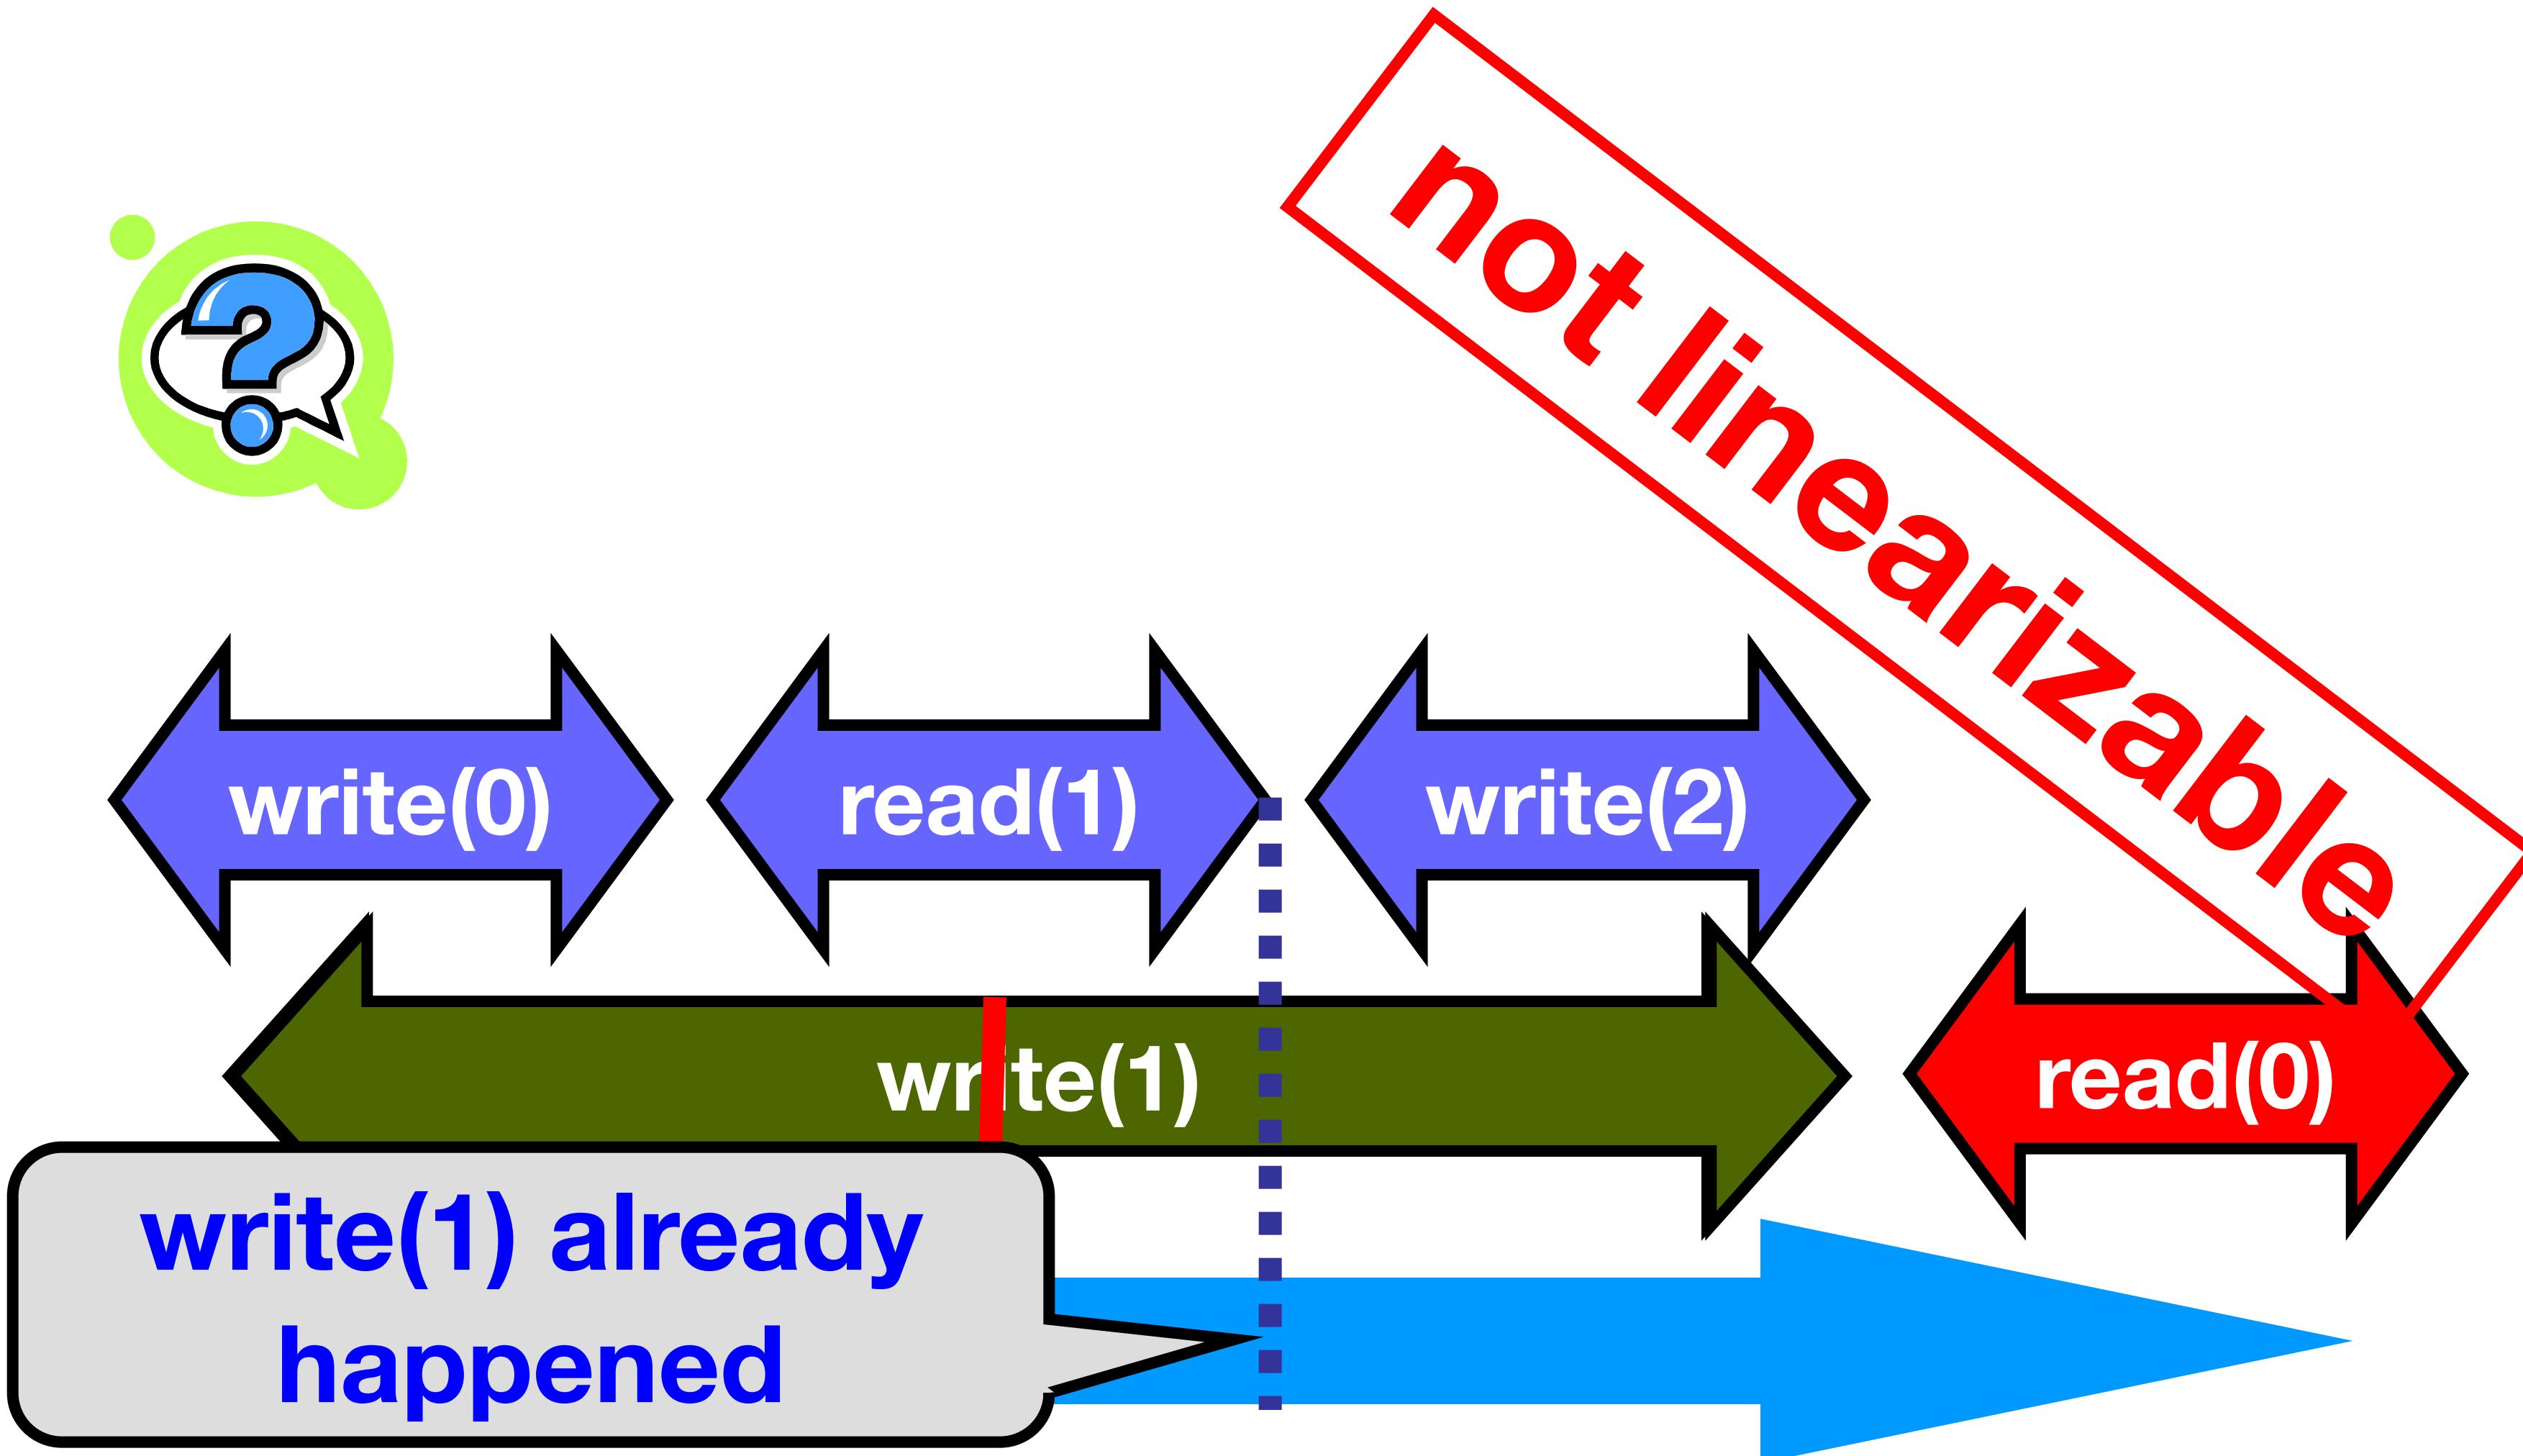
# Read/Write Register Example



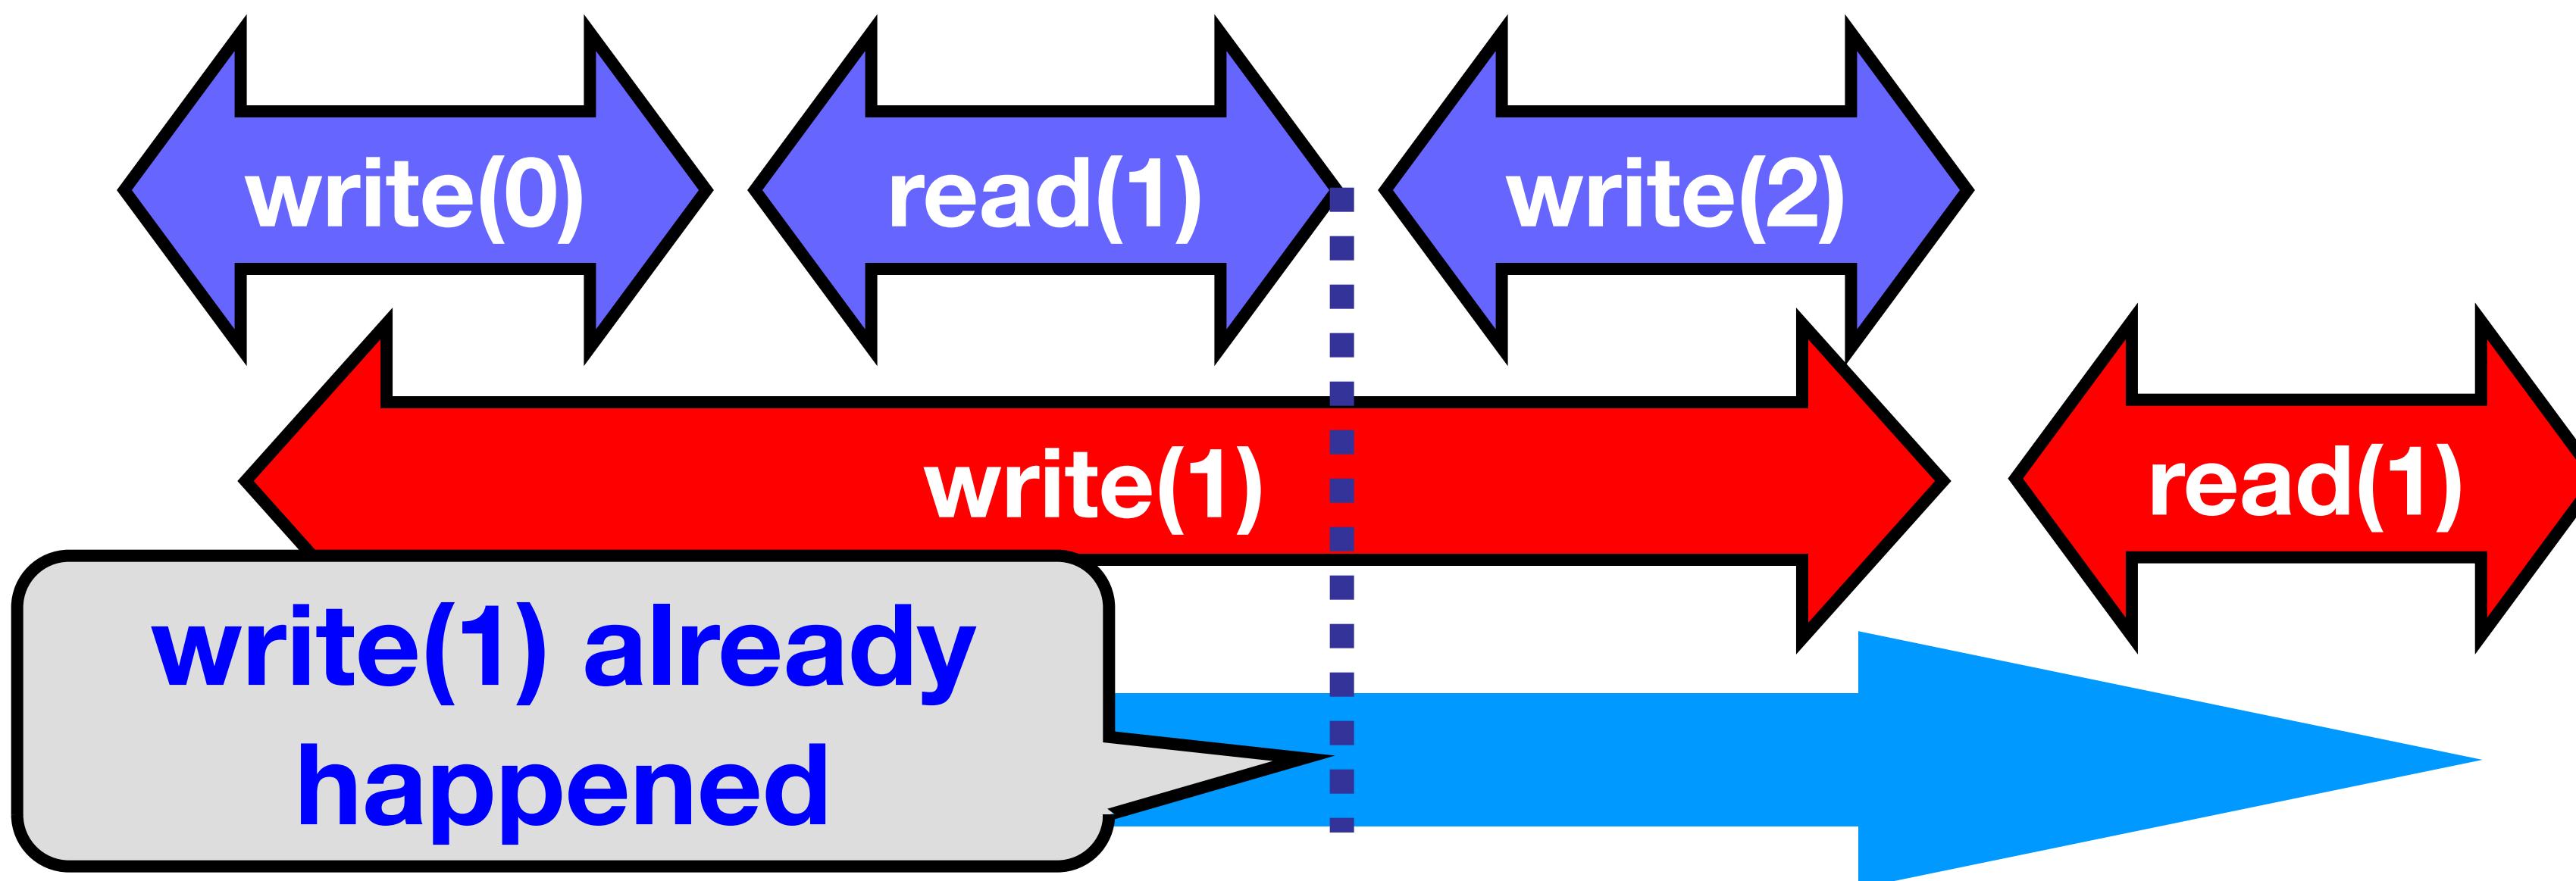
# Read/Write Register Example



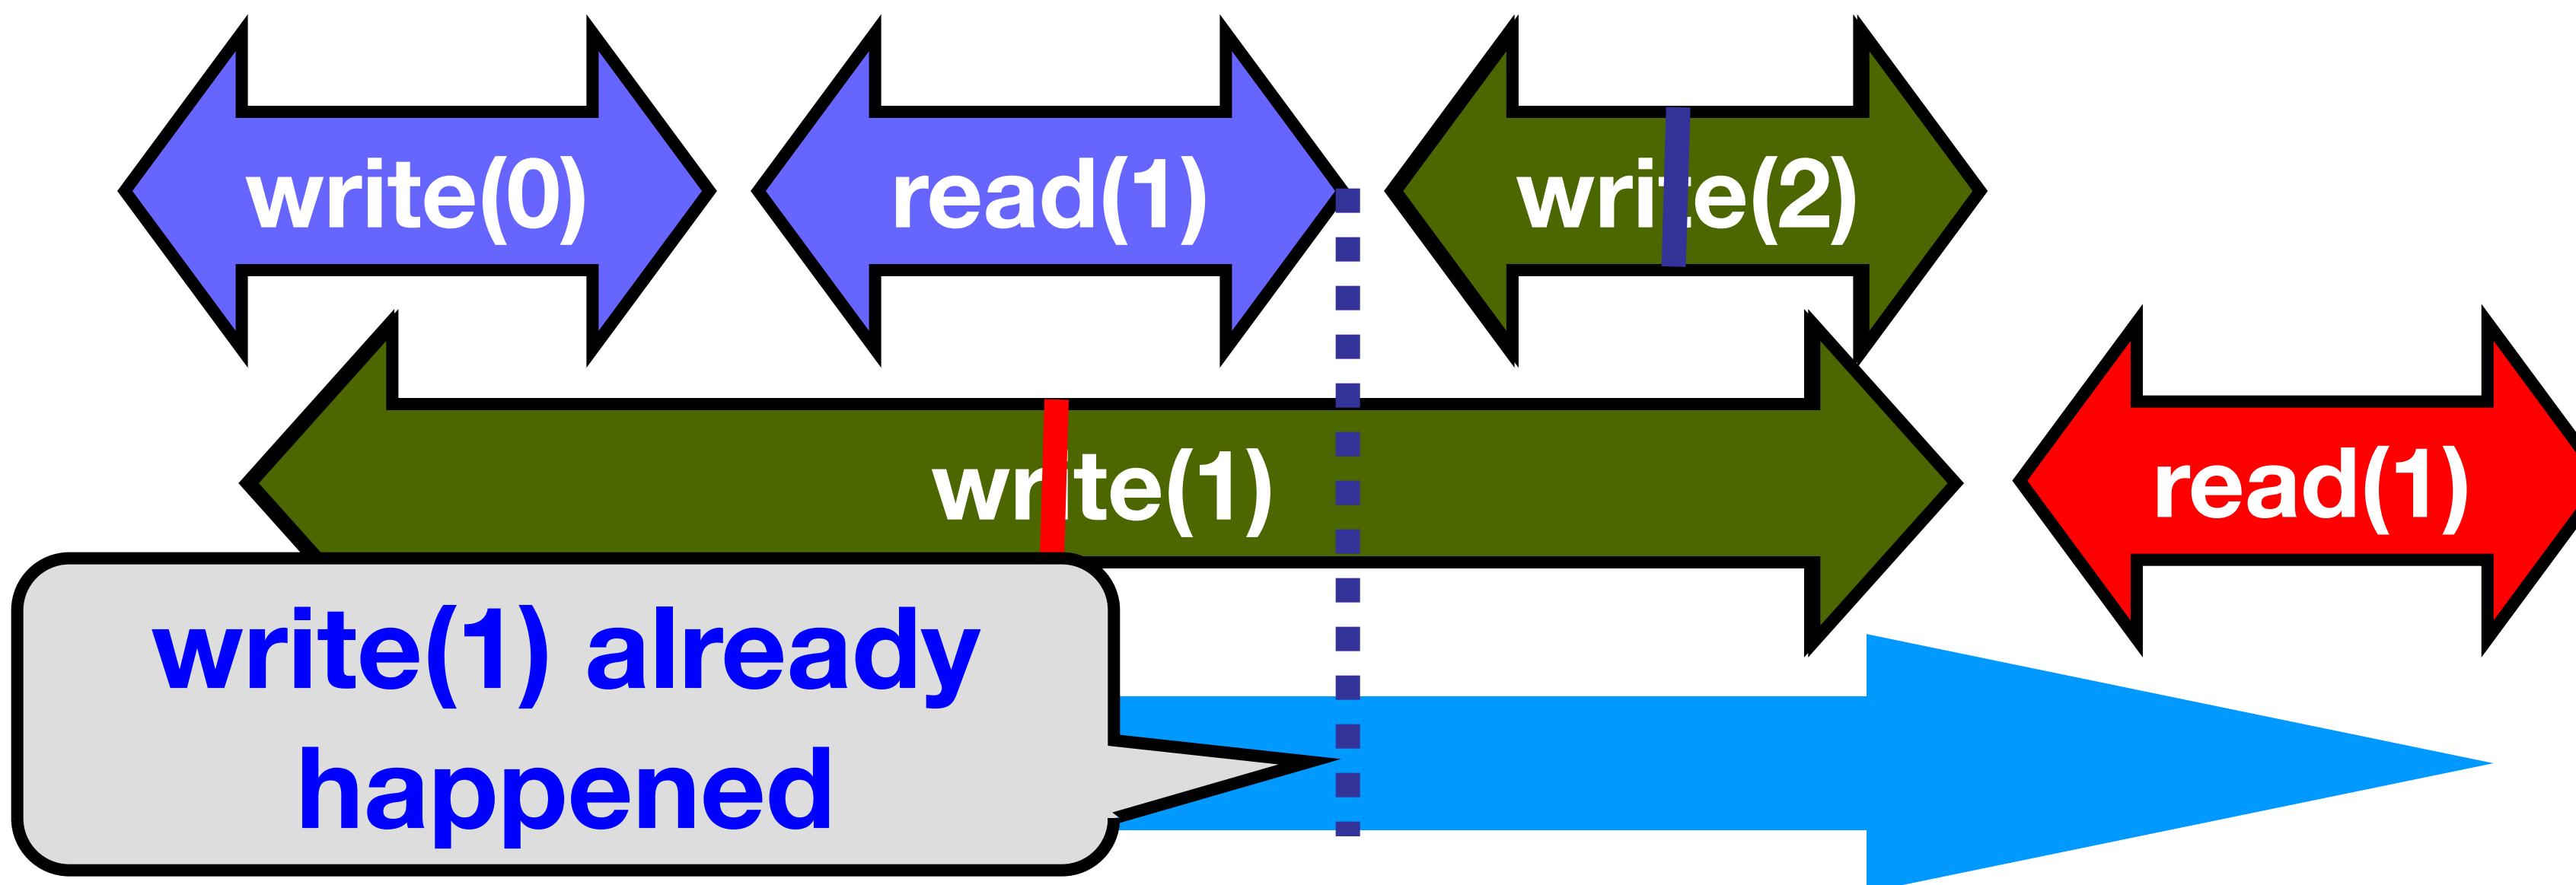
# Read/Write Register Example



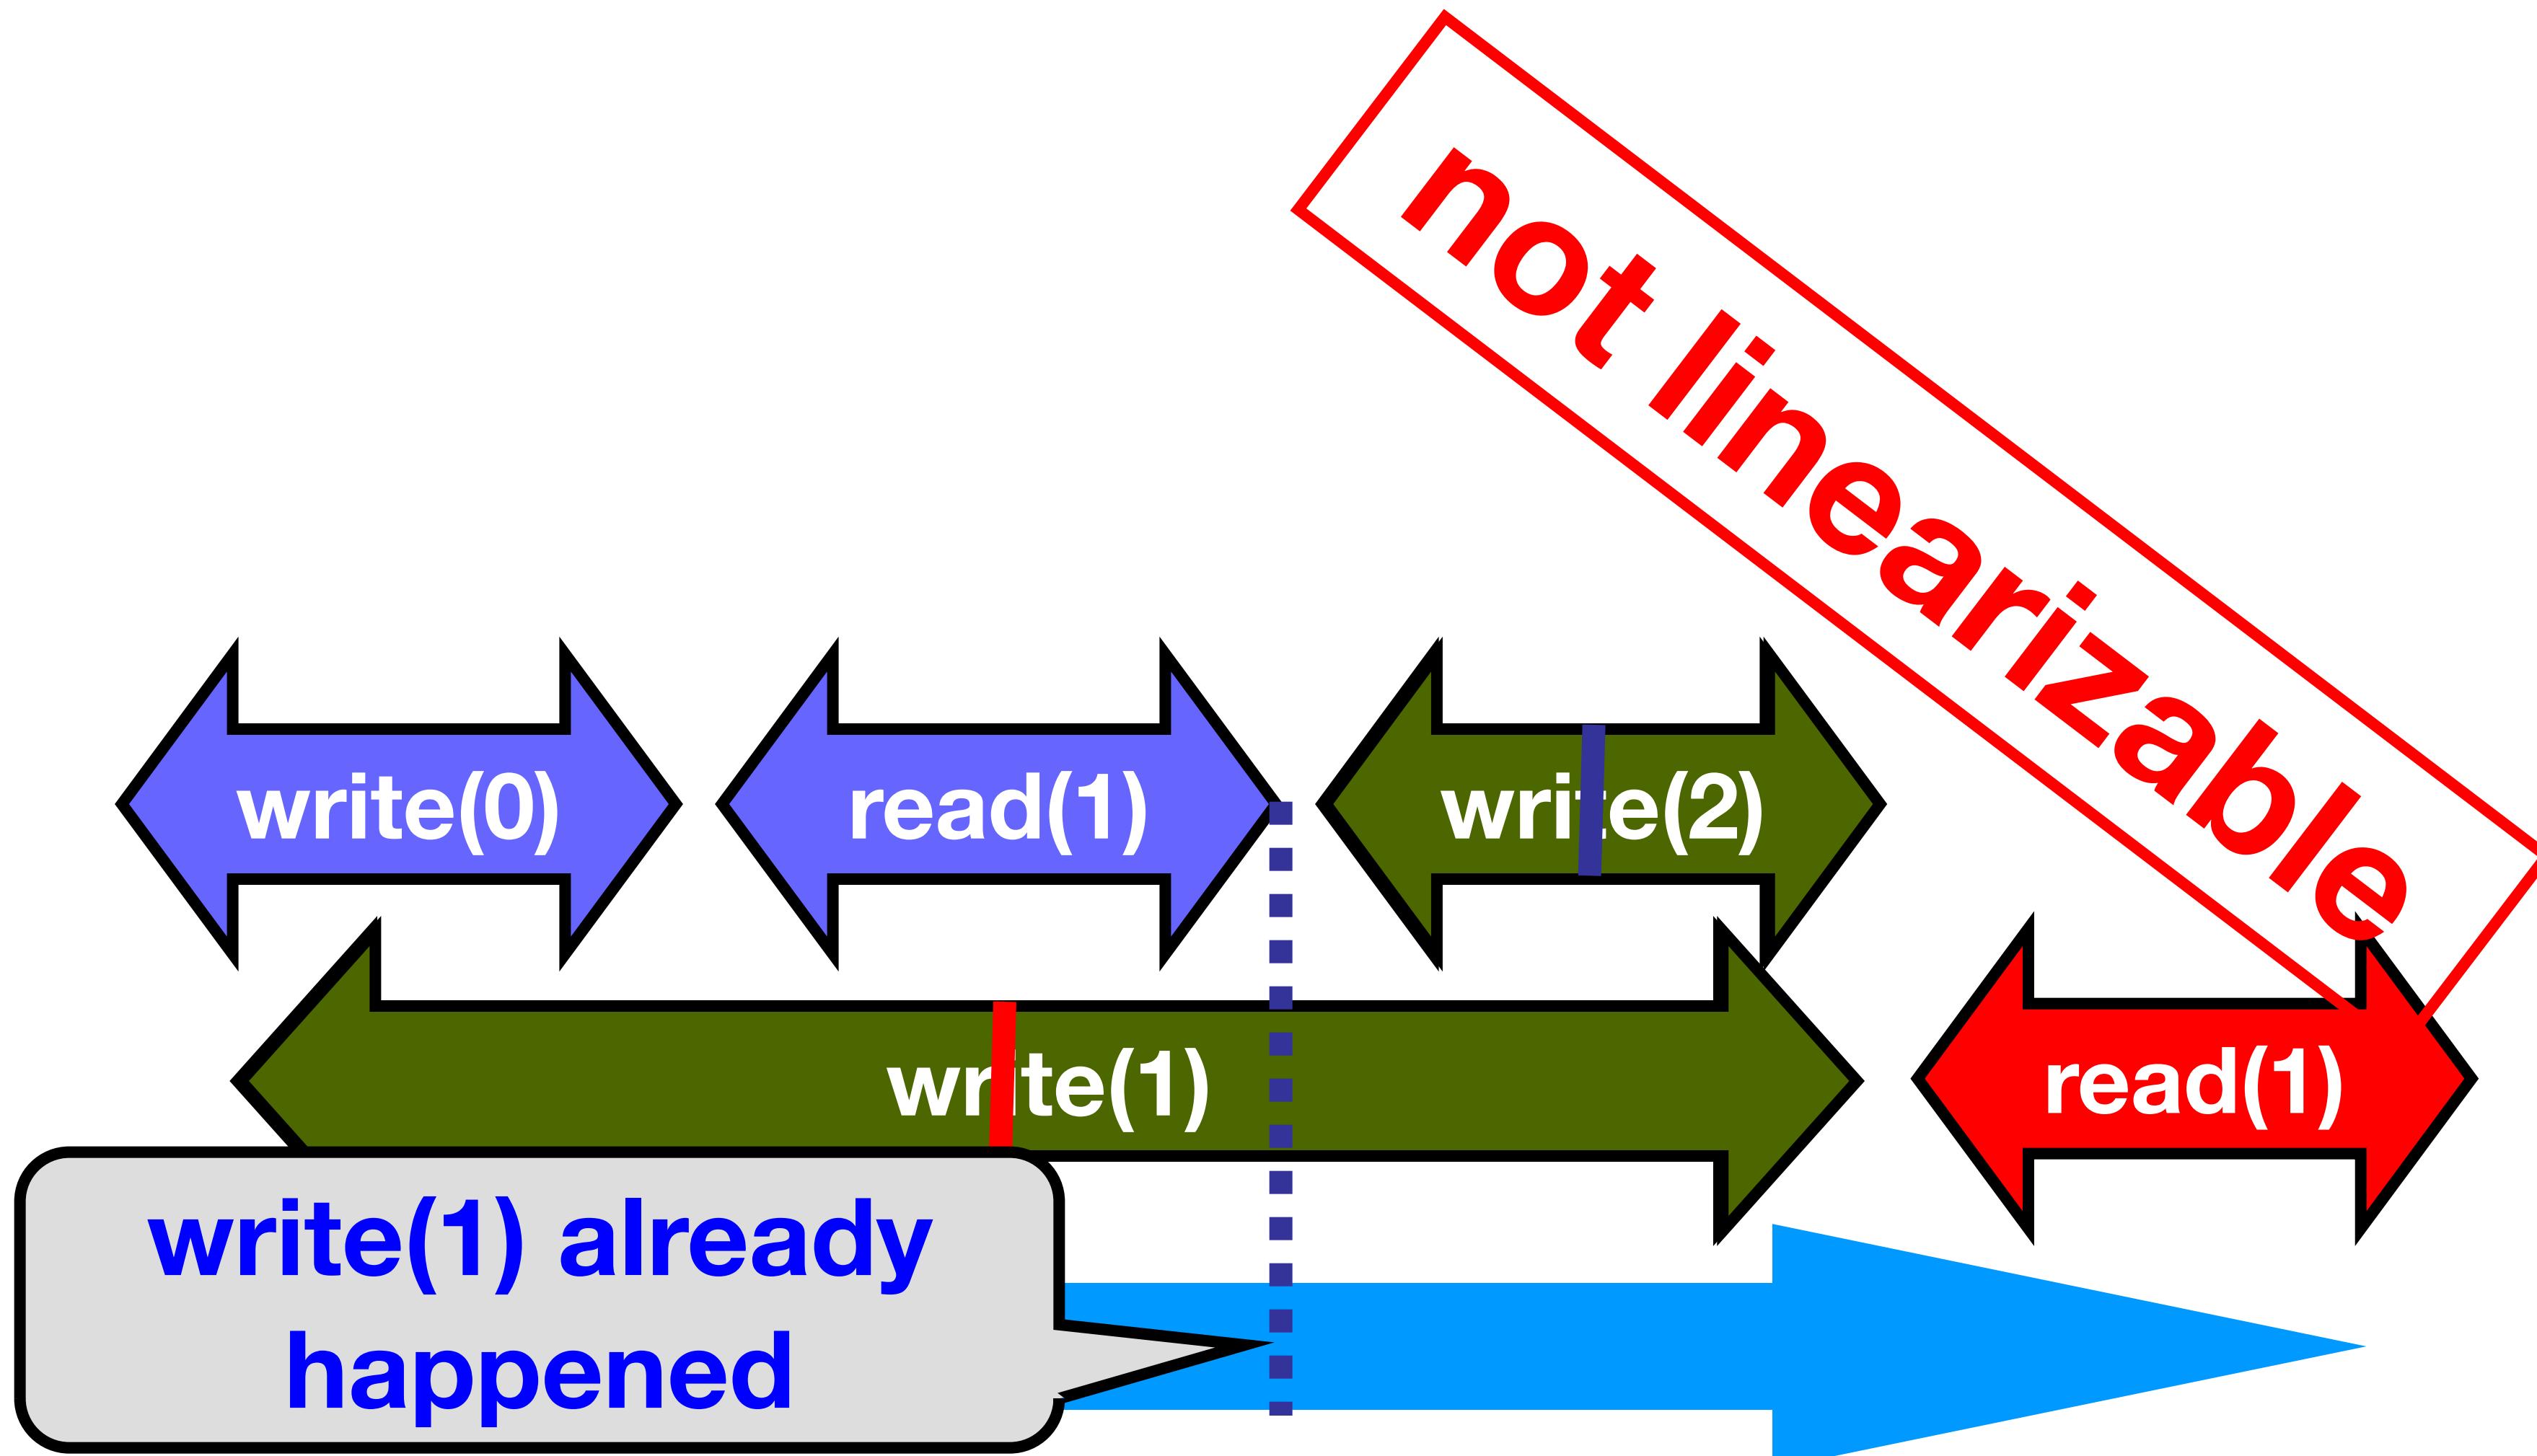
# Read/Write Register Example



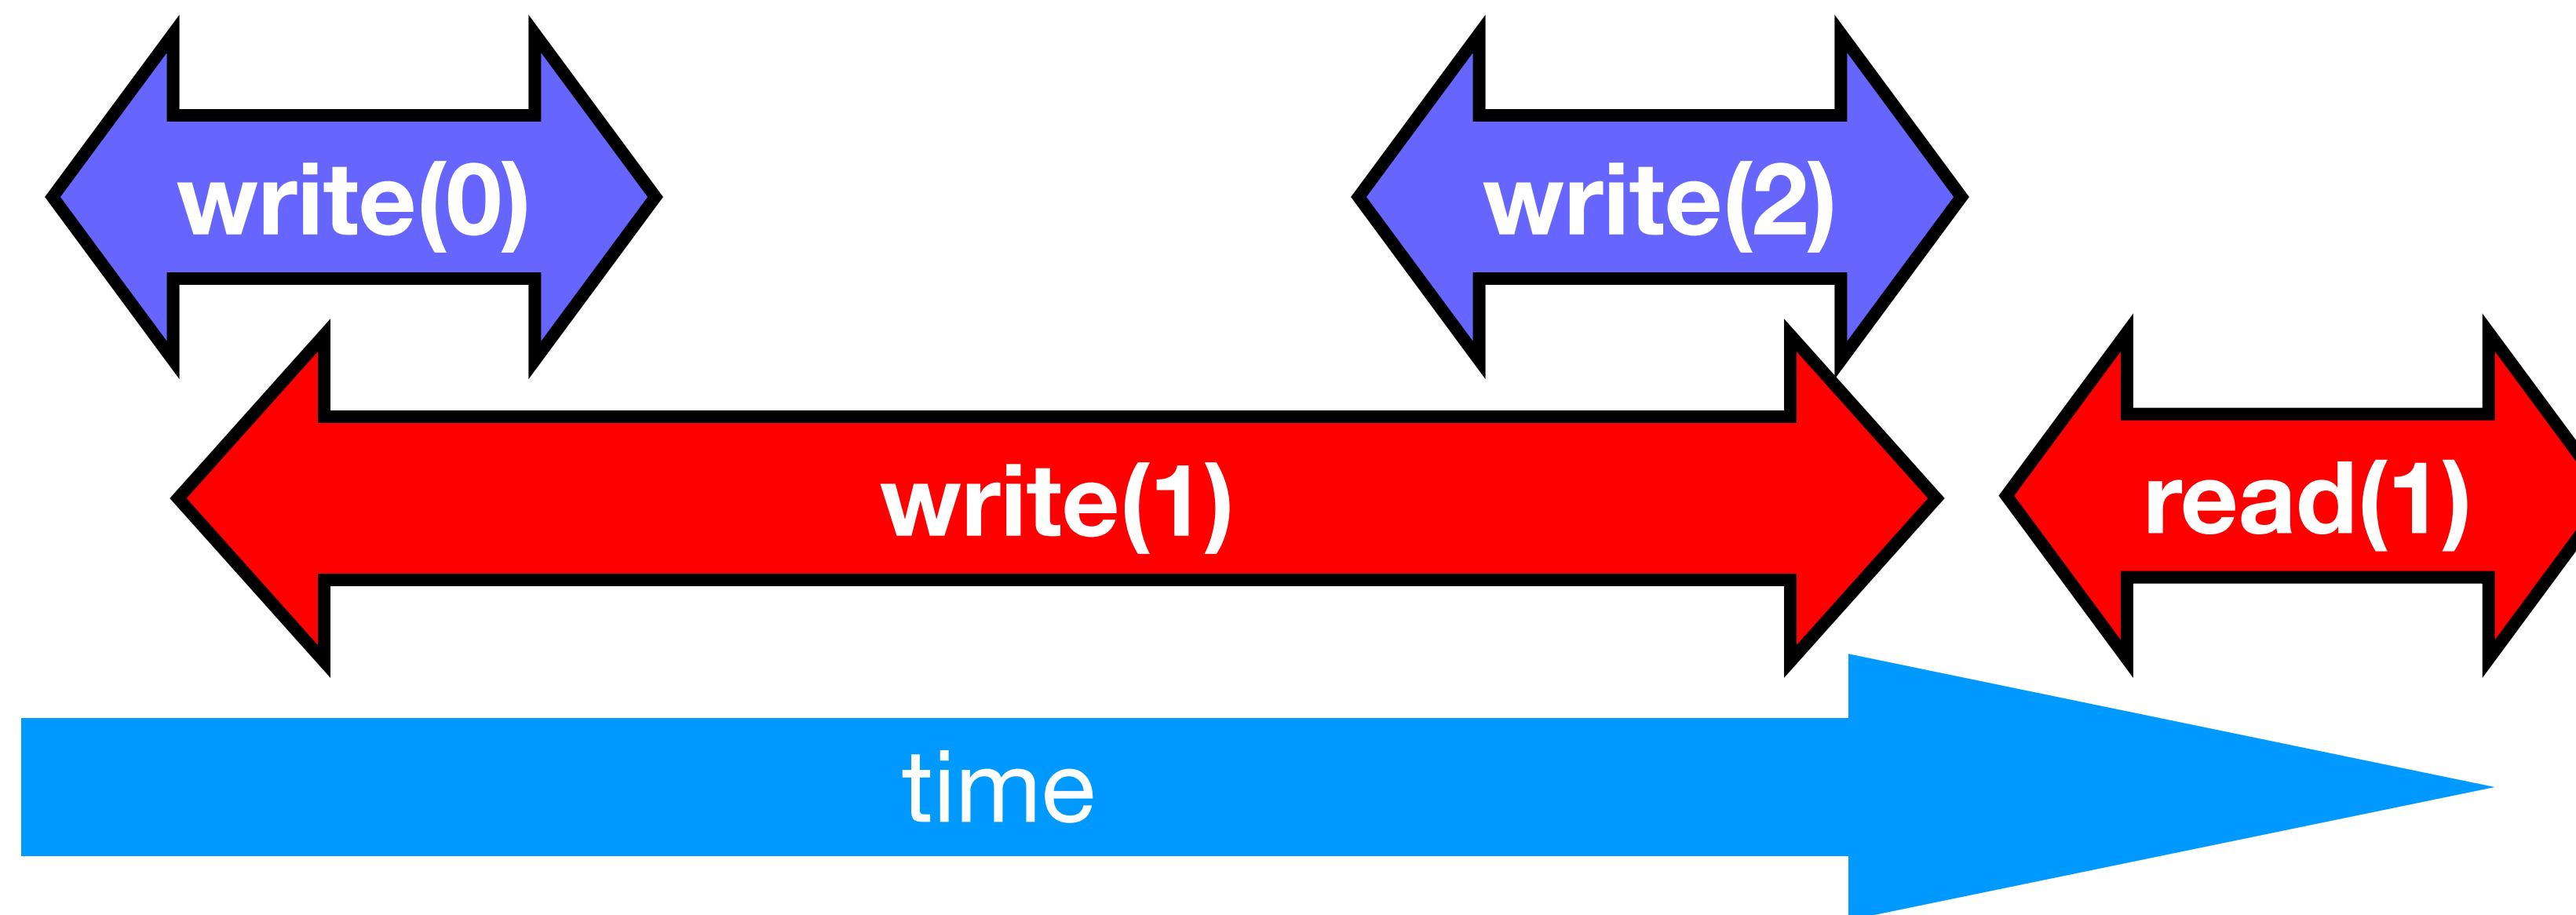
# Read/Write Register Example



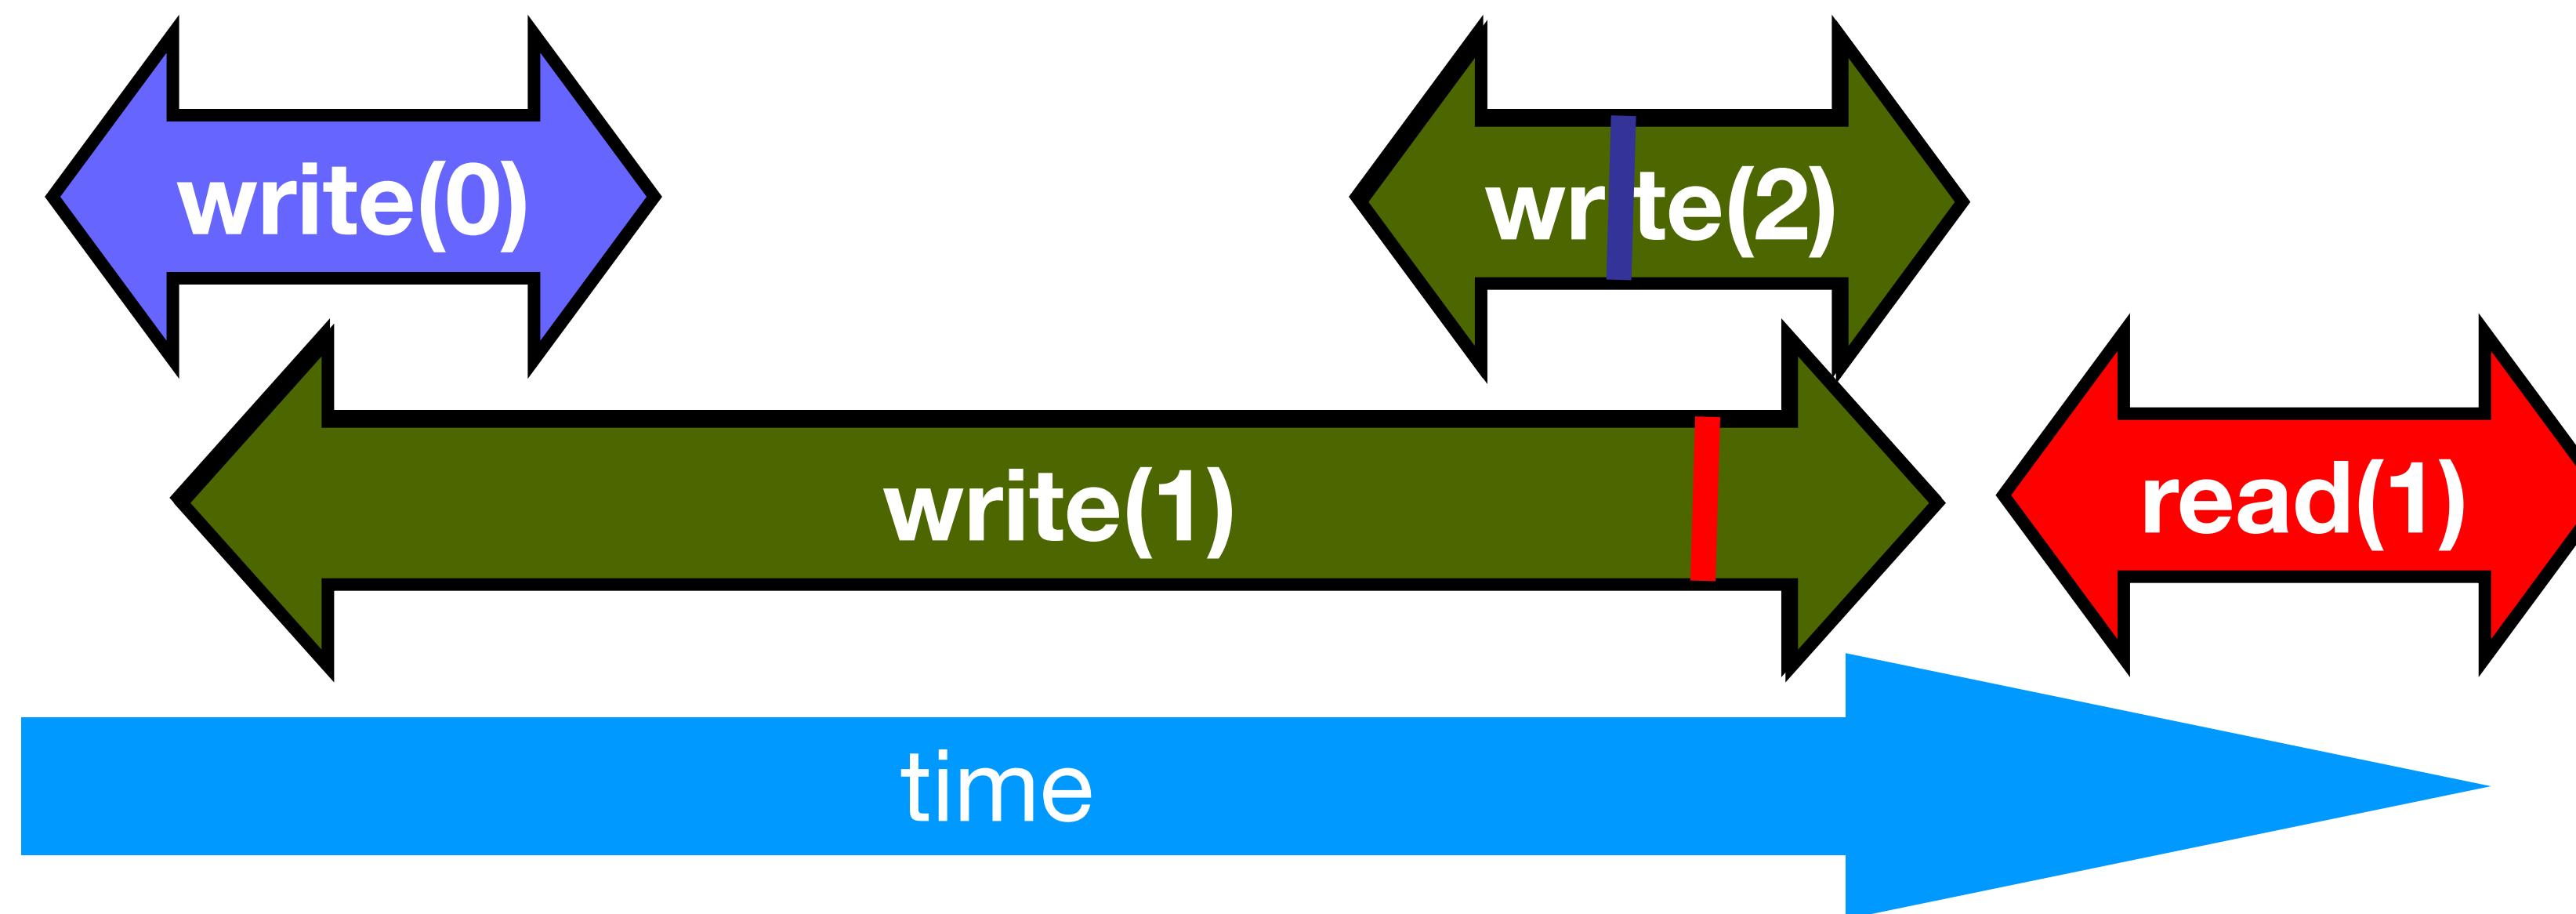
# Read/Write Register Example



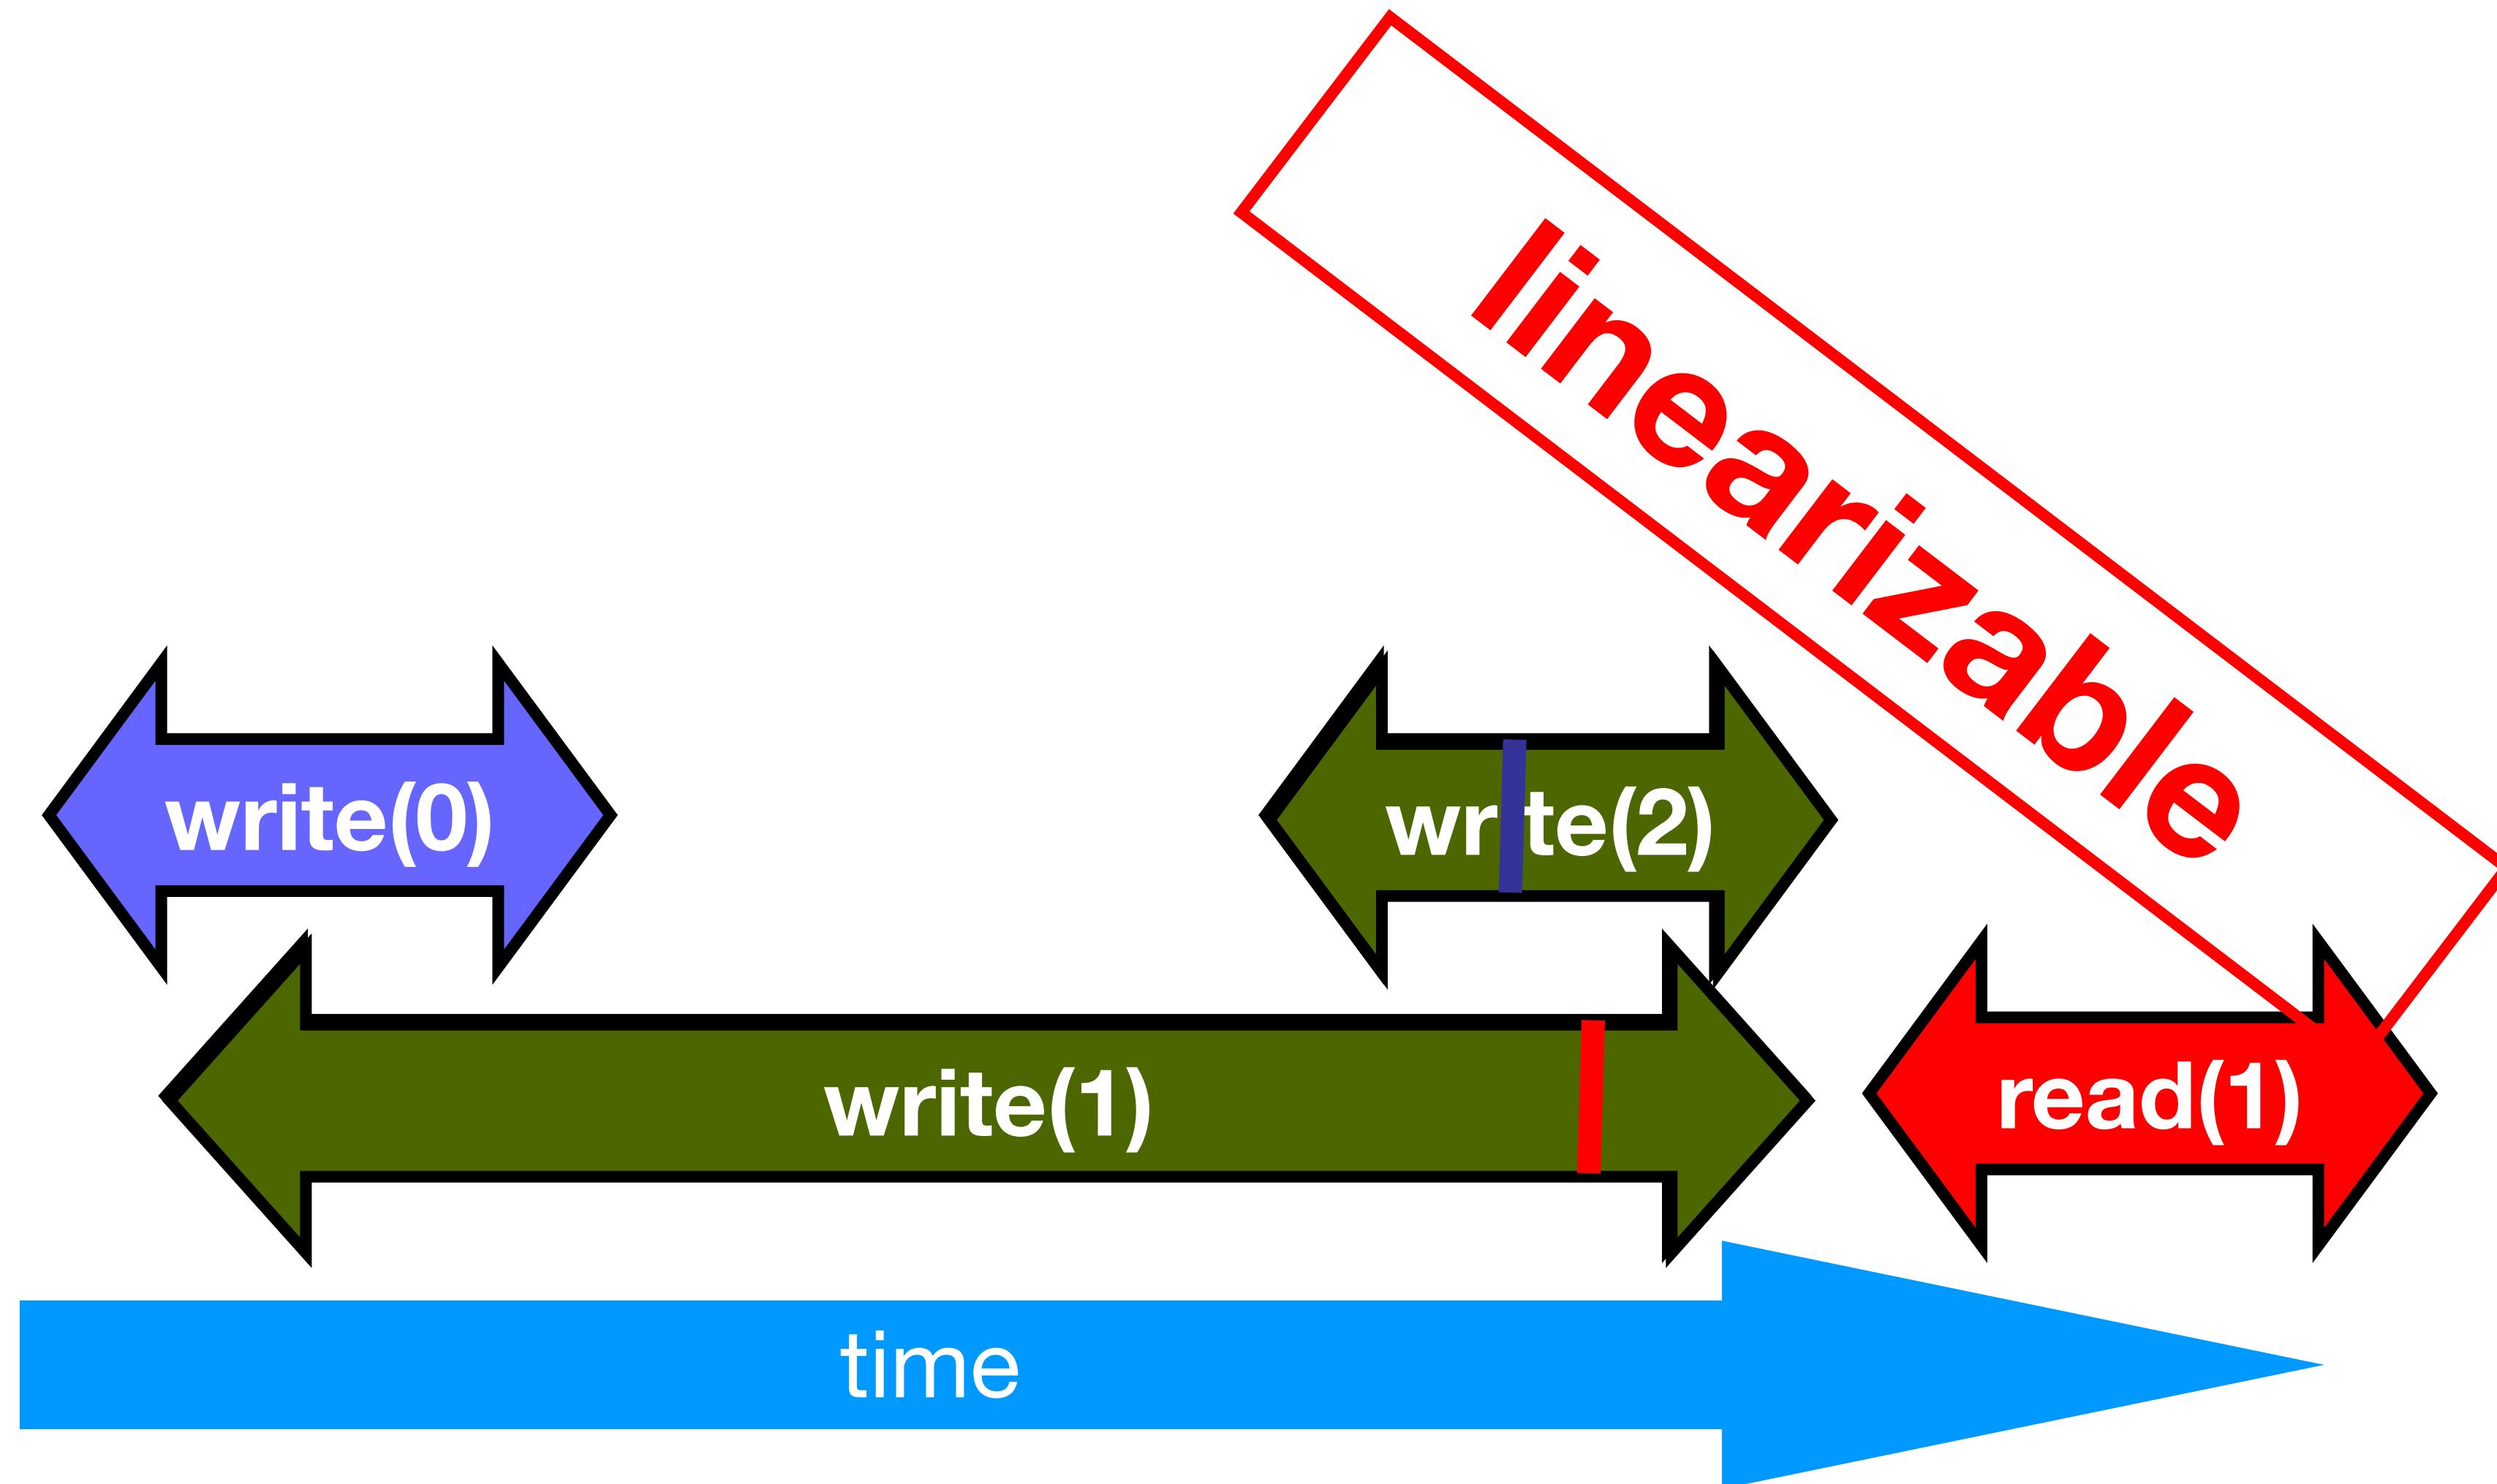
# Read/Write Register Example



# Read/Write Register Example



# Read/Write Register Example



# Talking About Executions

- Why?
  - Can't we specify the linearization point of each operation without describing an execution?
- Not Always
  - In some cases, linearization point ***depends on the execution***

# Formal Model of Executions

- Define precisely what we mean
  - Ambiguity is bad when intuition is weak
- Allow reasoning
  - Formal
  - But mostly informal
    - In the long run, actually, more important

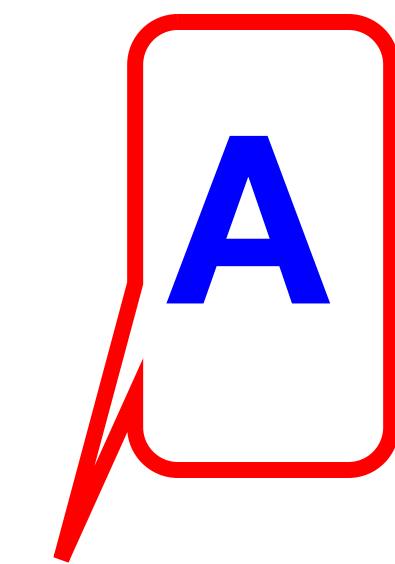
# Split Method Calls into Two Events

- Invocation
  - method name & args
  - `q.enq(x)`
- Response
  - result or exception
  - `q.enq(x)` returns **void**
  - `q.deq()` returns **x**
  - `q.deq()` throws **empty**
- Note that I'm following the convention of the book
  - Book uses OO
  - Code in this course uses FP
  - Note that we're still reasoning using **objectivism**
  - For the current discussion, distinction doesn't matter
    - `q.enq(x)` is read as **enq q x** in code
    - Returns **void** is read as returns **()**
    - Throws **empty** is read as raises **Empty**

# Invocation Notation

A q.enq(x)

# Invocation Notation

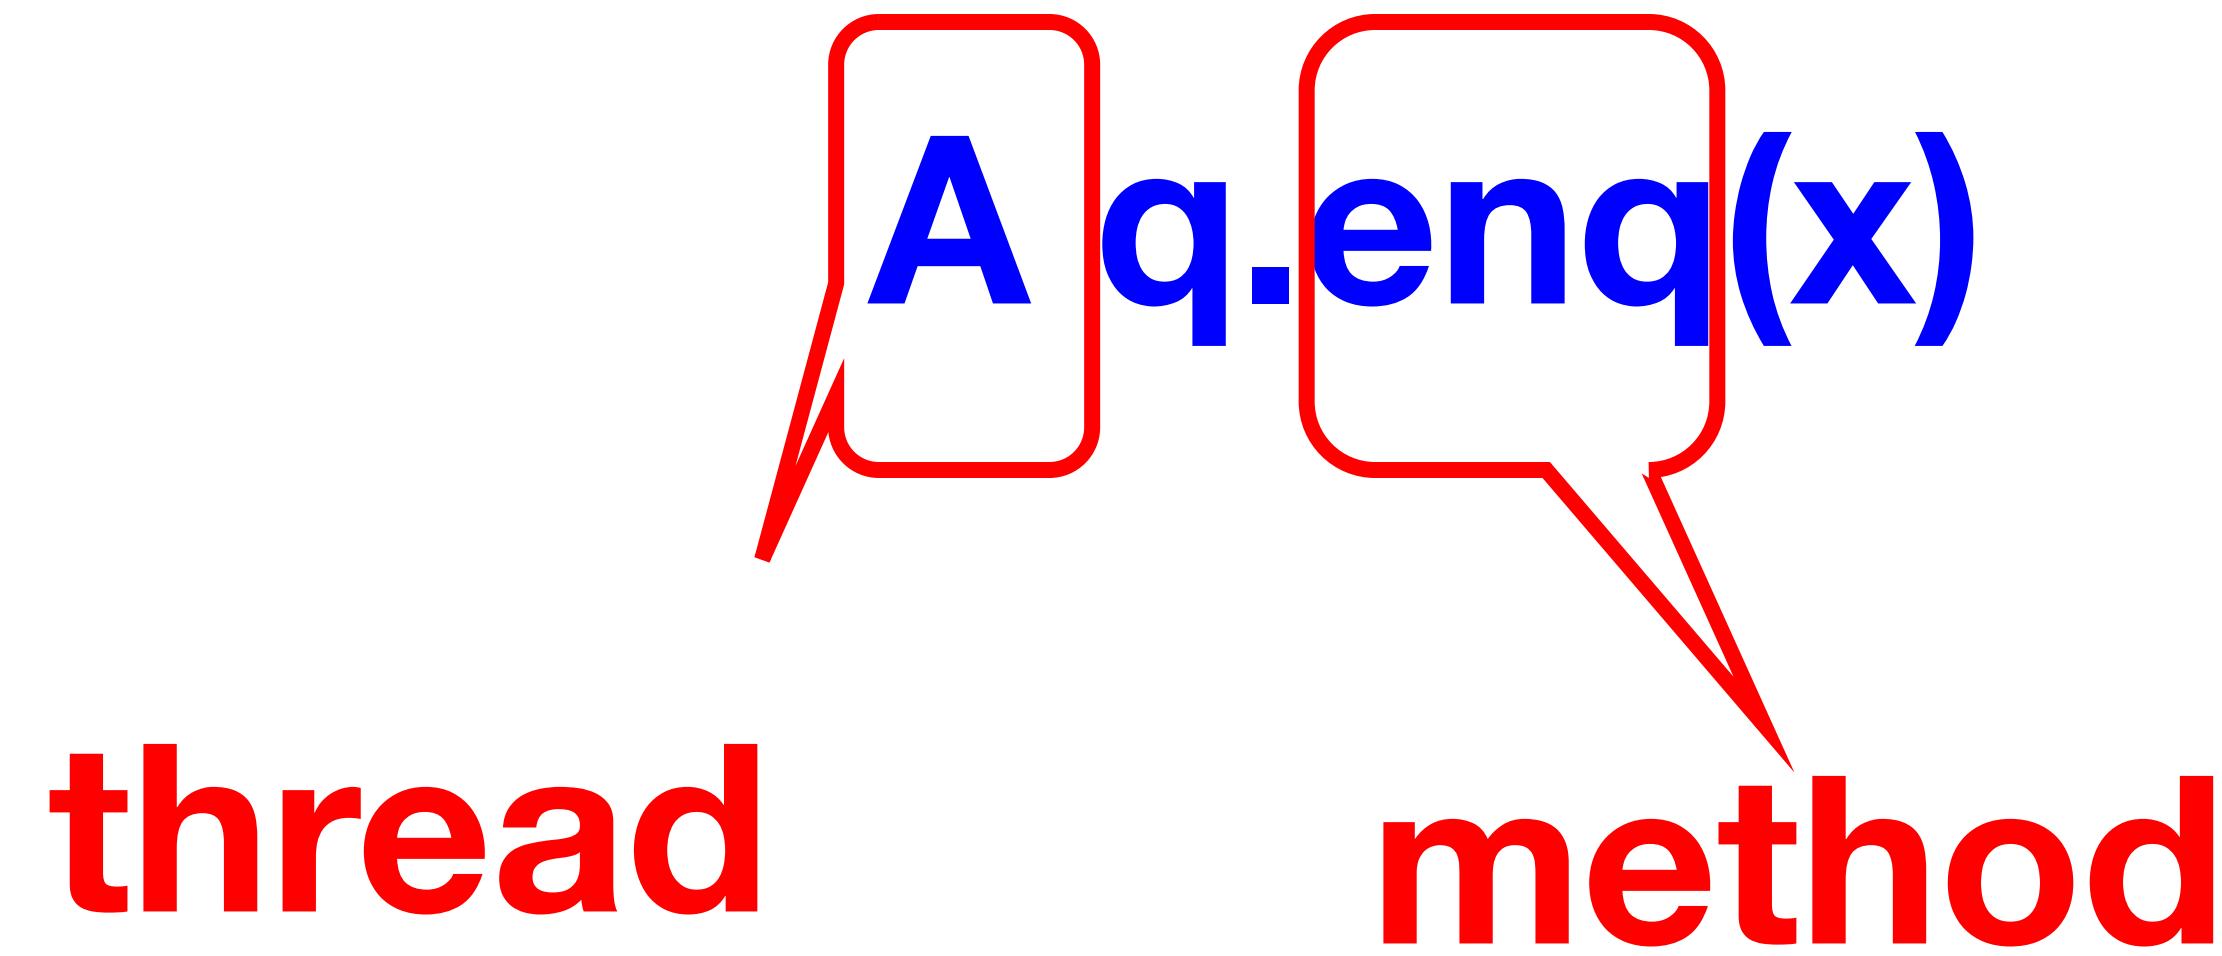


A red speech-like bracket points from the word "thread" below to the letter "A" in the text "A q.enq(x)" above.

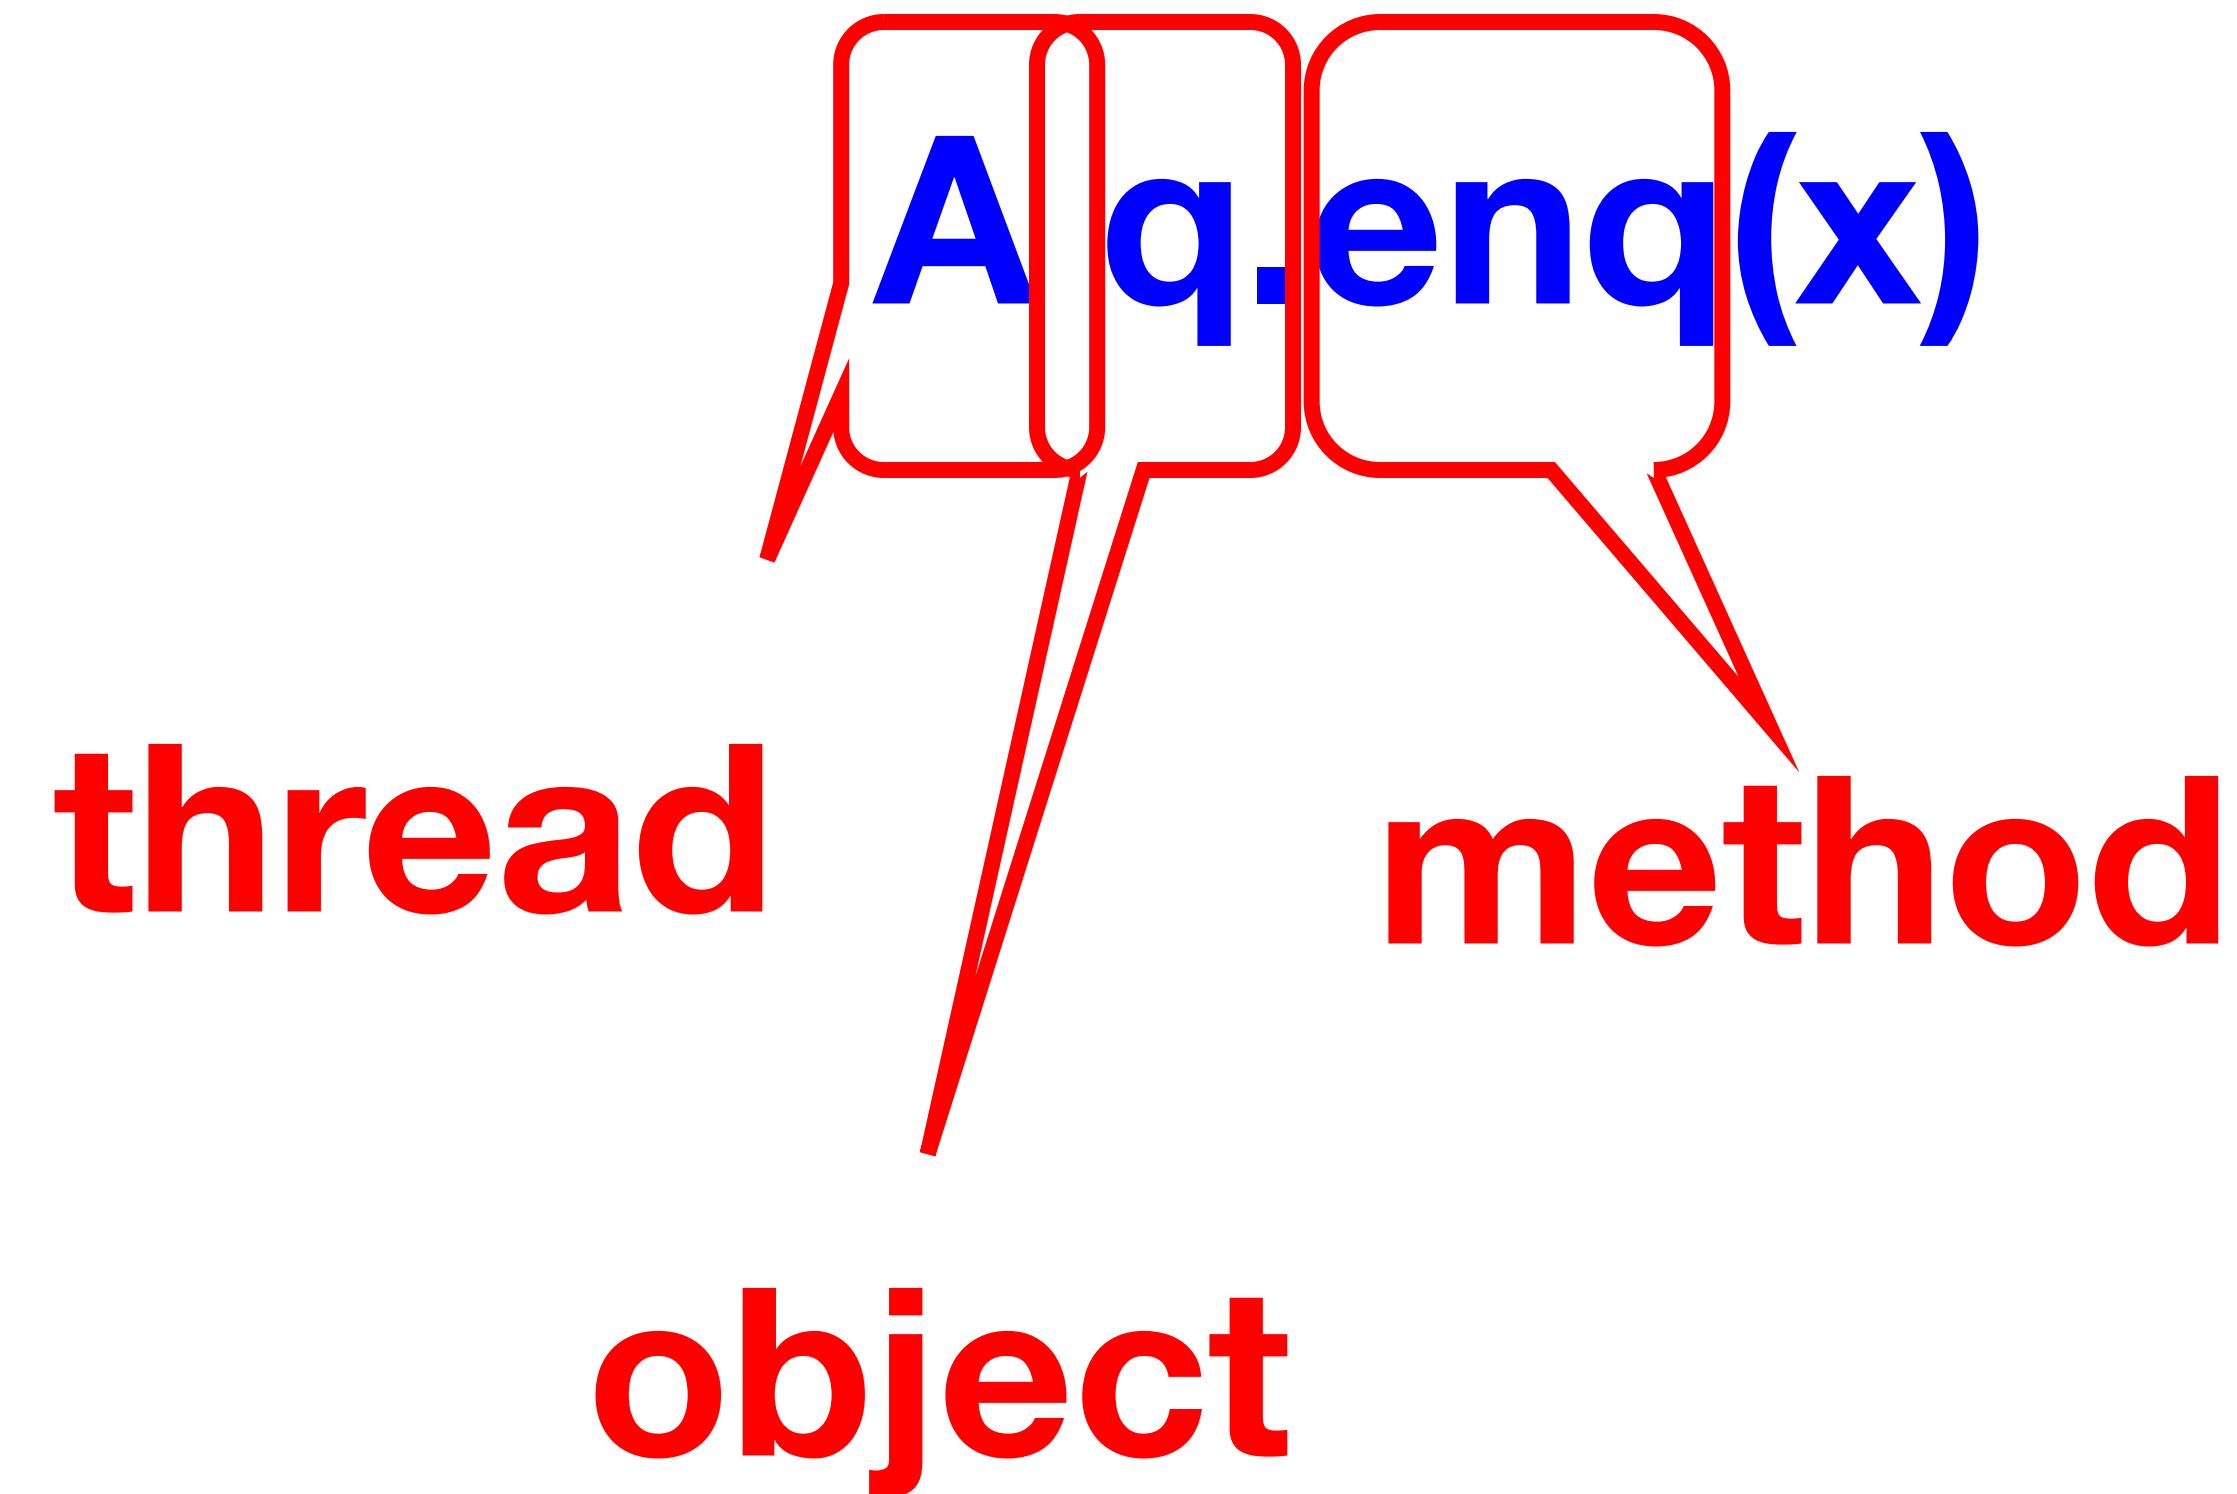
**A** q.enq(x)

**thread**

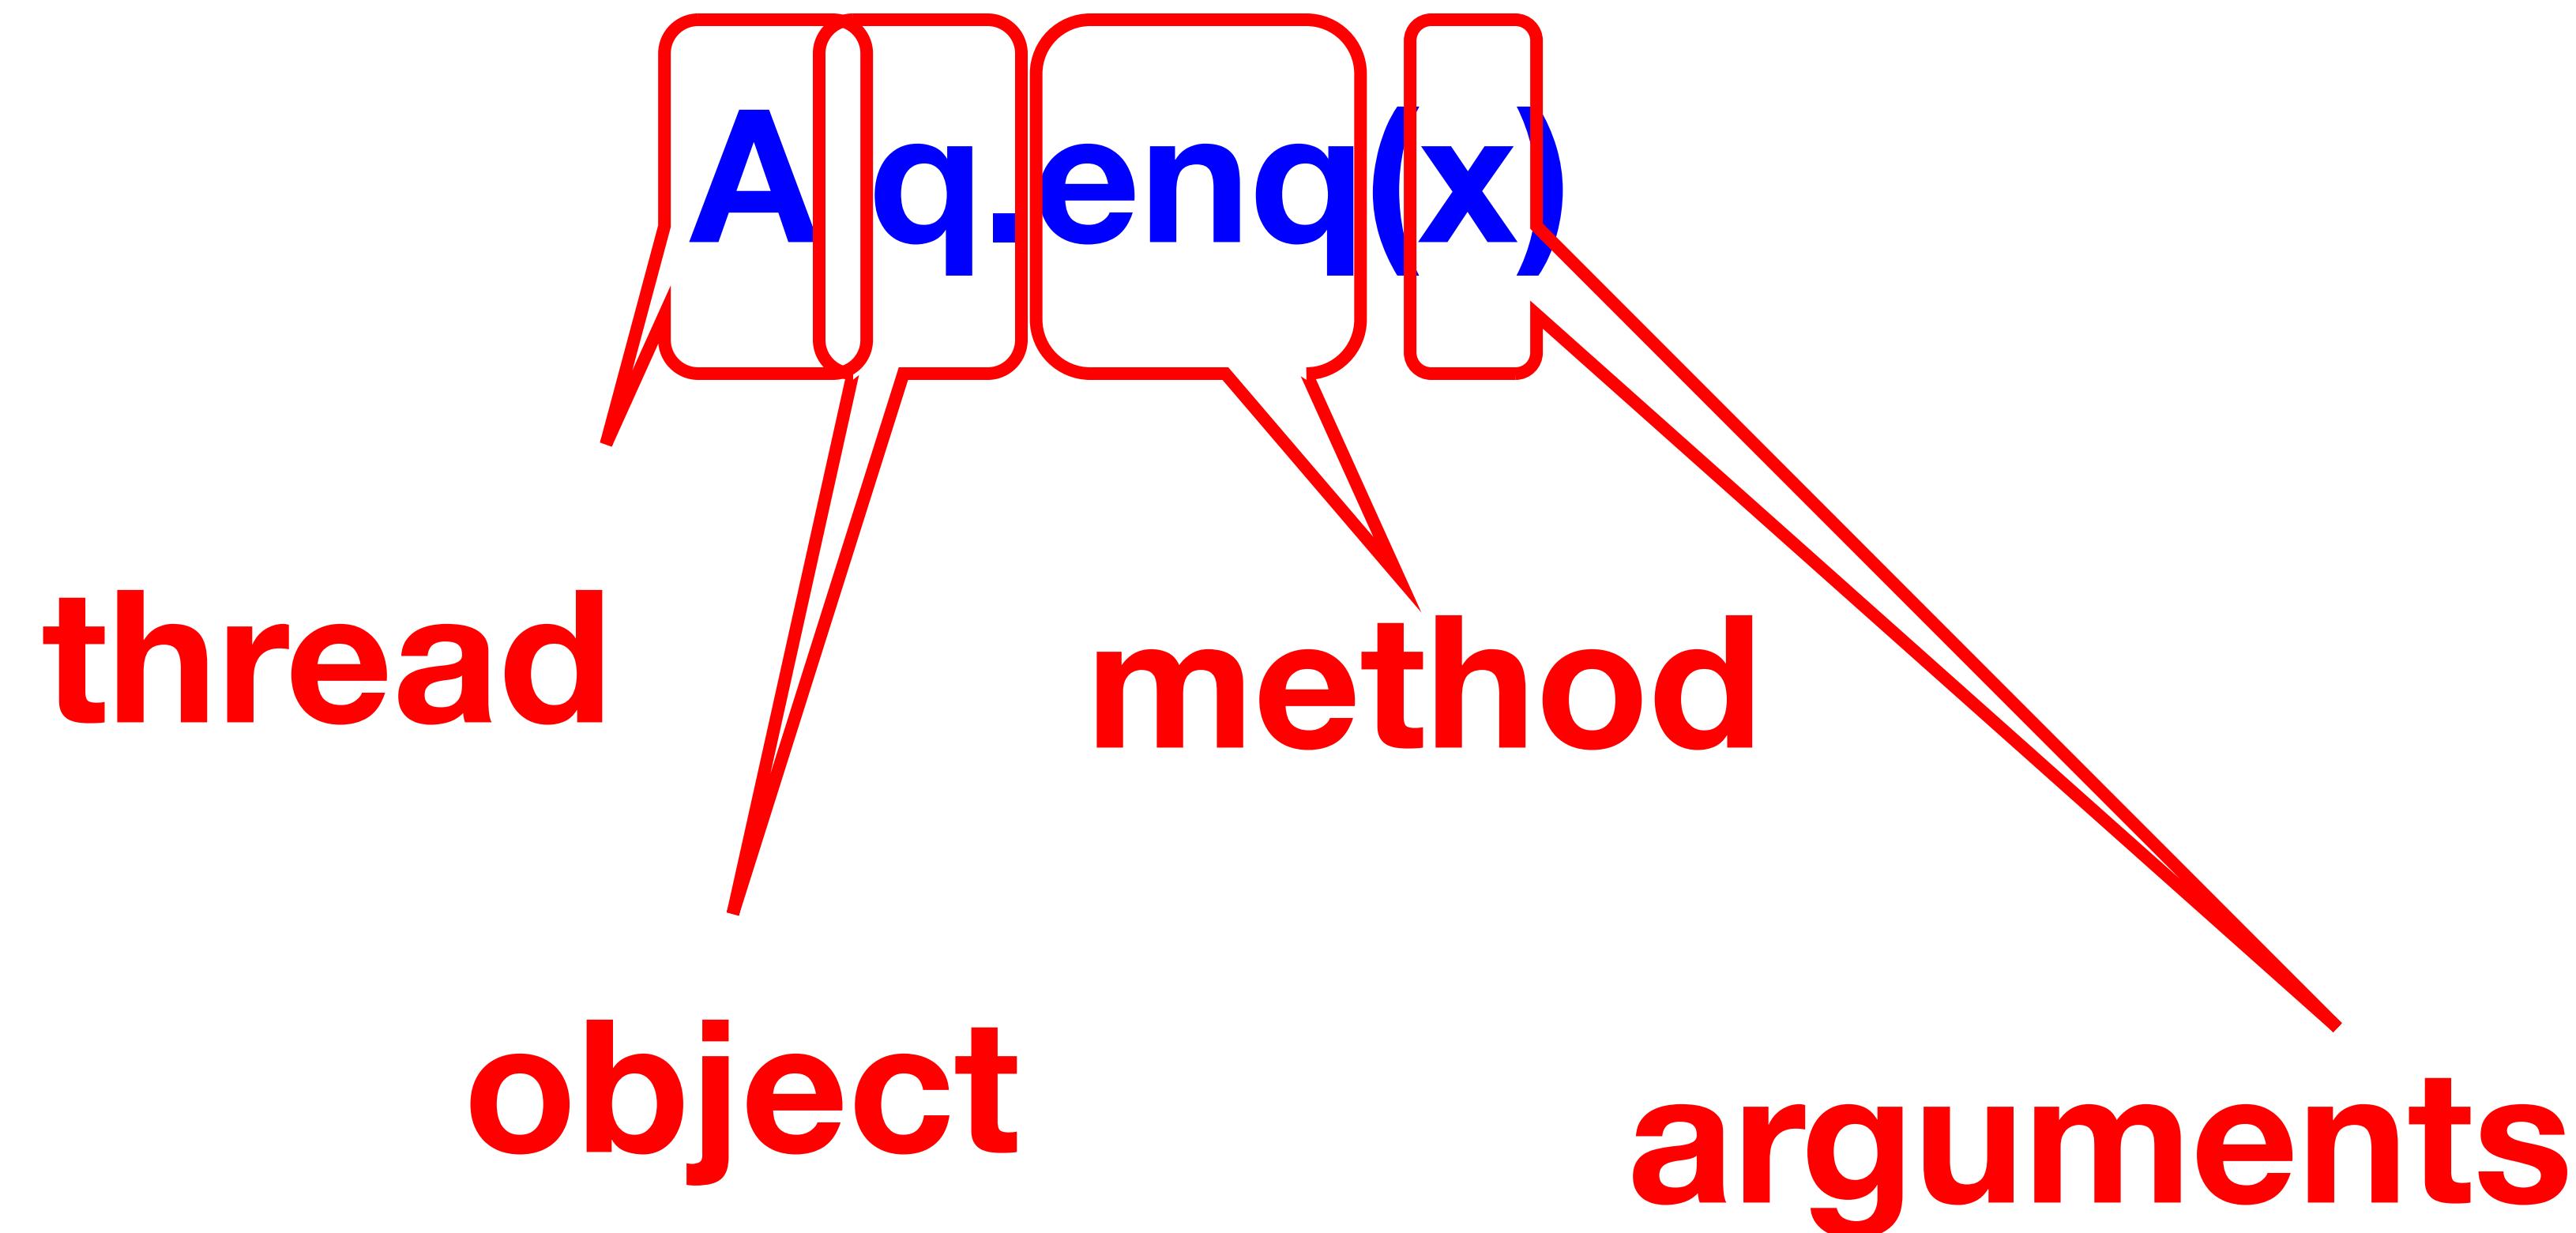
# Invocation Notation



# Invocation Notation



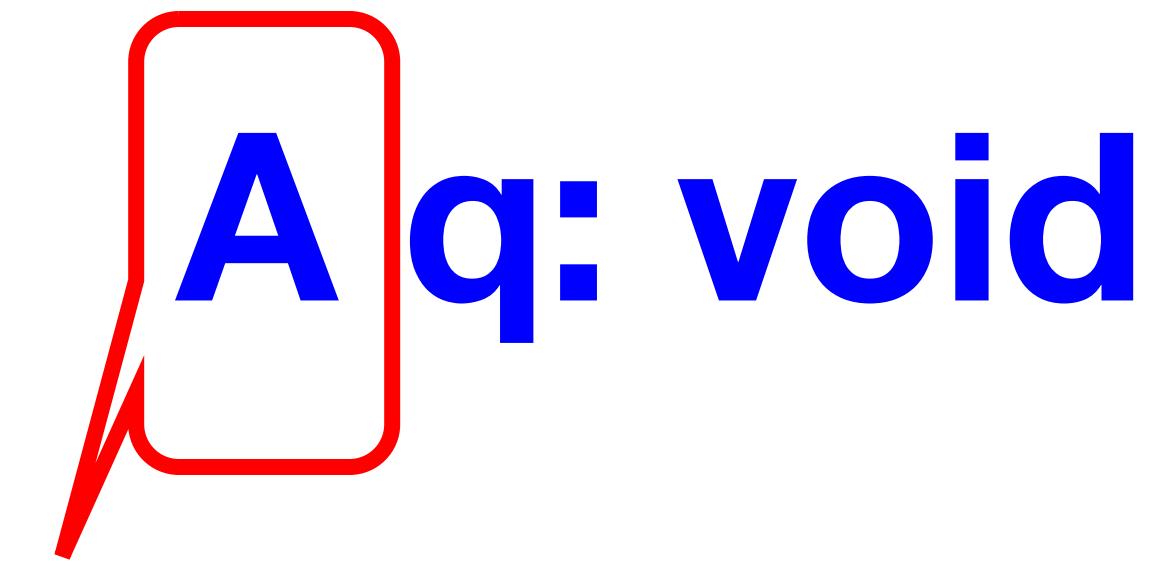
# Invocation Notation



# Response Notation

A q: void

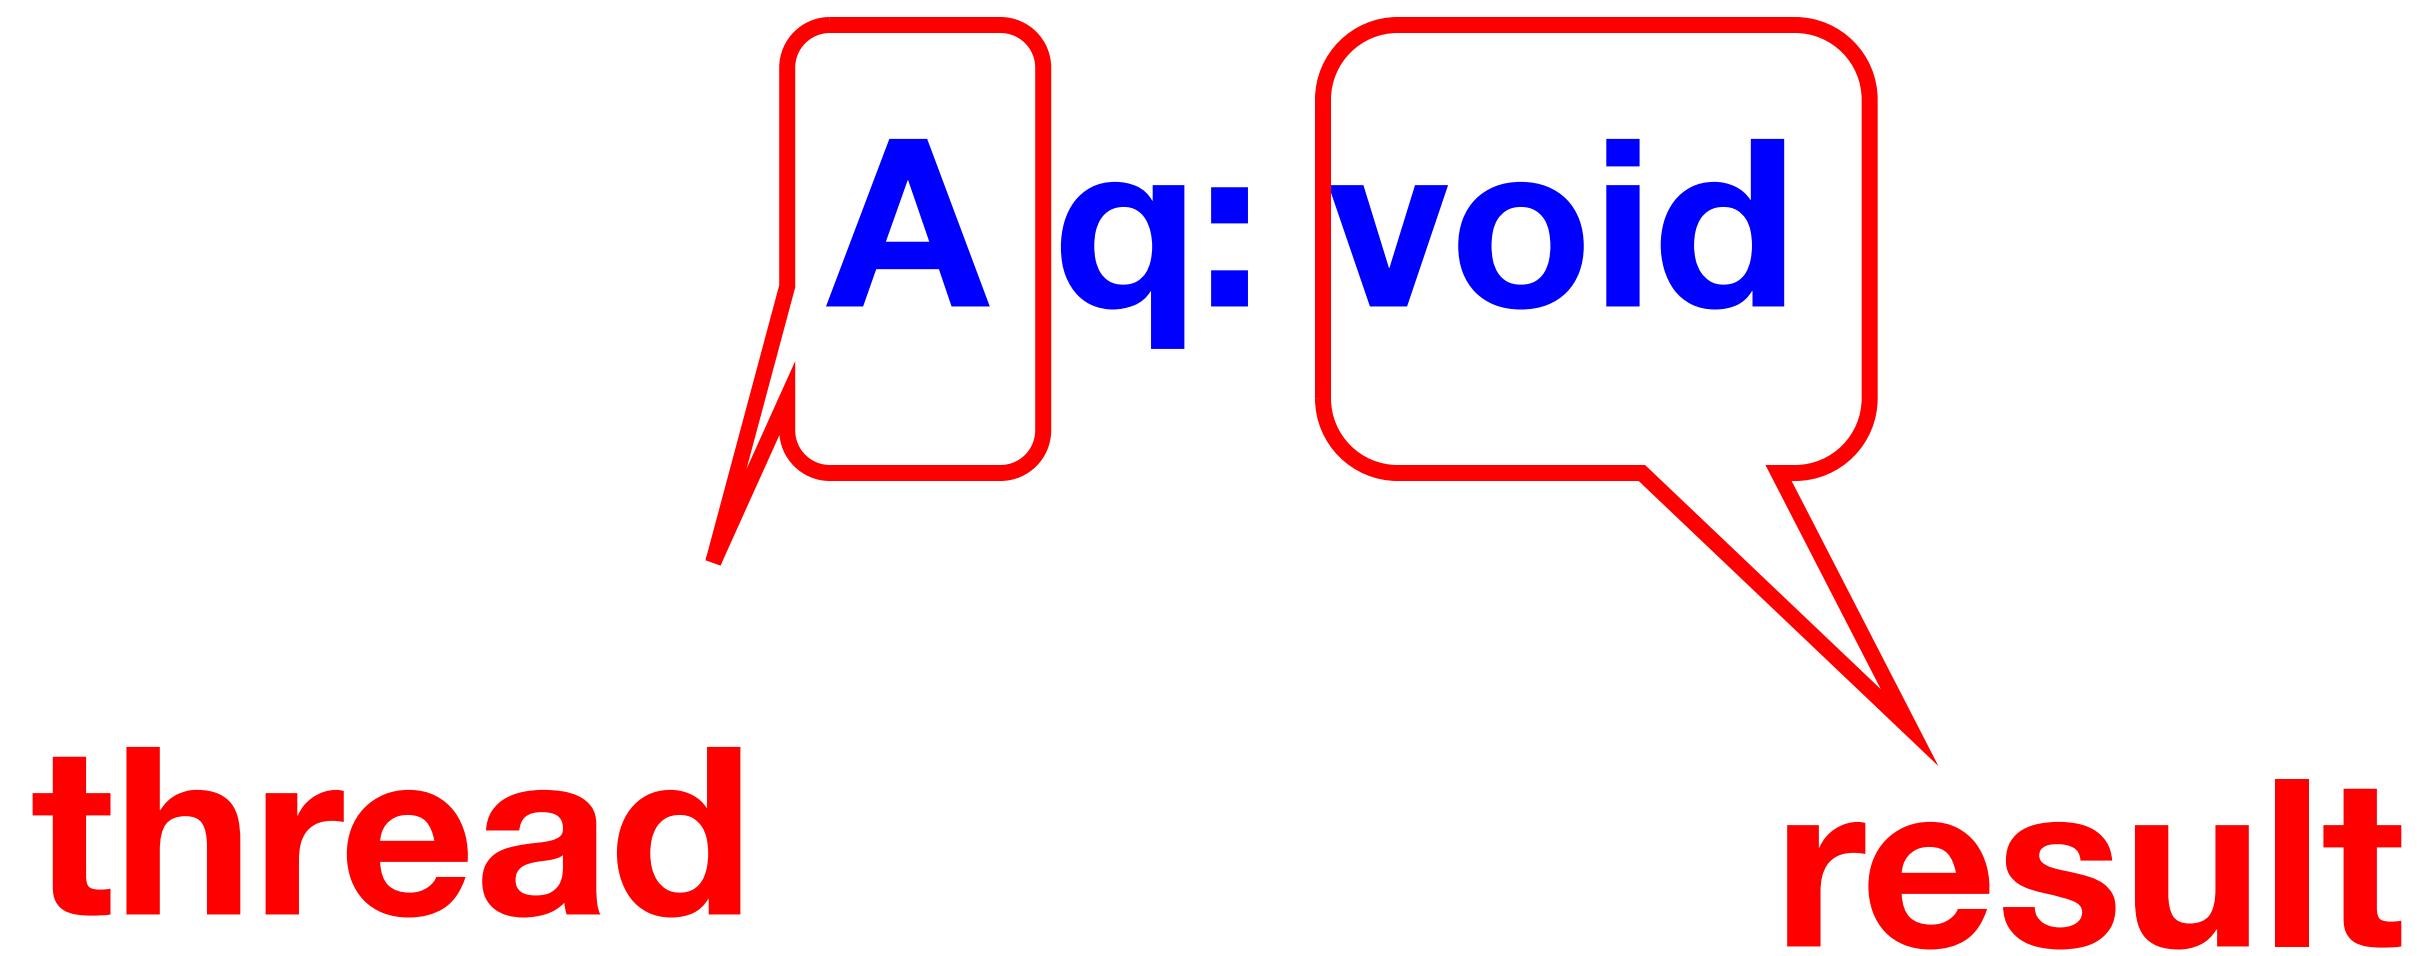
# Response Notation



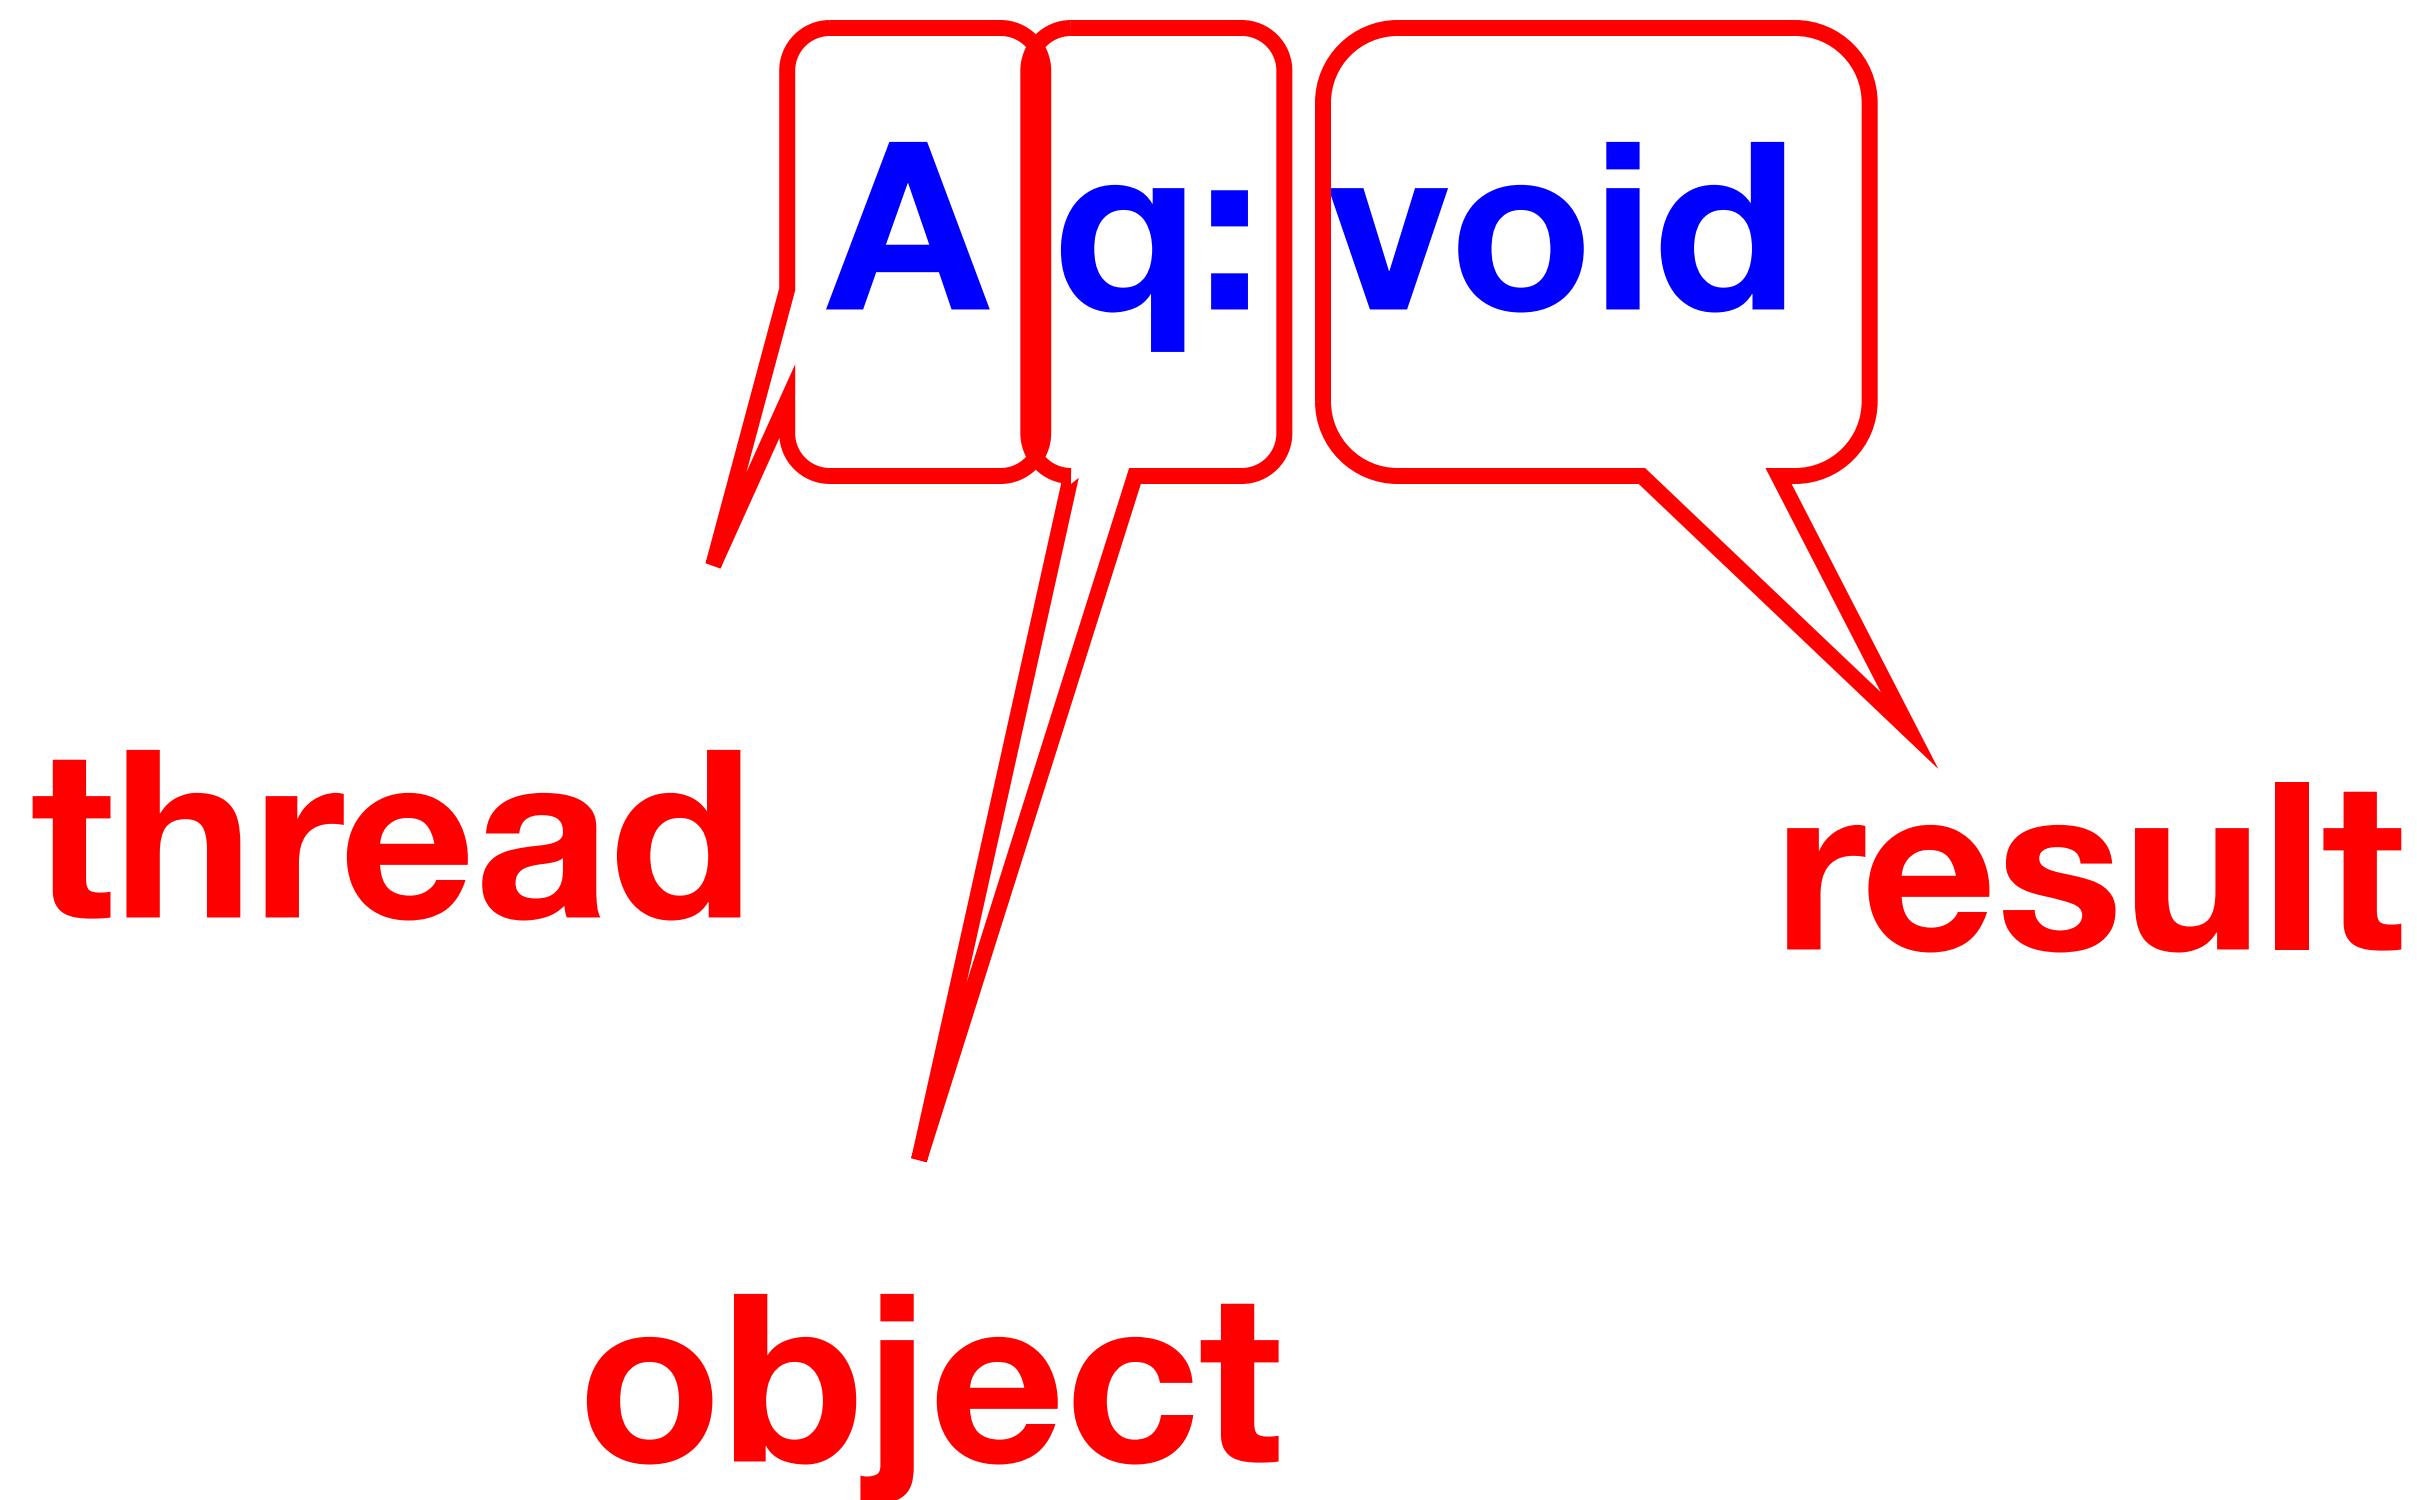
**q: void**

**thread**

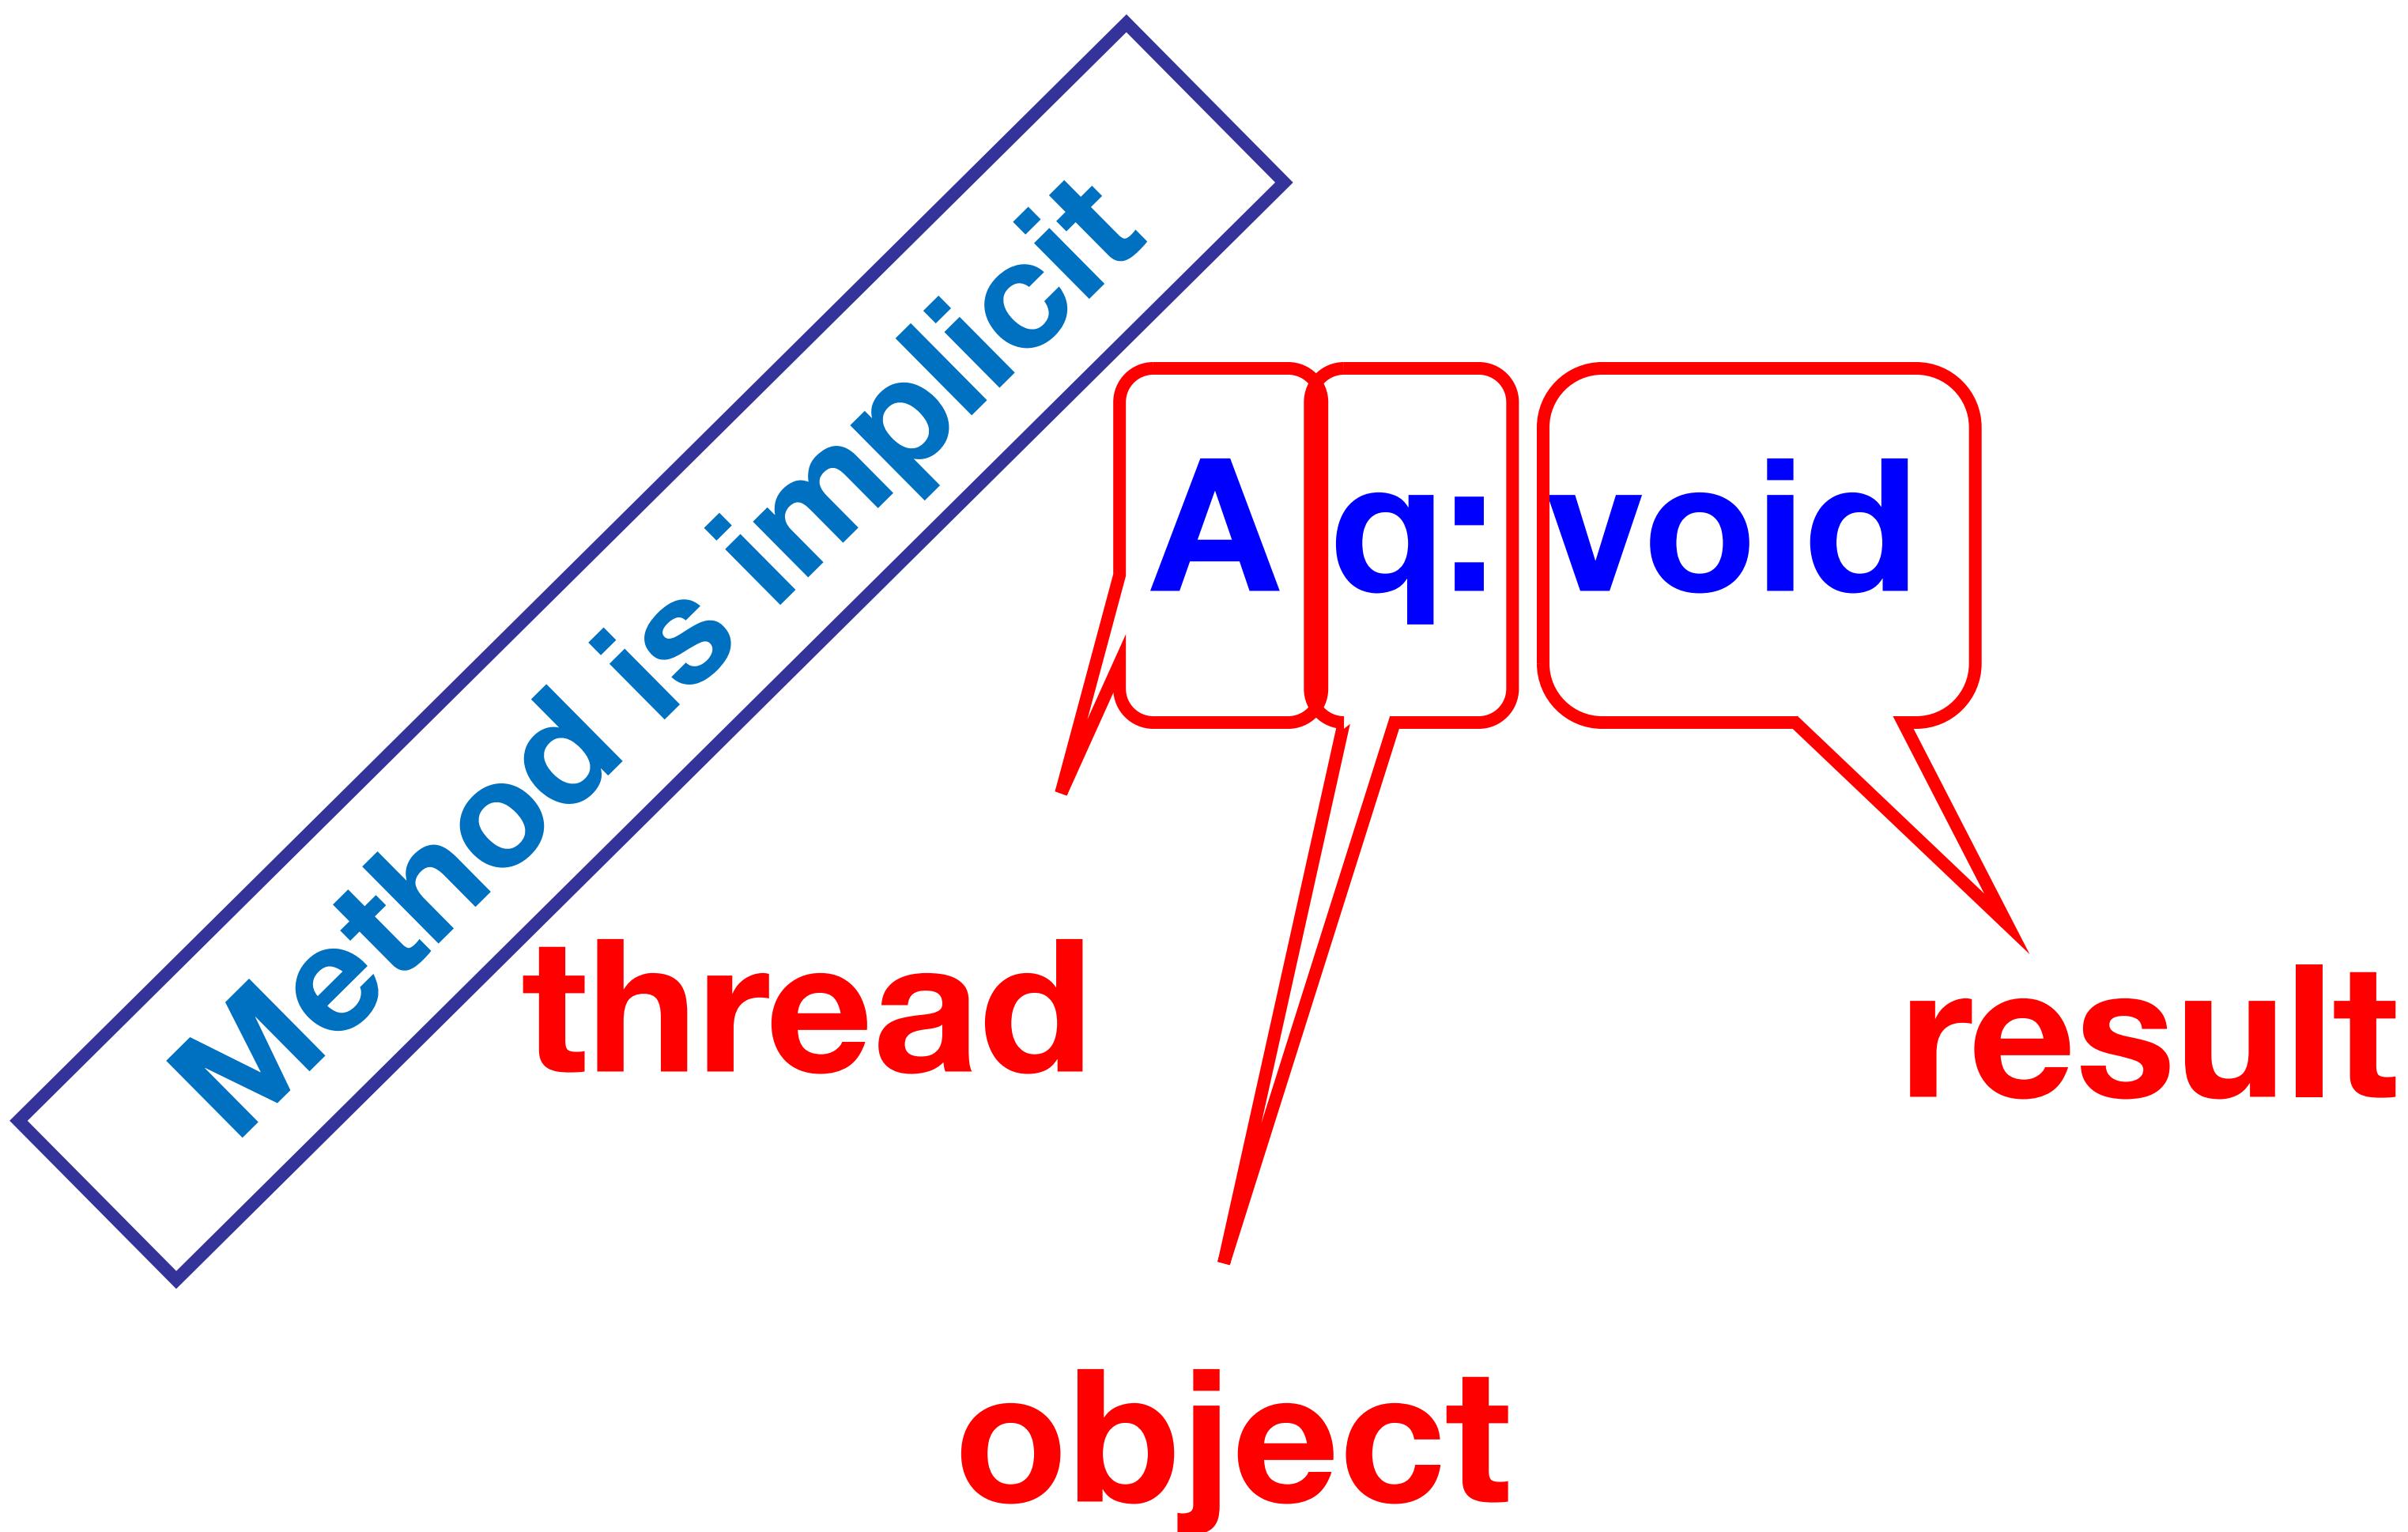
# Response Notation



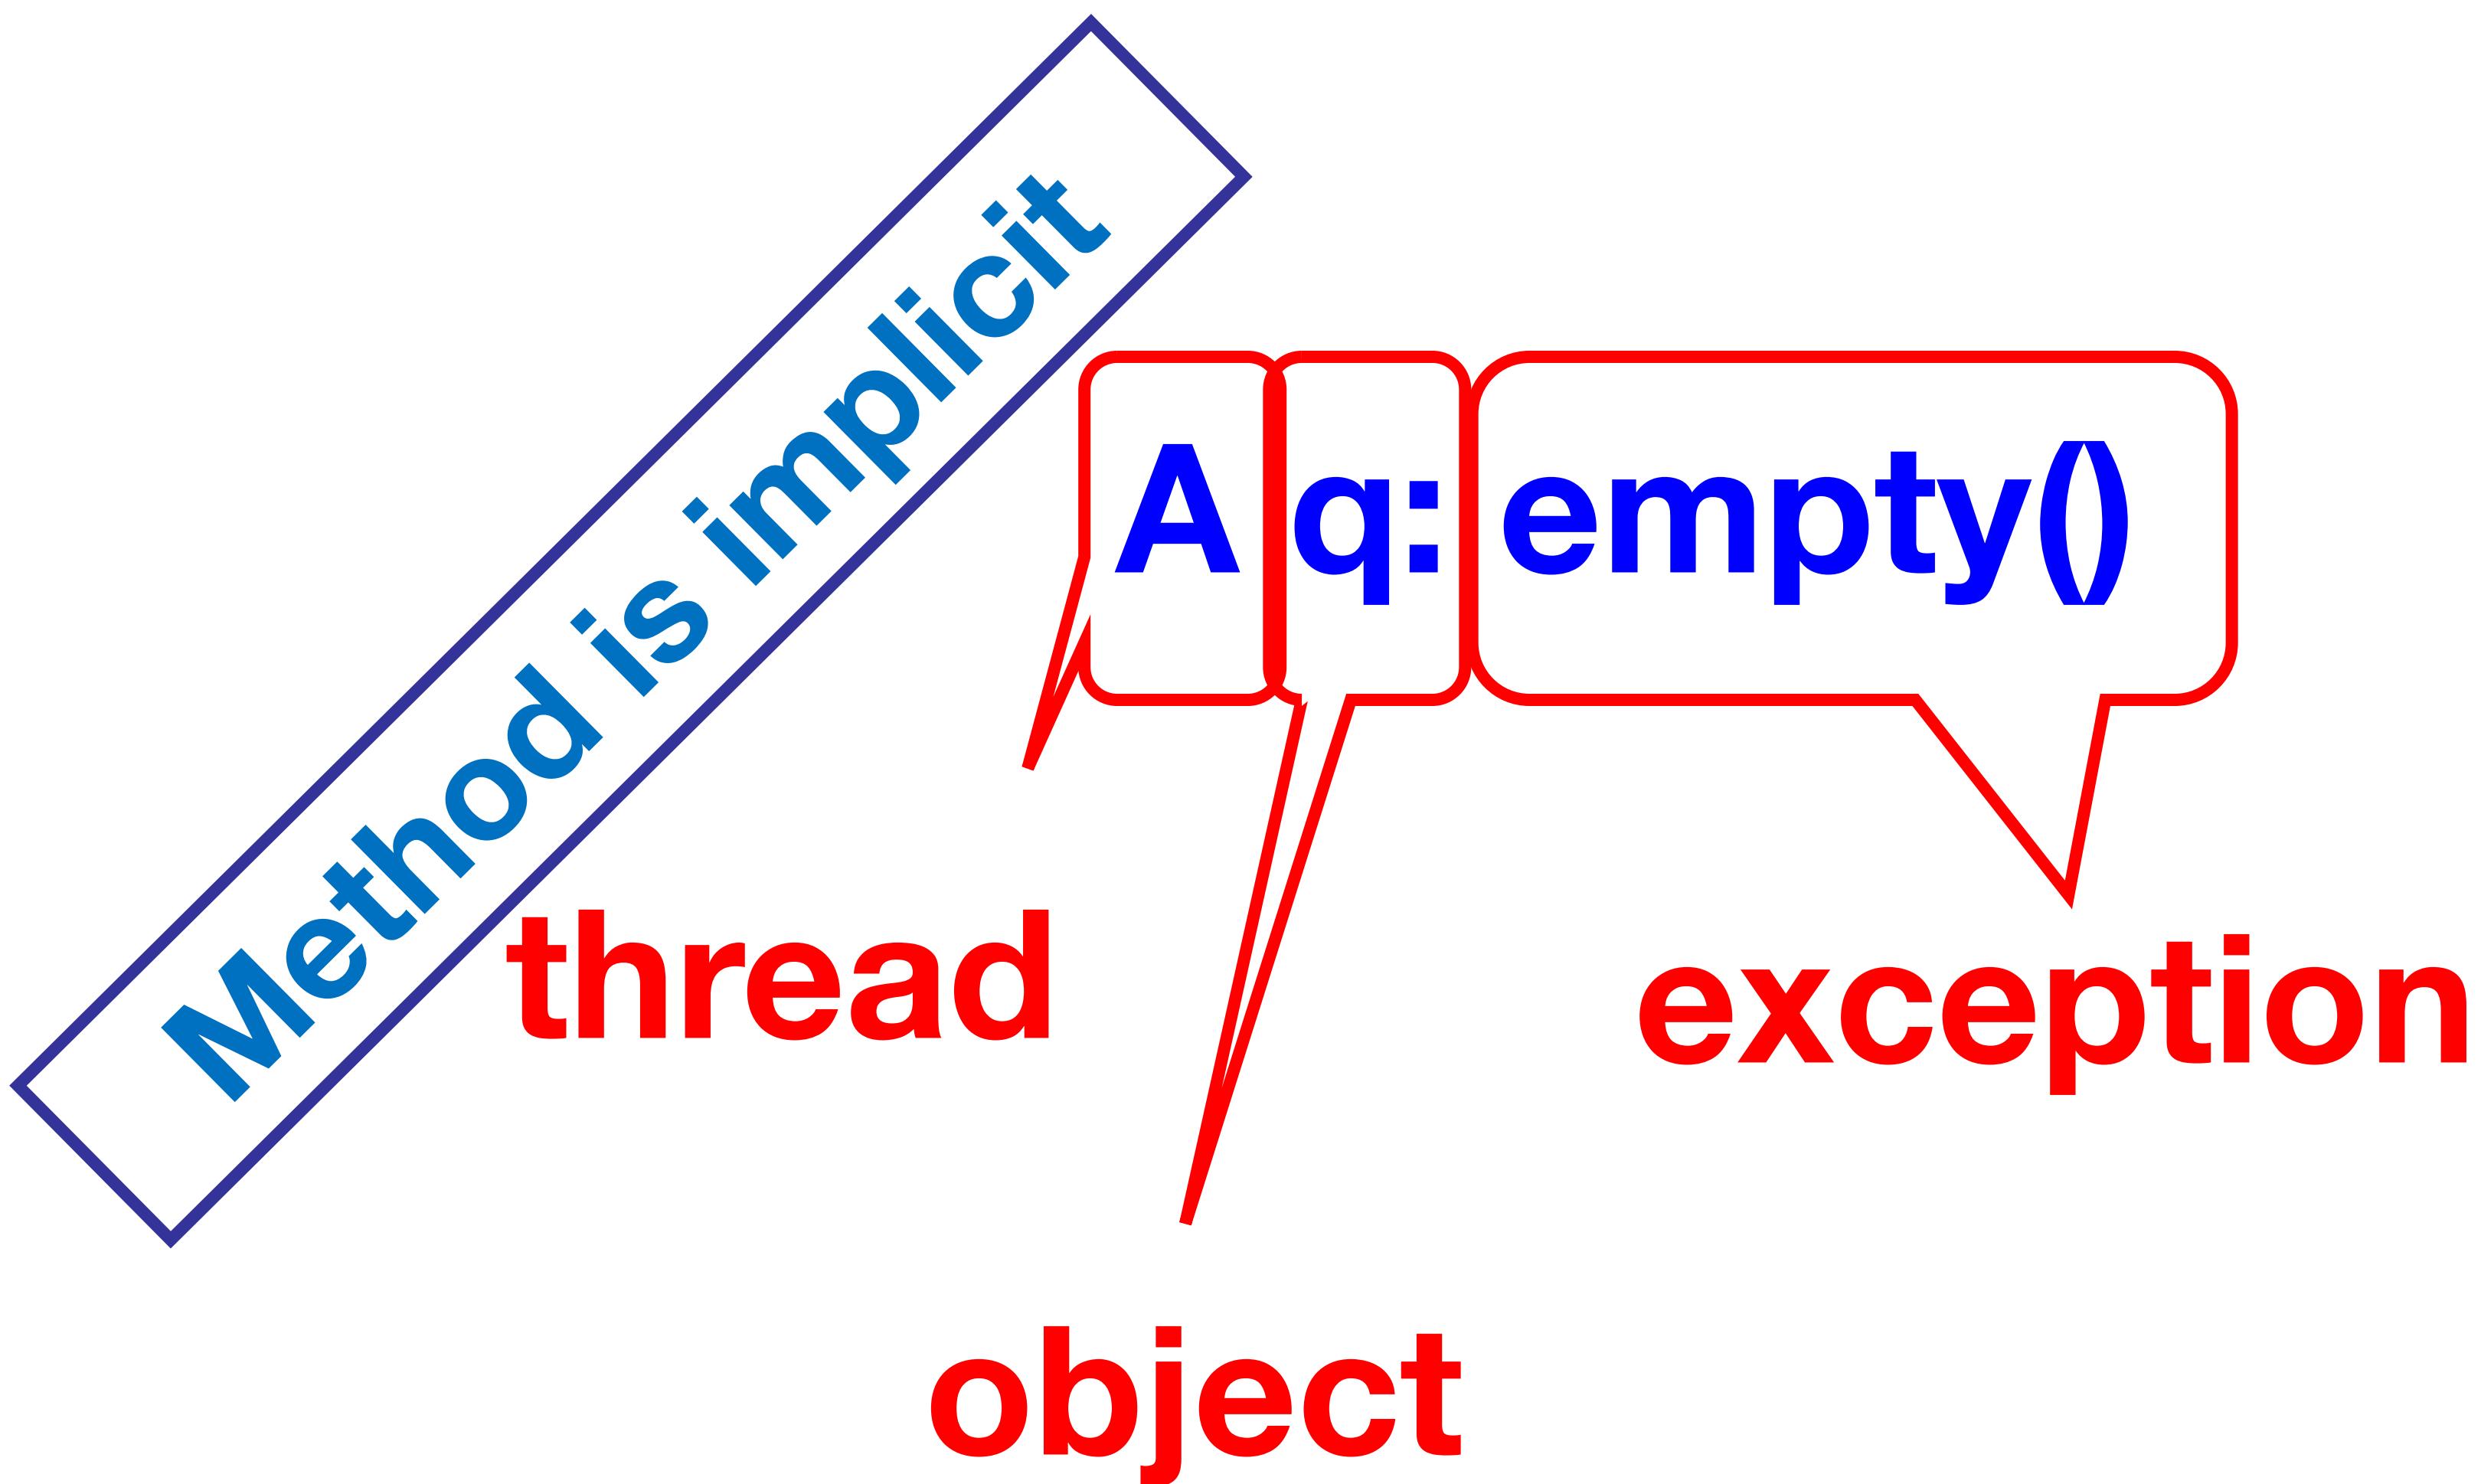
# Response Notation



# Response Notation



# Response Notation



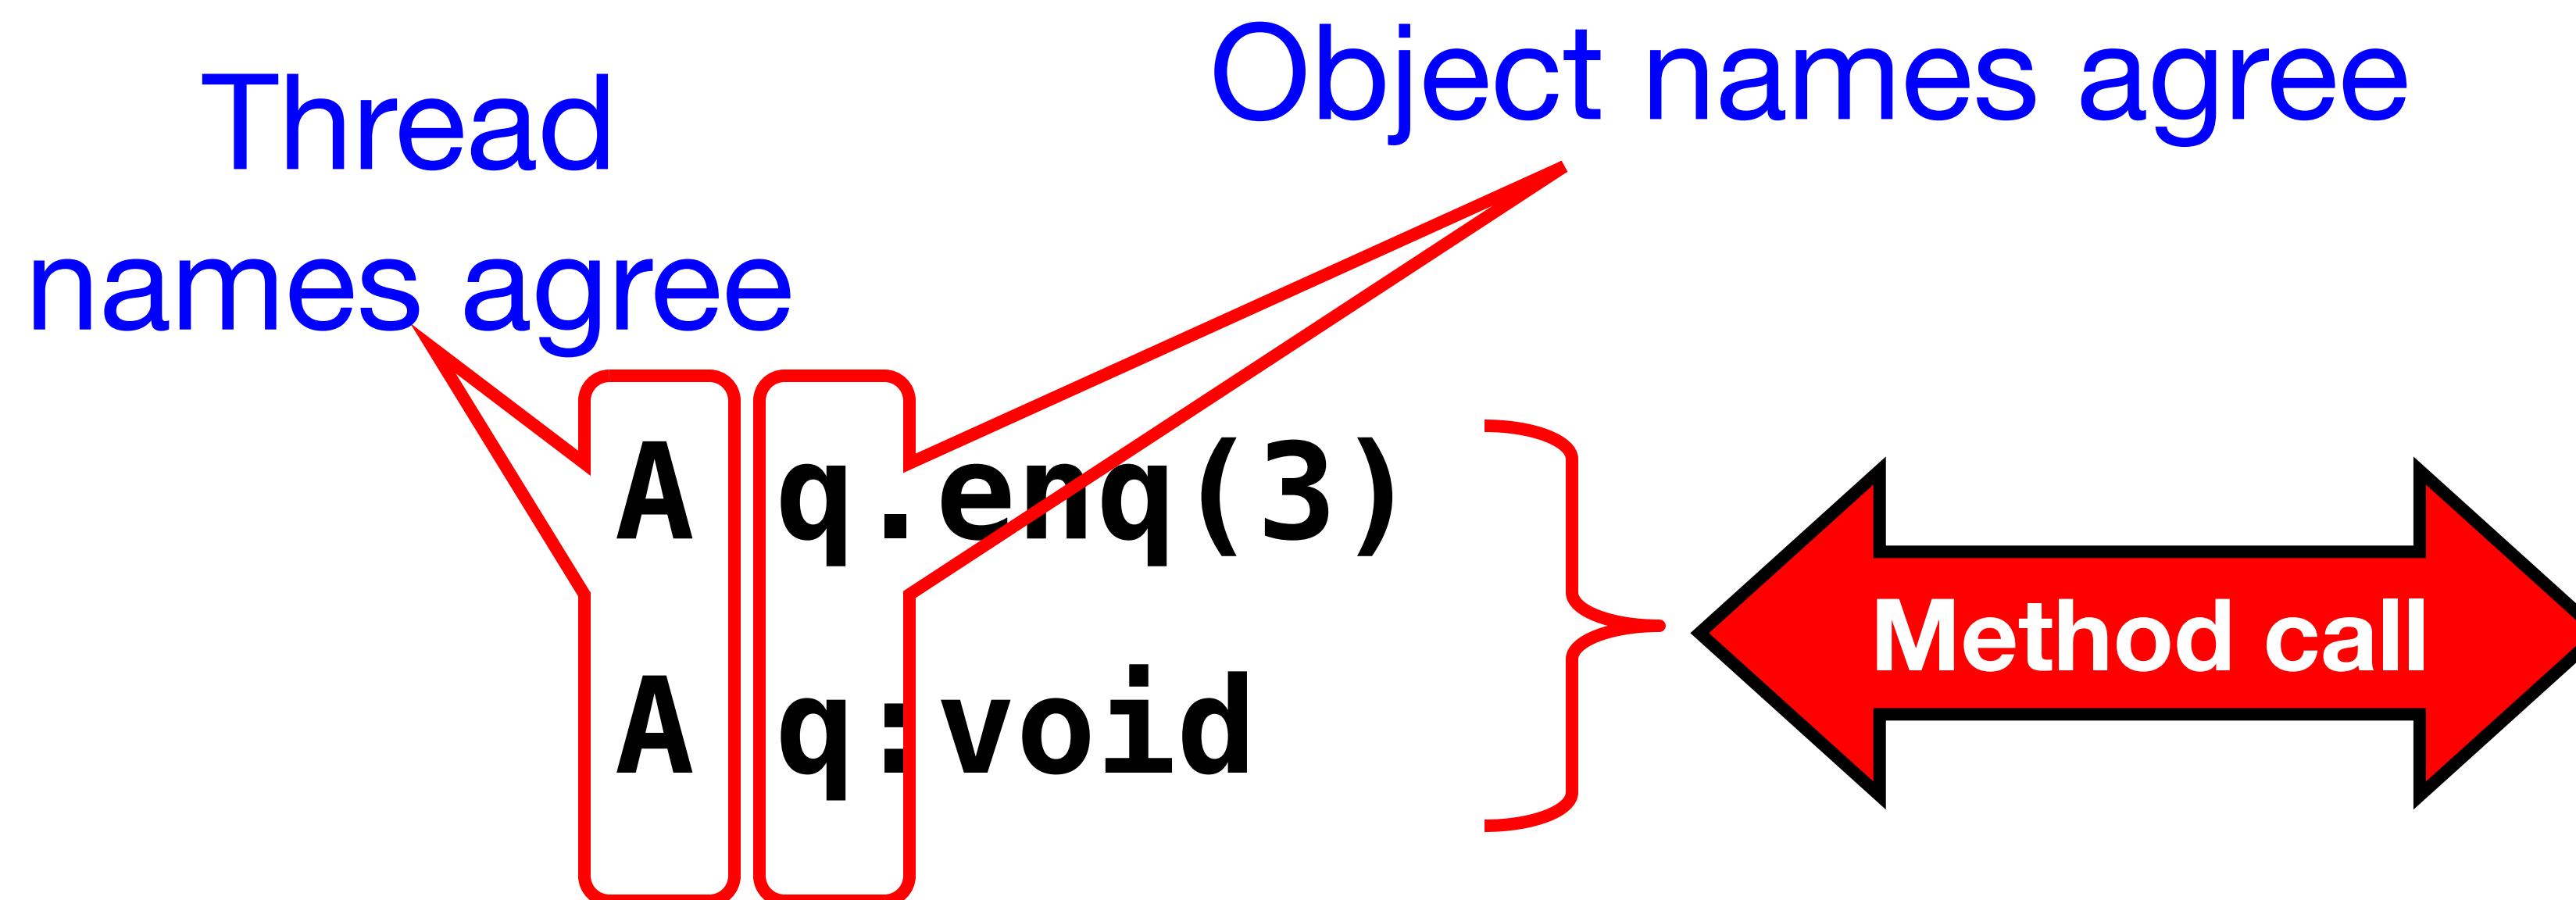
# History – Describing an execution

A q.enq(3)  
A q:void  
A q.enq(5)  
**H =** B p.enq(4)  
B p:void  
B q.deq()  
B q:3

**Sequence of  
invocations and  
responses**

# History – Describing an execution

- Invocation & response *match* if



# Object Projections

$H =$

A	q.enq(3)
A	q:void
B	p.enq(4)
B	p:void
B	q.deq()
B	q:3

# Object Projections

A q.enq(3)  
A q:void

H|q =

B q.deq()  
B q:3

# Thread Projections

$H =$

A	q . enq(3)
A	q : void
B	p . enq(4)
B	p : void
B	q . deq()
B	q : 3

# Thread Projections

$$H|B = \begin{array}{l} B \ p.\text{enq}(4) \\ B \ p:\text{void} \\ B \ q.\text{deq}() \\ B \ q:3 \end{array}$$

# Complete Subhistory

$H =$

A q.enq(3)
A q:void
<b>A q.enq(5)</b>
<del>B p.enq(4)</del>
B p:void
B q.deq()
B q:3

An invocation is  
*pending* if it has no  
matching response

# Complete Subhistory

$H =$

A q.enq(3)	
A q:void	
A q.enq(5)	
B p.enq(4)	
B p:void	
B q.deq()	
B q:3	

**May or may not have taken effect**

# Complete Subhistory

$H =$

A q.enq(3)	
A q:void	
A q.enq(5)	
B p.enq(4)	
B p:void	
B q.deq()	
B q:3	

may discard  
pending invocations

# Complete Subhistory

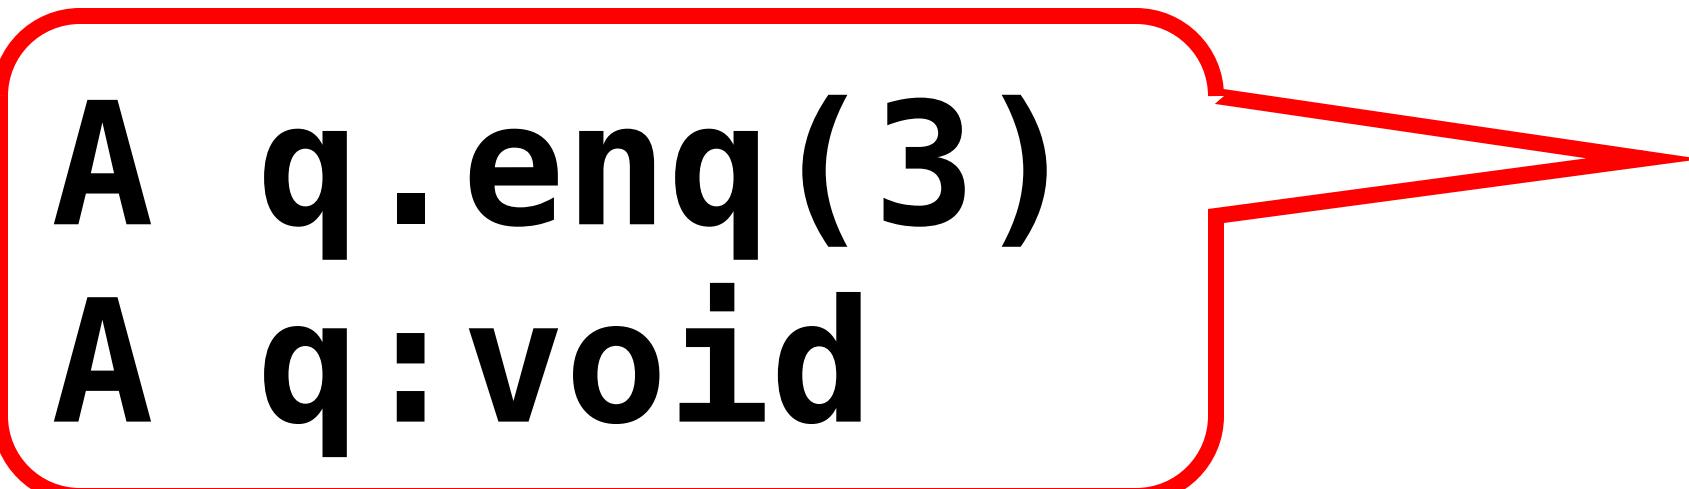
A q.enq(3)  
A q:void

**Complete(H) =** B p.enq(4)  
B p:void  
B q.deq()  
B q:3

# Sequential Histories

```
A q.enq(3)
A q:void
B p.enq(4)
B p:void
B q.deq()
B q:3
A q:enq(5)
```

# Sequential Histories



A q . enq (3)  
A q : void

B p . enq (4)  
B p : void  
B q . deq ()  
B q : 3  
A q : enq (5)

match

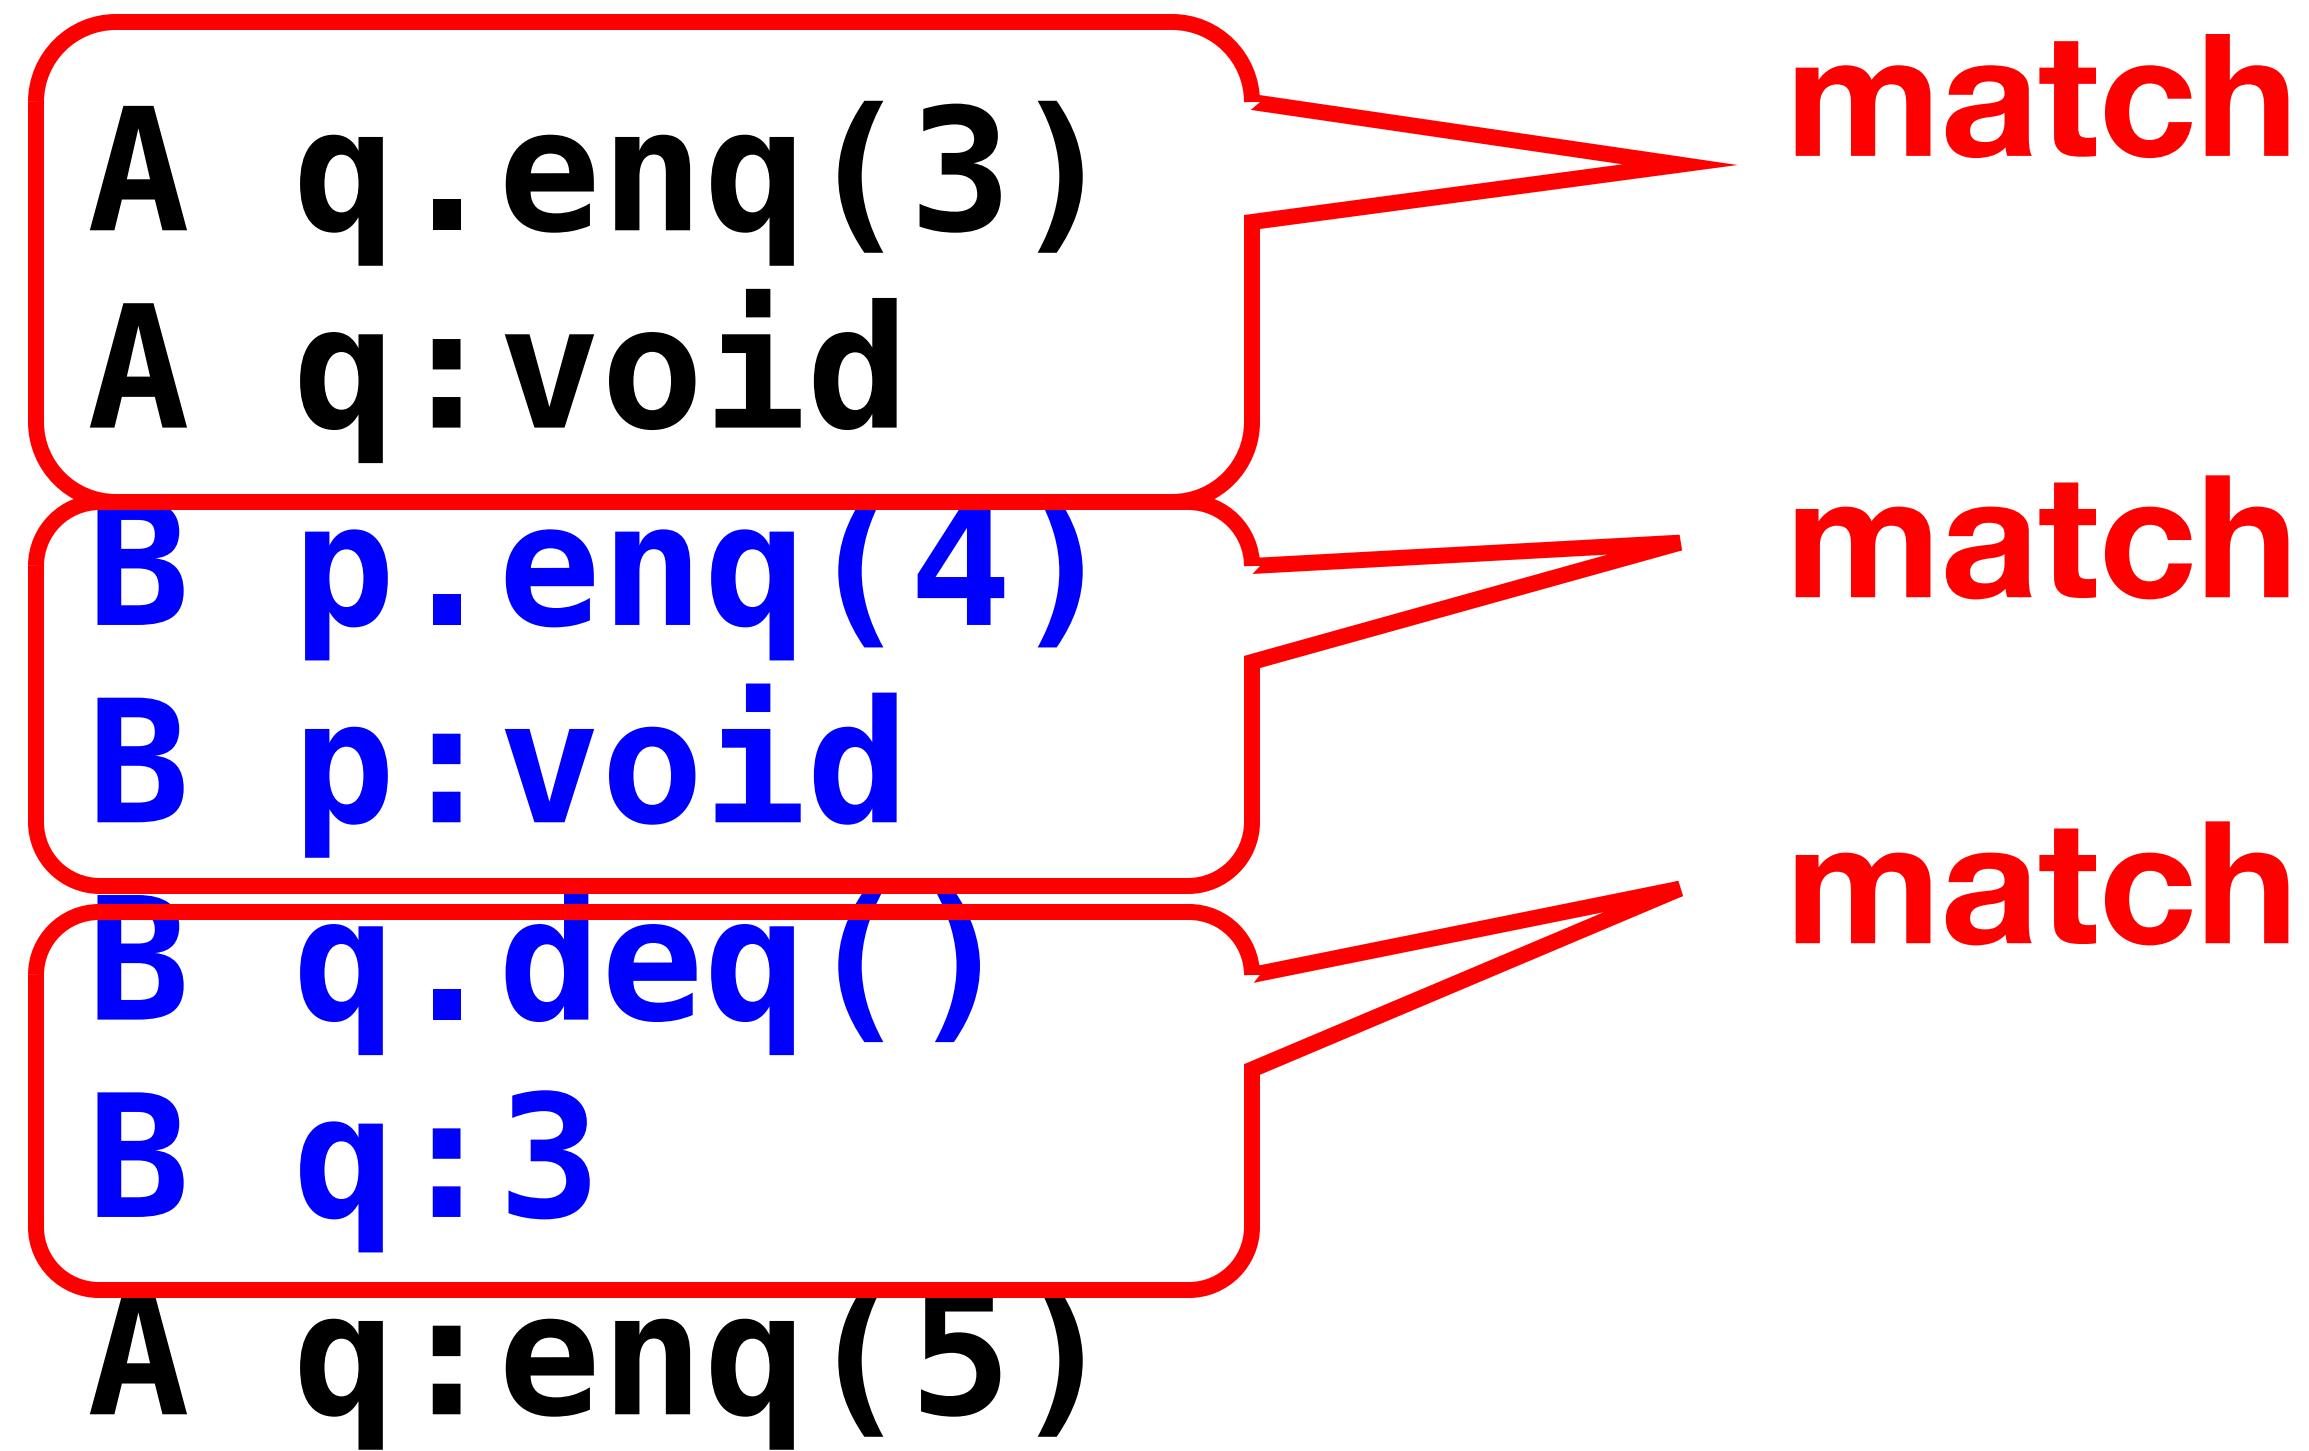
# Sequential Histories

```
A q.enq(3)
A q:void
B p.enq(4)
B p:void
B q.deq()
B q:3
A q:enq(5)
```

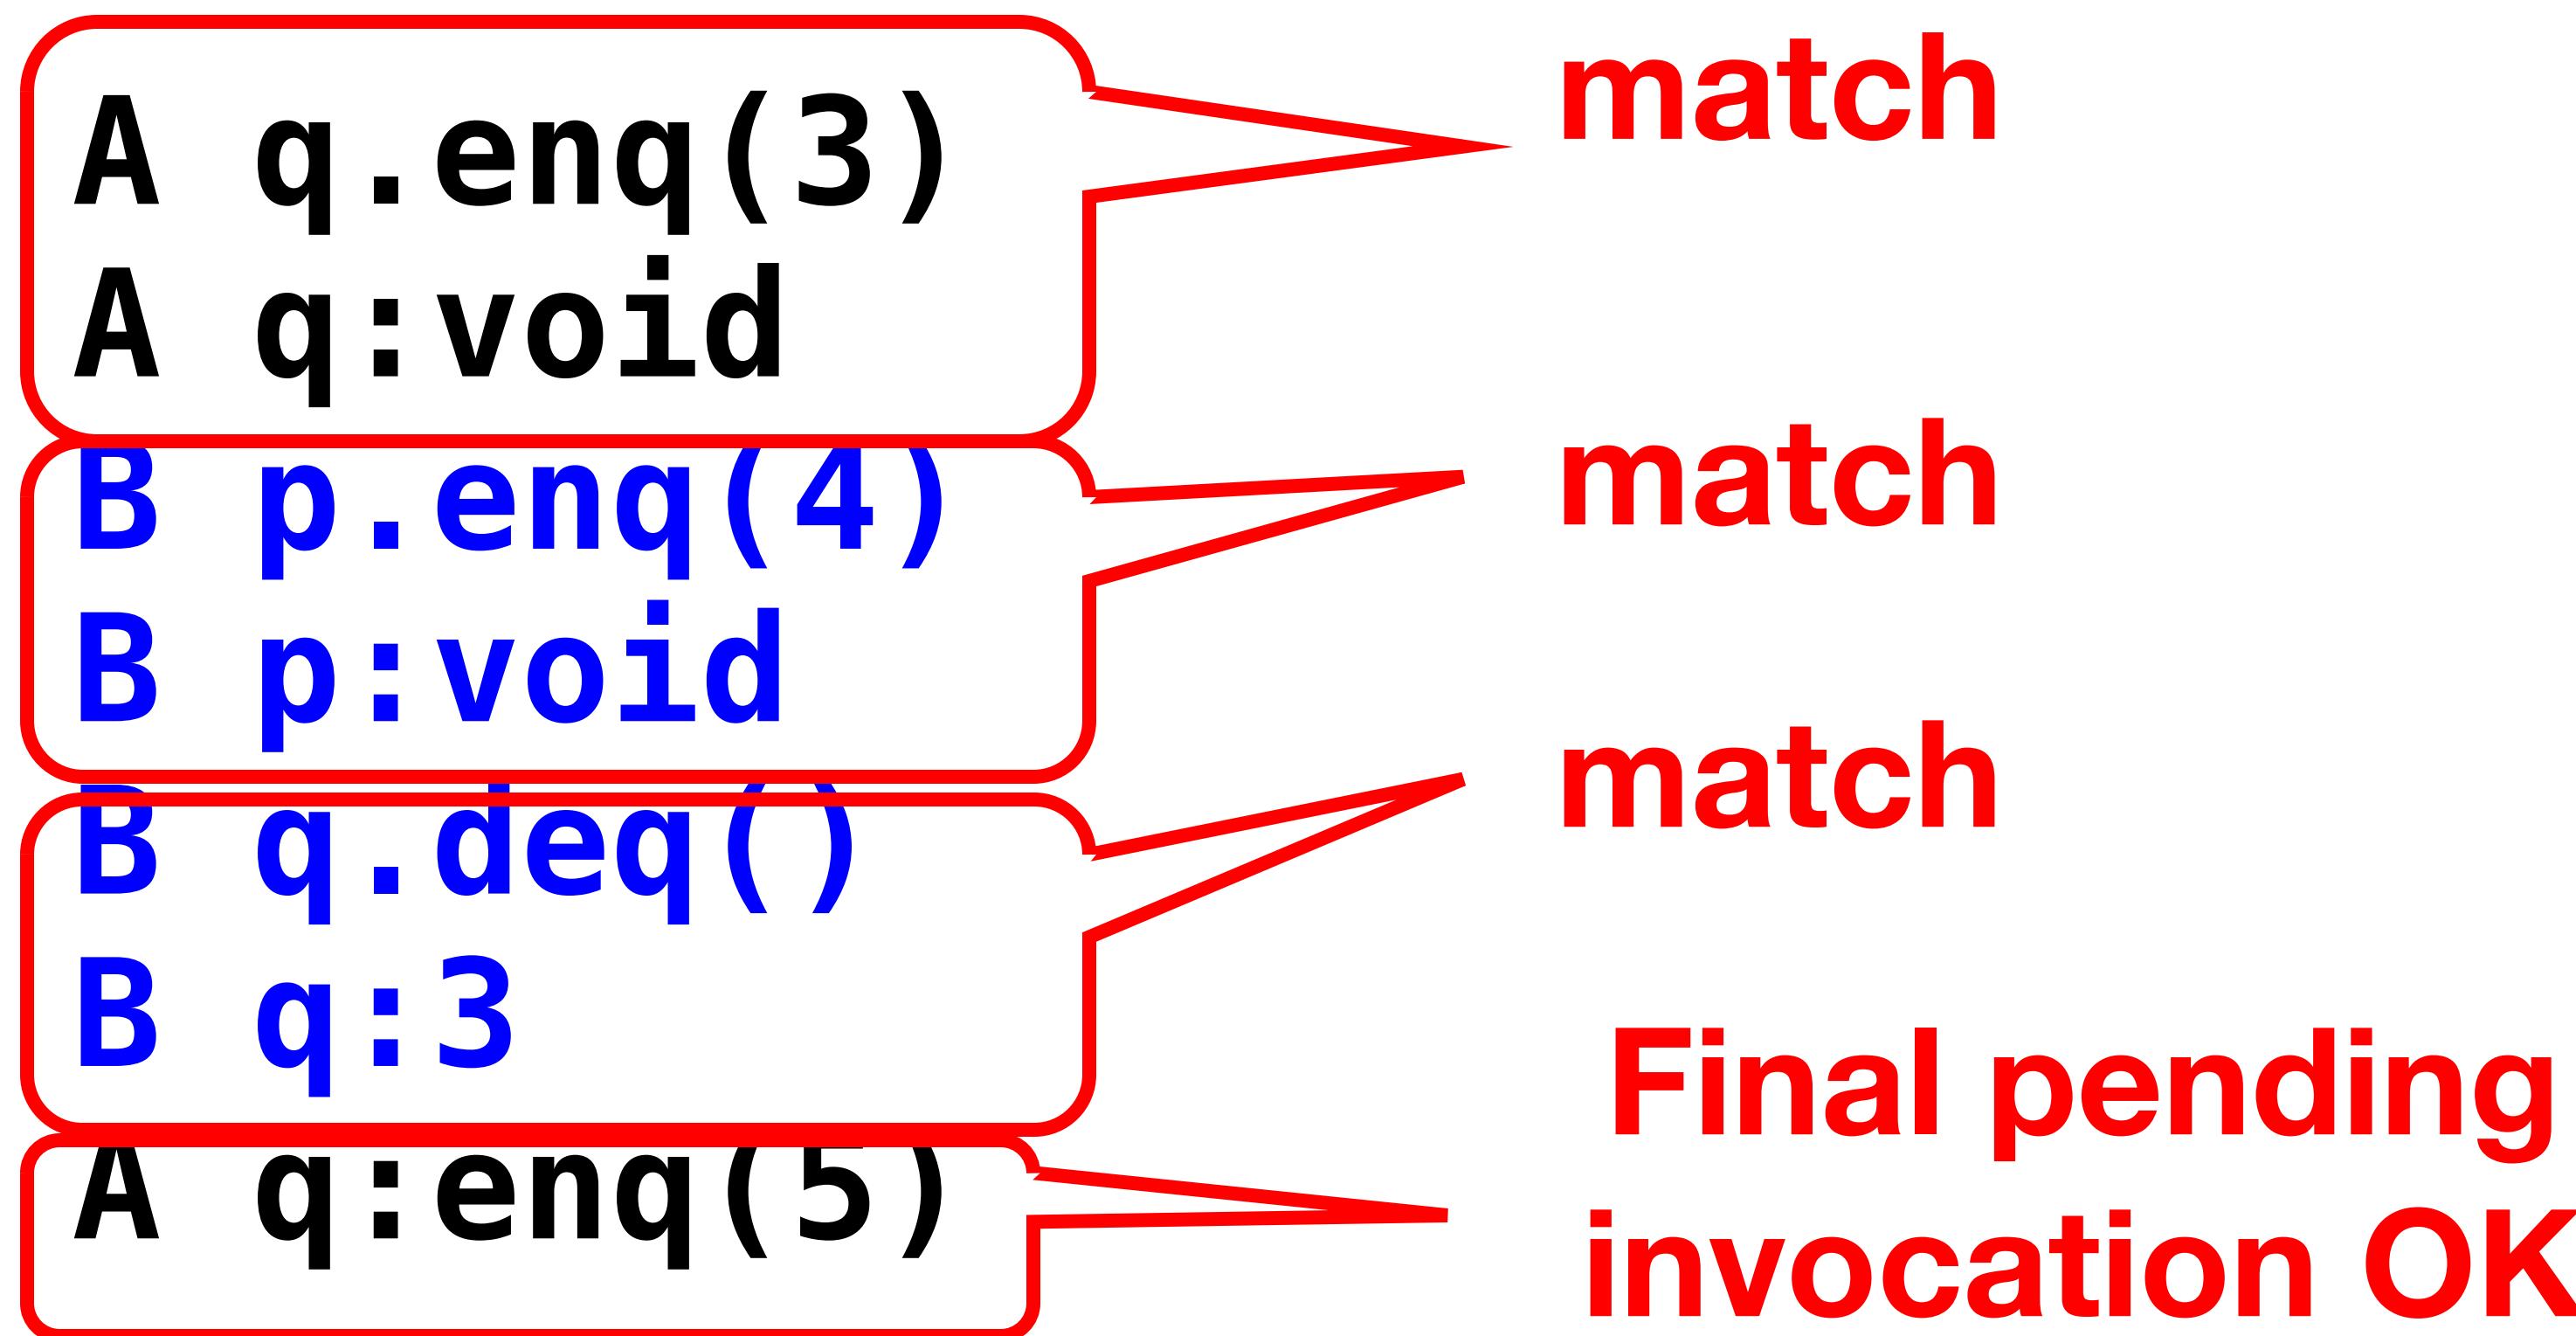
match

match

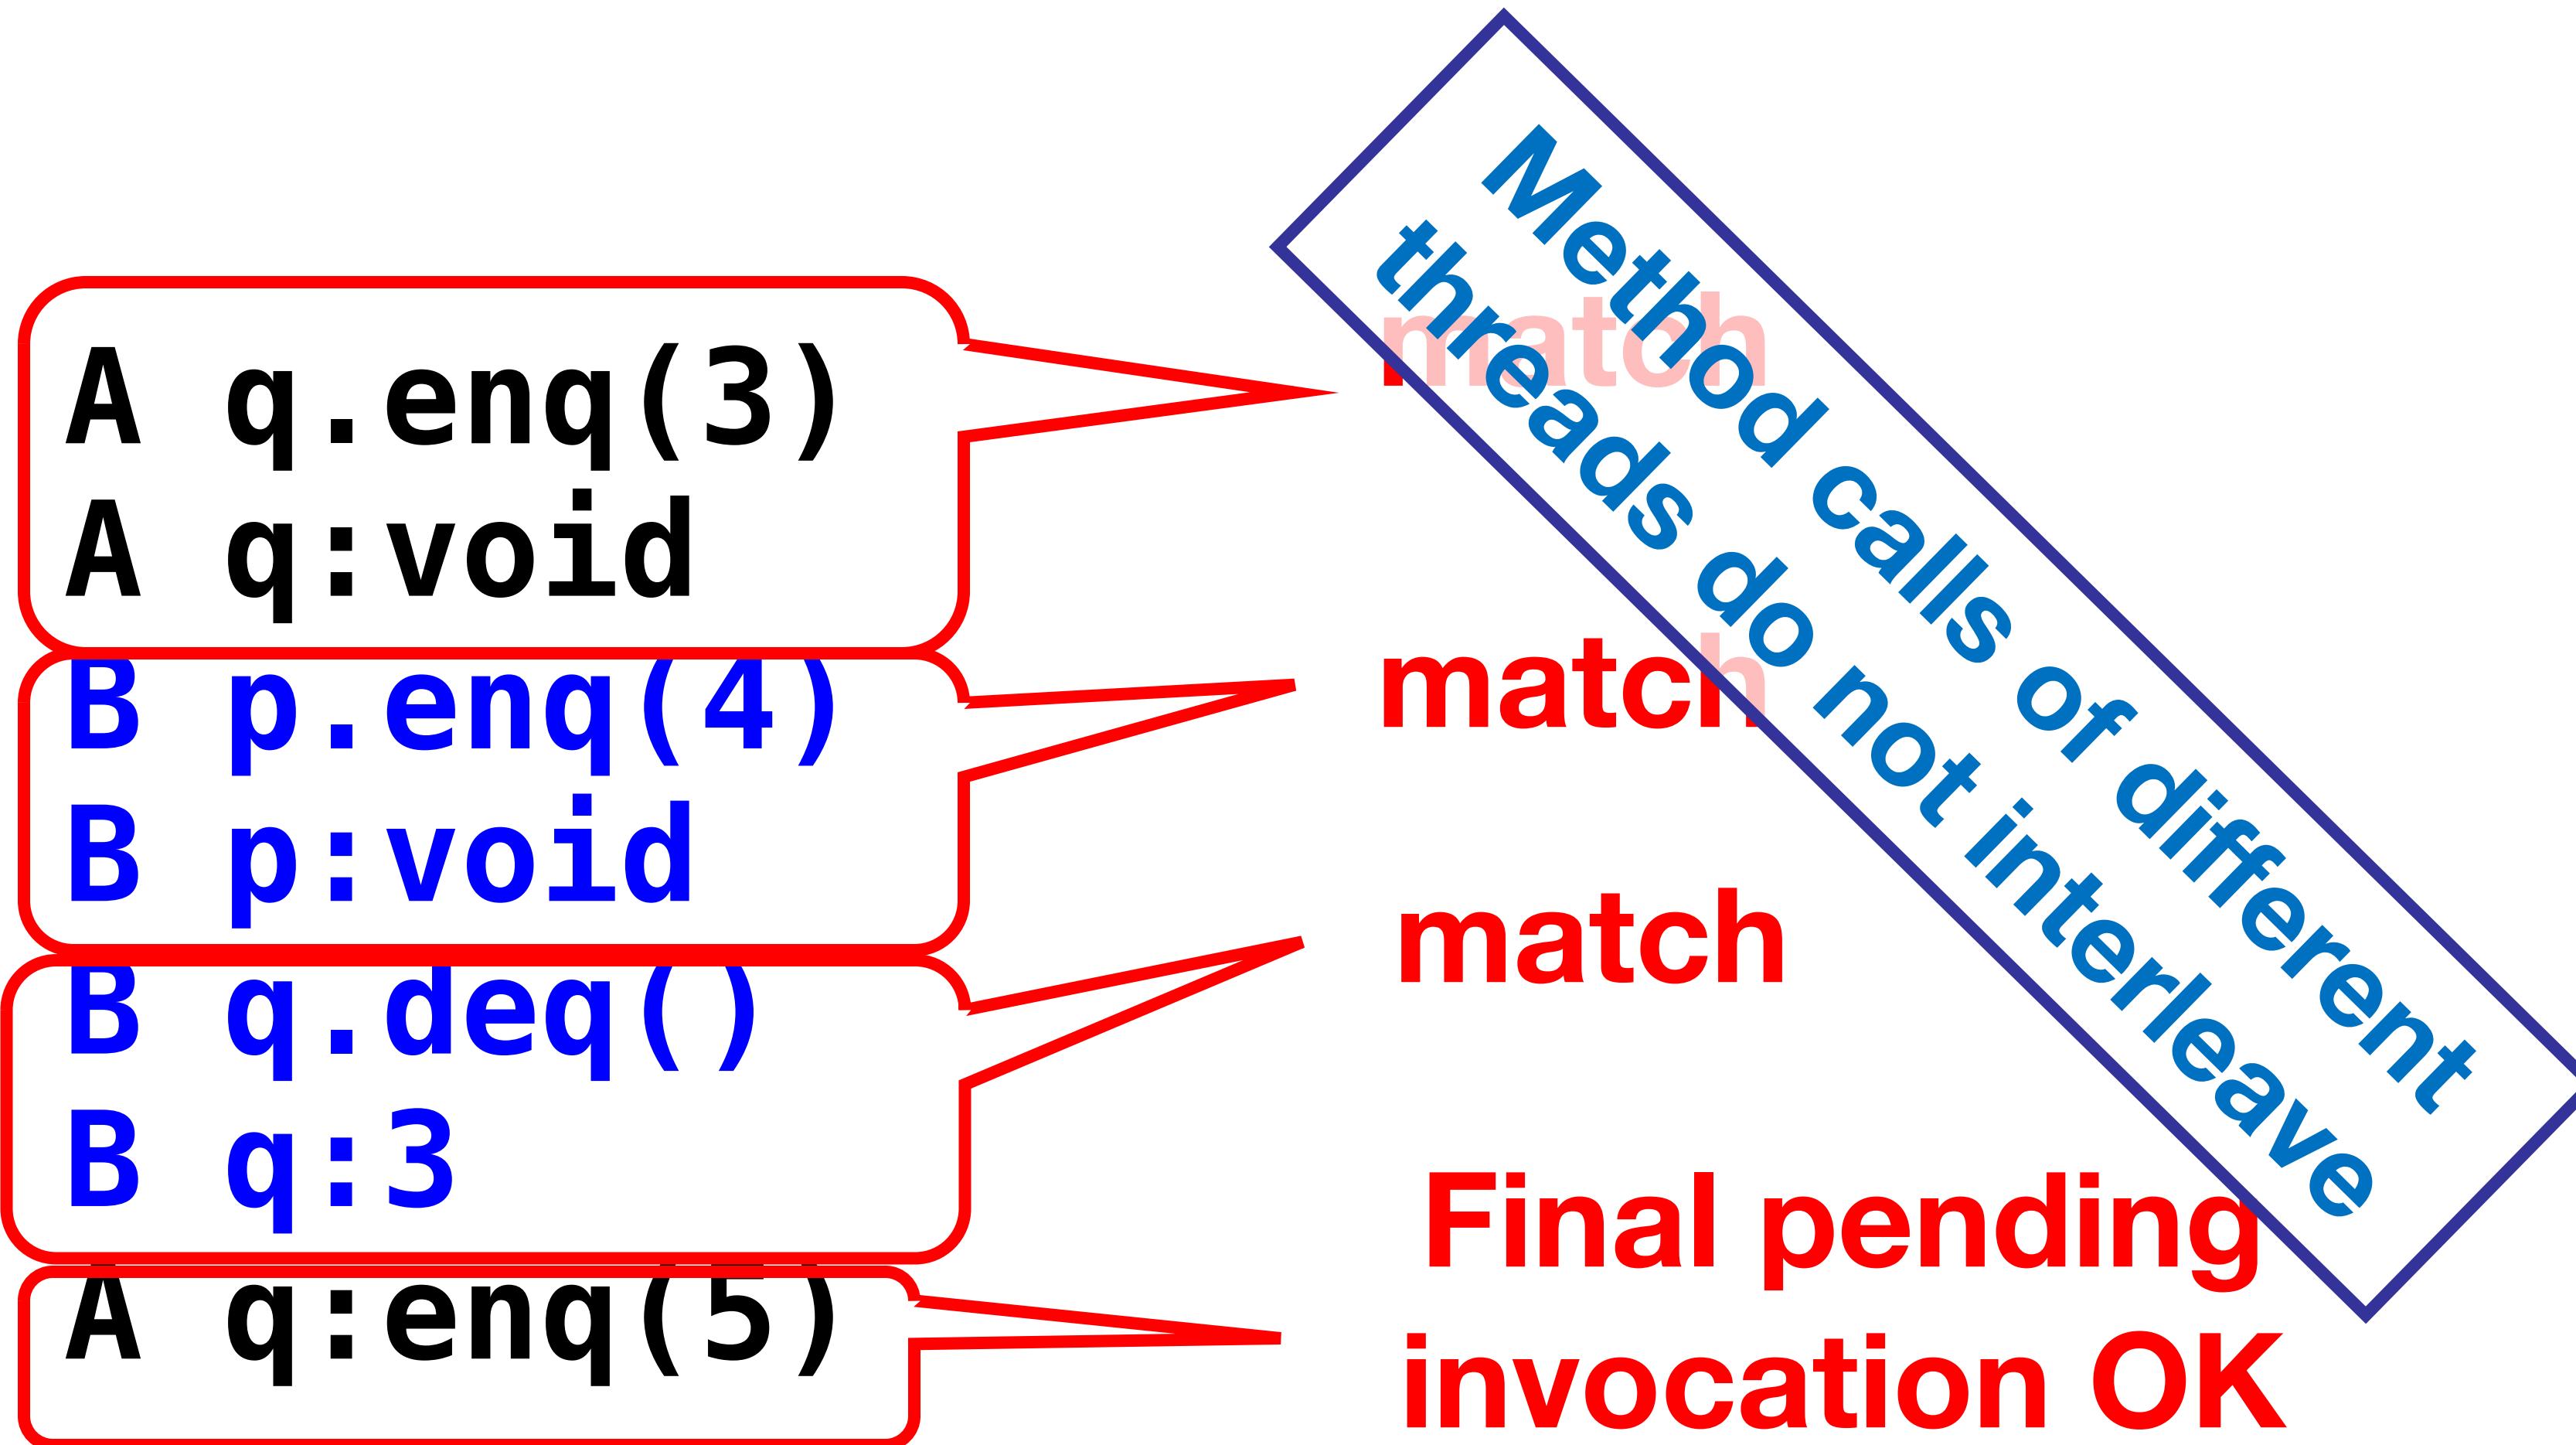
# Sequential Histories



# Sequential Histories



# Sequential Histories



# Well-formed Histories

A q.enq(3)  
B p.enq(4)  
B p:void  
H= B q.deq()  
A q:void  
B q:3

# Well-formed Histories

Per-thread projections  
sequential

H= A q.enq(3)  
B p.enq(4)  
B p:void  
B q.deq()  
A q:void  
B q:3

H | B= B p.enq(4)  
B p:void  
B q.deq()  
B q:3

# Well-formed Histories

Per-thread projections  
sequential

$H =$

- A q.enq(3)
- B p.enq(4)
- B p:void
- B q.deq()
- A q:void
- B q:3

$H | B =$

- B p.enq(4)
- B p:void
- B q.deq()
- B q:3

$H | A =$

- A q.enq(3)
- A q:void

# Equivalent Histories

Threads see the same  
thing in both

$$\left\{ \begin{array}{l} H|A = G|A \\ H|B = G|B \end{array} \right.$$

$H =$

```
A q.enq(3)
B p.enq(4)
B p:void
B q.deq()
A q:void
B q:3
```

$G =$

```
A q.enq(3)
A q:void
B p.enq(4)
B p:void
B q.deq()
B q:3
```

# Sequential Specifications

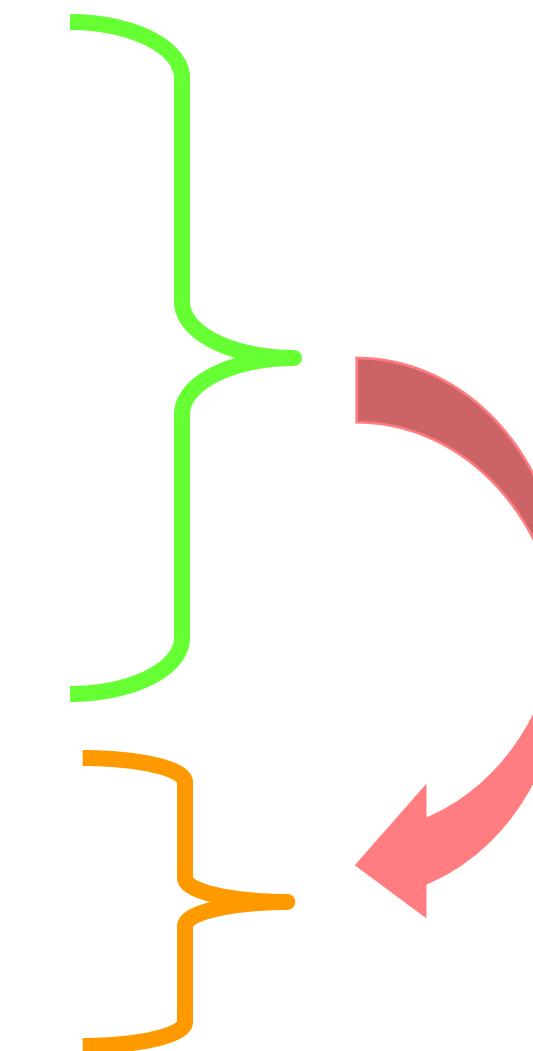
- A sequential ***specification*** is some way of telling whether a
  - Single-thread, single-object history
  - Is ***legal***
- For example:
  - Pre and post-conditions
  - But plenty of other techniques exist ...

# Legal Histories

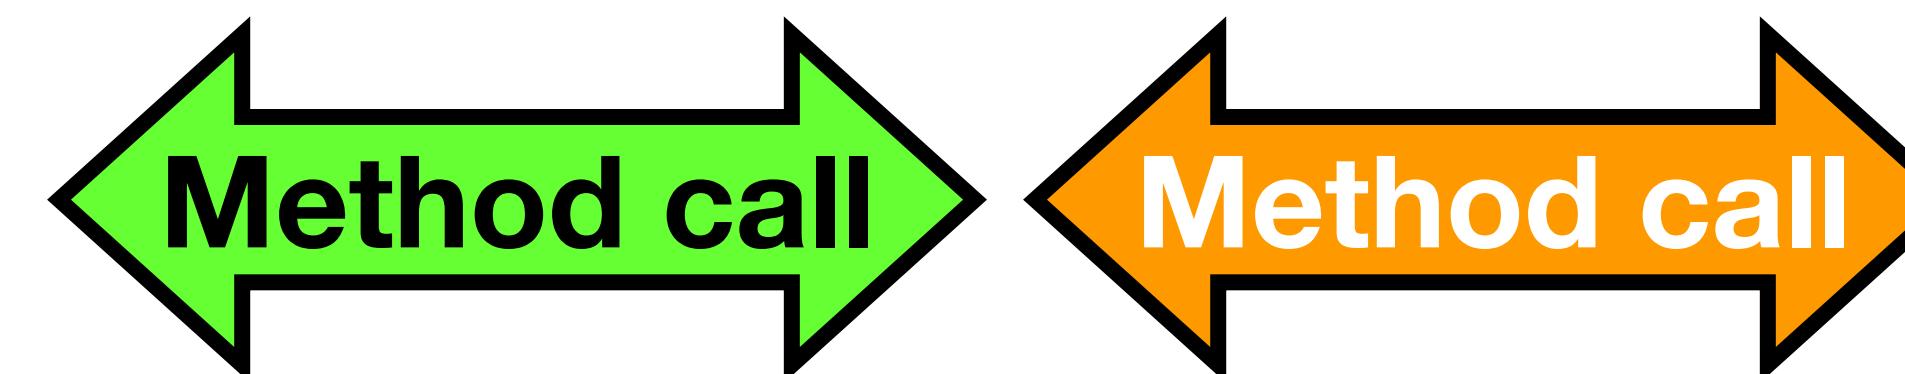
- A sequential (*multi-object*) history  $H$  is *legal* if
  - For every object  $x$
  - $H|x$  is in the sequential spec for  $x$

# Precedence

A q.enq(3)  
B p.enq(4)  
B p.void  
A q:void  
B q.deq()  
B q:3

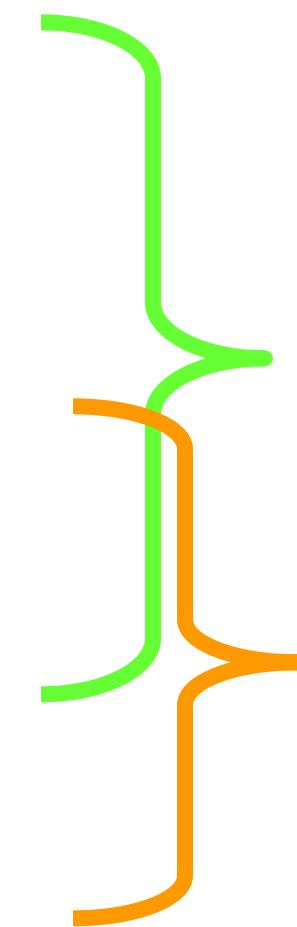


A method call **precedes** another if response event precedes invocation event

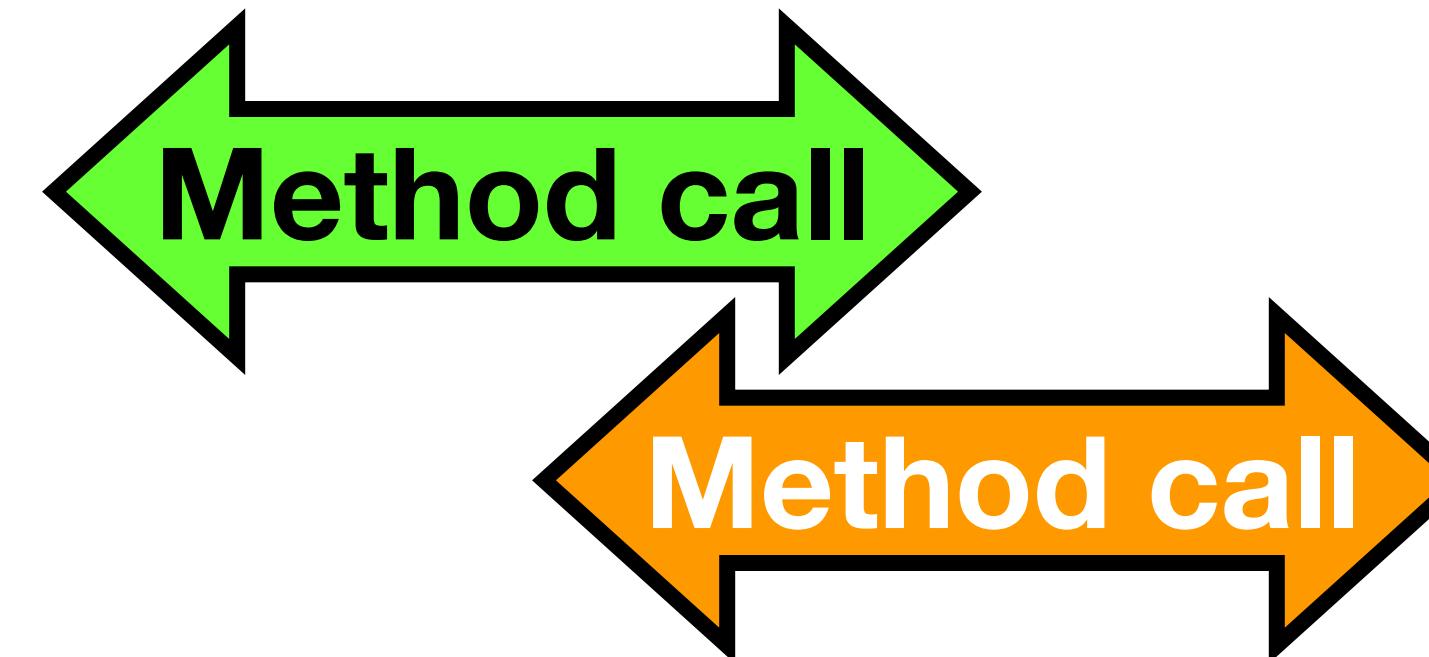


# Non-Precedence

```
A q.enq(3)
B p.enq(4)
B p_void
B q.deq()
A q_void
B q:3
```

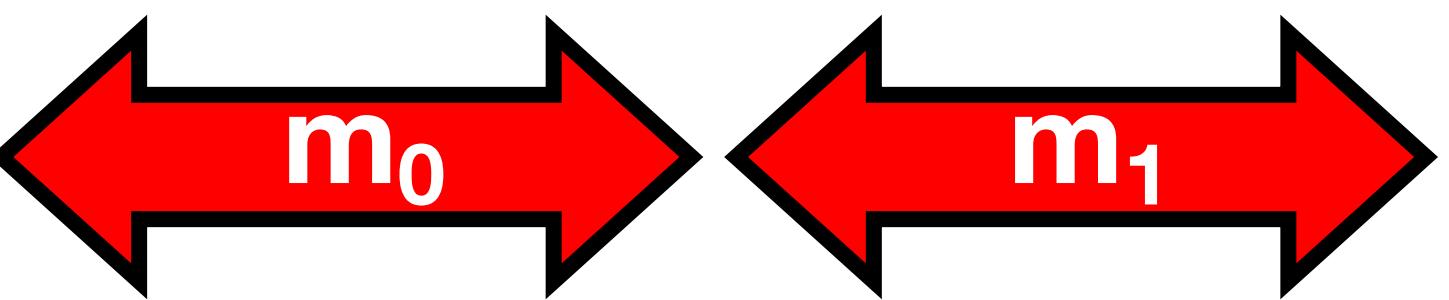


**Some method calls  
overlap one another**



# Notation

- Given
  - History  $\mathbf{H}$
  - method executions  $m_0$  and  $m_1$  in  $\mathbf{H}$
- We say  $m_0 \rightarrow_{\mathbf{H}} m_1$ , if
  - $m_0$  precedes  $m_1$
- Relation  $m_0 \rightarrow_{\mathbf{H}} m_1$  is a
  - Partial order
  - Total order if  $\mathbf{H}$  is sequential



# Linearizability

- History  $H$  is **linearizable** if it can be extended to  $G$  by
  - Appending zero or more responses to pending invocations
  - Discarding other pending invocations
- So that  $G$  is equivalent to
  - Legal sequential history  $S$
  - where  $\rightarrow_G \subset \rightarrow_S$

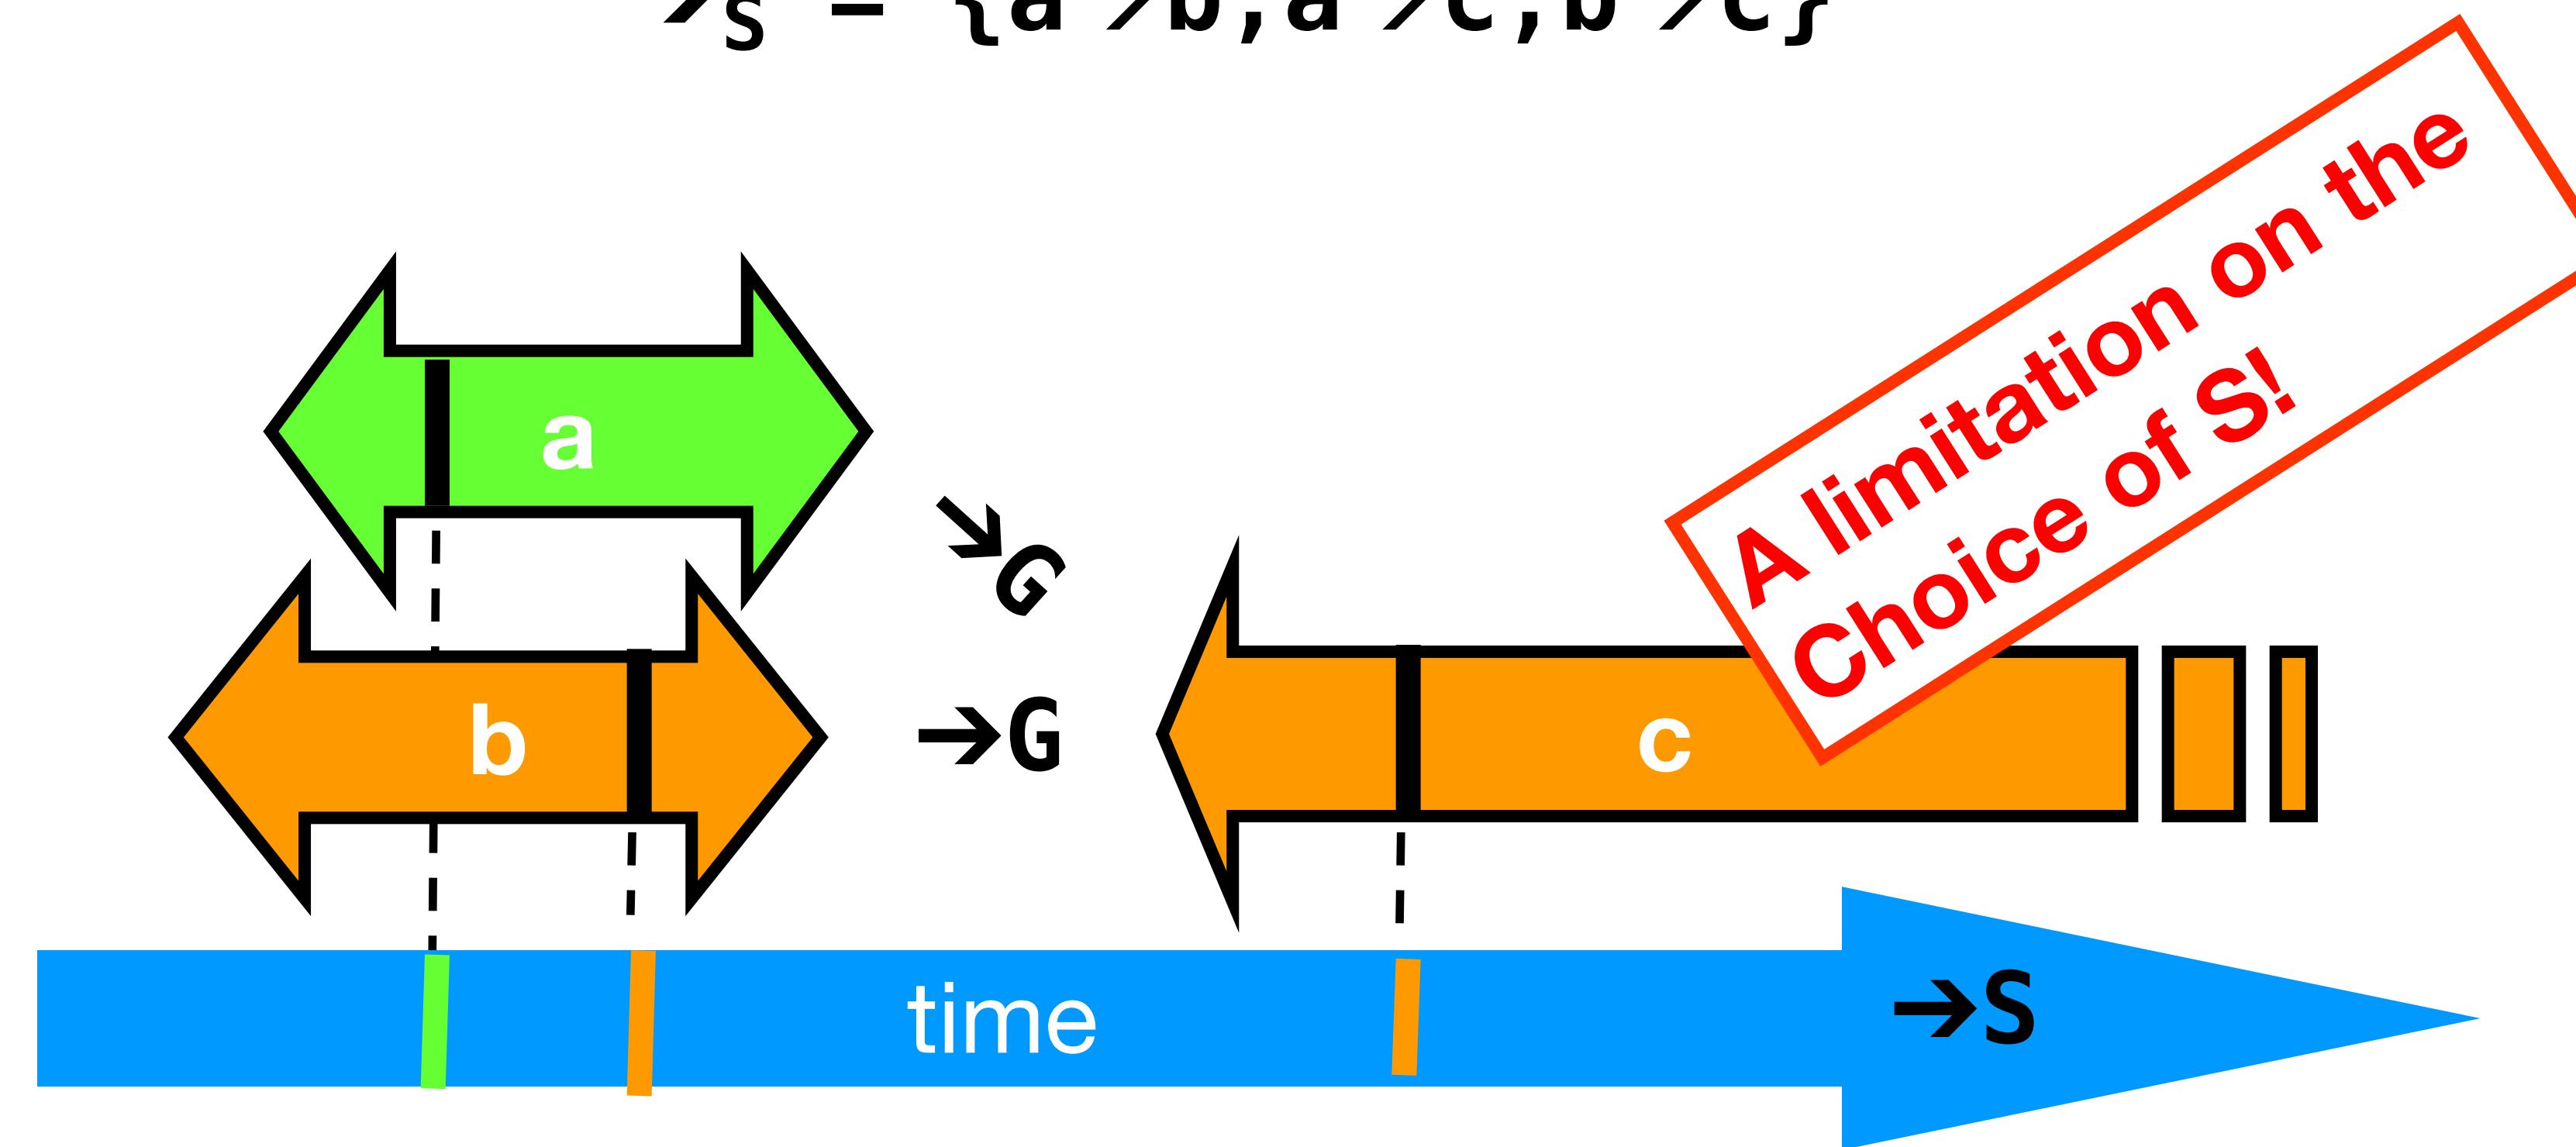
# Remarks on Linearizability

- Some pending invocations
  - Took effect, so keep them
  - Discard the rest
- Condition  $\rightarrow_G \subset \rightarrow_S$ 
  - Means that **S** respects “real-time order” of **G**

# Ensuring $\rightarrow_G \subset \rightarrow_S$

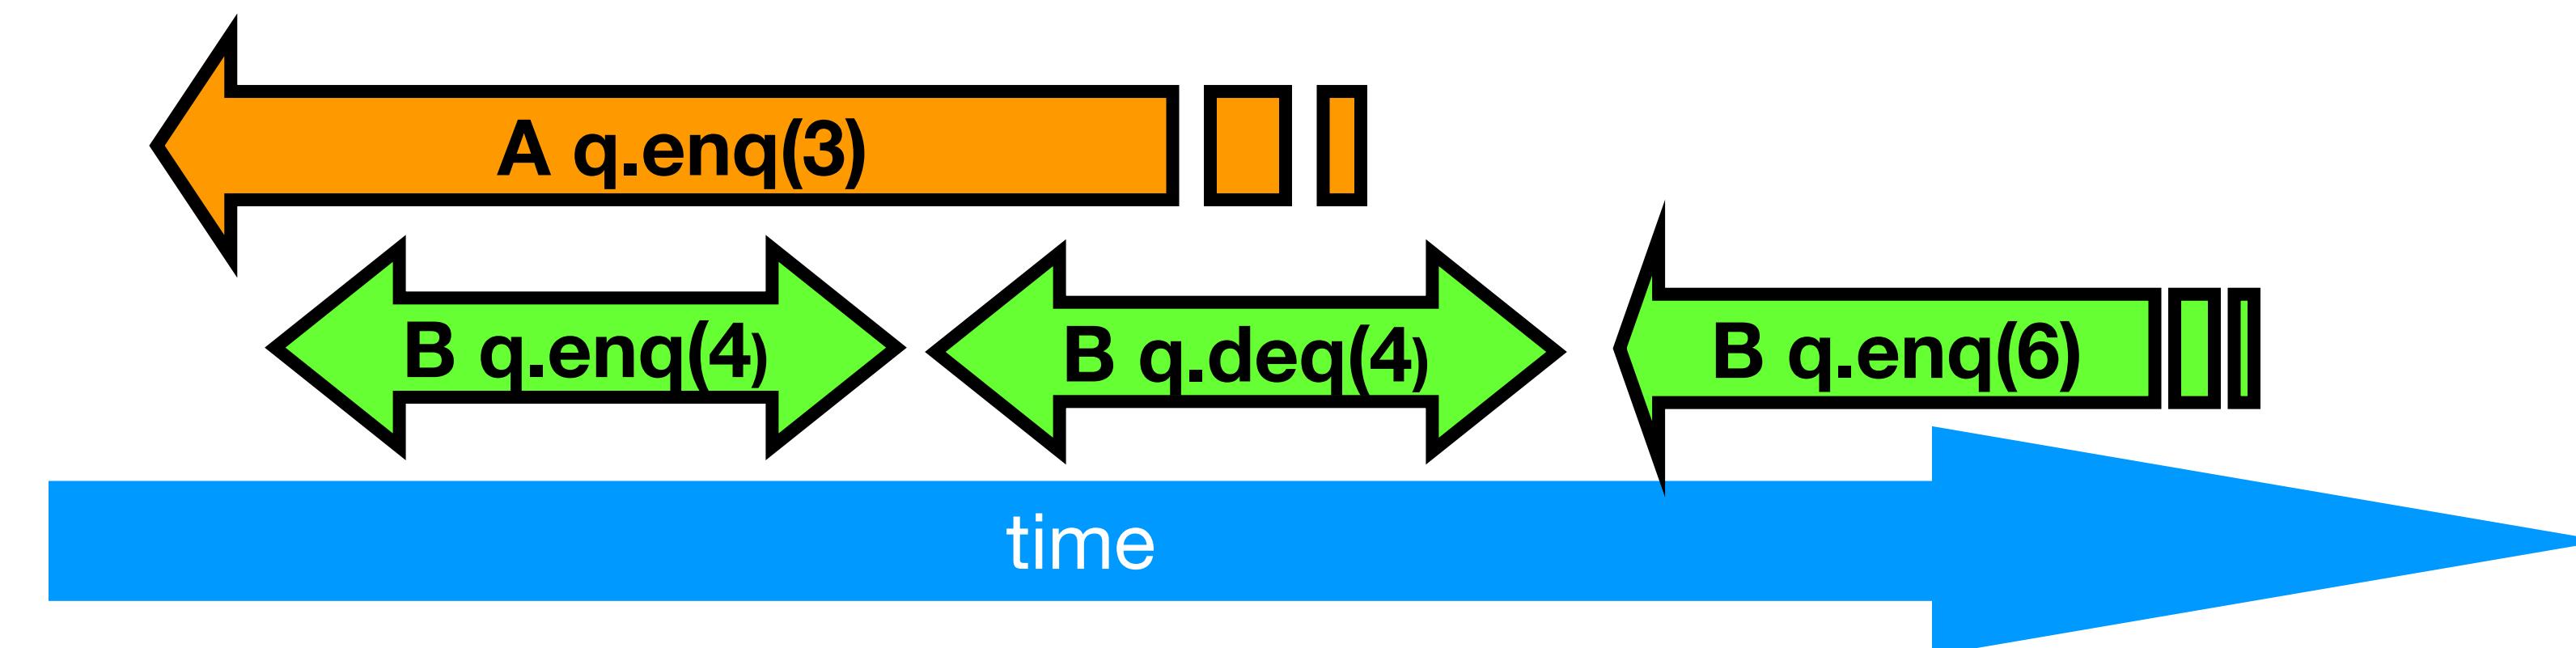
$$\rightarrow_G = \{a \rightarrow c, b \rightarrow c\}$$

$$\rightarrow_S = \{a \rightarrow b, a \rightarrow c, b \rightarrow c\}$$

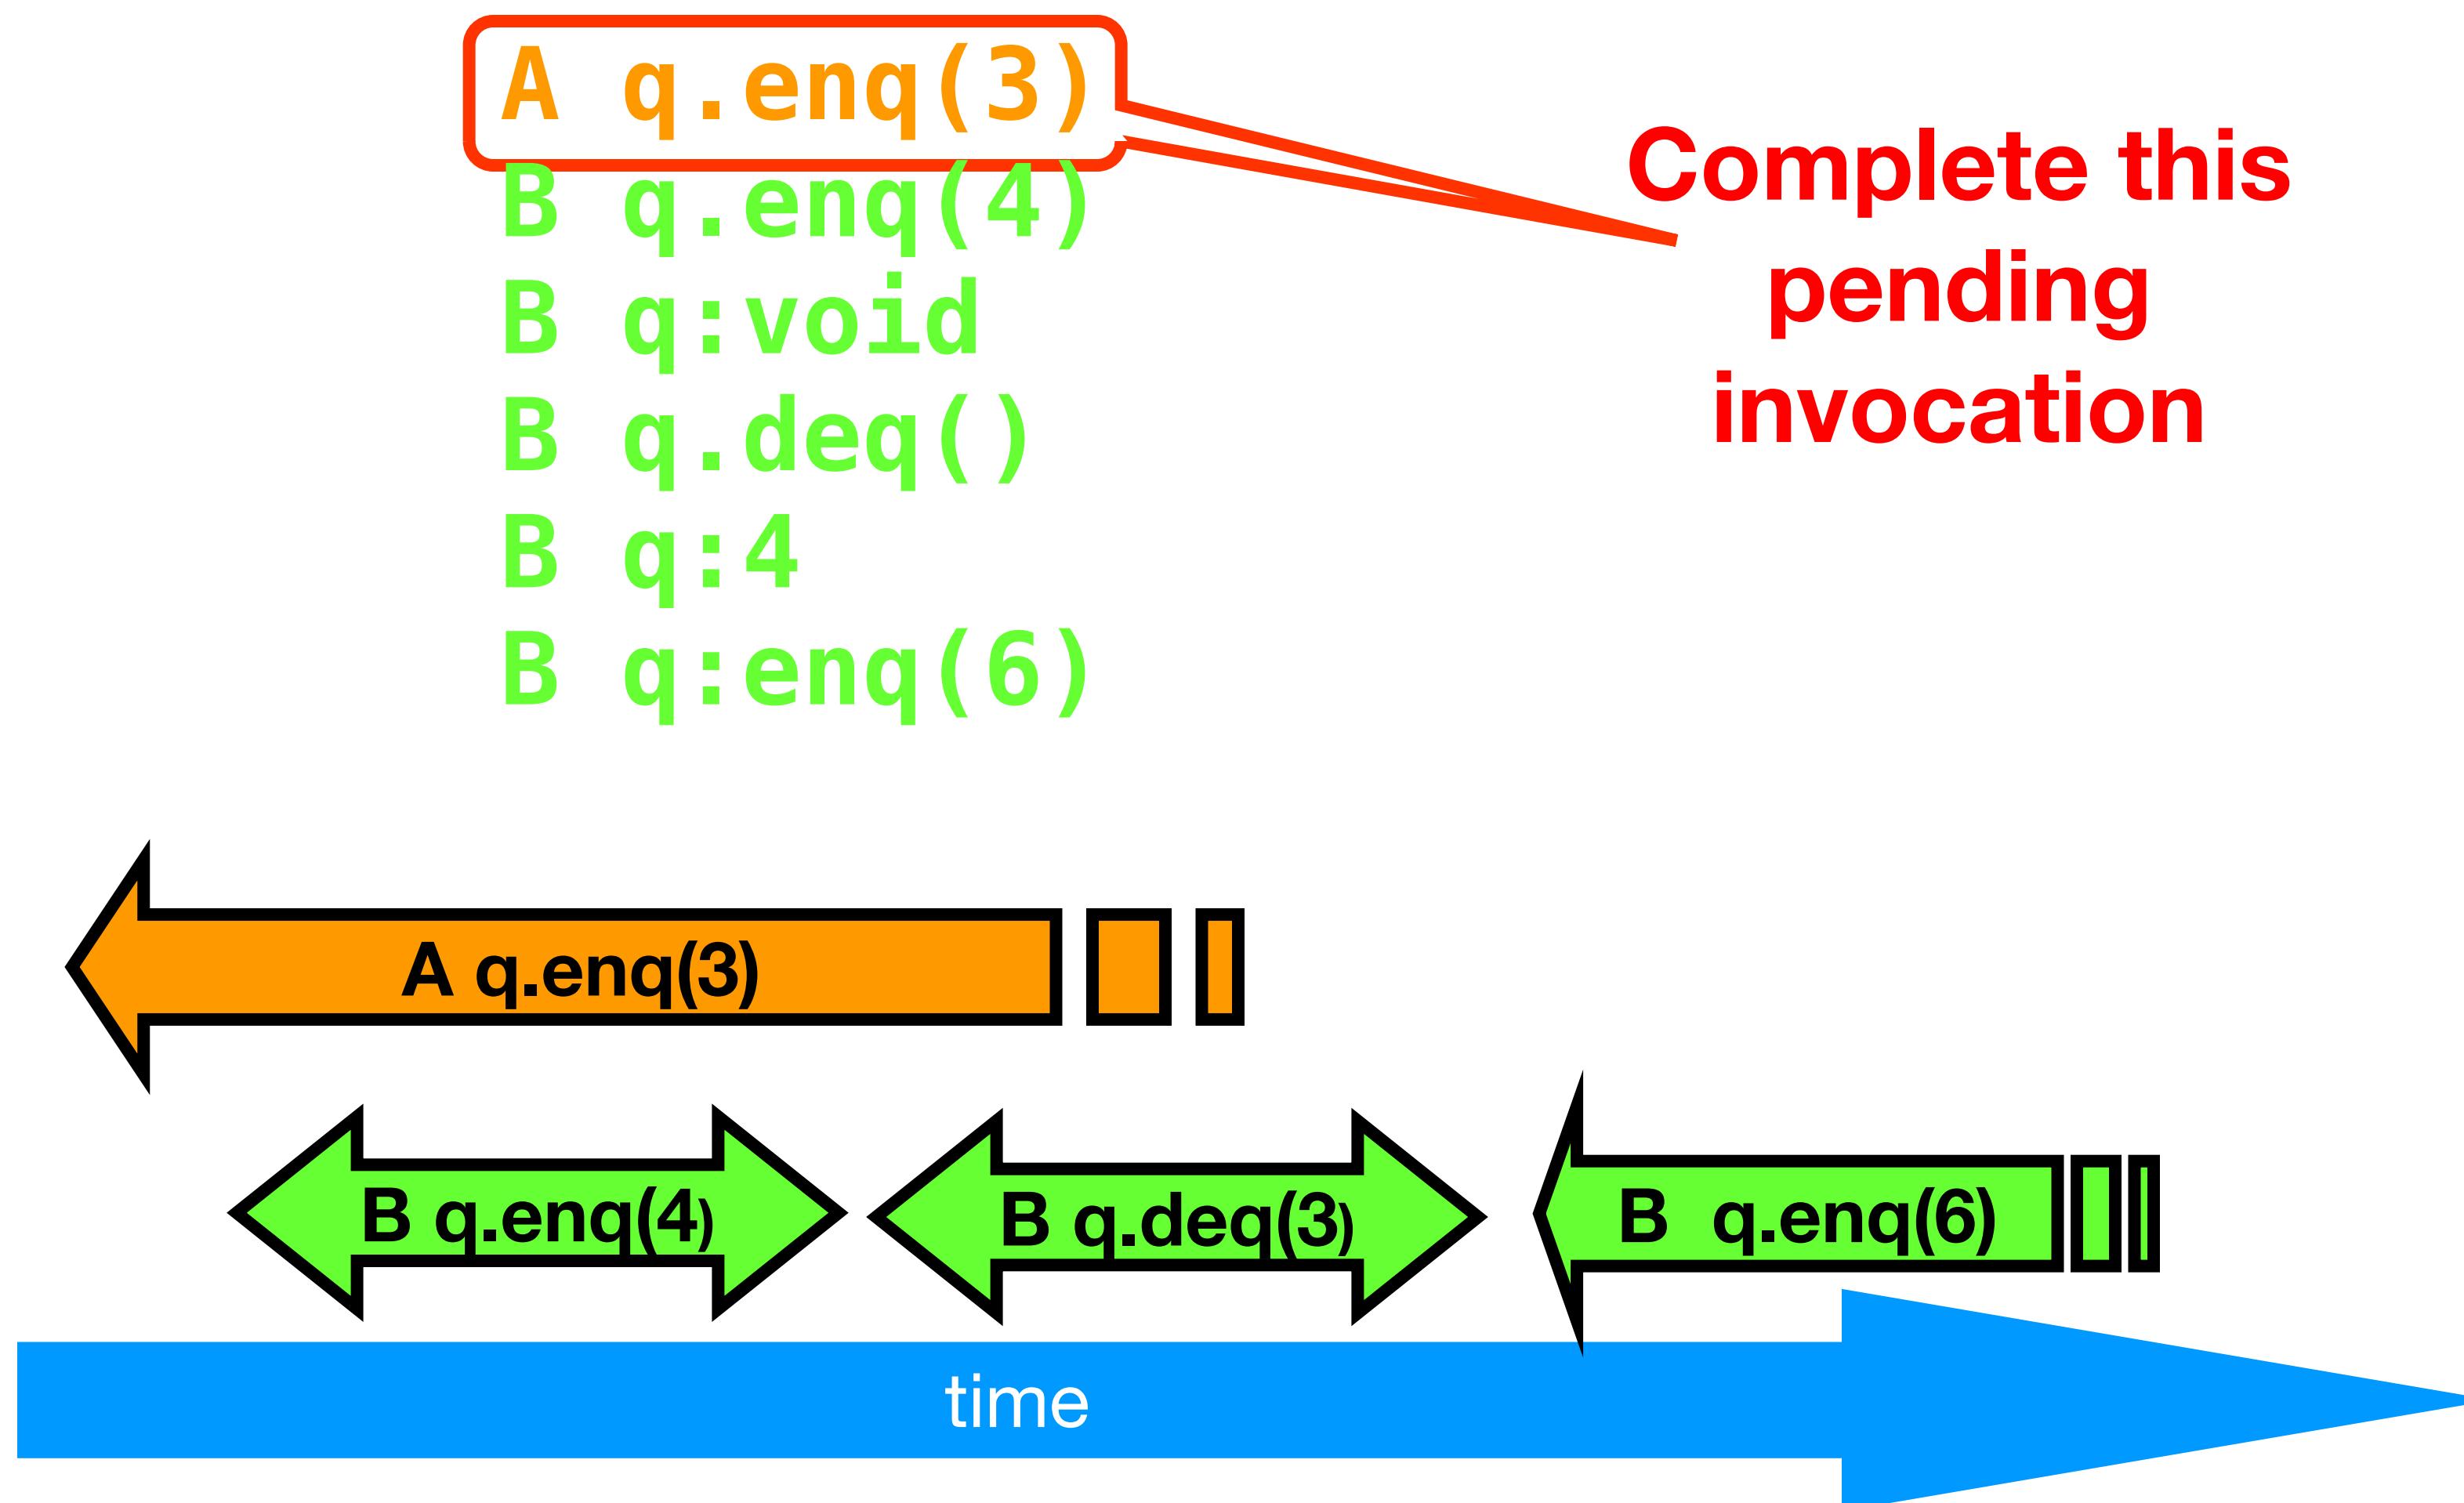


# Example

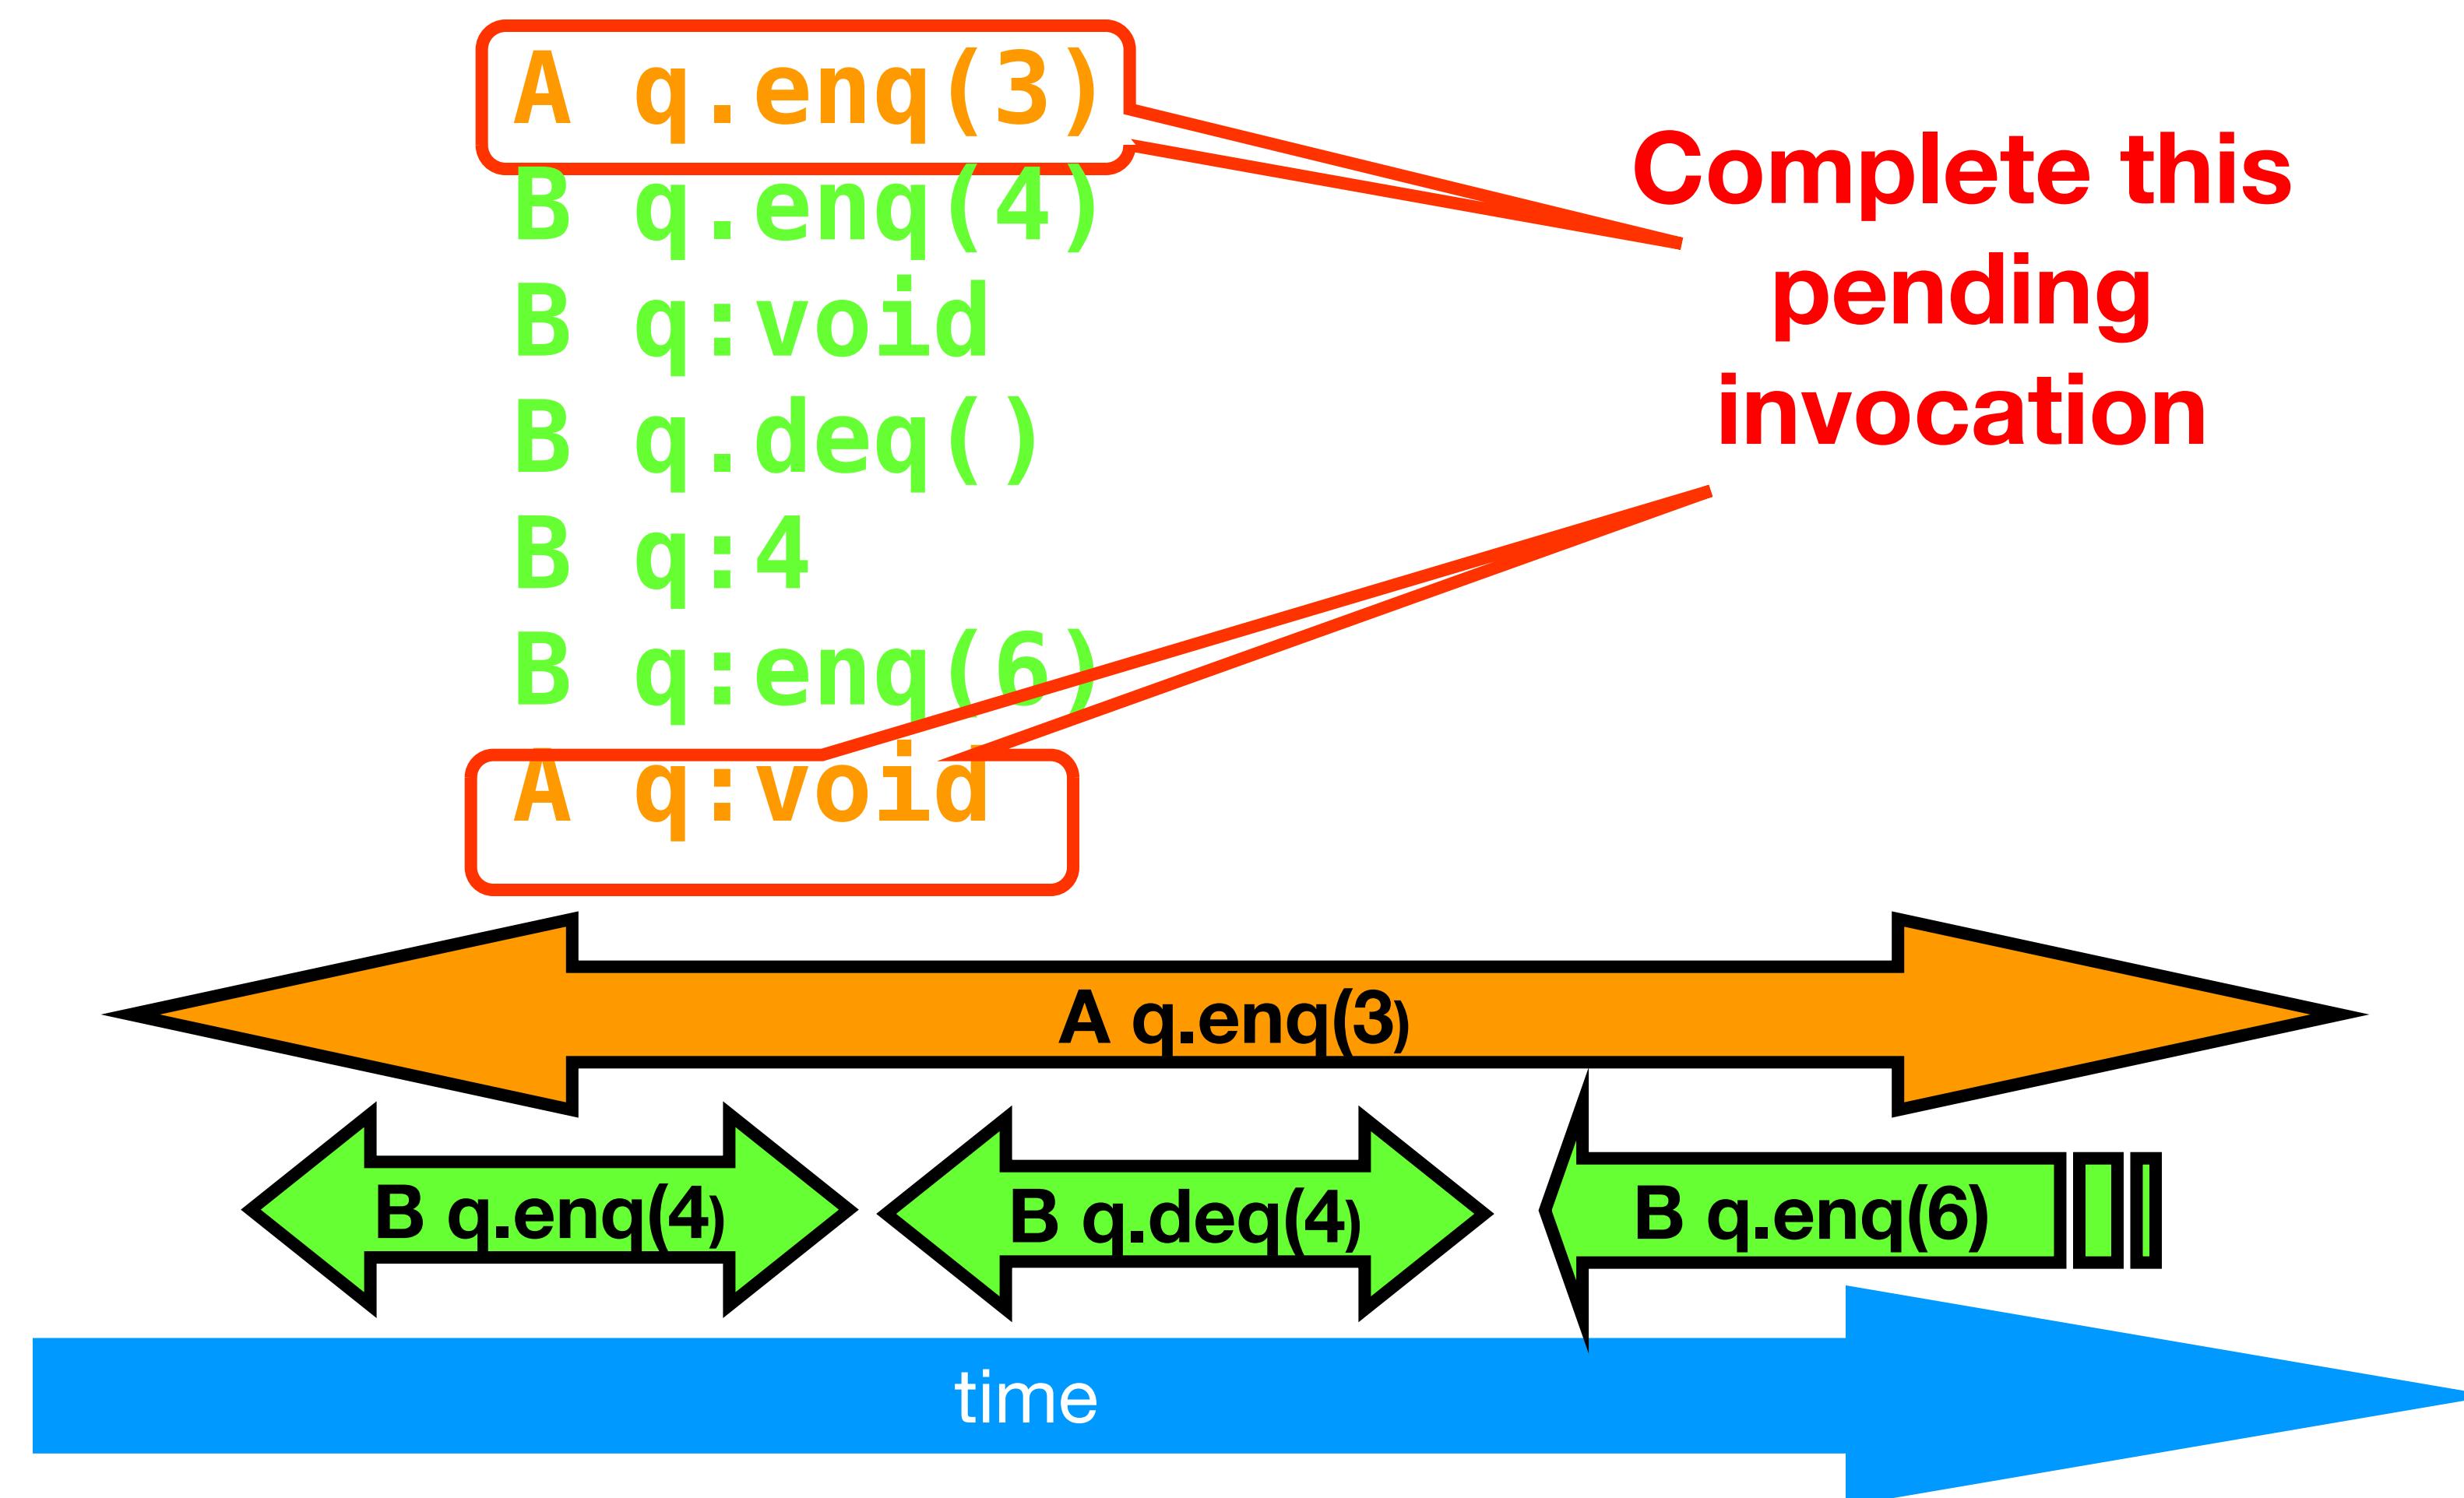
```
A q.enq(3)
B q.enq(4)
B q: void
B q.deq()
B q: 4
B q.enq(6)
```



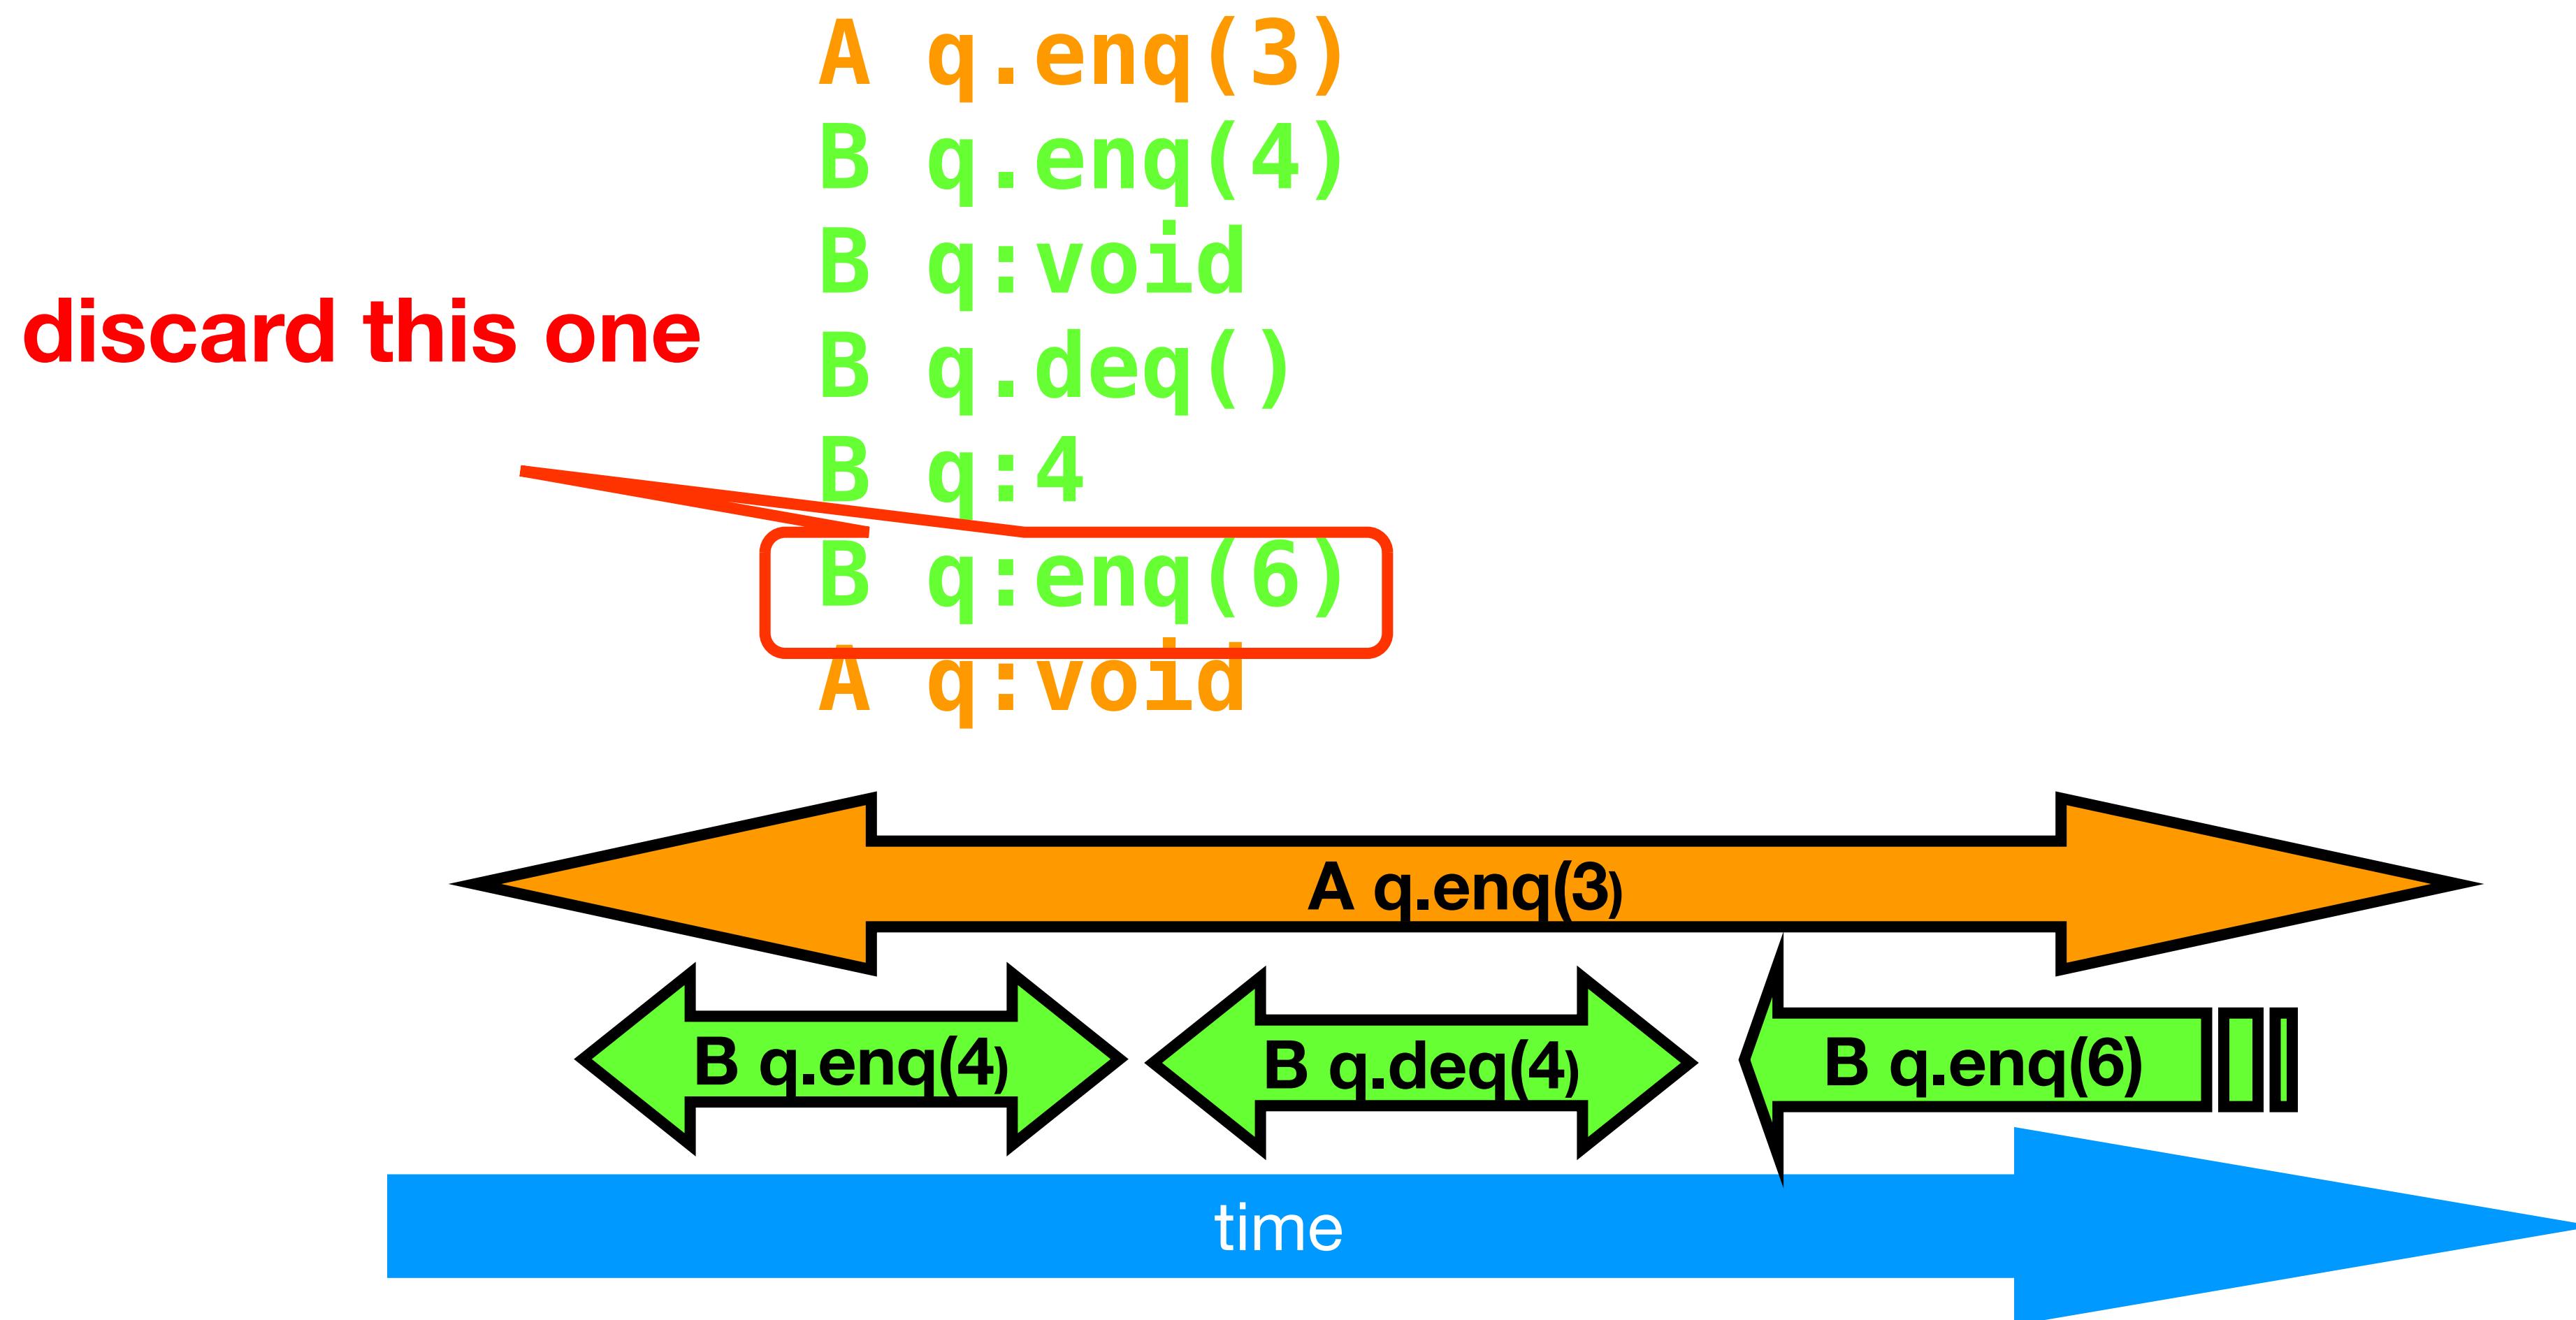
# Example



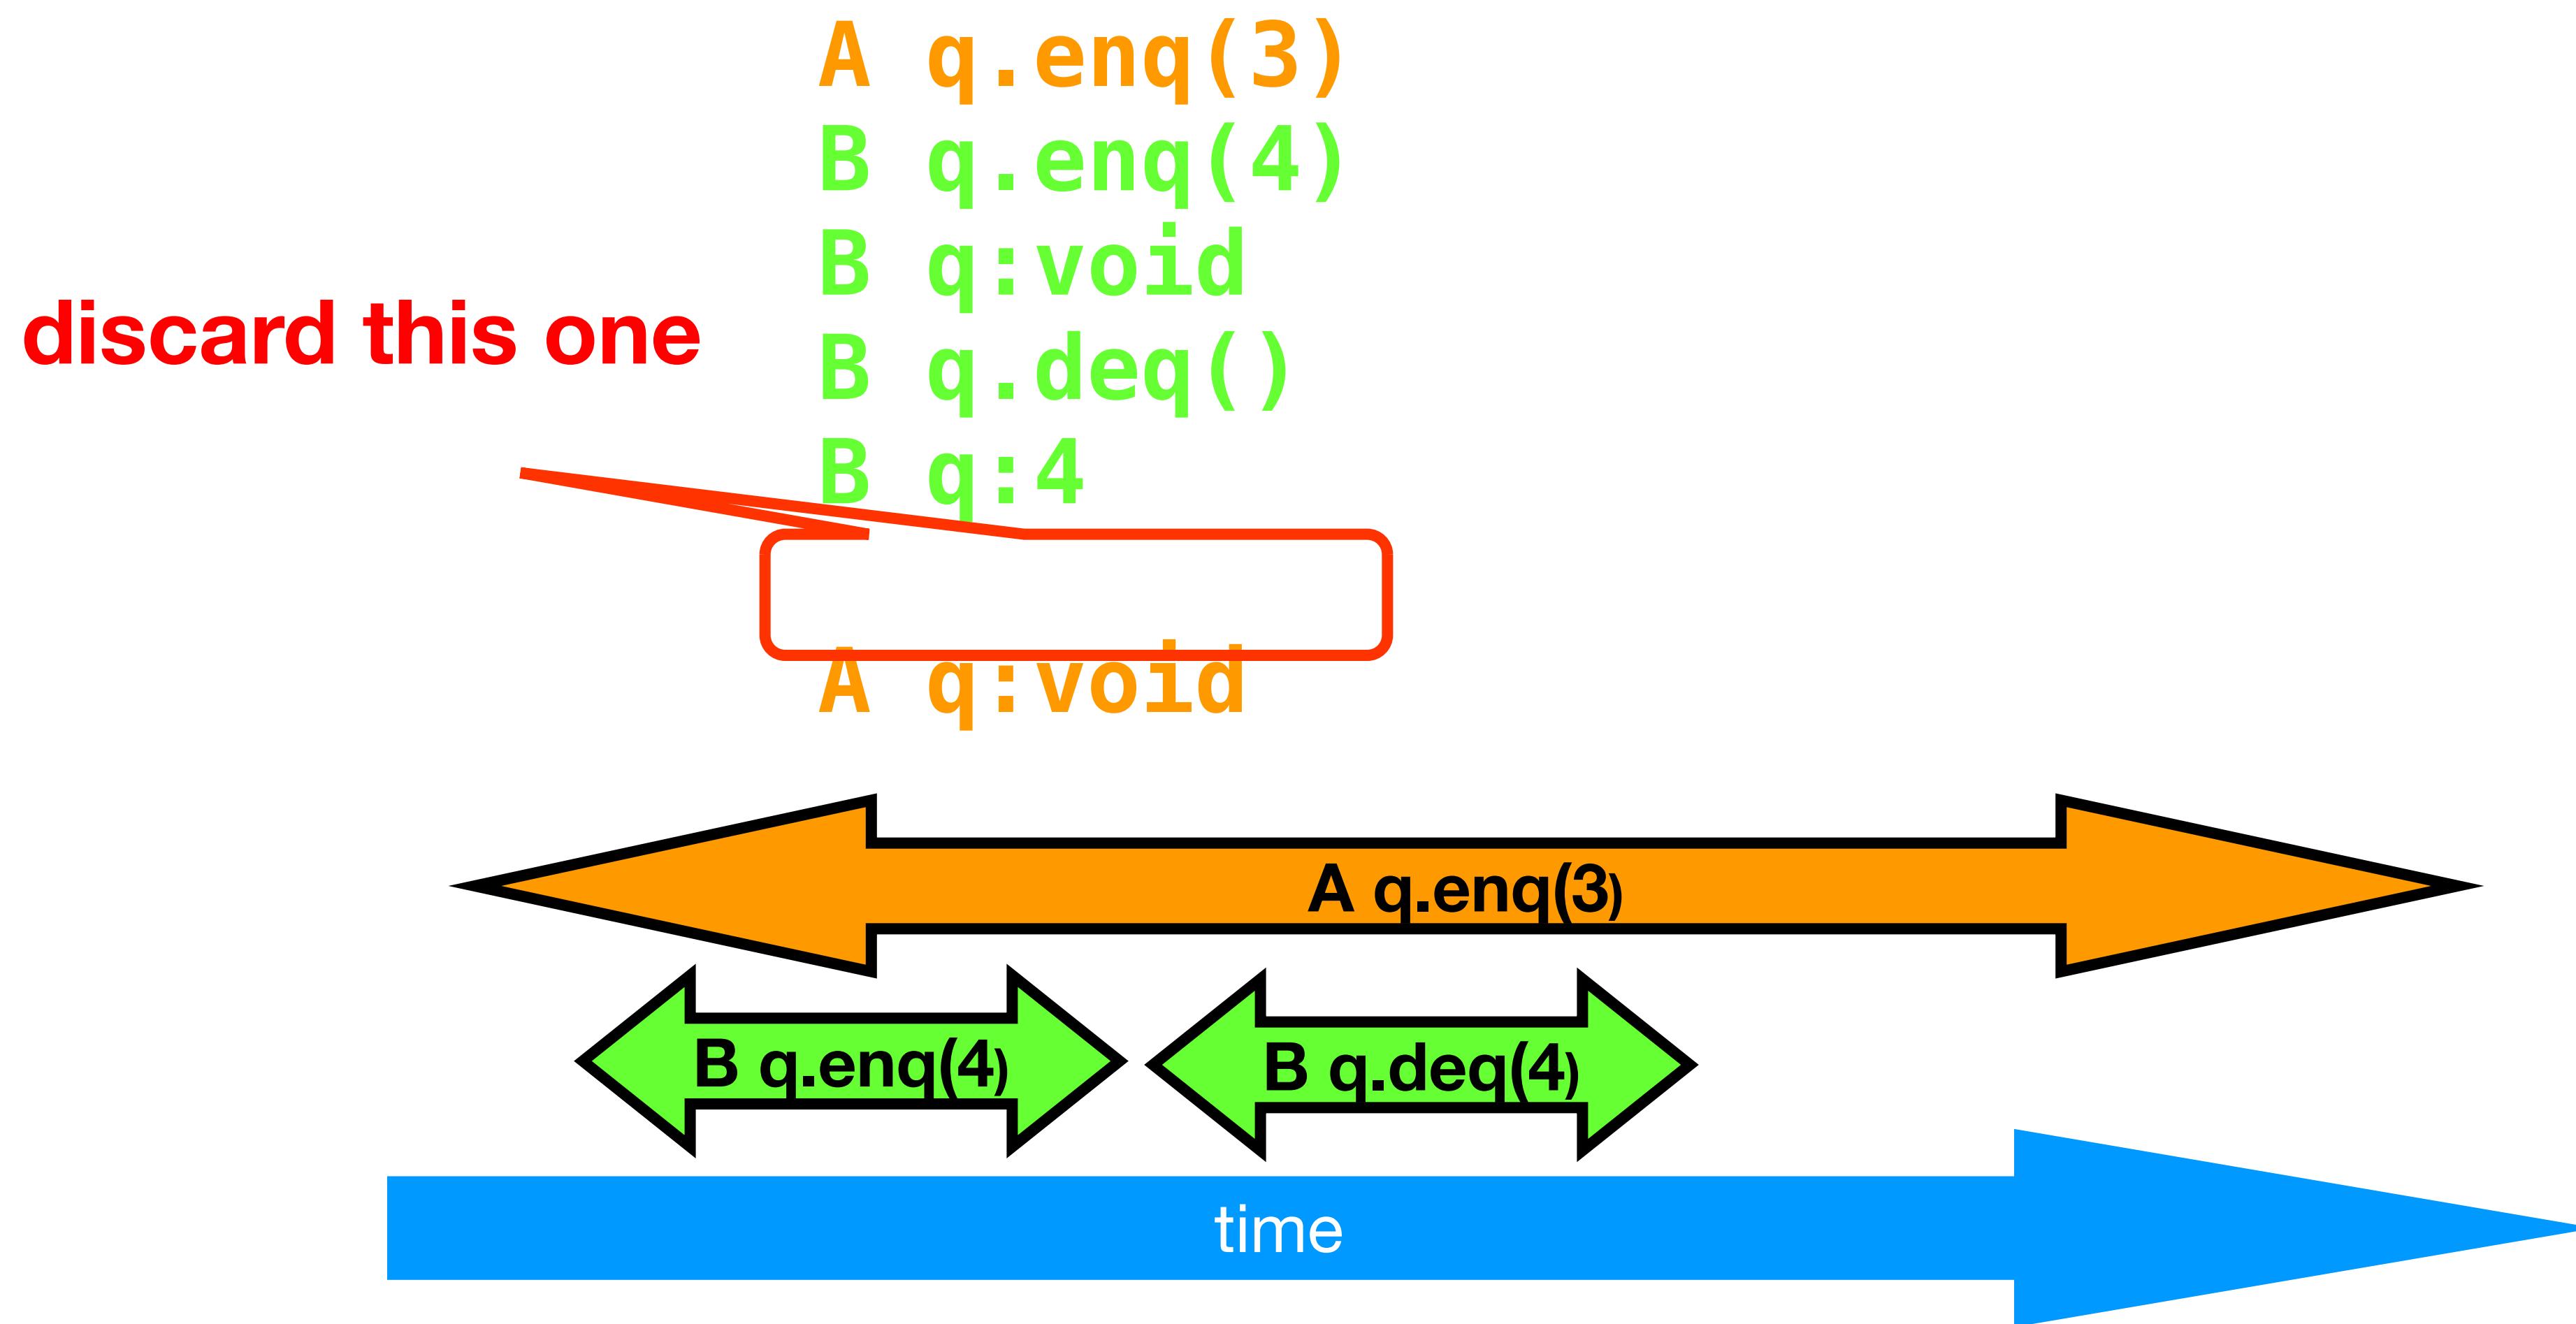
# Example



# Example

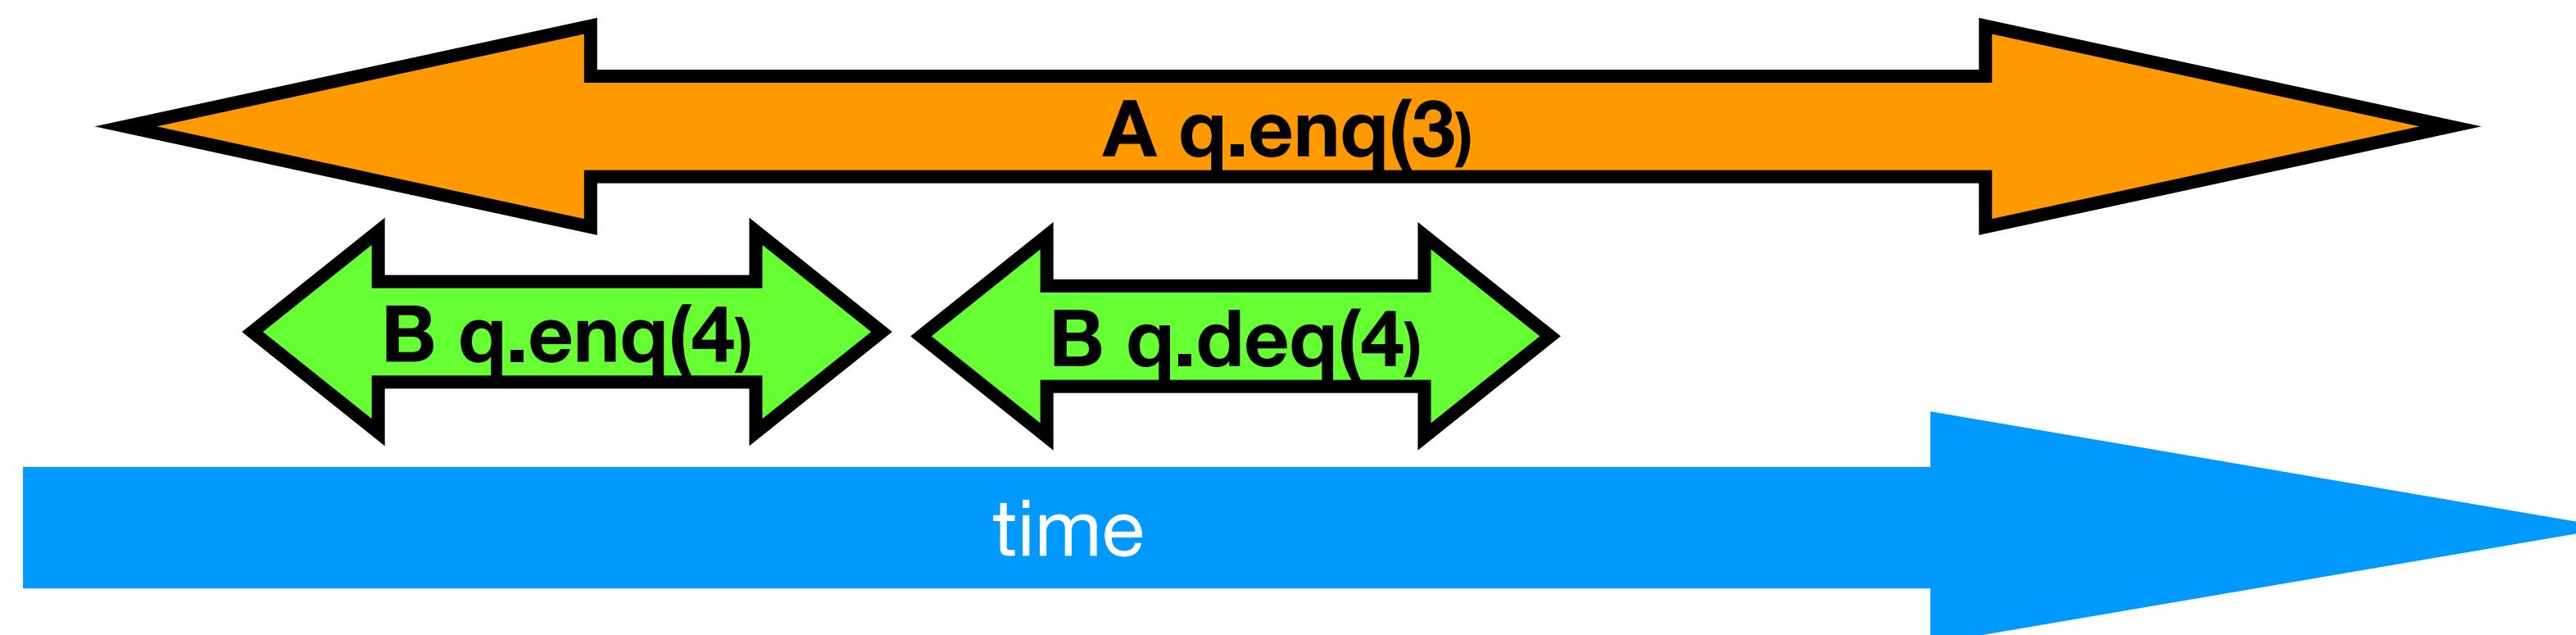


# Example



# Example

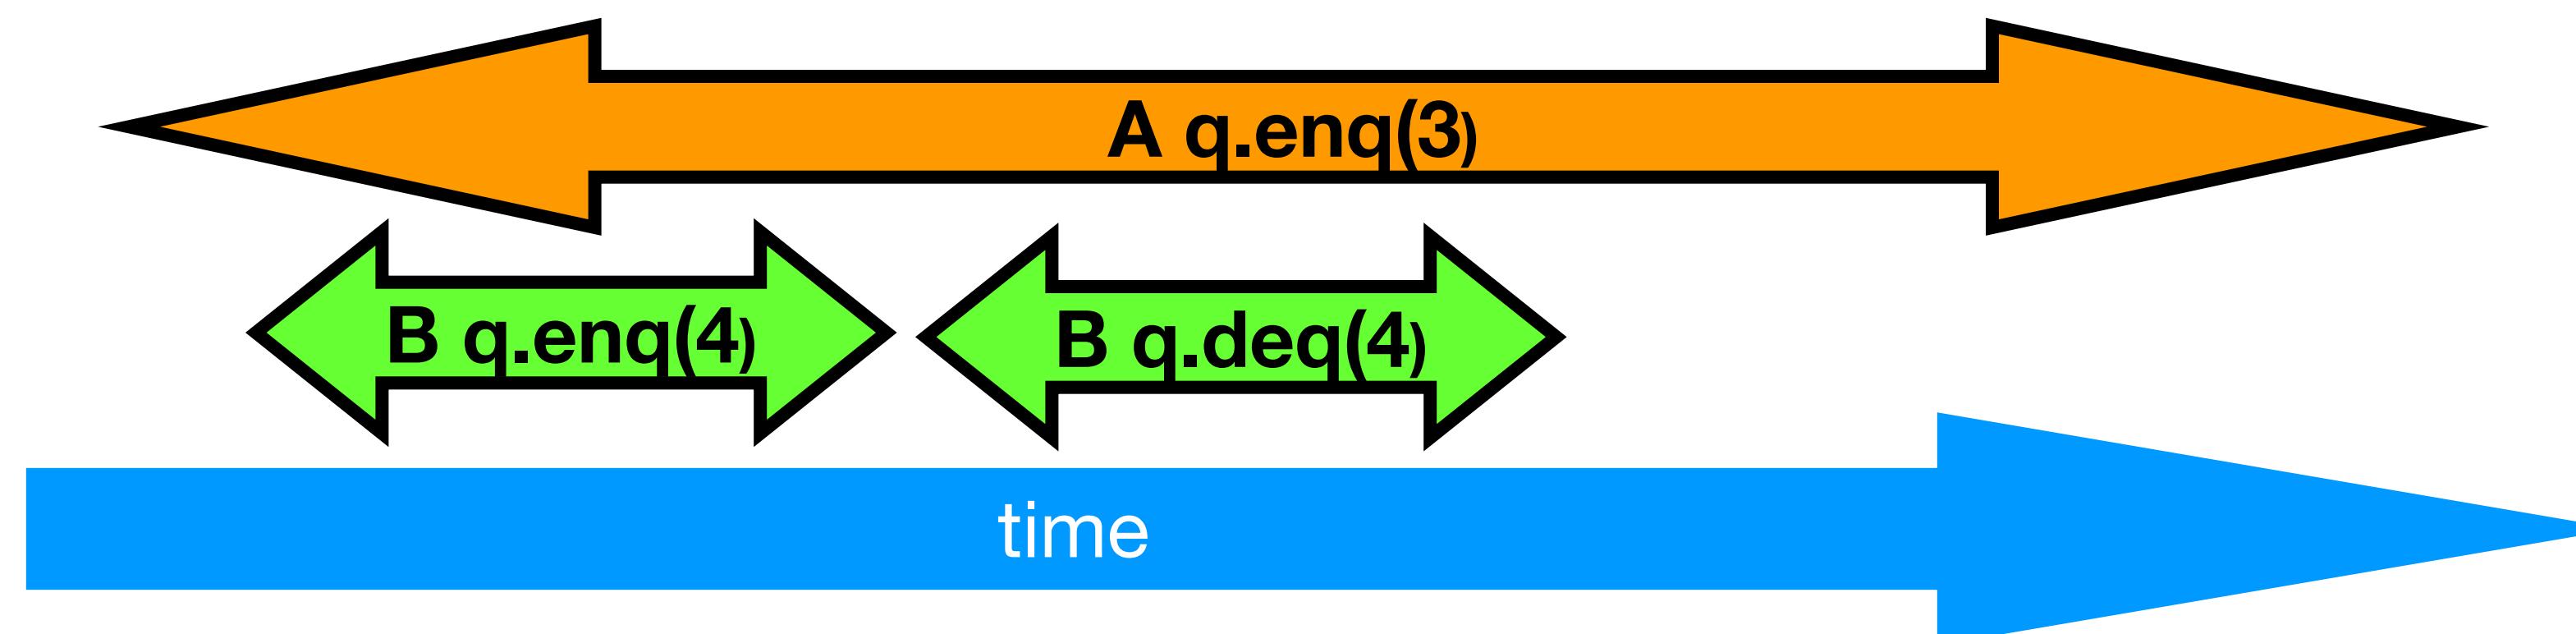
```
A q.enq(3)
B q.enq(4)
B q: void
B q.deq()
B q: 4
A q: void
```



# Example

A q.enq(3)  
B q.enq(4)  
B q: void  
B q.deq()  
B q: 4  
A q: void

B q.enq(4)  
B q: void  
A q.enq(3)  
A q: void  
B q.deq()  
B q: 4

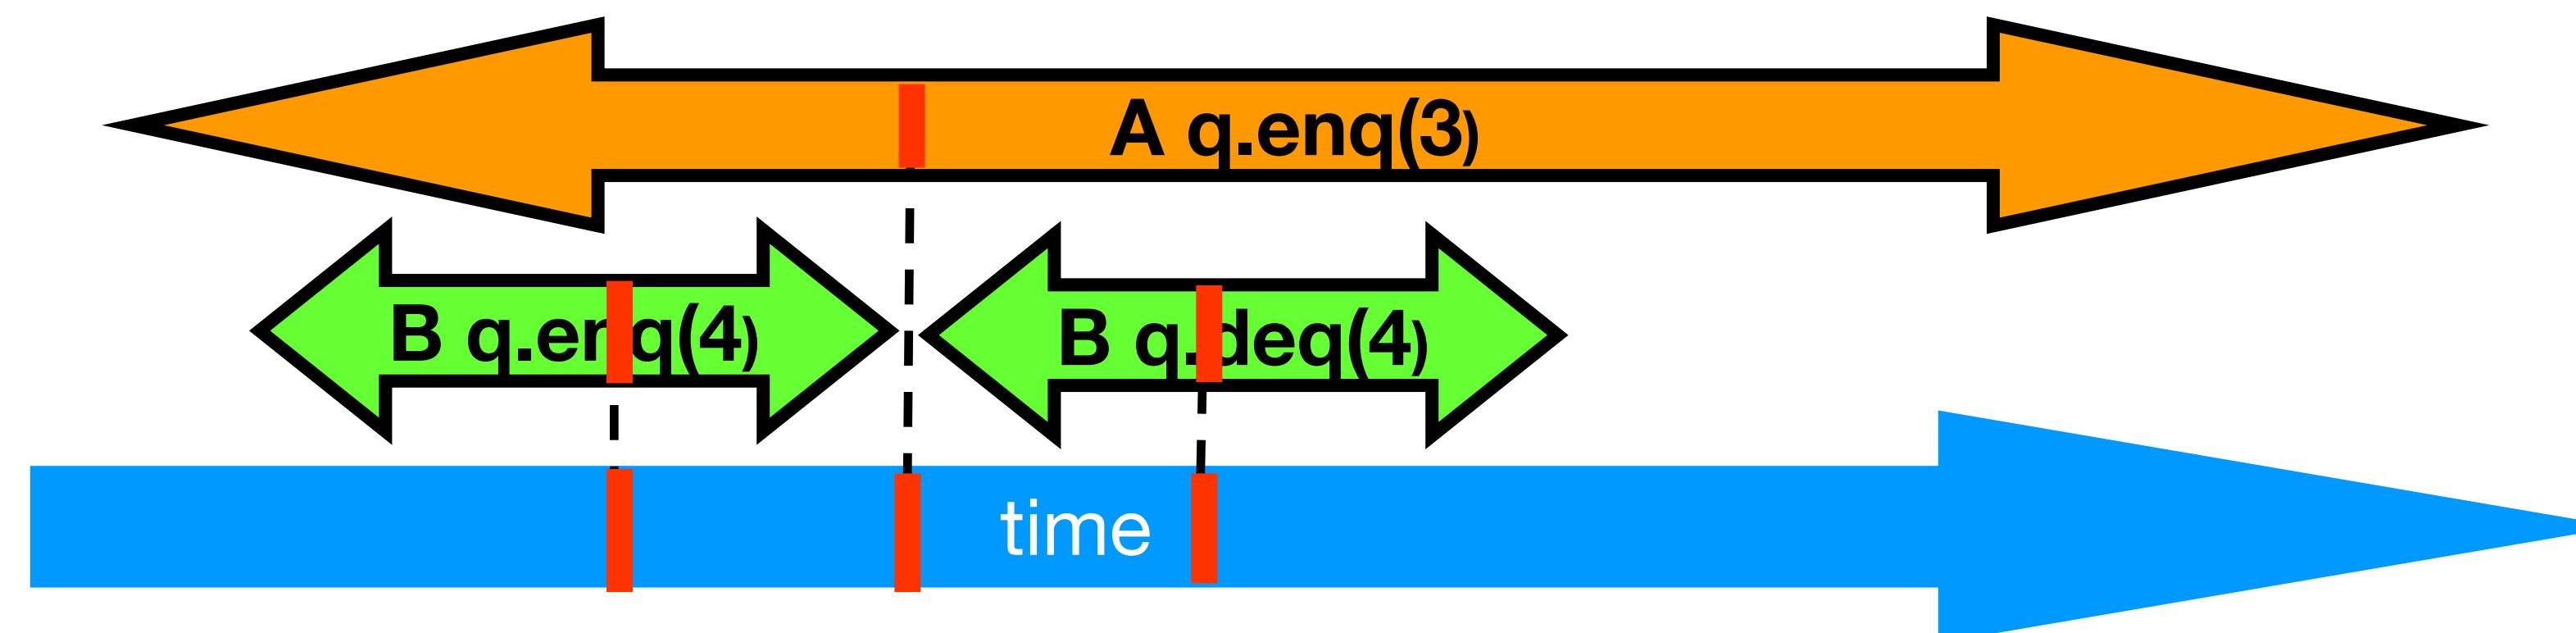


# Example

A q.enq(3)  
B q.enq(4)  
B q:void  
B q.deq()  
B q:4  
A q:void

Equivalent sequential history

B q.enq(4)  
B q:void  
A q.enq(3)  
A q:void  
B q.deq()  
B q:4



# Composability Theorem

- History  $H$  is linearizable if and only if
  - For every object  $x$
  - $H|x$  is linearizable
- We care about objects only!
  - (Materialism?)

# Why does composability matter?

- Modularity
- Can prove linearizability of objects in isolation
- Can compose independently-implemented objects

# Reasoning about Linearizability: Locking

```
let deq q =
  Mutex.lock q.lock;
  try
    if q.tail = q.head then
      raise Empty;
    match q.items.(q.head mod q.capacity) with
    | None -> assert false
    | Some x ->
        q.head <- q.head + 1;
        Mutex.unlock q.lock;
        x
  with e ->
    Mutex.unlock q.lock;
    raise e
```

Linearization points  
are when locks are  
released

# More Reasoning – Wait-free

```
(** Enqueue – should be called by only ONE thread *)    (** Dequeue – should be called by only ONE thread *)
let enq q x =
  (* Check if queue is full *)
  if q.tail - q.head = q.capacity then
    raise Full;
  (* Write to the array *)
  q.items.(q.tail mod q.capacity) <- Some x;
  (* Advance tail *)
  q.tail <- q.tail + 1

let deq q =
  (* Check if queue is empty *)
  if q.tail = q.head then
    raise Empty;
  (* Read from the array *)
  match q.items.(q.head mod q.capacity) with
  | None -> assert false (* Should never happen *)
  | Some x ->
    (* Advance head *)
    q.head <- q.head + 1;
    x
```

# More Reasoning – Wait-free

```
(** Enqueue – should be called by only ONE thread *)
let enq q x =
  (* Check if queue is full *)
  if q.tail - q.head = q.capacity then
    raise Full;
  (* Write to the array *)
  q.items.(q.tail mod q.capacity)
  (* Advance tail *)
  q.tail <- q.tail + 1

(** Dequeue – should be called by only ONE thread *)
let deque q =
  (* Check if queue is empty *)
  if q.tail = q.head then
    raise Empty;
  (* Read from the array *)
  match q.items.(q.head mod q.capacity) with
  | None -> assert false (* Should never happen *)
  | Some x ->
    (* Advance head *)
    q.head <- q.head + 1;
```

Remember that there  
is only one enqueue  
and only one dequeuer

Linearization order is  
order head and tail  
fields modified

# Finding linearisation points

- Identify one atomic step where the method “happens”
  - Critical section
  - Machine instruction
- Doesn’t always work
  - Might need to define several different steps for a given method
  - We will see this phenomenon in future lectures

# Linearizability: Summary

- Powerful specification tool for shared objects
- Allows us to capture the notion of objects being “atomic”
- Don’t leave home without it

# Alternative: Sequential Consistency

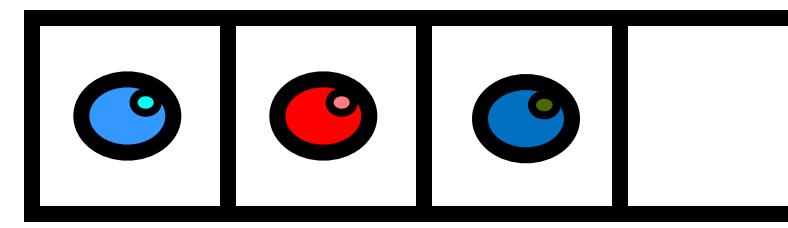
- History  $H$  is ***Sequentially Consistent*** if it can be extended to  $G$  by
  - Appending zero or more responses to pending invocations
  - Discarding other pending invocations
- So that  $G$  is equivalent to
  - Legal sequential history  $S$
  - ~~where  $\rightarrow_G \subset \rightarrow_S$~~

*Differs from  
Linearizability*

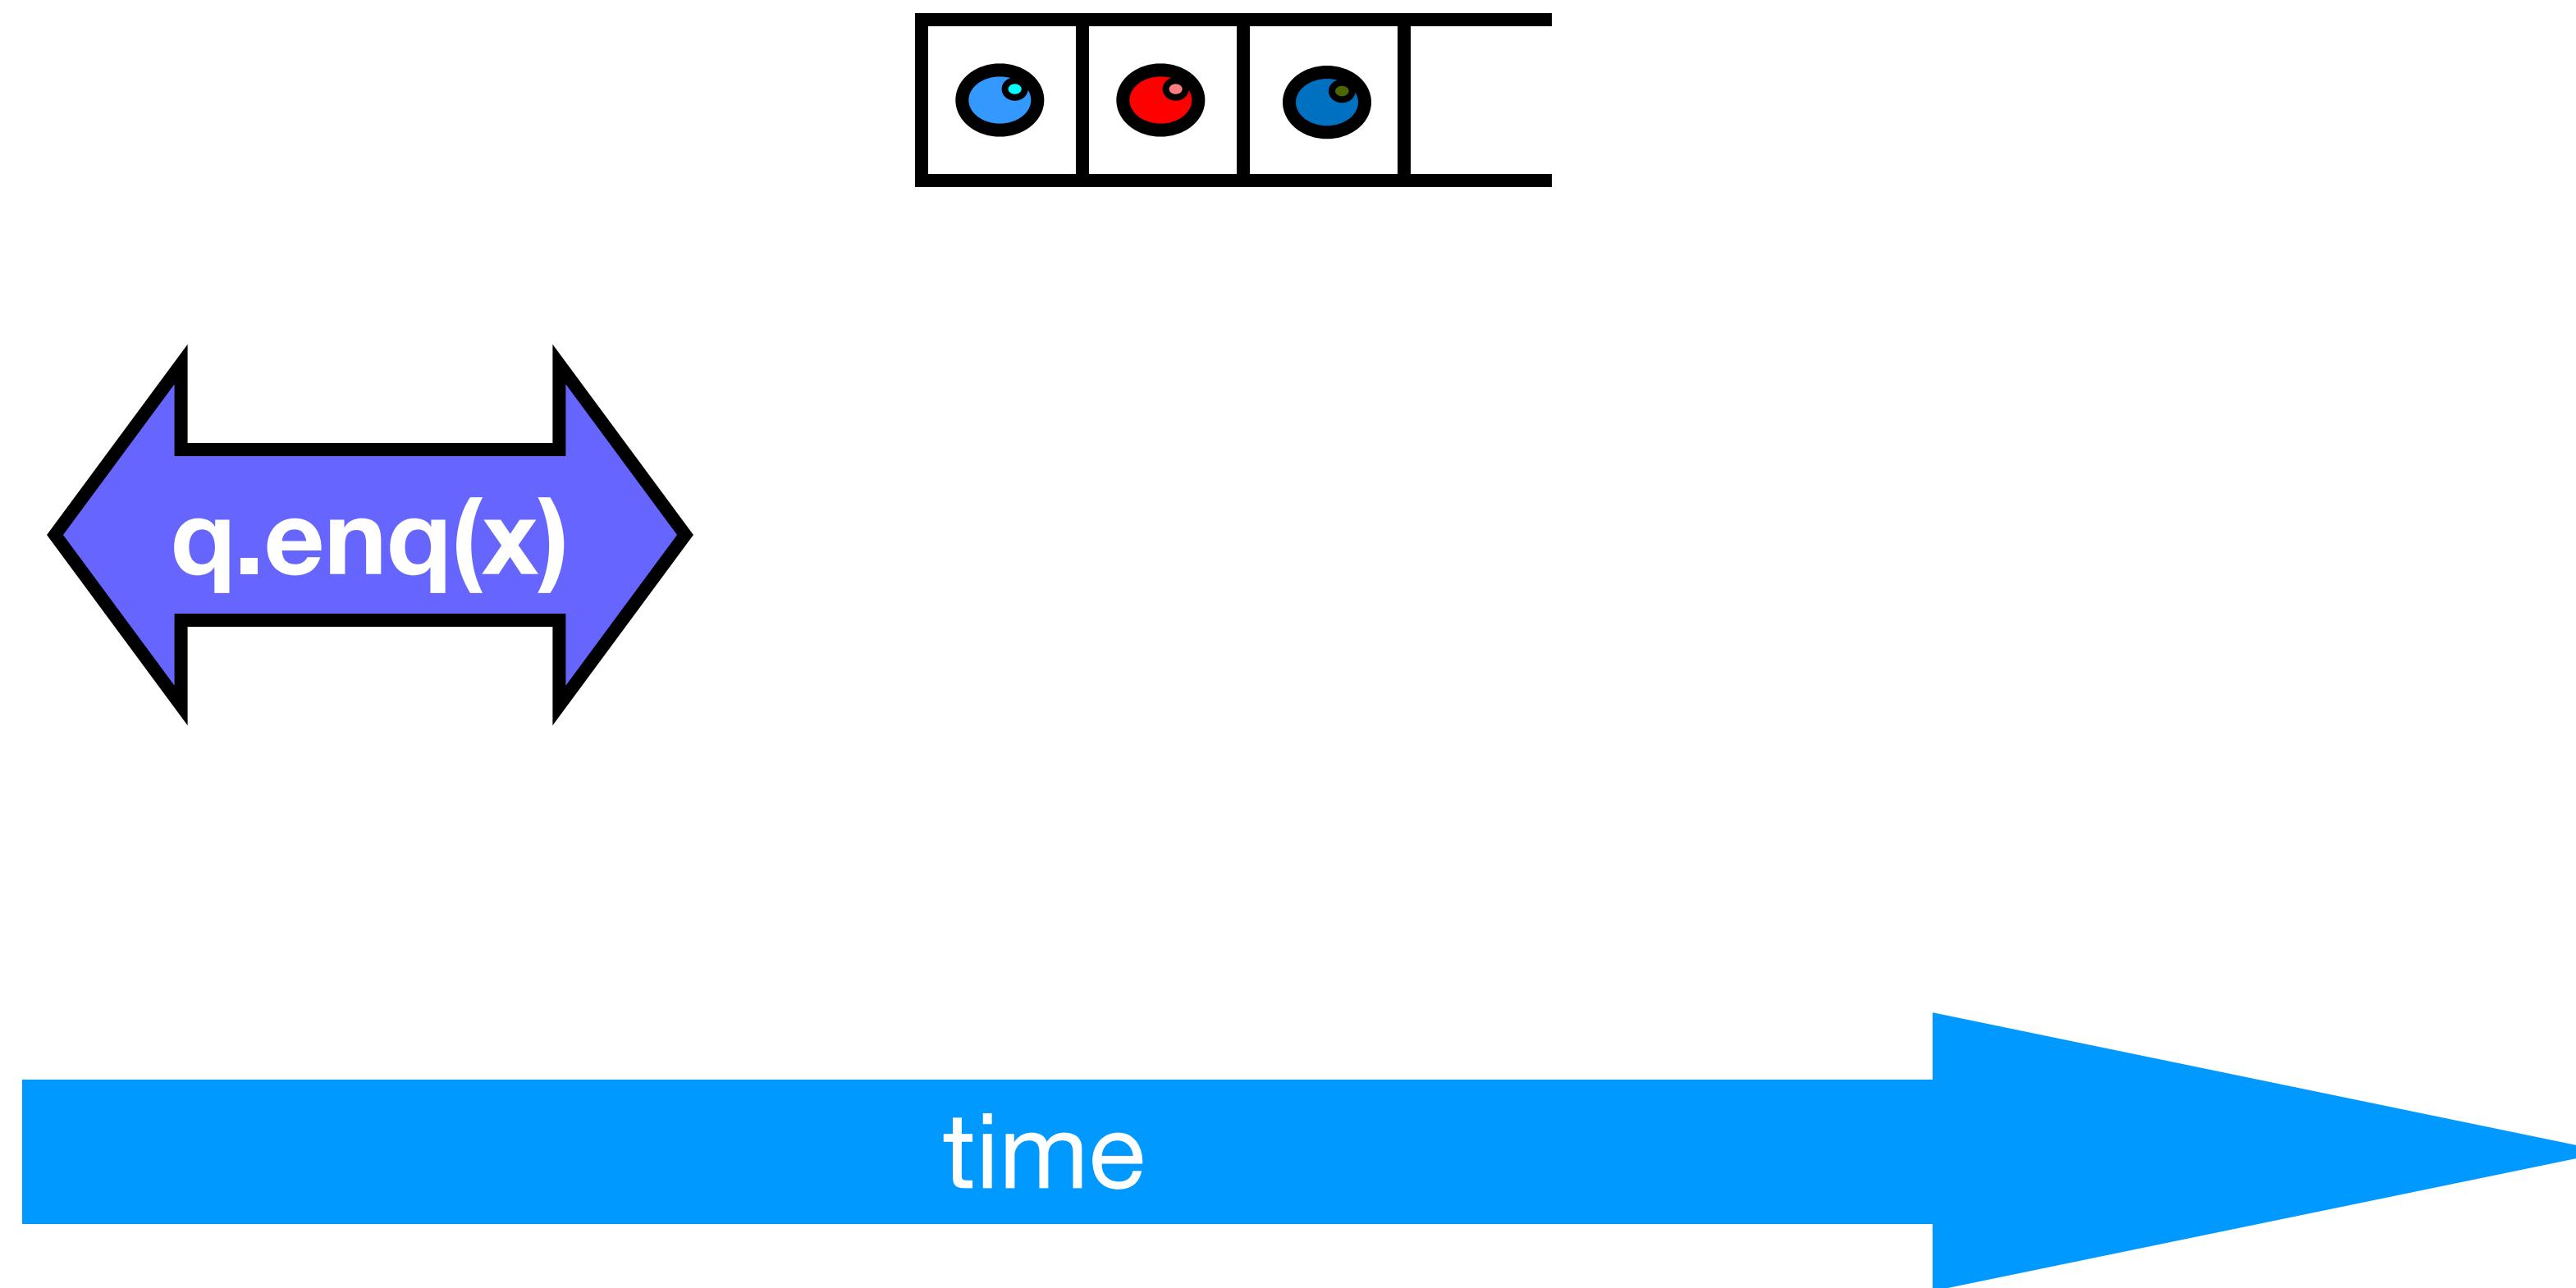
# Sequential Consistency

- No need to preserve real-time order
  - **Cannot** re-order operations done by the same thread
  - **Can** re-order non-overlapping operations done by different threads
- Often used to describe multiprocessor memory architectures

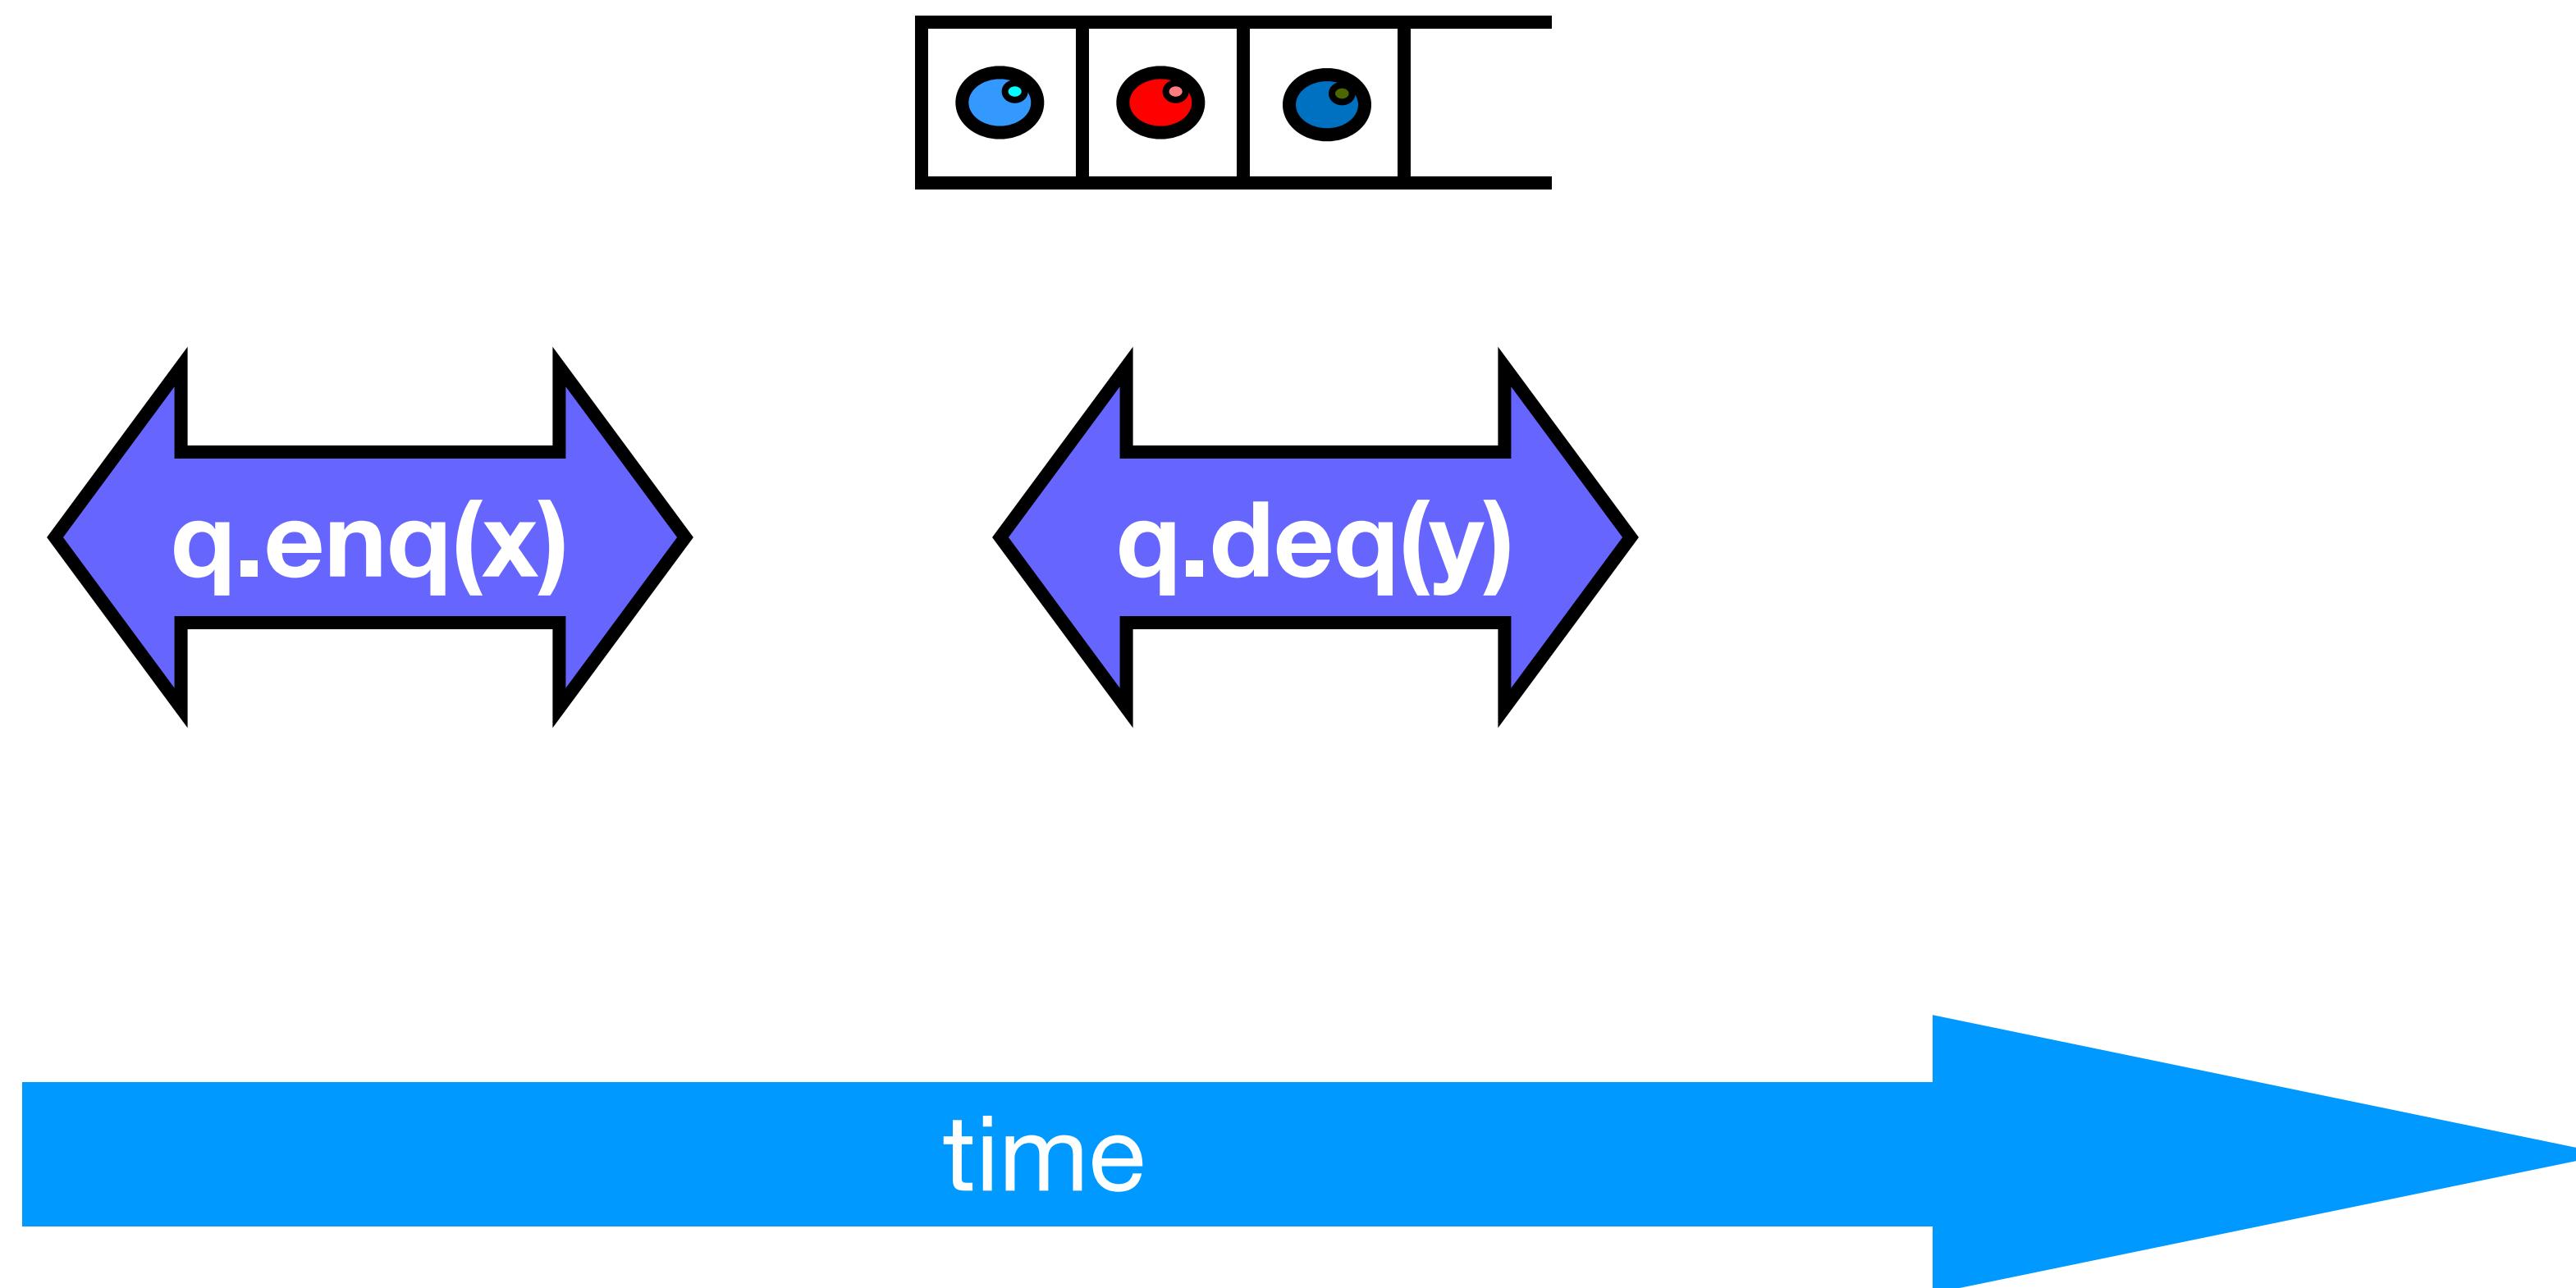
# Example



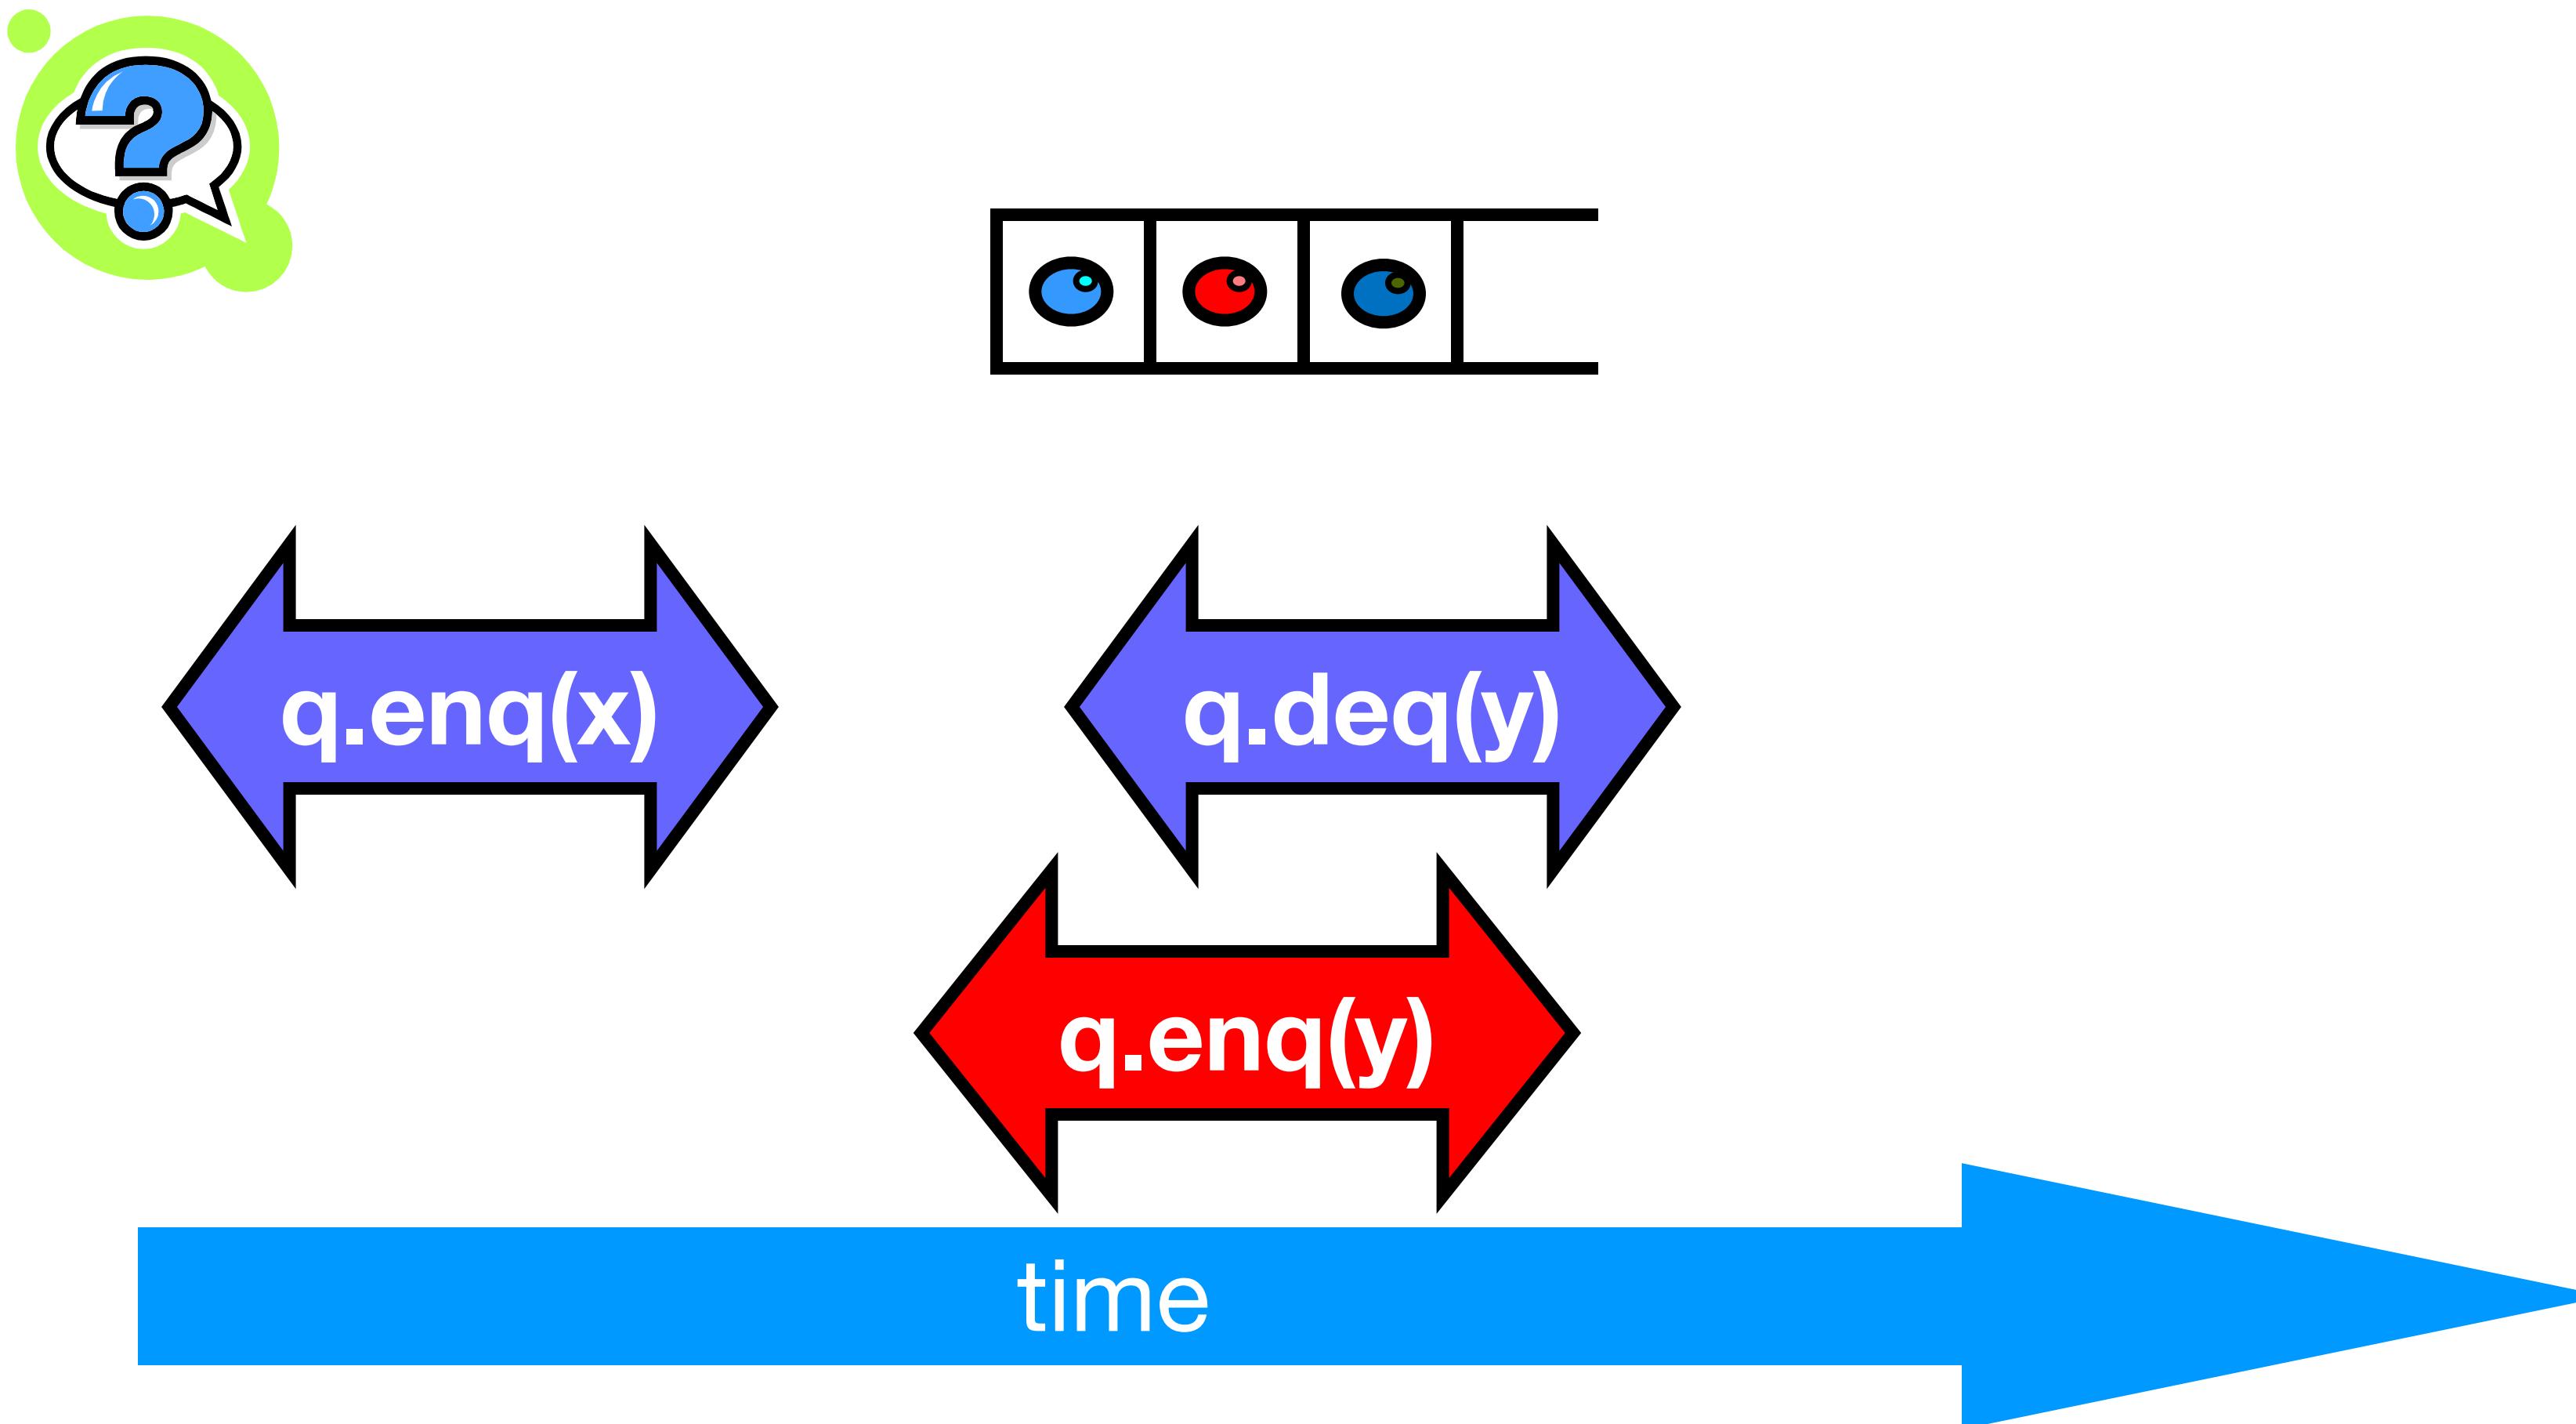
# Example



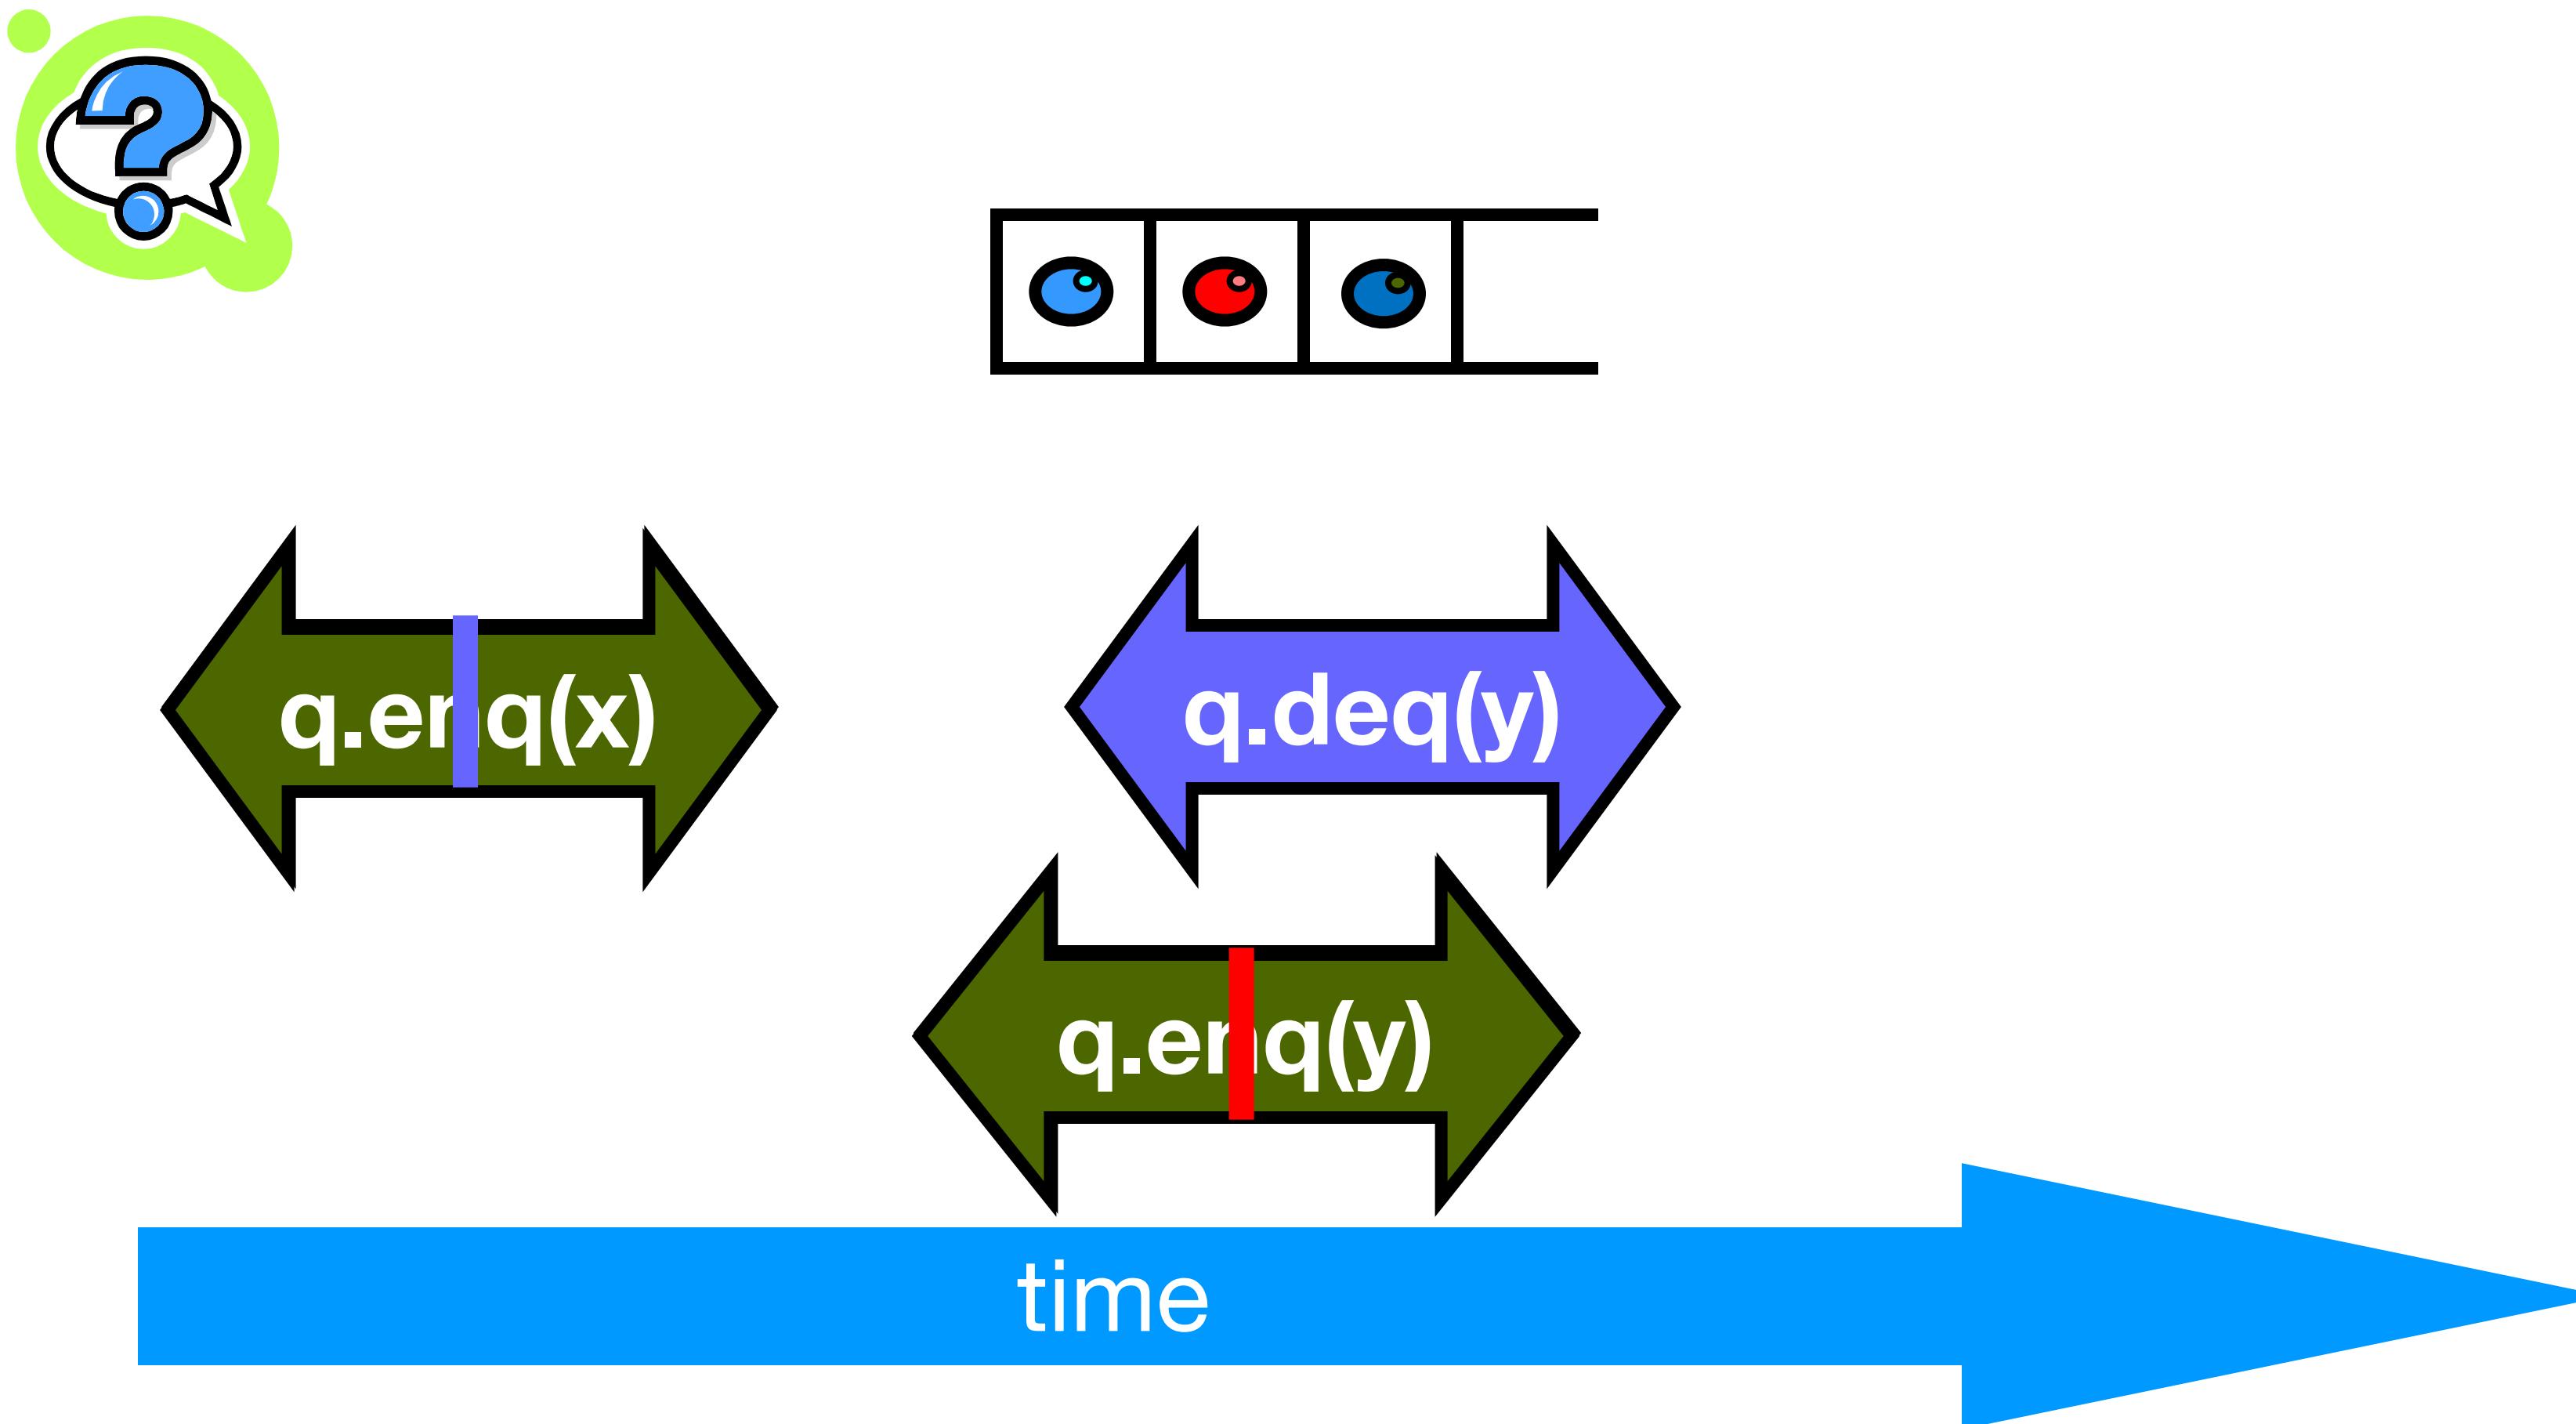
# Example



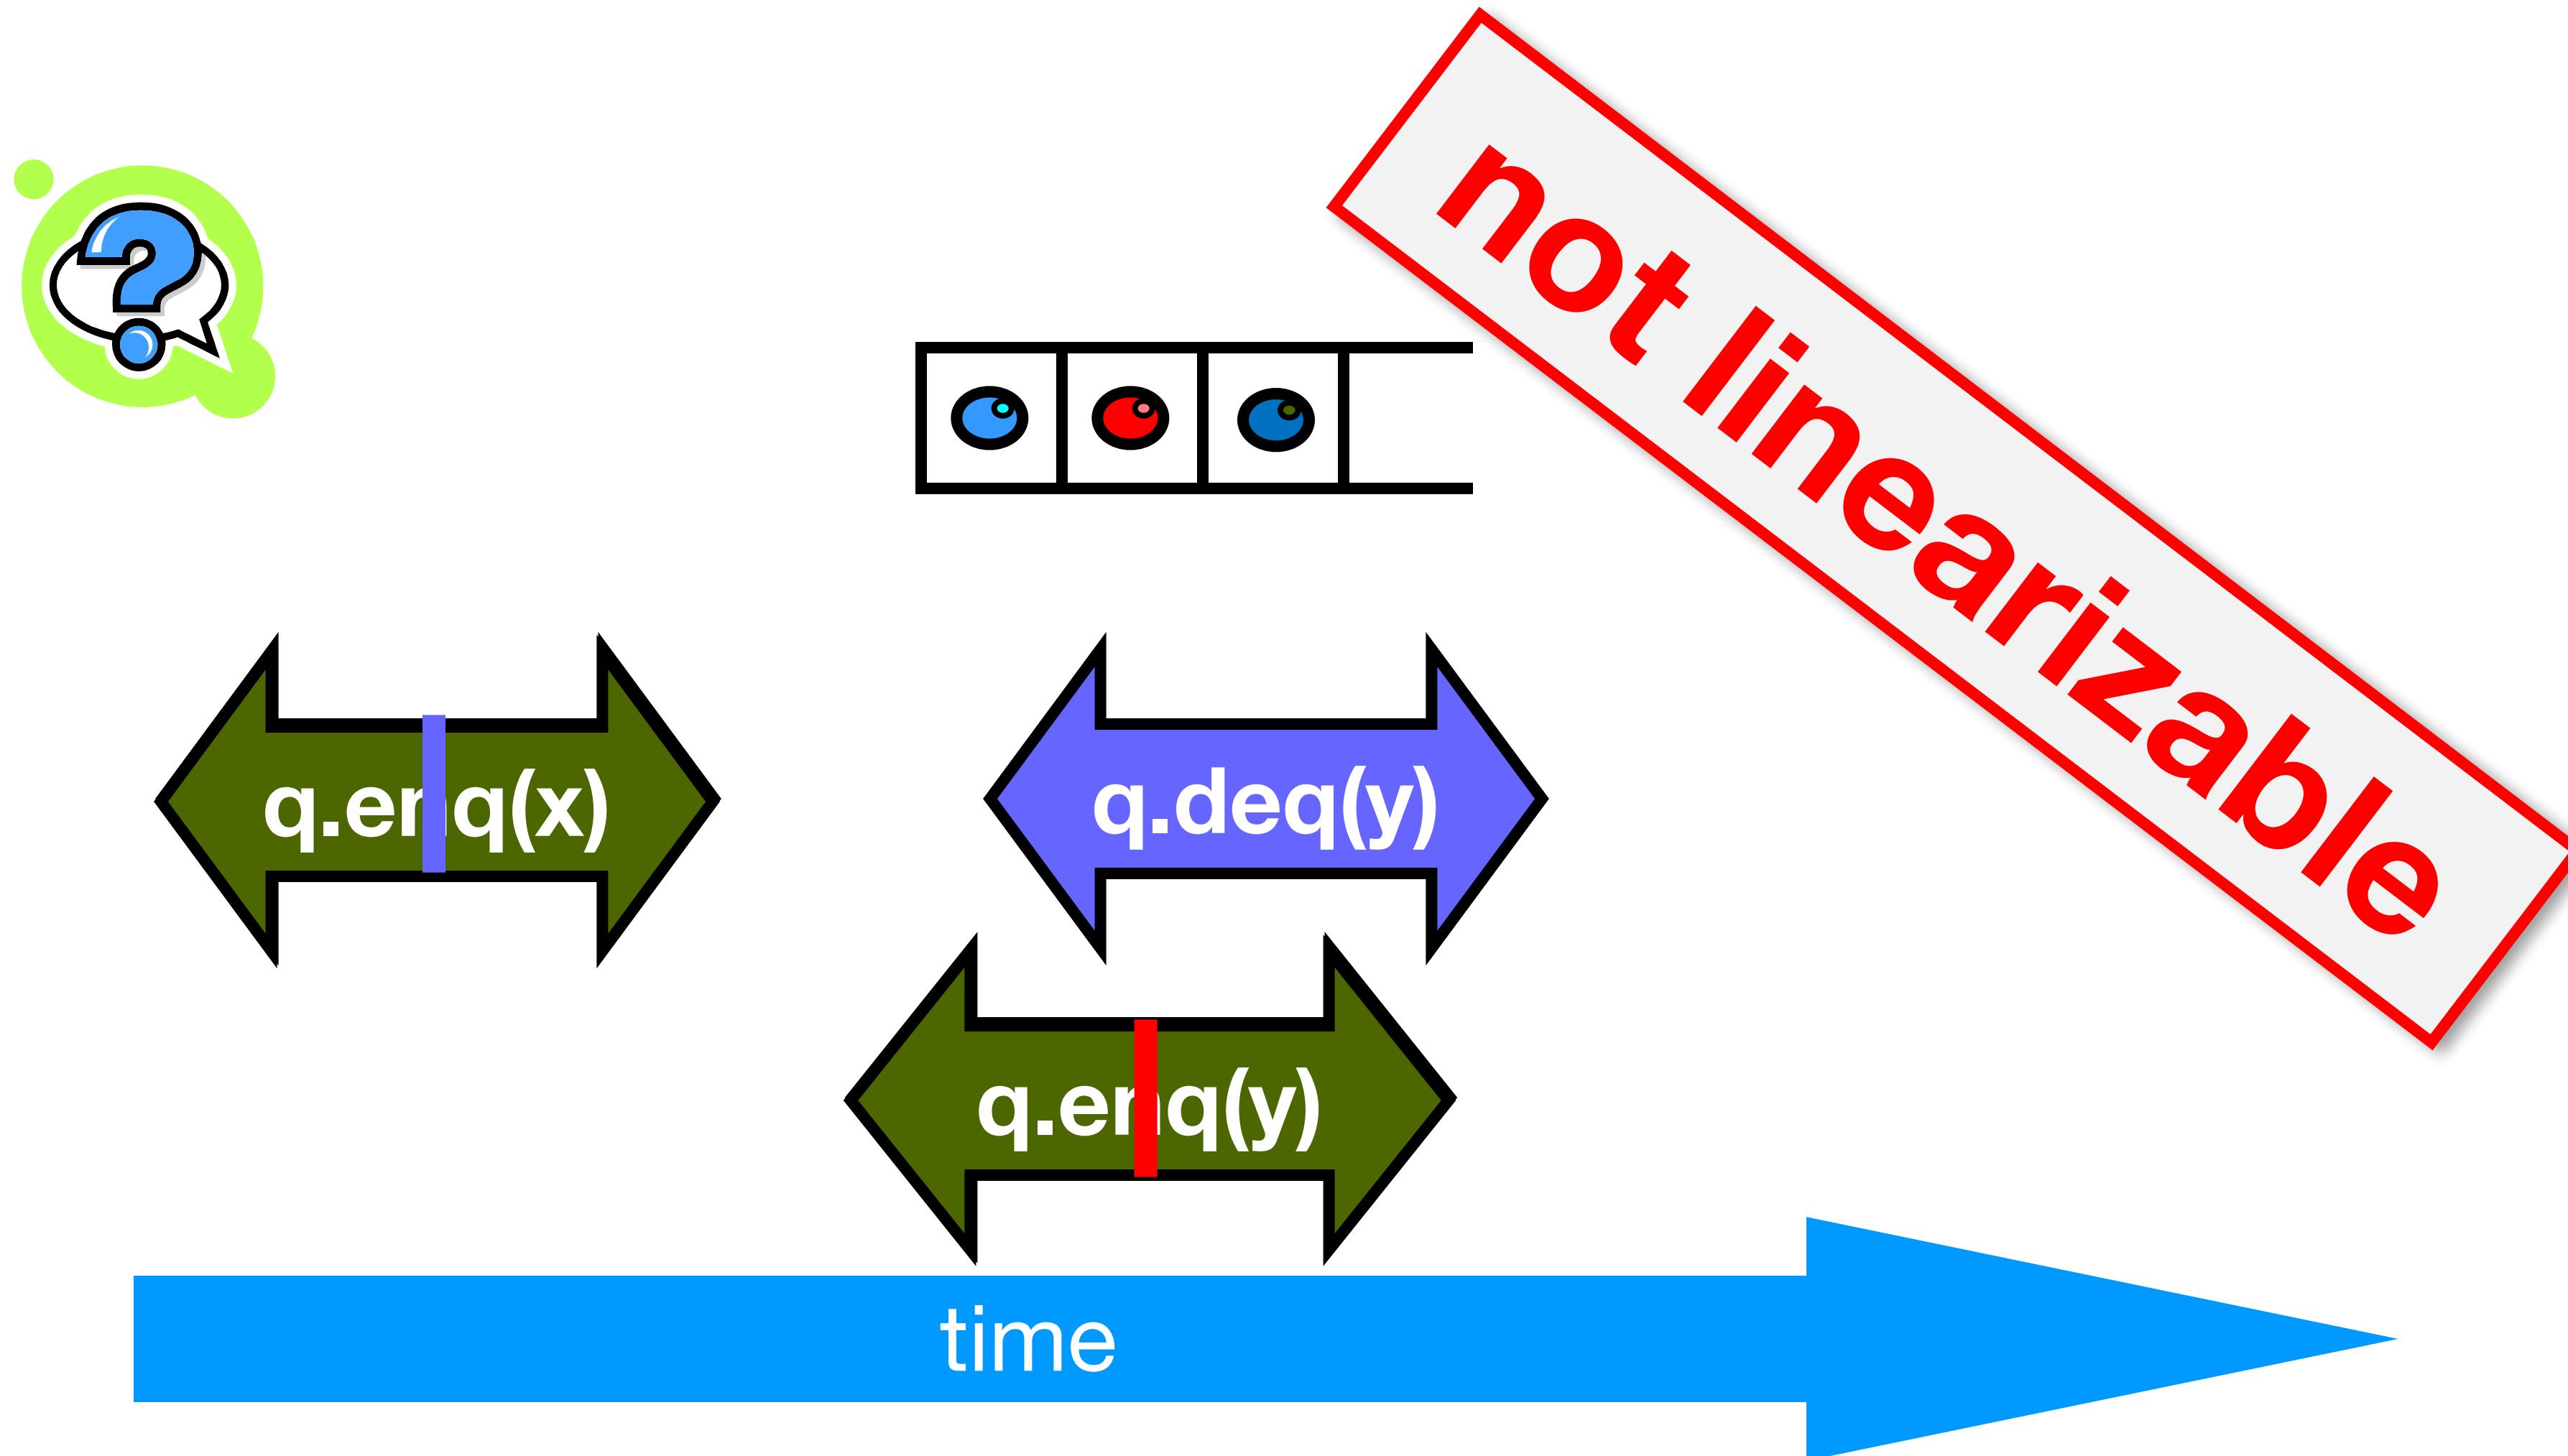
# Example



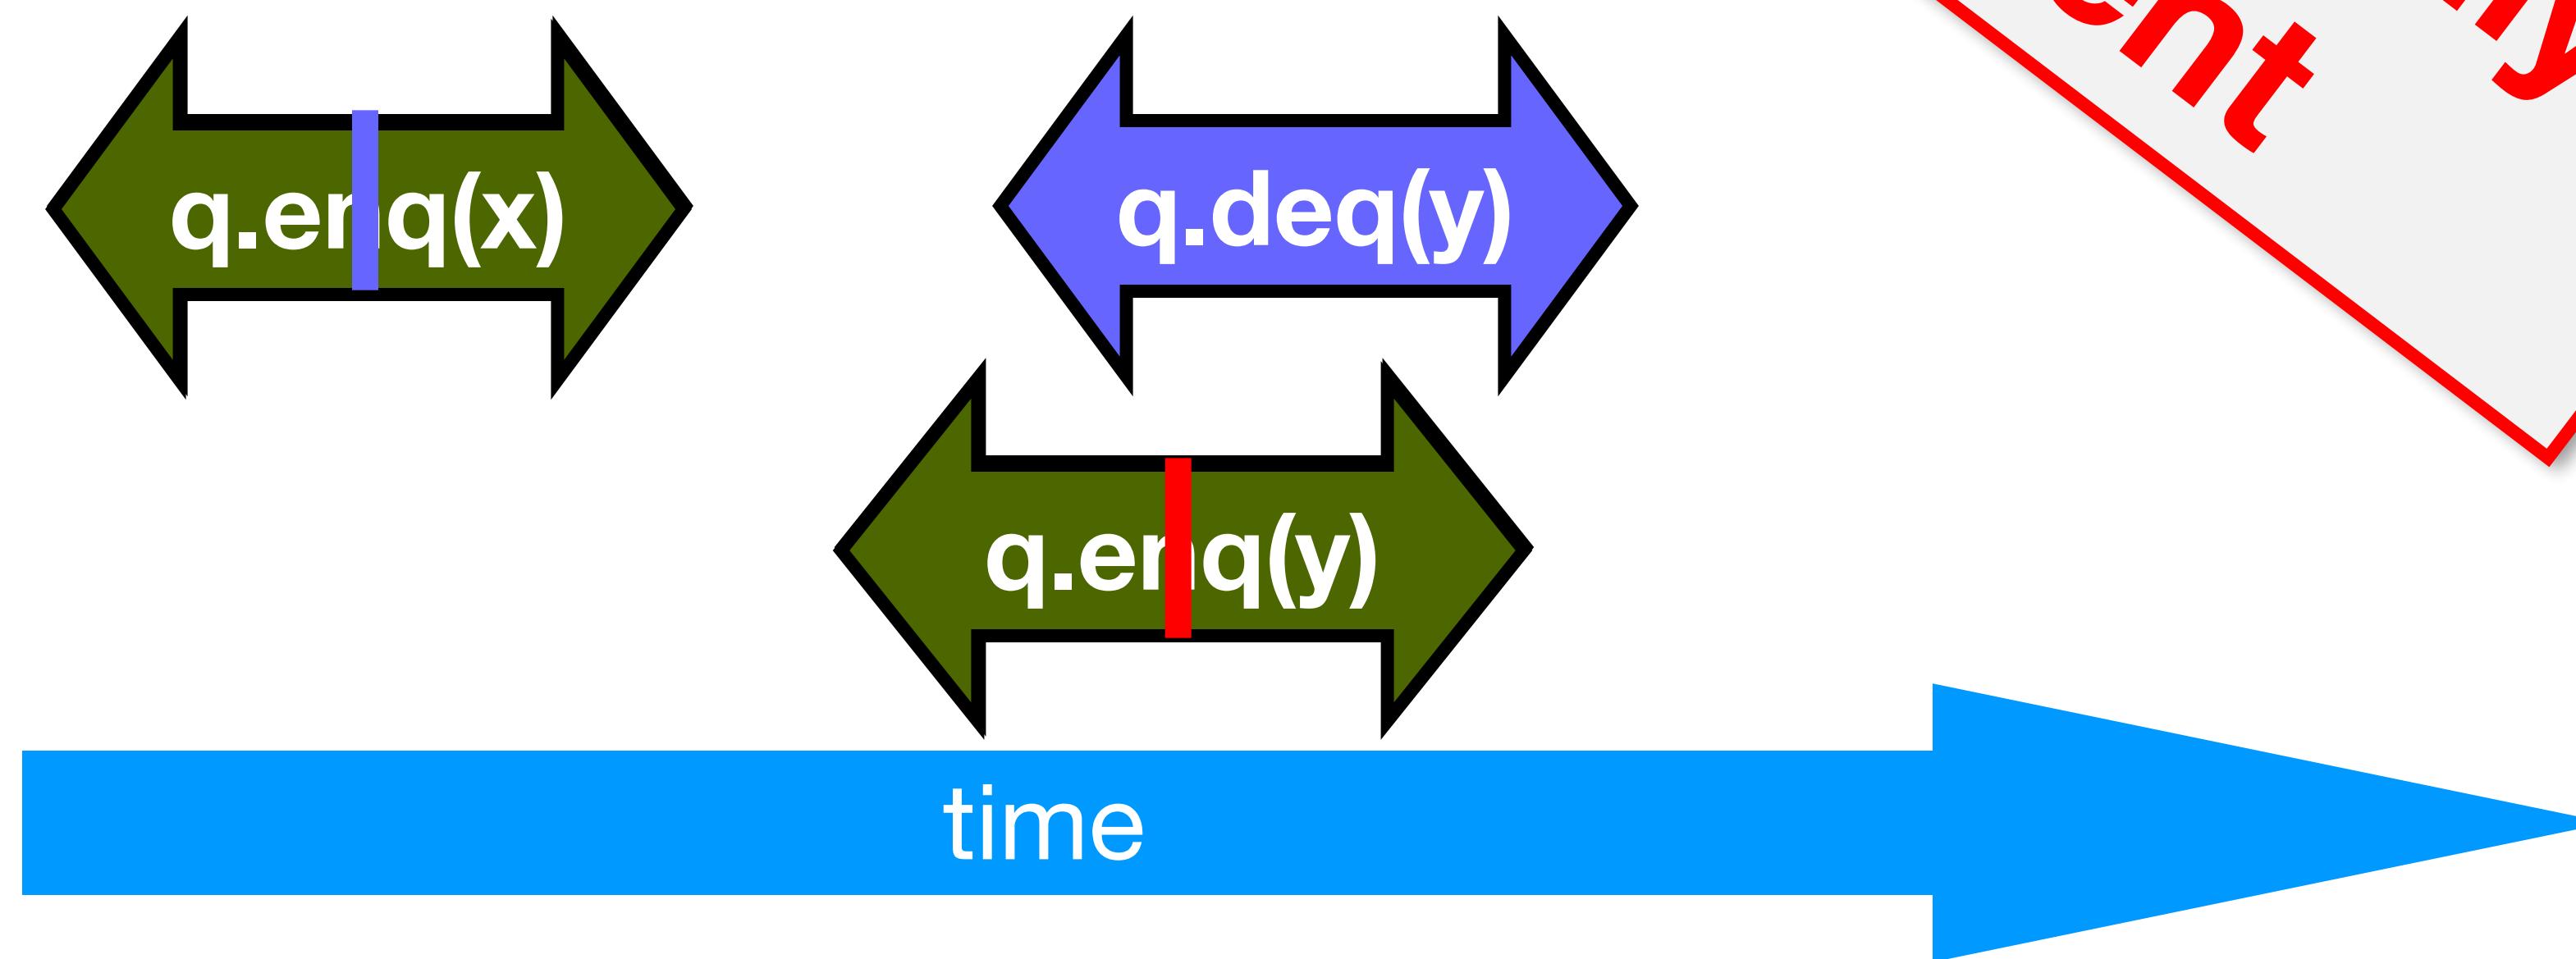
# Example



# Example



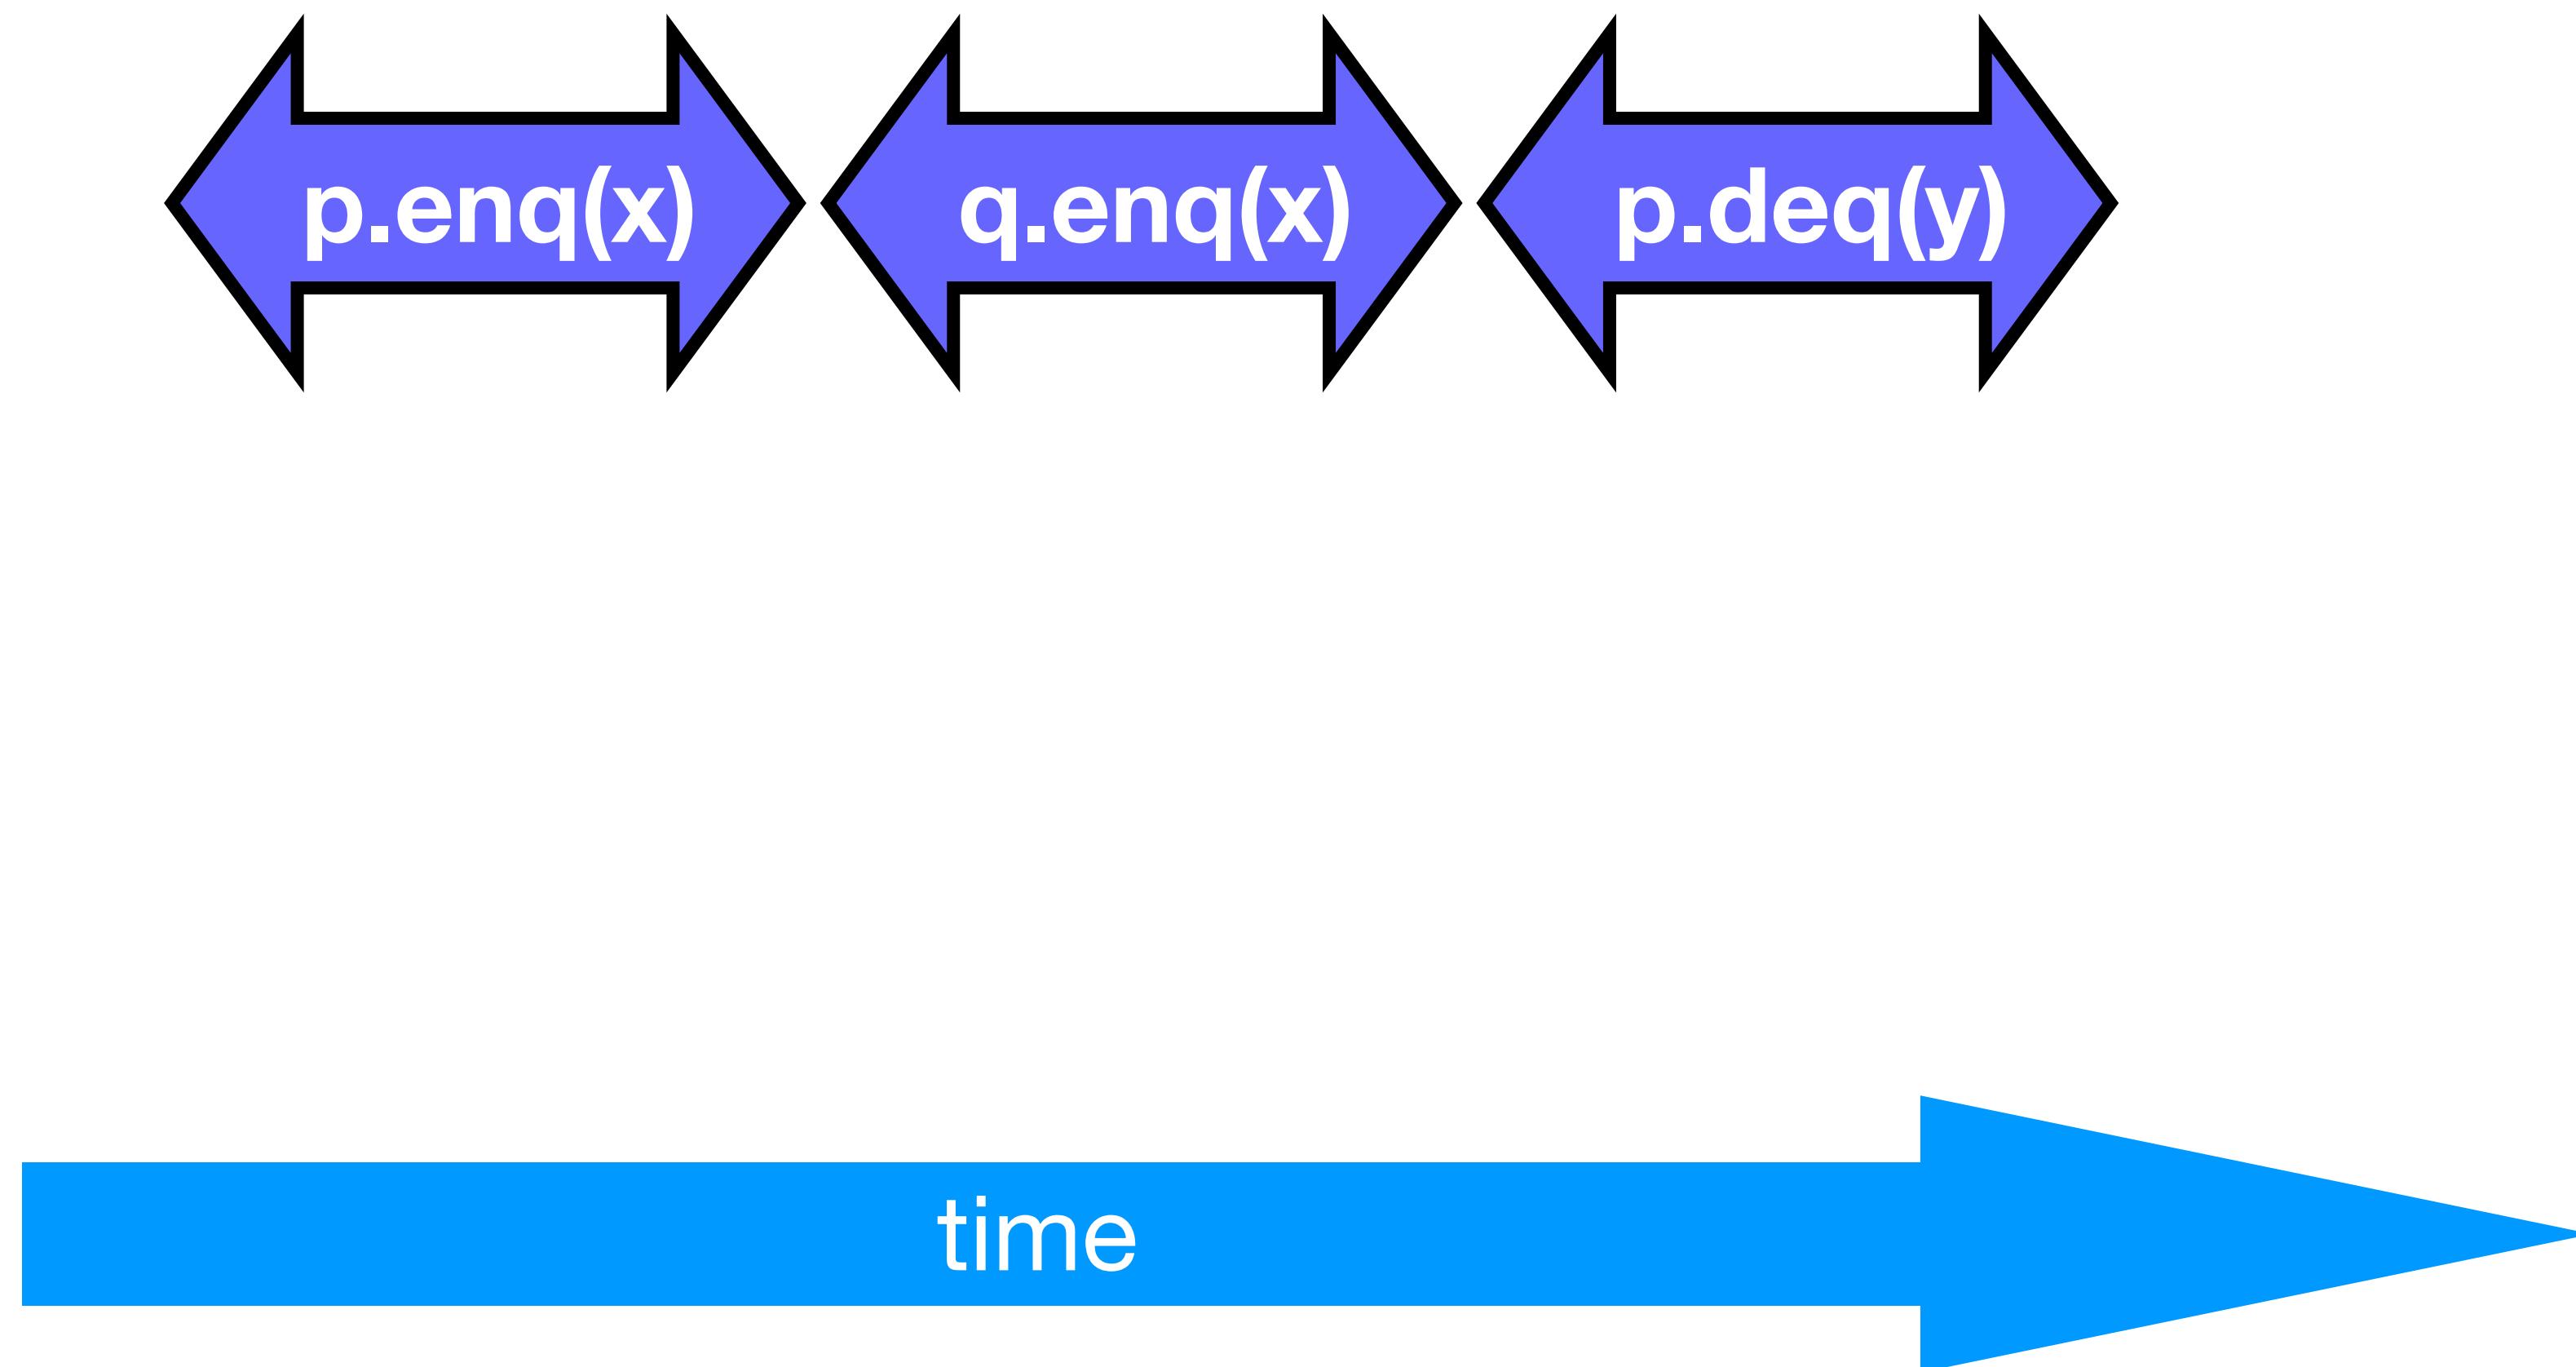
# Example



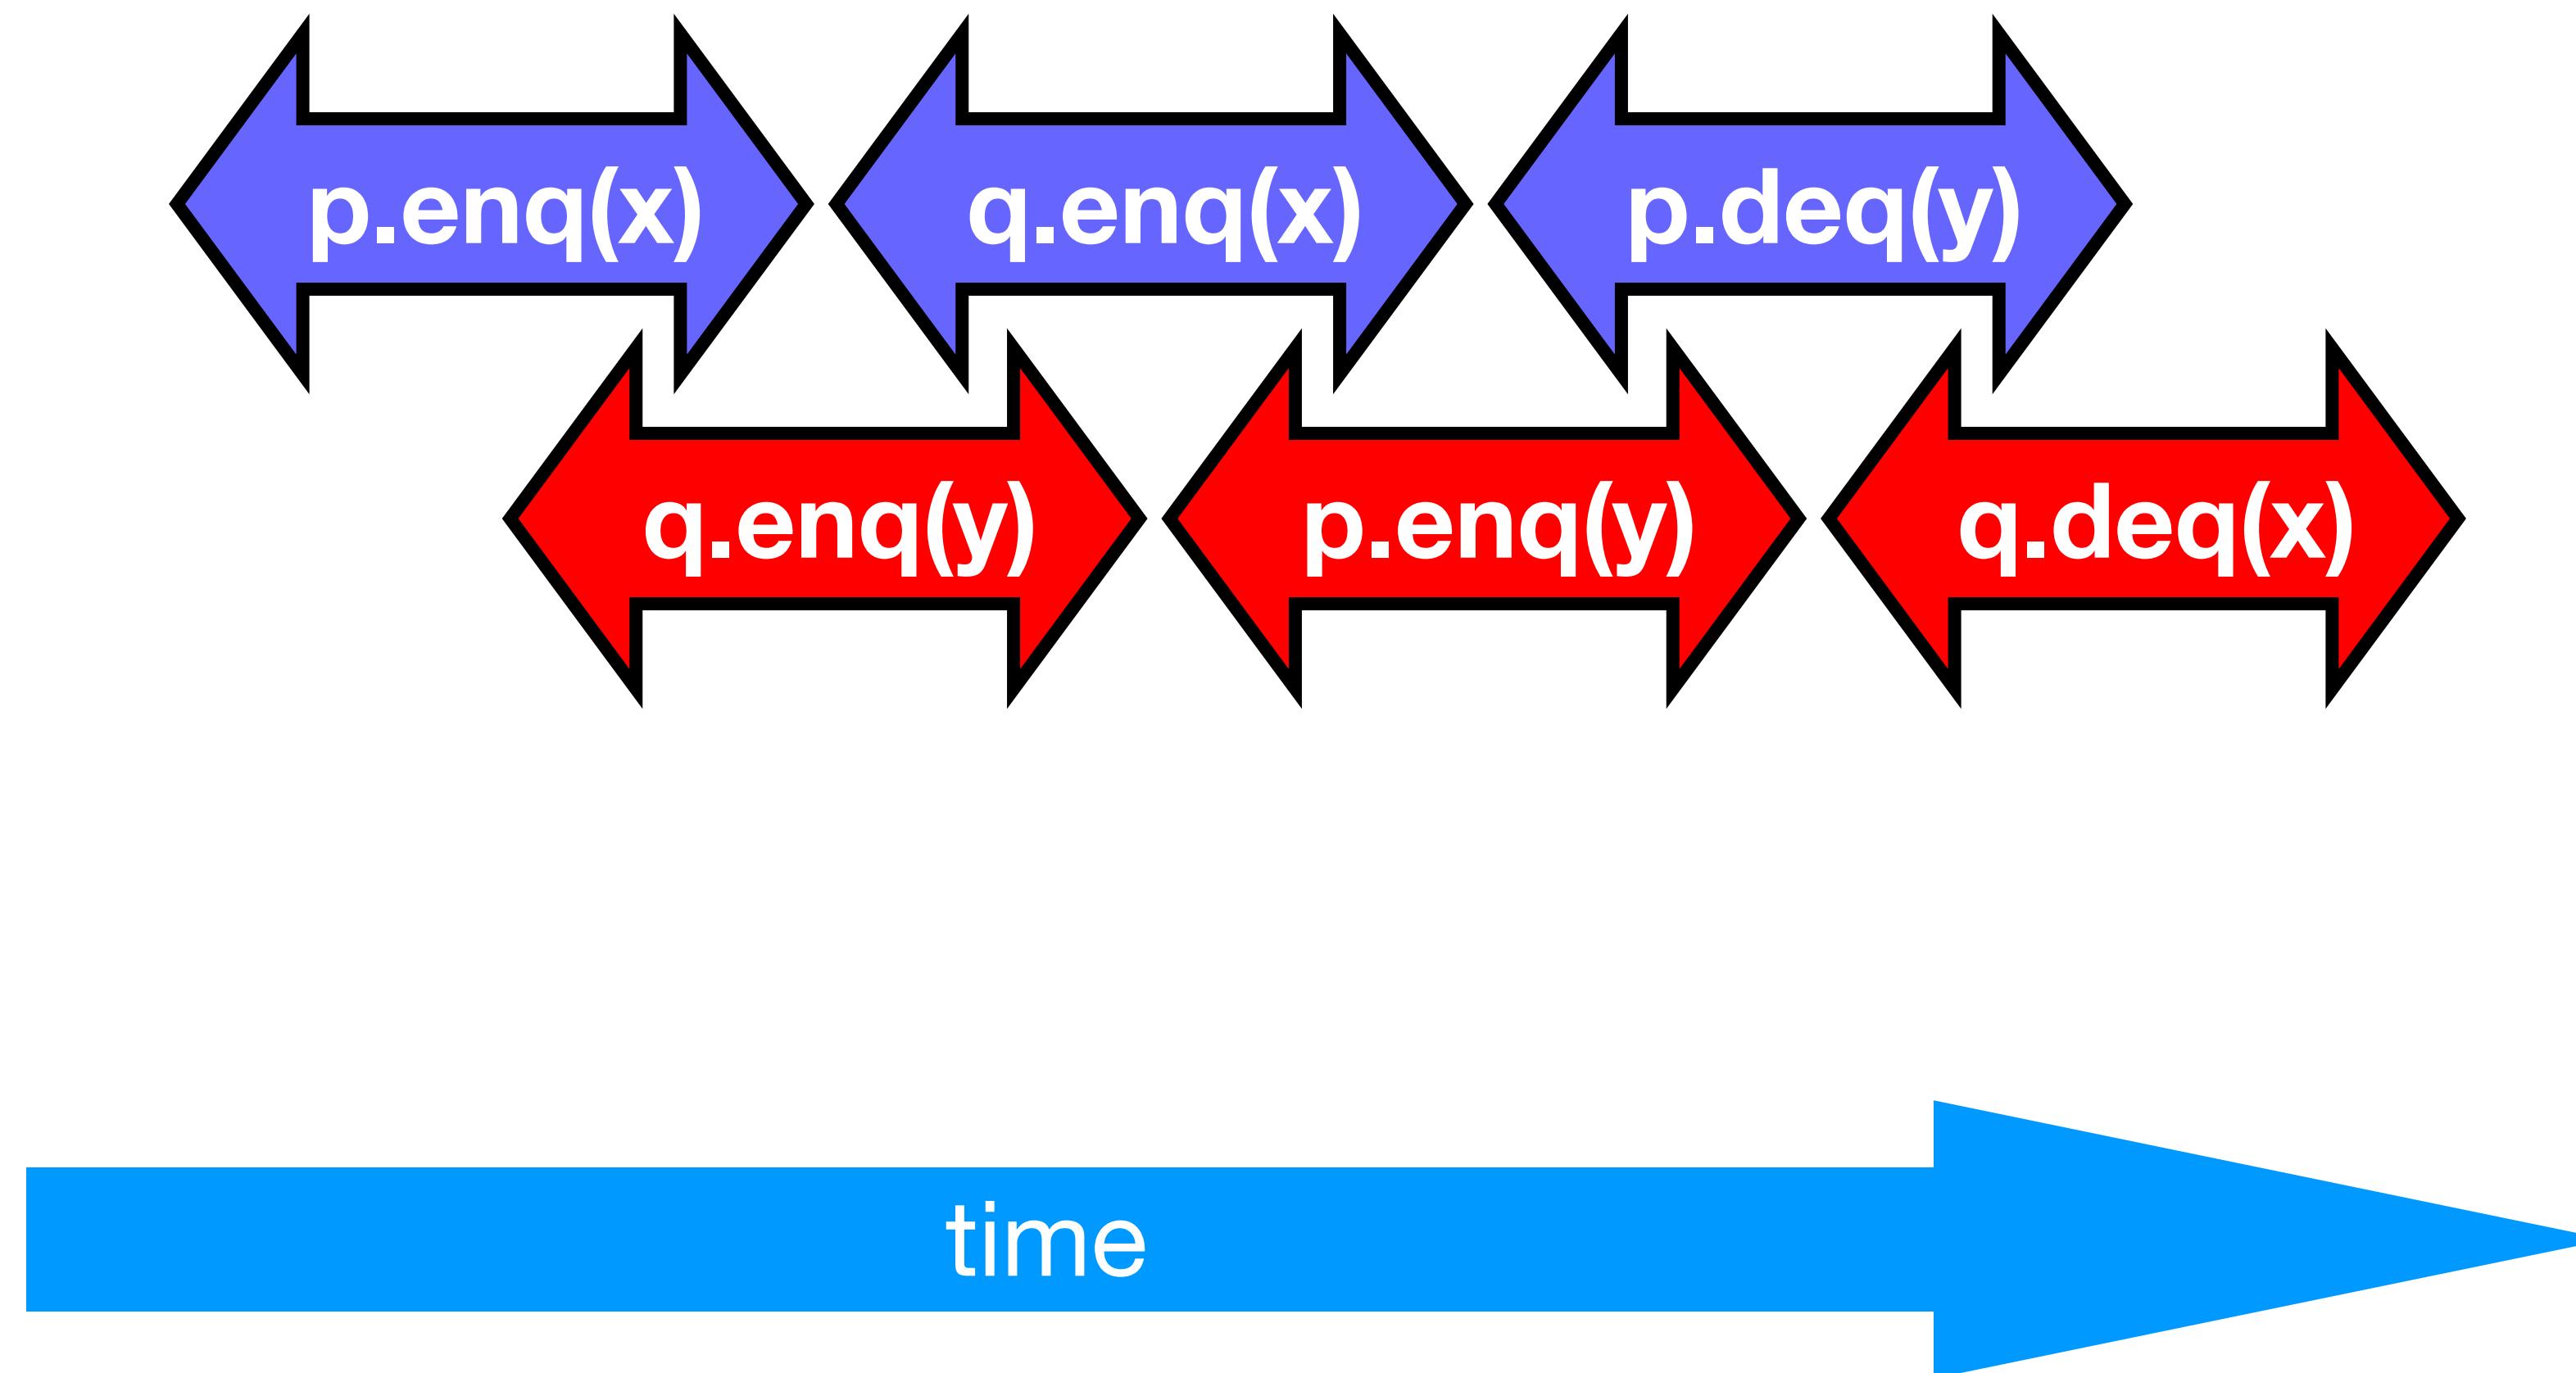
# Theorem

*Sequential Consistency is not Composable*

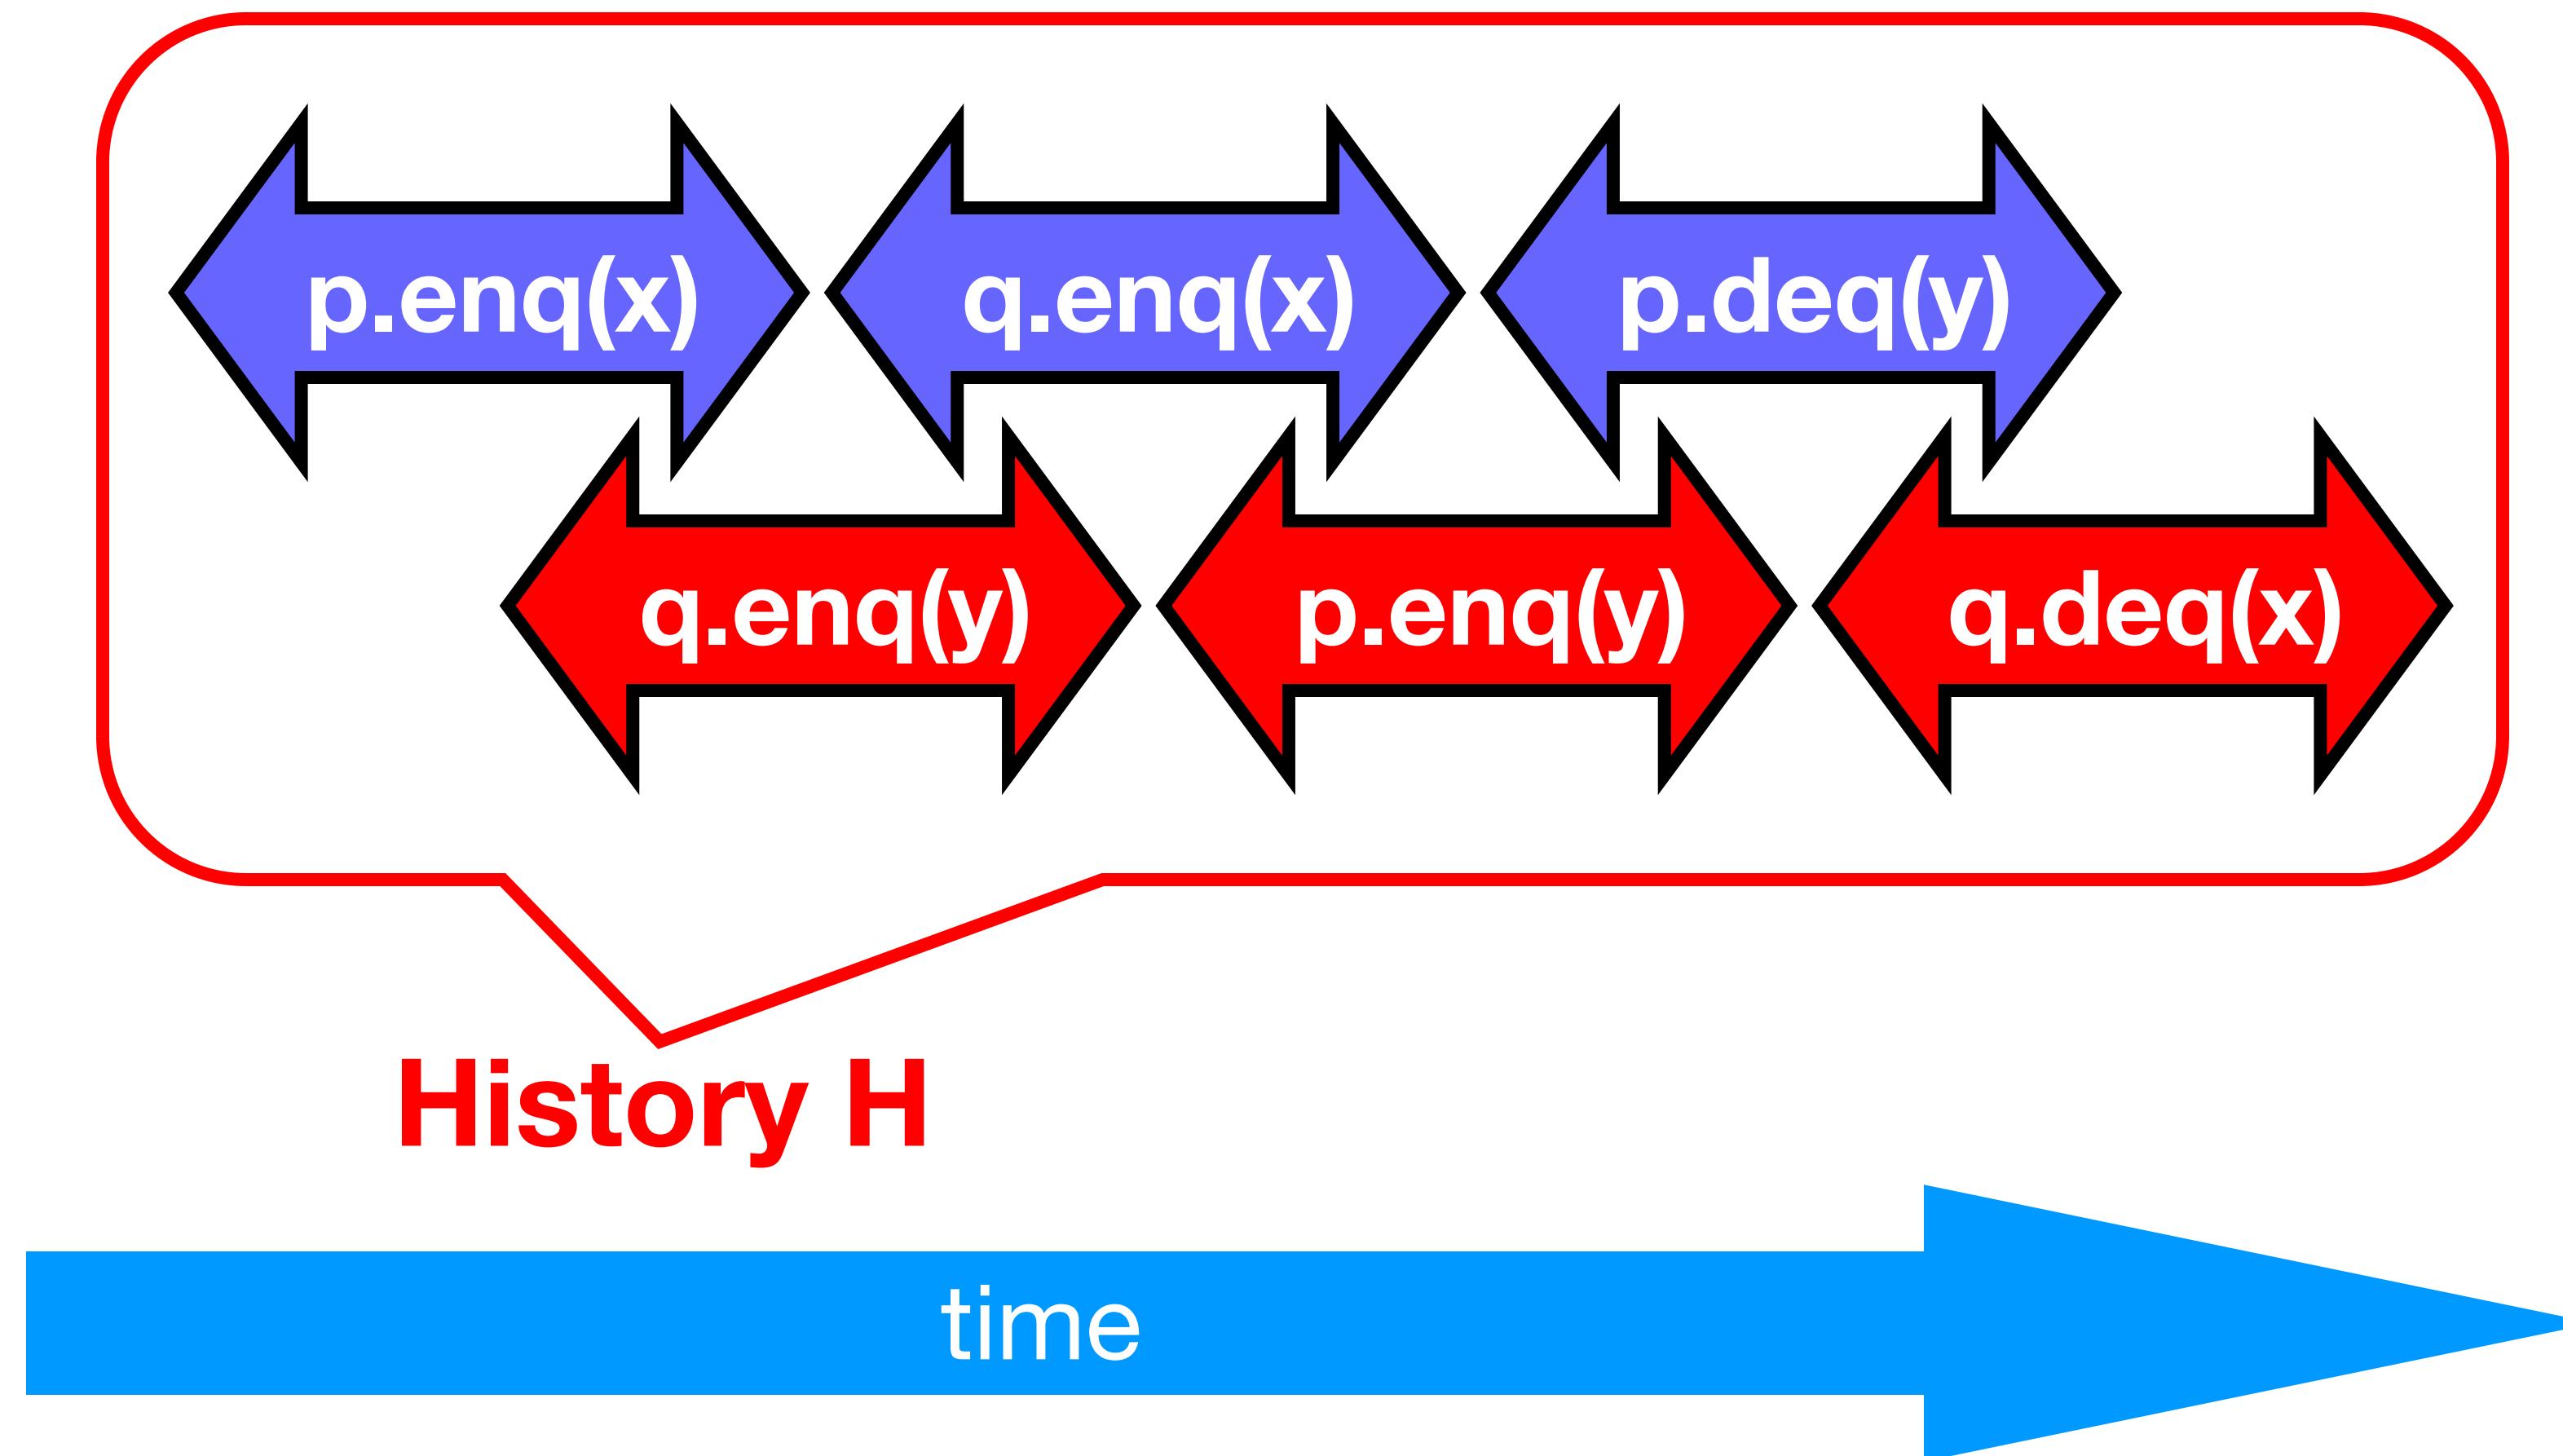
# FIFO Queue Example



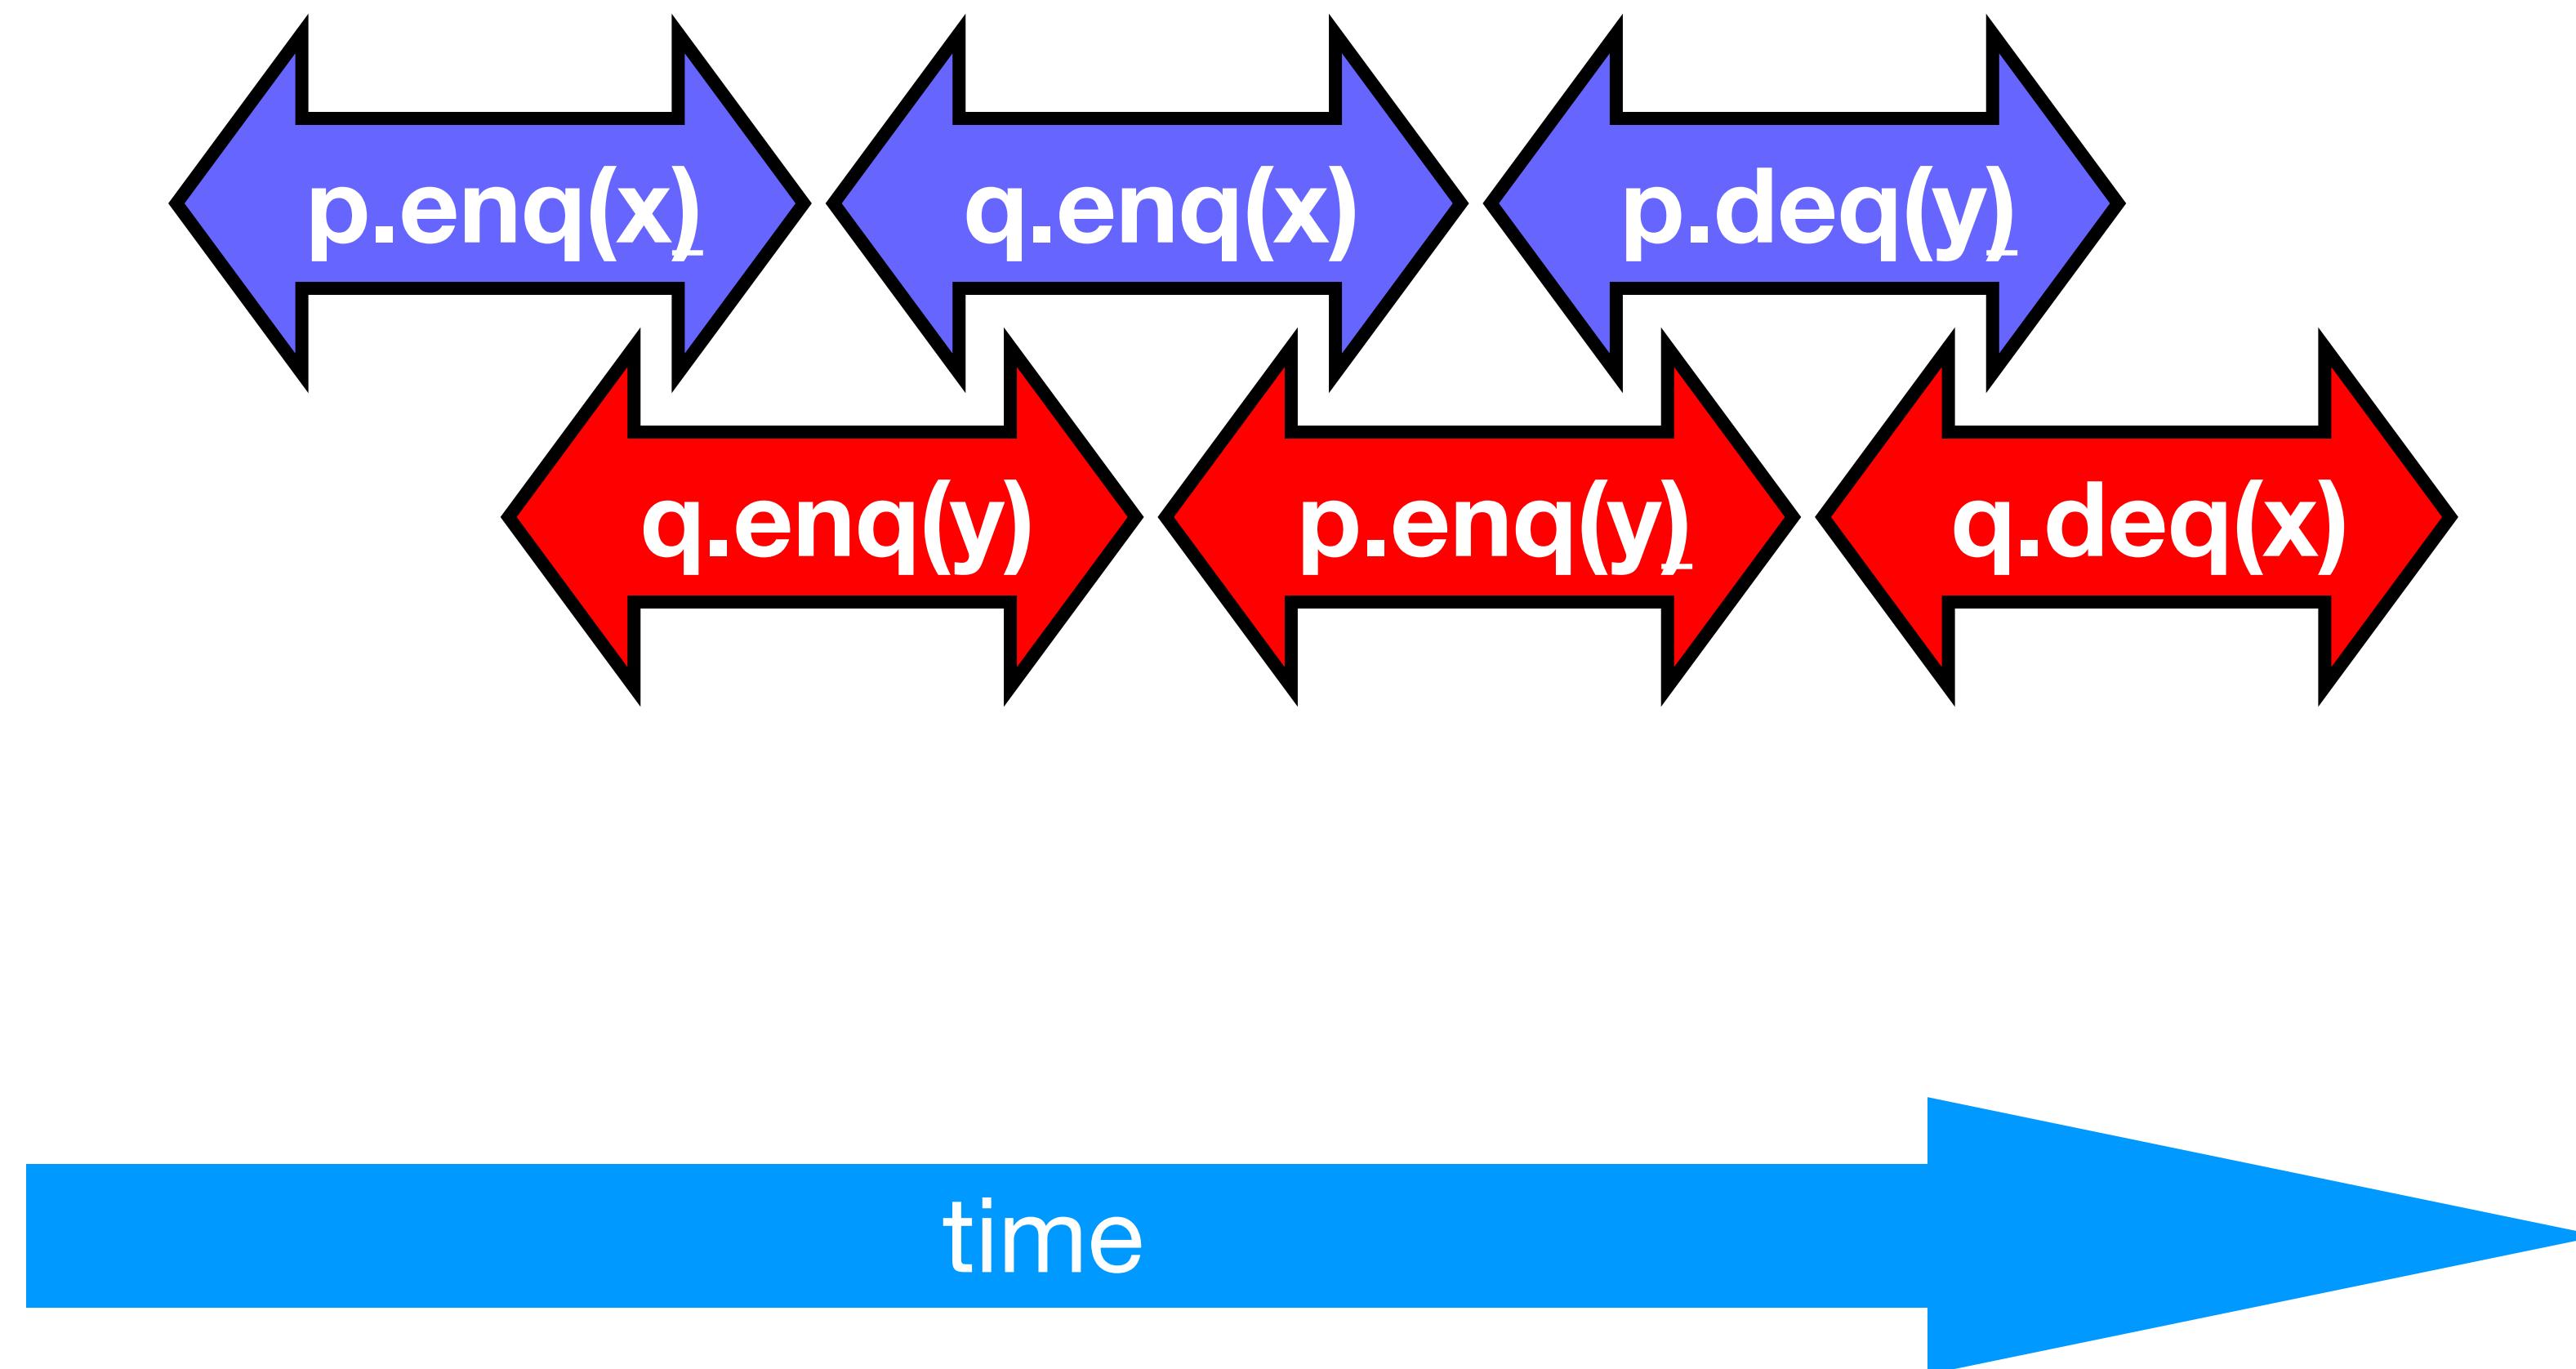
# FIFO Queue Example



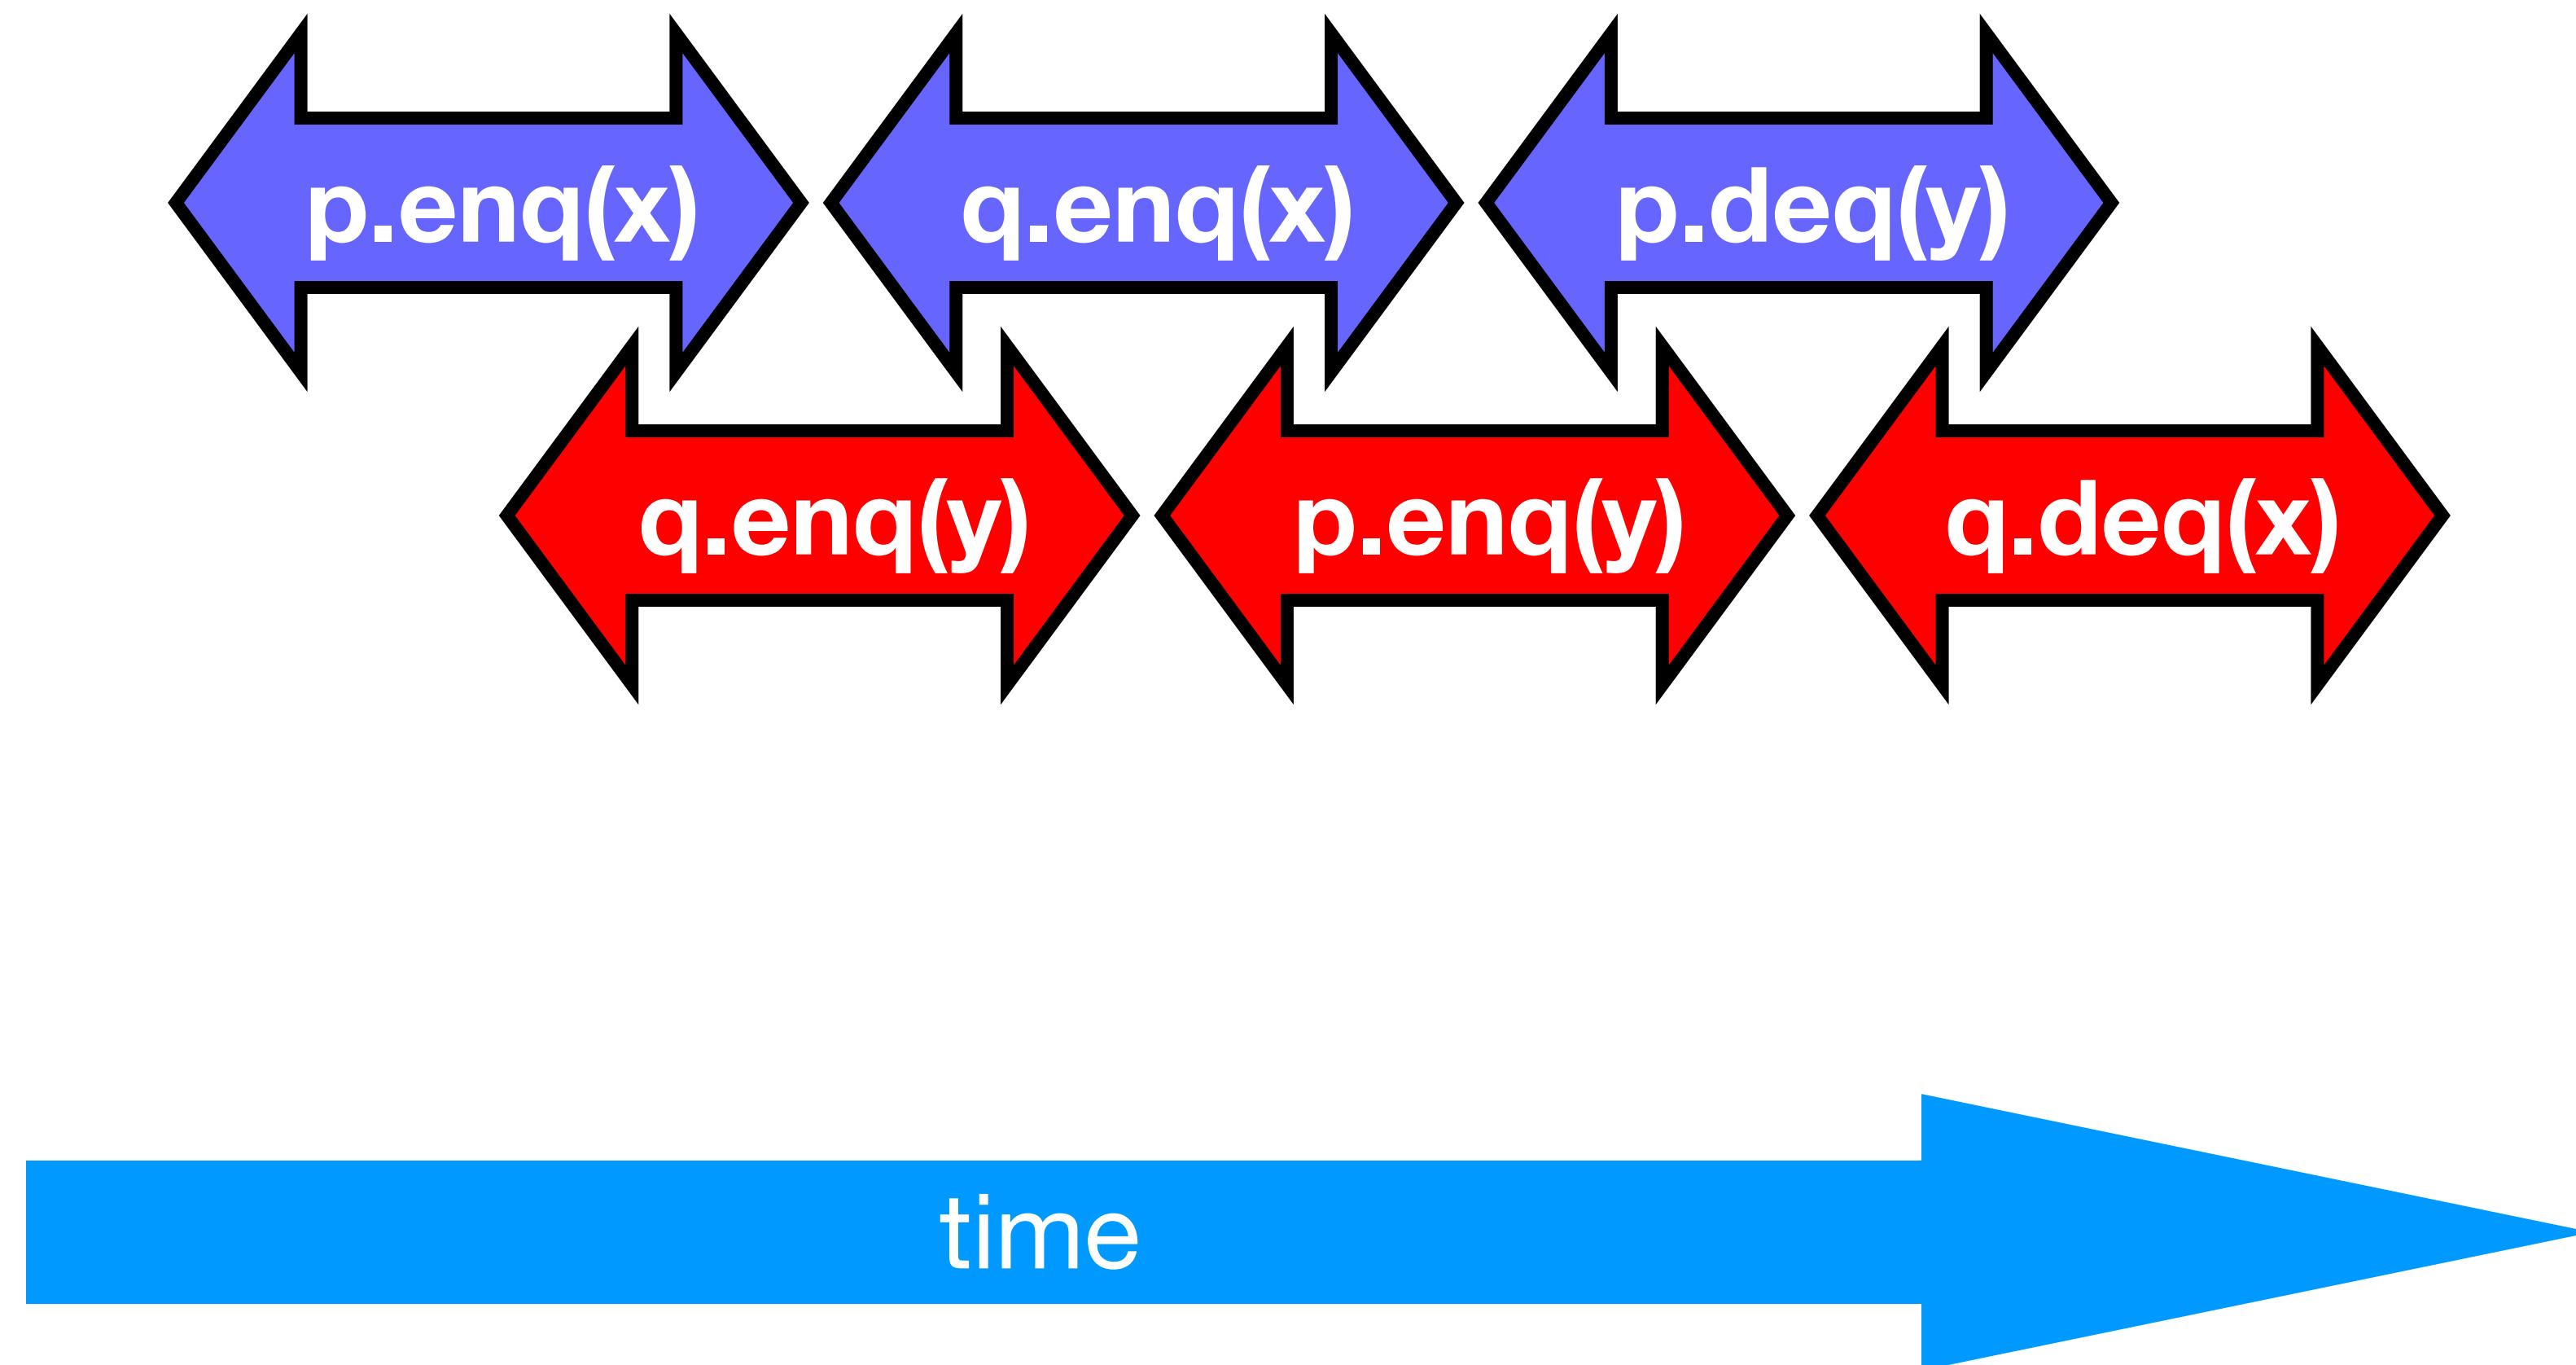
# FIFO Queue Example



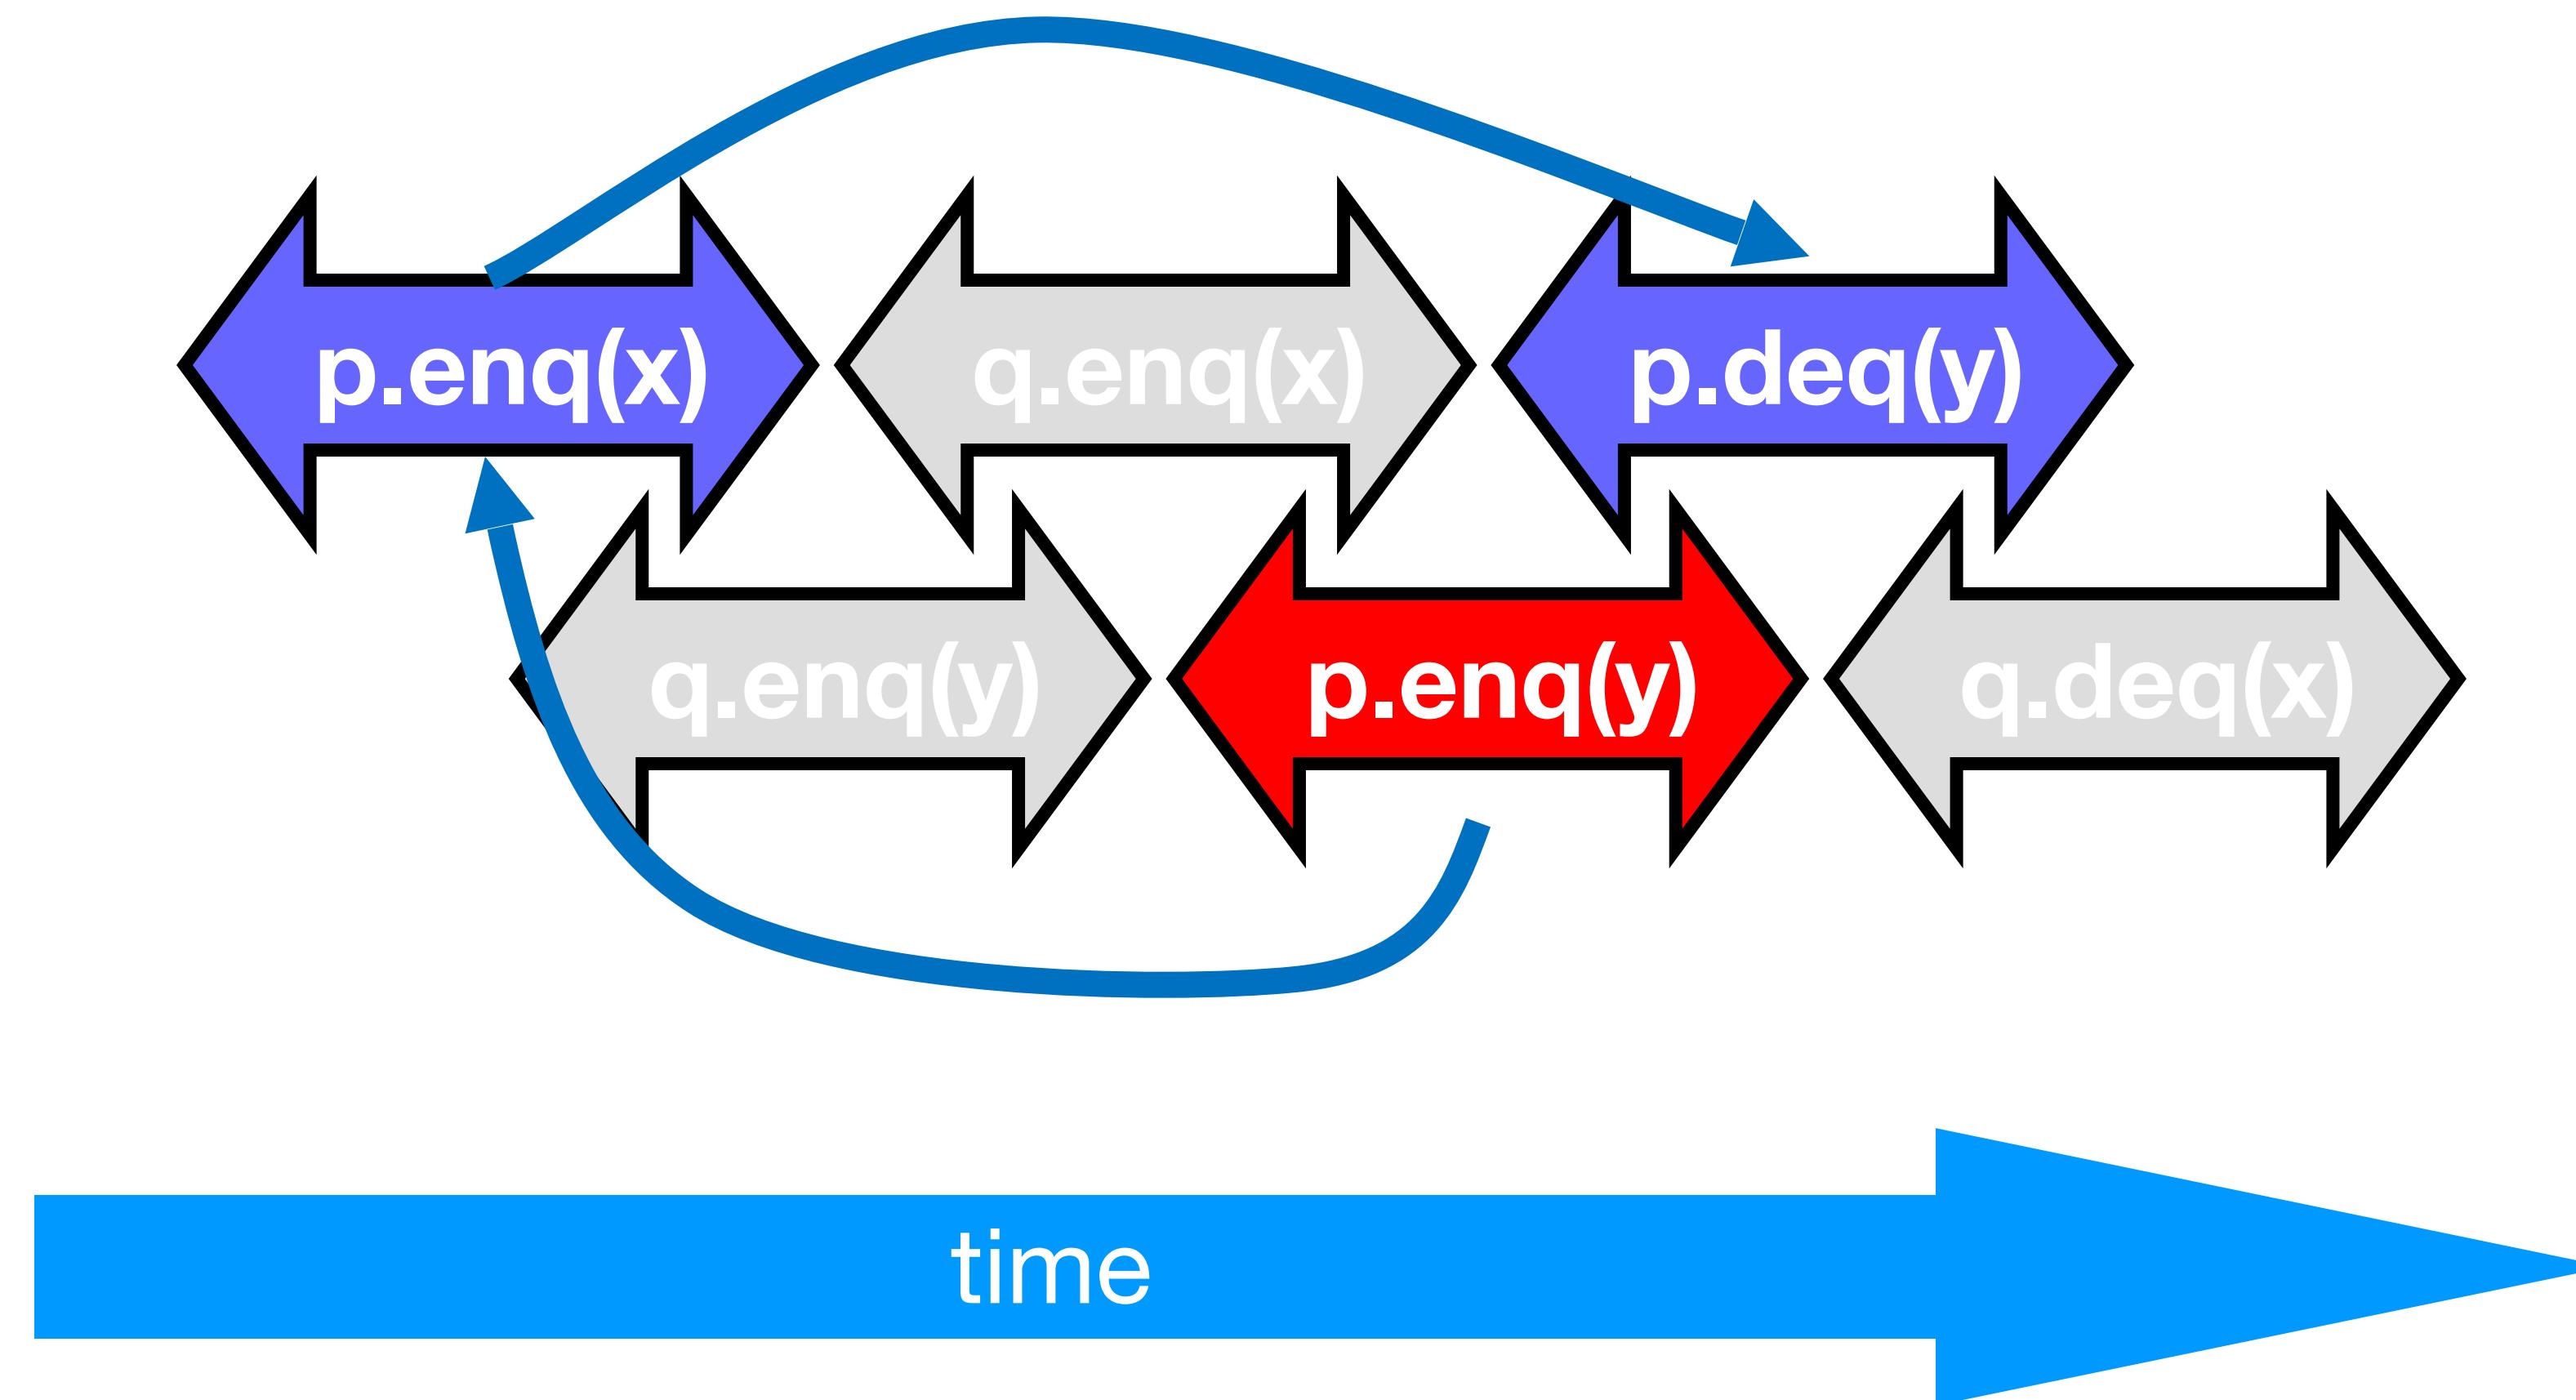
# H|p Sequentially Consistent



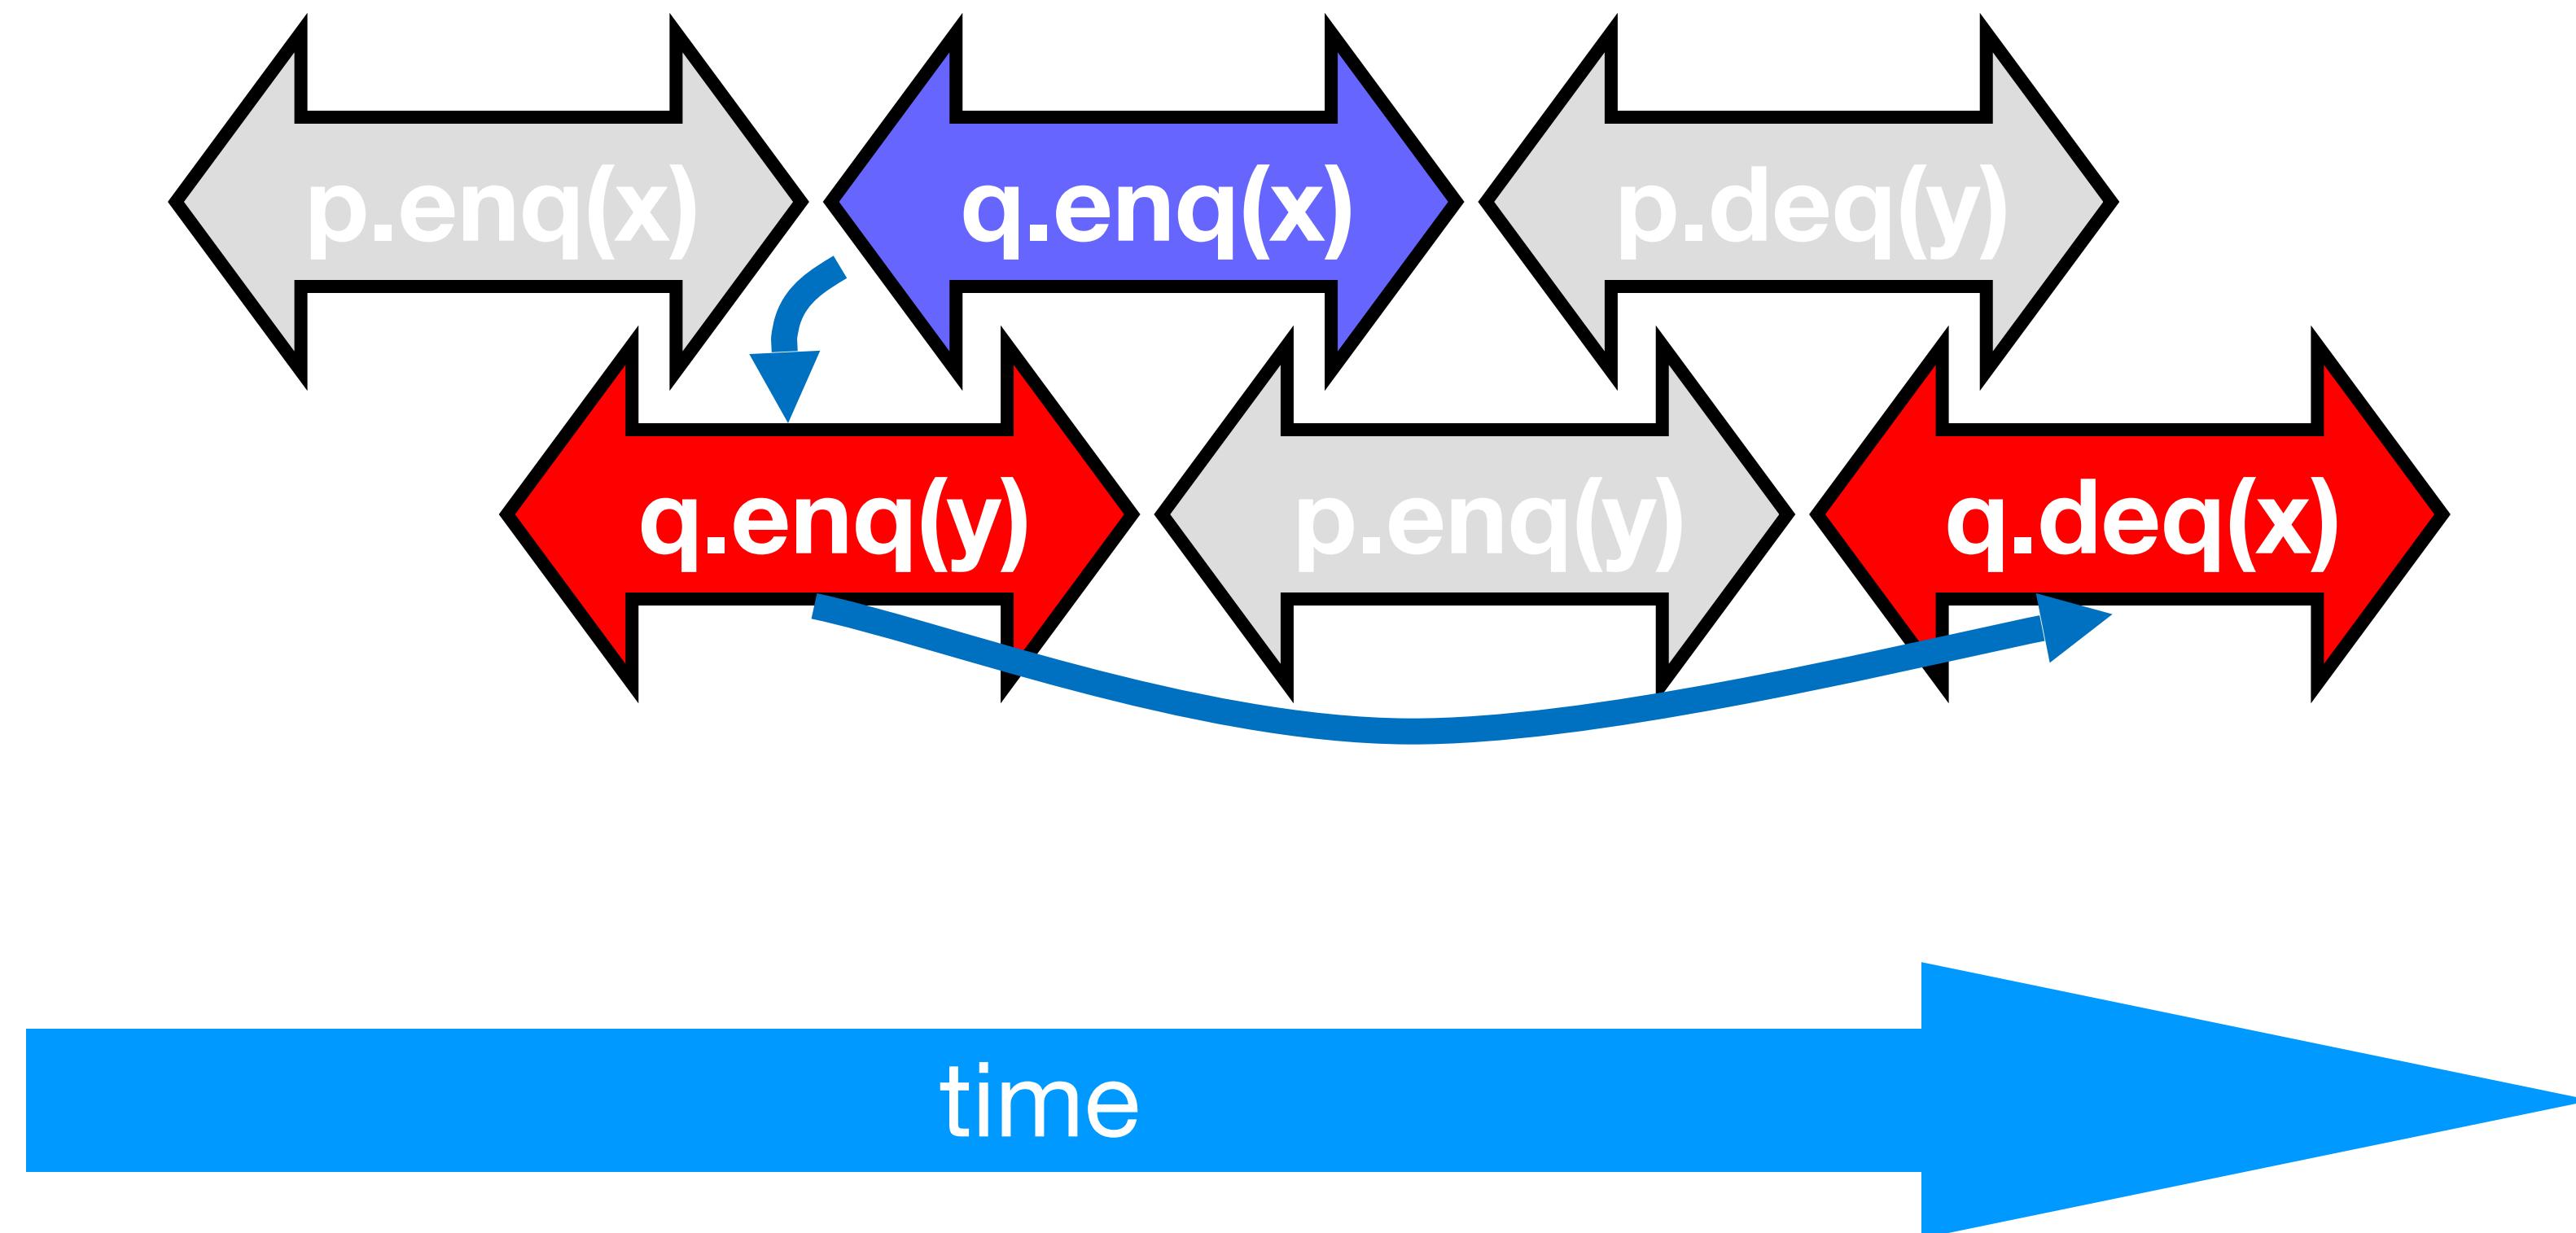
# H|q Sequentially Consistent



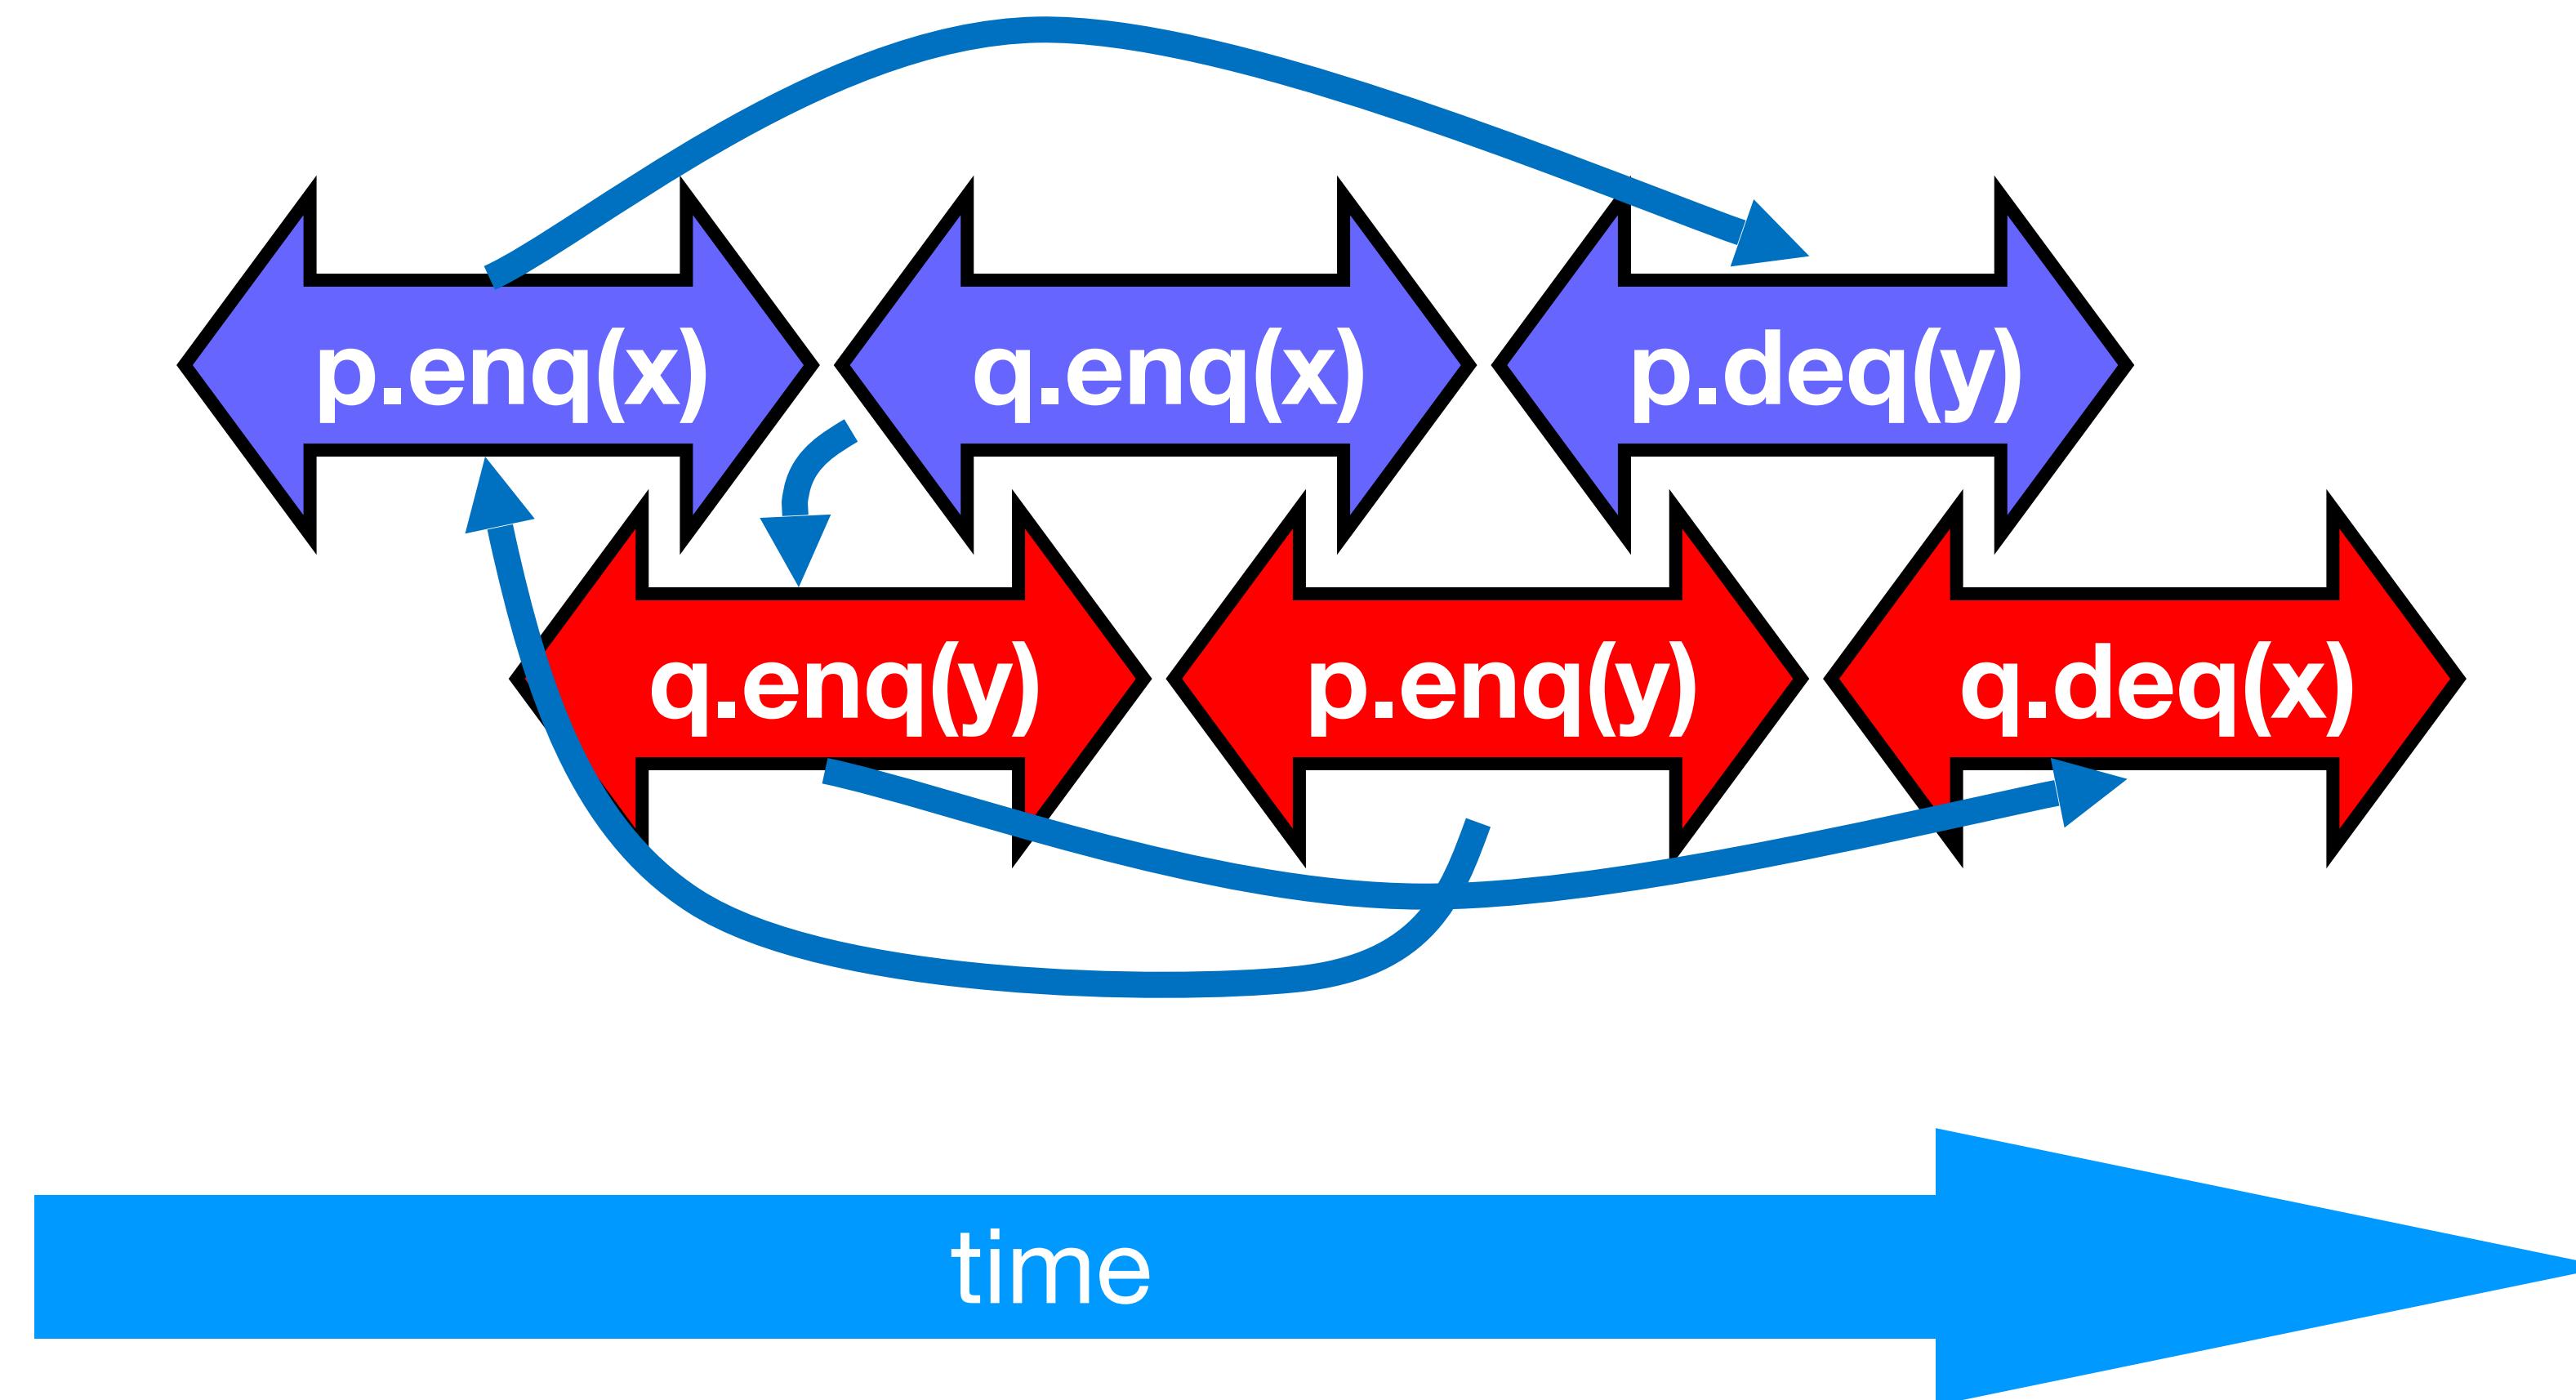
# Ordering imposed by p



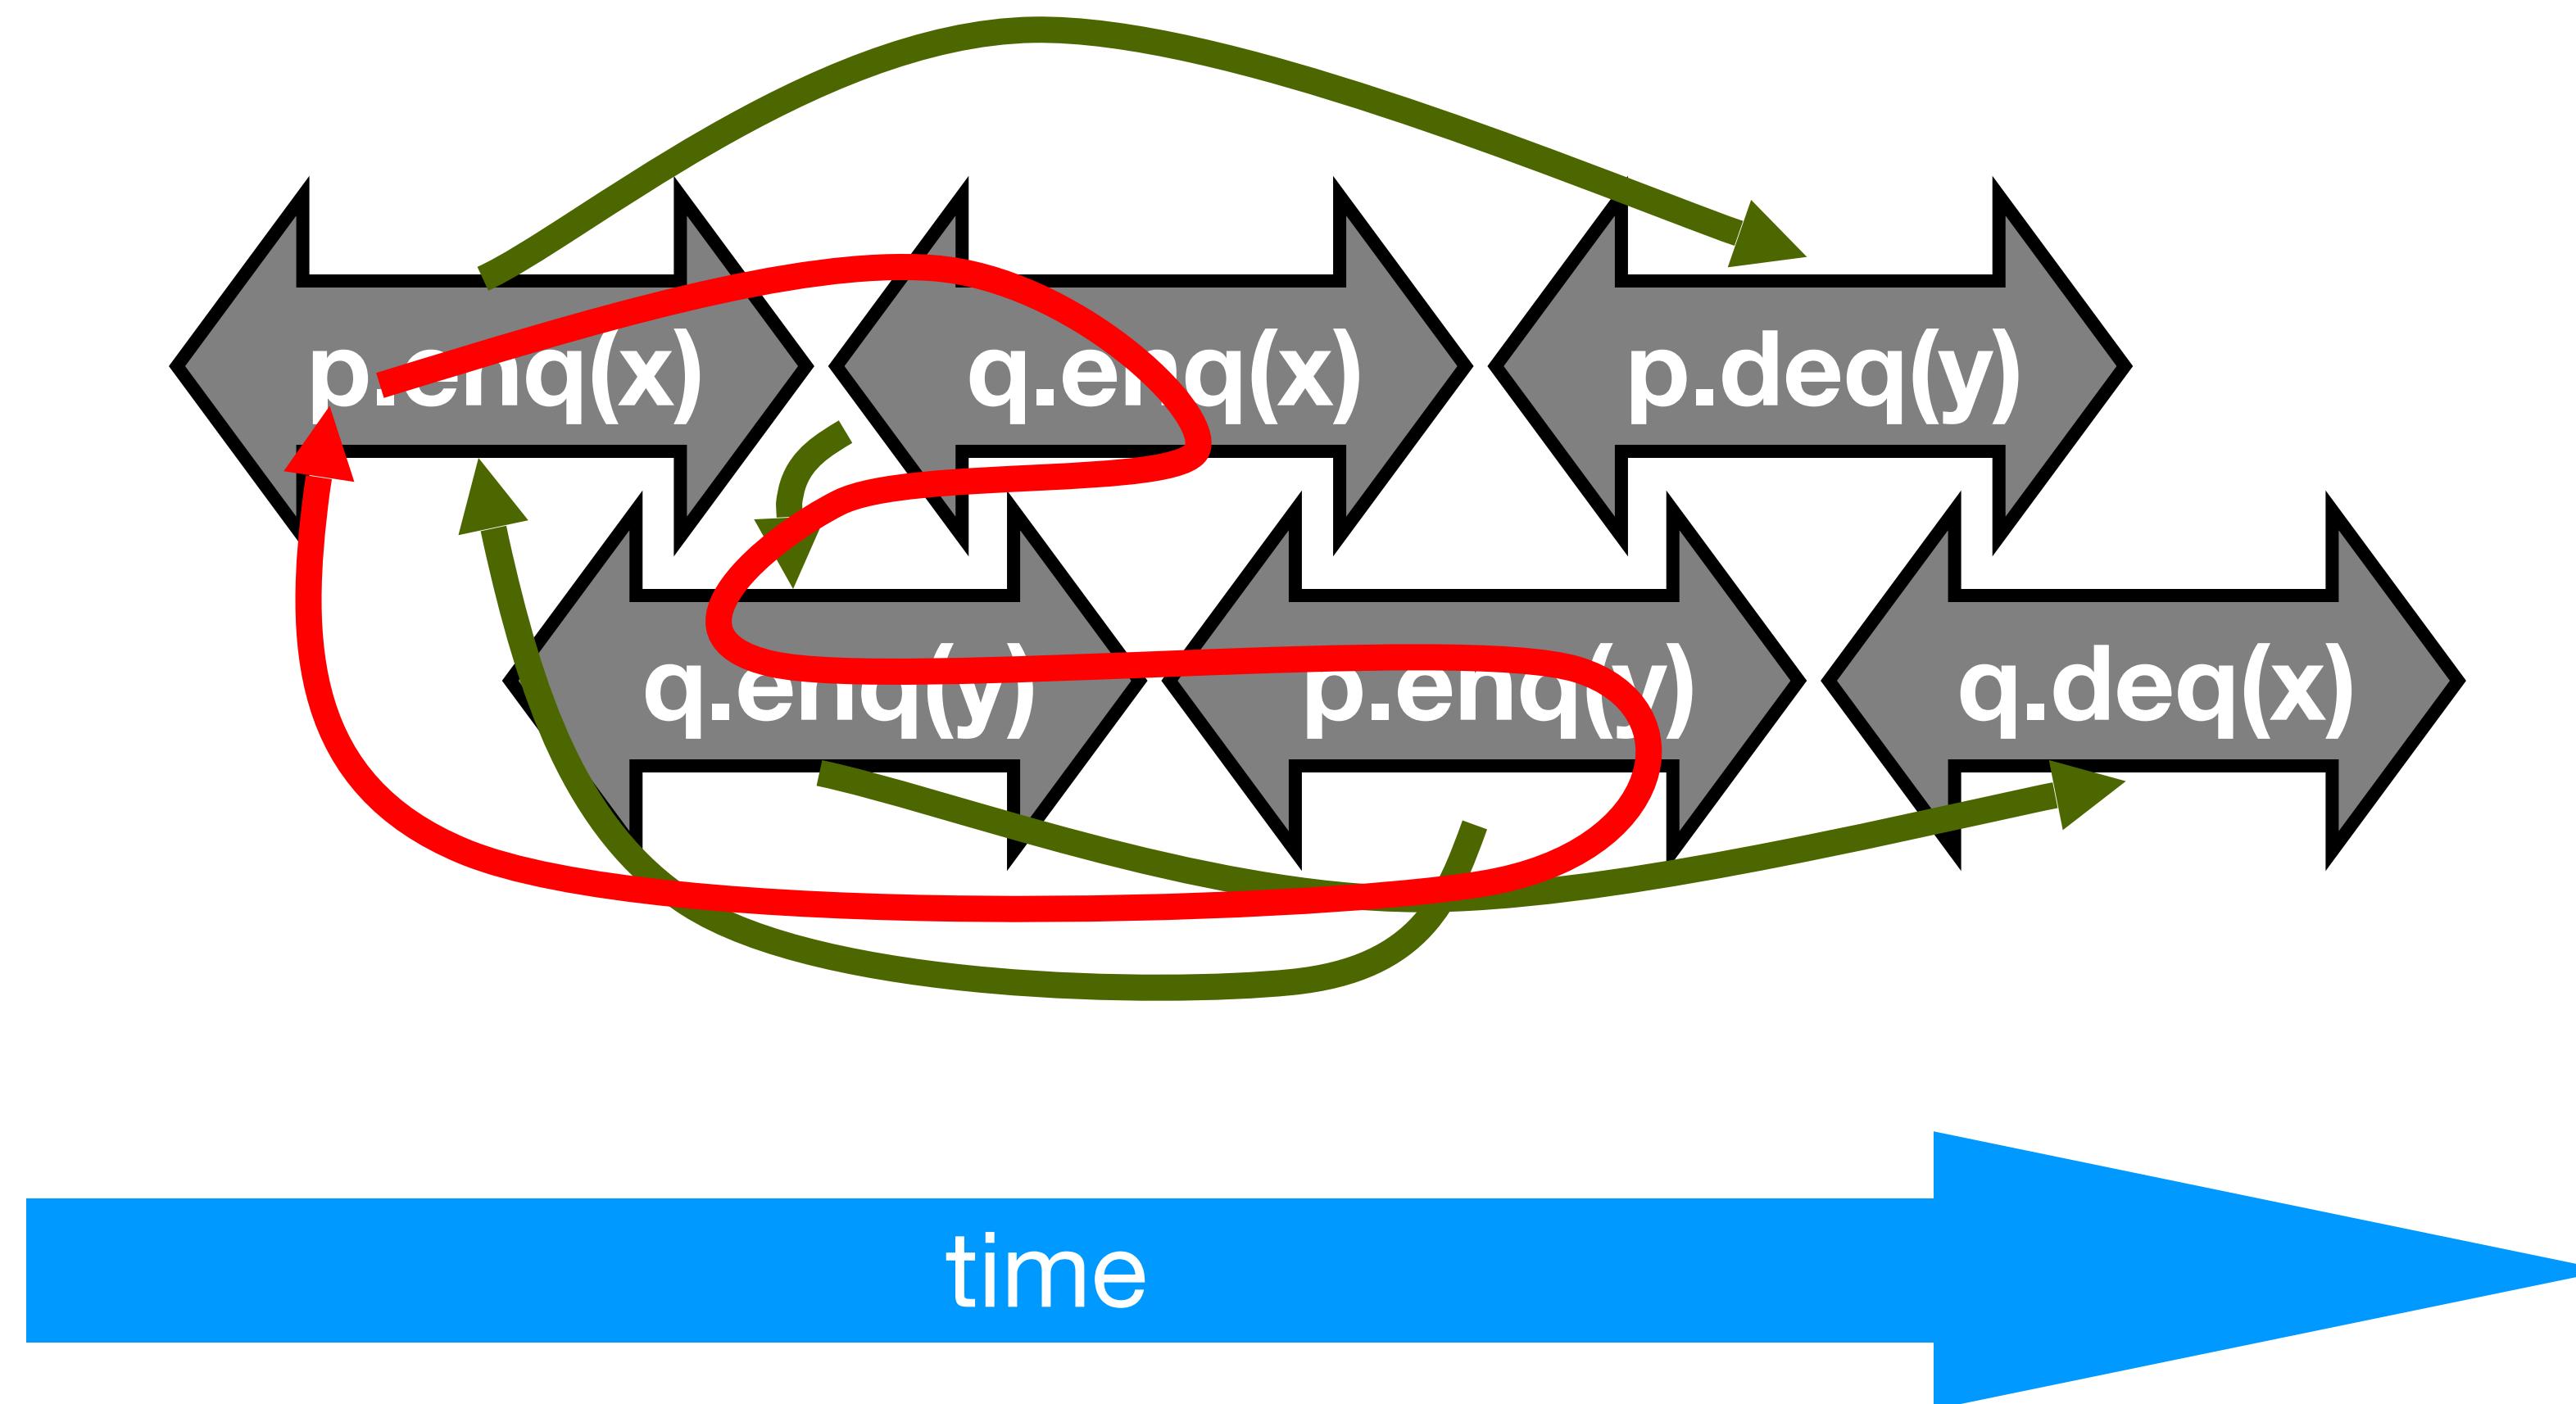
# Ordering imposed by q



# Ordering imposed by both



# Combining orders



# Concurrency Testing

- Linearizability and Sequential Consistency are good specifications for *testing* the correctness of concurrent data structures.
  - Any observed execution must match a sequential execution
  - Can be exploited for pragmatic testing
- See <https://github.com/ocaml-multicore/multicoretests>

# qcheck-lin

- Checks for sequential consistency violations (despite what the name says)
- ***Every sequential consistency violation is a linearizability violation***
- Check that the observed result of a parallel implementation can be observed with a sequential run

**Demo**

# qcheck-stm

- In *qcheck-lin*, what if the implementation is wrong for the sequential program itself?
  - We're only comparing equivalence.
  - Sequential run of a buggy implementation ≡ Parallel run of a buggy implementation
    - ***Is not useful!***
- qcheck-stm
  - Write a state-machine ***model*** of the concurrent object
  - Compare the sequential and parallel executions of the implementation against the state machine model
  - ***More work!***

**Demo**

# Summary

- ***Linearizability***
  - The operation takes effect instantaneously between the invocation and the response
  - Uses sequential specification, locality implies composability
- ***Sequential Consistency***
  - Linearizability without real-time ordering
  - Not composable
  - Harder to work with
  - Useful to reason about hardware models (next lecture)
- We will use ***linearizability*** as our consistency condition for reasoning about objects

# Progress

- We saw an implementation whose methods were lock-based (deadlock-free)
- We saw an implementation whose methods did not use locks (lock-free)
- How do they relate?

# Progress Conditions

- **Deadlock-free:** some thread trying to acquire the lock eventually succeeds.
- **Starvation-free:** every thread trying to acquire the lock eventually succeeds.
- **Lock-free:** some thread calling a method eventually returns.
- **Wait-free:** every thread calling a method eventually returns.

	Non-Blocking	Blocking	
Everyone makes progress	Wait-free	Starvation-free	<i>We will look at linearizable blocking and non-blocking implementations of objects.</i>
Someone makes progress	Lock-free	Deadlock-free	



## This work is licensed under a [Creative Commons Attribution-ShareAlike 2.5 License](#).

- **You are free:**
  - **to Share** – to copy, distribute and transmit the work
  - **to Remix** – to adapt the work
- **Under the following conditions:**
  - **Attribution.** You must attribute the work to “The Art of Multiprocessor Programming” (but not in any way that suggests that the authors endorse you or your use of the work).
  - **Share Alike.** If you alter, transform, or build upon this work, you may distribute the resulting work only under the same, similar or a compatible license.
- For any reuse or distribution, you must make clear to others the license terms of this work. The best way to do this is with a link to
  - <http://creativecommons.org/licenses/by-sa/3.0/>.
- Any of the above conditions can be waived if you get permission from the copyright holder.
- Nothing in this license impairs or restricts the author's moral rights.