

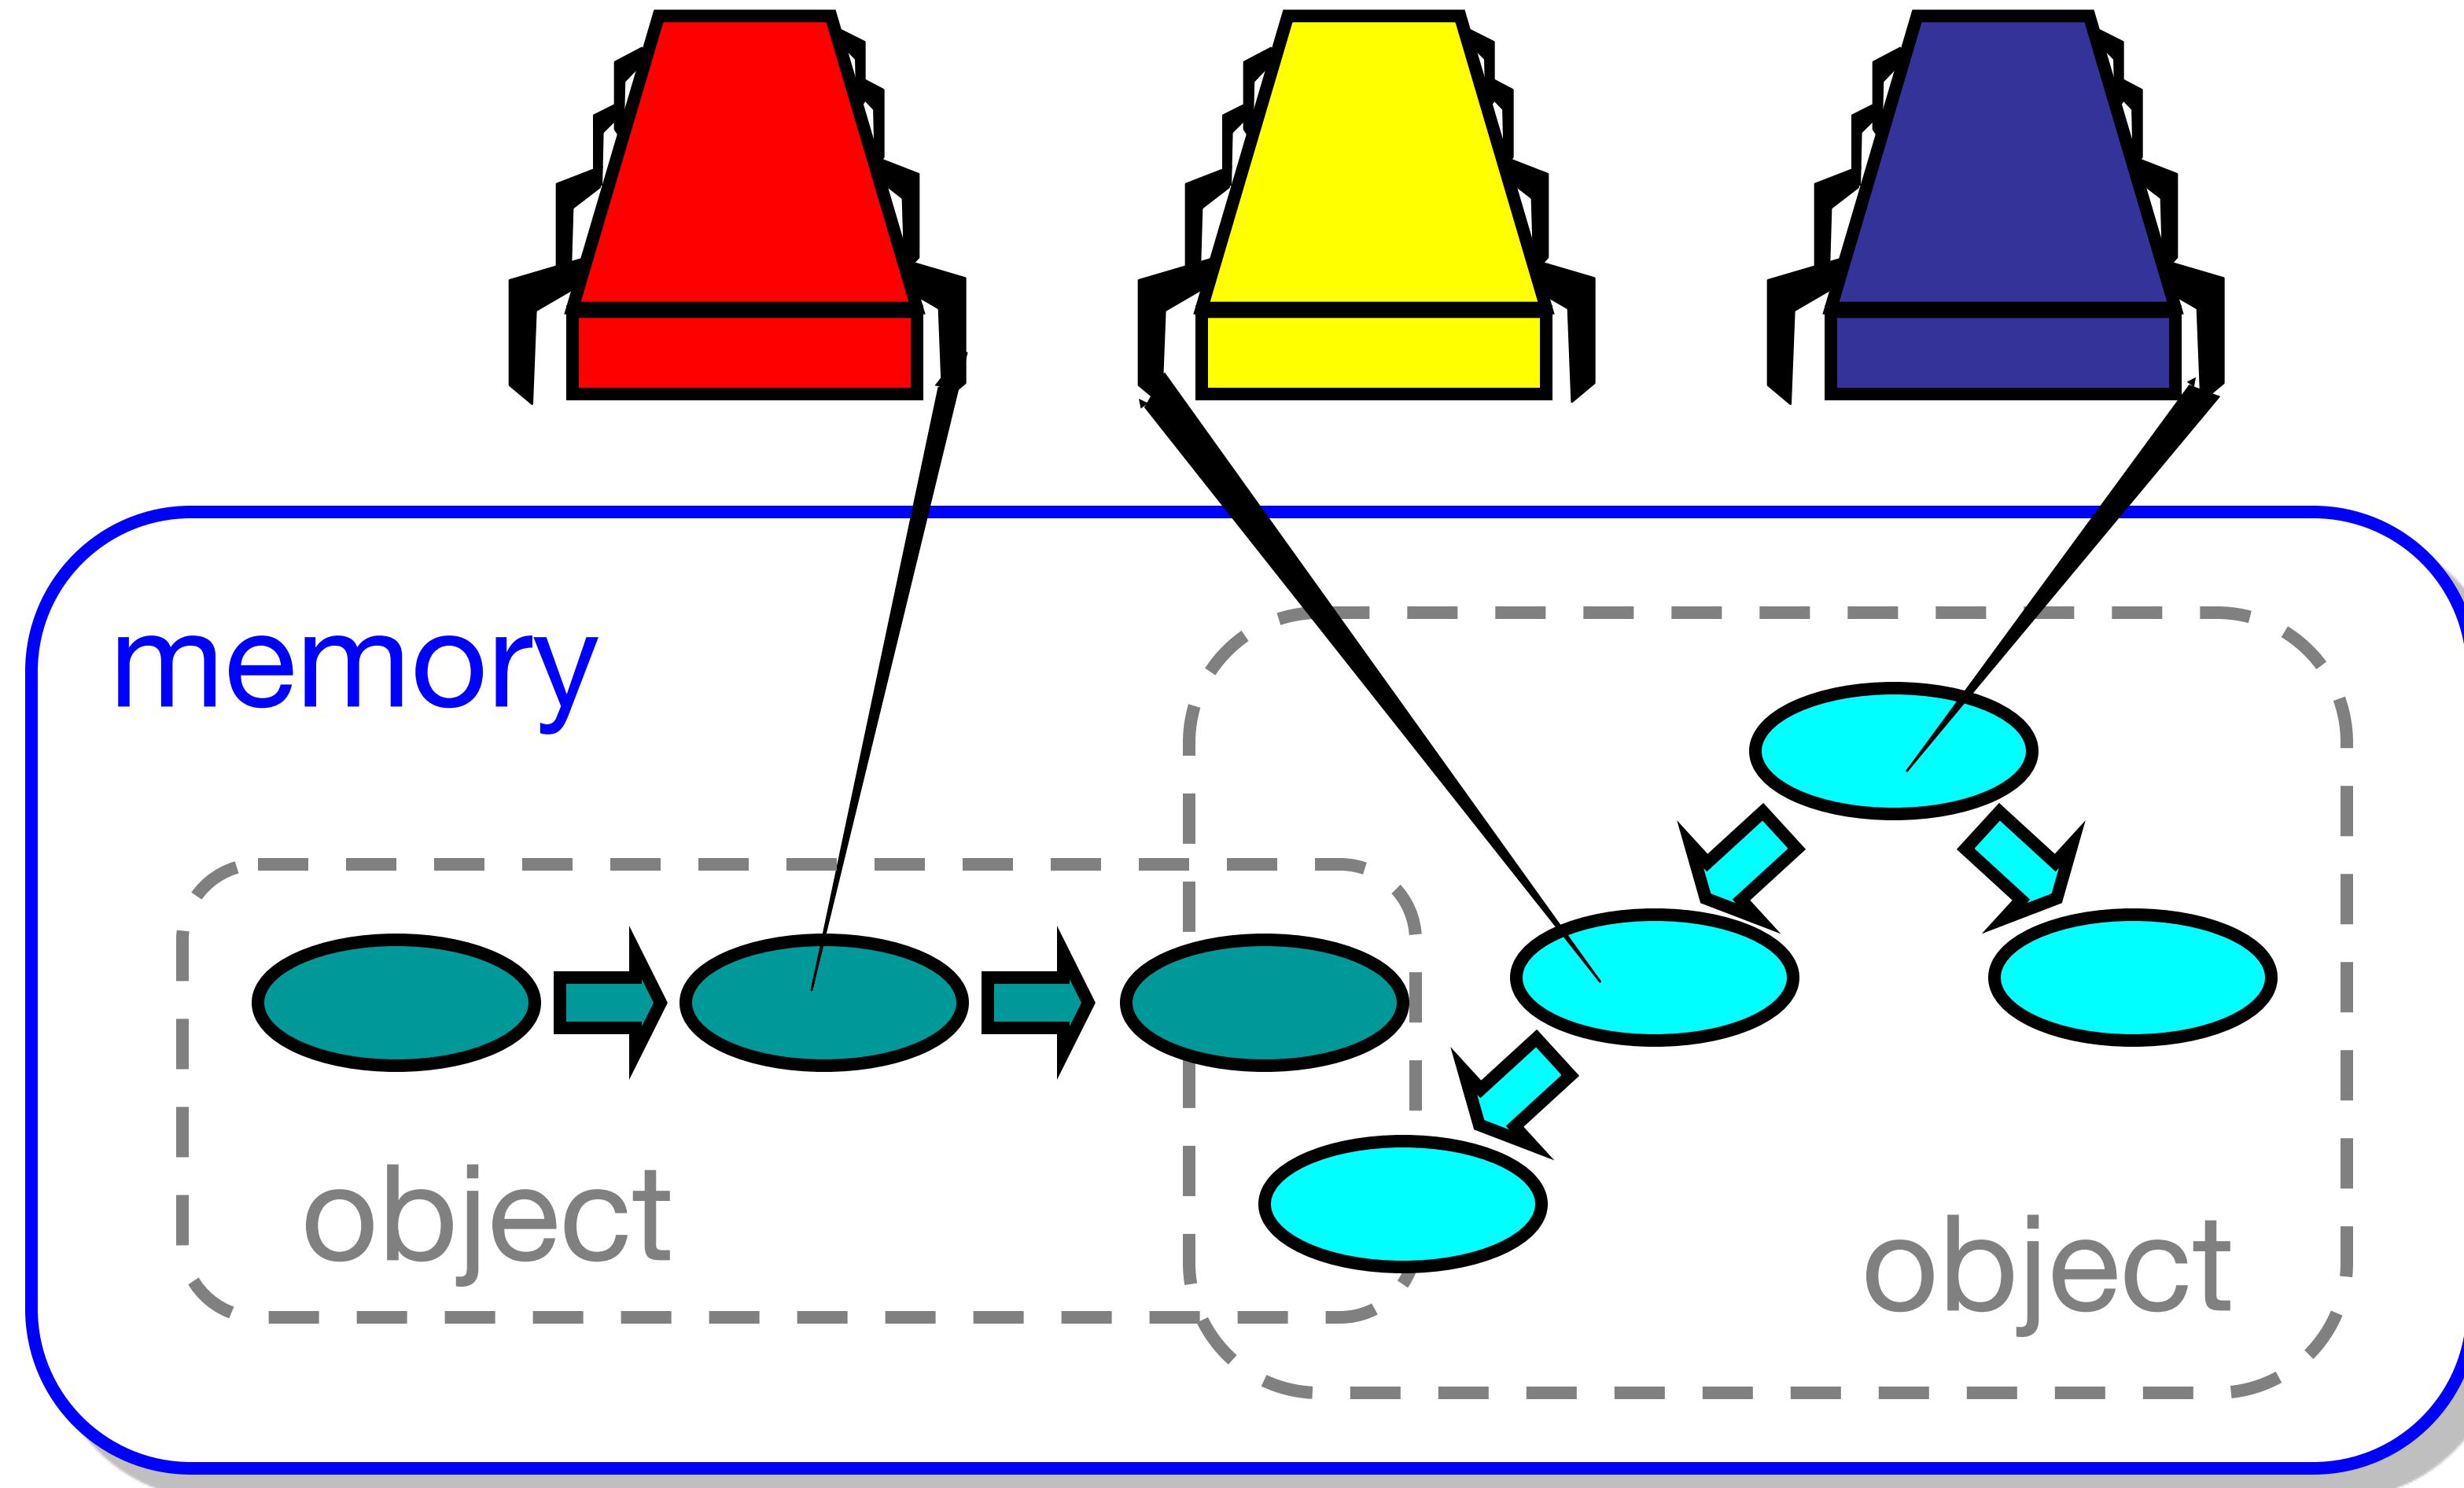
# **03 Concurrent Objects**

**CS 6868: Concurrent Programming**

**KC Sivaramakrishnan**

**Spring 2026, IIT Madras**

# Concurrent Computation



# Objectivism

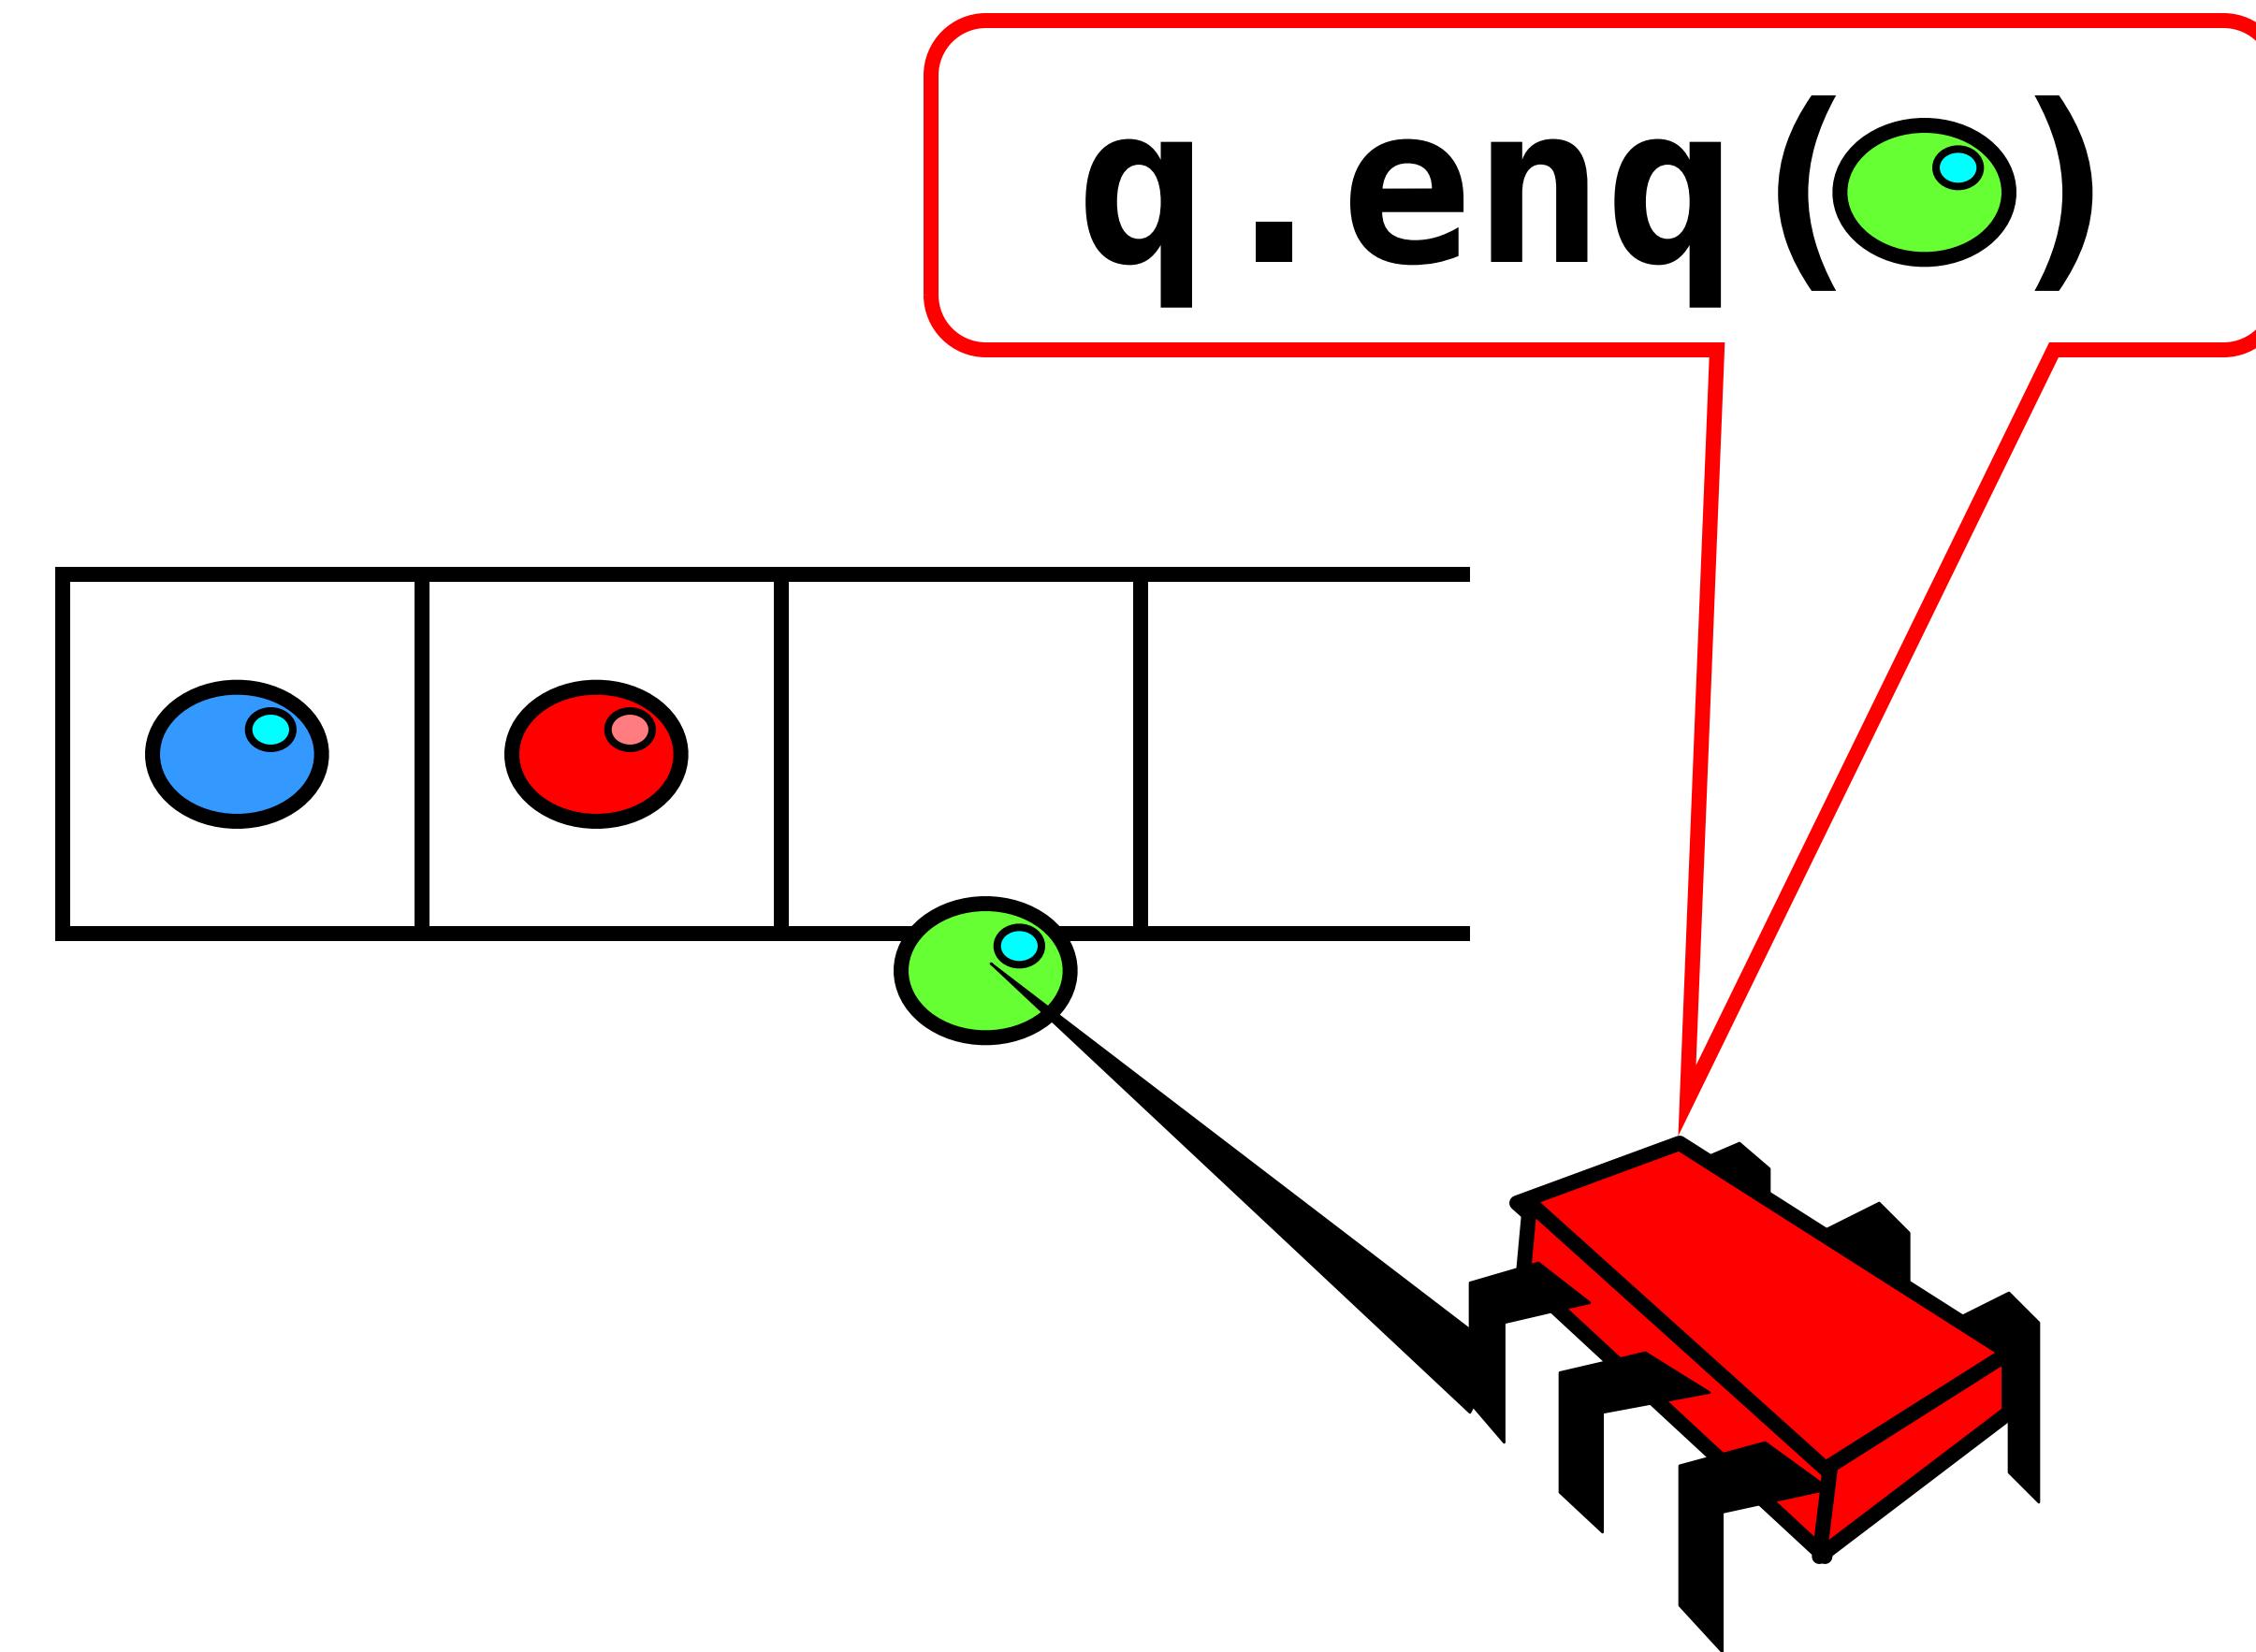
- What is a concurrent object?
  - How do we **describe** one?
  - How do we **implement** one?
  - How do we **tell if we're right**?

# Objectivism


- What is a concurrent object?
  - How do we **describe** one?
  - ~~How do we **implement** one?~~
  - How do we **tell if we're right?**

# Concurrent Queues

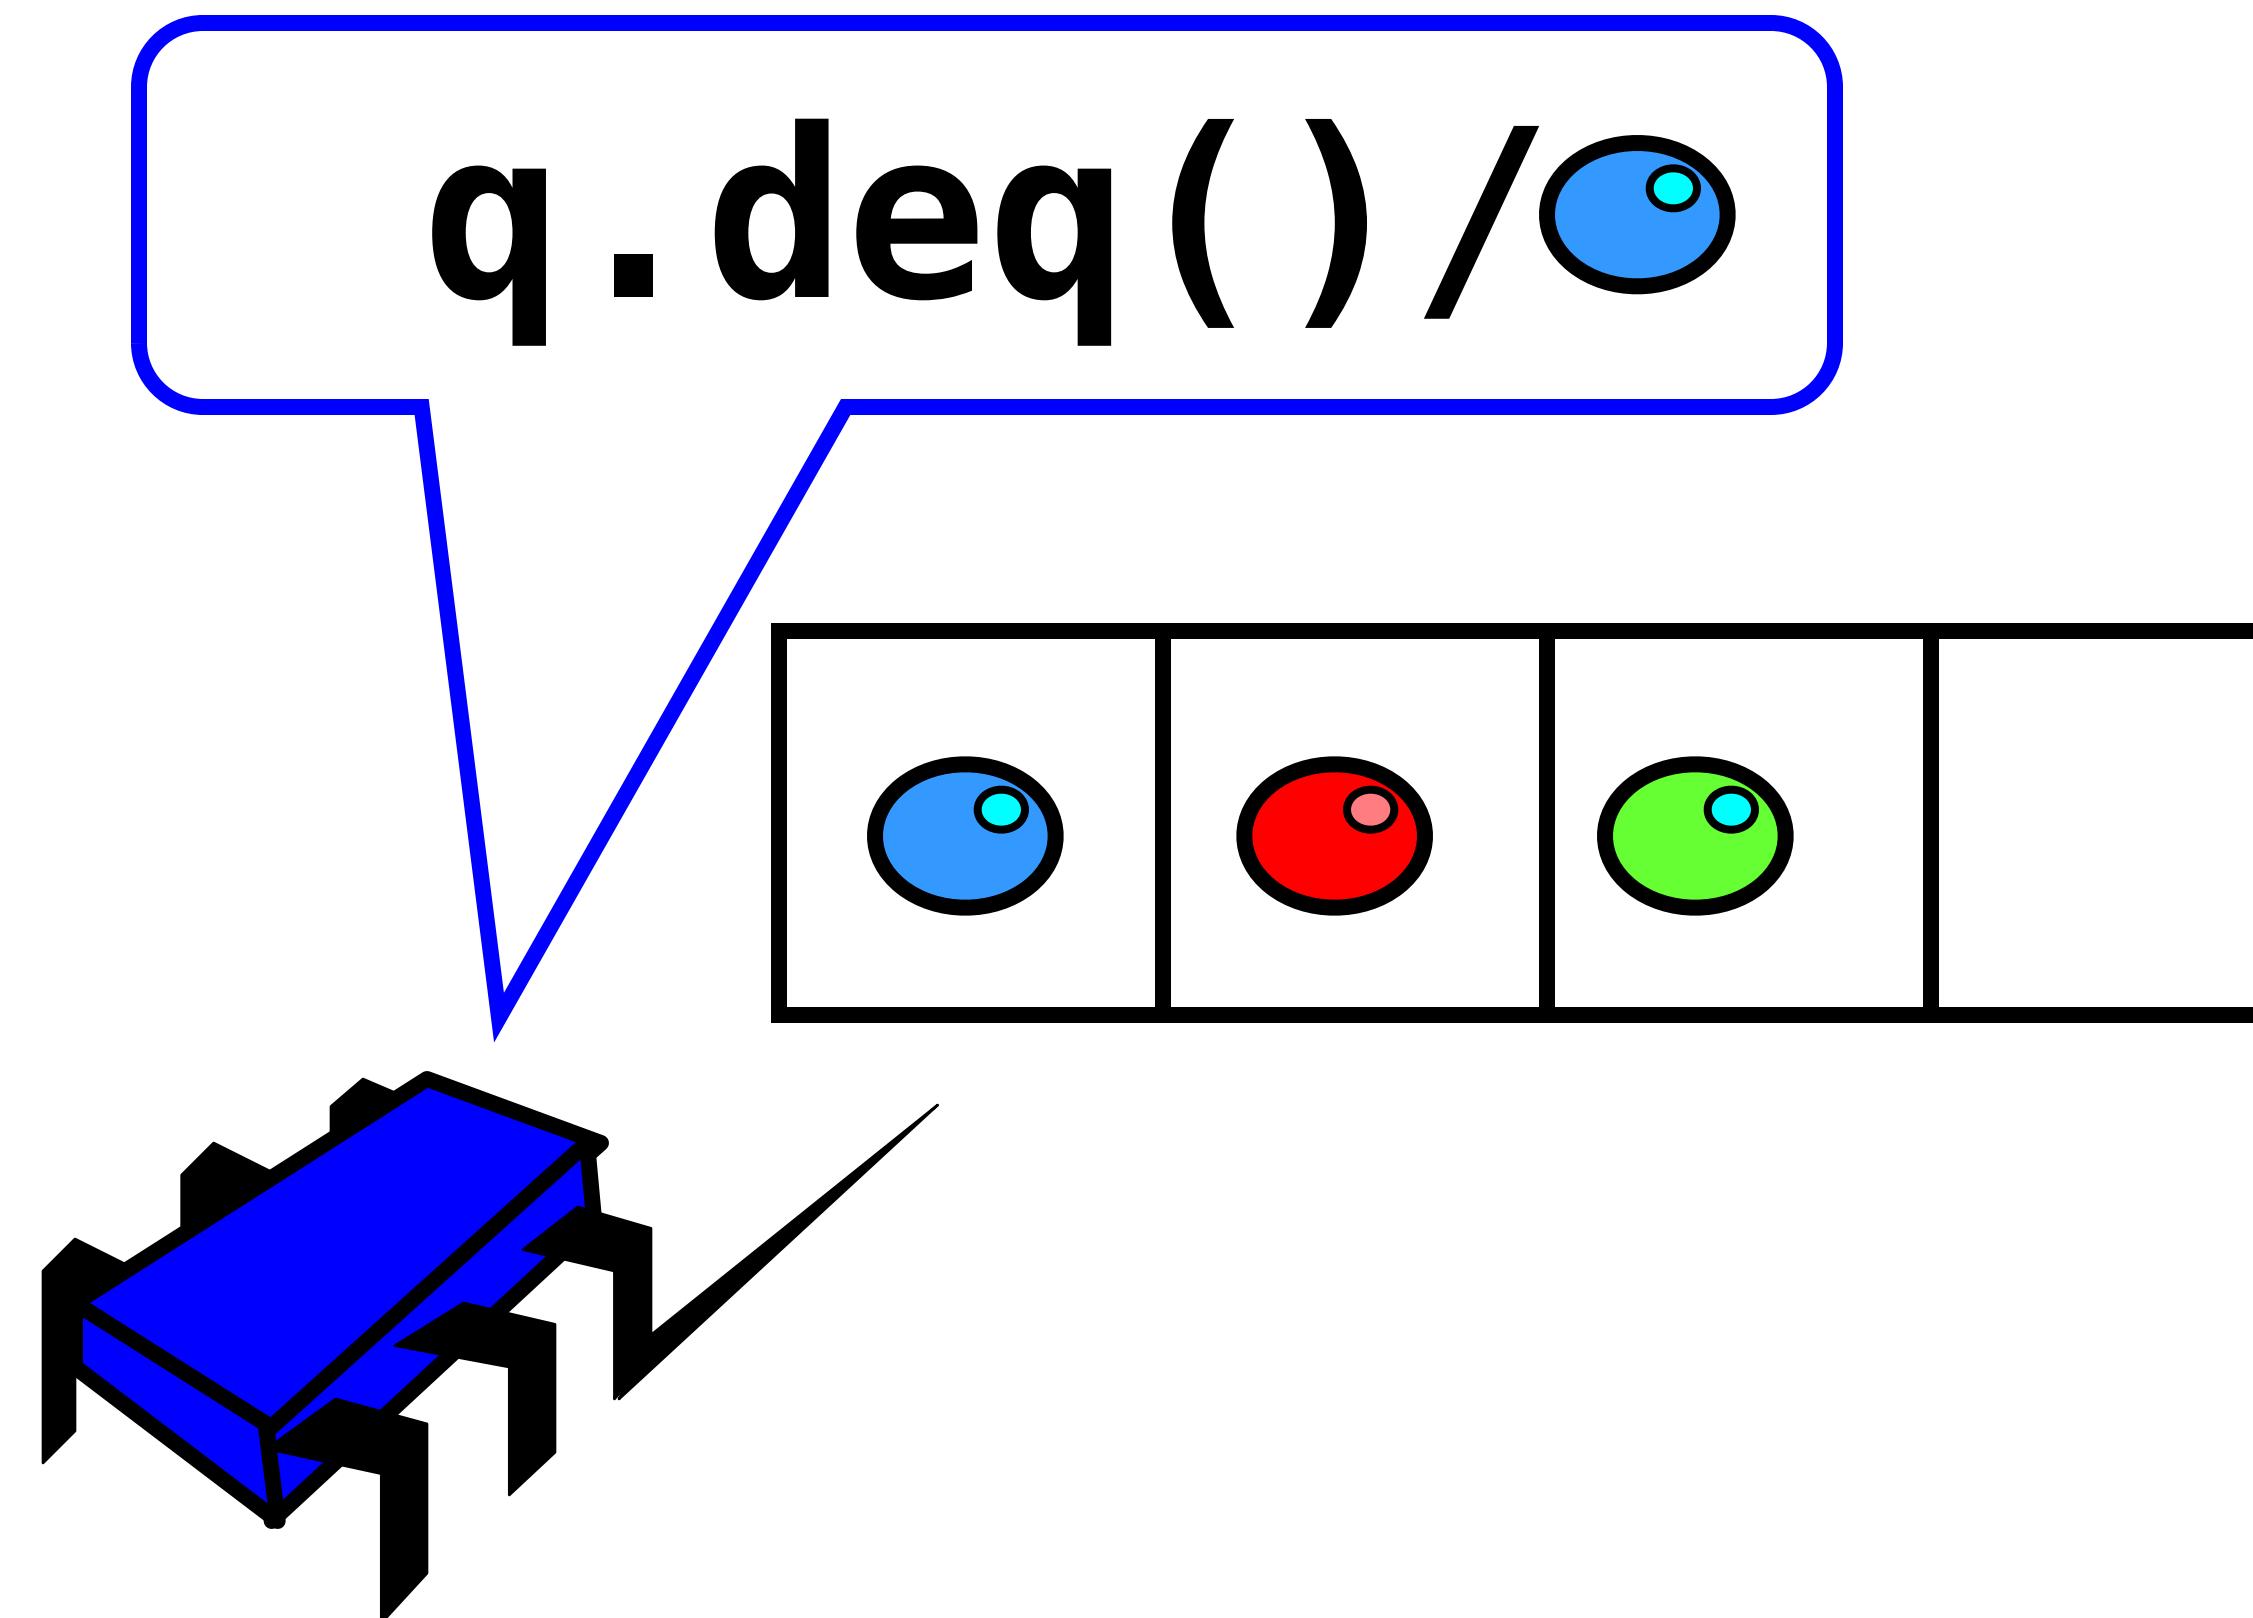
# FIFO Queue – Enqueue method




# FIFO Queue – Enqueue method




# FIFO Queue – Dequeue method



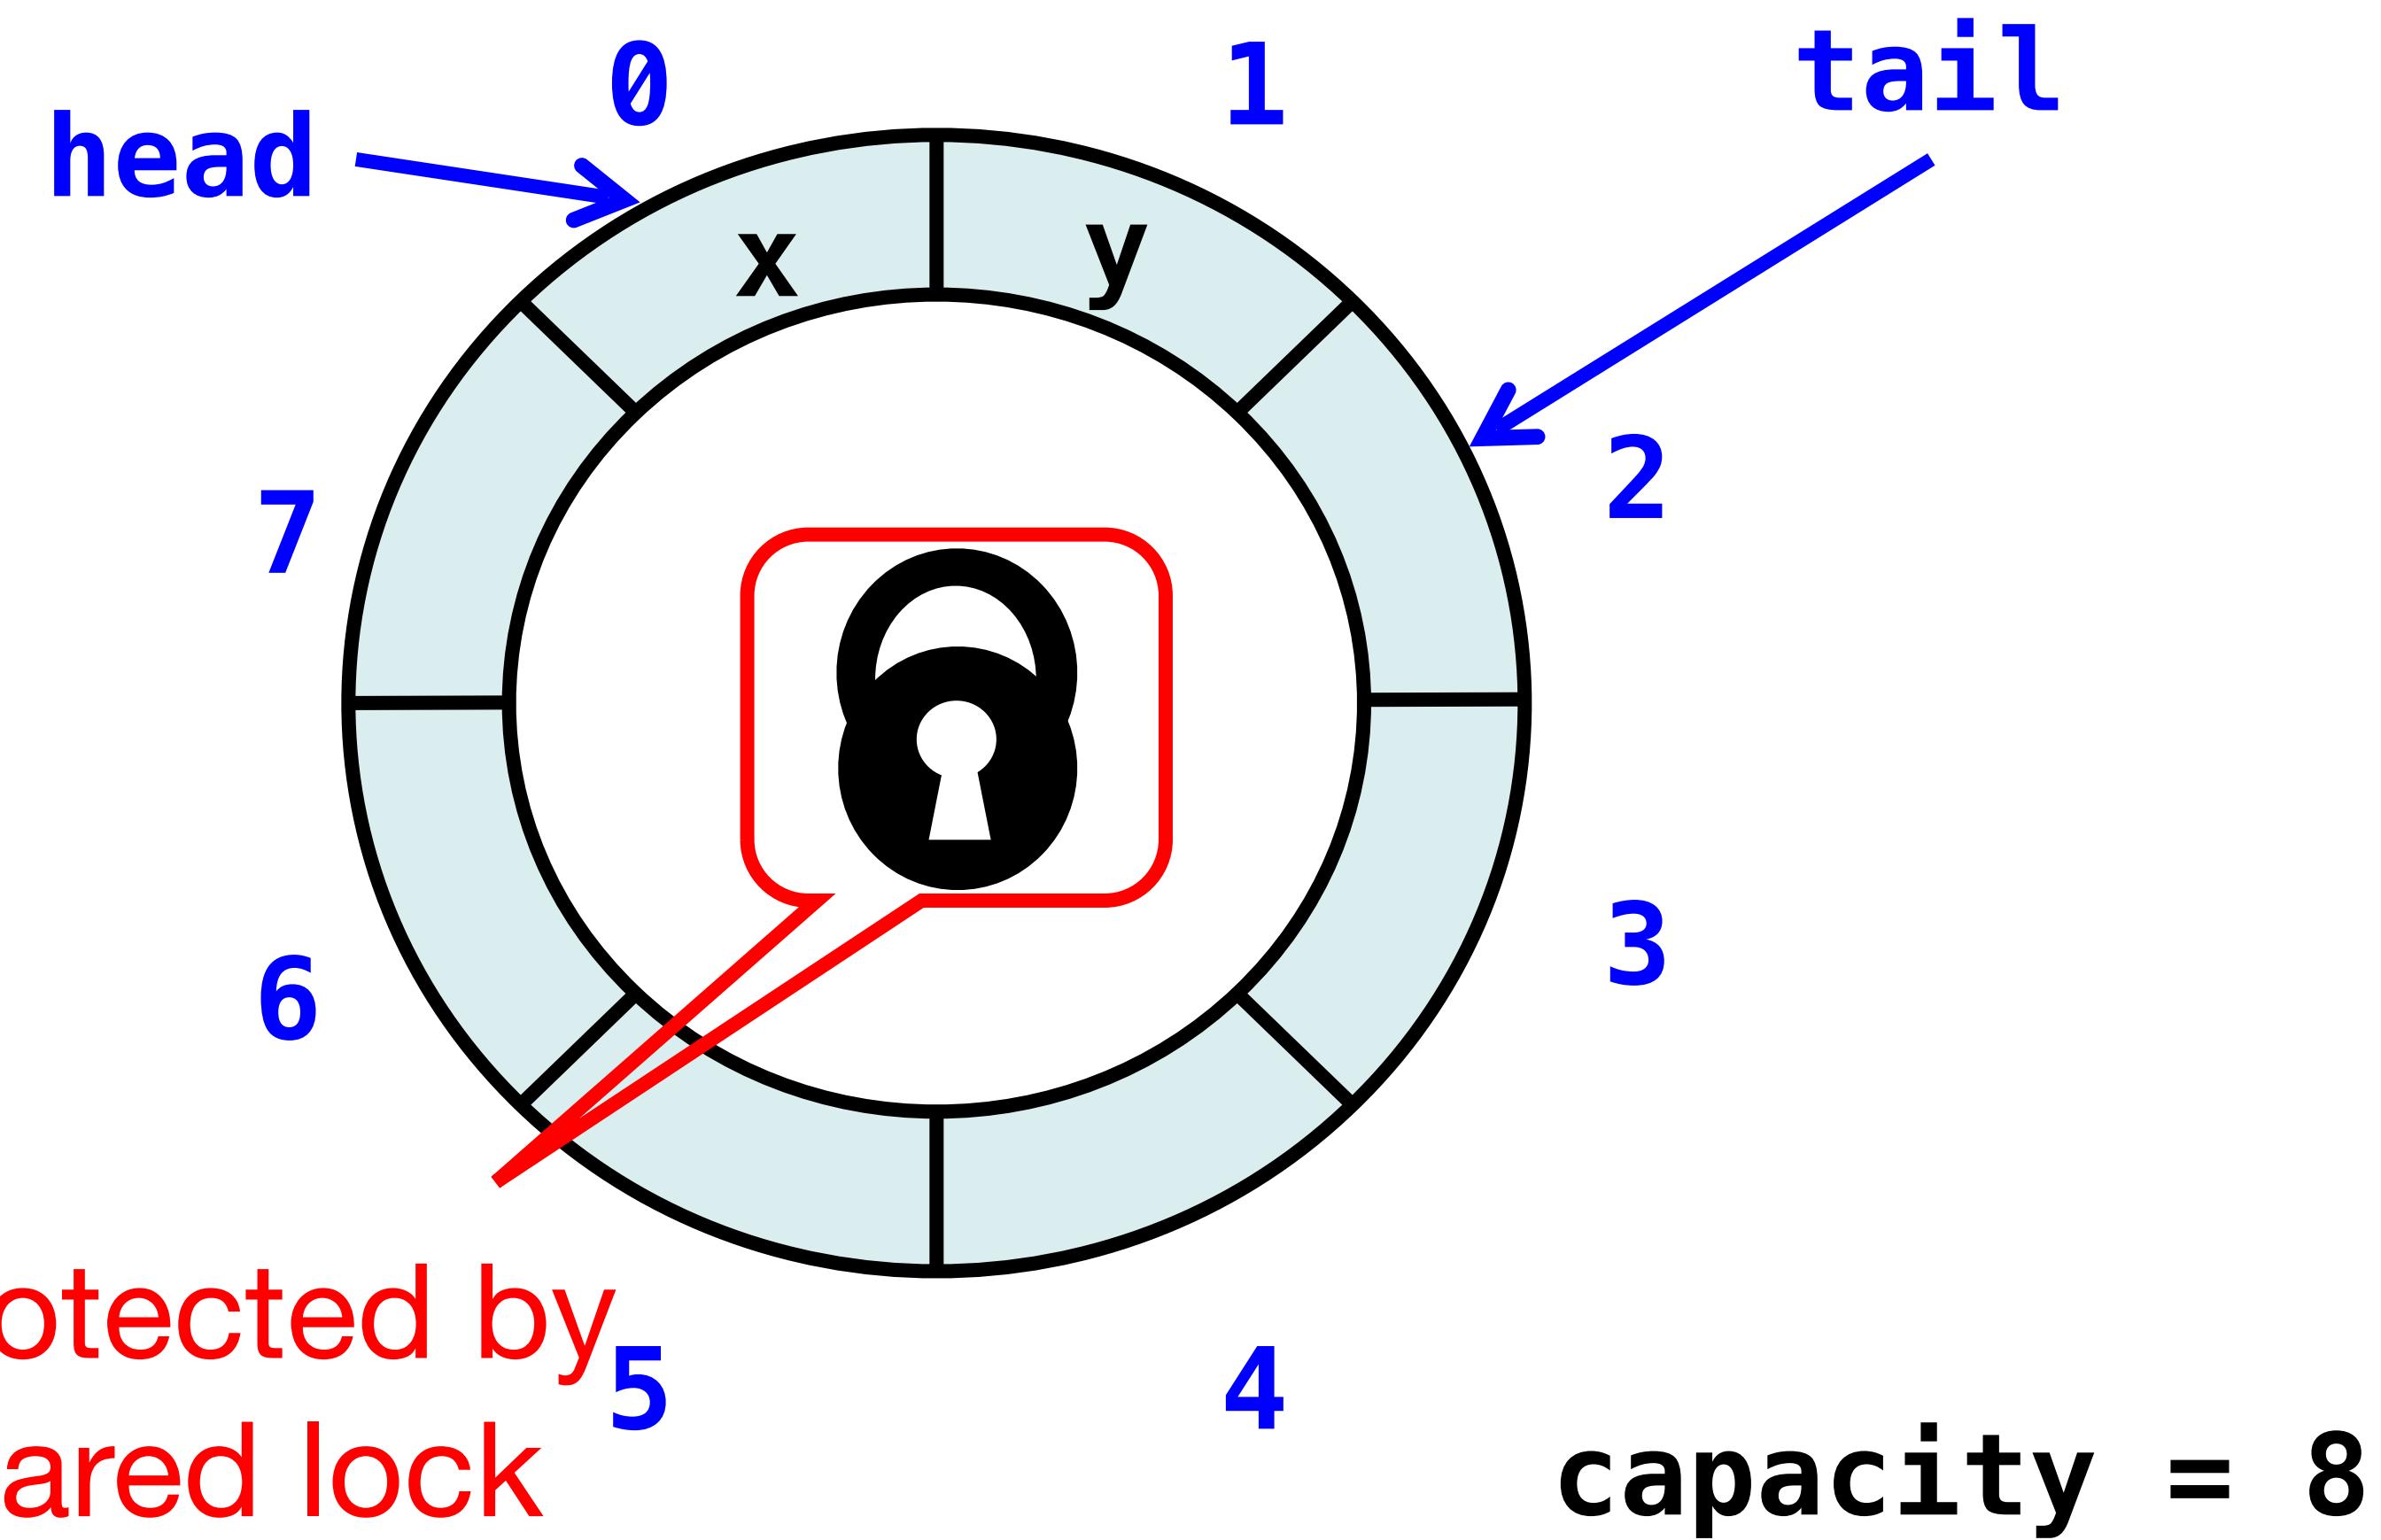
# FIFO Queue – Dequeue method



# Lock-based Queue



# Lock-based Queue



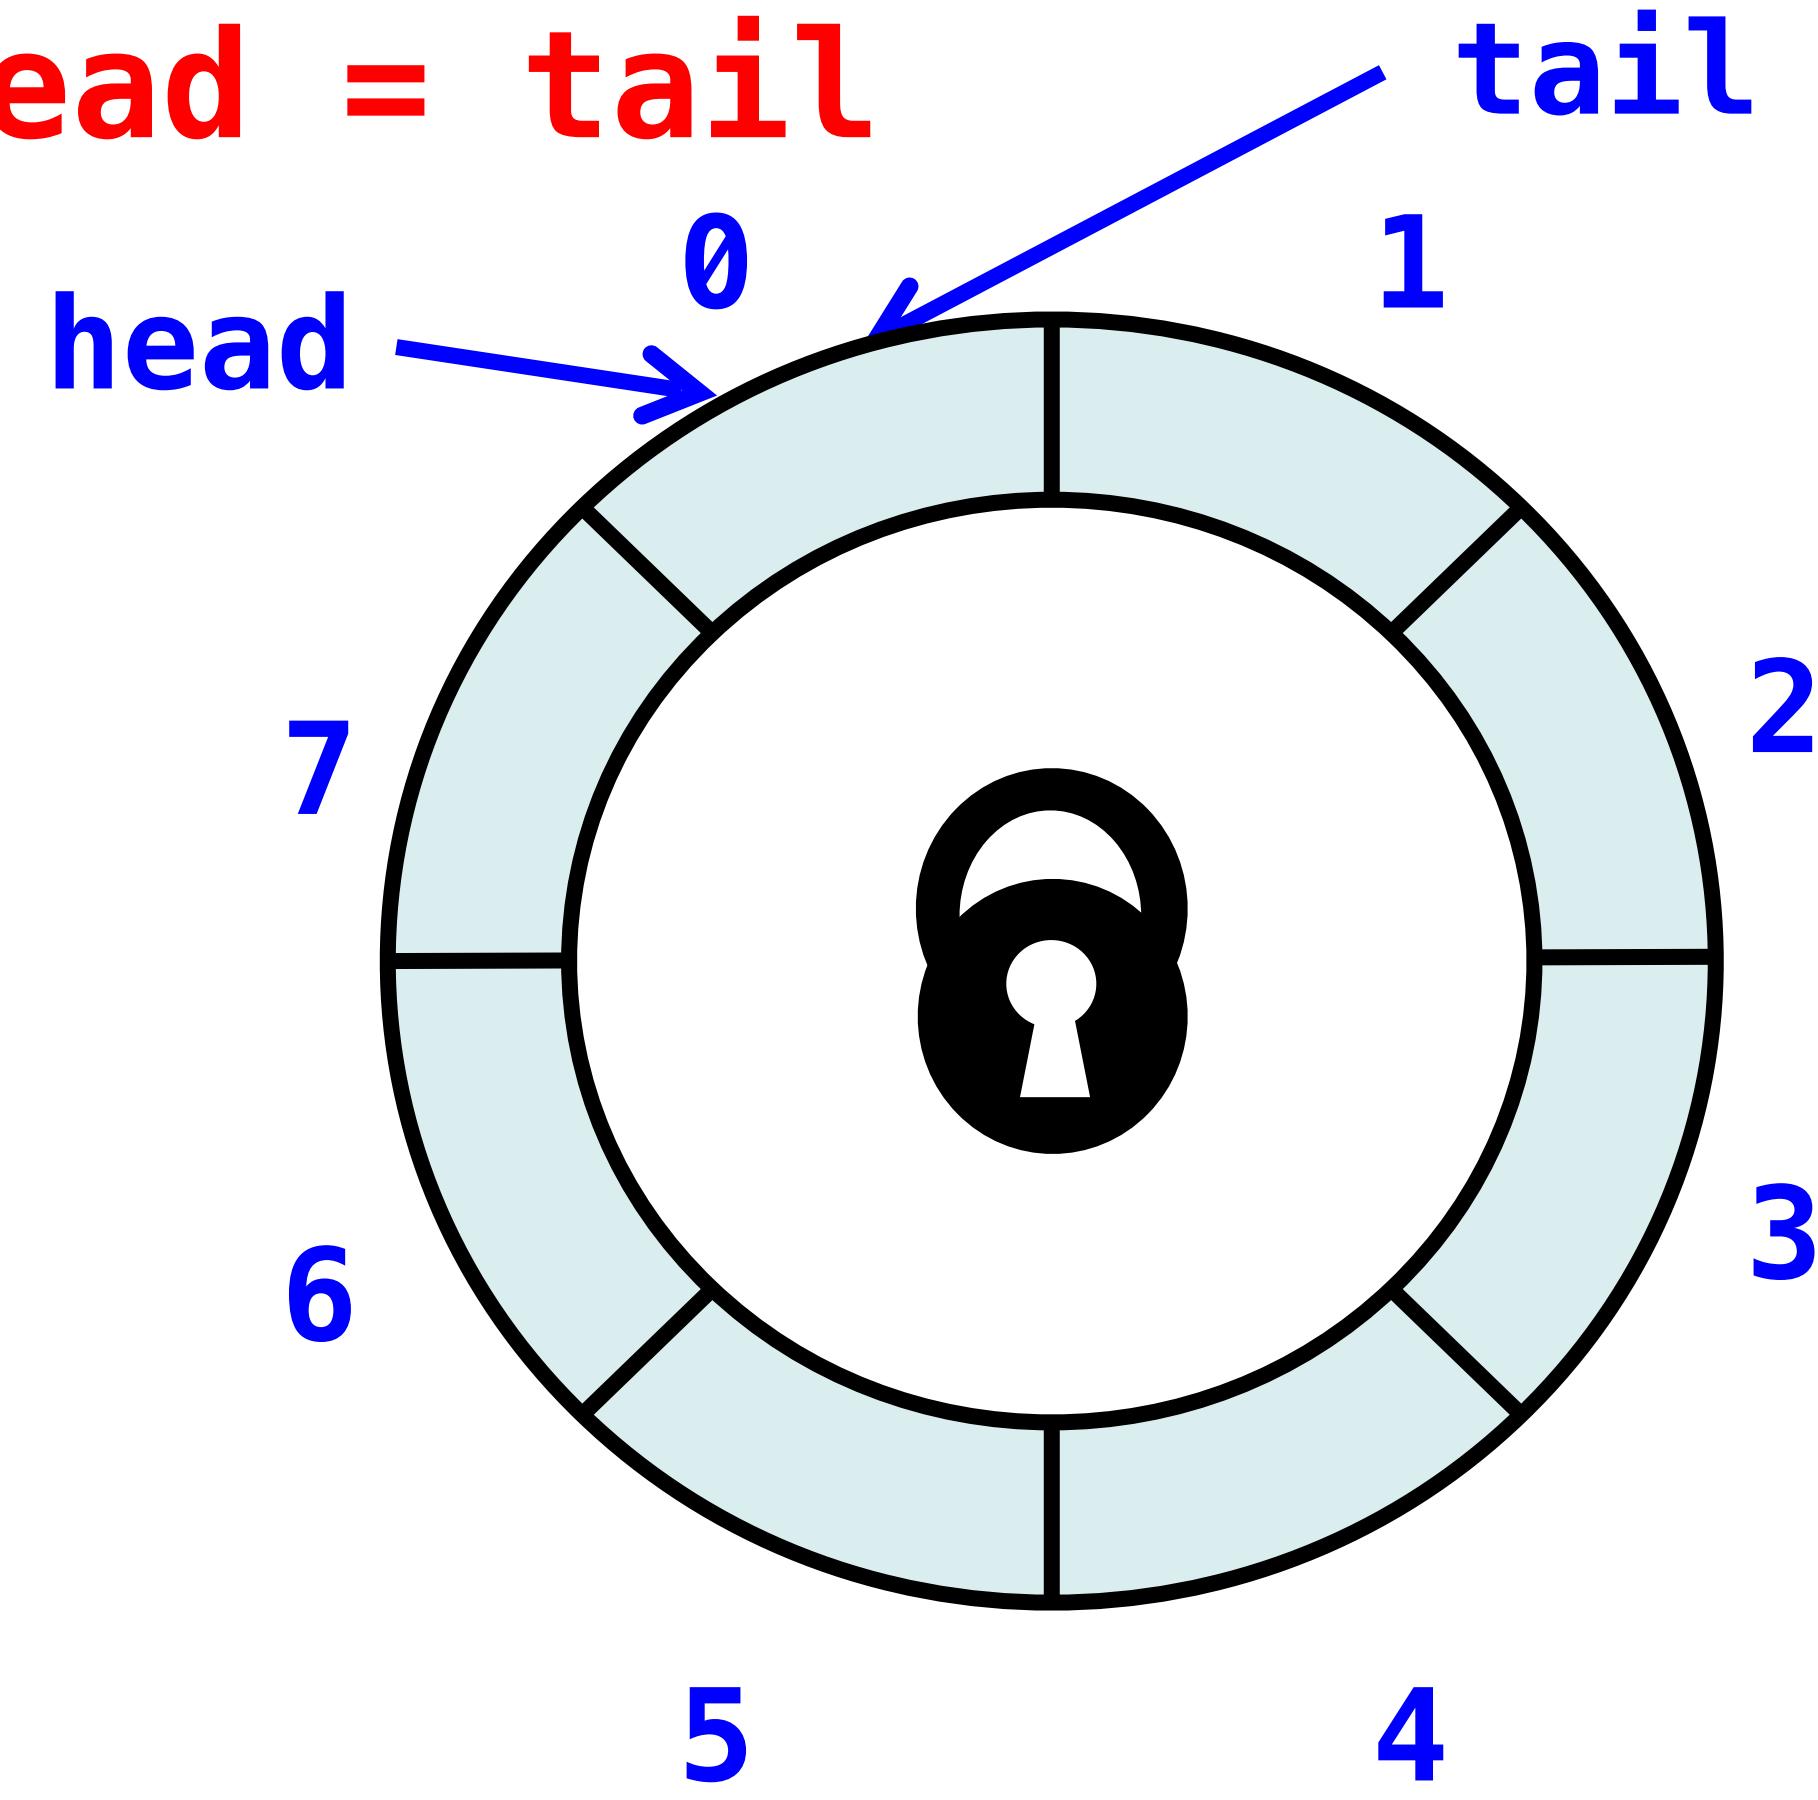
# Lock-based Queue

```
exception Full  
exception Empty
```


```
type 'a t = {  
    items : 'a option array;  
    capacity : int;  
    mutable head : int;  
    mutable tail : int;  
    lock : Mutex.t;  
}
```

```
let create capacity =  
{  
    items = Array.make capacity None;  
    capacity;  
    head = 0;  
    tail = 0;  
    lock = Mutex.create ();  
}
```

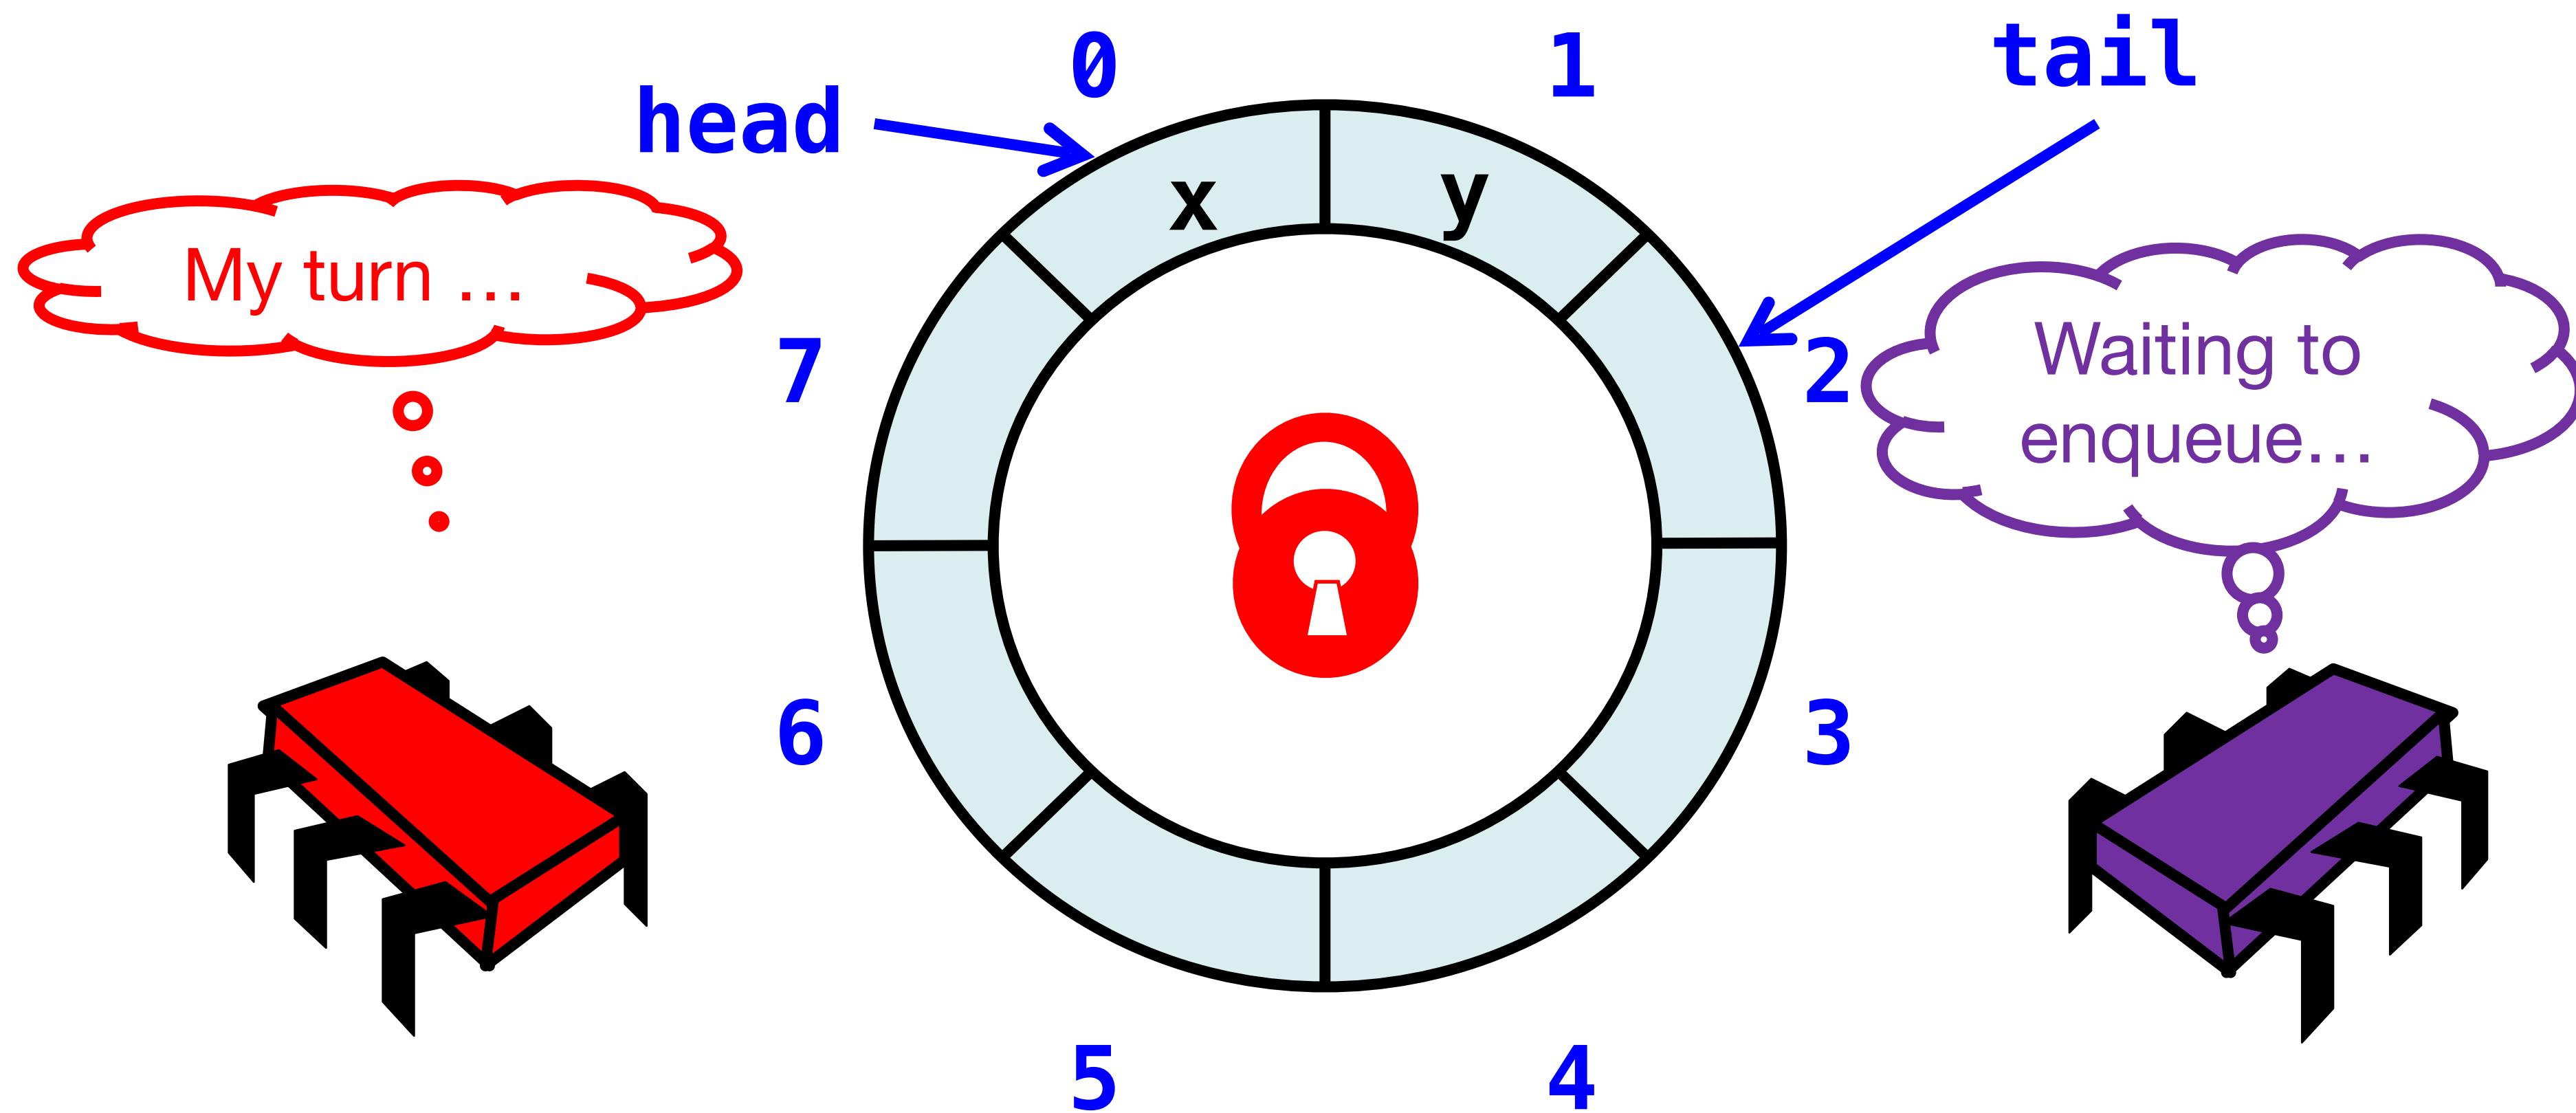
Initially: **head = tail**



# Lock-based Queue – deq( ) operation



# Acquire Lock




# Acquire Lock

```
let deq q =
  Mutex.lock q.lock;
  try
    if q.tail = q.head then
      raise Empty;
    match q.items.(q.head mod q.capacity) with
    | None -> assert false
    | Some x ->
        q.head <- q.head + 1;
        Mutex.unlock q.lock;
        x
  with e ->
    Mutex.unlock q.lock;
    raise e
```

Acquire lock at  
method start

# Check if non empty




# Check if non empty

```
let deq q =
  Mutex.lock q.lock;
  try
    if q.tail = q.head then
      raise Empty;
    match q.items.(q.head mod q.capacity) with
    | None -> assert false
    | Some x ->
        q.head <- q.head + 1;
        Mutex.unlock q.lock;
        x
  with e ->
    Mutex.unlock q.lock;
    raise e
```


If queue empty  
throw exception

In case of  
exceptions,  
lock released  
here

# Modify the queue



# Modify the queue




# Modify the queue


```
let deq q =
  Mutex.lock q.lock;
  try
    if q.tail = q.head then
      raise Empty;
    match q.items.(q.head mod q.capacity) with
    | None -> assert false
    | Some x ->
        q.head <- q.head + 1;
        Mutex.unlock q.lock;
        x
      with e ->
        Mutex.unlock q.lock;
        raise e
```

Queue not empty?  
Remove item “x” and update head

# Release the lock and return item



# Release the lock and return item



# Release the lock and return item

```
let deq q =
  Mutex.lock q.lock;
  try
    if q.tail = q.head then
      raise Empty;
    match q.items.(q.head mod q.capacity) with
    | None -> assert false
    | Some x ->
        q.head <- q.head + 1;
        Mutex.unlock q.lock;
        x
  with e ->
    Mutex.unlock q.lock;
    raise e
```

Unlock and return item “x”

# Implementation – deq( )

```
let deq q =
  Mutex.lock q.lock;
  try
    if q.tail = q.head then
      raise Empty;
    match q.items.(q.head mod q.capacity) with
    | None -> assert false
    | Some x ->
        q.head <- q.head + 1;
        Mutex.unlock q.lock;
        x
  with e ->
    Mutex.unlock q.lock;
    raise e
```

# Implementation – deq()

```
let deq q =
  Mutex.lock q.lock;
  try
    if q.tail = q.head then
      raise Empty;
    match q.items.(q.head mod q.capacity) with
    | None -> assert false
    | Some x ->
        q.head <- q.head + 1;
        Mutex.unlock q.lock;
        x
  with e ->
    Mutex.unlock q.lock;
    raise e
```

Should be correct because  
modifications are mutually exclusive...

# Implementation – deq()

```
let deq q =
  Mutex.lock q.lock;
  try
    if q.tail = q.head then
      raise Empty;
    match q.items.(q.head mod q.capacity) with
    | None -> assert false
    | Some x ->
        q.head <- q.head + 1;
        Mutex.unlock q.lock;
        x
  with e ->
    Mutex.unlock q.lock;
    raise e
```


Should be correct because  
modifications are mutually exclusive...

Demo

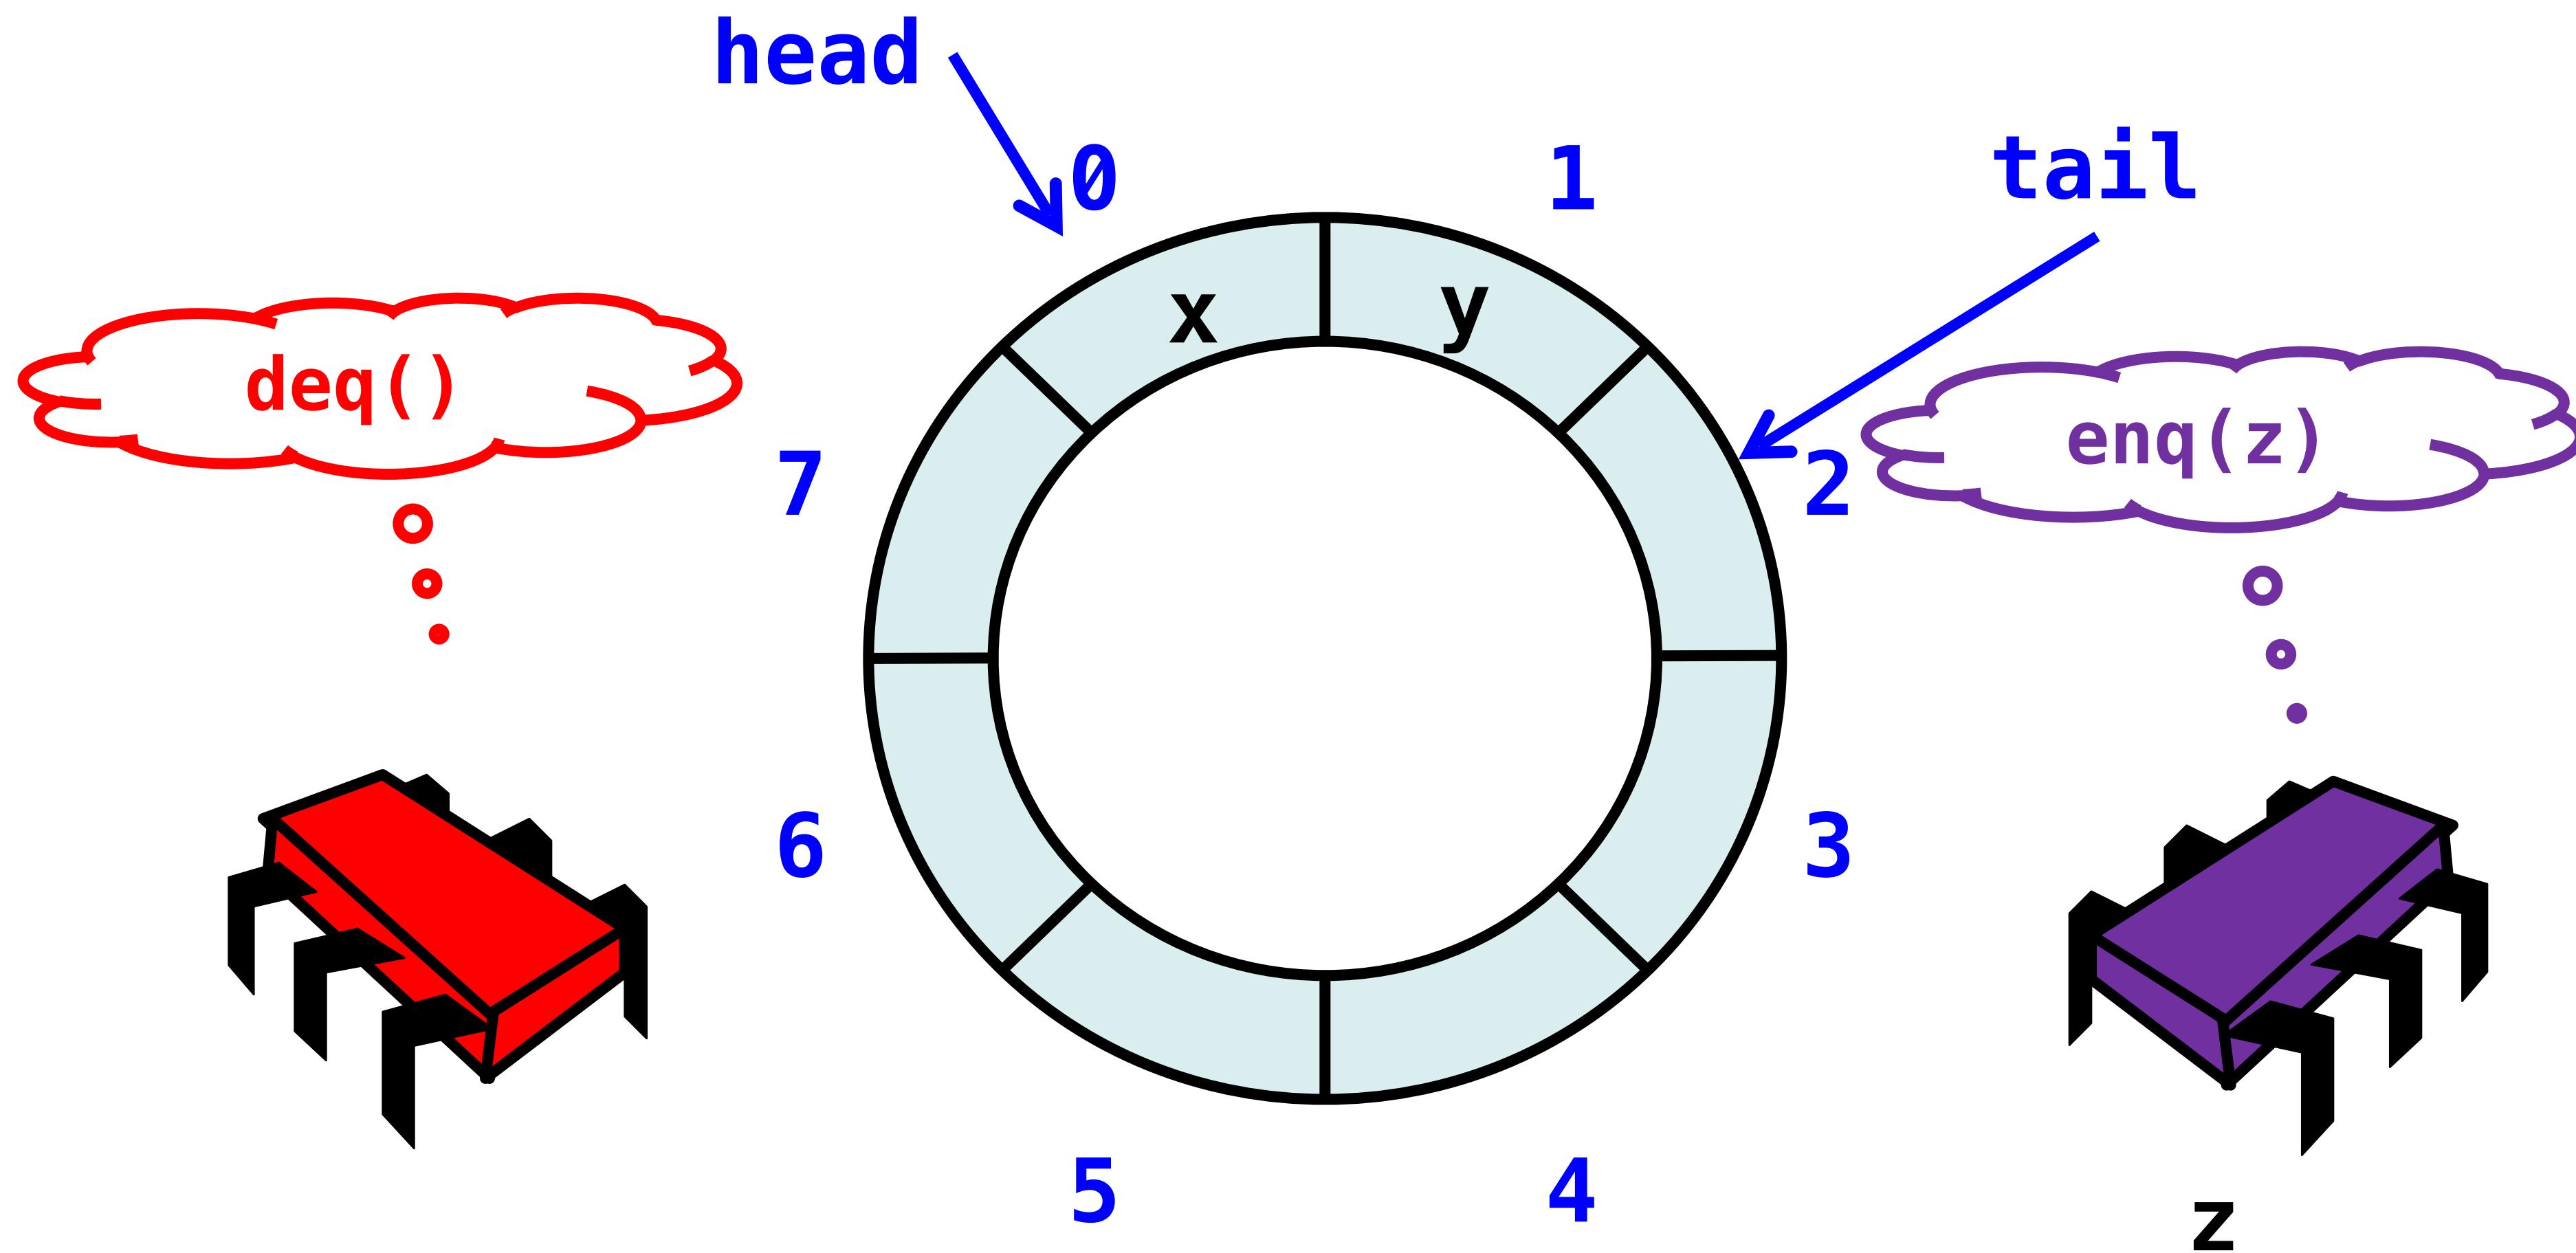
# Consider the following implementation

- The same thing without mutual exclusion
- For simplicity, only **two** threads
  - One thread **enq only**
  - The other **deq only**


# Wait-free 2-thread queue



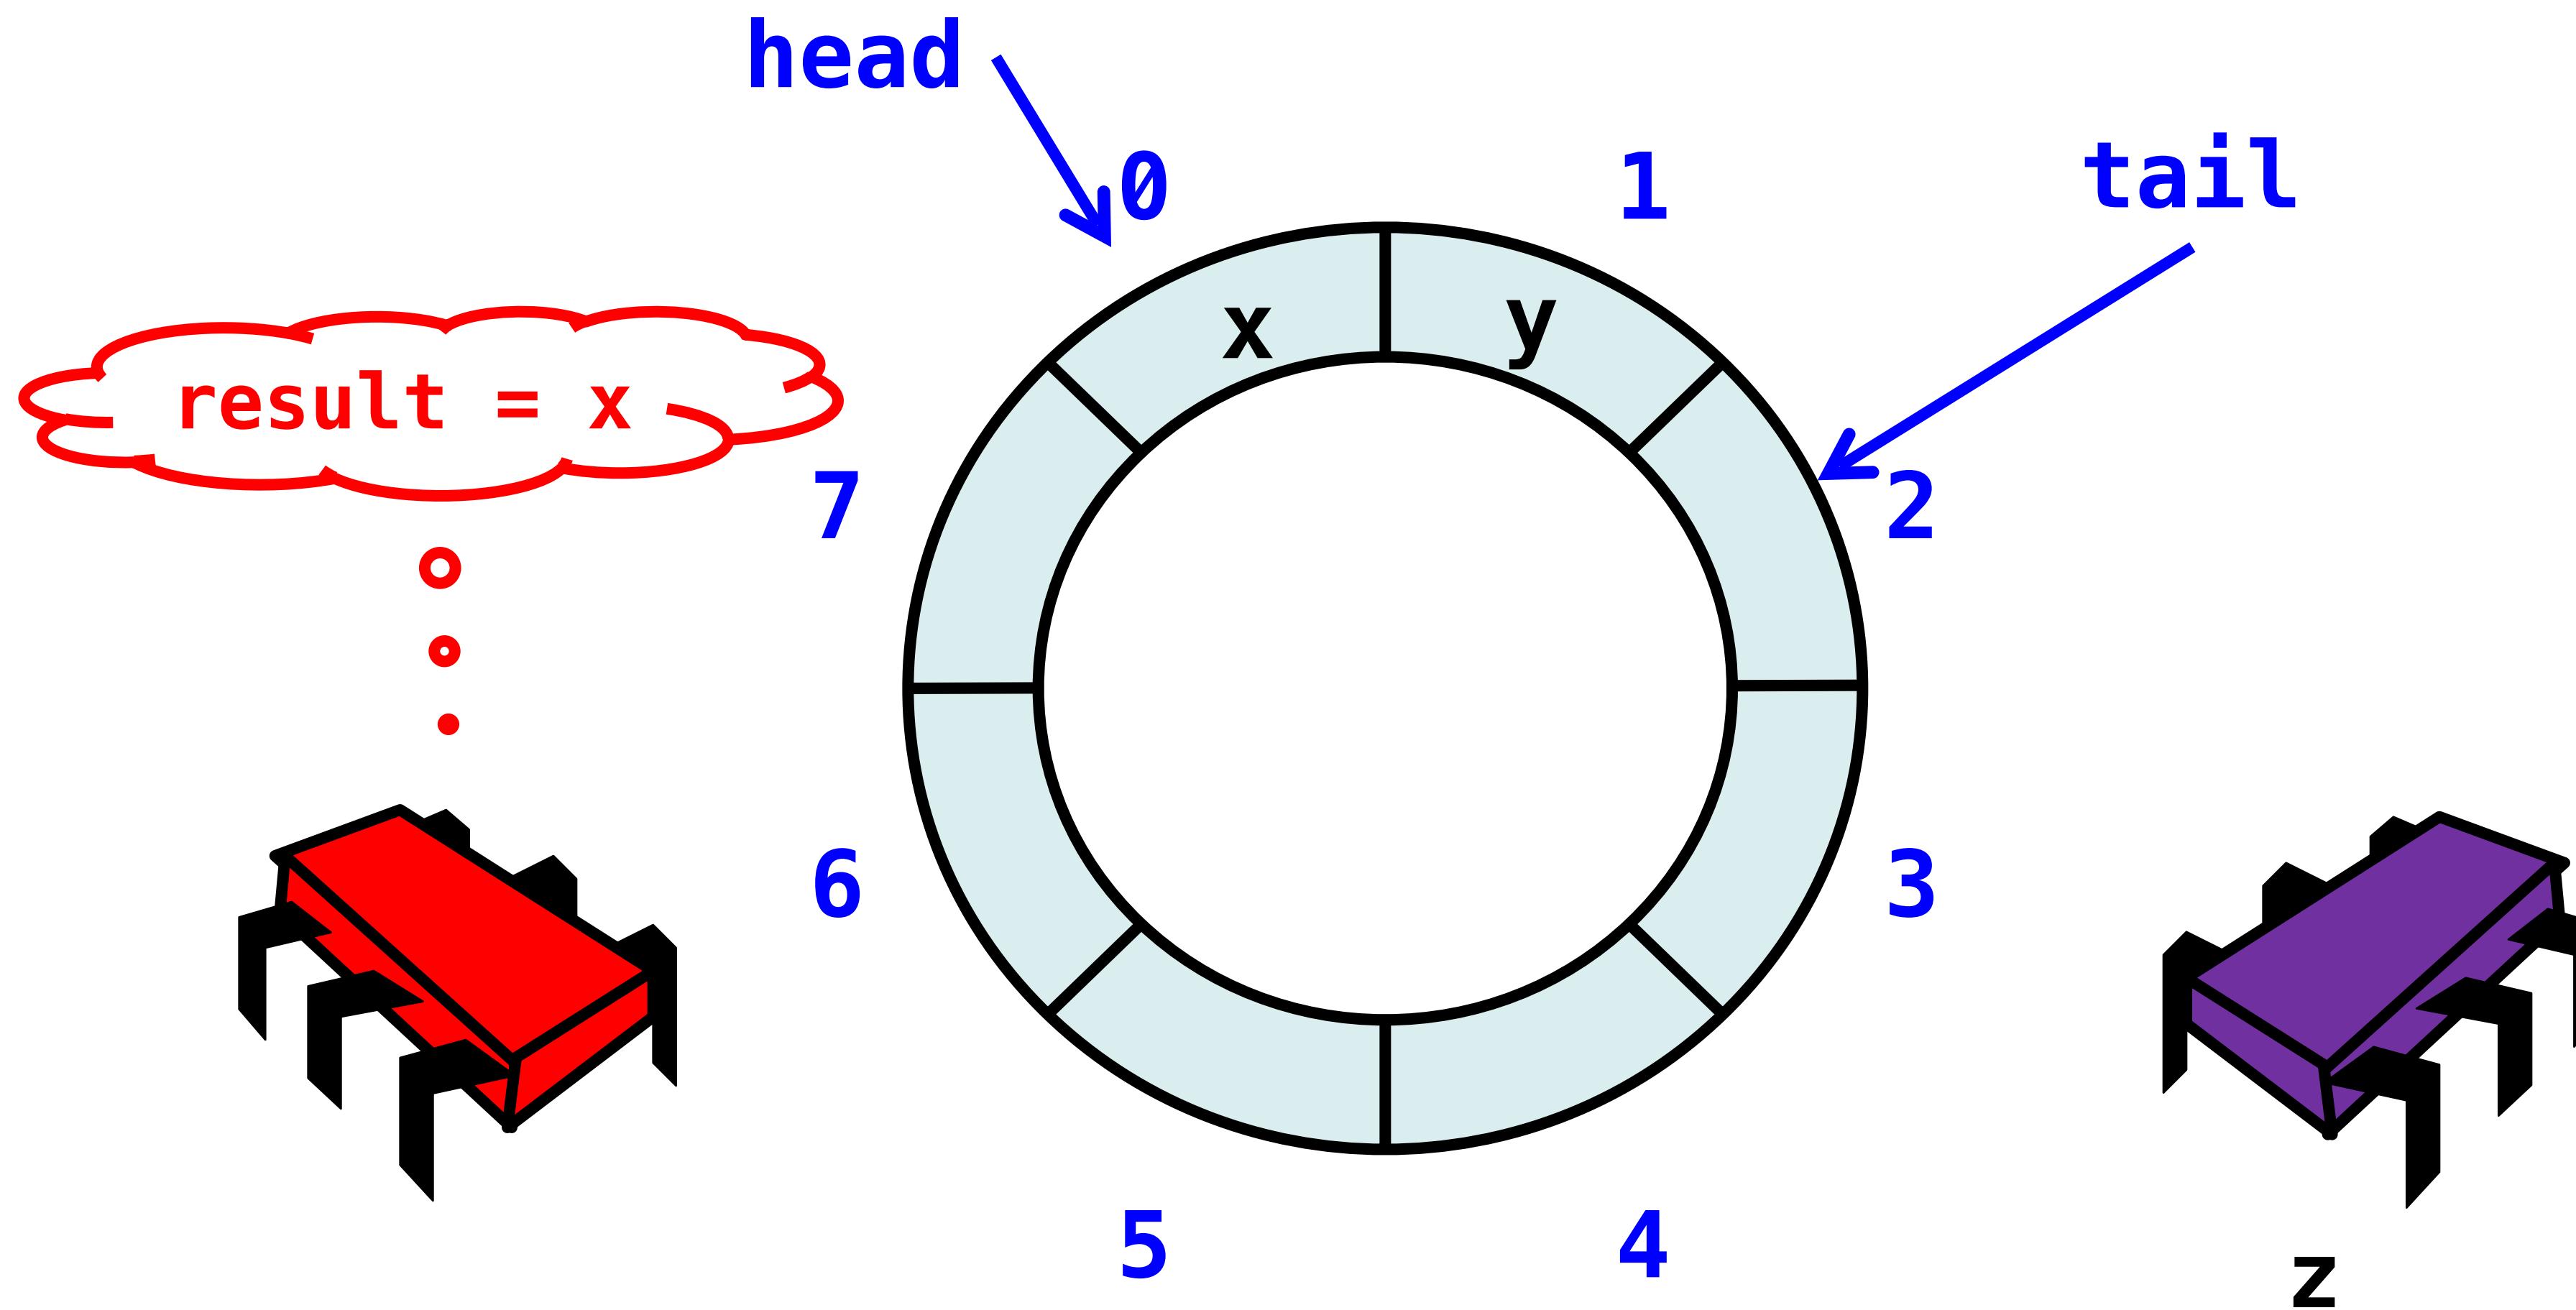
# Wait-free 2-thread queue




# Wait-free 2-thread queue




# Wait-free 2-thread queue



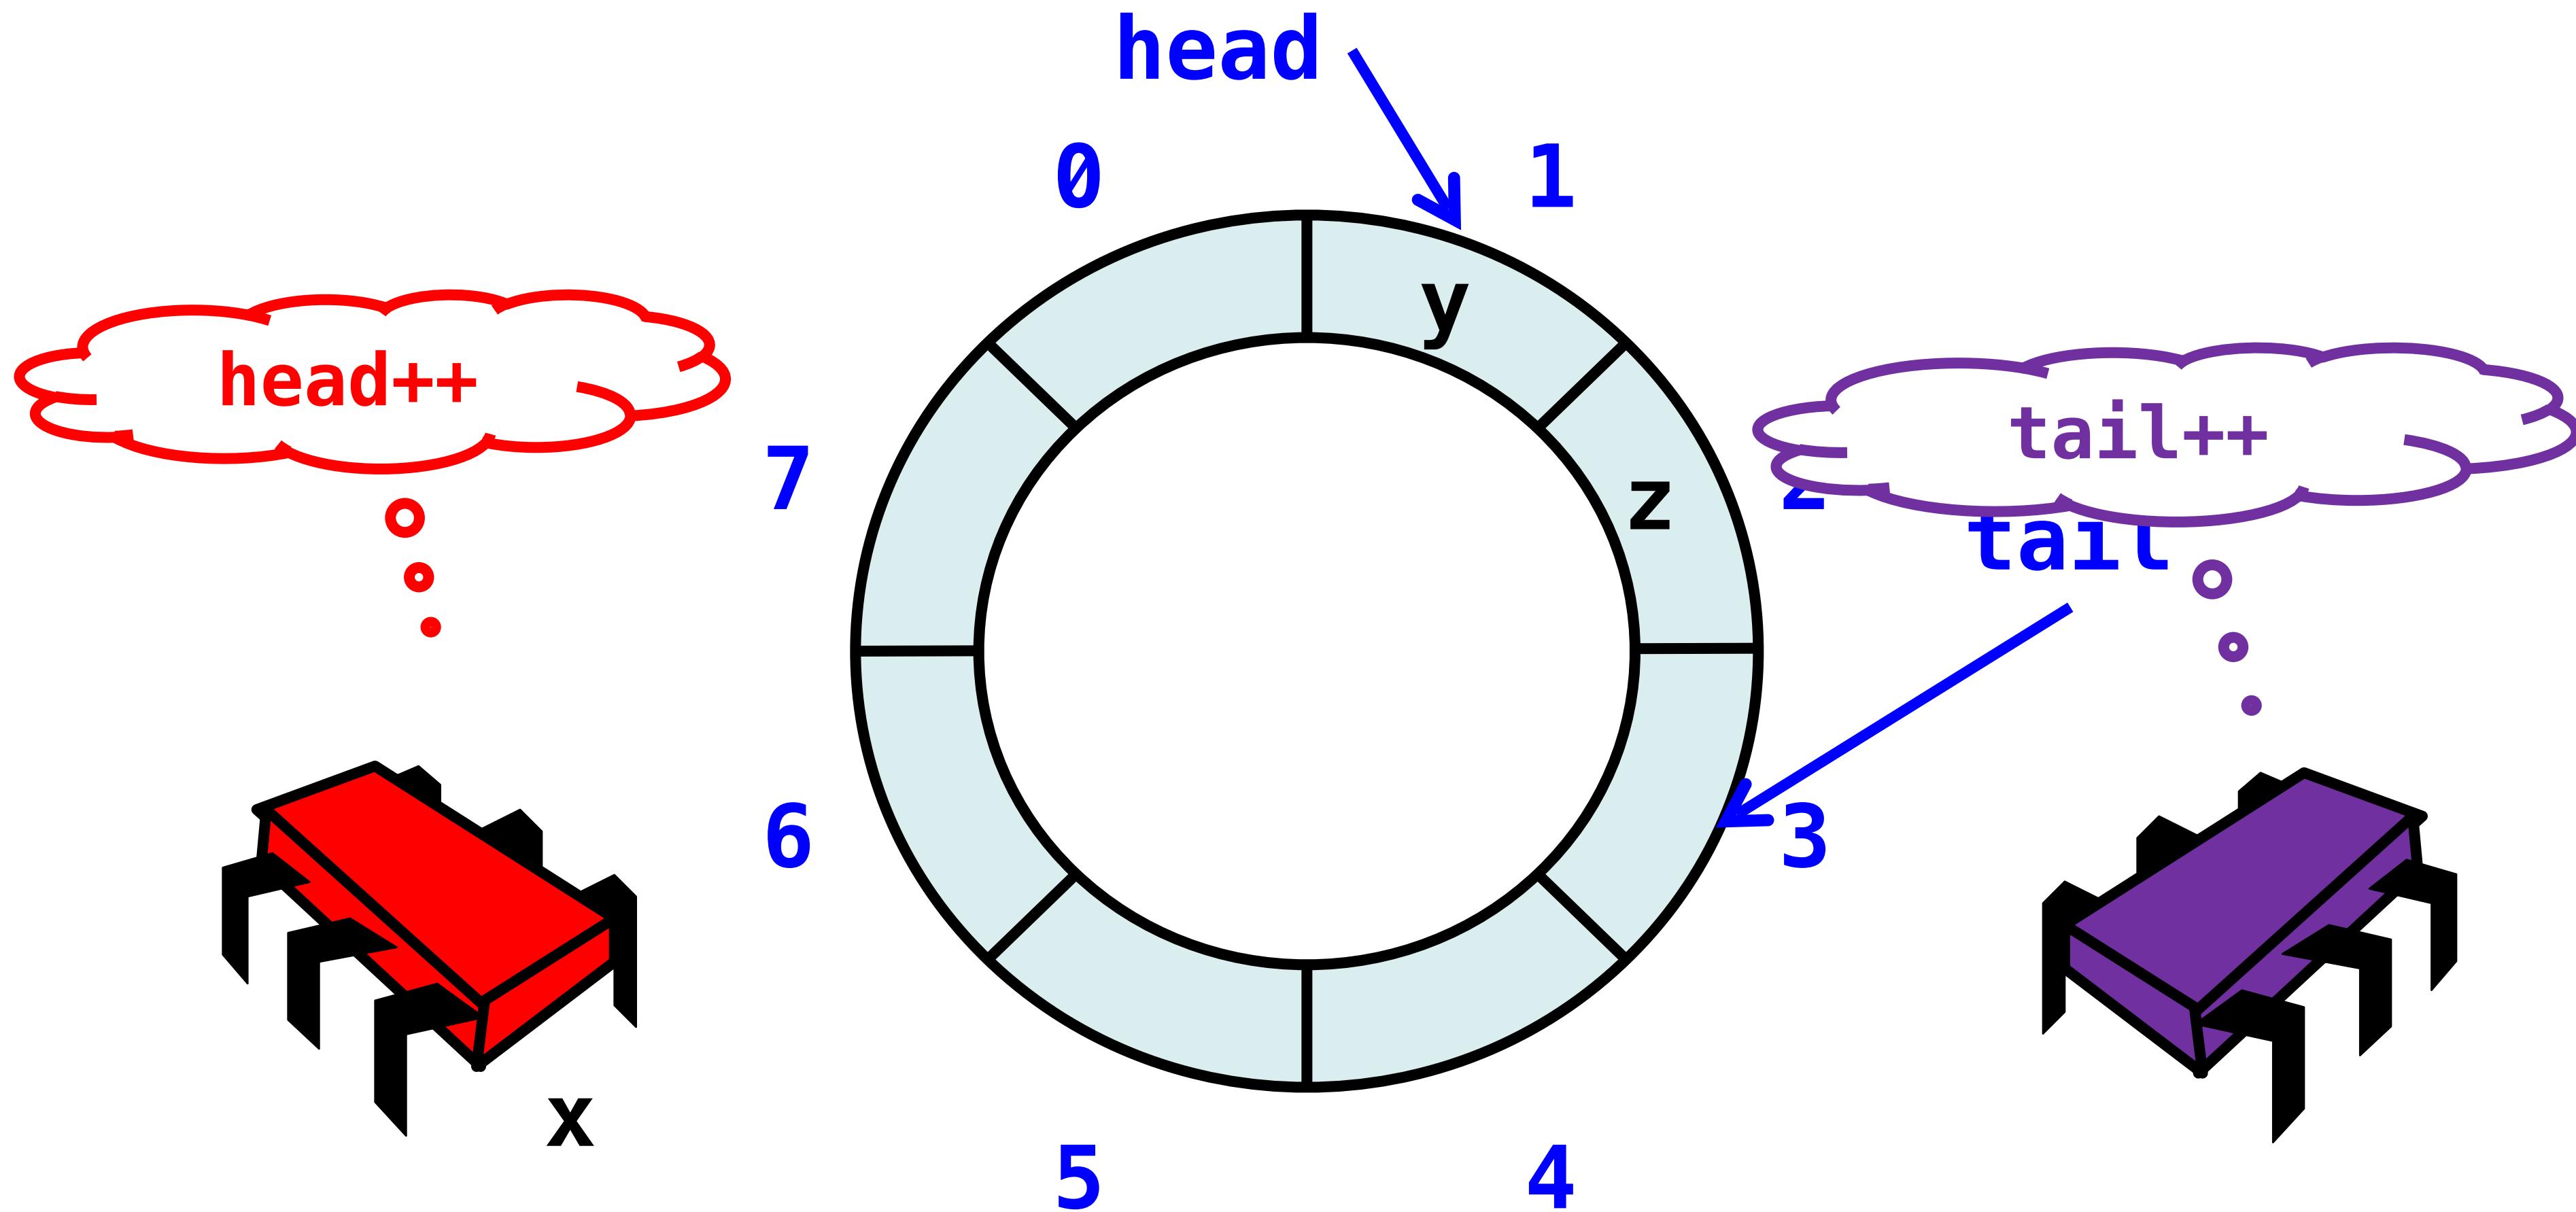
# Wait-free 2-thread queue



# Wait-free 2-thread queue



# Wait-free 2-thread queue



# Wait-free 2-thread queue

**No locks needed!**

```
(** Enqueue – should be called by only ONE thread *)    (** Dequeue – should be called by only ONE thread *)
let enq q x =
  (* Check if queue is full *)
  if q.tail - q.head = q.capacity then
    raise Full;
  (* Write to the array *)
  q.items.(q.tail mod q.capacity) <- Some x;
  (* Advance tail *)
  q.tail <- q.tail + 1

let deq q =
  (* Check if queue is empty *)
  if q.tail = q.head then
    raise Empty;
  (* Read from the array *)
  match q.items.(q.head mod q.capacity) with
  | None -> assert false (* Should never happen *)
  | Some x ->
    (* Advance head *)
    q.head <- q.head + 1;
    x
```

# Wait-free 2-thread queue

**No locks needed!**

```
(** Enqueue – should be called by only ONE thread *)    (** Dequeue – should be called by only ONE thread *)
let enq q x =
  (* Check if queue is full *)
  if q.tail - q.head = q.capacity then
    raise Full;
  (* Write to the array *)
  q.items.(q.tail mod q.capacity) <- Some x;
  (* Advance tail *)
  q.tail <- q.tail + 1
let deq q =
  (* Check if queue is empty *)
  if q.tail = q.head then
    raise Empty;
  (* Read from the array *)
  match q.items.(q.head mod q.capacity) with
  | None -> assert false (* Should never happen *)
  | Some x ->
    (* Advance head *)
    q.head <- q.head + 1;
    x
```

**Demo**

# Wait-free 2-thread queue

# *No locks needed!*

How do we define “correct” when modifications are not mutually exclusive?

## Demo

# **Concurrency Specification**

# What *is* a concurrent queue?

- Need a way to **specify** a concurrent queue object
- Need a way to **prove** that an algorithm implements the object's specification
- Lets talk about object specifications ...

# Correctness and Progress

- In a concurrent setting, we need to specify both the **safety** and the **liveness** properties of an object
- Need a way to define
  - when an implementation is **correct**
  - the conditions under which it guarantees **progress**

# Correctness and Progress

- In a concurrent setting, we need to specify both the **safety** and the **liveness** properties of an object
- Need a way to define
  - when an implementation is **correct**
  - the conditions under which it guarantees **progress**

Lets begin with correctness

# Sequential Objects

- Each object has a ***state***
  - Usually given by a set of ***fields***
  - Queue example: sequence of items
- Each object has a set of ***methods***
  - Only way to manipulate state
  - Queue example: **enq** and **deq** methods

# Sequential Specifications

- If (*precondition*)
  - the object is in such-and-such a state
  - before you call the method,

# Sequential Specifications

- If (*precondition*)
  - the object is in such-and-such a state
  - before you call the method,
- Then (*postcondition*)
  - the method will return a particular value
  - or throw a particular exception.

# Sequential Specifications

- If (*precondition*)
  - the object is in such-and-such a state
  - before you call the method,
- Then (*postcondition*)
  - the method will return a particular value
  - or throw a particular exception.
- and (*postcondition, cont*)
  - the object will be in some other state
  - when the method returns,

# Pre and Post Conditions for Dequeue

- Precondition:
  - Queue is *non-empty*
- Postcondition:
  - Returns first item in queue
- Postcondition:
  - Removes first item in queue

# Pre and Post Conditions for Dequeue

- Precondition:
  - Queue is ***non-empty***
- Postcondition:
  - Returns first item in queue
- Postcondition:
  - Removes first item in queue
- Precondition:
  - Queue is ***empty***
- Postcondition:
  - Raises Empty exception
- Postcondition:
  - Queue state is unchanged

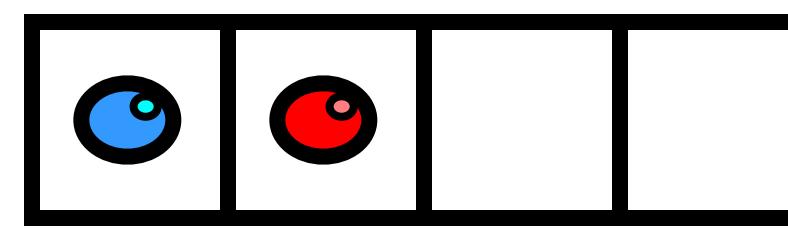
# Why Sequential Specifications Totally Rock

- Interactions among *methods* captured by side-effects on object state
  - State meaningful between method calls
- *Documentation* size is linear in the number of methods
  - Each method described in isolation
- Can add *new methods*
  - Without changing descriptions of old methods

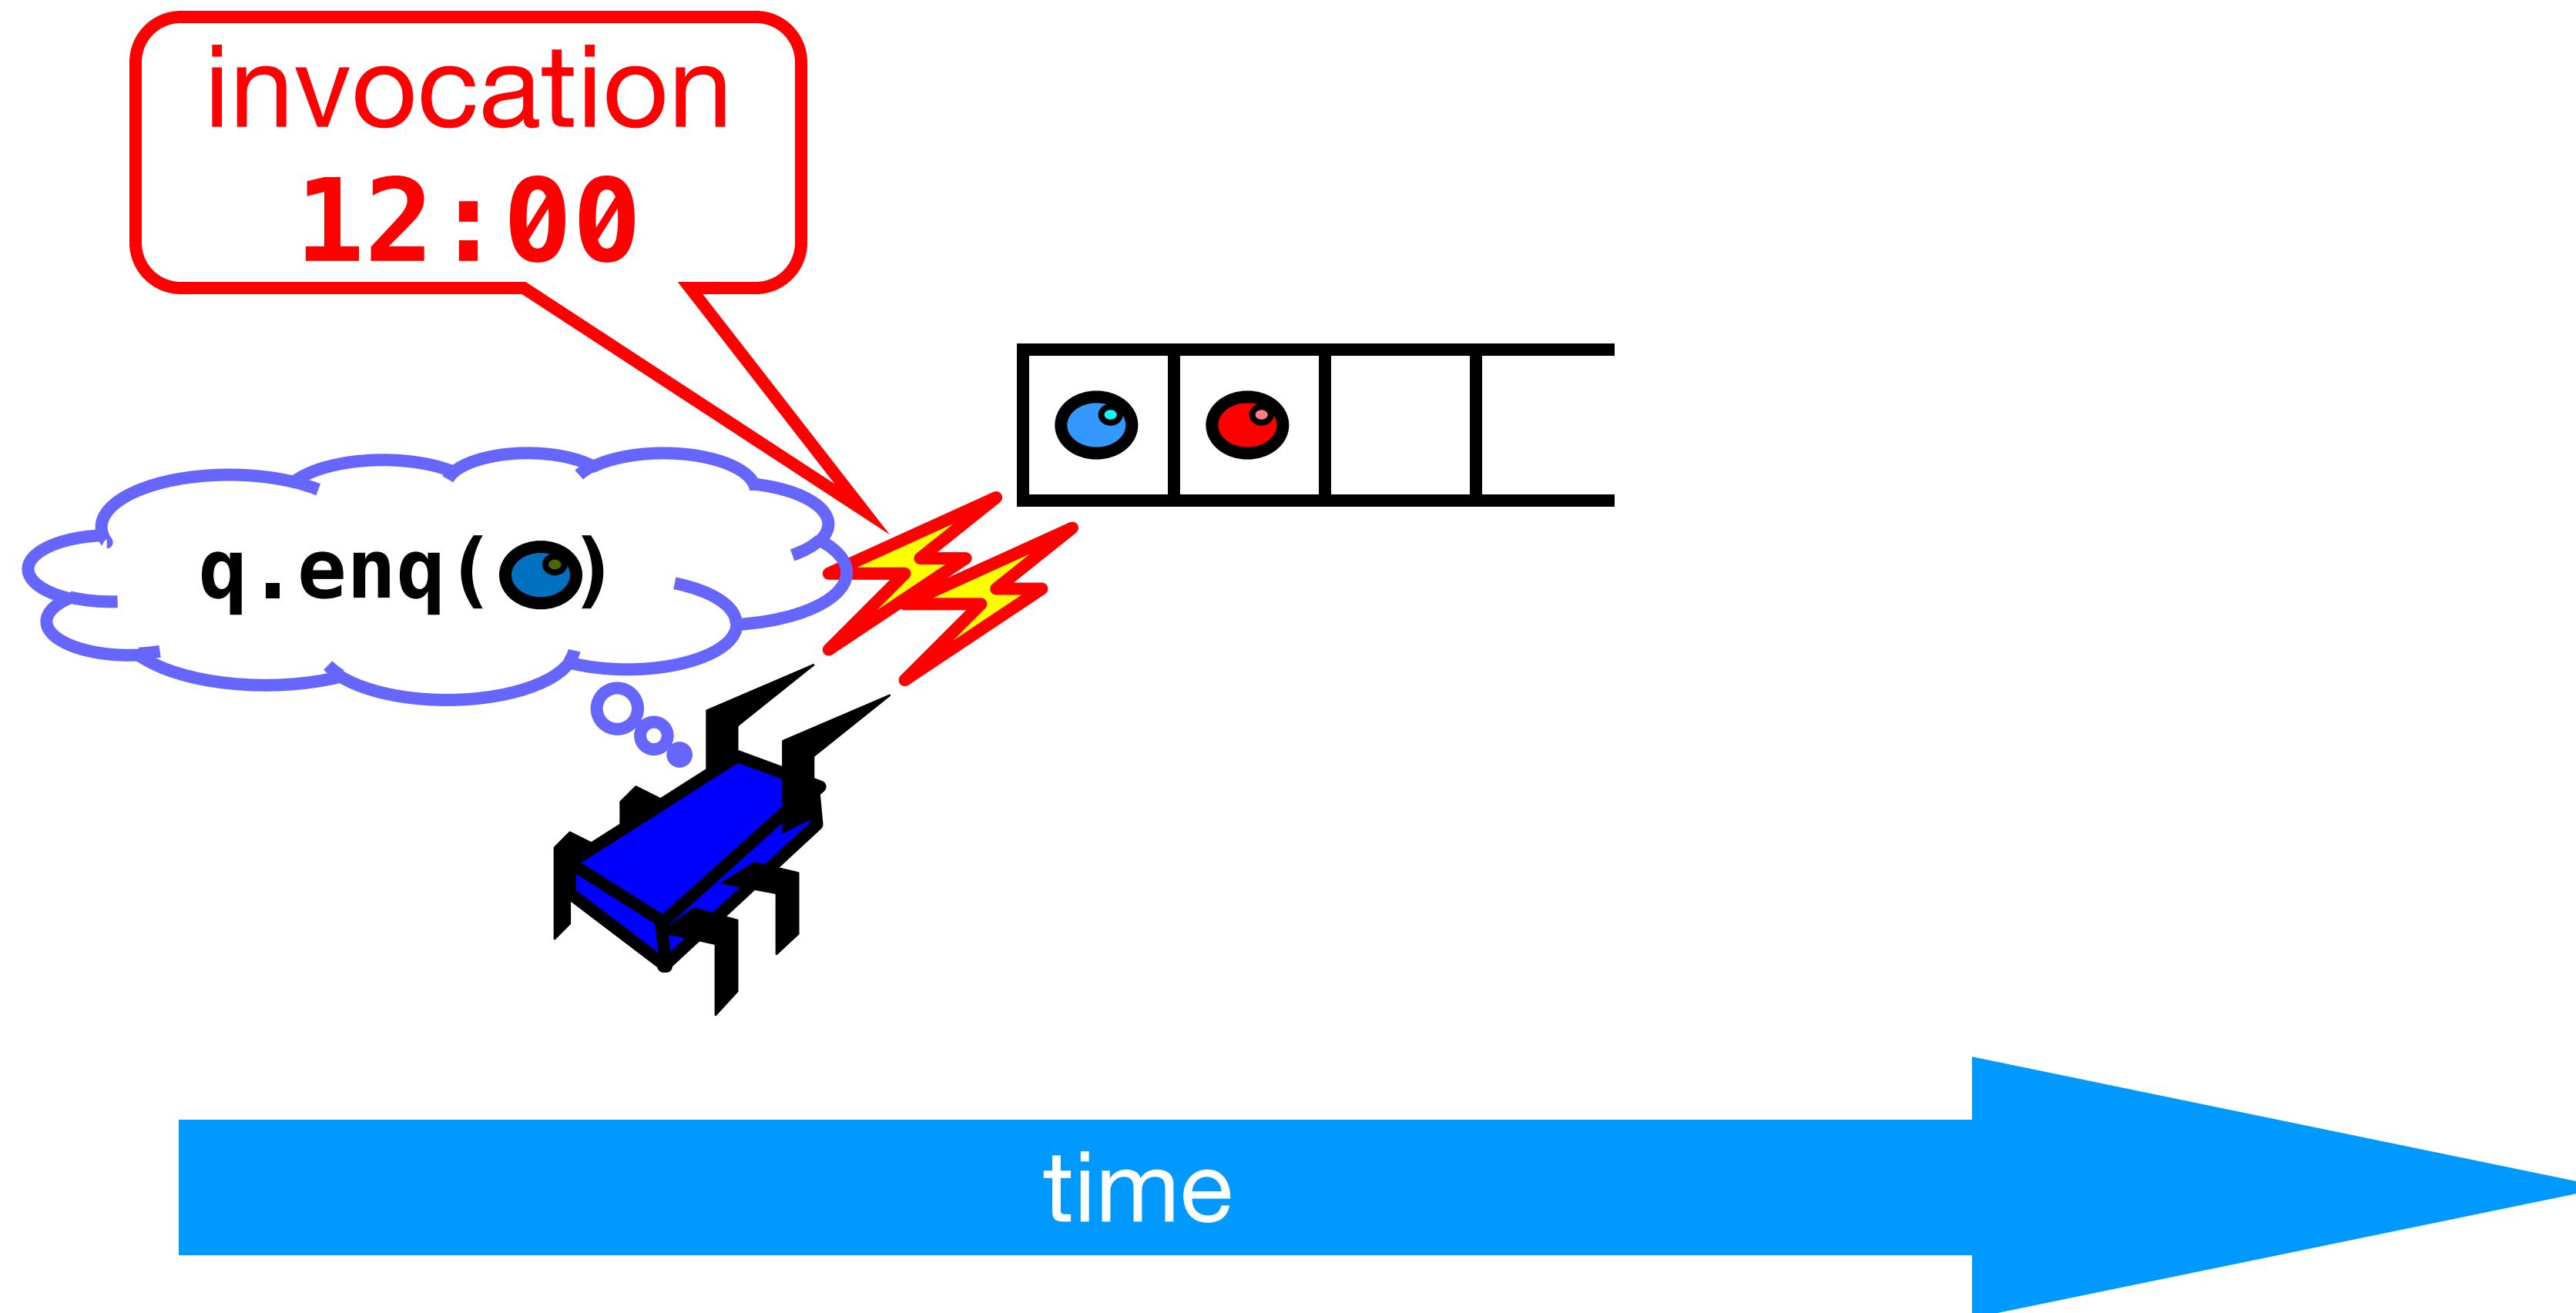
# What about concurrent Specifications?

- Methods?
- Documentation?
- Adding new methods?

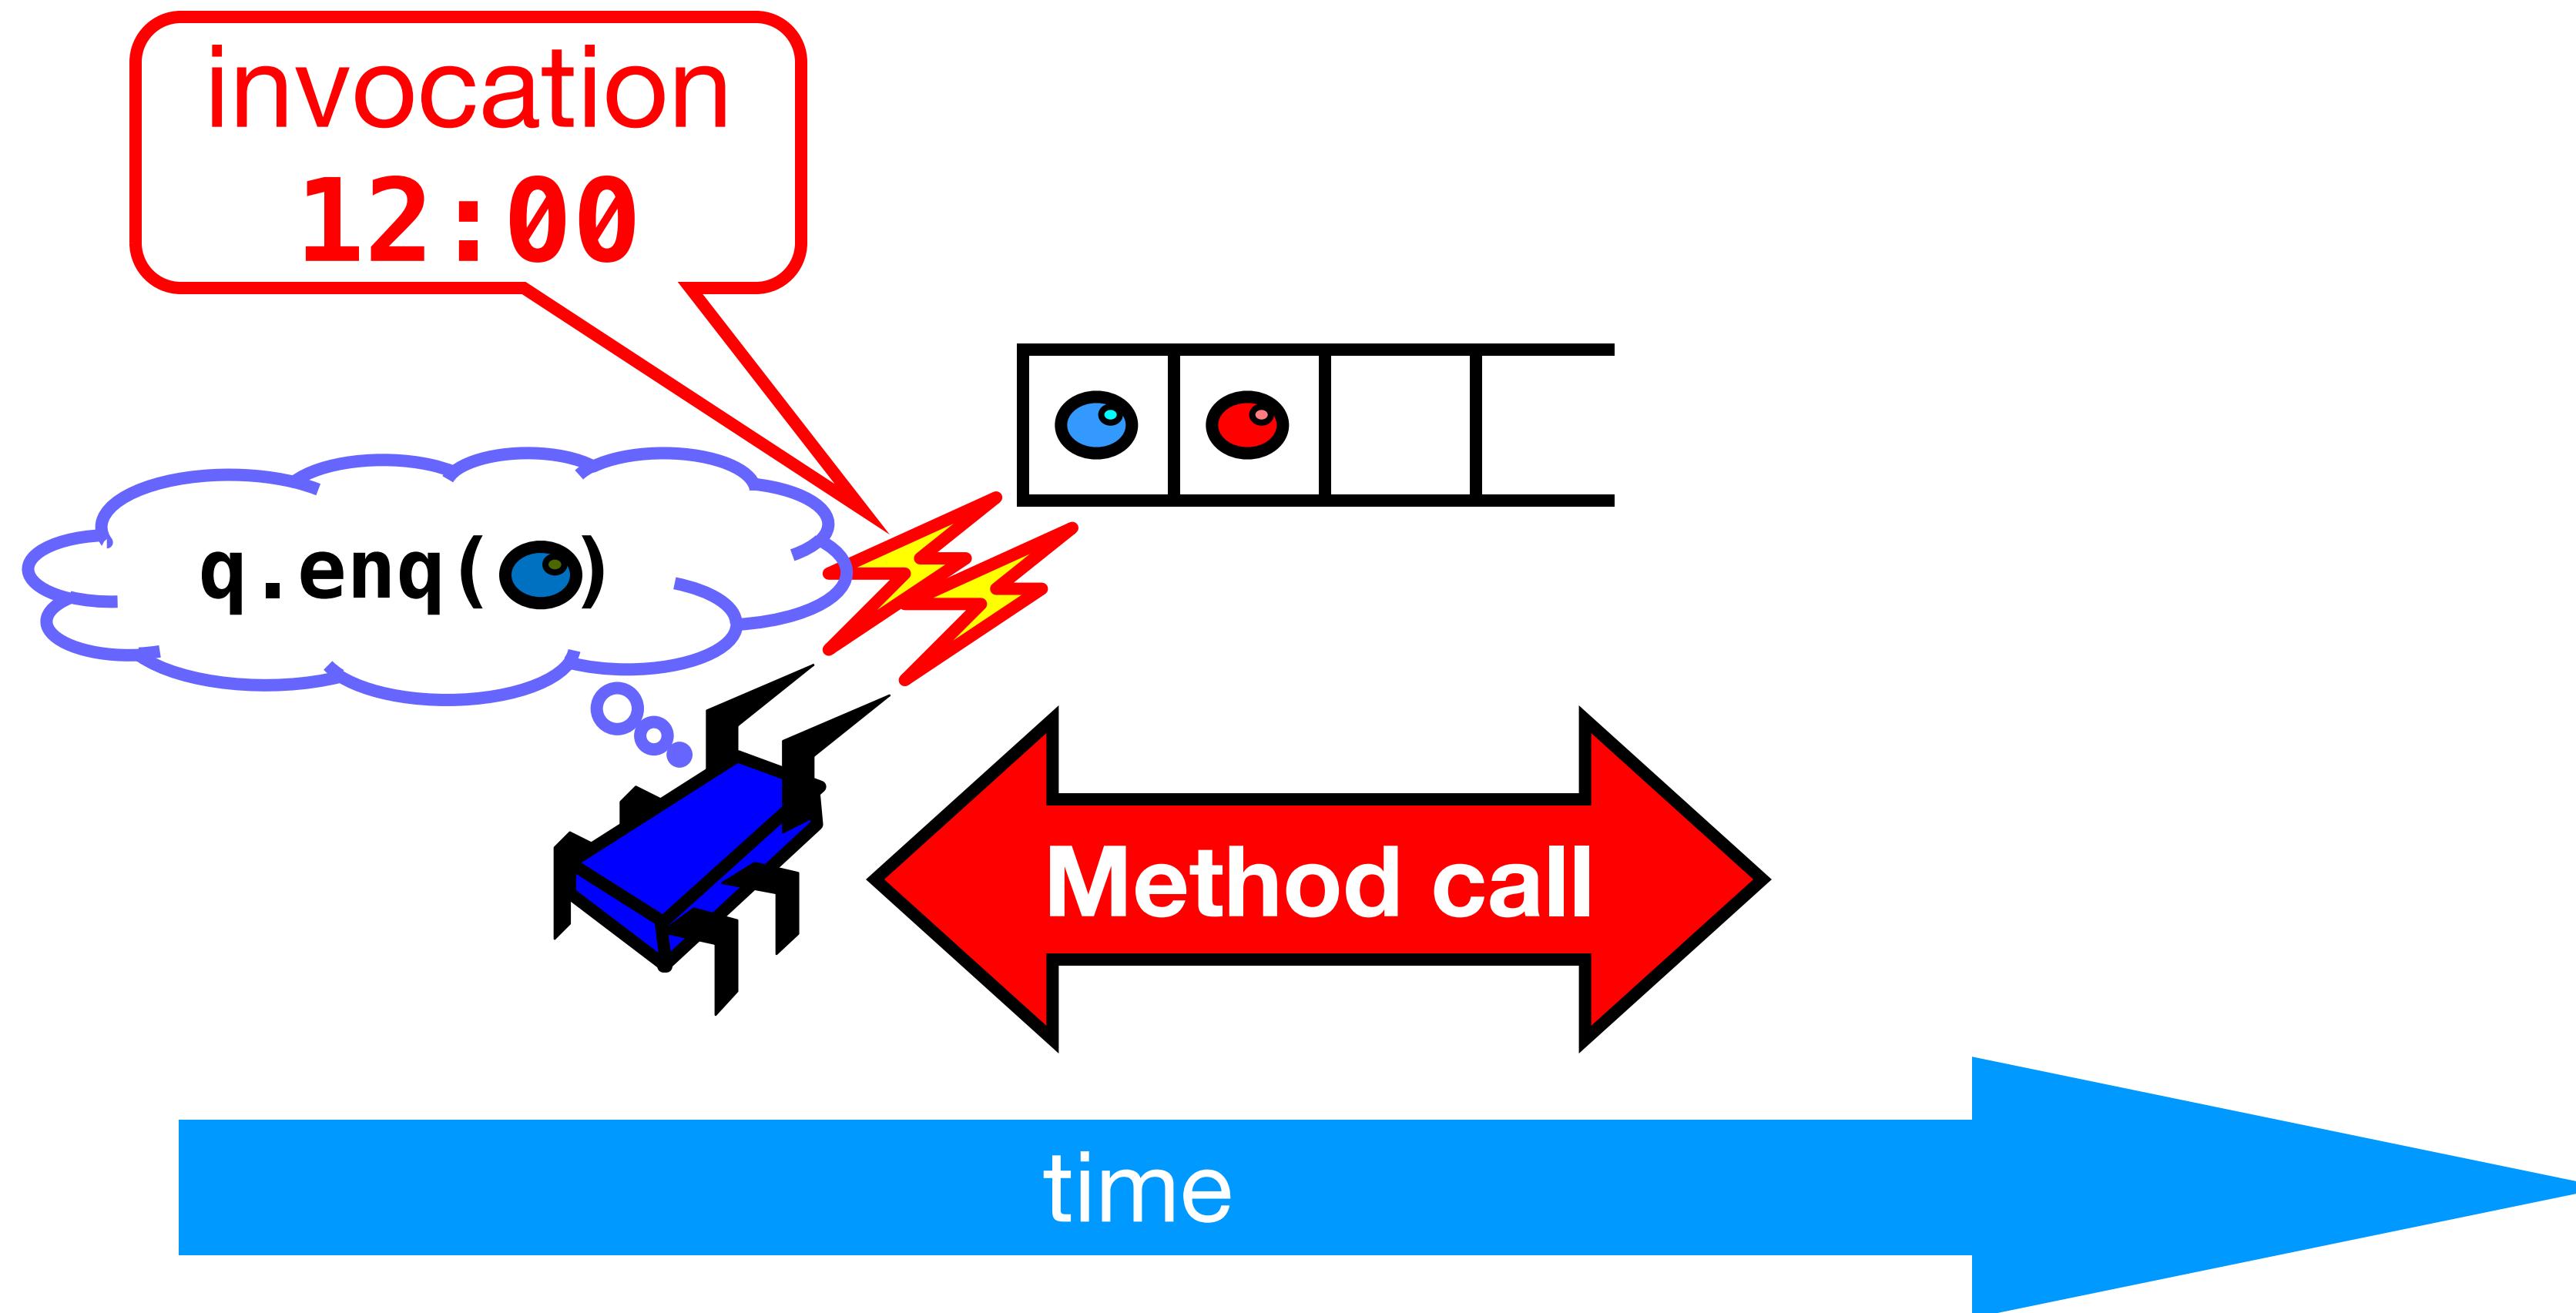
# Methods take time



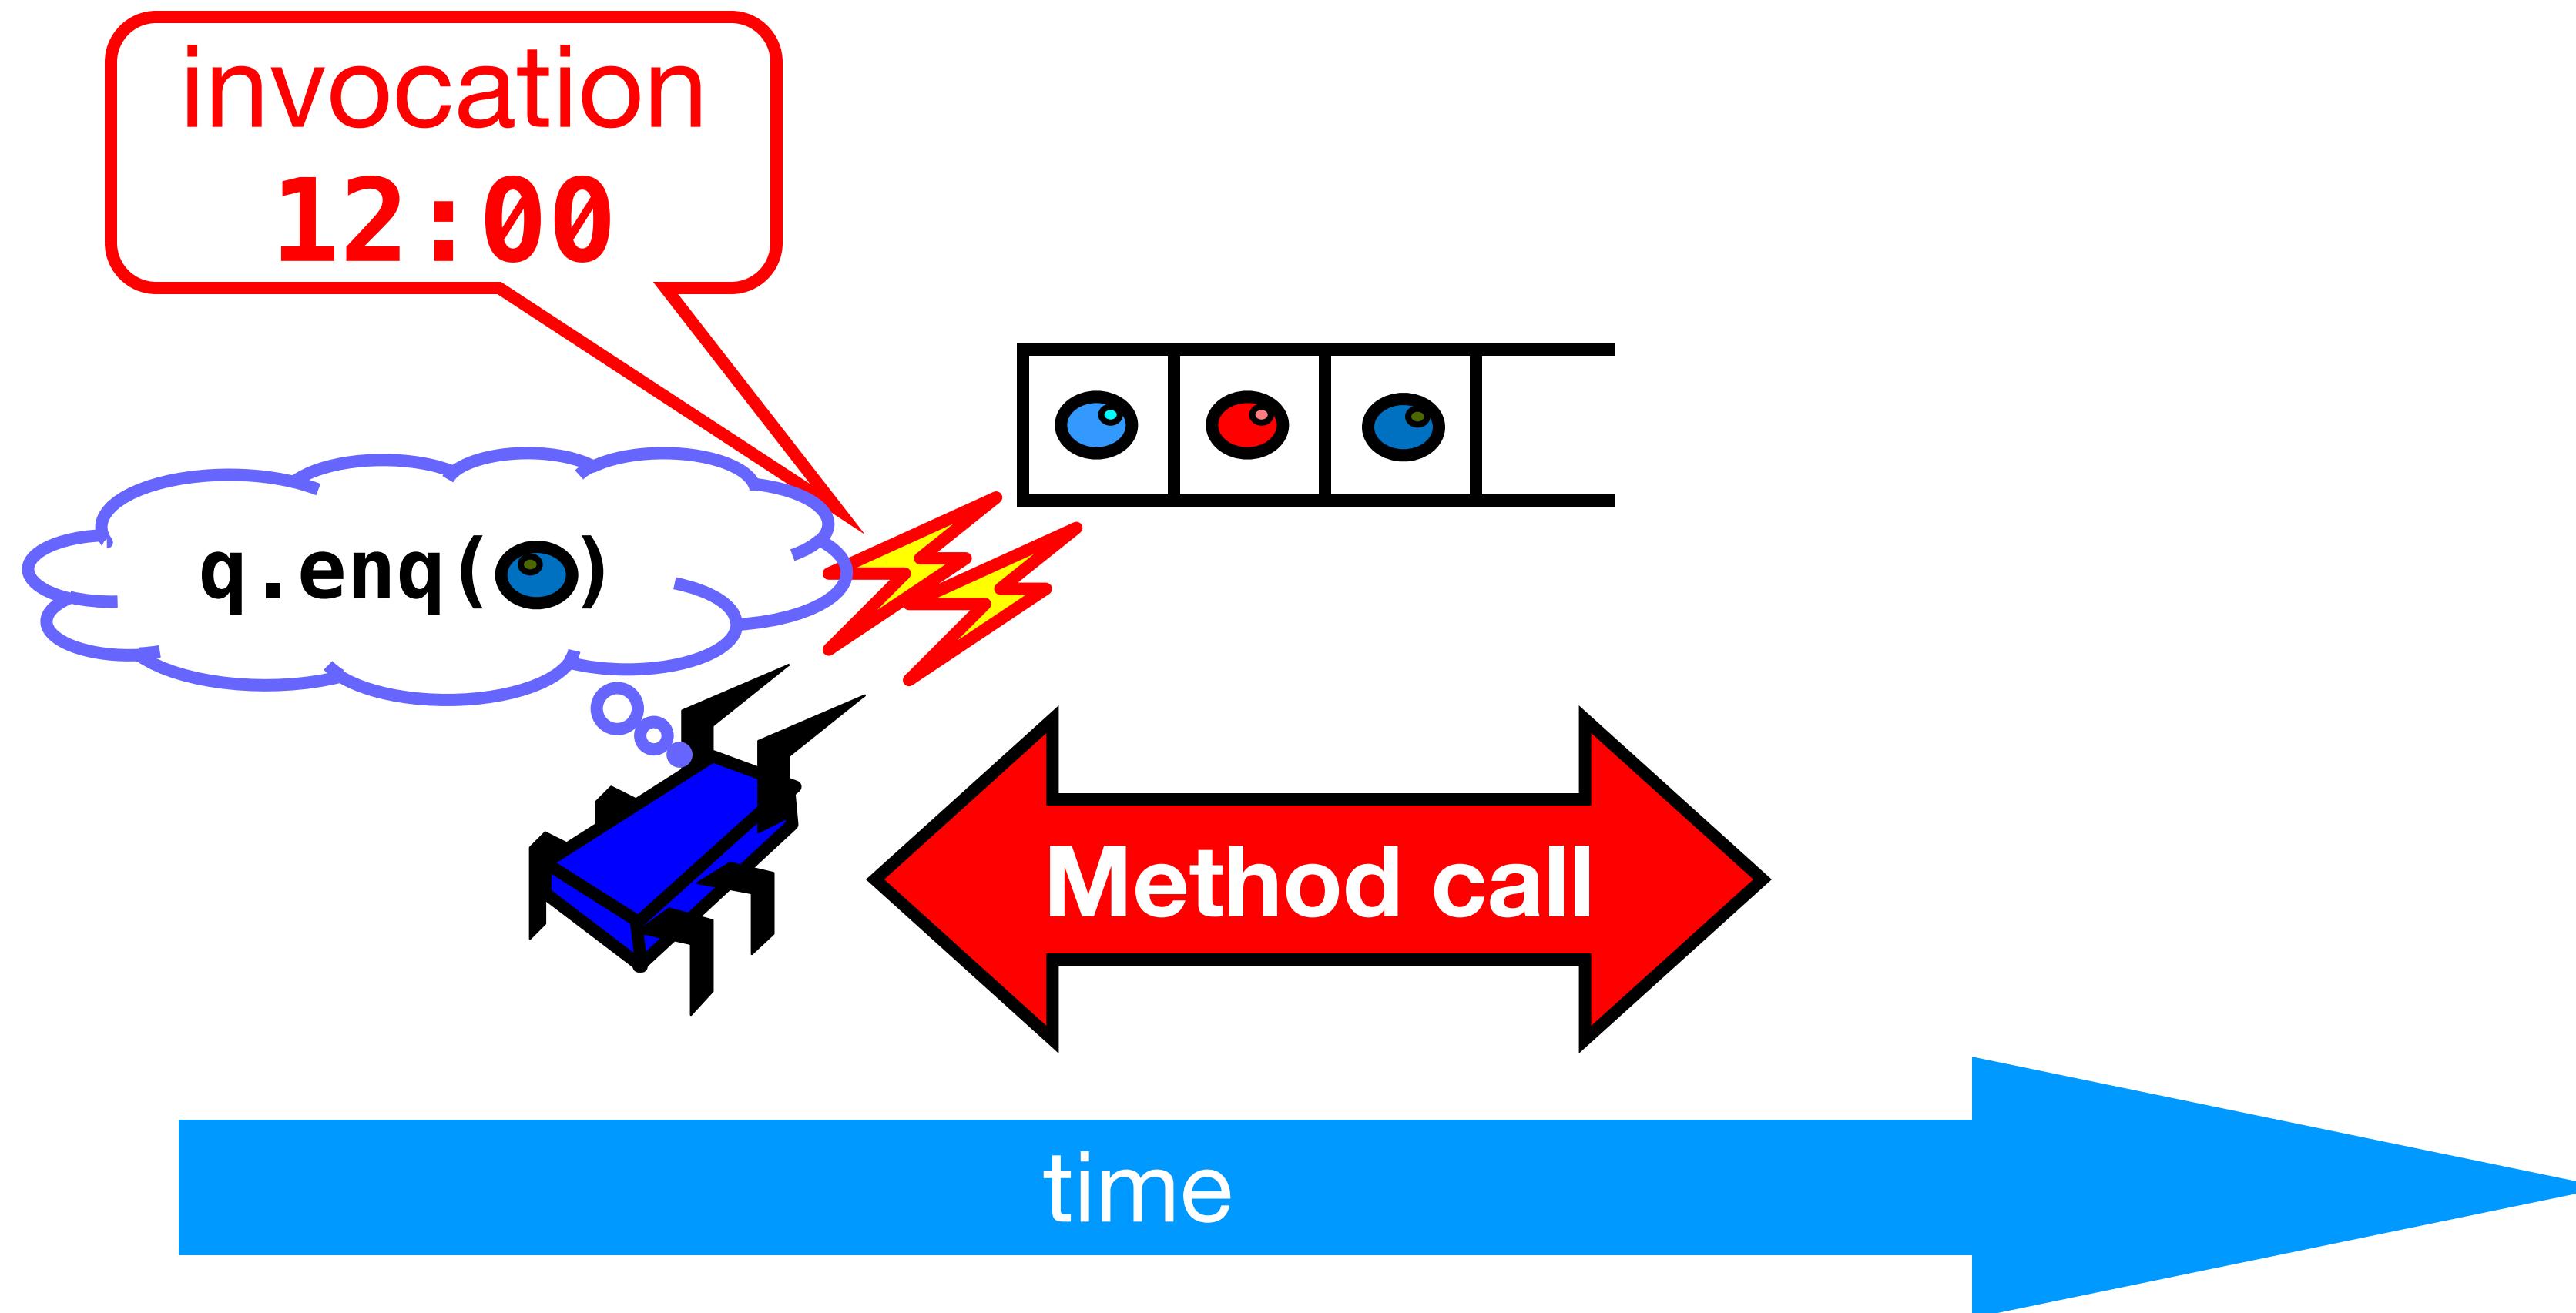
# Methods take time



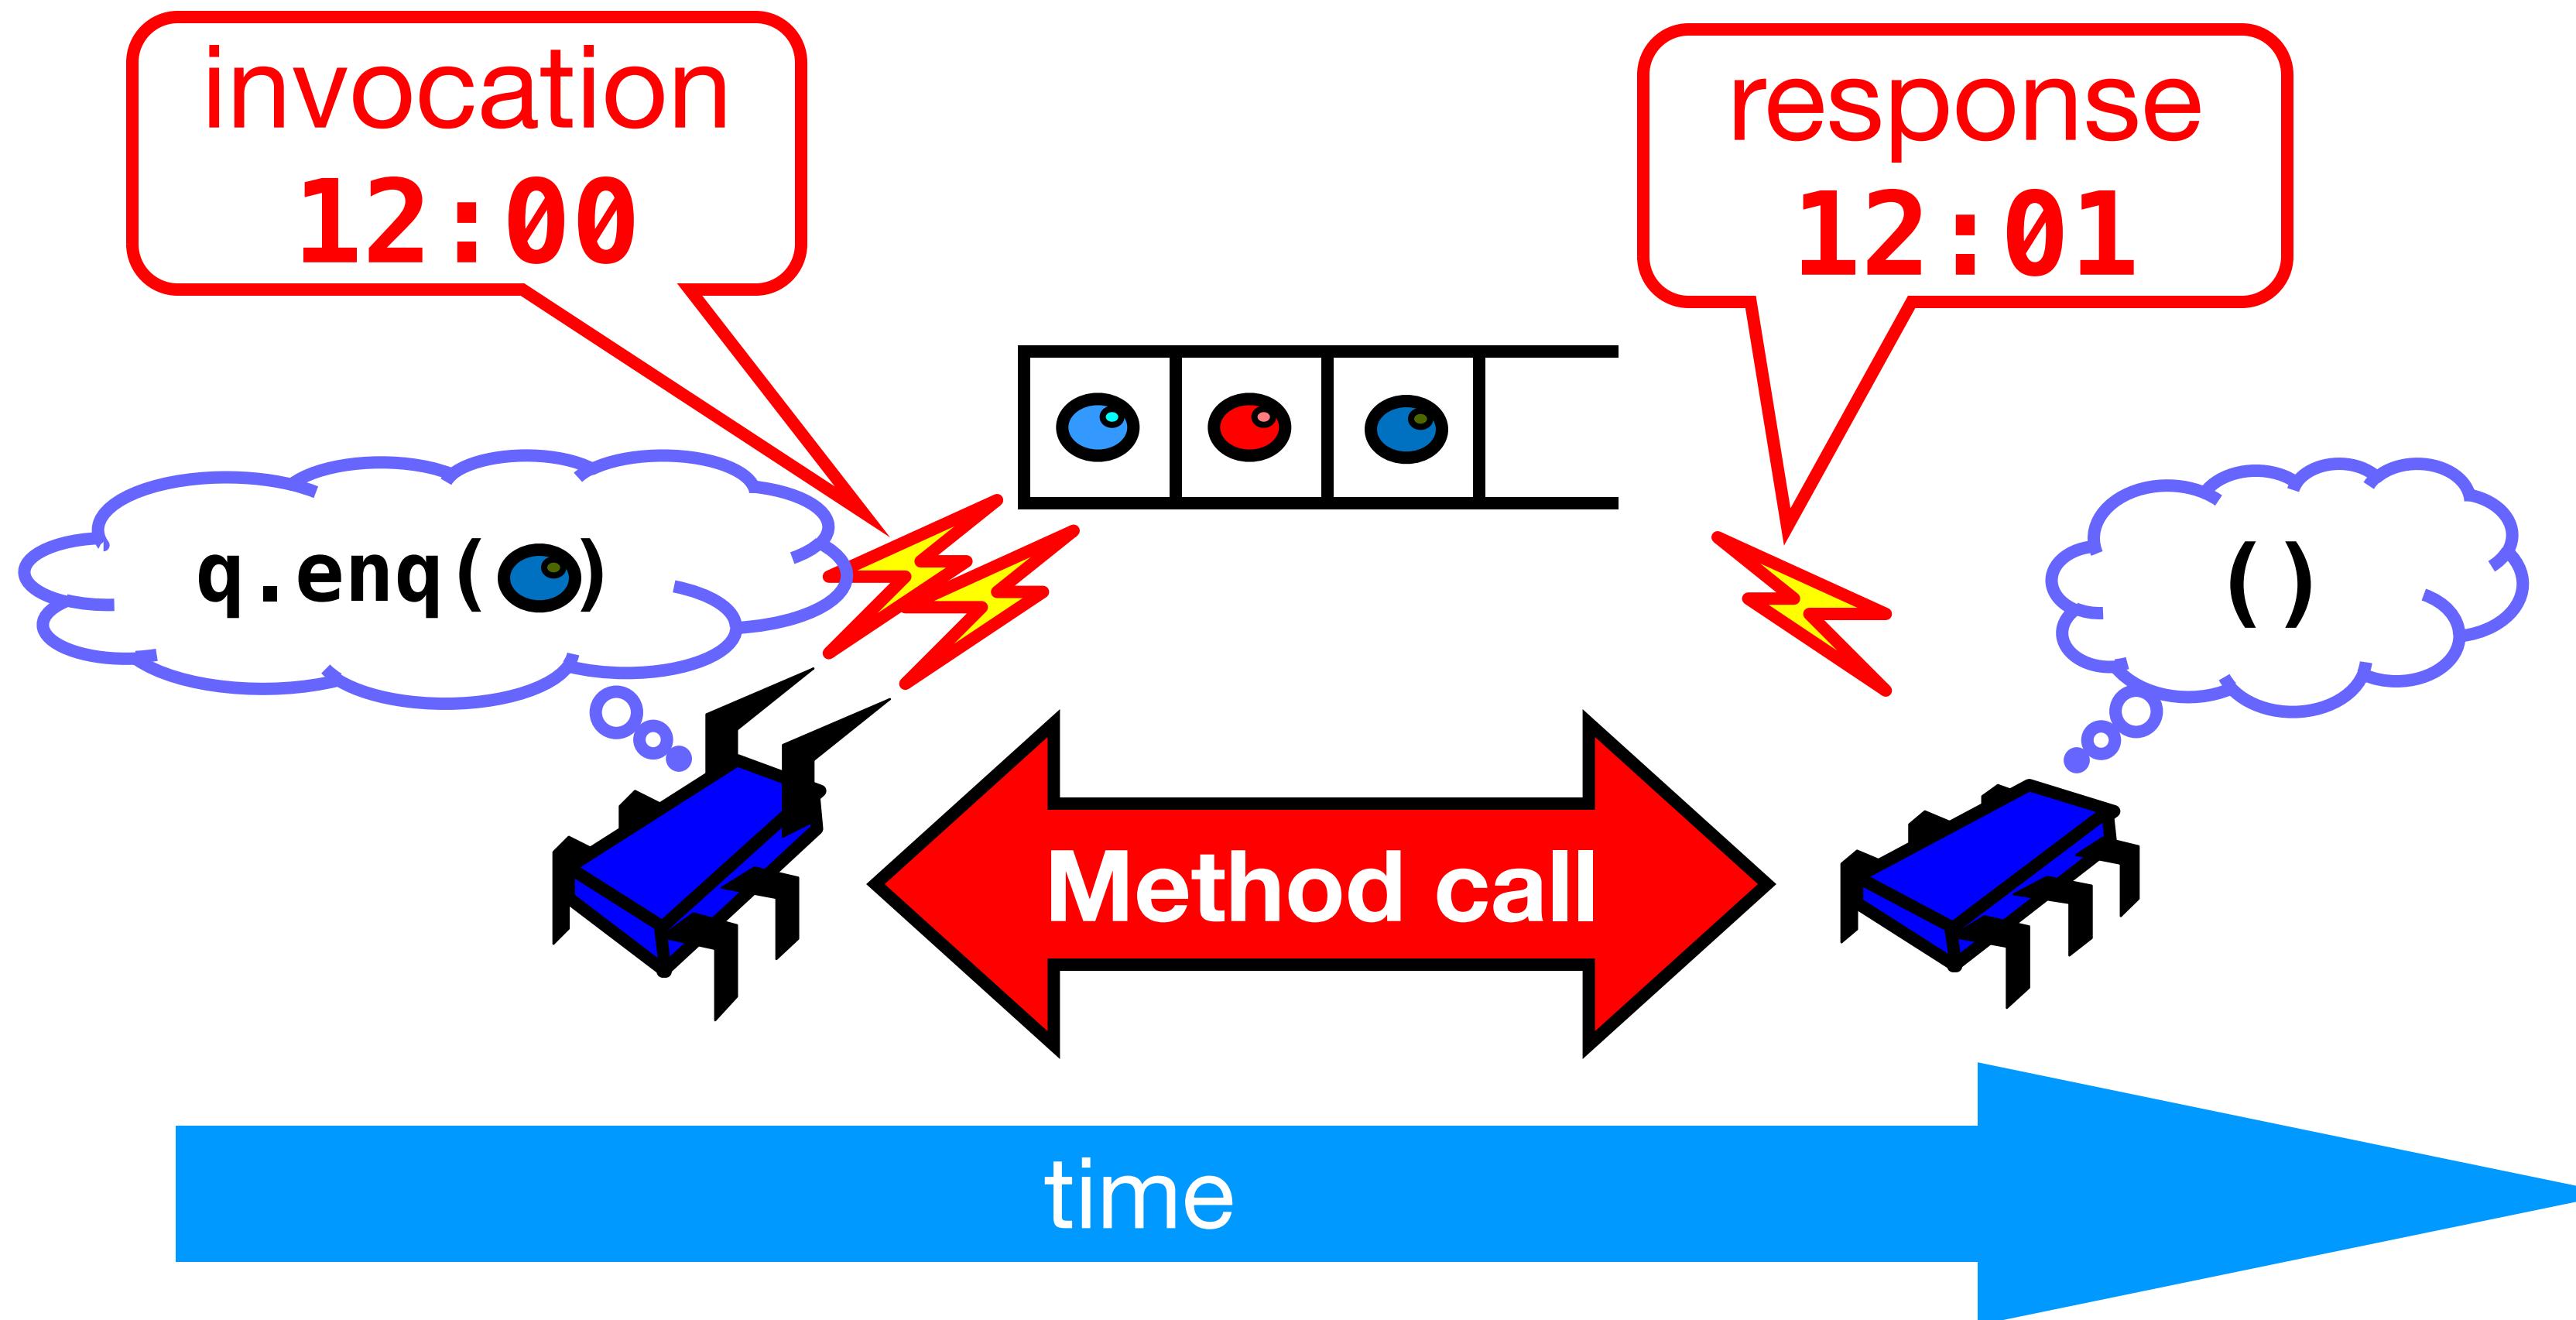
# Methods take time



# Methods take time



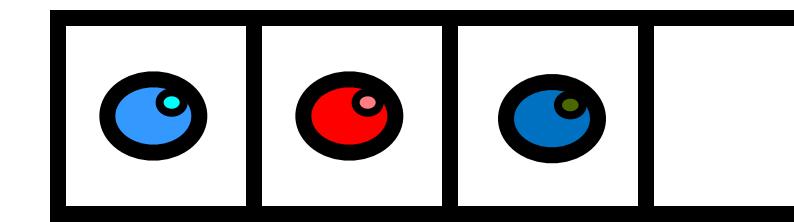
# Methods take time



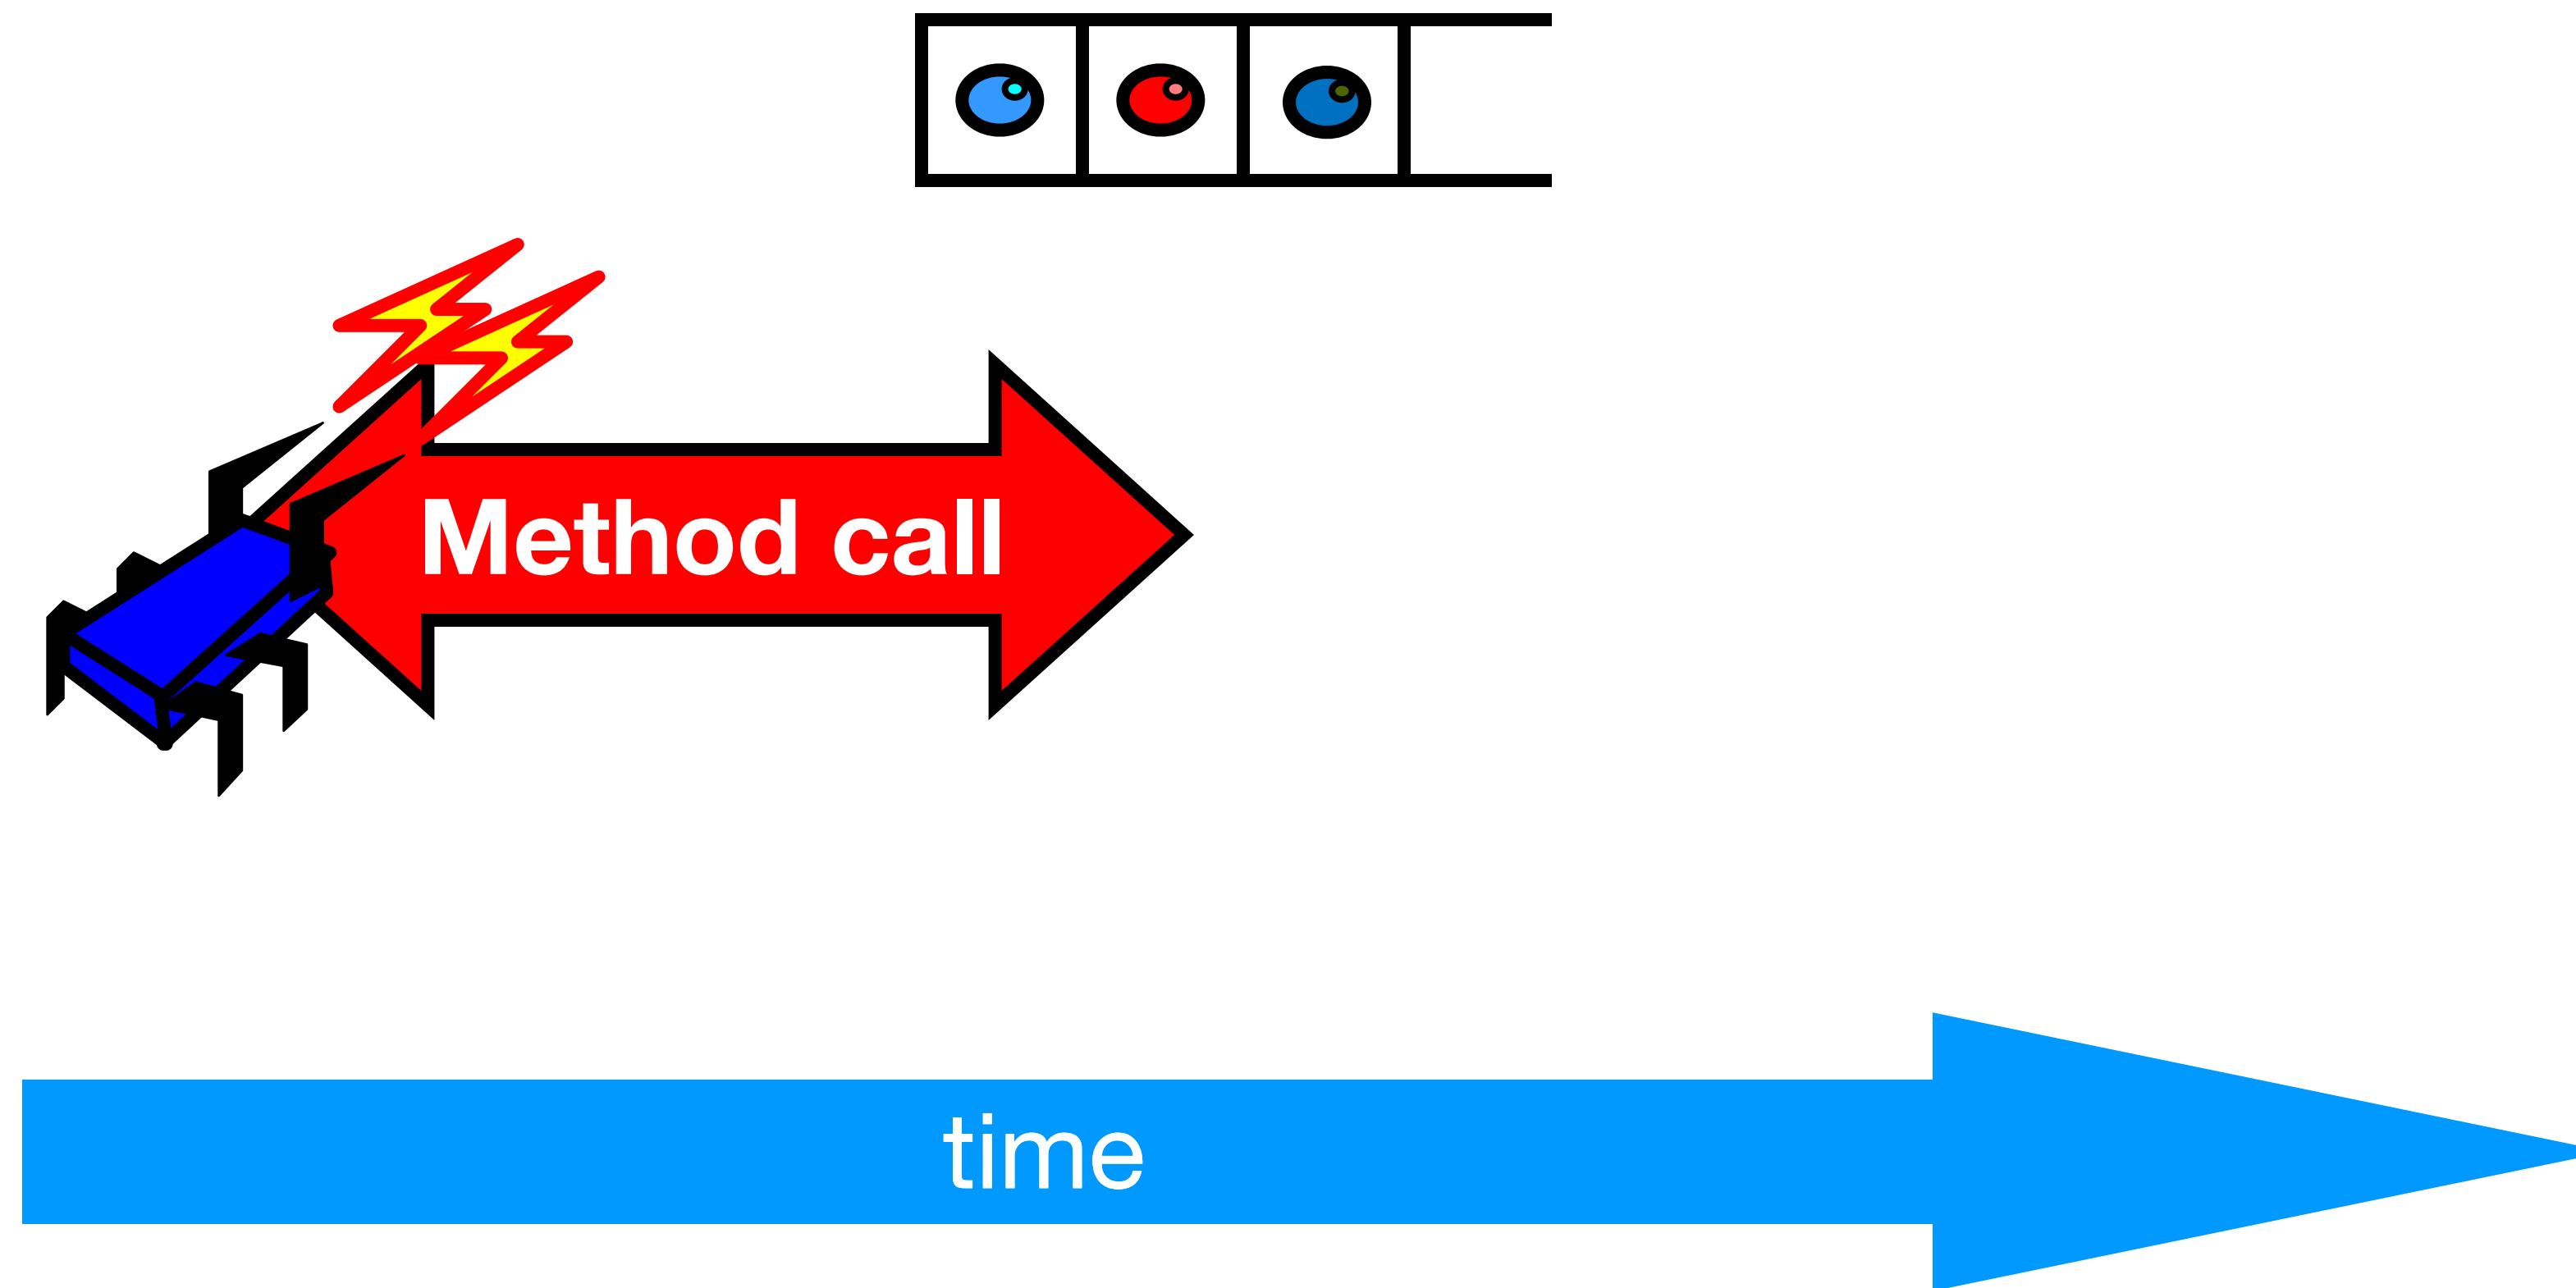
# Sequential vs Concurrent

- Sequential
  - Methods take time? Who knew?
- Concurrent
  - Method call is not an ***event***
  - Method call is an ***interval***

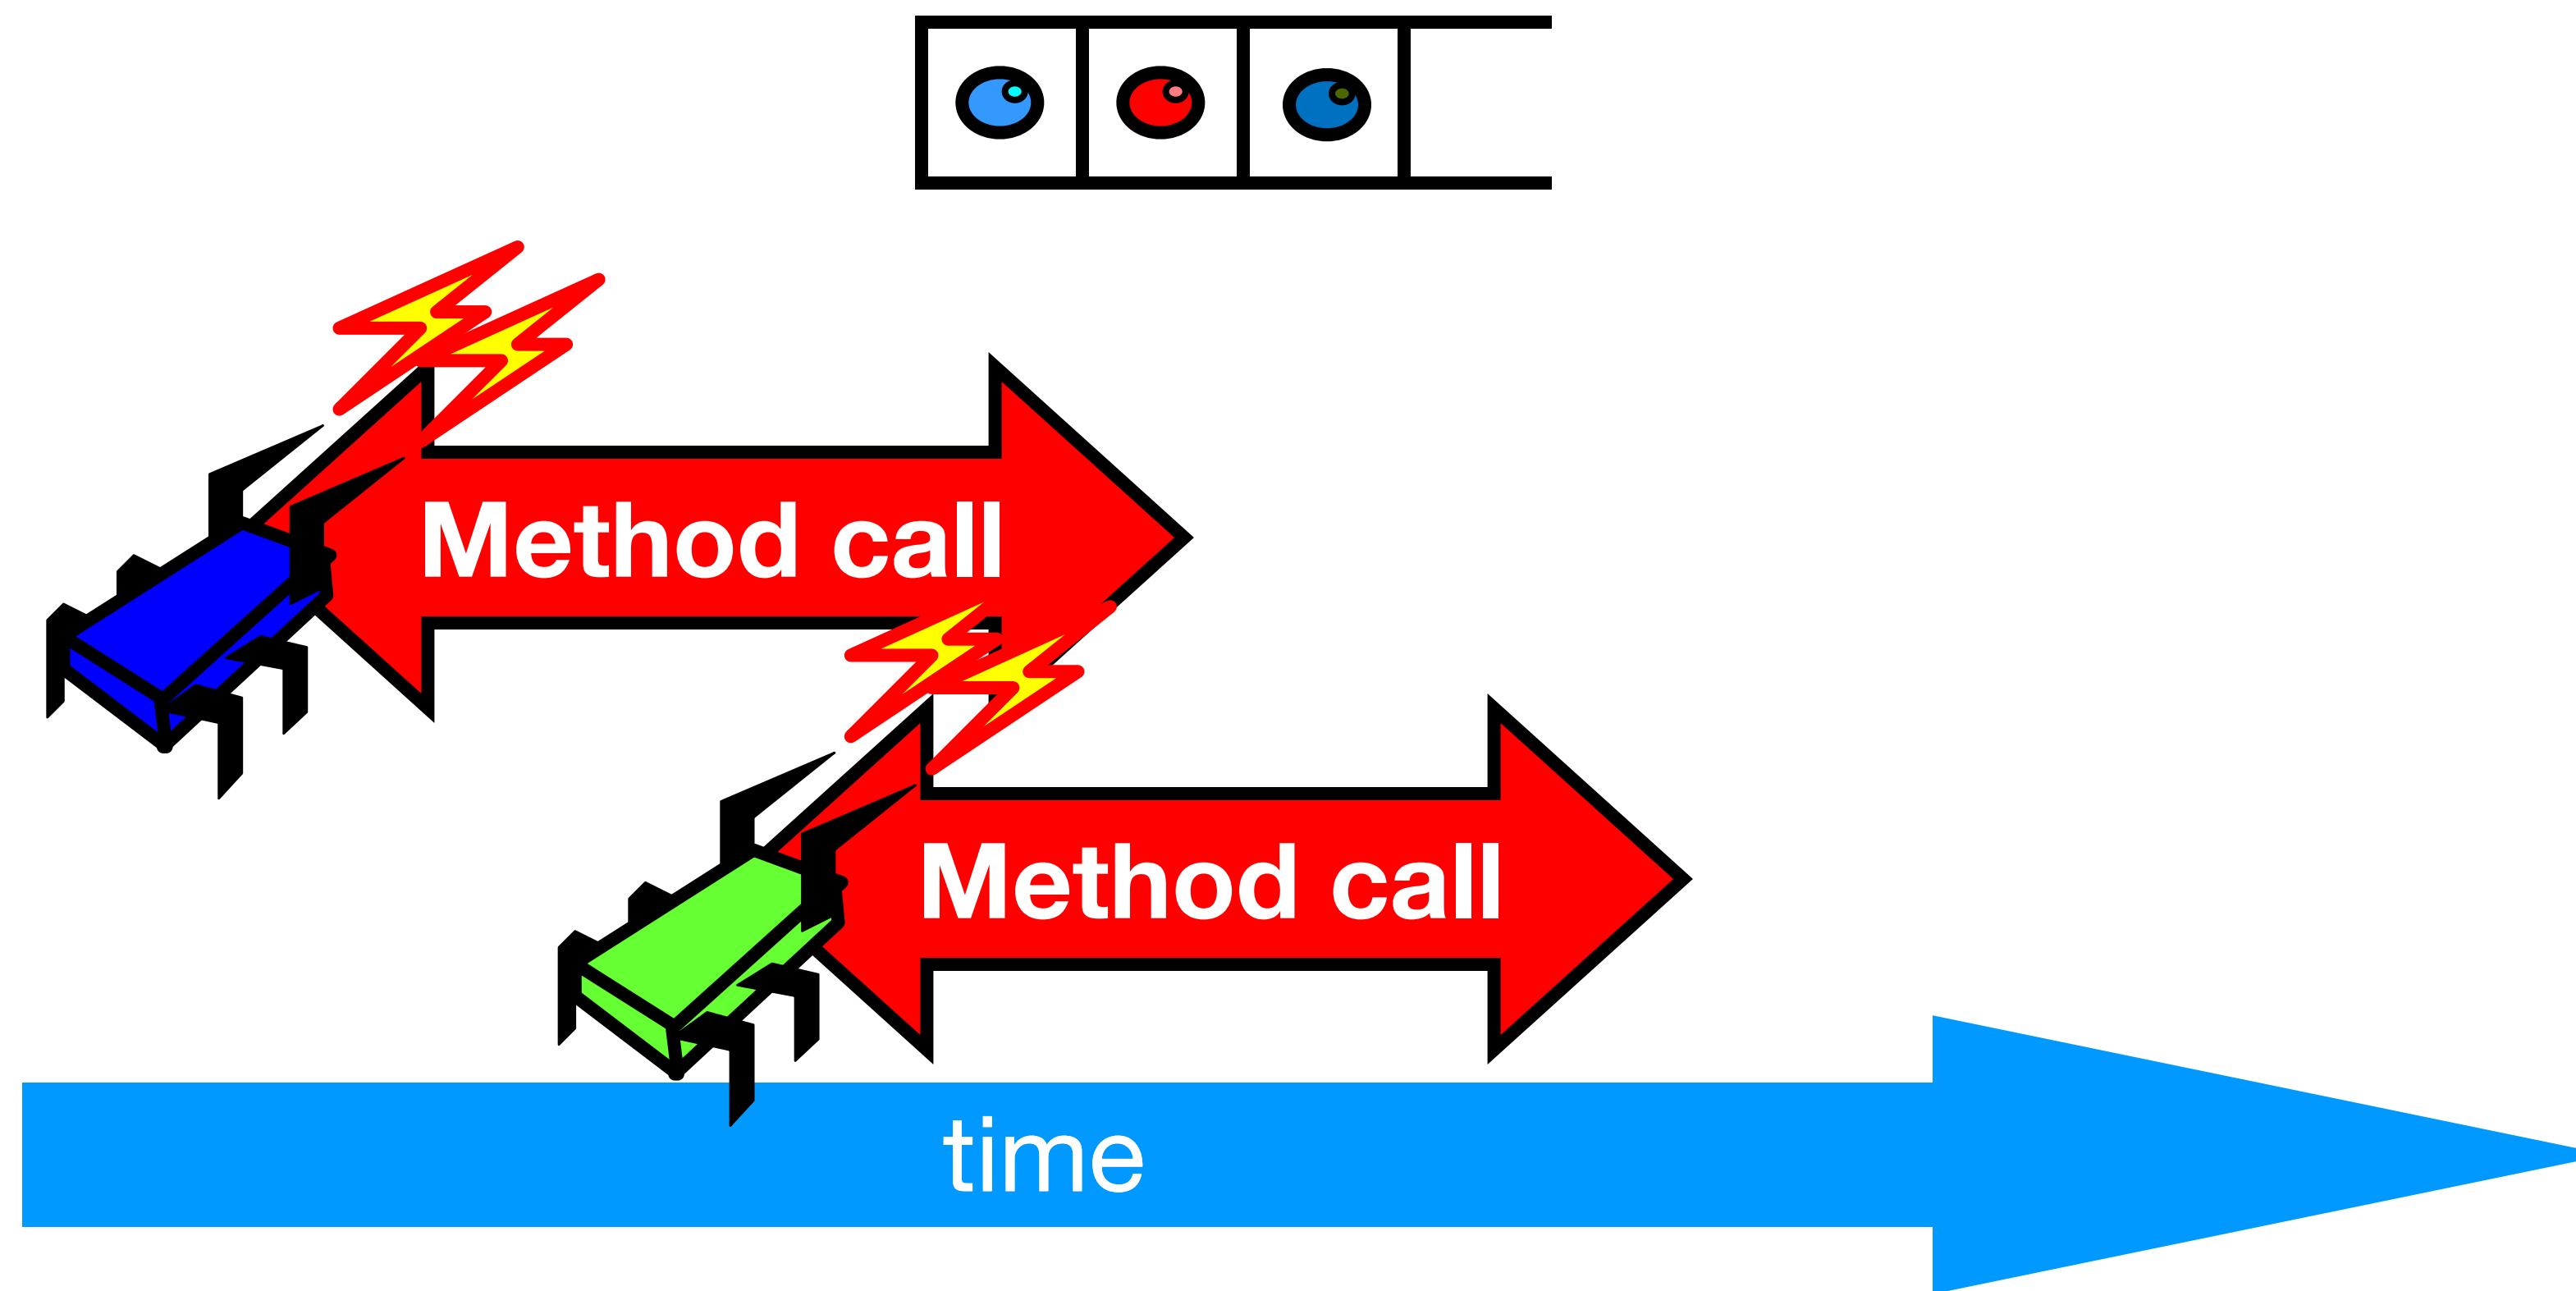
# Concurrent Methods Take Overlapping Time



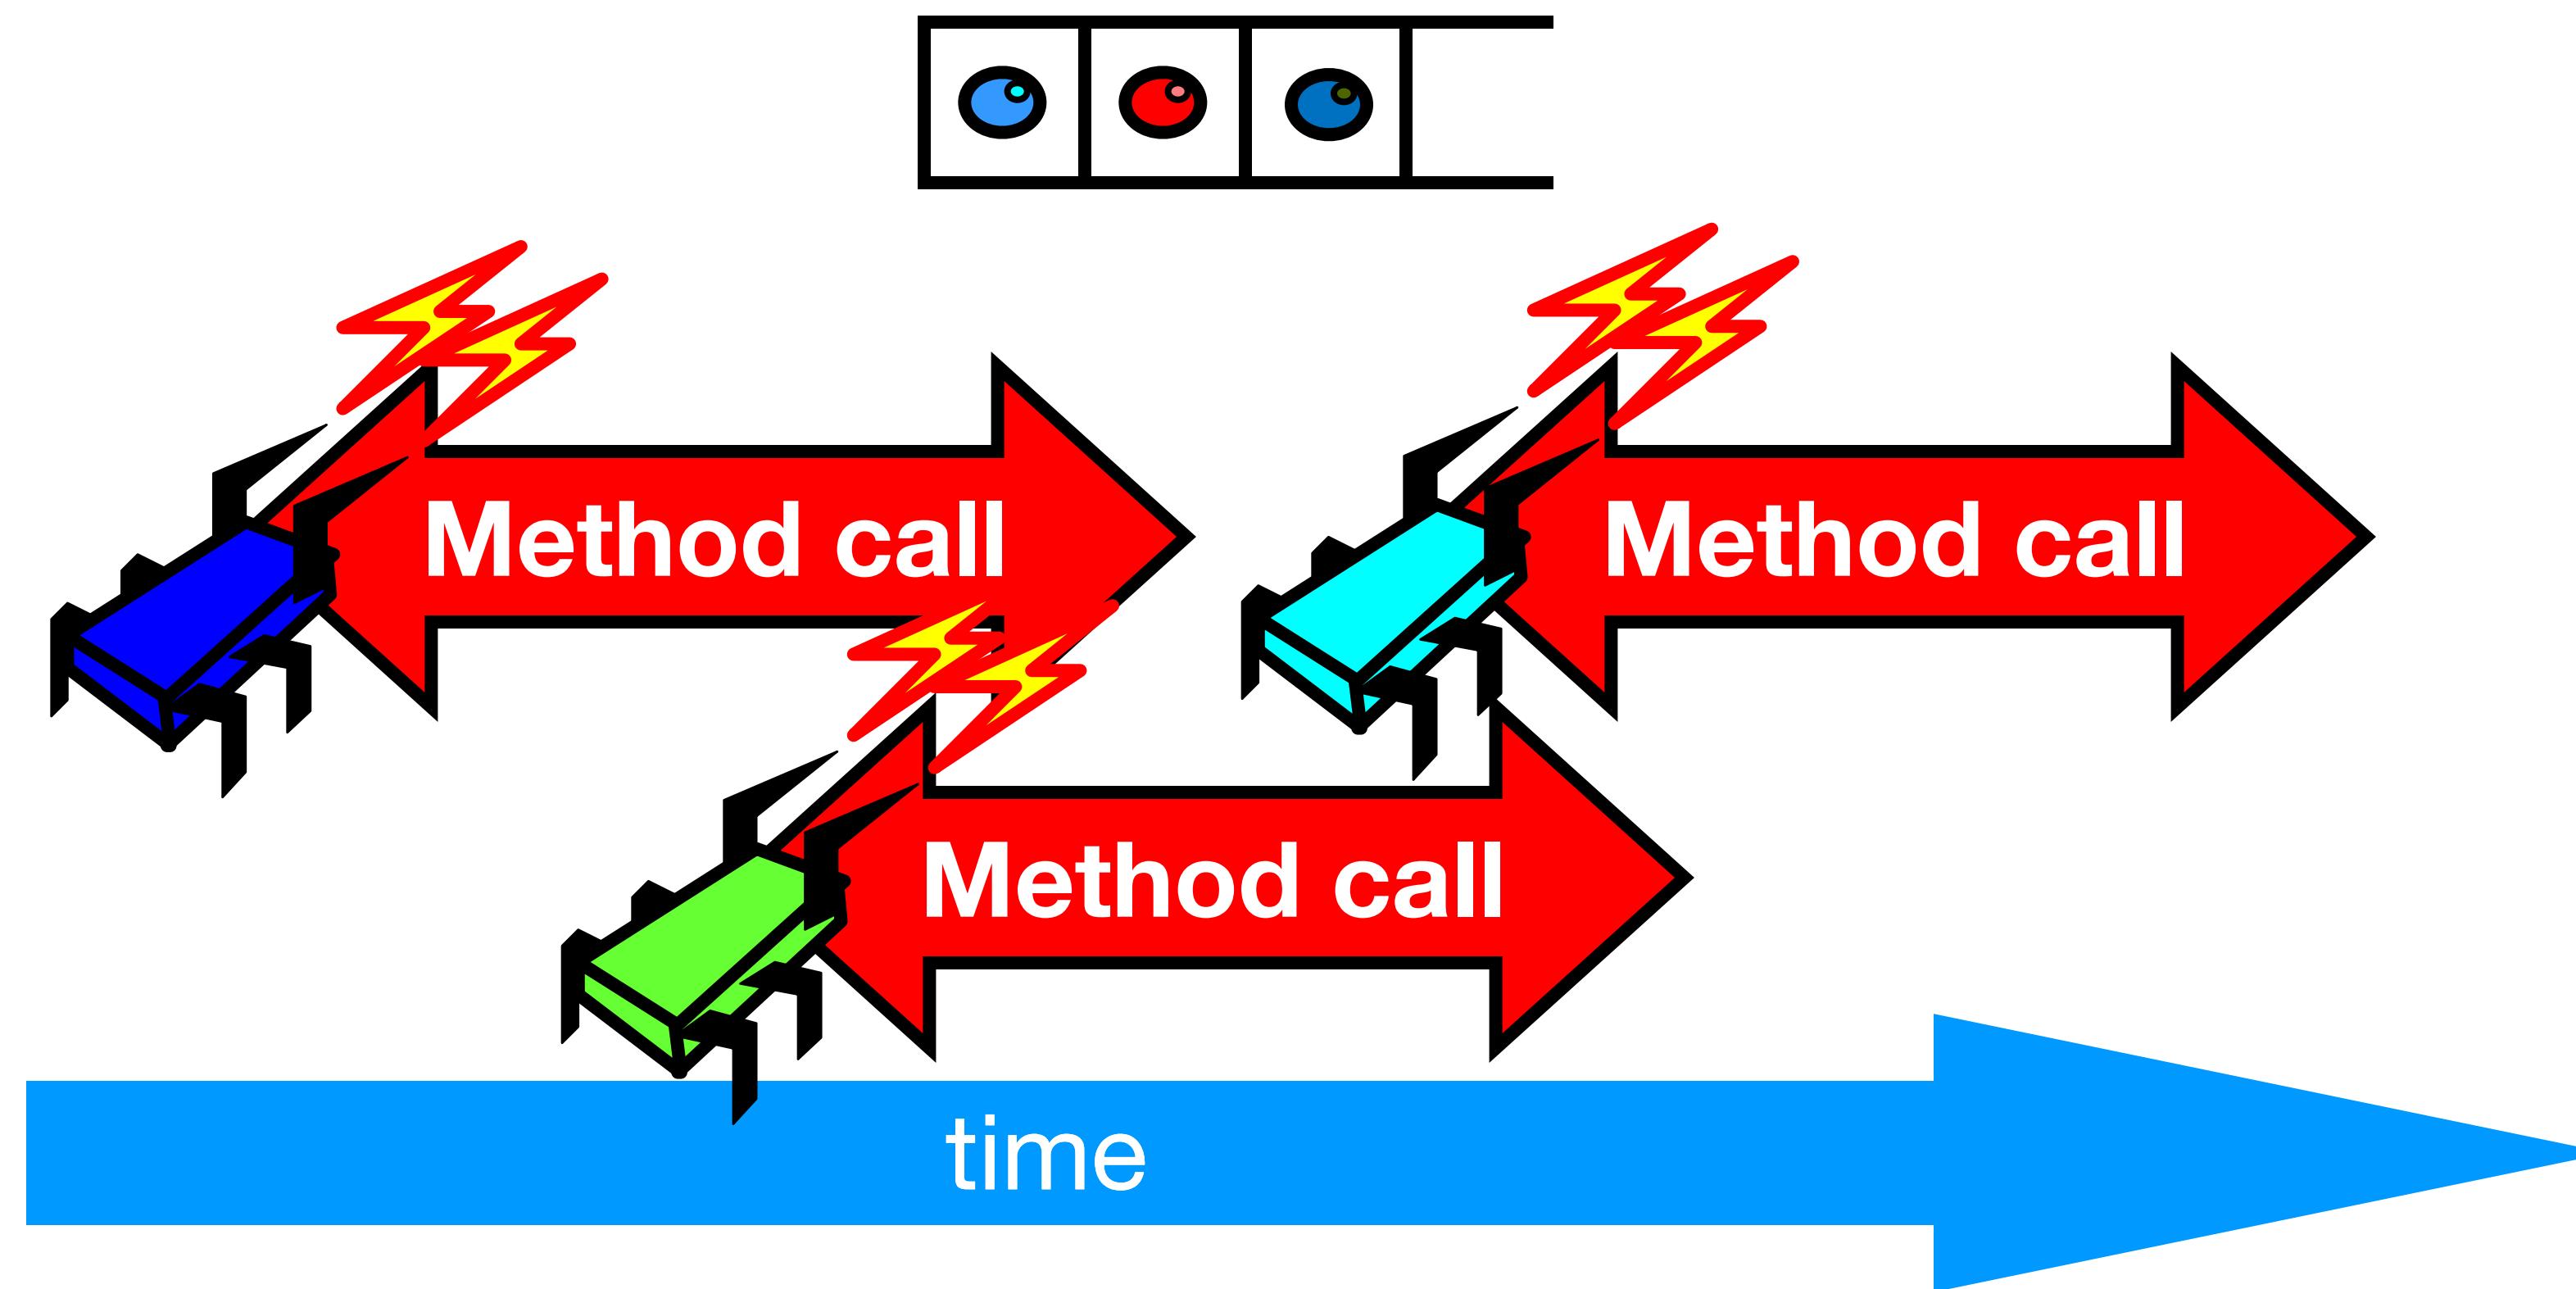
# Concurrent Methods Take Overlapping Time



# Concurrent Methods Take Overlapping Time



# Concurrent Methods Take Overlapping Time



# Sequential vs Concurrent

- Sequential
  - Object needs a meaningful state only ***between*** method calls
- Concurrent
  - Because method calls overlap, the object might ***never*** be between method calls

# Sequential vs Concurrent

- Sequential:
  - Each method described in isolation
- Concurrent
  - Must characterize ***all*** possible interactions with concurrent calls
    - What if two **enq()** calls overlap?
    - Two **deq()** calls? **enq()** and **deq()**? ...

# Sequential vs Concurrent

- Sequential:
  - Can add new methods without affecting older methods
- Concurrent:
  - Everything can potentially interact with everything else

# Sequential vs Concurrent

- Sequential:
  - Can add new methods without affecting older methods
- Concurrent:
  - Everything can potentially interact with everything else



# The Big Question

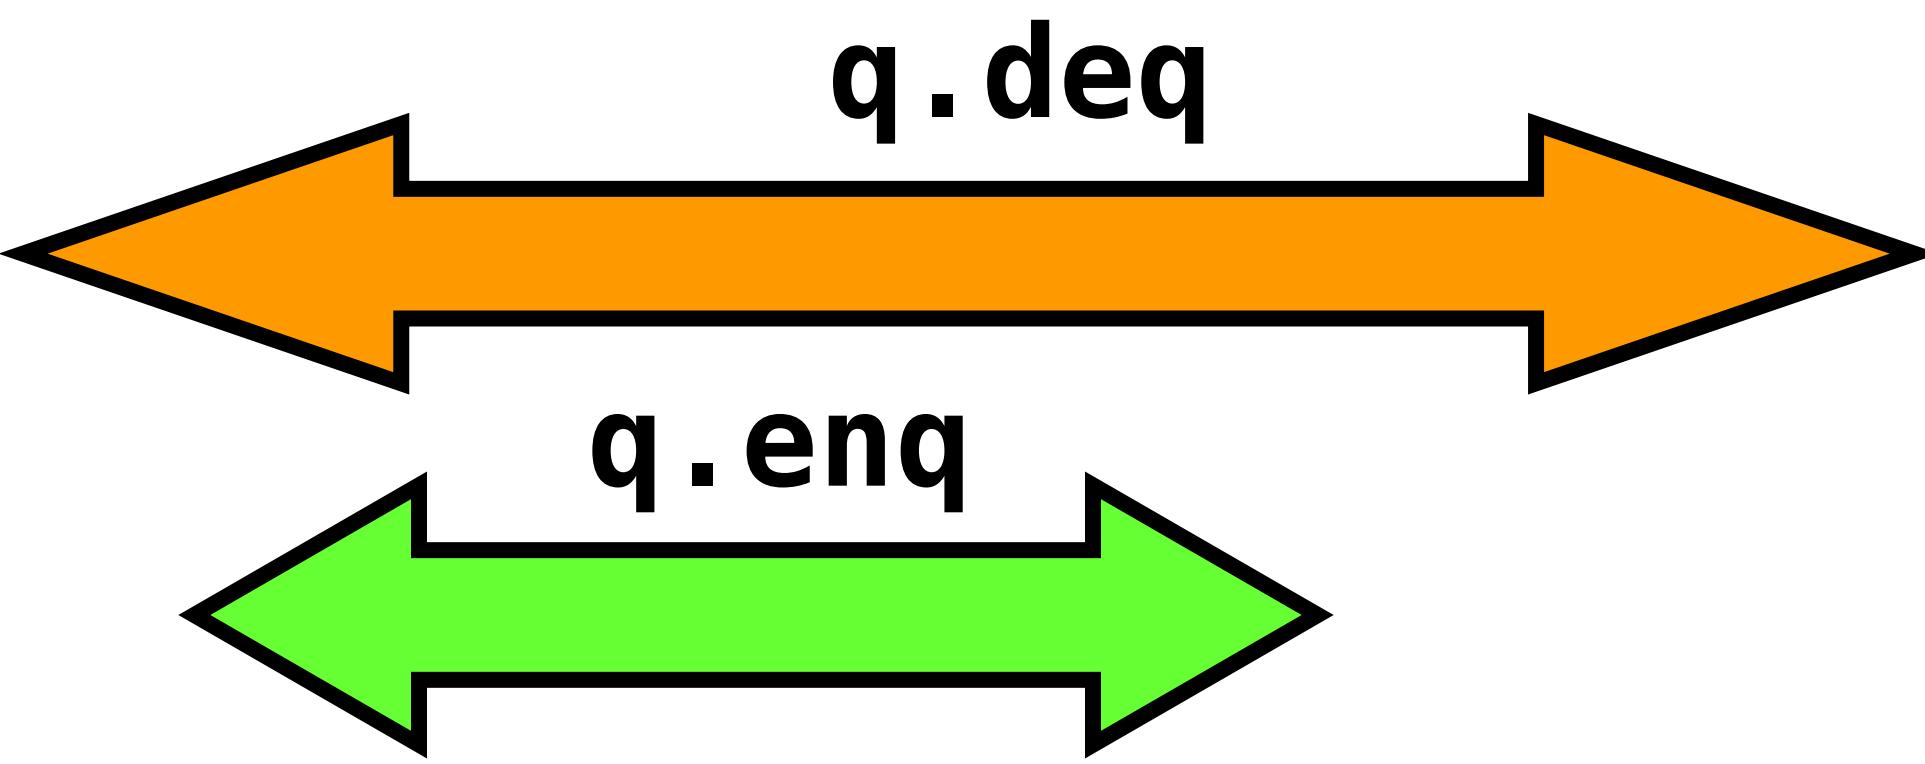
- What does it *mean* for a ***concurrent*** object to be correct?
  - What *is* a concurrent FIFO queue?
  - FIFO means ***strict temporal order***
  - Concurrent means ***ambiguous temporal order***

# Intuitively


```
let deq q =
  Mutex.lock q.lock;
  try
    if q.tail = q.head then
      raise Empty;
    match q.items.(q.head mod q.capacity) with
    | None -> assert false
    | Some x ->
        q.head <- q.head + 1;
        Mutex.unlock q.lock;
        x
  with e ->
    Mutex.unlock q.lock;
    raise e
```

*All queue modifications are mutually exclusive*


# Intuitively




# Intuitively




# Intuitively



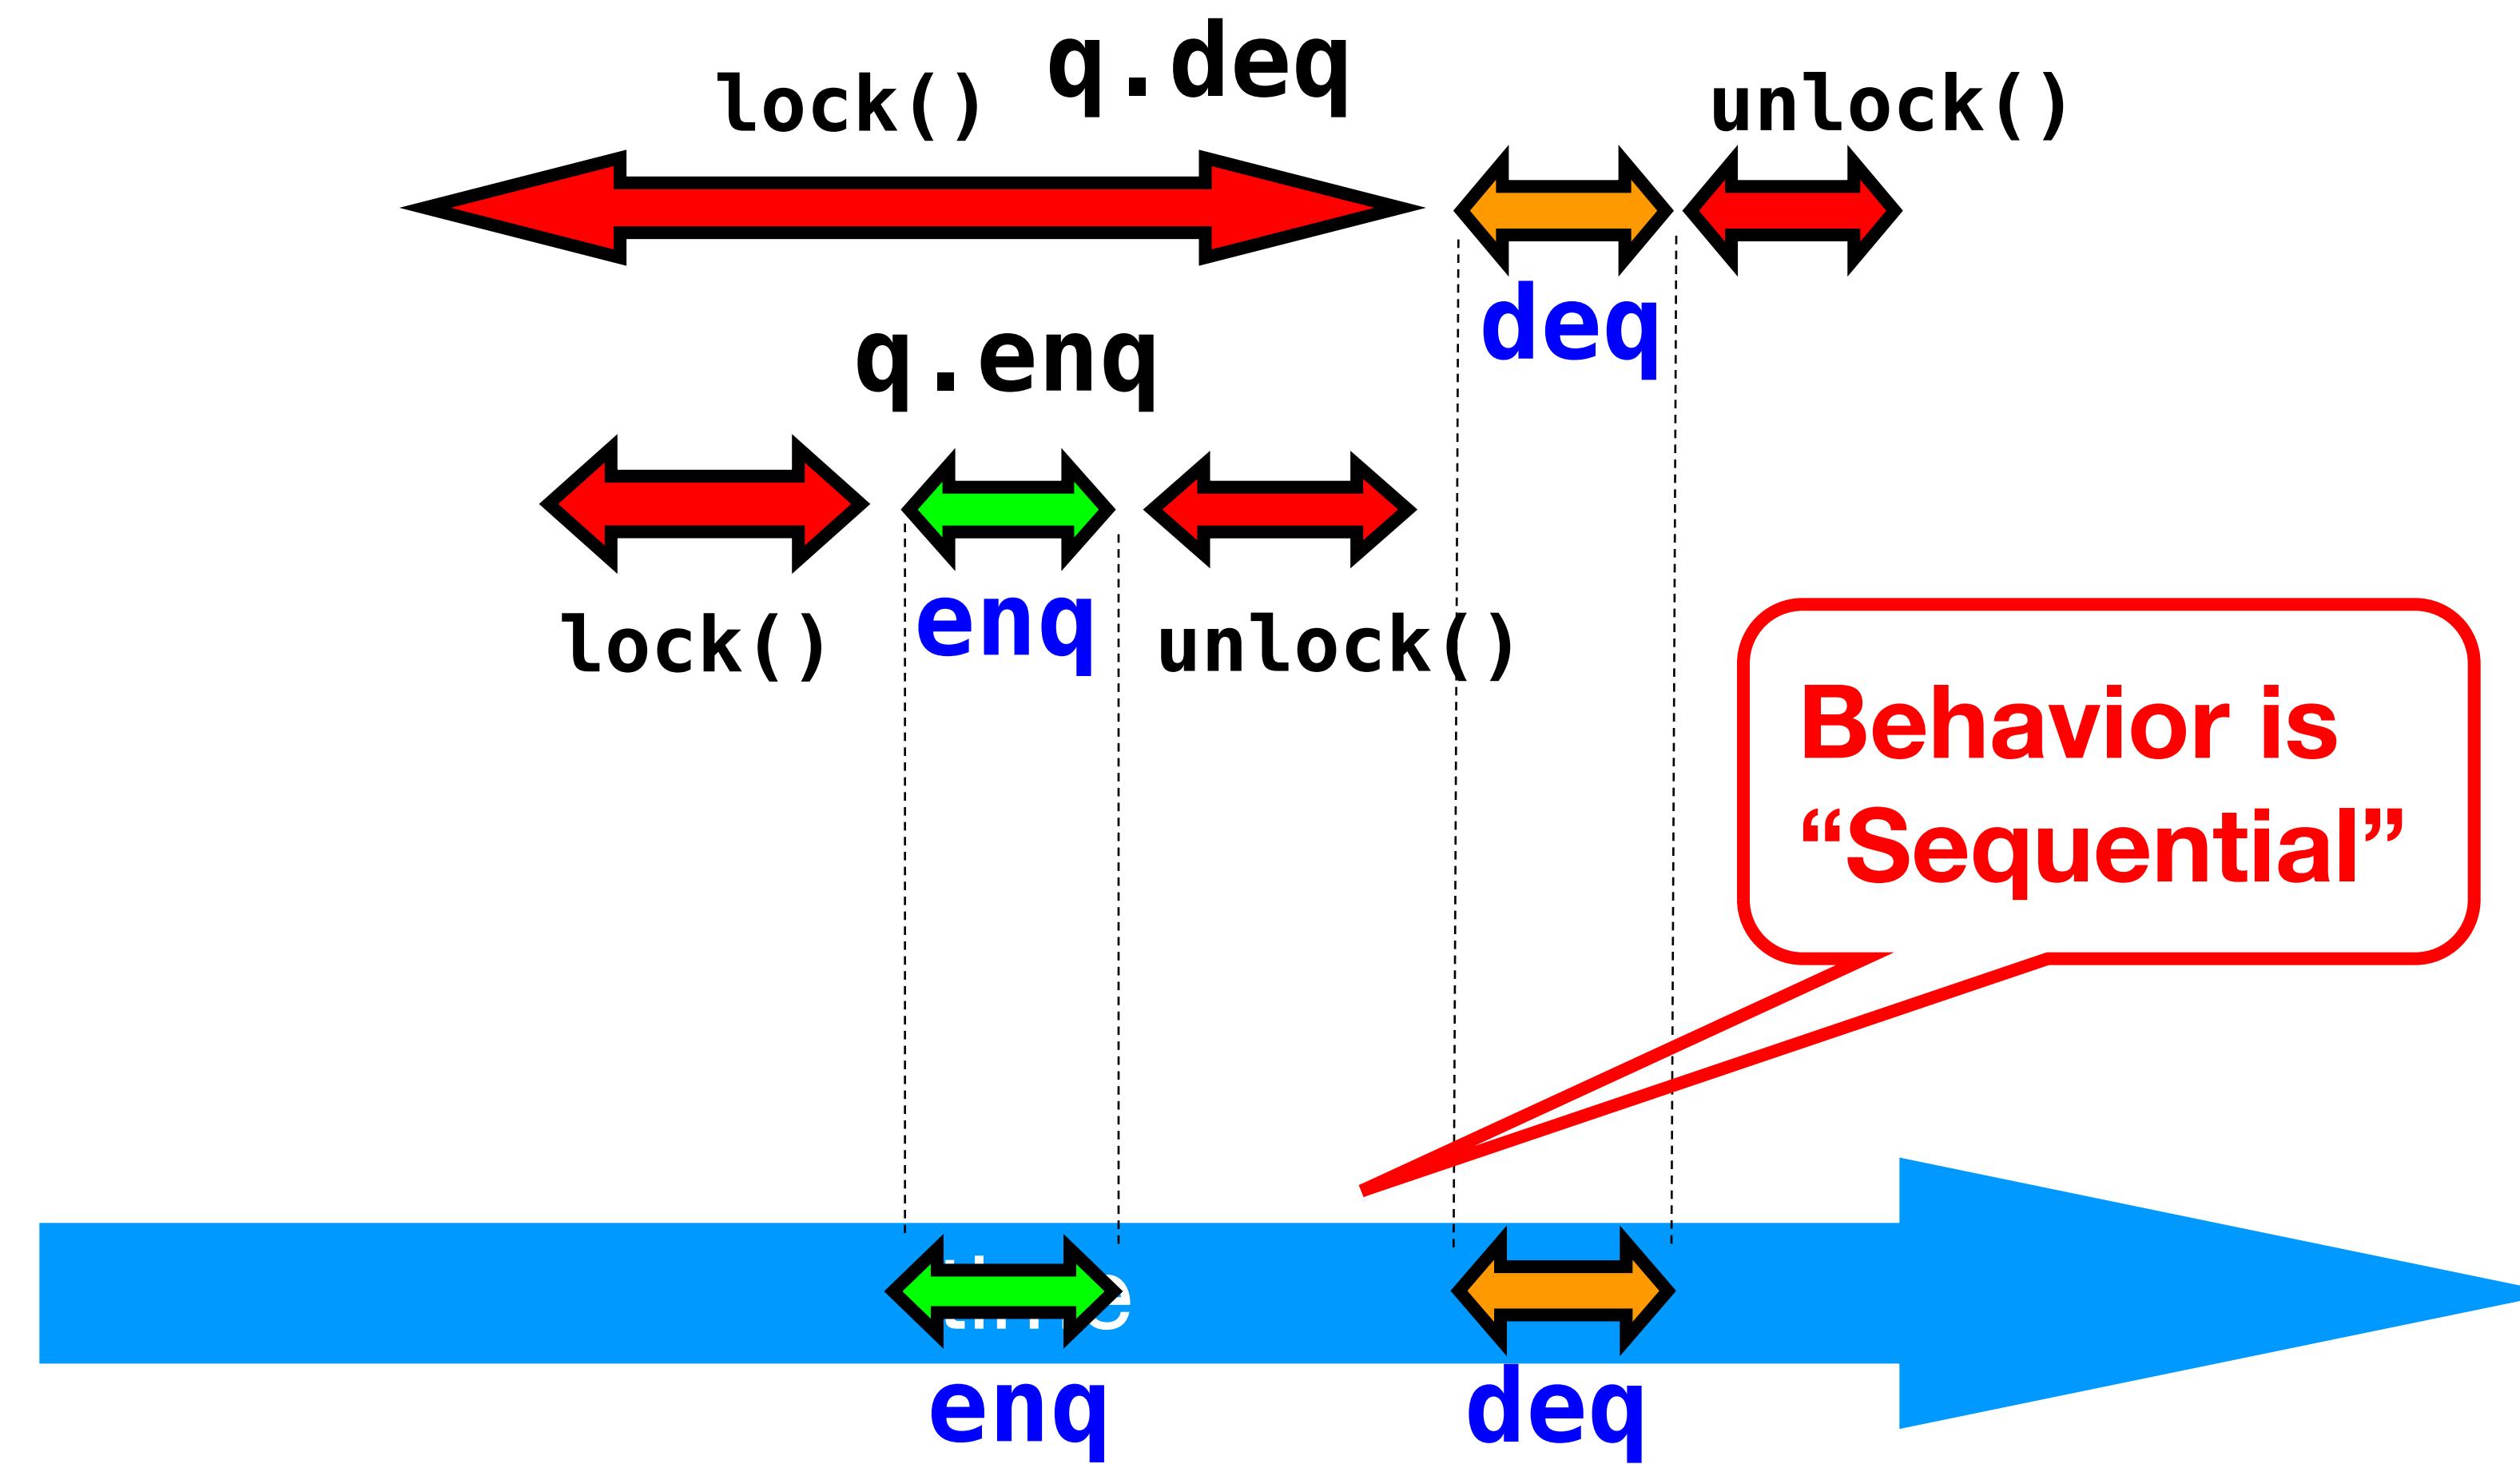
# Intuitively



# Intuitively




# Intuitively



# Intuitively

Lets capture the idea of describing  
the concurrent via the sequential



# Linearizability

# Linearizability

- Each method should
  - “take effect”
  - Instantaneously
  - Between *invocation* and *response* events

# Linearizability

- Each method should
  - “take effect”
  - Instantaneously
  - Between *invocation* and *response* events
- Object is correct if this “sequentialised” behaviour is correct

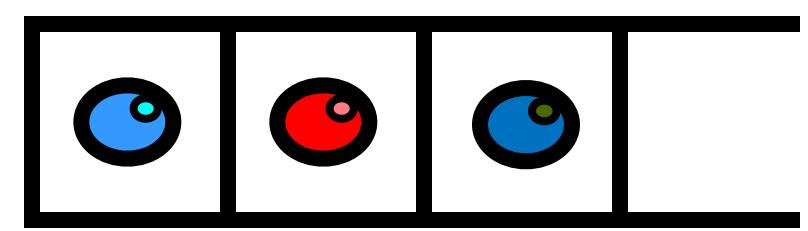
# Linearizability

- Each method should
  - “take effect”
  - Instantaneously
  - Between ***invocation*** and ***response*** events
- Object is correct if this “sequentialised” behaviour is correct
- Any such concurrent object is
  - **Linearizable™**


# Is it really about the object?

- Each method should
  - “take effect”
  - Instantaneously
  - Between *invocation* and *response* events
- Sounds like a property of *an execution...*
- A linearizable *object*
  - One of whose all possible executions are linearizable

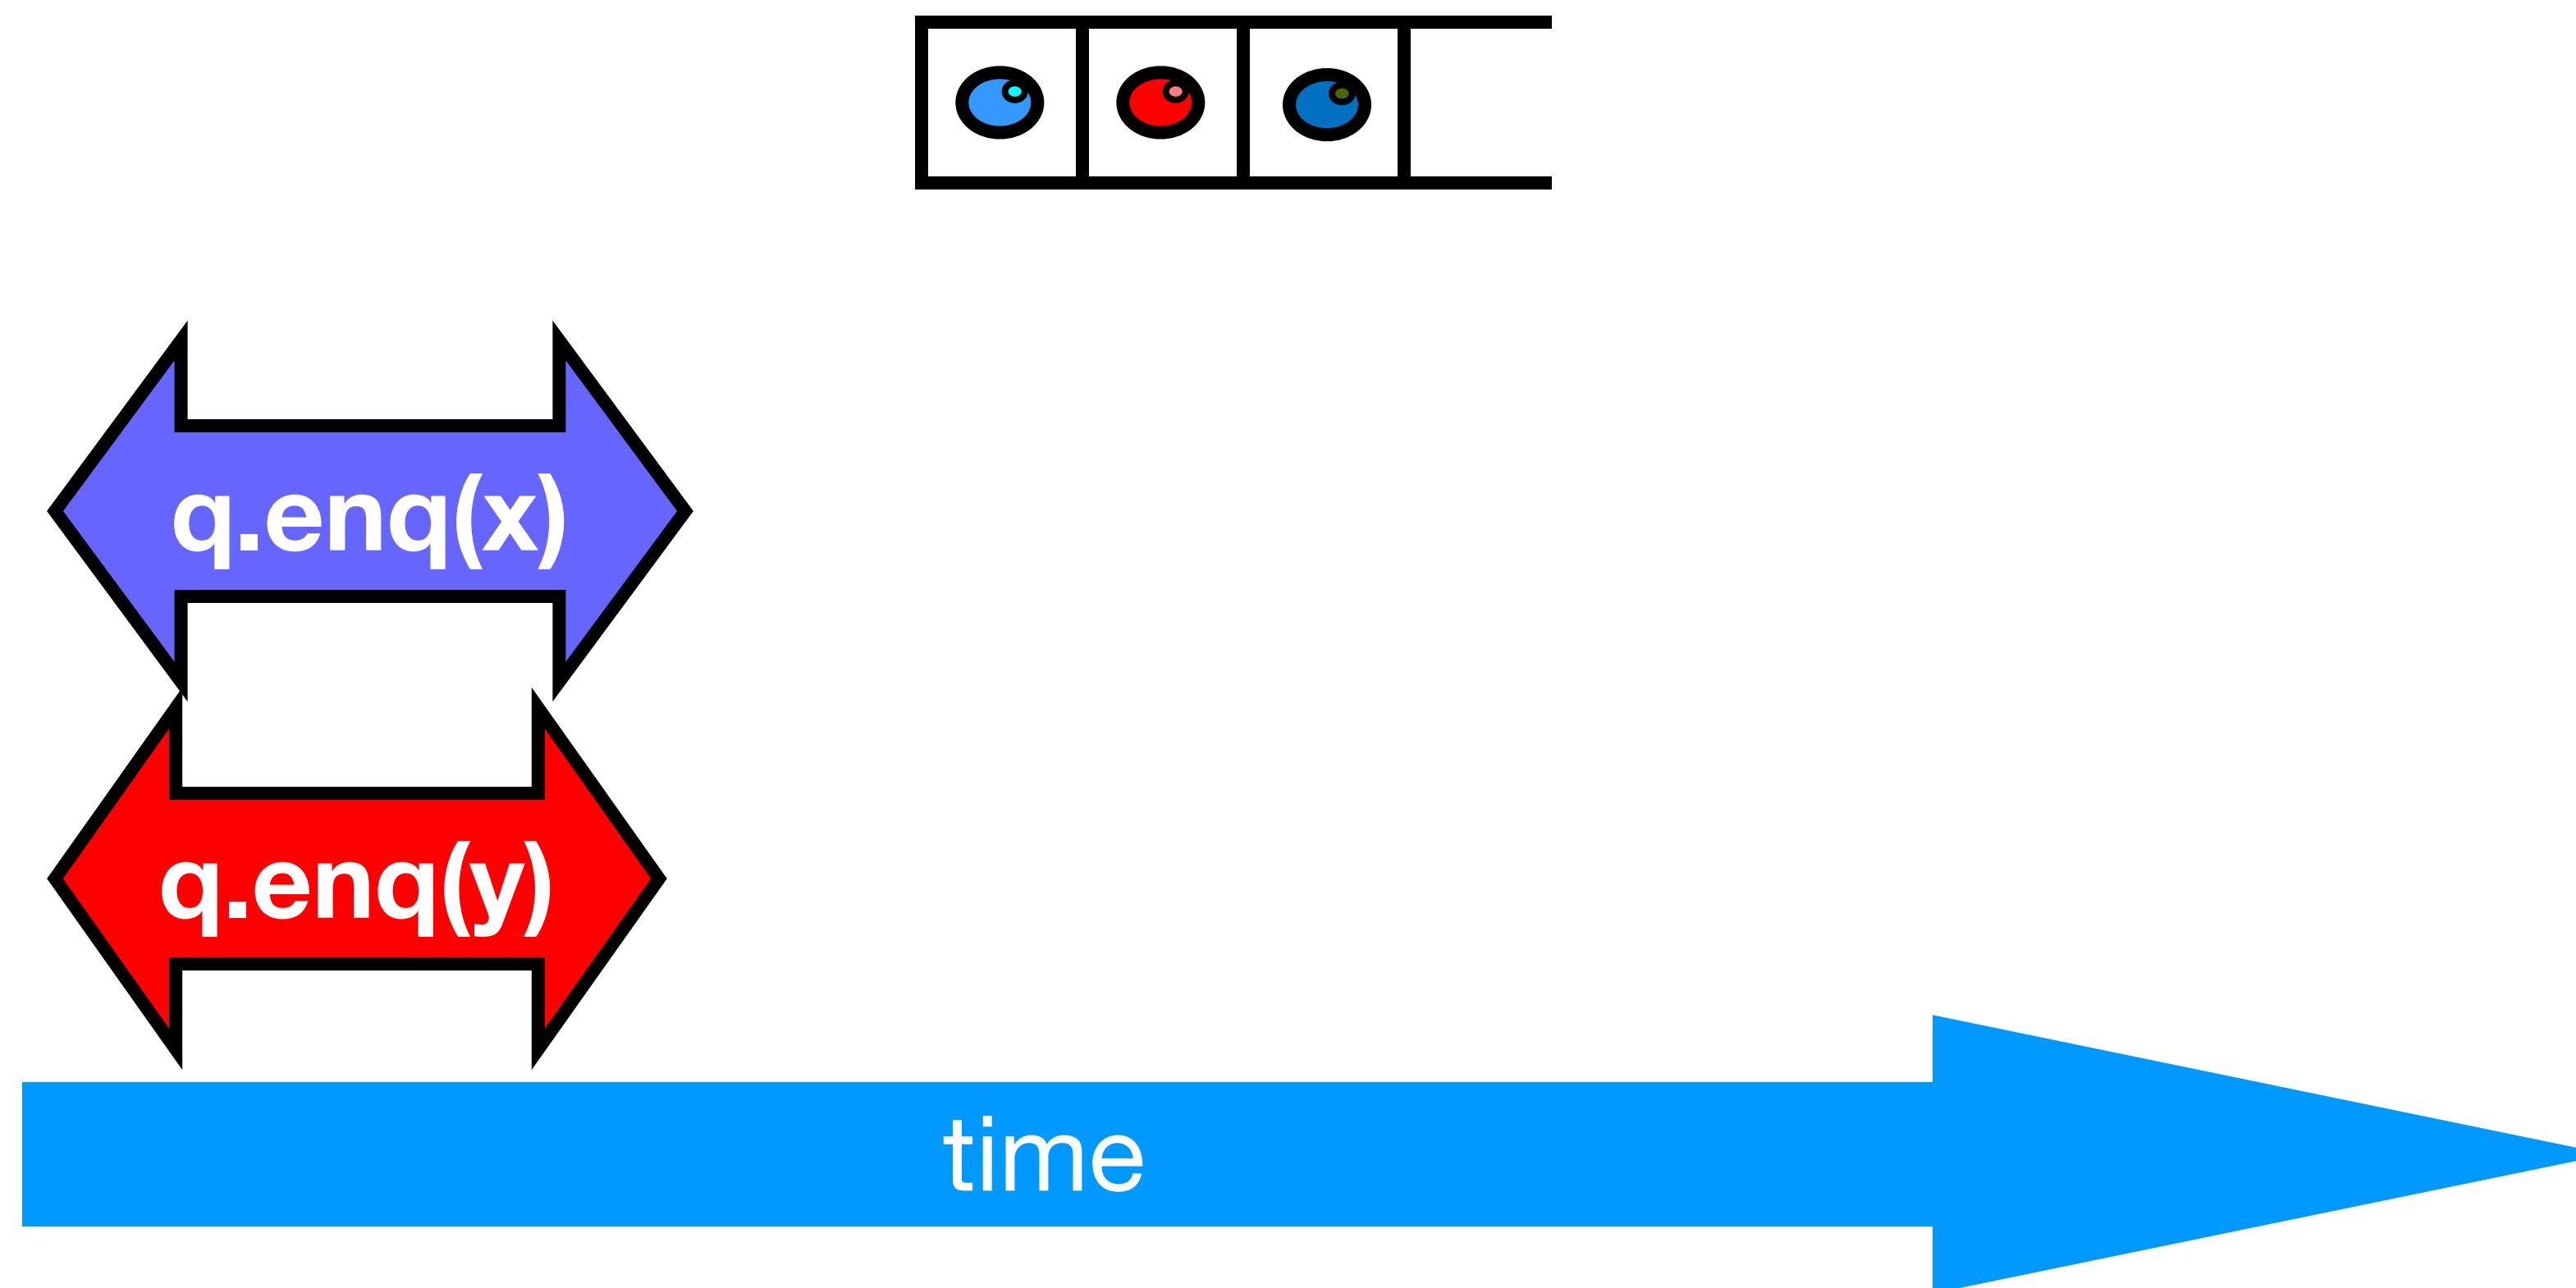
# Example




# Example



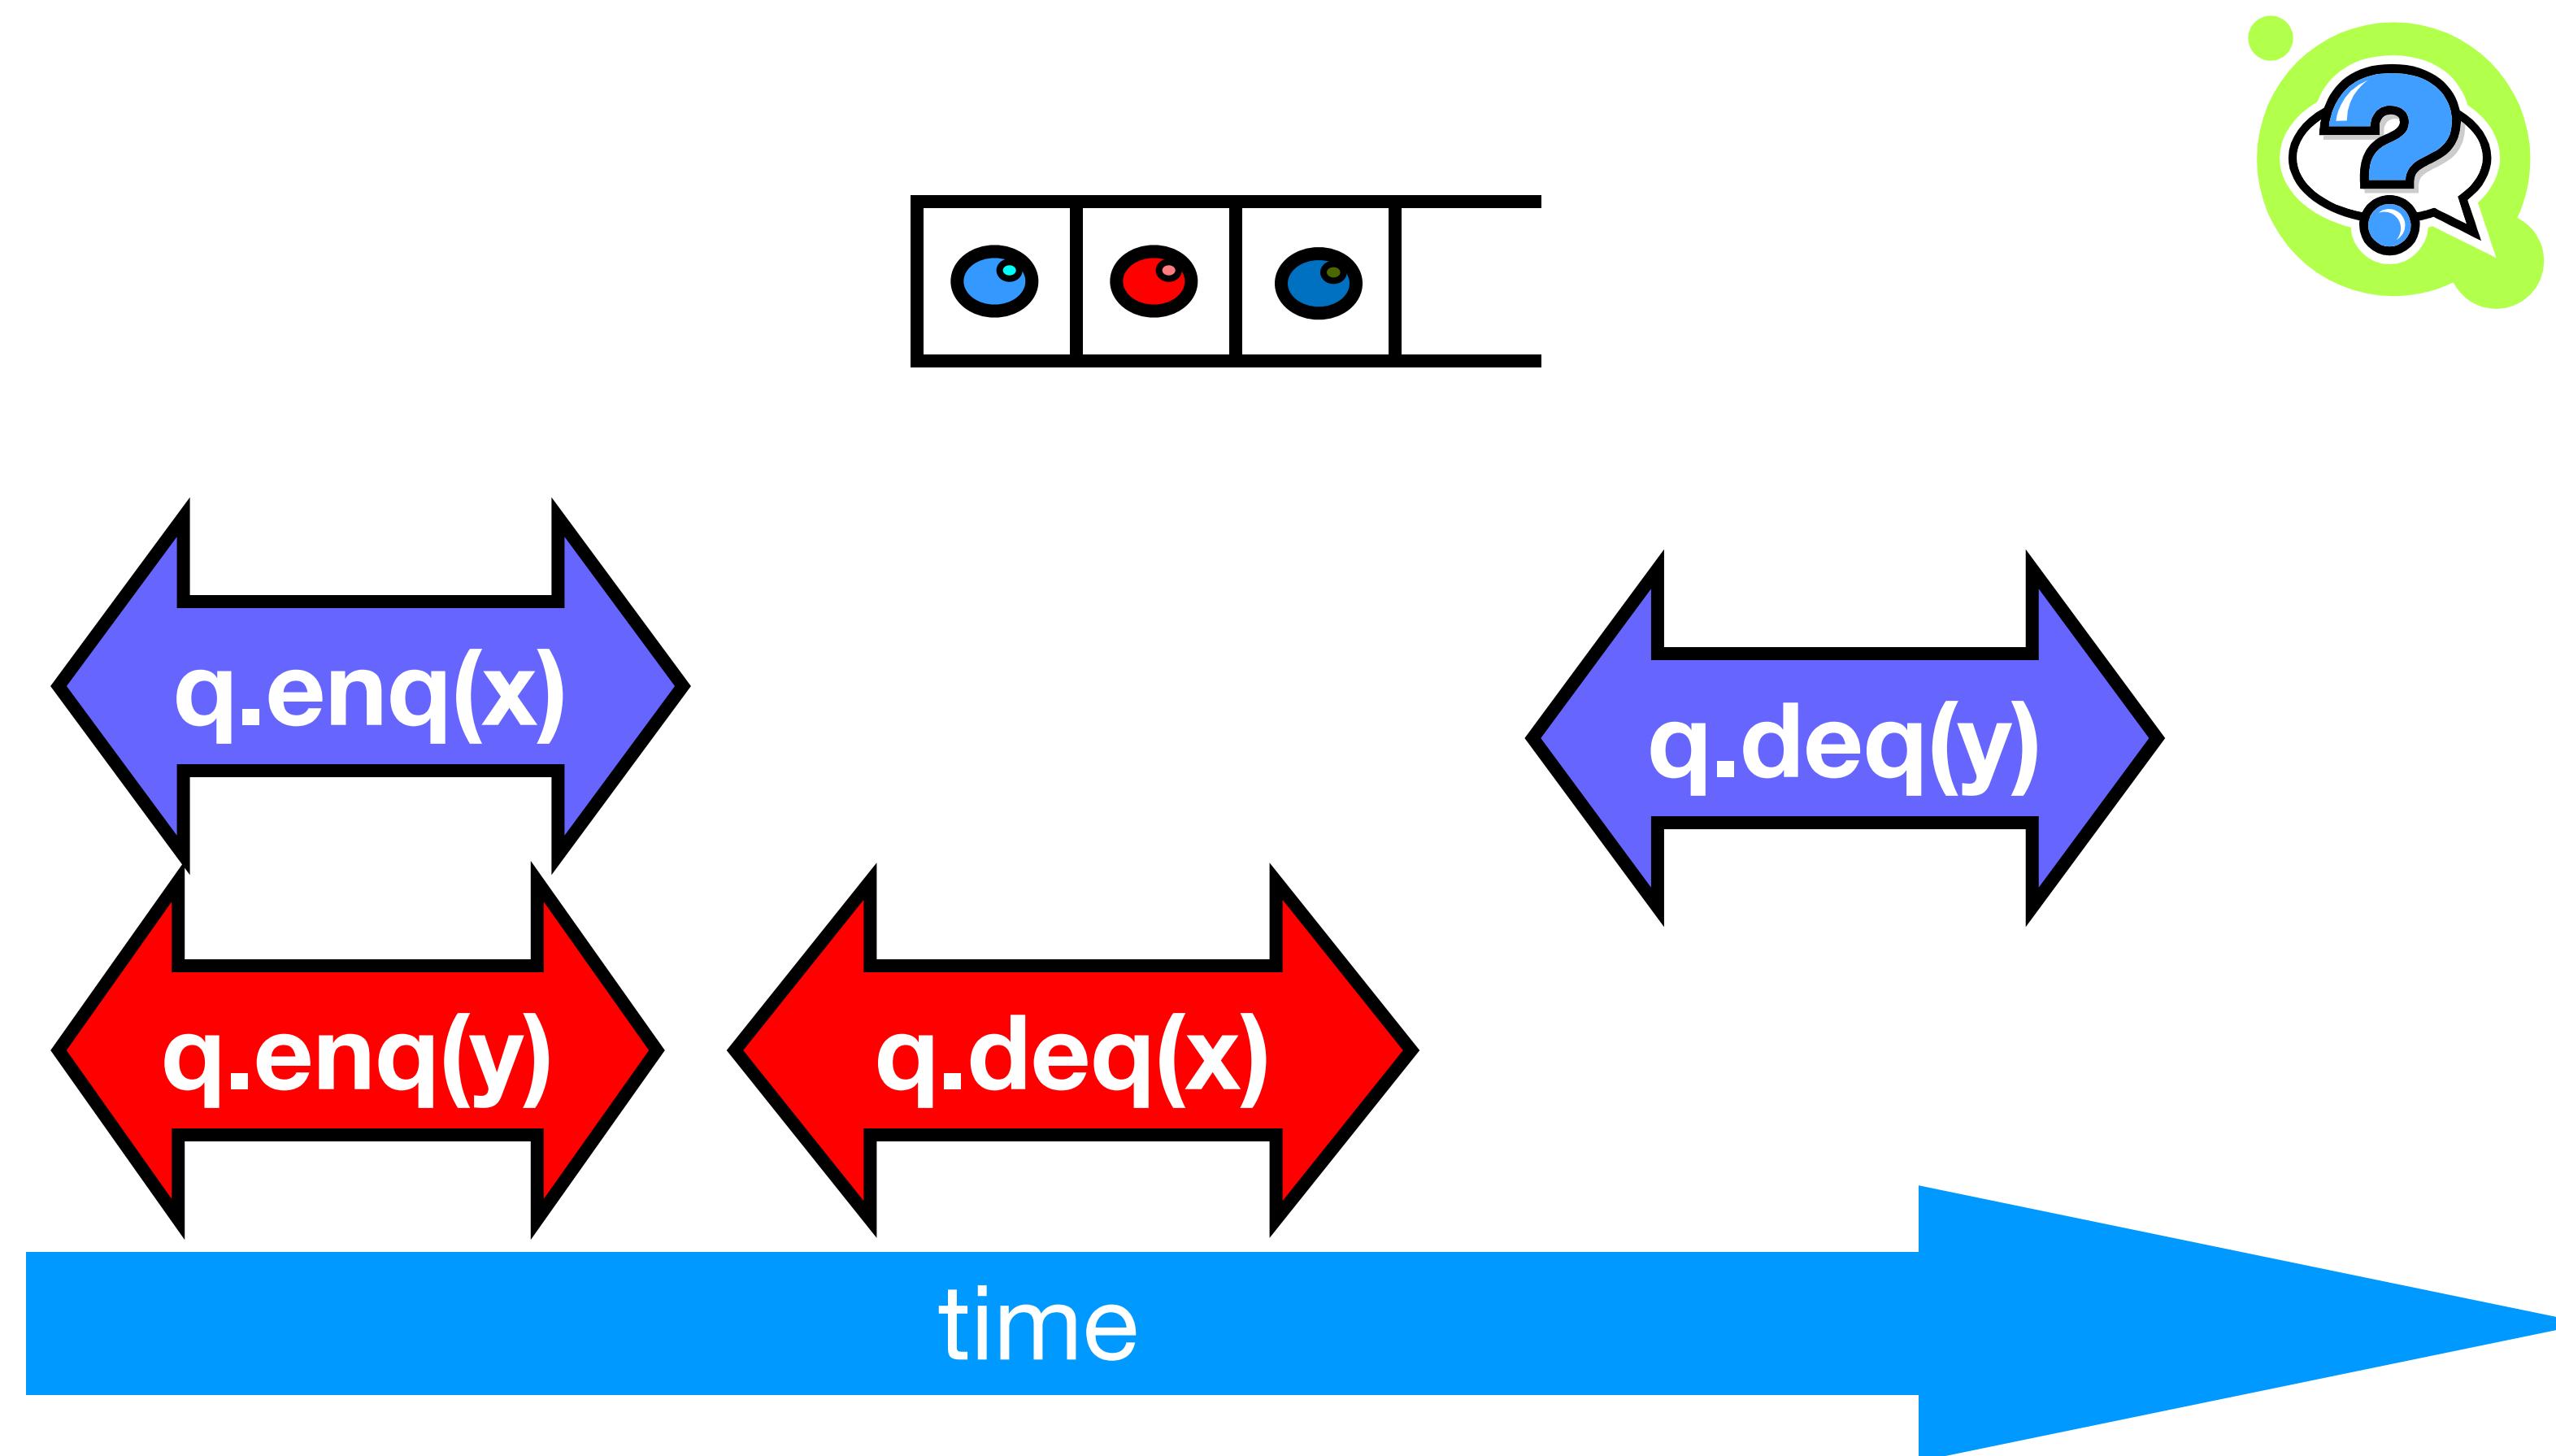
# Example



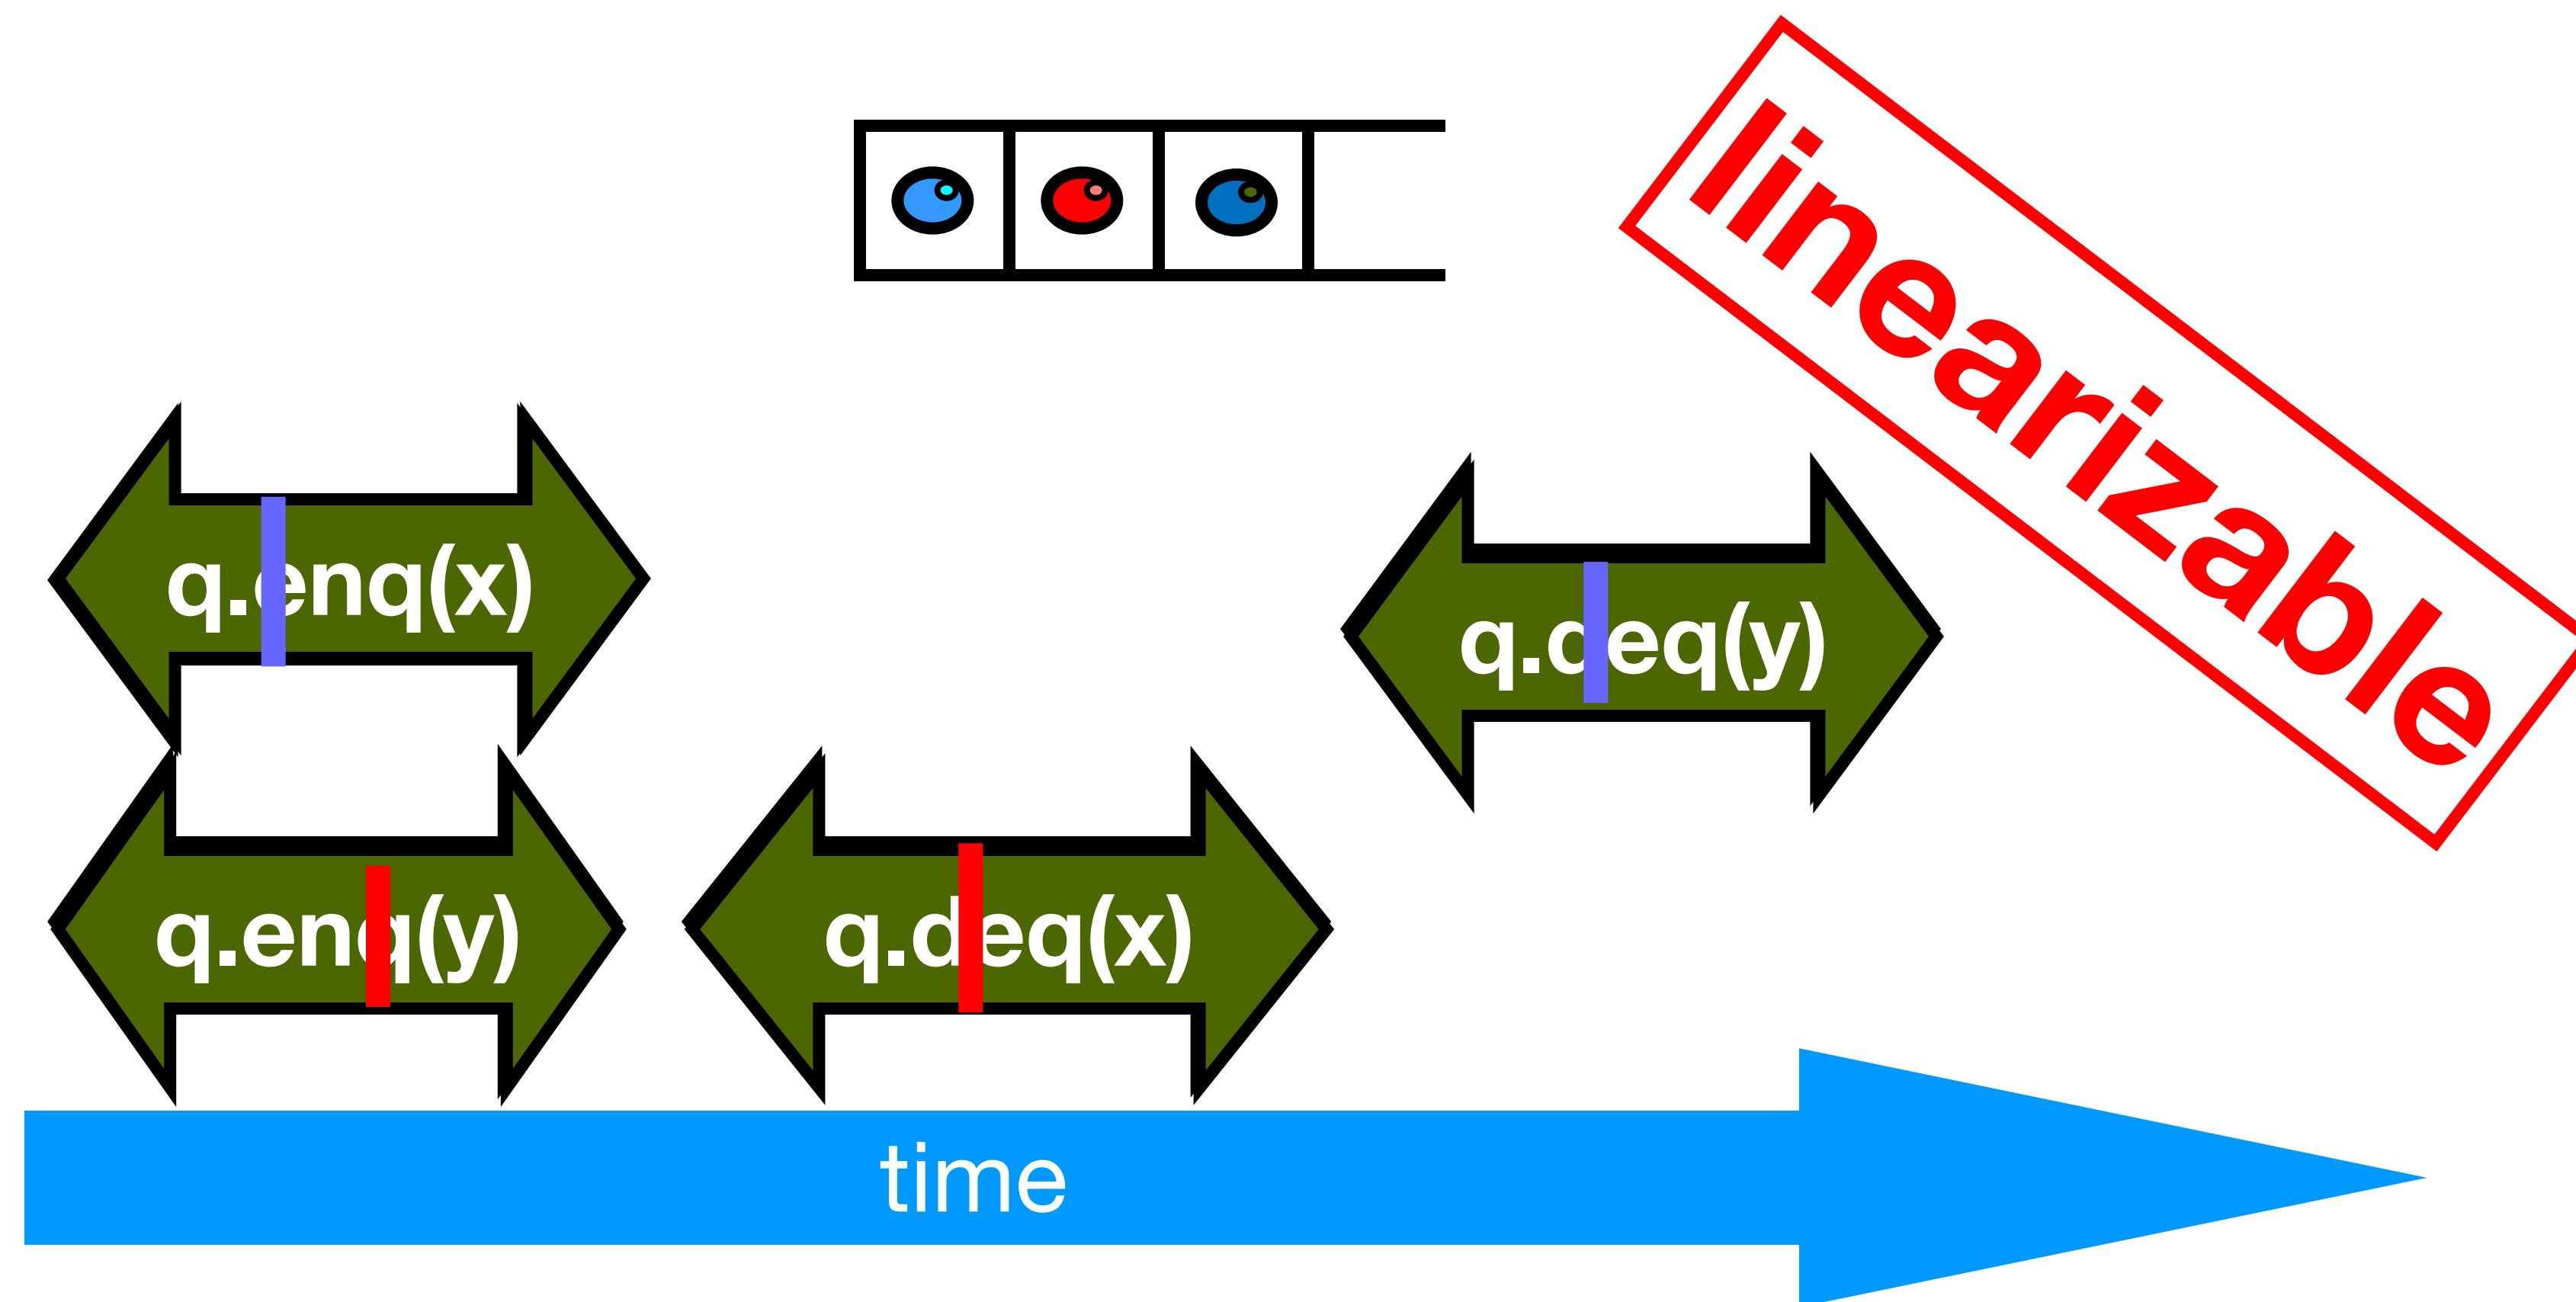
# Example



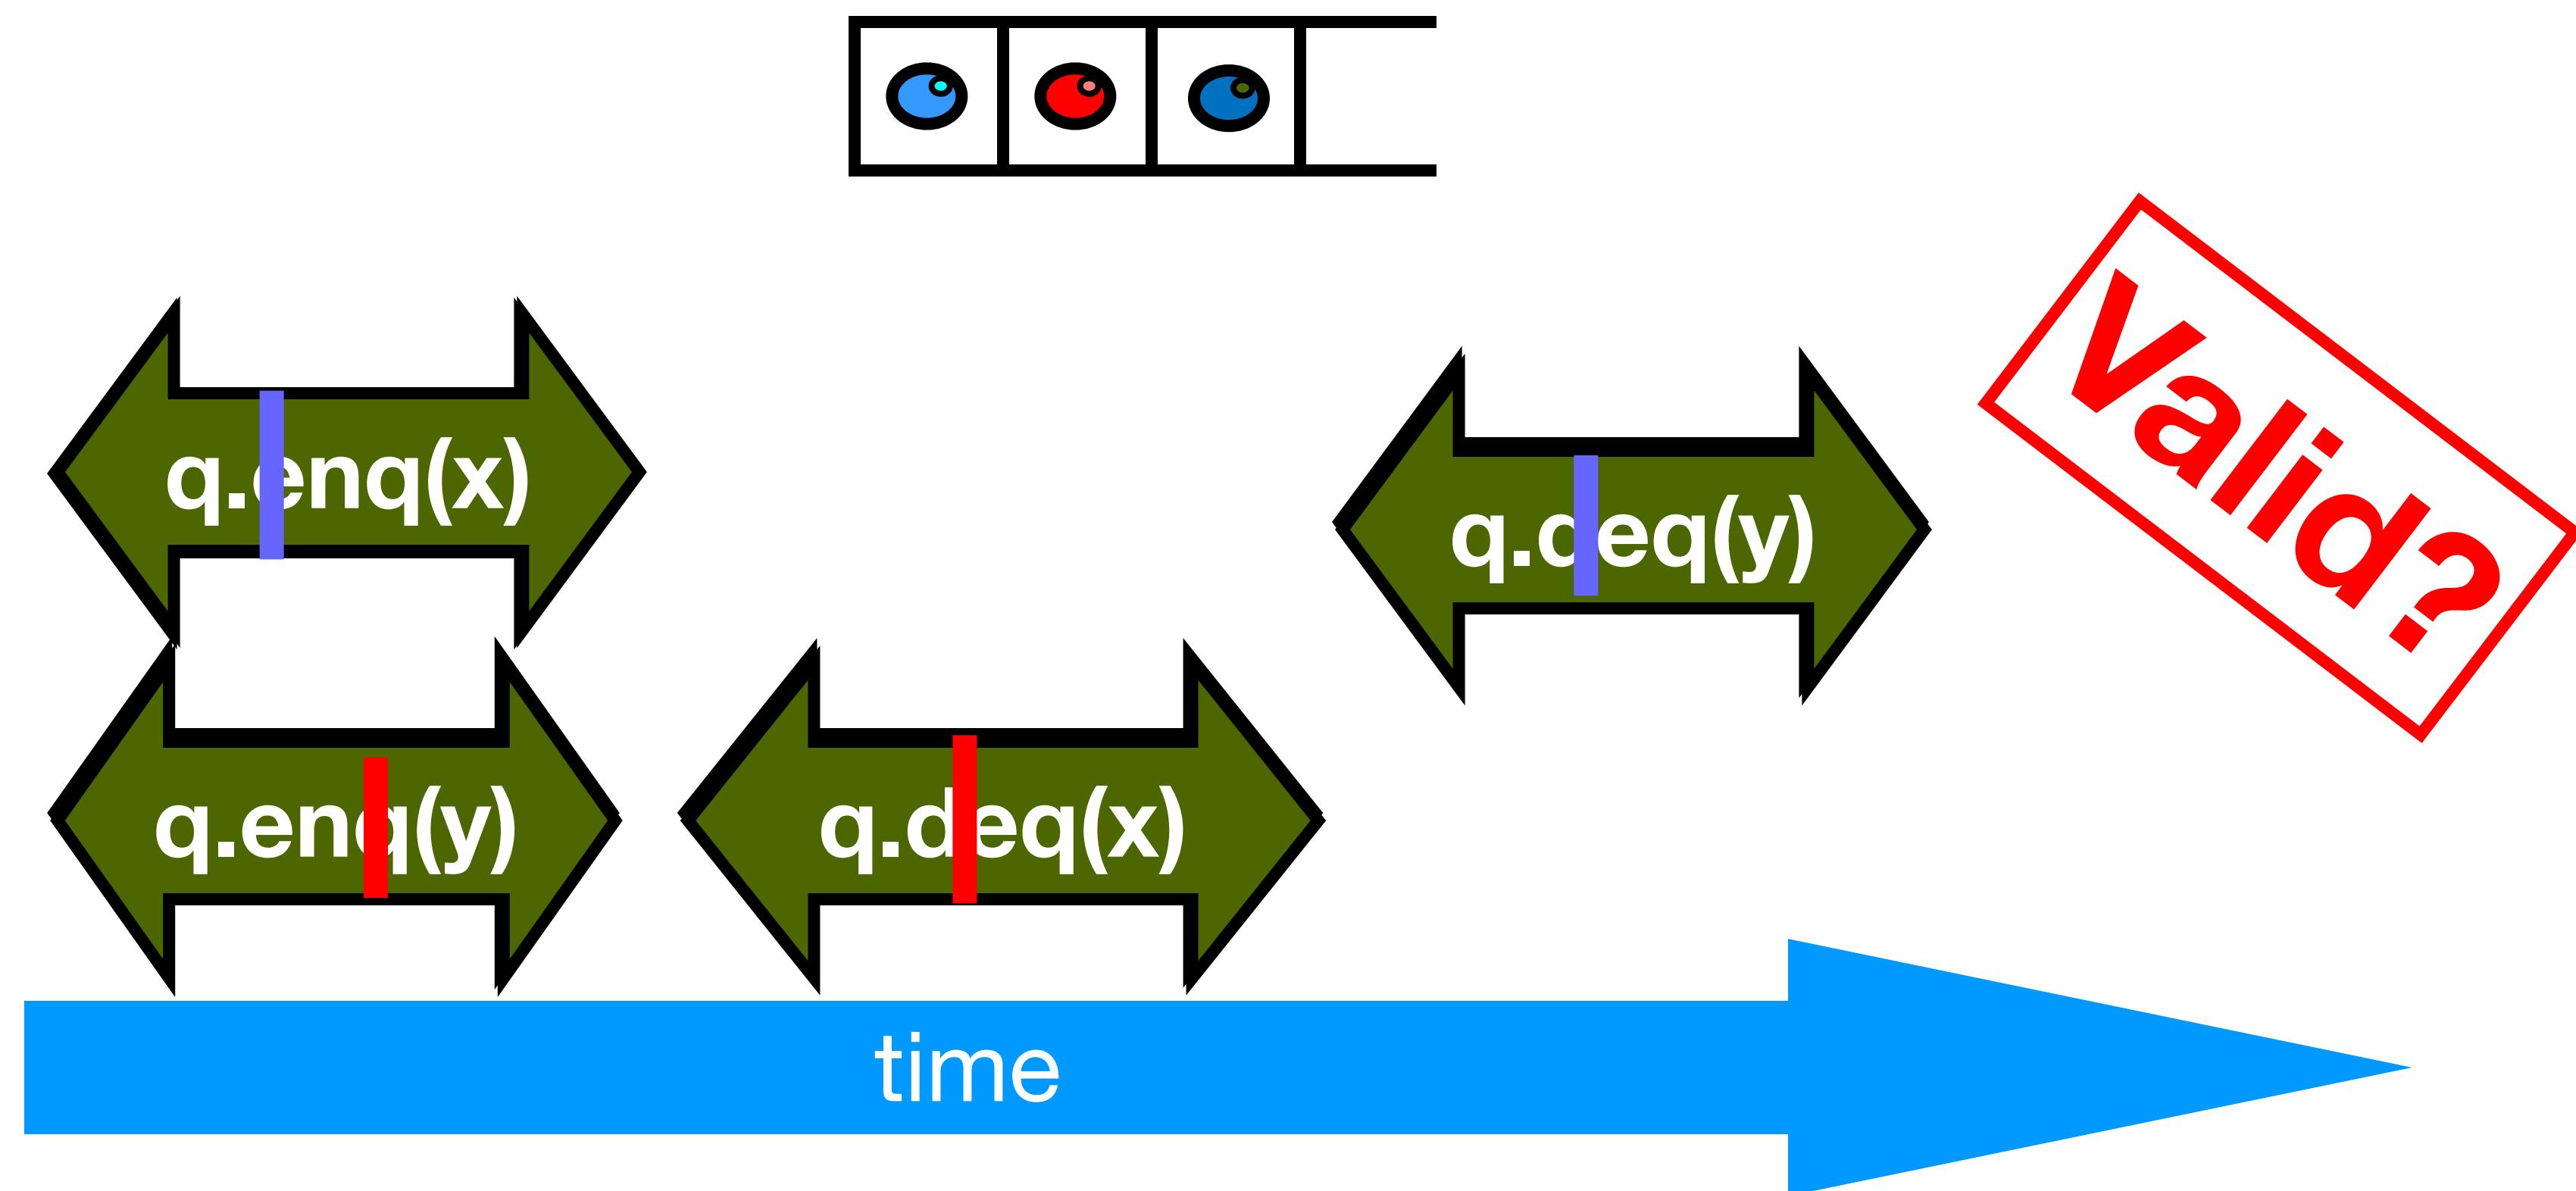
# Example




# Example



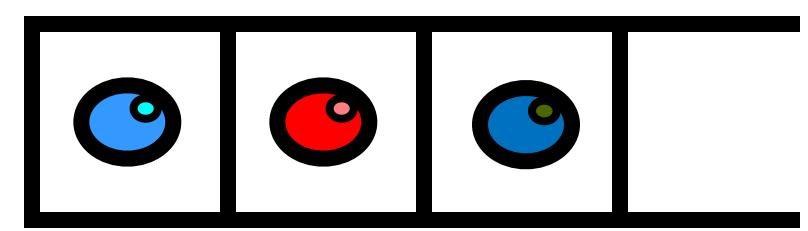
# Example



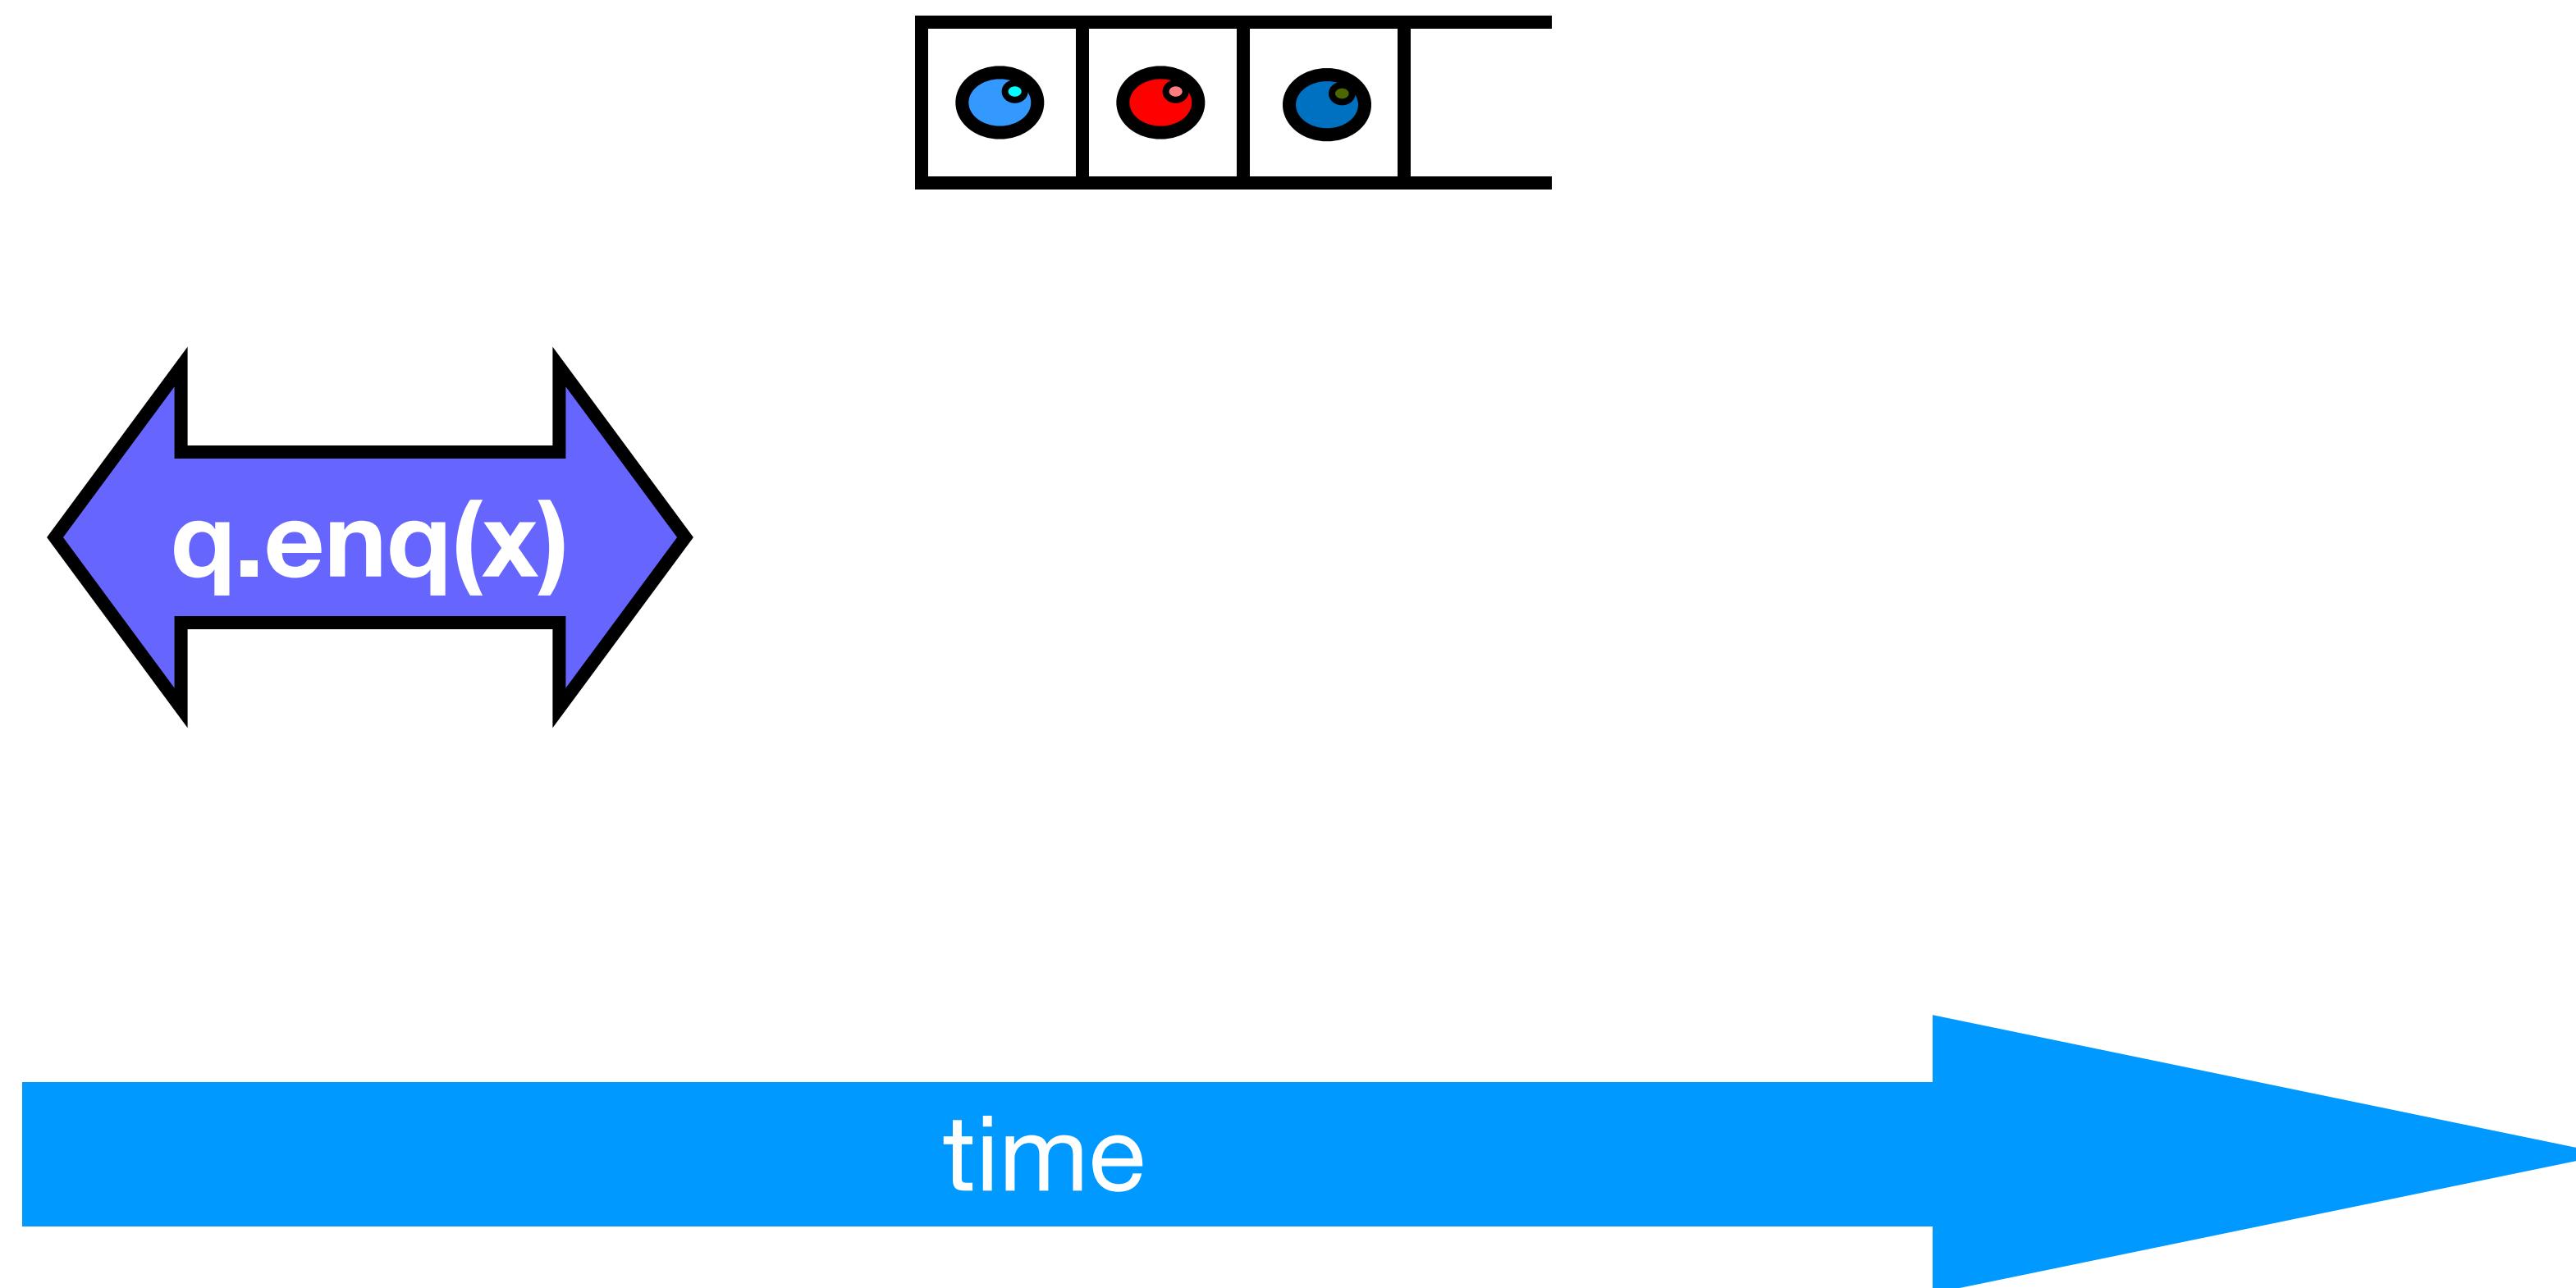
# Example



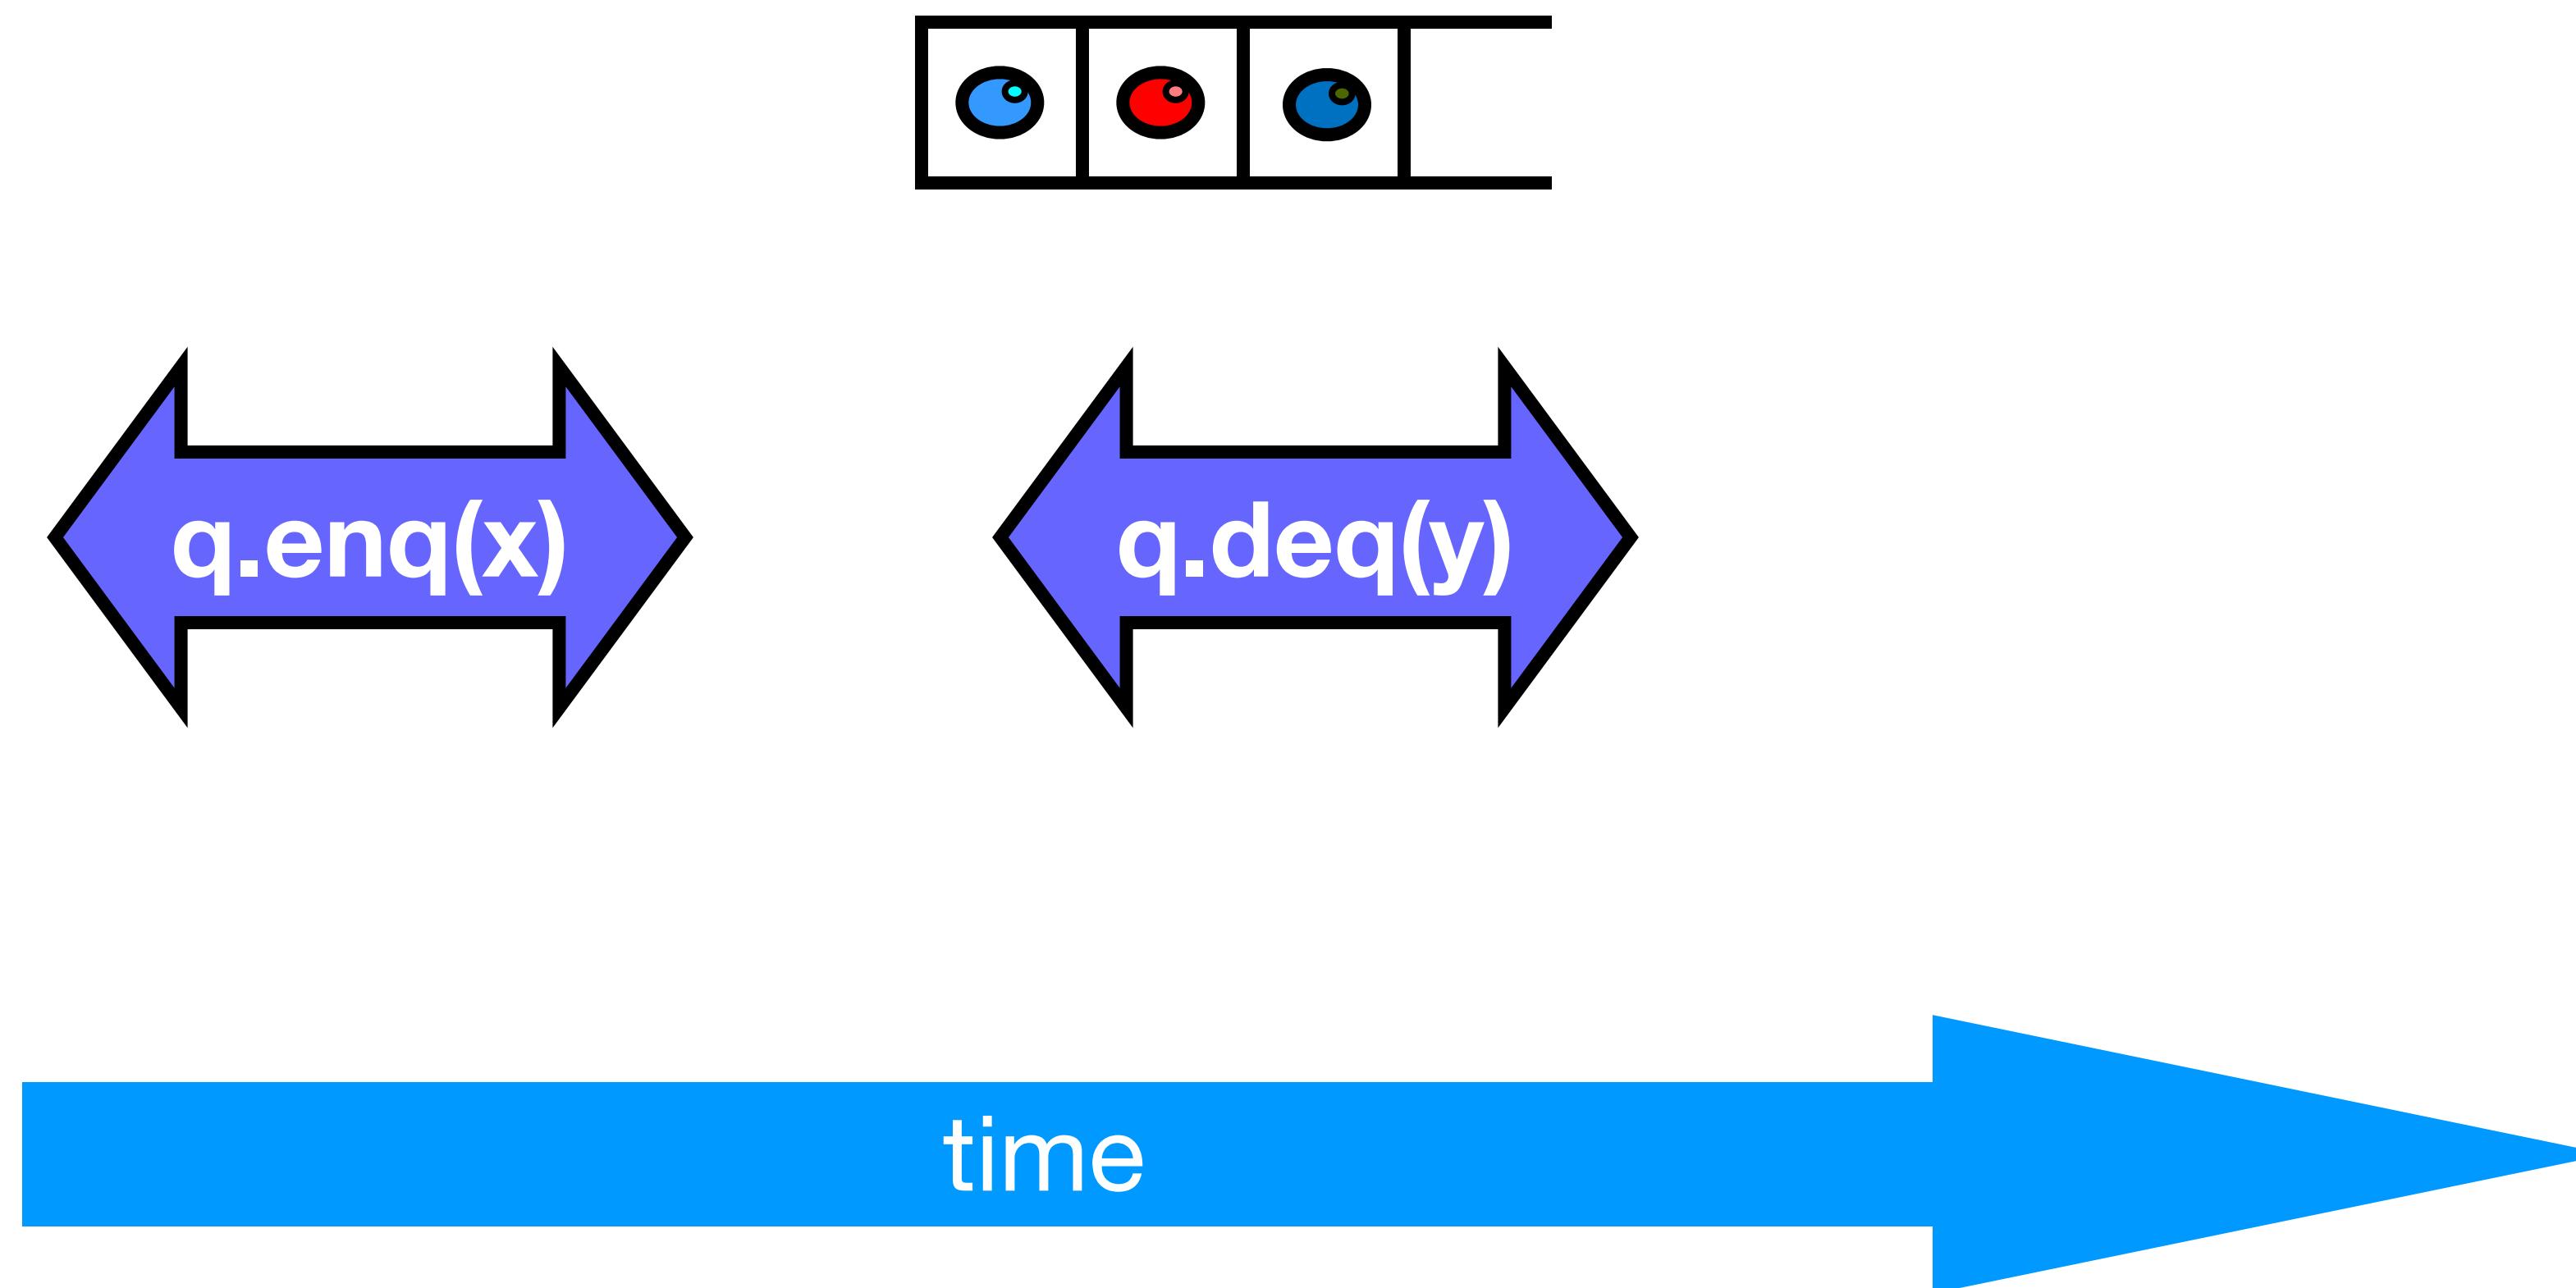
# Example



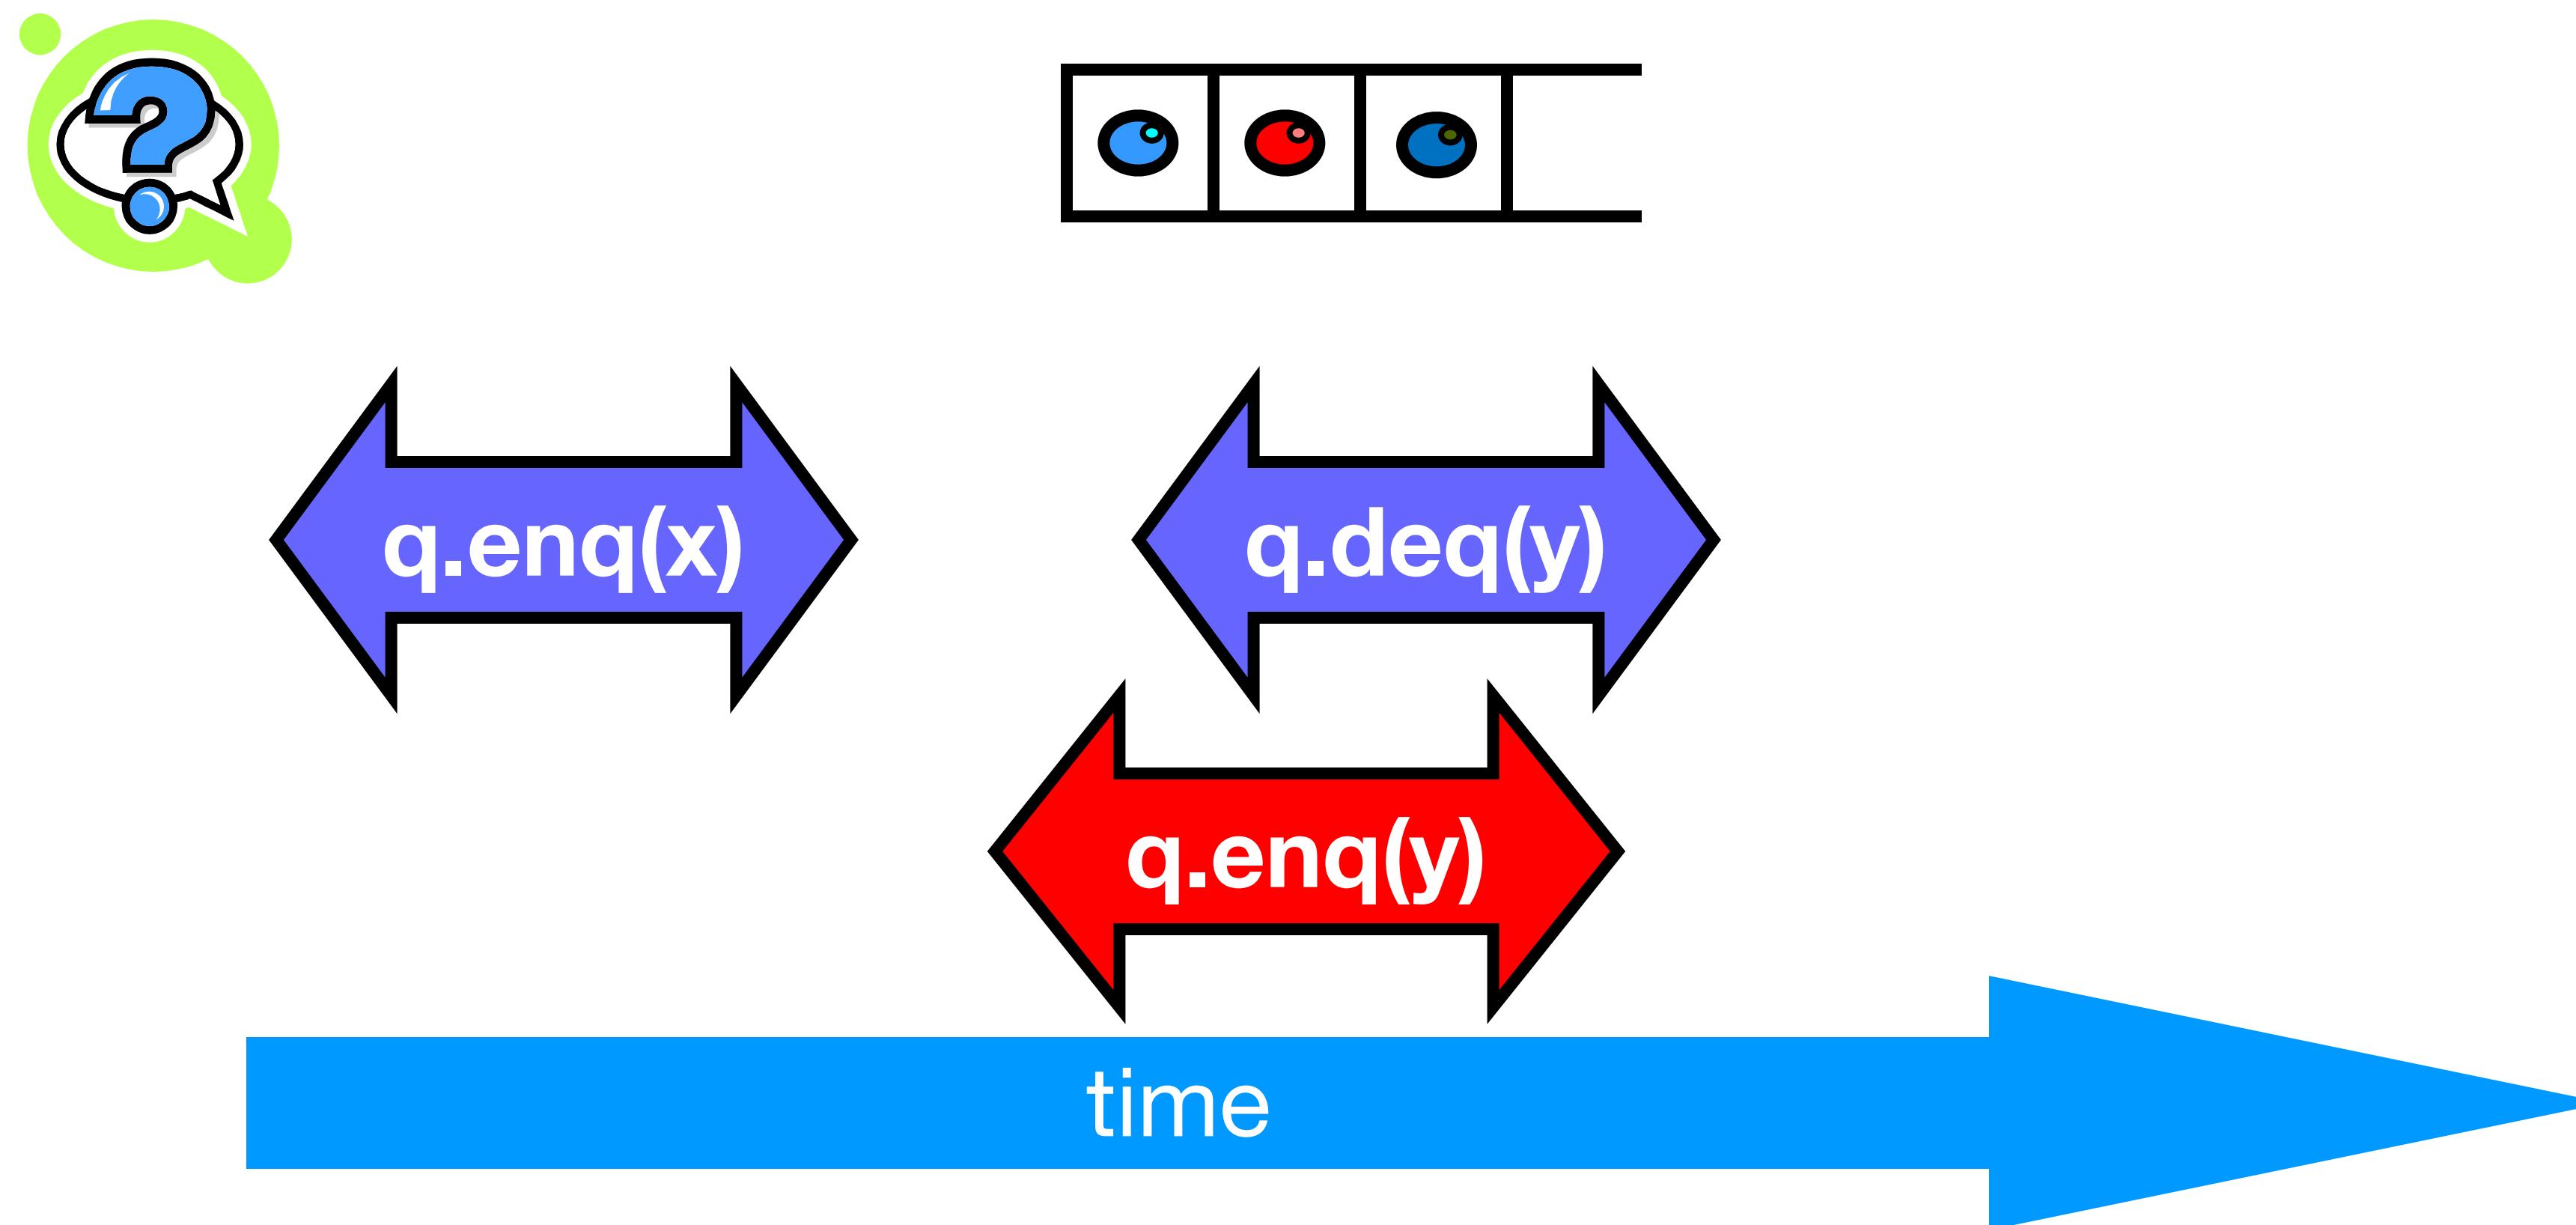
# Example



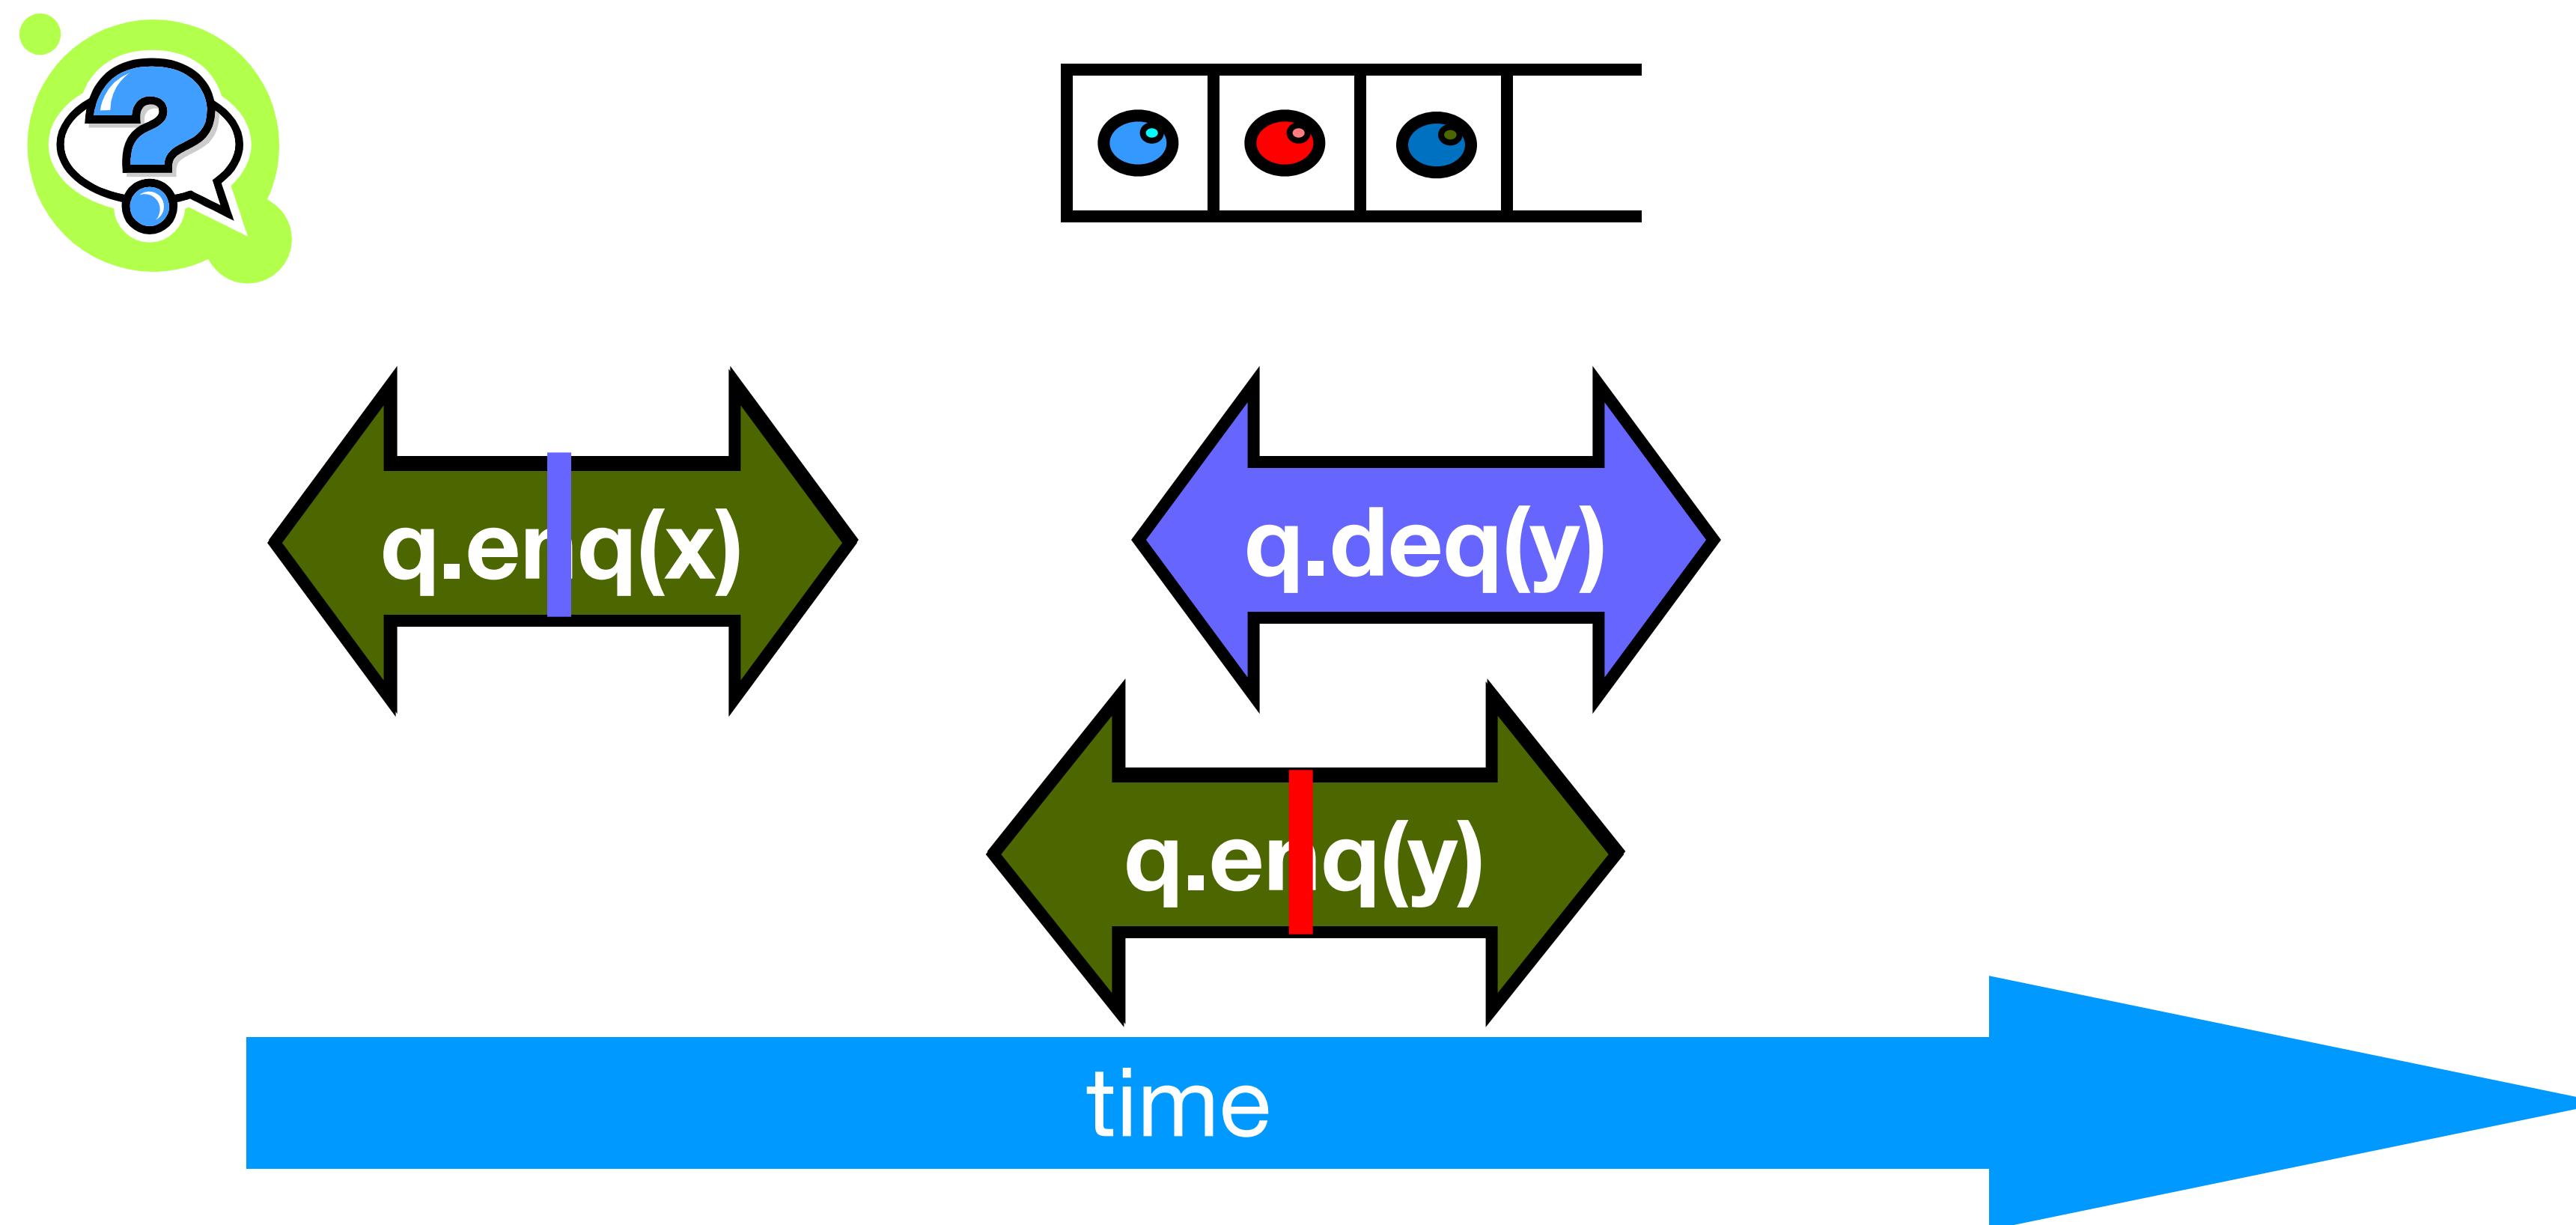
# Example



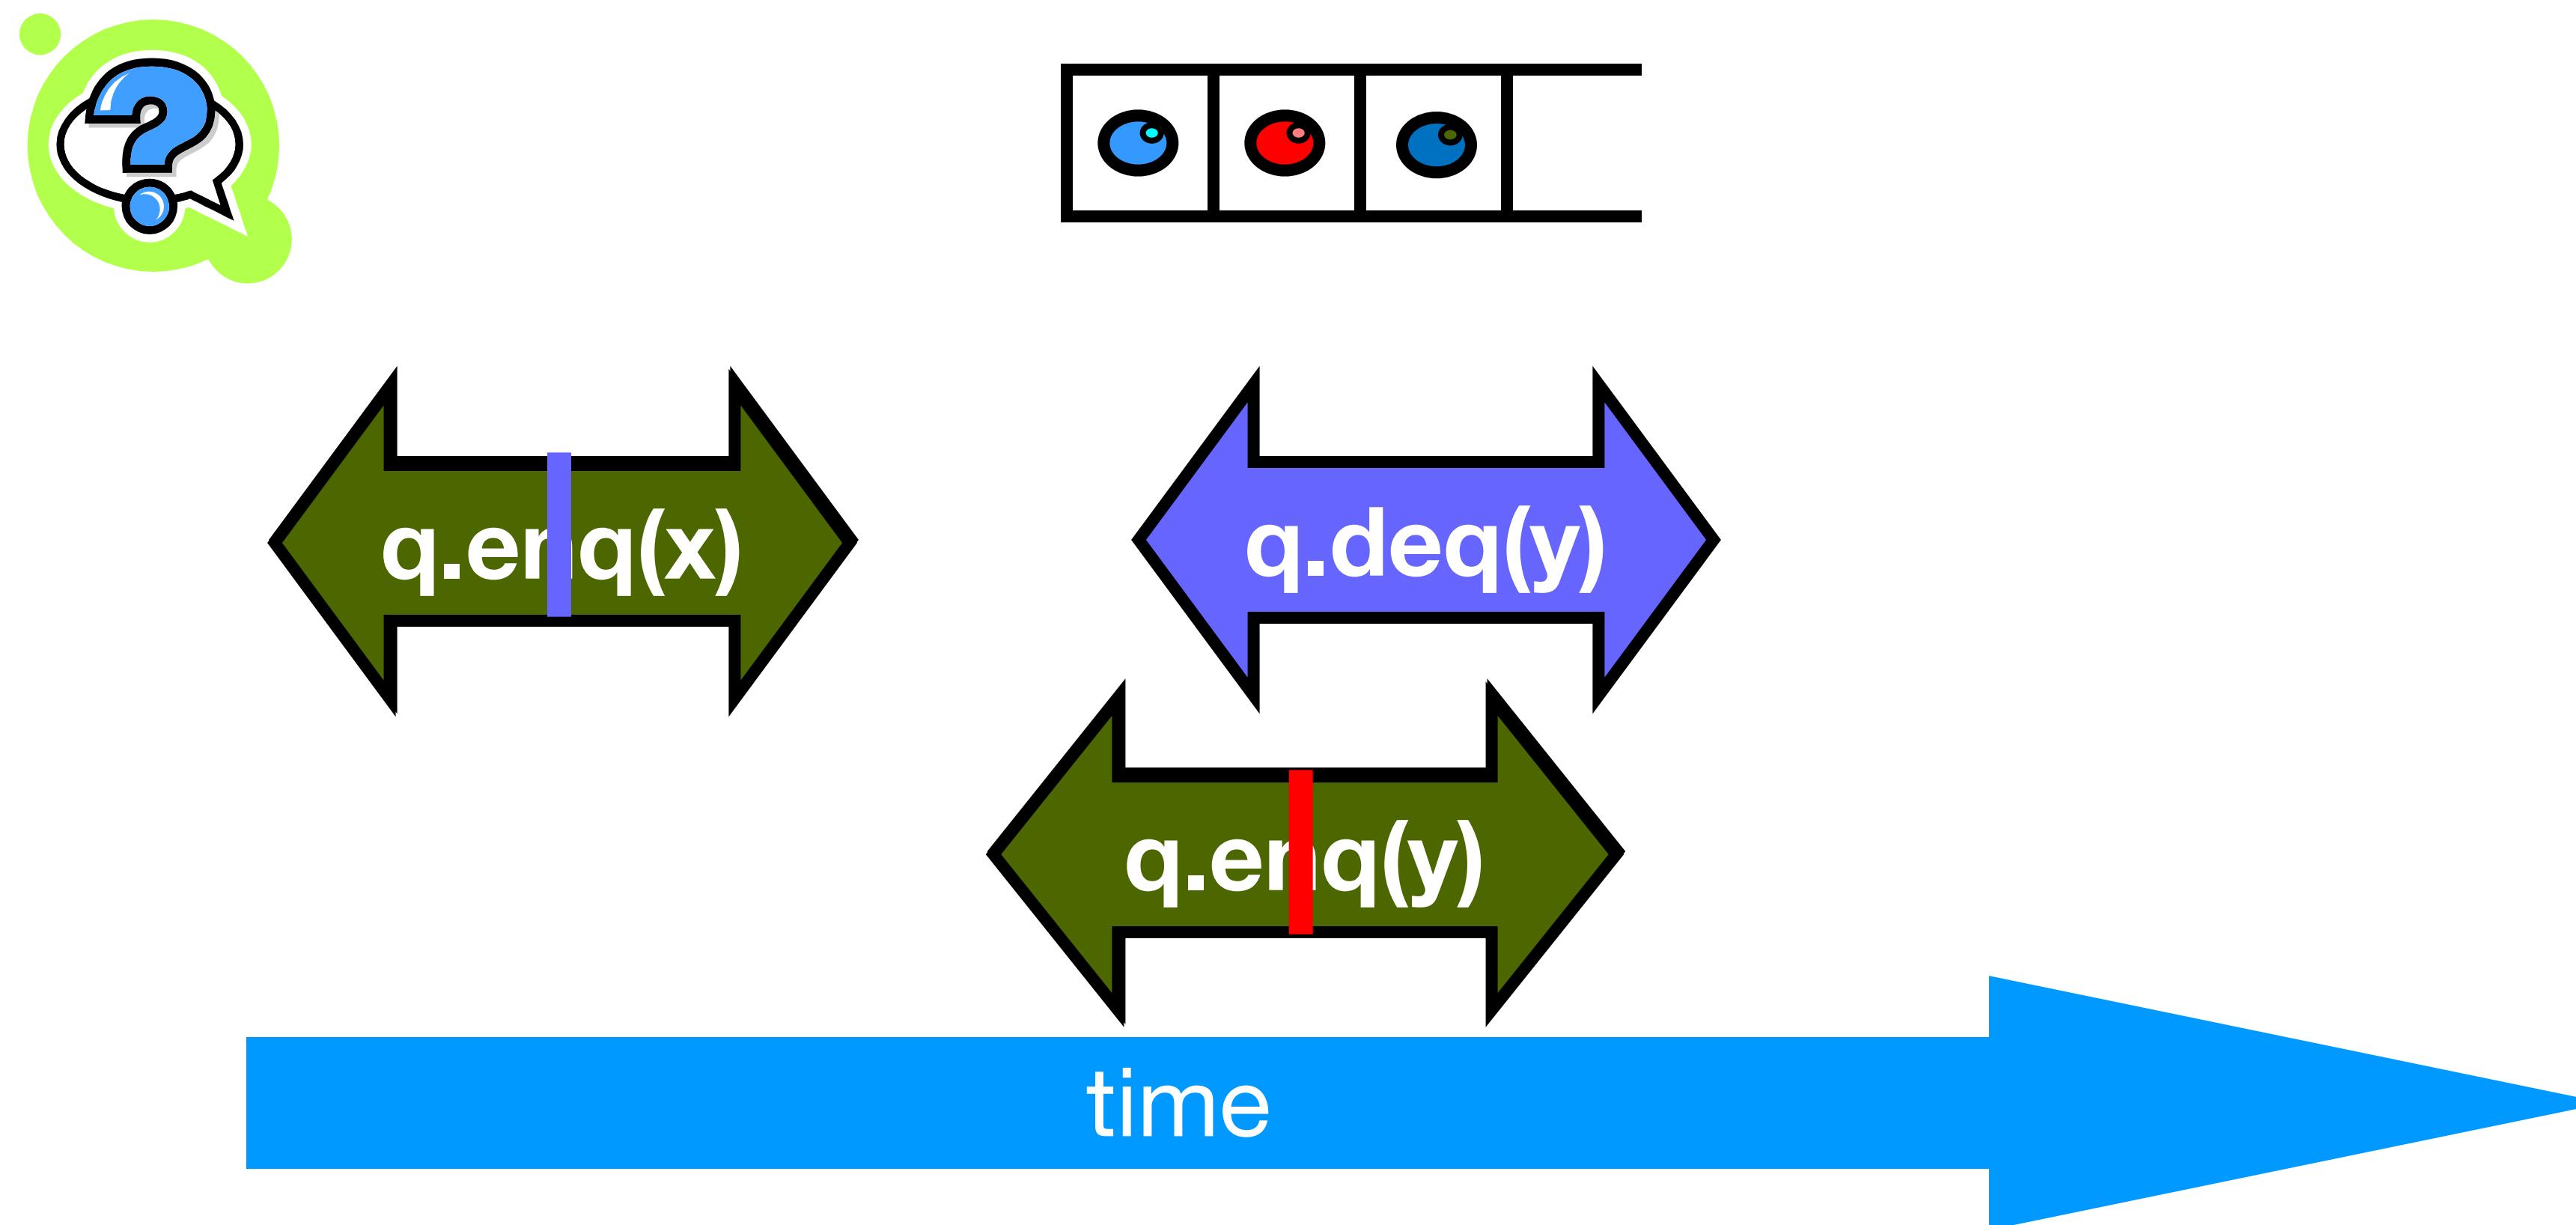
# Example



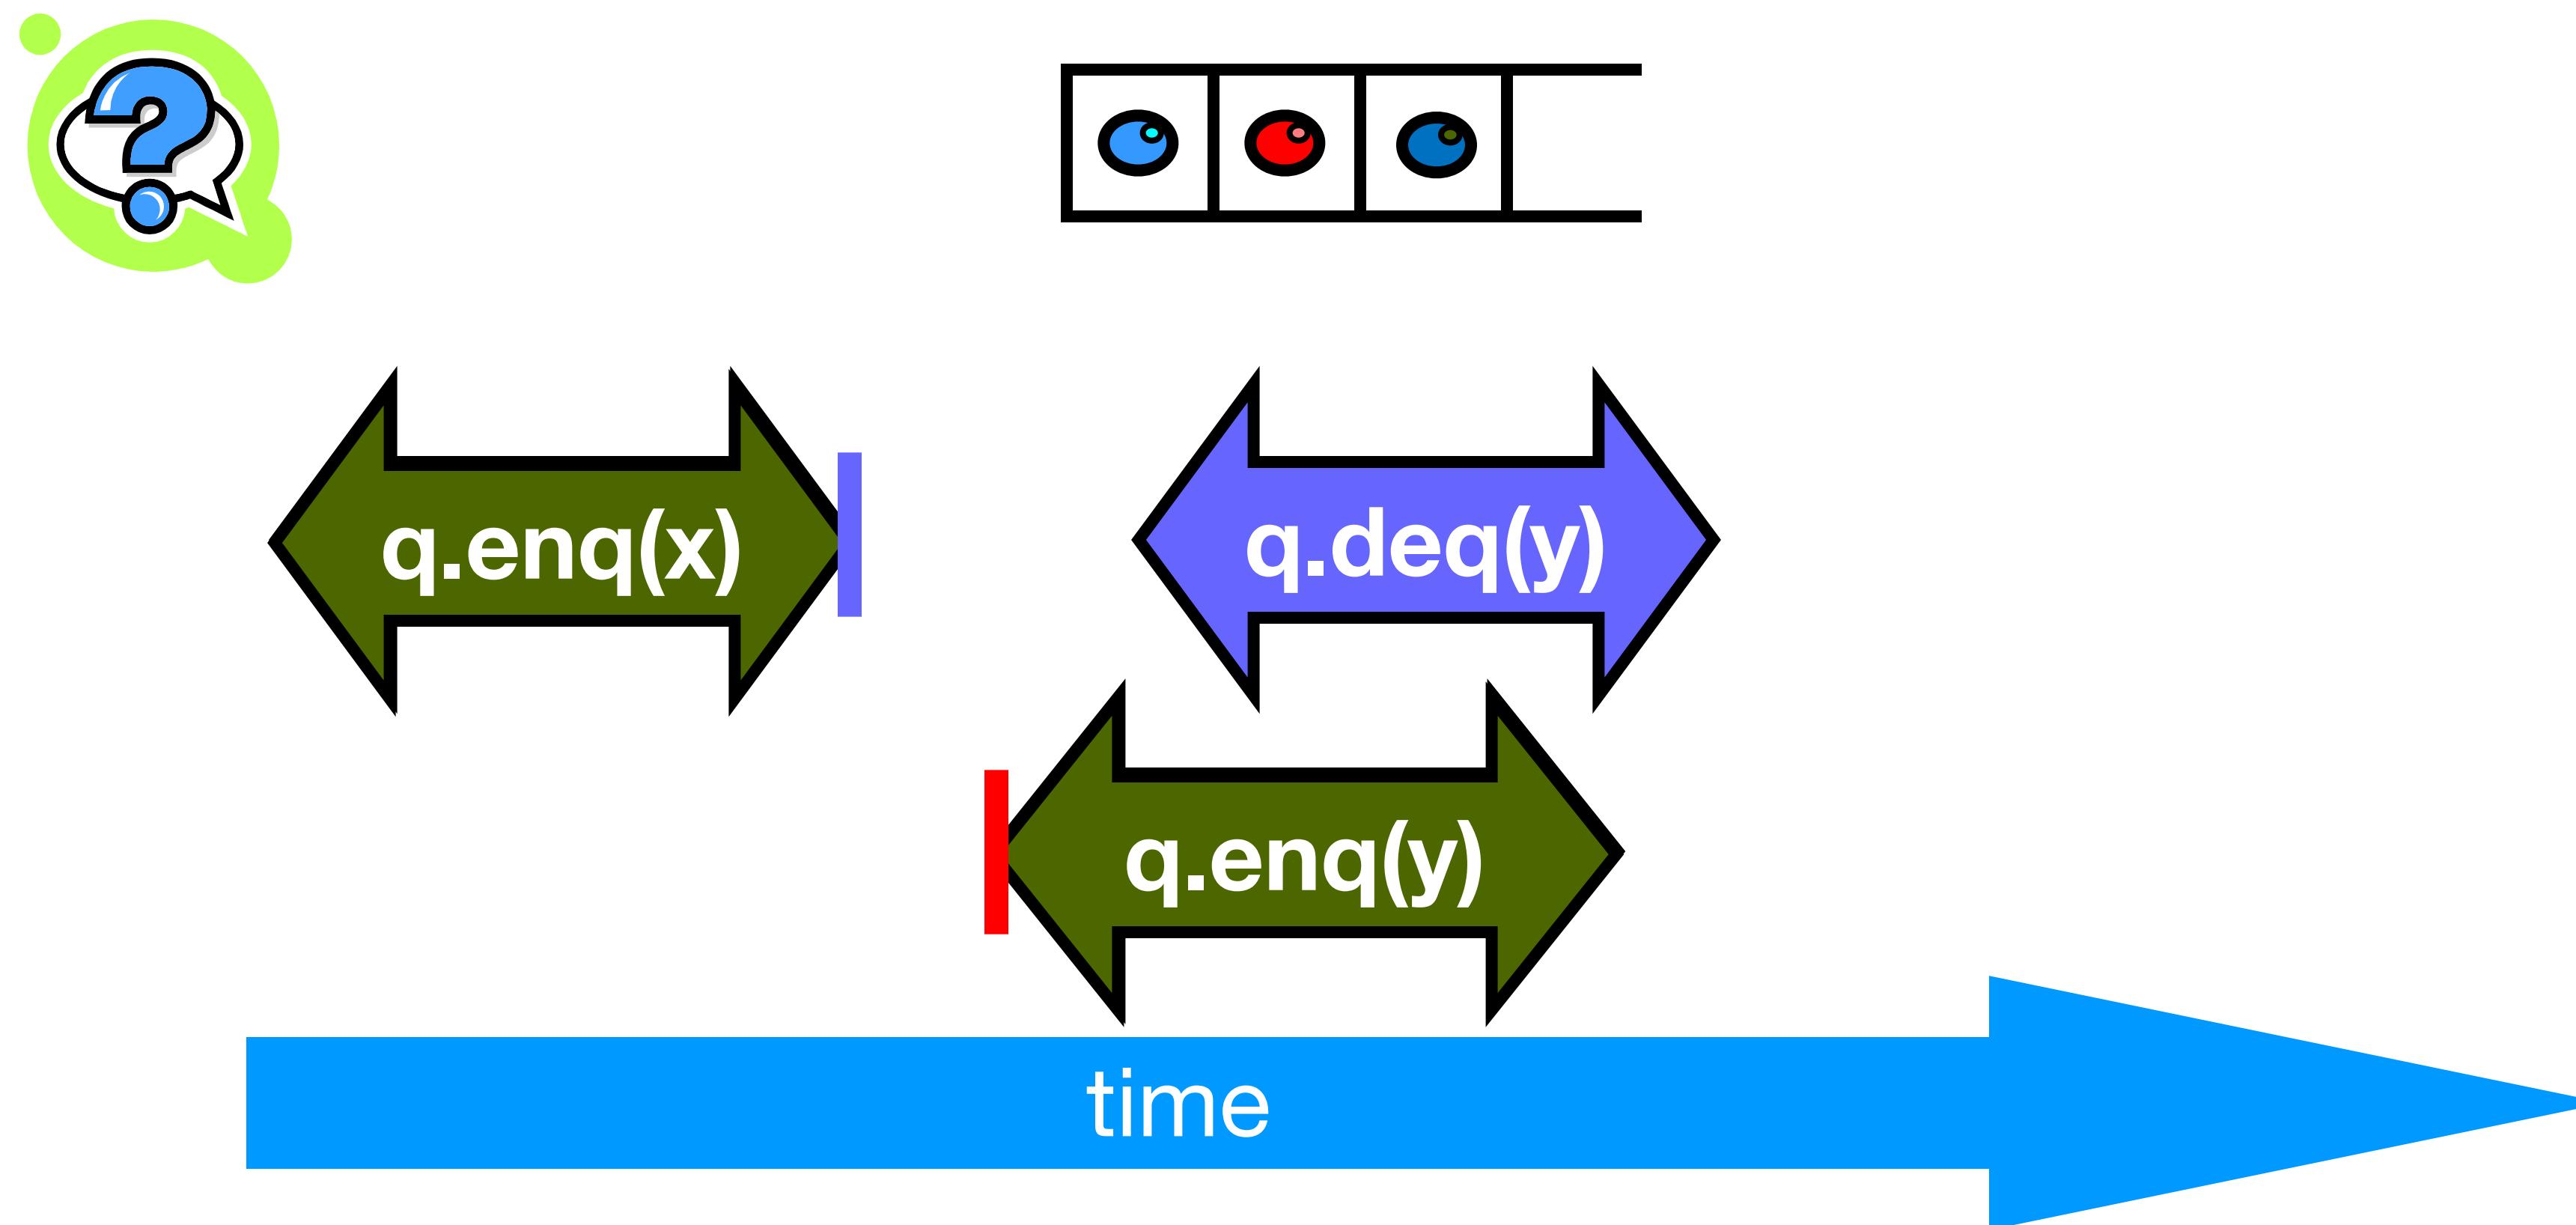
# Example



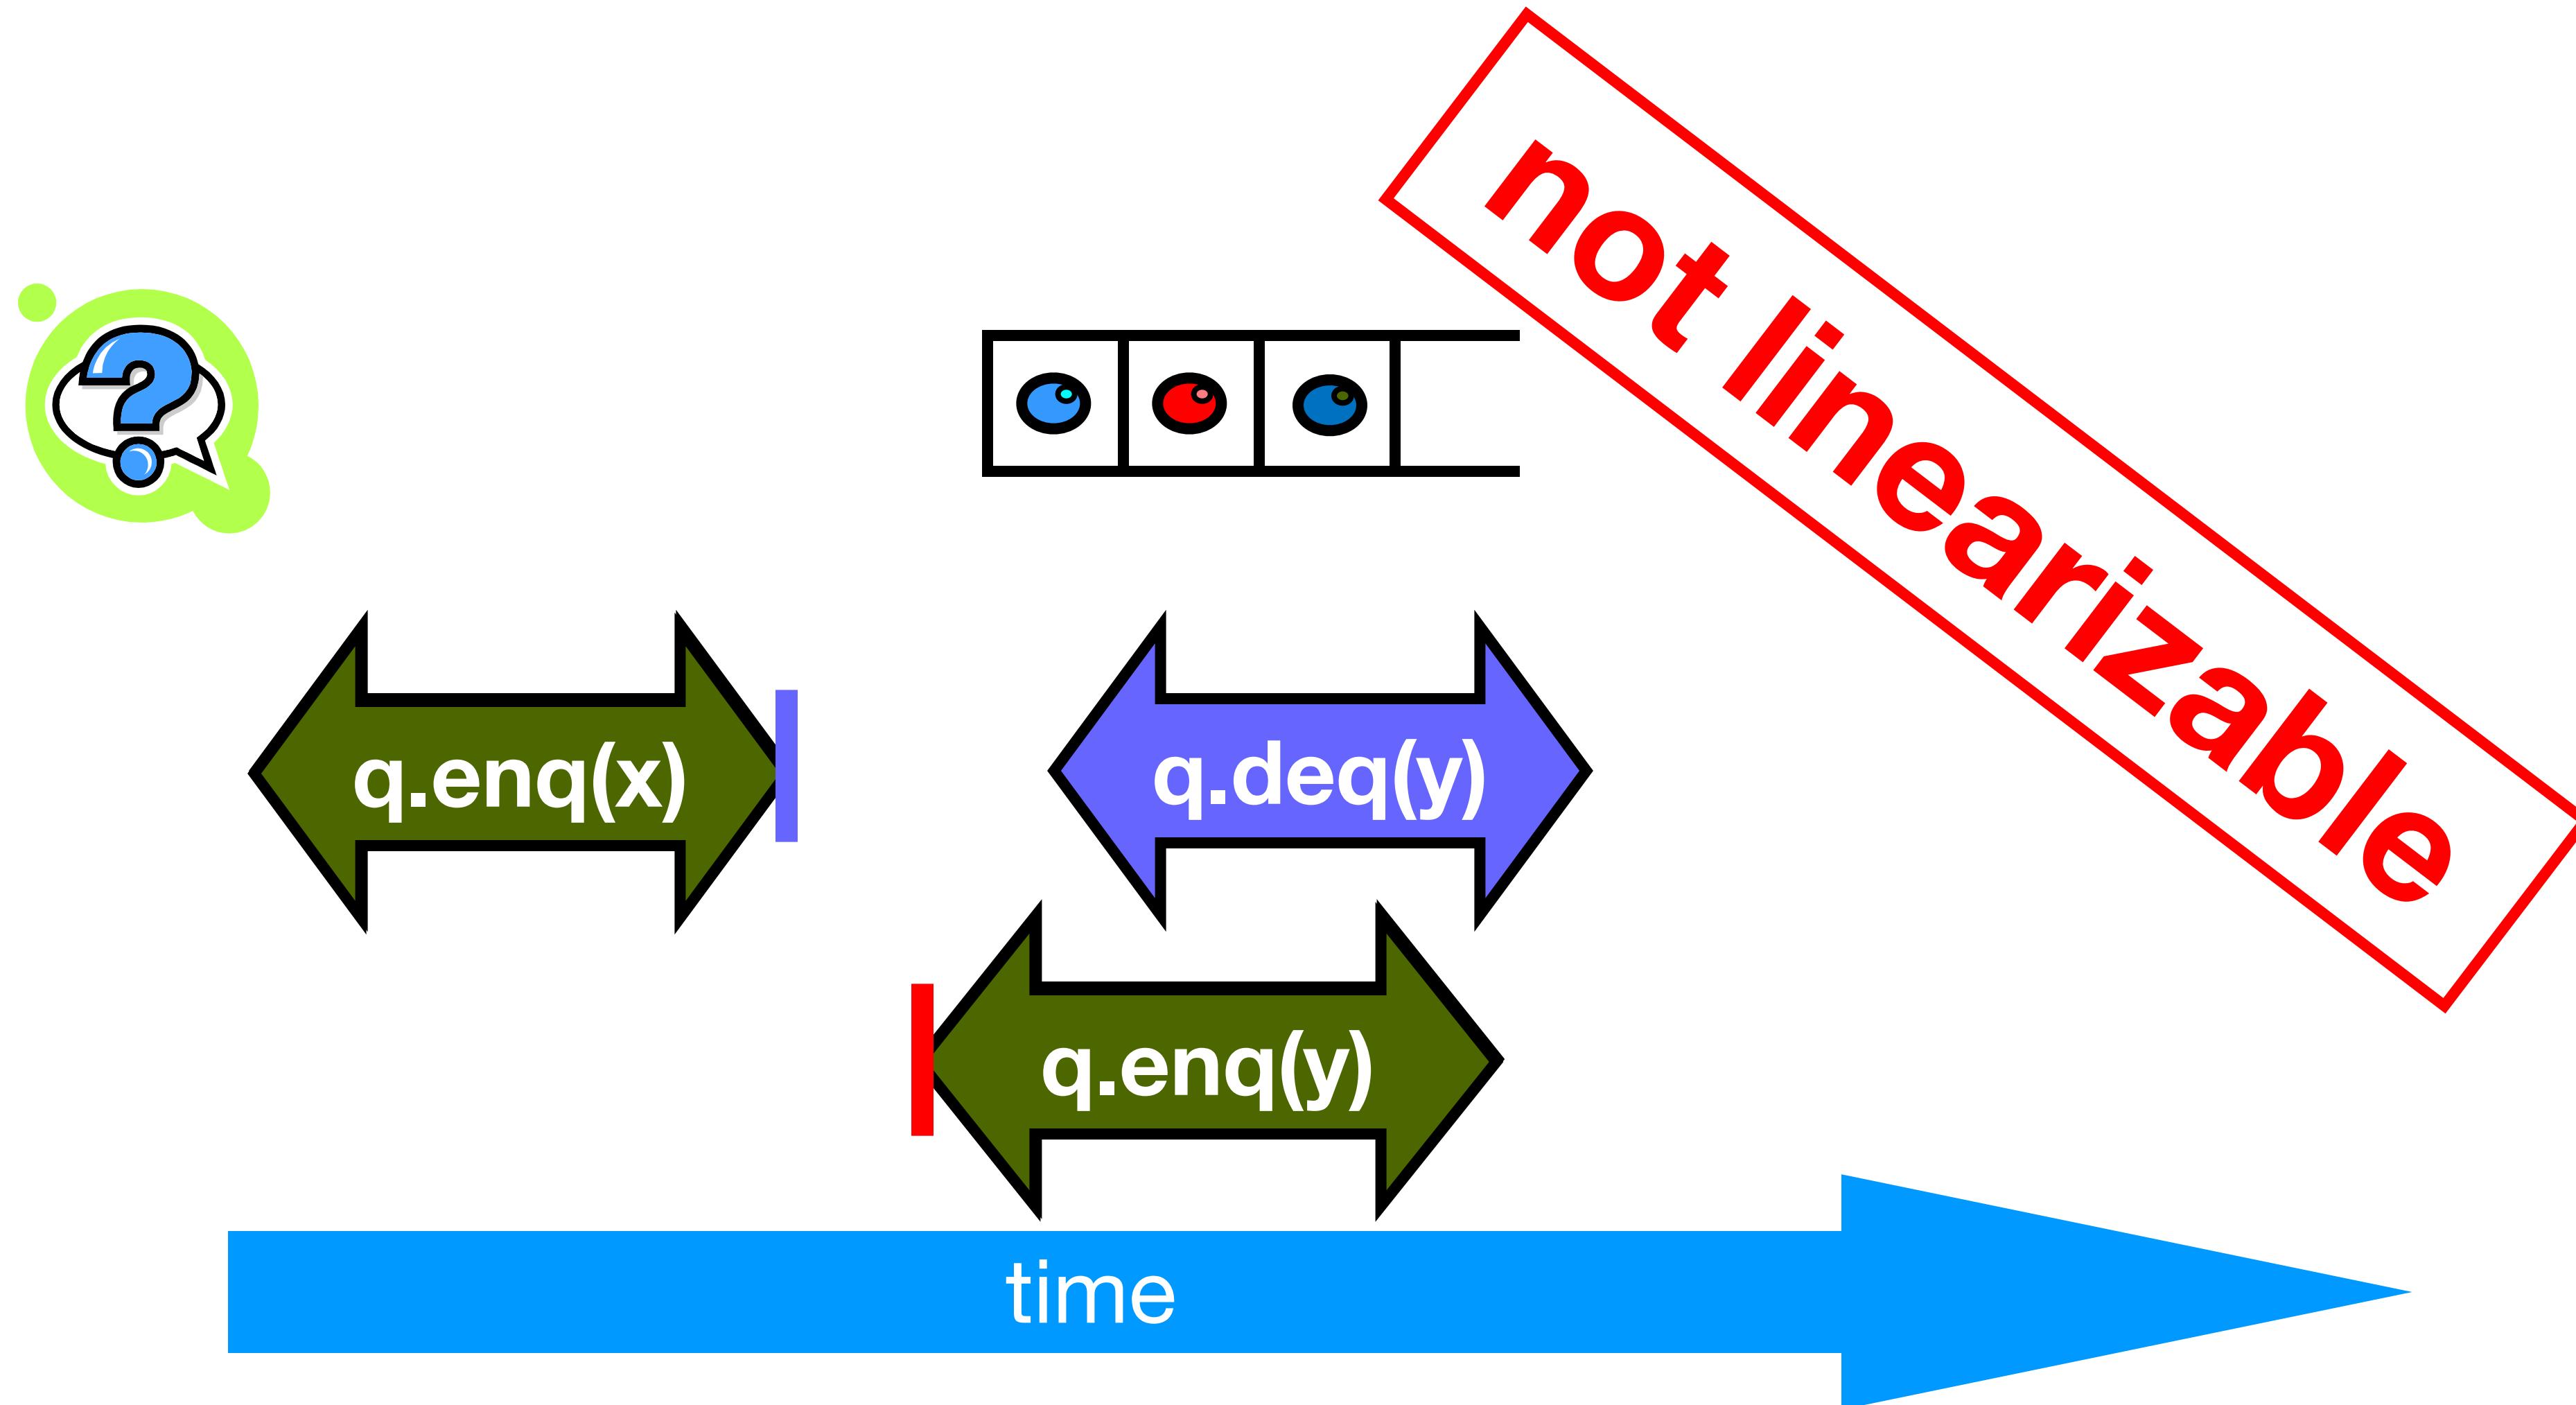
# Example



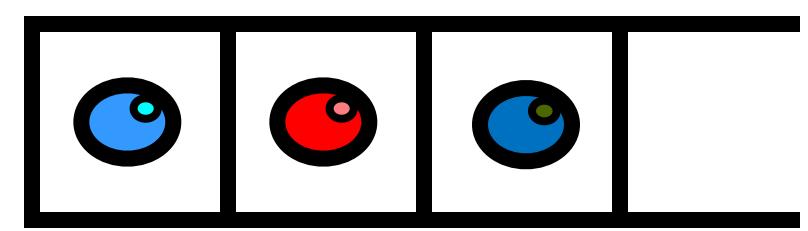
# Example




# Example





# Example



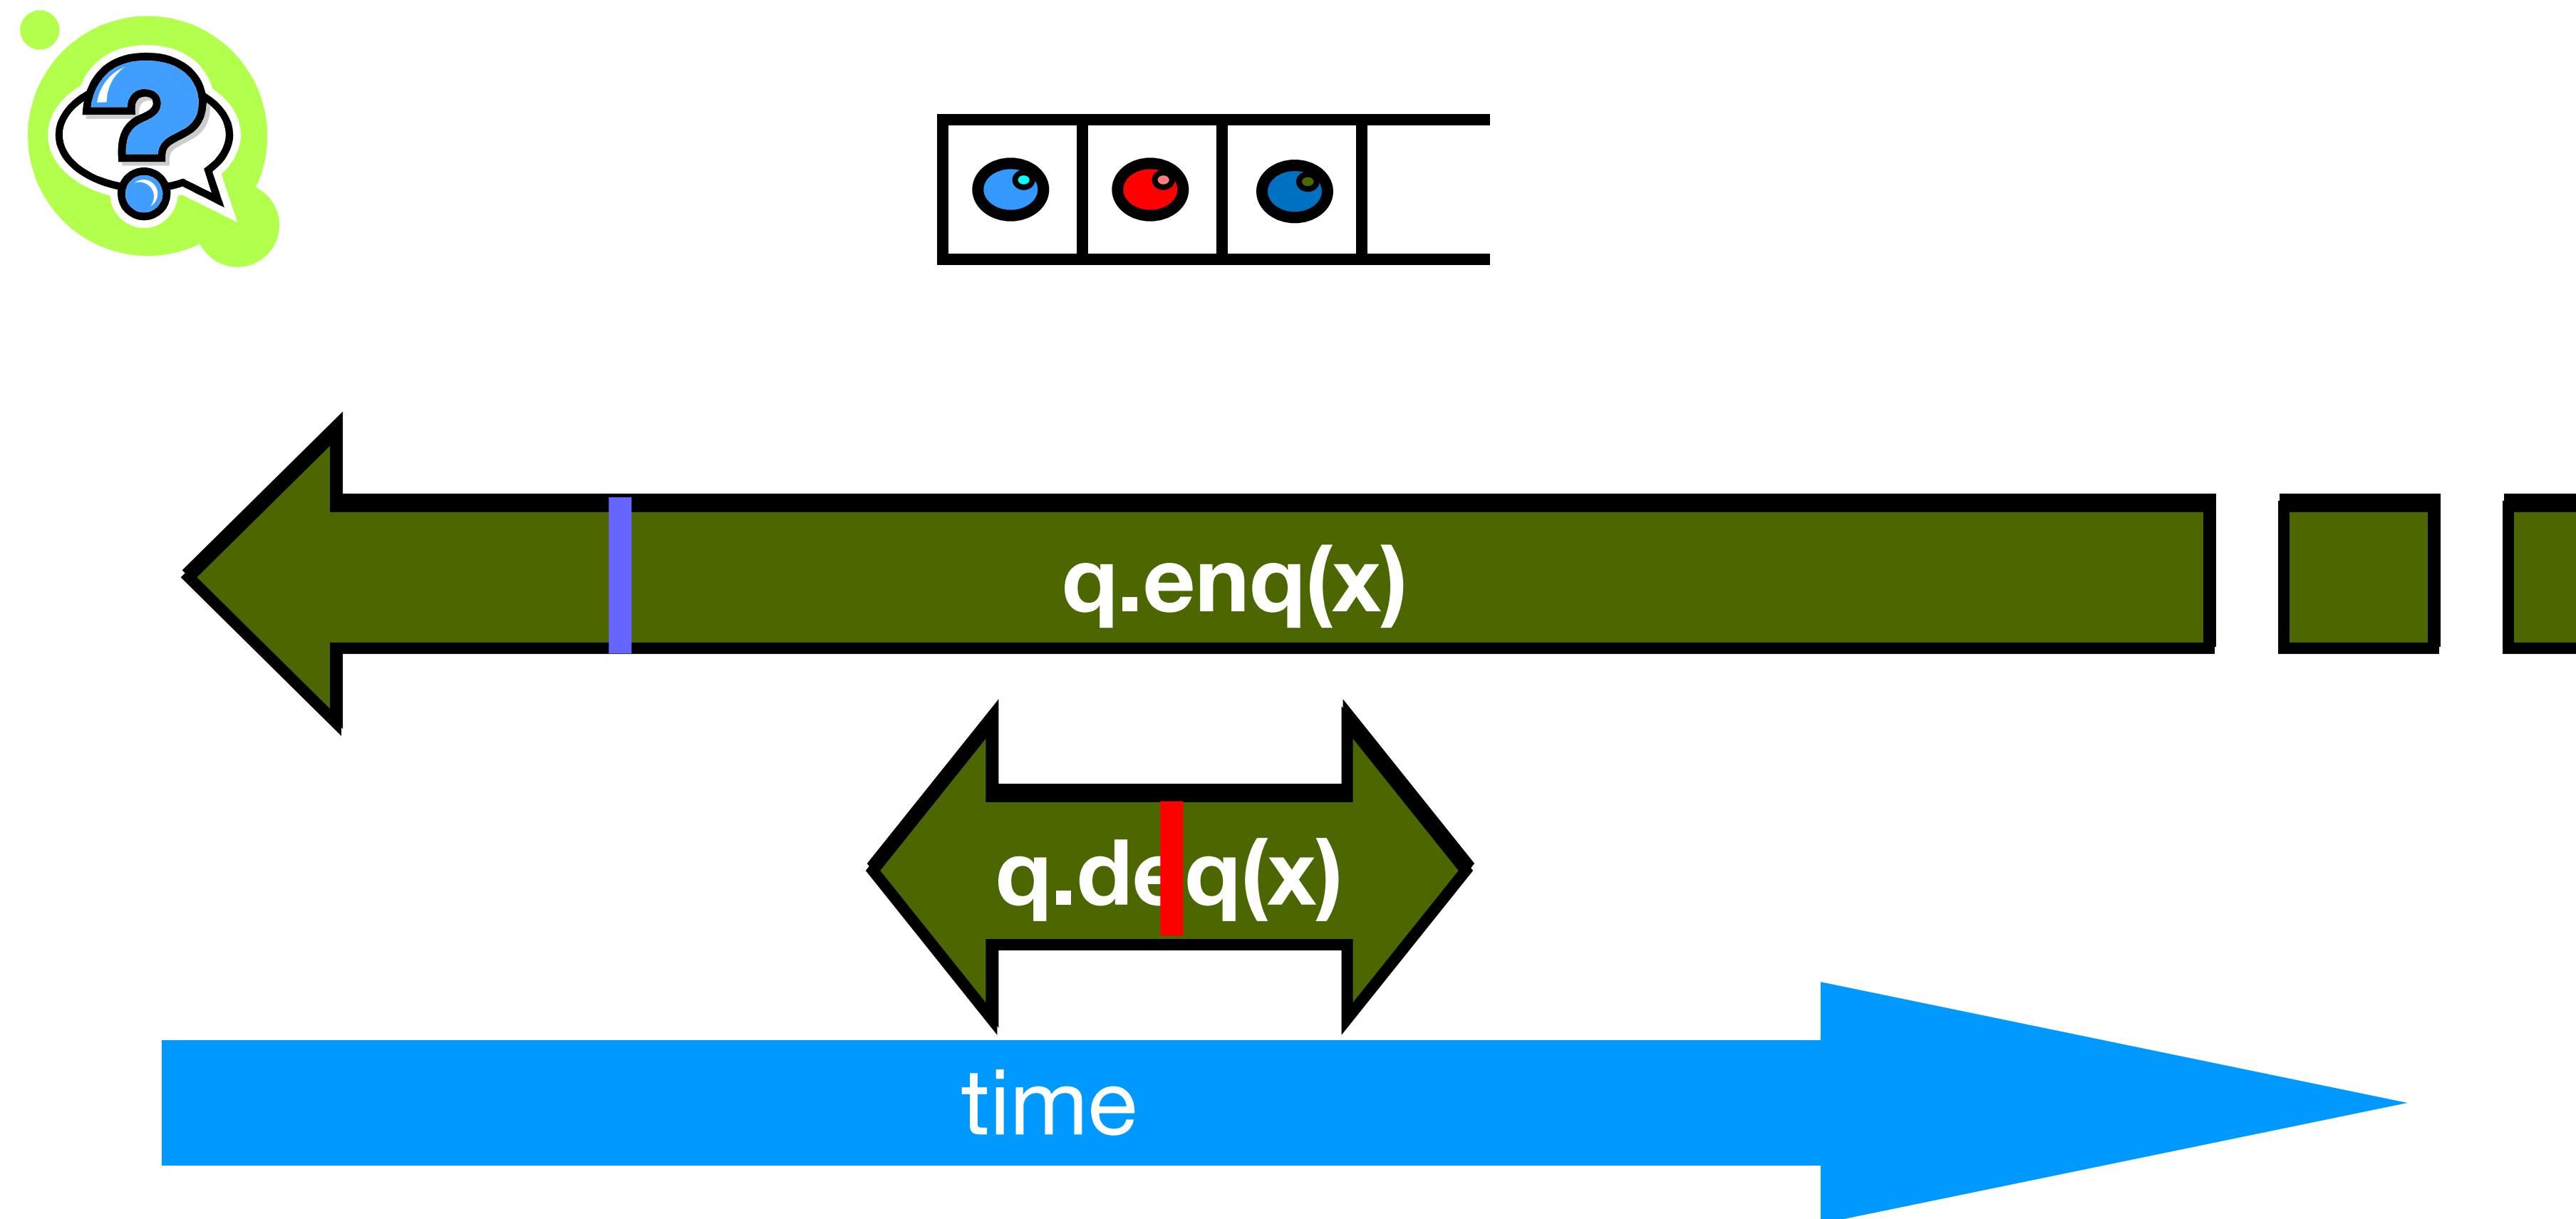
# Example



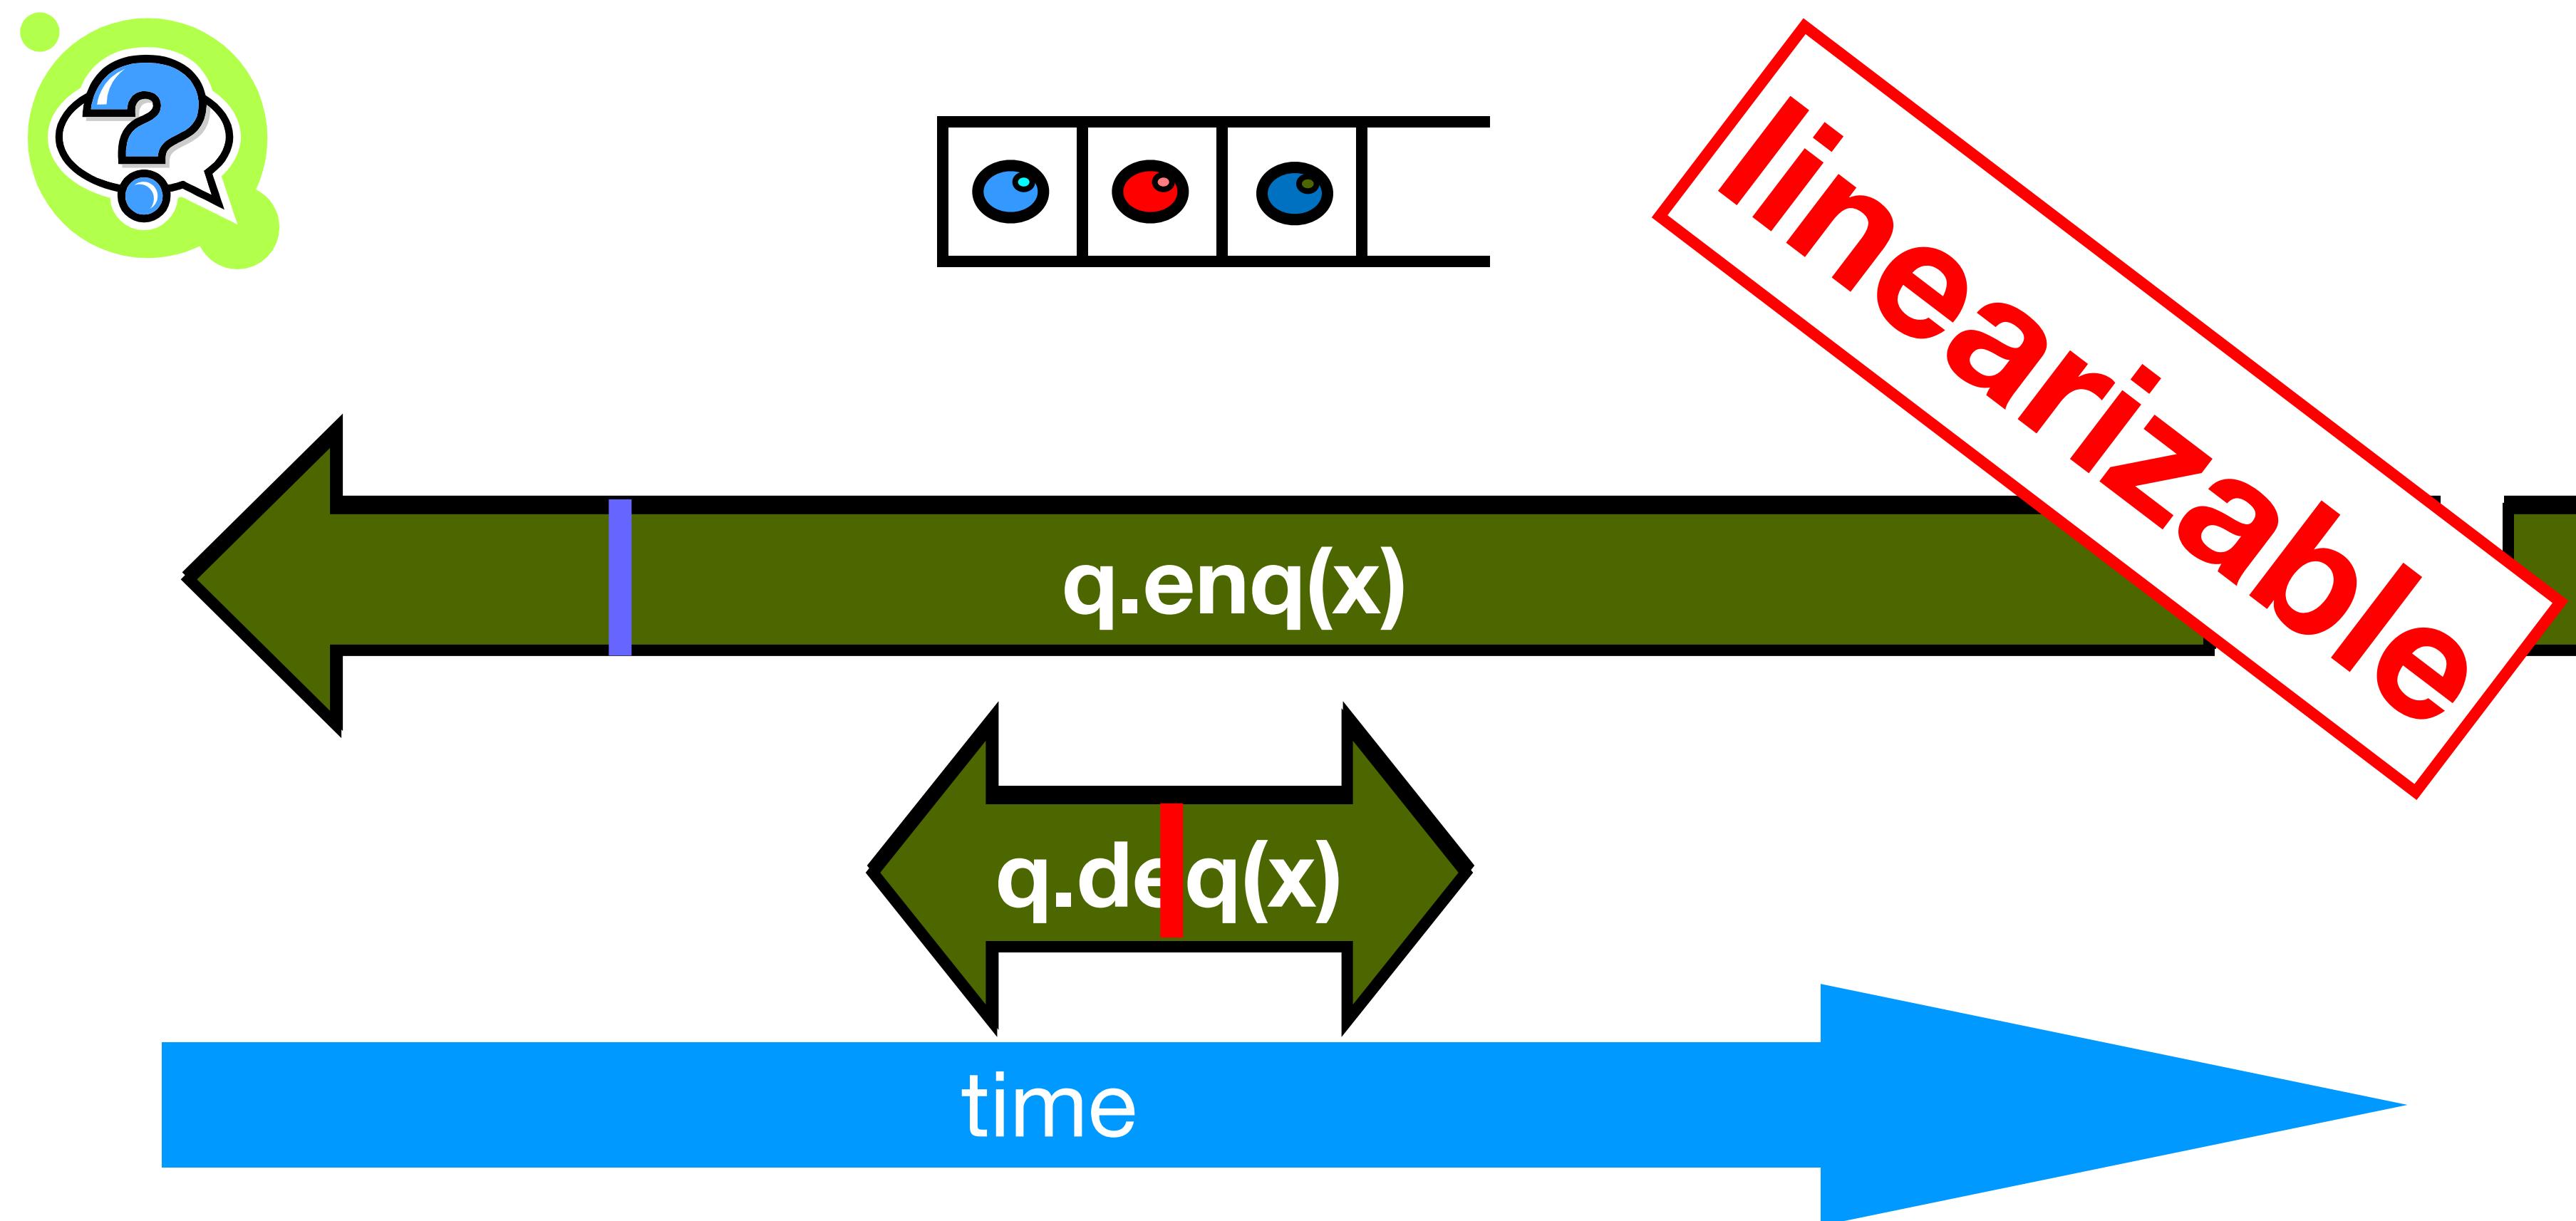
# Example



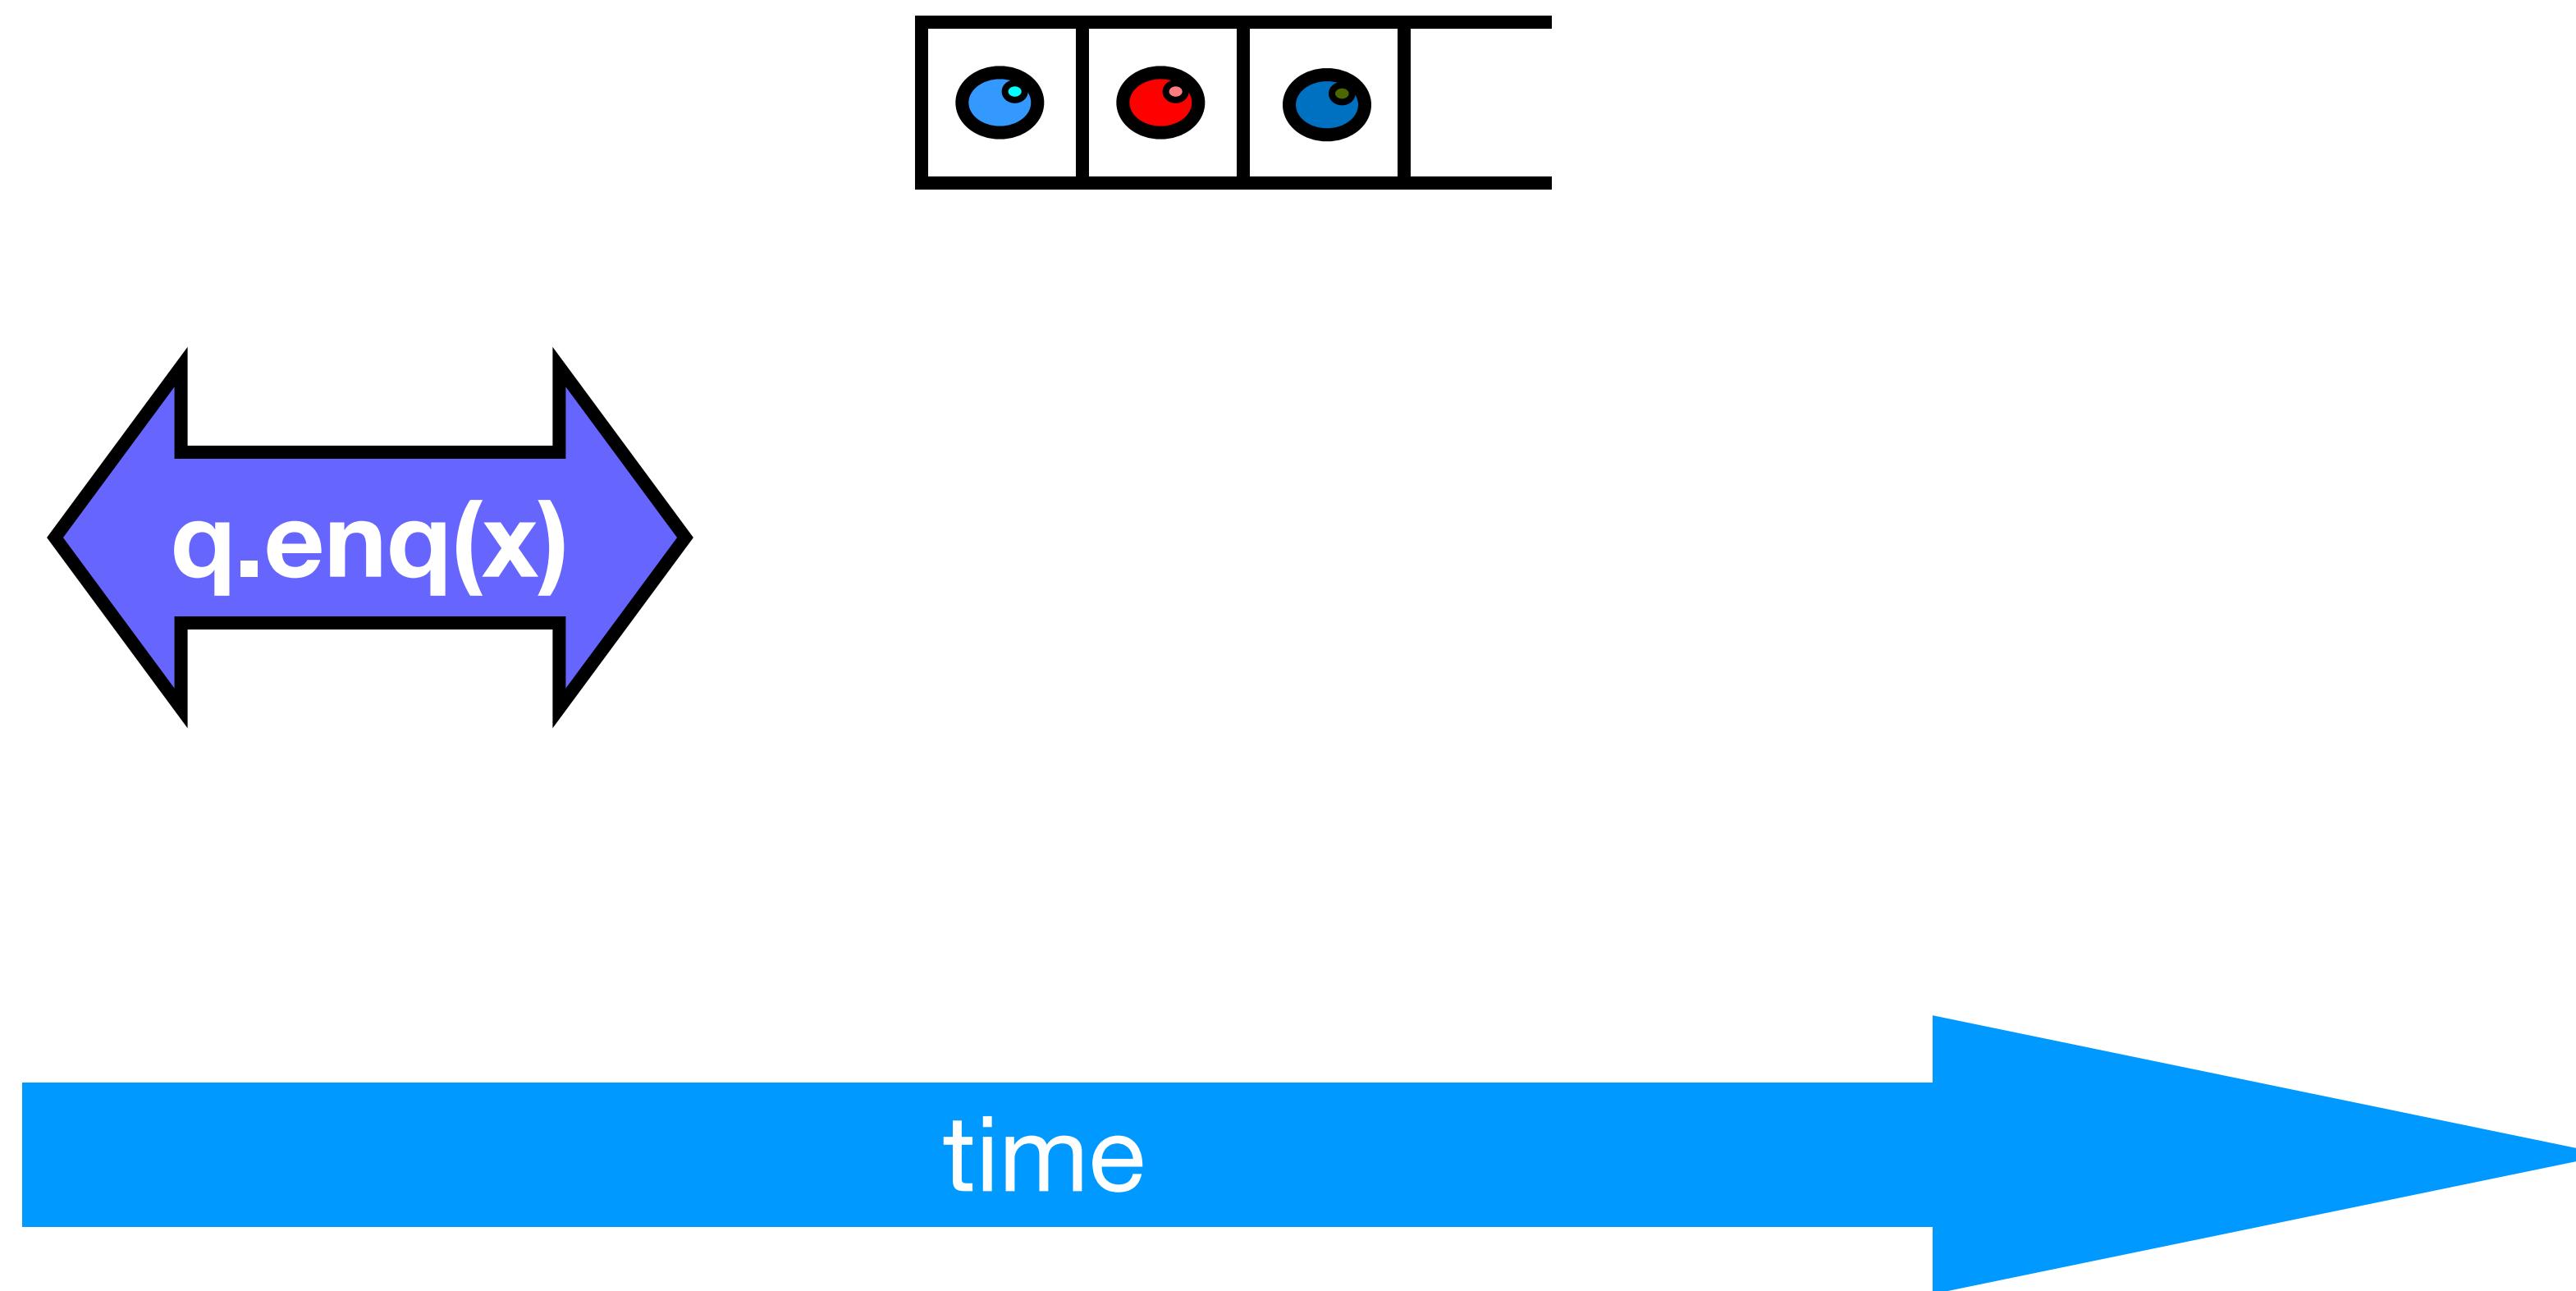
# Example




# Example



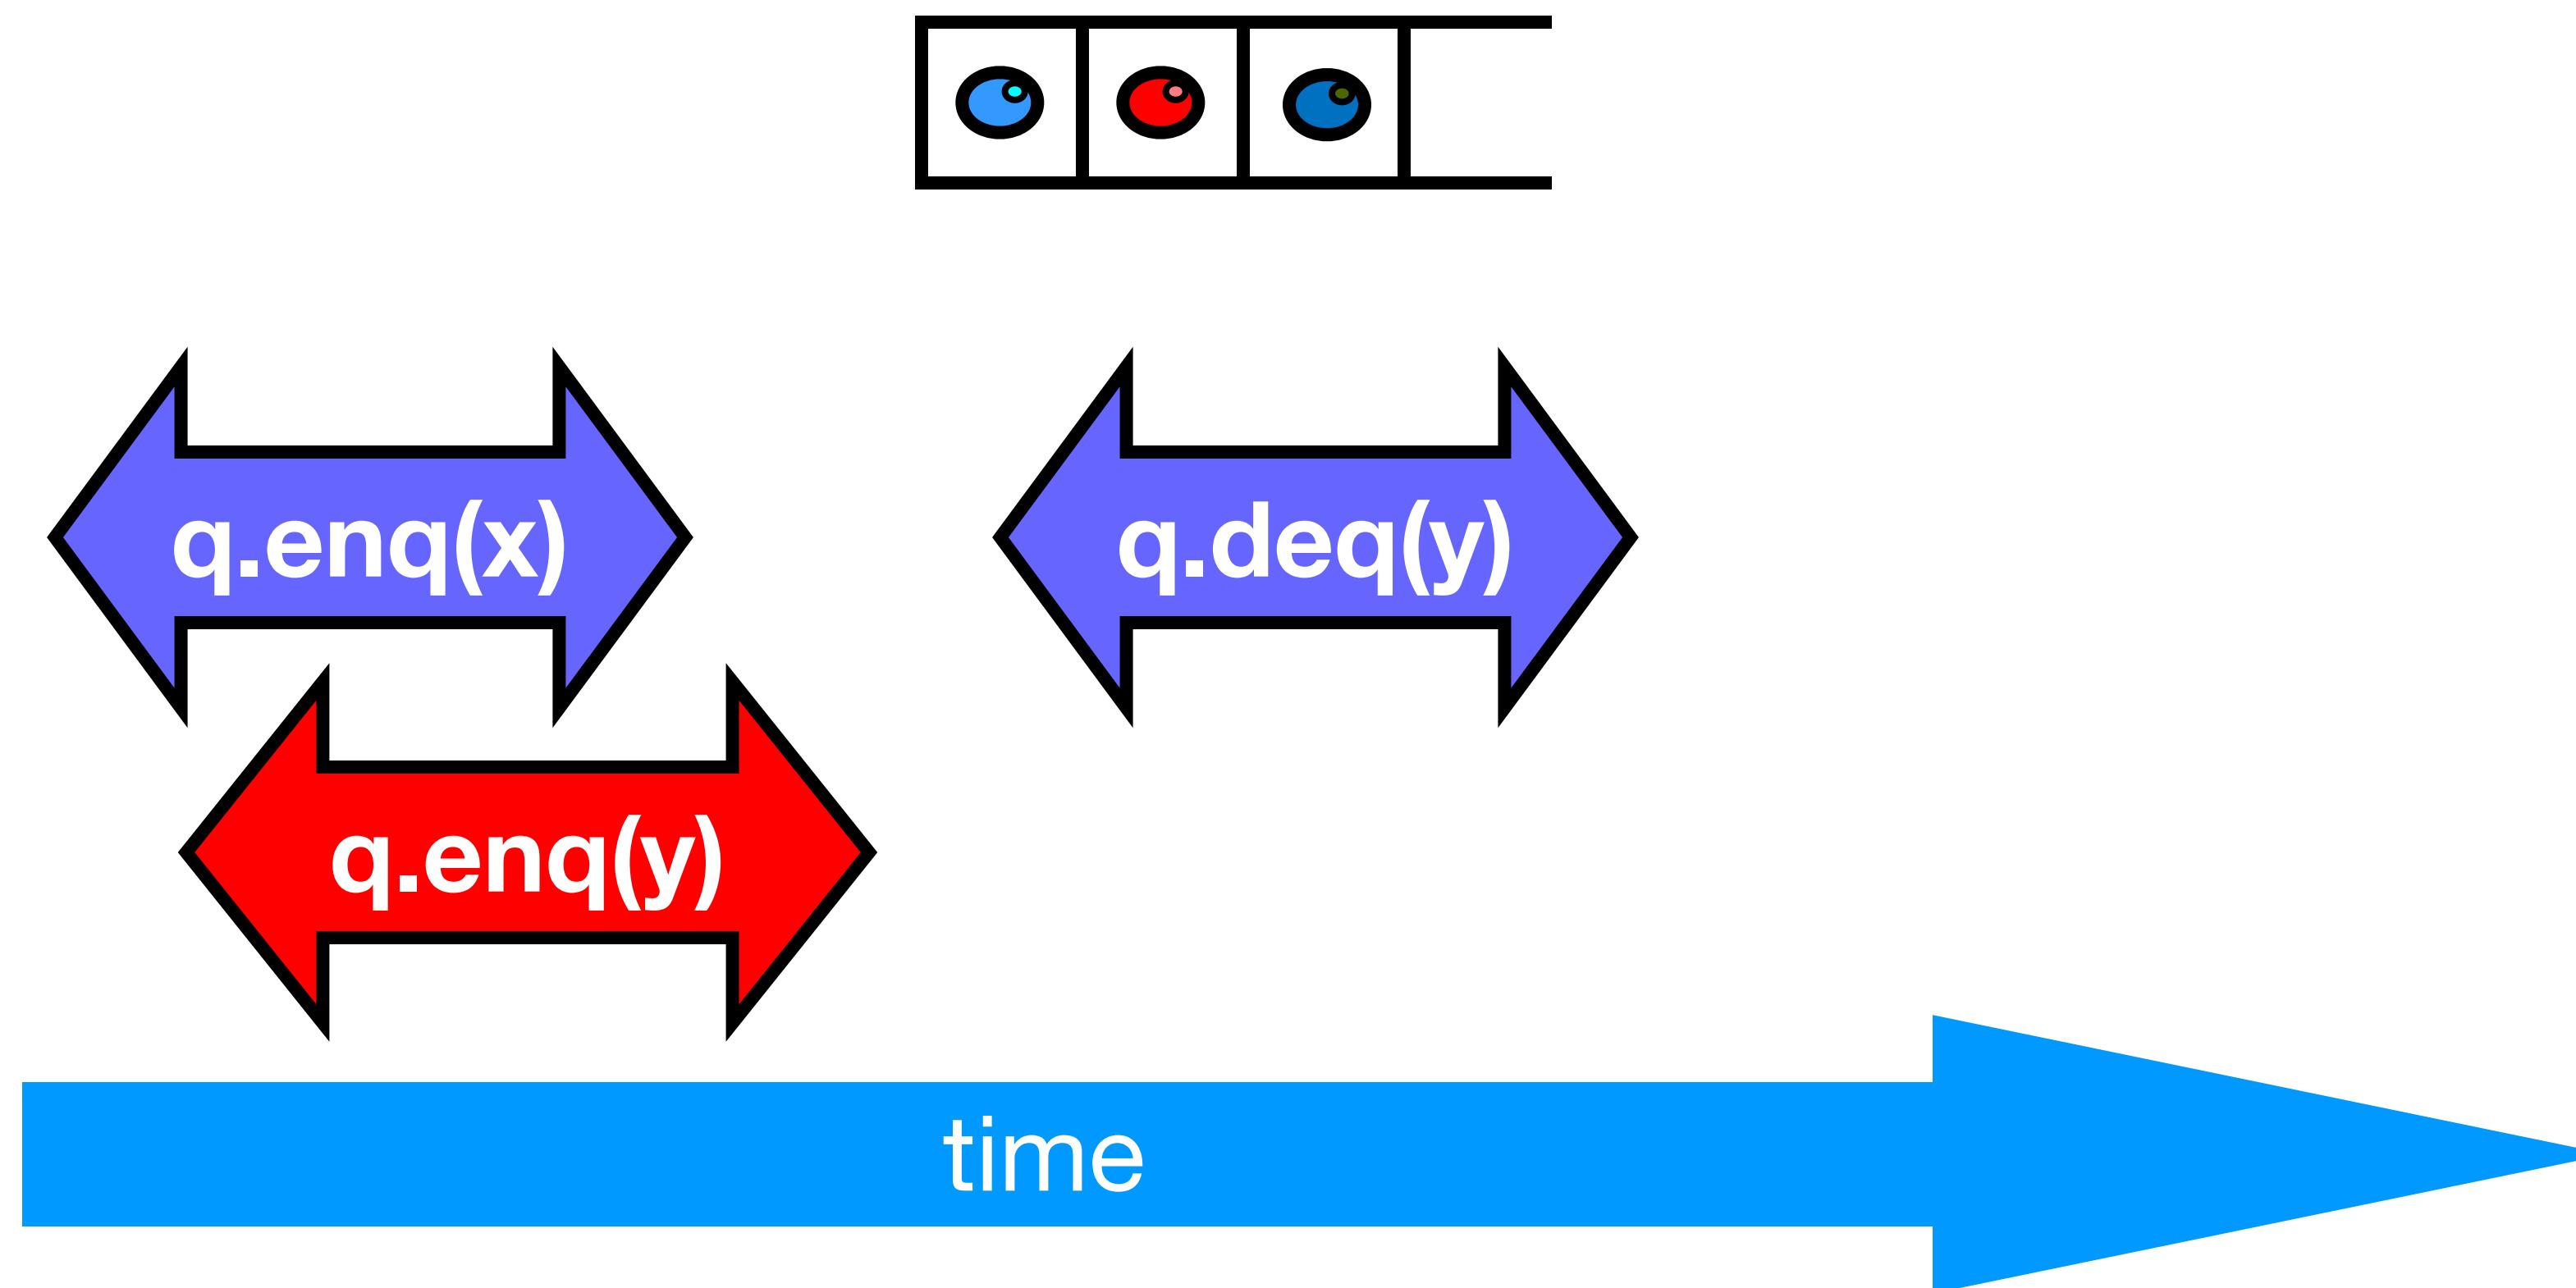
# Example




# Example



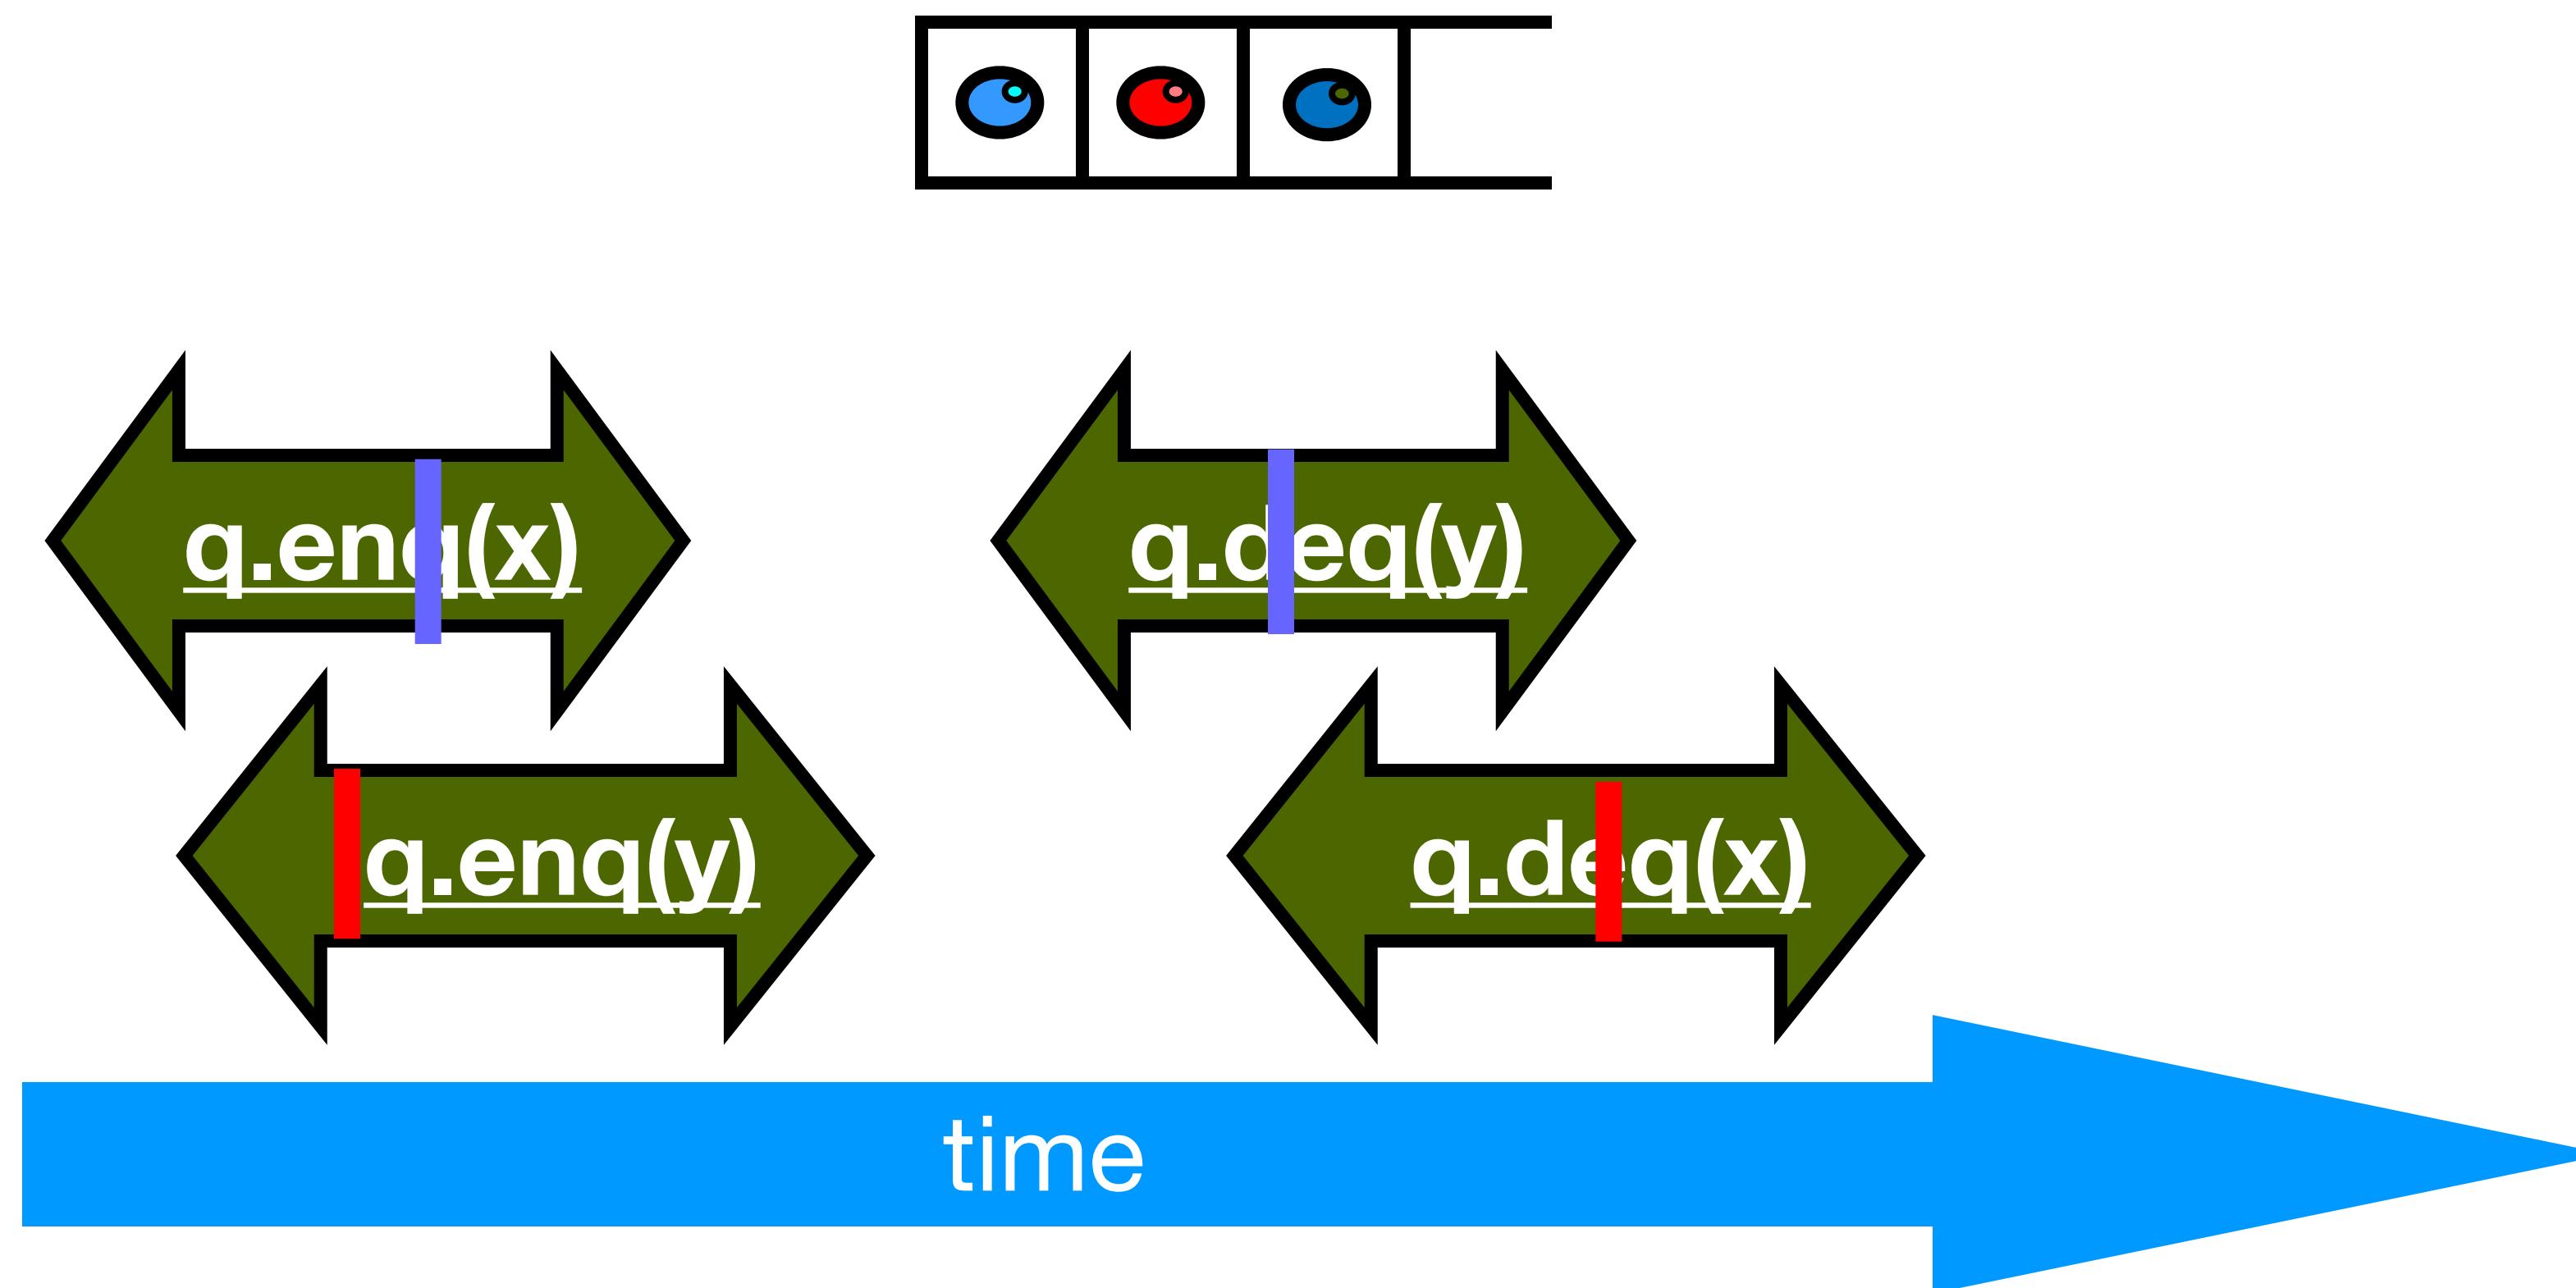
# Example



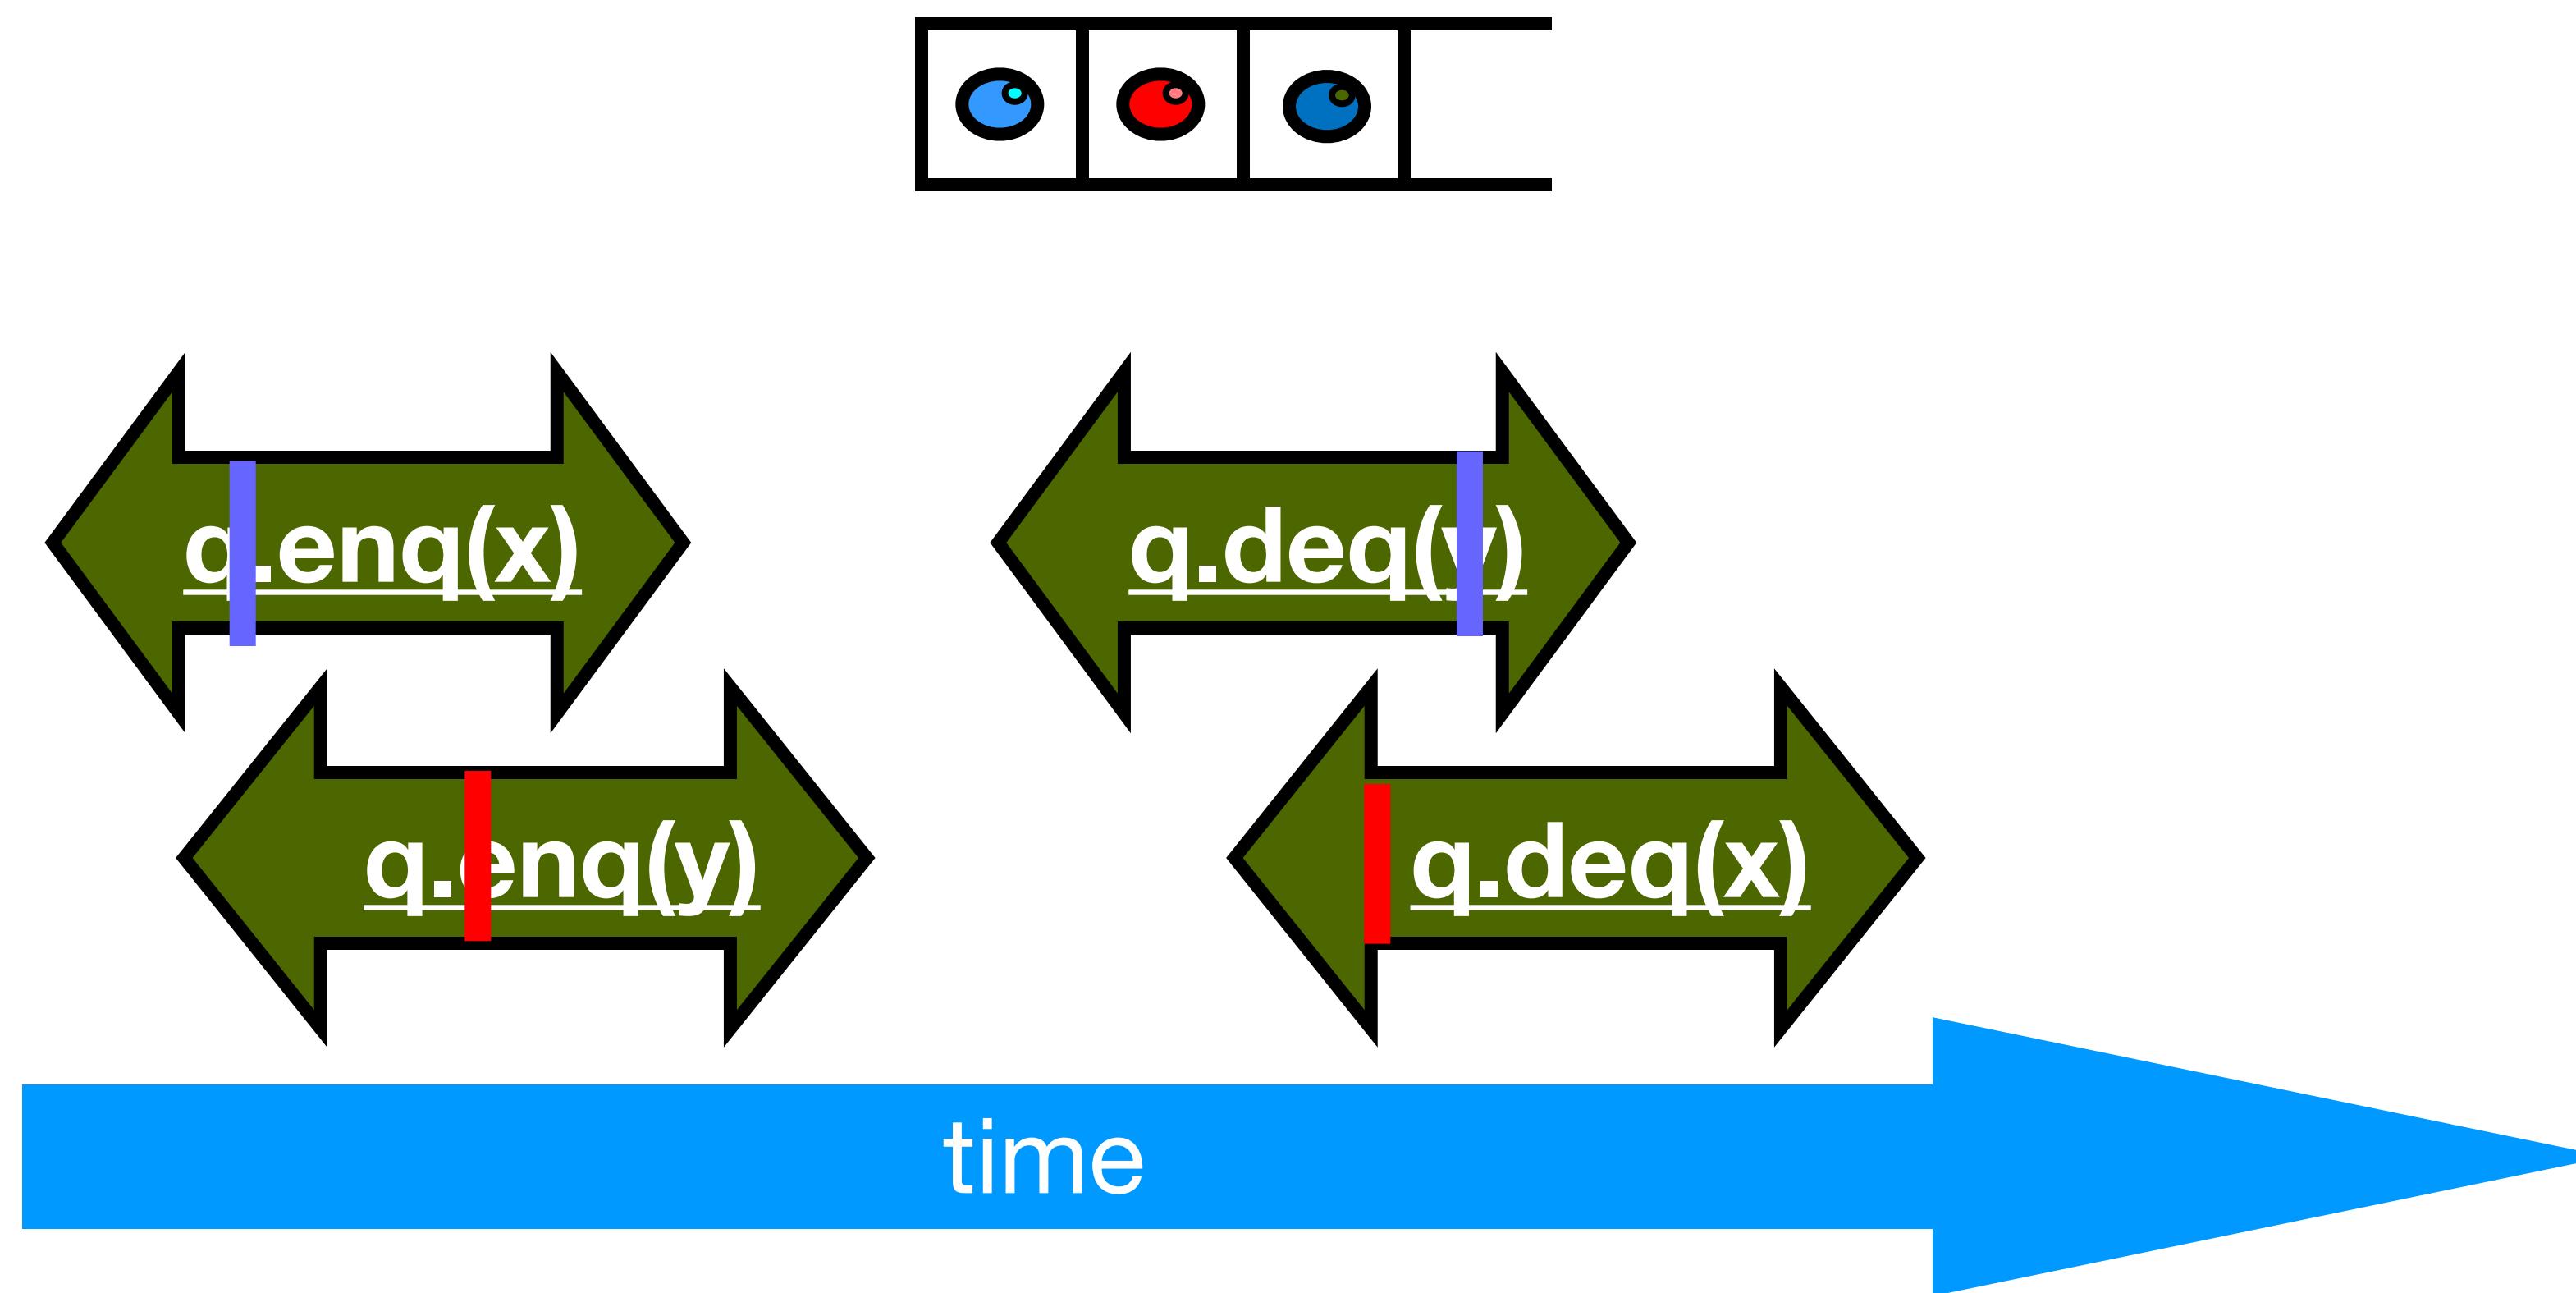
# Example




# Example



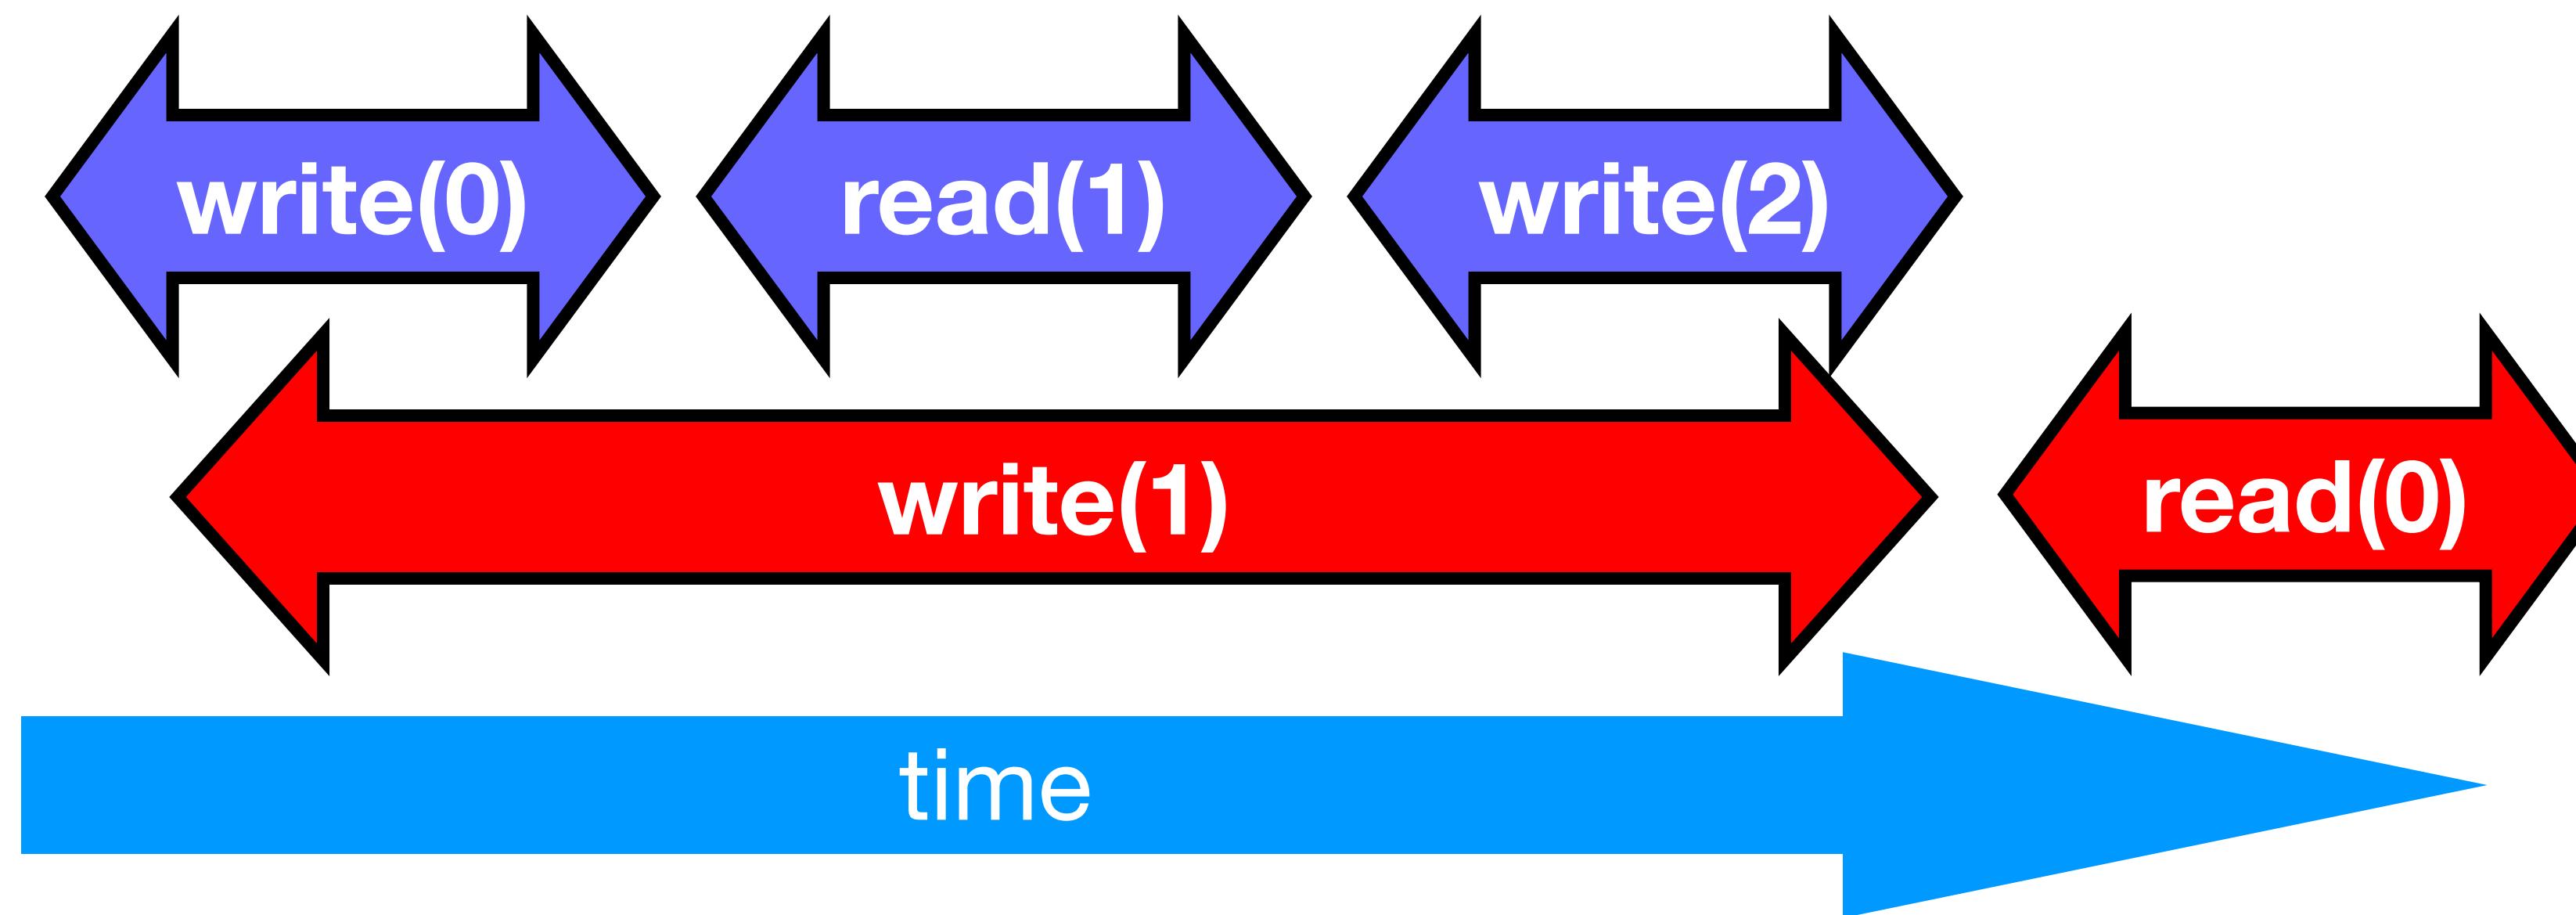
# Example



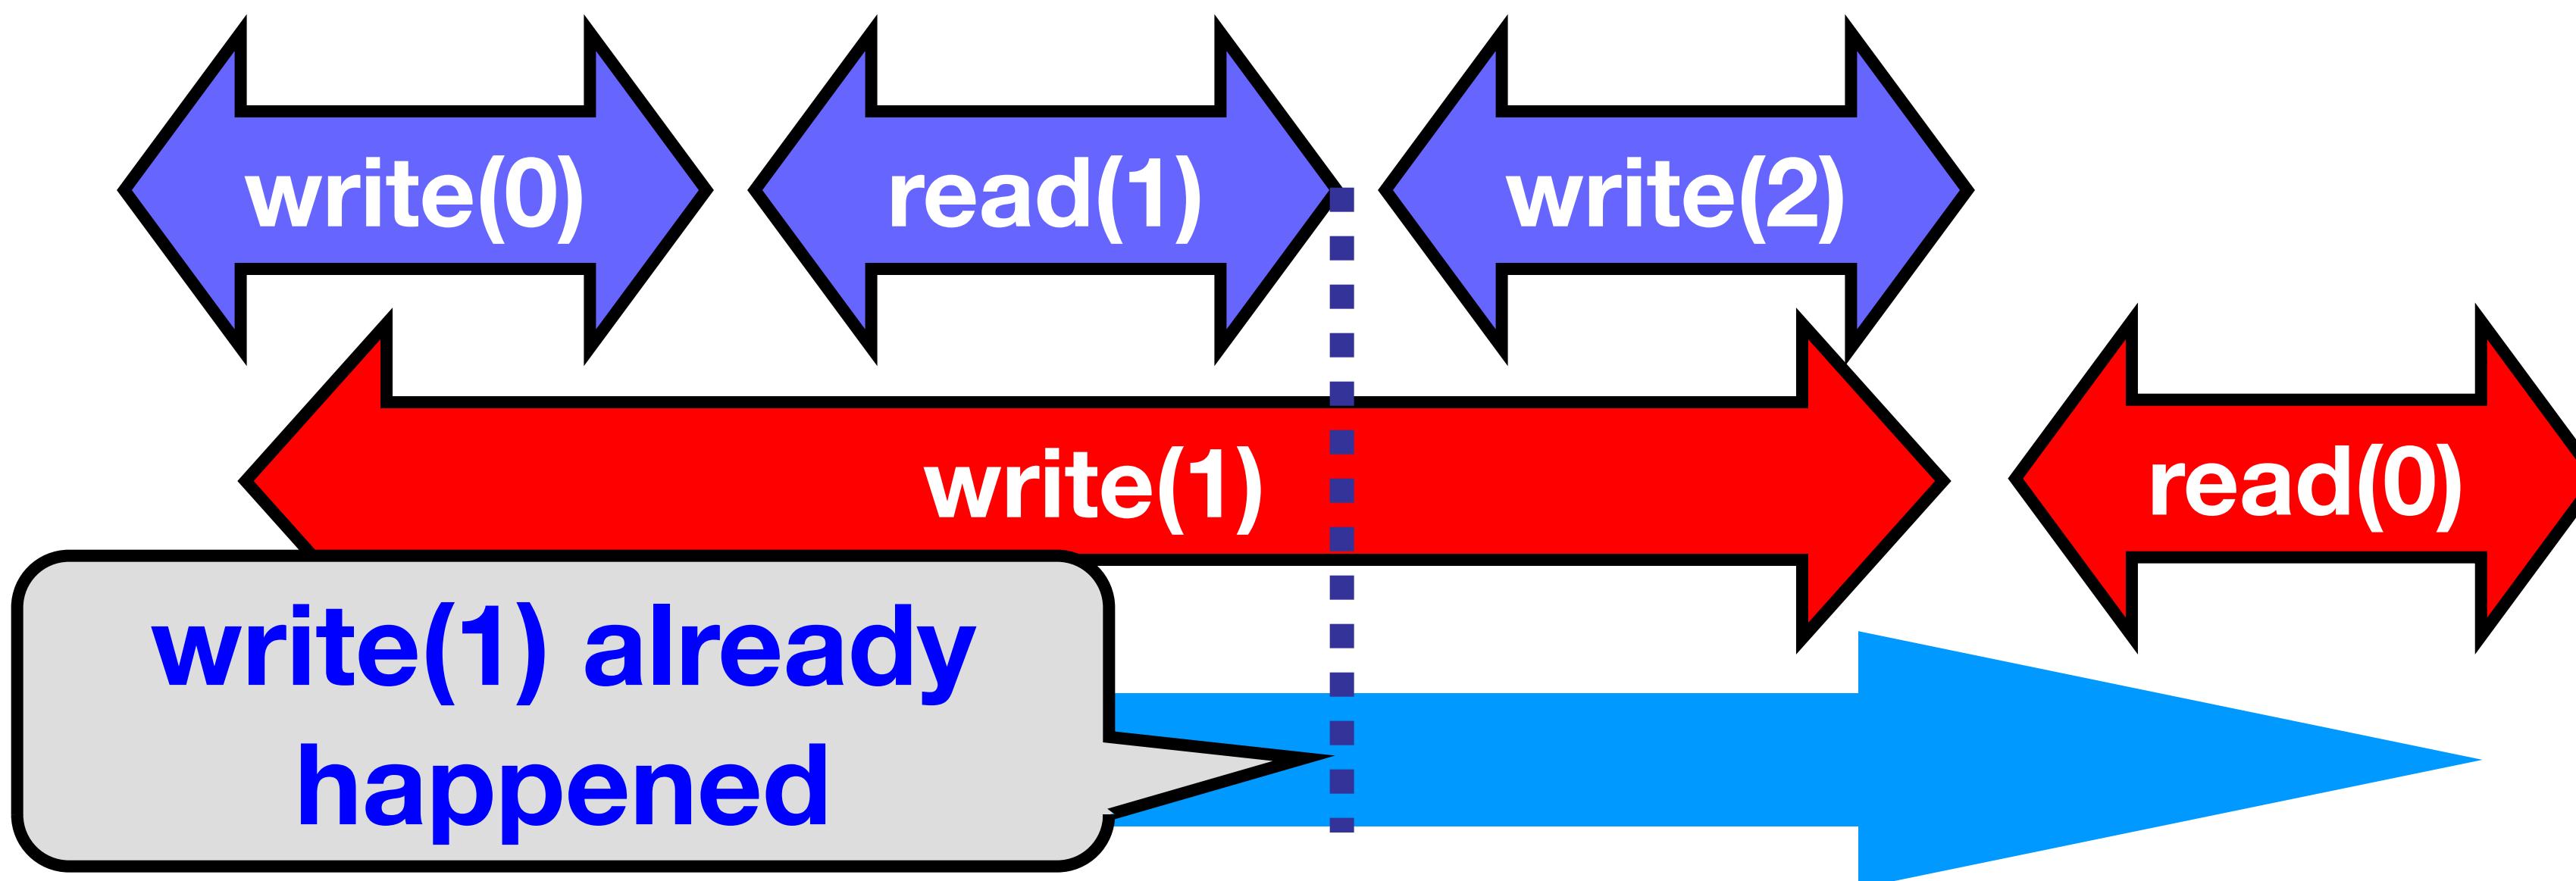
# Example




# Read/Write Register Example




# Read/Write Register Example




# Read/Write Register Example



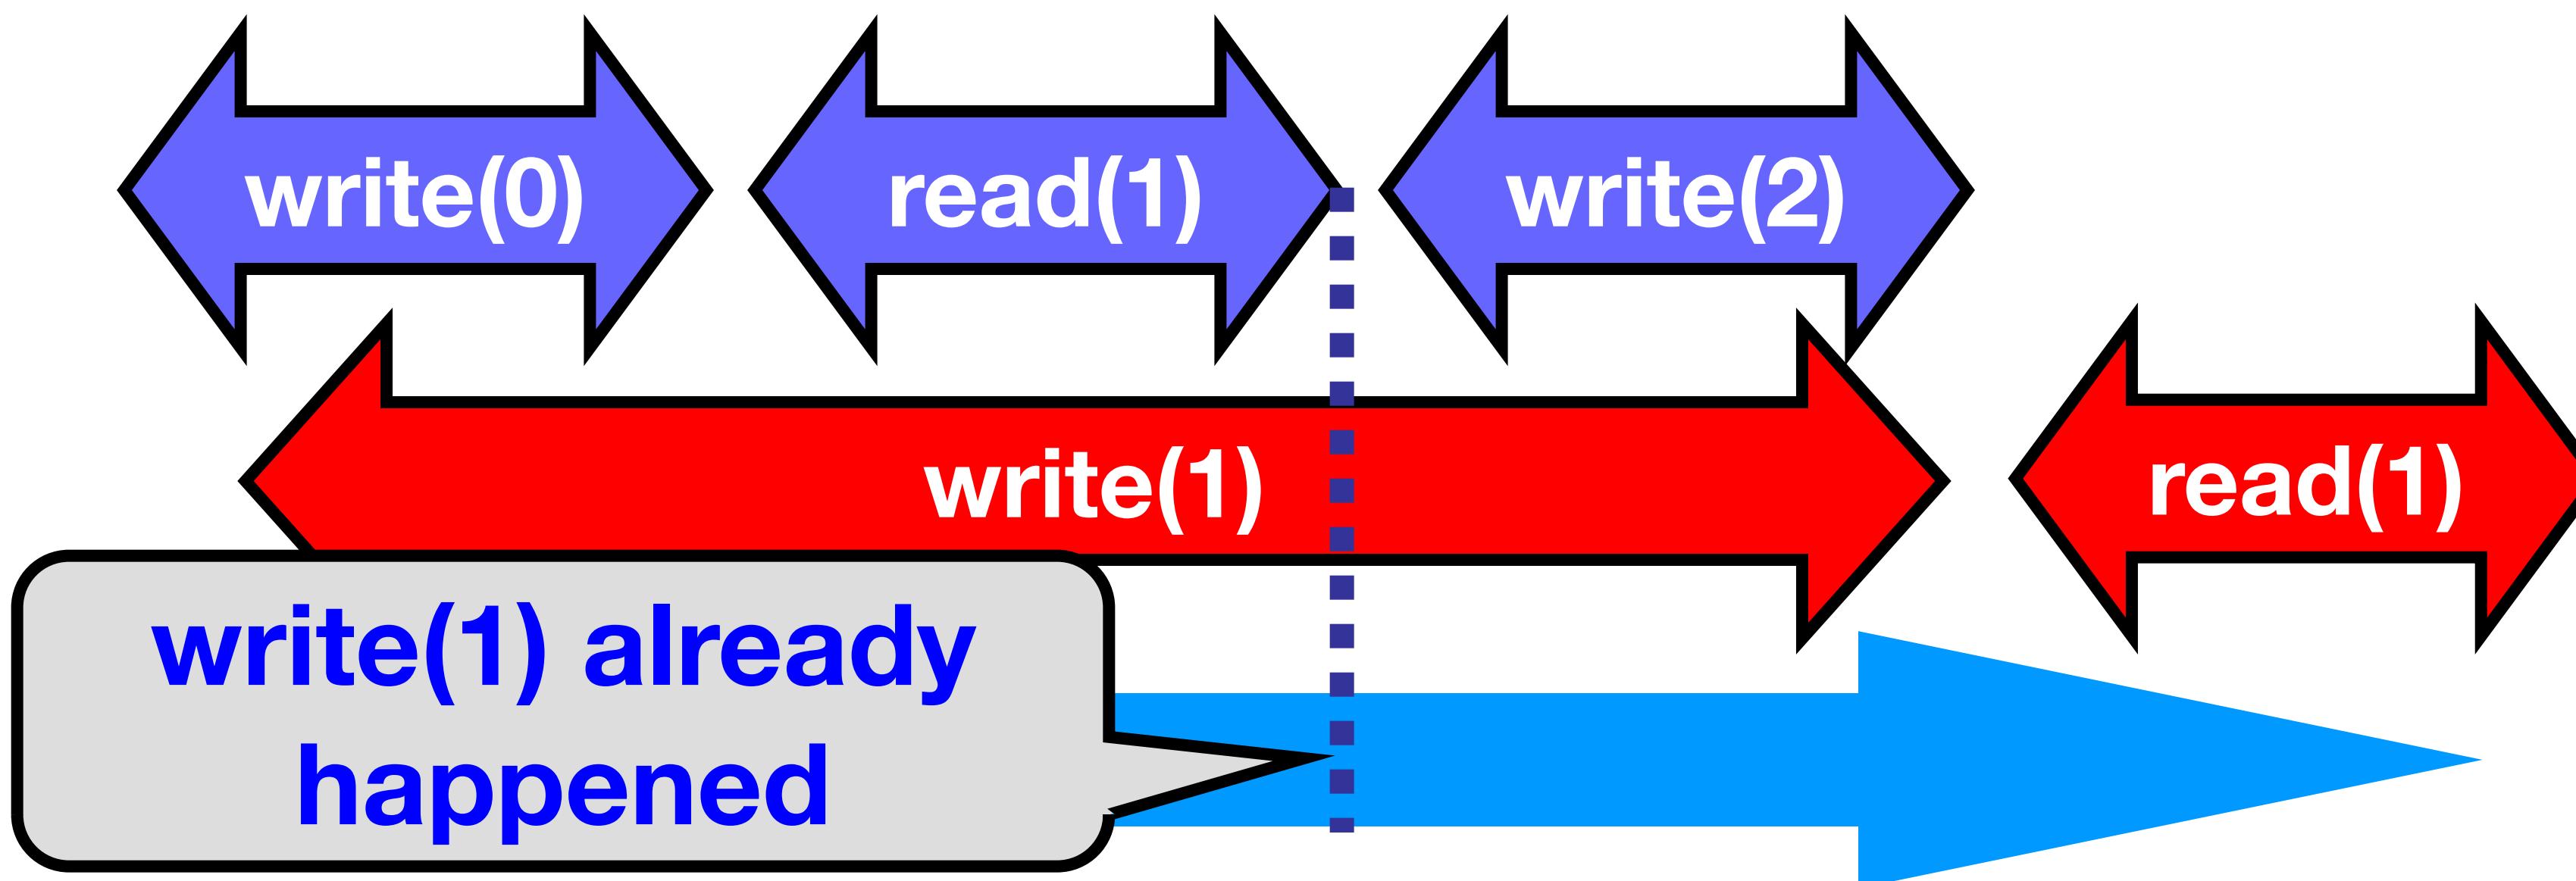
# Read/Write Register Example



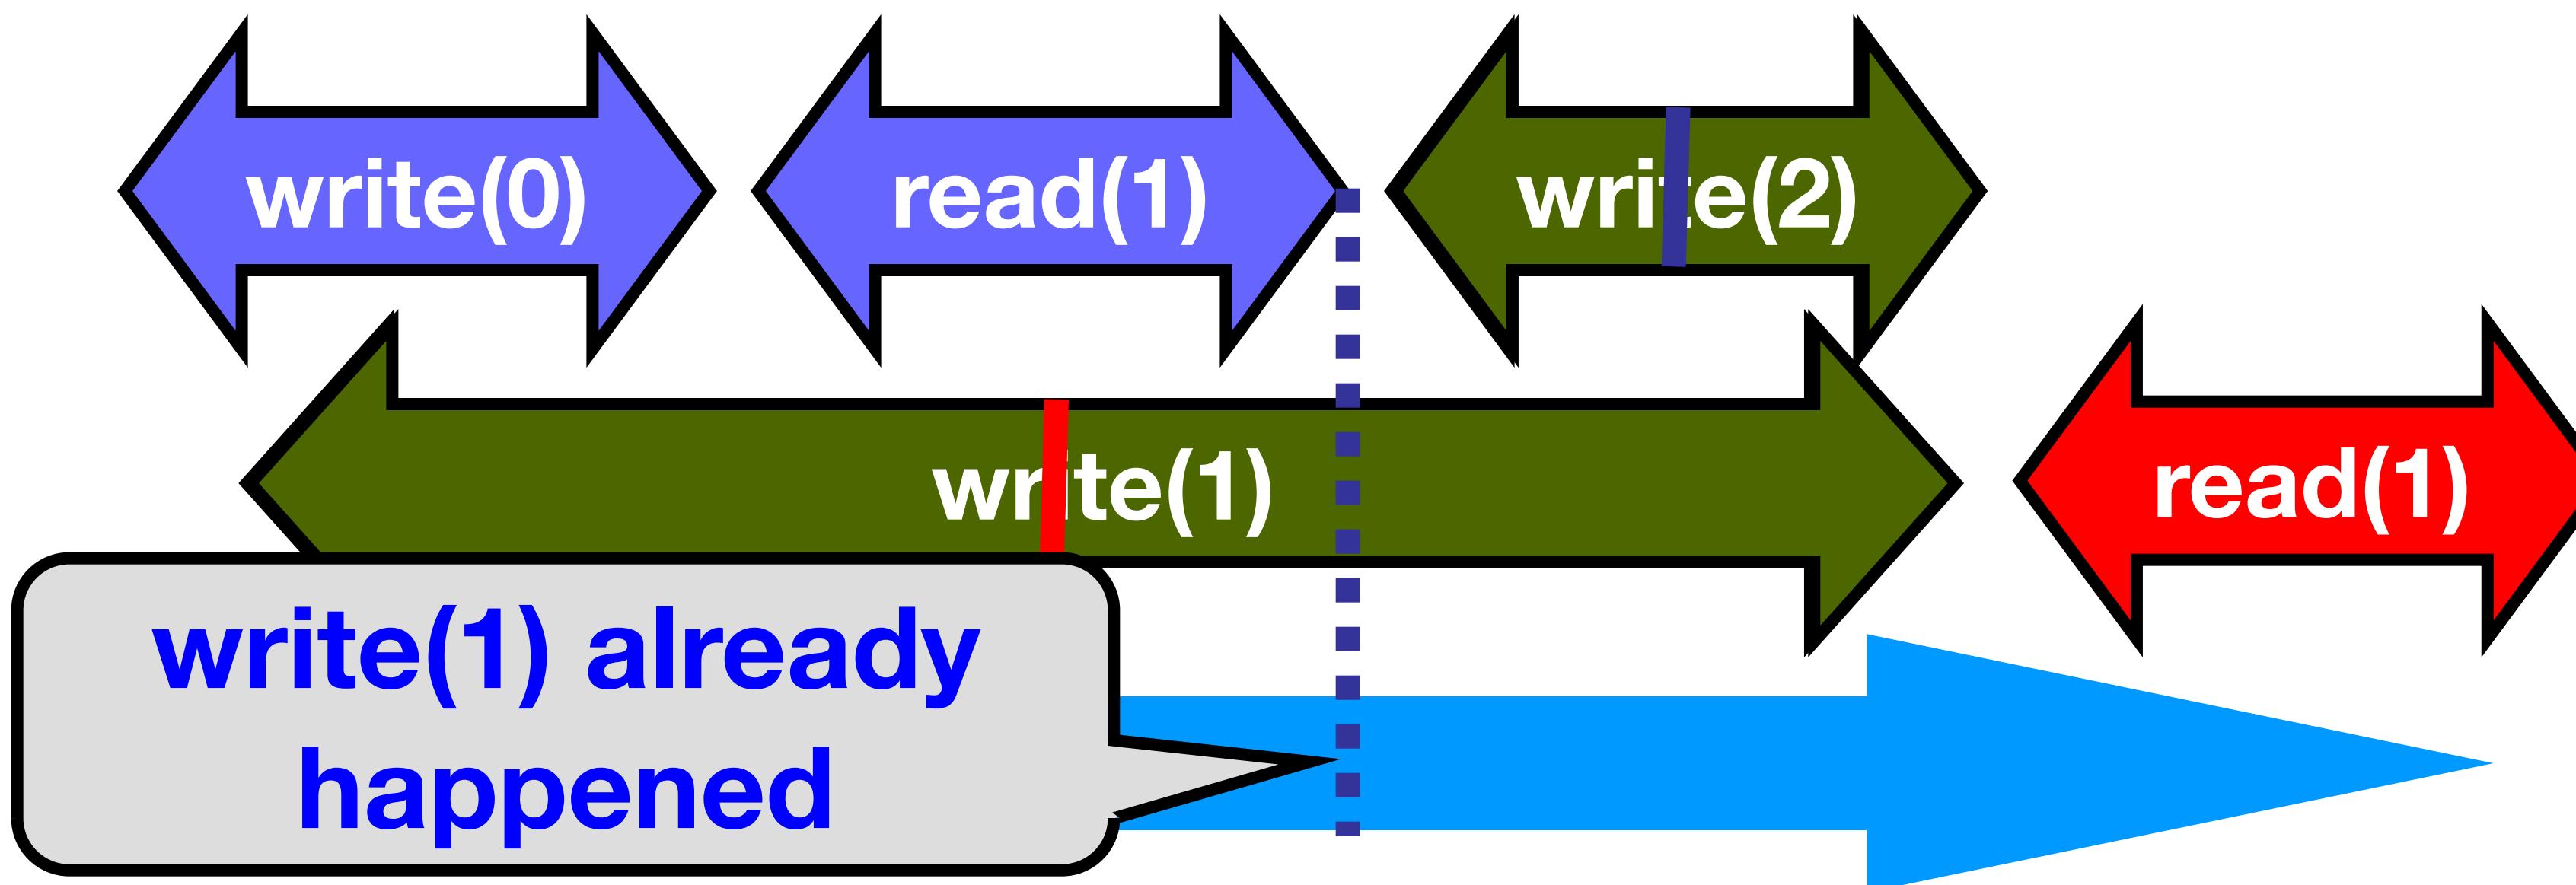
# Read/Write Register Example




# Read/Write Register Example




# Read/Write Register Example



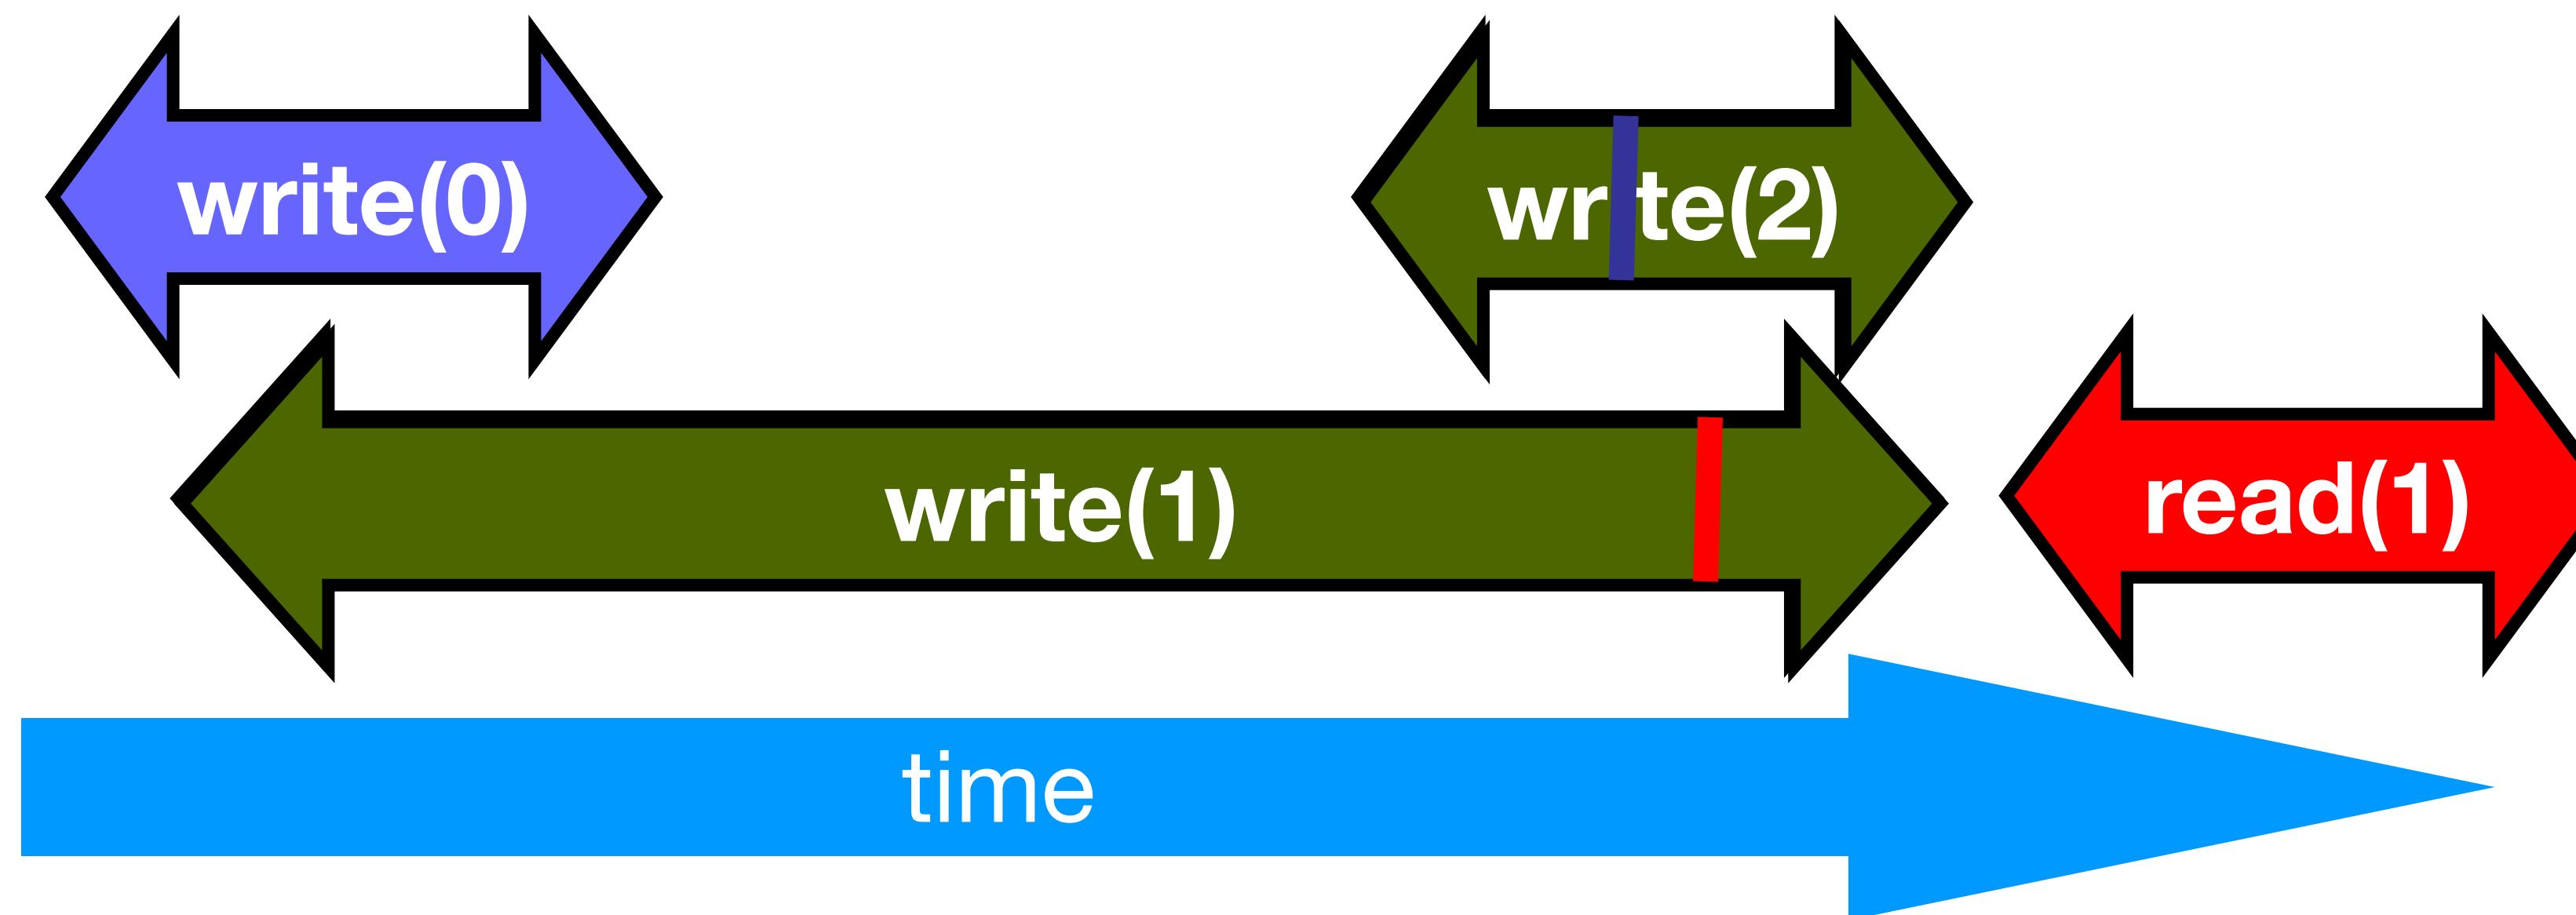
# Read/Write Register Example



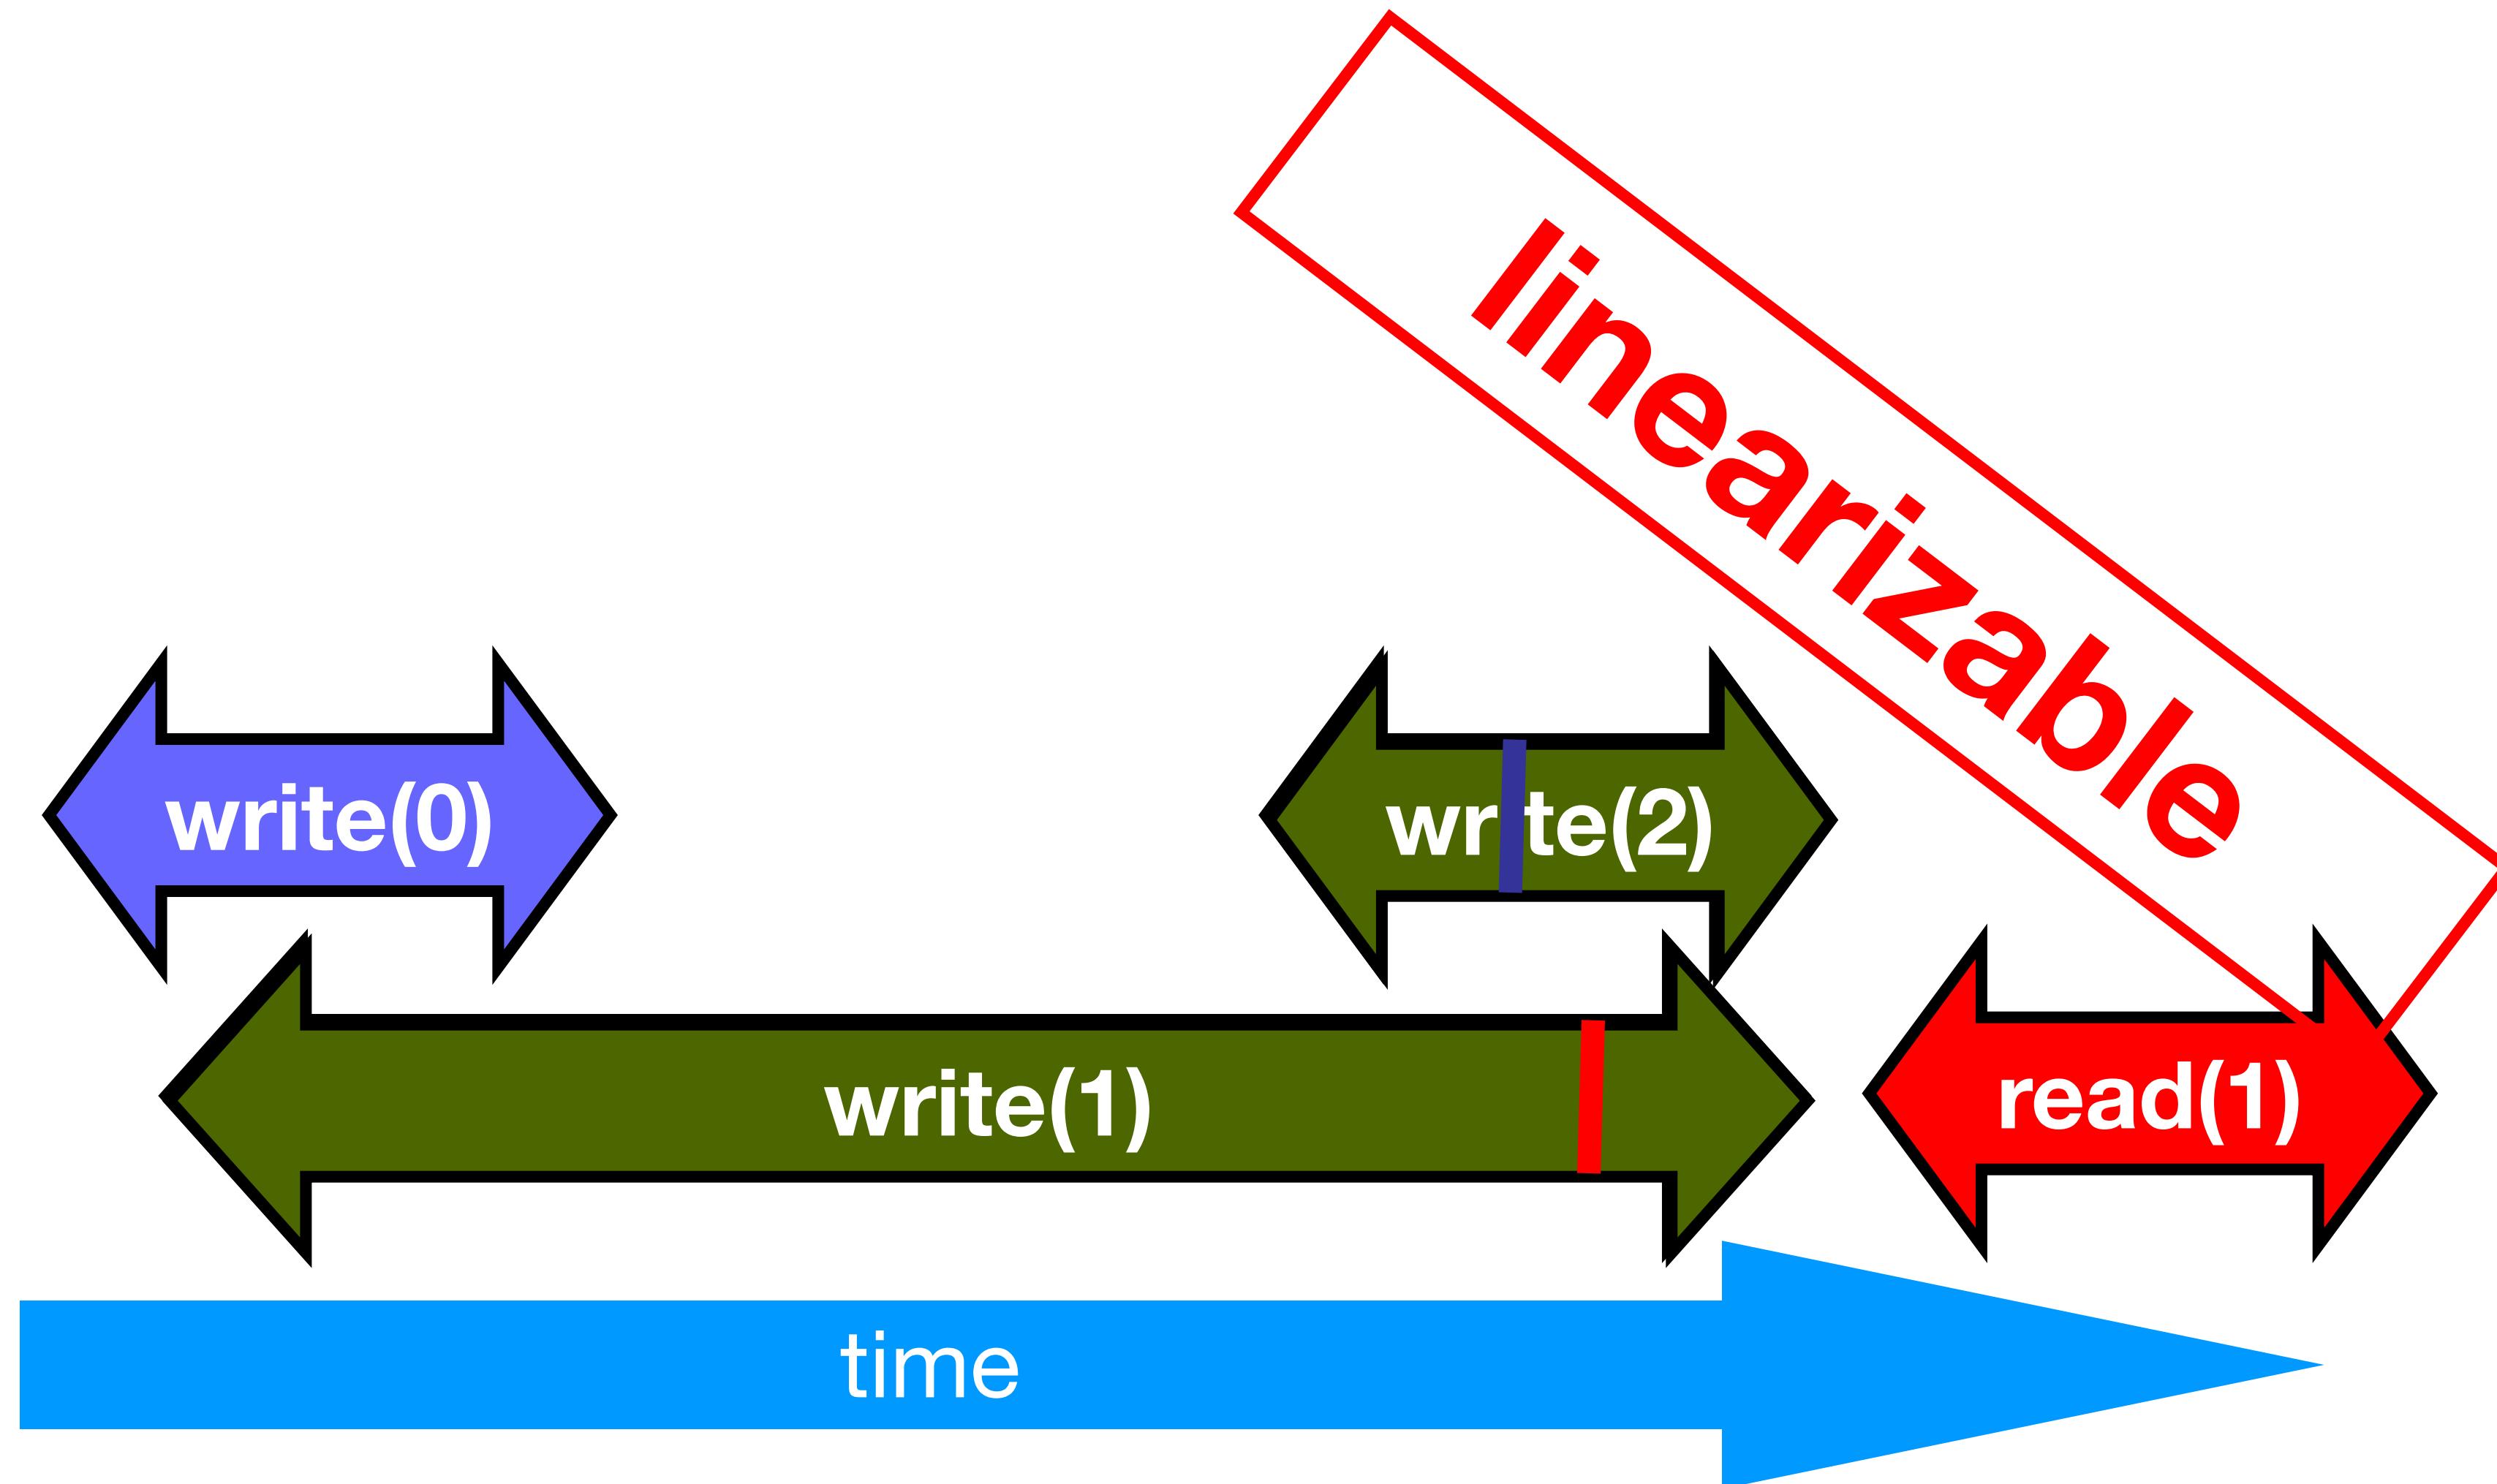
# Read/Write Register Example



# Read/Write Register Example



# Read/Write Register Example



# Talking About Executions

- Why?
  - Can't we specify the linearization point of each operation without describing an execution?
- Not Always
  - In some cases, linearization point ***depends on the execution***

# Formal Model of Executions

- Define precisely what we mean
  - Ambiguity is bad when intuition is weak
- Allow reasoning
  - Formal
  - But mostly informal
    - In the long run, actually, more important

# Split Method Calls into Two Events

- Invocation
  - method name & args
  - `q.enq(x)`
- Response
  - result or exception
  - `q.enq(x)` returns **void**
  - `q.deq()` returns `x`
  - `q.deq()` throws `empty`


# Split Method Calls into Two Events

- Invocation
  - method name & args
  - `q.enq(x)`
- Response
  - result or exception
  - `q.enq(x)` returns **void**
  - `q.deq()` returns **x**
  - `q.deq()` throws **empty**
- Note that I'm following the convention of the book
  - Book uses OO
  - Code in this course uses FP
- Note that we're still reasoning using **objectivism**
- For the current discussion, distinction doesn't matter
  - `q.enq(x)` is read as **enq q x** in code
  - Returns **void** is read as returns **()**
  - Throws **empty** is read as raises **Empty**

# Invocation Notation

A q.enq(x)

# Invocation Notation




A red bracket is positioned above the text "A q.enq(x)". The bracket starts from the left side of the letter "A" and extends upwards and to the right, ending under the closing parenthesis "(x)".

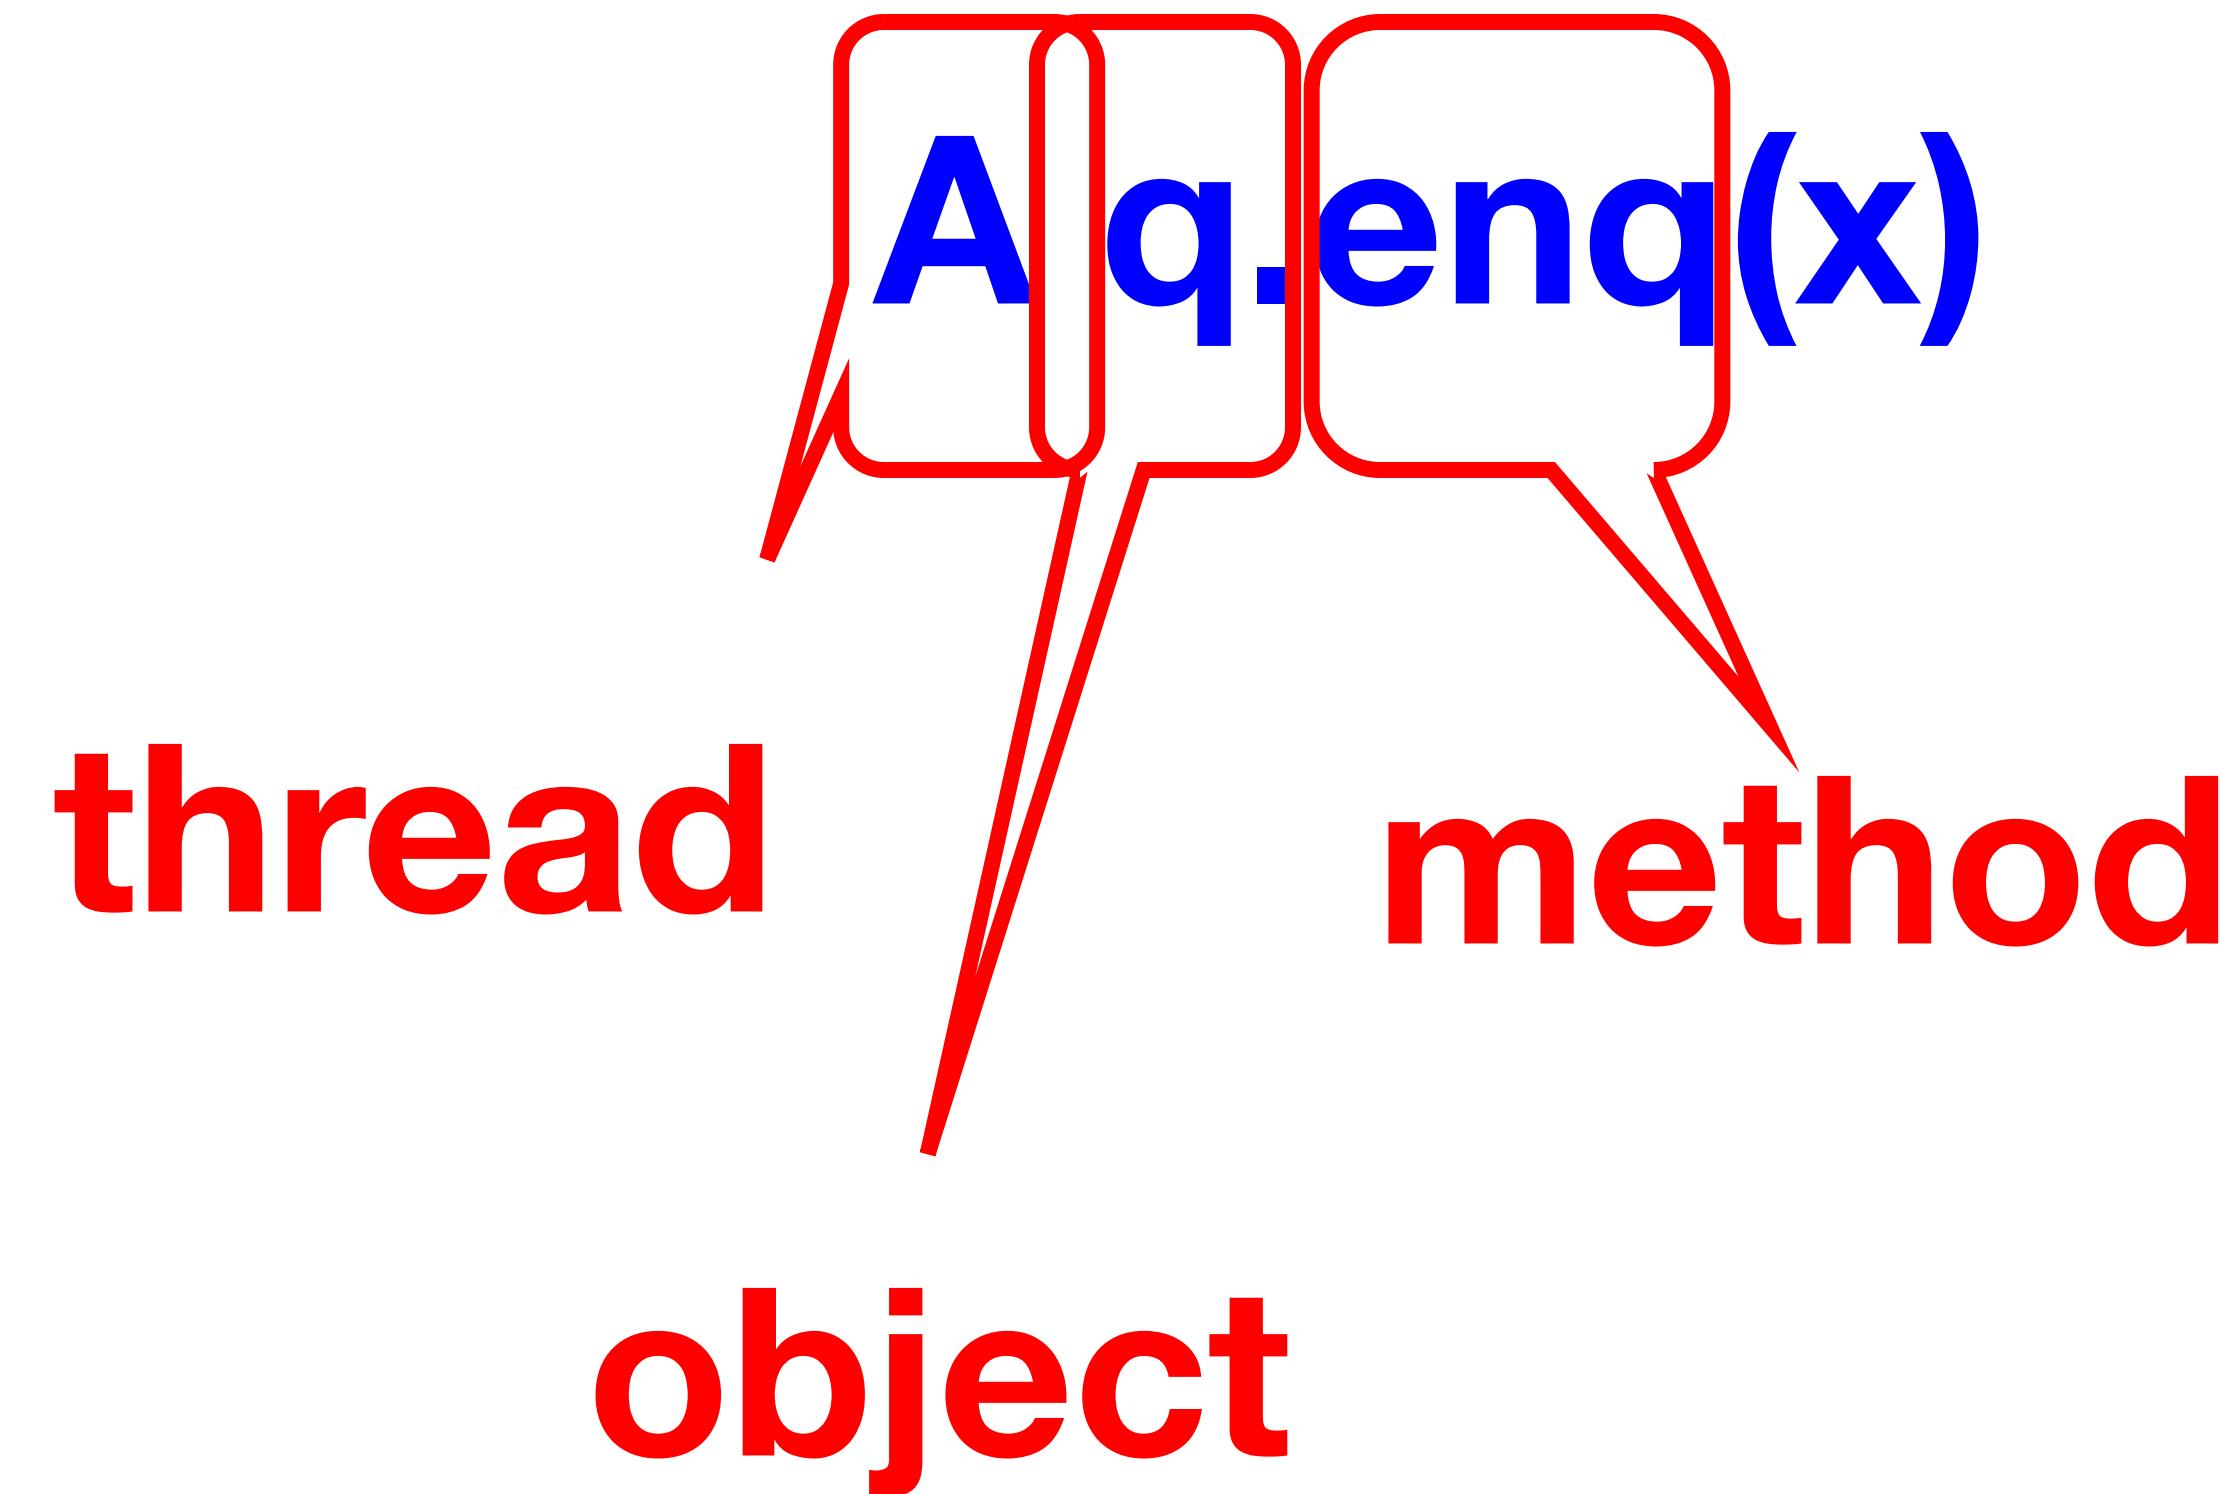
A q.enq(x)

thread

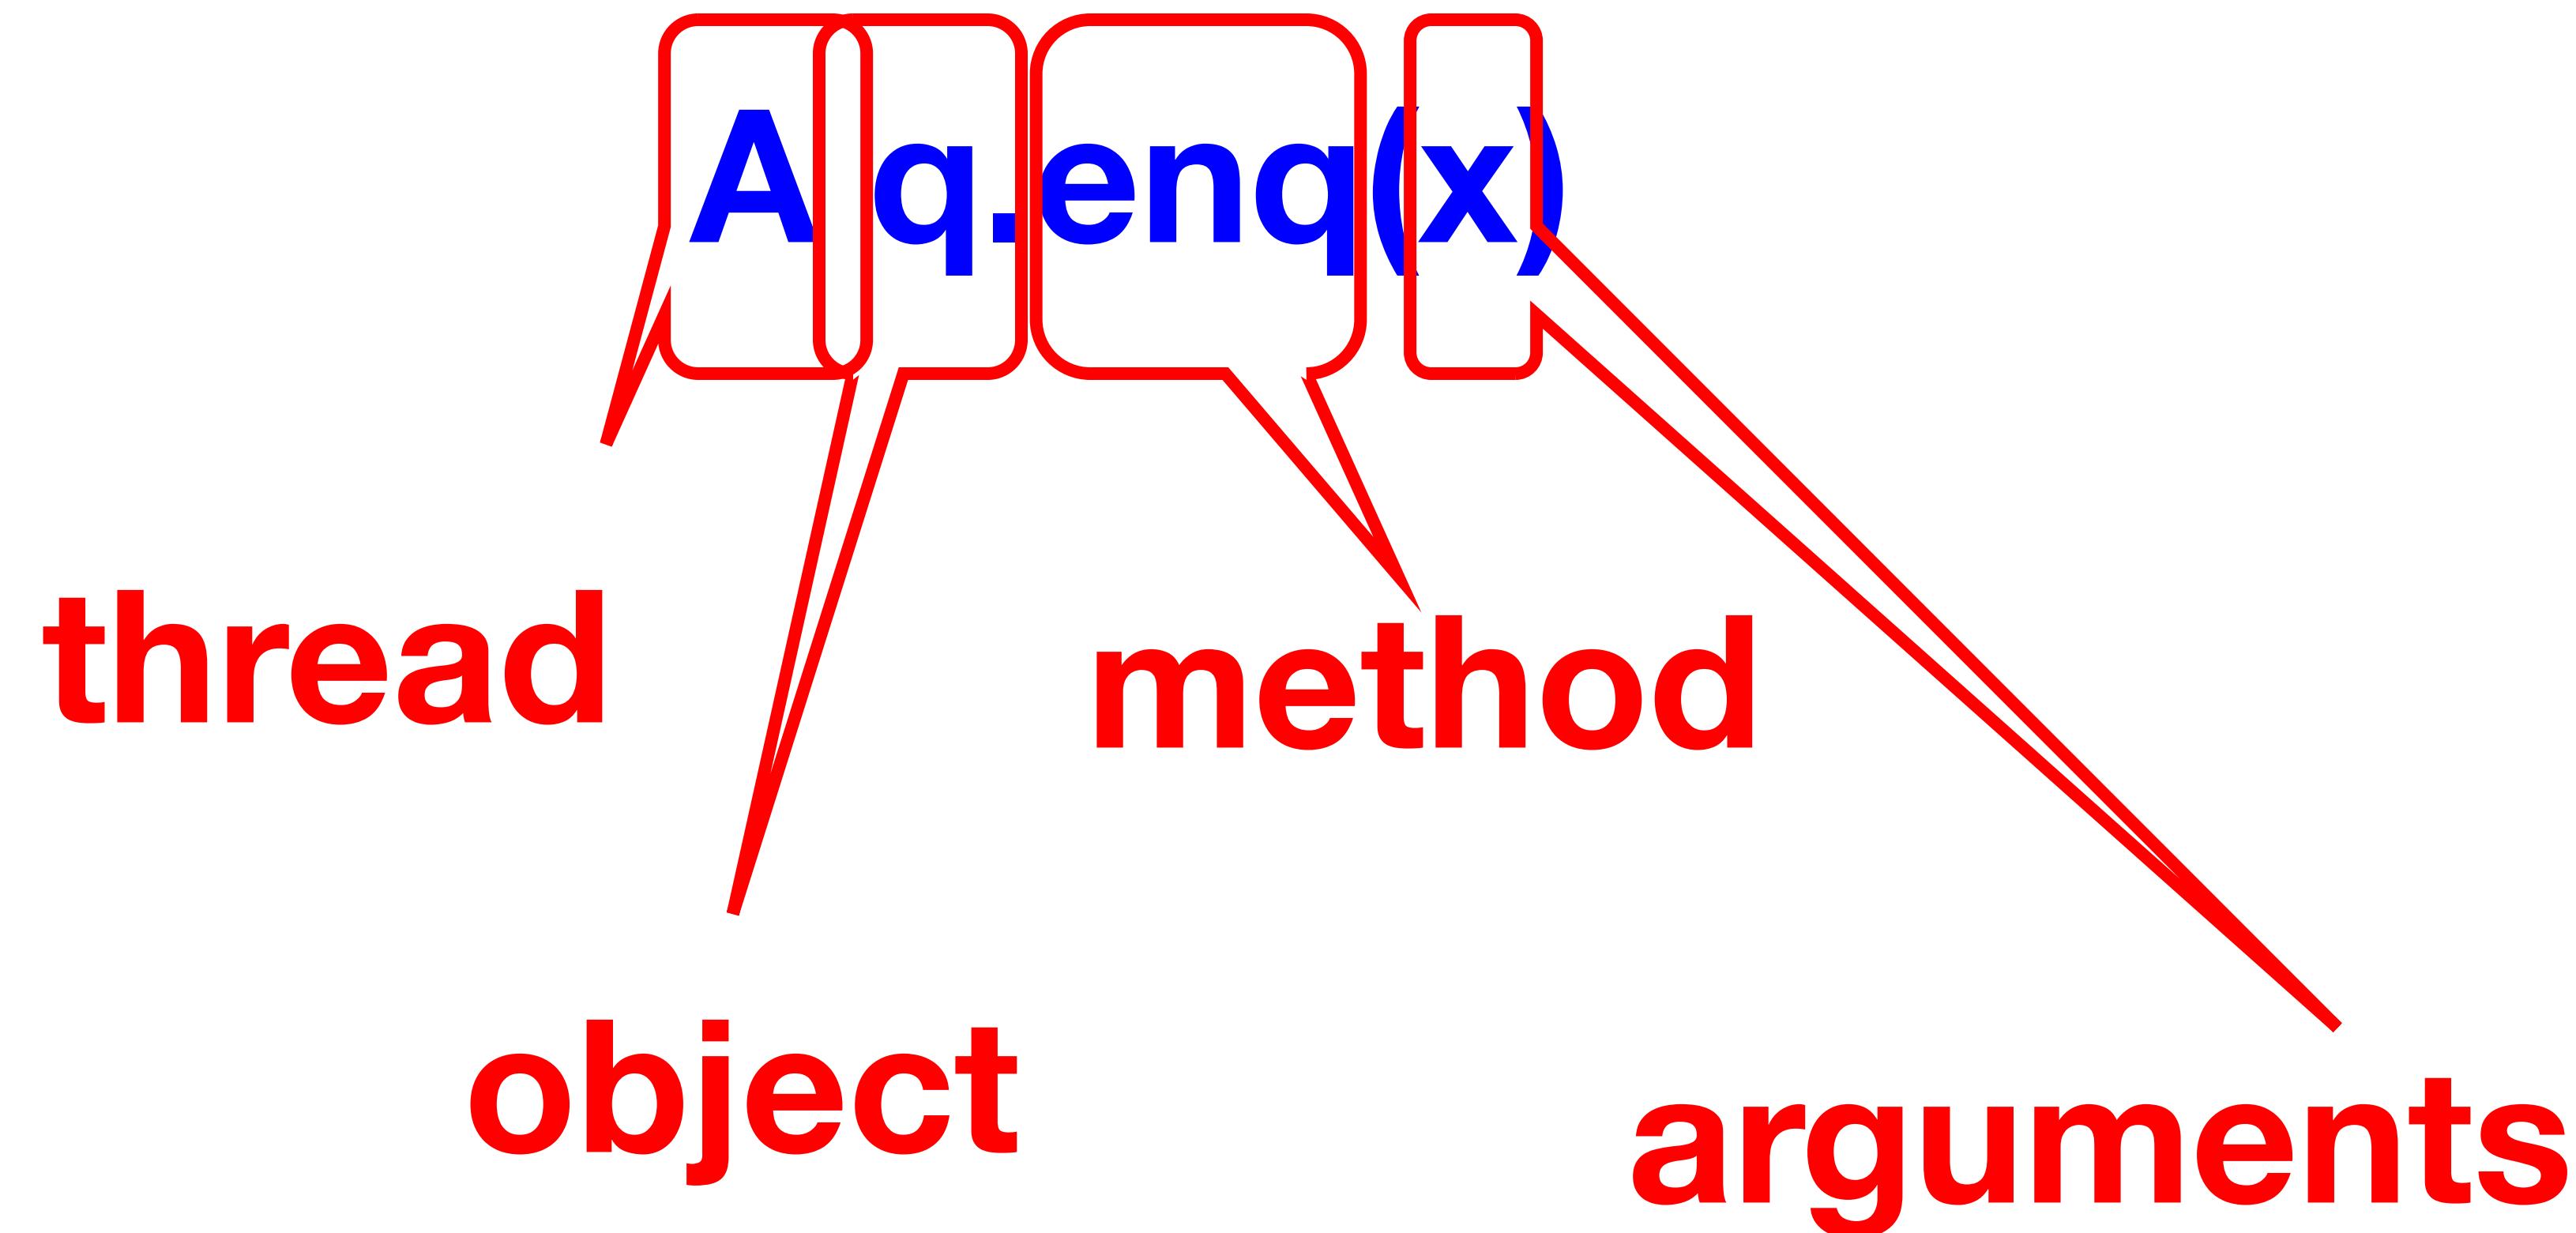
# Invocation Notation



# Invocation Notation



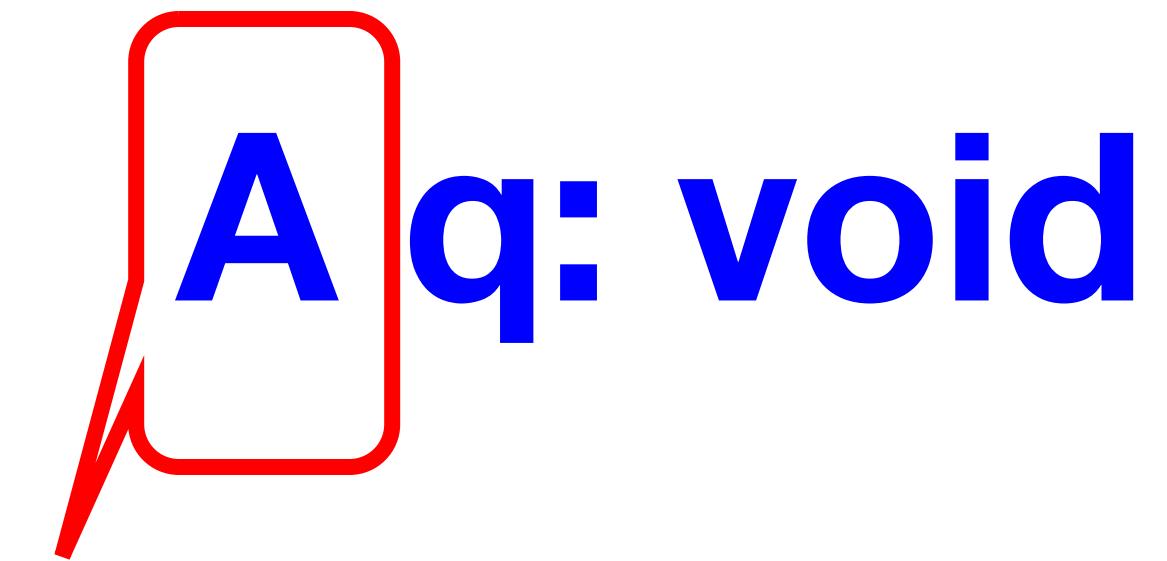
# Invocation Notation



# Response Notation

A q: void

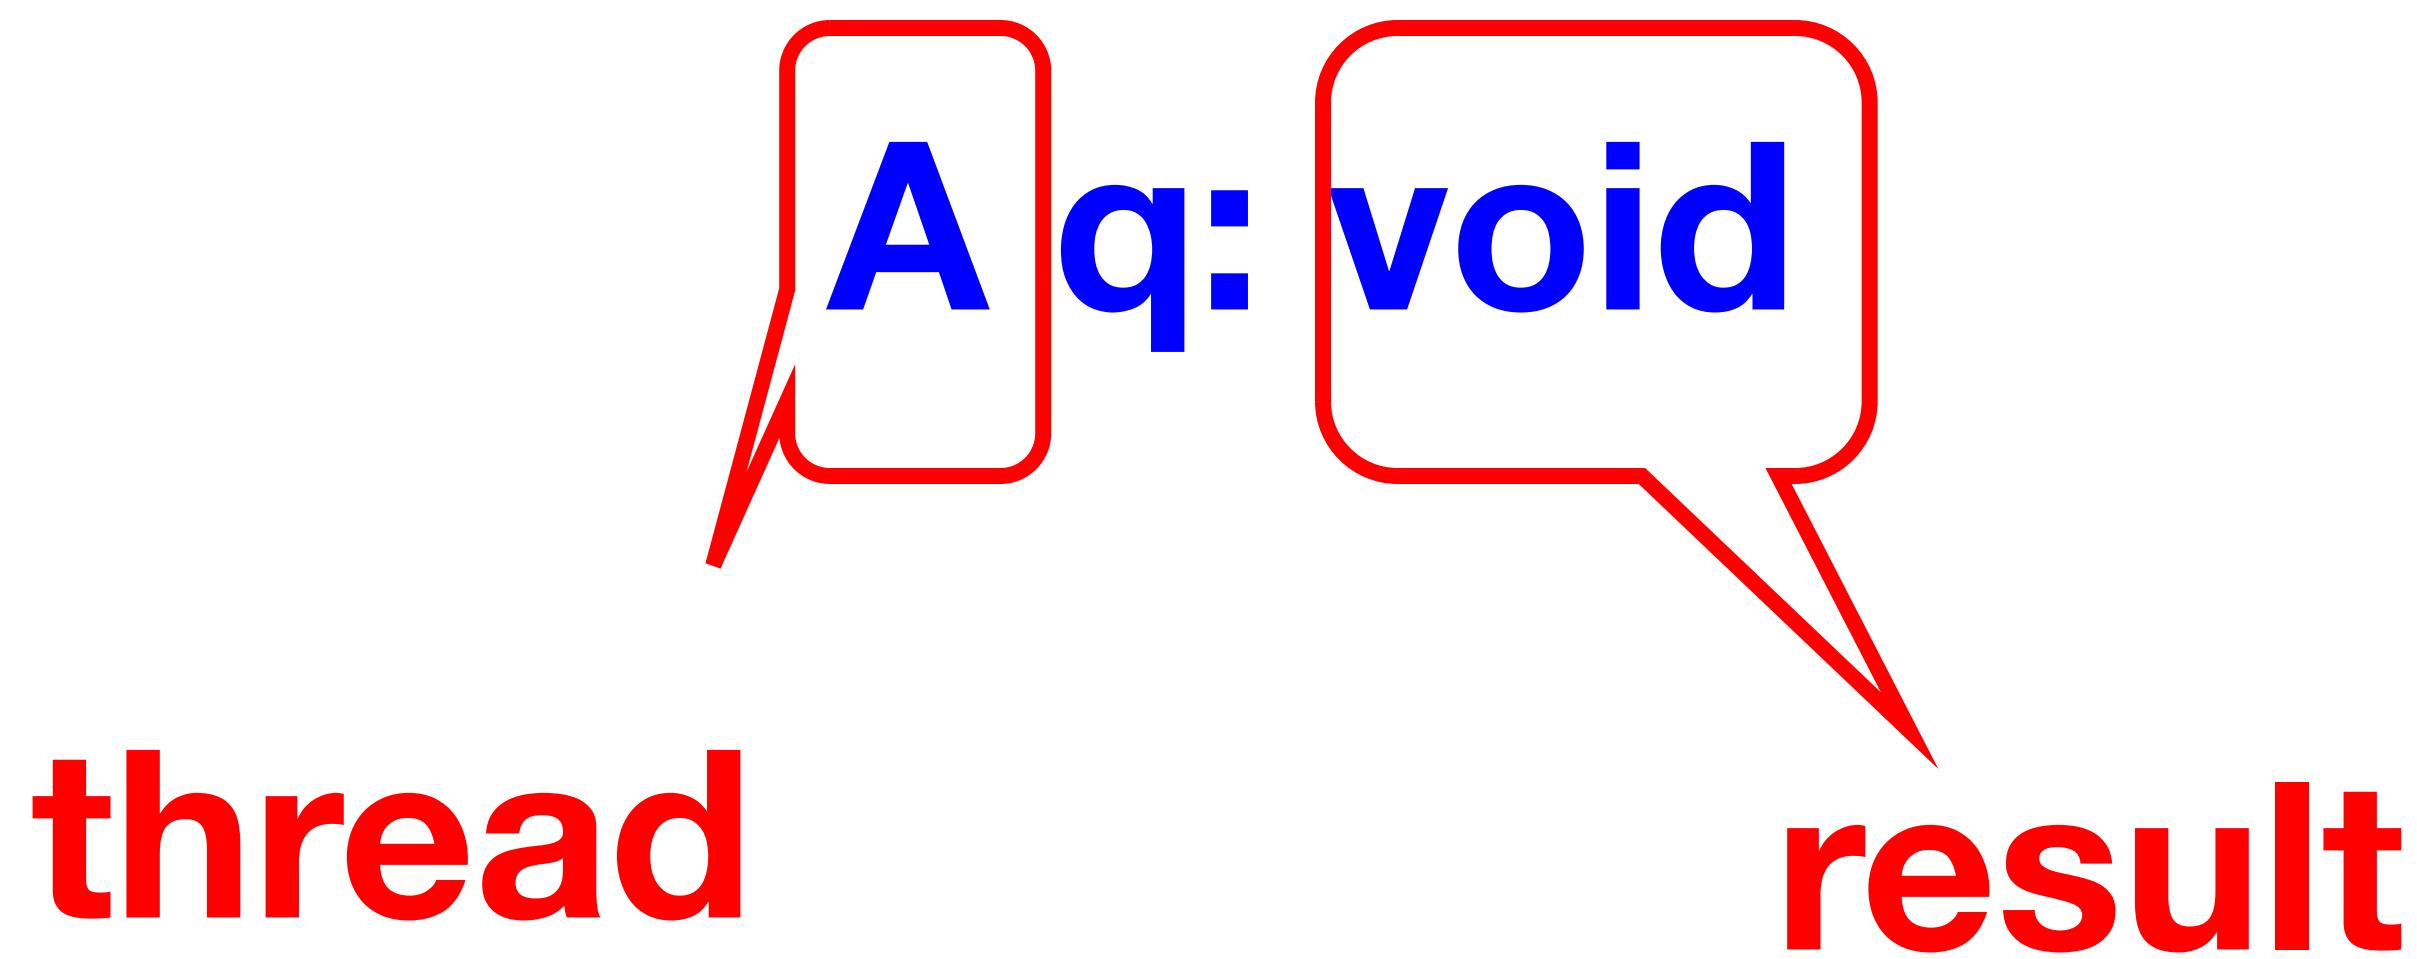
# Response Notation




**q: void**

**thread**

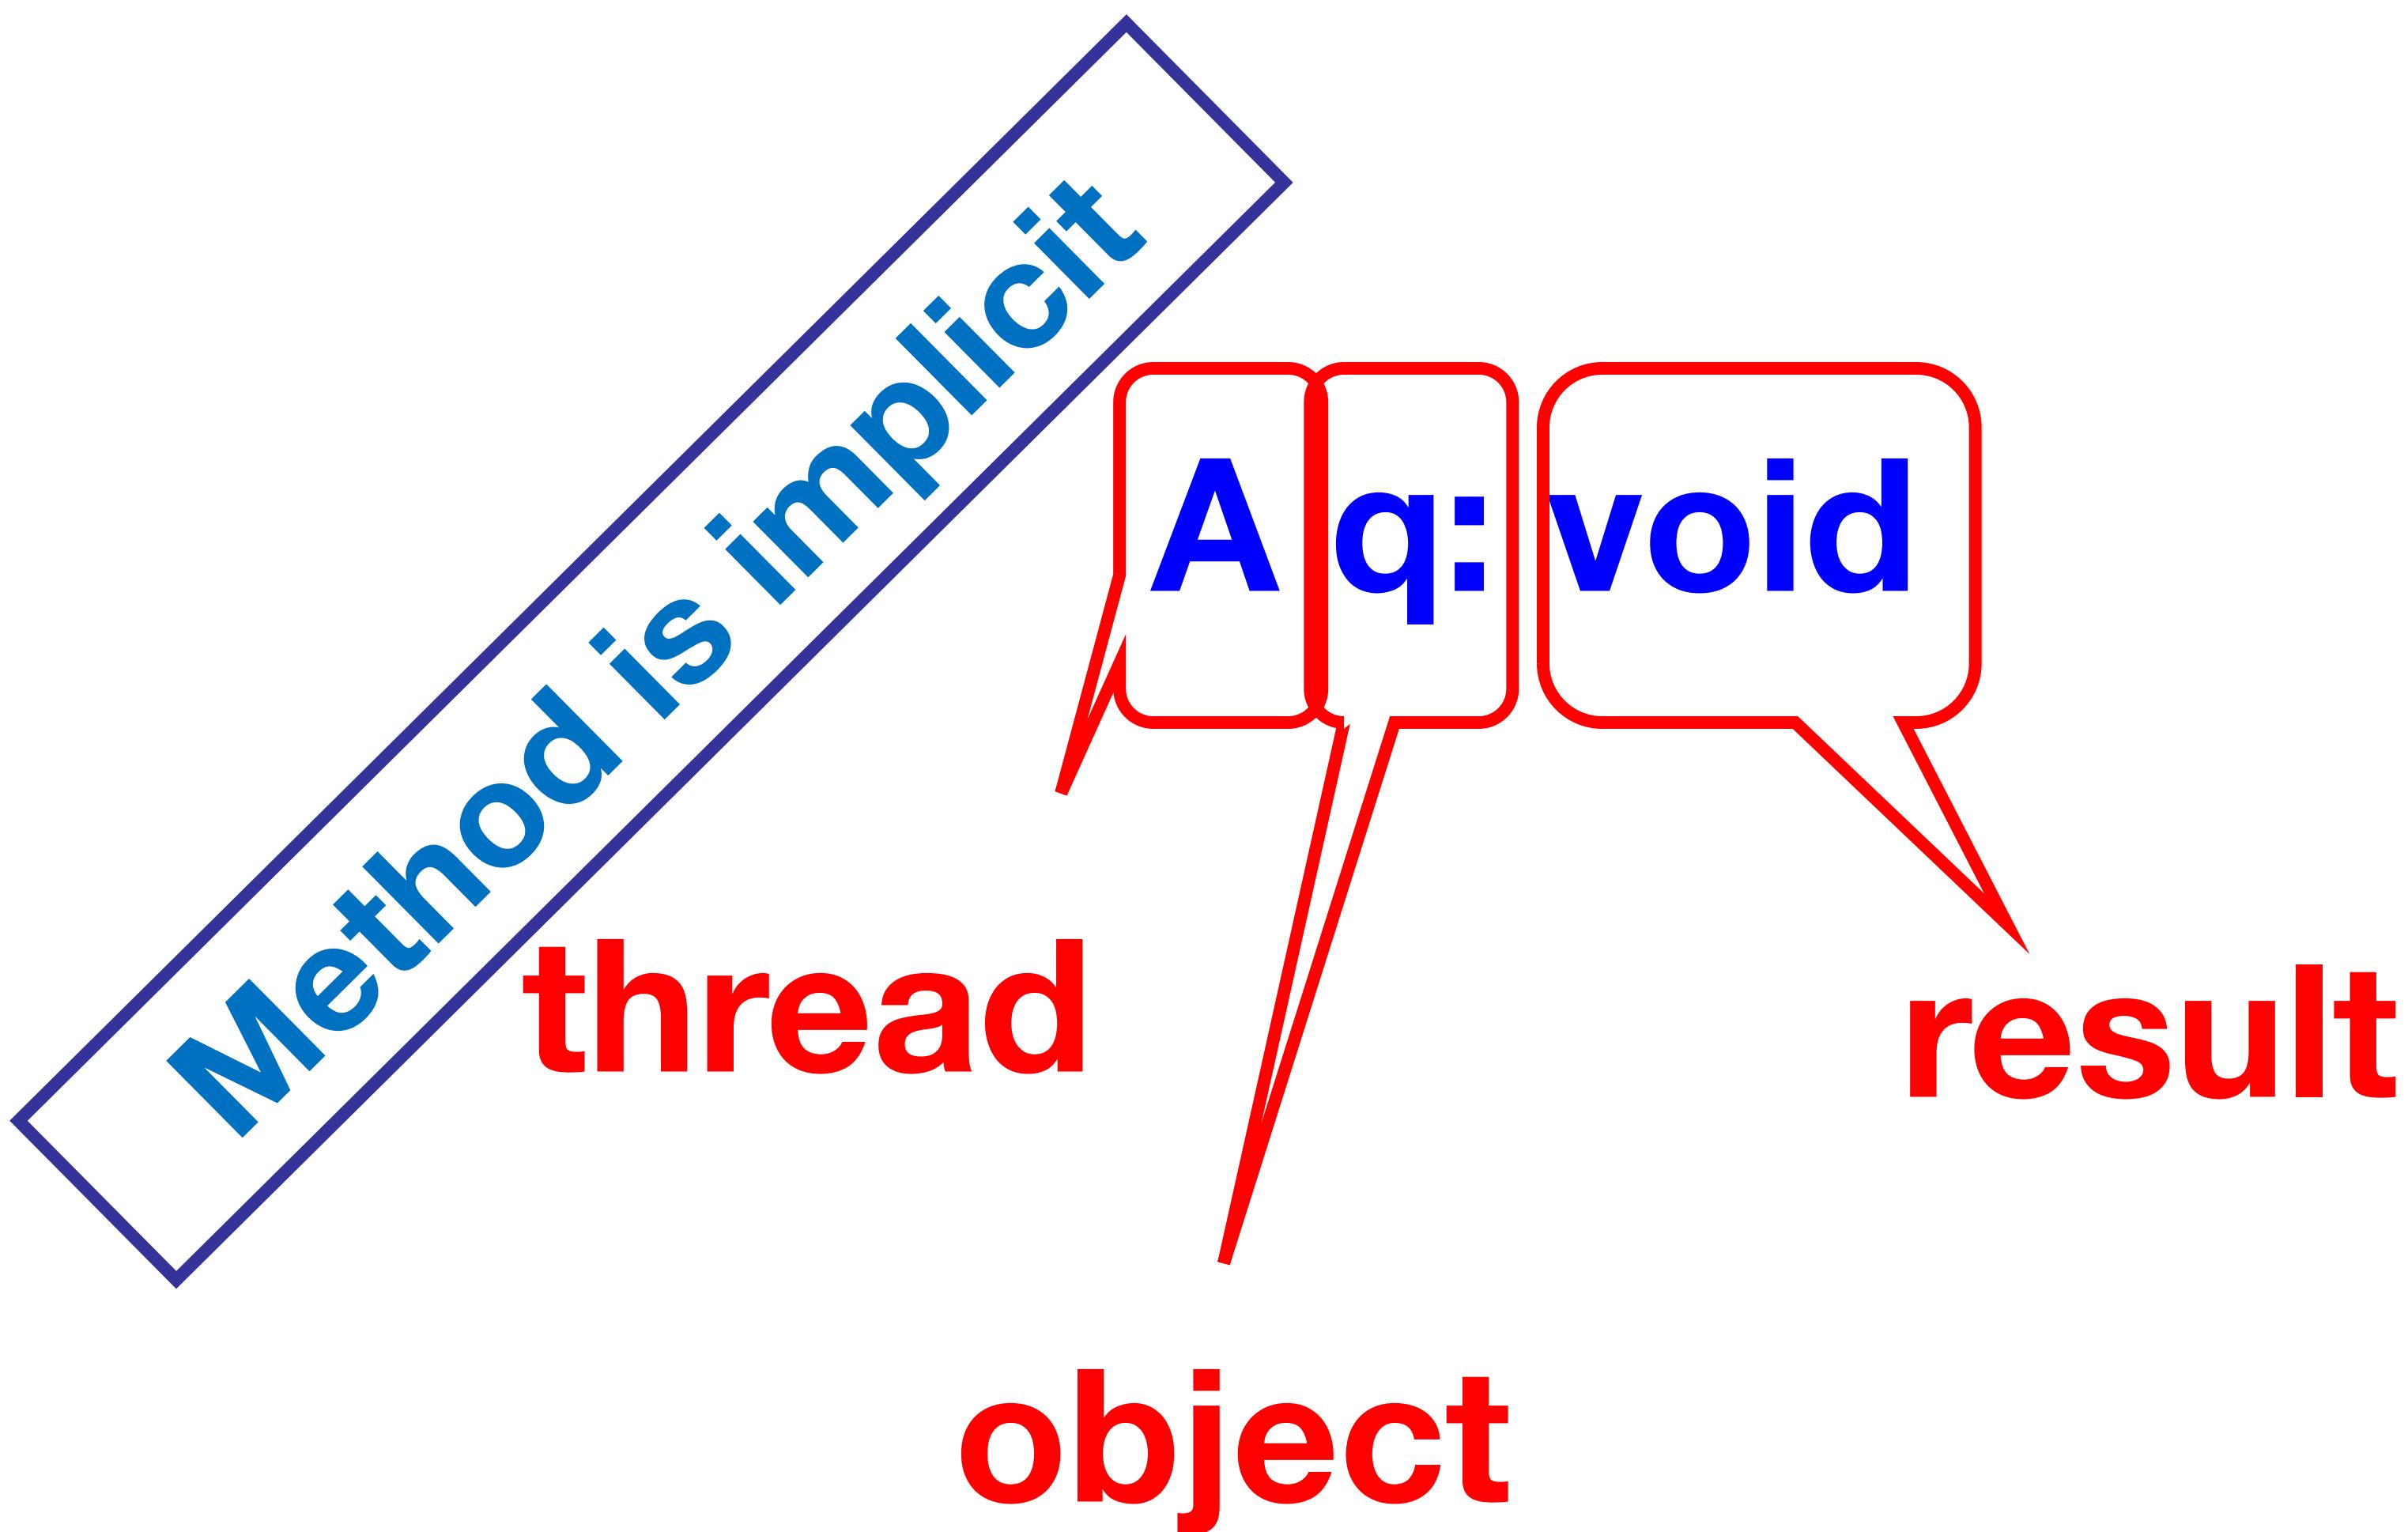
# Response Notation



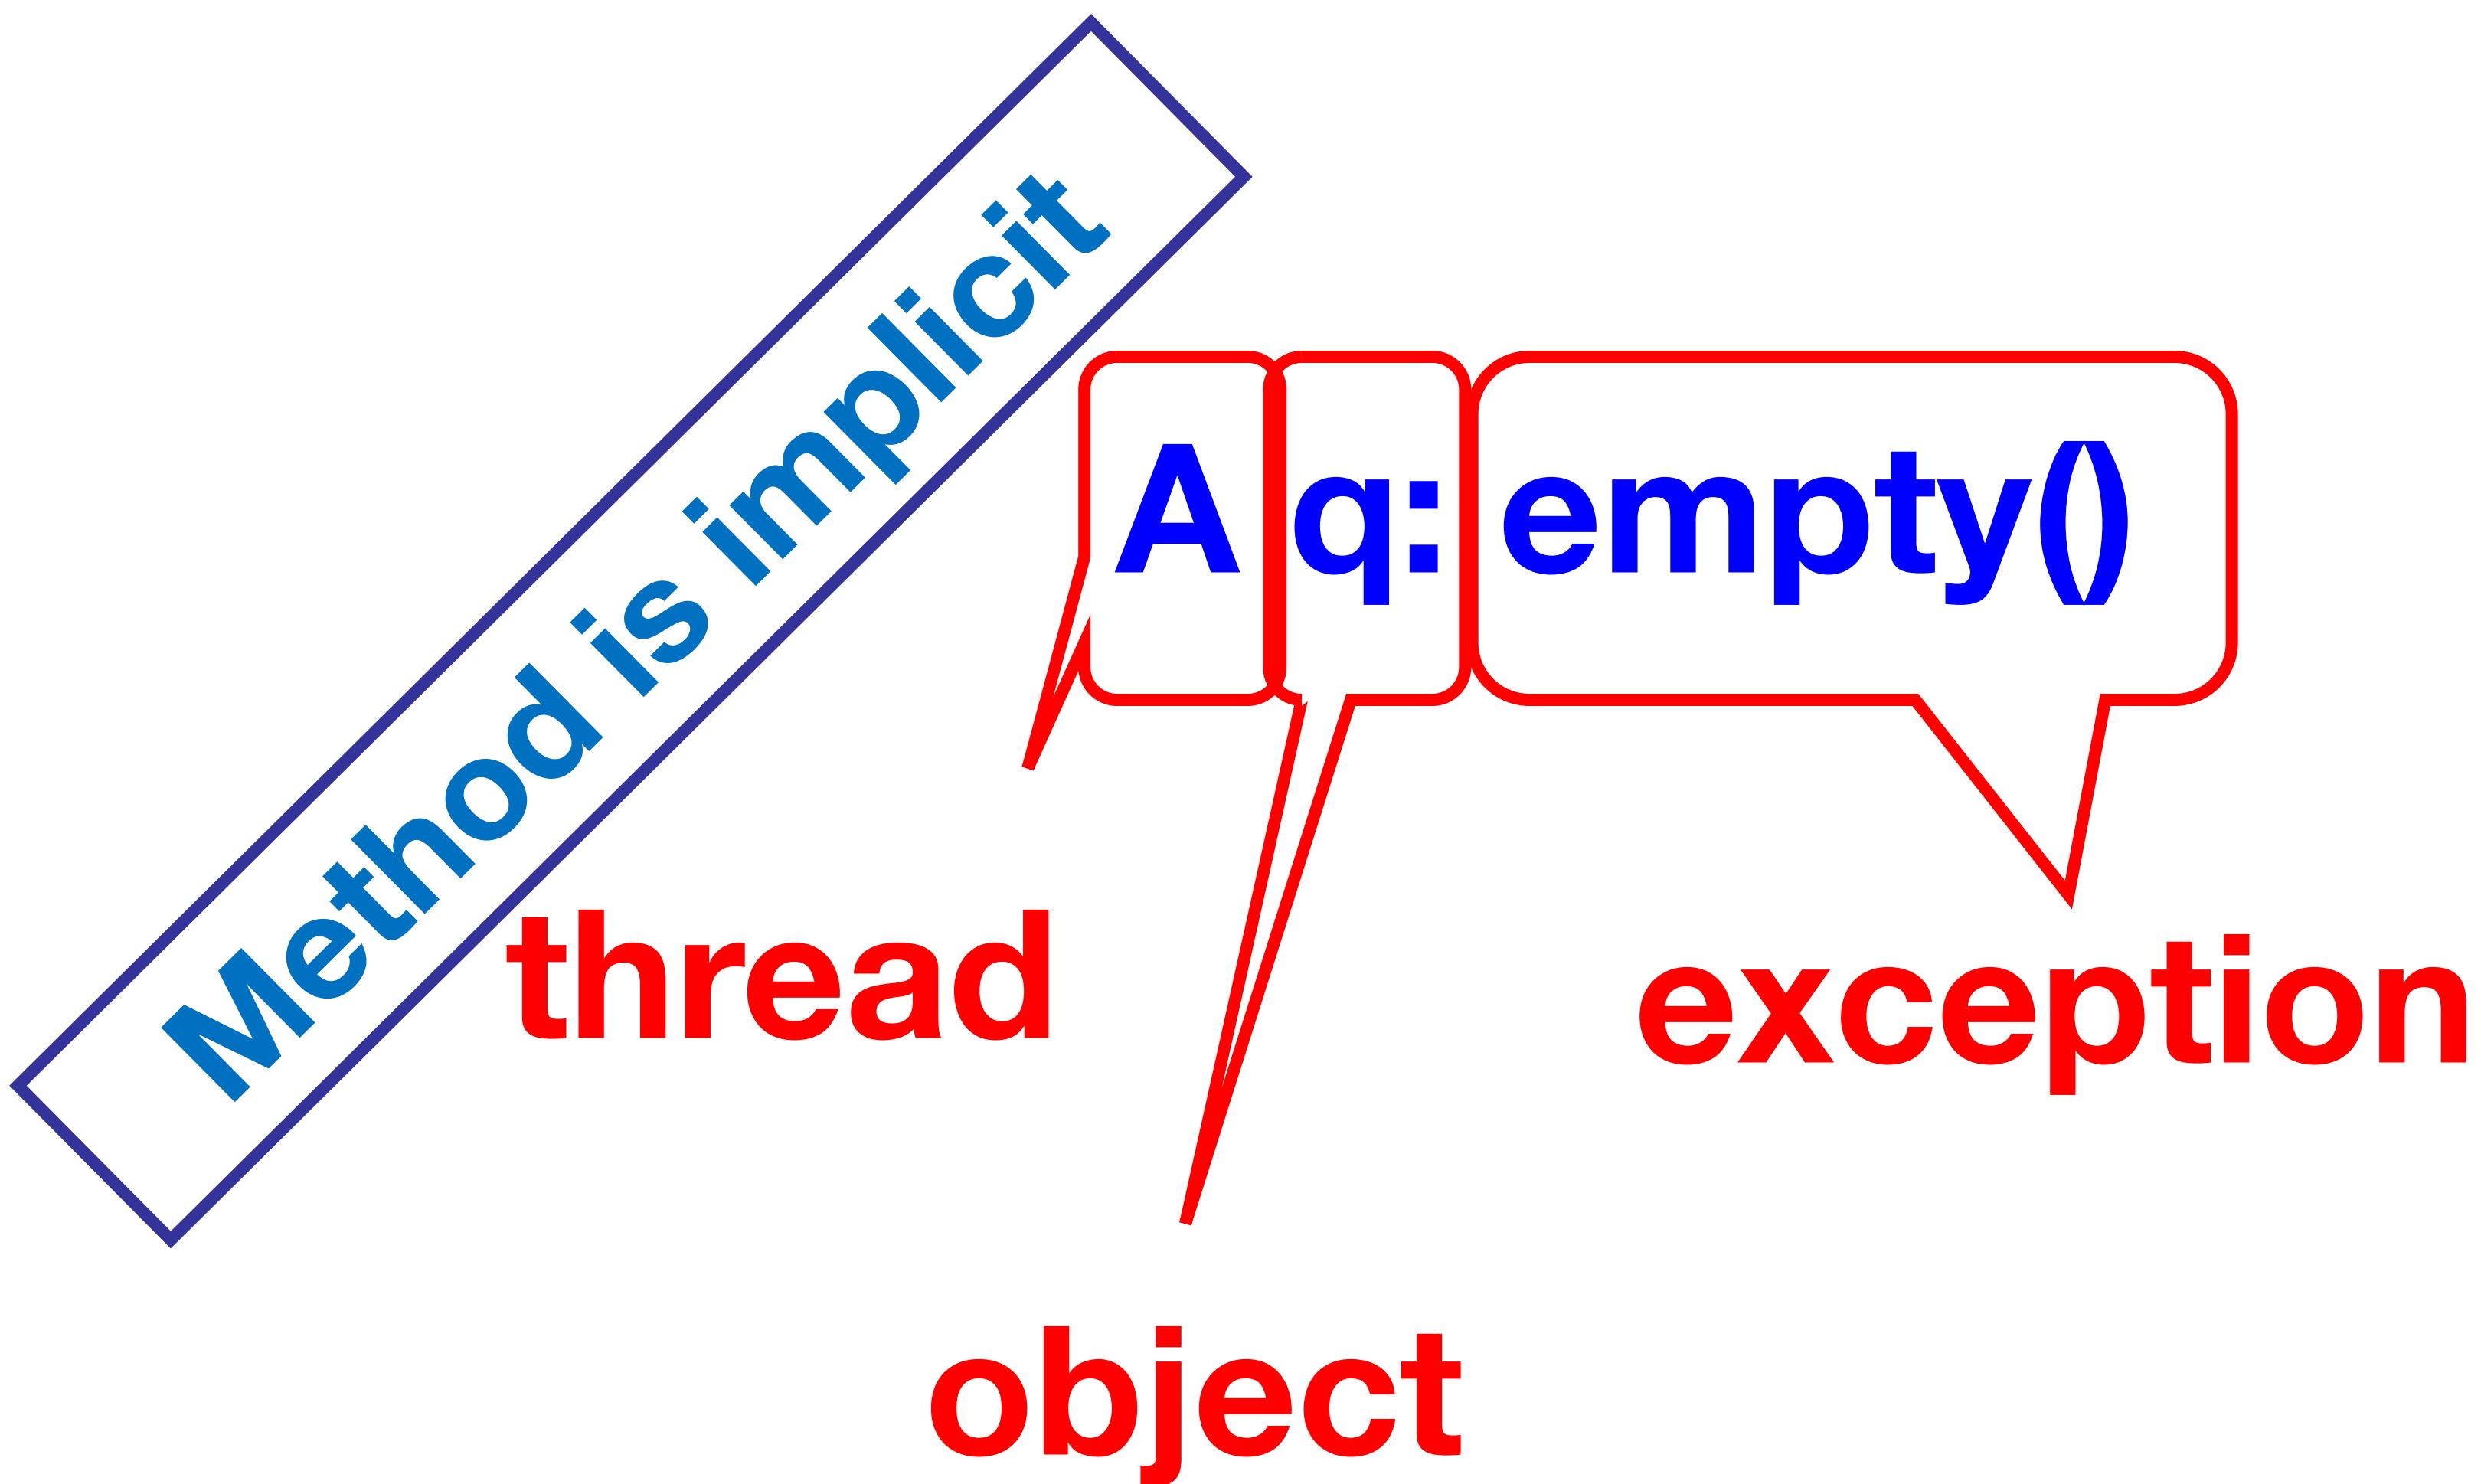
# Response Notation



# Response Notation



# Response Notation




# History – Describing an execution

A q.enq(3)  
A q:void  
A q.enq(5)  
**H =** B p.enq(4)  
B p:void  
B q.deq()  
B q:3

**Sequence of  
invocations and  
responses**

# History – Describing an execution

- Invocation & response *match* if



# Object Projections

$H =$

A	q.enq(3)
A	q:void
B	p.enq(4)
B	p:void
B	q.deq()
B	q:3

# Object Projections

A q.enq(3)  
A q:void

H|q =

B q.deq()  
B q:3

# Thread Projections

$H =$

A	q . enq(3)
A	q : void
B	p . enq(4)
B	p : void
B	q . deq()
B	q : 3

# Thread Projections

$$H|B = \begin{array}{l} B \ p.\text{enq}(4) \\ B \ p:\text{void} \\ B \ q.\text{deq}() \\ B \ q:3 \end{array}$$

# Complete Subhistory

$H =$

A q.enq(3)
A q:void
<b>A q.enq(5)</b>
<del>B p.enq(4)</del>
B p:void
B q.deq()
B q:3

An invocation is  
*pending* if it has no  
matching response

# Complete Subhistory

$H =$

A q.enq(3)	
A q:void	
A q.enq(5)	
B p.enq(4)	
B p:void	
B q.deq()	
B q:3	

**May or may not have taken effect**

# Complete Subhistory

$H =$

A q.enq(3)	
A q:void	
A q.enq(5)	
B p.enq(4)	
B p:void	
B q.deq()	
B q:3	

may discard  
pending invocations

# Complete Subhistory

A q.enq(3)  
A q:void

**Complete(H) =** B p.enq(4)  
B p:void  
B q.deq()  
B q:3

# Sequential Histories

```
A q.enq(3)
A q:void
B p.enq(4)
B p:void
B q.deq()
B q:3
A q:enq(5)
```

# Sequential Histories

A q . enq (3)  
A q : void

B p . enq (4)  
B p : void  
B q . deq ()  
B q : 3  
A q : enq (5)

match

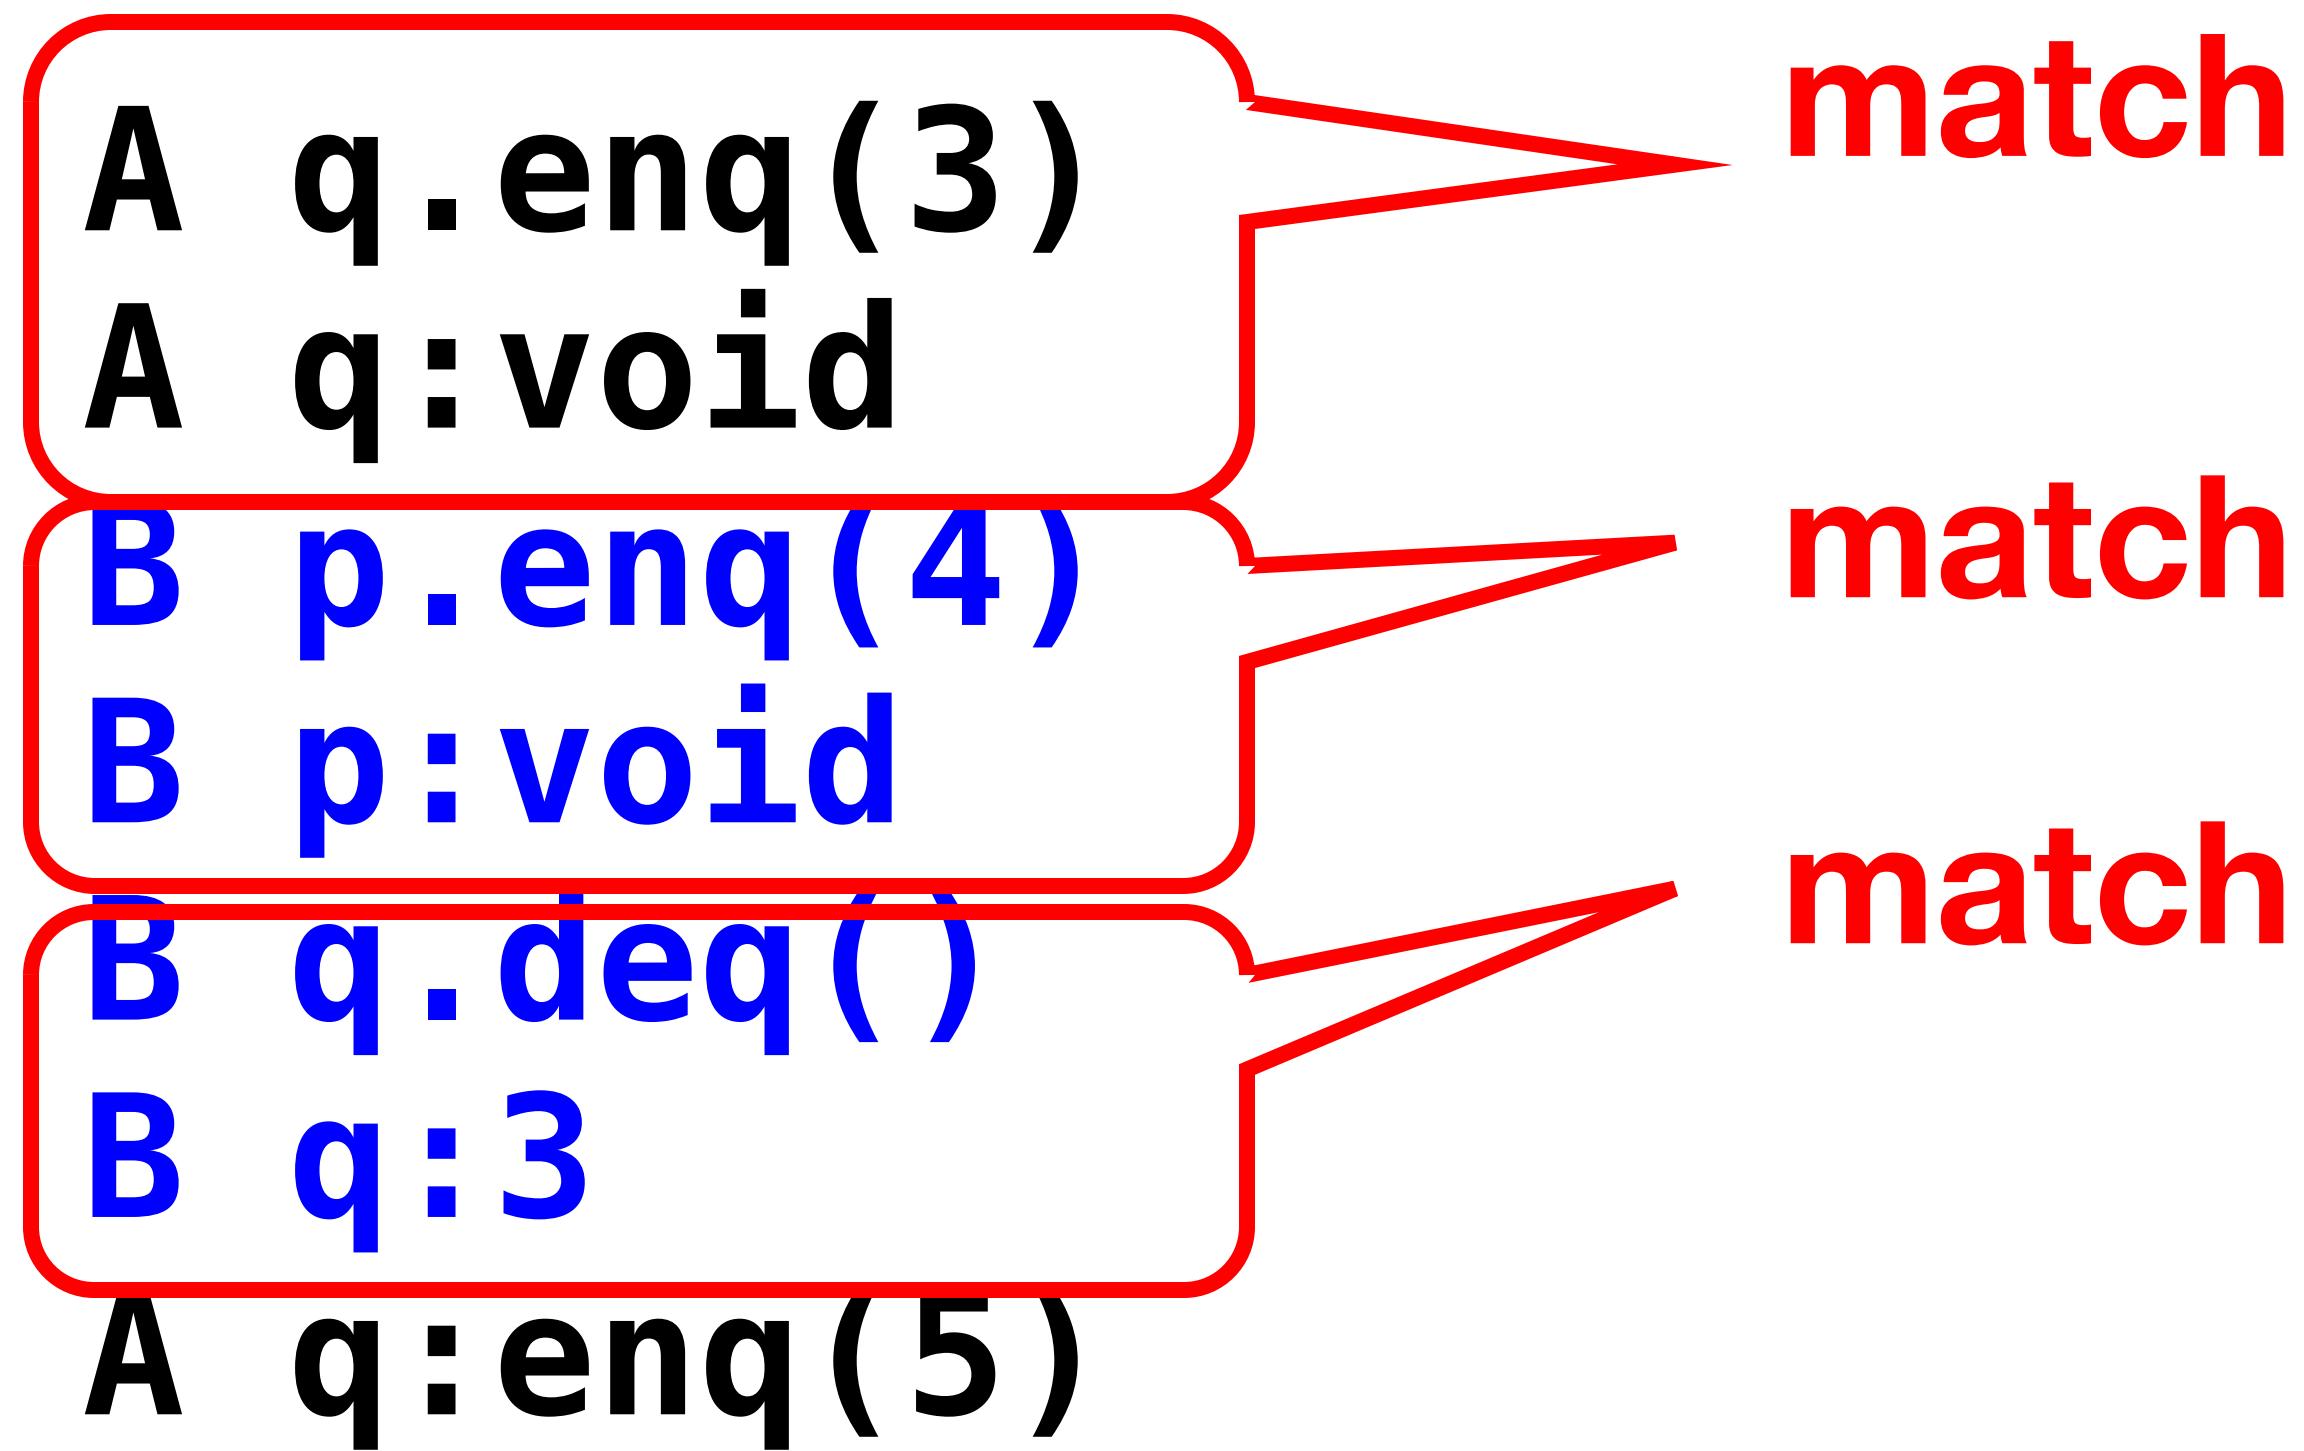
# Sequential Histories

```
A q.enq(3)
A q:void
B p.enq(4)
B p:void
B q.deq()
B q:3
A q:enq(5)
```


match

match


# Sequential Histories



# Sequential Histories



# Sequential Histories



# Well-formed Histories

A q.enq(3)  
B p.enq(4)  
B p:void  
H= B q.deq()  
A q:void  
B q:3

# Well-formed Histories

Per-thread projections  
sequential

H= A q.enq(3)  
B p.enq(4)  
B p:void  
B q.deq()  
A q:void  
B q:3

H | B= B p.enq(4)  
B p:void  
B q.deq()  
B q:3

# Well-formed Histories

Per-thread projections  
sequential

$H =$

- A q.enq(3)
- B p.enq(4)
- B p:void
- B q.deq()
- A q:void
- B q:3

$H | B =$

- B p.enq(4)
- B p:void
- B q.deq()
- B q:3

$H | A =$

- A q.enq(3)
- A q:void

# Equivalent Histories

Threads see the same  
thing in both

$$\left\{ \begin{array}{l} H|A = G|A \\ H|B = G|B \end{array} \right.$$

$H =$

```
A q.enq(3)
B p.enq(4)
B p:void
B q.deq()
A q:void
B q:3
```

$G =$

```
A q.enq(3)
A q:void
B p.enq(4)
B p:void
B q.deq()
B q:3
```

# Sequential Specifications

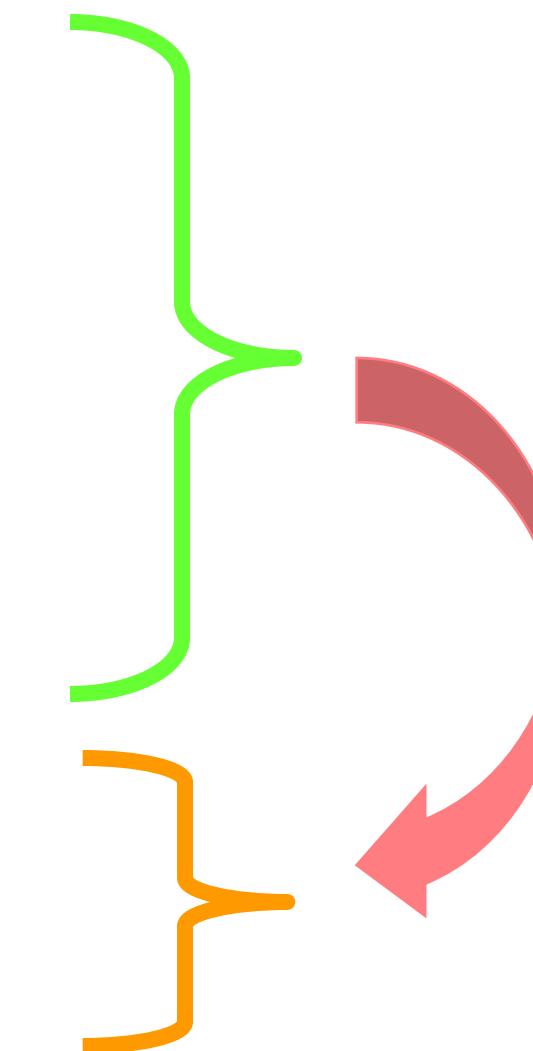
- A sequential ***specification*** is some way of telling whether a
  - Single-thread, single-object history is ***legal***
- For example:
  - Pre and post-conditions
  - But plenty of other techniques exist ...

# Legal Histories


- A sequential (*multi-object*) history  $H$  is *legal* if
  - For every object  $x$
  - $H|x$  is in the ***sequential spec*** for  $x$

# Precedence

A q.enq(3)  
B p.enq(4)  
B p.void  
A q:void  
B q.deq()  
B q:3




A method call **precedes** another if response event precedes invocation event




# Non-Precedence

```
A q.enq(3)
B p.enq(4)
B p_void
B q.deq()
A q_void
B q:3
```




**Some method calls  
overlap one another**



# Notation

- Given
  - History  $H$
  - method executions  $m_0$  and  $m_1$  in  $H$
- We say  $m_0 \rightarrow_H m_1$ , if
  - $m_0$  precedes  $m_1$
- Relation  $m_0 \rightarrow_H m_1$  is a
  - Partial order
  - Total order if  $H$  is sequential



# Linearizability

- History  $H$  is **linearizable** if it can be extended to  $G$  by
  - Appending zero or more responses to pending invocations
  - Discarding other pending invocations
- So that  $G$  is equivalent to
  - Legal sequential history  $S$
  - where  $\rightarrow_G \subset \rightarrow_S$


# Remarks on Linearizability

- Some pending invocations
  - Took effect, so keep them
  - Discard the rest
- Condition  $\rightarrow_G \subset \rightarrow_S$ 
  - Means that **S** respects “real-time order” of **G**

# Ensuring $\rightarrow_G \subset \rightarrow_S$

$$\rightarrow_G = \{a \rightarrow c, b \rightarrow c\}$$


$$\rightarrow_S = \{a \rightarrow b, a \rightarrow c, b \rightarrow c\}$$



# Ensuring $\rightarrow_G \subset \rightarrow_S$

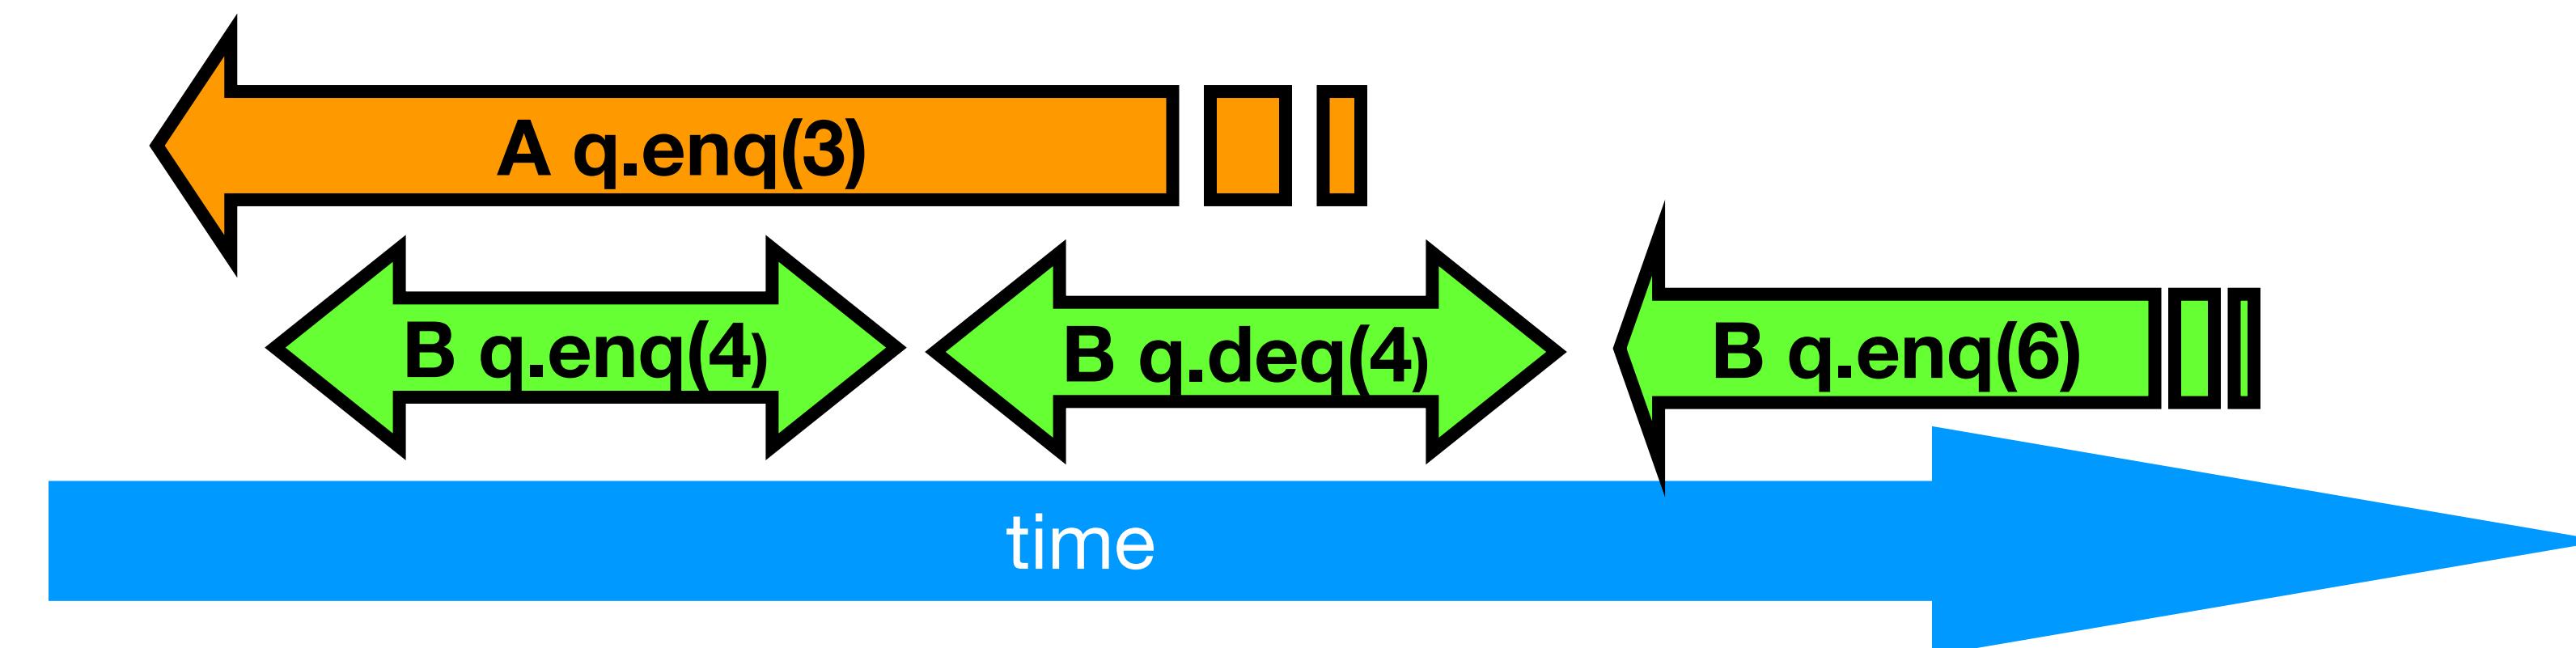
$$\rightarrow_G = \{a \rightarrow c, b \rightarrow c\}$$

$$\rightarrow_S = \{a \rightarrow b, a \rightarrow c, b \rightarrow c\}$$




# Example


```
A q.enq(3)
B q.enq(4)
B q: void
B q.deq()
B q: 4
B q.enq(6)
```




# Example



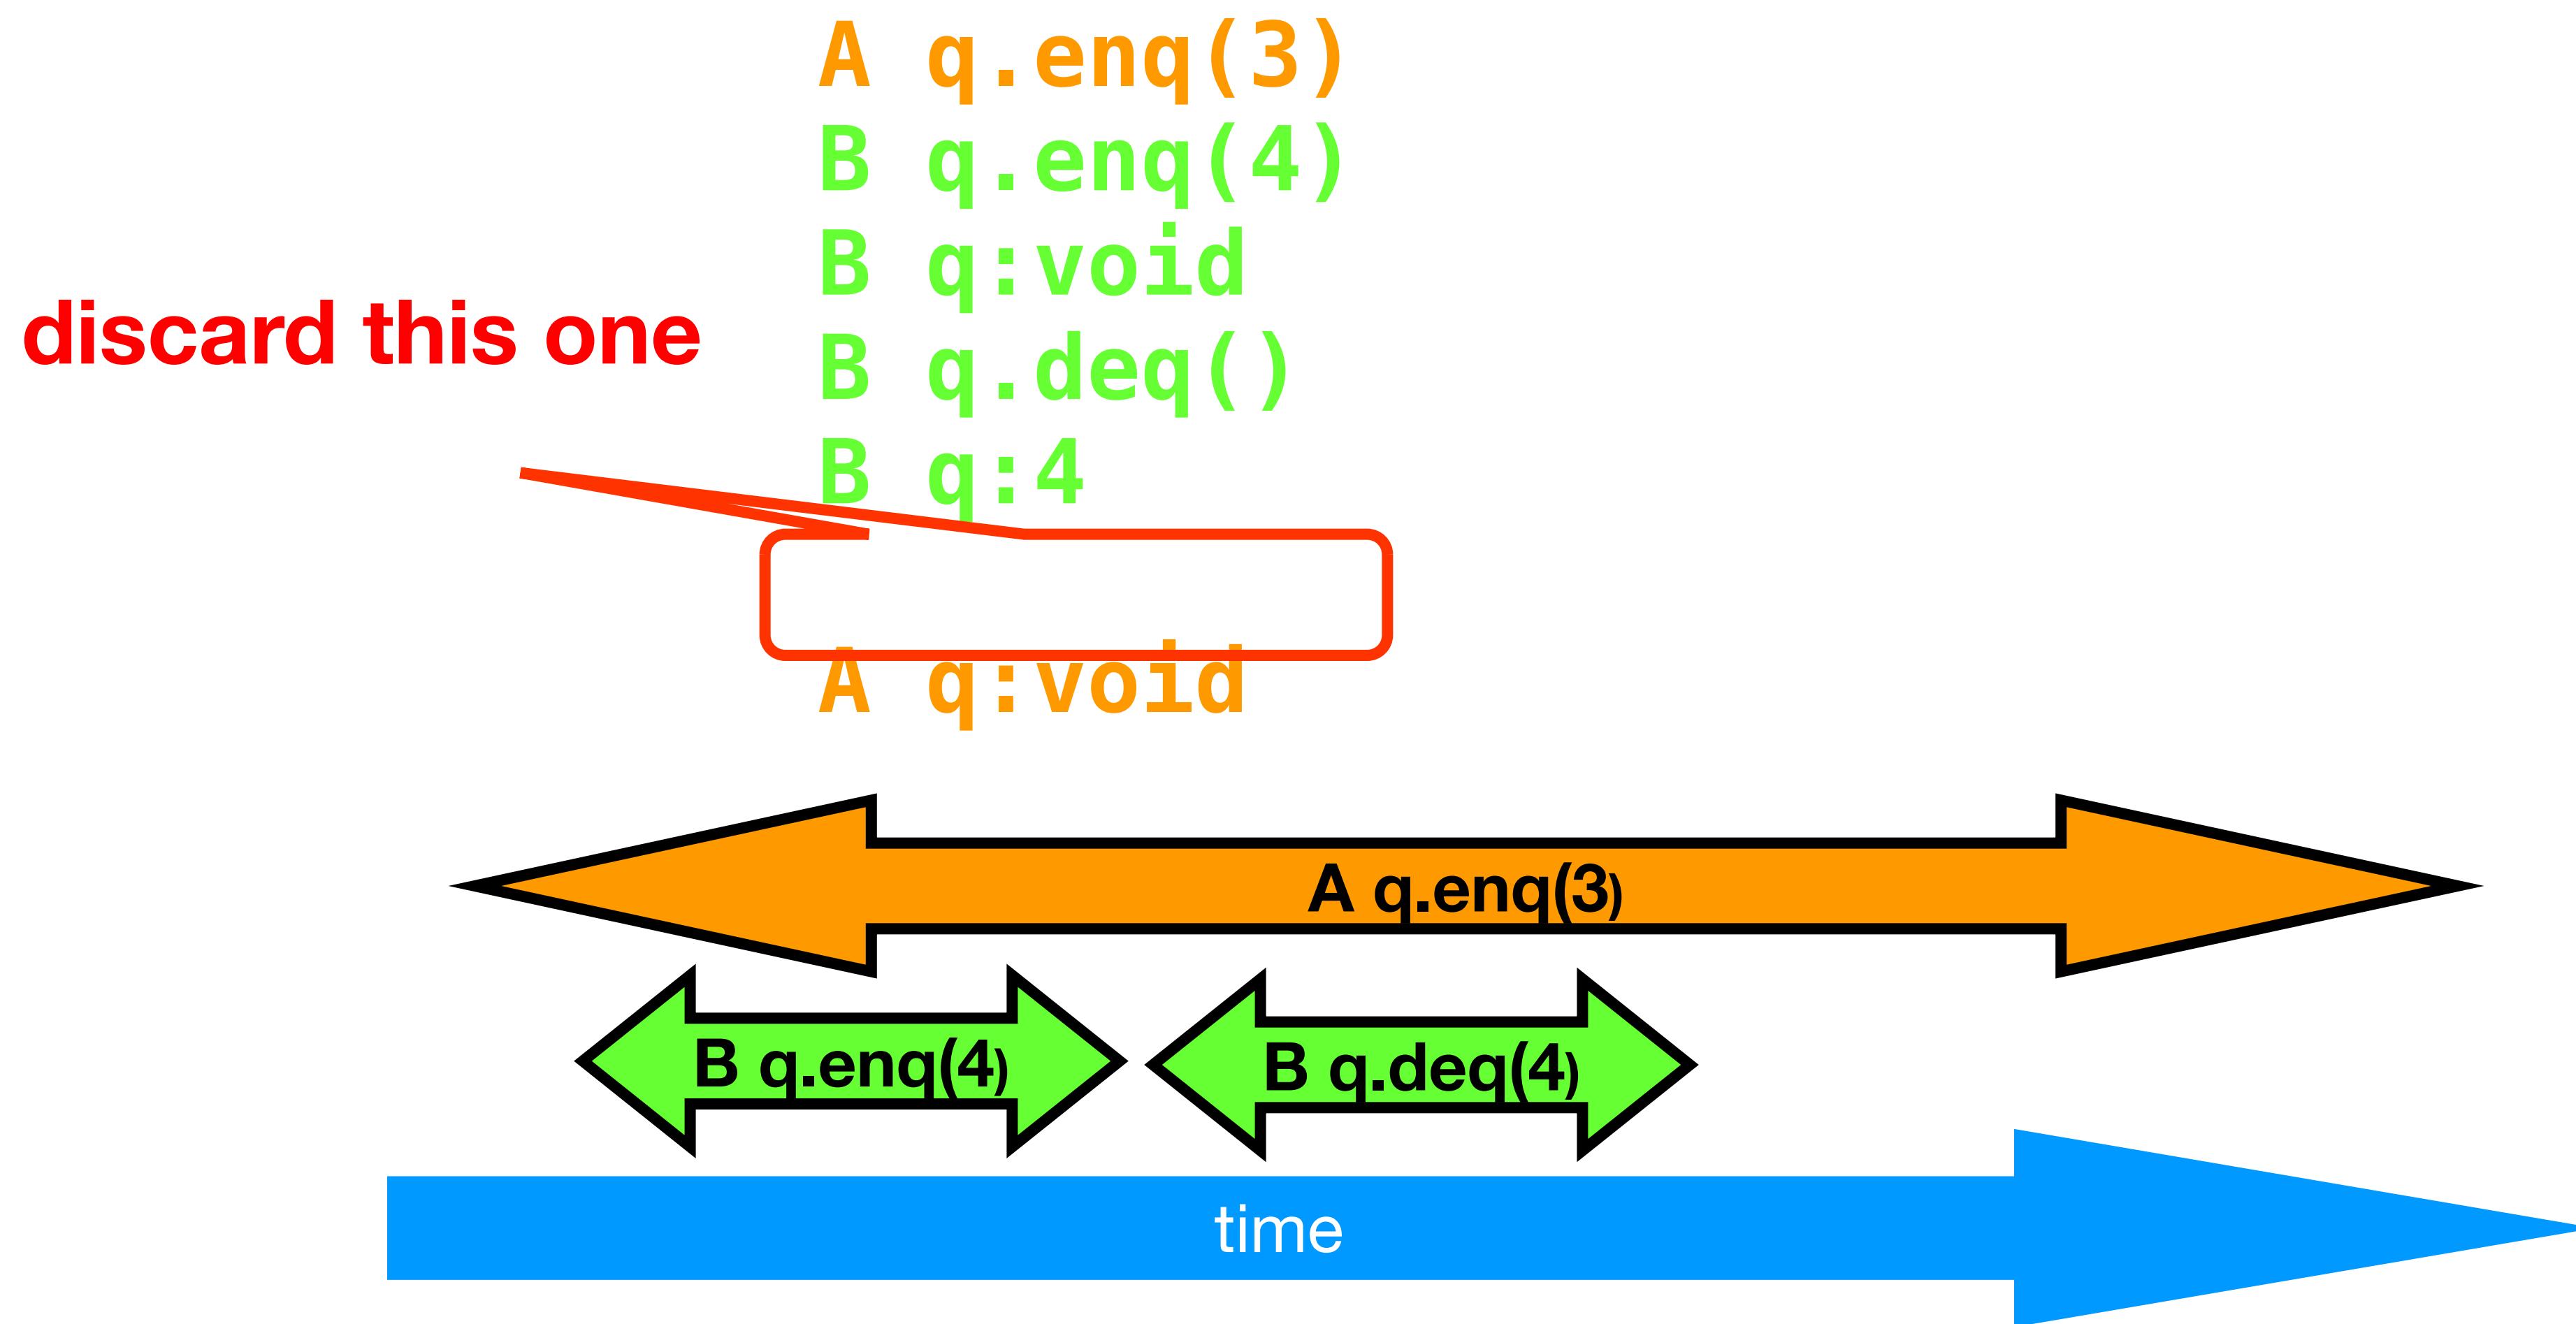
# Example



# Example




# Example



# Example

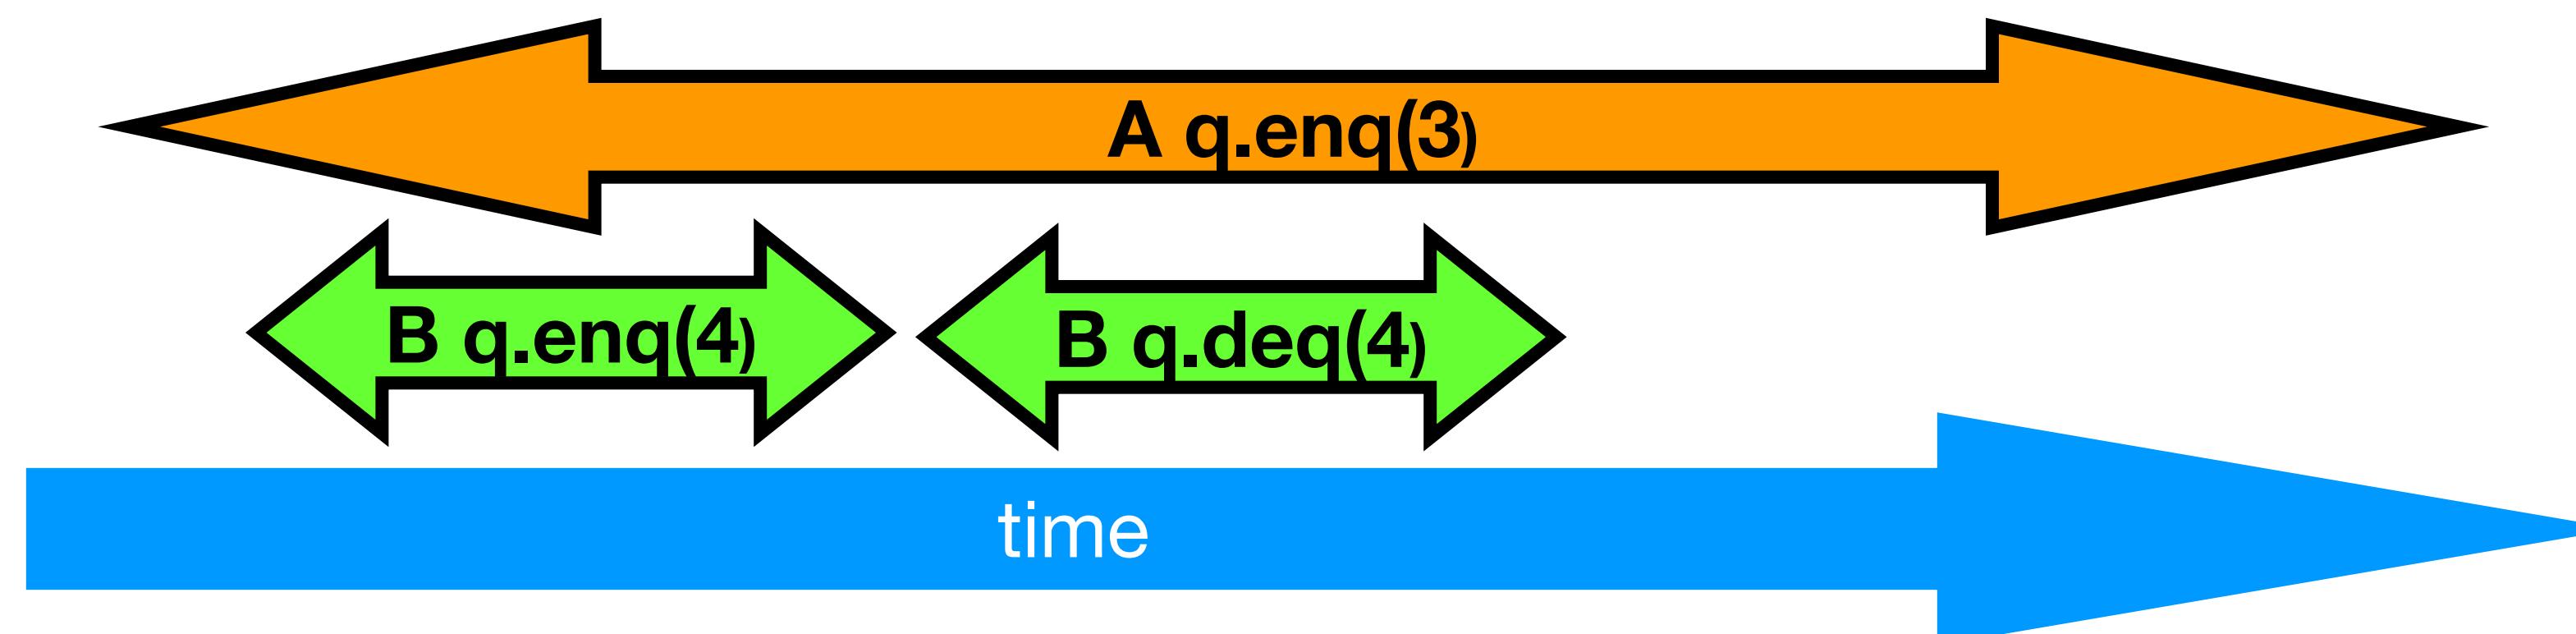
```
A q.enq(3)
B q.enq(4)
B q: void
B q.deq()
B q: 4
A q: void
```



# Example

A q.enq(3)  
B q.enq(4)  
B q: void  
B q.deq()  
B q: 4  
A q: void

B q.enq(4)  
B q: void  
A q.enq(3)  
A q: void  
B q.deq()  
B q: 4

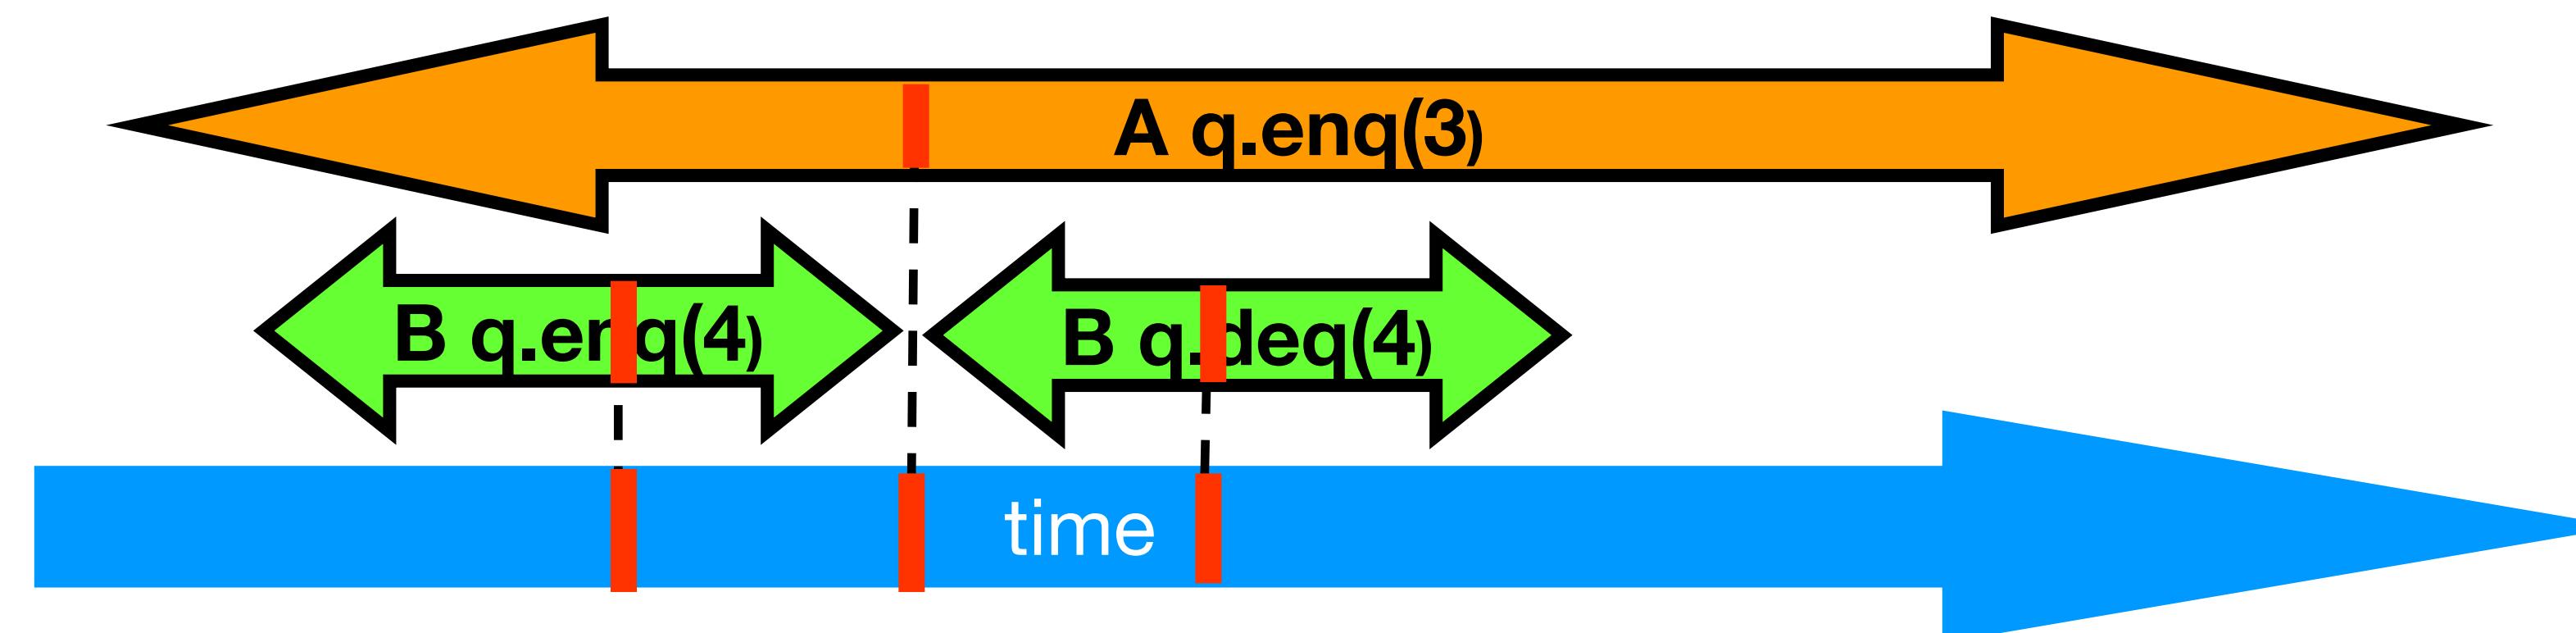


# Example

A q.enq(3)  
B q.enq(4)  
B q:void  
B q.deq()  
B q:4  
A q:void

Equivalent sequential history

B q.enq(4)  
B q:void  
A q.enq(3)  
A q:void  
B q.deq()  
B q:4



# Composability Theorem

- History  $H$  is linearizable if and only if
  - For every object  $x$
  - $H|x$  is linearizable
- We care about objects only!
  - (Materialism?)

# Why does composability matter?

- Modularity
- Can prove linearizability of objects in isolation
- Can compose independently-implemented objects

# Reasoning about Linearizability: Locking

```
let deq q =
  Mutex.lock q.lock;
  try
    if q.tail = q.head then
      raise Empty;
    match q.items.(q.head mod q.capacity) with
    | None -> assert false
    | Some x ->
        q.head <- q.head + 1;
        Mutex.unlock q.lock;
        x
  with e ->
    Mutex.unlock q.lock;
    raise e
```

Linearization points  
are when locks are  
released

# More Reasoning – Wait-free

```
(** Enqueue – should be called by only ONE thread *)    (** Dequeue – should be called by only ONE thread *)
let enq q x =
  (* Check if queue is full *)
  if q.tail - q.head = q.capacity then
    raise Full;
  (* Write to the array *)
  q.items.(q.tail mod q.capacity) <- Some x;
  (* Advance tail *)
  q.tail <- q.tail + 1

let deq q =
  (* Check if queue is empty *)
  if q.tail = q.head then
    raise Empty;
  (* Read from the array *)
  match q.items.(q.head mod q.capacity) with
  | None -> assert false (* Should never happen *)
  | Some x ->
    (* Advance head *)
    q.head <- q.head + 1;
    x
```

# More Reasoning – Wait-free

```
(** Enqueue – should be called by only ONE thread *)
let enq q x =
  (* Check if queue is full *)
  if q.tail - q.head = q.capacity then
    raise Full;
  (* Write to the array *)
  q.items.(q.tail mod q.capacity) <- Some x;
  (* Advance tail *)
  q.tail <- q.tail + 1
```

```
(** Dequeue – should be called by only ONE thread *)
let deq q =
  (* Check if queue is empty *)
  if q.tail = q.head then
    raise Empty;
  (* Read from the array *)
  match q.items.(q.head mod q.capacity) with
  | None -> assert false (* Should never happen *)
  | Some x ->
    (* Advance head *)
    q.head <- q.head + 1;
```

x

Linearization order is  
order head and tail  
fields modified

# More Reasoning – Wait-free

```
(** Enqueue – should be called by only ONE thread *)  
let enq q x =  
  (* Check if queue is full *)  
  if q.tail - q.head = q.capacity then  
    raise Full;  
  (* Write to the array *)  
  q.items.(q.tail mod q.capacity)  
  (* Advance tail *)  
  q.tail <- q.tail + 1  
  
(** Dequeue – should be called by only ONE thread *)  
let deque q =  
  (* Check if queue is empty *)  
  if q.tail = q.head then  
    raise Empty;  
  (* Read from the array *)  
  match q.items.(q.head mod q.capacity) with  
  | None -> assert false (* Should never happen *)  
  | Some x ->  
    (* Advance head *)  
    q.head <- q.head + 1;
```

Remember that there  
is only one enqueue  
and only one dequeue

Linearization order is  
order head and tail  
fields modified

# Finding linearisation points

- Identify one atomic step where the method “happens”
  - Critical section
  - Machine instruction
- Doesn’t always work
  - Might need to define several different steps for a given method
  - We will see this phenomenon in future lectures

# Linearizability: Summary

- Powerful specification tool for shared objects
- Allows us to capture the notion of objects being “atomic”
- Don’t leave home without it

# Alternative: Sequential Consistency


- History  $H$  is ***Sequentially Consistent*** if it can be extended to  $G$  by
  - Appending zero or more responses to pending invocations
  - Discarding other pending invocations
- So that  $G$  is equivalent to
  - Legal sequential history  $S$
  - ~~where  $\rightarrow_G \subset \rightarrow_S$~~

*Differs from  
Linearizability*

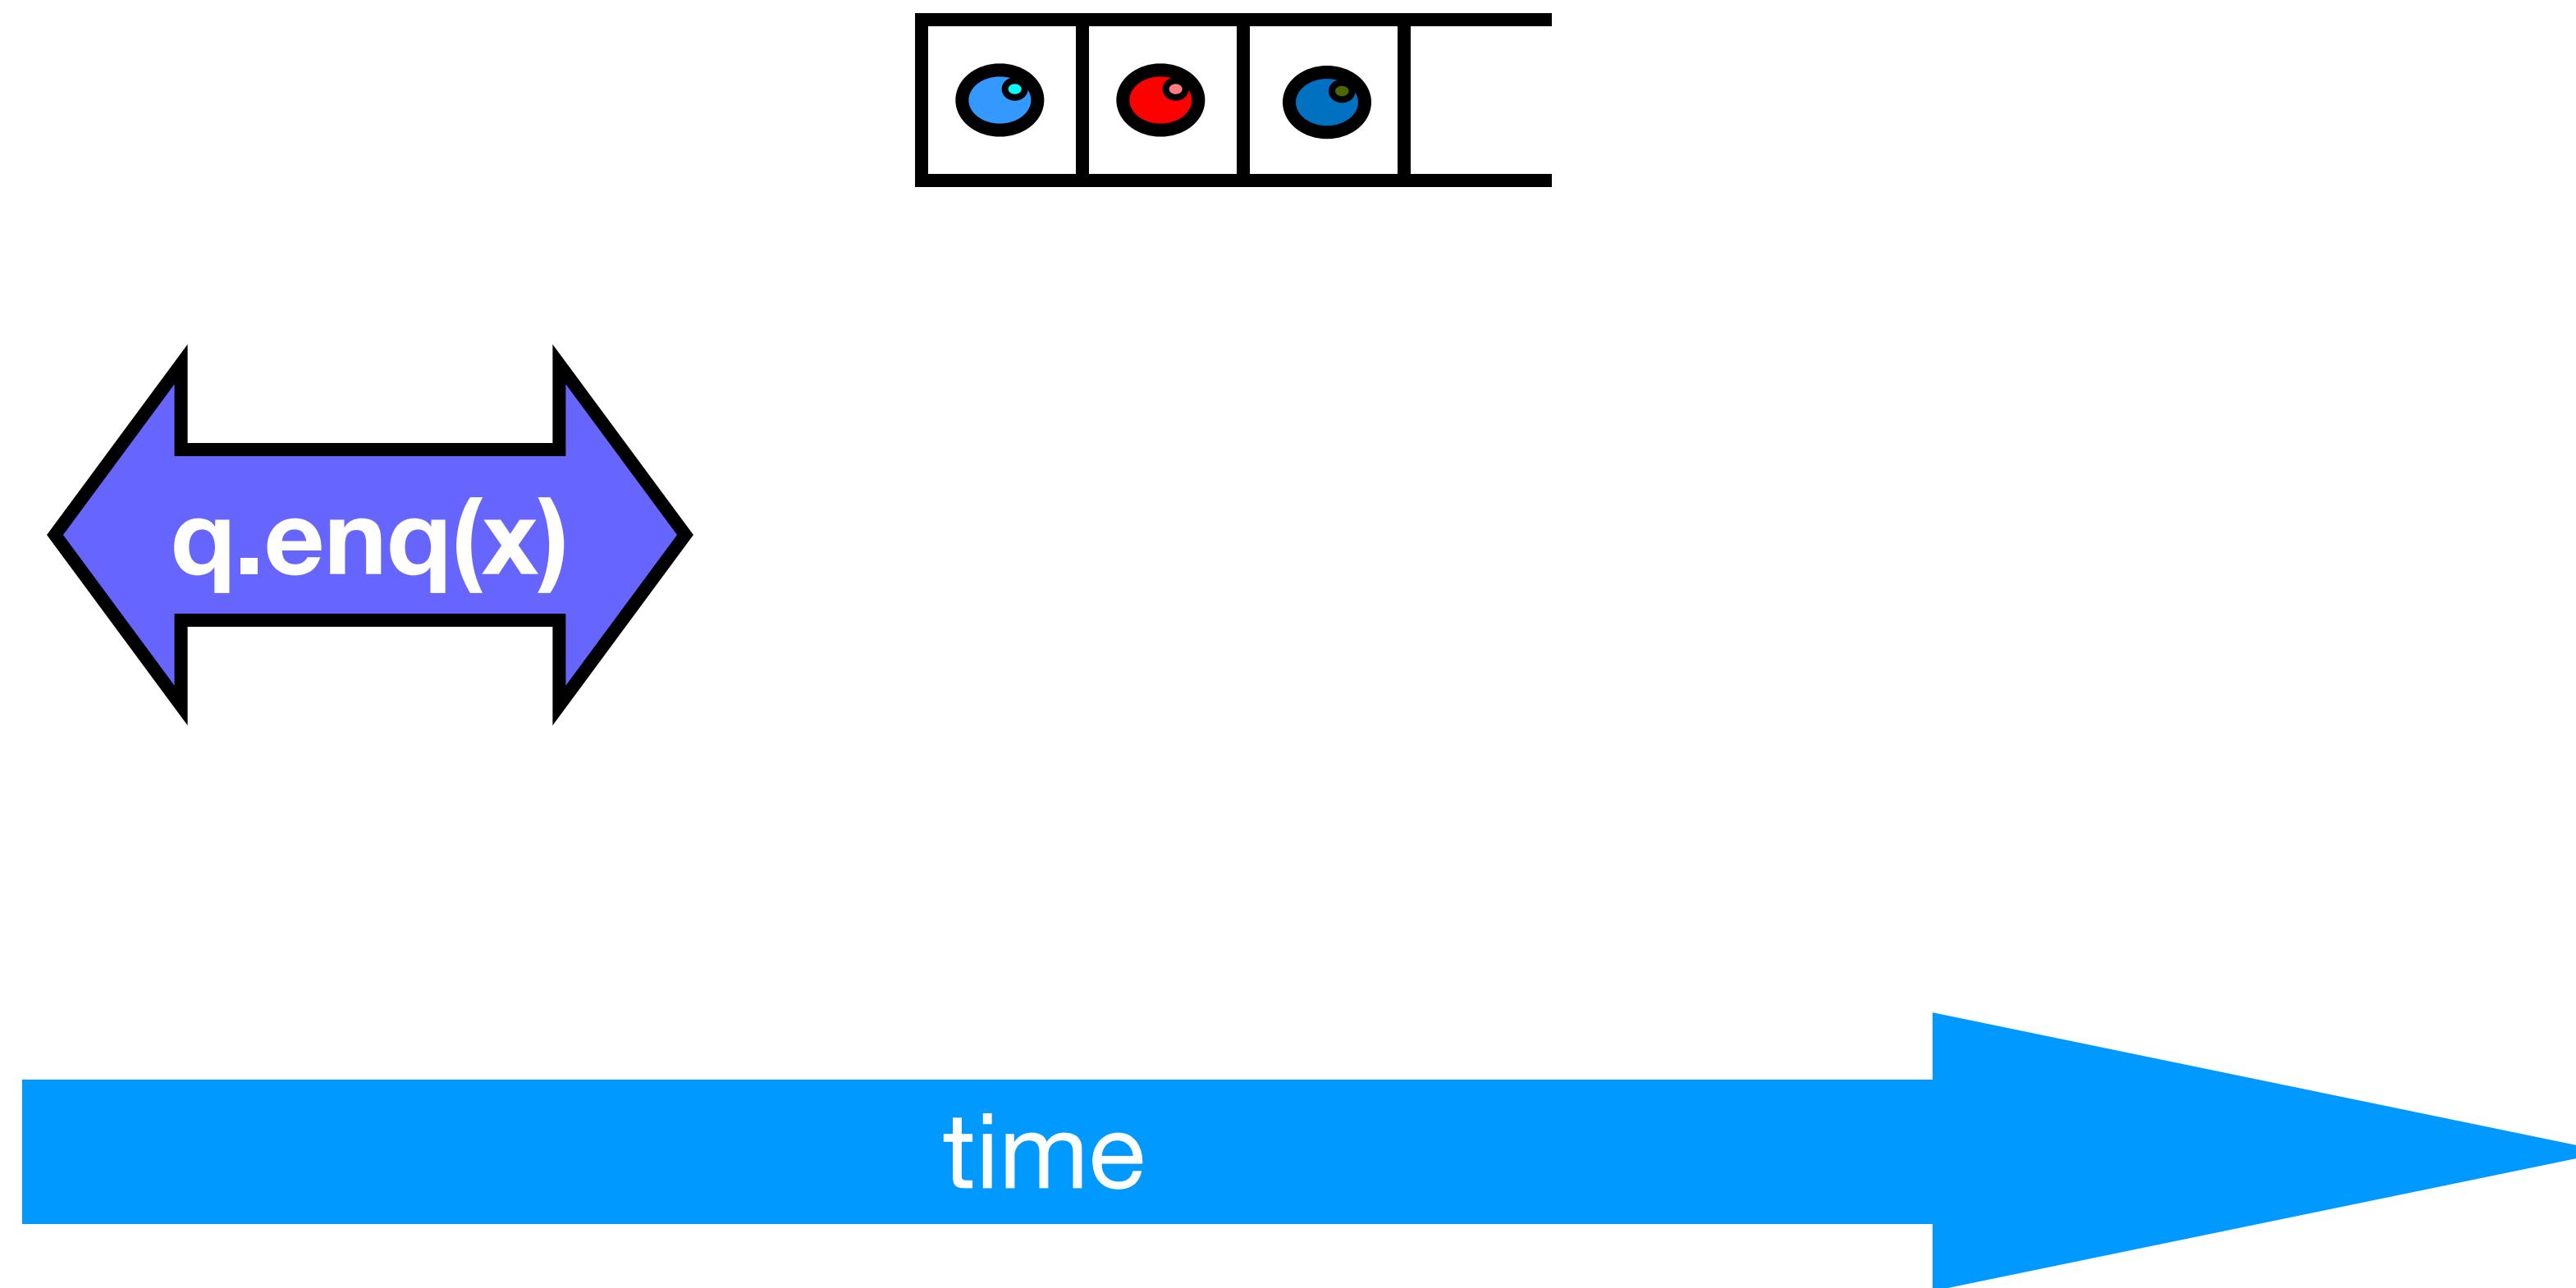
# Sequential Consistency

- No need to preserve real-time order
  - **Cannot** re-order operations done by the same thread
  - **Can** re-order non-overlapping operations done by different threads
- Often used to describe multiprocessor memory architectures

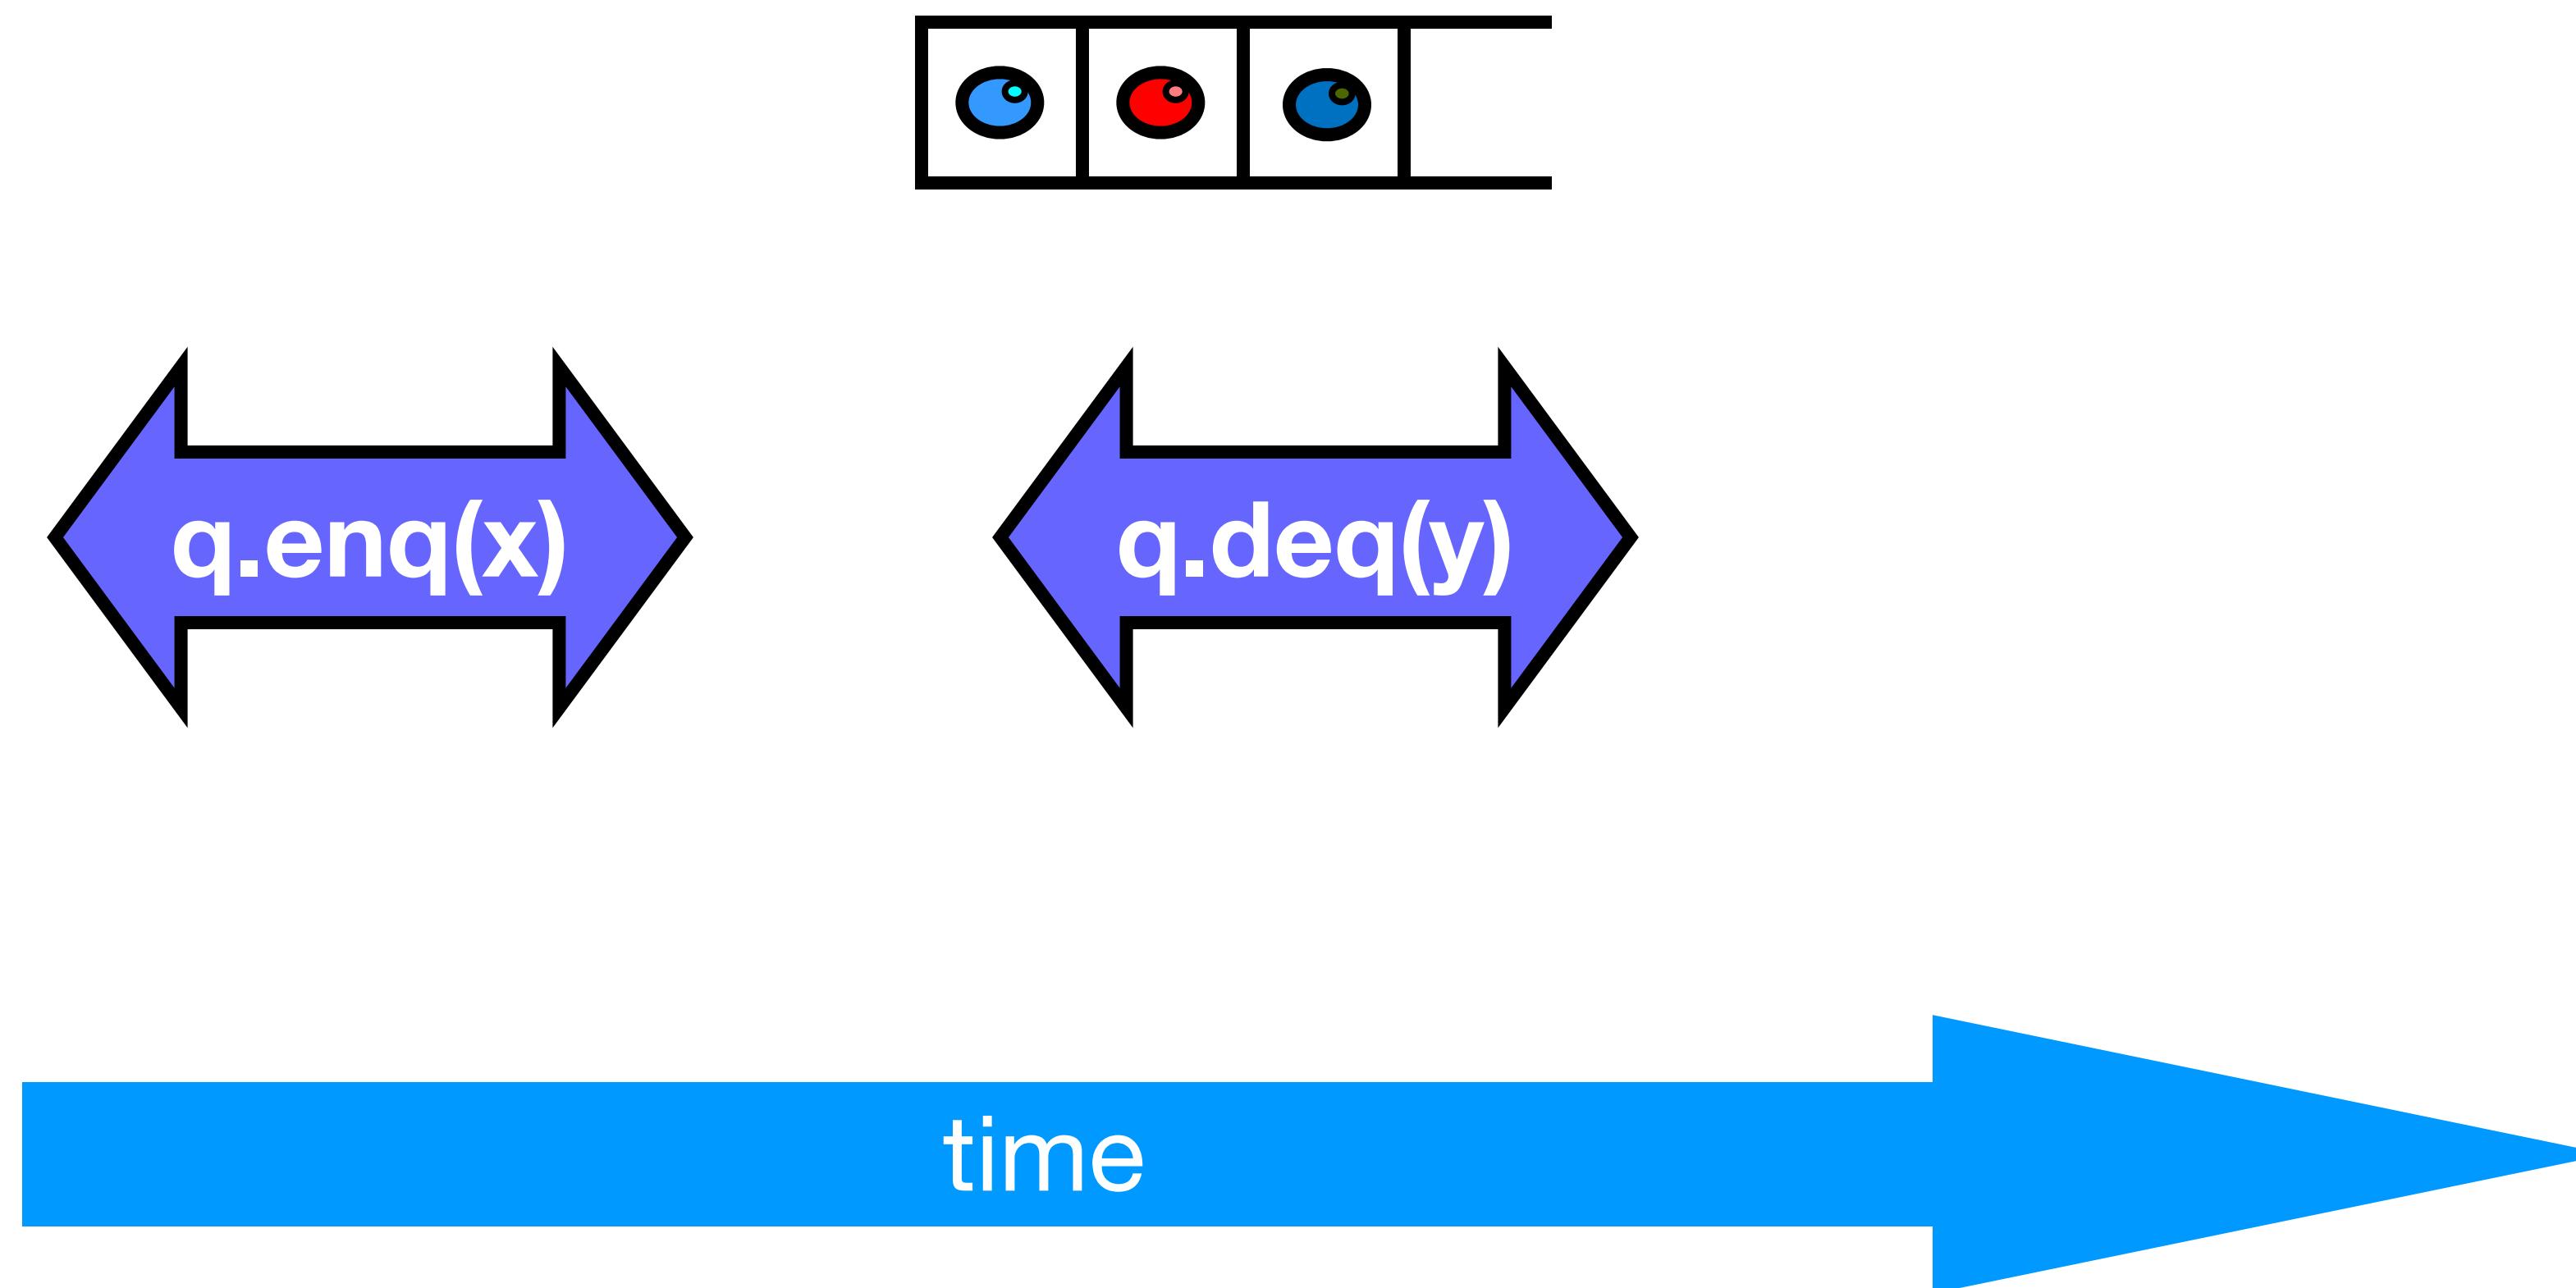
# Example



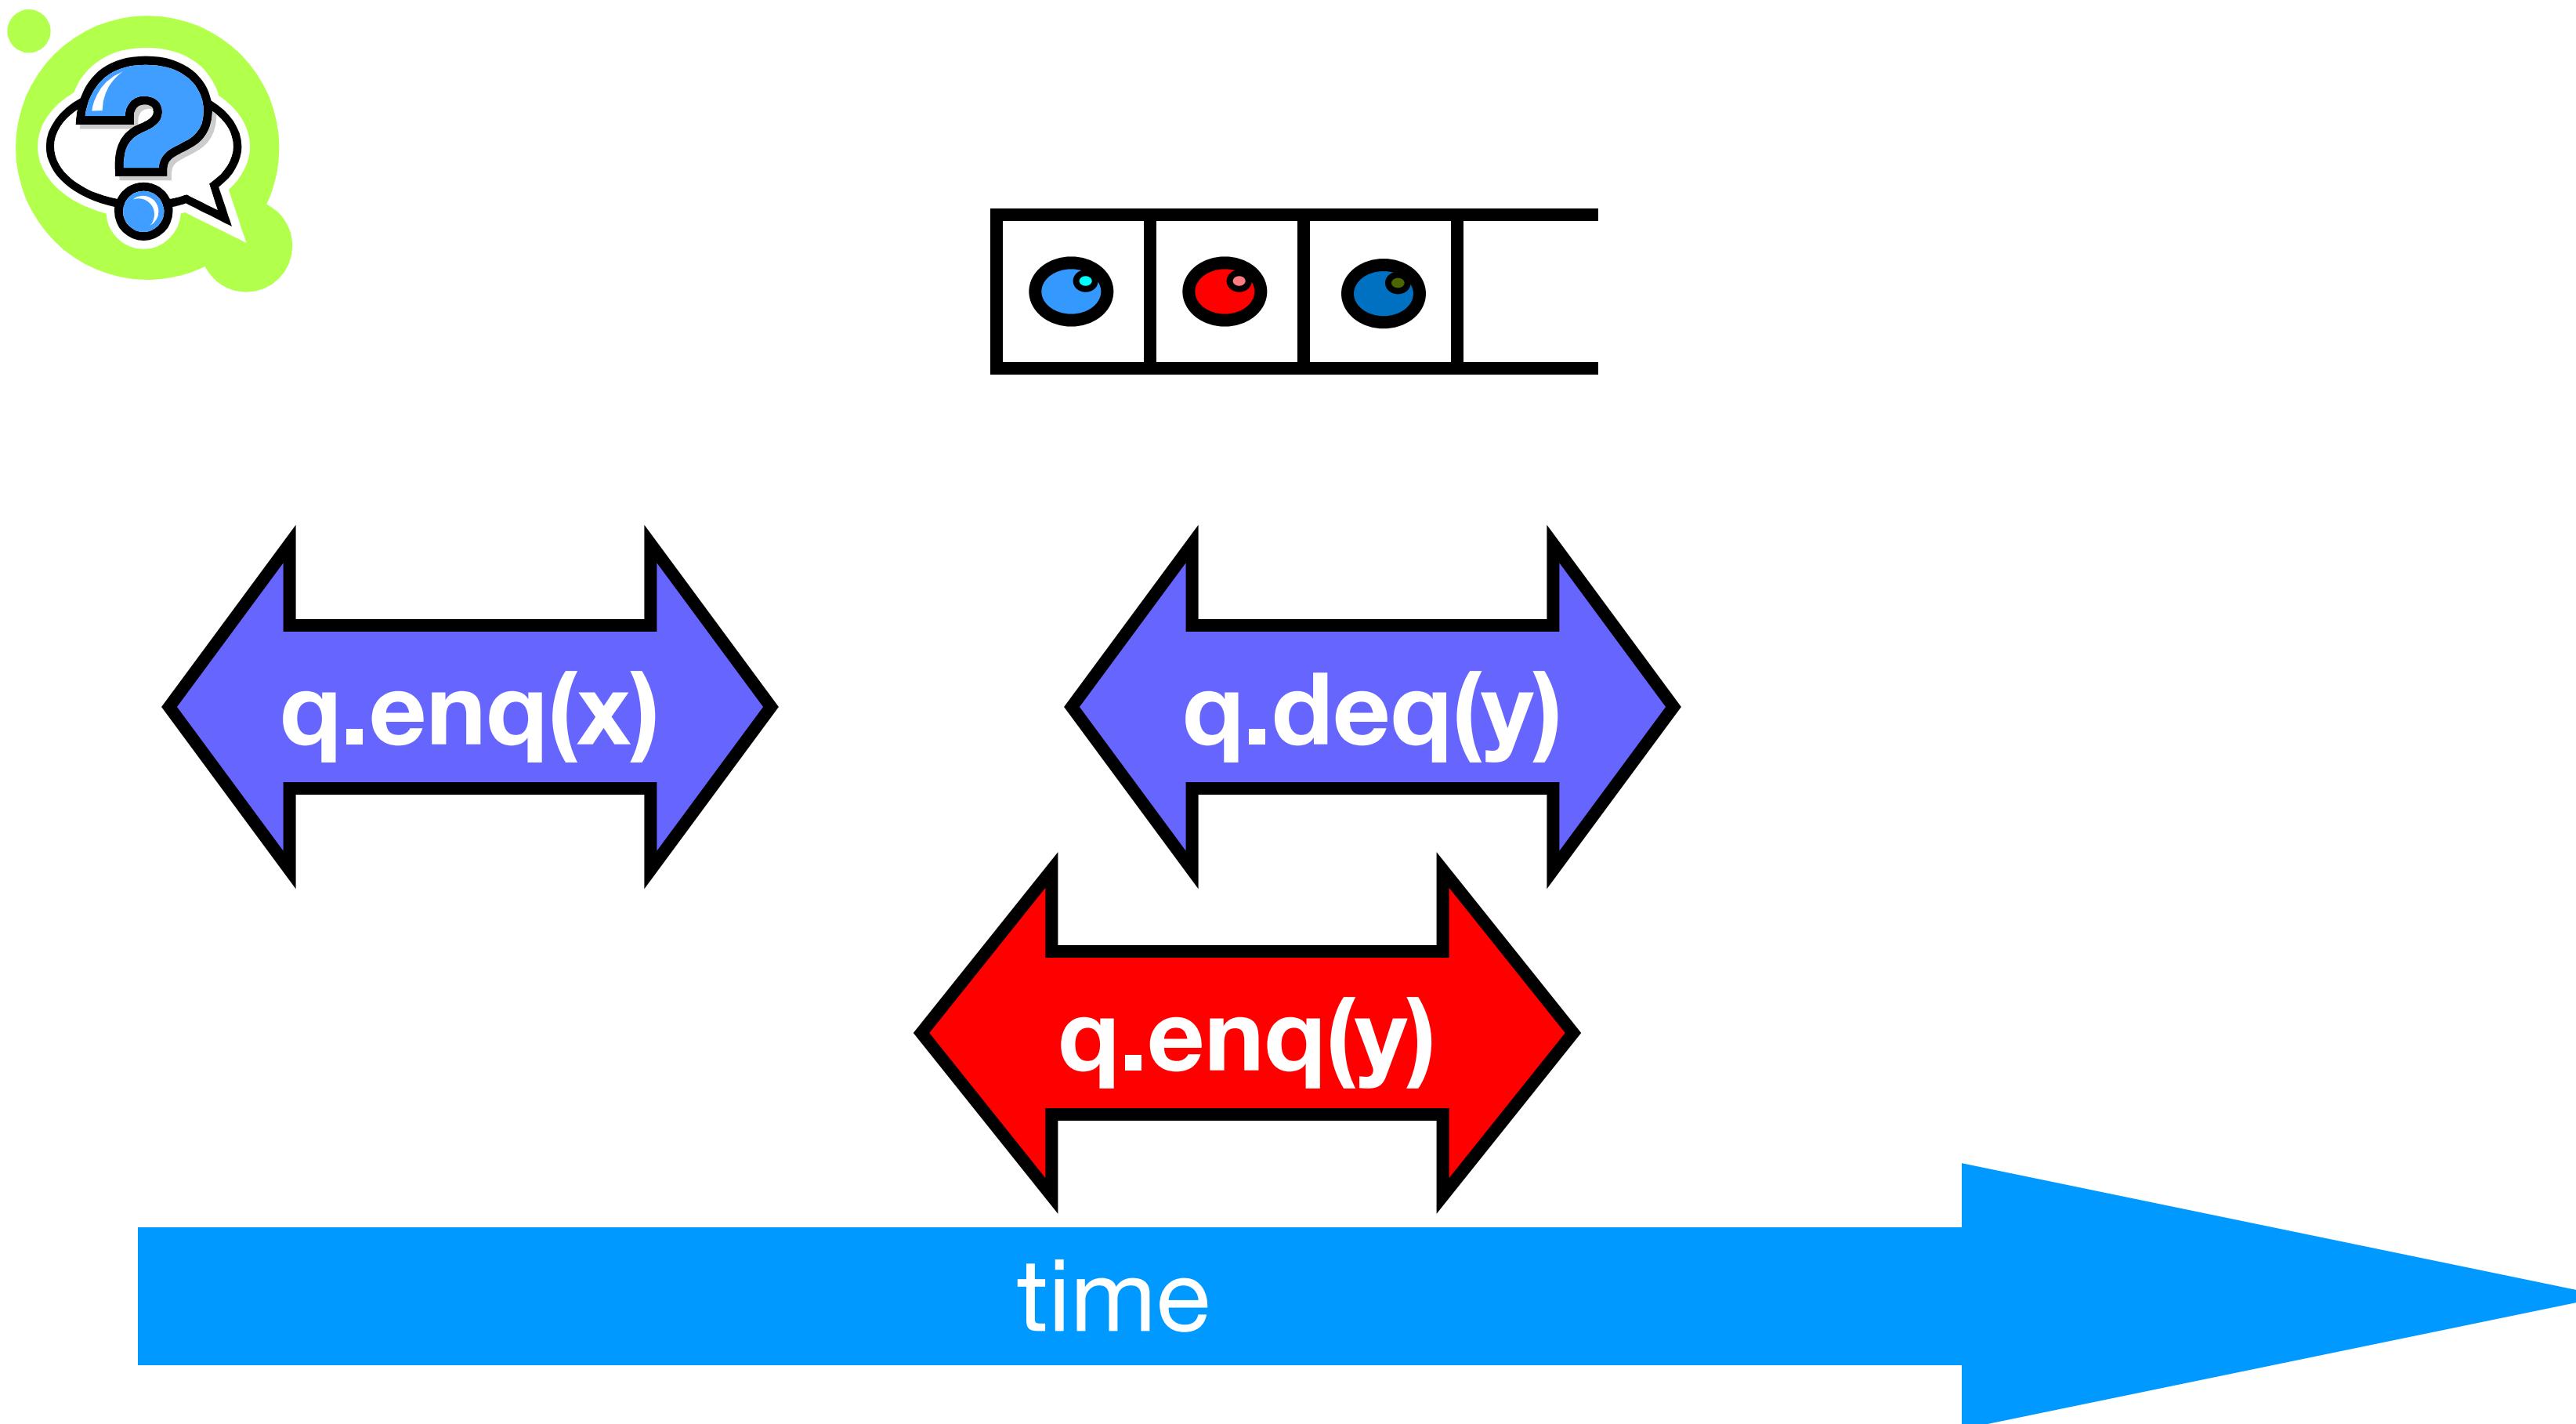
# Example



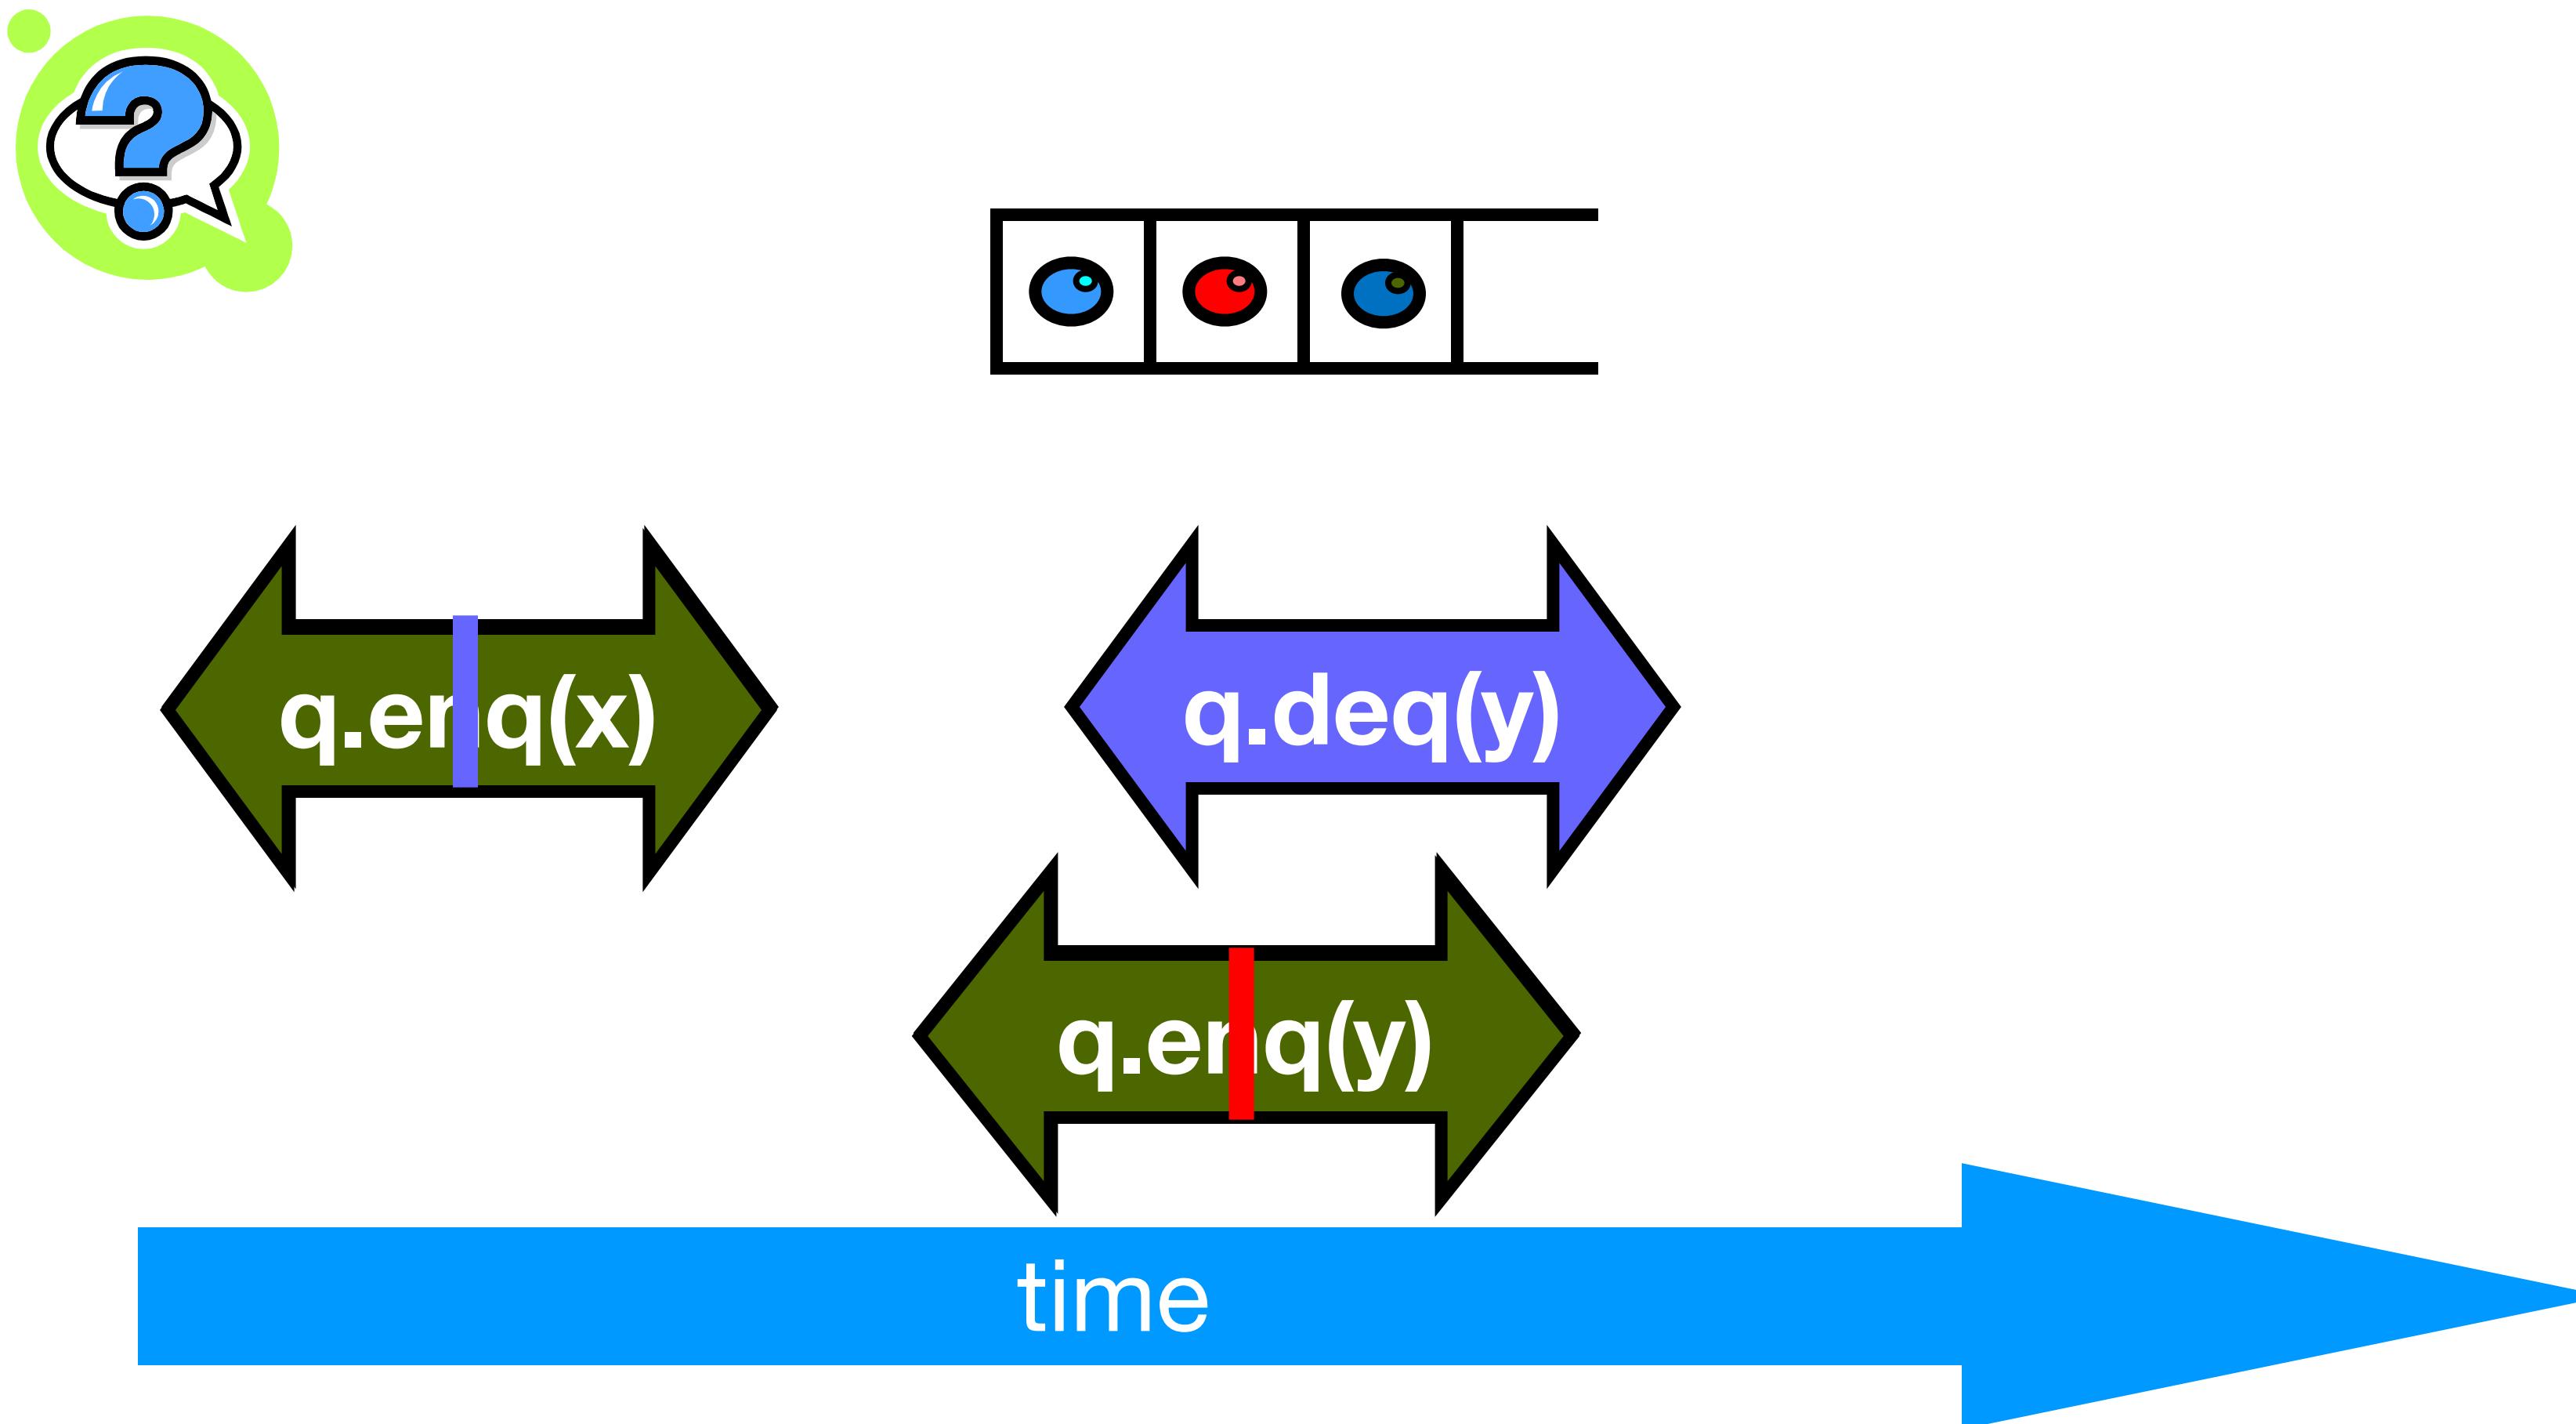
# Example



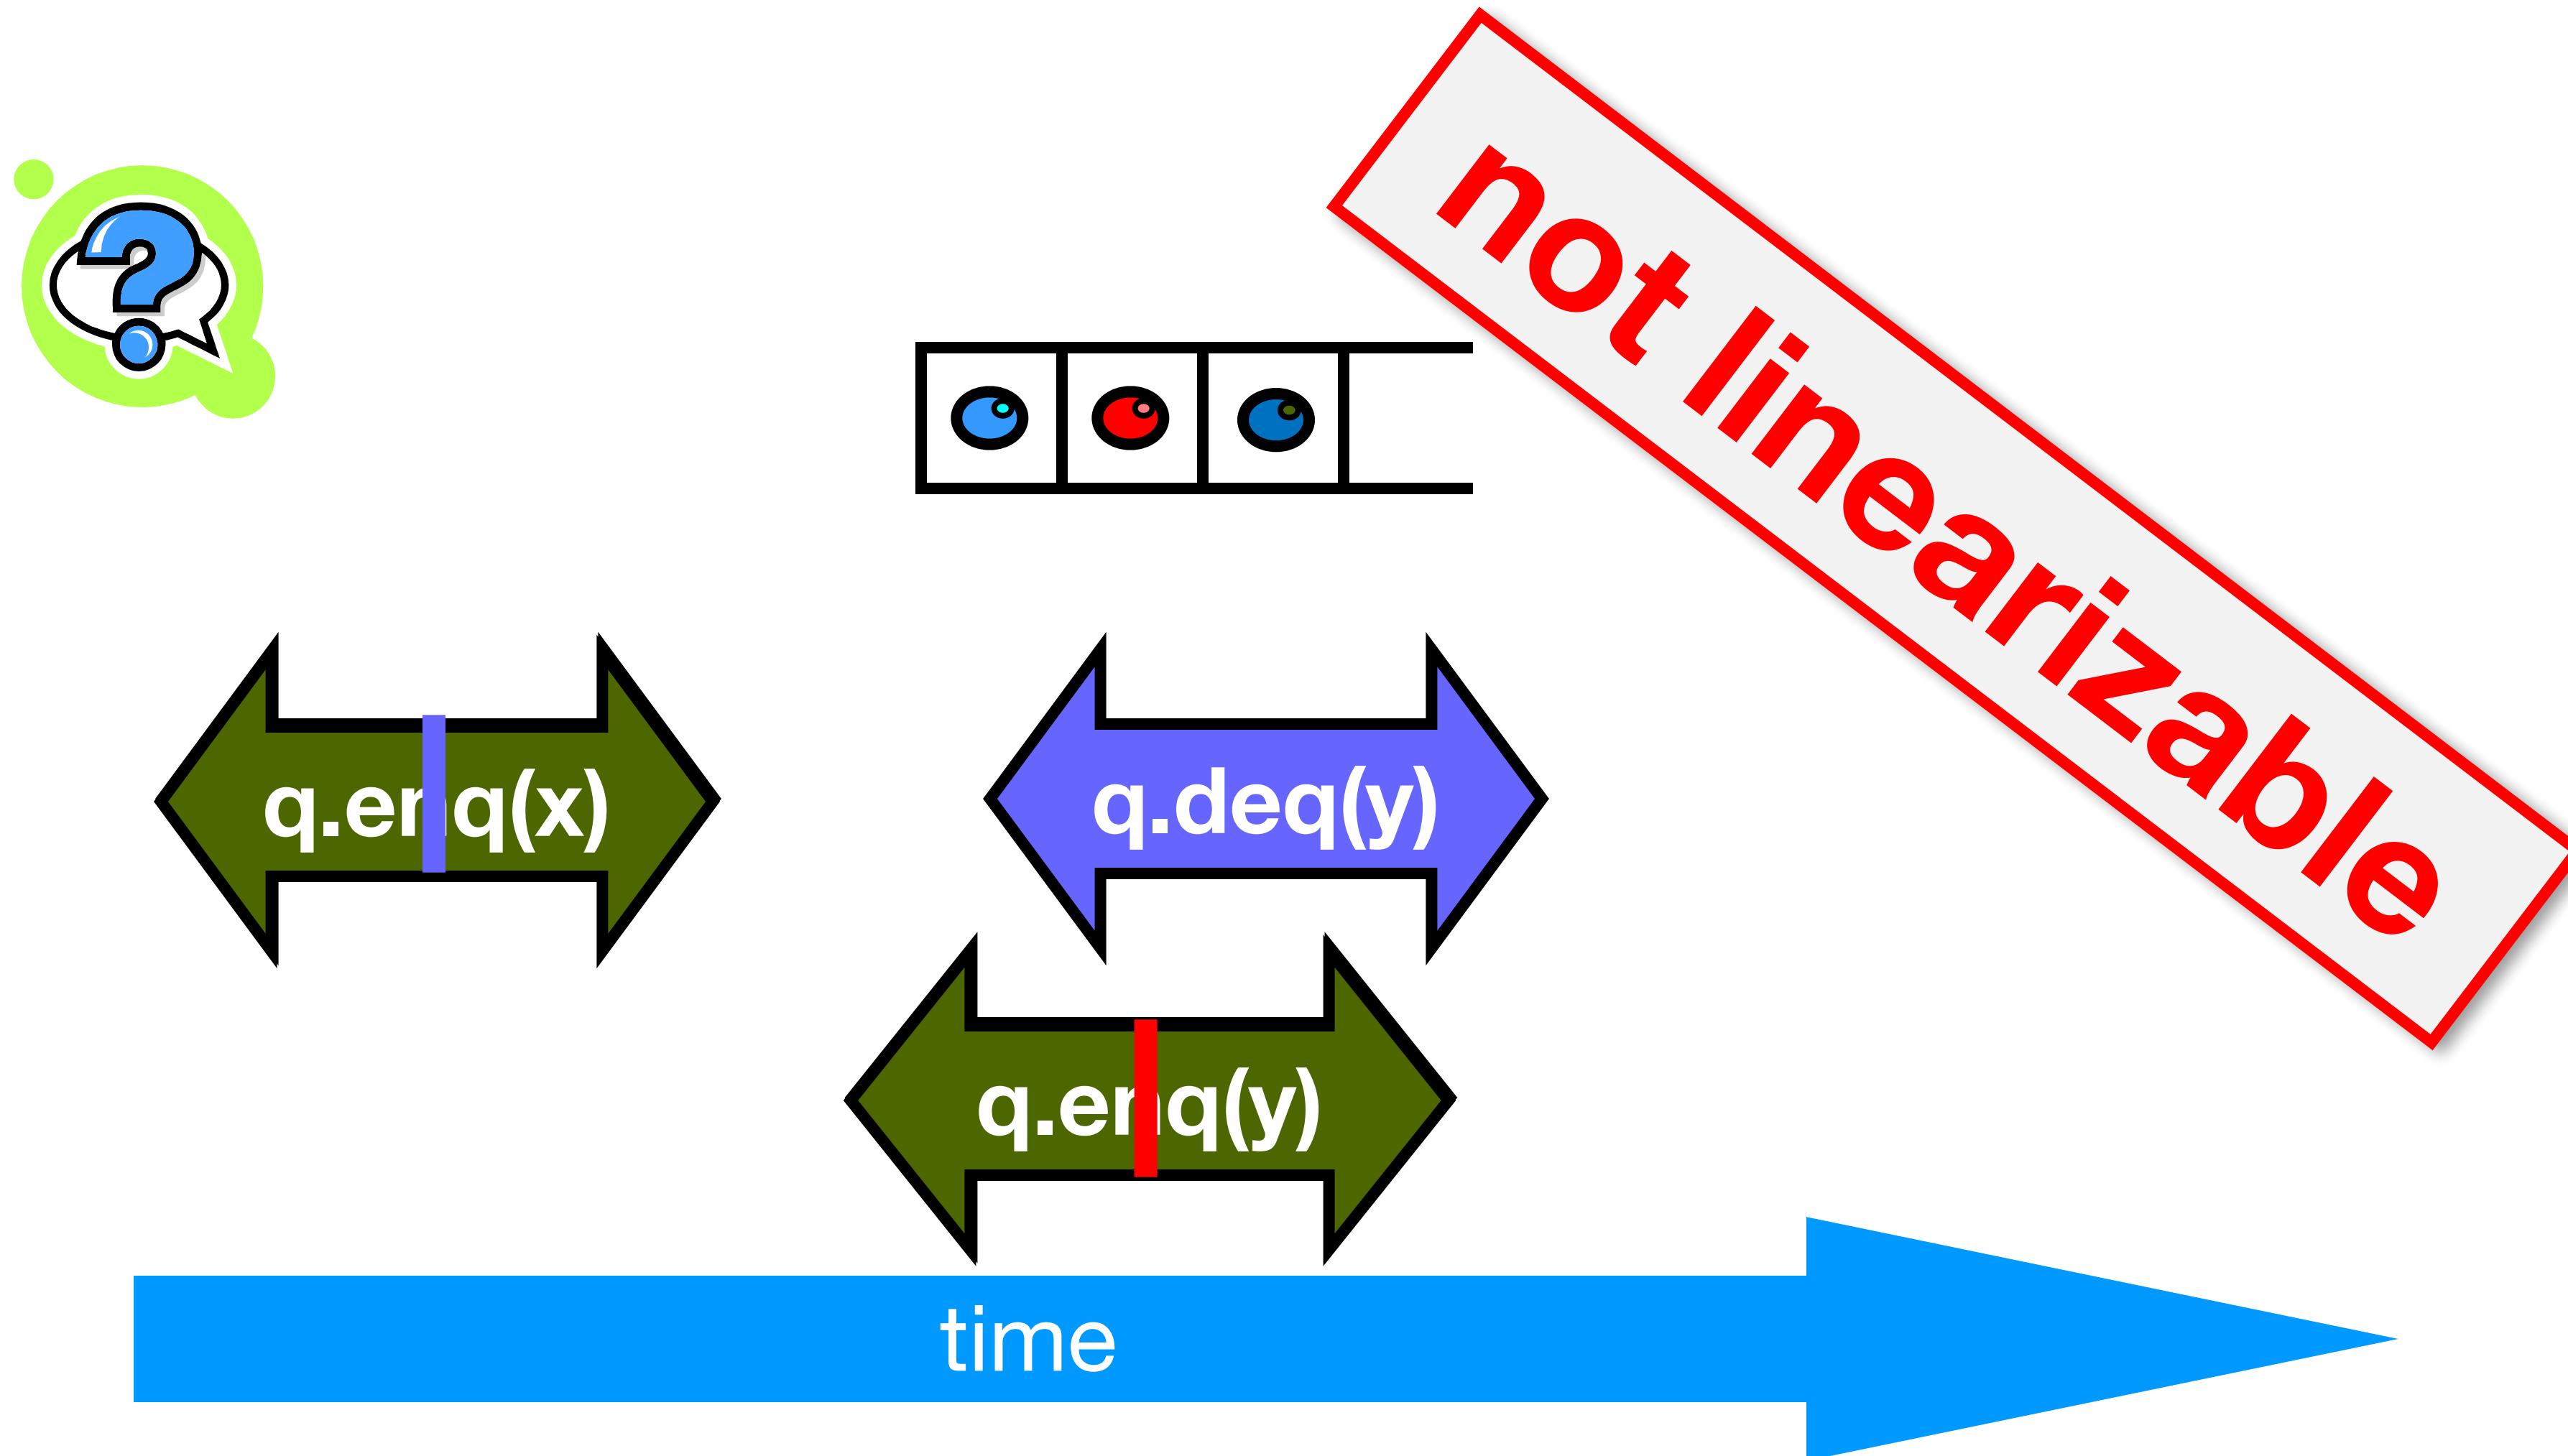
# Example



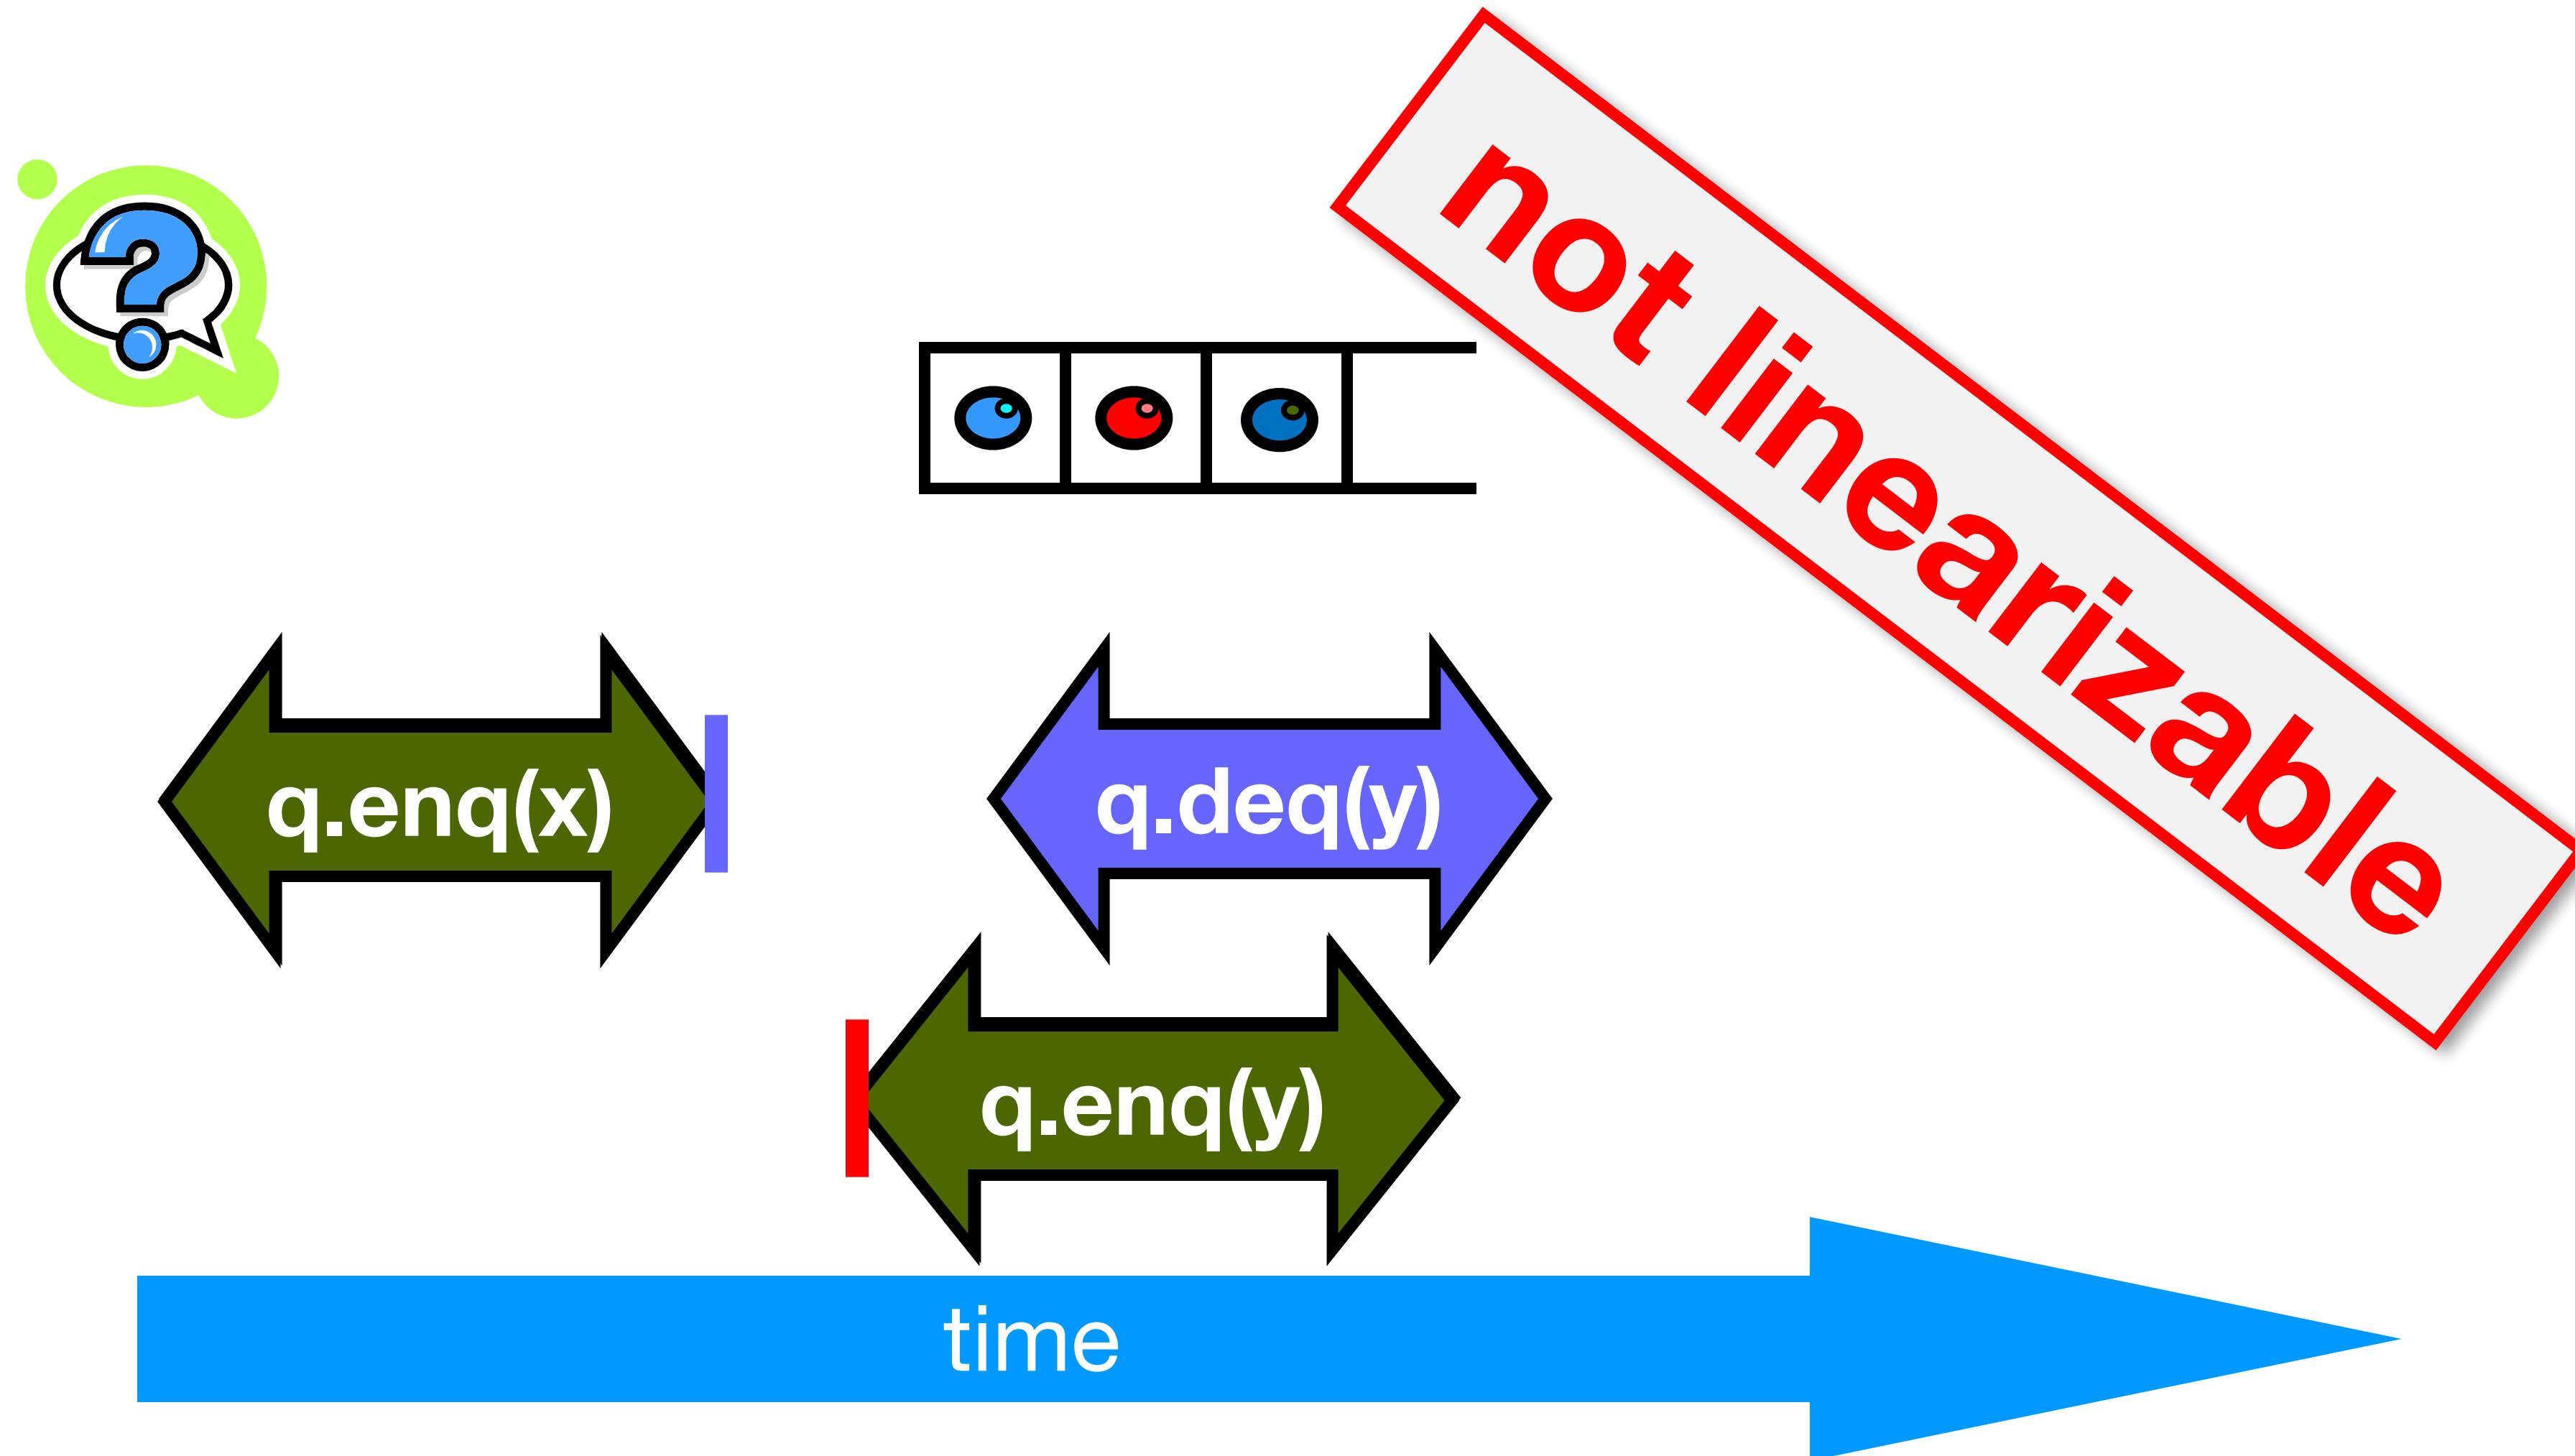
# Example




# Example



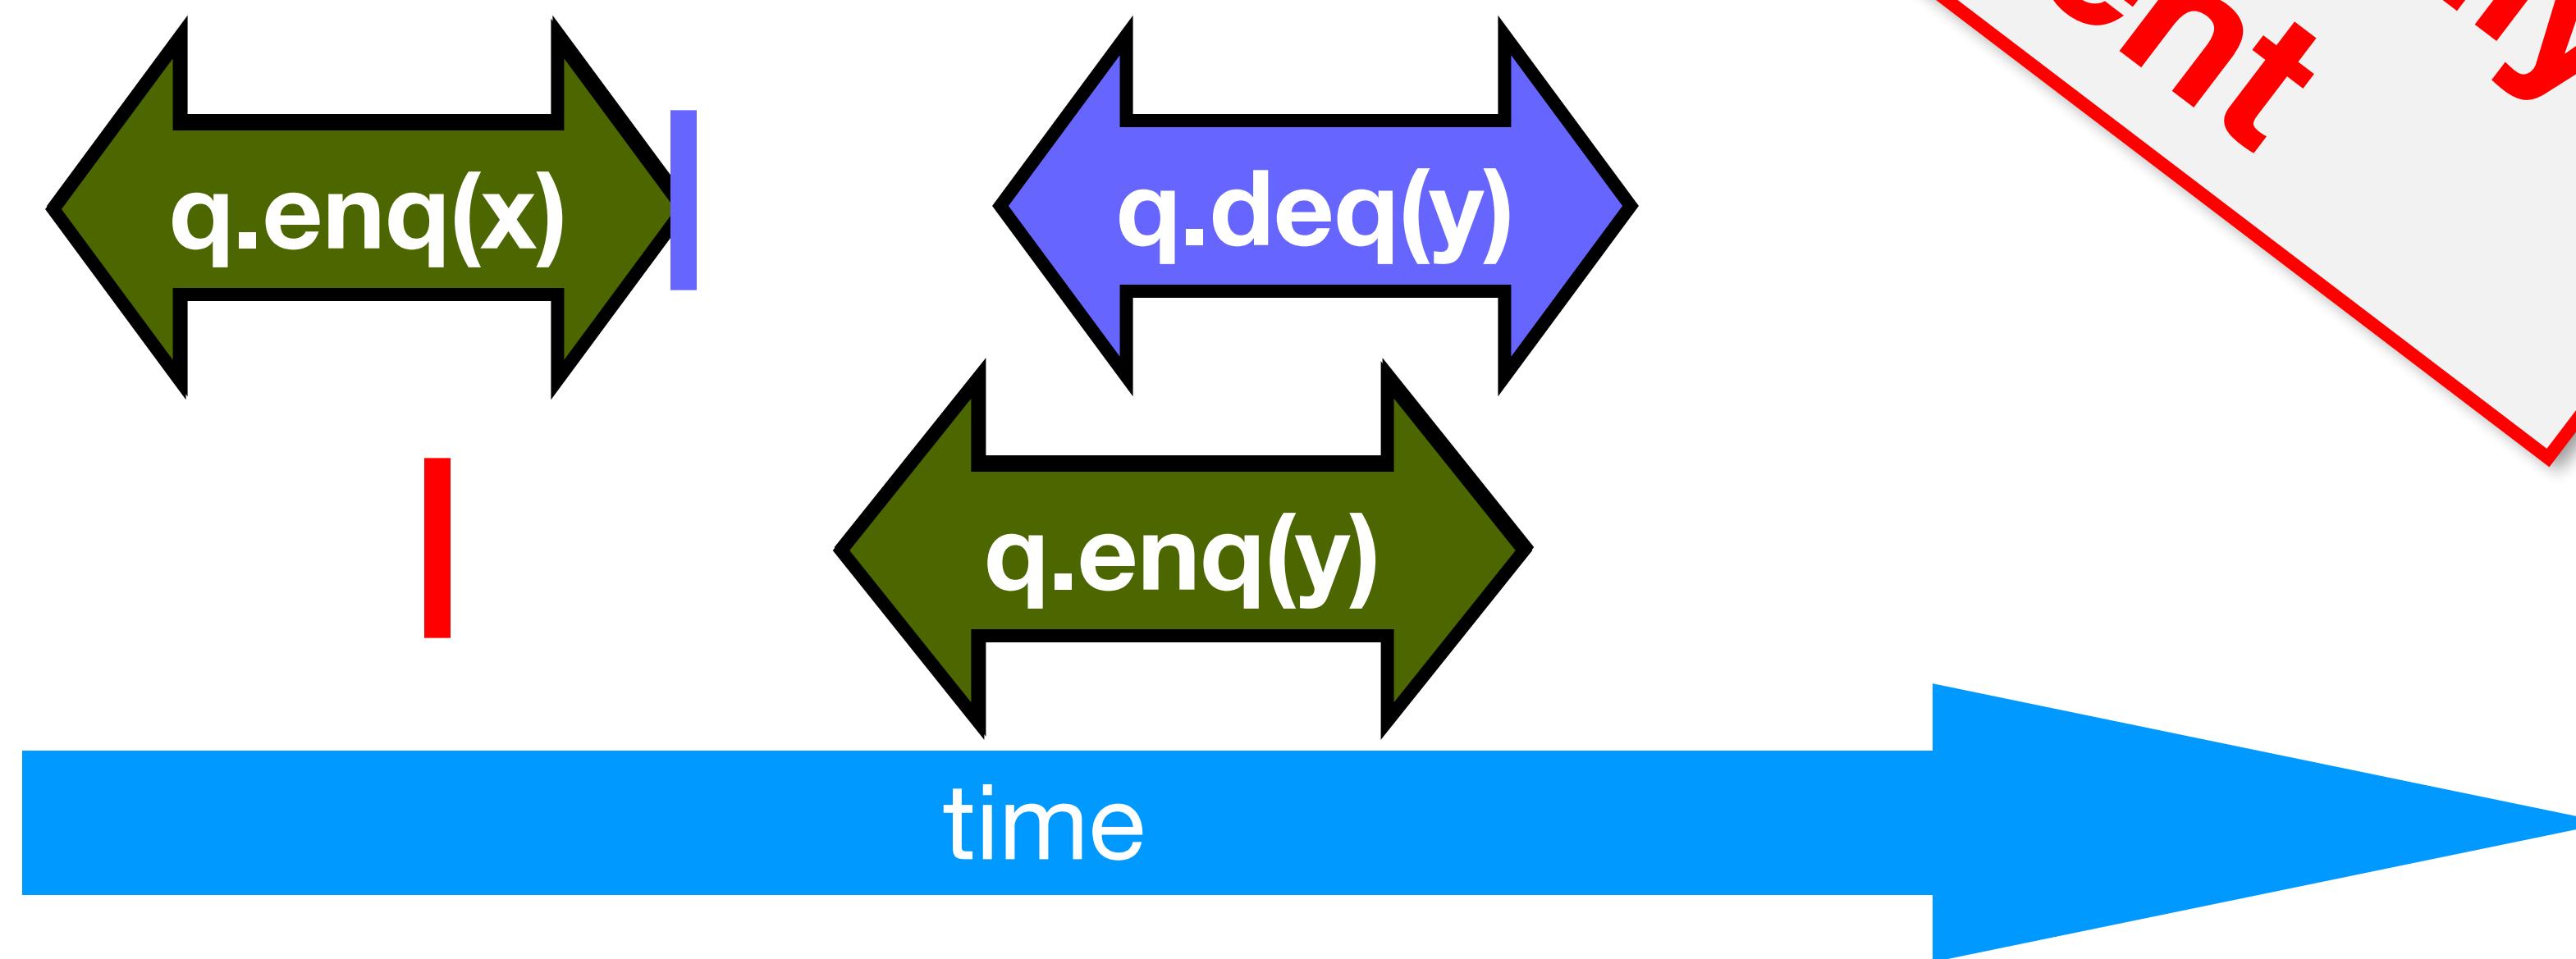
# Example



# Example



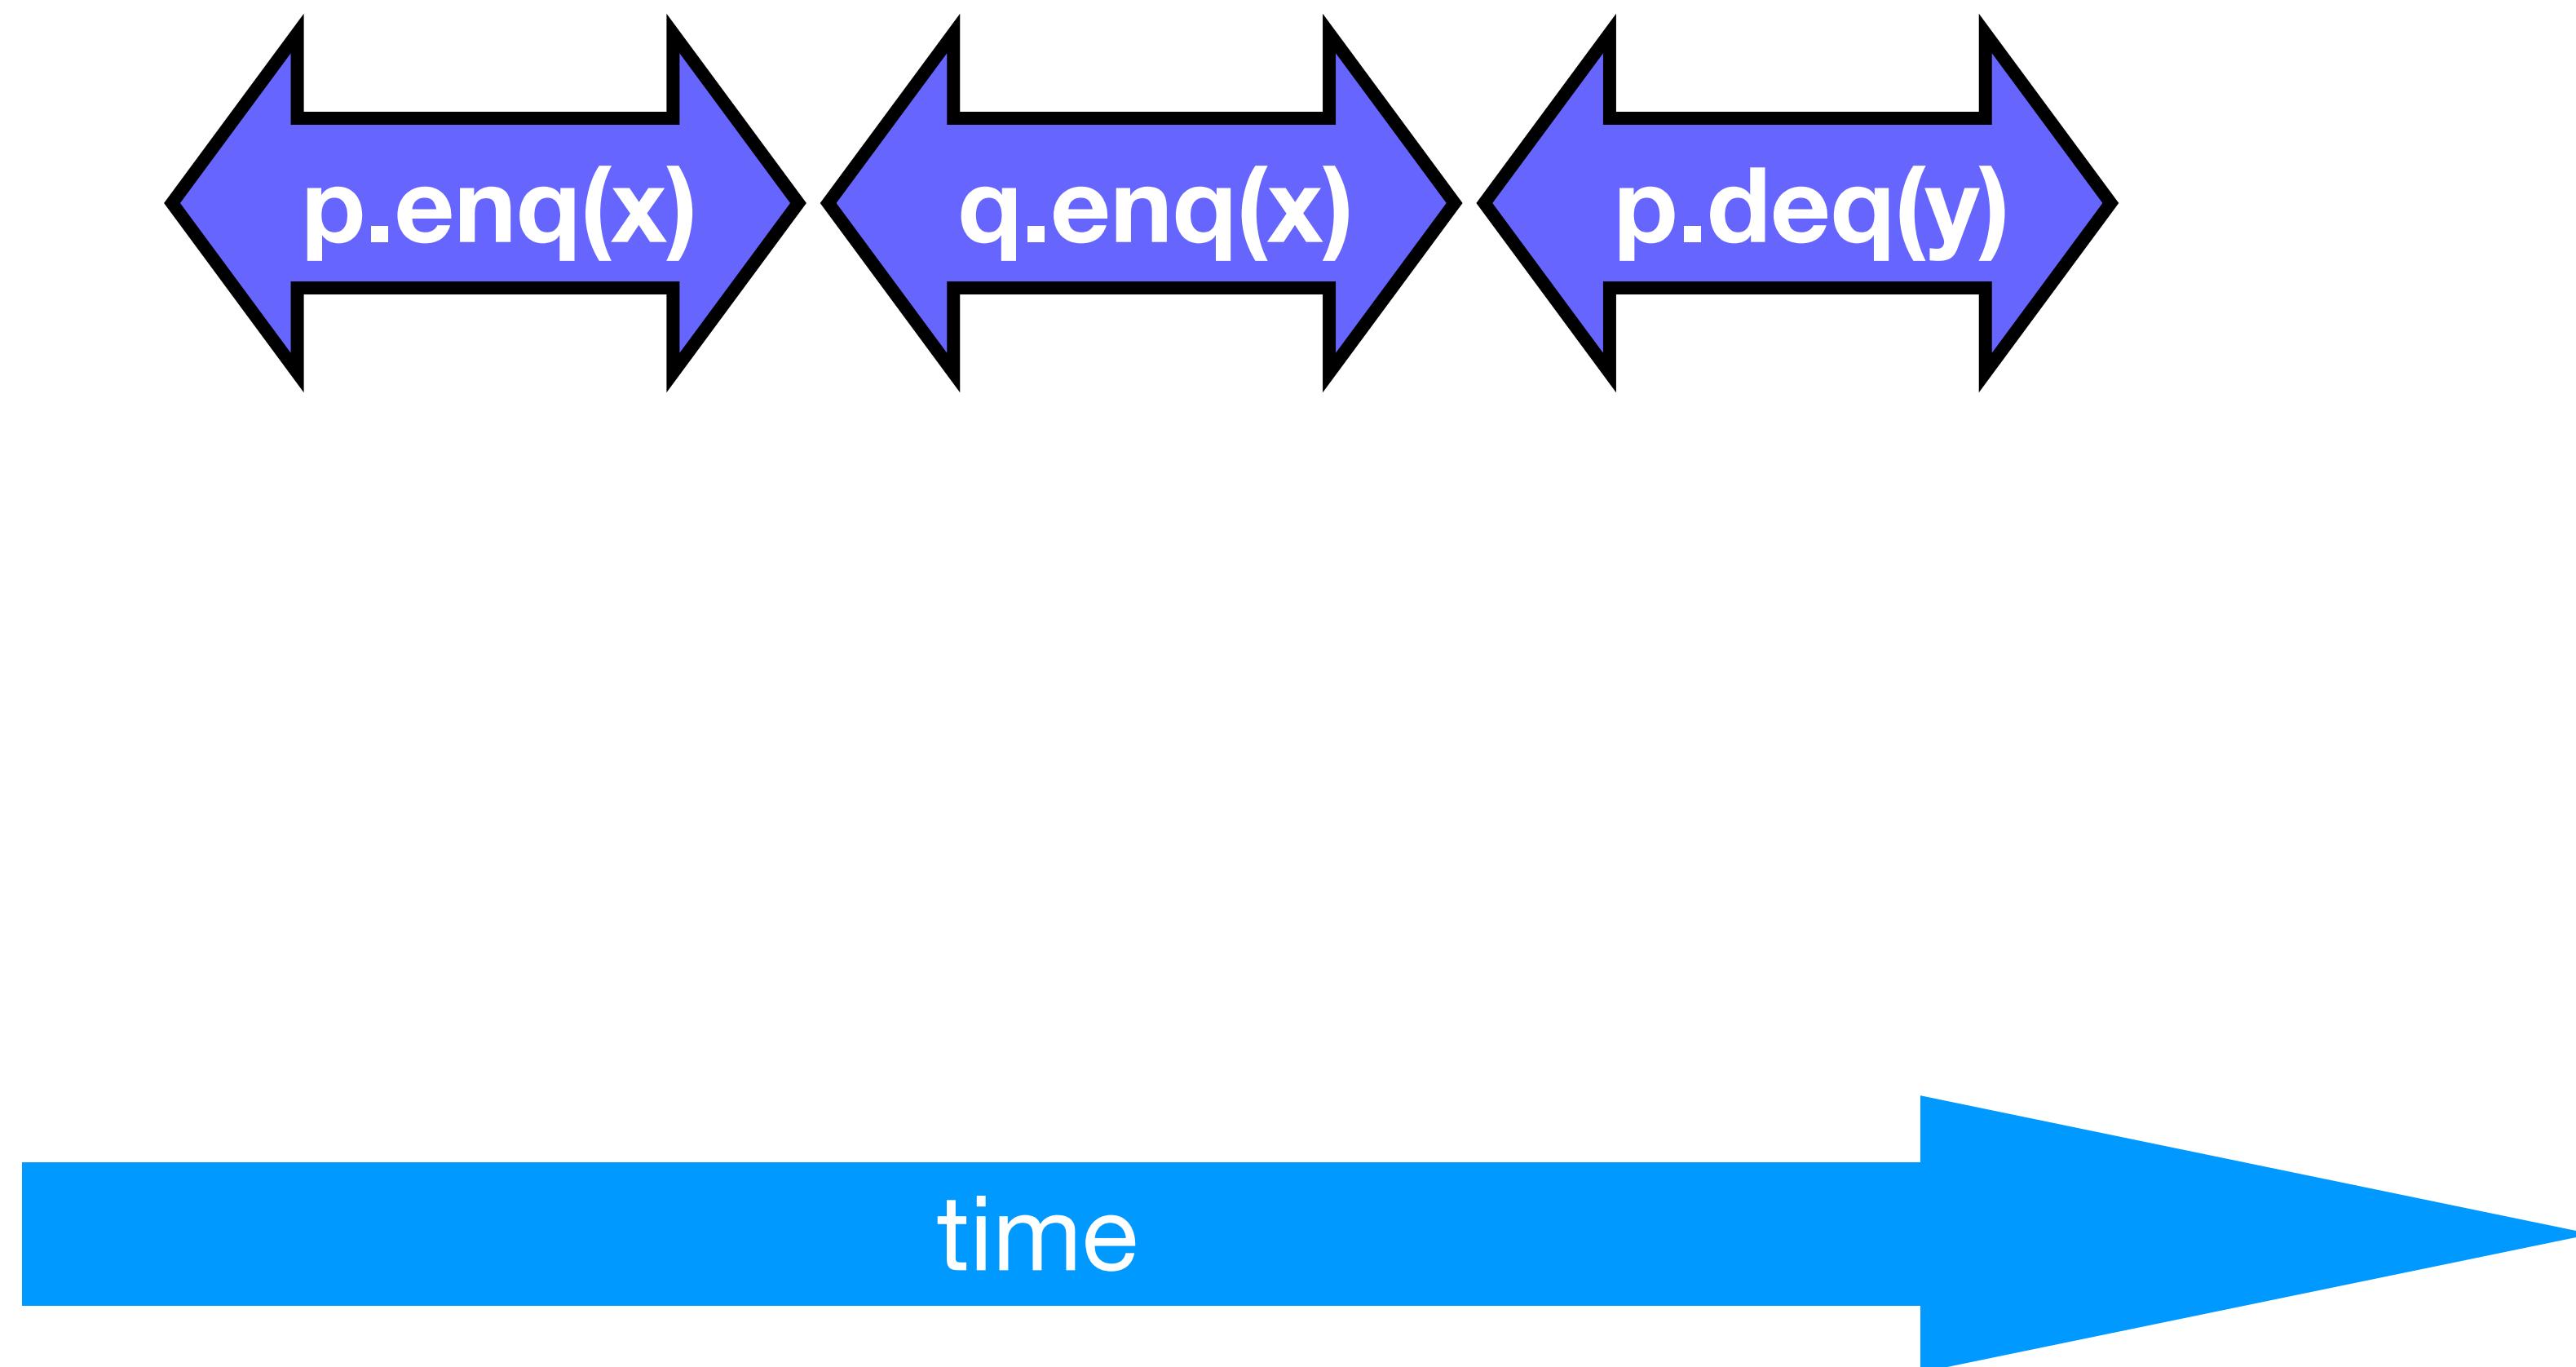
# Example



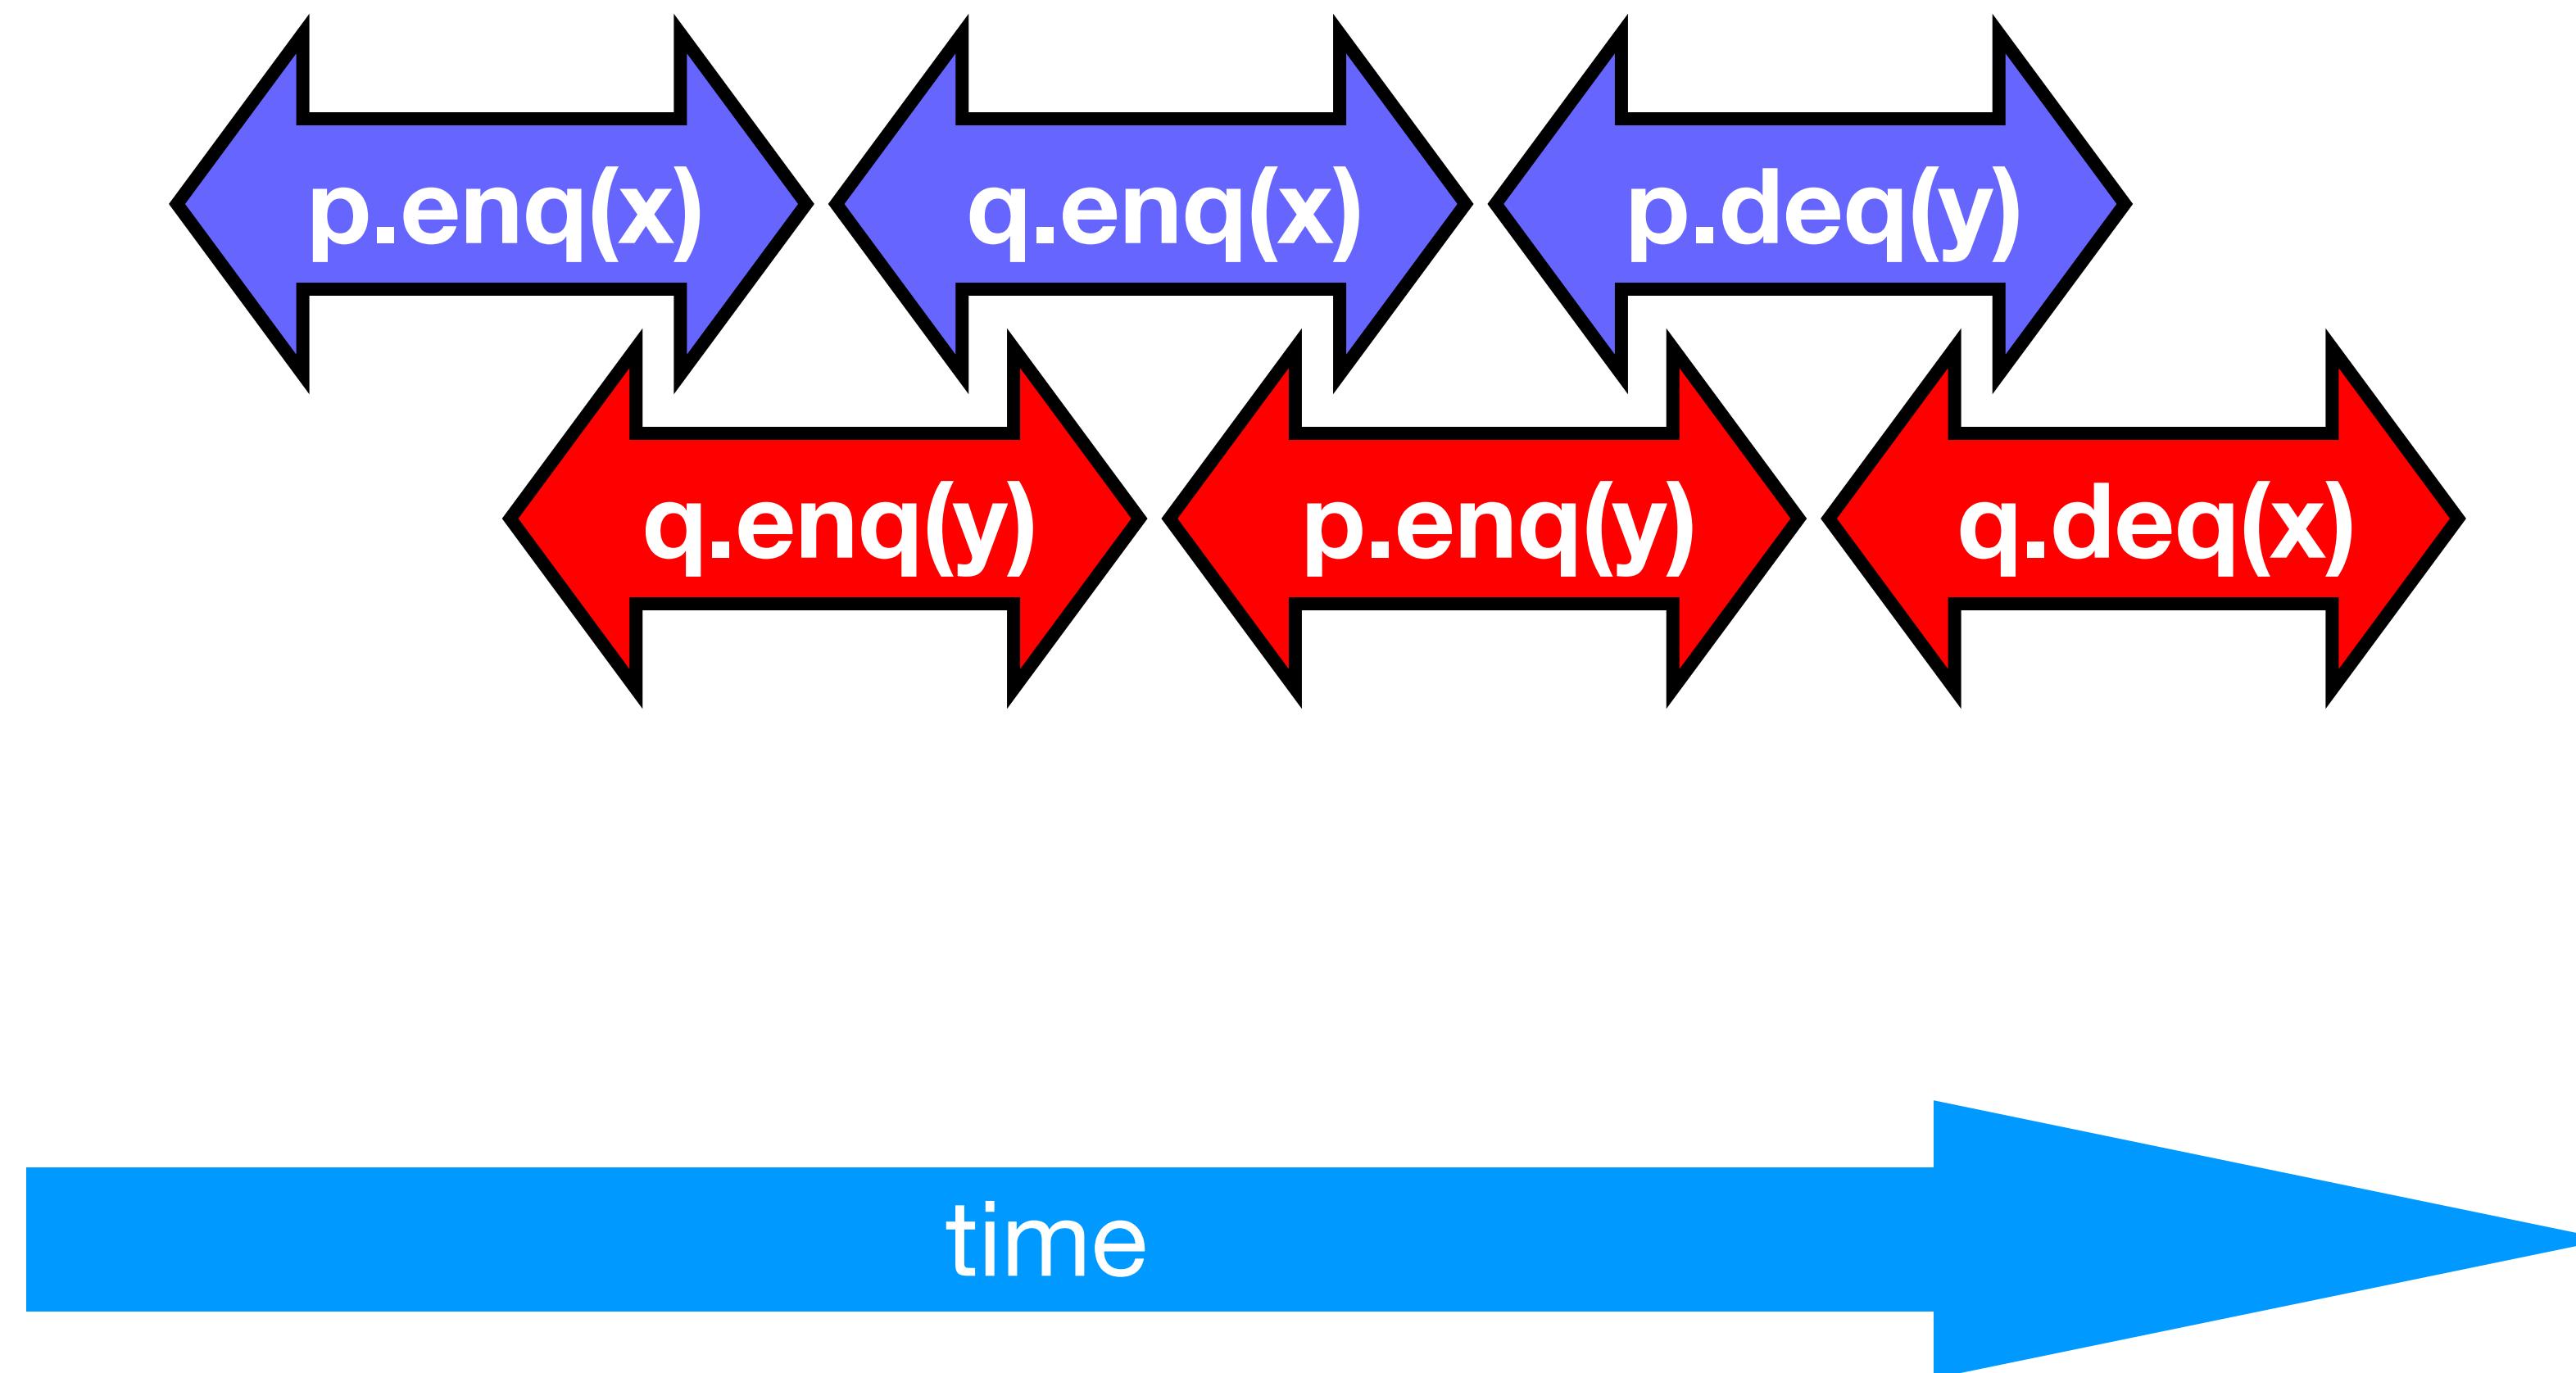
# Theorem

*Sequential Consistency is not Composable*

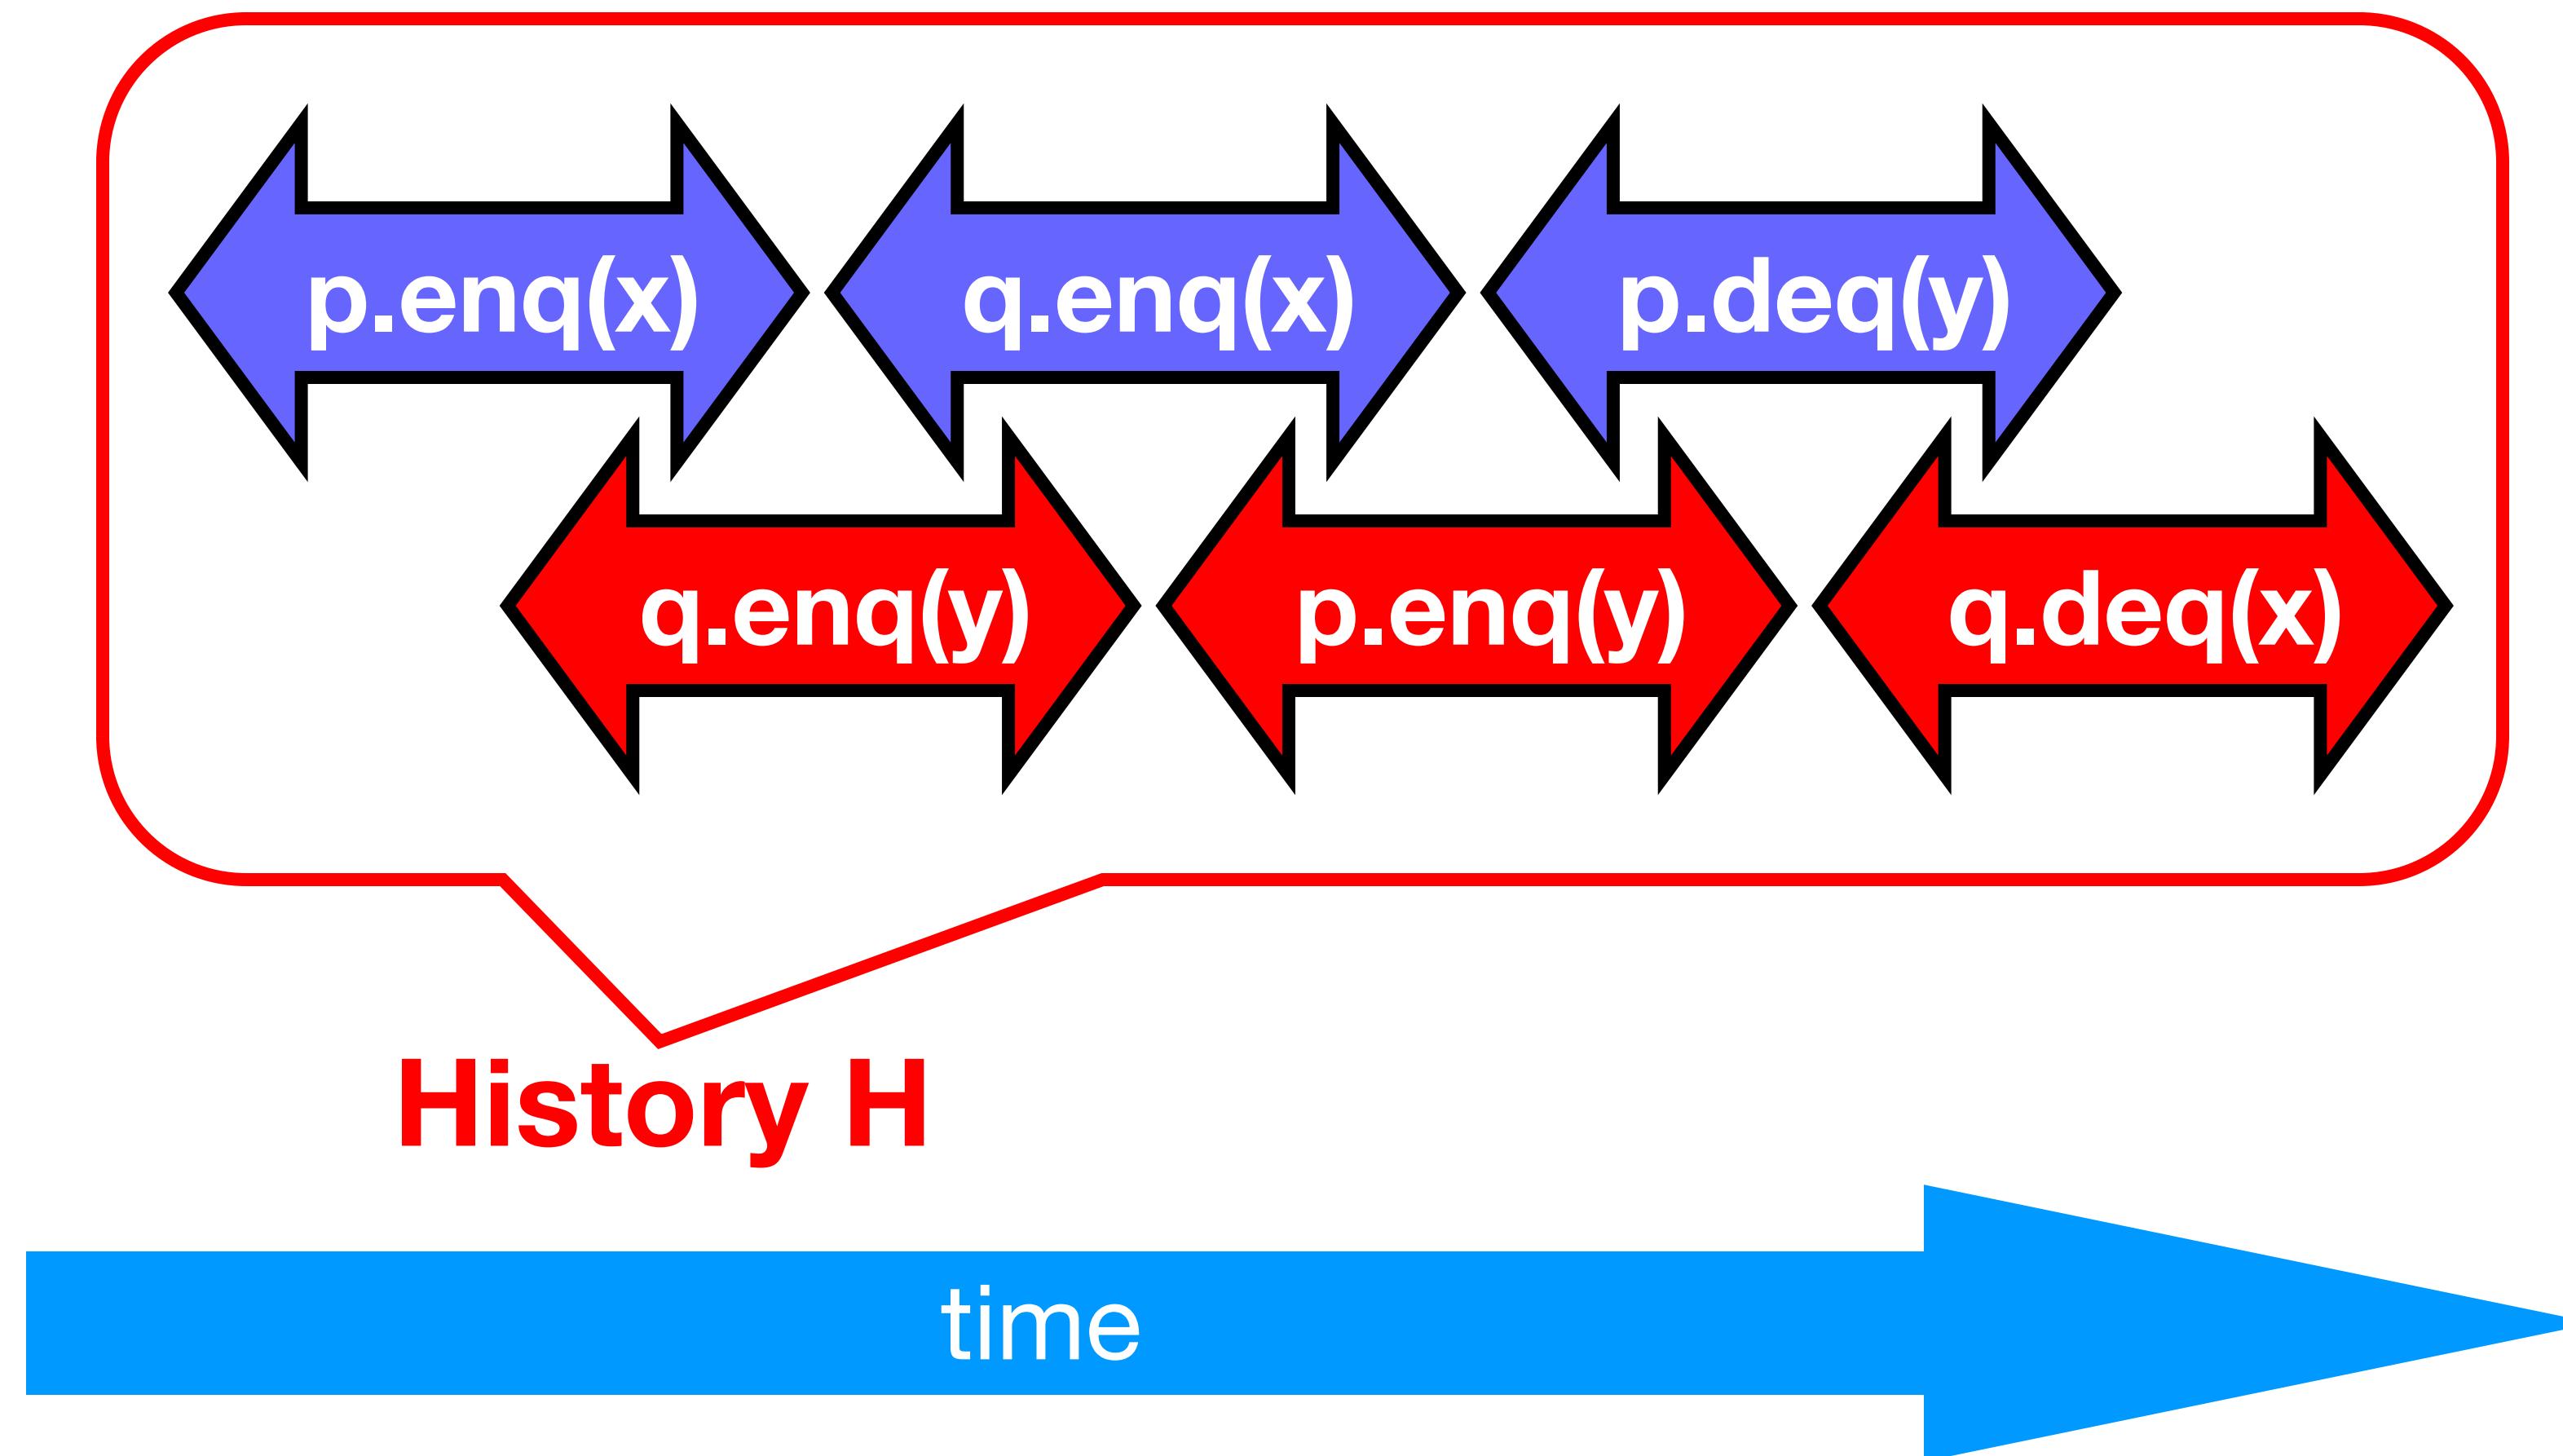
# FIFO Queue Example




# FIFO Queue Example



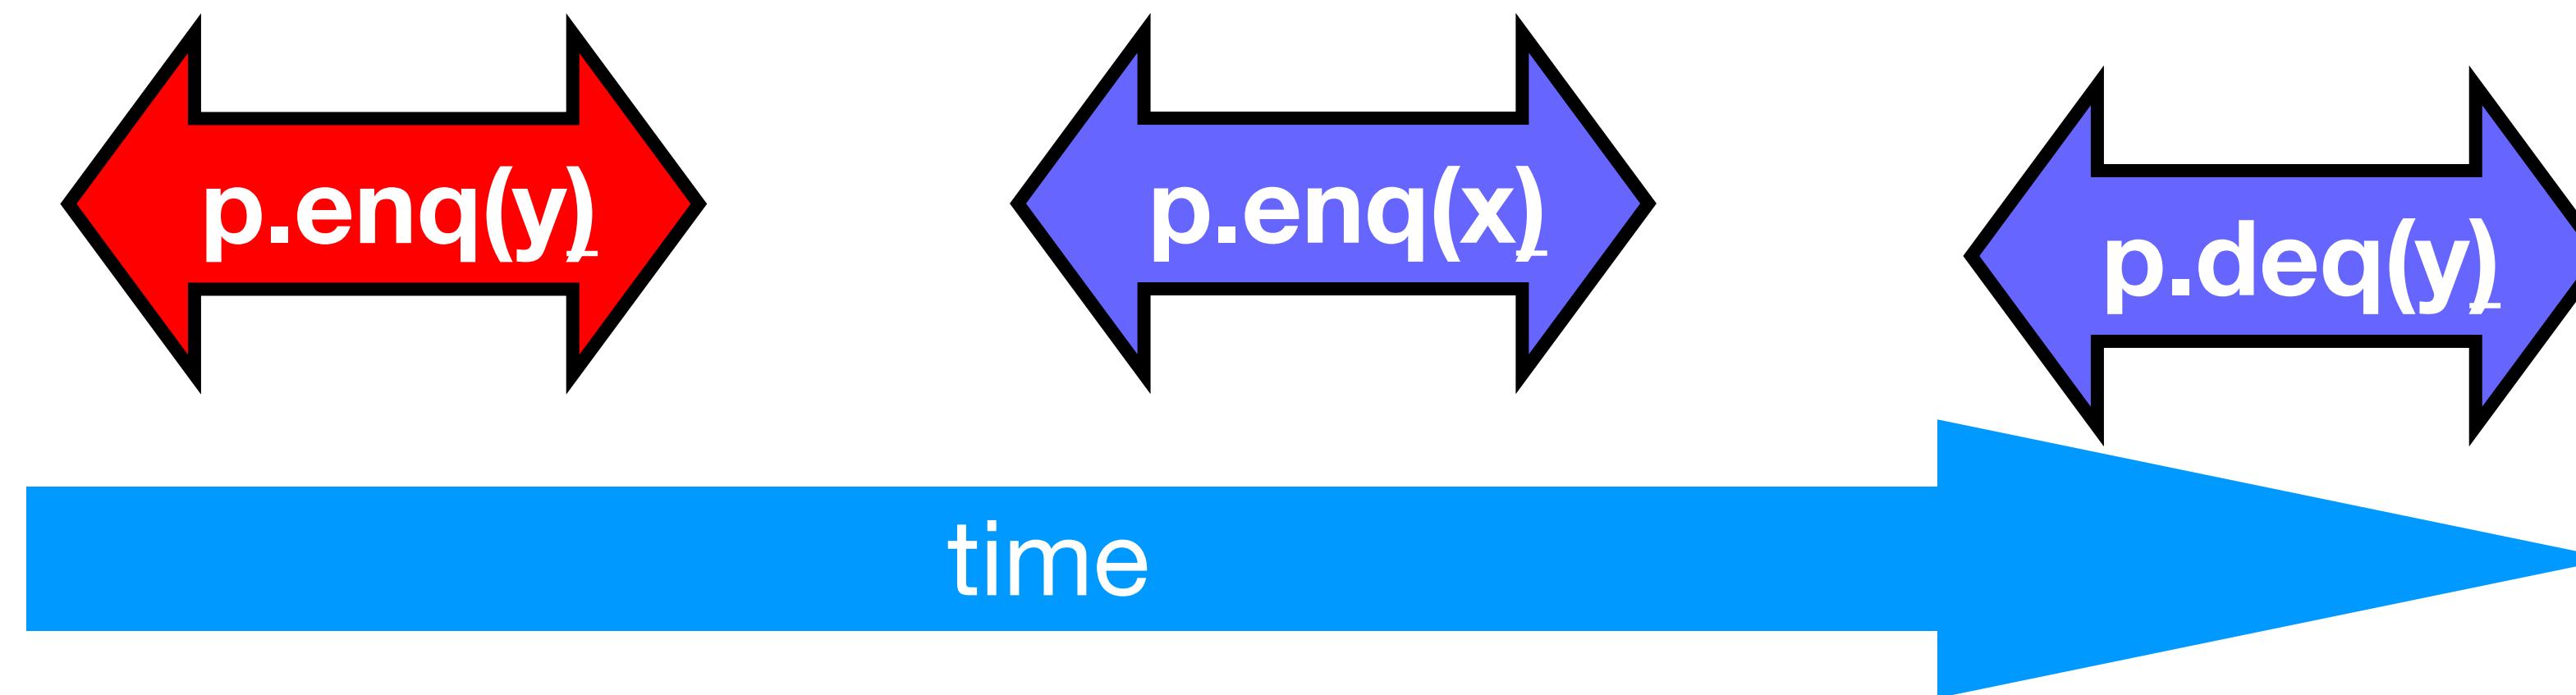
# FIFO Queue Example



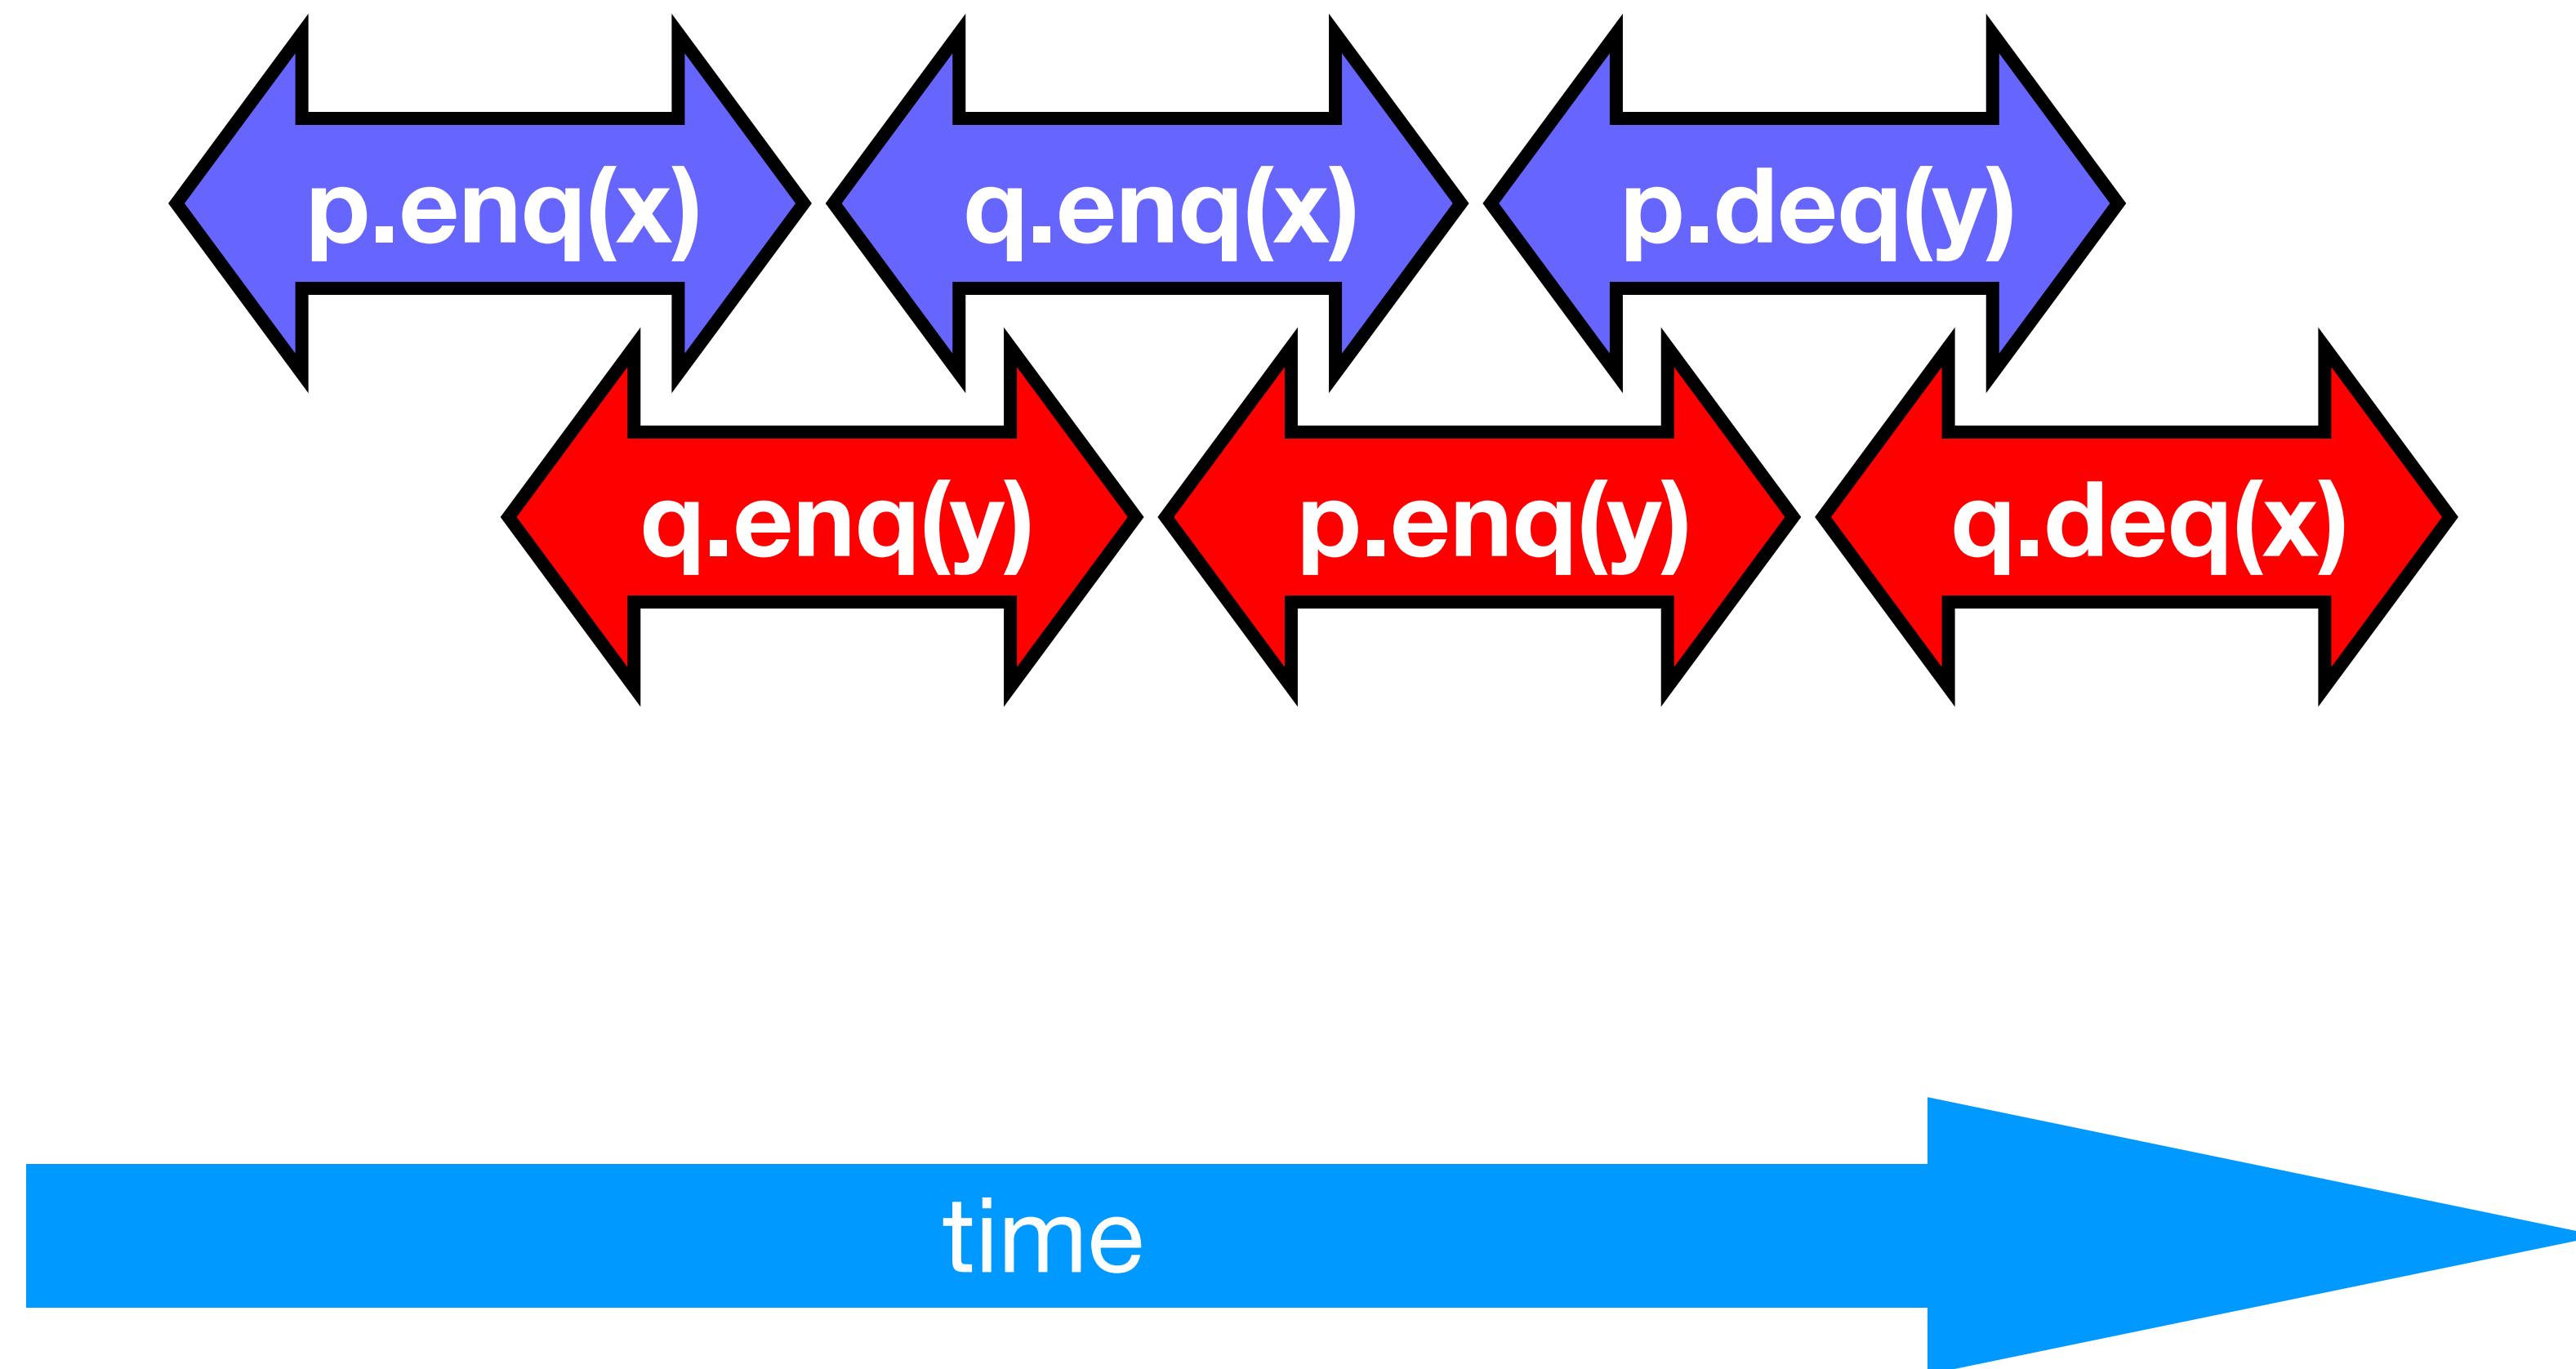
# H|p Sequentially Consistent



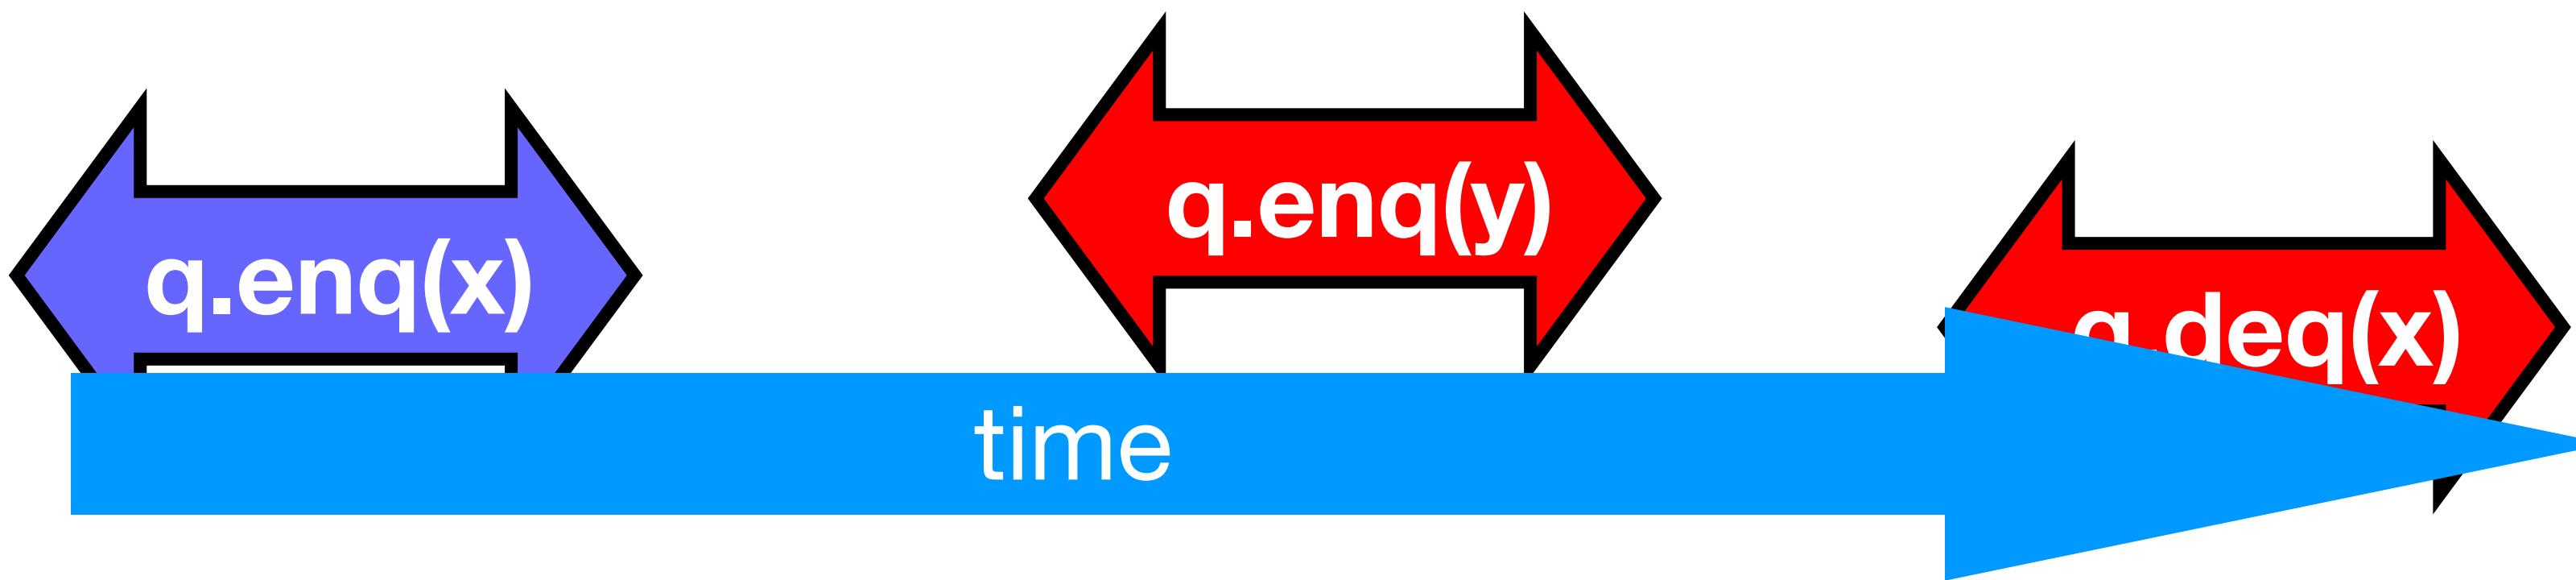
# H|p Sequentially Consistent



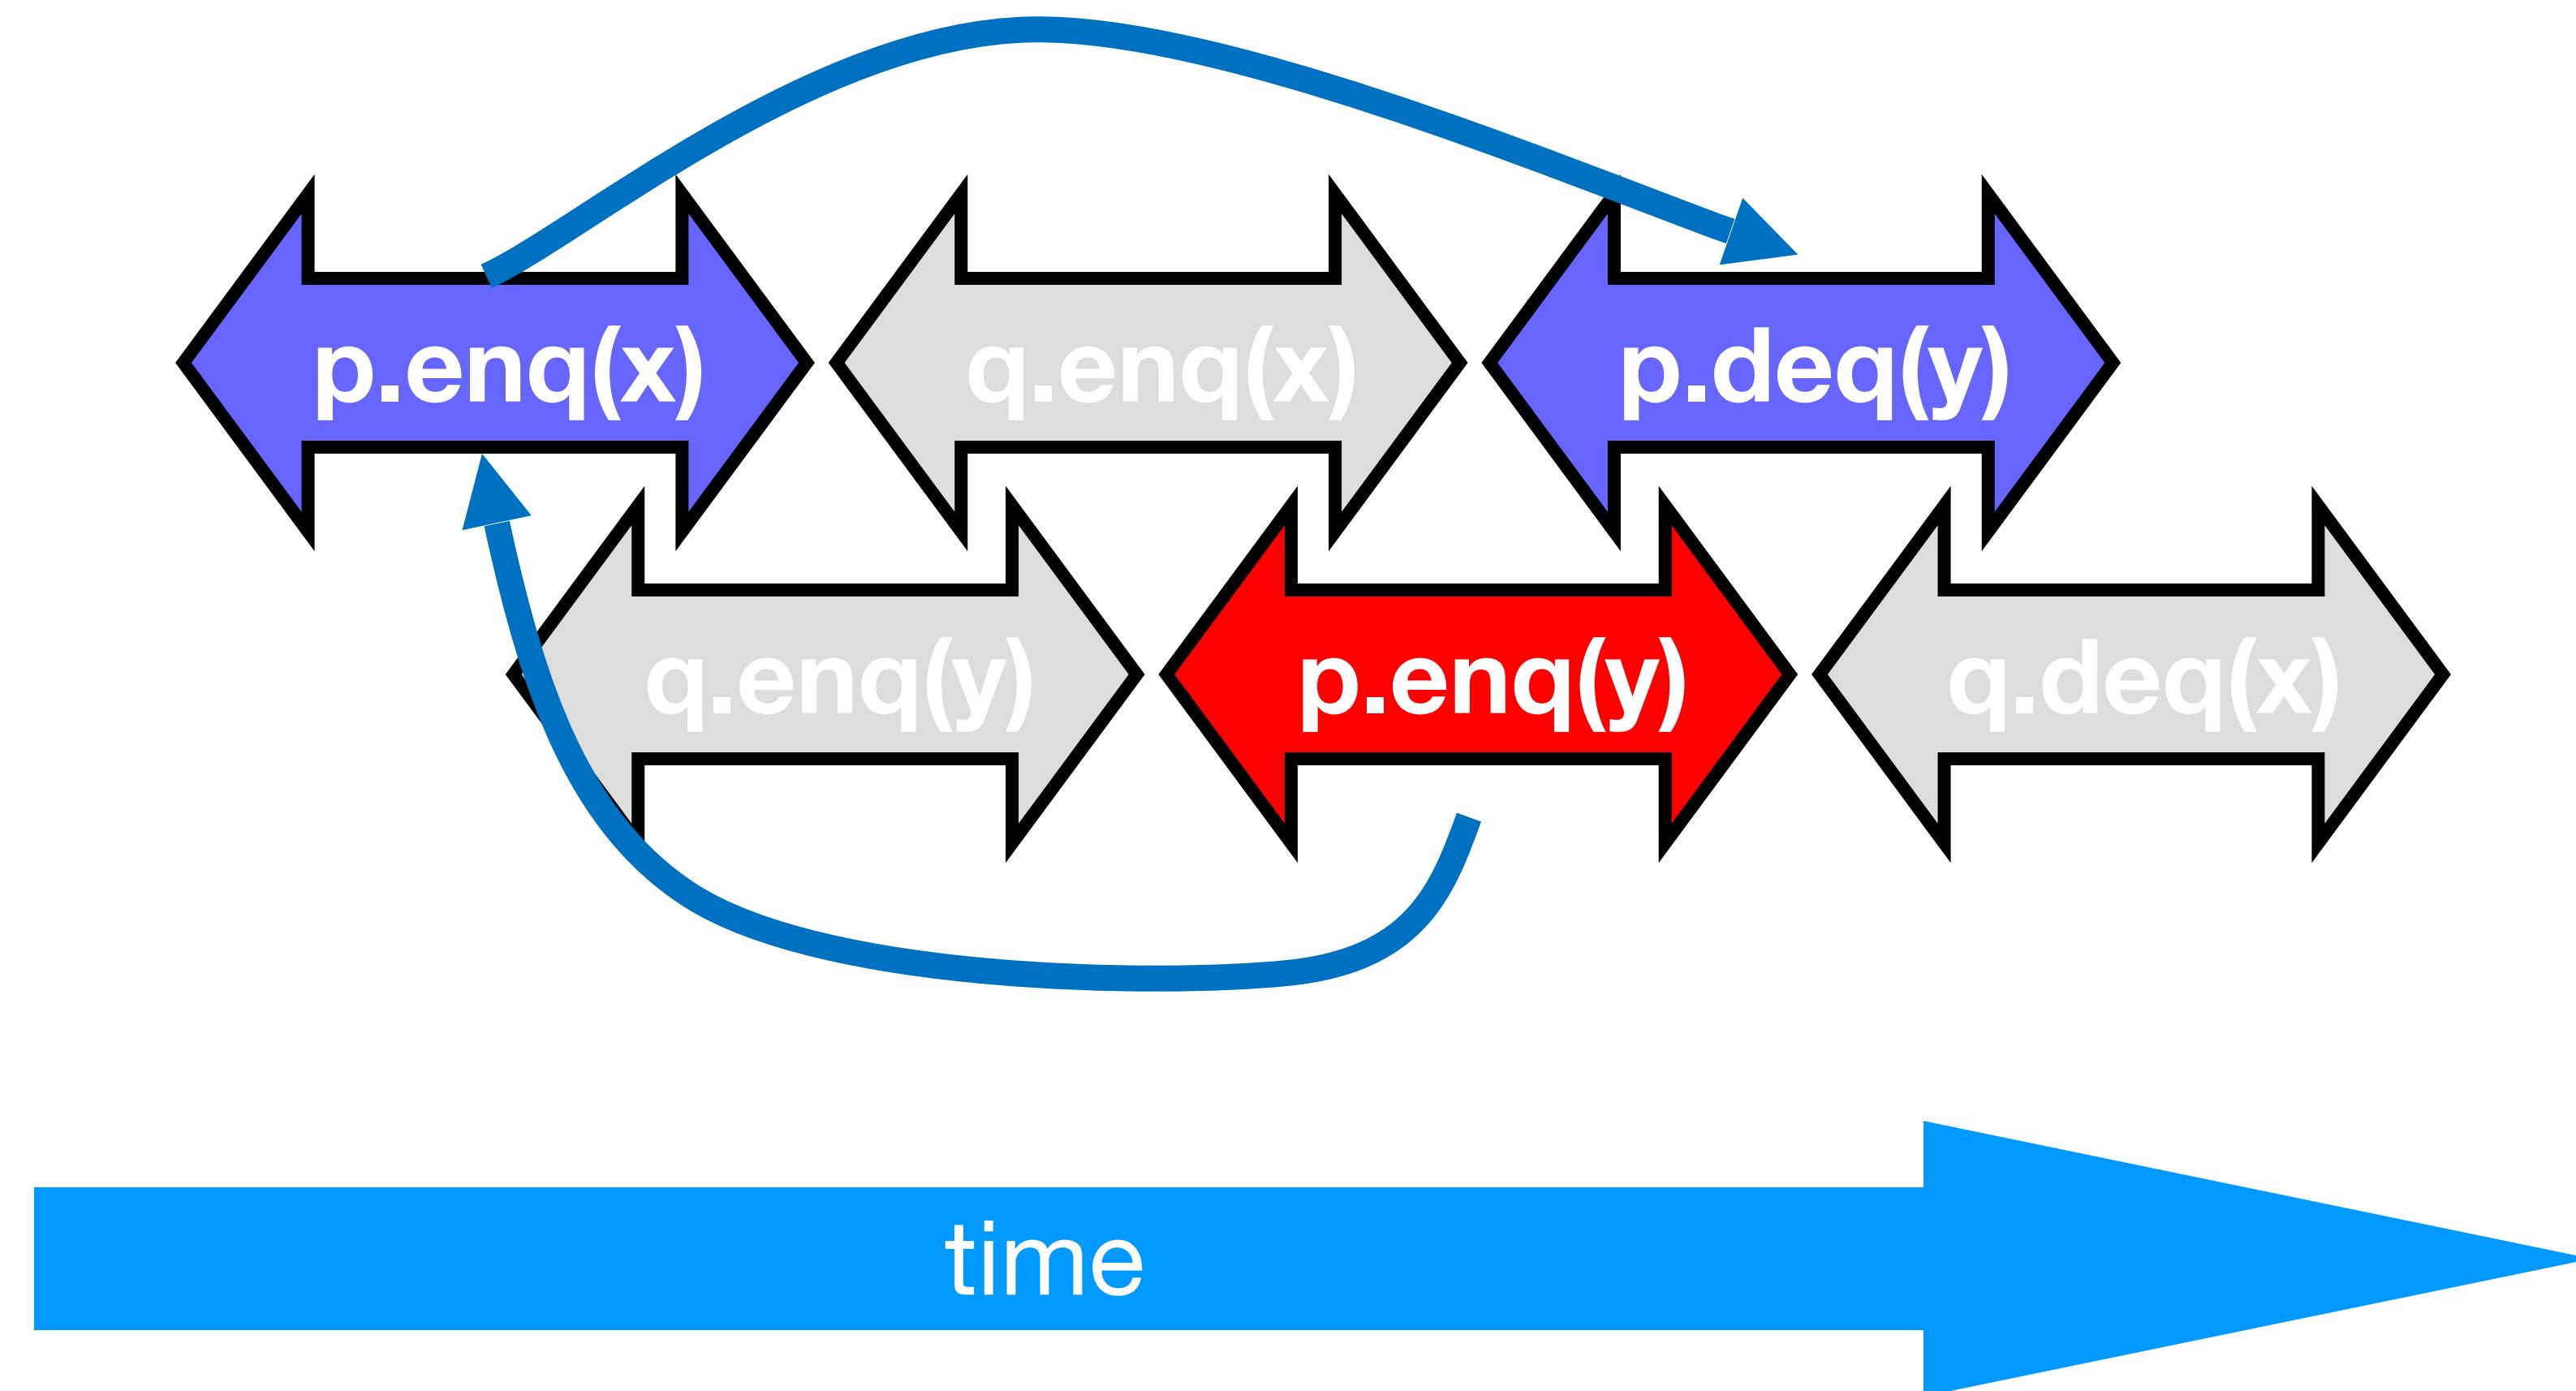
# H|q Sequentially Consistent



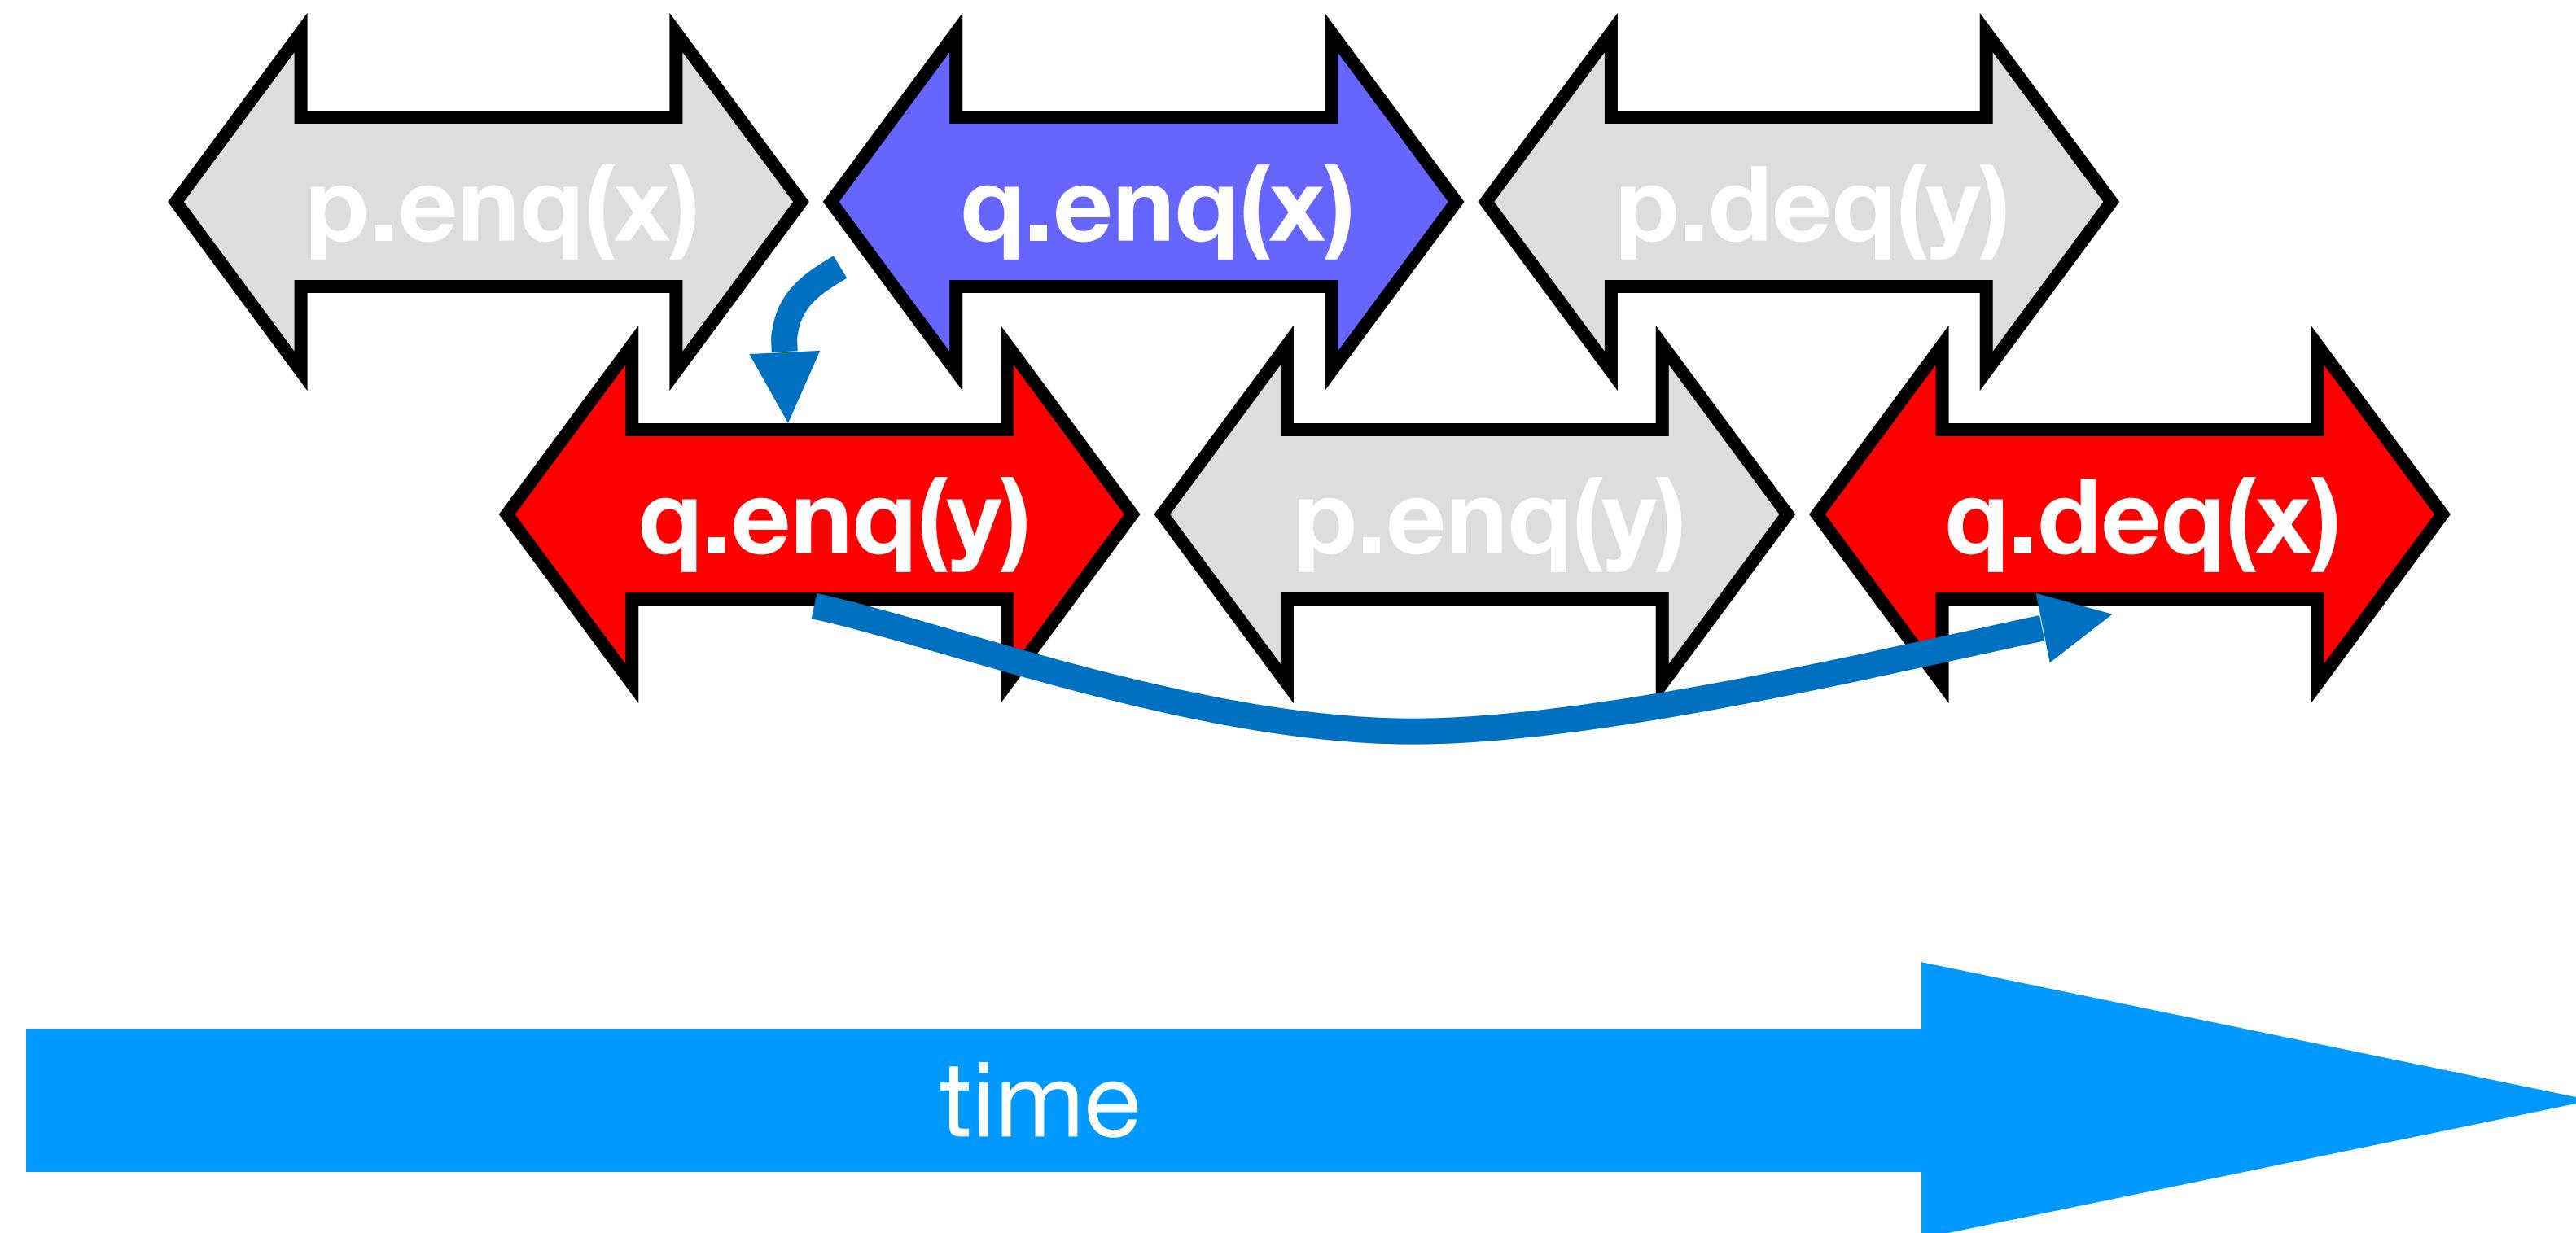
# $H|q$ Sequentially Consistent




# Ordering imposed by p




# Ordering imposed by q



# Ordering imposed by both



# Combining orders



# Concurrency Testing

- Linearizability and Sequential Consistency are good specifications for *testing* the correctness of concurrent data structures.
  - Any observed execution must match a sequential execution
  - Can be exploited for pragmatic testing
- See <https://github.com/ocaml-multicore/multicoretests>

# qcheck-lin

- Checks for sequential consistency violations (despite what the name says)
- ***Every sequential consistency violation is a linearizability violation***
- Check that the observed result of a parallel implementation can be observed with a sequential run

# qcheck-lin

- Checks for sequential consistency violations (despite what the name says)
- ***Every sequential consistency violation is a linearizability violation***
- Check that the observed result of a parallel implementation can be observed with a sequential run

**Demo**

# qcheck-stm

- In *qcheck-lin*, what if the implementation is wrong for the sequential program itself?
  - We're only comparing equivalence.
  - Sequential run of a buggy implementation  $\equiv$  Parallel run of a buggy implementation
    - ***Is not useful!***

# qcheck-stm

- In *qcheck-lin*, what if the implementation is wrong for the sequential program itself?
  - We're only comparing equivalence.
  - Sequential run of a buggy implementation ≡ Parallel run of a buggy implementation
    - ***Is not useful!***
- qcheck-stm
  - Write a state-machine ***model*** of the concurrent object
  - Compare the sequential and parallel executions of the implementation against the state machine model
  - ***More work!***

# qcheck-stm

- In *qcheck-lin*, what if the implementation is wrong for the sequential program itself?
  - We're only comparing equivalence.
  - Sequential run of a buggy implementation ≡ Parallel run of a buggy implementation
    - ***Is not useful!***
- qcheck-stm
  - Write a state-machine ***model*** of the concurrent object
  - Compare the sequential and parallel executions of the implementation against the state machine model
  - ***More work!***

**Demo**

# Summary

- ***Linearizability***
  - The operation takes effect instantaneously between the invocation and the response
  - Uses sequential specification, locality implies composability
- ***Sequential Consistency***
  - Linearizability without real-time ordering
  - Not composable
  - Harder to work with
  - Useful to reason about hardware models (next lecture)
- We will use ***linearizability*** as our consistency condition for reasoning about objects

# Progress

- We saw an implementation whose methods were lock-based (deadlock-free)
- We saw an implementation whose methods did not use locks (lock-free)
- How do they relate?

# Progress Conditions

- **Deadlock-free:** some thread trying to acquire the lock eventually succeeds.
- **Starvation-free:** every thread trying to acquire the lock eventually succeeds.
- **Lock-free:** some thread calling a method eventually returns.
- **Wait-free:** every thread calling a method eventually returns.

# Progress Conditions

- **Deadlock-free:** some thread trying to acquire the lock eventually succeeds.
- **Starvation-free:** every thread trying to acquire the lock eventually succeeds.
- **Lock-free:** some thread calling a method eventually returns.
- **Wait-free:** every thread calling a method eventually returns.

	Non-Blocking	Blocking
Everyone makes progress	Wait-free	Starvation-free
Someone makes progress	Lock-free	Deadlock-free

# Progress Conditions

- **Deadlock-free:** some thread trying to acquire the lock eventually succeeds.
- **Starvation-free:** every thread trying to acquire the lock eventually succeeds.
- **Lock-free:** some thread calling a method eventually returns.
- **Wait-free:** every thread calling a method eventually returns.

	Non-Blocking	Blocking	
Everyone makes progress	Wait-free	Starvation-free	<i>We will look at linearizable blocking and non-blocking implementations of objects.</i>
Someone makes progress	Lock-free	Deadlock-free	



## This work is licensed under a [Creative Commons Attribution-ShareAlike 2.5 License](#).

- **You are free:**
  - **to Share** – to copy, distribute and transmit the work
  - **to Remix** – to adapt the work
- **Under the following conditions:**
  - **Attribution.** You must attribute the work to “The Art of Multiprocessor Programming” (but not in any way that suggests that the authors endorse you or your use of the work).
  - **Share Alike.** If you alter, transform, or build upon this work, you may distribute the resulting work only under the same, similar or a compatible license.
- For any reuse or distribution, you must make clear to others the license terms of this work. The best way to do this is with a link to
  - <http://creativecommons.org/licenses/by-sa/3.0/>.
- Any of the above conditions can be waived if you get permission from the copyright holder.
- Nothing in this license impairs or restricts the author's moral rights.