

04 Memory Consistency Models

CS 6868: Concurrent Programming

KC Sivaramakrishnan

Spring 2026, IIT Madras


Memory Consistency Models

- A memory consistency model is
 - A set of **rules** governing how memory operations from multiple threads are processed
- Affects the **order** in which memory operations will appear to execute
 - determines **what value a read should return**
- A **contract** between a programmer and a system
 - Determines what optimisations can be performed for correct programs
- **Affects:** Ease of programming and performance


Fact

- Most hardware architectures don't support sequential consistency
- Because they think it's ***too strong***
- Here's a story ...

Flag Example




Flag Example



- Each thread's view is sequentially consistent
 - It went first




Flag Example




- Entire history isn't sequentially consistent
 - Can't both go first



Flag Example



- Is this behavior really so wrong?
 - We can argue either way ...



Opinion: It's Wrong

- This pattern
 - Write mine, read yours
- Is exactly the flag principle
 - Beloved of Alice and Bob
 - Heart of mutual exclusion
 - Peterson
 - Bakery, etc.
- *It's non-negotiable!*

Peterson's Algorithm


- Combine ideas from LockOne and LockTwo

```
module Peterson : LOCK = struct
  (* Two boolean flags (from LockOne) and one victim variable (from LockTwo) *)
  let flag = [| false; false |]
  let victim = ref 0

  let lock () =
    let i = (Domain.self () :> int) - 1 in
    let j = 1 - i in
    flag.(i) <- true; (* Announce I'm interested *)
    victim := i; (* Defer to the other *)
    (* Wait while the other thread wants to enter AND we're the victim *)
    while flag.(j) && !victim = i do
      ()
    done

  let unlock () =
    let i = (Domain.self () :> int) - 1 in
    flag.(i) <- false
end
```

Crux of Peterson's proof



WLOG assume that A wrote to victim last

$$\text{write}_B(\text{victim} := B) \rightarrow \text{write}_A(\text{victim} := A) \quad (3)$$

Observation: proof relied on fact that if a location is stored, a later load by some thread will return this or a later stored value.

Unfortunately...

- Many hardware architects think that sequential consistency (SC) is too **strong**
 - Too expensive to implement on modern hardware
- Many standard compiler optimisations are **incorrect** in the presence of multiple threads!
 - Too expensive not to optimise code
- OK if flag principle
 - violated by **default**
 - Honoured by explicit **request**
- ***Relaxed memory models***
 - Models weaker than sequential consistency (SC)

Relaxed memory models

- Models weaker than sequential consistency (SC)
- Both hardware and programming languages go for weaker-than-SC by default
 - Still provide practical *recipes* for recovering SC
- In this lecture,
 - Study one hardware (relaxed) memory model – **x86**
 - Study one programming language (relaxed) memory model – **OCaml**

Hardware Memory Consistency Models

Hardware Consistency

Initially, $a = b = 0$.

Processor 0

```
mov 1, a      ;Store  
mov b, %ebx ;Load
```

Processor 1

```
mov 1, b      ;Store  
mov a, %eax ;Load
```


What are the final possible values of **%eax** and
%ebx after both processors have executed?

Sequential consistency implies that no execution
ends with $\%eax = \%ebx = 0$

Hardware Consistency


- **No** modern-day processor implements sequential consistency.
 - Hardware actively reorders instructions.
 - Compilers may reorder instructions, too.
 - Why?
 - Most of the performance is derived from a single thread's unsynchronised execution of code.

Instruction Reordering




- Q. Why might the hardware or compiler decide to reorder these instructions?
- A. To obtain higher performance by covering load latency – ***instruction-level parallelism***.

Instruction Reordering



- Q. When is it safe for the hardware or compiler to perform this reordering?
- A. When $a \neq b$.
- A'. And there's no concurrency.


Hardware Reordering



- Processor can issue stores faster than the network can handle them \Rightarrow store buffer
- Loads take priority, bypassing the store buffer.
- Except if a load address matches an address in the store buffer, the store buffer returns the result.

X86 Memory Consistency


Thread's Code



1. Loads **are *not* reordered with loads.**
2. Stores **are *not* reordered with stores.**
3. Stores **are *not* reordered with prior loads.**
4. A load ***may* be reordered with a prior store to a different location but *not* with a prior store to the same location.**
5. Stores to the same location **respect a global total order.**

X86 Memory Consistency

Thread's Code



1. Loads **are not reordered with loads.**
2. Stores **are not reordered with stores.**
3. **Total Store Ordering (TSO)...weaker than sequential consistency**
or
4. ~~Stores to different location but not with a prior store to the same location.~~
- OK!** 5. Stores to the same location **respect a global total order.**

Memory Barriers (Fences)

- A *memory barrier* (or *memory fence*)
 - is a hardware action that
 - enforces an ordering constraint between the instructions before and after the fence.
- A memory barrier can be issued explicitly as an instruction
 - (x86: mfence)
- The typical cost of a memory fence is comparable to that of an L2-cache access.

X86 Memory Consistency

Thread's Code

```
Store1  
Store2  
Load1  
Load2  
Store3  
Store4  
Barrier  
Load3  
Load4  
Load5
```

1. Loads **are not reordered with loads.**
2. Stores **Total Store Ordering + properly placed memory barriers = sequential consistency**
3. Stores **to the same location respect a global total order.**
4. A load ~~can be reordered with a prior store to the same location but not with a prior store to the same location.~~
5. Stores ~~to the same location respect a global total order.~~

Memory Barriers

- Explicit Synchronization
- Memory **barrier** will
 - Flush write buffer
 - Bring caches up to date
- Compilers often do this for you
 - Entering and leaving critical sections

Summary: Hardware Consistency

- Hardware memory models are weaker than sequential consistency
 - They show up as reordering of instruction sequences on a single thread
- X86 uses Total Store Order (TSO)
 - ARM, RISC-V and Power uses **even weaker models!**
- Instruction Set Architectures (ISAs) provide fences/memory barriers to restrict reordering

*What does it mean for concurrent
programming languages?*

PL Memory Consistency Models

<https://ocaml.org/manual/5.4/memorymodel.html>

PL Consistency

- We saw that hardware architectures do not provide SC
 - Programming languages compile to hardware
 - Hence, provide weaker than SC model
- In addition, ***compiler optimisations also violate SC!***

Example

Initially $\text{!a} = \text{!b} = 1$

```
let t1 a b =
  let r1 = !a * 2 in
  let r2 = !b in
  let r3 = !a * 2 in
  (r1, r2, r3)
```

```
let t2 b = b := 0
```

- What are the possible outcomes for $r1$, $r2$ and $r3$?
- $r1 = 2$, $r2 = 1$, $r3 = 2$
 - *if the t1 reads b before the t2 writes 0 to it*
- $r1 = 2$, $r2 = 0$, $r3 = 2$
 - *if the t1 reads b after the t2 writes 0 to it*

Example

```
let t1 a b =
  let r1 = !a * 2 in
  let r2 = !b in
  let r3 = !a * 2 in
  (r1, r2, r3)
```

```
let t2 b = b := 0
```

a and b are aliases

```
let main () =
  let ab = ref 1 in
  let h = Domain.spawn (fun _ ->
    let r1, r2, r3 = t1 ab ab in
    assert (not (r1 = 2 && r2 = 0 && r3 = 2)))
  in
  t2 ab;
  Domain.join h
```

Will this assertion ever fail?

- What are the possible outcomes for r1, r2 and r3?
- r1 = 2, r2 = 1, r3 = 2
 - All reads happen before the write ($ab = 1$ throughout)
- r1 = 2, r2 = 0, r3 = 0
 - First read before write, then write happens, and remaining reads after write
- r1 = 0, r2 = 0, r3 = 0
 - All reads happen after the write ($ab = 0$ throughout)

Common Subexpression Elimination (CSE)

```
let t1 a b =
  let r1 = !a * 2 in
  let r2 = !b in
  let r3 = r1 in (* CSE: !a * 2 ==> r1 *)
  (r1, r2, r3)
```

- $r1 = 2, r2 = 0, r3 = 2$
- First read before write, then write happens, and the remaining read after write

```
let t2 b = b := 0
```

*a and b are
aliases*

```
let main () =
  let ab = ref 1 in
  let h = Domain.spawn (fun _ ->
    let r1, r2, r3 = t1 ab ab in
    assert (not (r1 = 2 && r2 = 0 && r3 = 2)))
  in
  t2 ab;
  Domain.join h
```

*Will this assertion
ever fail?*

Explaining this behaviour

```
let t1 a b =  
  let r1 = !a * 2 in  
  let r3 = !a * 2 in  
  let r2 = !b in  
(r1, r2, r3)
```

Reordering

- $r1 = 2, r2 = 0, r3 = 2$
- Write happens after the first two reads

```
let t2 b = b := 0
```

```
let main () =  
  let ab = ref 1 in  
  let h = Domain.spawn (fun _ ->  
    let r1, r2, r3 = t1 ab ab in  
    assert (not (r1 = 2 && r2 = 0 && r3 = 2)))  
  in  
  t2 ab;  
  Domain.join h
```

Common Subexpression Elimination (CSE)

```
let t1 a b =
  let r1 = !a * 2 in
  let r3 = !a * 2 in
  let r2 = !b in
(r1, r2, r3)
```

```
L104: ldr  x2, [x0, #0] ; Load !a (ONCE)
L106: add  x4, x3, x2, lsl #1 ; r1 = !a * 2
L107: ldr  x5, [x1, #0] ; Load !b (r2)
L114: str  x4, [x0, #0] ; Store r1
L115: str  x5, [x0, #8] ; Store r2
L116: str  x4, [x0, #16] ; Store r3 (SAME x4!)
```

; CSE: !a loaded once, x4 reused for r1 and r3
; No second load or computation for r3

OCaml

Assembly

OCaml Relaxed Memory Model

- Precisely describes what behaviours can be *observed*
 - Analogous to x86 TSO, but for the source language
- Compiler and hardware can *optimise* given they respect OCaml memory model
- OCaml also provides *language constructs* to enforce SC

Atomic locations

- **Non-atomic** locations are default
 - Reference cells, array fields, mutable record fields, ...
 - Immutable locations only have an *initialising write* and only reads afterwards
- **Atomic** locations are created using the Atomic module
- They're like reference cells in OCaml, but something more...
 - Include memory fences to enforce ordering against hardware memory model
 - Disables unsafe compiler optimisations that break SC

Module Atomic

```
module Atomic: sig .. end
```

Atomic references.

See [the examples](#) below. See 'Memory model: The hard bits' chapter in the manual.

Since 4.12

```
type !'a t
```

An atomic (mutable) reference to a value of type '`'a`'.

```
val make : 'a -> 'a t
```

Create an atomic reference.

```
val get : 'a t -> 'a
```

Get the current value of the atomic reference.

```
val set : 'a t -> 'a -> unit
```

Set a new value for the atomic reference.

Reasoning about OCaml memory model

- How do we understand the behaviour of OCaml programs under the relaxed memory model?
- **Events**
 - Reads and writes of atomic and non-atomic locations
 - Spawn and join of domains
 - Operations on mutexes

Happens-before relation

- **Happens-before relation** is a **reflexive, transitive closure** of
 - **Program order** – order in which events appear in a given domain

```
1: let t1 a b =
2:   let r1 = !a * 2 in
3:   let r2 = !b in
4:   let r3 = !a * 2 in
5:   (r1, r2, r3)
```

- $R(a) @ 2 \xrightarrow{hb} R(b) @ 3$
- $R(b) @ 3 \xrightarrow{hb} R(a) @ 4$
- $R(a) @ 2 \xrightarrow{hb} R(a) @ 4$

Happens-before relation

- **Happens-before relation** is a **reflexive, transitive closure** of
 - **Program order** – order in which events appear in a given domain
 - On an atomic location, a write and a subsequent
 - Read that sees the written value or
 - Write that overwrites the original write

```
let a = Atomic.make 1
let t1 () = Atomic.set a 2
let t2 () = Atomic.get a
let t3 () = Atomic.get a

• R(a) @ t1  $\xrightarrow{hb}$  R(a) @ t2
• R(a) @ t1  $\xrightarrow{hb}$  R(a) @ t3
```

Happens-before relation

- **Happens-before relation** is a **reflexive, transitive closure** of
 - **Program order** – order in which events appear in a given domain
 - On an atomic location, a write and a subsequent
 - Read that sees the written value or
 - Write that overwrites the original write
 - **Domain.spawn** and the first event of the newly spawned domain
 - Last event of domain and **Domain.join**


```
let d1 = Domain.spawn t1
      ↘ hb
let t1 () =
  Atomic.set a 2
...
...
...
Domain.join d1;
```

```
... ↗ hb
```

Happens-before relation

- **Happens-before relation** is a **reflexive, transitive closure** of
 - **Program order** – order in which events appear in a given domain
 - On an atomic location, a write and a subsequent
 - Read that sees the written value or
 - Write that overwrites the original write
 - **Domain.spawn** and the first event of the newly spawned domain
 - Last event of domain and **Domain.join**
 - **Unlock** of a mutex and subsequent **Lock**

```
let t1 m =
  Mutex.lock m;
(* .... *)
(* .... *)
Mutex.unlock m;
```



```
let t2 m =
  Mutex.lock m;
(* .... *)
(* .... *)
Mutex.unlock m;
```

Data race

- In a given execution, two actions are said to be ***conflicting*** if
 - At least one of them is a ***write***
 - Neither is an ***initialising write***
- An execution is said to have a ***data race***
 - There is a pair of ***conflicting*** actions
 - There is no ***happens-before*** relationship between them
- Conflicting memory accesses without happens-before are said to be ***unsynchronised***
- A program is said to have a data race if there exists at least one execution which has a data race
 - A program is said to be ***correctly synchronised*** if it doesn't have a data race

Why does this matter?

- **Programs with data races**
 - In C/C++, programs with data races have ***undefined behaviour***
 - Can do anything at all – crash, leak secrets, launch missiles, explode, ...
 - In OCaml, programs with data races have ***well-defined behaviour***
 - Type-safety preserved
 - Semantics are more complicated; you can see stale writes
- **Programs without data races**
 - In C/C++, Java, OCaml, etc, programs without data races are sequentially consistent.
 - Data-race-free sequential consistency (DRF-SC) model
 - ***No need to reason about reorderings!***

Does this have a data race?

- Consider that the functions d* run in parallel

```
let r = ref 0
let d1 () = r := 1
let d2 () = !r
```

Data race

```
let a = [| 0; 1 |]
let d1 () = a.(0) <- 42
let d2 () = a.(1) <- 42
```

No data race

```
type t = {
    mutable a : int;
    mutable b : int
}
let r = {a = 0; b = 1}
let d1 () = r.a <- 42
let d2 () = r.b <- 42
```

No data race

```
let r = Atomic.make 0
let d1 () = Atomic.set r 1
let d2 () = Atomic.get r
```

No data race

Does this have a data race?

- Consider that the functions d* run in parallel

```
let msg = ref 0
let flag = Atomic.make false
let d1 () =
  msg := 42; (* a *)
  Atomic.set flag true (* b *)
let d2 () =
  if Atomic.get flag (* c *) then
    !msg (* d *)
  else 0
```

No data race

```
let msg = ref 0
let flag = Atomic.make false
let d1 () =
  msg := 42; (* a *)
  Atomic.set flag true (* b *)
let d2 () =
  ignore (Atomic.get flag); (* c *)
  !msg (* d *)
```

Data race

Does this have a data race?

- Consider that the functions d* run in parallel

```
(* A counter is just a reference *)
let create_counter initial_value =
  ref initial_value

let get_and_increment counter =
  let v = !counter in
  counter := v + 1;
  v

let c = create_counter 0
let d1 () = get_and_increment c
let d2 () = get_and_increment c
```

Data race

```
type t = { counter : int ref; mutex : Mutex.t }

let create_counter () =
  { counter = ref 0; mutex = Mutex.create () }

let get_and_increment t =
  Mutex.lock t.mutex;
  let old_value = !(t.counter) in
  t.counter := old_value + 1;
  Mutex.unlock t.mutex;
  old_value

let c = create_counter 0
let d1 () = get_and_increment c
let d2 () = get_and_increment c
```

No data race

Does this have a data race?

- Consider that the functions d* run in parallel

```
module Peterson : LOCK = struct
  (* Two boolean flags (from LockOne) and one
     victim variable (from LockTwo) *)
  let flag = [| false; false |]
  let victim = ref 0

  let lock () =
    let i = (Domain.self () :> int) - 1 in
    let j = 1 - i in
    flag.(i) <- true; (* Announce I'm interested *)
    victim := i; (* Defer to the other *)
    (* Wait while the other thread wants to enter
       AND we're the victim *)
    while flag.(j) && !victim = i do
      ()
    done

  let unlock () =
    let i = (Domain.self () :> int) - 1 in
    flag.(i) <- false
end
```

```
let d1 () = Peterson.lock ()
let d2 () = Peterson.lock ()
```

Data race

- Where are the races?

Does this have a data race?

- Consider that the functions d* run in parallel

```
module Peterson : LOCK = struct
  (* Two boolean flags (from LockOne) and one
   * victim variable (from LockTwo) *)
  let flag = [| false; false |]
  let victim = ref 0
```

```
let lock () =
  let i = (Domain.self () :> int) - 1 in
  let j = 1 - i in
  flag.(i) <- true; (* Announce I'm interested *)
  victim := i; (* Defer to the other *)
  (* Wait while the other thread wants to enter
   * AND we're the victim *)
  while flag.(j) && !victim = i do
    ()
  done

let unlock () =
  let i = (Domain.self () :> int) - 1 in
  flag.(i) <- false
end
```

```
let d1 () = Peterson.lock ()
let d2 () = Peterson.lock ()
```

Data race

- Where are the races?
 - There are multiple races!
 - Shockingly, this Peterson's lock does not provide mutual exclusion! (**DEMO**)
 - Data races lead to unintuitive behaviour
 - Can we detect them automatically?

ThreadSanitizer

- Runtime data race detector (dynamic analysis, not static!)
 - Initially developed for C++ by Google, now supported in
 - C, C++ with GCC and clang, Go, Swift and OCaml
- Battle-tested, already found:
 - 1200+ races in Google's codebase
 - ~100 in the Go stdlib
 - 100+ in Chromium
 - LLVM, GCC, OpenSSL, WebRTC, Firefox

ThreadSanitizer

- Requires compiling your program specially
 - Works by instrumenting your program and detecting unsynchronised accesses
- Limitations
 - Only detects data races that appear during execution
 - “Testing only shows the presence of bugs, never the absence of them” – Edsger W. Dijkstra
 - Bookkeeping overheads
 - A finite number of memory accesses is remembered for every memory location
 - Misses out data races separated by many synchronised accesses in between
 - 2x to 7x slower execution time of programs

ThreadSanitizer Demo

```
opam switch create 5.4.0+tsan ocaml-option-tsan  
opam install dune
```

- Need to install a separate switch with TSAN enabled
- Let's detect data races on Peterson's Lock
- macOS detects the race, but no location info :-(
 - Location info reported on Linux

```
module Peterson : LOCK = struct  
  (* Two boolean flags (from LockOne) and one  
   victim variable (from LockTwo) *)  
  let flag = [| false; false |]  
  let victim = ref 0  
  
  let lock () =  
    let i = (Domain.self () :> int) - 1 in  
    let j = 1 - i in  
    flag.(i) <- true; (* Announce I'm interested *)  
    victim := i; (* Defer to the other *)  
    (* Wait while the other thread wants to enter  
     AND we're the victim *)  
    while flag.(j) && !victim = i do  
      ()  
    done  
  
  let unlock () =  
    let i = (Domain.self () :> int) - 1 in  
    flag.(i) <- false  
end
```

DRF Peterson's Lock

```
module Peterson = struct
  let flag = [| Atomic.make false; Atomic.make false |]
  let victim = Atomic.make 0

  let lock () =
    let i = (Domain.self () :> int) - 1 in
    let j = 1 - i in
    Atomic.set flag.(i) true;
    Atomic.set victim i;
    while Atomic.get flag.(j) && Atomic.get victim = i do
      ()
    done

  let unlock () =
    let i = (Domain.self () :> int) - 1 in
    Atomic.set flag.(i) false
end
```

No data race

```
module Peterson : LOCK = struct
  let flag = [| false; false |]
  let victim = ref 0

  let lock () =
    let i = (Domain.self () :> int) - 1 in
    let j = 1 - i in
    flag.(i) <- true;
    victim := i;
    while flag.(j) && !victim = i do
      ()
    done

  let unlock () =
    let i = (Domain.self () :> int) - 1 in
    flag.(i) <- false
end
```

Data race

Programs from previous lectures

- Most of the concurrent programs from the previous lecture have data races
 - LockOne, LockTwo, Peterson's, Bakery Lock, Wait-free SPSC queue, ...
- All of them assume sequential consistency
- Real-world implementations need to use **Atomic** module
 - You can fix them like we did the Peterson's lock

Summary

- ***Relaxed memory models*** are those which are weaker than SC
 - Behaviours must be explained with ***reorderings*** and not just ***interleaving***
- ***Hardware Models*** are usually weaker than SC
 - ***Hardware optimisations*** such as store buffers, out-of-order executions, speculation, cache-coherence protocols, ...
- ***PL Models*** are usually weaker than SC
 - ***Compiler optimisations***
- Data races are unsynchronised, conflicting memory accesses
 - PL models offer SC for data-race-free programs (DRF SC)
- Data races complicate program reasoning
 - OCaml ThreadSanitizer (TSAN) can detect data races dynamically



This work is licensed under a [Creative Commons Attribution-ShareAlike 2.5 License](#).

- **You are free:**
 - **to Share** – to copy, distribute and transmit the work
 - **to Remix** – to adapt the work
- **Under the following conditions:**
 - **Attribution.** You must attribute the work to “The Art of Multiprocessor Programming” (but not in any way that suggests that the authors endorse you or your use of the work).
 - **Share Alike.** If you alter, transform, or build upon this work, you may distribute the resulting work only under the same, similar or a compatible license.
- For any reuse or distribution, you must make clear to others the license terms of this work. The best way to do this is with a link to
 - <http://creativecommons.org/licenses/by-sa/3.0/>.
- Any of the above conditions can be waived if you get permission from the copyright holder.
- Nothing in this license impairs or restricts the author's moral rights.