

# 04 Spin Locks and Contention

CS 6868: Concurrent Programming

KC Sivaramakrishnan

Spring 2026, IIT Madras

# Focus so far: Correctness

- Models
  - Accurate (*we never lied to you*)
  - But idealised (*so we forgot to mention a few things*)
- Protocols
  - Elegant
  - Important
  - But naïve

# New Focus: Performance

- Models
  - More complicated (*not the same as complex!*)
  - Still focus on principles (*not soon obsolete*)
- Protocols
  - Elegant (*in their fashion*)
  - Important (*why else would we pay attention*)
  - And realistic (*your mileage may vary*)

# Kinds of Architectures

- **SISD (Uniprocessor)**
  - Single instruction stream
  - Single data stream
- **SIMD (Vector)**
  - Single instruction
  - Multiple data
- **MIMD (Multiprocessors)**
  - Multiple instruction
  - Multiple data.


# Kinds of Architectures

- **SISD (Uniprocessor)**
  - Single instruction stream
  - Single data stream
- **SIMD (Vector)**
  - Single instruction
  - Multiple data
- **MIMD (Multiprocessors)**
  - Multiple instruction
  - Multiple data.




*Our Focus*

# MIMD Architectures



**Shared Bus**



**Distributed**

- Memory Contention
- Communication Contention
- Communication Latency

# Today: Revisit Mutual Exclusion

- Performance, not just correctness
- Proper use of multiprocessor architectures
- A collection of locking algorithms...

# What should you do if you can't get a lock?


- Keep trying
  - “spin” or “busy-wait”
  - Good if delays are short
- Give up the processor
  - Good if delays are long
  - Always good on uniprocessor

# What should you do if you can't get a lock?


- Keep trying
  - “spin” or “busy-wait”
  - Good if delays are short
- Give up the processor
  - Good if delays are long
  - Always good on uniprocessor

*Our Focus*


# Basic Spin Lock



# Basic Spin Lock




# Basic Spin Lock




# Basic Spin Lock


...lock suffers from contention




# Basic Spin Lock



# Basic Spin Lock



# Basic Spin Lock



# Review: Test-and-set

- Boolean value
- Test-and-set (TAS)
  - Swap **true** with current value
  - Return value tells if prior value was **true** or **false**
- Can reset just by writing **false**
- TAS aka “exchange”

# Review: Test-and-set

- Boolean value
- Test-and-set (TAS)
  - Swap **true** with current value
  - Return value tells if prior value was **true** or **false**
- Can reset just by writing **false**
- TAS aka “exchange”

`Atomic.exchange : 'a Atomic.t -> 'a -> 'a`

*is semantically equivalent to*

```
let exchange (r : 'a Atomic.t) (v : 'a) : 'a =
  let old = Atomic.get r in
  Atomic.set r v;
  old
```

*but atomic*

# Test-and-Set Locks

- Locking
  - Lock is free: value is ***false***
  - Lock is taken: value is ***true***
- Acquire lock by calling TAS
  - If result is ***false***, you win
  - If result is ***true***, you lose
- Release lock by writing ***false***

# Test-and-Set Lock (TASLock)

```
module TASLock : Lock.LOCK = struct
  type t = { state : bool Atomic.t }

  let create () = { state = Atomic.make false }

  let lock t =
    (* Keep spinning until we successfully acquire the lock *)
    (* Atomic.exchange sets the value and returns the old value *)
    while Atomic.exchange t.state true do
      (* Spin – this is where the thread wastes CPU cycles *)
      ()
    done

    let unlock t = Atomic.set t.state false
  end
```


# Space Complexity

- TAS spin-lock has a small “footprint”
- N thread spin-lock uses **O(1)** space
- As opposed to **O(n)** Peterson/Bakery
- How did we overcome the  **$\Omega(n)$**  lower bound?
- We used an RMW operation...


# Space Complexity

- Experiment
  - $n$  threads
  - Increment shared counter ***1 million*** times
- How long ***should*** it take?
- How long ***does*** it take?


# Time taken




# Time taken




# Time taken



# Mystery #1



# Mystery #1



# Test-and-Test-and-Set Locks

- Lurking stage
  - Wait until lock “looks” free
  - Spin while read returns **true** (lock taken)
- Pouncing state
  - As soon as lock “looks” available
  - Read returns **false** (lock free)
  - Call TAS to acquire lock
  - If TAS loses, back to lurking

# Test-and-Test-and-Set Lock

```
module TTASLock : Lock.LOCK = struct
  type t = { state : bool Atomic.t }

  let create () = { state = Atomic.make false }

  let unlock t = Atomic.set t.state false
```

```
let lock t =
  (* Outer loop: keep trying until we get the lock *)
  while
    (* Inner loop: spin-read until lock appears free *)
    while Atomic.get t.state do () done;
    (* Lock looks free, try to acquire with exchange *)
    Atomic.exchange t.state true
  do
    (* If we're here, exchange returned true *)
    (* Somebody has the lock; spin-read again *)
    ()
  done
  (* If we exit the while loop, exchange returned false
   => we have the lock! *)
end
```

# Test-and-Test-and-Set Lock

```
module TTASLock : Lock.LOCK = struct
  type t = { state : bool Atomic.t }


  let create () = { state = Atomic.make false }

  let unlock t = Atomic.set t.state false
```

```
let lock t =
  (* Outer loop: keep trying until we get the lock *)
  while
    (* Inner loop: spin-read until lock appears free *)
    while Atomic.get t.state do () done;
    (* Lock looks free, try to acquire with exchange *)
    Atomic.exchange t.state true
  do
    (* If we're here, exchange returned true *)
    (* Somebody has the lock; spin-read again *)
    ()
  done
  (* If we exit the while loop, exchange returned false
     => we have the lock! *)
end
```

## Demo

# Mystery #2




# Mystery

- Both
  - TAS and TTAS
  - Do the same thing (in our model)
- Except that
  - TTAS performs much better than TAS
  - Neither approaches ideal


# Opinion

- Our memory abstraction is *broken*
- TAS & TTAS methods
  - Are provably the same (in our model)
  - Except they aren't (in field tests)
- Need a more detailed model ...

# Bus-based Architectures




# Bus-based Architectures



# Bus-based Architectures

## Shared Bus


- Broadcast medium
- One broadcaster at a time
- Processors and memory all “snoop”



# Bus-based Architecture

## Per-Processor Caches

- Small
- Fast: 1 or 2 cycles
- Address & state information




# Granularity

- Caches operate at a larger granularity than a word
- **Cache line:** fixed-size block containing the address (today 64 or 128 bytes)


# Locality

- If you use an address now, you will probably use it again soon
  - Fetch from cache, not memory
- If you use an address now, you will probably use a nearby address soon
  - In the same cache line

# L1 and L2 caches




# L1 and L2 caches



# L1 and L2 caches

**Larger and slower  
10s of cycles  
~128 byte line**



# Jargon Watch

- ***Cache hit***
  - “I found what I wanted in my cache”
  - Good Thing™
- ***Cache miss***
  - “I had to shlep all the way to memory for that data”
  - Bad Thing™

# Cave Canem


- This model is *still* a simplification
  - But not in any essential way
  - Illustrates basic principles
- Will discuss complexities later

# When a cache becomes full...

- Need to make room for new entry
- By ***evicting*** an existing entry
- Need a replacement policy
  - Usually some kind of ***least recently used*** heuristic


# Fully Associative Cache

- Any line can be anywhere in the cache
  - Advantage: can replace any line
  - Disadvantage: hard to find lines




# Direct Mapped Cache

- Every address has exactly 1 slot
  - Advantage: easy to find a line
  - Disadvantage: must replace fixed line




# K-way set associative cache

- Each slot holds  $k$  lines
  - Advantage: pretty easy to find a line
  - Advantage: some choice in replacing line



# Multicore Set Associativity

- $k$  is 8 or even 16 and growing...
  - Why? Because cores share sets
  - Threads cut effective size if accessing different data




# Cache Coherence

- A and B both cache address x
- A writes to x
  - Updates cache
- How does B find out?
- Many ***cache coherence*** protocols in literature


# MESI

- **Modified**
  - Have modified cached data, must write back to memory
- **Exclusive**
  - Not modified, I have only copy
- **Shared**
  - Not modified, may be cached elsewhere
- **Invalid**
  - Cache contents not meaningful


# Processor Issues Load Request




# Processor Issues Load Request




# Memory Responds




# Memory Responds




# Processor Issues Load Request




# Processor Issues Load Request




# Other Processor Responds




# Other Processor Responds




# Processor Issues Load Request




# Processor Issues Load Request



# Write-through Cache




# Write-through Cache




# Write-through Cache

- Immediately broadcast changes
- Good
  - Memory, caches always agree
  - More read hits, maybe
- Bad
  - Bus traffic on all writes
  - Most writes to unshared data
  - For example, loop indexes ...


# Write-through Cache

- Immediately broadcast changes
  - Good
    - Memory, caches always agree
    - More read hits, maybe
  - Bad
    - Bus traffic on all writes
    - Most writes to unshared data
    - For example, loop indexes ...
- “show stoppers”**
- 


# Invalidate




# Invalidate



# Invalidate




# Invalidate



# Invalidate

Memory provides data only if not present in any cache, so no need to change it now (expensive)



# Mutual Exclusion

- What do we want to optimize?
  - Bus bandwidth used by spinning threads
  - Release/Acquire latency
  - Acquire latency for idle lock


# Test-and-set Lock (TASLock)

- TAS *invalidates* (I) cache lines
- Spinners
  - Miss in cache
  - Go to bus
- Cache line **bounces** between all the spinners in modified (M) state
- Thread wants to release lock
  - delayed behind spinners


# Test-and-test-and-set Lock (TTASLock)

- Wait until lock “looks” free
  - Spin on local cache
  - Cache line is in Shared (S) state
  - No bus use while lock busy
- Problem: when lock is released
  - Invalidation storm ...

# Local spinning while lock is busy




# On Release




# On Release

Everyone misses,  
rereads



# On Release


Everyone tries TAS




# Problems

- Everyone misses
  - Reads satisfied sequentially
- Everyone does TAS
  - Invalidates others' caches
- Eventually ***quiesces*** after lock acquired
- ***Quiescence duration*** is the time between lock release and lock acquire
  - How long does this take?


# Quiescence Time



# Mystery Explained




# Mystery Explained




# Solution: Introduce Delay

- If the lock looks free
  - But I fail to get it
- There must be contention
  - Better to back off than to collide again



# Dynamic Solution – Exponential backoff



If I fail to get lock

- Wait random duration before retry
- Each subsequent failure doubles expected wait

# Exponential Backoff

```
module Backoff = struct
  type t = { min_delay : int; max_delay : int; mutable limit : int }

  let create min_delay max_delay = { min_delay; max_delay; limit = min_delay }

  let backoff t =
    (* Backoff for a random duration between 0 and 'limit' using cpu_relax *)
    (* Randomization prevents synchronized collisions between threads *)
    let delay = Random.int (t.limit + 1) in
    for _ = 1 to delay do
      Domain.cpu_relax ()
      (* Instructs the core to slow down fast(er than sleep()) *)
    done;
    (* Exponentially increase the limit, capped at max_delay *)
    t.limit <- min t.max_delay (t.limit * 2)
end
```

# BackoffLock


```
module MakeBackoffLock (P : BACKOFF_PARAMS) : Lock.LOCK = struct
  type t = { state : bool Atomic.t }

  let create () = { state = Atomic.make false }


  let lock t =
    let backoff = Backoff.create P.min_delay P.max_delay in
    (* Outer loop: keep trying until we get the lock *)
    while
      (* Inner loop: spin-read until lock appears free *)
      while Atomic.get t.state do
        ()
      done;
      (* Lock looks free, try to acquire *)
      Atomic.exchange t.state true
    do
      (* Failed to acquire – back off before trying again *)
      Backoff.backoff backoff
    done

    let unlock t = Atomic.set t.state false
  end
```

# Spin-waiting Overhead



# Results on 8-core M2 Apple Silicon




# Backoff: Other issues

- Good
  - Easy to implement
  - Beats TTAS lock
- Bad
  - Must choose parameters carefully
  - Not portable across platforms


# Idea

- Avoid useless invalidations
  - By keeping a **queue** of threads
- Each thread
  - Notifies next in line
  - Without bothering the others


# Anderson Queue Lock




# Anderson Queue Lock




# Anderson Queue Lock




# Anderson Queue Lock




# Anderson Queue Lock




# Anderson Queue Lock




# Anderson Queue Lock




# Anderson Queue Lock




# Anderson Queue Lock




# Anderson Queue Lock




# Local Spinning




# Local Spinning




# False Sharing




# False Sharing



# False Sharing



# OCaml Solution: Heap-separated allocation + padding



# OCaml ALock

```
type t = {  
  flags : bool Atomic.t array;  
  tail : int Atomic.t;  
  capacity : int;  
  my_slot : int Domain.DLS.key;  
}
```

# OCaml ALock

```
type t = {  
  flags : bool Atomic.t array;  
  tail : int Atomic.t;  
  capacity : int;  
  my_slot : int Domain.DLS.key;  
}
```

```
let create_with_capacity capacity =  
  if capacity <= 0 then  
    invalid_arg "ALock capacity must be positive";  
  (* Create array of atomic booleans using  
   * make_contended to prevent false sharing *)  
  let flags =  
    Array.init capacity (fun i ->  
      (* Only slot 0 starts as true (available) *)  
      Atomic.make_contended (i = 0))  
  in  
  {  
    flags;  
    tail = Atomic.make 0;  
    capacity;  
    (* DLS for remembering slot taken by domain *)  
    my_slot = Domain.DLS.new_key (fun () -> -1);  
  }
```

# OCaml ALock

```
let lock t =
  (* Get my slot using atomic fetch-and-increment *)
  let slot = (Atomic.fetch_and_add t.tail 1) mod t.capacity in

  (* Store slot in domain-local storage for unlock *)
  Domain.DLS.set t.my_slot slot;

  (* Cache the flag reference to avoid repeated array indexing *)
  let my_flag = t.flags.(slot) in

  (* Spin on MY flag until it becomes true *)
  (* This is the key: each thread spins on a DIFFERENT location *)
  while not (Atomic.get my_flag) do
    Domain.cpu_relax ()
  done
```

# OCaml ALock

```
let unlock t =
  (* Get my slot from domain-local storage *)
  let slot = Domain.DLS.get t.my_slot in


  if slot = -1 then
    failwith "unlock called without corresponding lock";

  (* Cache flag references *)
  let my_flag = t.flags.(slot) in
  let next_slot = (slot + 1) mod t.capacity in
  let next_flag = t.flags.(next_slot) in

  (* Clear my flag *)
  Atomic.set my_flag false;


  (* Signal the next thread *)
  Atomic.set next_flag true
```

# Performance



- Shorter handover than backoff
- Curve is practically flat
- Scalable performance

# Results on 28-core Intel Xeon Gold 5120




# Anderson Queue Lock

- Good
  - First truly scalable lock
  - Simple, easy to implement
  - Back to FCFS order (like Bakery)
- Bad
  - Space hog...
  - One bit per thread → one cache line per thread
    - What if unknown number of threads?
    - What if small number of actual contenders?

# More Spinlocks in the Book

- CHL Lock
- MCS Lock
- Fast-path composite locks
- Hierarchical backoff locks
- ...
- No silver bullet!



Chapter 7



## This work is licensed under a [Creative Commons Attribution-ShareAlike 2.5 License](#).

- **You are free:**
  - **to Share** – to copy, distribute and transmit the work
  - **to Remix** – to adapt the work
- **Under the following conditions:**
  - **Attribution.** You must attribute the work to “The Art of Multiprocessor Programming” (but not in any way that suggests that the authors endorse you or your use of the work).
  - **Share Alike.** If you alter, transform, or build upon this work, you may distribute the resulting work only under the same, similar or a compatible license.
- For any reuse or distribution, you must make clear to others the license terms of this work. The best way to do this is with a link to
  - <http://creativecommons.org/licenses/by-sa/3.0/>.
- Any of the above conditions can be waived if you get permission from the copyright holder.
- Nothing in this license impairs or restricts the author's moral rights.